

The Basic Practice

of

STATISTICS

Sixth Edition

MOORE / NOTZ / FLIGNER

Send



The Basic Practice of Statistics

SIXTH EDITION

DAVID S. MOORE

Purdue University

WILLIAM I. NOTZ

The Ohio State University

MICHAEL A. FLIGNER

The Ohio State University



W. H. Freeman and Company
New York

Publisher: **Ruth Baruth**

Acquisitions Editor: **Karen Carson**

Executive Marketing Manager: **Jennifer Somerville**

Developmental Editors: **Andrew Sylvester and Leslie Lahr**

Senior Media Acquisitions Editor: **Roland Cheyney**

Senior Media Editor: **Laura Capuano**

Associate Editor: **Katrina Wilhelm**

Assistant Media Editor: **Catriona Kaplan**

Editorial Assistant: **Tyler Holzer**

Photo Editor: **Cecilia Varas**

Photo Researcher: **Elyse Rieder**

Cover and Text Designer: **Blake Logan**

Senior Project Editor: **Mary Louise Byrd**

Illustrations: **Macmillan Solutions**

Production Coordinator: **Susan Wein**

Composition: **Aptara®, Inc.**

Printing and Binding: **Quad Graphics**

Library of Congress Control Number:

2011934674

Student Edition (Hardcover w/cd) Student Edition (Paperback w/cd) Student Edition (Looseleaf w/cd)

ISBN-13: 978-1-4641-0254-7 ISBN-13: 978-1-4641-0434-3 ISBN-13: 978-1-4641-0433-6

ISBN-10: 1-4641-0254-6 ISBN-10: 1-4641-0434-4 ISBN-10: 1-4641-0433-6

© 2013, 2010, 2007, 2004 by W. H. Freeman and Company

All rights reserved

Printed in the United States of America

First printing

W. H. Freeman and Company

41 Madison Avenue

New York, NY 10010

Houndsborough RG21 6XS, England

www.whfreeman.com



Brief Contents

Part I Exploring Data*Exploring Data: Variables and Distributions***CHAPTER 1** Picturing Distributions with Graphs 3**CHAPTER 2** Describing Distributions with Numbers 39**CHAPTER 3** The Normal Distributions 69*Exploring Data: Relationships***CHAPTER 4** Scatterplots and Correlation 97**CHAPTER 5** Regression 125**CHAPTER 6** Two-Way Tables* 159**CHAPTER 7** Exploring Data: Part I Review 175**Part II From Exploration to Inference**

197

*Producing Data***CHAPTER 8** Producing Data: Sampling 199**CHAPTER 9** Producing Data: Experiments 223*Commentary: Data Ethics***Probability and Sampling Distributions* 246**CHAPTER 10** Introducing Probability 259**CHAPTER 11** Sampling Distributions 285**CHAPTER 12** General Rules of Probability* 307**CHAPTER 13** Binomial Distributions* 331*Foundations of Inference***CHAPTER 14** Confidence Intervals: The Basics 351**CHAPTER 15** Tests of Significance: The Basics 369**CHAPTER 16** Inference in Practice 391**CHAPTER 17** From Exploration to Inference: Part II Review 417**Part III Inference about Variables**

435

*Quantitative Response Variable***CHAPTER 18** Inference about a Population Mean 437**CHAPTER 19** Two-Sample Problems 465*Categorical Response Variable***CHAPTER 20** Inference about a Population Proportion 493**CHAPTER 21** Comparing Two Proportions 515**CHAPTER 22** Inference about Variables: Part III Review 533**Part IV Inference about Relationships**

551

CHAPTER 23 Two Categorical Variables: The Chi-Square Test 553**CHAPTER 24** Inference for Regression 587**CHAPTER 25** One-Way Analysis of Variance: Comparing Several Means 623**Part V Optional Companion Chapters***(available on the BPS CD and online)***CHAPTER 26** Nonparametrics Tests 26-3**CHAPTER 27** Statistical Process Control 27-3**CHAPTER 28** Multiple Regression* 28-3**CHAPTER 29** More about Analysis of Variance 29-3

*Starred material is not required for later parts of the text.



Detailed Table of Contents

To the Instructor	viii
Media and Supplements	xix
About the Authors	xxiv
To the Student	xxvi

Part I Exploring Data

CHAPTER 1

Picturing Distributions with Graphs 3

Individuals and variables	3
Categorical variables: pie charts and bar graphs	6
Quantitative variables: histograms	11
Interpreting histograms	15
Quantitative variables: stemplots	20
Time plots	23

CHAPTER 2

Describing Distributions with Numbers 39

Measuring center: the mean	40
Measuring center: the median	41
Comparing the mean and the median	42
Measuring spread: the quartiles	43
The five-number summary and boxplots	45
Spotting suspected outliers*	48
Measuring spread: the standard deviation	49
Choosing measures of center and spread	51
Using technology	53
Organizing a statistical problem	55

CHAPTER 3

The Normal Distributions 69

Density curves	69
Describing density curves	73
Normal distributions	75
The 68–95–99.7 rule	77
The standard Normal distribution	80
Finding Normal proportions	81
Using the standard Normal table	83
Finding a value given a proportion	86

1

CHAPTER 4

Scatterplots and Correlation	97
Explanatory and response variables	97
Displaying relationships: scatterplots	99
Interpreting scatterplots	101
Adding categorical variables to scatterplots	104
Measuring linear association: correlation	106
Facts about correlation	108

CHAPTER 5

Regression	125
Regression lines	125
The least-squares regression line	128
Using technology	130
Facts about least-squares regression	132
Residuals	135
Influential observations	139
Cautions about correlation and regression	142
Association does not imply causation	144

CHAPTER 6

Two-Way Tables*	159
Marginal distributions	160
Conditional distributions	162
Simpson's paradox	166

CHAPTER 7

Exploring Data: Part I Review	175
Part I summary	177
Test yourself	180
Supplementary exercises	191

Part II From Exploration to Inference

197

CHAPTER 8

Producing Data: Sampling	199
Population versus sample	199
How to sample badly	202
Simple random samples	203

*Starred material is not required for later parts of the text.

Inference about the population	208
Other sampling designs	209
Cautions about sample surveys	210
The impact of technology	213

CHAPTER 9**Producing Data: Experiments 223**

Observation versus experiment	223
Subjects, factors, treatments	225
How to experiment badly	228
Randomized comparative experiments	229
The logic of randomized comparative experiments	232
Cautions about experimentation	234
Matched pairs and other block designs	236

Commentary: Data Ethics* 246

Institutional review boards	248
Informed consent	248
Confidentiality	250
Clinical trials	252
Behavioral and social science experiments	253

CHAPTER 10**Introducing Probability 259**

The idea of probability	260
The search for randomness*	262
Probability models	264
Probability rules	266
Finite and discrete probability models	268
Continuous probability models	271
Random variables	275
Personal probability*	276

CHAPTER 11**Sampling Distributions 285**

Parameters and statistics	285
Statistical estimation and the law of large numbers	287
Sampling distributions	290
The sampling distribution of \bar{x}	293
The central limit theorem	295

CHAPTER 12**General Rules of Probability* 307**

Independence and the multiplication rule	308
The general addition rule	312
Conditional probability	314
The general multiplication rule	316
Independence again	318
Tree diagrams	318

CHAPTER 13**Binomial Distributions* 331**

The binomial setting and binomial distributions	331
Binomial distributions in statistical sampling	333
Binomial probabilities	334
Using technology	336
Binomial mean and standard deviation	338
The Normal approximation to binomial distributions	340

CHAPTER 14**Confidence Intervals: The Basics 351**

The reasoning of statistical estimation	352
Margin of error and confidence level	354
Confidence intervals for a population mean	357
How confidence intervals behave	361

CHAPTER 15**Tests of Significance: The Basics 369**

The reasoning of tests of significance	370
Stating hypotheses	372
P-value and statistical significance	374
Tests for a population mean	378
Significance from a table*	382

CHAPTER 16**Inference in Practice 391**

Conditions for inference in practice	392
Cautions about confidence intervals	395
Cautions about significance tests	397
Planning studies: sample size for confidence intervals	401
Planning studies: the power of a statistical test*	402

CHAPTER 17**From Exploration to Inference: Part II Review 417**

Part II summary	419
Test yourself	423
Supplementary exercises	431

Part III Inference about Variables

435

CHAPTER 18**Inference about a Population Mean 437**

Conditions for inference about a mean	437
The <i>t</i> distributions	438
The one-sample <i>t</i> confidence interval	440

The one-sample t test	443
Using technology	446
Matched pairs t procedures	449
Robustness of t procedures	452

CHAPTER 19**Two-Sample Problems 465**

Two-sample problems	465
Comparing two population means	466
Two-sample t procedures	469
Using technology	474
Robustness again	477
Details of the t approximation*	480
Avoid the pooled two-sample t procedures*	481
Avoid inference about standard deviations*	482

CHAPTER 20**Inference about a Population Proportion 493**

The sample proportion \hat{p}	494
Large-sample confidence intervals for a proportion	496
Accurate confidence intervals for a proportion	499
Choosing the sample size	502
Significance tests for a proportion	504

CHAPTER 21**Comparing Two Proportions 515**

Two-sample problems: proportions	515
The sampling distribution of a difference between proportions	516
Large-sample confidence intervals for comparing proportions	517
Using technology	518
Accurate confidence intervals for comparing proportions	520
Significance tests for comparing proportions	522

CHAPTER 22**Inference about Variables: Part III Review 533**

Part III summary	536
Test yourself	538
Supplementary exercises	545

Part IV Inference about Relationships

551

CHAPTER 23**Two Categorical Variables: The Chi-Square Test 553**

Two-way tables	553
The problem of multiple comparisons	556
Expected counts in two-way tables	558

The chi-square test statistic	560
Cell counts required for the chi-square test	561
Using technology	562
Uses of the chi-square test	567
The chi-square distributions	570
The chi-square test for goodness of fit*	572

CHAPTER 24**Inference for Regression 587**

Conditions for regression inference	589
Estimating the parameters	590
Using technology	593
Testing the hypothesis of no linear relationship	597
Testing lack of correlation	598
Confidence intervals for the regression slope	600
Inference about prediction	602
Checking the conditions for inference	607

CHAPTER 25**One-Way Analysis of Variance: Comparing Several Means 623**

Comparing several means	625
The analysis of variance F test	625
Using technology	628
The idea of analysis of variance	631
Conditions for ANOVA	633
F distributions and degrees of freedom	637
Some details of ANOVA*	640

Notes and Data Sources 655**Tables 675**

TABLE A Standard Normal probabilities	676
TABLE B Random digits	678
TABLE C t distribution critical values	679
TABLE D Chi-square distribution critical values	680
TABLE E Critical values of the correlation r	681

Answers to Selected Exercises 682**Index 733****Part V Optional Companion Chapters***(available on the BPS CD and online)***CHAPTER 26****Nonparametric Tests 26-3**

Comparing two samples: the Wilcoxon rank sum test	26-4
The Normal approximation for W	26-8

Using technology	26-10
What hypotheses does Wilcoxon test?	26-13
Dealing with ties in rank tests	26-14
Matched pairs: the Wilcoxon signed rank test	26-19
The Normal approximation for W^+	26-22
Dealing with ties in the signed rank test	26-24
Comparing several samples: the Kruskal-Wallis test	26-27
Hypotheses and conditions for the Kruskal-Wallis test	26-29
The Kruskal-Wallis test statistic	26-29

CHAPTER 27

Statistical Process Control 27-3

Processes	27-4
Describing processes	27-4
The idea of statistical process control	27-9
\bar{x} charts for process monitoring	27-10
s charts for process monitoring	27-16
Using control charts	27-23
Setting up control charts	27-25
Comments on statistical control	27-32
Don't confuse control with capability!	27-34
Control charts for sample proportions	27-36
Control limits for p charts	27-37

CHAPTER 28

Multiple Regression* 28-3

Parallel regression lines	28-4
Estimating parameters	28-8
Using technology	28-13
Inference for multiple regression	28-16
Interaction	28-26
The multiple linear regression model	28-32
The woes of regression coefficients	28-39
A case study for multiple regression	28-41
Inference for regression parameters	28-53
Checking the conditions for inference	28-58

CHAPTER 29

More about Analysis of Variance 29-3

Beyond one-way ANOVA	29-3
Follow-up analysis: Tukey pairwise multiple comparisons	29-8
Follow-up analysis: contrasts*	29-12
Two-way ANOVA: conditions, main effects, and interaction	29-16
Inference for two-way ANOVA	29-23
Some details of two-way ANOVA*	29-32



To the Instructor: *About this Book*

Welcome to the sixth edition of *The Basic Practice of Statistics*. This book is the cumulation of 40 years of teaching undergraduates and 20 years of writing texts. Previous editions have been very successful, and we think that this new edition is the best yet. In this preface we describe for instructors the nature and features of the book and the changes in this sixth edition.

BPS is designed to be accessible to college and university students with limited quantitative background—“just algebra” in the sense of being able to read and use simple equations. It is usable with almost any level of technology for calculating and graphing—from a \$15 “two-variable statistics” calculator through a graphing calculator or spreadsheet program through full statistical software. Of course, graphs and calculations are less tedious with good technology, so we recommend making available to students the most effective technology that circumstances permit.

Despite its rather low mathematical level, BPS is a “serious” text in the sense that it wants students to do more than master the mechanics of statistical calculations and graphs. Even quite basic statistics is very useful in many fields of study and in everyday life, but only if the student has learned to move from a real-world setting to choose and carry out statistical methods and then carry conclusions back to the original setting. These translations require some conceptual understanding of such issues as the distinction between data analysis and inference, the critical role of where the data come from, the reasoning of inference, and the conditions under which we can trust the conclusions of inference. BPS tries to teach both the mechanics and the concepts needed for practical statistical work, at a level appropriate for beginners.

BPS is designed to reflect the actual practice of statistics, where data analysis and design of data production join with probability-based inference to form a coherent science of data. There are good pedagogical reasons for beginning with data analysis (Chapters 1 to 7), then moving to data production (Chapters 8 and 9), and then to probability (Chapters 10 to 13) and inference (Chapters 14 to 29). In studying data analysis, students learn useful skills immediately and get over some of their fear of statistics. Data analysis is a necessary preliminary to inference in practice, because inference requires clean data. Designed data production is the surest foundation for inference, and the deliberate use of chance in random sampling and randomized comparative experiments motivates the study of probability in a course that emphasizes data-oriented statistics. BPS gives a full presentation of basic probability and inference (20 of the 29 chapters) but places it in the context of statistics as a whole.

GUIDING PRINCIPLES AND THE GAISE GUIDELINES

David Moore has based BPS on three principles: balanced content, experience with data, and the importance of ideas. These principles are widely accepted by statisticians concerned about teaching and are directly connected to and reflected by the

themes of the College Report of the Guidelines for Assessment and Instruction in Statistics Education (GAISE) Project.

The GAISE Guidelines includes six recommendations for the introductory statistics course. The content, coverage, and features of BPS are closely aligned to these recommendations:

1. Emphasize statistical literacy and develop statistical thinking.

The intent of BPS is to be modern and accessible. The exposition is straightforward and concentrates on major ideas and skills. One principle of writing for beginners is not to try to tell students everything you know. Another principle is to offer frequent stopping points, marking off digestible bites of material. Statistical literacy is promoted throughout BPS in the many examples and exercises drawn from the popular press and from many fields of study. Statistical thinking is promoted in examples and exercises that give enough background to allow students to consider the meaning of their calculations. Exercises often ask for conclusions that are more than a number (or “reject H_0 ”). Some exercises require judgment in addition to right-or-wrong calculations and conclusions. Statistics, more than mathematics, depends on judgment for effective use. BPS begins to develop students’ judgment about statistical studies.

2. Use real data. The study of statistics is supposed to help students work with data in their varied academic disciplines and in their unpredictable later employment. Students learn to work with data by working with data. BPS is full of data from many fields of study and from everyday life. Data are more than mere numbers—they are numbers with a context that should play a role in making sense of the numbers and in stating conclusions. Examples and exercises in BPS, though intended for beginners, use real data and give enough background to allow students to consider the meaning of their calculations.

3. Stress conceptual understanding rather than mere knowledge of procedures. A first course in statistics introduces many skills, from making a stemplot and calculating a correlation to choosing and carrying out a significance test. In practice (even if not always in the course), calculations and graphs are automated. Moreover, anyone who makes serious use of statistics will need some specific procedures not taught in her college stat course. BPS therefore tries to make clear the larger patterns and big ideas of statistics, not in the abstract, but in the context of learning specific skills and working with specific data. Many of the big ideas are summarized in graphical outlines. Three of the most useful appear inside the front cover. Formulas without guiding principles do students little good once the final exam is past, so it is worth the time to slow down a bit and explain the ideas.

4. Foster active learning in the classroom. Fostering active learning is the business of the teacher, though an emphasis on working with data helps. To this end, we have created interactive applets to our specifications and made them available online and on the text CD. The applets are designed primarily to help in learning statistics rather than in doing statistics. An icon calls



attention to comments and exercises based on the applets. We suggest using selected applets for classroom demonstrations even if you do not ask students to work with them. The *Correlation and Regression*, *Confidence Interval*, and *P-value applets*, for example, convey core ideas more clearly than any amount of chalk and talk.

We also provide Web exercises at the end of each chapter. Our intent is to take advantage of the fact that most undergraduates are “Web savvy.” These exercises require students to search the Web for either data or statistical examples and then evaluate what they find. Teachers can use these as classroom activities or assign them as homework projects.

5. Use technology for developing conceptual understanding and analyzing data.

Automating calculations increases students’ ability to complete problems, reduces their frustration, and helps them concentrate on ideas and problem recognition rather than mechanics. At a minimum, students should have a “two-variable statistics” calculator with functions for correlation and the least-squares regression line as well as for the mean and standard deviation.

Many instructors will take advantage of more elaborate technology, as ASA/MAA and GAISE recommend. And many students who don’t use technology in their college statistics course will find themselves using (for example) Excel on the job. BPS does not assume or require use of software except in Parts IV and V, where the work is otherwise too tedious. It does accommodate software use and tries to convince students that they are gaining knowledge that will enable them to read and use output from almost any source. There are regular “Using Technology” sections throughout the text. Each one displays and comments on output from the same three technologies, representing graphing calculators (the Texas Instruments TI-83 or TI-84), spreadsheets (Microsoft Excel), and statistical software (Minitab and CrunchIt!). The output always concerns one of the main teaching examples so that students can compare text and output.

6. Use assessments to improve and evaluate student learning.

Within chapters, a few “Apply Your Knowledge” exercises follow each new idea or skill for a quick check of basic mastery and also to mark off digestible bites of material. Each of the first three parts of the book ends with a review chapter that includes a point-by-point outline of skills learned, problems students can use to test themselves, and several supplementary exercises. (Instructors can choose to cover any or none of the chapters in Parts IV and V, so each of these chapters includes a skills outline.) The review chapters present supplemental exercises without the “I just studied that” context, thus asking for another level of learning. We think it is helpful to assign some supplemental exercises. Many instructors will find that the review chapters appear at the right points for preexamination review. The “Test Yourself” questions can be used by students to review, self-assess, and prepare for such an examination. In addition, assessment materials in the form of a test bank and quizzes are available online.

WHAT'S NEW?

As always, a new edition of BPS brings many **new examples and exercises**. There are new data sets from a variety of sources, including marketing (methods for waitpersons to improve the size of tips), exercise physiology (improving swimming performance), psychology (do colors affect performance on cognitive tasks), and the environment (water quality in Ohio State Parks). The old favorite Florida manatee regression example returns to Chapters 4, 5, and 24 now that current data are available. These are just a few of a large number of new data settings in this edition.

A new edition is also an opportunity to introduce new features and polish the exposition in ways intended to help students learn. Here are some of the changes.

- Chapter summaries consist of two sections. One, titled “Chapter Specifics,” summarizes the material presented in the chapter. The second section titled “Link It,” relates the chapter content to material in previous and upcoming chapters. Together, Chapter Specifics and Link it help students understand how individual chapters relate to each other and to the overall practice of statistics.
- Each chapter now contains a few exercises that have students investigate data or statistical issues that can be found on the Web. Problems include locating data on the Web and exploring the data to answer questions about it, discussing statistical statements found online (much like we do in exercises asking students to evaluate statistical statements found in newspapers and magazines), or learning about special applications of statistics (for example, the Six Sigma movement in industry). These exercises appear in a section called “Exploring the Web” at the end of each chapter.
- The Part Reviews have been revised to include a section titled “Test Yourself.” This section provides multiple-choice, calculations, and short-answer questions to help students review the basic ideas and skills presented in that Part’s chapters. These questions replace the review exercises in previous editions.
- Data icons are located next to examples and exercises involving data sets. Data sets are now given descriptive names and are available in a variety of formats online.
- In response to reviewer requests, we have returned to the three-chapter coverage of inference used in the fourth edition of BPS. The sixth edition introduces this topic with separate chapters for confidence intervals (Chapter 14), tests of significance (Chapter 15), and inference in practice (Chapter 16). The revised organization allows students to focus on the basics of each procedure before exploring their uses in practice. The approach adopted in the fifth edition was intended to emphasize the relation between confidence intervals and tests of significance for means. Although confidence intervals and hypothesis tests have a natural connection, we realize it is simpler for students to first learn them as separate ideas and then explore the connection. We do not discuss one-sided intervals, so it may be prudent to avoid overemphasizing the connection and raising questions about how one-sided tests connect to confidence intervals.

FEATURES OF THE BASIC PRACTICE OF STATISTICS, Sixth Edition

In this chapter we cover... Each chapter opener gives a brief look at where the chapter is heading, often with references to previous chapters, and includes this bulleted list of the major topics covered.

4-Step Examples

In Chapter 2, students learn how to use the four-step process for working through statistical problems: State, Plan, Solve, and Conclude. By observing this process at work in selected examples throughout the text and practicing it in selected exercises, students develop the ability to solve and write reports on real statistical problems encountered outside of the classroom setting.

Chapter 2

IN THIS CHAPTER WE COVER...

- Measuring center: the mean
- Measuring center: the median
- Comparing the mean and the median
- Measuring spread: the quartiles
- The five-number summary and boxplots
- Spotting suspected outliers*
- Measuring spread: the standard deviation
- Choosing measures of center and spread
- Using technology
- Organizing a statistical problem



Art Wolfe/Getty Images



TROPICALFLOWER

EXAMPLE 2.9 Comparing tropical flowers

STATE: Ethan Temeles of Amherst College, with his colleague W. John Kress, studied the relationship between varieties of the tropical flower *Heliconia* on the island of Dominica and the different species of hummingbirds that fertilize the flowers.⁹ Over time, the researchers believe, the lengths of the flowers and the forms of the hummingbirds' beaks have evolved to match each other. If that is true, flower varieties fertilized by different hummingbird species should have distinct distributions of length.

Table 2.1 gives length measurements (in millimeters) for samples of three varieties of *Heliconia*, each fertilized by a different species of hummingbird. Do the three varieties display distinct distributions of length? How do the mean lengths compare?

PLAN: Use graphs and numerical descriptions to describe and compare these three distributions of flower length.

SOLVE: We might use boxplots to compare the distributions, but stemplots preserve more detail and work well for data sets of these sizes. Figure 2.4 displays stemplots with the stems lined up for easy comparison. The lengths have been rounded to the nearest tenth of a millimeter. The *bihai* and red varieties have somewhat skewed distributions, so we might choose to compare the five-number summaries. But because the researchers plan to use \bar{x} and s for further analysis, we instead calculate these measures:

Variety	Mean length	Standard deviation
<i>bihai</i>	47.60	1.213
red	39.71	1.799
yellow	36.18	0.975

CONCLUDE: The three varieties differ so much in flower length that there is little overlap among them. In particular, the flowers of *bihai* are longer than either red or yellow. The mean lengths are 47.6 mm for *H. bihai*, 39.7 mm for *H. caribaea* red, and 36.2 mm for *H. caribaea* yellow. ■

**APPLY YOUR KNOWLEDGE**

4.10 Coral reefs. Exercise 4.2 discusses a study in which scientists examined data on mean sea surface temperatures (in degrees Celsius) and mean coral growth (in millimeters per year) over a several-year period at locations in the Red Sea. Here are the data:⁷



Sea surface temperature	29.68	29.87	30.16	30.22	30.48	30.65	30.90
Growth	2.63	2.58	2.60	2.48	2.26	2.38	2.26

- (a) Make a scatterplot. Which is the explanatory variable? The plot shows a negative linear pattern.
- (b) Find the correlation r step-by-step. You may wish to round off to two decimal places in each step. First find the mean and standard deviation of each variable. Then find the seven standardized values for each variable. Finally, use the formula for r . Explain how your value for r matches your graph in (a).
- (c) Enter these data into your calculator or software and use the correlation function to find r . Check that you get the same result as in (b), up to roundoff error.

Using Technology Located throughout the text, these special sections display and comment on the output from graphing calculators, spreadsheets, and statistical software in the context of examples from the text.

**The vital few**

Skewed distributions can show us where to concentrate our efforts. Ten percent of the cars on the road account for half of all carbon dioxide emissions. A histogram of CO₂ emissions would show many cars with small or moderate values and a few with very high values. Cleaning up or replacing these cars would reduce pollution at a cost much lower than that of programs aimed at all cars. Statisticians who work at improving quality in industry make a principle of this: distinguish "the vital few" from "the trivial many."

USING TECHNOLOGY

Although a calculator with "two-variable statistics" functions will do the basic calculations we need, more elaborate tools are helpful. Graphing calculators and computer software will do calculations and make graphs as you command, freeing you to concentrate on choosing the right methods and interpreting your results. Figure 2.3 displays output describing the travel times to work of 20 people in New York State (Example 2.3). Can you find \bar{x} , s , and the five-number summary in each output? The big message of this section is: *once you know what to look for, you can read output from any technological tool.*

Statistics in Your World These brief asides, found in each chapter, illustrate major concepts or present cautionary tales through entertaining and relevant stories, allowing students to take a break from the exposition while staying engaged.

LINK IT

The methods of Chapters 1 to 6 can be used to describe data regardless of how the data were obtained. However, if we want to reason from data to give answers to specific questions or to draw conclusions about the larger population, then the method that was used to collect the data is important. Sampling is one way to collect data, but it does not guarantee that we can draw meaningful conclusions. Biased sampling methods, such as convenience sampling and voluntary response samples, produce data that can be misleading, resulting in incorrect conclusions. Simple random sampling avoids bias and produces data that can lead to valid conclusions regarding the population. Even with perfect sampling methods, there is still sample-to-sample variation; we will begin our study of the connection between sampling variation and drawing conclusions in Chapter 11.

Chapter Summary and Link It Each chapter concludes with a summary of the chapter specifics, including major terms and processes, followed by a brief discussion of how the chapter "links" to material in the previous and upcoming chapters.



CHAPTER 21 EXERCISES

21.34 Does preschool help? To study the long-term effects of preschool programs for poor children, the High/Scope Educational Research Foundation has followed two groups of Michigan children since early childhood.²⁷ One group of 62 attended preschool as three- and four-year-olds. A control group of 61 children from the same area and similar backgrounds did not attend preschool. Over a 10-year period as adults, 38 of the preschool sample and 49 of the control sample needed social services (mainly welfare). Does the study provide significant evidence that children who attend preschool have less need for social services as adults? How large is the difference between the proportions of the preschool and no-preschool populations that require social services? Do inference to answer both questions. Be sure to explain exactly what inference you choose to do.

21.35 Hand sanitizers. Hand disinfection is frequently recommended for prevention of transmission of the rhinovirus that causes the common cold. In particular, hand lotion containing 2% citric acid and 2% malic acid in 70% ethanol (HL+) has been found to have both immediate

Check Your Skills and Chapter Exercises Each chapter ends with a series of multiple-choice problems that test students' understanding of basic concepts and students' ability to apply the concepts to real-world statistical situations. The multiple-choice problems are followed by a set of more in-depth exercises that allow students to make judgments and draw conclusions based on real data and real scenarios.



EXPLORING THE WEB

2.55 Crime rates and outliers. The *Statistical Abstract of the United States* is a comprehensive summary of statistics on the social, political, and economic organization of the United States. It can be found at the Web site www.census.gov/compendia/statab/. Go to the section Law Enforcement, Courts and Prisons, and then to the subsection Crimes and Crime Rates. Several tables of data will be available.

- Open the Table on Crime Rates by State and Type for the latest year given. Why do you think they use rates per 100,000 population rather than the number of crimes committed? The District of Columbia is a high outlier in almost every crime category.
- Open the Table on Crime Rates by Type for Selected Large Cities. This table includes the District of Columbia, which is listed as Washington, DC. Without doing any formal calculations, does the District of Columbia look like a high outlier in the table for large cities? Whether or not the District of Columbia is an outlier depends on more than its crime rate. It also depends on the other observations included in the data set. Which data set do you feel is more appropriate for the District of Columbia?

WHY DID YOU DO THAT?

There is no single best way to organize our presentation of statistics to beginners. That said, our choices reflect thinking about both content and pedagogy. Here are comments on several “frequently asked questions” about the order and selection of material in BPS.

Why does the distinction between population and sample not appear in Part I? There is more to statistics than inference. In fact, statistical inference is appropriate only in rather special circumstances. The chapters in Part I present tools and tactics for describing data—any data. These tools and tactics do not depend on the idea of inference from sample to population. Many data sets in these chapters (for example, the several sets of data about the 50 states) do not lend themselves to inference because they represent an entire population. John Tukey of Bell Labs and Princeton, the philosopher of modern data analysis, insisted that the population-sample distinction be avoided when it is not relevant. He used the word “batch” for data sets in general. We see no need for a special word, but we think Tukey was right.

Why not begin with data production? It is certainly reasonable to do so—the natural flow of a planned study is from design to data analysis to inference. But in their future employment most students will use statistics mainly in settings other than planned research studies. We place the design of data production (Chapters 8 and 9) after data analysis to emphasize that data-analytic techniques apply to any data. One of the primary purposes of statistical designs for producing data is to make inference possible, so the discussion in Chapters 8 and 9 opens Part II and motivates the study of probability.

Why do Normal distributions appear in Part I? Density curves such as the Normal curves are just another tool to describe the distribution of a quantitative variable, along with stemplots, histograms, and boxplots. Professional statistical software offers to make density curves from data just as it offers histograms. We prefer not to suggest that this material is essentially tied to probability, as the traditional order does. And we find it very helpful to break up the indigestible lump of probability that troubles students so much. Meeting Normal distributions early does this and strengthens the “probability distributions are like data distributions” way of approaching probability.

Why not delay correlation and regression until late in the course, as was traditional? BPS begins by offering experience working with data and gives a conceptual structure for this nonmathematical but essential part of statistics. Students profit from more experience with data and from seeing the conceptual structure worked out in relations among variables as well as in describing single-variable data. Correlation and least-squares regression are very important descriptive tools and are often used in settings where there is no population-sample distinction, such as studies of all a firm’s employees. Perhaps most important, the BPS approach asks students to think about what kind of relationship lies behind the data (confounding, lurking variables, association doesn’t imply causation, and so on), without overwhelming them with the demands of formal inference methods. Inference in the correlation and regression setting is a bit complex, demands software, and often comes right at the end of the course. We find that delaying all mention of correlation and regression to that point means that students often don’t master the basic uses and properties of these methods. We consider Chapters 4 and 5 (correlation and regression) essential and Chapter 24 (regression inference) optional.

What about probability? Much of the usual formal probability appears in the optional Chapters 12 and 13. Chapters 10 and 11 present in a less formal way the ideas of probability and sampling distributions that are needed to understand inference. These two chapters follow a straight line from the idea of probability as long-term regularity, through concrete ways of assigning probabilities, to the central idea of the sampling distribution of a statistic. The law of large numbers and the central limit theorem appear in the context of discussing the sampling distribution of a sample mean. What is left to Chapters 12 and 13 is mostly “general probability rules” (including conditional probability) and the binomial distributions.

We suggest that you omit Chapters 12 and 13 unless you are constrained by external forces. Experienced teachers recognize that students find probability difficult. Research on learning confirms our experience. Even students who can do formally posed probability problems often have a very fragile conceptual grasp of probability ideas. Attempting to present a substantial introduction to probability in a data-oriented statistics course for students who are not mathematically trained is in our opinion unwise. Formal probability does not help these students master the ideas of inference (at least not as much as we teachers often imagine), and it depletes reserves of mental energy that might better be applied to essentially statistical ideas.

Why use the z procedures for a population mean to introduce the reasoning of inference? This is a pedagogical issue, not a question of statistics in practice. Sometime in the golden future we will start with resampling methods. We think that permutation tests make the reasoning of tests clearer than any traditional approach. For now the main choices are z for a mean and z for a proportion.

We find z for means quite a bit more accessible to students. Positively, we can say up front that we are going to explore the reasoning of inference in the overly simple setting described in the box on page 352 titled “Simple Conditions for Inference about a Mean.” As this box suggests, exactly Normal population and true simple random sample are as unrealistic as known σ . All the issues of practice—robustness against lack of Normality and application when the data aren’t an SRS as well as the need to estimate σ —are put off until, with the reasoning in hand, we discuss the practically useful t procedures. This separation of initial reasoning from messier practice works well. And it allows us to discuss optional topics such as Type I and II errors and the power of a statistical test.

Negatively, starting with inference for p introduces many side issues: no exact Normal sampling distribution, but a Normal approximation to a discrete distribution; use of \hat{p} in both the numerator and denominator of the test statistic to estimate both the parameter p and \hat{p} ’s own standard deviation; loss of the direct link between test and confidence interval. Once upon a time we had at least the compensation of developing practically useful procedures. Now the often gross inaccuracy of the traditional z confidence interval for p is better understood. See the following explanation.

Why does the presentation of inference for proportions go beyond the traditional methods? Computational and theoretical work has demonstrated convincingly that the standard confidence intervals for proportions can be trusted only for very large sample sizes. It is hard to abandon old friends, but we think that a look at the graphs in Section 2 of the paper by Brown, Cai, and DasGupta in the May 2001 issue of *Statistical Science* is both distressing and persuasive.¹ The standard intervals often have a true confidence level much less than what was requested, and requiring larger samples encounters a maze of “lucky” and “unlucky” sample sizes until very large samples are reached. Fortunately, there is a simple cure: just add two successes and two failures to your data. We present these “plus four intervals” in Chapters 20 and 21, along with guidelines for use.

Why didn't you cover Topic X? Introductory texts ought not to be encyclopedic. Including each reader's favorite special topic results in a text that is formidable in size and intimidating to students. We chose topics on two grounds: they are the most commonly used in practice, and they are suitable vehicles for learning broader statistical ideas. Students who have completed the core of BPS, Chapters 1–11 and 14–22, will have little difficulty moving on to more elaborate methods. There are, of course, seven chapters that discuss more advanced topics, three in this volume and four available on CD and online, to begin the next stages of learning.

ACKNOWLEDGMENTS

We are grateful to colleagues from two-year and four-year colleges and universities who commented on successive drafts of the manuscript.

Thanks to Professor Patricia Humphrey of Georgia Southern University who read the manuscript in detail and offered advice.

Thanks to Professor Jackie Miller of The Ohio State University for carefully reading all chapters, spotting many errors, and making numerous useful suggestions.

Others who provided comments:

James Adcock, University of Western Ontario
Brad Bailey, North Georgia College and State University
E. N. Barron, Loyola University, Chicago
Jennifer Beineke, Western New England College
Diane Benner, Harrisburg Area Community College
Zoubir Benzaid, University of Wisconsin, Oshkosh
Jennifer Borrello, Baylor University
Smiley Cheng, University of Manitoba, Winnipeg
Patti Collings, Brigham Young University
Tadd Colver, Purdue University
Patti Costello, Eastern Kentucky University
James Curl, Modesto Junior College
DeWayne Derryberry, Idaho State University
Jonathan Duggins, Virginia Tech
Lacy Echols, Butler University
Chris Edwards, University of Wisconsin, Oshkosh
Margaret Elrich, Georgia Perimeter College
Karen Estes, St. Petersburg College
Eugene Galperin, East Stroudsburg University
Mark Gebert, Eastern Kentucky University
Kim Gilbert, Clayton State University
Aaron Gladish, Austin Community College
Ellen Gundlach, Purdue University
Arjun Gupta, Bowling Green State University
Brenda Gunderson, University of Michigan

Leslie Hendrix, University of South Carolina
Jeanne Hill, Baylor University
Dawn Holmes, University of California, Santa Barbara
Patricia Humphrey, Georgia Southern University
Thomas Ilvento, University of Delaware
Mark Jacobson, University of Northern Iowa
Marc Kirschenbaum, John Carroll University
Greg Knofczynski, Armstrong Atlantic State University
Mike Kowalski, University of Alberta
Zhongshan Li, Georgia State University
Michael Licher, University of Buffalo
Tom Linton, Central College
William Liu, Bowling Green—Firelands College
Amy Maddox, Grand Rapids Community College
Steve Marsden, Glendale Community College
Catherine Matos, Clayton State University
Darcy Mays, Virginia Commonwealth University
Andrew McDougall, Montclair State University
Bill Meisel, Florida Community College—Jacksonville
Nancy Mendell, State University of New York, Stony Brook
Scott McClintock, West Chester University
Lynne Nielsen, Brigham Young University
Helen Noble, San Diego State University
Richard Numrich, College of Southern Nevada
Melvin Nyman, Alma College

Darlene Olsen, *Norwich University*
Eric Packard, *Mesa State University*
Mary Parker, *Austin Community College, Rio Grande Campus*
Don Porter, *Beloit College*
Bob Price, *East Tennessee State University*
Asoka Ramanayake, *University of Wisconsin, Oshkosh*
Eric Rudrud, *St. Cloud State University*
Christopher Richter, *Queens University*
Scott Richter, *University of North Carolina, Greensboro*
Corlis Robe, *East Tennessee State University*
Deborah Rumsey, *The Ohio State University*
Therese Shelton, *Southwestern University*
Rob Sinn, *North Georgia College*

Eugenia Skirta, *East Stroudsburg University*
Jim Smart, *Tallahassee Community College*
Dianna Spence, *North Georgia College and State University*
Tim Swartz, *Simon Fraser University*
Suzhong Tian, *Husson College*
Suzanne Tourville, *Columbia College*
Christopher Tripler, *Endicott College*
Gail Tudor, *Husson College*
Ramin Vakilian, *California State University, Northridge*
David Vlieger, *Northwest Missouri State University*
Joseph Walker, *Georgia State University*
Steve Waters, *Pacific Union College*
Yuanhui Xiao, *Georgia State University*
Yichuan Zhao, *Georgia State University*

We are also grateful to Craig Bleyer, Ruth Baruth, Karen Carson, Bruce Kaplan, Shona Burke, Mary Louise Byrd, Andrew Sylvester, Leslie Lahr, R. Scott Linder, Blake Logan, Pamela Bruton, and the other editorial and design professionals who have contributed greatly to the attractiveness of this book. Special thanks to Denise Showers at Aptara for her patience and careful attention.

Finally, we are indebted to the many statistics teachers with whom we have discussed the teaching of our subject over many years; to people from diverse fields with whom we have worked to understand data; and especially to students whose compliments and complaints have changed and improved our teaching. Working with teachers, colleagues in other disciplines, and students constantly reminds us of the importance of hands-on experience with data and of statistical thinking in an era when computer routines quickly handle statistical details.

David S. Moore, Michael Fligner, and William I. Notz



Media and Supplements

FOR STUDENTS



www.yourstatsportal.com (Access code or online purchase required.)

StatsPortal is the digital gateway to *The Basic Practice of Statistics*, Sixth Edition, designed to enrich the course and enhance students' study skills through a collection of Web-based tools. StatsPortal integrates a rich suite of diagnostic, assessment, tutorial, and enrichment features, enabling students to master statistics at their own pace. StatsPortal is organized around three main teaching and learning components:

1. Interactive eBook offers a complete and customizable online version of the text, fully integrated with all the media resources available with the sixth edition of BPS. The eBook allows students to quickly search the text, highlight key areas, and add notes about what they're reading. Instructors can customize the eBook to add, hide, and reorder content, add their own material, and highlight key text for students.

2. Resources organizes all the resources for BPS into one location for ease of use.

Student Resources

- **NEW! Statistical Video Series** consisting of StatClips, StatClips Examples, and Statistically Speaking "Snapshots." View animated lecture videos, whiteboard lessons, and documentary-style footage that illustrate key statistical concepts and help students visualize statistics in real world scenarios.
- **StatTutor Tutorials** offer audio-multimedia tutorials tied directly to the textbook, containing videos, applets, and animations.
- **NEW! LEARNINGCurve** is a formative quizzing system that offers immediate feedback at the question level to help students master course material.
- **Statistical Applets** offer a series of interactive applets to help students master key statistical concepts and work exercises from the text.
- **CrunchIt![®] Statistical Software** allows users to analyze data from any Internet location. Designed with the novice user in mind, the software is not only easily accessible but also easy to use. CrunchIt![®] offers all the basic statistical routines covered in introductory statistics courses and more.
- **Stats@Work Simulations** put students in the role of statistical consultant, helping them better understand statistics interactively within the context of real-life scenarios.
- **EESEE Case Studies**, developed by The Ohio State University Statistics Department, teach students to apply their statistical skills by exploring actual case studies using real data.
- **Data sets** are available in ASCII, Excel, TI, Minitab, SPSS, an IBM Company*, S-PLUS, and JMP formats.

*SPSS was acquired by IBM in October 2009.

- **Student Solutions Manual** provides solutions to the odd-numbered exercises, with stepped out solutions to select problems.
- **Statistical Software Manuals** for TI-83/84, Minitab, Excel, JMP, and SPSS provide instruction, examples, and exercises using specific statistical software packages.
- **Interactive Table Reader** allows students to use statistical tables interactively to seek the information they need.
- **Tables**

Resources for instructors only

- **Instructor's Guide with Full Solutions** includes teaching suggestions, chapter comments, and detailed solutions to all exercises.
- **Test Bank** offers hundreds of multiple choice questions.
- **Lecture PowerPoint slides** offer a detailed lecture presentation of statistical concepts covered in each chapter of *BPS*.
- **NEW! SolutionMaster** is a Web-based version of the solutions in the Instructor's Guide with Full Solutions. This easy-to-use tool allows instructors to generate a solution file for any set of homework exercises. Solutions can be downloaded in PDF format for convenient printing and posting. For more information or a demonstration, contact your local W. H. Freeman sales representative.

3. Assignments organizes assignments and guides instructors through an easy-to-create assignment process providing access to questions from the Test Bank and exercises from the text, including many algorithmic problems. The Assignment Center enables instructors to create their own assignments from a variety of question types for machine-gradable assignments. This powerful assignment manager allows instructors to select their preferred policies for scheduling, maximum attempts, time limitations, feedback, and more!

Online Study Center

www.whfreeman.com/osc/bps6e

(Access code or online purchase required.)

The Online Study Center offers all the resources available in StatsPortal except the eBook and Assignment Center.

Companion Web site

www.whfreeman.com/bps6e

This open-access Web site includes statistical applets, data sets, supplementary exercises, statistical profiles, and self-quizzes. The Web site also offers four optional companion chapters covering bootstrap methods and permutation tests and statistics for quality control and capability.

Interactive Student CD-ROM

Included with every new copy of the sixth edition of BPS, the CD contains access to all the content available on the Companion Web site. CrunchIt!® statistical software and EESEE case studies are available via an access-code-protected Web site. (An access code is included with every new text.)

Special Software Packages

Student versions of JMP, Minitab, S-PLUS, and SPSS are available on a CD-ROM packaged with the textbook. This software is not sold separately; it must be packaged with a text or a manual. Contact your W. H. Freeman representative for information or visit www.whfreeman.com.

NEW! Video Tool Kit

(Access code or online purchase required.)

This new Statistical Video Series consists of three types of videos aimed to illustrate key statistical concepts and help students visualize statistics in real-world scenarios:

- **StatClips lecture videos**, created and presented by Alan Dabney, PhD, Texas A&M University, are innovative visual tutorials that illustrate key statistical concepts. In 3 to 5 minutes, each StatClips video combines dynamic animation, data sets, and interesting scenarios to help students understand the concepts in an introductory statistics course.
- In **StatClips Examples**, Alan Dabney walks students through step-by-step examples related to the StatClips lecture videos to reinforce the concepts through problem solving.
- **SnapShots** videos are abbreviated, student friendly versions of the **Statistically Speaking** video series, and they bring the world of statistics into the classroom. In the same vein as the successful PBS series *Against All Odds Statistics*, **Statistically Speaking** uses new and updated documentary footage and interviews that show real people using data analysis to make important decisions in their careers and in their daily lives. From business to medicine, from the environment to understanding the census, SnapShots focus on why statistics is important for students' careers, and how statistics can be a powerful tool to understand their world.

Printed Student Solutions Manual

This printed manual provides step-by-step solutions for all odd-numbered exercises in the text. ISBN: 1-4292-8000-X

Software Manuals

Software manuals covering Minitab, Excel, SPSS, TI-83/84, and JMP are offered within StatsPortal and the Online Study Center. These manuals are also available

in printed versions through custom publishing. They serve as basic introductions to popular statistical software options and guides to their use with the sixth edition of *BPS*.

FOR INSTRUCTORS

Instructor's Web site

www.whfreeman.com/bps6e

Requires user registration as an instructor and features all the student Web materials plus:

- **Instructor version of EESEE** (Electronic Encyclopedia of Statistical Examples and Exercises), with solutions to the exercises in the student version.
- **PowerPoint slides** containing all textbook figures and tables.
- **Lecture PowerPoint slides** offering a detailed lecture presentation of statistical concepts covered in each chapter of *BPS* 6e.
- **Full answers to the Supplementary Exercises** supplement on the student Web site.

Printed Instructor's Guide with Full Solutions

This printed guide includes full solutions to all exercises and provides additional examples and data sets for class use, Internet resources, and sample examinations. It also contains brief discussions of the *BPS* approach for each chapter.
ISBN: 1-4641-1406-4

Test Bank

The test bank contains hundreds of multiple-choice questions to use in generating quizzes and tests for each chapter of the text. Available in print as well as electronically on CD-ROM (for Windows and Mac), where questions can be downloaded, edited, and resequenced to suit each instructor's needs.

Printed Version, ISBN: 1-4641-1492-7

Computerized (CD) Version, ISBN: 1-4641-0891-9

Enhanced Instructor's Resource CD-ROM

The CD allows instructors to search and export (by key term or chapter) all the material from the student CD, plus:

- All text images and tables
- Instructor's Guide with full solutions
- PowerPoint files and lecture slides
- Test bank files

ISBN: 1-4641-0893-5

Course Management Systems

W. H. Freeman and Company provides courses for Blackboard, WebCT (Campus Edition and Vista), Angel, Desire2Learn, Moodle, and Sakai course management systems. These solutions are completely integrated, and you can easily customize and adapt them to meet your teaching goals and course objectives. Visit www.bfwpub.com/lms for more information.

i-clicker

i-clicker is a two-way radio frequency classroom response solution developed by educators for educators. University of Illinois physicists Tim Stelzer, Gary Gladding, Mats Selen, and Benny Brown created the i-clicker system after using competing classroom response solutions and discovering that they were neither classroom-appropriate nor student-friendly. Each step of i-clicker's development has been informed by teaching and learning. i-clicker is superior to other systems from both a pedagogical and technical standpoint. To learn more about packaging i-clicker with this textbook, please contact your local sales rep or visit www.iclicker.com.





About the Authors

David S. Moore is Shanti S. Gupta Distinguished Professor of Statistics, Emeritus, at Purdue University and was the 1998 president of the American Statistical Association. He received his A.B. from Princeton and his Ph.D. from Cornell, both in mathematics. He has written many research papers in statistical theory and served on the editorial boards of several major journals. Professor Moore is an elected fellow of the American Statistical Association and of the Institute of Mathematical Statistics and an elected member of the International Statistical Institute. He has served as program director for statistics and probability at the National Science Foundation.

In recent years, Professor Moore has devoted his attention to the teaching of statistics. He was the content developer for the Annenberg/Corporation for Public Broadcasting college-level telecourse *Against All Odds: Inside Statistics* and for the series of video modules *Statistics: Decisions through Data*, intended to aid the teaching of statistics in schools. He is the author of influential articles on statistics education and of several leading texts. Professor Moore has served as president of the International Association for Statistical Education and has received the Mathematical Association of America's national award for distinguished college or university teaching of mathematics.

William I. Notz is Professor of Statistics at The Ohio State University. He received his B.S. in physics from the Johns Hopkins University and his Ph.D. in mathematics from Cornell University. His first academic job was as an assistant professor in the Department of Statistics at Purdue University. While there, he taught the introductory concepts course with Professor Moore and as a result of this experience he developed an interest in statistical education. Professor Notz is a coauthor of EESEE (the Electronic Encyclopedia of Statistical Examples and Exercises) and coauthor of *Statistics: Concepts and Controversies*.

Professor Notz's research interests have focused on experimental design and computer experiments. He is the author of several research papers and of a book on the design and analysis of computer experiments. He is an elected fellow of the American Statistical Association. He has served as the editor of the journal *Technometrics* and as editor of the *Journal of Statistics Education*. He has served as the director of the Statistical Consulting Service, as acting chair of the Department of Statistics for a year, and as an associate dean in the College of Mathematical and Physical Sciences at The Ohio State University. He is a winner of The Ohio State University's Alumni Distinguished Teaching Award.

Michael A. Fligner is Professor Emeritus at The Ohio State University. He received his B.S. in mathematics from the State University of New York at Stony Brook and his Ph.D. from the University of Connecticut. He has spent his entire professional career at The Ohio State University where he was vice chair of the

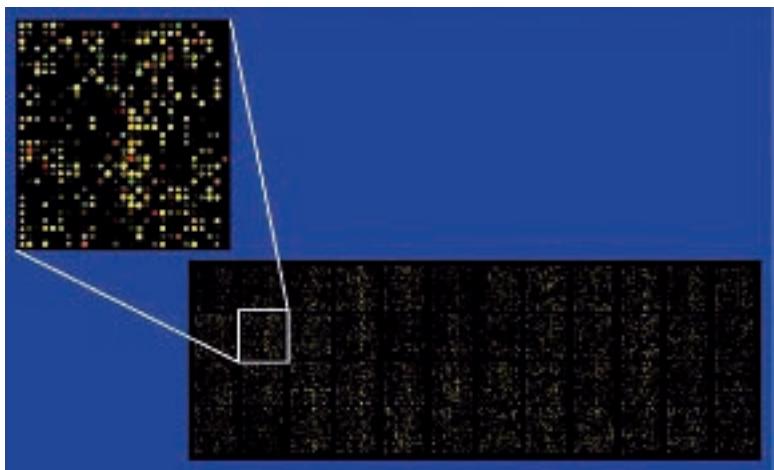
Department for over 10 years and also served as Director of the Statistical Consulting Service. He has done consulting work with several large corporations in Central Ohio.

Professor Fligner's research interests are in nonparametric statistical methods, and he received the Statistics in Chemistry award from the American Statistical Association for work on detecting biologically active compounds. He is a coauthor of the book *Statistical Methods for Behavioral Ecology* and received a Fulbright scholarship under the American Republics Research program to work at the Charles Darwin Research Station in the Galápagos Islands. He has been an associate editor of the *Journal of Statistical Education*.



To the Student

STATISTICAL THINKING



What genes are active in a tissue? Answering this question can unravel basic questions in biology, distinguish cancer cells from normal cells, and distinguish between closely related types of cancer. To learn the answer, apply the tissue to a “microarray” that contains thousands of snippets of DNA arranged in a grid on a chip about the size of your thumb. As DNA in the tissue binds to the snippets in the array, special recorders pick up spots of light of varying color and intensity across the grid and store what they see as numbers.

What's hot in popular music this week? SoundScan knows. SoundScan collects data electronically from the cash registers in more

than 14,000 retail outlets and also collects data on download sales from Web sites. When you buy a CD or download a digital track, the checkout scanner or Web site is probably telling SoundScan what you bought. SoundScan provides this information to *Billboard* magazine, MTV, and VH1, as well as to record companies and artists' agents.

Should women take hormones such as estrogen after menopause, when natural production of these hormones ends? In 1992, several major medical organizations said “Yes.” In particular, women who took hormones seemed to reduce their risk of a heart attack by 35% to 50%. The risks of taking hormones appeared small compared with the benefits. But in 2002, the National Institutes of Health declared these findings wrong. Use of hormones after menopause immediately plummeted.

Both recommendations were based on extensive studies. What happened? DNA microarrays, SoundScan, and medical studies all produce data (numerical facts), and lots of them. Using data effectively is a large and growing part of most professions. Reacting to data is part of everyday life. That's why statistics is important:

STATISTICS IS THE SCIENCE OF LEARNING FROM DATA

Data are numbers, but they are not “just numbers.” **Data are numbers with a context.** The number 10.5, for example, carries no information by itself. But if we hear that a friend's new baby weighed 10.5 pounds at birth, we congratulate her on the healthy size of the child. The context engages our background knowledge and allows us to make judgments. We know that a baby weighing 10.5 pounds is quite large, and that a human baby is unlikely to weigh 10.5 ounces or 10.5 kilograms. The context makes the number informative.

To gain insight from data, we make graphs and do calculations. But graphs and calculations are guided by ways of thinking that amount to educated common-sense. Let's begin our study of statistics with an informal look at some principles of statistical thinking.¹

WHERE THE DATA COME FROM MATTERS

What's behind the flip-flop in the advice offered to women about hormone replacement? The evidence in favor of hormone replacement came from a number of *observational studies* that compared women who were taking hormones with others who were not. But women who choose to take hormones are very different from women who do not: they are richer and better educated and see doctors more often. These women do many things to maintain their health. It isn't surprising that they have fewer heart attacks.

Large and careful observational studies are expensive, but they are easier to arrange than careful *experiments*. Experiments don't let women decide what to do. They assign women to either hormone replacement or to dummy pills that look and taste the same as the hormone pills. The assignment is done by a coin toss, so that all kinds of women are equally likely to get either treatment. Part of the difficulty of a good experiment is persuading women to accept the result—invisible to them—of the coin toss. By 2002, several experiments agreed that hormone replacement does not reduce the risk of heart attacks, at least for older women. Faced with this better evidence, medical authorities changed their recommendations.²

Of course, observational studies are often useful. We can learn from observational studies how chimpanzees behave in the wild or which popular songs sold best last week or what percent of workers were unemployed last month. SoundScan's data on popular music and the government's data on employment and unemployment come from sample surveys, an important kind of observational study that chooses a part(the sample) to represent a larger whole. Opinion polls interview perhaps 1000 of the 235 million adults in the United States to report the public's views on current issues. Can we trust the results? We'll see that this isn't a simple yes-or-no question. Let's just say that the government's published unemployment rate is much more trustworthy than opinion poll results, and not just because the Bureau of Labor Statistics interviews 60,000 people rather than 1000.

We can, however, say right away that some samples can't be trusted. The advice columnist Ann Landers once asked her readers, "If you had it to do over again, would you have children?" A few weeks later, her column was headlined "70% OF PARENTS SAY KIDS NOT WORTH IT." Indeed, 70% of the nearly 10,000 parents who wrote in said they would not have children if they could make the choice again. Those 10,000 parents were upset enough with their children to write Ann Landers. Most parents are happy with their kids and don't bother to write. Statistically designed samples, even opinion polls, don't let people choose themselves for the sample. They interview people selected by impersonal chance so that everyone has an equal opportunity to be in the sample. Such a poll showed

that 91% of parents would have children again. Where data come from matters a lot. If you are careless about how you get your data, you may announce 70% “No” when the truth is close to 90% “Yes.”

ALWAYS LOOK AT THE DATA

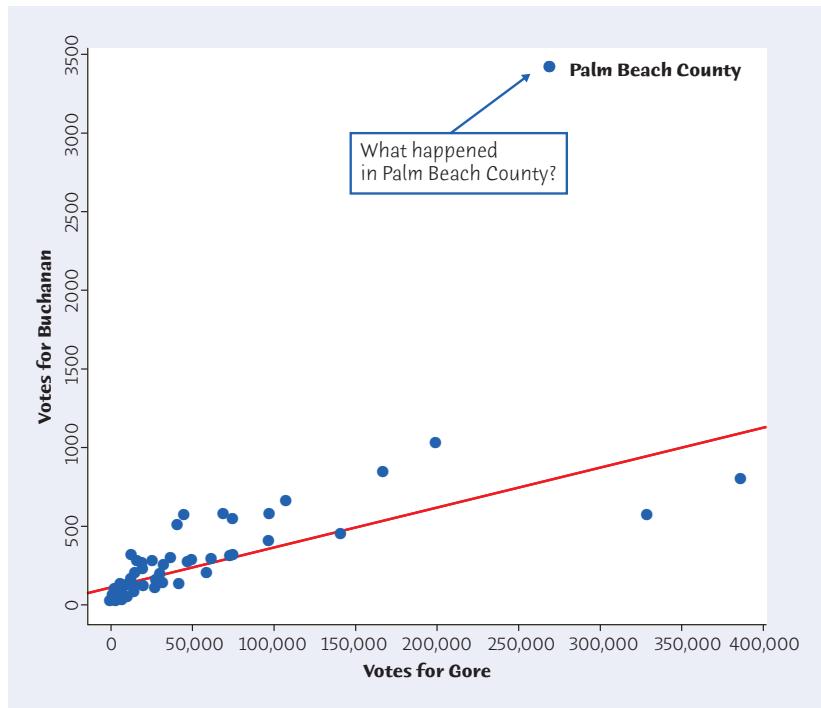
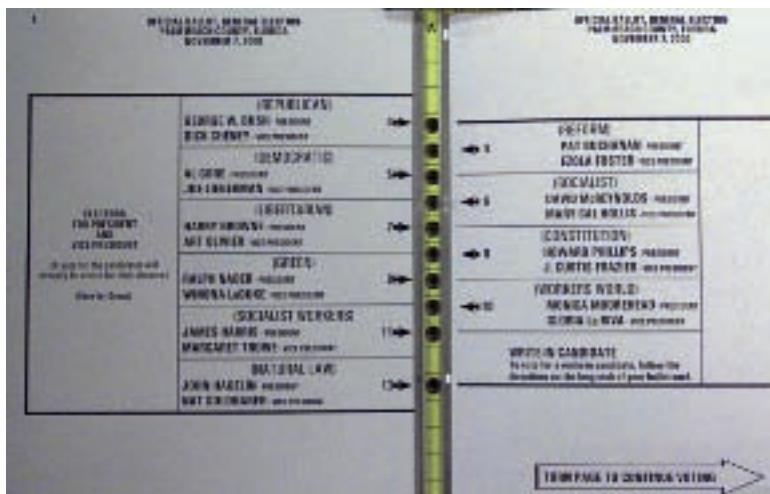


FIGURE 1

Votes in the 2000 presidential election for Al Gore and Patrick Buchanan in Florida's 67 counties. What happened in Palm Beach County?



Yogi Berra said it: “You can observe a lot by just watching.” That’s a motto for learning from data. A few *carefully chosen graphs are often more instructive than great piles of numbers*. Consider the outcome of the 2000 presidential election in Florida.

Elections don’t come much closer: after much recounting, state officials declared that George W. Bush had carried Florida by 537 votes out of almost 6 million votes cast. Florida’s vote decided the election and made George W. Bush, rather than Al Gore, president. Let’s look at some data. Figure 1 displays a graph that plots votes for the third-party candidate Pat Buchanan against votes for the Democratic candidate Al Gore in Florida’s 67 counties.

What happened in Palm Beach County? The question leaps from the graph. In this large and heavily Democratic county, a conservative third-party candidate did far better relative to the Democratic candidate than in any other county. The points for the other 66 counties show votes for both candidates increasing together in a roughly straight-line pattern. Both counts go up as county population goes up. Based on this pattern, we would expect Buchanan to receive around 800 votes in Palm Beach County. He actually received more than 3400 votes. That difference determined the election result in Florida and in the nation.

The graph demands an explanation. It turns out that Palm Beach County used a confusing “butterfly” ballot, in which candidate names on both left and right pages led to a voting column in the center. It would be easy for a voter who intended to vote for Gore to in fact cast a vote for Buchanan. The

graph is convincing evidence that this in fact happened, more convincing than the complaints of voters who (later) were unsure where their votes ended up.

BEWARE THE LURKING VARIABLE

Women who chose hormone replacement after menopause were on the average richer and better educated than those who didn't. No wonder they had fewer heart attacks. Children who play soccer tend to have prosperous and well-educated parents. No wonder they do better in school (on the average) than children who don't play soccer. We can't conclude that hormone replacement reduces heart attacks or that playing soccer increases school grades just because we see these relationships in data. In both examples, education and affluence are lurking variables, background factors that help explain the relationships between hormone replacement and good health and between soccer and good grades.

Almost all relationships between two variables are influenced by other variables lurking in the background. To understand the relationship between two variables, you must often look at other variables. Careful statistical studies try to think of and measure possible lurking variables in order to correct for their influence. As the hormone saga illustrates, this doesn't always work well. News reports often just ignore possible lurking variables that might ruin a good headline like "Playing soccer can improve your grades." The habit of asking, "What might lie behind this relationship?" is part of thinking statistically.

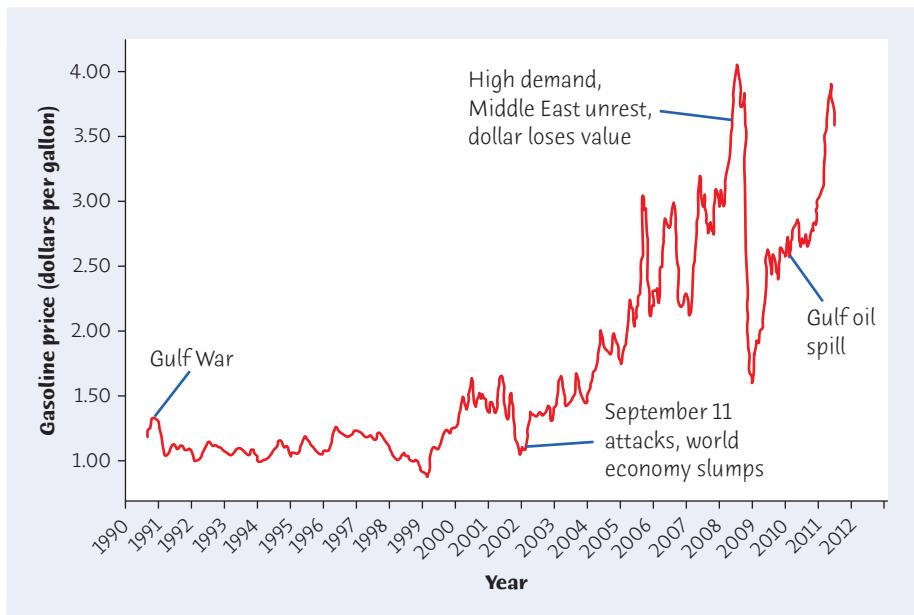
VARIATION IS EVERYWHERE

The company's sales reps file into their monthly meeting. The sales manager rises. "Congratulations! Our sales were up 2% last month, so we're all drinking champagne this morning. You remember that when sales were down 1% last month I fired half of our reps." This picture is only slightly exaggerated. Many managers overreact to small short-term variations in key figures. Here is Arthur Nielsen, head of the country's largest market research firm, describing his experience:

Too many business people assign equal validity to all numbers printed on paper. They accept numbers as representing Truth and find it difficult to work with the concept of probability. They do not see a number as a kind of shorthand for a range that describes our actual knowledge of the underlying condition.³

Business data such as sales and prices vary from month to month for reasons ranging from the weather to a customer's financial difficulties to the inevitable errors in gathering the data. The manager's challenge is to say when there is a real pattern behind the variation. We'll see that statistics provides tools for understanding variation and for seeking patterns behind the screen of variation.

Let's look at some more data. Figure 2 plots the average price of a gallon of regular unleaded gasoline each week from September 1990 to June 2011.⁴ There certainly is variation! But a close look shows a yearly pattern: gas prices go up during the summer driving season, then down as demand drops in the fall. On top of this regular pattern we see the effects of international events. For

**FIGURE 2**

Variation is everywhere: the average retail price of regular unleaded gasoline, 1990 to mid 2011.

Variation is everywhere. Individuals vary; repeated measurements on the same individual vary; almost everything varies over time. One reason we need to know some statistics is that statistics helps us deal with variation.

CONCLUSIONS ARE NOT CERTAIN

Cervical cancer is second only to breast cancer as a cause of cancer deaths in women. Almost all cervical cancers are caused by human papillomavirus (HPV). The first vaccine to protect against the most common varieties of HPV became available in 2006. The Centers for Disease Control and Prevention recommend that all girls be vaccinated at age 11 or 12.

How well does the vaccine work? Doctors rely on experiments (called “clinical Trials” in medicine) that give some women the new vaccine and others a dummy vaccine. (This is ethical when it is not yet known whether or not the vaccine is safe and effective.) The conclusion of the most important trial was that an estimated 98% of women up to age 26 who are vaccinated before they are infected with HPV will avoid cervical cancers over a 3-year period.

On the average, women who get the vaccine are much less likely to get cervical cancer. But because variation is everywhere, the results are different for different women. Some vaccinated women will get cancer, and many who are not vaccinated will escape. Statistical conclusions are “on the average” statements only. Well then, can we at least be certain that the vaccine reduces risk on the average? No. We can be very confident, but we can’t be certain.

example, prices rose when the 1990 Gulf War threatened oil supplies and dropped when the world economy turned down after the September 11, 2001, terrorist attacks in the United States. The years 2007 and 2008 brought the perfect storm: the ability to produce oil and refine gasoline was overwhelmed by high demand from China and the United States and continued turmoil in the oil-producing areas of the Middle East and Nigeria. Add a rapid fall in the value of the dollar, and prices at the pump skyrocketed to more than \$4 per gallon. In 2010 the Gulf oil spill also affected supply and hence prices. The data carry an important message: because the United States imports most of its oil, we can’t control the price we pay for gasoline.

Because variation is everywhere, conclusions are uncertain. Statistics gives us a language for talking about uncertainty that is used and understood by statistically literate people everywhere. In the case of HPV vaccine, the medical journal used that language to tell us: “Vaccine efficiency . . . was 98% (95 percent confidence interval 86% to 100%).”⁵ That “98% effective” is, in Arthur Nielsen’s words, “shorthand for a range that describes our actual knowledge of the underlying condition.” The range is 86% to 100%, and we are 95 percent confident that the truth lies in that range. We will soon learn to understand this language. We can’t escape variation and uncertainty. Learning statistics enables us to live more comfortably with these realities.

STATISTICAL THINKING AND YOU

What Lies Ahead in This Book

The purpose of *The Basic Practice of Statistics* (BPS) is to give you a working knowledge of the ideas and tools of practical statistics. We will divide practical statistics into three main areas:

1. **Data analysis** concerns methods and strategies for exploring, organizing, and describing data using graphs and numerical summaries. Only organized data can illuminate reality. Only thoughtful exploration of data can defeat the lurking variable. Part I of BPS (Chapters 1–7) discusses data analysis.
2. **Data production** provides methods for producing data that can give clear answers to specific questions. Where the data come from really is important. Basic concepts about how to select samples and design experiments are the most influential ideas in statistics. These concepts are the subject of Chapters 8 and 9.
3. **Statistical inference** moves beyond the data in hand to draw conclusions about some wider universe, taking into account that variation is everywhere and that conclusions are uncertain. To describe variation and uncertainty, inference uses the language of probability, introduced in Chapters 10 and 11. Because we are concerned with practice rather than theory, we need only a limited knowledge of probability. Chapters 12 and 13 offer more probability for students who want it. Chapters 14, 15, and 16 discuss the reasoning of statistical inference. These chapters are the key to the rest of the book. Chapters 18–21 present inference as used in practice in the most common settings. Chapters 23–25 and the Optional Companion Chapters 26–29 on the text CD and online concern more advanced or specialized kinds of inference.

Because data are numbers with a context, doing statistics means more than manipulating numbers. You must **state** a problem in its real-world context, **plan** your specific statistical work in detail, **solve** the problem by making the necessary graphs and calculations, and **conclude** by explaining what your findings say about the real-world setting. We’ll make regular use of this four-step process to encourage good habits that go beyond graphs and calculations to ask, “What do the data tell me?”

Statistics does involve lots of calculating and graphing. The text presents the techniques you need, but you should use technology to automate calculations and graphs as much as possible. Because the big ideas of statistics don't depend on any particular level of access to technology, BPS does not require software or a graphing calculator until we reach the more advanced methods in Part IV of the text. Even if you make little use of technology, you should look at the "Using Technology" sections throughout the book. You will see at once that you can read and apply the output from almost any technology used for statistical calculations. The ideas really are more important than the details of how to do the calculations.

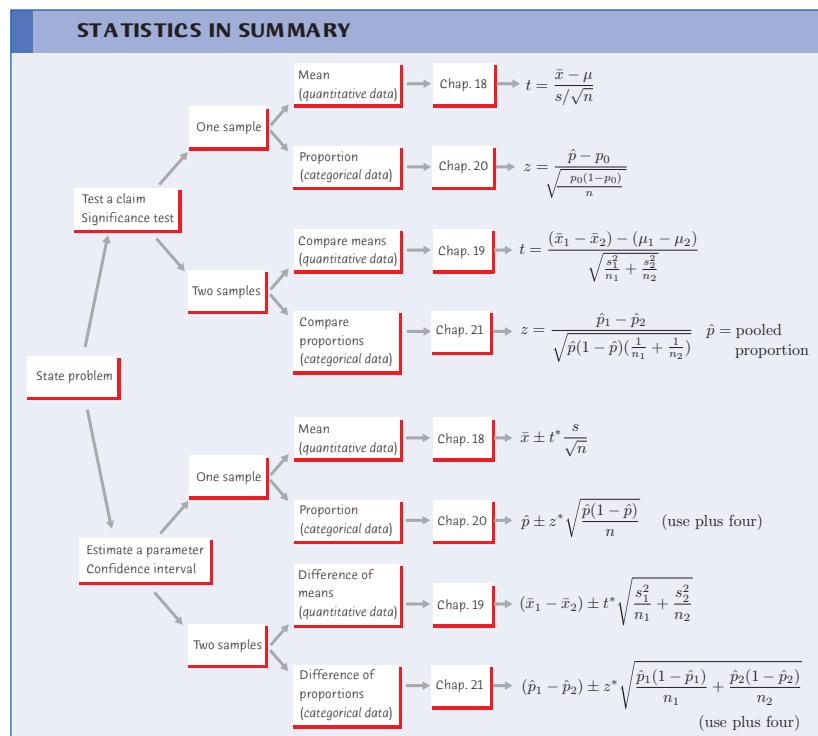
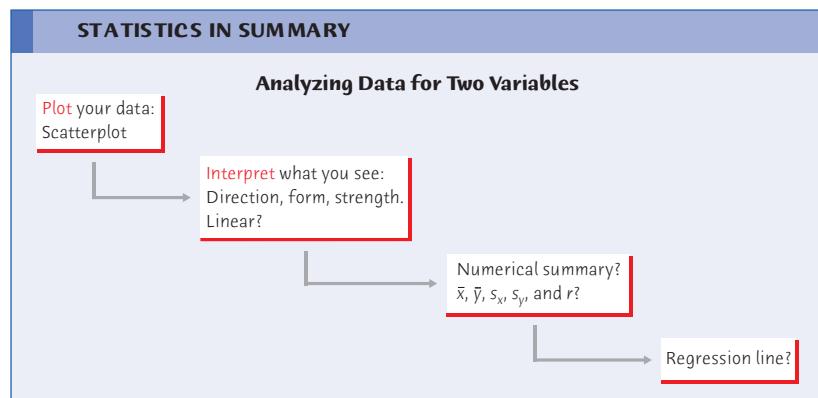
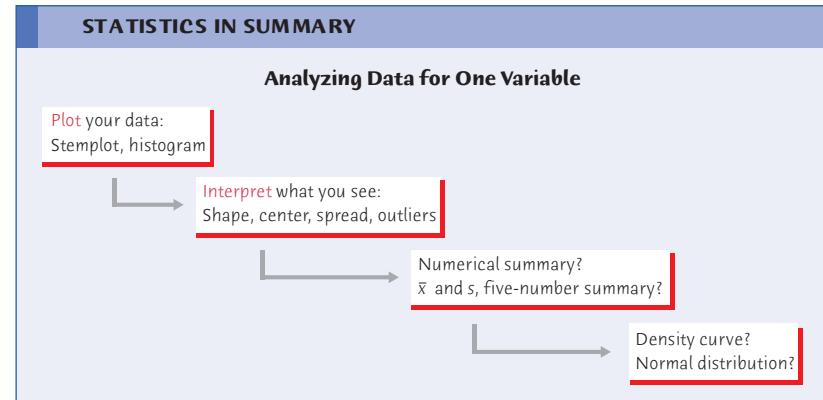
Unless you have constant access to software or a graphing calculator, *you will need a basic calculator with some built-in statistical functions*. Specifically, your calculator should find means and standard deviations and calculate correlations and regression lines. Look for a calculator that claims to do "two-variable statistics" or "regression."

Because graphing and calculating are automated in statistical practice, the most important assets you can gain from the study of statistics are an understanding of the big ideas and the beginnings of good judgment in working with data. BPS tries to explain the most important ideas of statistics, not just teach methods. Some examples of big ideas that you will meet (one from each of the three areas of statistics) are "always plot your data," "randomized comparative experiments," and "statistical significance."

You learn statistics by doing statistical problems. As you read, you will see several levels of exercises, arranged to help you learn. Short "Apply Your Knowledge" problem sets appear after each major idea. These straightforward exercises help you solidify the main points as you read. Be sure you can do these exercises before going on. The end-of-chapter exercises begin with multiple-choice "Check Your Skills" exercises (all answers are in the back of the book). Use them to check your grasp of the basics. The regular "Chapter Exercises" help you combine all the ideas of a chapter. Finally, the three part review chapters (Chapters 7, 17, and 22) look back over major blocks of learning. They begin with a section of multiple-choice, basic calculations and short-answer questions to help you "Test Yourself" on this set of chapters. The supplemental review exercises in the review chapters often require combining ideas learned in several chapters. For each type of exercise, you are given different information regarding the specific statistical ideas and skills the problem will require, so each type of exercise requires a different level of understanding.

The part review chapters (and the individual chapters in Part IV) include point-by-point lists of specific things you should be able to do. Go through the lists, and be sure you can say "I can do that" to each item. Then try some of the "Test Yourself" problems.

The key to learning is persistence. The main ideas of statistics, like the main ideas of any important subject, took a long time to discover and take some time to master. The gain will be worth the pain.



ORGANIZING A STATISTICAL PROBLEM: A Four-Step Process

STATE: What is the practical question, in the context of the real-world setting?

PLAN: What specific statistical operations does this problem call for?

SOLVE: Make the graphs and carry out the calculations needed for this problem.

CONCLUDE: Give your practical conclusion in the setting of the real-world problem.



CONFIDENCE INTERVALS: The Four-Step Process

STATE: What is the practical question that requires estimating a parameter?

PLAN: Identify the parameter, choose a level of confidence, and select the type of confidence interval that fits your situation.

SOLVE: Carry out the work in two phases:

1. Check the conditions for the interval you plan to use.
2. Calculate the confidence interval.

CONCLUDE: Return to the practical question to describe your results in this setting.



TESTS OF SIGNIFICANCE: A Four-Step Process

STATE: What is the practical question that requires a statistical test?

PLAN: Identify the parameter, state null and alternative hypotheses, and choose the type of test that fits your situation.

SOLVE: Carry out the test in three phases:

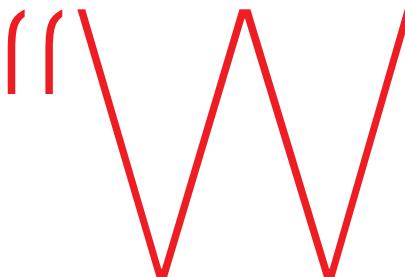
1. Check the conditions for the test you plan to use.
2. Calculate the test statistic.
3. Find the P-value.

CONCLUDE: Return to the practical question to describe your results in this setting.



Exploring Data

Part I



such as histograms and scatterplots and numerical measures such as means and correlations. At least as important as the tools are principles that organize our thinking as we examine data. The seven chapters in Part I present the principles and tools of statistical data analysis. They equip you with skills that are immediately useful whenever you deal with numbers.

These chapters reflect the strong emphasis on exploring data that characterizes modern statistics. Sometimes we hope to draw conclusions that apply to a setting that goes beyond the data in hand. This is *statistical inference*, the topic of much of the rest of the book. Data analysis is essential if we are to trust the results of inference, but data analysis isn't just preparation for inference. Roughly speaking, you can always do data analysis but inference requires rather special conditions.

One of the organizing principles of data analysis is to first look at one thing at a time and then at relationships. Our presentation follows this principle. In Chapters 1, 2, and 3 you will study *variables and their distributions*. Chapters 4, 5, and 6 concern *relationships among variables*. Chapter 7 reviews this part of the text.

“What do the data say?” is the first question we ask in any statistical study. *Data analysis* answers this question by open-ended exploration of the data. The tools of data analysis are graphs

EXPLORING DATA: Variables and Distributions

CHAPTER 1 Picturing Distributions with Graphs

CHAPTER 2 Describing Distributions with Numbers

CHAPTER 3 The Normal Distributions

EXPLORING DATA: Relationships

CHAPTER 4 Scatterplots and Correlation

CHAPTER 5 Regression

CHAPTER 6 Two-Way Tables*

CHAPTER 7 Exploring Data: Part I Review



U.S. Cens

US CENSUS BUREAU

Picturing Distributions with Graphs

Chapter 1

Statistics is the science of data. The volume of data available to us is overwhelming. For example, the U.S. Census Bureau's American Community Survey collects data from about 3,000,000 housing units each year. Astronomers work with data on tens of millions of galaxies. The checkout scanners at Walmart's 8000 stores in 15 countries record hundreds of millions of transactions every week, all saved to inform both Walmart and its suppliers. The first step in dealing with such a flood of data is to organize our thinking about data. Fortunately, we can do this without looking at millions of data points.

INDIVIDUALS AND VARIABLES

Any set of data contains information about some group of *individuals*. The information is organized in *variables*.

INDIVIDUALS AND VARIABLES

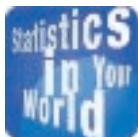
Individuals are the objects described by a set of data. Individuals may be people, but they may also be animals or things.

A **variable** is any characteristic of an individual. A variable can take different values for different individuals.

IN THIS CHAPTER WE COVER...

- Individuals and variables
- Categorical variables: pie charts and bar graphs
- Quantitative variables: histograms
- Interpreting histograms
- Quantitative variables: stemplots
- Time plots

Courtesy U.S. Census



What's that number?

You might think that numbers, unlike words, are universal. Think again. A “billion” in the United States means 1,000,000,000 (nine zeros). In Europe, a “billion” is 1,000,000,000,000 (twelve zeros). OK, those are words that describe numbers. But those commas in big numbers are periods in many other languages. This is so confusing that international standards call for spaces instead, so that an American billion is written 1 000 000 000. And the decimal point of the English-speaking world is the decimal comma in many other languages, so that 3.1416 in the United States becomes 3,1416 in Europe. So what is the number 10,642.389? Depends on where you are.

any set of data is accompanied by background information that helps us understand the data. When you plan a statistical study or explore data from someone else’s work, ask yourself the following questions:

- 1. Who?** What **individuals** do the data describe? How **many** individuals appear in the data?
- 2. What?** How many **variables** do the data contain? What are the **exact definitions** of these variables? In what **unit of measurement** is each variable recorded? Weights, for example, might be recorded in pounds, in thousands of pounds, or in kilograms.
- 3. Where?** Student GPAs and SAT scores (or lack of them) will vary from college to college depending on many variables, including admissions “selectivity” for the college.
- 4. When?** Students change from year to year, as do prices, salaries, etc.
- 5. Why?** What **purpose** do the data have? Do we hope to answer some specific questions? Do we want answers for just these individuals or for some larger group that these individuals are supposed to represent? Are the individuals and variables suitable for the intended purpose?

Some variables, like a person’s sex or college major, simply place individuals into categories. Others, like height and grade point average, take numerical values for which we can do arithmetic. It makes sense to give an average income for a company’s employees, but it does not make sense to give an “average” sex. We can, however, count the numbers of female and male employees and do arithmetic with these counts.

CATEGORICAL AND QUANTITATIVE VARIABLES

A **categorical variable** places an individual into one of several groups or categories.

A **quantitative variable** takes numerical values for which arithmetic operations such as adding and averaging make sense. The values of a quantitative variable are usually recorded in a **unit of measurement** such as seconds or kilograms.

EXAMPLE 1.1 The American Community Survey

At the U.S. Census Bureau Web site, you can view the detailed data collected by the American Community Survey, though of course the identities of people and housing units are protected. If you choose the file of data on people, the *individuals* are the people living in the housing units contacted by the survey. Over 100 variables are recorded for each individual. Figure 1.1 displays a very small part of the data.

Each row records data on one individual. Each column contains the values of one *variable* for all the individuals. Translated from the U.S. Census Bureau’s abbreviations, the variables are

	A	B	C	D	E	F	G
1	SERIALNO	PWGTP	AGEP	JWMNP	SCHL	SEX	WAGP
2		283	187		6	1	24000
3		283	158	66		2	0
4		323	176	54	10	12	11900
5		346	339	37	10	11	1
6		346	91	27	10	10	30000
7		370	234	53	10	13	1
8		370	181	46	15	10	2
9		370	155	18		9	2
10		487	233	26		14	2
11		487	146	23		12	2
12		511	236	53		9	2
13		511	131	53		11	1
14		515	213	38		11	2
15		515	194	40		9	1
16		515	221	18	20	9	1
17		515	193	11		3	1

FIGURE 1.1

A spreadsheet displaying data from the American Community Survey, for Example 1.1.

Each row in the spreadsheet contains data on one individual.

SERIALNO	An identifying number for the household.
PWGTP	Weight in pounds.
AGEP	Age in years.
JWMNP	Travel time to work in minutes.
SCHL	Highest level of education. The numbers designate categories, <i>not</i> specific grades. For example, 9 = high school graduate, 10 = some college but no degree, and 13 = bachelor's degree.
SEX	Sex, designated by 1 = male and 2 = female.
WAGP	Wage and salary income last year, in dollars.

Look at the highlighted row in Figure 1.1. This individual is a 53-year-old man who weighs 234 pounds, travels 10 minutes to work, has a bachelor's degree, and earned \$83,000 last year.

In addition to the household serial number, there are six variables. Education and sex are categorical variables. The values for education and sex are stored as numbers, but these numbers are just labels for the categories and have no units of measurement. The other four variables are quantitative. Their values do have units. These variables are weight in pounds, age in years, travel time in minutes, and income in dollars.

The *purpose* of the American Community Survey is to collect data that represent the entire nation in order to guide government policy and business decisions. To do this, the households contacted are chosen at random from all households in the country. We will see in Chapter 8 why choosing at random is a good idea. ■

Most data tables follow this format—each row is an individual, and each column is a variable. The data set in Figure 1.1 appears in a **spreadsheet** program that has rows and columns ready for your use. Spreadsheets are commonly used to enter and transmit data and to do simple calculations.

spreadsheet



APPLY YOUR KNOWLEDGE

1.1 Fuel economy. Here is a small part of a data set that describes the fuel economy (in miles per gallon) of model year 2010 motor vehicles:

Make and model	Vehicle type	Transmission type	Number of cylinders	City mpg	Highway mpg	Carbon footprint (tons)
:						
Aston Martin Vantage	Two-seater	Manual	8	12	19	13.1
Honda Civic	Subcompact	Automatic	4	25	36	6.3
Toyota Prius	Midsized	Automatic	4	51	48	3.7
Chevrolet Impala	Large	Automatic	6	18	29	8.3
:						

The *carbon footprint* measures a vehicle's impact on climate change in tons of carbon dioxide emitted annually.

- (a) What are the individuals in this data set?
- (b) For each individual, what variables are given? Which of these variables are categorical and which are quantitative?

1.2 Students and exercise. You are preparing to study the exercise habits of college students. Describe two categorical variables and two quantitative variables that you might measure for each student. Give the units of measurement for the quantitative variables.

CATEGORICAL VARIABLES: Pie Charts and Bar Graphs

exploratory data analysis

Statistical tools and ideas help us examine data in order to describe their main features. This examination is called **exploratory data analysis**. Like an explorer crossing unknown lands, we want first to simply describe what we see. Here are two principles that help us organize our exploration of a set of data.

EXPLORING DATA

1. Begin by examining each variable by itself. Then move on to study the relationships among the variables.
2. Begin with a graph or graphs. Then add numerical summaries of specific aspects of the data.

We will follow these principles in organizing our learning. Chapters 1 to 3 present methods for describing a single variable. We study relationships among

several variables in Chapters 4 to 6. In each case, we begin with graphical displays, then add numerical summaries for more complete description.

The proper choice of graph depends on the nature of the variable. To examine a single variable, we usually want to display its *distribution*.

DISTRIBUTION OF A VARIABLE

The **distribution** of a variable tells us what values it takes and how often it takes these values.

The values of a categorical variable are labels for the categories. The **distribution of a categorical variable** lists the categories and gives either the count or the percent of individuals who fall in each category.

EXAMPLE 1.2 Which major?

About 1.8 million students enroll in colleges and universities each year. What do they plan to study? Here are data from 2008 on the percents of first-year students who plan to major in several discipline areas:¹



Field of study	Percent of students
Arts and humanities	13.5
Biological sciences	9.3
Business	16.8
Education	8.2
Engineering	9.4
Physical sciences	3.2
Professional	13.8
Social science	11.5
Technical	1.0
Other majors and undeclared	13.2
Total	99.9

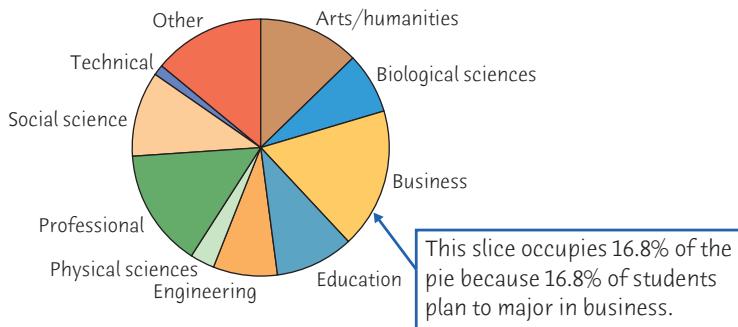
It's a good idea to check data for consistency. The percents should add to 100%. In fact, they add to 99.9%. What happened? Each percent is rounded to the nearest tenth. The exact percents would add to 100, but the rounded percents only come close. This is **roundoff error**. Roundoff errors don't point to mistakes in our work, just to the effect of rounding off results. ■

roundoff error

Columns of numbers take time to read. You can use a pie chart or a bar graph to display the distribution of a categorical variable more vividly. Figures 1.2 and 1.3 illustrate these displays for the distribution of intended college majors.

FIGURE 1.2

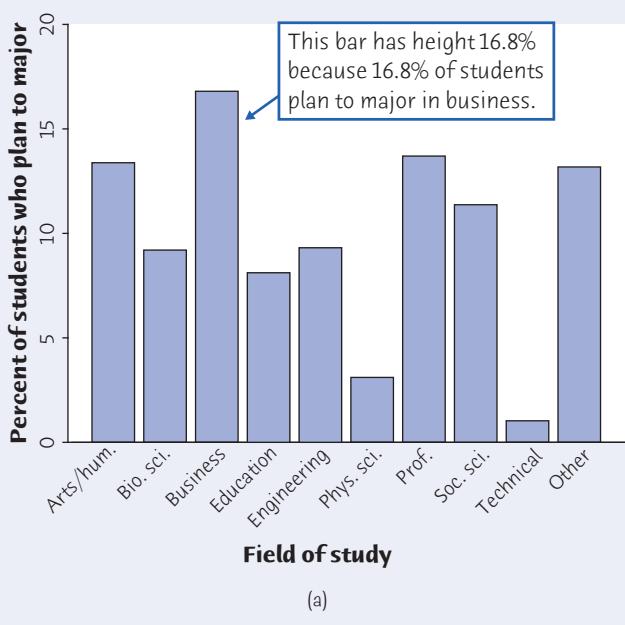
You can use a pie chart to display the distribution of a categorical variable. Here is a pie chart of the distribution of intended majors of students entering college.

**pie chart**

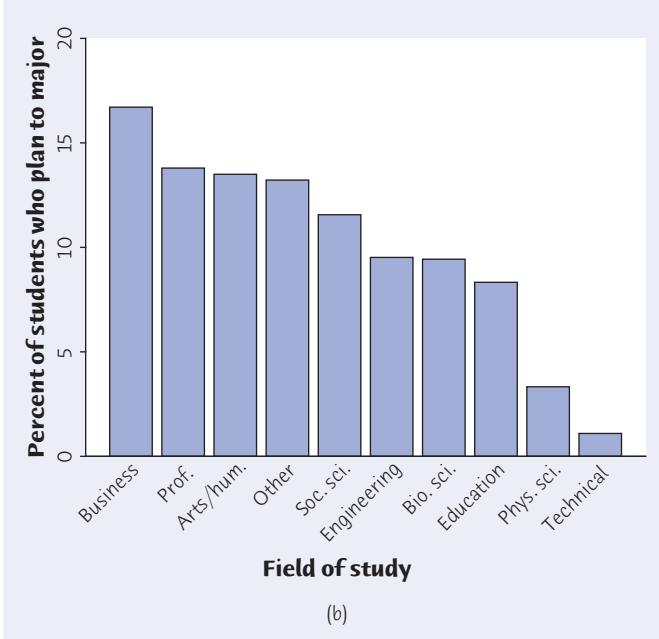
Pie charts show the distribution of a categorical variable as a “pie” whose slices are sized by the counts or percents for the categories. Pie charts are awkward to make by hand, but software will do the job for you. A *pie chart must include all the categories that make up a whole*. Use a pie chart only when you want to emphasize each category’s relation to the whole. We need the “Other majors” category in Example 1.2 to complete the whole (all intended majors) and allow us to make the pie chart in Figure 1.2.

**bar graph**

Bar graphs represent each category as a bar. The bar heights show the category counts or percents. Bar graphs are easier to make than pie charts and also easier to read. Figure 1.3 displays two bar graphs of the data on intended majors.



(a)



(b)

FIGURE 1.3

Bar graphs of the distribution of intended majors of students entering college. In (a), the bars follow the alphabetical order of fields of study. In (b), the same bars appear in order of height.

The first orders the bars alphabetically by field of study (with “Other” at the end). It is often better to arrange the bars in order of height, as in Figure 1.3(b). This helps us immediately see which majors appear most often.

Bar graphs are more flexible than pie charts. Both graphs can display the distribution of a categorical variable, but a bar graph can also compare any set of quantities that are measured in the same units.

EXAMPLE 1.3 Cell phones have biggest impact!

The rating service Arbitron asked Americans over 12 years old who used several high-tech platforms/devices to answer the question “How much of an impact on your life has (platform/device) had?” with 5 = Big impact and 1 = No impact at all. Here are the percents in decreasing order who said, “Big impact.” Only those platforms/devices for which 20% or more said, “Big impact” are reported.²



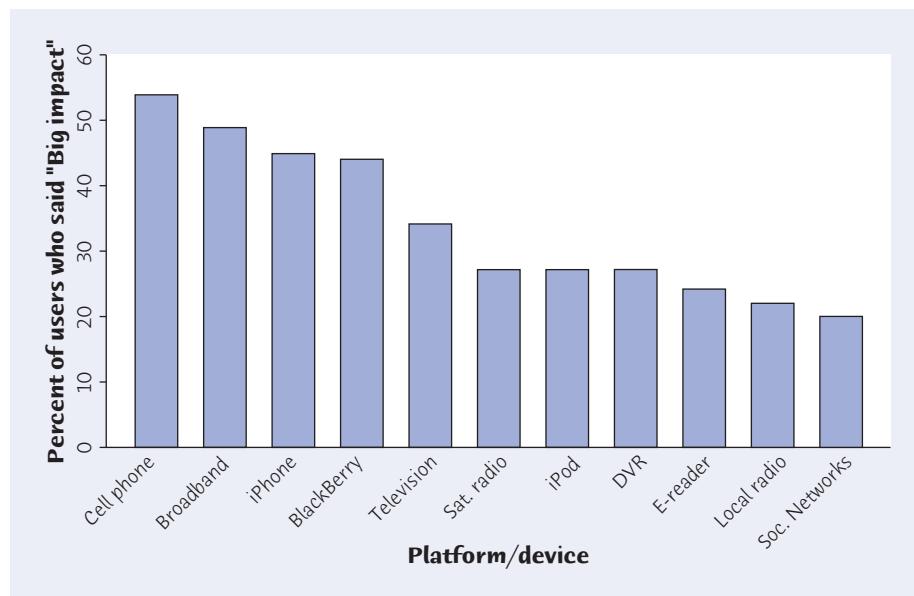
Platform/device	Percent of users who said “Big impact”
Cell phone	54
Broadband Internet	49
iPhone	45
BlackBerry	44
Television	34
Satellite radio	27
iPod	27
Digital video recorder	27
E-reader	24
Local AM/FM radio	22
Social-networking Web site	20

We can’t make a pie chart to display these data. Each percent in the table refers to a different platform/device, not to parts of a single whole. Figure 1.4 is a bar graph comparing the 11 platforms/devices. We have again arranged the bars in order of height. ■

Bar graphs and pie charts are mainly tools for presenting data: they help your audience grasp data quickly. Since it is easy to understand data on a single categorical variable without a graph, bar graphs and pie charts are of limited use for data analysis. We will move on to quantitative variables, where graphs are essential tools.

FIGURE 1.4

You can use a bar graph to compare quantities that are not part of a whole. This bar graph compares the percents of users who say various platforms/devices had a “Big impact” on their lives, for Example 1.3.



APPLY YOUR KNOWLEDGE

1.3 Do you listen to talk radio? The rating service Arbitron places U.S. radio stations into more than 50 categories that describe the kind of programs they broadcast. Which formats attract the largest audiences? Here are Arbitron’s measurements of the share of the listening audience (aged 12 and over) at a given time for the most popular formats:³ 

Format	Audience share
News/Talk/Information	12.6%
Country	12.5%
Adult Contemporary	8.2%
Pop Contemporary Hit	5.9%
Classic Rock	4.7%
Classic Hits	3.9%
Rhythmic Contemporary Hit	3.7%
Urban Adult Contemporary	3.6%
Hot Adult Contemporary	3.5%
Urban Contemporary	3.3%
Mexican Regional	2.9%
All Sports	2.5%

- (a) What is the sum of the audience shares for these formats? What percent of the radio audience listens to stations with other formats?

- (b) Make a bar graph to display these data. Be sure to include an “Other format” category.
- (c) Would it be correct to display these data in a pie chart? Why?

1.4 How do students pay for college? The Higher Education Research Institute’s Freshman Survey includes over 200,000 first-time full-time freshmen who entered college in 2009.⁴ The survey reports the following data on the sources students use to pay for college expenses.  COLLEGEPAY

Source for college expenses	Students
Family resources	78.2%
Student resources	62.8%
Aid—not to be repaid	70.0%
Aid—to be repaid	53.4%
Other	6.5%

- (a) Explain why it is *not* correct to use a pie chart to display these data.
- (b) Make a bar graph of the data. Notice that because the data contrast groups such as family and student resources it is better to keep these bars next to each other rather than to arrange the bars in order of height.

1.5 Never on Sunday? Births are not, as you might think, evenly distributed across the days of the week. Here are the average numbers of babies born on each day of the week in 2008:⁵  BIRTHS

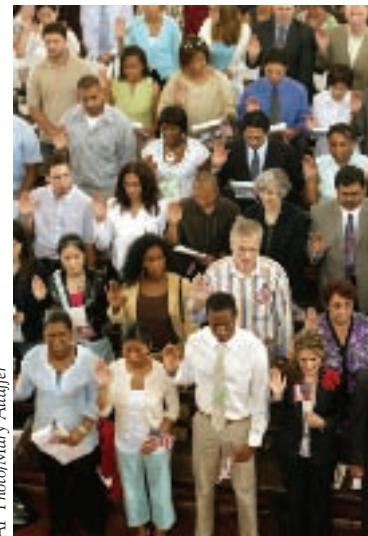
Day	Births
Sunday	7,534
Monday	12,371
Tuesday	13,415
Wednesday	13,171
Thursday	13,147
Friday	12,919
Saturday	8,617

Present these data in a well-labeled bar graph. Would it also be correct to make a pie chart? Suggest some possible reasons why there are fewer births on weekends.

QUANTITATIVE VARIABLES: Histograms

Quantitative variables often take many values. The distribution tells us what values the variable takes and how often it takes these values. A graph of the distribution is clearer if nearby values are grouped together. The most common graph of the distribution of one quantitative variable is a **histogram**.

histogram



AP Photo/Mary Altaffer



EXAMPLE 1.4 Making a histogram

What percent of your home state's residents were born outside the United States? The country as a whole has 12.5% foreign-born residents, but the states vary from 1.2% in West Virginia to 27.2% in California. Table 1.1 presents the data for all 50 states and the District of Columbia.⁶ The *individuals* in this data set are the states. The *variable* is the percent of a state's residents who are foreign-born. It's much easier to see how your state compares with other states from a graph than from the table. To make a histogram of the distribution of this variable, proceed as follows:

Step 1. Choose the classes. Divide the range of the data into classes of equal width. The data in Table 1.1 range from 1.2 to 27.2, so we decide to use these classes:

percent foreign-born between 0.1 and 5.0
percent foreign-born between 5.1 and 10.0
⋮
percent foreign-born between 25.1 and 30.0

It is equally correct to use classes 0.0 to 4.9, 5.0 to 9.9, and so on. Just be sure to specify the classes precisely so that each individual falls into exactly one class. Pennsylvania, with 5.1% foreign-born, falls into the second class, but a state with 5.0% would fall into the first.

TABLE 1.1 Percent of state population born outside the United States

STATE	PERCENT	STATE	PERCENT	STATE	PERCENT
Alabama	2.8	Louisiana	2.9	Ohio	3.6
Alaska	7.0	Maine	3.2	Oklahoma	4.9
Arizona	15.1	Maryland	12.2	Oregon	9.7
Arkansas	3.8	Massachusetts	14.1	Pennsylvania	5.1
California	27.2	Michigan	5.9	Rhode Island	12.6
Colorado	10.3	Minnesota	6.6	South Carolina	4.1
Connecticut	12.9	Mississippi	1.8	South Dakota	2.2
Delaware	8.1	Missouri	3.3	Tennessee	3.9
Florida	18.9	Montana	1.9	Texas	15.9
Georgia	9.2	Nebraska	5.6	Utah	8.3
Hawaii	16.3	Nevada	19.1	Vermont	3.9
Idaho	5.6	New Hampshire	5.4	Virginia	10.1
Illinois	13.8	New Jersey	20.1	Washington	12.4
Indiana	4.2	New Mexico	10.1	West Virginia	1.2
Iowa	3.8	New York	21.6	Wisconsin	4.4
Kansas	6.3	North Carolina	6.9	Wyoming	2.7
Kentucky	2.7	North Dakota	2.1	District of Columbia	12.7

Step 2. Count the individuals in each class. Here are the counts:

Class	Count
0.1 to 5.0	20
5.1 to 10.0	13
10.1 to 15.0	10
15.1 to 20.0	5
20.1 to 25.0	2
25.1 to 30.0	1

Check that the counts add to 51, the number of individuals in the data set (the 50 states and the District of Columbia).

Step 3. Draw the histogram. Mark the scale for the variable whose distribution you are displaying on the horizontal axis. That's the percent of a state's residents who are foreign-born. The scale runs from 0 to 30 because that is the span of the classes we chose. The vertical axis contains the scale of counts. Each bar represents a class. The base of the bar covers the class, and the bar height is the class count. Draw the bars with no horizontal space between them unless a class is empty, so that its bar has height zero. Figure 1.5 is our histogram. ■

Although histograms resemble bar graphs, their details and uses are different. A histogram displays the distribution of a quantitative variable. The horizontal axis of a histogram is marked in the units of measurement for

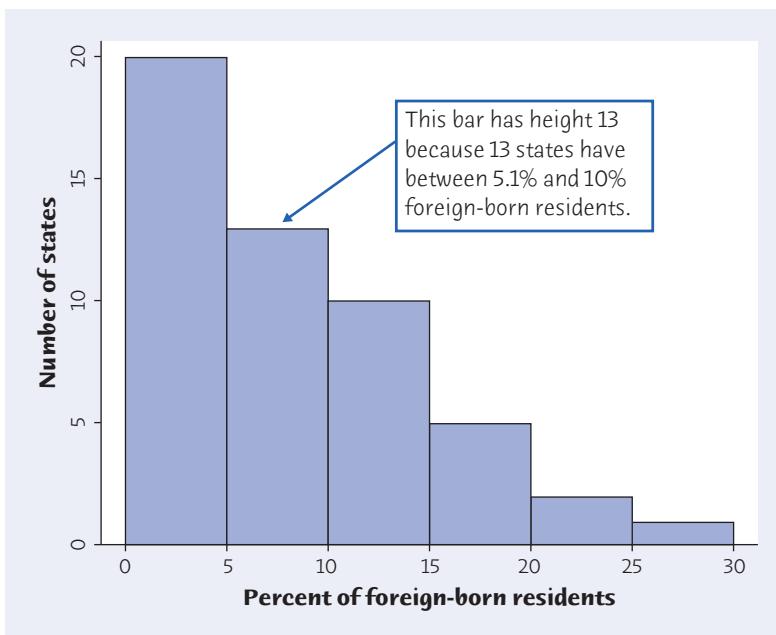


FIGURE 1.5

Histogram of the distribution of the percent of foreign-born residents in the 50 states and the District of Columbia, for Example 1.4.

the variable. A bar graph compares the sizes of different quantities. The horizontal axis of a bar graph need not have any measurement scale but simply identifies the quantities being compared. These may be the values of a categorical variable, but they may also be unrelated, like the high-tech devices in Example 1.3. Draw bar graphs with blank space between the bars to separate the quantities being compared. Draw histograms with no space, to indicate that all values of the variable are covered.

Our eyes respond to the *area* of the bars in a histogram.⁷ Because the classes are all the same width, area is determined by height and all classes are fairly represented. There is no one right choice of the classes in a histogram. Too few classes will give a “skyscraper” graph, with all values in a few classes with tall bars. Too many will produce a “pancake” graph, with most classes having one or no observations. Neither choice will give a good picture of the shape of the distribution. You must use your judgment in choosing classes to display the shape. Statistics software will choose the classes for you. The software’s choice is usually a good one, but you can change it if you want. The histogram function in the *One-Variable Statistical Calculator* applet on the text CD and Web site allows you to change the number of classes by dragging with the mouse, so that it is easy to see how the choice of classes affects the histogram.



APPLY YOUR KNOWLEDGE

1.6 Older Americans. Where are older Americans more likely to live? Table 1.2 gives the percent of residents aged 65 years and over in each of the 50 states and the District of Columbia.⁸ Make a histogram of the percent using classes of width 1% starting at 7%. That is, the first bar covers 7.0% to 7.9%, the second covers 8.0% to 8.9%, and so on. (Make this histogram by hand even if you have software, to be sure you understand the process. You may then want to compare your histogram with your software’s choice.) 

1.7 Choosing classes in a histogram. The data set menu that accompanies the *One-Variable Statistical Calculator* applet includes the data on foreign-born residents in the states from Table 1.1. Choose these data, then click on the “Histogram” tab to see a histogram.

- How many classes does the applet choose to use? (You can click on the graph outside the bars to get a count of classes.)
- Click on the graph and drag to the left. What is the smallest number of classes you can get? What are the lower and upper bounds of each class? (Click on the bar to find out.) Make a rough sketch of this histogram.
- Click and drag to the right. What is the greatest number of classes you can get? How many observations does the largest class have?
- You see that the choice of classes changes the appearance of a histogram. Drag back and forth until you get the histogram that you think best displays the distribution. How many classes did you use?

TABLE 1.2 Percent of residents aged 65 and over, 2009

STATE	PERCENT	STATE	PERCENT	STATE	PERCENT
Alabama	13.5	Louisiana	12.1	Ohio	13.6
Alaska	7.0	Maine	15.0	Oklahoma	13.3
Arizona	12.9	Maryland	11.8	Oregon	13.2
Arkansas	14.0	Massachusetts	13.4	Pennsylvania	15.3
California	10.9	Michigan	12.9	Rhode Island	14.0
Colorado	10.3	Minnesota	12.4	South Carolina	13.1
Connecticut	13.6	Mississippi	12.5	South Dakota	14.3
Delaware	13.8	Missouri	13.5	Tennessee	12.9
Florida	16.9	Montana	14.1	Texas	10.1
Georgia	10.0	Nebraska	13.3	Utah	8.8
Hawaii	14.1	Nevada	11.3	Vermont	13.8
Idaho	11.7	New Hampshire	12.8	Virginia	11.8
Illinois	12.1	New Jersey	13.2	Washington	11.8
Indiana	12.6	New Mexico	12.6	West Virginia	15.5
Iowa	14.7	New York	13.2	Wisconsin	13.2
Kansas	13.0	North Carolina	12.4	Wyoming	12.1
Kentucky	12.9	North Dakota	14.6	District of Columbia	11.8

INTERPRETING HISTOGRAMS

Making a statistical graph is not an end in itself. *The purpose of graphs is to help us understand the data.* After you make a graph, always ask, “What do I see?” Once you have displayed a distribution, you can see its important features as follows.

EXAMINING A HISTOGRAM

In any graph of data, look for the **overall pattern** and for striking **deviations** from that pattern.

You can describe the overall pattern of a histogram by its **shape, center, and spread**.

An important kind of deviation is an **outlier**, an individual value that falls outside the overall pattern.

One way to describe the center of a distribution is by its *midpoint*, the value with roughly half the observations taking smaller values and half taking larger values. For now, we will describe the spread of a distribution by giving the *smallest and largest values*. We will learn better ways to describe center and spread in Chapter 2. The overall shape of a distribution can often be described in terms of symmetry or skewness, defined as follows.

SYMMETRIC AND SKEWED DISTRIBUTIONS

A distribution is **symmetric** if the right and left sides of the histogram are approximately mirror images of each other.

A distribution is **skewed to the right** if the right side of the histogram (containing the half of the observations with larger values) extends much farther out than the left side. It is **skewed to the left** if the left side of the histogram extends much farther out than the right side.

EXAMPLE 1.5 Describing a distribution

Look again at the histogram in Figure 1.5. **Shape:** The distribution has a *single peak* at the left, which represents states in which between 0% and 5% of residents are foreign-born. The distribution is *skewed to the right*. A majority of states have no more than 10% foreign-born residents, but several states have much higher percents, so that the graph extends quite far to the right of its peak. **Center:** Arranging the observations from Table 1.1 in order of size shows that 6.3% (Kansas) is the midpoint of the distribution. There are 25 states with smaller percents foreign-born and 25 with larger. **Spread:** The spread is from 1.2% to 27.2%.

Outliers: Figure 1.5 shows no observations outside the overall single-peaked, right-skewed pattern of the distribution. Figure 1.6 is another histogram of the same distribution, with classes half as wide. Now California, at 27.2%, stands a bit apart to the right of the rest of the distribution. Is California an outlier or just the largest observation in a strongly skewed distribution? Unfortunately, there is no rule. Let's agree to call attention to only strong outliers that suggest something special about an observation—or an error such as typing 10.1 as 101. California is certainly not a strong outlier. ■

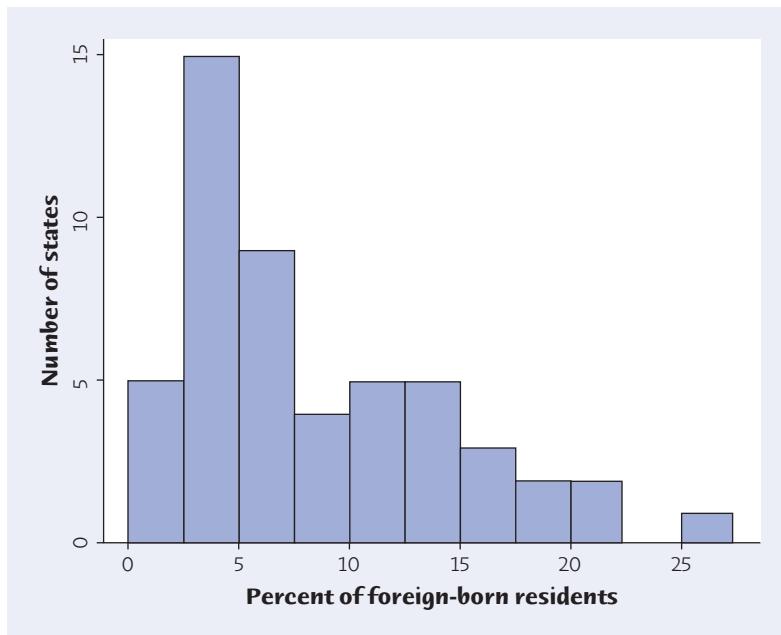


FIGURE 1.6

Another histogram of the distribution of the percent of foreign-born residents, with classes half as wide as in Figure 1.5. Histograms with more classes show more detail but may have a less clear pattern.

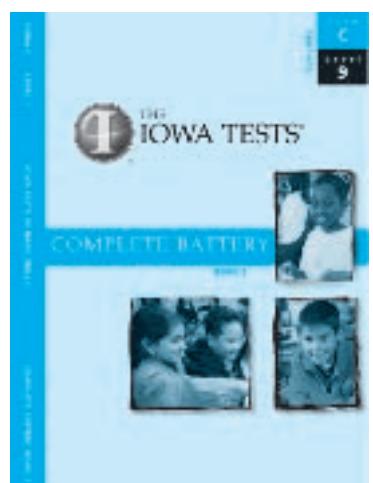
Figures 1.5 and 1.6 remind us that interpreting graphs calls for judgment. We also see that the choice of classes in a histogram can influence the appearance of a distribution. Because of this, and to avoid worrying about minor details, concentrate on the main features of a distribution. Look for major peaks, not for minor ups and downs, in the bars of the histogram. (For example, don't conclude that Figure 1.6 shows a second peak between 10% and 15%.) Look for clear outliers, not just for the smallest and largest observations. Look for rough symmetry or clear skewness.



Here are more examples of describing the overall pattern of a histogram.

EXAMPLE 1.6 Iowa Test scores

Figure 1.7 displays the scores of all 947 seventh-grade students in the public schools of Gary, Indiana, on the vocabulary part of the Iowa Test of Basic Skills.⁹ The distribution is *single-peaked* and *symmetric*. In mathematics, the two sides of symmetric patterns are exact mirror images. Real data are almost never exactly symmetric. We are content to describe Figure 1.7 as symmetric. The center (half above, half below) is close to 7. This is seventh-grade reading level. The scores range from 2.0 (second-grade level) to 12.1 (twelfth-grade level).



©2011 Houghton Mifflin Company. All rights reserved.

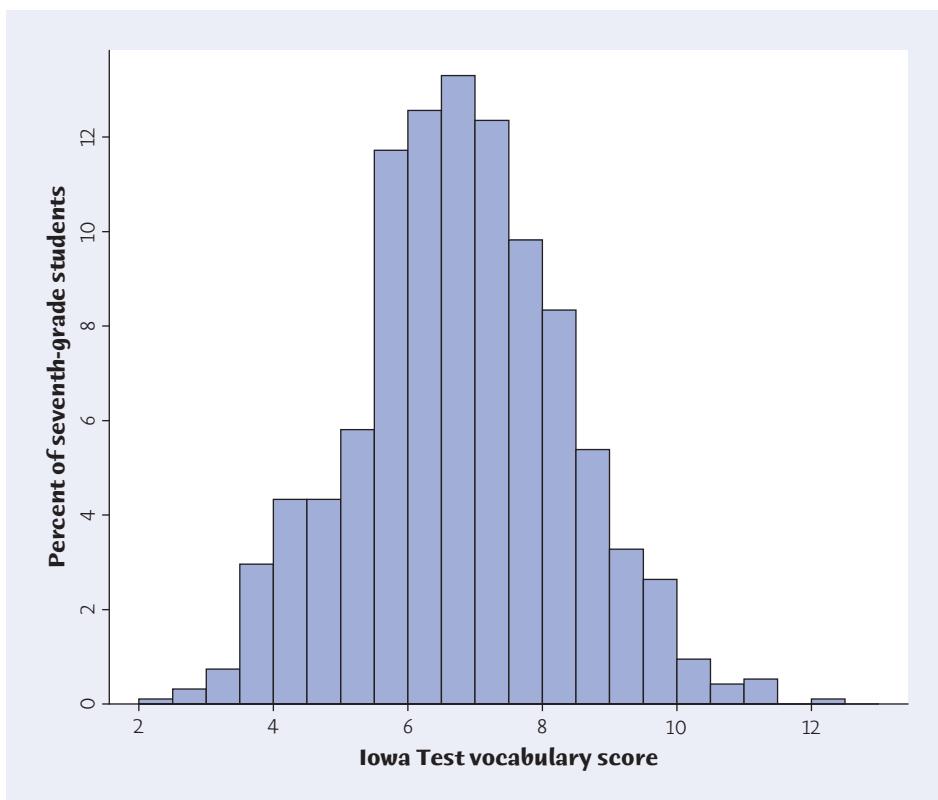


FIGURE 1.7

Histogram of the Iowa Test vocabulary scores of all seventh-grade students in Gary, Indiana, for Example 1.6. This distribution is single-peaked and symmetric.

Notice that the vertical scale in Figure 1.7 is not the *count* of students but the *percent* of students in each histogram class. A histogram of percents rather than counts is convenient when we want to compare several distributions. To compare Gary with Los Angeles, a much bigger city, we would use percents so that both histograms have the same vertical scale. ■



SATTAKERS

EXAMPLE 1.7 Who takes the SAT?

Depending on where you went to high school, the answer to this question may be “almost everybody” or “almost nobody.” Figure 1.8 is a histogram of the percent of high school graduates in each state who took the SAT Reasoning test.¹⁰

The histogram shows two peaks, a high peak at the left and a lower but broader peak centered in the 60% to 80% class. The presence of more than one peak suggests that a distribution mixes several kinds of individuals. That is the case here. There are two major tests of readiness for college, the ACT and the SAT. Most states have a strong preference for one or the other. In some states, many students take the ACT exam and few take the SAT—these states form the peak on the left. In other states, many students take the SAT and few choose the ACT—these states form the broader peak at the right.

Giving the center and spread of this distribution is not very useful. The midpoint falls in the 20% to 40% class, between the two peaks. The story told by the histogram is in the two peaks corresponding to ACT states and SAT states. ■

The overall shape of a distribution is important information about a variable. Some variables have distributions with predictable shapes. Many biological measurements on specimens from the same species and sex—lengths of bird bills,

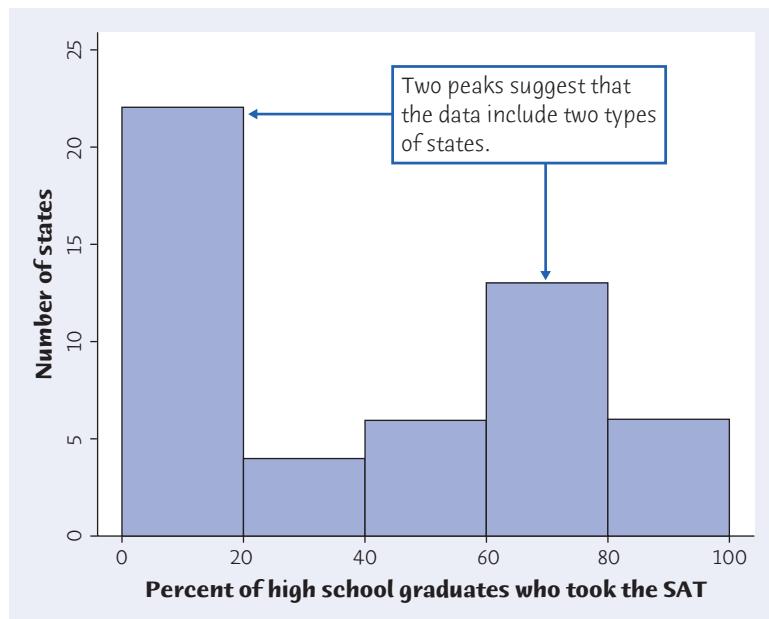


FIGURE 1.8

Histogram of the percent of high school graduates in each state who took the SAT Reasoning test, for Example 1.7. The graph shows two groups of states: ACT states (where few students take the SAT) at the left and SAT states at the right.

heights of young women—have symmetric distributions. On the other hand, data on people's incomes are usually strongly skewed to the right. There are many moderate incomes, some large incomes, and a few enormous incomes. Many distributions have irregular shapes that are neither symmetric nor skewed. Some data show other patterns, such as the two peaks in Figure 1.8. Use your eyes, describe the pattern you see, and then try to explain the pattern.

APPLY YOUR KNOWLEDGE

- 1.8 Older Americans.** In Exercise 1.6, you made a histogram of the percent of residents aged 65 years and over in each of the 50 states and the District of Columbia, given in Table 1.2. Describe the shape of the distribution. Is it closer to symmetric or skewed? About where is the center (midpoint) of the data? What is the spread in terms of the smallest and largest values? Are there any states with unusually large or small percent of residents over 65?  OVER65
- 1.9 Unmarried women.** Figure 1.9 shows the distribution of the state percents of women aged 15 and over who have never been married.  UNMARRIED
- The main body of the distribution is slightly skewed to the right. There is one clear outlier, the District of Columbia. Why is it not surprising that the percent of never-married women is higher in DC than in the 50 states?
 - The midpoint of the distribution is the 26th state in order of percent of never-married women. In which class does the midpoint fall? About what is the spread (smallest to largest) of the distribution?

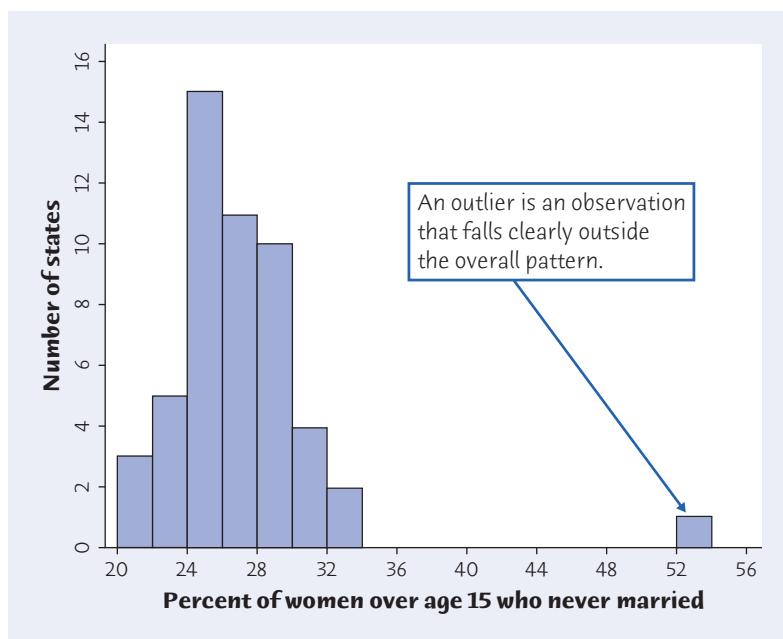


FIGURE 1.9

Histogram of the state percents of women aged 15 and over who have never been married, for Exercise 1.9.

QUANTITATIVE VARIABLES: Stemplots

Histograms are not the only graphical display of distributions. For small data sets, a *stemplot* is quicker to make and presents more detailed information.

STEMPLOT

To make a *stemplot*:

1. Separate each observation into a **stem**, consisting of all but the final (rightmost) digit, and a **leaf**, the final digit. Stems may have as many digits as needed, but each leaf contains only a single digit.
2. Write the stems in a vertical column with the smallest at the top, and draw a vertical line at the right of this column. Be sure to include all the stems needed to span the data, even when some stems will have no leaves.
3. Write each leaf in the row to the right of its stem, in increasing order out from the stem.

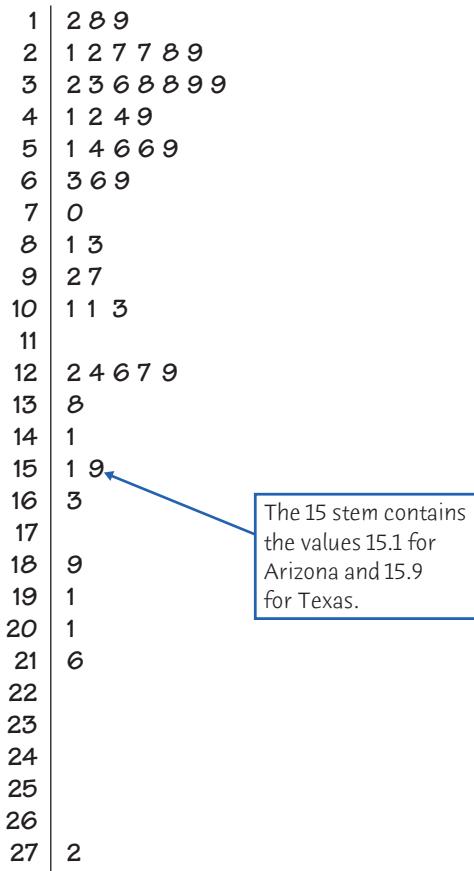
The vital few

 Skewed distributions can show us where to concentrate our efforts. Ten percent of the cars on the road account for half of all carbon dioxide emissions. A histogram of CO₂ emissions would show many cars with small or moderate values and a few with very high values. Cleaning up or replacing these cars would reduce pollution at a cost much lower than that of programs aimed at all cars. Statisticians who work at improving quality in industry make a principle of this: distinguish “the vital few” from “the trivial many.”

EXAMPLE 1.8 Making a stemplot

Table 1.1 presents the percents of state residents who were born outside the United States. To make a stemplot of these data, take the whole-number part of the percent as the stem and the final digit (tenths) as the leaf. Write stems from 1 for Mississippi, Montana, and West Virginia to 27 for California. Now add leaves. Arizona, 15.1%, has leaf 1 on the 15 stem. Texas, at 15.9%, places leaf 9 on the same stem. These are the only observations on this stem. Arrange the leaves in order, so that 15 | 19 is one row in the stemplot. Figure 1.10 is the complete stemplot for the data in Table 1.1. ■

A stemplot looks like a histogram turned on end. Compare the stemplot in Figure 1.10 with the histograms of the same data in Figures 1.5 and 1.6. The stemplot is like a histogram with many classes. You can choose the classes in a histogram. The classes (the stems) of a stemplot are given to you. All three graphs show a distribution that has one peak and is right-skewed. Figures 1.6 and 1.10 have enough classes to show that California (27.2%) stands slightly apart from the long right tail of the skewed distribution. Histograms are more flexible than stemplots because you can choose the classes. But the stemplot, unlike the  histogram, preserves the actual value of each observation. *Stemplots do not work well for large data sets, where each stem must hold a large number of leaves.* Don’t try to make a stemplot of a large data set, such as the 947 Iowa Test scores in Figure 1.7.

**FIGURE 1.10**

Stemplot of the percents of foreign-born residents in the states, for Example 1.8. Each stem is a percent and leaves are tenths of a percent.

EXAMPLE 1.9 Pulling wood apart

Student engineers learn that, although handbooks give the strength of a material as a single number, in fact the strength varies from piece to piece. A vital lesson in all fields of study is that “variation is everywhere.” Here are data from a typical student laboratory exercise: the load in pounds needed to pull apart pieces of Douglas fir 4 inches long and 1.5 inches square.

33,190	31,860	32,590	26,520	33,280
32,320	33,020	32,030	30,460	32,700
23,040	30,930	32,720	33,650	32,340
24,050	30,170	31,300	28,730	31,920

A stemplot of these data would have very many stems and no leaves or just one leaf on most stems. So we first **round** the data to the nearest hundred pounds. The rounded data are

332	319	326	265	333	323	330	320	305	327
230	309	327	337	323	241	302	313	287	319



Courtesy Department of Civil Engineering,
University of New Mexico

rounding



FIGURE 1.11

Stemplot of the breaking strength of pieces of wood, rounded to the nearest hundred pounds, for Example 1.9. Stems are thousands of pounds and leaves are hundreds of pounds.

23	0
24	1
25	
26	5
27	
28	7
29	
30	2 5 9
31	3 9 9
32	0 3 3 6 7 7
33	0 2 3 7

Now we can make a stemplot with the first two digits (thousands of pounds) as stems and the third digit (hundreds of pounds) as leaves. Figure 1.11 is the stemplot. Rotate the stemplot counterclockwise so that it resembles a histogram, with 230 at the left end of the scale. This makes it clear that the distribution is *skewed to the left*. The *midpoint* is around 320 (32,000 pounds) and the *spread* is from 230 to 337. Because of the strong skew, we are reluctant to call the smallest observations outliers. They appear to be part of the long left tail of the distribution. Before using wood like this in construction, we should ask why some pieces are much weaker than the rest. ■



Comparing Figures 1.10 (right-skewed) and 1.11 (left-skewed) reminds us that *the direction of skewness is the direction of the long tail, not the direction where most observations are clustered*.

splitting stems

You can also **split stems** in a stemplot to double the number of stems when all the leaves would otherwise fall on just a few stems. Each stem then appears twice. Leaves 0 to 4 go on the upper stem, and leaves 5 to 9 go on the lower stem. If you split the stems in the stemplot of Figure 1.11, for example, the 32 and 33 stems become

32	0 3 3
32	6 7 7
33	0 2 3
33	7

Rounding and splitting stems are matters for judgment, like choosing the classes in a histogram. The wood strength data require rounding but don't require splitting stems. The *One-Variable Statistical Calculator* applet on the text CD and Web site allows you to decide whether to split stems, so that it is easy to see the effect.



APPLY YOUR KNOWLEDGE

- 1.10 Older Americans.** Make a stemplot of percent of residents aged 65 years and over in each of the 50 states and the District of Columbia in Table 1.2. Use whole percents as your stems. Because the stemplot preserves the actual value of the

TABLE 1.3 Annual health expenditure per capita (PPP, international \$)

COUNTRY	DOLLARS	COUNTRY	DOLLARS	COUNTRY	DOLLARS
Argentina	1322	India	109	Russia	797
Australia	3357	Indonesia	81	Saudi Arabia	768
Austria	3763	Iran	689	South Africa	819
Belgium	3323	Ireland	3424	Spain	2671
Brazil	837	Italy	2686	Sweden	3323
Canada	3900	Japan	2696	Switzerland	4417
China	233	Korea, South	1688	Thailand	286
Denmark	3513	Mexico	819	Turkey	677
Finland	2840	Netherlands	3509	United Kingdom	2992
France	3709	Norway	4763	United States	7285
Germany	3588	Poland	1035	Venezuela	697
Greece	2727	Portugal	2284		

observations, it is easy to find the midpoint (26th of the 51 observations in order) and the spread. What are they? 

1.11 Health care spending. Table 1.3 shows the 2007 health care expenditure per capita in 35 countries with the highest gross domestic product in 2007.¹¹ Health expenditure per capita is the sum of public and private health expenditure (in PPP, international \$) divided by population. Health expenditures include the provision of health services, family-planning activities, nutrition activities, and emergency aid designated for health but exclude the provision of water and sanitation. Make a stemplot of the data after rounding to the nearest \$100 (so that stems are thousands of dollars and leaves are hundreds of dollars). Split the stems, placing leaves 0 to 4 on the first stem and leaves 5 to 9 on the second stem of the same value. Describe the shape, center, and spread of the distribution. Which country is the high outlier? 

TIME PLOTS

Many variables are measured at intervals over time. We might, for example, measure the height of a growing child or the price of a stock at the end of each month. In these examples, our main interest is change over time. To display change over time, make a *time plot*.

TIME PLOT

A time plot of a variable plots each observation against the time at which it was measured. Always put time on the horizontal scale of your plot and the variable you are measuring on the vertical scale. Connecting the data points by lines helps emphasize any change over time.

Courtesy U.S. Geological Survey



cycles

trend

EXAMPLE 1.10 Water levels in the Everglades

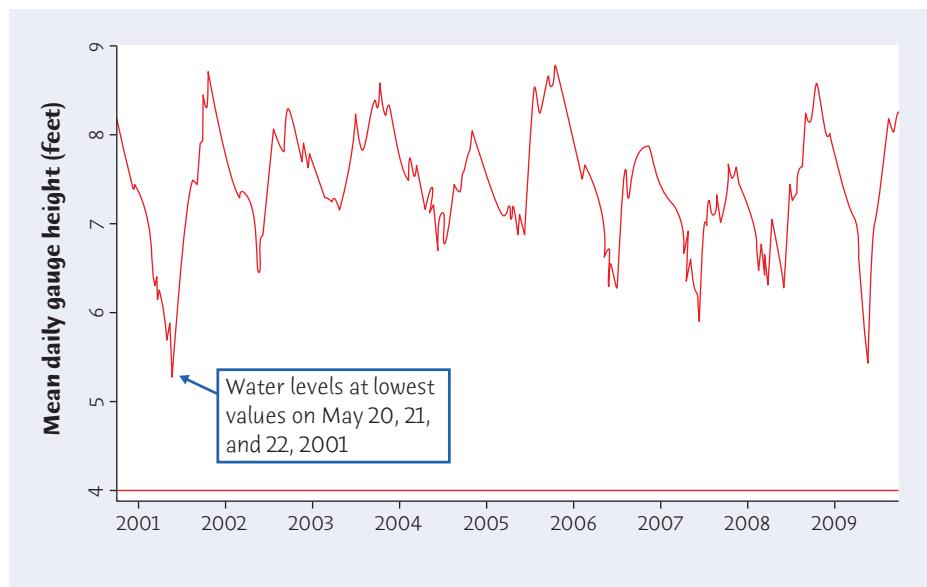
Water levels in Everglades National Park are critical to the survival of this unique region. The photo shows a water-monitoring station in Shark River Slough, the main path for surface water moving through the “river of grass” that is the Everglades. Each day the mean gauge height, the height of the water surface above the gauge datum, is measured at the Shark River Slough monitoring station. (The gauge datum is a vertical control measure established in 1929 and is used as a reference for establishing varying elevations. It establishes a zero point from which to measure the gauge height.) Figure 1.12 is a time plot of mean daily gauge height at this station from October 1, 2000 to September 30, 2009.¹² ■

When you examine a time plot, look once again for an overall pattern and for strong deviations from the pattern. Figure 1.12 shows strong **cycles**, regular up-and-down movements in water level. The cycles show the effects of Florida’s wet season (roughly June to November) and dry season (roughly December to May). Water levels are highest in late fall. If you look closely, you can see the year-to-year variation. The dry season in 2003 ended early, with the first-ever April tropical storm. In consequence, the dry-season water level in 2003 did not dip as low as in other years. The drought in the southeastern portion of the country in 2008 and 2009 shows up in the steep drop in the mean gauge height in 2009, while the lower peaks in 2006 and 2007 reflect lower water levels during the wet seasons in these years.

Another common overall pattern in a time plot is a **trend**, a long-term upward or downward movement over time. Many economic variables show an upward trend. Incomes, house prices, and (alas) college tuitions tend to move generally upward over time.

FIGURE 1.12

Time plot of average gauge height at a monitoring station in Everglades National Park over a nine-year period, for Example 1.10. The yearly cycles reflect Florida’s wet and dry seasons.



Histograms and time plots give different kinds of information about a variable. The time plot in Figure 1.12 presents **time series data** that show the change in water level at one location over time. A histogram displays **cross-sectional data**, such as water levels at many locations in the Everglades at the same time.

time series
cross-sectional

APPLY YOUR KNOWLEDGE

1.12 The cost of college. Here are data on the average tuition and fees charged to in-state students by public four-year colleges and universities for the 1980 to 2010 academic years. Because almost any variable measured in dollars increases over time due to inflation (the falling buying power of a dollar), the values are given in “constant dollars,” adjusted to have the same buying power that a dollar had in 2010.¹³  COLLEGECOST

Year	Tuition	Year	Tuition	Year	Tuition	Year	Tuition
1980	\$2119	1988	\$2903	1996	\$4131	2004	\$5900
1981	\$2163	1989	\$2972	1997	\$4226	2005	\$6128
1982	\$2305	1990	\$3190	1998	\$4338	2006	\$6218
1983	\$2505	1991	\$3373	1999	\$4397	2007	\$6480
1984	\$2572	1992	\$3622	2000	\$4426	2008	\$6532
1985	\$2665	1993	\$3827	2001	\$4626	2009	\$7137
1986	\$2815	1994	\$3974	2002	\$4961	2010	\$7605
1987	\$2845	1995	\$4019	2003	\$5507		

- Make a time plot of average tuition and fees.
- What overall pattern does your plot show?
- Some possible deviations from the overall pattern are outliers, periods when charges went down (in 2010 dollars) and periods of particularly rapid increase. Which are present in your plot, and during which years?
- In looking for patterns, do you think that it would be better to study a time series of the tuition for each year or the percent increase for each year? Why?

CHAPTER 1 SUMMARY

CHAPTER SPECIFICS

- A data set contains information on a number of **individuals**. Individuals may be people, animals, or things. For each individual, the data give values for one or more **variables**. A variable describes some characteristic of an individual, such as a person's height, sex, or salary.
- Some variables are **categorical** and others are **quantitative**. A categorical variable places each individual into a category, such as male or female. A quantitative variable has numerical values that measure some characteristic of each individual, such as height in centimeters or salary in dollars.

- Exploratory data analysis uses graphs and numerical summaries to describe the variables in a data set and the relations among them.
- After you understand the background of your data (individuals, variables, units of measurement), the first thing to do is almost always **plot your data**.
- The **distribution** of a variable describes what values the variable takes and how often it takes these values. **Pie charts** and **bar graphs** display the distribution of a categorical variable. Bar graphs can also compare any set of quantities measured in the same units. **Histograms** and **stemplots** graph the distribution of a quantitative variable.
- When examining any graph, look for an **overall pattern** and for notable **deviations** from the pattern.
- **Shape, center, and spread** describe the overall pattern of the distribution of a quantitative variable. Some distributions have simple shapes, such as **symmetric** or **skewed**. Not all distributions have a simple overall shape, especially when there are few observations.
- **Outliers** are observations that lie outside the overall pattern of a distribution. Always look for outliers and try to explain them.
- When observations on a variable are taken over time, make a **time plot** that graphs time horizontally and the values of the variable vertically. A time plot can reveal **trends, cycles**, or other changes over time.

LINK IT

Practical statistics uses data to draw conclusions about some wider universe. You should reread Example 1.1 of this chapter, as it will help you understand this basic idea. For the American Community Survey described in the example, the data are the responses from those households responding to the survey, although the wider universe of interest is the entire nation.

In our study of practical statistics, we will divide the subject into three main areas. In exploratory data analysis, graphs and numerical summaries are used for exploring, organizing, and describing data so that the patterns become apparent. Data production concerns where the data come from and helps us to understand whether what we learn from our data can be generalized to a wider universe. And statistical inference provides tools for generalizing what we learn to a wider universe.

In this chapter we have begun to learn about data analysis. A data set can consist of hundreds of observations on many variables. Even if we consider only one variable at a time, it is difficult to see what the data have to say by scanning a list containing many data values. Graphs provide a visual tool for organizing and identifying patterns in data and are a good starting point in the exploration of the distribution of a variable. Pie charts and bar graphs can summarize the information provided by a categorical variable by giving us the percent of the distribution in the various categories. Although a table containing the categories and percents gives the same information as a bar graph, a substantial advantage of the bar graph over a tabular presentation is that the bar graph allows us to visually compare percents among all categories simultaneously by means of the heights of the bars. Histograms and stemplots are graphical tools for summarizing the information provided by a quantitative variable. The overall pattern in a histogram or stemplot illustrates some of the important features of the distribution of a variable that will be of interest as we continue our study of practical statistics. The center of the histogram tells us about the value of a “typical” observation on this variable, while the spread gives us a sense of how

close most of the observations are to this value. Other interesting features are the presence of outliers and the general shape of the plot. For data collected over time, time plots can show patterns such as seasonal variation and trends in the variable. In the next chapter we will see how the information about the distribution of a variable can also be described using numerical summaries.

CHECK YOUR SKILLS

The multiple-choice exercises in Check Your Skills ask straightforward questions about basic facts from the chapter. Answers to all these exercises appear in the back of the book. You should expect almost all your answers to be correct.

1.13 Here are the first lines of a professor's data set at the end of a statistics course:

Name	Major	Points	Grade
ADVANI, SURA	COMM	397	B
BARTON, DAVID	HIST	323	C
BROWN, ANNENETTE	BIOL	446	A
CHIU, SUN	PSYC	405	B
CORTEZ, MARIA	PSYC	461	A

The individuals in these data are

- (a) the students. (b) the total points.
- (c) the course grades.

1.14 To display the distribution of grades (A, B, C, D, F) in the course, it would be correct to use

- (a) a pie chart but not a bar graph.
- (b) a bar graph but not a pie chart.
- (c) either a pie chart or a bar graph.

1.15 A description of different houses on the market includes the variables square footage of the house and the average monthly gas bill.

- (a) Square footage and average monthly gas bill are both categorical variables.
- (b) Square footage and average monthly gas bill are both quantitative variables.
- (c) Square footage is a categorical variable and average monthly gas bill is a quantitative variable.

1.16 A political party's data bank includes the zip codes of past donors, such as

47906	34236	53075	10010
90210	75204	30304	99709

Zip code is a

- (a) quantitative variable. (b) categorical variable.
- (c) unit of measurement.

1.17 Figure 1.9 (page 19) is a histogram of the percent of women in each state aged 15 and over who have never been married. The leftmost bar in the histogram covers percents of never-married women ranging from about

- (a) 20% to 24%. (b) 20% to 22%. (c) 0% to 20%.

1.18 Here are the amounts of money (cents) in coins carried by 10 students in a statistics class:

50 35 0 97 76 0 0 87 23 65

To make a stemplot of these data, you would use stems

- (a) 0, 1, 2, 3, 4, 5, 6, 7, 8, 9.
- (b) 0, 2, 3, 5, 6, 7, 8, 9.
- (c) 00, 10, 20, 30, 40, 50, 60, 70, 80, 90.

1.19 How long must you travel each day to get to work? Here is a stemplot of the average travel times to work for workers in the 50 states and the District of Columbia who are at least 16 years of age and don't work at home.¹⁴ The stems are whole minutes and the leaves are tenths of a minute. 

15	59
16	
17	6779
18	25
19	
20	017889
21	28
22	01333499
23	445669
24	01266
25	0012569
26	689
27	39
28	
29	12
30	69

The state with the longest average travel time is New York. On average, how long does it take New Yorkers to travel to work each day?

- (a) 30.69 minutes (b) 309 minutes (c) 30.9 minutes

1.20 The shape of the distribution in Exercise 1.19 is

- (a) clearly skewed to the right.
- (b) roughly symmetric.
- (c) clearly skewed to the left.

1.21 The center of the distribution in Exercise 1.19 is close to

- (a) 22 minutes.
- (b) 23.4 minutes.
- (c) 15.5 to 30.9 minutes.

1.22 You look at real estate ads for houses in Naples, Florida. There are many houses ranging from \$200,000 to \$500,000 in price. The few houses on the water, however, have prices up to \$15 million. The distribution of house prices will be

- (a) skewed to the left.
- (b) roughly symmetric.
- (c) skewed to the right.

CHAPTER 1 EXERCISES

1.23 Medical students. Students who have finished medical school are assigned to residencies in hospitals to receive further training in a medical specialty. Here is part of a hypothetical data base of students seeking residency positions. USMLE is the student's score on Step 1 of the national medical licensing examination.

Name	Medical school	Sex	Age	USMLE	Specialty sought
Abrams, Laurie	Florida	F	28	238	Family medicine
Brown, Gordon	Meharry	M	25	205	Radiology
Cabrera, Maria	Tufts	F	26	191	Pediatrics
Ismael, Miranda	Indiana	F	32	245	Internal medicine

(a) What individuals does this data set describe?

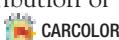
(b) In addition to the student's name, how many variables does the data set contain? Which of these variables are categorical and which are quantitative?

1.24 Protecting wood. How can we help wood surfaces resist weathering, especially when restoring historic wooden buildings? In a study of this question, researchers prepared

wooden panels and then exposed them to the weather. Here are some of the variables recorded. Which of these variables are categorical and which are quantitative?

- (a) Type of wood (yellow poplar, pine, cedar)
- (b) Type of water repellent (solvent-based, water-based)
- (c) Paint thickness (millimeters)
- (d) Paint color (white, gray, light blue)
- (e) Weathering time (months)

1.25 What color is your car? The most popular colors for cars and light trucks vary with region and over time. In North America white remains the top color choice, with black the top choice in Europe and silver the top choice in South America. Here is the distribution of the top colors for vehicles sold globally in 2010:¹⁵



Color	Popularity
Silver	26%
Black	24%
White	16%
Gray	16%
Red	6%
Blue	5%
Beige, brown	3%
Other colors	



© Photo 24/age Fotostock

Fill in the percent of vehicles that are in other colors. Make a graph to display the distribution of color popularity.

1.26 Facebook and MySpace audience. Although most social-networking Web sites in the United States have fairly short histories, the growth of these sites has been exponential. By far, the two most visited social-networking sites are Facebook.com and MySpace.com. Here is the age distribution of the audience for the two sites in December 2009:¹⁶



Age group	Facebook visitors	MySpace visitors
Under 25 years	26.8%	44.4%
25 to 34 years	23.0%	22.7%
35 to 49 years	31.6%	23.5%
Over 49 years	18.7%	9.4%

(a) Draw a bar graph for the age distribution of Facebook visitors. The leftmost bar should correspond to “under 25,” the next bar to “25 to 34,” and so on. Do the same for MySpace, using the same scale for the percent axis.

(b) Describe the most important difference in the age distribution of the audience for Facebook and MySpace. How does this difference show up in the bar graphs? Do you think it was important to order the bars by age to make the comparison easier?

(c) Explain why it is appropriate to use a pie chart to display either of these distributions. Draw a pie chart for each distribution. Do you think it is easier to compare the two distributions with bar graphs or pie charts? Explain your reasoning.

1.27 Deaths among young people. Among persons aged 15 to 24 years in the United States, the leading causes of death and number of deaths in 2008 were: accidents, 14,020; homicide, 5285; suicide, 4297; cancer, 1659; heart disease, 1059; congenital defects, 466.¹⁷

- (a) Make a bar graph to display these data.
 (b) To make a pie chart, you need one additional piece of information. What is it?

1.28 Hispanic origins. Figure 1.13 is a pie chart prepared by the U.S. Census Bureau to show the origin of the more than 43 million Hispanics in the United States in 2006.¹⁸ About what percent of Hispanics are Mexican? Puerto Rican? You see that it is hard to determine numbers from a pie chart. Bar graphs are much easier to use. (The U.S. Census Bureau did include the percents in its pie chart.)

Percent Distribution of Hispanics by Type: 2006

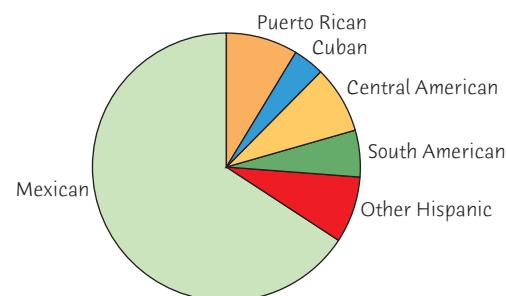


FIGURE 1.13

Pie chart of the national origins of Hispanic residents of the United States, for Exercise 1.28.

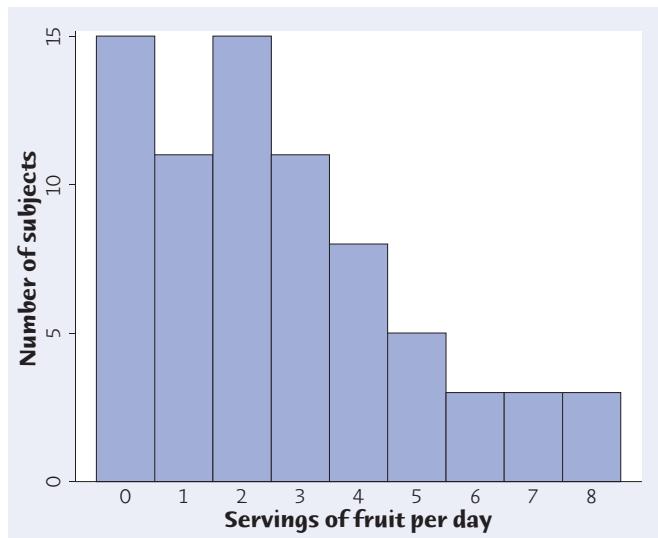
1.29 Canadian students rate their universities. The National Survey of Student Engagement asked students at many universities, “How would you evaluate your entire educational experience at this university?” Here are the percents of senior-year students at Canada’s 10 largest primarily English-speaking universities who responded “Excellent”:¹⁹



University	Excellent rating
Toronto	21%
York	18%
Alberta	23%
Ottawa	11%
Western Ontario	38%
British Columbia	18%
Calgary	14%
McGill	26%
Waterloo	36%
Concordia	21%

- (a) The list is arranged in order of undergraduate enrollment. Make a bar graph with the bars in order of student rating.
 (b) Explain carefully why it is not correct to make a pie chart of these data.

1.30 Do adolescent girls eat fruit? We all know that fruit is good for us. Many of us don’t eat enough. Figure 1.14 is a histogram of the number of servings of fruit per day claimed by 74 seventeen-year-old girls in a study in Pennsylvania.²⁰ Describe the shape, center, and spread of this distribution.

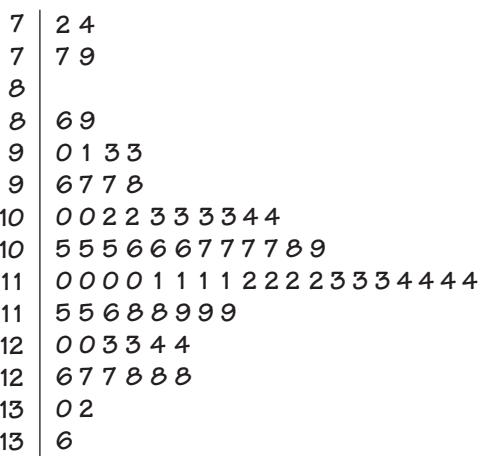
**FIGURE 1.14**

The distribution of fruit consumption in a sample of 74 seventeen-year-old girls, for Exercise 1.30.

What percent of these girls ate six or more servings per day? How many of these girls ate fewer than two servings per day? Are there any outliers?

1.31 IQ test scores. Figure 1.15 is a stemplot of the IQ test scores of 78 seventh-grade students in a rural midwestern school.²¹ 

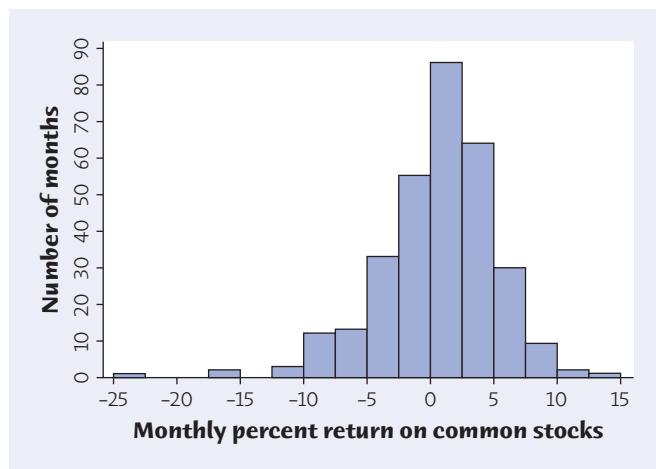
(a) Four students had low scores that might be considered outliers. Ignoring these, describe the shape, center, and spread of the remainder of the distribution.

**FIGURE 1.15**

The distribution of IQ scores for 78 seventh-grade students, for Exercise 1.31.

- (b) We often read that IQ scores for large populations are centered at 100. What percent of these 78 students have scores above 100?

1.32 Returns on common stocks. The return on a stock is the change in its market price plus any dividend payments made. Total return is usually expressed as a percent of the beginning price. Figure 1.16 is a histogram of the distribution of the monthly returns for all stocks listed on U.S. markets from January 1985 to November 2010 (311 months).²² The extreme low outlier is the market crash of October 1987, when stocks lost 23% of their value in one month. The other two low outliers are 16% during August 1998, a month when the Dow Jones Industrial Average experienced its second largest drop in history to that time, and the financial crisis in October 2008 when stocks lost 17% of their value. 

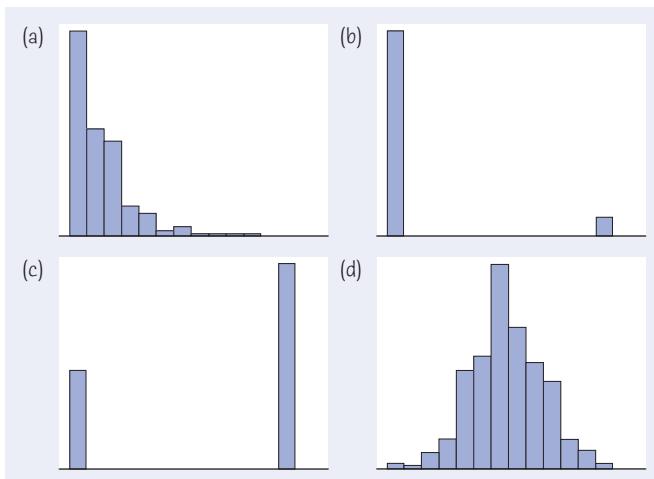
**FIGURE 1.16**

The distribution of monthly percent returns on U.S. common stocks from January 1985 to November 2010, for Exercise 1.32.

- (a) Ignoring the outliers, describe the overall shape of the distribution of monthly returns.
 (b) What is the approximate center of this distribution? (For now, take the center to be the value with roughly half the months having lower returns and half having higher returns.)
 (c) Approximately what were the smallest and largest monthly returns, leaving out the outliers? (This is one way to describe the spread of the distribution.)
 (d) A return less than zero means that stocks lost value in that month. About what percent of all months had returns less than zero?

1.33 Name that variable. A survey of a large college class asked the following questions:

- Are you female or male? (In the data, male = 0, female = 1.)

**FIGURE 1.17**

Histograms of four distributions, for Exercise 1.33.

2. Are you right-handed or left-handed? (In the data, right = 0, left = 1.)
3. What is your height in inches?
4. How many minutes do you study on a typical weeknight?

Figure 1.17 shows histograms of the student responses, in scrambled order and without scale markings. Which histogram goes with each variable? Explain your reasoning.

1.34 Food oils and health. Fatty acids, despite their unpleasant name, are necessary for human health. Two types of essential fatty acids, called omega-3 and omega-6, are not produced by our bodies and so must be obtained from our food. Food oils, widely used in food processing and cooking, are major sources of these compounds. There is some evidence that a healthy diet should have more omega-3 than omega-6. Table 1.4 (on the following page) gives the ratio of omega-3 to omega-6 in some common food oils.²³ Values greater than 1 show that an oil has more omega-3 than omega-6. 

- (a) Make a histogram of these data, using classes bounded by the whole numbers from 0 to 6.
- (b) What is the shape of the distribution? How many of the 30 food oils have more omega-3 than omega-6? What does this distribution suggest about the possible health effects of modern food oils?
- (c) Table 1.4 contains entries for several fish oils (cod, herring, menhaden, salmon, sardine). How do these values support the idea that eating fish is healthy?

1.35 Where are the nurses? Table 1.5 (on the following page) gives the number of active nurses per 100,000 people in each state.²⁴ 

- (a) Why is the number of nurses per 100,000 people a better measure of the availability of nurses than a simple count of the number of nurses in a state?

TABLE 1.4 Omega-3 fatty acids as a fraction of omega-6 fatty acids in food oils

OIL	RATIO	OIL	RATIO
Perilla	5.33	Flaxseed	3.56
Walnut	0.20	Canola	0.46
Wheat germ	0.13	Soybean	0.13
Mustard	0.38	Grape seed	0.00
Sardine	2.16	Menhaden	1.96
Salmon	2.50	Herring	2.67
Mayonnaise	0.06	Soybean, hydrogenated	0.07
Cod liver	2.00	Rice bran	0.05
Shortening (household)	0.11	Butter	0.64
Shortening (industrial)	0.06	Sunflower	0.03
Margarine	0.05	Corn	0.01
Olive	0.08	Sesame	0.01
Shea nut	0.06	Cottonseed	0.00
Sunflower (oleic)	0.05	Palm	0.02
Sunflower (linoleic)	0.00	Cocoa butter	0.04

TABLE 1.5 Nurses per 100,000 people, by state

STATE	NURSES	STATE	NURSES	STATE	NURSES
Alabama	912	Louisiana	894	Ohio	1001
Alaska	756	Maine	1053	Oklahoma	712
Arizona	544	Maryland	869	Oregon	795
Arkansas	774	Massachusetts	1210	Pennsylvania	1017
California	641	Michigan	841	Rhode Island	1007
Colorado	761	Minnesota	1017	South Carolina	795
Connecticut	994	Mississippi	868	South Dakota	1215
Delaware	977	Missouri	958	Tennessee	894
Florida	814	Montana	748	Texas	662
Georgia	653	Nebraska	1010	Utah	625
Hawaii	753	Nevada	574	Vermont	912
Idaho	642	New Hampshire	970	Virginia	750
Illinois	812	New Jersey	907	Washington	774
Indiana	864	New Mexico	580	West Virginia	938
Iowa	990	New York	859	Wisconsin	905
Kansas	867	North Carolina	886	Wyoming	812
Kentucky	923	North Dakota	1097	District of Columbia	1380

- (b) Make a histogram that displays the distribution of nurses per 100,000 people. Write a brief description of the distribution. Are there any outliers? If so, can you explain them?

1.36 Carbon dioxide emissions. Burning fuels in power plants and motor vehicles emits carbon dioxide (CO_2), which contributes to global warming. Table 1.6 (on the following page) displays the 2007 CO_2 emissions per person from countries with populations of at least 30 million in that year.²⁵ 

- (a) Why do you think we choose to measure emissions per person rather than total CO_2 emissions for each country?
 (b) Make a stemplot to display the data of Table 1.6. The data will first need to be rounded. What units are you going to use for the stems? The leaves? You should round the data to the units you are planning to use for the leaves before drawing the stemplot. Describe the shape, center, and spread of the distribution. Which countries are outliers?

1.37 Fur seals on St. George Island. Every year hundreds of thousands of northern fur seals return to their haul-outs in the Pribilof Islands in Alaska to breed, give birth, and teach their pups to swim, hunt, and survive in the Bering Sea. U.S. commercial fur sealing operations continued until 1983, but despite a reduction in harvest, the population of fur seals has continued to decline. Here are data on the number of fur seal pups born on St. George Island (in thousands) from 1975 to 2006:²⁶ 

Year	Pups born (thousands)	Year	Pups born (thousands)
1975	53.70	1991	24.28
1976	56.16	1992	25.16
1977	43.41	1993	23.70
1978	47.25	1994	22.24
1979	47.47	1995	24.82
1980	39.34	1996	27.39
1981	38.15	1997	24.74
1982	39.29	1998	22.09
1983	31.44	1999	21.13
1984	33.44	2000	20.18
1985	28.87	2001	18.89
1986	32.36	2002	17.59
1987	33.12	2003	17.24
1988	24.82	2004	16.88
1989	33.11	2005	16.97
1990	23.40	2006	17.07

TABLE 1.6 Annual carbon dioxide emissions in 2007 (metric tons per person)

COUNTRY	CO ₂	COUNTRY	CO ₂	COUNTRY	CO ₂
Afghanistan	0.0272	India	1.3844	Poland	8.3231
Algeria	4.1384	Indonesia	1.7677	Russia	10.8309
Argentina	4.6525	Iran	6.8472	South Africa	8.8163
Bangladesh	0.2773	Italy	7.6923	Spain	8.1555
Brazil	1.9373	Japan	9.8476	Sudan	0.2850
Canada	16.9171	Kenya	0.2976	Tanzania	0.1464
China	4.9194	Korea, South	10.4941	Thailand	4.1432
Colombia	1.4301	Mexico	4.3862	Turkey	3.9543
Congo	0.0389	Morocco	1.4862	Uganda	0.1046
Egypt	2.3065	Myanmar	0.2685	Ukraine	6.8598
Ethiopia	0.0828	Nigeria	0.6449	United Kingdom	8.8608
France	6.0207	Pakistan	0.9031	United States	18.9144
Germany	9.5690	Philippines	0.7993	Vietnam	1.2935

Make a stemplot to display the distribution of pups born per year. (Round to the nearest whole number and split the stems.) Describe the shape, center, and spread of the distribution. Are there any outliers?



Harry Walker/Photolibrary

1.38 Do women study more than men? We asked the students in a large first-year college class how many minutes they studied on a typical weeknight. Here are the responses of random samples of 30 women and 30 men from the class:  STUDYTIMES

(a) Examine the data. Why are you not surprised that most responses are multiples of 10 minutes? What is the other common multiple found in the data? We eliminated one student

who claimed to study 10,000 minutes per night. Are there any other responses you consider suspicious?

(b) Make a **back-to-back stemplot** to compare the two samples. That is, use one set of stems with two sets of leaves, one to the right and one to the left of the stems. (Draw a line on either side of the stems to separate stems and leaves.) Order both sets of leaves from smallest at the stem to largest away from the stem. Report the approximate midpoints of both groups. Does it appear that women study more than men (or at least claim that they do)?

back-to-back stemplot

Women					Men				
270	150	180	360	180	120	120	30	45	200
120	180	120	240	170	90	90	30	120	75
150	120	180	180	150	150	90	60	240	300
200	150	180	120	240	240	60	150	60	30
120	60	120	180	180	30	230	120	95	150
90	240	180	115	120	0	200	120	120	180

1.39 Fur seals on St. George Island. Make a time plot of the number of fur seals born per year from Exercise 1.37. What does the time plot show that your stemplot in Exercise 1.37 did not show? When you have data collected over time, a time plot is often needed to understand what is happening.  FURSEALS

1.40 Marijuana and traffic accidents. Researchers in New Zealand interviewed 907 drivers at age 21. They had data on traffic accidents and they asked the drivers about marijuana use. Here are data on the numbers of accidents caused by these drivers at age 19, broken down by marijuana use at the same age:²⁷

Marijuana Use per Year				
	Never	1–10 times	11–50 times	51+ times
Drivers	452	229	70	156
Accidents caused	59	36	15	50

(a) Explain carefully why a useful graph must compare *rates* (accidents per driver) rather than *counts* of accidents in the four marijuana use classes.

(b) Compute the accident rates in the four marijuana use classes. After you have done this, make a graph that displays the accident rate for each class. What do you conclude? (You can't conclude that marijuana use *causes* accidents, because risk takers are more likely both to drive aggressively and to use marijuana.)

1.41 Dates on coins. Sketch a histogram for a distribution that is skewed to the left. Suppose that you and your friends emptied your pockets of coins and recorded the year marked on each coin. The distribution of dates would be skewed to the left. Explain why.

1.42 El Niño and the monsoon. The earth is interconnected. For example, it appears that El Niño, the periodic warming of the Pacific Ocean west of South America, affects the monsoon rains that are essential for agriculture in India. Here are the monsoon rains (in millimeters) for the 23 strong El Niño years between 1871 and 2004:²⁸ 

628	669	740	651	710	736	717	698
653	604	781	784	790	811	830	858
858	896	806	790	792	957	872	

(a) To make a stemplot of these rainfall amounts, round the data to the nearest 10, so that stems are hundreds of millimeters and leaves are tens of millimeters. Make two stemplots, with and without splitting the stems. Which plot do you prefer?

(b) Describe the shape, center, and spread of the distribution. Are there any outliers?

(c) The average monsoon rainfall for all years from 1871 to 2004 is about 850 millimeters. What effect does El Niño appear to have on monsoon rains?

1.43 Watch those scales! Figures 1.18(a) and 1.18(b) both show time plots of tuition charged to in-state students from 1980 through 2010.²⁹

(a) Which graph appears to show the biggest increase in tuition between 2000 and 2010?

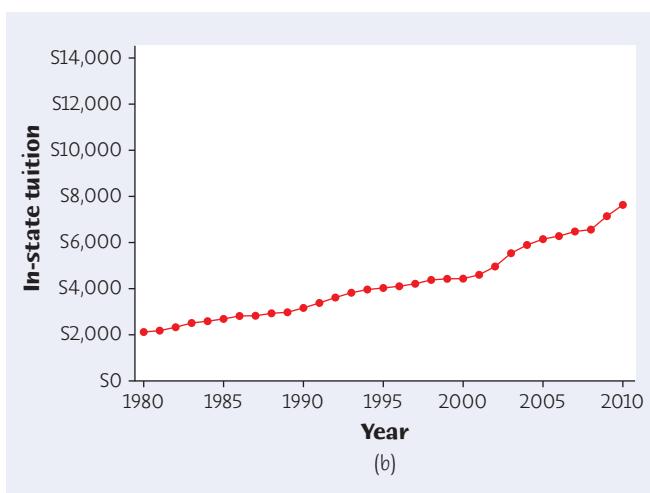
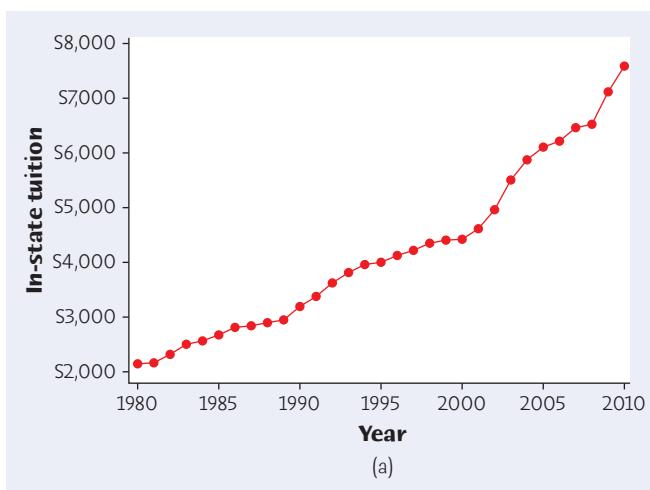


FIGURE 1.18

Time plots of in-state tuition between 1980 and 2010, for Exercise 1.43.

(b) Read the graphs and compute the actual increase in tuition between 2000 and 2010 in each graph. Do you think these graphs are for the same or different data sets? Why?

The impression that a time plot gives depends on the scales you use on the two axes. Changing the scales can make tuition appear to increase very rapidly or to have only a gentle increase. The moral of this exercise is: always pay close attention to the scales when you look at a time plot.

1.44 Housing starts. Figure 1.19 is a time plot of the number of single-family homes started by builders each month from January 1990 to August 2011.³⁰ The counts are in thousands of homes.  HOUSESTARTS

(a) The most notable pattern in this time plot is yearly up-and-down cycles. At what season of the year are housing starts highest? Lowest? The cycles are explained by the weather in the northern part of the country.

(b) Is there a longer-term trend visible in addition to the cycles? If so, describe it.

(c) The big economic news of 2007 was a severe downturn in housing that began in mid-2006. This was followed by the financial crisis in 2008. How are these economic events reflected in the time plot?

1.45 Ozone hole. The ozone hole is a region in the stratosphere over the Antarctic with exceptionally depleted ozone. The size of the hole is not constant over the year but is largest at the beginning of the Southern Hemisphere spring (August–October). The increase in

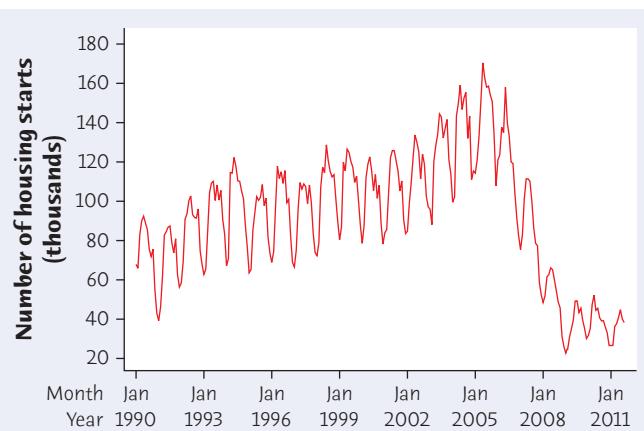


FIGURE 1.19

Time plot of the monthly count of new single-family homes started (in thousands) between January 1990 and August 2011, for Exercise 1.44.

the size of the ozone hole led to the Montreal Protocol in 1987, an international treaty designed to protect the ozone layer by phasing out the production of substances, such as chlorofluorocarbons (CFCs), believed to be responsible for ozone depletion. The table at the bottom of the page gives the average ozone hole size for the period September 7 to October 13 for each of the years from 1979 through

Year	Area (millions of km ²)	Year	Area (millions of km ²)	Year	Area (millions of km ²)
1979	0.1	1990	19.3	2001	25.0
1980	1.2	1991	19.0	2002	11.9
1981	0.6	1992	22.3	2003	25.8
1982	4.6	1993	24.2	2004	19.0
1983	7.7	1994	23.5	2005	23.9
1984	9.8	1995		2006	26.2
1985	14.1	1996	22.8	2007	21.6
1986	11.2	1997	22.1	2008	24.7
1987	19.3	1998	25.9	2009	21.6
1988	10.0	1999	23.3	2010	19.0
1989	18.8	2000	24.8		

2010 (note that no data were acquired in 1995).³¹ To get a better feel for the magnitude of the numbers, the area of North America is approximately 24.5 million square kilometers (km^2).  OZONEHOLE

The two parts of this exercise will have you draw two graphs of these data.

(a) First make a time plot of the data. The severity of the ozone hole will vary from year to year depending on the meteorology of the atmosphere above Antarctica. Does the time plot illustrate only year-to-year variation or are there other patterns apparent? Specifically, is there a trend over any period of years? What about cyclical fluctuation? Explain in words the change in the average size of the ozone hole over this 30-year period.

(b) Now make a stemplot of the data. What is the midpoint of the distribution of ozone hole size? Do you think that the stemplot and the midpoint are a good description of this data set? Is there important information in the time plot that is not contained in the stemplot? When data are collected over time, you should always make a time plot.

1.46 To split or not to split. The data sets in the *One-Variable Statistical Calculator* applet on the text CD and Web site include the “pulling wood apart” data from Example 1.9. The applet rounds the data in the same way as Figure 1.11 (page 22). Use the applet to make a stemplot with split stems. Do you prefer this stemplot or that in Figure 1.11? Explain your choice.



EXPLORING THE WEB

1.47 Natural Gas Prices. The Department of Energy Web site contains information about monthly wholesale and retail prices for natural gas in each state. Go to www.eia.doe.gov/naturalgas/data.cfm and then click on the link Monthly Wholesale and Retail Prices. Under Area, choose a state of interest to you, make sure the Period is monthly, and then under Residential Price click on View History. A window will open with a time plot covering approximately a 20-year period, along with a table of the monthly residential prices for each year.

- (a) If you have access to statistical software, you should use the *Download Data (XLS file)* link to save the data as an Excel (.xls) file on your computer. Then enter the data into your software package, and reproduce the time series plot using the graphical capabilities of your software package. Be sure you use an appropriate title and axis labels. If you do not have access to appropriate software, provide a rough sketch of the time plot that is given on the Web site.
- (b) Is there a regular pattern of seasonal variation that repeats each year? Describe it. Are the prices increasing over time?

1.48 Hank Aaron’s home run record. The all-time home run leader prior to 2007 was Hank Aaron. You can find his career statistics by going to the Web site www.baseball-reference.com and then clicking on the Players tab at the top of the page and going to Hank Aaron.

(a) Make a stemplot or a histogram of the number of home runs that Hank Aaron hit in each year during his career. Is the distribution roughly symmetric, clearly skewed, or neither? About how many home runs did Aaron hit in a typical year? Are there any outliers?

(b) Would a time plot be appropriate for these data? If so, what information would be included in the time plot that is not in the stemplot?

TO
SOUTH
I-40
Raleigh
Wilmington



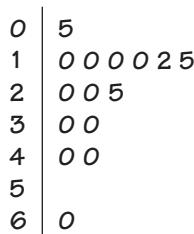
Describing Distributions with Numbers

Chapter 2

We saw in Chapter 1 (page 4) that the American Community Survey asks, among much else, workers' travel times to work. Here are the travel times in minutes for 15 workers in North Carolina, chosen at random by the Census Bureau:¹

30 20 10 40 25 20 10 60 15 40 5 30 12 10 10

We aren't surprised that most people estimate their travel time in multiples of 5 minutes. Here is a stemplot of these data:



The distribution is single-peaked and right-skewed. The longest travel time (60 minutes) may be an outlier. Our goal in this chapter is to describe with numbers the center and spread of this and other distributions.

IN THIS CHAPTER WE COVER...

- Measuring center: the mean
- Measuring center: the median
- Comparing the mean and the median
- Measuring spread: the quartiles
- The five-number summary and boxplots
- Spotting suspected outliers*
- Measuring spread: the standard deviation
- Choosing measures of center and spread
- Using technology
- Organizing a statistical problem

Logan Mock-Bunting/Getty Images

MEASURING CENTER: The Mean

The most common measure of center is the ordinary arithmetic average, or *mean*.

THE MEAN \bar{x}

To find the **mean** of a set of observations, add their values and divide by the number of observations. If the n observations are x_1, x_2, \dots, x_n , their mean is

$$\bar{x} = \frac{x_1 + x_2 + \cdots + x_n}{n}$$

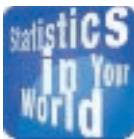
or, in more compact notation,

$$\bar{x} = \frac{1}{n} \sum x_i$$

The Σ (capital Greek sigma) in the formula for the mean is short for “add them all up.” The subscripts on the observations x_i are just a way of keeping the n observations distinct. They do not necessarily indicate order or any other special facts about the data. The bar over the x indicates the mean of all the x -values. Pronounce the mean \bar{x} as “ x -bar.” This notation is very common. When writers who are discussing data use \bar{x} or \bar{y} , they are talking about a mean.



NCTRAVELTIME



Don't hide the outliers

Data from an airliner's control surfaces, such

as the vertical tail rudder, go to cockpit instruments and then to the “black box” flight data recorder. To avoid confusing the pilots, short erratic movements in the data are “smoothed” so that the instruments show overall patterns. When a crash killed 260 people, investigators suspected a catastrophic movement of the tail rudder. But the black box contained only the smoothed data. Sometimes outliers are more important than the overall pattern.

resistant measure

EXAMPLE 2.1 Travel times to work

The mean travel time of our 15 North Carolina workers is

$$\begin{aligned}\bar{x} &= \frac{x_1 + x_2 + \cdots + x_n}{n} \\ &= \frac{30 + 20 + \cdots + 10}{15} \\ &= \frac{337}{15} = 22.5 \text{ minutes}\end{aligned}$$

In practice, you can enter the data into your calculator and ask for the mean. You don't have to actually add and divide. But you should know that this is what the calculator is doing.

Notice that only 6 of the 15 travel times are larger than the mean. If we leave out the longest single travel time, 60 minutes, the mean for the remaining 14 people is 19.8 minutes. That one observation raises the mean by 2.7 minutes. ■

Example 2.1 illustrates an important fact about the mean as a measure of center: it is sensitive to the influence of a few extreme observations. These may be outliers, but a skewed distribution that has no outliers will also pull the mean toward its long tail. Because the mean cannot resist the influence of extreme observations, we say that it is not a **resistant measure** of center.

APPLY YOUR KNOWLEDGE

- 2.1 Pulling wood apart.** Example 1.9 (page 21) gives the breaking strength in pounds of 20 pieces of Douglas fir. Find the mean breaking strength. How many of the pieces of wood have strengths less than the mean? What feature of the stemplot (Figure 1.11, page 22) explains the fact that the mean is smaller than most of the observations?  WOOD
- 2.2 Health care spending.** Table 1.3 (page 23) gives the 2007 health care expenditure per capita in 35 countries with the highest gross domestic product in 2007. The United States, at \$7285 (PPP, international \$) per person, is a high outlier. Find the mean health care spending in these nations with and without the United States. How much does the one outlier increase the mean?  HEALTHCARE

MEASURING CENTER: The Median

In Chapter 1, we used the midpoint of a distribution as an informal measure of center. The *median* is the formal version of the midpoint, with a specific rule for calculation.

THE MEDIAN M

The **median M** is the midpoint of a distribution, the number such that half the observations are smaller and the other half are larger. To find the median of a distribution:

1. Arrange all observations in order of size, from smallest to largest.
2. If the number of observations n is odd, the median M is the center observation in the ordered list. If the number of observations n is even, the median M is midway between the two center observations in the ordered list.
3. You can always locate the median in the ordered list of observations by counting up $(n + 1)/2$ observations from the start of the list.

Note that the formula $(n + 1)/2$ does not give the median, just the location of the median in the ordered list. Medians require little arithmetic, so they are easy to find by hand for small sets of data. Arranging even a moderate number of observations in order is very tedious, however, so that finding the median by hand for larger sets of data is unpleasant. Even simple calculators have an \bar{x} button, but you will need to use software or a graphing calculator to automate finding the median.

EXAMPLE 2.2 Finding the median: odd n

What is the median travel time for our 15 North Carolina workers? Here are the data arranged in order:

5 10 10 10 10 12 15 **20** 20 20 25 30 30 40 40 60

The count of observations $n = 15$ is odd. The bold 20 is the center observation in the ordered list, with 7 observations to its left and 7 to its right. This is the median, $M = 20$ minutes.

Because $n = 15$, our rule for the location of the median gives

$$\text{location of } M = \frac{n+1}{2} = \frac{16}{2} = 8$$

That is, the median is the 8th observation in the ordered list. It is faster to use this rule than to locate the center by eye. ■

Mitchell Funk/Getty Images



EXAMPLE 2.3 Finding the median: even n

Travel times to work in New York State are (on the average) longer than in North Carolina. Here are the travel times in minutes of 20 randomly chosen New York workers:

10 30 5 25 40 20 10 15 30 20 15 20 85 15 15 65 15 60 60 40 45

A stemplot not only displays the distribution but makes finding the median easy because it arranges the observations in order:

0	5
1	0 0 5 5 5 5
2	0 0 0 5
3	0 0
4	0 0 5
5	
6	0 0 5
7	
8	5

The distribution is single-peaked and right-skewed, with several travel times of an hour or more. There is no center observation, but there is a center pair. These are the bold 20 and 25 in the stemplot, which have 9 observations before them in the ordered list and 9 after them. The median is midway between these two observations:

$$M = \frac{20 + 25}{2} = 22.5 \text{ minutes}$$

With $n = 20$, the rule for locating the median in the list gives

$$\text{location of } M = \frac{n+1}{2} = \frac{21}{2} = 10.5$$

The location 10.5 means “halfway between the 10th and 11th observations in the ordered list.” That agrees with what we found by eye. ■

COMPARING THE MEAN AND THE MEDIAN

Examples 2.1 and 2.2 illustrate an important difference between the mean and the median. The median travel time (the midpoint of the distribution) is 20 minutes. The mean travel time is higher, 22.5 minutes. The mean is pulled toward the right tail of this right-skewed distribution. The median, unlike the mean, is *resistant*. If the longest travel time were 600 minutes rather than 60 minutes, the

mean would increase to more than 58 minutes but the median would not change at all. The outlier just counts as one observation above the center, no matter how far above the center it lies. The mean uses the actual value of each observation and so will chase a single large observation upward. The *Mean and Median* applet is an excellent way to compare the resistance of M and \bar{x} .



COMPARING THE MEAN AND THE MEDIAN

The mean and median of a roughly symmetric distribution are close together. If the distribution is exactly symmetric, the mean and median are exactly the same. In a skewed distribution, the mean is usually farther out in the long tail than is the median.²

Many economic variables have distributions that are skewed to the right. For example, the median endowment of colleges and universities in the United States and Canada in 2009 was about \$67 million—but the mean endowment was almost \$371 million. Most institutions have modest endowments, but a few are very wealthy. Harvard's endowment was over \$35 billion.³ The few wealthy institutions pull the mean up but do not affect the median. Reports about incomes and other strongly skewed distributions usually give the median (“midpoint”) rather than the mean (“arithmetic average”). However, a county that is about to impose a tax of 1% on the incomes of its residents cares about the mean income, not the median. The tax revenue will be 1% of total income, and the total is the mean times the number of residents. The mean and median measure center in different ways, and both are useful. Don't confuse the “average” value of a variable (the mean) with its “typical” value, which we might describe by the median.



APPLY YOUR KNOWLEDGE

- 2.3 New York travel times.** Find the mean of the travel times to work for the 20 New York workers in Example 2.3. Compare the mean and median for these data. What general fact does your comparison illustrate?  

- 2.4 New-house prices.** The mean and median sales prices of new homes sold in the United States in November 2010 were \$213,000 and \$268,700.⁴ Which of these numbers is the mean and which is the median? Explain how you know.

- 2.5 Carbon dioxide emissions.** Table 1.6 (page 33) gives the 2007 carbon dioxide (CO_2) emissions per person for countries with populations of at least 30 million. Find the mean and the median for these data. Make a histogram of the data. What features of the distribution explain why the mean is larger than the median?  



Jose Antonio Sanchez/PhotoAlamy

MEASURING SPREAD: The Quartiles

The mean and median provide two different measures of the center of a distribution. But a measure of center alone can be misleading. The Census Bureau reports that in 2009 the median income of American households was \$49,777. Half of all

households had incomes below \$49,777, and half had higher incomes. The mean was much higher, \$67,976, because the distribution of incomes is skewed to the right. But the median and mean don't tell the whole story. The bottom 10% of households had incomes less than \$12,120, and households in the top 5% took in more than \$180,001.⁵ We are interested in the spread or variability of incomes as well as their center. *The simplest useful numerical description of a distribution requires both a measure of center and a measure of spread.*

 One way to measure spread is to give the smallest and largest observations. For example, the travel times of our 15 North Carolina workers range from 5 minutes to 60 minutes. These single observations show the full spread of the data, but they may be outliers. We can improve our description of spread by also looking at the spread of the middle half of the data. The *quartiles* mark out the middle half. Count up the ordered list of observations, starting from the smallest. The *first quartile* lies one-quarter of the way up the list. The *third quartile* lies three-quarters of the way up the list. In other words, the first quartile is larger than 25% of the observations, and the third quartile is larger than 75% of the observations. The second quartile is the median, which is larger than 50% of the observations. That is the idea of quartiles. We need a rule to make the idea exact. The rule for calculating the quartiles uses the rule for the median.

THE QUARTILES Q_1 AND Q_3

To calculate the quartiles:

1. Arrange the observations in increasing order and locate the median M in the ordered list of observations.
2. The **first quartile** Q_1 is the median of the observations whose position in the ordered list is to the left of the location of the overall median.
3. The **third quartile** Q_3 is the median of the observations whose position in the ordered list is to the right of the location of the overall median.

Here are examples that show how the rules for the quartiles work for both odd and even numbers of observations.

EXAMPLE 2.4 Finding the quartiles: odd n

Our North Carolina sample of 15 workers' travel times, arranged in increasing order, is

5 10 10 10 10 12 15 20 20 25 30 30 40 40 60

There is an odd number of observations, so the median is the middle one, the bold 20 in the list. The first quartile is the median of the 7 observations to the left of the median. This is the 4th of these 7 observations, so $Q_1 = 10$ minutes. If you want, you can use the rule for the location of the median with $n = 7$:

$$\text{location of } Q_1 = \frac{n+1}{2} = \frac{7+1}{2} = 4$$

The third quartile is the median of the 7 observations to the right of the median, $Q_3 = 30$ minutes. When there is an odd number of observations, leave out the overall median when you locate the quartiles in the ordered list.

The quartiles are resistant because they are not affected by a few extreme observations. For example, Q_3 would still be 30 if the outlier were 600 rather than 60. ■

EXAMPLE 2.5 Finding the quartiles: even n

Here are the travel times to work of the 20 New York workers from Example 2.3, arranged in increasing order:

5 10 10 15 15 15 15 20 20 20 | 25 30 30 40 40 45 60 60 65 85

There is an even number of observations, so the median lies midway between the middle pair, the 10th and 11th in the list. Its value is $M = 22.5$ minutes. We have marked the location of the median by |. The first quartile is the median of the first 10 observations, because these are the observations to the left of the location of the median. Check that $Q_1 = 15$ minutes and $Q_3 = 42.5$ minutes. When the number of observations is even, include all the observations when you locate the quartiles. ■

Be careful when, as in these examples, several observations take the same numerical value. Write down all of the observations, arrange them in order, and apply the rules just as if they all had distinct values.

THE FIVE-NUMBER SUMMARY AND BOXPLOTS

The smallest and largest observations tell us little about the distribution as a whole, but they give information about the tails of the distribution that is missing if we know only the median and the quartiles. To get a quick summary of both center and spread, combine all five numbers.

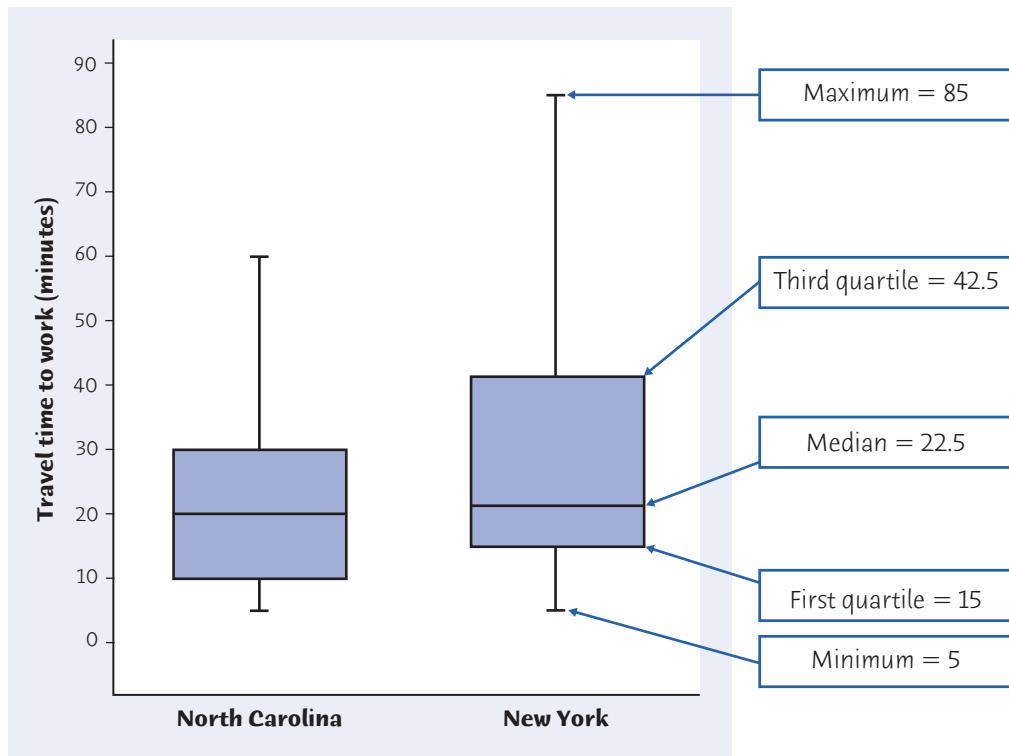
THE FIVE-NUMBER SUMMARY

The **five-number summary** of a distribution consists of the smallest observation, the first quartile, the median, the third quartile, and the largest observation, written in order from smallest to largest. In symbols, the five-number summary is

Minimum Q_1 M Q_3 Maximum

These five numbers offer a reasonably complete description of center and spread. The five-number summaries of travel times to work from Examples 2.4 and 2.5 are

North Carolina	5	10	20	30	60
New York	5	15	22.5	42.5	85

**FIGURE 2.1**

Boxplots comparing the travel times to work of samples of workers in North Carolina and New York.

The five-number summary of a distribution leads to a new graph, the *boxplot*. Figure 2.1 shows boxplots comparing travel times to work in North Carolina and New York.

BOXPLOT

A boxplot is a graph of the five-number summary.

- A central box spans the quartiles Q_1 and Q_3 .
- A line in the box marks the median M.
- Lines extend from the box out to the smallest and largest observations.

Because boxplots show less detail than histograms or stemplots, they are best used for side-by-side comparison of more than one distribution, as in Figure 2.1. Be sure to include a numerical scale in the graph. When you look at a boxplot, first locate the median, which marks the center of the distribution. Then look at the spread. The span of the central box shows the spread of the middle half of the data, and the extremes (the smallest and largest observations) show the spread of the entire data set. We see from Figure 2.1 that travel times to work are in general a bit longer in New York than in North Carolina. The median, both

quartiles, and the maximum are all larger in New York. New York travel times are also more variable, as shown by the span of the box and the spread between the extremes. Note that the boxes with arrows in Figure 2.1 that indicate the location of the five-number summary are *not* part of the boxplot, but are included purely for illustration.

Finally, the New York data are more strongly right-skewed. In a symmetric distribution, the first and third quartiles are equally distant from the median. In most distributions that are skewed to the right, on the other hand, the third quartile will be farther above the median than the first quartile is below it. The extremes behave the same way, but remember that they are just single observations and may say little about the distribution as a whole.

APPLY YOUR KNOWLEDGE

2.6 The Pittsburgh Steelers. The 2010 roster of the Pittsburgh Steelers professional football team included 7 defensive linemen and 9 offensive linemen. The weights in pounds of the defensive linemen were  STEELERS

305 325 305 300 285 280 298

and the weights of the offensive linemen were

338 324 325 304 344 315 304 319 318

- Make a stemplot of the weights of the defensive linemen and find the five-number summary.
- Make a stemplot of the weights of the offensive linemen and find the five-number summary.
- Does either group contain one or more clear outliers? Which group of players tends to be heavier?

2.7 Fuel economy for midsize cars. The Department of Energy provides fuel economy ratings for all cars and light trucks sold in the United States. Here are the estimated miles per gallon for city driving for the 129 cars classified as midsize in 2010, arranged in increasing order:⁶  MIDSIZECARS

9	10	10	11	11	11	12	13	14	14	15	15	15	15	15	15
15	15	16	16	16	16	16	16	16	16	16	16	16	16	16	16
16	16	16	16	17	17	17	17	17	17	17	17	17	17	17	17
17	17	17	17	18	18	18	18	18	18	18	18	18	18	18	18
18	18	18	18	18	18	18	18	18	18	18	18	19	19	19	19
19	19	19	19	19	19	19	19	20	20	20	21	21	21	21	21
21	22	22	22	22	22	22	22	22	22	22	22	22	22	22	22
22	22	22	23	23	23	24	24	24	25	26	26	26	26	26	26
26	26	26	28	33	35	41	41	51							

- Give the five-number summary of this distribution.
- Draw a boxplot of these data. What is the shape of the distribution shown by the boxplot? Which features of the boxplot led you to this conclusion? Are any observations unusually small or large?



AP Photo/Greg Trott



How much is that house worth?

The town of Manhattan, Kansas, is sometimes called “the Little Apple” to distinguish it from that other Manhattan, “the Big Apple.” A few years ago, a house there appeared in the county appraiser’s records valued at \$200,059,000. That would be quite a house even on Manhattan Island. As you might guess, the entry was wrong: the true value was \$59,500. But before the error was discovered, the county, the city, and the school board had based their budgets on the total appraised value of real estate, which the one outlier jacked up by 6.5%. It can pay to spot outliers before you trust your data.

SPOTTING SUSPECTED OUTLIERS*

Look again at the stemplot of travel times to work in New York in Example 2.3. The five-number summary for this distribution is

$$5 \quad 15 \quad 22.5 \quad 42.5 \quad 85$$

How shall we describe the spread of this distribution? The smallest and largest observations are extremes that don’t describe the spread of the majority of the data. The distance between the quartiles (the range of the center half of the data) is a more resistant measure of spread. This distance is called the *interquartile range*.

THE INTERQUARTILE RANGE IQR

The **interquartile range IQR** is the distance between the first and third quartiles,

$$IQR = Q_3 - Q_1$$



For our data on New York travel times, $IQR = 42.5 - 15 = 27.5$ minutes. However, *no single numerical measure of spread, such as IQR, is very useful for describing skewed distributions*. The two sides of a skewed distribution have different spreads, so one number can’t summarize them. That’s why we give the full five-number summary. The interquartile range is mainly used as the basis for a rule of thumb for identifying suspected outliers. In some software, suspected outliers are identified in a boxplot with a special plotting symbol such as *.

THE $1.5 \times IQR$ RULE FOR OUTLIERS

Call an observation a suspected outlier if it falls more than $1.5 \times IQR$ above the third quartile or below the first quartile.

EXAMPLE 2.6 Using the $1.5 \times IQR$ rule

For the New York travel time data, $IQR = 27.5$ and

$$1.5 \times IQR = 1.5 \times 27.5 = 41.25$$

Any values not falling between

$$Q_1 - (1.5 \times IQR) = 15.0 - 41.25 = -26.25 \quad \text{and} \\ Q_3 + (1.5 \times IQR) = 42.5 + 41.25 = 83.75$$

are flagged as suspected outliers. Look again at the stemplot in Example 2.3: the only suspected outlier is the longest travel time, 85 minutes. The $1.5 \times IQR$ rule suggests that the three next-longest travel times (60 and 65 minutes) are just part of the long right tail of this skewed distribution. ■

*This short section is optional.

The $1.5 \times IQR$ rule is not a replacement for looking at the data. It is most useful when large volumes of data are scanned automatically.

APPLY YOUR KNOWLEDGE

- 2.8 Travel time to work.** In Example 2.1, we noted the influence of one long travel time of 60 minutes in our sample of 15 North Carolina workers. Does the $1.5 \times IQR$ rule identify this travel time as a suspected outlier?
- 2.9 Fuel economy for midsize cars.** Exercise 2.7 gives the estimated miles per gallon (mpg) for city driving for the 129 cars classified as midsize in 2010. In that exercise we noted that several of the mpg values were unusually large. Which of these are suspected outliers by the $1.5 \times IQR$ rule? While outliers can be produced by errors or incorrectly recorded observations, they are often observations that differ from the others in some particular way. In this case, the cars producing the high outliers share a common feature. What do you think that is?  MIDSIZECARS

MEASURING SPREAD: The Standard Deviation

The five-number summary is not the most common numerical description of a distribution. That distinction belongs to the combination of the mean to measure center and the *standard deviation* to measure spread. The standard deviation and its close relative, the *variance*, measure spread by looking at how far the observations are from their mean.

THE STANDARD DEVIATION s

The **variance** s^2 of a set of observations is an average of the squares of the deviations of the observations from their mean. In symbols, the variance of n observations x_1, x_2, \dots, x_n is

$$s^2 = \frac{(x_1 - \bar{x})^2 + (x_2 - \bar{x})^2 + \cdots + (x_n - \bar{x})^2}{n - 1}$$

or, more compactly,

$$s^2 = \frac{1}{n - 1} \sum (x_i - \bar{x})^2$$

The **standard deviation** s is the square root of the variance s^2 :

$$s = \sqrt{\frac{1}{n - 1} \sum (x_i - \bar{x})^2}$$

In practice, use software or your calculator to obtain the standard deviation from keyed-in data. Doing an example step-by-step will help you understand how the variance and standard deviation work, however.



SATCR

EXAMPLE 2.7 Calculating the standard deviation

Georgia Southern University had 2417 students with regular admission in their freshman class of 2010. For each student, data are available on their SAT and ACT scores (if taken), high school GPA, and the college within the university to which they were admitted.⁷ In Exercise 3.49, the full data set for the SAT Critical Reading scores will be examined. Here are the first five observations from that data set:

650 490 580 450 570

We will compute \bar{x} and s for these students. First find the mean:

$$\begin{aligned}\bar{x} &= \frac{650 + 490 + 580 + 450 + 570}{5} \\ &= \frac{2740}{5} = 548\end{aligned}$$

Figure 2.2 displays the data as points above the number line, with their mean marked by an asterisk (*). The arrows mark two of the deviations from the mean. The deviations show how spread out the data are about their mean. They are the starting point for calculating the variance and the standard deviation.

Observations x_i	Deviations $x_i - \bar{x}$	Squared deviations $(x_i - \bar{x})^2$
650	$650 - 548 = 102$	$102^2 = 10,404$
490	$490 - 548 = -58$	$(-58)^2 = 3,364$
580	$580 - 548 = 32$	$32^2 = 1,024$
450	$450 - 548 = -98$	$(-98)^2 = 9,604$
570	$570 - 548 = 22$	$22^2 = 484$
	sum = 0	sum = 24,880

The variance is the sum of the squared deviations divided by one less than the number of observations:

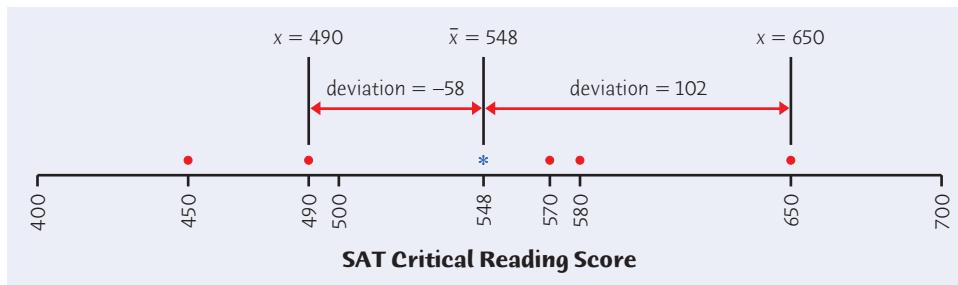
$$s^2 = \frac{1}{n-1} \sum (x_i - \bar{x})^2 = \frac{24,880}{4} = 6220$$

The standard deviation is the square root of the variance:

$$s = \sqrt{6220} = 78.87 \blacksquare$$

degrees of freedom

Notice that the “average” in the variance s^2 divides the sum by one fewer than the number of observations, that is, $n - 1$ rather than n . The reason is that the deviations $x_i - \bar{x}$ always sum to exactly 0, so that knowing $n - 1$ of them determines the last one. Only $n - 1$ of the squared deviations can vary freely, and we average by dividing the total by $n - 1$. The number $n - 1$ is called the **degrees of freedom** of the variance or standard deviation. Some calculators offer a choice between dividing by n and dividing by $n - 1$, so be sure to use $n - 1$.

**FIGURE 2.2**

SAT Critical Reading scores for five students, with their mean (*) and the deviations of two observations from the mean shown, for Example 2.7.

More important than the details of hand calculation are the properties that determine the usefulness of the standard deviation:

- s measures *spread about the mean* and should be used only when the mean is chosen as the measure of center.
- s is *always zero or greater than zero*. $s = 0$ only when there is no spread. This happens only when all observations have the same value. Otherwise, $s > 0$. As the observations become more spread out about their mean, s gets larger.
- s has the *same units of measurement as the original observations*. For example, if you measure weight in kilograms, both the mean \bar{x} and the standard deviation s are also in kilograms. This is one reason to prefer s to the variance s^2 , which would be in squared kilograms.
- Like the mean \bar{x} , s is *not resistant*. A few outliers can make s very large.

The use of squared deviations renders s even more sensitive than \bar{x} to a few extreme observations. For example, the standard deviation of the travel times for the 15 North Carolina workers in Example 2.1 is 15.23 minutes. (Use your calculator or software to verify this.) If we omit the high outlier, the standard deviation drops to 11.56 minutes.



If you feel that the importance of the standard deviation is not yet clear, you are right. We will see in Chapter 3 that the standard deviation is the natural measure of spread for a very important class of symmetric distributions, the Normal distributions. The usefulness of many statistical procedures is tied to distributions of particular shapes. This is certainly true of the standard deviation.

CHOOSING MEASURES OF CENTER AND SPREAD

We now have a choice between two descriptions of the center and spread of a distribution: the five-number summary, or \bar{x} and s . Because \bar{x} and s are sensitive to extreme observations, they can be misleading when a distribution is strongly skewed or has outliers. In fact, because the two sides of a skewed distribution have different spreads, no single number describes the spread well. The five-number summary, with its two quartiles and two extremes, does a better job.

CHOOSING A SUMMARY

The five-number summary is usually better than the mean and standard deviation for describing a skewed distribution or a distribution with strong outliers. Use \bar{x} and s only for reasonably symmetric distributions that are free of outliers.

Outliers can greatly affect the values of the mean \bar{x} and the standard deviation s , the most common measures of center and spread. Many more elaborate statistical procedures also can't be trusted when outliers are present. *Whenever you find outliers in your data, try to find an explanation for them.* Sometimes the explanation is as simple as a typing error, such as typing 10.1 as 101. Sometimes a measuring device broke down or a subject gave a frivolous response, like the student in a class survey who claimed to study 30,000 minutes per night. (Yes, that really happened.) In all these cases, you can simply remove the outlier from your data. When outliers are "real data," like the long travel times of some New York workers, you should choose statistical methods that are not greatly disturbed by the outliers. For example, use the five-number summary rather than \bar{x} and s to describe a distribution with extreme outliers. We will meet other examples later in the book.

 Remember that a graph gives the best overall picture of a distribution. If data have been entered into a calculator or statistical program, it is very simple and quick to create several graphs to see all the different features of a distribution. Numerical measures of center and spread report specific facts about a distribution, but they do not describe its entire shape. Numerical summaries do not disclose the presence of multiple peaks or clusters, for example. Exercise 2.11 shows how misleading numerical summaries can be. **Always plot your data.**

APPLY YOUR KNOWLEDGE



T. Jacobs/Custom Medical Stock Photo/Newscom

2.10 \bar{x} and s by hand. Radon is a naturally occurring gas and is the second leading cause of lung cancer in the United States.⁸ It comes from the natural breakdown of uranium in the soil and enters buildings through cracks and other holes in the foundations. Found throughout the United States, levels vary considerably from state to state. There are several methods to reduce the levels of radon in your home, and the Environmental Protection Agency recommends using one of these if the measured level in your home is above 4 picocuries per liter. Four readings from Franklin County, Ohio, where the county average is 9.32 picocuries per liter, were 5.2, 13.8, 8.6, and 16.8.

- Find the mean step-by-step. That is, find the sum of the 4 observations and divide by 4.
- Find the standard deviation step-by-step. That is, find the deviation of each observation from the mean, square the deviations, then obtain the variance and the standard deviation. Example 2.7 shows the method.

- (c) Now enter the data into your calculator and use the mean and standard deviation buttons to obtain \bar{x} and s . Do the results agree with your hand calculations?

2.11 \bar{x} and s are not enough. The mean \bar{x} and standard deviation s measure center and spread but are not a complete description of a distribution. Data sets with different shapes can have the same mean and standard deviation. To demonstrate this fact, use your calculator to find \bar{x} and s for these two small data sets. Then make a stemplot of each and comment on the shape of each distribution.  **DATASETS**

Data A	9.14	8.14	8.74	8.77	9.26	8.10	6.13	3.10	9.13	7.26	4.74
Data B	6.58	5.76	7.71	8.84	8.47	7.04	5.25	5.56	7.91	6.89	12.50

2.12 Choose a summary. The shape of a distribution is a rough guide to whether the mean and standard deviation are a helpful summary of center and spread. For which of the following distributions would \bar{x} and s be useful? In each case, give a reason for your decision.

- Percents of high school graduates in the states taking the SAT, Figure 1.8 (page 18)
- Iowa Test scores, Figure 1.7 (page 17)
- New York travel times, Figure 2.1 (page 46)

USING TECHNOLOGY

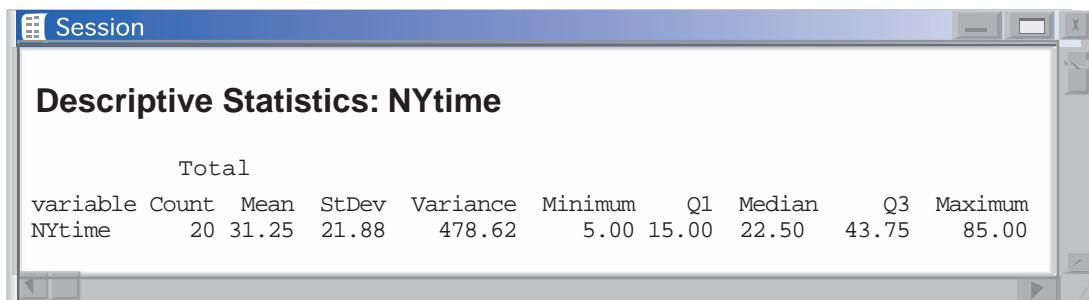
Although a calculator with “two-variable statistics” functions will do the basic calculations we need, more elaborate tools are helpful. Graphing calculators and computer software will do calculations and make graphs as you command, freeing you to concentrate on choosing the right methods and interpreting your results. Figure 2.3 displays output describing the travel times to work of 20 people in New York State (Example 2.3). Can you find \bar{x} , s , and the five-number summary in each output? The big message of this section is: *once you know what to look for, you can read output from any technological tool.*

The displays in Figure 2.3 come from a Texas Instruments graphing calculator, the Minitab and CrunchIt! statistical programs, and the Microsoft Excel spreadsheet program. Minitab allows you to choose what descriptive measures you want, while the descriptive measures in the CrunchIt! output are provided by default. Excel and the calculator give some things we don’t need. Just ignore the extras. Excel’s “Descriptive Statistics” menu item doesn’t give the quartiles. We used the spreadsheet’s separate quartile function to get Q_1 and Q_3 .

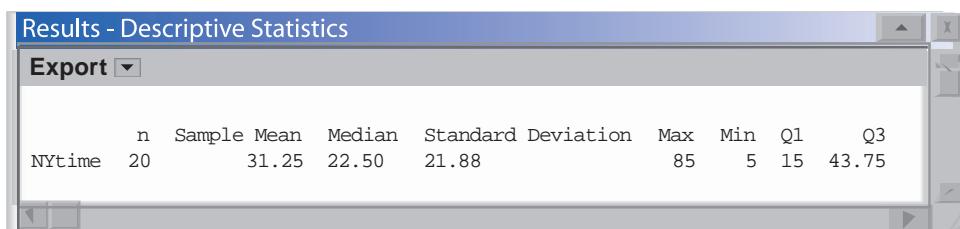
Texas Instruments Graphing Calculator

1-Var Stats
$\bar{x}=31.25$
$\sum x=625$
$\sum x^2=28625$
$Sx=21.8773495$
$sx=21.32348254$
$n=20$
1-Var Stats
$n=20$
$minX=5$
$Q_1=15$
$Med=22.5$
$Q_3=42.5$
$maxX=85$

Minitab



CrunchIt!



Microsoft Excel

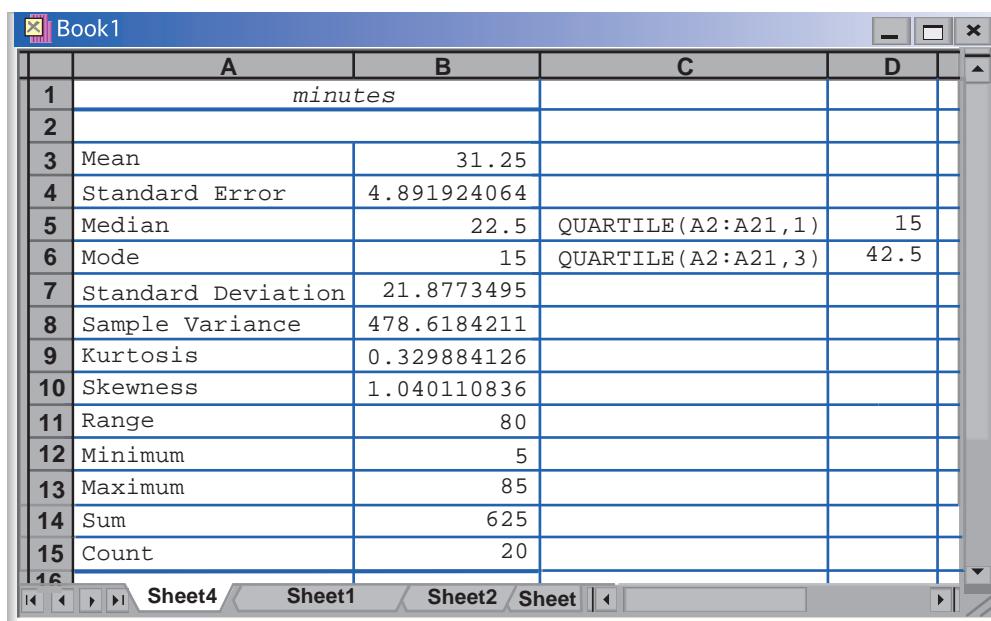


FIGURE 2.3

Output from a graphing calculator, two statistical software packages, and a spreadsheet program describing the data on travel times to work in New York State.

EXAMPLE 2.8 What is the third quartile?

In Example 2.5, we saw that the quartiles of the New York travel times are $Q_1 = 15$ and $Q_3 = 42.5$. Look at the output displays in Figure 2.3. The calculator and Excel agree with our work. Minitab and CrunchIt! say that $Q_3 = 43.75$. What happened? There are several rules for finding the quartiles. Some calculators and software use rules that give results different from ours for some sets of data. This is true of Minitab, CrunchIt!, and also Excel, though Excel agrees with our work in this example. Results from the various rules are always close to each other, so the differences are never important in practice. Our rule is the simplest for hand calculation. ■



ORGANIZING A STATISTICAL PROBLEM

Most of our examples and exercises have aimed to help you learn basic tools (graphs and calculations) for describing and comparing distributions. You have also learned principles that guide use of these tools, such as “start with a graph” and “look for the overall pattern and striking deviations from the pattern.” The data you work with are not just numbers—they describe specific settings such as water depth in the Everglades or travel time to work. Because data come from a specific setting, the final step in examining data is a *conclusion for that setting*. Water depth in the Everglades has a yearly cycle that reflects Florida’s wet and dry seasons. Travel times to work are generally longer in New York than in North Carolina.

As you learn more statistical tools and principles, you will face more complex statistical problems. Although no framework accommodates all the varied issues that arise in applying statistics to real settings, the following four-step thought process gives useful guidance. In particular, the first and last steps emphasize that statistical problems are tied to specific real-world settings and therefore involve more than doing calculations and making graphs.

ORGANIZING A STATISTICAL PROBLEM: A Four-Step Process

STATE: What is the practical question, in the context of the real-world setting?

PLAN: What specific statistical operations does this problem call for?

SOLVE: Make the graphs and carry out the calculations needed for this problem.

CONCLUDE: Give your practical conclusion in the setting of the real-world problem.

To help you master the basics, many exercises will continue to tell you what to do—make a histogram, find the five-number summary, and so on. Real statistical problems don’t come with detailed instructions. From now on, especially in the later chapters of the book, you will meet some exercises that are more realistic. Use the four-step process as a guide to solving and reporting these problems. They are marked with the four-step icon, as the following example illustrates.



EXAMPLE 2.9 Comparing tropical flowers

STATE: Ethan Temeles of Amherst College, with his colleague W. John Kress, studied the relationship between varieties of the tropical flower *Heliconia* on the island of Dominica and the different species of hummingbirds that fertilize the flowers.⁹ Over time, the researchers believe, the lengths of the flowers and the forms of the hummingbirds' beaks have evolved to match each other. If that is true, flower varieties fertilized by different hummingbird species should have distinct distributions of length.

Table 2.1 gives length measurements (in millimeters) for samples of three varieties of *Heliconia*, each fertilized by a different species of hummingbird. Do the three varieties display distinct distributions of length? How do the mean lengths compare?

PLAN: Use graphs and numerical descriptions to describe and compare these three distributions of flower length.

SOLVE: We might use boxplots to compare the distributions, but stemplots preserve more detail and work well for data sets of these sizes. Figure 2.4 displays stemplots with the stems lined up for easy comparison. The lengths have been rounded to the nearest tenth of a millimeter. The *bihai* and red varieties have somewhat skewed distributions, so we might choose to compare the five-number summaries. But because the researchers plan to use \bar{x} and s for further analysis, we instead calculate these measures:

Variety	Mean length	Standard deviation
<i>bihai</i>	47.60	1.213
red	39.71	1.799
yellow	36.18	0.975

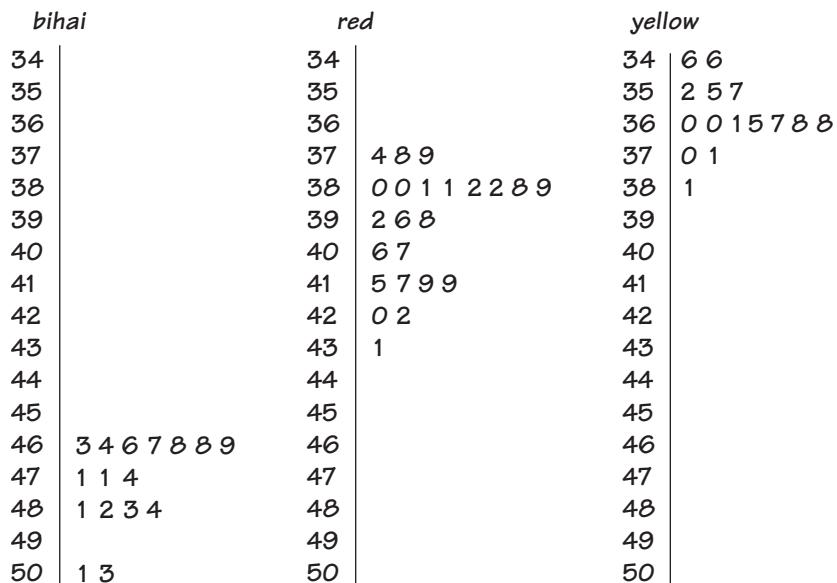
CONCLUDE: The three varieties differ so much in flower length that there is little overlap among them. In particular, the flowers of *bihai* are longer than either red or yellow. The mean lengths are 47.6 mm for *H. bihai*, 39.7 mm for *H. caribaea* red, and 36.2 mm for *H. caribaea* yellow. ■

TABLE 2.1 Flower lengths (millimeters) for three *Heliconia* varieties

<i>H. BIHAI</i>							
47.12	46.75	46.81	47.12	46.67	47.43	46.44	46.64
48.07	48.34	48.15	50.26	50.12	46.34	46.94	48.36
<i>H. CARIBAEA RED</i>							
41.90	42.01	41.93	43.09	41.47	41.69	39.78	40.57
39.63	42.18	40.66	37.87	39.16	37.40	38.20	38.07
38.10	37.97	38.79	38.23	38.87	37.78	38.01	
<i>H. CARIBAEA YELLOW</i>							
36.78	37.02	36.52	36.11	36.03	35.45	38.13	37.10
35.17	36.82	36.66	35.68	36.03	34.57	34.63	



TROPICALFLOWER

**FIGURE 2.4**

Stemplots comparing the distributions of flower lengths from Table 2.1, for Example 2.9. The stems are whole millimeters and the leaves are tenths of a millimeter.

APPLY YOUR KNOWLEDGE

2.13 Logging in the rain forest. “Conservationists have despaired over destruction of tropical rain forest by logging, clearing, and burning.” These words begin a report on a statistical study of the effects of logging in Borneo.¹⁰ Charles Cannon of Duke University and his coworkers compared forest plots that had never been logged (Group 1) with similar plots nearby that had been logged 1 year earlier (Group 2) and 8 years earlier (Group 3). All plots were 0.1 hectare in area. Here are the counts of trees for plots in each group:  **LOGGING**

Group 1	27	22	29	21	19	33	16	20	24	27	28	19
Group 2	12	12	15	9	20	18	17	14	14	2	17	19
Group 3	18	4	22	15	18	19	22	12	12			

To what extent has logging affected the count of trees? Follow the four-step process in reporting your work.

2.14 Diplomatic scofflaws. Until Congress allowed some enforcement in 2002, the thousands of foreign diplomats in New York City could freely violate parking laws. Two economists looked at the number of unpaid parking tickets per diplomat over a five-year period ending when enforcement reduced the problem.¹¹ They concluded that large numbers of unpaid tickets indicated a “culture of corruption” in a country and lined up well with more elaborate measures of corruption. The data set for 145 countries is too large to print here, but look at the data file on the text Web site and CD. The first 32 countries in the list (Australia to Trinidad and Tobago) are classified by the World Bank as “developed.” The remaining countries (Albania to Zimbabwe) are “developing.” The World Bank classification is based only on national income and does not take into account measures of social development.  **SCOFFLAWS**



Give a full description of the distribution of unpaid tickets for both groups of countries and identify any high outliers. Compare the two groups. Does national income alone do a good job of distinguishing countries whose diplomats do and do not obey parking laws?

CHAPTER 2 SUMMARY

CHAPTER SPECIFICS

- A numerical summary of a distribution should report at least its **center** and its **spread** or **variability**.
- The **mean** \bar{x} and the **median** M describe the center of a distribution in different ways. The mean is the arithmetic average of the observations, and the median is the midpoint of the values.
- When you use the median to indicate the center of the distribution, describe its spread by giving the **quartiles**. The **first quartile** Q_1 has one-fourth of the observations below it, and the **third quartile** Q_3 has three-fourths of the observations below it.
- The **five-number summary** consisting of the median, the quartiles, and the smallest and largest individual observations provides a quick overall description of a distribution. The median describes the center, and the quartiles and extremes show the spread.
- **Boxplots** based on the five-number summary are useful for comparing several distributions. The box spans the quartiles and shows the spread of the central half of the distribution. The median is marked within the box. Lines extend from the box to the extremes and show the full spread of the data.
- The **variance** s^2 and especially its square root, the **standard deviation** s , are common measures of spread about the mean as center. The standard deviation s is zero when there is no spread and gets larger as the spread increases.
- A **resistant measure** of any aspect of a distribution is relatively unaffected by changes in the numerical value of a small proportion of the total number of observations, no matter how large these changes are. The median and quartiles are resistant, but the mean and the standard deviation are not.
- The mean and standard deviation are good descriptions for symmetric distributions without outliers. They are most useful for the Normal distributions introduced in the next chapter. The five-number summary is a better description for skewed distributions.
- Numerical summaries do not fully describe the shape of a distribution. Always plot your data.
- A statistical problem has a real-world setting. You can organize many problems using the following four steps: **state**, **plan**, **solve**, and **conclude**.

LINK IT

In this chapter we have continued our study of exploratory data analysis. Graphs are an important visual tool for organizing and identifying patterns in data. They give a fairly complete description of a distribution, although for many problems the important

information in your data can be described by a few numbers. These numerical summaries can be useful for describing a single distribution as well as for comparing the distributions from several groups of observations.

Two important features of a distribution are the center and the spread. For distributions that are approximately symmetric without outliers, the mean and standard deviation are important numeric summaries for describing and comparing distributions. But if the distribution is not symmetric and/or has outliers, the five-number summary often provides a better description.

The boxplot gives a picture of the five-number summary that is useful for a simple comparison of several distributions. Remember that the boxplot is based only on the five-number summary and does not have any information beyond these five numbers. Certain features of a distribution that are revealed in histograms and stemplots will not be evident from a boxplot alone. These include gaps in the data and the presence of several peaks. You must be careful when reducing a distribution to a few numbers to make sure that important information has not been lost in the process.

CHECK YOUR SKILLS

2.15 The respiratory system can be a limiting factor in maximal exercise performance. Researchers from the United Kingdom studied the effect of two breathing frequencies on both performance times and several physiological parameters in swimming.¹² Subjects were 10 male collegiate swimmers. Here are their times in seconds to swim 200 meters at 90% of race pace when breathing every second stroke in front-crawl swimming:  SWIMTIMES

151.6	165.1	159.2	163.5	174.8
173.2	177.6	174.3	164.1	171.4

The mean of these data is

- (a) 165.10. (b) 167.48. (c) 168.25.

2.16 The median of the data in Exercise 2.15 is

- (a) 167.48. (b) 168.25. (c) 174.00.

2.17 The five-number summary of the data in Exercise 2.15 is

- (a) 151.6, 159.2, 167.48, 174.8, 177.6.
(b) 151.6, 163.5, 168.25, 174.3, 177.6.
(c) 151.6, 159.2, 168.25, 174.8, 177.6.

2.18 If a distribution is skewed to the right,

- (a) the mean is less than the median.
(b) the mean and median are equal.
(c) the mean is greater than the median.

2.19 What percent of the observations in a distribution lie between the first quartile and the third quartile?

- (a) 25% (b) 50% (c) 75%

2.20 To make a boxplot of a distribution, you must know

- (a) all of the individual observations.
(b) the mean and the standard deviation.
(c) the five-number summary.

2.21 The standard deviation of the 10 swim times in Exercise 2.15 (use your calculator) is about

- (a) 7.4. (b) 7.8. (c) 8.2.

2.22 What are all the values that a standard deviation s can possibly take?

- (a) $0 \leq s$ (b) $0 \leq s \leq 1$ (c) $-1 \leq s \leq 1$

2.23 The correct units for the standard deviation in Exercise 2.21 are

- (a) no units—it's just a number.
(b) seconds.
(c) seconds squared.

2.24 Which of the following is least affected if an extreme high outlier is added to your data?

- (a) The median
(b) The mean
(c) The standard deviation

CHAPTER 2 EXERCISES

2.25 Incomes of college grads. According to the Census Bureau's 2010 Current Population Survey, the mean and median 2009 income of people at least 25 years old who had a bachelor's degree but no higher degree were \$46,931 and \$58,762. Which of these numbers is the mean and which is the median? Explain your reasoning.

2.26 Saving for retirement. Retirement seems a long way off and we need money now, so saving for retirement is hard. Once every three years, the Board of Governors of the Federal Reserve System collects data on household assets and liabilities through the Survey of Consumer Finances (SCF). The most recent such survey was conducted in 2007, and the survey results were released to the public in April 2009. The survey presents data on household ownership of, and balances in, retirement savings accounts. Only 53.6% of households own retirement accounts. The mean value per household is \$148,579, but the median value is just \$45,000. For households in which the head of household is under 35, 42.6% own retirement accounts, the mean is \$25,279, and the median is \$9600.¹³ What explains the differences between the two measures of center, both for all households and for the under-35 age group?

2.27 University endowments. The National Association of College and University Business Officers collects data on college endowments. In 2009, 842 colleges and universities reported the value of their endowments. When the endowment values are arranged in order, what are the locations of the median and the quartiles in this ordered list?

2.28 Pulling wood apart. Example 1.9 (page 21) gives the breaking strengths of 20 pieces of Douglas fir. 

(a) Give the five-number summary of the distribution of breaking strengths. (The stemplot, Figure 1.11, helps because it arranges the data in order, but you should use the unrounded values in numerical work.)

(b) The stemplot shows that the distribution is skewed to the left. Does the five-number summary show the skew? Remember that only a graph gives a clear picture of the shape of a distribution.

2.29 Comparing tropical flowers. An alternative presentation of the flower length data in Table 2.1 reports the five-number summary and uses boxplots to display the distributions. Do this. Do the boxplots fail to reveal any important information visible in the stemplots in Figure 2.4? 

2.30 How much fruit do adolescent girls eat? Figure 1.14 (page 30) is a histogram of the number of servings of fruit per day claimed by 74 seventeen-year-old girls.

(a) With a little care, you can find the median and the quartiles from the histogram. What are these numbers? How did you find them?

(b) With a little care, you can also find the mean number of servings of fruit claimed per day. First use the information in the histogram to compute the sum of the 74 observations, and then use this to compute the mean. What is the relationship between the mean and median? Is this what you expected?

2.31 Guinea pig survival times. Here are the survival times in days of 72 guinea pigs after they were injected with infectious bacteria in a medical experiment.¹⁴ Survival times, whether of machines under stress or cancer patients after treatment, usually have distributions that are skewed to the right.  GUINEAPIGS

43	45	53	56	56	57	58	66	67	73
74	79	80	80	81	81	81	82	83	83
84	88	89	91	91	92	92	97	99	99
100	100	101	102	102	102	103	104	107	108
109	113	114	118	121	123	126	128	137	138
139	144	145	147	156	162	174	178	179	184
191	198	211	214	243	249	329	380	403	511
522	598								

(a) Graph the distribution and describe its main features. Does it show the expected right-skew?

(b) Which numerical summary would you choose for these data? Calculate your chosen summary. How does it reflect the skewness of the distribution?

2.32 Weight of newborns. Page 61 gives the distribution of the weight at birth for all babies born in the United States in 2008:¹⁵



PhotoDisc Red/Getty Images

Weight (grams)	Count	Weight (grams)	Count
Less than 500	6,581	3,000 to 3,499	1,663,512
500 to 999	23,292	3,500 to 3,999	1,120,642
1,000 to 1,499	31,900	4,000 to 4,499	280,270
1,500 to 1,999	67,140	4,500 to 4,999	39,109
2,000 to 2,499	218,296	5,000 to 5,499	4,443
2,500 to 2,999	788,148		

- (a) For comparison with other years and with other countries, we prefer a histogram of the percents in each weight class rather than the counts. Explain why.
- (b) How many babies were there?
- (c) Make a histogram of the distribution, using percents on the vertical scale.
- (d) What are the locations of the median and quartiles in the ordered list of all birth weights? In which weight classes do the median and quartiles fall?

2.33 More on study times. In Exercise 1.38 (page 34) you examined the nightly study time claimed by first-year college men and women. The most common methods for formal comparison of two groups use \bar{x} and s to summarize the data. 

- (a) What kinds of distributions are best summarized by \bar{x} and s ? Do you think these summary measures are appropriate in this case?
- (b) One student in each group claimed to study at least 300 minutes (five hours) per night. How much does removing these observations change \bar{x} and s for each group? You will need to compute \bar{x} and s for each group, both with and without the high outlier.

2.34 Making resistance visible. In the *Mean and Median* applet, place three observations on the line by clicking below it: two close together near the center of the line and one somewhat to the right of these two.

- (a) Pull the single rightmost observation out to the right. (Place the cursor on the point, hold down a mouse button, and drag the point.) How does the mean behave? How does the median behave? Explain briefly why each measure acts as it does.
- (b) Now drag the single rightmost point to the left as far as you can. What happens to the mean? What happens to the median as you drag this point past the other two (watch carefully)?

2.35 Behavior of the median. Place five observations on the line in the *Mean and Median* applet by clicking below it.

(a) Add one additional observation *without changing the median*. Where is your new point?

(b) Use the applet to convince yourself that when you add yet another observation (there are now seven in all), the median does not change no matter where you put the seventh point. Explain why this must be true.

2.36 Never on Sunday: also in Canada? Exercise 1.5 (page 11) gives the number of births in the United States on each day of the week during an entire year. The boxplots in Figure 2.5 (page 62) are based on more detailed data from Toronto, Canada: the number of births on each of the 365 days in a year, grouped by day of the week.¹⁶ Based on these plots, compare the day-of-the-week distributions using shape, center, and spread. Summarize your findings.

2.37 Thinking about means. Table 1.1 (page 12) gives the percent of foreign-born residents in each of the states. For the nation as a whole, 12.5% of residents are foreign-born. Find the mean of the 51 entries in Table 1.1. It is not 12.5%. Explain carefully why this happens. (*Hint:* The states with the largest populations are California, Texas, New York, and Florida. Look at their entries in Table 1.1.)

2.38 Thinking about medians. A report says that “the median credit card debt of American households is zero.” We know that many households have large amounts of credit card debt. In fact, the mean household credit card debt is close to \$8000. Explain how the median debt can nonetheless be zero.

2.39 A standard deviation contest. This is a standard deviation contest. You must choose four numbers from the whole numbers 0 to 10, with repeats allowed.

- (a) Choose four numbers that have the smallest possible standard deviation.
- (b) Choose four numbers that have the largest possible standard deviation.
- (c) Is more than one choice possible in either (a) or (b)? Explain.

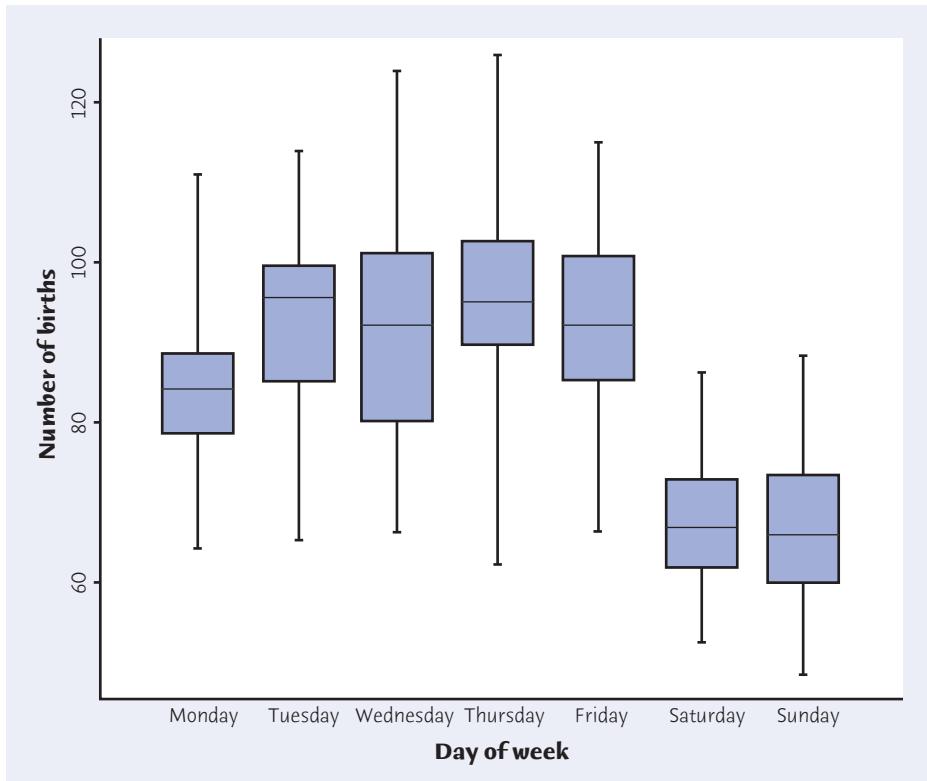
2.40 Test your technology. This exercise requires a calculator with a standard deviation button or statistical software on a computer. The observations

$$10,001 \quad 10,002 \quad 10,003$$

have mean $\bar{x} = 10,002$ and standard deviation $s = 1$. Adding a 0 in the center of each number, the next set becomes

$$100,001 \quad 100,002 \quad 100,003$$

The standard deviation remains $s = 1$ as more 0s are added. Use your calculator or software to find the standard deviation of these numbers, adding extra 0s until you get an incorrect answer. How soon did you go wrong? This demonstrates that

**FIGURE 2.5**

Boxplots of the distributions of numbers of births in Toronto, Canada, on each day of the week during a year, for Exercise 2.36.

calculators and software cannot handle an arbitrary number of digits correctly.

2.41 You create the data. Create a set of 5 positive numbers (repeats allowed) that have median 7 and mean 10. What thought process did you use to create your numbers?

2.42 You create the data. Give an example of a small set of data for which the mean is smaller than the first quartile.

2.43 Adolescent obesity. Adolescent obesity is a serious health risk affecting more than 5 million young people in the United States alone. Laparoscopic adjustable gastric banding has the potential to provide a safe and effective treatment. Fifty adolescents between 14 and 18 years old with a body mass index (BMI) higher than 35 were recruited from the Melbourne, Australia, community for the study.¹⁷ Twenty-five were randomly selected to undergo gastric banding, and the remaining twenty-five were assigned to a supervised lifestyle intervention program involving diet, exercise, and

behavior modification. All subjects were followed for two years. Here are the weight losses in kilograms for the subjects who completed the study:  GASTRICBANDS

Gastric banding						
35.6	81.4	57.6	32.8	31.0	37.6	
36.5	-5.4	27.9	49.0	64.8	39.0	
43.0	33.9	29.7	20.2	15.2	41.7	
53.4	13.4	24.8	19.4	32.3	22.0	

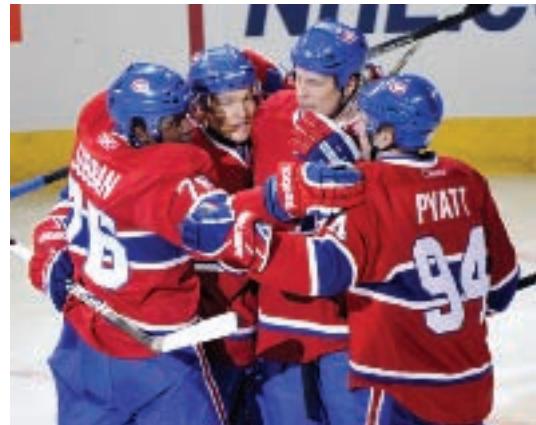
Lifestyle intervention						
6.0	2.0	-3.0	20.6	11.6	15.5	
-17.0	1.4	4.0	-4.6	15.8	34.6	
6.0	-3.1	-4.3	-16.7	-1.8	-12.8	

- (a) In the context of this study, what do the negative values in the data set mean?
- (b) Give a graphical comparison of the weight loss distribution for both groups using side-by-side boxplots. Provide appropriate numerical summaries for the two distributions and identify any high outliers in either group. What can you say about the effects of gastric banding versus lifestyle intervention on weight loss for the subjects in this study?
- (c) The measured variable was weight loss in kilograms. Would two subjects with the same weight loss always have similar benefits from a weight reduction program? Does it depend on their initial weights? Other variables considered in this study were the percent of excess weight lost and the reduction in BMI. Do you see any advantages to either of these variables when comparing weight loss for two groups?
- (d) One subject from the gastric-banding group dropped out of the study and seven subjects from the lifestyle group dropped out. Of the seven dropouts in the lifestyle group, six had gained weight at the time they dropped out. If all subjects had completed the study, how do you think it would have affected the comparison between the two groups?

*Exercises 2.44 to 2.49 ask you to analyze data without having the details outlined for you. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan**, **Solve**, and **Conclude** steps as illustrated in Example 2.9.*

2.44 Athletes' salaries. The Montreal Canadiens were founded in 1909 and are the longest continuously operating professional ice hockey team. They have won 24 Stanley Cups, making them one of the most suc-

cessful professional sports teams of the traditional four major sports of Canada and the United States. Table 2.2 gives the salaries of the 2010–2011 roster.¹⁸ Provide the team owner with a full description of the distribution of salaries and a brief summary of its most important features.  HOCKEYSALARIES



AP Photo/The Canadian Press, Ryan Remiorz

2.45 Returns on stocks. How well have stocks done over the past generation? The Wilshire 5000 index describes the average performance of all U.S. stocks. The average is weighted by the total market value of each company's stock, so think of the index as measuring the performance of the average investor. Page 64 gives the percent returns on the Wilshire 5000 index for the years from 1971 to 2010:  WILSHIRE5000

TABLE 2.2 Salaries for the 2010–2011 Montreal Canadiens

PLAYER	SALARY	PLAYER	SALARY	PLAYER	SALARY
Scott Gomez	\$8,000,000	Andrei Markov	\$5,750,000	Roman Hamrlik	\$5,500,000
Mike Cammalleri	\$5,000,000	Brian Gionta	\$5,000,000	Tomas Plekanec	\$5,000,000
Jaroslav Spacek	\$3,833,000	Andrei Kostitsyn	\$3,250,000	James Wisniewski	\$3,250,000
Carey Price	\$2,500,000	Hal Gill	\$2,250,000	Travis Moen	\$1,500,000
Benoit Pouliot	\$1,350,000	Josh Gorges	\$1,300,000	Alex Auld	\$1,000,000
Max Pacioretty	\$875,000	Lars Eller	\$875,000	P. K. Subban	\$875,000
Yannick Weber	\$637,500	Jeff Halpern	\$600,000	Alexandre Picard	\$600,000
David Desharnais	\$550,000	Mathieu Darche	\$500,000	Tom Pyatt	\$500,000

Wilshire index for the years 1971 to 2010			
Year	Return	Year	Return
1971	16.19	1991	33.58
1972	17.34	1992	9.02
1973	-18.78	1993	10.67
1974	-27.87	1994	0.06
1975	37.38	1995	36.41
1976	26.77	1996	21.56
1977	-2.97	1997	31.48
1978	8.54	1998	24.31
1979	24.40	1999	24.23
1980	33.21	2000	-10.89
1981	-3.98	2001	-10.97
1982	20.43	2002	-20.86
1983	22.71	2003	31.64
1984	3.27	2004	12.48
1985	31.46	2005	6.38
1986	15.61	2006	15.77
1987	1.75	2007	5.62
1988	17.59	2008	-37.23
1989	28.53	2009	28.30
1990	-6.03	2010	17.16

What can you say about the distribution of yearly returns on stocks?

2.46 Do good smells bring good business?

Businesses know that customers often respond to background music. Do they also respond to odors? Nicolas Guéguen and his colleagues studied this question in a small pizza restaurant in France on Saturday evenings in May. On one evening, a relaxing lavender odor was spread through the restaurant; on another evening, a stimulating lemon odor; a third evening served as a control, with no odor. Table 2.3 shows the amounts (in euros) that customers spent on each of these evenings.¹⁹ Compare the three distributions. Were both odors associated with increased customer spending?  ODORS

2.47 Daily activity and obesity.

People gain weight when they take in more energy from food than they expend. Table 2.4 (page 65) compares volunteer subjects who were lean with others who were mildly obese. None of the subjects followed an exercise program. The subjects wore sensors that recorded every move for 10 days. The table shows the average minutes per day spent in activity (standing and walking) and in lying down.²⁰ Compare the distributions of time spent actively for lean and obese subjects and also the distributions of time spent lying down. How does the behavior of lean and mildly obese people differ?  OBESITY

TABLE 2.3 Amount spent (euros) by customers in a restaurant when exposed to odors

NO ODOR									
15.9	18.5	15.9	18.5	18.5	21.9	15.9	15.9	15.9	15.9
15.9	18.5	18.5	18.5	20.5	18.5	18.5	15.9	15.9	15.9
18.5	18.5	15.9	18.5	15.9	18.5	15.9	25.5	12.9	15.9
LEMON ODOR									
18.5	15.9	18.5	18.5	18.5	15.9	18.5	15.9	18.5	18.5
15.9	18.5	21.5	15.9	21.9	15.9	18.5	18.5	18.5	18.5
25.9	15.9	15.9	15.9	18.5	18.5	18.5	18.5		
LAVENDER ODOR									
21.9	18.5	22.3	21.9	18.5	24.9	18.5	22.5	21.5	21.9
21.5	18.5	25.5	18.5	18.5	21.9	18.5	18.5	24.9	21.9
25.9	21.9	18.5	18.5	22.8	18.5	21.9	20.7	21.9	22.5

TABLE 2.4 Time (minutes per day) active and lying down by lean and obese subjects

LEAN SUBJECTS			OBESesubjects		
SUBJECT	STAND/WALK	LIE	SUBJECT	STAND/WALK	LIE
1	511.100	555.500	11	260.244	521.044
2	607.925	450.650	12	464.756	514.931
3	319.212	537.362	13	367.138	563.300
4	584.644	489.269	14	413.667	532.208
5	578.869	514.081	15	347.375	504.931
6	543.388	506.500	16	416.531	448.856
7	677.188	467.700	17	358.650	460.550
8	555.656	567.006	18	267.344	509.981
9	374.831	531.431	19	410.631	448.706
10	504.700	396.962	20	426.356	412.919

2.48 Good weather and tipping. Favorable weather has been shown to be associated with increased tipping. Will just the belief that future weather will be favorable lead to higher tips? The researchers gave 60 index cards to a waitress at an Italian restaurant in New Jersey. Before delivering the bill to each customer, the waitress randomly selected a card and wrote on the bill the same message that was printed on the index card. Twenty of the cards had the message “The weather is supposed to be really good tomorrow. I hope you enjoy the day!” Another 20 cards contained the message “The weather is supposed to be not so good tomorrow. I hope you enjoy the day anyway!” The remaining 20 cards were blank, indicating that the waitress was not supposed to write any message. Choosing a card at random ensured that there was a random assignment of the customers to the three experimental conditions. Here are the tip percents for the three messages:²¹ 

Good weather report:	20.8	18.7	19.9	20.6	22.0	23.4	22.8	24.9	22.2	20.3
	24.9	22.3	27.0	20.4	22.2	24.0	21.2	22.1	22.0	22.7
Bad weather report:	18.0	19.0	19.2	18.8	18.4	19.0	18.5	16.1	16.8	14.0
	17.0	13.6	17.5	19.9	20.2	18.8	18.0	23.2	18.2	19.4
No weather report:	19.9	16.0	15.0	20.1	19.3	19.2	18.0	19.2	21.2	18.8
	18.5	19.3	19.3	19.4	10.8	19.1	19.7	19.8	21.3	20.6

Compare the three distributions. How did the tip percents vary with the weather report information?

2.49 Canadians' earnings in 1901. Table 2.5 (page 66) presents the “earnings from occupation or trade” in the year 1901 of a random sample of those Canadians who did

have earnings in that year. The amounts are in Canadian dollars. If they seem low, remember that a loaf of bread cost 4 cents and a pound of beef 14 cents.²² Of course, Canadians today have much higher incomes, even after adjusting for inflation. Give a complete graphical and numerical description of these data, and briefly describe your findings. 

Exercises 2.50 to 2.53 make use of the optional material on the $1.5 \times \text{IQR}$ rule for suspected outliers.

2.50 Older Americans. In Exercise 1.10 you were asked to use a stemplot to display the distribution of the percents of residents aged 65 and older in the states. Stemplots help you find the five-number summary because they arrange the observations in increasing order. 

- Give the five-number summary of this distribution.
- Use the five-number summary to draw a boxplot of the data. What is the shape of the distribution?

(c) Which observations does the $1.5 \times \text{IQR}$ rule flag as suspected outliers? (The rule flags several observations that are not that extreme. The reason is that the center half of the observations are close together, so that the IQR is small. This example reminds us to use our eyes, not a rule, to spot outliers.)

2.51 Carbon dioxide emissions. Table 1.6 (page 33) gives the 2007 carbon dioxide (CO_2) emissions per person for countries with populations of at least 30 million in that year. A stemplot or histogram shows that the distribution is strongly

TABLE 2.5 Earnings (dollars) of a sample of 200 Canadians in 1901

400	360	400	200	150	175	60	96	600	100
125	525	500	200	175	360	50	300	175	480
300	500	160	960	300	96	500	170	380	600
300	10	250	200	124	350	150	225	600	100
250	250	220	100	200	250	100	256	399	125
100	360	250	600	400	450	150	550	250	240
100	970	240	300	100	400	450	125	250	225
480	350	300	100	144	500	200	150	225	240
250	300	250	150	100	225	96	200	500	20
250	500	100	165	350	200	200	500	300	57
1000	200	800	415	450	190	360	180	1200	100
320	240	100	208	8	142	420	300	700	600
205	150	75	205	300	300	350	150	450	350
450	462	350	1500	120	200	300	225	475	2000
150	325	300	150	720	140	350	150	85	400
184	125	300	500	340	150	160	480	300	376
220	1000	300	2200	350	290	700	120	700	600
375	150	450	360	500	575	500	400	350	180
600	300	1000	600	300	80	1000	300	390	499
210	500	550	450	520	200	300	200	540	1200

skewed to the right. The United States and several other countries appear to be high outliers. 

- (a) Give the five-number summary. Explain why this summary suggests that the distribution is right-skewed.
- (b) Which countries are outliers according to the $1.5 \times IQR$ rule? Make a stemplot of the data or look at your stemplot from Exercise 1.36. Do you agree with the rule's suggestions about which countries are and are not outliers?

2.52 Athletes' salaries. Which members of the Montreal Canadiens (Table 2.2) have salaries that are suspected outliers by the $1.5 \times IQR$ rule? 

2.53 Canadians' earnings in 1901. The Canadians' earnings in Exercise 2.49 are right-skewed, with a few large incomes. Which incomes are suspected outliers by the $1.5 \times IQR$ rule?



EXPLORING THE WEB

2.54 Home run leaders. The three top players on the career home run list are Barry Bonds, Hank Aaron, and Babe Ruth. You can find their home run statistics by going to the Web site www.baseball-reference.com and then clicking on the Players tab at the top of the page. Construct three side-by-side boxplots comparing the yearly home run production of Barry Bonds, Hank Aaron, and Babe Ruth. Describe any differences that you observe. It is worth noting that in his first four seasons, Babe Ruth was primarily a pitcher. If these four seasons are ignored, how does Babe Ruth compare with Barry Bonds and Hank Aaron?

2.55 Crime rates and outliers. The *Statistical Abstract of the United States* is a comprehensive summary of statistics on the social, political, and economic organization of the United States. It can be found at the Web site www.census.gov/compendia/statab/. Go to the section Law Enforcement, Courts and Prisons, and then to the subsection Crimes and Crime Rates. Several tables of data will be available.

- (a) Open the Table on Crime Rates by State and Type for the latest year given. Why do you think they use rates per 100,000 population rather than the number of crimes committed? The District of Columbia is a high outlier in almost every crime category.
- (b) Open the Table on Crime Rates by Type for Selected Large Cities. This table includes the District of Columbia, which is listed as Washington, DC. Without doing any formal calculations, does the District of Columbia look like a high outlier in the table for large cities? Whether or not the District of Columbia is an outlier depends on more than its crime rate. It also depends on the other observations included in the data set. Which data set do you feel is more appropriate for the District of Columbia?
- (c) Using the Table on Crime Rates by State and Type for the latest year given, choose a crime category, and give a full description of its distribution over the 50 states, omitting the District of Columbia. Your description of the distribution should include appropriate graphical and numerical summaries and a brief report describing the main features of the distribution. You can open the data as an Excel file and import it into your statistical software.



The Normal Distributions

Chapter 3

We now have a tool box of graphical and numerical methods for describing distributions. What is more, we have a clear strategy for exploring data on a single quantitative variable.

EXPLORING A DISTRIBUTION

1. Always plot your data: make a graph, usually a histogram or a stemplot.
2. Look for the overall pattern (shape, center, spread) and for striking deviations such as outliers.
3. Calculate a numerical summary to briefly describe center and spread.

In this chapter, we add one more step to this strategy:

4. Sometimes the overall pattern of a large number of observations is so regular that we can describe it by a smooth curve.

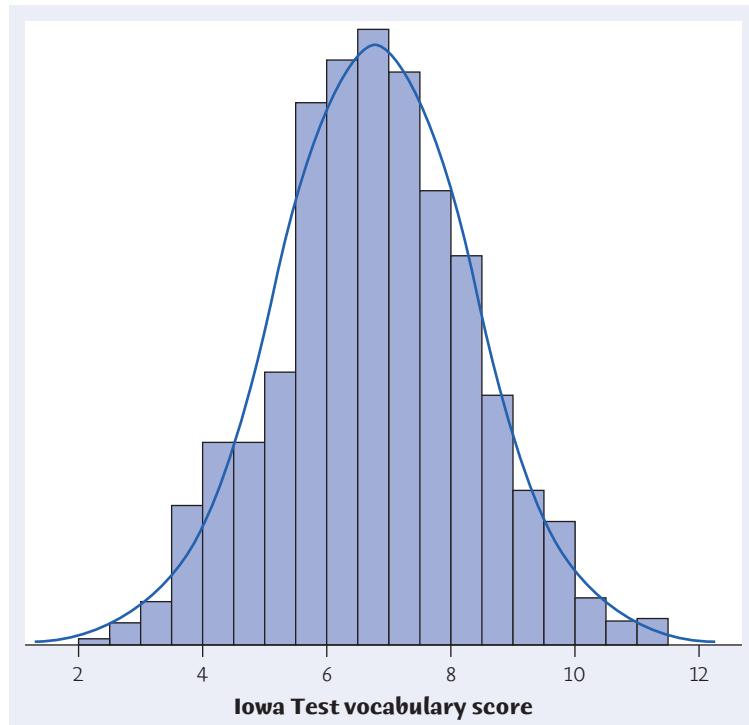
IN THIS CHAPTER WE COVER...

- Density curves
- Describing density curves
- Normal distributions
- The 68–95–99.7 rule
- The standard Normal distribution
- Finding Normal proportions
- Using the standard Normal table
- Finding a value when given a proportion

DENSITY CURVES

Figure 3.1 is a histogram of the scores of all 947 seventh-grade students in Gary, Indiana, on the vocabulary part of the Iowa Test of Basic Skills.¹ Scores of many students on this national test have a quite regular distribution. The histogram is symmetric, and both tails fall off smoothly from a



**FIGURE 3.1**

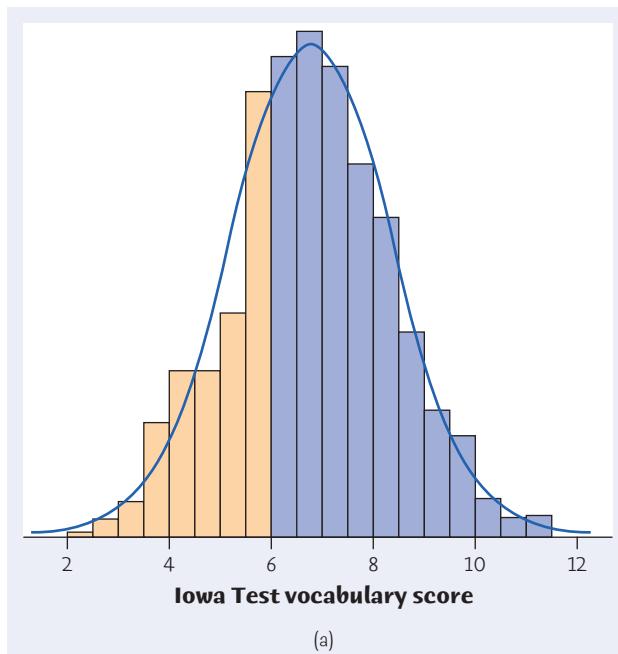
Histogram of the Iowa Test vocabulary scores of all seventh-grade students in Gary, Indiana. The smooth curve shows the overall shape of the distribution.

single center peak. There are no large gaps or obvious outliers. The smooth curve drawn through the tops of the histogram bars in Figure 3.1 is a good description of the overall pattern of the data.

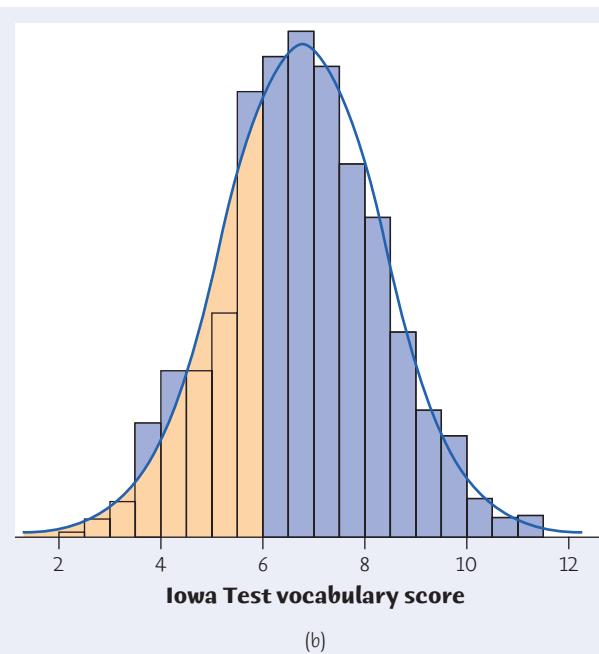
EXAMPLE 3.1 From histogram to density curve

Our eyes respond to the *areas* of the bars in a histogram. The bar areas represent proportions of the observations. Figure 3.2(a) is a copy of Figure 3.1 with the leftmost bars shaded. The area of the shaded bars in Figure 3.2(a) represents the students with vocabulary scores of 6.0 or lower. There are 287 such students, who make up the proportion $287/947 = 0.303$ of all Gary seventh-graders.

Now look at the curve drawn through the bars. In Figure 3.2(b), the area under the curve to the left of 6.0 is shaded. We can draw histogram bars taller or shorter by adjusting the vertical scale. In moving from histogram bars to a smooth curve, we make a specific choice: we adjust the scale of the graph so that *the total area under the curve is exactly 1*. The total area represents the proportion 1, that is, all the observations. We can then interpret areas under the curve as proportions of the observations. The curve is now a *density curve*. The shaded area under the density



(a)



(b)

FIGURE 3.2(a)

The proportion of scores less than or equal to 6.0 in the actual data is 0.303.

FIGURE 3.2(b)

The proportion of scores less than or equal to 6.0 from the density curve is 0.293. The density curve is a good approximation to the distribution of the data.

curve in Figure 3.2(b) represents the proportion of students with scores of 6.0 or lower. This area is 0.293, only 0.010 away from the actual proportion 0.303. The method for finding this area will be presented shortly. For now, note that the areas under the density curve give quite good approximations to the actual distribution of the 947 test scores. ■

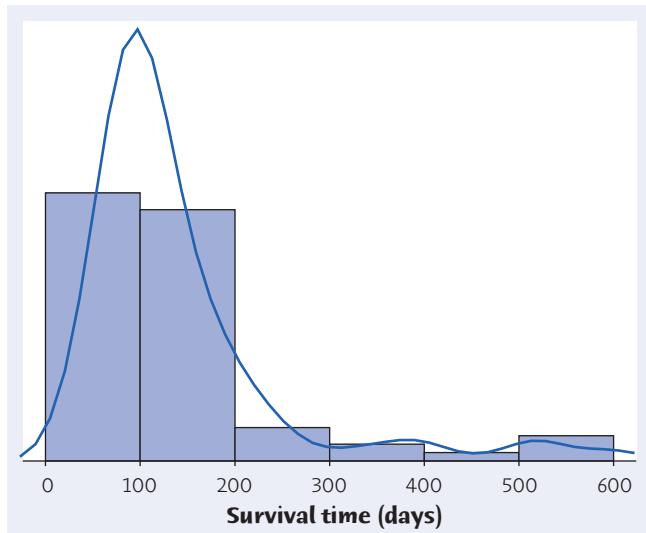
DENSITY CURVE

A **density curve** is a curve that

- is always on or above the horizontal axis, and
- has area exactly 1 underneath it.

A density curve describes the overall pattern of a distribution. The area under the curve and above any range of values is the proportion of all observations that fall in that range.

Density curves, like distributions, come in many shapes. Figure 3.3 shows a strongly skewed distribution, the survival times of guinea pigs from Exercise 2.31 (page 60). The histogram and density curve were both created from the data by software. Both show the overall shape and the “bumps” in the long right tail. The density curve shows a single high peak as a main feature of the distribution. The histogram divides the observations near the peak between two bars, thus reducing

**FIGURE 3.3**

A right-skewed distribution pictured by both a histogram and a density curve.

the height of the peak. A density curve is often a good description of the overall pattern of a distribution. Outliers, which are deviations from the overall pattern,

 are not described by the curve. Of course, no set of real data is exactly described by a density curve. The curve is an idealized description that is easy to use and accurate enough for practical use.

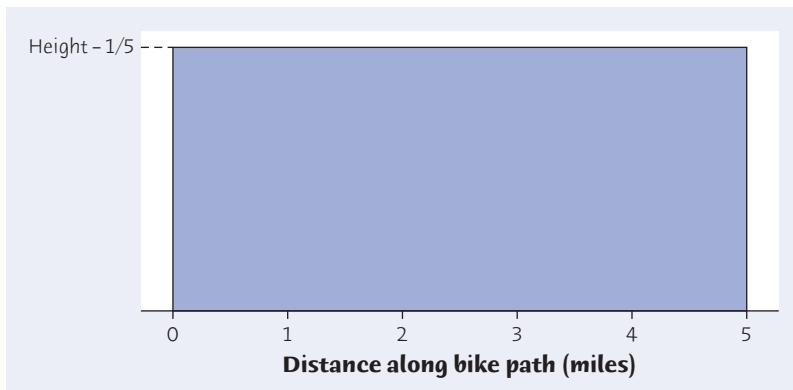
APPLY YOUR KNOWLEDGE

3.1 Sketch density curves. Sketch density curves that describe distributions with the following shapes:

- Symmetric, but with two peaks (that is, two strong clusters of observations)
- Single peak and skewed to the left

3.2 Accidents on a bike path. Examining the location of accidents on a level, 5-mile bike path shows that they occur uniformly along the length of the path. Figure 3.4 displays the density curve that describes the distribution of accidents.

- Explain why this curve satisfies the two requirements for a density curve.
- The proportion of accidents that occur in the first mile of the path is the area under the density curve between 0 miles and 1 mile. What is this area?
- There is a stream alongside the bike path between the 0.8-mile mark and the 1.3-mile mark. What proportion of accidents happen on the bike path alongside the stream?
- The bike path is a paved path through the woods, and there is a road at each end. What proportion of accidents happen more than 1 mile from either road? (Hint: First determine where on the bike path the accident needs to occur to be more than 1 mile from either road, and then find the area.)

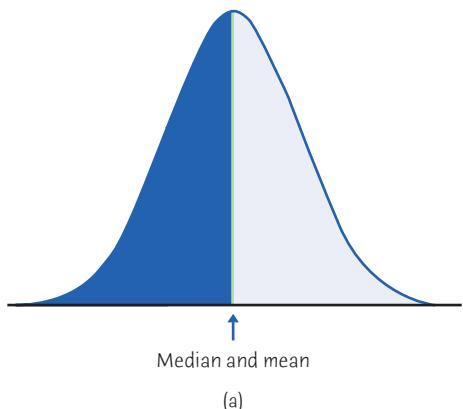
**FIGURE 3.4**

The density curve for the location of accidents along a 5-mile bike path, for Exercise 3.2.

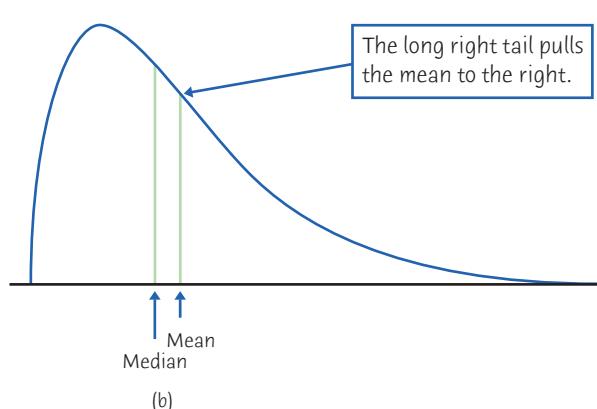
DESCRIBING DENSITY CURVES

Our measures of center and spread apply to density curves as well as to actual sets of observations. The median and quartiles are easy. Areas under a density curve represent proportions of the total number of observations. The median is the point with half the observations on either side. So the *median of a density curve is the equal-areas point*, the point with half the area under the curve to its left and the remaining half of the area to its right. The quartiles divide the area under the curve into quarters. One-fourth of the area under the curve is to the left of the first quartile, and three-fourths of the area is to the left of the third quartile. You can roughly locate the median and quartiles of any density curve by eye by dividing the area under the curve into four equal parts.

Because density curves are idealized patterns, a symmetric density curve is exactly symmetric. The median of a symmetric density curve is therefore at its center. Figure 3.5(a) shows a symmetric density curve with the median marked. It isn't so easy to spot the equal-areas point on a skewed curve. There are mathematical ways of finding the median for any density curve. That's how we marked the median on the skewed curve in Figure 3.5(b).

**FIGURE 3.5(a)**

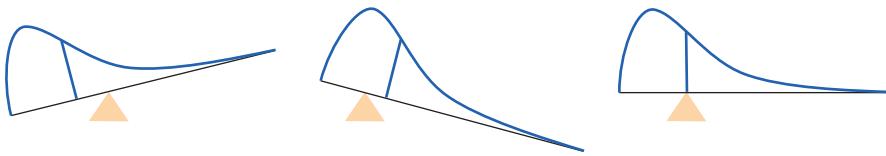
The median and mean of a symmetric density curve both lie at the center of symmetry.

**FIGURE 3.5(b)**

The median and mean of a right-skewed density curve. The mean is pulled away from the median toward the long tail.

FIGURE 3.6

The mean is the balance point of a density curve.



What about the mean? The mean of a set of observations is their arithmetic average. If we think of the observations as weights strung out along a thin rod, the mean is the point at which the rod would balance. This fact is also true of density curves. *The mean is the point at which the curve would balance if made of solid material.* Figure 3.6 illustrates this fact about the mean. A symmetric curve balances at its center because the two sides are identical. *The mean and median of a symmetric density curve are equal*, as in Figure 3.5(a). We know that the mean of a skewed distribution is pulled toward the long tail. Figure 3.5(b) shows how the mean of a skewed density curve is pulled toward the long tail more than is the median. It's hard to locate the balance point by eye on a skewed curve. There are mathematical ways of calculating the mean for any density curve, so we are able to mark the mean as well as the median in Figure 3.5(b).

MEDIAN AND MEAN OF A DENSITY CURVE

The **median** of a density curve is the equal-areas point, the point that divides the area under the curve in half.

The **mean** of a density curve is the balance point, at which the curve would balance if made of solid material.

The median and mean are the same for a symmetric density curve. They both lie at the center of the curve. The mean of a skewed curve is pulled away from the median in the direction of the long tail.

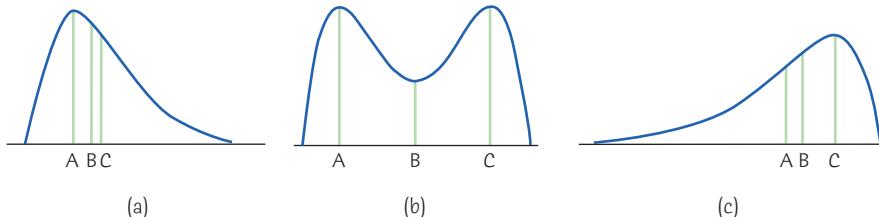
*mean μ
standard deviation σ*

Because a density curve is an idealized description of a distribution of data, we need to distinguish between the mean and standard deviation of the density curve and the mean \bar{x} and standard deviation s computed from the actual observations. The usual notation for the **mean of a density curve** is μ (the Greek letter mu). We write the **standard deviation of a density curve** as σ (the Greek letter sigma). We can roughly locate the mean μ of any density curve by eye, as the balance point. There is no easy way to locate the standard deviation σ by eye for density curves in general.

APPLY YOUR KNOWLEDGE

3.3 Mean and median. What is the mean μ of the density curve pictured in Figure 3.4 on page 73? (That is, where would the curve balance?) What is the median? (That is, where is the point with area 0.5 on either side?)

3.4 Mean and median. Figure 3.7 displays three density curves, each with three points marked on them. At which of these points on each curve do the mean and the median fall?

**FIGURE 3.7**

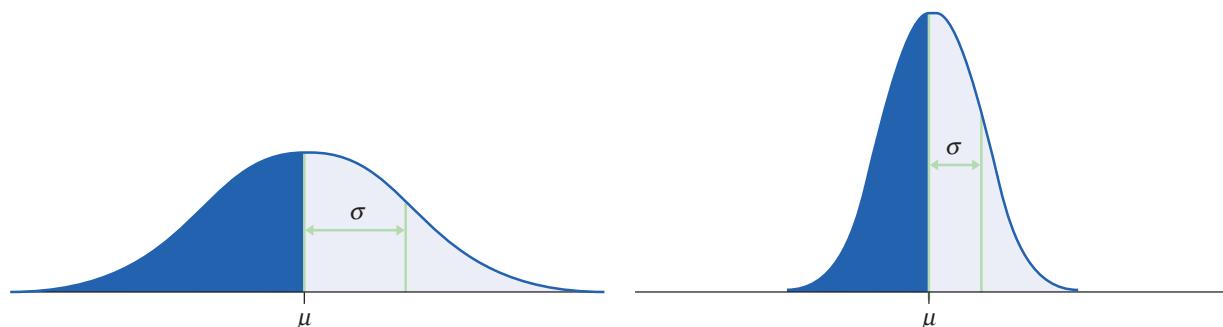
Three density curves, for Exercise 3.4.

NORMAL DISTRIBUTIONS

One particularly important class of density curves has already appeared in Figures 3.1 and 3.2. They are called **Normal curves**. The distributions they describe are called **Normal distributions**. Normal distributions play a large role in statistics, but they are rather special and not at all “normal” in the sense of being usual or average. We capitalize Normal to remind you that these curves are special. Look at the two Normal curves in Figure 3.8. They illustrate several important facts:

- All Normal curves have the same overall shape: symmetric, single-peaked, bell-shaped.
- Any specific Normal curve is completely described by giving its mean μ and its standard deviation σ .
- The mean is located at the center of the symmetric curve and is the same as the median. Changing μ without changing σ moves the Normal curve along the horizontal axis without changing its spread.
- The standard deviation σ controls the spread of a Normal curve. Curves with larger standard deviations are more spread out.

The standard deviation σ is the natural measure of spread for Normal distributions. Not only do μ and σ completely determine the shape of a Normal curve, but we can locate σ by eye on a Normal curve. Here's how. Imagine that you

Normal curve**Normal distribution****FIGURE 3.8**Two Normal curves, showing the mean μ and standard deviation σ .

are skiing down a mountain that has the shape of a Normal curve. At first, you descend at an ever-steepener angle as you go out from the peak:



Fortunately, before you find yourself going straight down, the slope begins to grow flatter rather than steeper as you go out and down:



The points at which this change of curvature takes place are located at distance σ on either side of the mean μ . You can feel the change as you run a pencil along a Normal curve, and so find the standard deviation. Remember that μ and σ alone do not specify the shape of most distributions, and that the shape of density curves in general does not reveal σ . These are special properties of Normal distributions.

NORMAL DISTRIBUTIONS

A **Normal distribution** is described by a Normal density curve. Any particular Normal distribution is completely specified by two numbers, its mean μ and standard deviation σ .

The mean of a Normal distribution is at the center of the symmetric Normal curve. The standard deviation is the distance from the center to the change-of-curvature points on either side.

Why are the Normal distributions important in statistics? Here are three reasons. First, Normal distributions are good descriptions for some distributions of *real data*. Distributions that are often close to Normal include scores on tests taken by many people (such as Iowa Tests and SAT exams), repeated careful measurements of the same quantity, and characteristics of biological populations (such as lengths of crickets and yields of corn). Second, Normal distributions are good approximations to the results of many kinds of *chance outcomes*, such as the proportion of heads in many tosses of a coin. Third, we will see that many *statistical inference* procedures based on Normal distributions work well for other roughly symmetric distributions. However, many sets of data do not follow a Normal distribution. Most income distributions, for example, are skewed to the right and so are not Normal. Non-Normal data, like nonnormal people, not only are common but are sometimes more interesting than their Normal counterparts.

THE 68–95–99.7 RULE

Although there are many Normal curves, they all have common properties. In particular, all Normal distributions obey the following rule.

THE 68–95–99.7 RULE

In the Normal distribution with mean μ and standard deviation σ :

- Approximately 68% of the observations fall within σ of the mean μ .
- Approximately 95% of the observations fall within 2σ of μ .
- Approximately 99.7% of the observations fall within 3σ of μ .

Figure 3.9 illustrates the 68–95–99.7 rule. By remembering these three numbers, you can think about Normal distributions without constantly making detailed calculations.

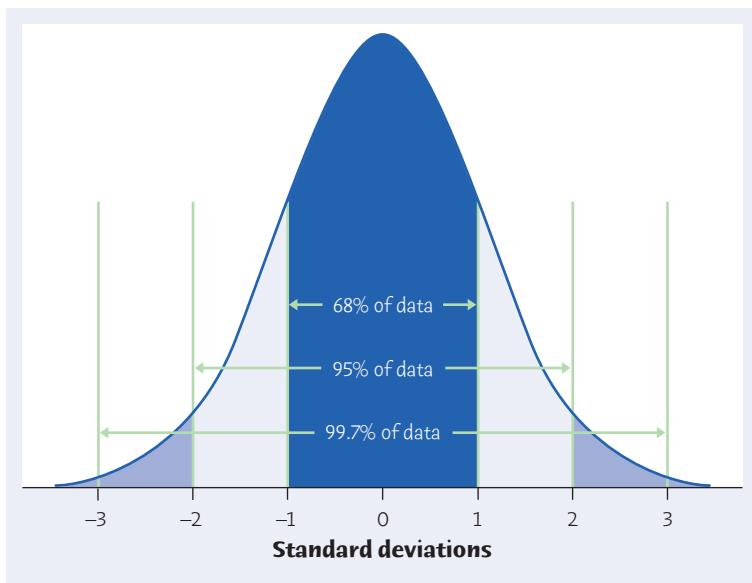


FIGURE 3.9

The 68–95–99.7 rule for Normal distributions.

EXAMPLE 3.2 Iowa Test scores

Figures 3.1 and 3.2 (see pages 70 and 71) show that the distribution of Iowa Test vocabulary scores for seventh-grade students in Gary, Indiana, is close to Normal. Suppose that the distribution is exactly Normal with mean $\mu = 6.84$ and standard deviation $\sigma = 1.55$. (These are the mean and standard deviation of the 947 actual scores.)

FIGURE 3.10

The 68–95–99.7 rule applied to the distribution of Iowa Test scores for seventh-grade students in Gary, Indiana, for Example 3.2. The mean and standard deviation are $\mu = 6.84$ and $\sigma = 1.55$.

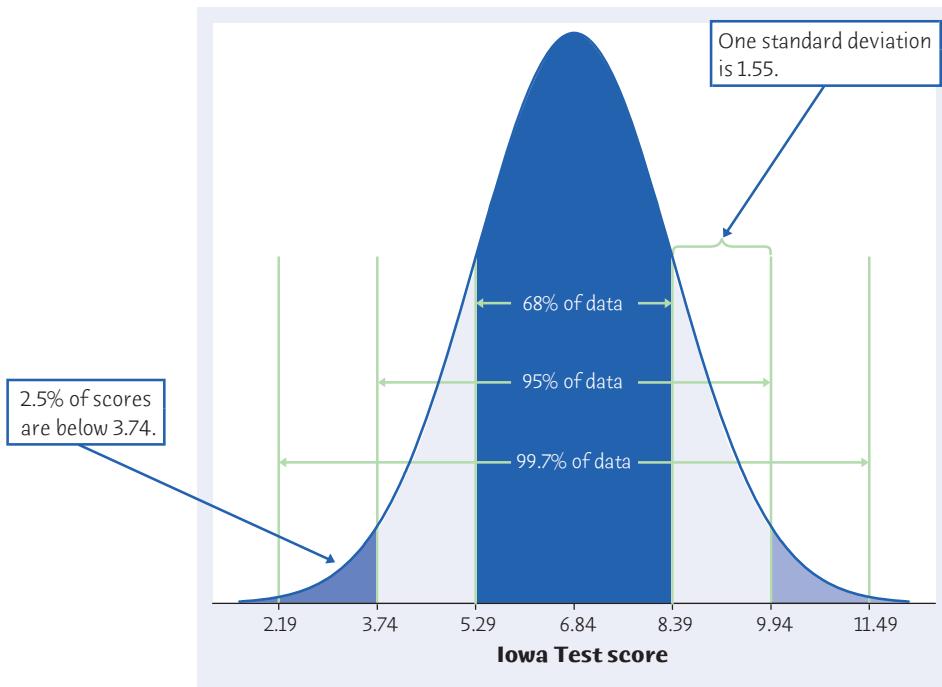


Figure 3.10 applies the 68–95–99.7 rule to the Iowa Test scores. The 95 part of the rule says that 95% of all scores are between

$$\mu - 2\sigma = 6.84 - (2)(1.55) = 6.84 - 3.10 = 3.74$$

and

$$\mu + 2\sigma = 6.84 + (2)(1.55) = 6.84 + 3.10 = 9.94$$

The other 5% of scores are outside this range. Because Normal distributions are symmetric, half of these scores are lower than 3.74 and half are higher than 9.94. That is, 2.5% of the scores are below 3.74 and 2.5% are above 9.94. ■

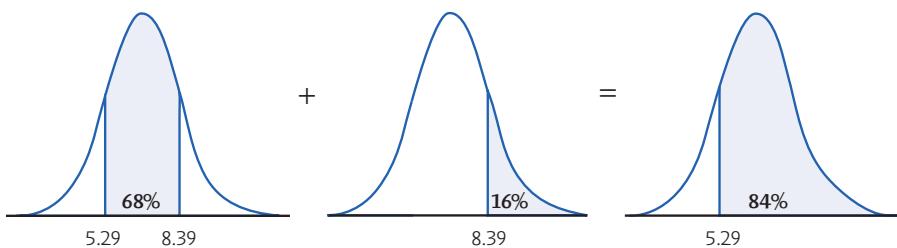


The 68–95–99.7 rule describes distributions that are exactly Normal. Real data such as the actual Gary scores are never exactly Normal. For one thing, Iowa Test scores are reported only to the nearest tenth. A score can be 9.9 or 10.0, but not 9.94. We use a Normal distribution because it's a good approximation, and because we think the knowledge that the test measures is continuous rather than stopping at tenths.

How well does our work in Example 3.2 describe the actual Iowa Test scores? Well, 900 of the 947 scores are between 3.74 and 9.94. That's 95.04%, very accurate indeed. Of the remaining 47 scores, 20 are below 3.74 and 27 are above 9.94. The tails of the actual data are not quite equal, as they would be in an exactly Normal distribution. Normal distributions often describe real data better in the center of the distribution than in the extreme high and low tails.

EXAMPLE 3.3 Iowa Test scores

Look again at Figure 3.10. A score of 5.29 is one standard deviation below the mean. What percent of scores are higher than 5.29? Find the answer by adding areas in the figure. Here is the calculation in pictures:



$$\text{percent between } 5.29 \text{ and } 8.39 + \text{percent above } 8.39 = \text{percent above } 5.29$$

$$68\% + 16\% = 84\%$$

Be sure you understand where the 16% came from. We know that 68% of scores are between 5.29 and 8.39, so 32% of scores are outside that range. These are equally split between the two tails, 16% below 5.29 and 16% above 8.39. ■

Because we will mention Normal distributions often, a short notation is helpful. We abbreviate the Normal distribution with mean μ and standard deviation σ as $N(\mu, \sigma)$. For example, the distribution of Gary Iowa Test scores is approximately $N(6.84, 1.55)$.

APPLY YOUR KNOWLEDGE

3.5 Fruit flies. The common fruit fly *Drosophila melanogaster* is the most studied organism in genetic research because it is small, easy to grow, and reproduces rapidly. The length of the thorax (where the wings and legs attach) in a population of male fruit flies is approximately Normal with mean 0.800 millimeters (mm) and standard deviation 0.078 mm. Draw a Normal curve on which this mean and standard deviation are correctly located. (*Hint:* Draw an unlabeled Normal curve, locate the points where the curvature changes, then add number labels on the horizontal axis.)

3.6 Fruit flies. The length of the thorax in a population of male fruit flies is approximately Normal with mean 0.800 millimeters (mm) and standard deviation 0.078 mm. Use the 68–95–99.7 rule to answer the following questions. (Start by making a sketch like Figure 3.10.)

- (a) What range of lengths covers almost all (99.7%) of this distribution?
- (b) What percent of male fruit flies have a thorax length exceeding 0.878 mm?

3.7 Monsoon rains. The summer monsoon brings 80% of India's rainfall and is essential for the country's agriculture. Records going back more than a century show that the amount of monsoon rainfall varies from year to year according to a distribution



Plastique1/Dreamtime.com

that is approximately Normal with mean 852 millimeters (mm) and standard deviation 82 mm.² Use the 68–95–99.7 rule to answer the following questions.

- Between what values do the monsoon rains fall in 95% of all years?
- How small are the monsoon rains in the driest 2.5% of all years?

THE STANDARD NORMAL DISTRIBUTION

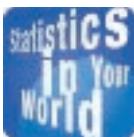
As the 68–95–99.7 rule suggests, all Normal distributions share many properties. In fact, all Normal distributions are the same if we measure in units of size σ about the mean μ as center. Changing to these units is called *standardizing*. To standardize a value, subtract the mean of the distribution and then divide by the standard deviation.

STANDARDIZING AND z -SCORES

If x is an observation from a distribution that has mean μ and standard deviation σ , the **standardized value** of x is

$$z = \frac{x - \mu}{\sigma}$$

A standardized value is often called a **z -score**.



He said, she said

Height, weight, and body mass distributions

in this book come from actual measurements by a government survey. Good thing that is. When asked their weight, almost all women say they weigh less than they really do. Heavier men also underreport their weight—but lighter men claim to weigh more than the scale shows. We leave you to ponder the psychology of the two sexes. Just remember that “say so” is no substitute for measuring.

A z -score tells us how many standard deviations the original observation falls away from the mean, and in which direction. Observations larger than the mean are positive when standardized, and observations smaller than the mean are negative.

EXAMPLE 3.4 Standardizing women's heights

The heights of women aged 20 to 29 are approximately Normal with $\mu = 64.3$ inches and $\sigma = 2.7$ inches.³ The standardized height is

$$z = \frac{\text{height} - 64.3}{2.7}$$

A woman's standardized height is the number of standard deviations by which her height differs from the mean height of all young women. A woman 70 inches tall, for example, has standardized height

$$z = \frac{70 - 64.3}{2.7} = 2.11$$

or 2.11 standard deviations above the mean. Similarly, a woman 5 feet (60 inches) tall has standardized height

$$z = \frac{60 - 64.3}{2.7} = -1.59$$

or 1.59 standard deviations less than the mean height. ■

We often standardize observations from symmetric distributions to express them in a common scale. We might, for example, compare the heights of two children of different ages by calculating their z -scores. The standardized heights tell us where each child stands in the distribution for his or her age group.

If the variable we standardize has a Normal distribution, standardizing does more than give a common scale. It makes all Normal distributions into a single distribution, and this distribution is still Normal. Standardizing a variable that has any Normal distribution produces a new variable that has the *standard Normal distribution*.

STANDARD NORMAL DISTRIBUTION

The **standard Normal distribution** is the Normal distribution $N(0, 1)$ with mean 0 and standard deviation 1.

If a variable x has any Normal distribution $N(\mu, \sigma)$ with mean μ and standard deviation σ , then the standardized variable

$$z = \frac{x - \mu}{\sigma}$$

has the standard Normal distribution.

APPLY YOUR KNOWLEDGE

3.8 SAT versus ACT. In 2010, when she was a high school senior, Alysha scored 670 on the Mathematics part of the SAT.⁴ The distribution of SAT Math scores in 2010 was Normal with mean 516 and standard deviation 116. John took the ACT and scored 26 on the Mathematics portion. ACT Math scores for 2010 were Normally distributed with mean 21.0 and standard deviation 5.3. Find the standardized scores for both students. Assuming that both tests measure the same kind of ability, who had the higher score?

3.9 Men's and women's heights. The heights of women aged 20 to 29 are approximately Normal with mean 64.3 inches and standard deviation 2.7 inches. Men the same age have mean height 69.9 inches with standard deviation 3.1 inches.⁵ What are the z -scores for a woman 6 feet tall and a man 6 feet tall? Say in simple language what information the z -scores give that the original nonstandardized heights do not.

FINDING NORMAL PROPORTIONS

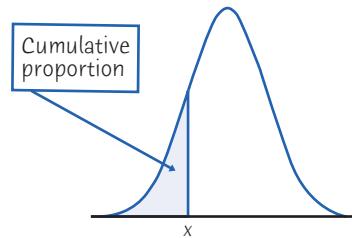
Areas under a Normal curve represent proportions of observations from that Normal distribution. There is no formula for areas under a Normal curve. Calculations use either software that calculates areas or a table of areas. Most tables and software calculate one kind of area, *cumulative proportions*. The idea of “cumulative” is “everything that came before.” Here is the exact statement.



Spencer Grant/PhotoEdit

CUMULATIVE PROPORTIONS

The **cumulative proportion** for a value x in a distribution is the proportion of observations in the distribution that are less than or equal to x .

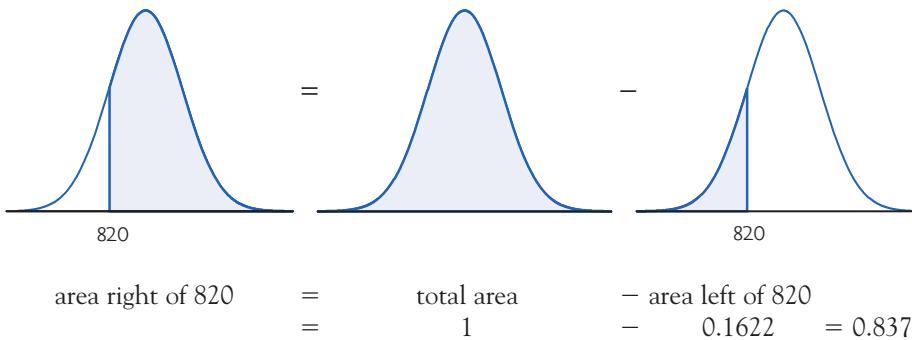


The key to calculating Normal proportions is to match the area you want with areas that represent cumulative proportions. If you make a sketch of the area you want, you will almost never go wrong. Find areas for cumulative proportions either from software or (with an extra step) from a table. The following example shows the method in a picture. The numerical answer in the example was obtained using technology and will be explained following the example.

EXAMPLE 3.5 Who qualifies for college sports?

The National Collegiate Athletic Association (NCAA) uses a sliding scale for eligibility for Division I athletes.⁶ Those students with a 2.5 high school GPA must have a combined score of at least 820 on the Mathematics and Reading parts of the SAT in order to compete in their first college year. The scores of the 1.5 million high school seniors taking the SAT this year are approximately Normal with mean 1026 and standard deviation 209. What percent of high school seniors meet this SAT requirement of a combined score of 820 or better?

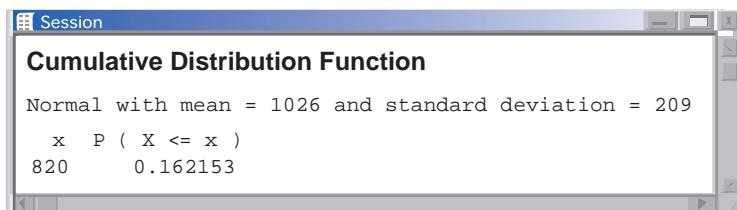
Here is the calculation in a picture: the proportion of scores above 820 is the area under the curve to the right of 820. That's the total area under the curve (which is always 1) minus the cumulative proportion up to 820.



About 84% of all high school seniors meet this SAT requirement of a combined math and reading score of 820 or higher. ■

There is no area under a smooth curve and exactly over the point 820. Consequently, the area to the right of 820 (the proportion of scores > 820) is the same as the area at or to the right of this point (the proportion of scores ≥ 820). The actual data may contain a student who scored exactly 820 on the SAT. That the proportion of scores exactly equal to 820 is 0 for a Normal distribution is a consequence of the idealized smoothing of Normal distributions for data.

To find the numerical value 0.1622 of the cumulative proportion in Example 3.5 using software, plug in mean 1026 and standard deviation 209 and ask for the cumulative proportion for 820. Software often uses terms such as “cumulative distribution” or “cumulative probability.” We will learn in Chapter 10 why the language of probability fits. Here, for example, is Minitab’s output:



The screenshot shows a Minitab session window titled "Session". It displays the following output for the "Cumulative Distribution Function" command:

```
Normal with mean = 1026 and standard deviation = 209
      x   P ( X <= x )
    820     0.162153
```

The P in the output stands for “probability,” but we can read it as “proportion of the observations.” The *Normal Curve* applet is even handier because it draws pictures as well as finding areas. If you are not using software, you can find cumulative proportions for Normal curves from a table. This requires an extra step.



USING THE STANDARD NORMAL TABLE

The extra step in finding cumulative proportions from a table is that we must first standardize to express the problem in the standard scale of z -scores. This allows us to get by with just one table, a table of *standard Normal cumulative proportions*. Table A in the back of the book gives cumulative proportions for the standard Normal distribution. The pictures at the top of the table remind us that the entries are cumulative proportions, areas under the curve to the left of a value z .

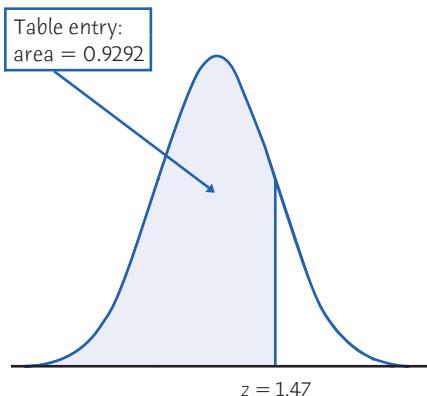
EXAMPLE 3.6 The standard Normal table

What proportion of observations on a standard Normal variable z take values less than 1.47?

Solution: To find the area to the left of 1.47, locate 1.4 in the left-hand column of Table A, then locate the remaining digit 7 as .07 in the top row. The entry opposite 1.4 and under .07 is 0.9292. This is the cumulative proportion we seek. Figure 3.11 illustrates this area. ■

FIGURE 3.11

The area under a standard Normal curve to the left of the point $z = 1.47$ is 0.9292. Table A gives areas under the standard Normal curve.



Now that you see how Table A works, let's redo Example 3.5 using the table. We can break Normal calculations using the table into three steps.

EXAMPLE 3.7 Who qualifies for college sports?

Scores of high school seniors on the SAT follow the Normal distribution with mean $\mu = 1026$ and standard deviation $\sigma = 209$. What proportion of seniors score at least 820?

Step 1. Draw a picture. The picture is exactly as in Example 3.5. It shows that

$$\text{area to the right of } 820 = 1 - \text{area to the left of } 820$$

Step 2. Standardize. Call the SAT score x . Subtract the mean and then divide by the standard deviation to transform the problem about x into a problem about a standard Normal z :

$$\begin{aligned} x &\geq 820 \\ \frac{x - 1026}{209} &\geq \frac{820 - 1026}{209} \\ z &\geq -0.99 \end{aligned}$$

Step 3. Use the table. The picture shows that we need the cumulative proportion for $x = 820$. Step 2 says that this is the same as the cumulative proportion for $z = -0.99$. The Table A entry for $z = -0.99$ says that this cumulative proportion is 0.1611. The area to the right of -0.99 is therefore $1 - 0.1611 = 0.8389$. ■

The area from the table in Example 3.7 (0.8389) is slightly less accurate than the area from software in Example 3.5 (0.8378) because we must round z to two decimal places when we use Table A. The difference is rarely important in practice. Here's the method in outline form.

USING TABLE A TO FIND NORMAL PROPORTIONS

Step 1. State the problem in terms of the observed variable x . **Draw a picture** that shows the proportion you want in terms of cumulative proportions.

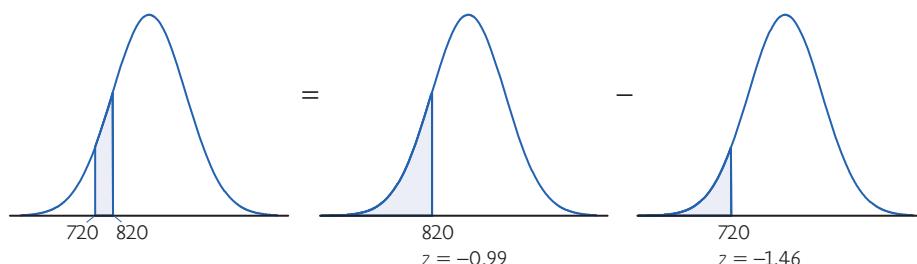
Step 2. Standardize x to restate the problem in terms of a standard Normal variable z .

Step 3. Use Table A and the fact that the total area under the curve is 1 to find the required area under the standard Normal curve.

EXAMPLE 3.8 Who qualifies for college sports?

Recall that the NCAA uses a sliding scale for eligibility for Division I athletics. Students with a 2.5 GPA must have a combined SAT score of 820 or higher to be eligible. Students with lower GPAs will require higher SAT scores for eligibility, while students with higher GPAs can have a lower SAT score and still be eligible. For example, students with a 2.75 GPA are required to have a combined SAT score that is at least 720. What proportion of all students who take the SAT would meet an SAT requirement of at least 720, but not 820?

Step 1. State the problem and draw a picture. Call the SAT score x . The variable x has the $N(1026, 209)$ distribution. What proportion of SAT scores fall between 720 and 820? Here is the picture:



Step 2. Standardize. Subtract the mean and then divide by the standard deviation to turn x into a standard Normal z :

$$\begin{aligned} 720 &\leq x < 820 \\ \frac{720 - 1026}{209} &\leq \frac{x - 1026}{209} < \frac{820 - 1026}{209} \\ -1.46 &\leq z < -0.99 \end{aligned}$$

Step 3. Use the table. Follow the picture (we added the z -scores to the picture to help you):

$$\begin{aligned} \text{area between } -1.46 \text{ and } -0.99 &= (\text{area left of } -0.99) - (\text{area left of } -1.46) \\ &= 0.1611 - 0.0721 = 0.0890 \end{aligned}$$

About 9% of high school seniors have SAT scores between 720 and 820. ■

Sometimes we encounter a value of z more extreme than those appearing in Table A. For example, the area to the left of $z = -4$ is not given directly in the table. The z -values in Table A leave only area 0.0002 in each tail unaccounted for. For practical purposes, we can act as if there is zero area outside the range of Table A. Specifically, we act as if there is zero area below $z = -3.5$ and zero area above $z = 3.5$.



APPLY YOUR KNOWLEDGE

3.10 Use the Normal table. Use Table A to find the proportion of observations from a standard Normal distribution that satisfies each of the following statements. In each case, sketch a standard Normal curve and shade the area under the curve that is the answer to the question.

- (a) $z < -1.42$ (b) $z > -1.42$ (c) $z < 2.35$ (d) $-1.42 < z < 2.35$

3.11 Monsoon rains. The summer monsoon rains in India follow approximately a Normal distribution with mean 852 millimeters (mm) of rainfall and standard deviation 82 mm.

- (a) In the drought year 1987, 697 mm of rain fell. In what percent of all years will India have 697 mm or less of monsoon rain?
 (b) “Normal rainfall” means within 20% of the long-term average, or between 683 mm and 1022 mm. In what percent of all years is the rainfall normal?

3.12 The Medical College Admission Test. Almost all medical schools in the United States require students to take the Medical College Admission Test (MCAT).⁷ The exam is composed of three multiple-choice sections (Physical Sciences, Verbal Reasoning, and Biological Sciences). The score on each section is converted to a 15-point scale so that the total score has a maximum value of 45. The total scores follow a Normal distribution, and in 2010 the mean was 25.0 with a standard deviation of 6.4. There is little change in the distribution of scores from year to year.

- (a) What proportion of students taking the MCAT had a score over 30?
 (b) What proportion had scores between 20 and 25?

FINDING A VALUE WHEN GIVEN A PROPORTION

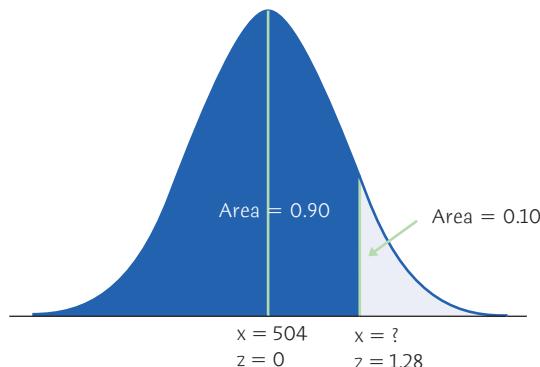
Examples 3.5 to 3.8 illustrate the use of software or Table A to find what proportion of the observations satisfies some condition, such as “SAT score above 820.” We may instead want to find the observed value with a given proportion of the observations above or below it. Statistical software will do this directly.



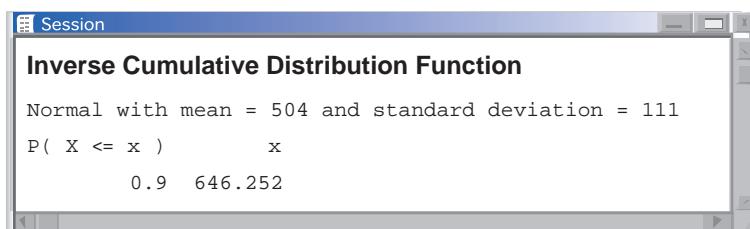
EXAMPLE 3.9 Find the top 10% using software

Scores on the SAT Reading test in recent years follow approximately the $N(504, 111)$ distribution. How high must a student score to place in the top 10% of all students taking the SAT?

We want to find the SAT score x with area 0.1 to its right under the Normal curve with mean $\mu = 504$ and standard deviation $\sigma = 111$. That’s the same as finding the SAT score x with area 0.9 to its left. Figure 3.12 poses the question in graphical form. Most software will tell you x when you plug in mean 504, standard deviation 111, and cumulative proportion 0.9. Here is Minitab’s output:

**FIGURE 3.12**

Locating the point on a Normal curve with area 0.10 to its right, for Examples 3.9 and 3.10.



Minitab gives $x = 646.252$. So scores above 647 are in the top 10%. (Round up because SAT scores can only be whole numbers.) ■

Without software, use Table A backward. Find the given proportion in the body of the table and then read the corresponding z from the left column and top row. There are again three steps.

EXAMPLE 3.10 Find the top 10% using Table A

Scores on the SAT Reading test in recent years follow approximately the $N(504, 111)$ distribution. How high must a student score to place in the top 10% of all students taking the SAT?

Step 1. State the problem and draw a picture. This step is exactly as in Example 3.9. The picture is Figure 3.12. The x -value that puts a student in the top 10% is the same as the x -value for which 90% of the area is to the left of x .

Step 2. Use the table. Look in the body of Table A for the entry closest to 0.9. It is 0.8997. This is the entry corresponding to $z = 1.28$. So $z = 1.28$ is the standardized value with area 0.9 to its left.

Step 3. Unstandardize to transform z back to the original x scale. We know that the standardized value of the unknown x is $z = 1.28$. This means that x itself lies 1.28 standard deviations above the mean on this particular Normal curve. That is,

$$\begin{aligned}x &= \text{mean} + (1.28)(\text{standard deviation}) \\&= 504 + (1.28)(111) = 646.08\end{aligned}$$

A student must score at least 647 to place in the highest 10%. ■



Will & Deni McIntyre/Photo Researchers

EXAMPLE 3.11 Find the first quartile

High levels of cholesterol in the blood increase the risk of heart disease. For 14-year-old boys, the distribution of blood cholesterol is approximately Normal with mean $\mu = 170$ milligrams of cholesterol per deciliter of blood (mg/dl) and standard deviation $\sigma = 30$ mg/dl.⁸ What is the first quartile of the distribution of blood cholesterol?

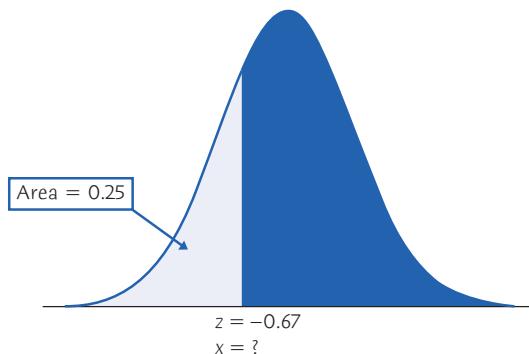
Step 1. State the problem and draw a picture. Call the cholesterol level x . The variable x has the $N(170, 30)$ distribution. The first quartile is the value with 25% of the distribution to its left. Figure 3.13 is the picture.

Step 2. Use the table. Look in the body of Table A for the entry closest to 0.25. It is 0.2514. This is the entry corresponding to $z = -0.67$. So $z = -0.67$ is the standardized value with area 0.25 to its left.

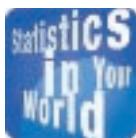
Step 3. Unstandardize. The cholesterol level corresponding to $z = -0.67$ lies 0.67 standard deviations below the mean, so

$$\begin{aligned}x &= \text{mean} - (0.67)(\text{standard deviation}) \\&= 170 - (0.67)(30) = 149.9\end{aligned}$$

The first quartile of blood cholesterol levels in 14-year-old boys is about 150 mg/dl. ■

**FIGURE 3.13**

Locating the first quartile of a Normal curve, for Example 3.11.



The bell curve?

Does the distribution of human

intelligence follow the “bell curve” of a Normal distribution? Scores on IQ tests do roughly follow a Normal distribution. That is because a test score is calculated from a person’s answers in a way that is designed to produce a Normal distribution. To conclude that intelligence follows a bell curve, we must agree that the test scores directly measure intelligence. Many psychologists don’t think there is one human characteristic that we can call “intelligence” and can measure by a single test score.

APPLY YOUR KNOWLEDGE

3.13 Table A. Use Table A to find the value z of a standard Normal variable that satisfies each of the following conditions. (Use the value of z from Table A that comes closest to satisfying the condition.) In each case, sketch a standard Normal curve with your value of z marked on the axis.

- (a) The point z with 15% of the observations falling below it
- (b) The point z with 70% of the observations falling above it

3.14 The Medical College Admission Test. The total scores on the Medical College Admission Test (MCAT) follow a Normal distribution with mean 25.0 and standard deviation 6.4. What are the median and the first and third quartiles of the MCAT scores?

CHAPTER 3 SUMMARY

CHAPTER SPECIFICS

- We can sometimes describe the overall pattern of a distribution by a **density curve**. A density curve has total area 1 underneath it. An area under a density curve gives the proportion of observations that fall in a range of values.
- A density curve is an idealized description of the overall pattern of a distribution that smooths out the irregularities in the actual data. We write the **mean of a density curve** as μ and the **standard deviation of a density curve** as σ to distinguish them from the mean \bar{x} and standard deviation s of the actual data.
- The mean, the median, and the quartiles of a density curve can be located by eye. The **mean μ** is the balance point of the curve. The **median** divides the area under the curve in half. The **quartiles** and the median divide the area under the curve into quarters. The **standard deviation σ** cannot be located by eye on most density curves.
- The mean and median are equal for symmetric density curves. The mean of a skewed curve is located farther toward the long tail than is the median.
- The **Normal distributions** are described by a special family of bell-shaped, symmetric density curves, called **Normal curves**. The mean μ and standard deviation σ completely specify a Normal distribution $N(\mu, \sigma)$. The mean is the center of the curve, and σ is the distance from μ to the change-of-curvature points on either side.
- To **standardize** any observation x , subtract the mean of the distribution and then divide by the standard deviation. The resulting **z -score**

$$z = \frac{x - \mu}{\sigma}$$

says how many standard deviations x lies from the distribution mean.

- All Normal distributions are the same when measurements are transformed to the standardized scale. In particular, all Normal distributions satisfy the **68–95–99.7 rule**, which describes what percent of observations lie within one, two, and three standard deviations of the mean, respectively.
- If x has the $N(\mu, \sigma)$ distribution, then the **standardized variable** $z = (x - \mu)/\sigma$ has the **standard Normal distribution** $N(0, 1)$ with mean 0 and standard deviation 1. Table A gives the **cumulative proportions** of standard Normal observations that are less than z for many values of z . By standardizing, we can use Table A for any Normal distribution.

LINK IT

When exploring data, some data sets can be shown to closely follow the Normal distribution. When this is true, the description of the data can be greatly simplified without much loss of information. We can calculate the percent of the distribution in an interval for *any* Normal distribution if we know its mean and standard deviation. This also shows why the mean and standard deviation can be important numerical summaries. For distributions that are approximately Normal,

these two numerical summaries give a complete description of the distribution of our data. It is important to remember that not all distributions can be well approximated by a Normal curve. In these cases, calculations based on the Normal distribution can be misleading.

Normal distributions are also good approximations to many kinds of chance outcomes such as the proportion of heads in many tosses of a coin (this setting will be described in more detail in Chapter 20). And when we discuss statistical inference in Part III of the text, we will find that many procedures based on Normal distributions work well for other roughly symmetric distributions.

CHECK YOUR SKILLS

3.15 Which of these variables is most likely to have a Normal distribution?

- (a) Income per person for 150 different countries.
- (b) Sale prices of 200 homes in a suburb of Chicago.
- (c) Heights of 100 white pine trees in a forest.

3.16 To completely specify the shape of a Normal distribution, you must give

- (a) the mean and the standard deviation.
- (b) the five-number summary.
- (c) the median and the quartiles.

3.17 Figure 3.14 shows a Normal curve. The mean of this distribution is

- (a) 0. (b) 2. (c) 3.

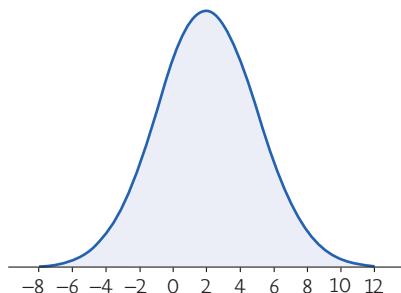


FIGURE 3.14

A Normal curve, for Exercises 3.17 and 3.18.

3.18 The standard deviation of the Normal distribution in Figure 3.14 is

- (a) 2. (b) 3. (c) 5.

3.19 The length of human pregnancies from conception to birth varies according to a distribution that is approximately Normal with mean 266 days and standard deviation 16 days. About 95% of all pregnancies last between

- (a) 250 and 282 days. (b) 234 and 298 days.
- (c) 218 and 314 days.

3.20 The scores of adults on an IQ test are approximately Normal with mean 100 and standard deviation 15. The organization MENSA, which calls itself “the high-IQ society,” requires an IQ score of 130 or higher for membership. What percent of adults would qualify for membership?

- (a) 95% (b) 5% (c) 2.5%

3.21 The scores of adults on an IQ test are approximately Normal with mean 100 and standard deviation 15. Clara scores 127 on such a test. Her z -score is about

- (a) 1.27. (b) 1.80. (c) 8.47.

3.22 The proportion of observations from a standard Normal distribution that take values greater than 1.83 is about

- (a) 0.9641. (b) 0.0359. (c) 0.0336.

3.23 The proportion of observations from a standard Normal distribution that take values less than -0.75 is about

- (a) 0.2266. (b) 0.7734. (c) 0.8023.

3.24 The scores of adults on an IQ test are approximately Normal with mean 100 and standard deviation 15. Clara scores 127 on such a test. She scores higher than what percent of all adults?

- (a) About 10%. (b) About 90%. (c) About 96%.

CHAPTER 3 EXERCISES

3.25 Understanding density curves. Remember that it is areas under a density curve, not the height of the curve, that give proportions in a distribution. To illustrate this, sketch a density curve that has a tall, thin peak at 0 on the horizontal axis but has most of its area close to 1 on the horizontal axis without a high peak at 1.

3.26 Daily activity. It appears that people who are mildly obese are less active than leaner people. One study looked at the average number of minutes per day that people spend standing or walking.⁹ Among mildly obese people, minutes of activity varied according to the $N(373, 67)$ distribution. Minutes of activity for lean people had the $N(526, 107)$ distribution. Within what limits do the active minutes for about 95% of the people in each group fall? Use the 68–95–99.7 rule.

3.27 Low IQ test scores. Scores on the Wechsler Adult Intelligence Scale (WAIS) are approximately Normal with mean 100 and standard deviation 15. People with WAIS scores below 70 are considered mentally retarded when, for example, applying for Social Security disability benefits. According to the 68–95–99.7 rule, about what percent of adults are retarded by this criterion?

3.28 Standard Normal drill. Use Table A to find the proportion of observations from a standard Normal distribution that fall in each of the following regions. In each case, sketch a standard Normal curve and shade the area representing the region.

- | | |
|--------------------|------------------------|
| (a) $z \leq -1.25$ | (b) $z \geq -1.25$ |
| (c) $z > 2.17$ | (d) $-1.25 < z < 2.17$ |

3.29 Standard Normal drill.

- (a) Find the number z such that the proportion of observations that are less than z in a standard Normal distribution is 0.6.
 (b) Find the number z such that 15% of all observations from a standard Normal distribution are greater than z .

3.30 Fruit flies. The thorax lengths in a population of male fruit flies follow a Normal distribution with mean 0.800 millimeters (mm) and standard deviation 0.078 mm.

- (a) What proportion of flies have thorax lengths less than 0.7 mm?
 (b) What proportion have thorax lengths greater than 1 mm?
 (c) What proportion have thorax lengths between 0.7 and 1 mm?

3.31 Acid rain? Emissions of sulfur dioxide by industry set off chemical changes in the atmosphere that result in “acid rain.” The acidity of liquids is measured by pH on a scale of 0 to 14.

Distilled water has pH 7.0, and lower pH values indicate acidity. Normal rain is somewhat acidic, so acid rain is sometimes defined as rainfall with a pH below 5.0. The pH of rain at one location varies among rainy days according to a Normal distribution with mean 5.43 and standard deviation 0.54. What proportion of rainy days have rainfall with pH below 5.0?

3.32 Runners. In a study of exercise, a large group of male runners walk on a treadmill for 6 minutes. Their heart rates in beats per minute at the end vary from runner to runner according to the $N(104, 12.5)$ distribution. The heart rates for male nonrunners after the same exercise have the $N(130, 17)$ distribution.

- (a) What percent of the runners have heart rates above 130?
 (b) What percent of the nonrunners have heart rates above 130?

3.33 A milling machine. Automated manufacturing operations are quite precise but still vary, often with distributions that are close to Normal.

The width in inches of slots cut by a milling machine follows approximately the $N(0.8750, 0.0012)$ distribution. The specifications allow slot widths between 0.8720 and 0.8780 inch. What proportion of slots meet these specifications?

3.34 Body mass index.

Your body mass index (BMI) is your weight in kilograms divided by the square of your height in meters. Many online BMI calculators allow you to enter weight in pounds and height in inches. High BMI is a common but controversial indicator of overweight or obesity. A study by the National Center for Health Statistics found that the BMI of American young women (ages 20 to 29) is approximately Normal with mean 26.5 and standard deviation 6.4.¹⁰

- (a) People with BMI less than 18.5 are often classified as “underweight.” What percent of young women are underweight by this criterion?
 (b) People with BMI over 30 are often classified as “obese.” What percent of young women are obese by this criterion?



Metalpix/Alamy

Miles per gallon. In its Fuel Economy Guide for model year 2010 vehicles, the Environmental Protection Agency gives data on 1101 vehicles. There are a number of high outliers, mainly hybrid gas-electric vehicles. If we ignore the vehicles identified as outliers, however, the combined city and highway gas mileage of the other 1082 vehicles is approximately Normal with mean 20.3 miles per gallon (mpg) and standard deviation 4.3 mpg. Exercises 3.35 to 3.38 concern this distribution.

3.35 In my Chevrolet. The 2010 Chevrolet Camaro with an eight-cylinder engine and automatic transmission has a combined gas mileage of 19 mpg. What percent of all vehicles have better gas mileage than the Camaro?

3.36 The bottom 10%. How low must a 2010 vehicle's gas mileage be in order to fall in the bottom 10% of all vehicles?

3.37 The middle half. The quartiles of any distribution are the values with cumulative proportions 0.25 and 0.75. They span the middle half of the distribution. What are the quartiles of the distribution of gas mileage?

3.38 Quintiles. The quintiles of any distribution are the values with cumulative proportions 0.20, 0.40, 0.60, and 0.80. What are the quintiles of the distribution of gas mileage?

3.39 What's your percentile? Reports on a student's ACT, SAT, or MCAT usually give the percentile as well as the actual score. The percentile is just the cumulative proportion stated as a percent: the percent of all scores that were lower than this one. In 2010, the total MCAT scores were close to Normal with mean 25.0 and standard deviation 6.4. William scored 32. What was his percentile?

3.40 Perfect SAT scores. It is possible to score higher than 1600 on the combined Mathematics and Reading portions of the SAT, but scores of 1600 and above are reported as 1600. The distribution of SAT scores (combining Mathematics and Reading) was close to Normal with mean 1021 and standard deviation 211. What proportion of SAT scores for these two parts were reported as 1600?

3.41 Heights of women. The heights of women aged 20 to 29 follow approximately the $N(64.3, 2.7)$ distribution. Men the same age have heights distributed as $N(69.9, 3.1)$. What percent of young women are taller than the mean height of young men?

3.42 Weights aren't Normal. The heights of people of the same sex and similar ages follow a Normal distribution reasonably closely. Weights, on the other hand, are not Normally distributed. The weights of women aged 20 to 29 have mean 155.9 pounds and median 144.0 pounds. The first and third quartiles are 124.1 pounds and 173.7 pounds. What can you say about the shape of the weight distribution? Why?

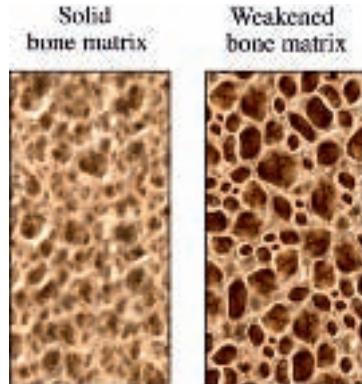
3.43 A surprising calculation. Changing the mean and standard deviation of a Normal distribution by a moderate amount can greatly change the percent of observations in the tails. Suppose that a college is looking for applicants with SAT Math scores of 750 and above.

(a) In 2010, the scores of men on the SAT Math test followed the $N(534, 118)$ distribution. What percent of men scored 750 or better?

(b) Women's SAT Math scores that year had the $N(500, 112)$ distribution. What percent of women scored 750 or better? You see that the percent of men above 750 is almost three times the percent of women with such high scores. Why this is true is controversial. (On the other hand, women score higher than men on the new SAT Writing test, though by a smaller amount.)

3.44 Grading managers. Some companies "grade on a bell curve" to compare the performance of their managers and professional workers. This forces the use of some low performance ratings so that not all workers are listed as "above average." Ford Motor Company's "performance management process" for this year assigned 10% A grades, 80% B grades, and 10% C grades to the company's managers. Suppose that Ford's performance scores really are Normally distributed. This year, managers with scores less than 25 received C's and those with scores above 475 received A's. What are the mean and standard deviation of the scores?

3.45 Osteoporosis. Osteoporosis is a condition in which the bones become brittle due to loss of minerals. To diagnose osteoporosis, an elaborate apparatus measures bone mineral density (BMD). BMD is usually reported in standardized form. The standardization is based on a population of healthy young adults. The World Health Organization (WHO) criterion for osteoporosis is a BMD 2.5 standard deviations below the mean for young adults. BMD measurements in a population of people similar in age and sex roughly follow a Normal distribution.



Nucleus Medical Art, Inc/Phototake—All rights reserved.

(a) What percent of healthy young adults have osteoporosis by the WHO criterion?

(b) Women aged 70 to 79 are of course not young adults. The mean BMD in this age is about -2 on the standard scale

for young adults. Suppose that the standard deviation is the same as for young adults. What percent of this older population have osteoporosis?

In later chapters we will meet many statistical procedures that work well when the data are “close enough to Normal.” Exercises 3.46 to 3.50 concern data that are mostly close enough to Normal for statistical work, while Exercise 3.51 concerns data for which the data are not close to Normal. These exercises ask you to do data analysis and Normal calculations to investigate how close to Normal real data are.

3.46 Normal is only approximate: IQ test scores. Here are the IQ test scores of 31 seventh-grade girls in a Midwest school district:¹¹  MIDWESTIQ

114	100	104	89	102	91	114	114	103	105	108
130	120	132	111	128	118	119	86	72	111	103
74	112	107	103	98	96	112	112	93		

(a) We expect IQ scores to be approximately Normal. Make a stemplot to check that there are no major departures from Normality.

(b) Nonetheless, proportions calculated from a Normal distribution are not always very accurate for small numbers of observations. Find the mean \bar{x} and standard deviation s for these IQ scores. What proportion of the scores are within one standard deviation of the mean. Within two standard deviations of the mean? What would these proportions be in an exactly Normal distribution?

3.47 Normal is only approximate: ACT scores. Scores on the ACT test for the 2010 high school graduating class had mean 21.0 and standard deviation 5.2. In all, 1,568,835 students in this class took the test. Of these, 145,000 had scores higher than 28 and another 50,860 had scores exactly 28. ACT scores are always whole numbers. The exactly Normal $N(21.0, 5.2)$ distribution can include any value, not just whole numbers. What is more, there is no area exactly above 28 under the smooth Normal curve. So ACT scores can be only approximately Normal. To illustrate this fact, find

- (a) the percent of 2010 ACT scores greater than 28.
- (b) the percent of 2010 ACT scores greater than or equal to 28.
- (c) the percent of observations from the $N(21.0, 5.2)$ distribution that are greater than 28. (The percent greater than or equal to 28 is the same, because there is no area exactly above 28.)

3.48 Are the data Normal? Acidity of rainfall. Exercise 3.31 concerns the acidity (measured by pH) of rainfall. A

sample of 105 rainwater specimens had mean pH 5.43, standard deviation 0.54, and five-number summary 4.33, 5.05, 5.44, 5.79, 6.81.¹²

- (a) Compare the mean and median and also the distances of the two quartiles from the median. Does it appear that the distribution is quite symmetric? Why?
- (b) If the distribution is really $N(5.43, 0.54)$, what proportion of observations would be less than 5.05? Less than 5.79? Do these proportions suggest that the distribution is close to Normal? Why?

3.49 Are the data Normal? SAT Critical Reading scores. Georgia Southern University (GSU) had 2417 students with regular admission in their freshman class of 2010. For each student, data are available on their SAT and ACT scores, if taken, high school GPA, and the college within the university to which they were admitted.¹³ Here are the first 20 SAT Critical Reading scores from that data set:  SATCR

650	490	580	450	570	540	510	530	510	560
560	590	470	690	530	570	460	590	530	490

The complete data set is on the text Web site and CD, which contains both the original scores and the ordered scores.

- (a) Make a histogram of the distribution (if your software allows it, superimpose a Normal curve over the histogram as in Figure 3.1). Although the resulting histogram depends a bit on your choice of classes, the distribution appears roughly symmetric with no outliers.
- (b) Find the mean, median, standard deviation, and quartiles for these data. Comparing the mean and the median and comparing the distances of the two quartiles from the median suggest that the distribution is quite symmetric. Why?
- (c) In 2010, the mean score on the Critical Reading portion of the SAT for all college-bound seniors was 501. If the distribution were exactly Normal with the mean and standard deviation you found in part (b), what proportion of regularly admitted GSU freshmen scored above the mean for all college-bound seniors?
- (d) Compute the exact proportion of regularly admitted GSU freshmen who scored above the mean for all college-bound seniors. It will be simplest to use the ordered scores in the data file to calculate this. How does this percent compare with the percent calculated in part (c)? Despite the discrepancy, this distribution is “close enough to Normal” for statistical work in later chapters.

3.50 Are the data Normal? Monsoon rains. Here are the amounts of summer monsoon rainfall (millimeters) for India in the 100 years from 1901 to 2000:¹⁴  MONSOONS

722.4	792.2	861.3	750.6	716.8	885.5	777.9	897.5	889.6	935.4
736.8	806.4	784.8	898.5	781.0	951.1	1004.7	651.2	885.0	719.4
866.2	869.4	823.5	863.0	804.0	903.1	853.5	768.2	821.5	804.9
877.6	803.8	976.2	913.8	843.9	908.7	842.4	908.6	789.9	853.6
728.7	958.1	868.6	920.8	911.3	904.0	945.9	874.3	904.2	877.3
739.2	793.3	923.4	885.8	930.5	983.6	789.0	889.6	944.3	839.9
1020.5	810.0	858.1	922.8	709.6	740.2	860.3	754.8	831.3	940.0
887.0	653.1	913.6	748.3	963.0	857.0	883.4	909.5	708.0	882.9
852.4	735.6	955.9	836.9	760.0	743.2	697.4	961.7	866.9	908.8
784.7	785.0	896.6	938.4	826.4	857.3	870.5	873.8	827.0	770.2



John Henry Claude Wilson/Photolibrary

(a) Make a histogram of these rainfall amounts. Find the mean and the median.

(b) Although the distribution is reasonably Normal, your work shows some departure from Normality. In what way are the data not Normal?

3.51 Are the data Normal? Canadians' earnings in 1901.

Table 2.5 (page 66) gives data on the earnings of a sample of 200 Canadians in 1901. The mean of the earnings is \$350.30 and the standard deviation is \$292.20. Figure 3.15 gives a histogram of the data along with a smooth curve representing an $N(350.30, 292.20)$ distribution. From the figure, the Normal curve does not appear to follow the pattern in the histogram that closely. Because of this, the use of

areas under the Normal curve may not provide a good approximation to incomes in various intervals. CANADA EARN

(a) Referring to Table 2.5, what proportion of earnings are above \$375? What percent of the $N(350.30, 292.20)$ distribution is above \$375?

(b) There are no negative incomes. What percent of the $N(350.30, 292.20)$ distribution is below zero?

(c) Based on your answers in (a) and (b), do you think it is a good idea to summarize the distribution of incomes by an $N(350.30, 292.20)$ distribution?

The *Normal Curve applet* allows you to do Normal calculations quickly. It is somewhat limited by the number of pixels available for use, so that it can't hit every value exactly. In the exercises below, use the closest available values. In each case, make a sketch of the curve from the applet marked with the values you used to answer the questions.

3.52 How accurate is 68–95–99.7? The 68–95–99.7 rule for Normal distributions is a useful approximation.

To see how accurate the rule is, drag one flag across the other so that the applet shows the area under the curve between the two flags.

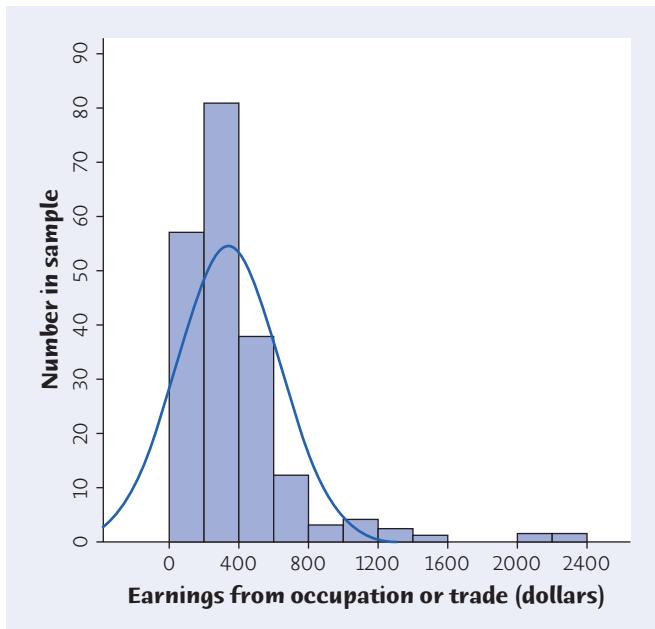
(a) Place the flags one standard deviation on either side of the mean. What is the area between these two values? What does the 68–95–99.7 rule say this area is?

(b) Repeat for locations two and three standard deviations on either side of the mean. Again compare the 68–95–99.7 rule with the area given by the applet.

3.53 Where are the quartiles? How many standard deviations above and below the mean do the quartiles of any Normal distribution lie? (Use the standard Normal distribution to answer this question.)

3.54 Grading managers. In Exercise 3.44, we saw that Ford

Motor Company once graded its managers in such a way that the top 10% received an A grade, the bottom 10% a C, and the middle 80% a B. Let's suppose that performance scores follow a Normal distribution. How many standard deviations above and below the mean do the A/B and B/C cutoffs lie? (Use the standard Normal distribution to answer this question.)

**FIGURE 3.15**

Histogram of the earnings of 200 Canadians in 1901, with a Normal curve superimposed, for Exercise 3.51.



EXPLORING THE WEB

3.55 Are the data Normal? Comparing quartiles. The Web site <http://professionals.collegeboard.com/data-reports-research/sat> presents data for high school seniors who participated in the SAT Program during the current year as well as previous years. Under *SAT Data & Reports*, click on the link for *College Bound Seniors* for the most recent year given. In the window that opens, click on the link for *Total Group Report: College Bound Seniors* for this year. The Total Group Profile will open and contains several tables, each giving different summary information. Go to the Overall Mean Scores table. How many students took the Critical Reading portion of the SAT? What were the mean and standard deviation of the scores? Assuming that the distribution of scores is Normal with the mean and standard deviation given in the Overall Mean Scores table, what are the first and third quartiles of the distribution? Now go to the Percentiles for the Total Group table, and compare the actual first and third quartiles from the data to the values obtained from the Normal curve. Does this give any evidence that the distribution of Critical Reading scores is not Normal?

3.56 Are the data normal? Comparing proportions. As in Exercise 3.55, this exercise compares the distribution of Critical Reading Scores to a Normal distribution, but uses a different criterion. Instead of comparing quartiles, you will compare the proportions of individuals obtaining scores in several intervals with the corresponding areas under the Normal curve. Follow the directions described in the previous exercise for accessing the College Board website and opening the Total Group Profile. Go to the Overall Mean Scores table. How many students took the Critical Reading portion of the SAT? What were the mean and standard deviation of the scores? Now go to the table on the Score Distribution. The scores are broken into intervals 200–290, 300–390, etc. Using the Total column, find the actual percentage of students who scored in each of the six intervals reported for the Critical Reading portion of the SAT. Assuming the distribution of scores is Normal with the mean and standard deviation given in the Overall Mean Scores table, find the area under the Normal curve for each interval. Because of the discreteness of the actual scores, take the interval 200–290 as the interval 200–300, the interval 300–390 as the interval 300–400, etc., when finding the areas under the Normal curve. How do the actual percentages compare to the areas under the Normal curve? Does this give any evidence that the distribution of Critical Reading scores is not Normal?



Scatterplots and Correlation

Chapter 4

A medical study finds that short women are more likely to have heart attacks than women of average height, while tall women have the fewest heart attacks. An insurance group reports that heavier cars have fewer deaths per 10,000 vehicles registered than do lighter cars. These and many other statistical studies look at the *relationship between two variables*. Statistical relationships are overall tendencies, not ironclad rules. They allow individual exceptions. Although smokers on the average die younger than nonsmokers, some people live to 90 while smoking three packs a day.

To understand a statistical relationship between two variables, we measure both variables on the same individuals. Often, we must examine other variables as well. To conclude that shorter women have higher risk from heart attacks, for example, the researchers had to eliminate the effect of other variables such as weight and exercise habits. In this and the following chapter we study relationships between variables. One of our main themes is that the relationship between two variables can be strongly influenced by other variables that are lurking in the background.

EXPLANATORY AND RESPONSE VARIABLES

We think that car weight helps explain accident deaths and that smoking influences life expectancy. In each of these relationships, the two variables play different roles: one explains or influences the other.

RESPONSE VARIABLE, EXPLANATORY VARIABLE

A **response variable** measures an outcome of a study. An **explanatory variable** may explain or influence changes in a response variable.

IN THIS CHAPTER WE COVER...

- Explanatory and response variables
- Displaying relationships: scatterplots
- Interpreting scatterplots
- Adding categorical variables to scatterplots
- Measuring linear association: correlation
- Facts about correlation

You will often find explanatory variables called *independent variables* or *predictor variables* and response variables called *dependent variables*. The idea behind this language is that the response variable depends on the explanatory variable. Because “independent” and “dependent” have other meanings in statistics that are unrelated to the explanatory-response distinction, we prefer to avoid those words.

It is easiest to identify explanatory and response variables when we actually set values of one variable in order to see how it affects another variable.

EXAMPLE 4.1 Beer and blood alcohol

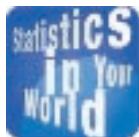
How does drinking beer affect the level of alcohol in our blood? The legal limit for driving in all states is 0.08%. Student volunteers at the Ohio State University drank different numbers of cans of beer. Thirty minutes later, a police officer measured their blood alcohol content. Number of beers consumed is the explanatory variable, and percent of alcohol in the blood is the response variable. ■

When we don’t set the values of either variable but just observe both variables, there may or may not be explanatory and response variables. Whether there are depends on how we plan to use the data.

EXAMPLE 4.2 College debts

A college student aid officer looks at the findings of the National Student Loan Survey. She notes data on the amount of debt of recent graduates, their current income, and how stressed they feel about college debt. She isn’t interested in predictions but is simply trying to understand the situation of recent college graduates. The distinction between explanatory and response variables does not apply.

A sociologist looks at the same data with an eye to using amount of debt and income, along with other variables, to explain the stress caused by college debt. Now amount of debt and income are explanatory variables and stress level is the response variable. ■



After you plot your data, think!

The statistician Abraham Wald (1902–1950) worked on war problems during World War II. Wald invented some statistical methods that were military secrets until the war ended. Here is one of his simpler ideas. Asked where extra armor should be added to airplanes, Wald studied the location of enemy bullet holes in planes returning from combat. He plotted the locations on an outline of the plane. As data accumulated, most of the outline filled up. Put the armor in the few spots with no bullet holes, said Wald. That’s where bullets hit the planes that didn’t make it back.

In many studies, the goal is to show that changes in one or more explanatory variables actually *cause* changes in a response variable. Other explanatory-response relationships do not involve direct causation. Nations with more television sets per person have greater life expectancies, but shipping many television sets to Botswana won’t *cause* life expectancy to increase.

Most statistical studies examine data on more than one variable. Fortunately, statistical analysis of several-variable data builds on the tools we used to examine individual variables. The principles that guide our work also remain the same:

- Plot your data. Look for overall patterns and deviations from those patterns.
- Based on what your plot shows, choose numerical summaries for some aspects of the data.

APPLY YOUR KNOWLEDGE

4.1 Explanatory and response variables? You have data on a large group of college students. Here are four pairs of variables measured on these students. For each pair, is it more reasonable to simply explore the relationship between the two variables or to view one of the variables as an explanatory variable and the other as a response variable? In the latter case, which is the explanatory variable and which is the response variable?

- (a) Amount of time spent studying for a statistics exam and grade on the exam.
- (b) Weight in kilograms and height in centimeters.
- (c) Hours per week spent online using Facebook and grade point average.
- (d) Score on the SAT Writing exam and score on the SAT Critical Reading exam.

4.2 Coral reefs. How sensitive to changes in water temperature are coral reefs? To find out, scientists examined data on sea surface temperatures and coral growth per year at locations in the Red Sea.¹ What are the explanatory and response variables? Are they categorical or quantitative?

4.3 Beer and blood alcohol. Example 4.1 describes a study in which college students drank different amounts of beer. The response variable was their blood alcohol content (BAC). BAC for the same amount of beer might depend on other facts about the students. Name two other variables that could influence BAC.



Georgette Douwma/Getty

DISPLAYING RELATIONSHIPS: SCATTERPLOTS

The most useful graph for displaying the relationship between two quantitative variables is a *scatterplot*.

EXAMPLE 4.3 State SAT Mathematics scores

Figure 1.8 (page 18) reminded us that in some states most high school graduates take the SAT test of readiness for college, and in other states most take the ACT. Who takes a test may influence the average score. Let's follow our four-step process (page 55) to examine this influence.²

STATE: The percent of high school students who take the SAT varies from state to state. Does this fact help explain differences among the states in average SAT Mathematics score?

PLAN: Examine the relationship between percent taking the SAT and state mean score on the Mathematics part of the SAT. Choose the explanatory and response variables. Make a scatterplot to display the relationship between the variables. Interpret the plot to understand the relationship.

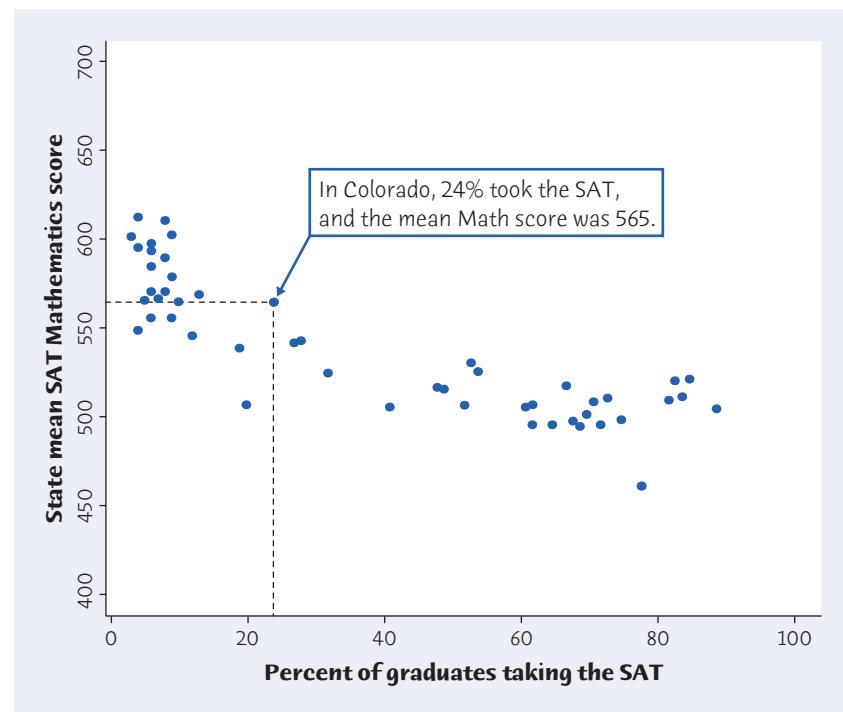
SOLVE (make the plot): We suspect that “percent taking” will help explain “mean score.” So “percent taking” is the explanatory variable and “mean score” is the response variable. We want to see how mean score changes when percent taking changes, so we put percent taking (the explanatory variable) on the horizontal axis. Figure 4.1 is the



MATHSAT

FIGURE 4.1

Scatterplot of the mean SAT Mathematics score in each state against the percent of that state's high school graduates who take the SAT, for Example 4.3. The dotted lines intersect at the point (24, 565), the data for Colorado.



scatterplot. Each point represents a single state. In Colorado, for example, 24% took the SAT, and their mean SAT Math score was 565. Find 24 on the x (horizontal) axis and 565 on the y (vertical) axis. Colorado appears as the point (24, 565) above 24 and to the right of 565.

CONCLUDE: We will explore conclusions in Example 4.4. ■

SCATTERPLOT

A **scatterplot** shows the relationship between two quantitative variables measured on the same individuals. The values of one variable appear on the horizontal axis, and the values of the other variable appear on the vertical axis. Each individual in the data appears as the point in the plot fixed by the values of both variables for that individual.

Always plot the explanatory variable, if there is one, on the horizontal axis (the x axis) of a scatterplot. As a reminder, we usually call the explanatory variable x and the response variable y . If there is no explanatory-response distinction, either variable can go on the horizontal axis.

APPLY YOUR KNOWLEDGE

4.4 Do heavier people burn more energy? Metabolic rate, the rate at which the body consumes energy, is important in studies of weight gain, dieting, and exercise. We have data on the lean body mass and resting metabolic rate for 12 women who are subjects in a study of dieting. Lean body mass, given in kilograms, is a person's

weight leaving out all fat. Metabolic rate is measured in calories burned per 24 hours, the same calories used to describe the energy content of foods.



Mass	36.1	54.6	48.5	42.0	50.6	42.0	40.3	33.1	42.4	34.5	51.1	41.2
Rate	995	1425	1396	1418	1502	1256	1189	913	1124	1052	1347	1204

The researchers believe that lean body mass is an important influence on metabolic rate. Make a scatterplot to examine this belief. (The *Two-Variable Statistical Calculator Applet* provides an easy way to make scatterplots. Click “Data” to enter your data, then “Scatterplot” to see the plot.)



4.5 Outsourcing by airlines. Airlines have increasingly outsourced the maintenance of their planes to other companies. A concern voiced by critics is that the maintenance may be less carefully done, so that outsourcing creates a safety hazard. In addition, flight delays are often due to maintenance problems, so one might look at government data on percent of major maintenance outsourced and percent of flight delays blamed on the airline to determine if these concerns are justified. This was done, and data from 2005 and 2006 appeared to justify the concerns of the critics. Do more recent data still support the concerns of the critics? Here are data from 2009:³



Airline	Outsource percent	Delay percent	Airline	Outsource percent	Delay percent
AirTran	52.8	24.21	Hawaiian	74.1	7.94
Alaska	56.8	17.09	JetBlue	53.7	22.55
American	23.3	22.71	Northwest	59.8	20.85
Continental	44.5	21.23	Southwest	61.7	17.00
Delta	25.6	21.44	United	40.6	19.02
Frontier	26.8	21.70	US Airways	60.4	19.13

Make a scatterplot that shows the relation between delays and outsourcing.

INTERPRETING SCATTERPLOTS

To interpret a scatterplot, adapt the strategies of data analysis learned in Chapters 1 and 2 to the new two-variable setting.

EXAMINING A SCATTERPLOT

In any graph of data, look for the **overall pattern** and for striking **deviations** from that pattern.

You can describe the overall pattern of a scatterplot by the **direction**, **form**, and **strength** of the relationship.

An important kind of deviation is an **outlier**, an individual value that falls outside the overall pattern of the relationship.



clusters

EXAMPLE 4.4 Understanding state SAT scores

We continue to explore the state SAT Mathematics scores by interpreting what the scatterplot tells us about the variation in scores from state to state.

SOLVE (INTERPRET THE PLOT): Figure 4.1 shows a clear *direction*: the overall pattern moves from upper left to lower right. That is, states in which higher percents of high school graduates take the SAT tend to have lower mean SAT Mathematics scores. We call this a *negative association* between the two variables.

The *form* of the relationship is roughly a straight line with a slight curve to the right as it moves down. What is more, most states fall into two distinct **clusters**. As in the histogram in Figure 1.8, the ACT states cluster at the left and the SAT states at the right. In 22 states, fewer than 20% of seniors took the SAT; in another 22 states, more than 50% took the SAT.

The *strength* of a relationship in a scatterplot is determined by how closely the points follow a clear form. The overall relationship in Figure 4.1 is moderately strong: states with similar percents taking the SAT tend to have roughly similar mean SAT Math scores.

CONCLUDE: Percent taking explains much of the variation among states in average SAT Mathematics score. States in which a higher percent of students take the SAT tend to have lower mean scores because the mean includes a broader group of students. SAT states as a group have lower mean SAT scores than ACT states. So average SAT score says almost nothing about the quality of education in a state. It is foolish to “rank” states by their average SAT scores. ■

POSITIVE ASSOCIATION, NEGATIVE ASSOCIATION

Two variables are **positively associated** when above-average values of one tend to accompany above-average values of the other, and below-average values also tend to occur together.

Two variables are **negatively associated** when above-average values of one tend to accompany below-average values of the other, and vice versa.



Douglas Faulkner/Photo Researchers



EXAMPLE 4.5 The endangered manatee

STATE: Manatees are large, gentle, slow-moving creatures found along the coast of Florida. Many manatees are injured or killed by boats. Table 4.1 contains data on the number of boats registered in Florida (in thousands) and the number of manatees killed by boats for the years between 1977 and 2009.⁴ Examine the relationship. Is it plausible that restricting the number of boats would help protect manatees?

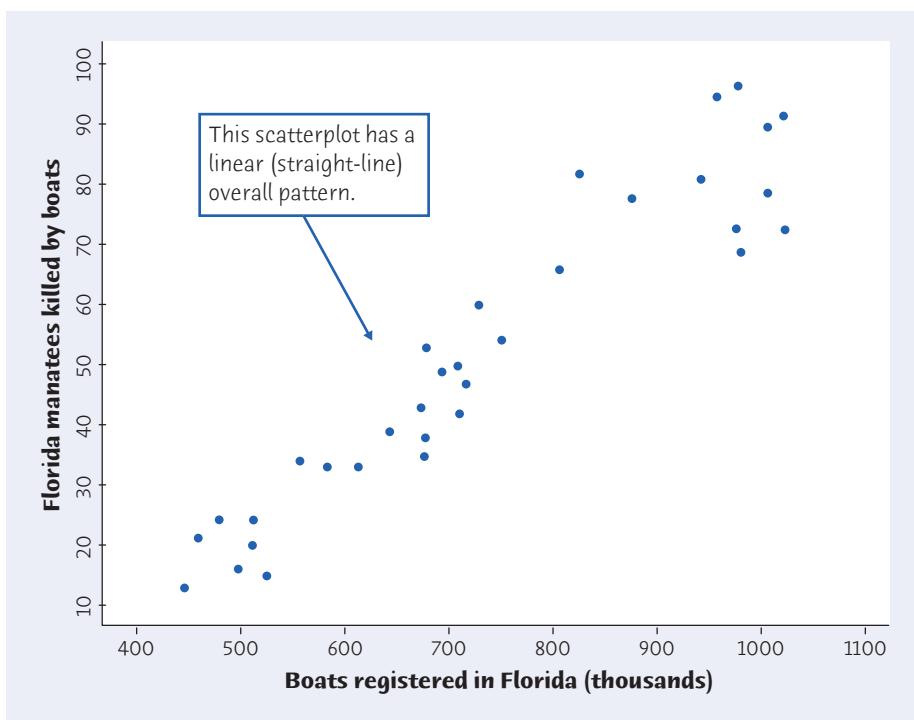
TABLE 4.1 Florida boat registrations (thousands) and manatees killed by boats

YEAR	BOATS	MANATEES	YEAR	BOATS	MANATEES	YEAR	BOATS	MANATEES
1977	447	13	1989	711	50	2001	944	81
1978	460	21	1990	719	47	2002	962	95
1979	481	24	1991	681	53	2003	978	73
1980	498	16	1992	679	38	2004	983	69
1981	513	24	1993	678	35	2005	1010	79
1982	512	20	1994	696	49	2006	1024	92
1983	526	15	1995	713	42	2007	1027	73
1984	559	34	1996	732	60	2008	1010	90
1985	585	33	1997	755	54	2009	982	97
1986	614	33	1998	809	66			
1987	645	39	1999	830	82			
1988	675	43	2000	880	78			

PLAN: Make a scatterplot with “boats registered” as the explanatory variable and “manatees killed” as the response variable. Describe the form, direction, and strength of the relationship.

SOLVE: Figure 4.2 is the scatterplot. There is a positive association—more boats goes with more manatees killed. The form of the relationship is **linear**. That is, the overall pattern follows a straight line from lower left to upper right. The relationship is strong because the points don’t deviate greatly from a line.

linear relationship

**FIGURE 4.2**

Scatterplot of the number of Florida manatees killed by boats in the years 1977 to 2009 against the number of boats registered in Florida that year, for Example 4.5. There is a strong linear (straight-line) pattern.

CONCLUDE: As more boats are registered, the number of manatees killed by boats goes up linearly. Data from the Florida Wildlife Commission indicate that in recent years boats accounted for 24% of manatee deaths and 31% of deaths whose causes could be determined. Although many manatees die from other causes, it appears that fewer boats would mean fewer manatee deaths. ■



As the following chapter will emphasize, *it is wise to always ask what other variables lurking in the background might contribute to the relationship displayed in a scatterplot.* Because both boats registered and manatees killed are recorded year by year, any change in conditions over time might affect the relationship. For example, if boats in Florida have tended to go faster over the years, that might result in more manatees killed by the same number of boats.

APPLY YOUR KNOWLEDGE

4.6 Do heavier people burn more energy? Describe the direction, form, and strength of the relationship between lean body mass and metabolic rate, as displayed in your plot for Exercise 4.4. METABOLIC

4.7 Outsourcing by airlines. Does your plot for Exercise 4.5 show a positive, negative, or no association between maintenance outsourcing and delays caused by the airline? One airline is a low outlier in delay percent. Which airline is this? Aside from the outlier, does the plot show a roughly linear form? If it does, is the relationship very strong? AIRLINES

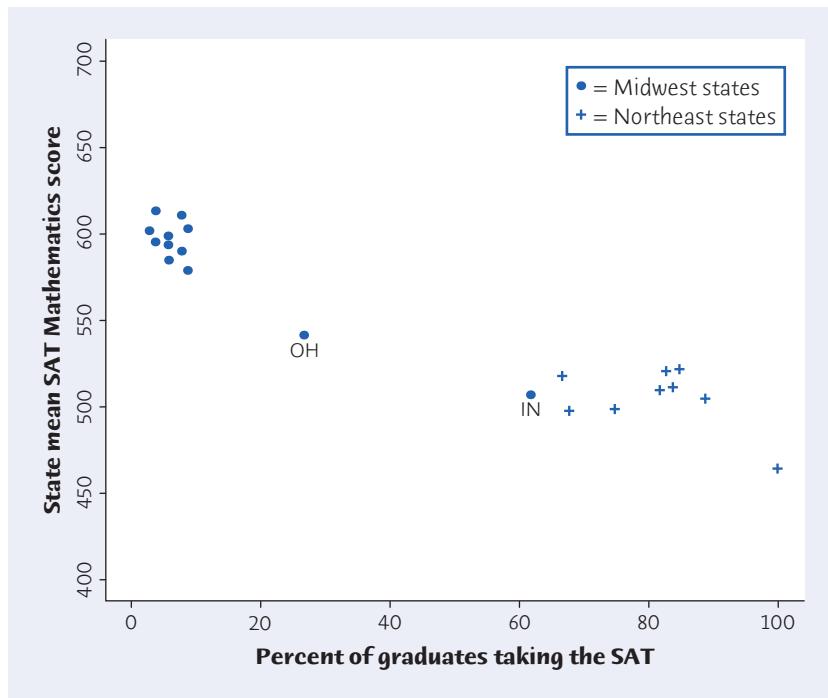
4.8 Does fast driving waste fuel? How does the fuel consumption of a car change as its speed increases? Here are data for a British Ford Escort. Speed is measured in kilometers per hour, and fuel consumption is measured in liters of gasoline used per 100 kilometers traveled.⁵ FASTDRIVE

Speed	10	20	30	40	50	60	70	80
Fuel	21.00	13.00	10.00	8.00	7.00	5.90	6.30	6.95
Speed	90	100	110	120	130	140	150	
Fuel	7.57	8.27	9.03	9.87	10.79	11.77	12.83	

- (a) Make a scatterplot. (Which is the explanatory variable?)
- (b) Describe the form of the relationship. It is not linear. Explain why the form of the relationship makes sense.
- (c) It does not make sense to describe the variables as either positively associated or negatively associated. Why?
- (d) Is the relationship reasonably strong or quite weak? Explain your answer.

ADDING CATEGORICAL VARIABLES TO SCATTERPLOTS

The Census Bureau groups the states into four broad regions, named Midwest, Northeast, South, and West. We might ask about regional patterns in SAT exam scores. Figure 4.3 repeats part of Figure 4.1, with an important difference. We

**FIGURE 4.3**

Mean SAT Mathematics score and percent of high school graduates who take the test for only the Midwest (•) and Northeast (+) states.

have plotted only the Midwest and Northeast groups of states, using the plot symbol “•” for the Midwest states and the symbol “+” for the Northeast states.

The regional comparison is striking. The 9 Northeast states are all SAT states—at least 67% of high school graduates in each of these states take the SAT. The 12 Midwest states are mostly ACT states. In 10 of these states, fewer than 10% of high school graduates take the SAT. One Midwest state is clearly an outlier within the region: Indiana is an SAT state (62% take the SAT) that falls close to the Northeast cluster. Ohio, where 27% take the SAT, also lies outside the Midwest cluster.

Dividing the states into regions introduces a third variable into the scatterplot. “Region” is a categorical variable that has four values, although we plotted data from only two of the four regions. The two regions are identified by the two different plotting symbols.

CATEGORICAL VARIABLES IN SCATTERPLOTS

To add a categorical variable to a scatterplot, use a different plot color or symbol for each category.

APPLY YOUR KNOWLEDGE

- 4.9 Do heavier people burn more energy?** The study of dieting described in Exercise 4.4 collected data on the lean body mass (in kilograms) and metabolic rate (in calories) for both female and male subjects: METABOLIC2

Sex	F	F	F	F	F	F	F	F	F	F	F
Mass	36.1	54.6	48.5	42.0	50.6	42.0	40.3	33.1	42.4	34.5	
Rate	995	1425	1396	1418	1502	1256	1189	913	1124	1052	
Sex	F	F	M	M	M	M	M	M	M	M	
Mass	51.1	41.2	51.9	46.9	62.0	62.9	47.4	48.7	51.9		
Rate	1347	1204	1867	1439	1792	1666	1322	1614	1460		

- (a) Make a scatterplot of metabolic rate versus lean body mass for all 19 subjects. Use separate symbols to distinguish women and men.
- (b) Does the same overall pattern hold for both women and men? What is the most important difference between women and men?

MEASURING LINEAR ASSOCIATION: CORRELATION

A scatterplot displays the direction, form, and strength of the relationship between two quantitative variables. Linear (straight-line) relations are particularly important because a straight line is a simple pattern that is quite common. A linear relation is strong if the points lie close to a straight line, and weak if they are widely scattered about a line. Our eyes are not good judges of how strong a linear relationship is. The two scatterplots in Figure 4.4 depict exactly the same data, but the lower plot is drawn smaller in a large field. The lower plot seems to show a stronger linear relationship. Our eyes can be fooled by changing the plotting scales or the amount of space around the cloud of points in a scatterplot.⁶ We need to follow our strategy for data analysis by using a numerical measure to supplement the graph. Correlation is the measure we use.

CORRELATION

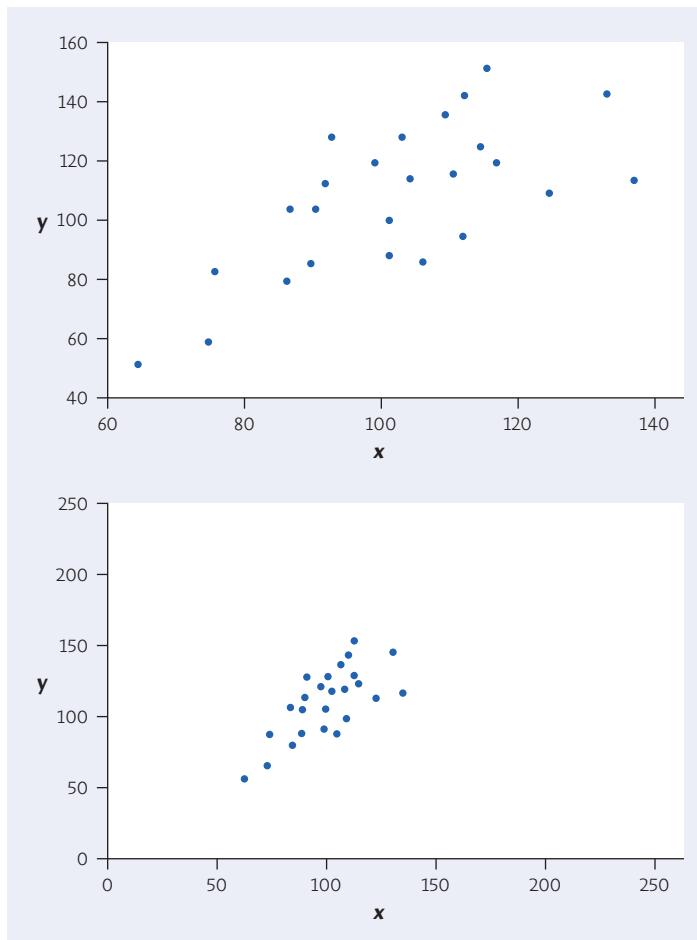
The **correlation** measures the direction and strength of the linear relationship between two quantitative variables. Correlation is usually written as r .

Suppose that we have data on variables x and y for n individuals. The values for the first individual are x_1 and y_1 , the values for the second individual are x_2 and y_2 , and so on. The means and standard deviations of the two variables are \bar{x} and s_x for the x -values, and \bar{y} and s_y for the y -values. The correlation r between x and y is

$$r = \frac{1}{n-1} \left[\left(\frac{x_1 - \bar{x}}{s_x} \right) \left(\frac{y_1 - \bar{y}}{s_y} \right) + \left(\frac{x_2 - \bar{x}}{s_x} \right) \left(\frac{y_2 - \bar{y}}{s_y} \right) + \dots + \left(\frac{x_n - \bar{x}}{s_x} \right) \left(\frac{y_n - \bar{y}}{s_y} \right) \right]$$

or, more compactly,

$$r = \frac{1}{n-1} \sum \left(\frac{x_i - \bar{x}}{s_x} \right) \left(\frac{y_i - \bar{y}}{s_y} \right)$$

**FIGURE 4.4**

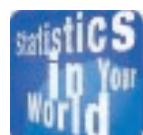
Two scatterplots of the same data. The straight-line pattern in the lower plot appears stronger because of the surrounding space.

The formula for the correlation r is a bit complex. It helps us see what correlation is, but in practice you should use software or a calculator that finds r from keyed-in values of two variables x and y . Exercise 4.10 asks you to calculate a correlation step-by-step from the definition to solidify its meaning.

The formula for r begins by standardizing the observations. Suppose, for example, that x is height in centimeters and y is weight in kilograms and that we have height and weight measurements for n people. Then \bar{x} and s_x are the mean and standard deviation of the n heights, both in centimeters. The value

$$\frac{x_i - \bar{x}}{s_x}$$

is the standardized height of the i th person, from Chapter 3. The standardized height says how many standard deviations above or below the mean a person's height lies. Standardized values have no units—in this example, they are no longer measured in centimeters. Standardize the weights also. The correlation r is an average of the products of the standardized height and the standardized weight for all the individuals. Just as in the case of the standard deviation s , the “average” here divides by one fewer than the number of individuals.



Death from superstition?

Is there a relationship between

superstitious beliefs and bad things happening? Apparently there is. Chinese and Japanese people think that the number 4 is unlucky because when pronounced it sounds like the word for “death.” Sociologists looked at 15 years’ worth of death certificates for Chinese and Japanese Americans and for white Americans. Deaths from heart disease were notably higher on the fourth day of the month among Chinese and Japanese but not among whites. The sociologists think the explanation is increased stress on “unlucky days.”



APPLY YOUR KNOWLEDGE

4.10 Coral reefs. Exercise 4.2 discusses a study in which scientists examined data on mean sea surface temperatures (in degrees Celsius) and mean coral growth (in millimeters per year) over a several-year period at locations in the Red Sea. Here are the data:⁷



Sea surface temperature	29.68	29.87	30.16	30.22	30.48	30.65	30.90
Growth	2.63	2.58	2.60	2.48	2.26	2.38	2.26

- (a) Make a scatterplot. Which is the explanatory variable? The plot shows a negative linear pattern.
- (b) Find the correlation r step-by-step. You may wish to round off to two decimal places in each step. First find the mean and standard deviation of each variable. Then find the seven standardized values for each variable. Finally, use the formula for r . Explain how your value for r matches your graph in (a).
- (c) Enter these data into your calculator or software and use the correlation function to find r . Check that you get the same result as in (b), up to roundoff error.

FACTS ABOUT CORRELATION

The formula for correlation helps us see that r is positive when there is a positive association between the variables. Height and weight, for example, have a positive association. People who are above average in height tend to also be above average in weight. Both the standardized height and the standardized weight are positive. People who are below average in height tend to also have below-average weight. Then both standardized height and standardized weight are negative. In both cases, the products in the formula for r are mostly positive and so r is positive. In the same way, we can see that r is negative when the association between x and y is negative. More detailed study of the formula gives more detailed properties of r . Here is what you need to know in order to interpret correlation.

1. *Correlation makes no distinction between explanatory and response variables.* It makes no difference which variable you call x and which you call y in calculating the correlation.
2. *Because r uses the standardized values of the observations, r does not change when we change the units of measurement of x , y , or both.* Measuring height in inches rather than centimeters and weight in pounds rather than kilograms does not change the correlation between height and weight. The correlation r itself has no unit of measurement; it is just a number.
3. *Positive r indicates positive association between the variables, and negative r indicates negative association.*
4. *The correlation r is always a number between -1 and 1 .* Values of r near 0 indicate a very weak linear relationship. The strength of the linear relationship increases as r moves away from 0 toward either -1 or 1 . Values of r close

to -1 or 1 indicate that the points in a scatterplot lie close to a straight line. The extreme values $r = -1$ and $r = 1$ occur only in the case of a perfect linear relationship, when the points lie exactly along a straight line.

EXAMPLE 4.6 From scatterplot to correlation

The scatterplots in Figure 4.5 illustrate how values of r closer to 1 or -1 correspond to stronger linear relationships. To make the meaning of r clearer, the standard deviations of both variables in these plots are equal, and the horizontal and vertical scales are the same. In general, it is not so easy to guess the value of r from the appearance of a scatterplot. Remember that changing the plotting scales in a scatterplot may mislead our eyes, but it does not change the correlation.

The scatterplots in Figure 4.6 show four sets of real data. The patterns are less regular than those in Figure 4.5, but they also illustrate how correlation measures the strength of linear relationships.⁸

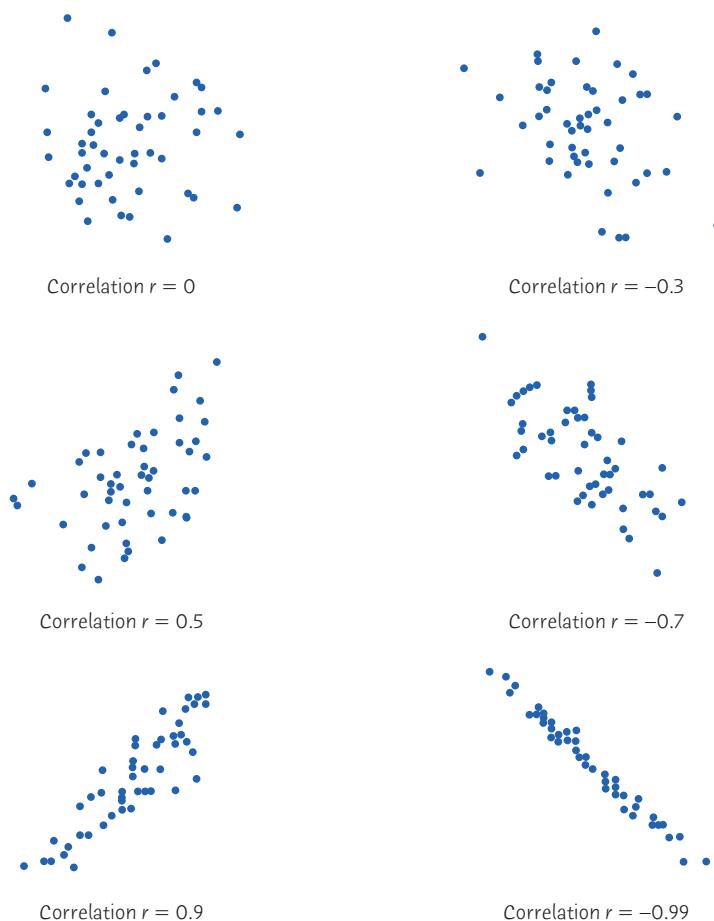
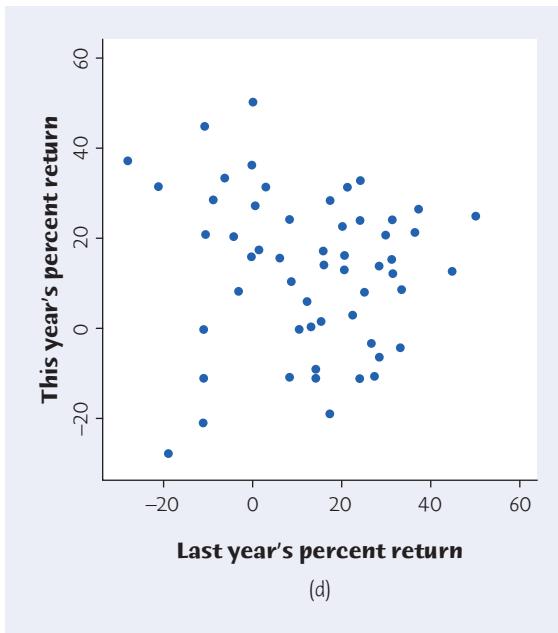
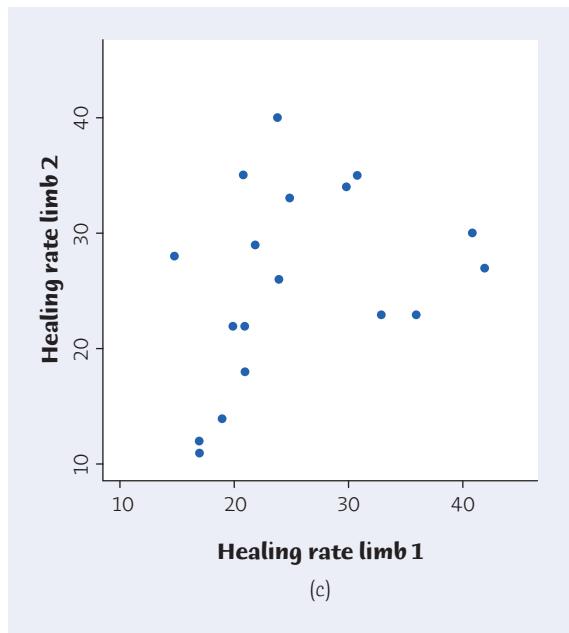
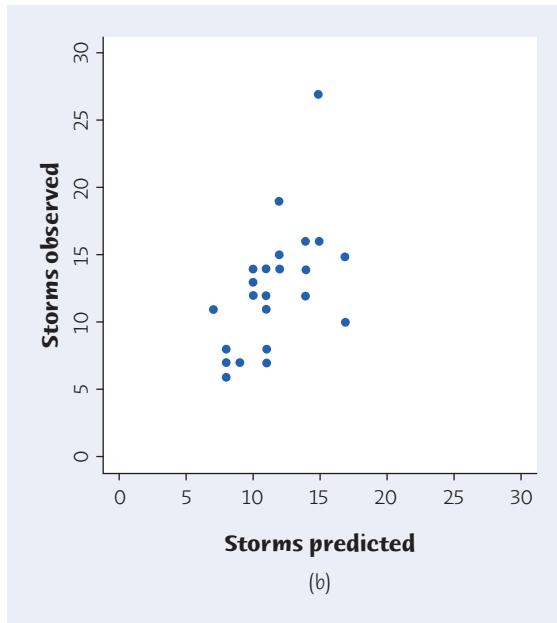
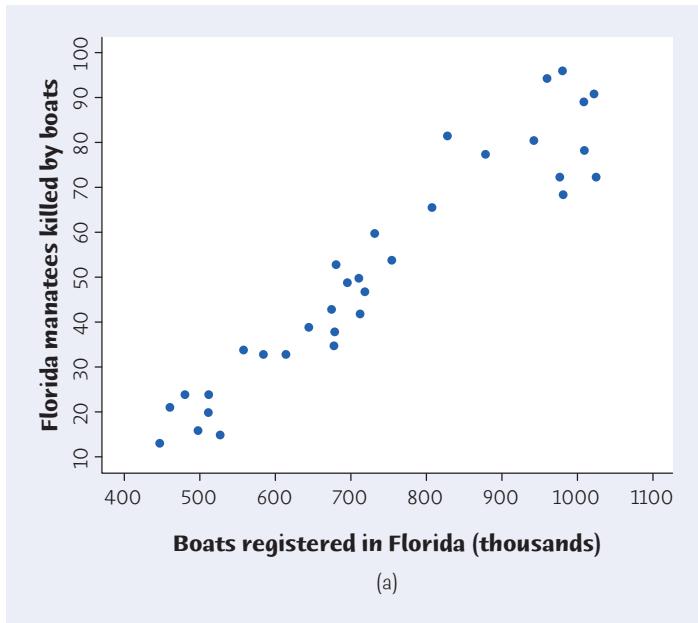


FIGURE 4.5

How correlation measures the strength of a linear relationship, for Example 4.6. Patterns closer to a straight line have correlations closer to 1 or -1 .

**FIGURE 4.6**

How correlation measures the strength of a linear relationship, for Example 4.6. Four sets of real data with (a) $r = 0.951$, (b) $r = 0.584$, (c) $r = 0.358$, and (d) $r = -0.081$.

- (a) This repeats the manatee plot in Figure 4.2. There is a strong positive linear relationship, $r = 0.951$.
- (b) Here are the number of named tropical storms each year between 1984 and 2008 plotted against the number predicted before the start of hurricane season by William Gray of Colorado State University. There is a moderate linear relationship, $r = 0.584$.
- (c) These data come from an experiment that studied how quickly cuts in the limbs of newts heal. Each point represents the healing rate in micrometers (millionths of a meter) per hour for the two front limbs of the same newt. This relationship is weaker than those in (a) and (b), with $r = 0.358$.
- (d) Does last year's stock market performance help predict how stocks will do this year? No. The correlation between last year's percent return and this year's percent return over 56 years is only $r = -0.081$. The scatterplot shows a cloud of points with no visible linear pattern. ■



Describing the relationship between two variables is a more complex task than describing the distribution of one variable. Here are some more facts about correlation, cautions to keep in mind when you use r :

- 1.** *Correlation requires that both variables be quantitative, so that it makes sense to do the arithmetic indicated by the formula for r .* We cannot calculate a correlation between the incomes of a group of people and what city they live in, because city is a categorical variable. 
- 2.** *Correlation measures the strength of only the linear relationship between two variables. Correlation does not describe curved relationships between variables, no matter how strong they are.* Exercise 4.13 (page 112) illustrates this important fact. 
- 3.** *Like the mean and standard deviation, the correlation is not resistant: r is strongly affected by a few outlying observations.* Use r with caution when outliers appear in the scatterplot. Figure 4.6(b) contains an outlier, the disastrous 2005 season, whose 27 named storms included Hurricane Katrina. Adding this one point to the other 22 increases the correlation from 0.564 to 0.584. Because the outlier extends the linear pattern, it increases the correlation. 
- 4.** *Correlation is not a complete summary of two-variable data, even when the relationship between the variables is linear.* You should give the means and standard deviations of both x and y along with the correlation. 

Because the formula for correlation uses the means and standard deviations, these measures are the proper choice to accompany a correlation. Here is an example in which understanding requires both means and correlation.

EXAMPLE 4.7 Scoring American Idol at home

One Web site recommends that fans of the television show *American Idol* score contestants at home on a scale from 1 to 10, with higher scores indicating a better performance. Two friends, Angela and Elizabeth, decide to follow this advice and score



Michael Becker/Fox/PictureGroup via AP
IMAGES

contestants over the course of a season. How well do they agree? We calculate that the correlation between their scores is $r = 0.9$, suggesting that they agree. But the mean of Angela's scores is 0.8 point lower than Elizabeth's mean. Does this suggest that the two friends disagree?

These facts do not contradict each other. They are simply different kinds of information. The mean scores show that Angela awards lower scores than Elizabeth. But because Angela gives *every* contestant a score about 0.8 point lower than Elizabeth, the correlation remains high. Adding the same number to all values of either x or y does not change the correlation. Angela and Elizabeth actually score consistently because they agree on which performances are better. The high r shows their agreement. ■

Of course, even giving means, standard deviations, and the correlation for state SAT scores and percent taking will not point out the clusters in Figure 4.1 (page 100). Numerical summaries complement plots of data, but they don't replace them.

APPLY YOUR KNOWLEDGE

4.11 Changing the units. The healing rates plotted in Figure 4.6(c) (page 110) are measured in micrometers (millionths of a meter) per hour. The correlation between healing rates for the two front limbs of newts is $r = 0.358$. If the measurements were made in inches per day, would the correlation change? Explain your answer.



4.12 Changing the correlation. Use your calculator, software, or the *Two-Variable Statistical Calculator* applet to demonstrate how outliers can affect correlation.

- What is the correlation between lean body mass and metabolic rate for the 12 women in Exercise 4.4?  METABOLIC
- Make a scatterplot of the data with two new points added. Point A: mass 65 kilograms, metabolic rate 1761 calories. Point B: mass 35 kilograms, metabolic rate 1400 calories. Find two new correlations: one for the original data plus Point A, and another for the original data plus Point B.
- By looking at your plot, explain why adding Point A makes the correlation stronger (closer to 1) and adding Point B makes the correlation weaker (closer to 0).  METABOLICS

4.13 Strong association but no correlation. The gas mileage of an automobile first increases and then decreases as the speed increases. Suppose that this relationship is very regular, as shown by the following data on speed (miles per hour) and mileage (miles per gallon):

Speed	30	40	50	60	70
Mileage	24	28	30	28	24

Make a scatterplot of mileage versus speed. Show that the correlation between speed and mileage is $r = 0$. Explain why the correlation is 0 even though there is a strong relationship between speed and mileage. 

CHAPTER 4 SUMMARY

CHAPTER SPECIFICS

- To study relationships between variables, we must measure the variables on the same group of individuals.
- If we think that a variable x may explain or even cause changes in another variable y , we call x an **explanatory variable** and y a **response variable**.
- A **scatterplot** displays the relationship between two quantitative variables measured on the same individuals. Mark values of one variable on the horizontal axis (x axis) and values of the other variable on the vertical axis (y axis). Plot each individual's data as a point on the graph. Always plot the explanatory variable, if there is one, on the x axis of a scatterplot.
- Plot points with different colors or symbols to see the effect of a categorical variable in a scatterplot.
- In examining a scatterplot, look for an overall pattern showing the **direction**, **form**, and **strength** of the relationship and then for **outliers** or other deviations from this pattern.
- **Direction:** If the relationship has a clear direction, we speak of either **positive association** (high values of the two variables tend to occur together) or **negative association** (high values of one variable tend to occur with low values of the other variable).
- **Form:** **Linear relationships**, where the points show a straight-line pattern, are an important form of relationship between two variables. Curved relationships and **clusters** are other forms to watch for.
- **Strength:** The **strength** of a relationship is determined by how close the points in the scatterplot lie to a simple form such as a line.
- The **correlation r** measures the direction and strength of the linear association between two quantitative variables x and y . Although you can calculate a correlation for any scatterplot, r measures only straight-line relationships.
- Correlation indicates the direction of a linear relationship by its sign: $r > 0$ for a positive association and $r < 0$ for a negative association. Correlation always satisfies $-1 \leq r \leq 1$ and indicates the strength of a relationship by how close it is to -1 or 1 . Perfect correlation, $r = \pm 1$, occurs only when the points on a scatterplot lie exactly on a straight line.
- Correlation ignores the distinction between explanatory and response variables. The value of r is not affected by changes in the unit of measurement of either variable. Correlation is not resistant, so outliers can greatly change the value of r .

LINK IT

In this chapter we continue our study of exploratory data analysis but for the purpose of examining relationships *between* variables. Scatterplots are a type of graph that can be used to visualize patterns in the relationship between two variables. We look for an overall pattern showing the direction, form, and strength of the relationship and then for outliers or other deviations from this pattern. The direction of the pattern is often

summarized as either a positive association (high values of the two variables tend to occur together) or a negative association (high values of one variable tend to occur with low values of the other variable). Forms to watch for are straight-line patterns, curved patterns, or clusters. The strength of a relationship is determined by how close the points in the scatterplot lie to a simple form such as a straight line or curve.

One of the simplest forms is a linear relationship, where the points suggest a straight-line pattern. Correlation is a number that summarizes the strength and direction of a linear relation. Positive values of the correlation correspond to a positive association. Negative values correspond to a negative association. The closer the absolute value of the correlation is to 1, the stronger the linear relationship (the more closely the points in the scatterplot come to lying on a straight line).

It is tempting to assume that the patterns we observe in our data hold for values of our variables that we have not observed—in other words, that additional data would continue to conform to these patterns. The process of identifying underlying patterns would seem to assume that this is the case. But is this assumption justified? We will return to this issue in Part IV of the book.

CHECK YOUR SKILLS

4.14 You have data for many years on the average price of a barrel of oil and the average retail price of a gallon of unleaded regular gasoline. When you make a scatterplot, the explanatory variable on the x axis

- (a) is the price of oil. (b) is the price of gasoline.
- (c) can be either oil price or gasoline price.

4.15 In a scatterplot of the average price of a barrel of oil and the average retail price of a gallon of gasoline, you expect to see

- (a) a positive association. (b) very little association.
- (c) a negative association.

4.16 Figure 4.7 is a scatterplot of school GPA against IQ test scores for 15 seventh-grade students. There is one low outlier in the plot. The IQ and GPA scores for this student are

- (a) IQ = 0.5, GPA = 103. (b) IQ = 103, GPA = 0.5.
- (c) IQ = 103, GPA = 7.6.

4.17 If we leave out the low outlier, the correlation for the remaining 14 points in Figure 4.7 is closest to

- (a) 0.9. (b) -0.9. (c) 0.1.

4.18 What are all the values that a correlation r can possibly take?

- (a) $r \geq 0$ (b) $0 \leq r \leq 1$ (c) $-1 \leq r \leq 1$

4.19 If the correlation between two variables is close to 0, you can conclude that a scatterplot would show

- (a) a strong straight-line pattern.
- (b) a cloud of points with no visible pattern.

(c) no straight-line pattern, but there might be a strong pattern of another form.

4.20 The points on a scatterplot lie very close to a straight line. The correlation between x and y is close to

- (a) -1. (b) 1. (c) either -1 or 1, we can't say which.

4.21 If men always married women who were two years younger than themselves, the correlation between the ages of husband and wife would be

- (a) 1. (b) -1. (c) Can't tell without seeing the data.

4.22 For a biology project, you measure the weight in grams and the tail length in millimeters of a group of mice. The correlation is $r = 0.7$. If you had measured tail length in centimeters instead of millimeters, what would be the correlation? (There are 10 millimeters in a centimeter.)

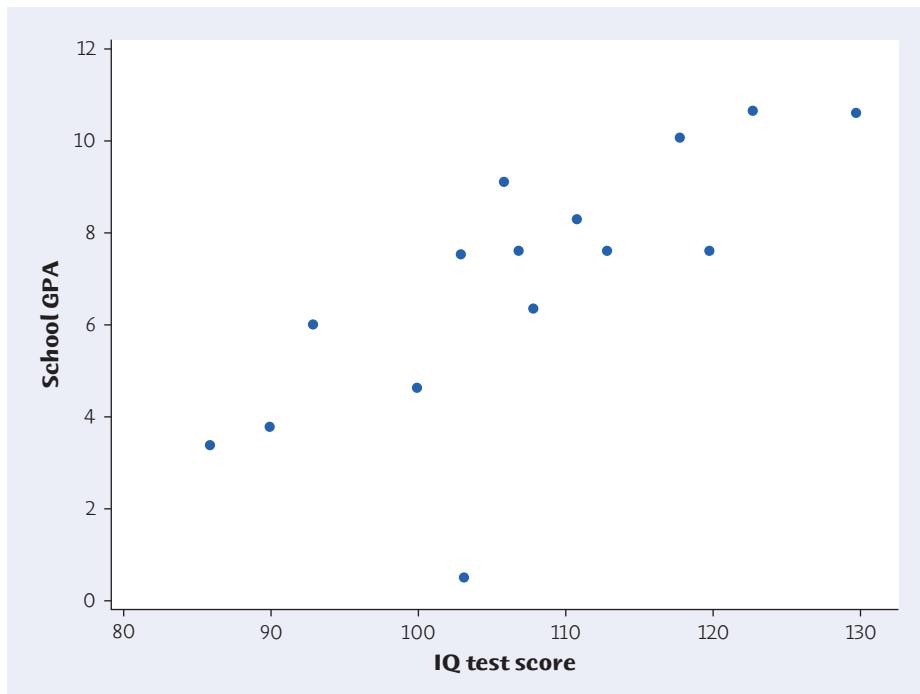
- (a) $0.7/10 = 0.07$ (b) 0.7 (c) $(0.7)(10) = 7$

4.23 Because elderly people may have difficulty standing to have their heights measured, a study looked at predicting overall height from height to the knee. Here are data (in centimeters) for six elderly men:  KNEEHT

Knee height x	57.7	47.4	43.5	44.8	55.2	54.6
Height y	192.1	153.3	146.4	162.7	169.1	177.8

Use your calculator or software: the correlation between knee height and overall height is about

- (a) $r = 0.08$. (b) $r = 0.89$. (c) $r = 0.74$.

**FIGURE 4.7**

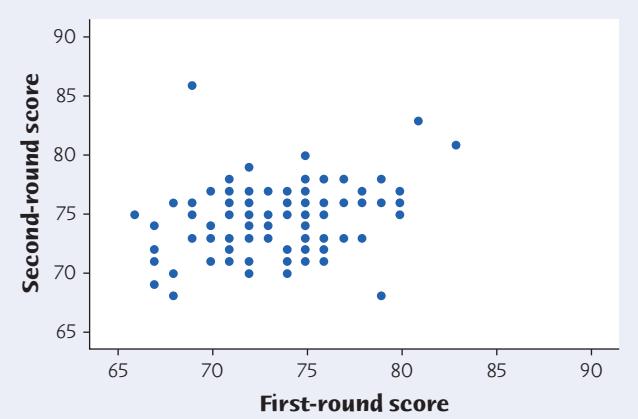
Scatterplot of school GPA against IQ test scores for seventh-grade students, for Exercises 4.16 and 4.17.

CHAPTER 4 EXERCISES

4.24 Scores at the Masters. The Masters is one of the four major golf tournaments. Figure 4.8 is a scatterplot of the scores for the first two rounds of the 2010 Masters for all the golfers entered. Only the 60 golfers with the lowest two-round total advance to the final two rounds. The plot has a grid pattern because golf scores must be whole numbers.⁹ 

- Read the graph: What was the lowest score in the first round of play? How many golfers had this low score? What were their scores in the second round?
- Read the graph: Sandy Lyle had the highest score in the second round. What was this score? What was Lyle's score in the first round?
- Is the correlation between first-round scores and second-round scores closest to $r = 0.1$, $r = 0.5$, or $r = 0.9$? Explain your choice. Does the graph suggest that knowing a professional golfer's score for one round is much help in predicting his score for another round on the same course?

4.25 Happy states. Human happiness or well-being can be assessed either subjectively or objectively. Subjective assessment can be accomplished by listening to what people say. Objective assessment can be made from data related to well-being such as income, climate, availability of entertainment, housing prices, lack of traffic congestion, etc. Do subjective and objec-

**FIGURE 4.8**

Scatterplot of the scores in the first two rounds of the 2010 Masters Tournament, for Exercise 4.24.

tive assessments agree? To study this, investigators made both subjective and objective assessments of happiness for each of the 50 states. The subjective measurement was the mean score on a life-satisfaction question found on the Behavioral Risk Factor Surveillance System (BRFSS), which is a state-based

system of health surveys. Lower scores indicate a greater degree of happiness. To objectively assess happiness, the investigators computed a mean well-being score (called the compensating-differentials score) for each state, based on objective measures that have been found to be related to happiness or well-being. The states were then ranked according to this score (Rank 1 being the happiest). Figure 4.9 is a scatterplot of mean BRFSS scores (response) against the rank based on the compensating-differentials scores (explanatory).¹⁰

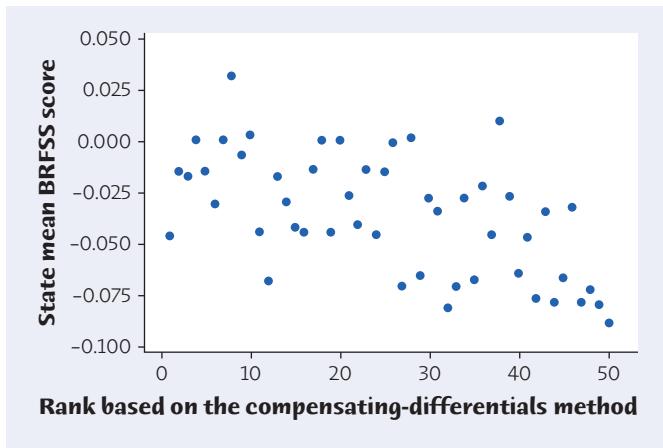


FIGURE 4.9

Scatterplot of mean BRFSS score in each state against each state's well-being rank, for Exercise 4.25.

- (a) Is there an overall positive association or an overall negative association between mean BRFSS score and rank based on the compensating-differentials method?
- (b) Does the overall association indicate agreement or disagreement between the mean subjective BRFSS score and the ranking based on objective data used in the compensating-differentials method?
- (c) Are there any outliers? If so, what are the BRFSS scores corresponding to these outliers?

4.26 Wine and cancer in women. Some studies have suggested that a nightly glass of wine may not only take the edge off a day but also improve health. Is wine good for your health? A study of nearly 1.3 million middle-aged British women examined wine consumption and the risk of breast cancer. The researchers were interested in how risk changed as wine consumption increased. Risk is based on breast cancer rates in drinkers relative to breast cancer rates in nondrinkers in the study, with higher values indicating greater risk. In particular, a value greater than 1 indicates a greater breast cancer rate than that of nondrinkers. Wine intake is the mean wine intake, in grams per day, of all women in the study who drank

approximately the same amount of wine per week. Here are the data (for drinkers only):¹¹



Wine intake x (grams per day)	2.5	8.5	15.5	26.5
Relative risk y	1.00	1.08	1.15	1.22

- (a) Make a scatterplot of these data. Based on the scatterplot, do you expect the correlation to be positive or negative? Near ± 1 or not?

- (b) Find the correlation r between wine intake and relative risk. Do the data show that women who consume more wine tend to have higher relative risks of breast cancer?

4.27 Ebola and gorillas. The deadly Ebola virus is a threat to both people and gorillas in Central Africa. An outbreak in 2002 and 2003 killed 91 of the 95 gorillas in 7 home ranges in the Congo. To study the spread of the virus, measure “distance” by the number of home ranges separating a group of gorillas from the first group infected. Here are data on distance and time in number of days until deaths began in each later group:¹²



Distance	1	3	4	4	4	5
Time	4	21	33	41	43	46

- (a) Make a scatterplot. Which is the explanatory variable? What kind of pattern does your plot show?

- (b) Find the correlation r between distance and time.

- (c) If time in days were replaced by time in number of weeks until death began in each later group (fractions allowed so that 4 days becomes $4/7$ weeks), would the correlation between distance and time change? Explain your answer.

4.28 Sparrowhawk colonies.

One of nature’s patterns connects the percent of adult birds in a colony that return from the previous year and the number of new adults that join the colony. Here are data for 13 colonies of sparrowhawks:¹³



Morales Morales/Photolibrary

Percent return	74	66	81	52	73	62	52	45	62	46	60	46	38
New adults	5	6	8	11	12	15	16	17	18	18	19	20	20

(a) Plot the count of new adults (response) against the percent of returning birds (explanatory). Describe the direction and form of the relationship. Is the correlation r an appropriate measure of the strength of this relationship? If so, find r .

(b) For short-lived birds, the association between these variables is positive: changes in weather and food supply drive the populations of new and returning birds up or down together. For long-lived territorial birds, on the other hand, the association is negative because returning birds claim their territories in the colony and don't leave room for new recruits. Which type of species is the sparrowhawk?

4.29 Our brains don't like losses. Most people dislike losses more than they like gains. In money terms, people are about as sensitive to a loss of \$10 as to a gain of \$20. To discover what parts of the brain are active in decisions about gain and loss, psychologists presented subjects with a series of gambles with different odds and different amounts of winnings and losses. From a subject's choices, they constructed a measure of "behavioral loss aversion." Higher scores show greater sensitivity to losses. Observing brain activity while subjects made their decisions pointed to specific brain regions. Here are data for 16 subjects on behavioral loss aversion and "neural loss aversion," a measure of activity in one region of the brain:¹⁴ 

Neural	-50.0	-39.1	-25.9	-26.7	-28.6	-19.8	-17.6	5.5	of alcohol other than wine per week. Here are the data for both women who drank wine and women who drank alcoholic beverages other than wine:
Behavioral	0.08	0.81	0.01	0.12	0.68	0.11	0.36	0.34	 CANCER
Neural	2.6	20.7	12.1	15.5	28.8	41.7	55.3	155.2	
Behavioral	0.53	0.68	0.99	1.04	0.66	0.86	1.29	1.94	

(a) Make a scatterplot that shows how behavior responds to brain activity.

(b) Describe the overall pattern of the data. There is one clear outlier. What is the behavioral score associated with this outlier?

(c) Find the correlation r between neural and behavioral loss aversion both with and without the outlier. Does the outlier have a strong influence on the value of r ? By looking at your plot, explain why adding the outlier to the other data points causes r to increase.

4.30 Sulfur, the ocean, and the sun.

Sulfur in the atmosphere affects climate by influencing formation of clouds. The main natural source of sulfur is dimethyl sulfide (DMS) produced by small organisms in the upper layers of the oceans. DMS production is in turn influenced by the amount of energy the upper ocean receives from sunlight. Here are monthly data on solar radiation dose (SRD, in watts per square meter) and surface DMS concentration (in nanomolars) for a region in the Mediterranean:¹⁵ 

SRD	12.55	12.91	14.34	19.72	21.52	22.41	37.65	48.41
DMS	0.796	0.692	1.744	1.062	0.682	1.517	0.736	0.720
SRD	74.41	94.14	109.38	157.79	262.67	268.96	289.23	
DMS	1.820	1.099	2.692	5.134	8.038	7.280	8.872	

(a) Make a scatterplot that shows how DMS responds to SRD.

(b) Describe the overall pattern of the data. Find the correlation r between DMS and SRD. Because SRD changes with the seasons of the year, the close relationship between SRD and DMS helps explain other seasonal patterns.

4.31 Alcohol and cancer in women. Exercise 4.26 discusses a study of the relationship between wine consumption and the risk of breast cancer in women. The researchers were also interested in how risk changed as consumption of alcoholic beverages other than wine increased. Intake of alcoholic beverages other than wine is the mean intake, in grams per day, of all women in the study who drank approximately the same amount

of alcohol other than wine per week. Here are the data for both women who drank wine and women who drank alcoholic beverages other than wine:  CANCER

Wine intake (grams per day)	2.5	8.5	15.5	26.5
Relative risk	1.00	1.08	1.15	1.22
Alcohol intake (grams per day)	2.0	7.0	13.0	24.0
Relative risk	0.96	1.06	1.11	1.20

(a) Make a scatterplot of the relative risk versus intake, using separate symbols for the two types of drinks.

(b) What does your plot show about the pattern of risk? What does it show about the effect of type of drink on risk?

4.32 Feed the birds. Canaries provide more food to their babies when the babies beg more intensely. Researchers wondered if begging was the main factor determining how much food baby canaries receive, or if parents also take into account whether the babies are theirs or not. To investigate, researchers conducted an experiment allowing canary parents to raise two broods: one of their own and one fostered from a different pair of parents. If begging determines how much food babies receive, then differences in the “begging intensities” of the broods should be strongly associated with differences in the amount of food the broods receive. The researchers decided to use the relative growth rates (the growth rate of the foster babies relative to that of the natural babies, with values greater than 1 indicating that the foster babies grew more rapidly than the natural babies) as a measure of the difference in the amount of food received. They recorded the difference in begging intensities (the begging intensity of the foster babies minus that of the natural babies) and relative growth rates. Here are data from the experiment:¹⁶

 CANARIES

Difference in begging intensity	-14.0	-12.5	-12.0	-8.0	-8.0	-6.5	-5.5
Relative growth rate	0.85	1.00	1.33	0.85	0.90	1.15	1.00
Difference in begging intensity	-3.5	-3.0	-2.0	-1.5	-1.5	0.0	0.0
Relative growth rate	1.30	1.33	1.03	0.95	1.15	1.13	1.00
Difference in begging intensity	2.00	2.00	3.00	4.50	7.00	8.00	8.50
Relative growth rate	1.07	1.14	1.00	0.83	1.15	0.93	0.70

- (a) Make a scatterplot that shows how relative growth rate responds to the difference in begging intensity.

(b) Describe the overall pattern of the relationship. Is it linear? Is there a positive or negative association, or neither? Find the correlation r . Is r a helpful description of this relationship?

(c) If begging intensity is the main factor determining food received, with higher intensity leading to more food, one would expect the relative growth rate to increase as the difference in begging intensity increases. However, if both begging intensity and a preference for their own babies determine the amount of food received (and hence the relative growth rate),

Weather report	
Good	20.
	24.
Bad	18.
	17.
None	19.
	18.

we might expect growth rate to increase initially as begging intensity increases but then to level off (or even decrease) as the parents begin to ignore increases in begging by the foster babies. Which of these theories do the data appear to support? Explain your answer.

4.33 Good weather and tipping. Favorable weather has been shown to be associated with increased tipping. Will just the belief that future weather will be favorable lead to higher tips? Researchers gave 60 index cards to a waitress at an Italian restaurant in New Jersey. Before delivering the bill to each customer, the waitress randomly selected a card and wrote on the bill the same message that was printed on the index card. Twenty of the cards had the message “The weather is supposed to be really good



tomorrow. I hope you enjoy the day!" Another 20 cards contained the message "The weather is supposed to be not so good tomorrow. I hope you enjoy the day anyway!" The remaining 20 cards were

blank, indicating that the waitress was not supposed to write any message. Choosing a card at random ensured that there was a random assignment of the diners to the three experimental conditions. Here are the percentage tips for the three messages:¹⁷

 TIPPING

Weather report	Percentage tip									
Good	20.8	18.7	19.9	20.6	22.0	23.4	22.8	24.9	22.2	20.3
	24.9	22.3	27.0	20.4	22.2	24.0	21.2	22.1	22.0	22.7
Bad	18.0	19.0	19.2	18.8	18.4	19.0	18.5	16.1	16.8	14.0
	17.0	13.6	17.5	19.9	20.2	18.8	18.0	23.2	18.2	19.4
None	19.9	16.0	15.0	20.1	19.3	19.2	18.0	19.2	21.1	18.8
	18.5	19.3	19.3	19.4	10.8	19.1	19.7	19.8	21.3	20.6

(a) Make a plot of percentage tip against the weather report on the bill (space the three weather reports equally on the horizontal axis). Which weather report appears to lead to the best tip?

(b) Does it make sense to speak of a positive or negative association between weather report and percentage tip? Why? Is correlation r a helpful description of the relationship? Why?

4.34 Thinking about correlation. Exercise 4.26 presents data on wine intake and the relative risk of breast cancer in women.

(a) If wine intake is measured in ounces per day rather than grams per day, how would the correlation change? (There are 0.035 ounces in a gram.)

(b) How would r change if all the relative risks were 0.25 less than the values given in the table? Does the correlation tell us that among women who drink, those who drink more wine tend to have a greater relative risk of cancer than women who don't drink at all?

(c) If drinking an additional gram of wine each day raised the relative risk of breast cancer by exactly 0.01, what would be the correlation between wine intake and relative risk of breast cancer? (Hint: Draw a scatterplot for several values of wine intake.)

4.35 The effect of changing units. Changing the units of measurement can dramatically alter the appearance of a scatterplot. Return to the data on knee height and overall height in Exercise 4.23:  KNEEHT2

Knee height x	57.7	47.4	43.5	44.8	55.2	54.6
Height y	192.1	153.3	146.4	162.7	169.1	177.8

Both heights are measured in centimeters. A mad scientist decides to measure knee height in millimeters and height in meters. The same data in these units are

Knee height x	577	474	435	448	552	546
Height y	1.921	1.533	1.464	1.627	1.691	1.778

(a) Make a plot with the x axis extending from 0 to 600 and the y axis from 0 to 250. Plot the original data on these axes. Then plot the new data using a different color or symbol. The two plots look very different.

(b) Nonetheless, the correlation is exactly the same for the two sets of measurements. Why do you know that this is true without doing any calculations? Find the two correlations to verify that they are the same.

4.36 Statistics for investing. Investment reports now often include correlations. Following a table of correlations among mutual funds, a report adds: "Two funds can have perfect correlation, yet different levels of risk. For example, Fund A and Fund B may be perfectly correlated, yet Fund A moves 20% whenever Fund B moves 10%." Write a brief explanation, for someone who knows no statistics, of how this can happen. Include a sketch to illustrate your explanation.

4.37 Statistics for investing. A mutual funds company's newsletter says, "A well-diversified portfolio includes assets with low correlations." The newsletter includes a table of correlations between the returns on various classes of investments. For example, the correlation between municipal bonds and large-cap stocks is 0.50, and the correlation between municipal bonds and small-cap stocks is 0.21.

(a) Rachel invests heavily in municipal bonds. She wants to diversify by adding an investment whose returns do not closely follow the returns on her bonds. Should she choose large-cap stocks or small-cap stocks for this purpose? Explain your answer.

(b) If Rachel wants an investment that tends to increase when the return on her bonds drops, what kind of correlation should she look for?

4.38 Teaching and research. A college newspaper interviews a psychologist about student ratings of the teaching of faculty members. The psychologist says, "The evidence indicates that the correlation between the research productivity and teaching rating of faculty members is close to zero." The paper reports this as "Professor McDaniel said that good researchers tend to be poor teachers, and vice versa." Explain why the paper's report is wrong. Write a statement in plain language (don't use the word "correlation") to explain the psychologist's meaning.

4.39 Sloppy writing about correlation. Each of the following statements contains a blunder. Explain in each case what is wrong.

(a) "There is a high correlation between the sex of American workers and their income."

(b) "We found a high correlation ($r = 1.09$) between students' ratings of faculty teaching and ratings made by other faculty members."

(c) "The correlation between height and weight of the subjects was $r = 0.63$ centimeter."

4.40 Correlation is not resistant. Go to the Correlation and Regression applet. Click on the scatterplot to create a group of 10 points in the lower-left corner of the scatterplot with a strong straight-line pattern (correlation about 0.9).

(a) Add one point at the upper right that is in line with the first 10. How does the correlation change?

(b) Drag this last point down until it is opposite the group of 10 points. How small can you make the correlation? Can you make the correlation negative? You see that a single outlier can greatly strengthen or weaken a correlation. Always plot your data to check for outlying points.

4.41 Match the correlation. You are going to use the *Correlation and Regression* applet to make scatterplots with 10 points that have correlation close to 0.7. The lesson is that many patterns can have the same correlation. Always plot your data before you trust a correlation.

(a) Click on the scatterplot to add the first 2 points. What is the value of the correlation? Why does it have this value? (b) Make a lower-left to upper-right pattern of 10 points with correlation about $r = 0.7$. (You can drag points up or down to adjust r after you have 10 points.) Make a rough sketch of your scatterplot.

(c) Make another scatterplot with 9 points in a vertical stack at the left of the plot. Add 1 point far to the right and move it until the correlation is close to 0.7. Make a rough sketch of your scatterplot.

(d) Make yet another scatterplot with 10 points in a curved pattern that starts at the lower left, rises to the right, then falls again at the far right. Adjust the points up or down until you have a quite smooth curve with correlation close to 0.7. Make a rough sketch of this scatterplot also.

The following exercises ask you to answer questions from data without having the details outlined for you. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan, Solve, and Conclude** steps of the process, described on page xxx.

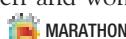
4.42 Brighter sunlight? The brightness of sunlight at the earth's surface changes over time depending on whether the earth's atmosphere is more or less clear. Sunlight dimmed between 1960 and 1990. After 1990, air pollution dropped in industrial countries. Did sunlight brighten? Here are annual averages computed by researchers from data from Ny Alesund, Spitsbergen, Norway, averaging over only clear days each year. (Other locations show simi-

lar trends.) The response variable is solar radiation in watts per square meter.¹⁸



Year	1993	1994	1995	1996	1997	1998	1999	2000
Sun	116.0	120.0	123.0	123.5	125.5	125.0	129.0	128.0

4.43 Will women outrun men? Does the physiology of women make them better suited than men to long-distance running? Will women eventually outperform men in long-distance races? Researchers examined data on world record times (in seconds) for men and women in the marathon. Here are data for women:¹⁹



Year	1926	1964	1967	1970	1971	1974	1975
Time	13,222.0	11,973.0	11,246.0	10,973.0	9990.0	9834.5	9499.0
Year	1977	1980	1981	1982	1983	1985	
Time	9287.5	9027.0	8806.0	8771.0	8563.0	8466.0	

Here are data for men:

Year	1908	1909	1913	1920	1925	1935	1947
Time	10,518.4	9751.0	9366.6	9155.8	8941.8	8802.0	8739.0
Year	1952	1953	1954	1958	1960	1963	1964
Time	8442.2	8314.8	8259.4	8117.0	8116.2	8068.0	7931.2
Year	1965	1967	1969	1981	1984	1985	1988
Time	7920.0	7776.4	7713.6	7698.0	7685.0	7632.0	7610.0

(a) What do the data show about women's and men's times in the marathon? (Start by plotting both sets of data on the same plot, using two different plotting symbols.)

(b) Based on these data, researchers (in 1992) predicted that women would outrun men in the marathon in 1998. How do you think they arrived at this date? Was their prediction accurate? (You may want to look on the Web; try doing a Google search on "women's world record marathon times.")

4.44 Toucan's beak.

The toco toucan, the largest member of the toucan family, possesses the largest beak relative to body size of



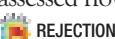
all birds. This exaggerated feature has received various interpretations, such as being a refined adaptation for feeding. However, the large surface area may also be an important mechanism for radiating heat (and hence cooling the bird) as outdoor temperature increases. Here are data for beak heat loss, as a percent of total body heat loss, at various temperatures in degrees Celsius:²⁰



Temperature (°C)	15	16	17	18	19	20	21	22
Percent heat loss from beak	32	34	35	33	37	46	55	51
Temperature (°C)	23	24	25	26	27	28	29	30
Percent heat loss from beak	43	52	45	53	58	60	62	62

Investigate the relationship between outdoor temperature and beak heat loss, as a percent of total body heat loss.

4.45 Does social rejection hurt? We often describe our emotional reaction to social rejection as “pain.” Does social rejection cause activity in areas of the brain that are known to be activated by physical pain? If it does, we really do experience social and physical pain in similar ways. Psychologists first included and then deliberately excluded individuals from a social activity while they measured changes in brain activity. After each activity, the subjects filled out questionnaires that assessed how excluded they felt. Here are data for 13 subjects:²¹



Subject	Social distress	Brain activity
1	1.26	-0.055
2	1.85	-0.040
3	1.10	-0.026
4	2.50	-0.017
5	2.17	-0.017
6	2.67	0.017
7	2.01	0.021
8	2.18	0.025
9	2.58	0.027
10	2.75	0.033
11	2.75	0.064
12	3.33	0.077
13	3.65	0.124

The explanatory variable is “social distress” measured by each subject’s questionnaire score after exclusion relative to the score after inclusion. (So values greater than 1 show the degree of distress caused by exclusion.) The response variable is change in activity in a region of the brain that is activated by physical pain. Negative values show a decrease in activity, suggesting less distress. Discuss what the data show.

4.46 Bushmeat. African peoples often eat “bushmeat,” the meat of wild animals. Bushmeat is widely traded in Africa, but its consumption threatens the survival of some animals in the wild. Bushmeat is often not the first choice of consumers—they eat bushmeat when other sources of protein are in short supply. Researchers looked at declines in 41 species of mammals in nature reserves in Ghana and at catches of fish (the primary source of animal protein) in the same region. The data appear in Table 4.2.²² Fish supply is measured in kilograms per person. The other variable is the percent change in the total “biomass” (weight in tons) for the 41 animal species in six nature reserves. Most of the yearly percent changes in wildlife mass are negative because most years saw fewer wild animals in West Africa. Discuss how the data support the idea that more animals are killed for bushmeat when the fish supply is low.



TABLE 4.2 Fish supply and wildlife decline in West Africa

YEAR	FISH SUPPLY (KG PER PERSON)	BIOMASS CHANGE (PERCENT)	YEAR	FISH SUPPLY (KG PER PERSON)	BIOMASS CHANGE (PERCENT)
1971	34.7	2.9	1985	21.3	-5.5
1972	39.3	3.1	1986	24.3	-0.7
1973	32.4	-1.2	1987	27.4	-5.1
1974	31.8	-1.1	1988	24.5	-7.1
1975	32.8	-3.3	1989	25.2	-4.2
1976	38.4	3.7	1990	25.9	0.9
1977	33.2	1.9	1991	23.0	-6.1
1978	29.7	-0.3	1992	27.1	-4.1
1979	25.0	-5.9	1993	23.4	-4.8
1980	21.8	-7.9	1994	18.9	-11.3
1981	20.8	-5.5	1995	19.6	-9.3
1982	19.7	-7.2	1996	25.3	-10.7
1983	20.8	-4.1	1997	22.0	-1.8
1984	21.1	-8.6	1998	21.0	-7.4

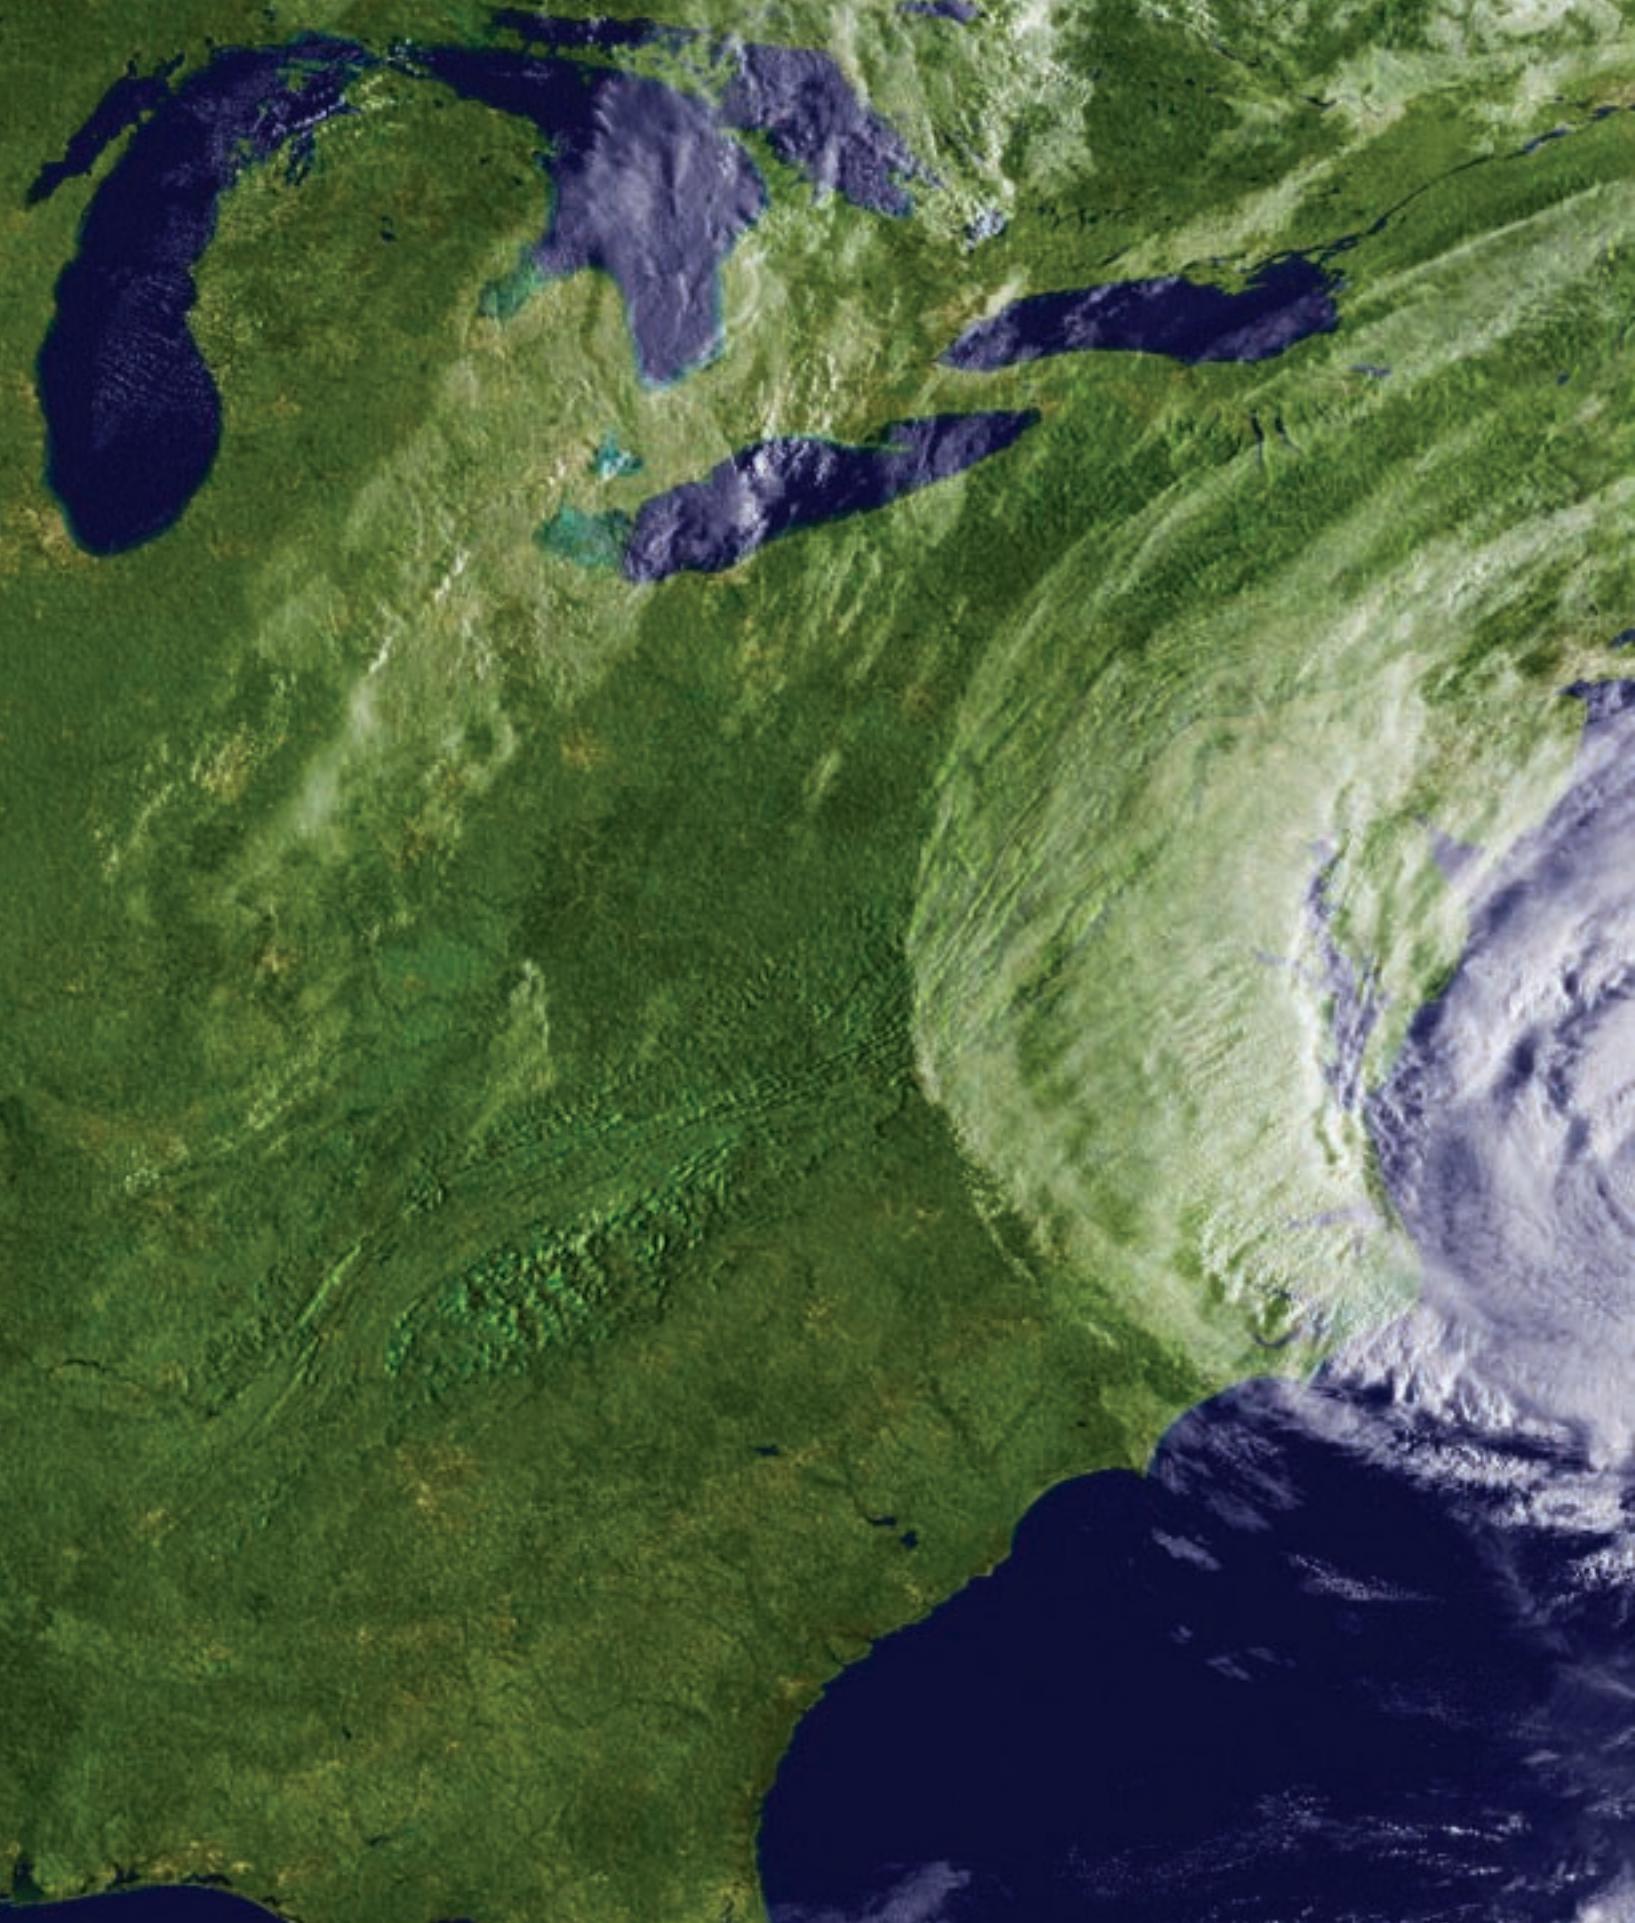


EXPLORING THE WEB

4.47 Drive for show, putt for dough. A popular saying in golf is “You drive for show but you putt for dough.” The point is that hitting the golf ball a long way with a driver looks impressive, but putting well is more important for the final score and hence the amount of money you win. You can find this season’s Professional Golfers Association (PGA) Tour statistics at the PGA Tour Web site: www.pgatour.com/r/stats (click on “View All” under any category displayed to see the statistics for all golfers). You can also find these statistics at the ESPN Web site: espn.go.com/golf/statistics/_/year/. Look at the most recent putting, driving, and money earnings data for the current season on the PGA Tour.

- (a) Make a scatterplot of earnings and putting average. Use earnings as the response variable. Describe the direction, form, and strength of the relationship in the plot. Are there any outliers?
- (b) Make a scatterplot of earnings and driving distance. Use earnings as the response variable. Describe the direction, form, and strength of the relationship in the plot. Are there any outliers?
- (c) Do your plots support the maxim “You drive for show but you putt for dough”?

4.48 Olympic medals. Go to the *Chance News* Web site at www.causeweb.org/wiki/chance/index.php/Chance_News_61#Predicting_medal_counts and read the article “Predicting Medal Counts.” Next, search the Web and locate the Winter Olympics medal counts for 2002 and 2006 (I found Winter Olympics medal counts on Wikipedia). Make a scatterplot that is similar to the one in the *Chance News* article but that uses the 2002 medal counts to predict the 2010 medal counts. How does your plot compare with the plot in the *Chance News* article?



Regression

Chapter 5

Linear (straight-line) relationships between two quantitative variables are easy to understand and quite common. In Chapter 4, we found linear relationships in settings as varied as Florida manatee deaths, the risk of cancer, and predicting tropical storms. Correlation measures the direction and strength of these linear relationships. When a scatterplot shows a linear relationship, we would like to summarize the overall pattern by drawing a line on the scatterplot.

REGRESSION LINES

A *regression line* summarizes the relationship between two variables, but only in a specific setting: one of the variables helps explain or predict the other. That is, regression describes a relationship between an explanatory variable and a response variable.

REGRESSION LINE

A **regression line** is a straight line that describes how a response variable y changes as an explanatory variable x changes. We often use a regression line to predict the value of y for a given value of x .

EXAMPLE 5.1 Does fidgeting keep you slim?

Why is it that some people find it easy to stay slim? Here, following our four-step process (page 55), is an account of a study that sheds some light on gaining weight.

STATE: Some people don't gain weight even when they overeat. Perhaps fidgeting and other "nonexercise activity" (NEA) explains why. Some people may

IN THIS CHAPTER WE COVER...

- Regression lines
- The least-squares regression line
- Using technology
- Facts about least-squares regression
- Residuals
- Influential observations
- Cautions about correlation and regression
- Association does not imply causation





spontaneously increase nonexercise activity when fed more. Researchers deliberately overfed 16 healthy young adults for 8 weeks. They measured fat gain (in kilograms) and, as an explanatory variable, change in energy use (in calories) from activity other than deliberate exercise—fidgeting, daily living, and the like. Change in energy use was energy use measured the last day of the 8-week period minus energy use measured the day before the overfeeding began. Here are the data:¹

NEA change (cal)	-94	-57	-29	135	143	151	245	355
Fat gain (kg)	4.2	3.0	3.7	2.7	3.2	3.6	2.4	1.3
NEA change (cal)	392	473	486	535	571	580	620	690
Fat gain (kg)	3.8	1.7	1.6	2.2	1.0	0.4	2.3	1.1

Do people with larger increases in NEA tend to gain less fat?

PLAN: Make a scatterplot of the data and examine the pattern. If it is linear, use correlation to measure its strength and draw a regression line on the scatterplot to predict fat gain from change in NEA.

SOLVE: Figure 5.1 is a scatterplot of these data. The plot shows a moderately strong negative linear association with no outliers. The correlation is $r = -0.7786$. The line on the plot is a regression line for predicting fat gain from change in NEA.

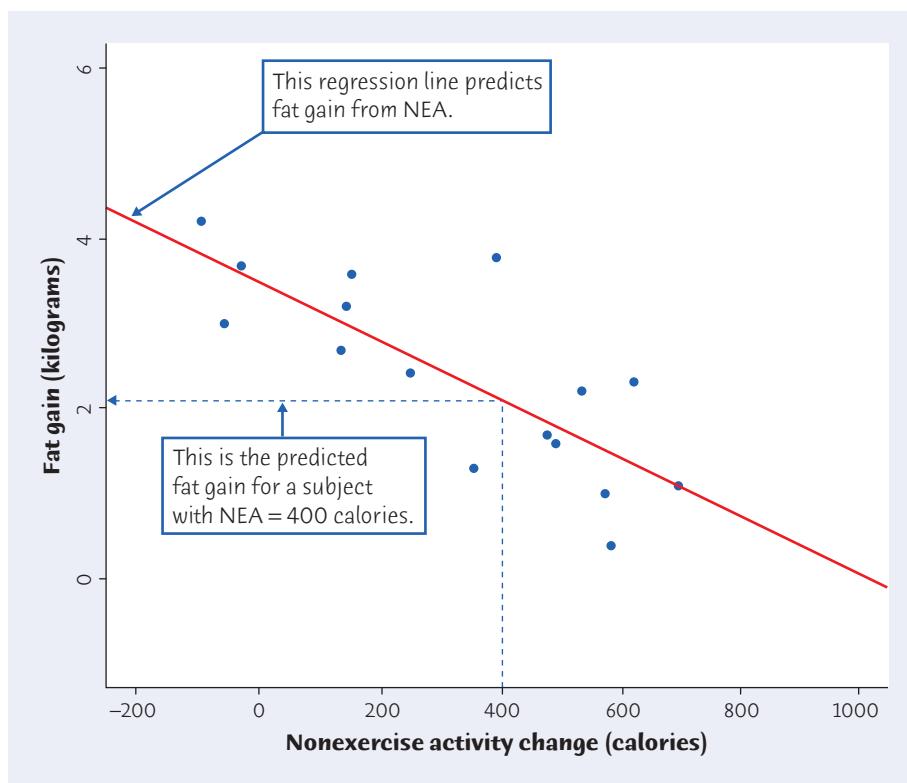


FIGURE 5.1

Fat gain after 8 weeks of overeating, plotted against increase in nonexercise activity over the same period, for Example 5.1.

CONCLUDE: People with larger increases in NEA do indeed gain less fat. To add to this conclusion, we must study regression lines in more detail.

We can, however, already use the regression line to predict fat gain from NEA. Suppose that an individual's NEA increases by 400 calories when she overeats. Go "up and over" on the graph in Figure 5.1. From 400 calories on the x axis, go up to the regression line and then over to the y axis. The graph shows that the predicted gain in fat is a bit more than 2 kilograms. ■

Many calculators and software programs will give you the equation of a regression line from keyed-in data. Understanding and using the line are more important than the details of where the equation comes from.

REVIEW OF STRAIGHT LINES

Suppose that y is a response variable (plotted on the vertical axis) and x is an explanatory variable (plotted on the horizontal axis). A straight line relating y to x has an equation of the form

$$y = a + bx$$

In this equation, b is the **slope**, the amount by which y changes when x increases by one unit. The number a is the **intercept**, the value of y when $x = 0$.



Regression toward the mean

To "regress"

means to go backward. Why are statistical methods for predicting a response from an explanatory variable called "regression"? Sir Francis Galton (1822-1911), who was the first to apply regression to biological and psychological data, looked at examples such as the heights of children versus the heights of their parents. He found that the taller-than-average parents tended to have children who were also taller than average but not as tall as their parents. Galton called this fact "regression toward the mean," and the name came to be applied to the statistical method.

EXAMPLE 5.2 Using a regression line

Any straight line describing the NEA data has the form

$$\text{fat gain} = a + (b \times \text{NEA change})$$

The line in Figure 5.1 is the regression line with the equation

$$\text{fat gain} = 3.505 - 0.00344 \times \text{NEA change}$$

Be sure you understand the role of the two numbers in this equation:

- The slope $b = -0.00344$ tells us that, on average, fat gained goes down by 0.00344 kilogram for each added calorie of NEA change. The slope of a regression line is the *rate of change* in the response, on average, as the explanatory variable changes.
- The intercept, $a = 3.505$ kilograms, is the estimated fat gain if NEA does not change when a person overeats.

The equation of the regression line makes it easy to predict fat gain. If a person's NEA change increases by 400 calories when she overeats, substitute $x = 400$ in the equation. The predicted fat gain is

$$\text{fat gain} = 3.505 - (0.00344 \times 400) = 2.13 \text{ kilograms}$$

This is a bit more than 2 kilograms, as we estimated directly from the plot in Example 5.1.

To **plot the line** on the scatterplot, use the equation to find the predicted y for two values of x , one near each end of the range of x in the data. Plot each y above its x -value and draw the line through the two points. ■

plotting a line

The slope of a regression line is an important numerical description of the relationship between the two variables. Although we need the value of the intercept to draw the line, this value is statistically meaningful only when, as in Example 5.2, the explanatory variable can actually take values close to zero. The slope $b = -0.00344$ in Example 5.2 is small. This does not mean that change in NEA has little effect on fat gain. The size of the slope depends on the units in which we measure the two variables. In this example, the slope is the change in fat gain in kilograms when NEA increases by one calorie. There are 1000 grams in a kilogram. If we measured fat gain in grams, the slope would be 1000 times larger, $b = 3.44$. You can't say how important a relationship is by looking at the size of the slope of the regression line.



APPLY YOUR KNOWLEDGE

- 5.1 City mileage, highway mileage.** We expect a car's highway gas mileage to be related to its city gas mileage. Data for all 1040 vehicles in the government's 2010 *Fuel Economy Guide* give the regression line

$$\text{highway mpg} = 6.554 + (1.016 \times \text{city mpg})$$

for predicting highway mileage from city mileage.

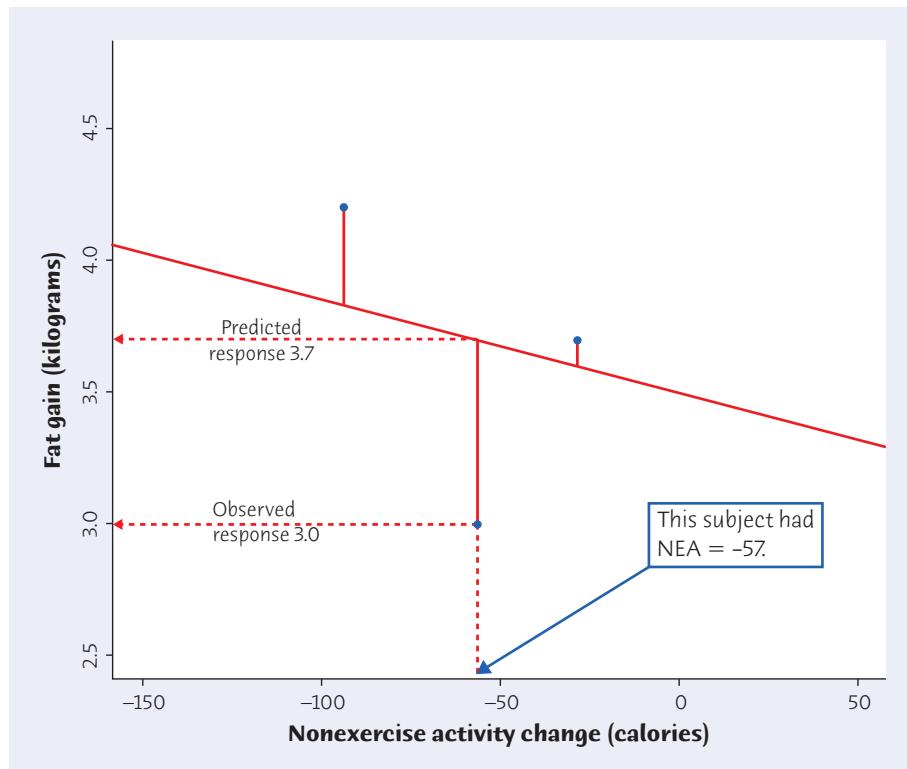
- (a) What is the slope of this line? Say in words what the numerical value of the slope tells you.
- (b) What is the intercept? Explain why the value of the intercept is not statistically meaningful.
- (c) Find the predicted highway mileage for a car that gets 16 miles per gallon in the city. Do the same for a car with city mileage 28 mpg.
- (d) Draw a graph of the regression line for city mileages between 10 and 50 mpg. (Be sure to show the scales for the x and y axes.)

- 5.2 What's the line?** You use the same bar of soap to shower each morning. The bar weighs 80 grams when it is new. Its weight goes down by 5 grams per day on the average. What is the equation of the regression line for predicting weight from days of use?

THE LEAST-SQUARES REGRESSION LINE

In most cases, no line will pass exactly through all the points in a scatterplot. Different people will draw different lines by eye. We need a way to draw a regression line that doesn't depend on our guess as to where the line should go. Because we use the line to predict y from x , the prediction errors we make are errors in y , the vertical direction in the scatterplot. A good regression line makes the vertical distances of the points from the line as small as possible.

Figure 5.2 illustrates the idea. This plot shows three of the points from Figure 5.1, along with the line, on an expanded scale. The line passes above one of the points and below two of them. The three prediction errors appear as vertical line segments. For example, one subject had $x = -57$, a decrease of 57 calories in NEA.

**FIGURE 5.2**

The least-squares idea. For each observation, find the vertical distance of each point on the scatterplot from a regression line. The least-squares regression line makes the sum of the squares of these distances as small as possible.

The line predicts a fat gain of 3.7 kilograms, but the actual fat gain for this subject was 3.0 kilograms. The prediction error is

$$\begin{aligned} \text{error} &= \text{observed response} - \text{predicted response} \\ &= 3.0 - 3.7 = -0.7 \text{ kilogram} \end{aligned}$$

There are many ways to make the collection of vertical distances “as small as possible.” The most common is the *least-squares* method.

LEAST-SQUARES REGRESSION LINE

The **least-squares regression line** of y on x is the line that makes the sum of the squares of the vertical distances of the data points from the line as small as possible.

One reason for the popularity of the least-squares regression line is that the problem of finding the line has a simple answer. We can give the equation for the least-squares line in terms of the means and standard deviations of the two variables and the correlation between them.

EQUATION OF THE LEAST-SQUARES REGRESSION LINE

We have data on an explanatory variable x and a response variable y for n individuals. From the data, calculate the means \bar{x} and \bar{y} and the standard deviations s_x and s_y of the two variables, and their correlation r . The least-squares regression line is the line

$$\hat{y} = a + bx$$

with slope

$$b = r \frac{s_y}{s_x}$$

and intercept

$$a = \bar{y} - b\bar{x}$$

We write \hat{y} (read “y hat”) in the equation of the regression line to emphasize that the line gives a *predicted* response \hat{y} for any x . Because of the scatter of points about the line, the predicted response will usually not be exactly the same as the actually *observed* response y . In practice, you don’t need to calculate the means, standard deviations, and correlation first. Software or your calculator will give the slope b and intercept a of the least-squares line from the values of the variables x and y . You can then concentrate on understanding and using the regression line.

USING TECHNOLOGY

Least-squares regression is one of the most common statistical procedures. Any technology you use for statistical calculations will give you the least-squares line and related information. Figure 5.3 displays the regression output for the data of Examples 5.1 and 5.2 from a graphing calculator, two statistical programs, and a spreadsheet program. Each output records the slope and intercept of the least-squares line. The software also provides information that we do not yet need, although we will use much of it later. (In fact, we left out part of the Minitab and Excel outputs.) Be sure that you can locate the slope and intercept on all four outputs. Once you understand the statistical ideas, you can read and work with almost any software output.

APPLY YOUR KNOWLEDGE

5.3 Coral reefs. Exercises 4.2 and 4.10 discuss a study in which scientists examined data on mean sea surface temperatures (in degrees Celsius) and mean coral growth (in millimeters per year) over a several-year period at locations in the Red Sea. Here are the data:² 

Sea surface temperature	29.68	29.87	30.16	30.22	30.48	30.65	30.90
Growth	2.63	2.58	2.60	2.48	2.26	2.38	2.26

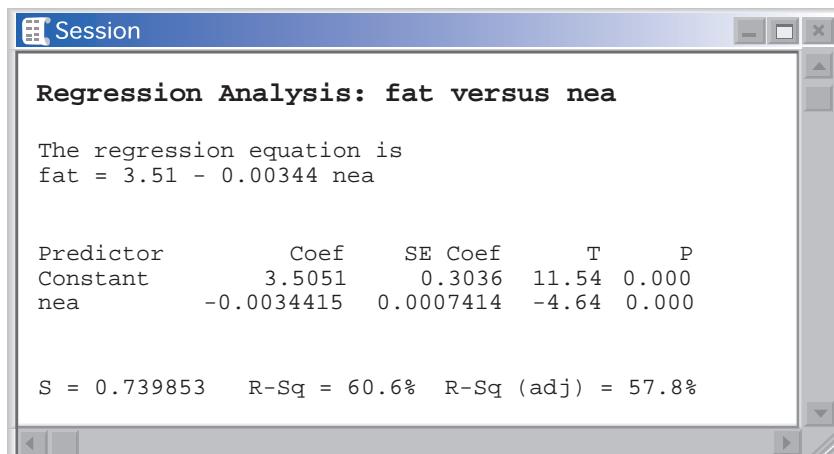
Texas Instruments Graphing Calculator

```

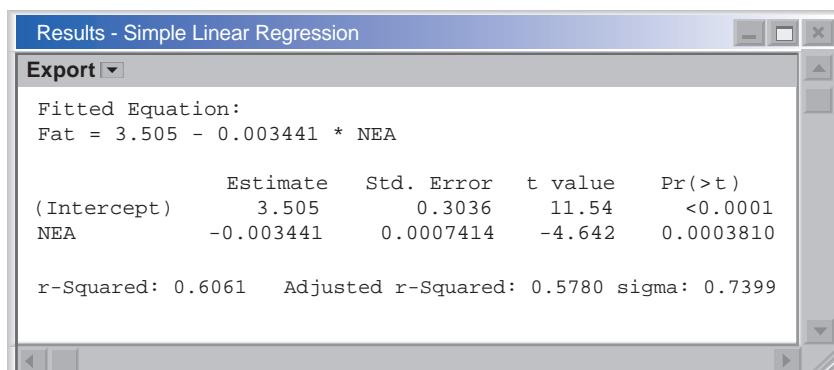
LinReg
y=a+bx
a=3.505122916
b=-.0003441487
r2=.6061492049
r=-.7785558457

```

Minitab



CrunchIt!



Microsoft Excel

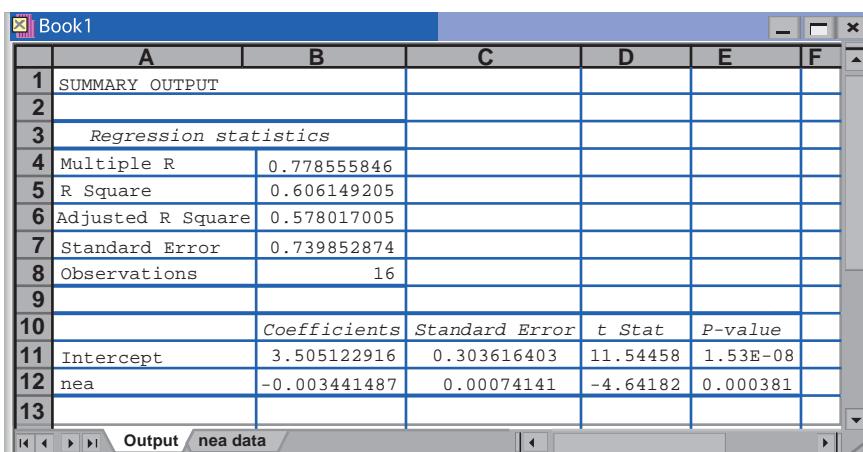


FIGURE 5.3

Least-squares regression for the non-exercise activity data: output from a graphing calculator, two statistical programs, and a spreadsheet program.

- Use your calculator to find the mean and standard deviation of both sea surface temperature x and growth y and the correlation r between x and y . Use these basic measures to find the equation of the least-squares line for predicting y from x .
- Enter the data into your software or calculator and use the regression function to find the least-squares line. The result should agree with your work in (a) up to roundoff error.

5.4 Do heavier people burn more energy? We have data on the lean body mass and resting metabolic rate for 12 women who are subjects in a study of dieting. Lean body mass, given in kilograms, is a person's weight leaving out all fat. Metabolic rate, in calories burned per 24 hours, is the rate at which the body consumes energy.  METABOLIC

Mass	36.1	54.6	48.5	42.0	50.6	42.0	40.3	33.1	42.4	34.5	51.1	41.2
Rate	995	1425	1396	1418	1502	1256	1189	913	1124	1052	1347	1204

- Make a scatterplot that shows how metabolic rate depends on body mass. There is a quite strong linear relationship, with correlation $r = 0.876$.
- Find the least-squares regression line for predicting metabolic rate from body mass. Add this line to your scatterplot.
- Explain in words what the slope of the regression line tells us.
- Another woman has a lean body mass of 45 kilograms. What is her predicted metabolic rate?

FACTS ABOUT LEAST-SQUARES REGRESSION

One reason for the popularity of least-squares regression lines is that they have many convenient properties. Here are some facts about least-squares regression lines.

Fact 1. The distinction between explanatory and response variables is essential in regression. Least-squares regression makes the distances of the data points from the line small only in the y direction. If we reverse the roles of the two variables, we get a different least-squares regression line.

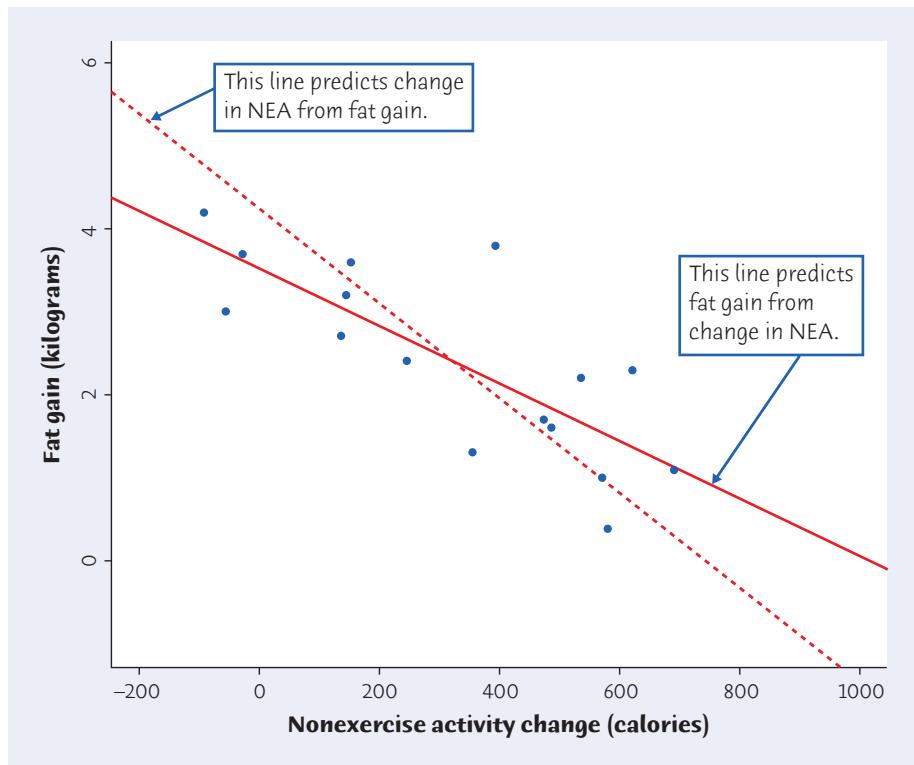
EXAMPLE 5.3 Predicting fat gain, predicting change in NEA

Figure 5.4 repeats the scatterplot of the NEA data in Figure 5.1, but with two least-squares regression lines. The solid line is the regression line for predicting fat gain from change in NEA. This is the line that appeared in Figure 5.1.

We might also use the data on these 16 subjects to predict the change in NEA for another subject from that subject's fat gain when overfed for 8 weeks. Now the roles of the variables are reversed: fat gain is the explanatory variable and change in NEA is the



response variable. The dashed line in Figure 5.4 is the least-squares line for predicting NEA change from fat gain. The two regression lines are not the same. *In the regression setting, you must know clearly which variable is explanatory.* ■

**FIGURE 5.4**

Two least-squares regression lines for the nonexercise activity data, for Example 5.3. The solid line predicts fat gain from change in nonexercise activity. The dashed line predicts change in nonexercise activity from fat gain.

Fact 2. There is a close connection between correlation and the slope of the least-squares line. The slope is

$$b = r \frac{s_y}{s_x}$$

You see that the slope and the correlation always have the same sign. For example, if a scatterplot shows a positive association, then both b and r are positive. The formula for the slope b says more: along the regression line, a **change of one standard deviation in x corresponds to a change of r standard deviations in y .** When the variables are perfectly correlated ($r = 1$ or $r = -1$), the change in the predicted response \hat{y} is the same (in standard deviation units) as the change in x . Otherwise, because $-1 \leq r \leq 1$, the change in \hat{y} (in standard deviation units) is less than the change in x . As the correlation grows less strong, the prediction \hat{y} moves less in response to changes in x .

Fact 3. The least-squares regression line always passes through the point (\bar{x}, \bar{y}) on the graph of y against x . This is a consequence of the equation of the least-squares regression line (box on page 130). In Exercise 5.48 we ask you to confirm this.

Fact 4. The correlation r describes the strength of a straight-line relationship. In the regression setting, this description takes a specific form: the square of the correlation, r^2 , is the fraction of the variation in the values of y that is explained by the least-squares regression of y on x .

The idea is that when there is a linear relationship, some of the variation in y is accounted for by the fact that as x changes, y changes along with it. Look again at Figure 5.1 (page 126), the scatterplot of the NEA data. The variation in y appears as the spread of fat gains from 0.4 to 4.2 kg. Some of this variation is explained by the fact that x (change in NEA) varies from a loss of 94 calories to a gain of 690 calories. As x changes from -94 to 690 , y changes along the line. You would predict a smaller fat gain for a subject whose NEA increased by 600 calories than for someone with 0 change in NEA. But the straight-line tie of y to x doesn't explain *all* of the variation in y . The remaining variation appears as the scatter of points above and below the line.

Although we won't do the algebra, it is possible to break the variation in the observed values of y into two parts. One part measures the variation in \hat{y} along the least-squares regression line as x varies. The other measures the vertical scatter of the data points above and below the line. The squared correlation r^2 is the first of these as a fraction of the whole:

$$r^2 = \frac{\text{variation in } \hat{y} \text{ along the regression line as } x \text{ varies}}{\text{total variation in observed values of } y}$$

EXAMPLE 5.4 Using r^2

For the NEA data, $r = -0.7786$ and $r^2 = (-0.7786)^2 = 0.6062$. About 61% of the variation in fat gained is accounted for by the linear relationship with change in NEA. The other 39% is individual variation among subjects that is not explained by the linear relationship.

Figure 4.2 (page 103) shows a stronger linear relationship between boat registrations in Florida and manatees killed by boats. The correlation is $r = 0.951$ and $r^2 = (0.951)^2 = 0.904$. Slightly more than 90% of the year-to-year variation in number of manatees killed by boats is explained by regression on number of boats registered. Only about 10% is variation among years with similar numbers of boats registered. ■



You can find a regression line for any relationship between two quantitative variables, but the usefulness of the line for prediction depends on the strength of the linear relationship. So r^2 is almost as important as the equation of the line in reporting a regression. All the outputs in Figure 5.3 (page 131) include r^2 , either in decimal form or as a percent. When you see a correlation, square it to get a better feel for the strength of the association. Perfect correlation ($r = -1$ or $r = 1$) means the points lie exactly on a line. Then $r^2 = 1$ and all the variation in one variable is accounted for by the linear relationship with the other variable. If $r = -0.7$ or $r = 0.7$, $r^2 = 0.49$ and about half the variation is accounted for by the linear relationship. In the r^2 scale, correlation ± 0.7 is about halfway between 0 and ± 1 .

Facts 2, 3, and 4 are special properties of least-squares regression. They are not true for other methods of fitting a line to data.

APPLY YOUR KNOWLEDGE

5.5 How useful is regression? Figure 4.8 (page 115) displays the relationship between golfers' scores on the first and second rounds of the 2010 Masters Tournament. The correlation is $r = 0.347$. Exercise 4.30 gives data on solar radiation (SRD) and concentration of dimethyl sulfide (DMS) over a region of the Mediterranean. The correlation is $r = 0.969$. Explain in simple language why knowing only these correlations enables you to say that prediction of DMS from SRD by a regression line will be much more accurate than prediction of a golfer's second-round score from his first-round score.

5.6 Feed the birds. Exercise 4.32 (page 118) gives data from a study in which canary parents cared for both their own babies and those of other parents. Investigators looked at how the growth rate of the foster babies relative to the growth rate of the natural babies changed as the begging intensity for food by the foster babies increased over the begging intensity of the natural babies. If begging intensity is the main factor determining food received, with higher intensity leading to more food, one would expect the relative growth rate to increase as the difference in begging intensity increases. However, if both begging intensity and a preference for their own babies determine the amount of food received (and hence the relative growth rate), we might expect growth rate to increase initially as begging intensity increases but then to level off (or even decrease) as the parents begin to ignore further increases in begging by the foster babies.  CANARIES

- Make a scatterplot of the data. Find the least-squares regression line for predicting relative growth rate of the foster brood from the difference in begging intensity between the foster brood and the actual babies of the parents and add this line to your plot. Should we not use the regression line for prediction in this setting?
- What is r^2 ? What does this value say about the success of the regression line in predicting relative growth rate?



Arco Images GmbH/Alamy

RESIDUALS

One of the first principles of data analysis is to look for an overall pattern and also for striking deviations from the pattern. A regression line describes the overall pattern of a linear relationship between an explanatory variable and a response variable. We see deviations from this pattern by looking at the scatter of the data points about the regression line. The vertical distances from the points to the least-squares regression line are as small as possible, in the sense that they have the smallest possible sum of squares. Because they represent "leftover" variation in the response after fitting the regression line, these distances are called *residuals*.

RESIDUALS

A **residual** is the difference between an observed value of the response variable and the value predicted by the regression line. That is, a residual is the prediction error that remains after we have chosen the regression line:

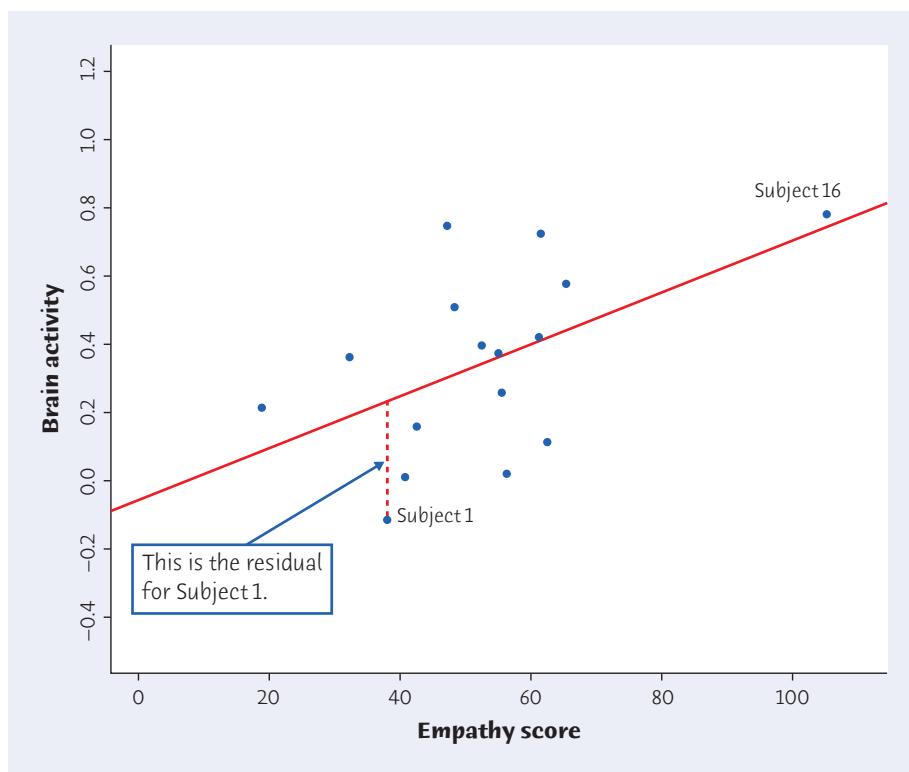
$$\begin{aligned}\text{residual} &= \text{observed } y - \text{predicted } y \\ &= y - \hat{y}\end{aligned}$$

**EMPATHY****EXAMPLE 5.5 I feel your pain**

“Empathy” means being able to understand what others feel. To see how the brain expresses empathy, researchers recruited 16 couples in their midtwenties who were married or had been dating for at least two years. They zapped the man’s hand with an electrode while the woman watched and measured the activity in several parts of the woman’s brain that would respond to her own pain. Brain activity was recorded as a fraction of the activity observed when the woman herself was zapped with the electrode. The women also completed a psychological test that measures empathy. Will women who score higher in empathy respond more strongly when their partner has a painful experience? Here are data for one brain region:³

Subject	1	2	3	4	5	6	7	8
Empathy score	38	53	41	55	56	61	62	48
Brain activity	-0.120	0.392	0.005	0.369	0.016	0.415	0.107	0.506
Subject	9	10	11	12	13	14	15	16
Empathy score	43	47	56	65	19	61	32	105
Brain activity	0.153	0.745	0.255	0.574	0.210	0.722	0.358	0.779

Figure 5.5 is a scatterplot, with empathy score as the explanatory variable x and brain activity as the response variable y . The plot shows a positive association. That is,

**FIGURE 5.5**

Scatterplot of activity in a region of the brain that responds to pain versus score on a test of empathy, for Example 5.5. Brain activity is measured as the subject watches her partner experience pain. The line is the least-squares regression line.

women who score higher in empathy do indeed react more strongly to their partner's pain. The overall pattern is moderately linear, with correlation $r = 0.515$.

The line on the plot is the least-squares regression line of brain activity on empathy score. Its equation is

$$\hat{y} = -0.0578 + 0.00761x$$

For Subject 1, with empathy score 38, we predict

$$\hat{y} = -0.0578 + (0.00761)(38) = 0.231$$

This subject's actual brain activity level was -0.120 . The residual is

$$\begin{aligned}\text{residual} &= \text{observed } y - \text{predicted } y \\ &= -0.120 - 0.231 = -0.351\end{aligned}$$

The residual is negative because the data point lies below the regression line. The dashed line segment in Figure 5.5 shows the size of the residual. ■

There is a residual for each data point. Finding the residuals is a bit unpleasant because you must first find the predicted response for every x . Software or a graphing calculator gives you the residuals all at once. Here are the 16 residuals for the empathy study data, from software:

```
residuals:  
-0.3515 -0.2494 -0.3526 -0.3072 -0.1166 -0.1136 0.1231 0.1721  
0.0463 0.0080 0.0084 0.1983 0.4449 0.1369 0.3154 0.0374
```

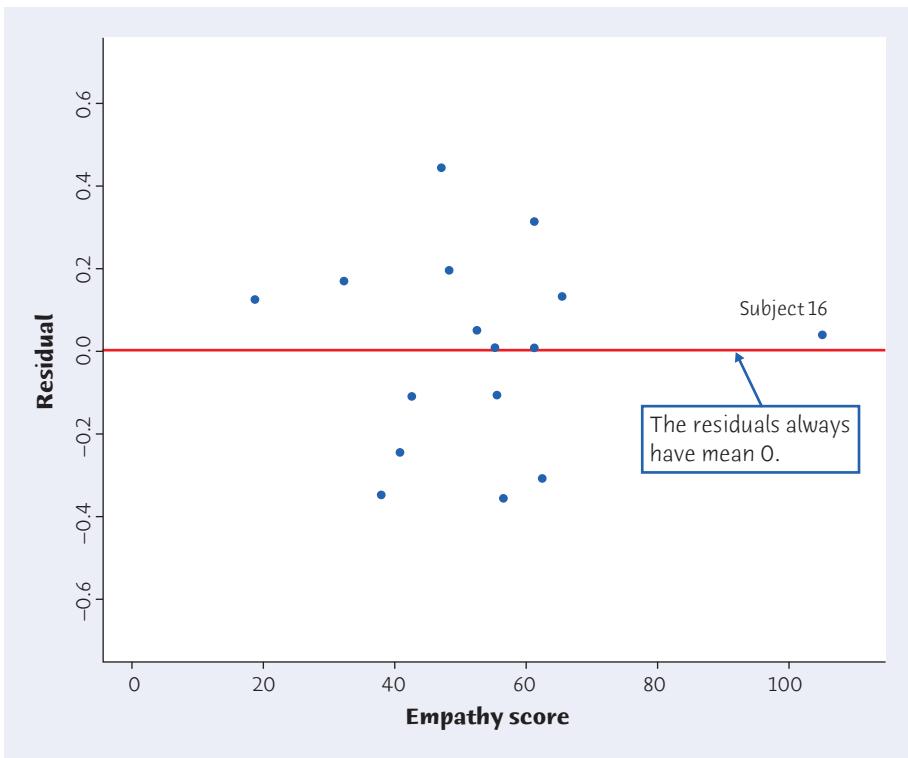
Because the residuals show how far the data fall from our regression line, examining the residuals helps us assess how well the line describes the data. Although residuals can be calculated from any curve or line fitted to the data, the residuals from the least-squares line have a special property: **the mean of the least-squares residuals is always zero.**

Compare the scatterplot in Figure 5.5 with the *residual plot* for the same data in Figure 5.6. The horizontal line at zero in Figure 5.6 helps orient us. This “residual = 0” line corresponds to the regression line in Figure 5.5.

RESIDUAL PLOTS

A **residual plot** is a scatterplot of the regression residuals against the explanatory variable. Residual plots help us assess how well a regression line fits the data.

A residual plot in effect turns the regression line horizontal. It magnifies the deviations of the points from the line and makes it easier to see unusual observations and patterns.

**FIGURE 5.6**

Residual plot for the data shown in Figure 5.5. The horizontal line at zero residual corresponds to the regression line in Figure 5.5.

APPLY YOUR KNOWLEDGE

5.7 Residuals by hand. In Exercise 5.3 (page 130) you found the equation of the least-squares line for predicting coral growth y from mean sea surface temperature x .

- Use the equation to obtain the 7 residuals step-by-step. That is, find the prediction \hat{y} for each observation and then find the residual $y - \hat{y}$.
- Check that (up to roundoff error) the residuals add to 0.
- The residuals are the part of the response y left over after the straight-line tie between y and x is removed. Show that the correlation between the residuals and x is 0 (up to roundoff error). That this correlation is always 0 is another special property of least-squares regression.

5.8 Does fast driving waste fuel? Exercise 4.8 (page 104) gives data on the fuel consumption y of a car at various speeds x . Fuel consumption is measured in liters of gasoline per 100 kilometers driven, and speed is measured in kilometers per hour. Software tells us that the equation of the least-squares regression line is

$$\hat{y} = 11.058 - 0.01466x$$

Using this equation we can add the residuals to the original data:  FASTDRIVE2

Speed	10	20	30	40	50	60	70	80
Fuel	21.00	13.00	10.00	8.00	7.00	5.90	6.30	6.95
Residual	10.09	2.24	-0.62	-2.47	-3.33	-4.28	-3.73	-2.94
Speed	90	100	110	120	130	140	150	
Fuel	7.57	8.27	9.03	9.87	10.79	11.77	12.83	
Residual	-2.17	-1.32	-0.42	0.57	1.64	2.76	3.97	

- (a) Make a scatterplot of the observations and draw the regression line on your plot.
- (b) Would you use the regression line to predict y from x ? Explain your answer.
- (c) Verify the value of the first residual, for $x = 10$. Verify that the residuals have sum zero (up to roundoff error).
- (d) Make a plot of the residuals against the values of x . Draw a horizontal line at height zero on your plot. How does the pattern of the residuals about this line compare with the pattern of the data points about the regression line in your scatterplot from (a)?

INFLUENTIAL OBSERVATIONS

Figures 5.5 and 5.6 show one unusual observation. Subject 16 is an outlier in the x direction, with empathy score 40 points higher than any other subject. Because of its extreme position on the empathy scale, this point has a strong influence on the correlation. Dropping Subject 16 reduces the correlation from $r = 0.515$ to $r = 0.331$. You can see that this point extends the linear pattern in Figure 5.5 and so increases the correlation. We say that Subject 16 is *influential* for calculating the correlation.

INFLUENTIAL OBSERVATIONS

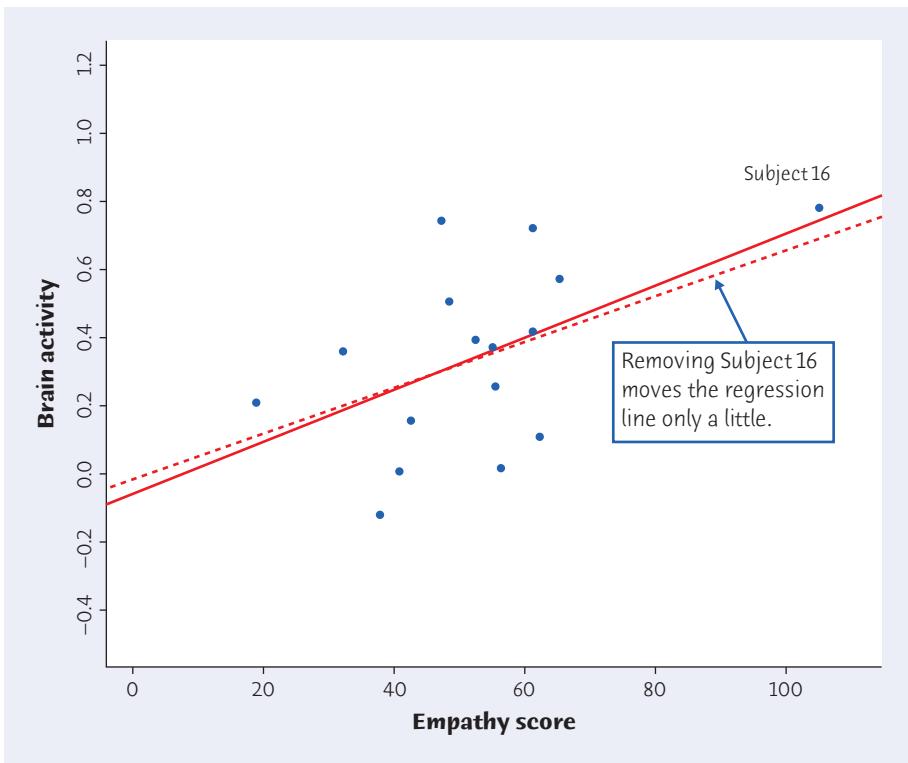
An observation is **influential** for a statistical calculation if removing it would markedly change the result of the calculation.

The result of a statistical calculation may be of little practical use if it depends strongly on a few influential observations.

Points that are outliers in either the x or the y direction of a scatterplot are often influential for the correlation. Points that are outliers in the x direction are often influential for the least-squares regression line.

EXAMPLE 5.6 An influential observation?

Subject 16 in Example 5.5 is influential for the correlation between empathy score and brain activity because removing it reduces r from 0.515 to 0.331. Calculating that $r = 0.515$ is not a very useful description of the data, because the value depends so strongly on just one of the 16 subjects.

**FIGURE 5.7**

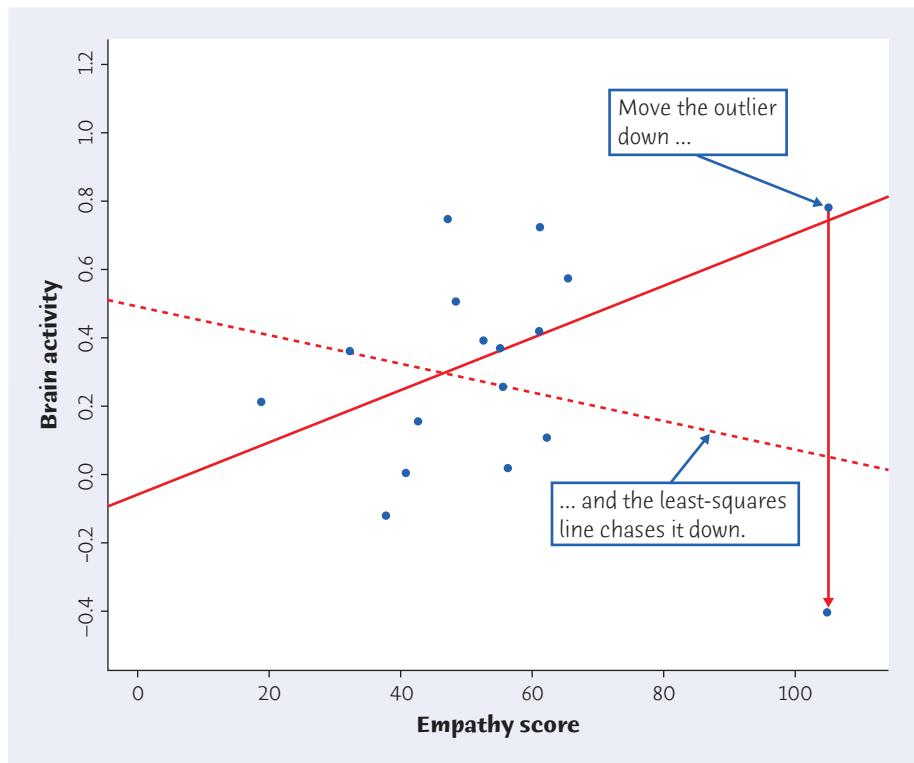
Subject 16 is an outlier in the x direction. The outlier is not influential for least-squares regression, because removing it moves the regression line only a little.

Is this observation also influential for the least-squares line? Figure 5.7 shows that it is not. The regression line calculated without Subject 16 (dashed) differs little from the line that uses all the observations (solid). The reason that the outlier has little influence on the regression line is that it lies close to the dashed regression line calculated from the other observations. ■

To see why points that are outliers in the x direction are often influential for regression, let's try an experiment. Suppose that Subject 16's point in the scatterplot moves straight down. What happens to the regression line? Figure 5.8 gives the answer. The dashed line is the regression line with the outlier in its new, lower position. Because there are no other points with similar x -values, the line chases the outlier. The *Correlation and Regression* applet allows you to try this experiment yourself—see Exercise 5.9. *An outlier in x pulls the least-squares line toward itself. If the outlier does not lie close to the line calculated from the other observations, it will be influential.*



We did not need the distinction between outliers and influential observations in Chapter 2. A single high salary that pulls up the mean salary \bar{x} for a group of workers is an outlier because it lies far above the other salaries. It is also influential, because the mean changes when it is removed. In the regression setting, however, not all outliers are influential.

**FIGURE 5.8**

An outlier in the x direction pulls the least-squares line to itself because there are no other observations with similar values of x to hold the line in place. When the outlier moves down, the regression line chases it down. The original regression line is solid, and the final position of the regression line is dashed.

APPLY YOUR KNOWLEDGE

5.9 Influence in regression. The *Correlation and Regression* applet allows you to animate Figure 5.8. Click to create a group of 10 points in the lower-left corner of the scatterplot with a strong straight-line pattern (correlation about 0.9). Click the “Show least-squares line” box to display the regression line.



- Add 1 point at the upper right that is far from the other 10 points but exactly on the regression line. Why does this outlier have no effect on the line even though it changes the correlation?
- Now use the mouse to drag this last point straight down. You see that one end of the least-squares line chases this single point, while the other end remains near the middle of the original group of 10. What makes the last point so influential?

5.10 Do heavier people burn more energy? Return to the data of Exercise 5.4 (page 132) on body mass and metabolic rate. We will use these data to illustrate influence.

- Make a scatterplot of the data that is suitable for predicting metabolic rate from body mass, with two new points added. Point A: mass 42 kilograms, metabolic rate 1500 calories. Point B: mass 70 kilograms, metabolic rate 1400 calories. In which direction is each of these points an outlier?
- Add three least-squares regression lines to your plot: for the original 12 women, for the original women plus Point A, and for the original women plus Point B. Which new point is more influential for the regression line? Explain in simple language why each new point moves the line in the way your graph shows. METABOLIC2

5.11 Outsourcing by airlines. Exercise 4.5 (page 101) gives data for 12 airlines on the percent of major maintenance outsourced and the percent of flight delays blamed on the airline.

- Make a scatterplot with outsourcing percent as x and delay percent as y . Would you consider Hawaiian Airlines to be influential?
- Find the correlation r with and without Hawaiian Airlines. How influential is the outlier for correlation?
- Find the least-squares line for predicting y from x with and without Hawaiian Airlines. Draw both lines on your scatterplot. Use both lines to predict the percent of delays blamed on an airline that has outsourced 74.1% of its major maintenance. How influential is the outlier for the least-squares line? 

CAUTIONS ABOUT CORRELATION AND REGRESSION

Correlation and regression are powerful tools for describing the relationship between two variables. When you use these tools, you must be aware of their limitations. You already know that

-  ■ *Correlation and regression lines describe only linear relationships.* You can do the calculations for any relationship between two quantitative variables, but the results are useful only if the scatterplot shows a linear pattern.
-  ■ *Correlation and least-squares regression lines are not resistant.* Always plot your data and look for observations that may be influential.

Here are three more things to keep in mind when you use correlation and regression.

Beware ecological correlation. There is a large positive correlation between *average* income and number of years of education. The correlation is smaller if we compare the incomes of *individuals* with number of years of education. The correlation based on average income ignores the large variation in the incomes of individuals having the same amount of education. The variation from individual to individual increases the scatter in a scatterplot, reducing the correlation. The correlation between average income and education overstates the strength of the relation between the incomes of individuals and number of years of education. *Correlations based on averages can be misleading if they are interpreted to be about individuals.*

ECOLOGICAL CORRELATION

A correlation based on averages rather than on individuals is called an **ecological correlation**.

Beware extrapolation. Suppose that you have data on a child's growth between 3 and 8 years of age. You find a strong linear relationship between age x and height y . If you fit a regression line to these data and use it to predict height at age 25 years, you will predict that the child will be 8 feet tall. Growth slows down and then stops at maturity, so extending the straight line to adult ages is foolish. Few relationships are linear for all values of x . Don't make predictions far outside the range of x that actually appears in your data.



EXTRAPOLATION

Extrapolation is the use of a regression line for prediction far outside the range of values of the explanatory variable x that you used to obtain the line. Such predictions are often not accurate.

Beware the lurking variable. Another caution is even more important: *the relationship between two variables can often be understood only by taking other variables into account. Lurking variables can make a correlation or regression misleading.*



LURKING VARIABLE

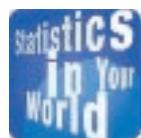
A **lurking variable** is a variable that is not among the explanatory or response variables in a study and yet may influence the interpretation of relationships among those variables.

You should always think about possible lurking variables before you draw conclusions based on correlation or regression.

EXAMPLE 5.7 Magic Mozart?

The Kalamazoo (Michigan) Symphony once advertised a "Mozart for Minors" program with this statement: "Question: Which students scored 51 points higher in verbal skills and 39 points higher in math? Answer: Students who had experience in music."⁴

We could as well answer "Students who played soccer." Why? Children with prosperous and well-educated parents are more likely than poorer children to have experience with music and also to play soccer. They are also likely to attend good schools, get good health care, and be encouraged to study hard. These advantages lead to high test scores. Family background is a lurking variable that explains why test scores are related to experience with music. ■



Do left-handers die early?

Yes, said a study of 1000 deaths in California. Left-handed people died at an average age of 66 years; right-handers, at 75 years of age. Should left-handed people fear an early death? No—the lurking variable has struck again. Older people grew up in an era when many natural left-handers were forced to use their right hands. So right-handers are more common among older people, and left-handers are more common among the young. When we look at deaths, the left-handers who die are younger on the average because left-handers in general are younger. Mystery solved.

APPLY YOUR KNOWLEDGE

5.12 One more inch, three more pounds. Data on the *average* weight of men who are between 5 feet 2 inches and 6 feet 4 inches tall (rounded to the nearest inch) show a very high positive correlation. Would the correlation be greater, smaller, or about

the same if you calculated the correlation between the weights of individual men and their heights (rounded to the nearest inch)? Explain your answer.

5.13 The endangered manatee. Table 4.1 gives 33 years of data on boats registered in Florida and manatees killed by boats. Figure 4.2 (page 103) shows a strong positive linear relationship. The correlation is $r = 0.951$.

- Find the equation of the least-squares line for predicting manatees killed from thousands of boats registered. Because the linear pattern is so strong, we expect predictions from this line to be quite accurate—but only if conditions in Florida remain similar to those of the past 33 years.
- In 2009, experts predicted that the number of boats registered in Florida would be 975,000 in 2010. What would you predict the number of manatees killed by boats to be if there are 975,000 boats registered? Explain why we can trust this prediction.
- Predict manatee deaths if there were no boats registered in Florida. Explain why the predicted count of deaths is impossible. (We use $x = 0$ to find the intercept of the regression line, but unless the explanatory variable x actually takes values near 0, prediction for $x = 0$ is an example of extrapolation.)  MANATEES

5.14 Is math the key to success in college? A College Board study of 15,941 high school graduates found a strong correlation between how much math minority students took in high school and their later success in college. News articles quoted the head of the College Board as saying that “math is the gatekeeper for success in college.”⁵ Maybe so, but we should also think about lurking variables. What might lead minority students to take more or fewer high school math courses? Would these same factors influence success in college?

ASSOCIATION DOES NOT IMPLY CAUSATION



Thinking about lurking variables leads to the most important caution about correlation and regression. When we study the relationship between two variables, we often hope to show that changes in the explanatory variable *cause* changes in the response variable. A *strong association between two variables is not enough to draw conclusions about cause and effect*. Sometimes an observed association really does reflect cause and effect. A household that heats with natural gas uses more gas in colder months because cold weather requires burning more gas to stay warm. In other cases, an association is explained by lurking variables, and the conclusion that x causes y is either wrong or not proved.

EXAMPLE 5.8 Does having more cars make you live longer?

A serious study once found that people with two cars live longer than people who own only one car.⁶ Owning three cars is even better, and so on. There is a substantial positive correlation between number of cars x and length of life y .

The basic meaning of causation is that by changing x we can bring about a change in y . Could we lengthen our lives by buying more cars? No. The study used number of cars as a quick indicator of affluence. Well-off people tend to have more cars. They

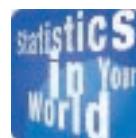


also tend to live longer, probably because they are better educated, take better care of themselves, and get better medical care. The cars have nothing to do with it. There is no cause-and-effect tie between number of cars and length of life. ■

Correlations such as that in Example 5.8 are sometimes called “nonsense correlations.” The correlation is real. What is nonsense is the conclusion that changing one of the variables causes changes in the other. A lurking variable—such as personal affluence in Example 5.8—that influences both x and y can create a high correlation even though there is no direct connection between x and y .

ASSOCIATION DOES NOT IMPLY CAUSATION

An association between an explanatory variable x and a response variable y , even if it is very strong, is not by itself good evidence that changes in x actually cause changes in y .



The Super Bowl effect

The Super Bowl is the most-watched TV broadcast in the United States. Data show that on Super Bowl Sunday we consume 3 times as many potato chips as on an average day, and 17 times as much beer. What's more, the number of fatal traffic accidents goes up in the hours after the game ends. Could that be celebration? Or catching up with tasks left undone? Or maybe it's the beer.

EXAMPLE 5.9 Overweight mothers, overweight daughters

Overweight parents tend to have overweight children. The results of a study of Mexican American girls aged 9 to 12 years are typical. The investigators measured body mass index (BMI), a measure of weight relative to height, for both the girls and their mothers. People with high BMI are overweight. The correlation between the BMI of daughters and the BMI of their mothers was $r = 0.506$.⁷

Body type is in part determined by heredity. Daughters inherit half their genes from their mothers. There is therefore a direct cause-and-effect link between the BMI of mothers and daughters. But perhaps mothers who are overweight also set an example of little exercise, poor eating habits, and lots of television. Their daughters may pick up these habits, so the influence of heredity is mixed up with influences from the girls' environment. Both contribute to the mother-daughter correlation. ■

The lesson of Example 5.9 is more subtle than just “association does not imply causation.” Even when direct causation is present, it may not be the whole explanation for a correlation. You must still worry about lurking variables. Careful statistical studies try to anticipate lurking variables and measure them. The mother-daughter study did measure TV viewing, exercise, and diet. Elaborate statistical analysis can remove the effects of these variables to come closer to the direct effect of mother's BMI on daughter's BMI. This remains a second-best approach to causation. The best way to get good evidence that x causes y is to do an **experiment** in which we change x and keep lurking variables under control. We will discuss experiments in Chapter 9.



When experiments cannot be done, explaining an observed association can be difficult and controversial. Many of the sharpest disputes in which statistics plays a role involve questions of causation that cannot be settled by experiment. Do gun control laws reduce violent crime? Does using cell phones cause brain tumors? Has increased free trade widened the gap between the incomes of more educated

experiment

and less educated American workers? All these questions have become public issues. All concern associations among variables. And all have this in common: they try to pinpoint cause and effect in a setting involving complex relations among many interacting variables.

James Leyendecker/CORBIS



EXAMPLE 5.10 Does smoking cause lung cancer?

Despite the difficulties, it is sometimes possible to build a strong case for causation in the absence of experiments. The evidence that smoking causes lung cancer is about as strong as nonexperimental evidence can be.

Doctors had long observed that most lung cancer patients were smokers. Comparison of smokers and “similar” nonsmokers showed a very strong association between smoking and death from lung cancer. Could the association be explained by lurking variables? Might there be, for example, a genetic factor that predisposes people both to nicotine addiction and to lung cancer? Smoking and lung cancer would then be positively associated even if smoking had no direct effect on the lungs. How were these objections overcome? ■

Let’s answer this question in general terms: what are the criteria for establishing causation when we cannot do an experiment?

- *The association is strong.* The association between smoking and lung cancer is very strong.
- *The association is consistent.* Many studies of different kinds of people in many countries link smoking to lung cancer. That reduces the chance that a lurking variable specific to one group or one study explains the association.
- *Higher doses are associated with stronger responses.* People who smoke more cigarettes per day or who smoke over a longer period get lung cancer more often. People who stop smoking reduce their risk.
- *The alleged cause precedes the effect in time.* Lung cancer develops after years of smoking. The number of men dying of lung cancer rose as smoking became more common, with a lag of about 30 years. Lung cancer kills more men than any other form of cancer. Lung cancer was rare among women until women began to smoke. Lung cancer in women rose along with smoking, again with a lag of about 30 years, and has now passed breast cancer as the leading cause of cancer death among women.
- *The alleged cause is plausible.* Experiments with animals show that tars from cigarette smoke do cause cancer.

Medical authorities do not hesitate to say that smoking causes lung cancer. The U.S. Surgeon General has long stated that cigarette smoking is “the largest avoidable cause of death and disability in the United States.”⁸ The evidence for causation is overwhelming—but it is not as strong as the evidence provided by well-designed experiments.



APPLY YOUR KNOWLEDGE

5.15 Another reason not to smoke? A stop-smoking booklet says, “Children of mothers who smoked during pregnancy scored nine points lower on intelligence tests at ages three and four than children of nonsmokers.” Suggest some lurking variables that may help explain the association between smoking during pregnancy and children’s later test scores. The association by itself is not good evidence that mothers’ smoking *causes* lower scores.

5.16 Education and income. There is a strong positive association between workers’ education and their income. For example, the U.S. Census Bureau reported in 2008 that the median income of young adults (ages 25 to 34) who worked full-time increased from \$20,260 for those with less than a ninth-grade education, to \$30,543 for high school graduates, to \$46,932 for holders of a bachelor’s degree, and on up for yet more education. In part, this association reflects causation—education helps people qualify for better jobs. Suggest several lurking variables that also contribute. (Ask yourself what kinds of people tend to get more education.)

5.17 To earn more, get married? Data show that men who are married, and also divorced or widowed men, earn quite a bit more than men the same age who have never been married. This does not mean that a man can raise his income by getting married, because men who have never been married are different from married men in many ways other than marital status. Suggest several lurking variables that might help explain the association between marital status and income.



CHAPTER 5 SUMMARY

CHAPTER SPECIFICS

- A **regression line** is a straight line that describes how a response variable y changes as an explanatory variable x changes. You can use a regression line to **predict** the value of y for any value of x by substituting this x into the equation of the line.
- The **slope** b of a regression line $\hat{y} = a + bx$ is the rate at which the predicted response \hat{y} changes along the line as the explanatory variable x changes. Specifically, b is the change in \hat{y} when x increases by 1.
- The **intercept** a of a regression line $\hat{y} = a + bx$ is the predicted response \hat{y} when the explanatory variable $x = 0$. This prediction is of no statistical interest unless x can actually take values near 0.
- The most common method of fitting a line to a scatterplot is least squares. The **least-squares regression line** is the straight line $\hat{y} = a + bx$ that minimizes the sum of the squares of the vertical distances of the observed points from the line.
- The least-squares regression line of y on x is the line with slope $b = rs_y/s_x$ and intercept $a = \bar{y} - b\bar{x}$. This line always passes through the point (\bar{x}, \bar{y}) .
- **Correlation and regression** are closely connected. The correlation r is the slope of the least-squares regression line when we measure both x and y in standardized units. The **square of the correlation** r^2 is the fraction of the variation in one variable that is explained by least-squares regression on the other variable.

- Correlation and regression must be interpreted with caution. Plot the data to be sure the relationship is roughly linear and to detect outliers and influential observations. A plot of the residuals makes these effects easier to see.
- Look for **influential observations**, individual points that substantially change the correlation or the regression line. Outliers in the x direction are often influential for the regression line.
- Be aware of **ecological correlation**, the tendency for correlations based on averages to be stronger than correlations based on individuals. Be careful not to misinterpret correlations based on averages as applying to individuals.
- Avoid **extrapolation**, the use of a regression line for prediction for values of the explanatory variable far outside the range of the data from which the line was calculated.
- **Lurking variables** may explain the relationship between the explanatory and response variables. Correlation and regression can be misleading if you ignore important lurking variables.
- Most of all, be careful not to conclude that there is a cause-and-effect relationship between two variables just because they are strongly associated. **High correlation does not imply causation.** The best evidence that an association is due to causation comes from an **experiment** in which the explanatory variable is directly changed and other influences on the response are controlled.

LINK IT

In this chapter we use the least-squares regression line to describe the straight-line relationship between two variables when such a pattern is seen in a scatterplot. The equation of the least-squares regression line is a numerical summary that makes precise the notion of “straight-line relationship.”

To help us assess whether the least-squares regression line is a sensible description of the relationship between two variables, we examine residual plots. Outliers and, in particular, influential observations may indicate that the least-squares regression line is not a good description of this relationship.

Even if the least-squares regression line is a good description of the relationship between the observed values of two variables, we must exercise caution in how we interpret this relationship. Such interpretations rest on the assumption that the relationship is valid in some broader sense. We will explore this more carefully later in this book, but in this chapter we have issued some cautions. Association, as indicated by a large correlation, does not imply that there is an underlying cause-and-effect relation between the response and explanatory variables. There may, in fact, be a lurking variable that influences the interpretation of any relation between the response and explanatory variables. Correlations based on averages of measurements tend to be higher than correlations based on the individual observations used to compute the averages. Be careful not to misinterpret correlations based on averages as applying to individuals. Finally, be careful not to use the least-squares regression line to make predictions outside the range of values of the explanatory variable that you used to obtain the line.


CHECK YOUR SKILLS

5.18 Figure 5.9 is a scatterplot of school GPA against IQ test scores for 15 seventh-grade students. The line is the least-squares regression line for predicting school GPA from IQ score. If another child in this class has IQ score 110, you predict the school GPA to be close to

- (a) 2. (b) 7.5. (c) 11.

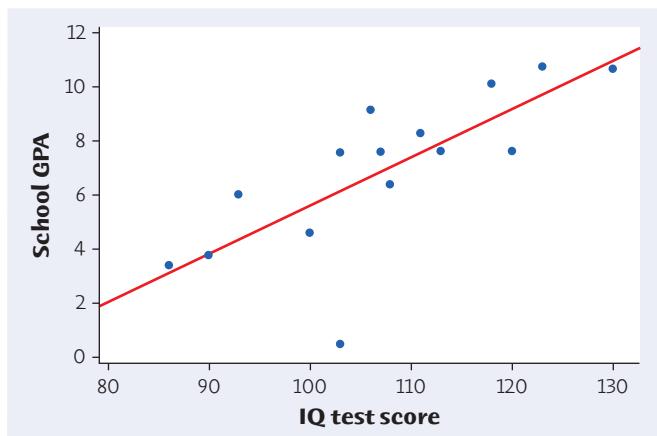


FIGURE 5.9

Scatterplot of IQ test scores and school GPA for 15 seventh-grade students, for Exercises 5.18 and 5.19.

5.19 The slope of the line in Figure 5.9 is closest to

- (a) -11. (b) 0.2. (c) 2.0.

5.20 The points on a scatterplot lie close to the line whose equation is $y = 4 - 3x$. The slope of this line is

- (a) 4. (b) 3. (c) -3.

5.21 Fred keeps his savings in his mattress. He began with \$1000 from his mother and adds \$100 each year. His total savings y after x years are given by the equation

- (a) $y = 1000 + 100x$. (b) $y = 100 + 1000x$.
(c) $y = 1000 + x$.

5.22 Smokers don't live as long (on the average) as nonsmokers, and heavy smokers don't live as long as light smokers. You regress the age at death of a group of male smokers on the number of packs per day they smoked. The slope of your regression line

- (a) will be greater than 0. (b) will be less than 0.
(c) Can't tell without seeing the data.

5.23 An owner of a home in the Midwest installed solar panels to reduce heating costs. After installing the solar panels, he measured the amount of natural gas used y (in cubic feet) to heat the home and outside temperature x (in degree-days, where a day's degree-days are the number of degrees its average temperature falls below 65° F) over a 23-month period. He then computed the least-squares regression line for predicting y from x and found it to be⁹

$$\hat{y} = 85 + 16x$$

How much, on average, does gas used increase for each additional degree-day?

- (a) 23 cubic feet (b) 85 cubic feet (c) 16 cubic feet

5.24 According to the regression line in Exercise 5.23, the predicted amount of gas used when the outside temperature is 20 degree-days is about

- (a) 405 cubic feet. (b) 320 cubic feet.
(c) 105 cubic feet.

5.25 By looking at the equation of the least-squares regression line in Exercise 5.23, you can see that the correlation between amount of gas used and degree-days is

- (a) greater than zero. (b) less than zero.
(c) Can't tell without seeing the data.

5.26 The software used to compute the least-squares regression line in Exercise 5.23 says that $r^2 = 0.98$. This suggests that

- (a) although degree-days and gas used are correlated, degree-days does not predict gas used very accurately.
(b) gas used increases by $\sqrt{0.98} = 0.99$ cubic feet for each additional degree-day.
(c) prediction of gas used from degree-days will be quite accurate.

5.27 Because elderly people may have difficulty standing to have their heights measured, a study looked at predicting overall height from height to the knee. Here are data (in centimeters) for six elderly men:



Knee height x	57.7	47.4	43.5	44.8	55.2	54.6
Height y	192.1	153.3	146.4	162.7	169.1	177.8

Use your calculator or software: what is the equation of the least-squares regression line for predicting overall height from knee height?

- (a) $\hat{y} = 42.9 + 2.5x$ (b) $\hat{y} = -3.4 + 0.3x$
(c) $\hat{y} = 2.5 + 42.9x$

CHAPTER 5 EXERCISES

5.28 Penguins diving. A study of king penguins looked for a relationship between how deep the penguins dive to seek food and how long they stay underwater.¹⁰ For all but the shallowest dives, there is a linear relationship that is different for different penguins. The study report gives a scatterplot for one penguin titled “The relation of dive duration (DD) to depth (D).” Duration DD is measured in minutes and depth D is in meters. The report then says, “The regression equation for this bird is: $DD = 2.69 + 0.0138D$.”



Paul A. Souders/CORBIS

(a) What is the slope of the regression line? Explain in specific language what this slope says about this penguin’s dives.

(b) According to the regression line, how long does a typical dive to a depth of 200 meters last?

(c) The dives varied from 40 meters to 300 meters in depth. Plot the regression line from $D = 40$ to $D = 300$.

5.29 The price of diamond rings. A newspaper advertisement in the Straits Times of Singapore contained pictures of diamond rings and listed their prices, diamond weight (in carats), and gold purity. Based on data for only the 20-carat gold ladies’ rings in the advertisement, the least-squares regression line for predicting price (in Singapore dollars) from the weight of the diamond (in carats) is¹¹

$$\text{price} = 259.63 + 3721.02 \text{ (carats)}$$

(a) What does the slope of this line say about the relationship between price and number of carats?

(b) What is the predicted price when number of carats = 0? How would you interpret this price?

5.30 Does social rejection hurt? Exercise 4.45 (page 121) gives data from a study that shows that social exclusion causes “real pain.” That is, activity in an area of the brain that responds to physical pain goes up as distress from social exclusion goes up. A scatterplot shows a moderately strong linear relationship. Figure 5.10 shows Minitab regression output for these data. 

(a) What is the equation of the least-squares regression line for predicting brain activity from social distress score? Use the equation to predict brain activity for a social distress score of 2.0.

(b) What percent of the variation in brain activity among these subjects is explained by the straight-line relationship with social distress score?

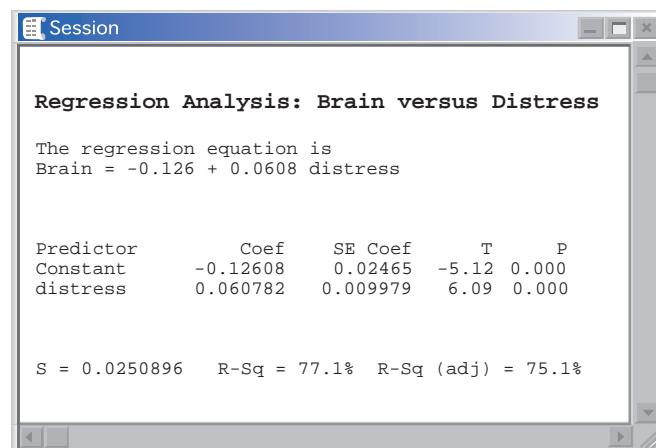


FIGURE 5.10

Minitab regression output for a study of the effects of social rejection on brain activity, for Exercise 5.30.

(c) Use the information in Figure 5.10 to find the correlation r between social distress score and brain activity. How do you know whether the sign of r is + or -?

5.31 Toucan’s beak. Exercise 4.44 (page 120) gives data on beak heat loss, as a percent of total body heat loss from all sources, at various temperatures. The data show that beak heat loss is higher at higher temperatures and that the relationship is roughly linear. Figure 5.11 shows Minitab regression output for these data. 

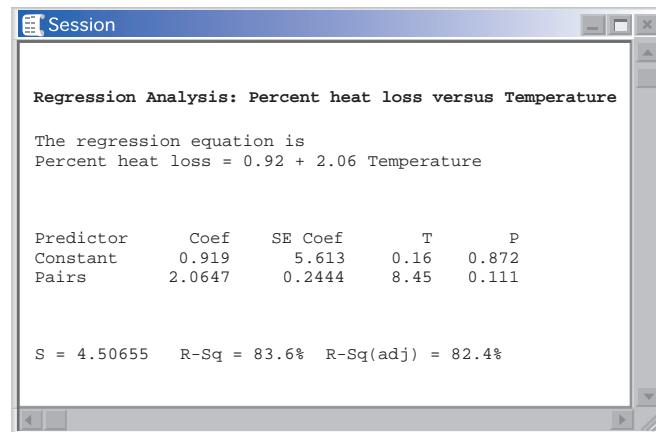


FIGURE 5.11

Minitab regression output for a study of how temperature affects beak heat loss in toucans, for Exercise 5.31.

- (a) What is the equation of the least-squares regression line for predicting beak heat loss, as a percent of total body heat loss from all sources, from temperature? Use the equation to predict beak heat loss, as a percent of total body heat loss from all sources, at a temperature of 25 degrees Celsius.
- (b) What percent of the variation in beak heat loss is explained by the straight-line relationship with temperature?
- (c) Use the information in Figure 5.11 to find the correlation r between beak heat loss and temperature. How do you know whether the sign of r is + or -?

5.32 Husbands and wives. The mean height of American women in their twenties is about 64.3 inches, and the standard deviation is about 3.9 inches. The mean height of men the same age is about 69.9 inches, with standard deviation about 3.1 inches. Suppose that the correlation between the heights of husbands and wives is about $r = 0.5$.

- (a) What are the slope and intercept of the regression line of the husband's height on the wife's height in young couples?
- (b) Draw a graph of this regression line for heights of wives between 56 and 72 inches. Predict the height of the husband of a woman who is 67 inches tall, and plot the wife's height and predicted husband's height on your graph.
- (c) You don't expect this prediction for a single couple to be very accurate. Why not?

5.33 What's my grade? In Professor Krugman's economics course the correlation between the students' total scores prior to the final examination and their final-examination scores is $r = 0.5$. The pre-exam totals for all students in the course have mean 280 and standard deviation 40. The final-exam scores have mean 75 and standard deviation 8. Professor Krugman has lost Julie's final exam but knows that her total before the exam was 300. He decides to predict her final-exam score from her pre-exam total.

(a) What is the slope of the least-squares regression line of final-exam scores on pre-exam total scores in this course? What is the intercept?

- (b) Use the regression line to predict Julie's final-exam score.
- (c) Julie doesn't think this method accurately predicts how well she did on the final exam. Use r^2 to argue that her actual score could have been much higher (or much lower) than the predicted value.

5.34 Going to class. A study of class attendance and grades among first-year students at a state university showed that, in general, students who attended a higher percent of their classes earned higher grades. Class attendance explained 16% of the variation in grade index among the students.

What is the numerical value of the correlation between percent of classes attended and grade index?

5.35 Sisters and brothers. How strongly do physical characteristics of sisters and brothers correlate? Here are data on the heights (in inches) of 12 adult pairs:¹²



Brother	71	68	66	67	70	71	70	73	72	65	66	70
Sister	69	64	65	63	65	62	65	64	66	59	62	64

(a) Use your calculator or software to find the correlation and the equation of the least-squares line for predicting sister's height from brother's height. Make a scatterplot of the data and add the regression line to your plot.

(b) Damien is 70 inches tall. Predict the height of his sister Tonya. Based on the scatterplot and the correlation r , do you expect your prediction to be very accurate? Why?

5.36 Keeping water clean. Keeping water supplies clean requires regular measurement of levels of pollutants. The measurements are indirect—a typical analysis involves forming a dye by a chemical reaction with the dissolved pollutant, then passing light through the solution and measuring its "absorbence." To calibrate such measurements, the laboratory measures known standard solutions and uses regression to relate absorbence and pollutant concentration. This is usually done every day. Here is one series of data on the absorbence for different levels of nitrates. Nitrates are measured in milligrams per liter of water.¹³



Nitrates	50	50	100	200	400	800	1200	1600	2000	2000
Absorbence	7.0	7.5	12.8	24.0	47.0	93.0	138.0	183.0	230.0	226.0

(a) Chemical theory says that these data should lie on a straight line. If the correlation is not at least 0.997, something went wrong and the calibration procedure is repeated. Plot the data and find the correlation. Must the calibration be done again?

(b) The calibration process sets nitrate level and measures absorbence. The linear relationship that results is used to estimate the nitrate level in water from a measurement of absorbence. What is the equation of the line used to estimate nitrate level? What does the slope of this line say about the relationship between nitrate level and absorbence? What is the estimated nitrate level in a water specimen with absorbence 40?

(c) Do you expect estimates of nitrate level from absorbence to be quite accurate? Why?

5.37 Sparrowhawk colonies. One of nature's patterns connects the percent of adult birds in a colony that return from the previous year and the number of new adults that join the colony. Here are data for 13 colonies of sparrowhawks:¹⁴

Percent return x	74	66	81	52	73	62	52	45	62	46	60	46	38
New adults y	5	6	8	11	12	15	16	17	18	18	19	20	20

You saw in Exercise 4.28 that there is a moderately strong linear relationship, with correlation $r = -0.748$. 

(a) Find the least-squares regression line for predicting y from x . Make a scatterplot and draw your line on the plot.

(b) Explain in words what the slope of the regression line tells us.

(c) An ecologist uses the line, based on 13 colonies, to predict how many new birds will join another colony, to which 60% of the adults from the previous year return. What is the prediction?

5.38 Our brains don't like losses. Exercise 4.29 (page 117) describes an experiment that showed a linear relationship between how sensitive people are to monetary losses ("behavioral loss aversion") and activity in one part of their brains ("neural loss aversion"). 

(a) Make a scatterplot with neural loss aversion as x and behavioral loss aversion as y . One point is a high outlier in both the x and y directions.

(b) Find the least-squares line for predicting y from x , leaving out the outlier, and add the line to your plot.

(c) The outlier lies very close to your regression line. Looking at the plot, you now expect that adding the outlier will increase the correlation but will have little effect on the least-squares line. Explain why.

tions from (c).

5.39 Always plot your data! Table 5.1 presents four sets of data prepared by the statistician Frank Anscombe to illustrate the dangers of calculating without first plotting the data.¹⁵ 

(a) Without making scatterplots, find the correlation and the least-squares regression line for all four data sets. What do you notice? Use the regression line to predict y for $x = 10$. 

(b) Make a scatterplot for each of the data sets and add the regression line to each plot. 

(c) In which of the four cases would you be willing to use the regression line to describe the dependence of y on x ? Explain your answer in each case. 

5.40 Managing diabetes. People with diabetes must man-



Glow Wellness/Alamy

TABLE 5.1 Four data sets for exploring correlation and regression

DATA SET A												
x	10	8	13	9	11	14	6	4	12	7	5	
y	8.04	6.95	7.58	8.81	8.33	9.96	7.24	4.26	10.84	4.82	5.68	
DATA SET B												
x	10	8	13	9	11	14	6	4	12	7	5	
y	9.14	8.14	8.74	8.77	9.26	8.10	6.13	3.10	9.13	7.26	4.74	
DATA SET C												
x	10	8	13	9	11	14	6	4	12	7	5	
y	7.46	6.77	12.74	7.11	7.81	8.84	6.08	5.39	8.15	6.42	5.73	
DATA SET D												
x	8	8	8	8	8	8	8	8	8	8	19	
y	6.58	5.76	7.71	8.84	8.47	7.04	5.25	5.56	7.91	6.89	12.50	

TABLE 5.2 Two measures of glucose level in diabetics

SUBJECT	HbA (%)	FPG (mg/ml)	SUBJECT	HbA (%)	FPG (mg/ml)	SUBJECT	HbA (%)	FPG (mg/ml)
1	6.1	141	7	7.5	96	13	10.6	103
2	6.3	158	8	7.7	78	14	10.7	172
3	6.4	112	9	7.9	148	15	10.7	359
4	6.8	153	10	8.7	172	16	11.2	145
5	7.0	134	11	9.4	200	17	13.7	147
6	7.1	95	12	10.4	271	18	19.3	255

checkups, is called HbA. This is roughly the percent of red blood cells that have a glucose molecule attached. It measures average exposure to glucose over a period of several months. Table 5.2 gives data on both HbA and FPG for 18 diabetics five months after they had completed a diabetes education class.¹⁶ 

(a) Make a scatterplot with HbA as the explanatory variable. There is a positive linear relationship, but it is surprisingly weak.

(b) Subject 15 is an outlier in the y direction. Subject 18 is an outlier in the x direction. Find the correlation for all 18 subjects, for all except Subject 15, and for all except Subject 18. Are either or both of these subjects influential for the correlation? Explain in simple language why r changes in opposite directions when we remove each of these points.

5.41 The effect of changing units. The equation of a regression line, unlike the correlation, depends on the units we use to measure the explanatory and response variables. Here are data on knee height and overall height (in centimeters) for six elderly men: 

Knee height x	57.7	47.4	43.5	44.8	55.2	54.6
Height y	192.1	153.3	146.4	162.7	169.1	177.8

- (a) Find the equation of the regression line for predicting overall height in centimeters from knee height in centimeters.
 (b) A mad scientist decides to measure knee height in millimeters and height in meters. The same data in these units are

Knee height x	577	474	435	448	552	546
Height y	1.921	1.533	1.464	1.627	1.691	1.778

Find the equation of the regression line for predicting overall height in meters from knee height in millimeters.

(c) Use both lines to predict the overall height of a man whose knee height is 50 centimeters, which is the same as 500 millimeters. Use the fact that there are 100 centimeters in a meter to show that the two predictions are the same (up to roundoff error).

5.42 Managing diabetes, continued. Add three regression lines for predicting FPG from HbA to your scatterplot from Exercise 5.40: for all 18 subjects, for all except Subject 15, and for all except Subject 18. Is either Subject 15 or Subject 18 strongly influential for the least-squares line? Explain in simple language what features of the scatterplot explain the degree of influence. 

5.43 Are you happy? Exercise 4.25 (page 115) discusses a study in which the mean BRFSS life-satisfaction score of individuals in each state was compared with the mean of an objective measure of well-being (based on the “compensating-differentials method”) for each state. Suppose that instead of the means for the states, the BRFSS life-satisfaction scores for individuals were compared with the corresponding measure of well-being (based on the compensating-differentials method) for these individuals. Would you expect the correlation between the mean state scores on these two measures to be lower, about the same, or higher than the correlation between the scores of individuals on these two measures? Explain your answer.

5.44 Do artificial sweeteners cause weight gain? People who use artificial sweeteners in place of sugar tend to be heavier than people who use sugar. Does this mean that artificial sweeteners cause weight gain? Give a more plausible explanation for this association.

5.45 Learning online. Many colleges offer online versions of courses that are also taught in the classroom. It often happens that the students who enroll in the online version do better than the classroom students on the course exams. This

does not show that online instruction is more effective than classroom teaching, because the people who sign up for online courses are often quite different from the classroom students. Suggest some differences between online and classroom students that might explain why online students do better.

5.46 Grade inflation and the SAT. The effect of a lurking variable can be surprising when individuals are divided into groups. In recent years, the mean SAT score of all high school seniors has increased. But the mean SAT score has decreased for students at each level of high school grades (A, B, C, and so on). Explain how grade inflation in high school (the lurking variable) can account for this pattern.

5.47 Workers' incomes. Here is another example of the group effect cautioned about in the previous exercise. Explain how, as a nation's population grows older, median income can go down for workers in each age group, yet still go up for all workers.

5.48 Some regression math. Use the equation of the least-squares regression line (box on page 130) to show that the regression line for predicting y from x always passes through the point (\bar{x}, \bar{y}) . That is, when $x = \bar{x}$, the equation gives $\hat{y} = \bar{y}$.

5.49 Regression to the mean. Figure 4.8 (page 115) displays the relationship between golfers' scores on the first and second rounds of the 2010 Masters Tournament. The least-squares line for predicting second-round scores from first-round scores has equation $\hat{y} = 52.74 + 0.297x$. Find the predicted second-round scores for a player who shot 80 in the first round and for a player who shot 70. The mean second-round score for all players was 74.48. So a player who does well in the first round is predicted to do less well, but still better than average, in the second round. And a player who does poorly in the first is predicted to do better, but still worse than average, in the second.

(Comment: This is **regression to the mean**. If you select individuals with extreme scores on some measure, they tend to have less extreme scores when measured again. That's because their extreme position is partly merit and partly luck. The luck will be different next time. Regression to the mean contributes to lots of "effects." The rookie of the year often doesn't do as well the next year; the best player in an orchestral audition may play less well once hired than the runners-up; a student who feels she needs coaching after taking the SAT often does better on the next try without coaching.)

5.50 Regression to the mean. We expect that students who do well on the midterm exam in a course will usually also do well on the final exam. Gary Smith of Pomona College looked at the exam scores of all 346 students who took his statistics class over a 10-year period.¹⁷ The least-squares line

for predicting final-exam score from midterm-exam score was $\hat{y} = 46.6 + 0.41x$. (Both exams have a 100-point scale.)

Octavio scores 10 points above the class mean on the midterm. How many points above the class mean do you predict that he will score on the final? (*Hint:* Use the fact that the least-squares line passes through the point (\bar{x}, \bar{y}) and the fact that Octavio's midterm score is $\bar{x} + 10$.) This is another example of regression to the mean: students who do well on the midterm will, in general, do less well, but still above average, on the final.

5.51 Is regression useful? In Exercise 4.41 (page 120)

 you used the *Correlation and Regression* applet to create three scatterplots having correlation about $r = 0.7$ between the horizontal variable x and the vertical variable y . Create three similar scatterplots again, and click the "Show least-squares line" box to display the regression lines. Correlation $r = 0.7$ is considered reasonably strong in many areas of work. Because there is a reasonably strong correlation, we might use a regression line to predict y from x . In which of your three scatterplots does it make sense to use a straight line for prediction?

5.52 Guessing a regression line. In the *Correlation and*

 *Regression* applet, click on the scatterplot to create a group of 15 to 20 points from lower left to upper right with a clear positive straight-line pattern (correlation around 0.7). Click the "Draw line" button and use the mouse (right-click and drag) to draw a line through the middle of the cloud of points from lower left to upper right. Note the "thermometer" above the plot. The red portion is the sum of the squared vertical distances from the points in the plot to the least-squares line. The green portion is the "extra" sum of squares for your line—it shows by how much your line misses the smallest possible sum of squares.

(a) You drew a line by eye through the middle of the pattern. Yet the right-hand part of the bar is probably almost entirely green. What does that tell you?

(b) Now click the "Show least-squares line" box. Is the slope of the least-squares line smaller (the new line is less steep) or larger (line is steeper) than that of your line? If you repeat this exercise several times, you will consistently get the same result. The least-squares line minimizes the *vertical* distances of the points from the line. It is *not* the line through the "middle" of the cloud of points. This is one reason why it is hard to draw a good regression line by eye.

The following exercises ask you to answer questions from data without having the details outlined for you. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan**, **Solve**, and **Conclude** steps of the process, described on page 55.

TABLE 5.3 Reaction times (in milliseconds) in a computer game

TIME	DISTANCE	HAND	TIME	DISTANCE	HAND
115	190.70	right	240	190.70	left
96	138.52	right	190	138.52	left
110	165.08	right	170	165.08	left
100	126.19	right	125	126.19	left
111	163.19	right	315	163.19	left
101	305.66	right	240	305.66	left
111	176.15	right	141	176.15	left
106	162.78	right	210	162.78	left
96	147.87	right	200	147.87	left
96	271.46	right	401	271.46	left
95	40.25	right	320	40.25	left
96	24.76	right	113	24.76	left
96	104.80	right	176	104.80	left
106	136.80	right	211	136.80	left
100	308.60	right	238	308.60	left
113	279.80	right	316	279.80	left
123	125.51	right	176	125.51	left
111	329.80	right	173	329.80	left
95	51.66	right	210	51.66	left
108	201.95	right	170	201.95	left

5.53 Beavers and beetles. Do beavers benefit beetles?

 Researchers laid out 23 circular plots, each 4 meters in diameter, in an area where beavers were cutting down cottonwood trees. In each plot, they counted the number of stumps from trees cut by beavers and the number of clusters of beetle larvae. Ecologists think that the new sprouts from stumps are more tender than other cottonwood growth, so that beetles prefer them. If so, more stumps should produce more beetle larvae. Here are the data:¹⁸

Stumps	2	2	1	3	3	4	3	1	2	5	1	3
Beetle larvae	10	30	12	24	36	40	43	11	27	56	18	40
Stumps	2	1	2	2	1	1	4	1	2	1	4	
Beetle larvae	25	8	21	14	16	6	54	9	13	14	50	

Analyze these data to see if they support the “beavers benefit beetles” idea.  BEAVERS

5.54 A computer game. A multimedia statistics learning system includes a test of skill in using the computer’s mouse. The software displays a circle at a random location on the computer screen. The subject clicks in the circle

with the mouse as quickly as possible. A new circle appears as soon as the subject clicks the old one. Table 5.3 gives data for one subject’s trials, 20 with each hand. Distance is the distance from the cursor location to the center of the new circle, in units whose actual size depends on the size of the screen. Time is the time required to click in the new circle, in milliseconds.¹⁹ We suspect that time depends on distance. We also suspect that performance will not be the same with the right and left hands. Analyze the data with a view to predicting performance separately for the two hands.  COMPUTERGAME

5.55 Predicting tropical storms. William Gray heads the Tropical Meteorology Project at Colorado State University (well away from the hurricane belt). His forecasts before each year’s hurricane season attract lots of attention. Here are data on the



NASA/GSFC

TABLE 5.4 Arctic river discharge (cubic kilometers), 1936 to 2008

YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE
1936	1721	1955	1656	1974	2000	1993	1845
1937	1713	1956	1721	1975	1928	1994	1902
1938	1860	1957	1762	1976	1653	1995	1842
1939	1739	1958	1936	1977	1698	1996	1849
1940	1615	1959	1906	1978	2008	1997	2007
1941	1838	1960	1736	1979	1970	1998	1903
1942	1762	1961	1970	1980	1758	1999	1970
1943	1709	1962	1849	1981	1774	2000	1905
1944	1921	1963	1774	1982	1728	2001	1890
1945	1581	1964	1606	1983	1920	2002	2085
1946	1834	1965	1735	1984	1823	2003	1780
1947	1890	1966	1883	1985	1822	2004	1900
1948	1898	1967	1642	1986	1860	2005	1930
1949	1958	1968	1713	1987	1732	2006	1910
1950	1830	1969	1742	1988	1906	2007	2270
1951	1864	1970	1751	1989	1932	2008	2078
1952	1829	1971	1879	1990	1861		
1953	1652	1972	1736	1991	1801		
1954	1589	1973	1861	1992	1793		

number of named Atlantic tropical storms predicted by Dr. Gray and the actual number of storms for the years 1984 to 2010.²⁰



Analyze these data. How accurate are Dr. Gray's forecasts? How many tropical storms would you expect in a year when his preseason forecast calls for 16 storms? What is the effect of the disastrous 2005 season on your answers?

Year	Forecast	Actual	Year	Forecast	Actual
1984	10	12	1998	10	14
1985	11	11	1999	14	12
1986	8	6	2000	12	14
1987	8	7	2001	12	15
1988	11	12	2002	11	12
1989	7	11	2003	14	16
1990	11	14	2004	14	14
1991	8	8	2005	15	27
1992	8	6	2006	17	10
1993	11	8	2007	17	14
1994	9	7	2008	15	16
1995	12	19	2009	11	9
1996	10	13	2010	18	19
1997	11	7			

5.56 Great Arctic rivers. One effect of global warming is to increase the flow of water into the Arctic Ocean from rivers. Such an increase may have major effects on the world's climate. Six rivers (Yenisey, Lena, Ob, Pechora, Kolyma, and Severnaya Dvina) drain two-thirds of the Arctic in Europe and Asia. Several of these are among the largest rivers on earth. Table 5.4 presents the total discharge from these rivers each year from 1936 to 2008.²¹ Discharge is measured in cubic kilometers of water. Analyze these data to uncover the nature and strength of the trend in total discharge over time.



5.57 Will women outrun men? Does the physiology of women make them better suited than men to long-distance running? Will women eventually outperform men in long-distance races? Researchers examined data on world record times (in seconds) for men and women in the marathon. Based on these data, researchers (in 1992) attempted to predict when women would outrun men in the marathon. Here are data for women:²²



Year	1926	1964	1967	1970	1971	1974	1975
Time	13,222.0	11,973.0	11,246.0	10,973.0	9990.0	9834.5	9499.0
Year	1977	1980	1981	1982	1983	1985	
Time	9287.5	9027.0	8806.0	8771.0	8563.0	8466.0	

Here are data for men:

Year	1908	1909	1913	1920	1925	1935	1947
Time	10,518.4	9751.0	9366.6	9155.8	8941.8	8802.0	8739.0
Year	1952	1953	1954	1958	1960	1963	1964
Time	8442.2	8314.8	8259.4	8117.0	8116.2	8068.0	7931.2
Year	1965	1967	1969	1981	1984	1985	1988
Time	7920.0	7776.4	7713.6	7698.0	7685.0	7632.0	7610.0

Analyze these data using least-squares regression to estimate when men and women's record times will be equal. How reliable is your estimate?



EXPLORING THE WEB

5.58 Association and causation. Find an example of a study in which the issue of association and causation is present. This can be either an example in which association is confused with causation or an example in which the association is not confused with causation. Summarize the study and its conclusions in your own words. Be sure to include either a copy of the actual article or at least the Web source, title, and where the article was published. The *Chance News* Web site at www.causeweb.org/wiki/chance/index.php/Main_Page is a good place to look for examples.

5.59 Predicting batting averages. Go to www.mlb.com/ and find the batting averages for a diverse set of 30 players for both the 2009 and 2010 seasons. You can click on the “Stats” tab to find the results for the current season as well as historical data. You should select only players who played in at least 50 games both seasons. Make a scatterplot of the batting averages using the 2009 season average as the explanatory variable and the 2010 season average as the response. Is it reasonable to fit a straight line to these data? If so, find the least-squares regression line for predicting batting average in 2010 from that in 2009 based on your sample of 30 players. In 2009, the major league leader in batting was Joe Mauer, who had a batting average of .365. What does your least-squares regression line predict for the 2010 batting average of someone who hit .365 in 2009? Is the 2010 predicted batting average higher or lower than .365?

5.60 Predicting the federal budget. Go to the Congressional Budget Office Web site, www.cbo.gov. Under the tab “Publications” (hover over “By Subject”) select “Budget and Economic Information” and then “Budget and Economic Outlook.” What is the current prediction for the federal budget in five years’ time? Is a surplus or a deficit predicted? Do you think this prediction is accurate? Why or why not?



Two-Way Tables*

Chapter 6

We have concentrated on relationships in which at least the response variable is quantitative. Now we will describe relationships between two categorical variables. Some variables—such as sex, race, and occupation—are categorical by nature. Other categorical variables are created by grouping values of a quantitative variable into classes. Published data often appear in grouped form to save space. To analyze categorical data, we use the counts or percents of individuals that fall into various categories.

EXAMPLE 6.1 I think I'll be rich by age 30

A sample survey of young adults (aged 19 to 25) asked, “What do you think are the chances you will have much more than a middle-class income at age 30?” Table 6.1 shows the responses, omitting a few people who refused to respond or who said they were already rich.¹ This is a **two-way table** because it describes two categorical variables: sex and opinion about becoming rich. Opinion is the **row variable** because each row in the table describes young adults who held one of the five opinions about their chances. Because the opinions have a natural order from “Almost no chance” to “Almost certain,” the rows are also in this order. Sex is the **column variable** because each

*This material is important in statistics, but it is needed later in this book only for Chapter 23. You may omit it if you do not plan to read Chapter 23 or delay reading it until you reach Chapter 23.

IN THIS CHAPTER WE COVER...

- Marginal distributions
- Conditional distributions
- Simpson’s paradox

two-way table

row and column variables

TABLE 6.1 Young adults by sex and chance of getting rich

OPINION	SEX		TOTAL
	FEMALE	MALE	
Almost no chance	96	98	194
Some chance but probably not	426	286	712
A 50-50 chance	696	720	1416
A good chance	663	758	1421
Almost certain	486	597	1083
Total	2367	2459	4826

column describes one sex. The entries in the table are the counts of individuals in each opinion-by-sex class. ■

MARGINAL DISTRIBUTIONS

How can we best grasp the information contained in Table 6.1? First, look at the distribution of each variable separately. The distribution of a categorical variable says how often each outcome occurred. The “Total” column at the right of the table contains the totals for each of the rows. These row totals give the distribution of opinions about becoming rich in the entire group of 4826 young adults: 194 felt that they had almost no chance, 712 thought they had just some chance, and so on.

marginal distribution

If the row and column totals are missing, the first thing to do in studying a two-way table is to calculate them. The distributions of opinion alone and sex alone are called **marginal distributions** because they appear at the right and bottom margins of the two-way table.

Percents are often more informative than counts. We can display the marginal distribution of opinions in percents by dividing each row total by the table total and converting to a percent.

EXAMPLE 6.2 Calculating a marginal distribution

The percent of these young adults who think they are almost certain to be rich by age 30 is

$$\frac{\text{almost certain total}}{\text{table total}} = \frac{1083}{4826} = 0.224 = 22.4\%$$

Do four more such calculations to obtain the marginal distribution of opinion in percents. Here is the complete distribution:

Response	Percent
Almost no chance	$\frac{194}{4826} = 4.0\%$
Some chance	$\frac{712}{4826} = 14.8\%$
A 50-50 chance	$\frac{1416}{4826} = 29.3\%$
A good chance	$\frac{1421}{4826} = 29.4\%$
Almost certain	$\frac{1083}{4826} = 22.4\%$

It seems that many young adults are optimistic about their future income. The total should be 100% because everyone holds one of the five opinions. In fact, the percents add to 99.9% because we rounded each one to the nearest tenth. This is **roundoff error**. ■

roundoff error

Each marginal distribution from a two-way table is a distribution for a single categorical variable. As we saw in Chapter 1, we can use a bar graph or a pie chart to display such a distribution. Figure 6.1 is a bar graph of the distribution of opinion among young adults.

In working with two-way tables, you must calculate lots of percents. Here's a tip to help you decide what fraction gives the percent you want. Ask, "What group represents the total of which I want a percent?" The count for that group is the denominator of the fraction that leads to the percent. In Example 6.2, we want a percent "of young adults," so the count of young adults (the table total) is the denominator.

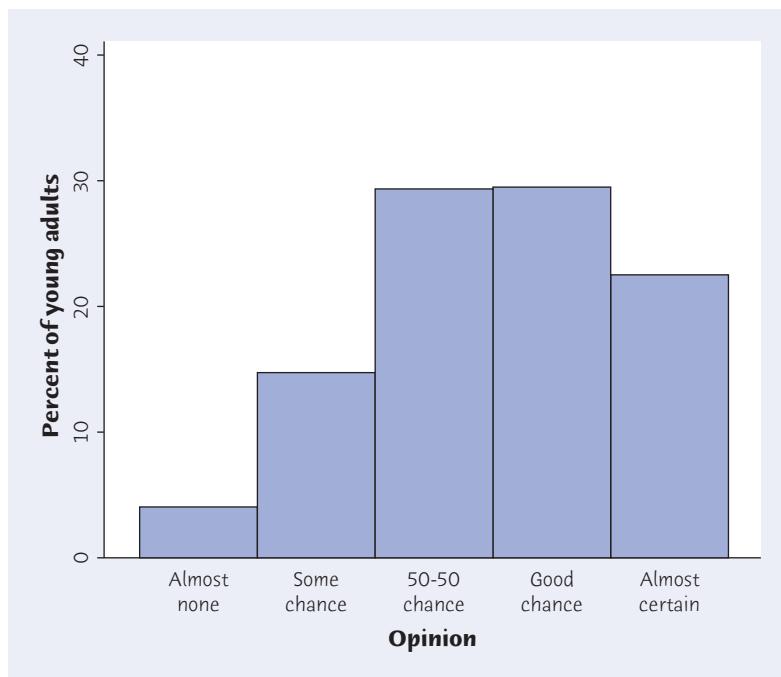


FIGURE 6.1

Bar graph of the distribution of opinions of young adults about becoming rich by age 30. This is one of the marginal distributions for Table 6.1



APPLY YOUR KNOWLEDGE

6.1 Video-gaming and grades. The popularity of computer, video, online, and virtual reality games has raised concerns about their ability to negatively impact youth. The data in this exercise are based on a recent survey of 14- to 18-year-olds in Connecticut high schools. Here are the grade distributions of boys who have and have not played video games.²  GAMING

	Grade Average		
	A's and B's	C's	D's and F's
Played games	736	450	193
Never played games	205	144	80

- (a) How many people does this table describe? How many of these have played video games?
- (b) Give the marginal distribution of the grades. What percent of the boys represented in the table received a grade of C or lower?

6.2 Undergraduates' ages. Here is a two-way table of U.S. Census Bureau data describing the age and sex of all American undergraduate college students. The table entries are counts in thousands of students.³  UNDERGRADAGES



Laif/Redux

Age group	Female	Male
15 to 19 years	2124	1876
20 to 24 years	2814	2648
25 to 34 years	703	536
35 years or older	518	159

- (a) How many college undergraduates are there?
- (b) Find the marginal distribution of age group. What percent of undergraduates are in the 20 to 24 college age group?

CONDITIONAL DISTRIBUTIONS

Table 6.1 contains much more information than the two marginal distributions of opinion alone and sex alone. *Marginal distributions tell us nothing about the relationship between two variables.* To describe a relationship between two categorical variables, we must calculate some well-chosen percents from the counts given in the body of the table. 

Let's say that we want to compare the opinions of women and men. To do this, compare percents for women alone with percents for men alone. To study the opinions of women, we look only at the "Female" column in Table 6.1. To find the percent of young women who think they are almost certain to be rich

by age 30, divide the count of such women by the total number of women (the column total):

$$\frac{\text{women who are almost certain}}{\text{column total}} = \frac{486}{2367} = 0.205 = 20.5\%$$

Doing this for all five entries in the “Female” column gives the *conditional distribution* of opinion among women. We use the term “conditional” because this distribution describes only young adults who satisfy the condition that they are female.

MARGINAL AND CONDITIONAL DISTRIBUTIONS

The **marginal distribution** of one of the categorical variables in a two-way table of counts is the distribution of values of that variable among all individuals described by the table.

A **conditional distribution** of a variable is the distribution of values of that variable among only individuals who have a given value of the other variable. There is a separate conditional distribution for each value of the other variable.



Smiling faces

Women smile

more than men.

The same data
that produce this

fact allow us to link smiling to other variables in two-way tables. For example, as the second variable add whether or not the person thinks they are being observed. If yes, that's when women smile more. If no, there's no difference between women and men. Next, take the second variable to be the person's social role (for example, is he or she the boss in an office?). Within each role, there is very little difference in smiling between women and men.

EXAMPLE 6.3 Comparing women and men

STATE: How do young men and young women differ in their responses to the question “What do you think are the chances you will have much more than a middle-class income at age 30?”

PLAN: Make a two-way table of response by sex. Find the two conditional distributions of response for men alone and for women alone. Compare these two distributions.

SOLVE: Table 6.1 is the two-way table we need. Look first at just the “Female” column to find the conditional distribution for women, then at just the “Male” column to find the conditional distribution for men. Here are the calculations and the two conditional distributions:

Response	Female	Male
Almost no chance	$\frac{96}{2367} = 4.1\%$	$\frac{98}{2459} = 4.0\%$
Some chance	$\frac{426}{2367} = 18.0\%$	$\frac{286}{2459} = 11.6\%$
A 50-50 chance	$\frac{696}{2367} = 29.4\%$	$\frac{720}{2459} = 29.3\%$
A good chance	$\frac{663}{2367} = 28.0\%$	$\frac{758}{2459} = 30.8\%$
Almost certain	$\frac{486}{2367} = 20.5\%$	$\frac{597}{2459} = 24.3\%$



Each set of percents adds to 100% because everyone holds one of the five opinions.

CONCLUDE: Men are somewhat more optimistic about their future income than are women. Men are less likely to say that they have “some chance but probably not” and more likely to say that they have “a good chance” or are “almost certain” to have much more than a middle-class income by age 30. ■

Software will do these calculations for you. Most programs allow you to choose which conditional distributions you want to compare. The output in Figure 6.2

FIGURE 6.2

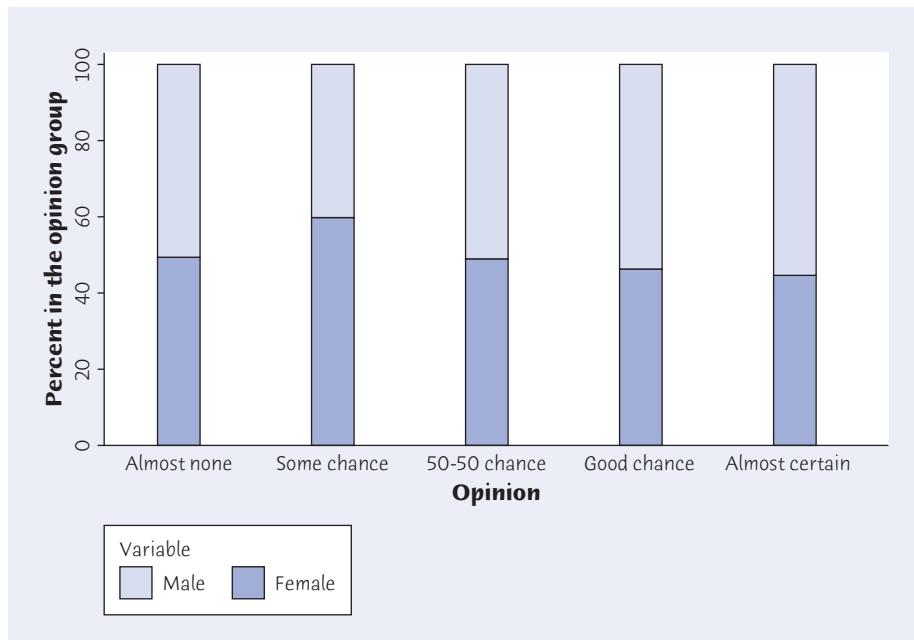
Minitab and CrunchIt! output for the two-way table of young adults by sex and chance of getting rich. Each entry in the Minitab output includes the percent of its column total. The “Female” and “Male” columns give the conditional distributions of responses for women and men, and the “All” column shows the marginal distribution of responses for all these young adults. Each entry in the CrunchIt! output includes the percent of its row total, the percent of its column total, and the percent of the entire table total. The second entry in each cell gives the conditional distribution of responses for the different opinions. The third entry in each cell for the “Female” and “Male” columns gives the conditional distributions of responses for women and men. The “All” row and column show the corresponding marginal distribution of responses for all these young adults.

Minitab

	Female	Male	All
A: Almost no chance	96	98	194
	4.06	3.99	4.02
B: Some chance but probably not	426	286	712
	18.00	11.63	14.75
C: A 50-50 chance	696	720	1416
	29.40	29.28	29.34
D: A good chance	663	758	1421
	28.01	30.83	29.44
E: Almost certain	486	597	1083
	20.53	24.28	22.44
All	2367	2459	4826
	100.00	100.00	100.00
Cell Contents: Count % of Column			

CrunchIt!

	Female	Male	All
A: Almost no chance	96	98	194
	49.48	50.52	100
	4.056	3.985	4.020
	1.989	2.031	4.020
B: Some chance	426	286	712
	59.83	40.17	100
	18.00	11.63	14.75
	8.827	5.926	14.75
C: A 50-50 chance	696	720	1416
	49.15	50.85	100
	29.40	29.28	29.34
	14.42	14.92	29.34
D: A good chance	663	758	1421
	46.66	53.34	100
	28.01	30.83	29.44
	13.74	15.71	29.44
E: Almost certain	486	597	1083
	44.88	55.12	100
	20.53	24.28	22.44
	10.07	12.37	22.44
All	2367	2459	4826
	49.05	50.95	100
	100	100	100
	49.05	50.95	100
Count % of Row % of Col % of Total			

**FIGURE 6.3**

Bar graph comparing the percents of females (darker shading) and the percents of males (lighter shading) among those who hold each opinion about their chances of getting rich by age 30.

presents the two conditional distributions of opinion, for women and for men, and also the marginal distribution of opinion for all the young adults. The distributions agree (up to roundoff) with the results in Examples 6.2 and 6.3.

Remember that there are two sets of conditional distributions for any two-way table. Example 6.3 looked at the conditional distributions of opinion for the two sexes. We could also examine the five conditional distributions of sex, one for each of the five opinions, by looking separately at the five rows in Table 6.1. Figure 6.3 makes this comparison in a bar graph. Each bar is divided (segmented) into two parts, represented by two colors. The lower portion of each bar represents the percent of women among young adults who hold each opinion. The upper portion represents the percent of men. Each bar has a height of 100%, because each bar represents all the young adults in each different group of people. Bar graphs like that in Figure 6.3 in which each bar is divided into parts, each part representing a different category, are sometimes called segmented bar graphs.

No single graph (such as a scatterplot) portrays the form of the relationship between categorical variables. No single numerical measure (such as the correlation) summarizes the strength of the association. Bar graphs are flexible enough to be helpful, but you must think about what comparisons you want to display. For numerical measures, we rely on well-chosen percents. You must decide which percents you need. Here is a hint: *if there is an explanatory-response relationship, compare the conditional distributions of the response variable for the separate values of the explanatory variable.* If you think that sex influences young adults' opinions about their chances of getting rich by age 30, compare the conditional distributions of opinion for women and for men, as in Example 6.3.





Keith Bedford/The New York Times/Redux



APPLY YOUR KNOWLEDGE

6.3 Video-gaming and grades. Exercise 6.1 (page 162) gives data on the grade distribution of boys who have and have not played video games. To see the relationship between grades and game-playing experience, find the conditional distributions of grades (the response variable) for players and nonplayers. What do you conclude? GAMING

6.4 Undergraduates' ages. Exercise 6.2 (page 162) gives U.S. Census Bureau data describing the age and sex of all American college undergraduates. We suspect that the percent of women is higher among students in the 25- to 34-year age group than in the 20- to 24-year age group. Do the data support this suspicion? Follow the four-step process as illustrated in Example 6.3. UNDERGRADUATES

6.5 Marginal distributions aren't the whole story. Here are the row and column totals for a two-way table with two rows and two columns:

a	b	50
c	d	50
60	40	100

Make up *two different* sets of counts a , b , c , and d for the body of the table that give these same totals. This shows that the relationship between two variables cannot be obtained from the two individual distributions of the variables.

SIMPSON'S PARADOX

As is the case with quantitative variables, the effects of lurking variables can change or even reverse relationships between two categorical variables. Here is an example that demonstrates the surprises that can await the unsuspecting user of data.



Ashley/Cooper/PICIMPACT/CORBIS

EXAMPLE 6.4 Do medical helicopters save lives?

Accident victims are sometimes taken by helicopter from the accident scene to a hospital. Helicopters save time. Do they also save lives? Let's compare the percents of accident victims who die with helicopter evacuation and with the usual transport to a hospital by road. Here are hypothetical data that illustrate a practical difficulty:⁴

	Helicopter	Road
Victim died	64	260
Victim survived	136	840
Total	200	1100

We see that 32% (64 out of 200) of helicopter patients died, but only 24% (260 out of 1100) of the others did. That seems discouraging.

The explanation is that the helicopter is sent mostly to serious accidents, so that the victims transported by helicopter are more often seriously injured. They are more likely to die with or without helicopter evacuation. Here are the same data broken down by the seriousness of the accident:

Serious Accidents			Less Serious Accidents		
	Helicopter	Road		Helicopter	Road
Died	48	60	Died	16	200
Survived	52	40	Survived	84	800
Total	100	100	Total	100	1000

Inspect these tables to convince yourself that they describe the same 1300 accident victims as the original two-way table. For example, 200 ($100 + 100$) were moved by helicopter, and 64 ($48 + 16$) of these died.

Among victims of serious accidents, the helicopter saves 52% (52 out of 100) compared with 40% for road transport. If we look only at less serious accidents, 84% of those transported by helicopter survive, versus 80% of those transported by road. Both groups of victims have a higher survival rate when evacuated by helicopter. ■

How can it happen that the helicopter does better for both groups of victims but worse when all victims are lumped together? Examining the data makes the explanation clear. Half the helicopter transport patients are from serious accidents, compared with only 100 of the 1100 road transport patients. So the helicopter carries patients who are more likely to die. The seriousness of the accident was a lurking variable that, until we uncovered it, hid the true relationship between survival and mode of transport to a hospital. Example 6.4 illustrates Simpson's paradox.

SIMPSON'S PARADOX

An association or comparison that holds for all of several groups can reverse direction when the data are combined to form a single group. This reversal is called Simpson's paradox.

The lurking variable in Simpson's paradox is categorical. That is, it breaks the individuals into groups, as when accident victims are classified as injured in a "serious accident" or a "less serious accident." Simpson's paradox is just an extreme form of the fact that observed associations can be misleading when there are lurking variables.

APPLY YOUR KNOWLEDGE

6.6 Field goal shooting. Here are data on field goal shooting for two members of the Kent State University 2002–2003 women's basketball team:⁵  BASKETBALL

	Jamie Rubis		Lindsay Shearer	
	Made	Missed	Made	Missed
Two-pointers	119	115	86	84
Three-pointers	36	61	5	16

- What percent of all field goal attempts did Jamie Rubis make? What percent of all field goal attempts did Lindsay Shearer make?
- Now find the percent of all two-point field goals and all three-point field goals that Jamie made. Do the same for Lindsay.
- Lindsay had a lower percent for *both* types of field goals but had a better overall percent. That sounds impossible. Explain carefully, referring to the data, how this can happen.

6.7 Bias in the jury pool? The New Zealand Department of Justice did a study of the composition of juries in court cases. Of interest was whether Maori, the indigenous people of New Zealand, were adequately represented in jury pools. Here are the results for two districts, Rotura and Nelson, in New Zealand (similar results were found in all districts):⁶  JURYPOOL

Rotura			Nelson		
	Maori	Non-Maori		Maori	Non-Maori
In jury pool	79	258	In jury pool	1	56
Not in jury pool	8810	23,751	Not in jury pool	1328	32,602
Total	8889	24,009	Total	1329	32,658

- Compare percents to show that the percent of all Maori in the jury pool in each district is less than the percent of non-Maori in the jury pool.
- Combine the data into a single two-way table of outcome (“in jury pool” or “not in jury pool”) by ethnicity (Maori or non-Maori). The original study only reported such an overall rate. Which ethnic group has a higher percent of its people in the jury pool?
- Explain from the data, in language that a reporter can understand, how Maori can have a higher percent overall even though non-Maori have higher percents for both districts.

CHAPTER 6 SUMMARY

CHAPTER SPECIFICS

- A **two-way table** of counts organizes data about two categorical variables. Values of the **row variable** label the rows that run across the table, and values of the **column variable** label the columns that run down the table. Two-way tables are often used to summarize large amounts of information by grouping outcomes into categories.
- The **row totals** and **column totals** in a two-way table give the **marginal distributions** of the two individual variables. It is clearer to present these distributions as percents of the table total. Marginal distributions tell us nothing about the relationship between the variables.
- There are two sets of **conditional distributions** for a two-way table: the distributions of the row variable for each fixed value of the column variable and the distributions of the column variable for each fixed value of the row variable. Comparing one set of conditional distributions is one way to describe the association between the row and the column variables.

- To find the **conditional distribution** of the row variable for one specific value of the column variable, look only at that one column in the table. Find each entry in the column as a percent of the column total.
- Bar graphs** are a flexible means of presenting categorical data. There is no single best way to describe an association between two categorical variables.
- A comparison between two variables that holds for each individual value of a third variable can be changed or even reversed when the data for all values of the third variable are combined. This is **Simpson's paradox**. Simpson's paradox is an example of the effect of lurking variables on an observed association.

LINK IT

In Chapters 4 and 5 we considered relationships between two quantitative variables. In this chapter we use two-way tables to describe relationships between two *categorical* variables. To explore relationships between two categorical variables, we examine their conditional distributions. The conditional distribution of one of the categorical variables is the distribution of that variable among only individuals who have a given value of the other variable. There is a separate conditional distribution for each value of this other variable. Changes in the pattern of the conditional distribution of one variable as the value of the other varies provide information about the relationship between the variables. No change in this pattern suggests that there is no relationship.

Although it provides no information about the relationship between two categorical variables, we can examine each variable separately by looking at their marginal distributions. The marginal distribution of one of the categorical variables is the distribution of values of that variable among all individuals described in the table. The marginal distributions allow us to see how frequently the values of each variable occur, ignoring the other variable.

As in Chapters 4 and 5, we must be careful not to assume that the patterns we observe would continue to hold for additional data or in a broader setting. Simpson's paradox is an example of how such an assumption could mislead us. Simpson's paradox occurs when the association or comparison that holds for all of several groups reverses direction when these groups are combined into a single group.

CHECK YOUR SKILLS

The Pew Internet and American Life Project interviewed several hundred teens (ages 12 to 17). One question asked was “How often do you take your cell phone to school?” Below is a two-way table of the responses by how permissive the school is with regard to cell phone use:⁷

Frequency	Forbid	Allow in school but not in class	Allow in class
Never	25	19	4
Less often	14	31	6
At least several times per week	14	23	8
Every day	97	314	57

Exercises 6.8 to 6.16 are based on this table.  CELLPHONE

6.8 How many individuals are described by this table?

(a) 468 (b) 612 (c) Need more information

6.9 How many teens from schools that forbid cell phones were among the respondents?

(a) 48 (b) 150 (c) Need more information



6.10 The percent of teens from schools that forbid cell phones among the respondents was

- (a) about 8%. (b) about 25%. (c) about 48%.

6.11 Your percent from the previous exercise is part of

- (a) the marginal distribution of school permissiveness.
 (b) the marginal distribution of the frequency that a teen brought a cell phone to school.
 (c) the conditional distribution of the frequency that a teen brought a cell phone to school among schools with a given level of permissiveness.

6.12 What percent of teens from schools that forbid cell phones brought their cell phone to school every day?

- (a) about 16% (b) about 21% (c) about 65%

6.13 Your percent from the previous exercise is part of

- (a) the marginal distribution of the frequency that a teen brought a cell phone to school.
 (b) the conditional distribution of school permissiveness among those who brought a cell phone to school every day.
 (c) the conditional distribution of the frequency that a teen brought a cell phone to school among schools that forbid cell phones.

6.14 What percent of those who brought their cell phone to school every day were from schools that forbid cell phones?

- (a) about 16% (b) about 21% (c) about 65%

6.15 Your percent from the previous exercise is part of

- (a) the marginal distribution of the frequency that a teen brought a cell phone to school.

- (b) the conditional distribution of school permissiveness among those who brought a cell phone to school every day.
 (c) the conditional distribution of the frequency that a teen brought a cell phone to school among schools with a given level of permissiveness.

6.16 A bar graph showing the conditional distribution of the frequency that a teen brought a cell phone to school among schools with a given level of permissiveness would have

- (a) 3 bars. (b) 4 bars. (c) 12 bars.

6.17 A college looks at the grade point average (GPA) of its full-time and part-time students. Grades in science courses are generally lower than grades in other courses. There are few science majors among part-time students but many science majors among full-time students. The college finds that full-time students who are science majors have higher GPAs than part-time students who are science majors. Full-time students who are not science majors also have higher GPAs than part-time students who are not science majors. Yet part-time students as a group have higher GPAs than full-time students. This finding is

(a) not possible: if both science and other majors who are full-time have higher GPAs than those who are part-time, then all full-time students together must have higher GPAs than all part-time students together.

(b) an example of Simpson's paradox: full-time students do better in both kinds of courses but worse overall because they take more science courses.

(c) due to comparing two conditional distributions that should not be compared.

CHAPTER 6 EXERCISES

6.18 Is astrology scientific? The University of Chicago's General Social Survey (GSS) is the nation's most important social science sample survey. The GSS asked a random sample of adults their opinion about whether astrology is very scientific, sort of scientific, or not at all scientific. Here is a two-way table of counts for people in the sample who had three levels of higher education degrees:⁸



Find the two conditional distributions of degree held, one for those who hold the opinion that astrology is not at all scientific and one for those who say astrology is very or sort of scientific. Based on your calculations, describe with a graph and in words the differences between those who say astrology is not at all scientific and those who say it is very or sort of scientific.

6.19 Weight-lifting injuries. Resistance training is a popular form of conditioning aimed at enhancing sports performance and is widely used among high school, college, and professional athletes, although its use for younger athletes is controversial. A random sample of 4111 patients between the ages of 8 and 30 admitted to U.S. emergency rooms with the injury code "weightlifting" was obtained. These injuries were classified as "accidental" if caused by dropped weight or improper equipment use. The patients were also classified

	Degree Held		
	Junior college	Bachelor	Graduate
Not at all scientific	87	198	111
Very or sort of scientific	43	57	28

into the four age categories 8 to 13 years, 14 to 18, 19 to 22, and 23 to 30. Here is a two-way table of the results:⁹

Age	Accidental	Not accidental
8–13	295	102
14–18	655	916
19–22	239	533
23–30	363	1008

Compare the distributions of ages for accidental and nonaccidental injuries. Use percents and draw a bar graph. What do you conclude?  WEIGHTLIFTING

Marital status and job level. We sometimes hear that getting married is good for your career. Table 6.2 presents data from one of the studies behind this generalization. To avoid gender effects, the investigators looked only at men. The data describe the marital status and the job level of all 8235 male managers and professionals employed by a large manufacturing firm.¹⁰ The firm assigns each position a grade that reflects the value of that particular job to the company. The authors of the study grouped the many job grades into quarters. Grade 1 contains jobs in the lowest quarter of the job grades, and Grade 4 contains those in the highest quarter. Exercises 6.20 to 6.24 are based on these data.  MARITALSTAT

TABLE 6.2 Marital status and job level

JOB GRADE	MARITAL STATUS				TOTAL
	SINGLE	MARRIED	DIVORCED	WIDOWED	
1	58	874	15	8	955
2	222	3927	70	20	4239
3	50	2396	34	10	2490
4	7	533	7	4	551
Total	337	7730	126	42	8235

6.20 Marginal distributions. Give (in percents) the two marginal distributions, for marital status and for job grade. Do each of your two sets of percents add to exactly 100%? If not, why not?

6.21 Percents. What percent of single men hold Grade 1 jobs? What percent of Grade 1 jobs are held by single men?

6.22 Conditional distribution. Give (in percents) the conditional distribution of job grade among single men. Should your percents add to 100% (up to roundoff error)?

6.23 Marital status and job grade. One way to see the relationship is to look at who holds Grade 1 jobs.

(a) There are 874 married men with Grade 1 jobs, and only 58 single men with such jobs. Explain why these counts by themselves don't describe the relationship between marital status and job grade.

(b) Find the percent of men in each marital status group who have Grade 1 jobs. Then find the percent in each marital group who have Grade 4 jobs. What do these percents say about the relationship?

6.24 Association is not causation. The data in Table 6.2 show that single men are more likely to hold lower-grade jobs than are married men. We should not conclude that single men can help their career by getting married. What lurking variables might help explain the association between marital status and job grade?

6.25 Race and the death penalty. Whether a convicted murderer gets the death penalty seems to be influenced by the race of the victim. Several researchers studied this issue in the 1970s and 1980s, resulting in several landmark, oft-cited, and controversial papers. Here are data on 326 cases in which the defendant was convicted of murder from one of these studies:¹¹

White Defendant		Black Defendant			
White victim	Black victim	White victim	Black victim		
Death	19	0	Death	11	6
Not	132	9	Not	52	97

(a) Use these data to make a two-way table of defendant's race (white or black) versus death penalty (Yes or No).

(b) Show that Simpson's paradox holds: a higher percent of white defendants are sentenced to death overall, but for both black and white victims a higher percent of black defendants are sentenced to death.

(c) Use the data to explain why the paradox holds in language that a judge could understand.  DISCRIM

6.26 Obesity and health. To estimate the health risks of obesity, we might compare how long obese and nonobese people live. Smoking is a lurking variable that may reduce the gap between the two groups, because smoking tends to both reduce weight and lead to earlier death. So if we ignore smoking, we may underestimate the health risks of obesity. Illustrate Simpson's paradox by a simplified version of this situation: make up two-way tables of obese (Yes or No) by early death (Yes or No) separately for smokers and nonsmokers such that

■ Obese smokers and obese nonsmokers are both more likely to die earlier than those who are not obese.

■ But when smokers and nonsmokers are combined into a two-way table of obese by early death, persons who are not obese are more likely to die earlier because more of them are smokers.

The following exercises ask you to answer questions from data without having the details outlined for you. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan**, **Solve**, and **Conclude** steps of the process as illustrated in Example 6.3 (page 163).

6.27 Smoking cessation. A large randomized trial was conducted to assess the efficacy of Chantix for smoking cessation compared with bupropion (more commonly known as Wellbutrin or Zyban) and a placebo. Chantix is different from most other quit-smoking products in that it targets nicotine receptors in the brain, attaches to them, and blocks nicotine from reaching them, while bupropion is an antidepressant often used to help people stop smoking. Generally healthy smokers who smoked at least 10 cigarettes per day were assigned at random to take Chantix ($n = 352$), bupropion ($n = 329$), or a placebo ($n = 344$). The response measure is continuous cessation from smoking for Weeks 9 through 12 of the study. Here is a two-way table of the results:¹²



Joe Raedle/Getty Images

	Treatment		
	Chantix	Bupropion	Placebo
No smoking in Weeks 9–12	155	97	61
Smoked in Weeks 9–12	197	232	283

How does whether a subject smoked in Weeks 9 to 12 depend on the treatment received?

6.28 Animal testing. “It is right to use animals for medical testing if it might save human lives.” The General Social Survey asked 1152 adults to react to this statement. Here is the two-way table of their responses:

Response	Male	Female
Strongly agree	76	59
Agree	270	247
Neither agree nor disagree	87	139
Disagree	61	123
Strongly disagree	22	68

How do the distributions of opinion differ between men and women? ANTESTING

6.29 College degrees. “Colleges and universities across the country are grappling with the case of the mysteriously vanishing male.” So said an article in the Washington Post. Here are data on the numbers of degrees earned in 2012–2013, as projected by the National Center for Education Statistics. The table entries are counts of degrees in thousands.¹³ DEGREES

Degree	Female	Male
Associate's	519	304
Bachelor's	989	731
Master's	418	266
Professional	49	50
Doctor's	40	35

Briefly contrast the counts and distributions of men and women in earning degrees. Are men “vanishing” from colleges and universities across the country?

6.30 Complications of bariatric surgery. Bariatric surgery, or weight-loss surgery, includes a variety of procedures performed on people who are obese. Weight loss is achieved by reducing the size of the stomach with an implanted medical device (gastric banding), by removing a portion of the stomach (sleeve gastrectomy), or by resecting and rerouting the small intestines to a small stomach pouch (gastric bypass surgery). Because there can be complications using any of these methods, the National Institute of Health recommends bariatric surgery for obese people with a body mass index (BMI) of at least 40, and for people with BMI 35 and serious coexisting medical conditions such as diabetes. Serious complications include potentially life-threatening, permanently disabling, and fatal outcomes. Here is a two-way table for data collected in Michigan over several years giving counts of non-life-threatening complications, serious complications, and no complications for these three types of surgeries:¹⁴ BARIATRIC

	Type of Complication				Total
	Non-life-threatening	Serious	None		
Gastric banding	81	46	5253	5380	
Sleeve gastrectomy	31	19	804	854	
Gastric bypass	606	325	8110	9041	

Health	Current Smoker	
	Yes	No
Excellent	25	484
Very good	115	1557
Good	145	1309
Fair	90	545
Poor	29	11

What do the data say about differences in complications for the three types of surgeries?

6.31 Smokers rate their health. The University of Michigan Health and Retirement Study (HRS) surveys more than 22,000 Americans over the age of 50 every two years. A subsample of the HRS participated in a 2009 Internet-based survey that collected information on a number of topical areas, including health (physical and mental health behaviors), psychosocial items, economics (income, assets, expectations, and consumption), and retirement.¹⁵ Two of the questions asked were “Would you say your health is excellent, very good, good, fair, or poor?” and “Do you smoke cigarettes now?” The two-way table summarizes the answers on these two questions.  **SMOKERRATING**

What do these data say about differences in self-evaluation of health for current smokers and nonsmokers?

6.32 Python eggs. How is the hatching of water python eggs influenced by the temperature of the snake’s nest? Researchers placed 104 newly laid eggs in a hot environment, 56 in a neutral environment, and 27 in a cold environment. Hot duplicates the warmth provided by the mother python. Neutral and cold are cooler, as when the mother is absent. The results: 75 of the hot eggs hatched, along with 38 of the neutral eggs and 16 of the cold eggs.¹⁶

- (a) Make a two-way table of “environment temperature” against “hatched or not.”
- (b) The researchers anticipated that eggs would hatch less well at cooler temperatures. Do the data support that anticipation?



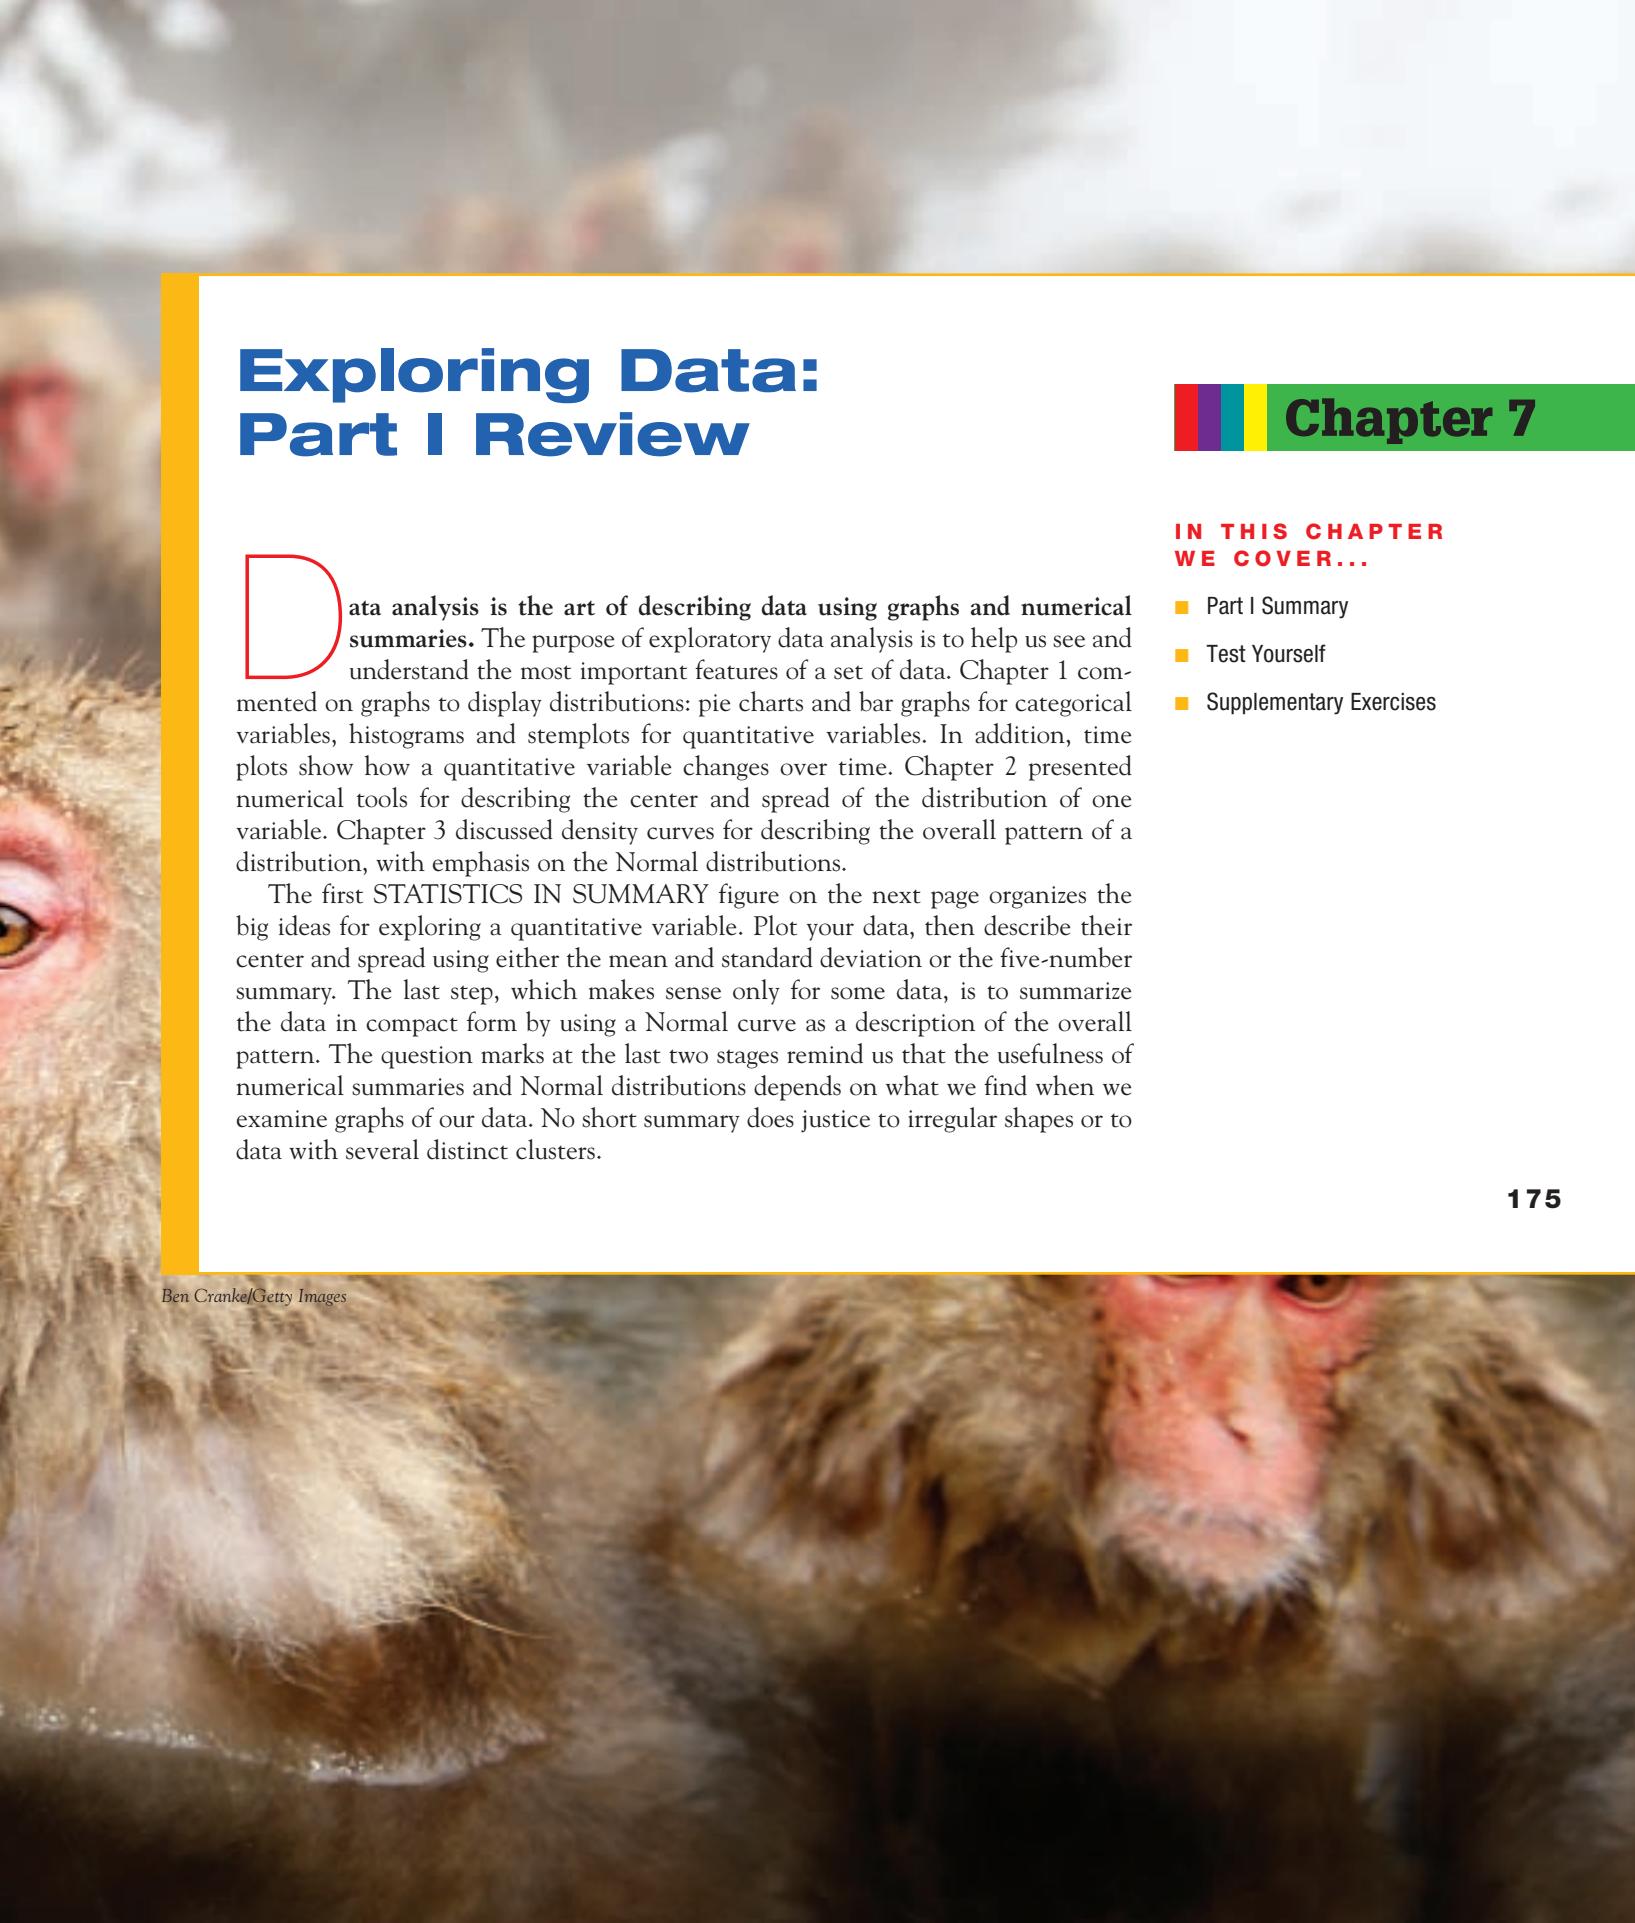
EXPLORING THE WEB

6.33 Promoting women. In academics, faculty typically start as assistant professors, are promoted to associate professor (and gain tenure), and finally reach the rank of full professor. Some have argued that women have a harder time gaining promotion to associate and full professor than do men. Do data support this argument? Search the Web to find the number of faculty by rank and gender at some university. Do you see a pattern that suggests that the proportion of women decreases as rank increases? We found several sources of data by doing a Google search on “faculty head count by rank and gender.” In addition to discussing the pattern you find, provide the data, the name of the school, and the source of the data.

6.34 Accidental deaths and ages. Accidental deaths are shocking and tragic. Do the ways in which people die by accident change with age? Look at the most recent *Statistical Abstract of the United States* (www.census.gov/compendia/statab/) and make a two-way table that provides the counts of deaths due to accidents from various causes for three different age groups. What do you conclude?

6.35 Simpson’s paradox. Find an example of Simpson’s paradox and discuss how your example illustrates the paradox. Two examples that we found (thanks to Patricia Humphrey at Georgia Southern University) are www.nytimes.com/2006/07/15/education/15report.html and online.wsj.com/article/SB125970744553071829.html.





Exploring Data: Part I Review



Chapter 7

IN THIS CHAPTER WE COVER...

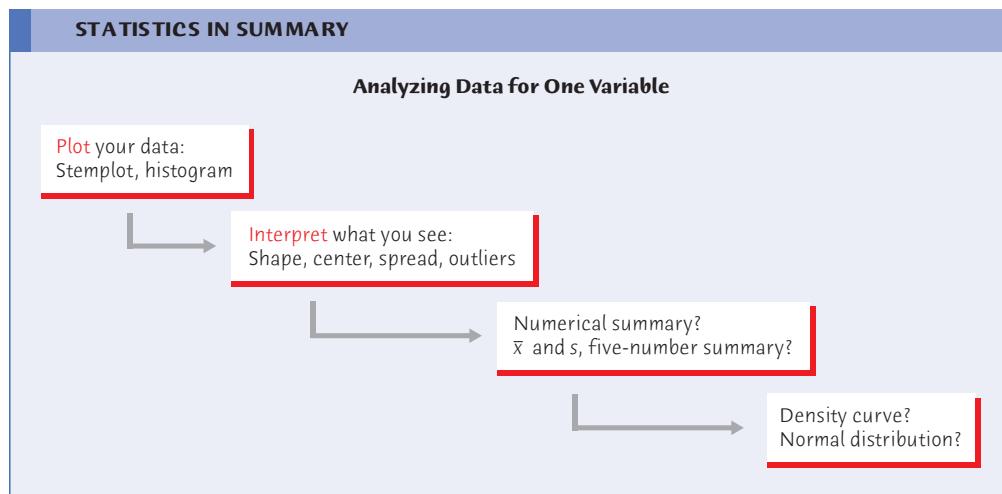
- Part I Summary
- Test Yourself
- Supplementary Exercises

Data analysis is the art of describing data using graphs and numerical summaries. The purpose of exploratory data analysis is to help us see and understand the most important features of a set of data. Chapter 1 commented on graphs to display distributions: pie charts and bar graphs for categorical variables, histograms and stemplots for quantitative variables. In addition, time plots show how a quantitative variable changes over time. Chapter 2 presented numerical tools for describing the center and spread of the distribution of one variable. Chapter 3 discussed density curves for describing the overall pattern of a distribution, with emphasis on the Normal distributions.

The first STATISTICS IN SUMMARY figure on the next page organizes the big ideas for exploring a quantitative variable. Plot your data, then describe their center and spread using either the mean and standard deviation or the five-number summary. The last step, which makes sense only for some data, is to summarize the data in compact form by using a Normal curve as a description of the overall pattern. The question marks at the last two stages remind us that the usefulness of numerical summaries and Normal distributions depends on what we find when we examine graphs of our data. No short summary does justice to irregular shapes or to data with several distinct clusters.

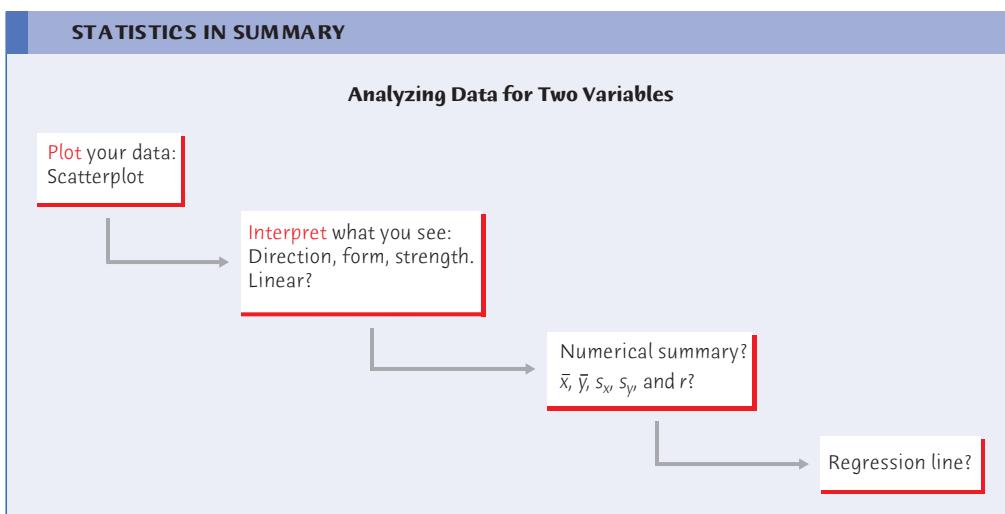
175

Ben Cranke/Getty Images



Chapters 4 and 5 applied the same ideas to relationships between two quantitative variables. The second STATISTICS IN SUMMARY figure retraces the big ideas, with details that fit the new setting. Always begin by making graphs of your data. In the case of a scatterplot, we have learned a numerical summary only for data that show a roughly linear pattern on the scatterplot. The summary is then the means and standard deviations of the two variables and their correlation. A regression line drawn on the plot gives a compact description of the overall pattern that we can use for prediction. Once again there are question marks at the last two stages to remind us that correlation and regression describe only straight-line relationships. Chapter 6 shows how to understand relationships between two categorical variables; comparing well-chosen percents is the key.

You can organize your work in any open-ended data analysis setting by following the four-step STATE, PLAN, SOLVE, and CONCLUDE process first introduced in Chapter 2.



After we have mastered the extra background needed for statistical inference, this process will also guide practical work on inference later in the book.

PART I SUMMARY

Here are the most important skills you should have acquired from reading Chapters 1 to 6.

A. Data

1. Identify the individuals and variables in a set of data.
2. Identify each variable as categorical or quantitative. Identify the units in which each quantitative variable is measured.
3. Identify the explanatory and response variables in situations where one variable explains or influences another.

B. Displaying Distributions

1. Recognize when a pie chart can and cannot be used.
2. Make a bar graph of the distribution of a categorical variable or, in general, to compare related quantities.
3. Interpret pie charts and bar graphs.
4. Make a histogram of the distribution of a quantitative variable.
5. Make a stemplot of the distribution of a small set of observations. Round leaves or split stems as needed to make an effective stemplot.
6. Make a time plot of a quantitative variable over time. Recognize patterns such as trends and cycles in time plots.

C. Describing Distributions (Quantitative Variable)

1. Look for the overall pattern and for major deviations from the pattern.
2. Assess from a histogram or stemplot whether the shape of a distribution is roughly symmetric, distinctly skewed, or neither. Assess whether the distribution has one or more major peaks.
3. Describe the overall pattern by giving numerical measures of center and spread in addition to a verbal description of shape.
4. Decide which measures of center and spread are more appropriate: the mean and standard deviation (especially for symmetric distributions) or the five-number summary (especially for skewed distributions).
5. Recognize outliers and give plausible explanations for them.

D. Numerical Summaries of Distributions

1. Find the median M and the quartiles Q_1 and Q_3 for a set of observations.
2. Find the five-number summary and draw a boxplot; assess center, spread, symmetry, and skewness from a boxplot.

3. Find the mean \bar{x} and the standard deviation s for a set of observations.
4. Understand that the median is more resistant than the mean. Recognize that skewness in a distribution moves the mean away from the median toward the long tail.
5. Know the basic properties of the standard deviation: $s \geq 0$ always; $s = 0$ only when all observations are identical and increases as the spread increases; s has the same units as the original measurements; s is pulled strongly up by outliers or skewness.

E. Density Curves and Normal Distributions

1. Know that areas under a density curve represent proportions of all observations and that the total area under a density curve is 1.
2. Approximately locate the median (equal-areas point) and the mean (balance point) on a density curve.
3. Know that the mean and median both lie at the center of a symmetric density curve and that the mean moves farther toward the long tail of a skewed curve.
4. Recognize the shape of Normal curves and estimate by eye both the mean and standard deviation from such a curve.
5. Use the 68–95–99.7 rule and symmetry to state what percent of the observations from a Normal distribution fall between two points when both points lie at the mean or one, two, or three standard deviations on either side of the mean.
6. Find the standardized value (z -score) of an observation. Interpret z -scores and understand that any Normal distribution becomes the standard Normal $N(0, 1)$ distribution when standardized.
7. Given that a variable has a Normal distribution with a stated mean μ and standard deviation σ , calculate the proportion of values above a stated number, below a stated number, or between two stated numbers.
8. Given that a variable has a Normal distribution with a stated mean μ and standard deviation σ , calculate the point having a stated proportion of all values above it or below it.

F. Scatterplots and Correlation

1. Make a scatterplot to display the relationship between two quantitative variables measured on the same subjects. Place the explanatory variable (if any) on the horizontal scale of the plot.
2. Add a categorical variable to a scatterplot by using a different plotting symbol or color.
3. Describe the direction, form, and strength of the overall pattern of a scatterplot. In particular, recognize positive or negative association and linear (straight-line) patterns. Recognize outliers in a scatterplot.
4. Judge whether it is appropriate to use correlation to describe the relationship between two quantitative variables. Find the correlation r .

5. Know the basic properties of correlation: r measures the direction and strength of only straight-line relationships; r is always a number between -1 and 1 ; $r = \pm 1$ only for perfect straight-line relationships; r moves away from 0 toward ± 1 as the straight-line relationship gets stronger.

G. Regression Lines

1. Understand that regression requires an explanatory variable and a response variable. Correctly identifying which variable is the explanatory variable and which is the response variable is important. Switching these will result in different regression lines. Use a calculator or software to find the least-squares regression line of a response variable y on an explanatory variable x from data.
2. Explain what the slope b and the intercept a mean in the equation $\hat{y} = a + bx$ of a regression line.
3. Draw a graph of a regression line when you are given its equation.
4. Use a regression line to predict y for a given x . Recognize extrapolation and be aware of its dangers.
5. Find the slope and intercept of the least-squares regression line from the means and standard deviations of x and y and their correlation.
6. Use r^2 , the square of the correlation, to describe how much of the variation in one variable can be accounted for by a straight-line relationship with another variable.
7. Recognize outliers and potentially influential observations from a scatterplot with the regression line drawn on it.
8. Calculate the residuals and plot them against the explanatory variable x . Recognize that a residual plot magnifies the pattern of the scatterplot of y versus x .



Driving in Canada

Canada is a civilized and restrained nation, at least in the eyes of Americans. A survey sponsored by the Canada Safety Council suggests that driving in Canada may be more adventurous than expected. Of the Canadian drivers surveyed, 88% admitted to aggressive driving in the past year, and 76% said that sleep-deprived drivers were common on Canadian roads. What really alarms us is the name of the survey: the Nerves of Steel Aggressive Driving Study.

H. Cautions about Correlation and Regression

1. Understand that both r and the least-squares regression line can be strongly influenced by a few extreme observations.
2. Recognize possible lurking variables that may explain the observed association between two variables x and y .
3. Understand that even a strong correlation does not mean that there is a cause-and-effect relationship between x and y .
4. Give plausible explanations for an observed association between two variables: direct cause and effect, the influence of lurking variables, or both.

I. Categorical Data (Optional)

1. From a two-way table of counts, find the marginal distributions of both variables by obtaining the row sums and column sums.
2. Express any distribution in percents by dividing the category counts by their total.

3. Describe the relationship between two categorical variables by computing and comparing percents. Often this involves comparing the conditional distributions of one variable for the different categories of the other variable.
4. Recognize Simpson's paradox and be able to explain it.

TEST YOURSELF

The questions below include multiple-choice, calculations, and short-answer questions. They will help you review the basic ideas and skills presented in Chapters 1 to 6.

- 7.1** As part of a data base on new births at a hospital some variables recorded are the age of the mother, marital status of the mother (single, married, divorced, other), weight of the baby, and sex of the baby. Of these variables
- (a) age, marital status, and weight are quantitative variables.
 - (b) age and weight are categorical variables.
 - (c) sex and marital status are categorical variables.
 - (d) sex, marital status, and age are categorical variables.
- 7.2** You are interested in obtaining information about the performance of students in your statistics class and seeing how this performance is affected by several factors such as sex. To do this you are going to give a questionnaire to all students in the class. Give two questions for which the response is categorical and two questions for which the response is quantitative. For the categorical variables, give the possible values.

Pocket change. In a statistics class with 136 students, the professor records how much money each student has in his or her possession during the first class of the semester. Figure 7.1 gives the histogram of the data collected. Use this histogram to help answer Questions 7.3 to 7.5.

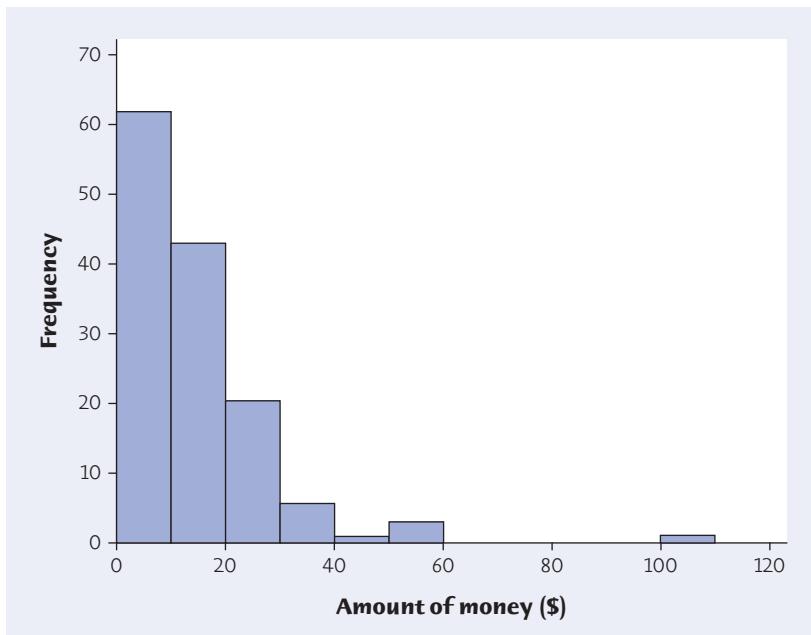


FIGURE 7.1

Histogram of the distribution of the amount of money carried by each student on the first day of class, for Questions 7.3 to 7.5.

- 7.3** The *number* of students with under \$10 in their possession is closest to
(a) 40. (b) 50. (c) 60. (d) 70.
- 7.4** The histogram
(a) is skewed right. (b) has an outlier.
(c) is asymmetric. (d) is all of the above.
- 7.5** The *percent* of students with \$20 or more in their possession is
(a) about 10%. (b) about 20%.
(c) about 30%. (d) over 40%.
- 7.6** A reporter wishes to portray baseball players as overpaid. Which measure of center should he report as the average salary of major league players?
(a) The mean (b) The median
(c) Either the mean or the median. It doesn't matter since they will be equal.
(d) Neither the mean nor the median. Both will be much lower than the actual average salary.

Genetic engineering for cancer treatment. Here's a new idea for treating advanced melanoma, the most serious kind of skin cancer. Genetically engineer white blood cells to better recognize and destroy cancer cells, then infuse these cells into patients. The subjects in a small initial study were 11 patients whose melanoma had not responded to existing treatments. One question was how rapidly the new cells would multiply after infusion, as measured by the doubling time in days.¹ Use the following doubling times in days to answer Questions 7.7 to 7.10.

1.4 1.0 1.3 1.0 1.3 2.0 0.6 0.8 0.7 0.9 1.9

- 7.7** What is the mean number of days for the 11 doubling times?
(a) 1.17 days (b) 1.0 day (c) 0.9 day (d) 0.46 day
- 7.8** What is the median number of days for the 11 doubling times?
(a) 2.0 days (b) 1.17 days (c) 1.0 day (d) 0.9 day
- 7.9** What is the first quartile for these data?
(a) 1.3 days (b) 1.0 day (c) 0.85 day (d) 0.8 day
- 7.10** What is the interquartile range for these data?
(a) 1.3 days (b) 0.8 day (c) 0.6 day (d) 0.4 day
- 7.11** Which of the following is likely to have a mean that is smaller than the median?
(a) The salaries of all National Football League players
(b) The scores of students (out of 100 points) on a very easy exam in which most students score perfectly, but a few do very poorly
(c) The prices of homes in a large city
(d) The scores of students (out of 100 points) on a very difficult exam in which most students score poorly, but a few do very well
- 7.12** For a biology project, you measure the tail length in centimeters and weight in grams of 12 mice of the same variety. What units of measurement do each of the following have?
(a) The mean length of the tails (b) The first quartile of the tail lengths
(c) The standard deviation of the tail lengths
(d) The variance of the weights



Zeva Oelbaum/Photolibrary

Employment times. A sample of 40 employees from the local Honda plant was obtained, and the length of time (in months) worked was recorded for each employee. A stemplot of these data follows. Use the stemplot to answer Questions 7.13 and 7.14. In the stemplot 5|2 represents 52 months.

5	2 2 3 3 4 5 7 8 9 9
6	0 0 0 2 3 4 4 4 5 6 7 7 8 8 8 9
7	3 4 5 5 6 6 7 7 7 8 8 9 9
8	
9	8

- 7.13** What would be a better way to represent this data set?
- Display the data in a time plot
 - Split the stems
 - Use a pie chart
 - Use a histogram with class width equal to 10
- 7.14** The percent of employees in the sample who have worked at the plant for less than 5 years is
- approximately zero.
 - 10%.
 - 15%.
 - 25%.

What color is your car? The bar graph in Figure 7.2 gives the distribution of the most popular colors for vehicles sold in North America in 2010.² Use the bar graph to help answer Questions 7.15 and 7.16.

- 7.15** Approximately what percent of vehicles sold in North America in 2010 were beige or brown?
- 5%
 - 10%
 - 15%
 - 20%
- 7.16** The total of the percents of the bars in the graph add to 96%.
- The percent of green cars sold must be less than 2%.
 - A pie chart could be drawn for the colors given in the bar graph.
 - The percent of vehicles sold that are either silver or white is just over 40%.
 - None of the above.

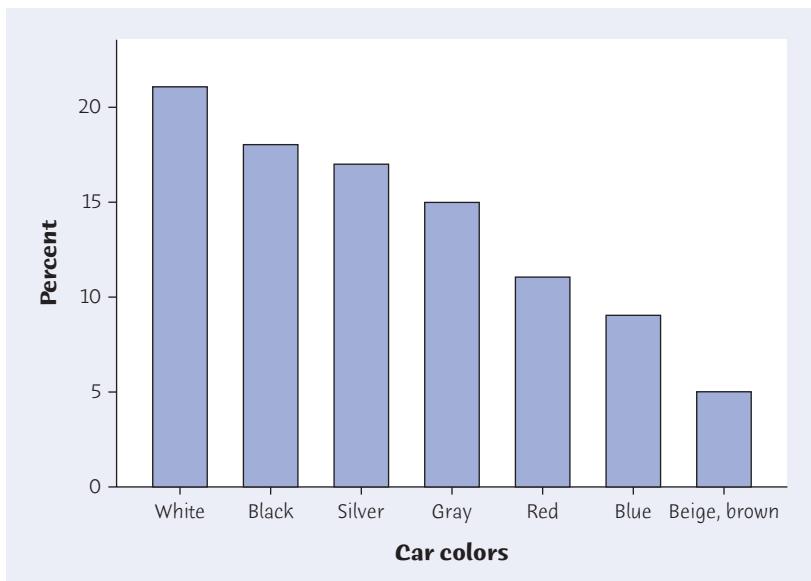


FIGURE 7.2

Bar graph of the distribution of the most popular colors for vehicles sold in North America in 2010, for Questions 7.15 and 7.16.

7.17 Mechanical measurements on supposedly identical objects usually vary. The variation often follows a Normal distribution. The stress required to break a type of bolt varies Normally with mean 75 kilopounds per square inch (ksi) and standard deviation 8.3 ksi.

- What percent of these bolts will withstand a stress of 90 ksi without breaking?
- What range covers the middle 50% of breaking strengths for these bolts?

7.18 A professor knows from past experience that the time for students to complete a quiz has an $N(19, 3)$ distribution.

- If he allows 20 minutes for the quiz, what percent of the students will not complete the quiz?
- Suppose that he wants to allow sufficient time so that 95% of the students will complete the quiz in the allotted time. How much time should he allow for the quiz?

7.19 The Aleppo pine and the Torrey pine are widely planted as ornamental trees in Southern California. Here are the lengths (centimeters) of 15 Aleppo pine needles:³

10.2 7.2 7.6 9.3 12.1 10.9 9.4 11.3 8.5 8.5 12.8 8.7 9.0 9.0 9.4

- Find the five-number summary for the distribution of Aleppo pine needles.

Figure 7.3 gives a boxplot for the distribution of the lengths (centimeters) of 18 Torrey pine needles. Use this information to help answer the remainder of this question.

- The median of the distribution of Torrey pine needles is closest to which of the following values?

24 25 27 30

- Twenty-five percent of the Torrey pine needles exceed what value?
- Given only the length of a needle, do you think you could say which pine species it comes from? Explain briefly.

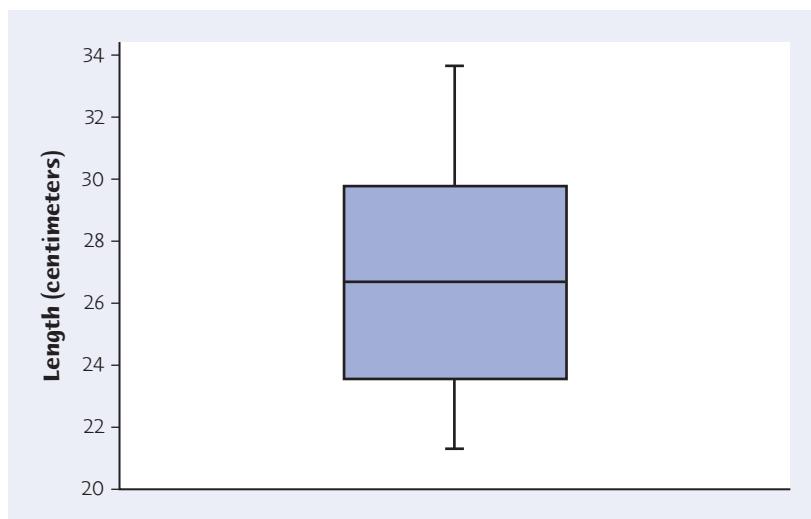


FIGURE 7.3

Boxplot for the distribution of the lengths (centimeters) of 18 Torrey pine needles, for Question 7.19.



Graig Tuttle/CORBIS



Beer in South Dakota

Take a break from doing exercises to apply your math

to beer cans in South Dakota. A newspaper there reported that every year an average of 650 beer cans per mile are tossed onto the state's highways. South Dakota has about 83,000 miles of roads. How many beer cans is that in all? The Census Bureau says that there are about 780,000 people in South Dakota. How many beer cans does each man, woman, and child in the state toss on the road each year? That's pretty impressive. Maybe the paper got its numbers wrong.

Soap in the shower. From Rex Boggs in Australia comes an unusual data set: before showering in the morning, he weighed the bar of soap in his shower stall. The weight goes down as the soap is used. The data appear below (weights in grams). Notice that Mr. Boggs forgot to weigh the soap on some days. Questions 7.20 to 7.23 are based on the soap data set. 

Day	Weight	Day	Weight	Day	Weight
1	124	8	84	16	27
2	121	9	78	18	16
5	103	10	71	19	12
6	96	12	58	20	8
7	90	13	50	21	6

7.20 Figure 7.4 is a scatterplot of the weight of the bar of soap against day. How would you describe the overall pattern?

- (a) Sharply curved.
- (b) There are two distinct clusters that are widely separated.
- (c) A very weak positive association.
- (d) A strong negative association.

7.21 The equation of the least-squares regression line for predicting soap weight from day is

$$\text{weight} = 133.2 - 6.3 \times \text{day}$$

What does this tell us about the rate at which the soap lost weight?

- (a) The soap lost about -6.3 grams per day.
- (b) The soap lost about 6.3 grams per day.
- (c) The soap lost about -133.2 grams per day.
- (d) The soap lost about 133.2 grams per day.

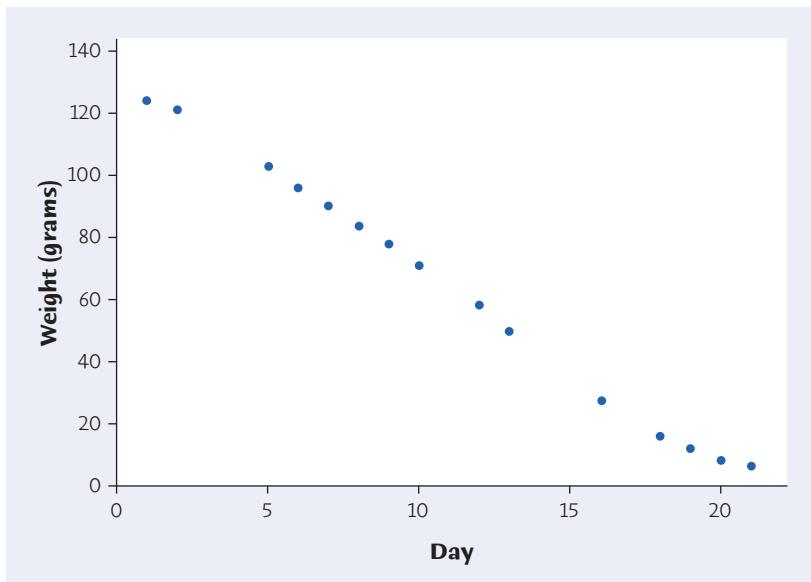


FIGURE 7.4

Scatterplot of the weight of a bar of soap against day, for Question 7.20.

- 7.22** The equation of the least-squares regression line for predicting soap weight from day is

$$\text{weight} = 133.2 - 6.3 \times \text{day}$$

Mr. Boggs did not measure the weight of the soap on Day 4. Use the regression equation to predict that weight.

- (a) 108 grams (b) 126.9 grams (c) 157.3 grams (d) 526.5 grams

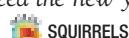
- 7.23** The equation of the least-squares regression line for predicting soap weight from day is

$$\text{weight} = 133.2 - 6.3 \times \text{day}$$

I use the regression equation to predict the weight of the soap on Day 30. I conclude that

- (a) the soap will last at least a month.
- (b) the prediction is not sensible, because the prediction is far outside the range of values of the response variable.
- (c) the prediction is not sensible, because 30 days is far outside the range of values of the explanatory variable.
- (d) the prediction is not sensible, because of the outlier present in the data.

Squirrels and their food supply. The fact that animal species produce more offspring when their supply of food goes up isn't surprising. The fact that some animals appear able to anticipate unusual food abundance is surprising. Red squirrels eat seeds from pine cones, a food source that occasionally has very large crops (called seed masting). Below are data on an index of the abundance of pine cones (larger values indicate greater abundance) and average number of offspring per female over 16 years.⁴ What makes these data interesting is that the offspring are conceived in the spring, before the cones mature in the fall to feed the new young squirrels through the winter. Questions 7.24 to 7.26 are based on these data.



SQUIRRELS

Cone index x	0.00	2.02	0.25	3.22	4.68	0.31	3.37	3.09
Offspring y	1.49	1.10	1.29	2.71	4.07	1.29	3.36	2.41
Cone index x	2.44	4.81	1.88	0.31	1.61	1.88	0.91	1.04
Offspring y	1.97	3.41	1.49	2.02	3.34	2.41	2.15	2.12

- 7.24** Figure 7.5 is a scatterplot of average number of offspring per female against cone index. Which of the following is a plausible value of the correlation, r , between average number of offspring per female and cone index?

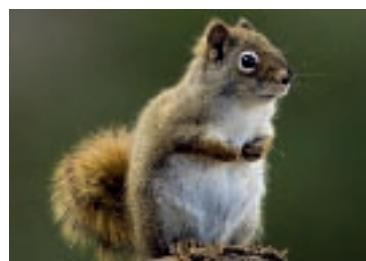
- (a) 0 (b) 1 (c) 0.75 (d) -0.75

- 7.25** The equation of the least-squares regression line for predicting average number of offspring per female from cone index is

$$\text{offspring} = 1.41 + 0.44 \times \text{cone index}$$

What does the intercept of 1.41 tell us?

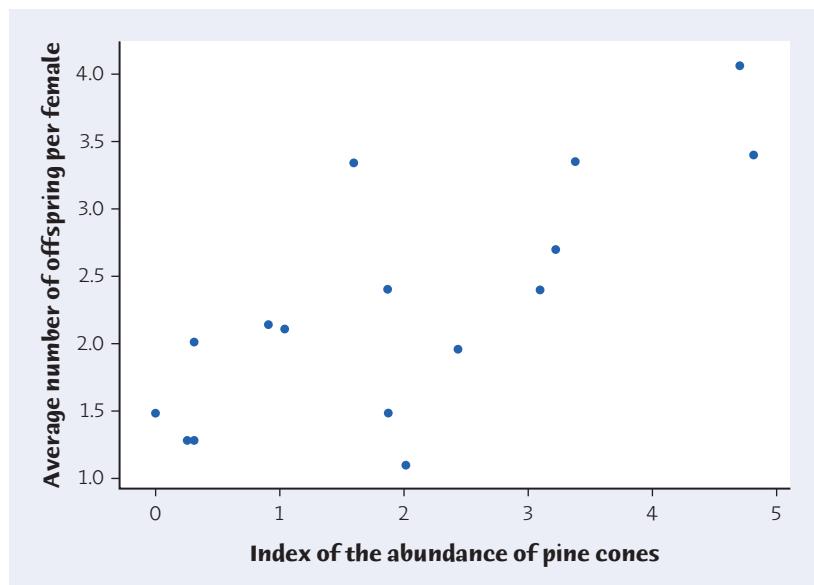
- (a) The average number of offspring per female is 1.41.
- (b) The predicted number of offspring per female is 1.41.
- (c) The predicted number of offspring per female when the cone index is 0 is 1.41.
- (d) All of the above.



© Don Johnston

FIGURE 7.5

Scatterplot of the average number of offspring per female against cone index, for Question 7.24.



- 7.26** The equation of the least-squares regression line for predicting average number of offspring per female from cone index is

$$\text{offspring} = 1.41 + 0.44 \times \text{cone index}$$

Use this to predict the average number of offspring per female for a year with a cone index of 0.25.

- (a) 1.52 (b) 1.29 (c) 0.44 (d) 0.11

- 7.27** How well do people remember their past diet? Data are available for 91 people who were asked about their diet when they were 18 years old. Researchers asked them at about age 55 to describe their eating habits at age 18. For each subject, the researchers calculated the correlation between actual intakes of many foods at age 18 and the intakes the subjects now remember. The median of the 91 correlations was $r = 0.217$.⁵ Which of the following conclusions is consistent with this correlation?

- (a) We conclude that subjects remember approximately 21.7% of their food intakes at age 18.
 (b) We conclude that subjects remember approximately $r^2 = 0.217^2 = 0.047$ of their food intakes at age 18.
 (c) We conclude that food intake at age 55 is about 21.7% that of food intake at age 18.
 (d) We conclude that memory of food intakes in the distant past is fair to poor.

- 7.28** Joe's retirement plan invests in stocks through an "index fund" that follows the behavior of the stock market as a whole, as measured by the Standard & Poor's (S&P) 500 stock index. Joe wants to buy a mutual fund that does not track the index closely. He reads that monthly returns from Fidelity Technology Fund have correlation $r = 0.77$ with the S&P 500 index and that Fidelity Real Estate Fund has correlation $r = 0.37$ with the index. Which of the following is correct?

- (a) The Fidelity Technology Fund has a closer relationship to returns from the stock market as a whole and also has higher returns than the Fidelity Real Estate Fund.

- (b) The Fidelity Technology Fund has a closer relationship to returns from the stock market as a whole, but we cannot say that it has higher returns than the Fidelity Real Estate Fund.
- (c) The Fidelity Real Estate Fund has a closer relationship to returns from the stock market as a whole and also has higher returns than the Fidelity Technology Fund.
- (d) The Fidelity Real Estate Fund has a closer relationship to returns from the stock market as a whole, but we cannot say that it has higher returns than the Fidelity Technology Fund.

Monkey calls. The usual way to study the brain's response to sounds is to have subjects listen to "pure tones." The response to recognizable sounds may differ. To compare responses, researchers anesthetized macaque monkeys. They fed pure tones and also monkey calls directly to their brains by inserting electrodes. Response to the stimulus was measured by the firing rate (electrical spikes per second) of neurons in various areas of the brain. Table 7.1 contains the responses for 37 neurons.⁶ Figure 7.6 is a scatterplot of monkey call response against pure-tone response (explanatory variable). Questions 7.29 and 7.30 refer to these data and the scatterplot.  MONKEYCALLS

7.29 We might expect some neurons to have strong responses to any stimulus and others to have consistently weak responses. There would then be a strong relationship between tone response and call response. From the scatterplot of monkey call response against pure-tone response in Figure 7.6 what would you estimate the correlation r to be?

- (a) -0.6 (b) -0.1 (c) 0.1 (d) 0.6

7.30 Which of the following statements about the scatterplot in Figure 7.6 is correct?

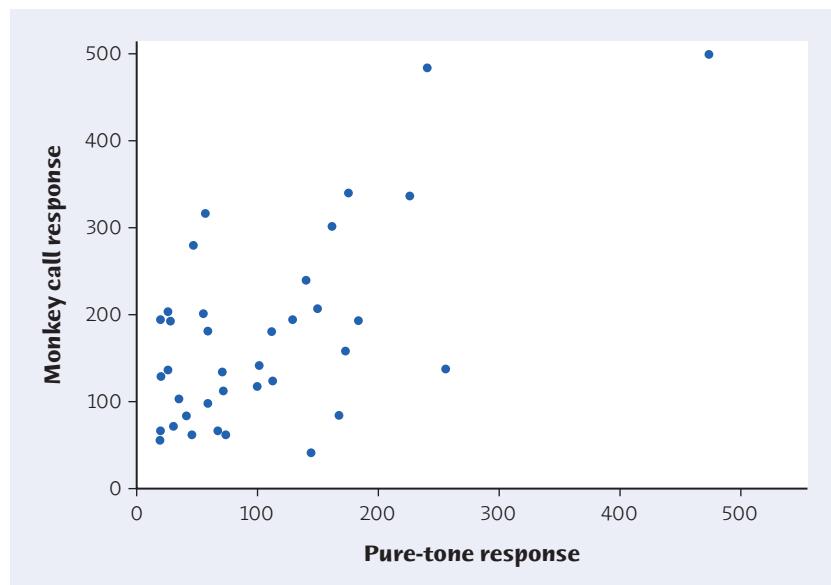
- (a) There is moderate evidence that pure-tone response causes monkey call response.
- (b) There is moderate evidence that monkey call response causes pure-tone response.
- (c) There are one or two outliers and at least one of these may also be influential.
- (d) None of the above.

TABLE 7.1 Neuron response (electrical firing rate per second) to pure tones and monkey calls

NEURON	TONE	CALL	NEURON	TONE	CALL	NEURON	TONE	CALL
1	474	500	14	145	42	26	71	134
2	256	138	15	141	241	27	68	65
3	241	485	16	129	194	28	59	182
4	226	338	17	113	123	29	59	97
5	185	194	18	112	182	30	57	318
6	174	159	19	102	141	31	56	201
7	176	341	20	100	118	32	47	279
8	168	85	21	74	62	33	46	62
9	161	303	22	72	112	34	41	84
10	150	208	23	20	193	35	26	203
11	19	66	24	21	129	36	28	192
12	20	54	25	26	135	37	31	70
13	35	103						

FIGURE 7.6

Scatterplot of monkey tone response against pure-tone response, for Questions 7.29 and 7.30.



Catalog shopping (optional). What is the most important reason that students buy from catalogs? The answer may differ for different groups of students. Here are counts for samples of American and East Asian students at a large midwestern university.⁷ Use these counts to answer Questions 7.31 to 7.33.  SHOPPING

Reason	American	Asian
Save time	29	10
Easy	28	11
Low price	17	34
Live far from stores	11	4
No pressure to buy	10	3
Other	20	7
Total	115	69

- 7.31** What percent of all students say that the most important reason to buy from a catalog is to save time?
 (a) 74% (b) 25% (c) 21% (d) 14%
- 7.32** What percent of East Asian students say that the most important reason to buy from a catalog is low price?
 (a) 67% (b) 49% (c) 28% (d) 18%
- 7.33** What are the most important differences between the two groups of students?
 (a) The most important reasons for American students to buy from a catalog are to save time and because it is easy, while for East Asian students it is low price.
 (b) American students appear to be almost three times more likely to live far from stores than East Asian students.

- (c) East Asian students are twice as likely to purchase from a catalog because of low price than American students.
 (d) All of the above.

Investment strategies. One reason to invest abroad is that markets in different countries don't move in step. When American stocks go down, foreign stocks may go up. So an investor who holds both bears less risk. That's the theory. But then we read in a magazine article that the correlation between changes in American and European stock prices rose from 0.4 in the mid-1990s to 0.8 in 2000.⁸ Questions 7.34 and 7.35 refer to this article.

- 7.34** Explain to an investor who knows no statistics why the fact stated in this article reduces the protection provided by buying European stocks.
- 7.35** The same article that claims that the correlation between changes in stock prices in Europe and the United States is 0.8 goes on to say: "Crudely, that means that movements on Wall Street can explain 80% of price movements in Europe."
 (a) Is this true? Circle your answer: Yes No
 (b) What is the correct percent explained if $r = 0.8$?
- 7.36** Researchers wished to determine whether individual differences in introspective ability are reflected in the anatomy of brain regions responsible for this function. They measured introspective ability (using a score on a test of introspective ability, with larger values indicating greater introspective ability) and gray-matter volume in milliliters (the Brodmann area) in the anterior prefrontal cortex of the brain of 29 subjects. Here are the data:  **INTROSPECTION**

Volume	0.55	0.58	0.59	0.59	0.59	0.61	0.62	0.63	0.63	0.63
Introspective ability	59	62	43	63	83	61	55	57	57	67
Volume	0.63	0.64	0.65	0.65	0.65	0.65	0.65	0.66	0.66	0.67
Introspective ability	72	62	58	62	65	70	75	60	63	71
Volume	0.67	0.67	0.68	0.69	0.70	0.70	0.71	0.72	0.75	
Introspective ability	71	80	68	72	66	73	61	80	75	

The researchers wished to determine the equation of the least-squares regression line for predicting introspective ability (y) from gray-matter volume (x). To do this they calculated the following summary statistics:

$$\bar{x} = 0.649, s_x = 0.045$$

$$\bar{y} = 65.86, s_y = 8.69$$

$$r = 0.448$$

- (a) Use this information to calculate the equation of the least-squares regression line.
 (b) Based on the least-squares regression line, what would you predict introspective ability to be for someone with gray-matter volume 0.60?
 (c) Based on the least-squares regression line, what would you predict introspective ability to be for someone with gray-matter volume 0.99? How reliable do you think this prediction is? Explain your answer.
- 7.37** Animals and people that take in more energy than they expend will get fatter. Here are data on 12 rhesus monkeys: 6 lean monkeys (4% to 9% body fat) and 6 obese

monkeys (13% to 44% body fat). The data report the energy expended in 24 hours (kilojoules per minute) and the lean body mass (kilograms, leaving out fat) for each monkey.⁹

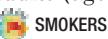


THINFATMONKEYS

Lean		Obese	
Mass	Energy	Mass	Energy
6.6	1.17	7.9	0.93
7.8	1.02	9.4	1.39
8.9	1.46	10.7	1.19
9.8	1.68	12.2	1.49
9.7	1.06	12.1	1.29
9.3	1.16	10.8	1.31

- (a) Compute the mean lean body mass of the lean monkeys.
- (b) Compute the mean lean body mass of the obese monkeys.
- (c) The goal of the study is to compare the energy expended in 24 hours by the lean monkeys with that of the obese monkeys. However, animals with higher lean mass usually expend more energy. Based on your calculations in parts (a) and (b), would it make sense to simply compute the mean energy expended by lean and obese monkeys and compare the means? Explain.
- (d) To investigate how energy expended is related to body mass, make a scatterplot of energy versus mass, using different plot symbols for lean and obese monkeys.
- (e) What do the trends in your scatterplot suggest about the monkeys?

- 7.38** The number of adult Americans who smoke continues to drop. Here are estimates of the percents of adults (aged 18 and over) who were smokers in the years between 1965 and 2009.¹⁰



Year x	1965	1974	1979	1983	1987	1990	1993	1997	2000	2002	2006	2009
Smokers y	41.9	37.0	33.3	31.9	28.6	25.3	24.8	24.6	23.1	22.5	20.8	20.6

- (a) Make a scatterplot of these data.
- (b) Describe the direction, form, and strength of the relationship between percent of smokers and year. Are there any outliers?
- (c) Here are the means and standard deviations for both variables and the correlation between percent of smokers and year:

$$\bar{x} = 1990.4, s_x = 13.4$$

$$\bar{y} = 27.9, s_y = 6.8$$

$$r = -0.98$$

Use this information to find the least-squares regression line for predicting percent of smokers from year and add the line to your plot.

- (d) According to your regression line, how much did smoking decline per year during this period, on the average?
- (e) What percent of the observed variation in percent of adults who smoke can be explained by linear change over time?

- (f) One of the government's national health objectives was to reduce smoking to no more than 12% of adults by 2010. Use your regression line to predict the percent of adults who smoked in 2010. In 2009 did the regression line suggest that this health objective would be met?
- (g) Use your regression line to predict the percent of adults who will smoke in 2050. Why is your result impossible? Why was it foolish to use the regression line for this prediction?

7.39 (Optional) People who get angry easily tend to have more heart disease. That's the conclusion of a study that followed a random sample of 12,986 people from three locations for about four years. All subjects were free of heart disease at the beginning of the study. The subjects took the Spielberger Trait Anger Scale test, which measures how prone a person is to sudden anger. Here are data for the 8474 people in the sample who had normal blood pressure. CHD stands for "coronary heart disease." This includes people who had heart attacks and those who needed medical treatment for heart disease.¹¹



	Low anger	Moderate anger	High anger	Total
CHD	53	110	27	190
No CHD	3057	4621	606	8284
Total	3110	4731	633	8474

- (a) What percent of all 8474 people with normal blood pressure had CHD?
- (b) What percent of all 8474 people were classified as having high anger?
- (c) What percent of those classified as having high anger had CHD?
- (d) What percent of those with no CHD were classified as having moderate anger?
- (e) Do these data provide any evidence that as anger score increases, the percent who suffer CHD increases? Explain.

SUPPLEMENTARY EXERCISES

Supplementary exercises apply the skills you have learned in ways that require more thought or more elaborate use of technology. Some of these exercises ask you to follow the **Plan**, **Solve**, and **Conclude** steps of the four-step process introduced on page 55.

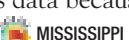
7.40 The Mississippi River. Table 7.2 gives the volume of water discharged by the Mississippi River into the Gulf of Mexico for each year from 1954 to 2001.¹² The units are cubic kilometers of water—the Mississippi is a big river.



- (a) Make a graph of the distribution of water volume. Describe the overall shape of the distribution and any outliers.
- (b) Based on the shape of the distribution, do you expect the mean to be close to the median, clearly less than the median, or clearly greater than the median? Why? Find the mean and the median to check your answer.
- (c) Based on the shape of the distribution, does it seem reasonable to use \bar{x} and s to describe the center and

spread of this distribution? Why? Find \bar{x} and s if you think they are a good choice. Otherwise, find the five-number summary.

7.41 More on the Mississippi River. The data in Table 7.2 are a time series. Make a time plot that shows how the volume of water in the Mississippi changed between 1954 and 2001. What does the time plot reveal that the histogram from the previous exercise does not? It is a good idea to always make a time plot of time series data because a histogram cannot show changes over time.



Falling through the ice. The Nenana Ice Classic is an annual contest to guess the exact time in the spring thaw when a tripod erected on the frozen Tanana River near Nenana, Alaska, will fall through the ice. The 2010 jackpot prize was \$279,030. The contest has been run since 1917. Table 7.3 gives simplified data that record only the date on which the tripod fell each year. The

TABLE 7.2 Yearly discharge (cubic kilometers of water) of the Mississippi River

YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE
1954	290	1966	410	1978	560	1990	680
1955	420	1967	460	1979	800	1991	700
1956	390	1968	510	1980	500	1992	510
1957	610	1969	560	1981	420	1993	900
1958	550	1970	540	1982	640	1994	640
1959	440	1971	480	1983	770	1995	590
1960	470	1972	600	1984	710	1996	670
1961	600	1973	880	1985	680	1997	680
1962	550	1974	710	1986	600	1998	690
1963	360	1975	670	1987	450	1999	580
1964	390	1976	420	1988	420	2000	390
1965	500	1977	430	1989	630	2001	580

TABLE 7.3 Days from April 20 for the Tanana River tripod to fall

YEAR	DAY								
1917	11	1933	19	1949	25	1965	18	1981	11
1918	22	1934	11	1950	17	1966	19	1982	21
1919	14	1935	26	1951	11	1967	15	1983	10
1920	22	1936	11	1952	23	1968	19	1984	20
1921	22	1937	23	1953	10	1969	9	1985	23
1922	23	1938	17	1954	17	1970	15	1986	19
1923	20	1939	10	1955	20	1971	19	1987	16
1924	22	1940	1	1956	12	1972	21	1988	8
1925	16	1941	14	1957	16	1973	15	1989	12
1926	7	1942	11	1958	10	1974	17	1990	5
1927	23	1943	9	1959	19	1975	21	1991	12
1928	17	1944	15	1960	13	1976	13	1992	25
1929	16	1945	27	1961	16	1977	17	1993	4
1930	19	1946	16	1962	23	1978	11	1994	10
1931	21	1947	14	1963	16	1979	11	1995	7
1932	12	1948	24	1964	31	1980	10	1996	16

earliest date so far is April 20. To make the data easier to use, the table gives the date each year in days starting with April 20. That is, April 20 is 1, April 21 is 2, and so on. Exercises 7.42 to 7.44 concern these data.¹³



7.42 When does the ice break up? We have 94 years of data on the date of ice breakup on the Tanana River. Describe the distribution of the breakup date with both a graph or graphs and appropriate numerical summaries. What is the median date (month and day) for ice breakup?



2006 Bill Watkins/Alaska Stock.com

7.43 Global warming? Because of the high stakes, the falling of the tripod has been carefully observed for many years. If the date the tripod falls has been getting earlier, that may be evidence for the effects of global warming.

- (a) Make a time plot of the date the tripod falls against year.
- (b) There is a great deal of year-to-year variation. Fitting a regression line to the data may help us see the trend. Fit the least-squares line and add it to your time plot. What do you conclude?
- (c) There is much variation about the line. Give a numerical description of how much of the year-to-year variation in ice breakup time is accounted for by the time trend represented by the regression line. (This simple example is typical of more complex evidence for the effects of global warming: large year-to-year variation requires many years of data to see a trend.)

7.44 More on global warming. Side-by-side boxplots offer a different look at the data. Group the data into periods of roughly equal length: 1917 to 1939, 1940 to 1962, 1963 to 1985, and 1986 to 2010. Make boxplots to compare ice breakup dates in these four time periods. Write a brief description of what the plots show.

7.45 Big government? The data file GDP on the text CD and Web site contains the percent of gross domestic product (GDP, the total value of all goods and services a country produces) taken by the government in 82 countries. For example, the government share of GDP is 12.28% in Canada and 10.54% in the United States.¹⁴



- (a) Make a stemplot or a histogram to display the distribution of government share of GDP.

(b) There are several high outliers. What countries are these? (In the most extreme case, the government took more than the total annual GDP!) What is the overall shape of the distribution if you ignore the outliers?

- (c) Based on your work in (b), give a numerical summary of the center and spread of the distribution, omitting the outliers.
- (d) Some Americans complain about big government and heavy taxes. Where does the United States (10.54%) stand in this international comparison?

7.46 Cicadas as fertilizer? Every 17 years, swarms of

cicadas emerge from the ground in the eastern United States, live for about six weeks, then die. (There are several “broods,” so we experience cicada eruptions more often than every 17 years.) There are so many cicadas that their dead bodies can serve as fertilizer and increase plant growth. In an experiment, a researcher added 10 cicadas under some plants in a natural plot of American bellflowers in a forest, leaving other plants undisturbed. One of the response variables was the size of seeds produced by the plants. Here are data (seed mass in milligrams) for 39 cicada plants and 33 undisturbed (control) plants:¹⁵



Alastair Shay; Papilio/CORBIS



Cicada plants				Control plants			
0.237	0.277	0.241	0.142	0.212	0.188	0.263	0.253
0.109	0.209	0.238	0.277	0.261	0.265	0.135	0.170
0.261	0.227	0.171	0.235	0.203	0.241	0.257	0.155
0.276	0.234	0.255	0.296	0.215	0.285	0.198	0.266
0.239	0.266	0.296	0.217	0.178	0.244	0.190	0.212
0.238	0.210	0.295	0.193	0.290	0.253	0.249	0.253
0.218	0.263	0.305	0.257	0.268	0.190	0.196	0.220
0.351	0.245	0.226	0.276	0.246	0.145	0.247	0.140
0.317	0.310	0.223	0.229	0.241			
0.192	0.201	0.211					

Describe and compare the two distributions. Do the data support the idea that dead cicadas can serve as fertilizer?

TABLE 7.4 Angle of deformity (degrees) for two types of foot deformity

HAV ANGLE	MA ANGLE	HAV ANGLE	MA ANGLE	HAV ANGLE	MA ANGLE
28	18	21	15	16	10
32	16	17	16	30	12
25	22	16	10	30	10
34	17	21	7	20	10
38	33	23	11	50	12
26	10	14	15	25	25
25	18	32	12	26	30
18	13	25	16	28	22
30	19	21	16	31	24
26	10	22	18	38	20
28	17	20	10	32	37
13	14	18	15	21	23
20	20	26	16		

7.47 A big-toe problem. Hallux abducto valgus (call it HAV) is a deformation of the big toe that is not common in youth and often requires surgery. Doctors used X-rays to measure the angle (in degrees) of deformity in 38 consecutive patients under the age of 21 who came to a medical center for surgery to correct HAV.¹⁶ The angle is a measure of the seriousness of the deformity. The data appear in Table 7.4 as “HAV angle.” Describe the distribution of the angle of deformity among young patients needing surgery for this condition. 

7.48 Prey attract predators. Here is one way in which nature regulates the size of animal populations: high population density attracts predators, who remove a higher proportion of the population than when the density of the prey is low. One study looked at kelp perch and their common predator, the kelp bass. The researcher set up four large circular pens on sandy ocean bottom in Southern California. He chose young perch at random from a large group and placed 10, 20, 40, and 60 perch in the four pens. Then he dropped the nets protecting the pens, allowing bass to swarm in, and counted the perch left after 2 hours. Here are data on the proportions of perch eaten in four repetitions of this setup:¹⁷ 

Perch	Proportion killed			
	10	20	40	60
0.0	0.1	0.3	0.3	
0.2	0.3	0.3	0.6	
0.075	0.3	0.6	0.725	
0.517	0.55	0.7	0.817	

Do the data support the principle that “more prey attract more predators, who drive down the number of prey”?

7.49 Predicting foot problems. Metatarsus adductus (call it MA) is a turning in of the front part of the foot that is common in adolescents and usually corrects itself. Table 7.4 gives the severity of MA (“MA angle”). Doctors speculate that the severity of MA can help predict the severity of HAV. Describe the relationship between MA and HAV. Do you think the data confirm the doctors’ speculation? Why or why not? 

7.50 Change in the Serengeti. Long-term records from the Serengeti National Park in Tanzania show interesting ecological relationships. When wildebeest are more abundant, they graze the grass more heavily, so there are fewer fires and more trees grow. Lions feed more successfully when there are more trees, so the lion population increases. Here are data on one part of this cycle, wildebeest abundance (in thousands of animals) and the percent of the grass area that burned in the same year:¹⁸ 



Gallo Image—Anthony Bannister/Getty Images

Wildebeest (1000s)	Percent burned	Wildebeest (1000s)	Percent burned	Wildebeest (1000s)	Percent burned
396	56	360	88	1147	32
476	50	444	88	1173	31
698	25	524	75	1178	24
1049	16	622	60	1253	24
1178	7	600	56	1249	53
1200	5	902	45		
1302	7	1440	21		

To what extent do these data support the claim that more wildebeest reduce the percent of grasslands that burn? How rapidly does burned area decrease as the number of wildebeest increases? Include a graph and suitable calculations.

7.51 Casting aluminum. In casting metal parts, molten metal flows through a “gate” into a die that shapes the part. The gate velocity (the speed at which metal is forced through the gate) plays a critical role in die casting. A firm that casts cylindrical aluminum pistons examined 12 types formed from the same alloy. How does the piston wall thickness (inches) influence the gate velocity (feet per second) chosen by the skilled workers who do the casting? If there is a clear pattern, it can be used to direct new workers or to automate the process. Analyze these data and report your findings.¹⁹ 

Thickness	Velocity	Thickness	Velocity	Thickness	Velocity
0.248	123.8	0.524	228.6	0.697	145.2
0.359	223.9	0.552	223.8	0.752	263.1
0.366	180.9	0.628	326.2	0.806	302.4
0.400	104.8	0.697	302.4	0.821	302.4

7.52 How are schools doing? (optional) The non-profit group Public Agenda conducted telephone interviews with parents of high school children. Interviewers chose equal numbers of black, Hispanic, and non-Hispanic white parents at random. One question asked was “Are the high schools in your state doing an excellent, good, fair or poor job, or don’t you know enough to say?” Here are the survey results.²⁰ 

Opinion	Black parents	Hispanic parents	White parents
Excellent	12	34	22
Good	69	55	81
Fair	75	61	60
Poor	24	24	24
Don’t know	22	28	14
Total	202	202	201

Write a brief analysis of these results that focuses on the relationship between parent group and opinions about schools.

7.53 Influence: hot stocks? Investment advertisements always warn that “past performance does not guarantee future results.” Here is an example that shows why you should pay attention to this warning. Stocks fell sharply in 2002, then rose sharply in 2003. The table below gives the percent returns from 23 Fidelity Investments “sector funds” in these two years. Sector funds invest in narrow segments of the stock market. They often rise and fall faster than the market as a whole.  MUTUALFUNDS

2002 return	2003 return	2002 return	2003 return	2002 return	2003 return
-17.1	23.9	-0.7	36.9	-37.8	59.4
-6.7	14.1	-5.6	27.5	-11.5	22.9
-21.1	41.8	-26.9	26.1	-0.7	36.9
-12.8	43.9	-42.0	62.7	64.3	32.1
-18.9	31.1	-47.8	68.1	-9.6	28.7
-7.7	32.3	-50.5	71.9	-11.7	29.5
-17.2	36.5	-49.5	57.0	-2.3	19.1
-11.4	30.6	-23.4	35.0		

- (a) Make a scatterplot of 2003 return (response) against 2002 return (explanatory). The funds with the best performance in 2002 tend to have the worst performance in 2003. Fidelity Gold Fund, the only fund with a positive return in both years, is an extreme outlier.

(b) To demonstrate that correlation is not resistant, find r for all 23 funds and then find r for the 22 funds other than Gold. Explain from Gold's position in your plot why omitting this point makes r more negative.

(c) Find the equations of two least-squares lines for predicting 2003 return from 2002 return, one for all 23 funds and one omitting Fidelity Gold Fund. Add both lines to your scatterplot. Starting with the least-squares idea, explain why adding Fidelity Gold Fund to the other 22 funds moves the line in the direction that your graph shows.

7.54 Influence: monkey calls. Table 7.1 (page 187) contains data on the response of 37 monkey neurons to pure-tones and to monkey calls. Figure 7.6 (page 188) is a scatterplot of these data.  MONKEYCALLS

(a) Find the least-squares line for predicting a neuron's call response from its pure-tone response. Add the line to your scatterplot. Mark on your plot the point (call it A) with the largest residual (either positive or negative) and also the point (call it B) that is an outlier in the x direction.

(b) How influential are each of these points for the correlation r ?

(c) How influential are each of these points for the regression line?

7.55 Influence: bushmeat. Table 4.2 (page 122) gives data on fish catches in a region of West Africa and the percent change in the biomass (total weight) of 41 animals in nature reserves. It appears that years with smaller fish catches see greater declines in animals, probably because local people turn to "bushmeat" when other sources of protein are not available. The next year (1999) had a fish catch of 23.0 kilograms per person and animal biomass change of -22.9% .  BUSHMEAT

(a) Make a scatterplot that shows how change in animal biomass depends on fish catch. Be sure to include the additional data point. Describe the overall pattern. The added point is a low outlier in the y direction.

(b) Find the correlation between fish catch and change in animal biomass both with and without the outlier. The outlier is influential for correlation. Explain from your plot why adding the outlier makes the correlation smaller.

(c) Find the least-squares line for predicting change in animal biomass from fish catch both with and without the additional data point for 1999. Add both lines to your scatterplot from (a). The outlier is not influential for the least-squares line. Explain from your plot why this is true.

From Exploration to Inference



The purpose of statistics is to gain understanding from data. We can seek understanding in different ways, depending on the circumstances. We have studied one approach to data, *exploratory data analysis*, in some detail. Now we move from data analysis toward *statistical inference*.

Both types of reasoning are essential to effective work with data. Here is a brief sketch of the differences between them:

EXPLORATORY DATA ANALYSIS	STATISTICAL INFERENCE
Purpose is unrestricted exploration of the data, searching for interesting patterns.	Purpose is to answer specific questions, posed before the data were produced.
Conclusions apply only to the individuals and circumstances for which we have data in hand.	Conclusions apply to a larger group of individuals or a broader class of circumstances.
Conclusions are informal, based on what we see in the data.	Conclusions are formal, backed by a statement of our confidence in them.

Our journey toward inference begins in Chapters 8 and 9, which describe statistical designs for *producing data* by samples and experiments. The conclusions of inference use the language of *probability*, the mathematics of chance. Chapters 10 and 11 present the ideas we need, and the optional Chapters 12 and 13 add more detail. Armed with designs for producing trustworthy data, data analysis techniques for examining the data, and the language of probability, we are prepared to understand the big ideas of inference in Chapters 14, 15, and 16. These chapters are the foundation for the discussion of inference in practice that occupies the rest of the book.

PRODUCING DATA

CHAPTER 8 Producing Data:
Sampling

CHAPTER 9 Producing Data:
Experiments

COMMENTARY: Data Ethics*

PROBABILITY AND SAMPLING DISTRIBUTIONS

CHAPTER 10 Introducing Probability

CHAPTER 11 Sampling Distributions

CHAPTER 12 General Rules of Probability*

CHAPTER 13 Binomial Distributions*

FOUNDATIONS OF INFERENCE

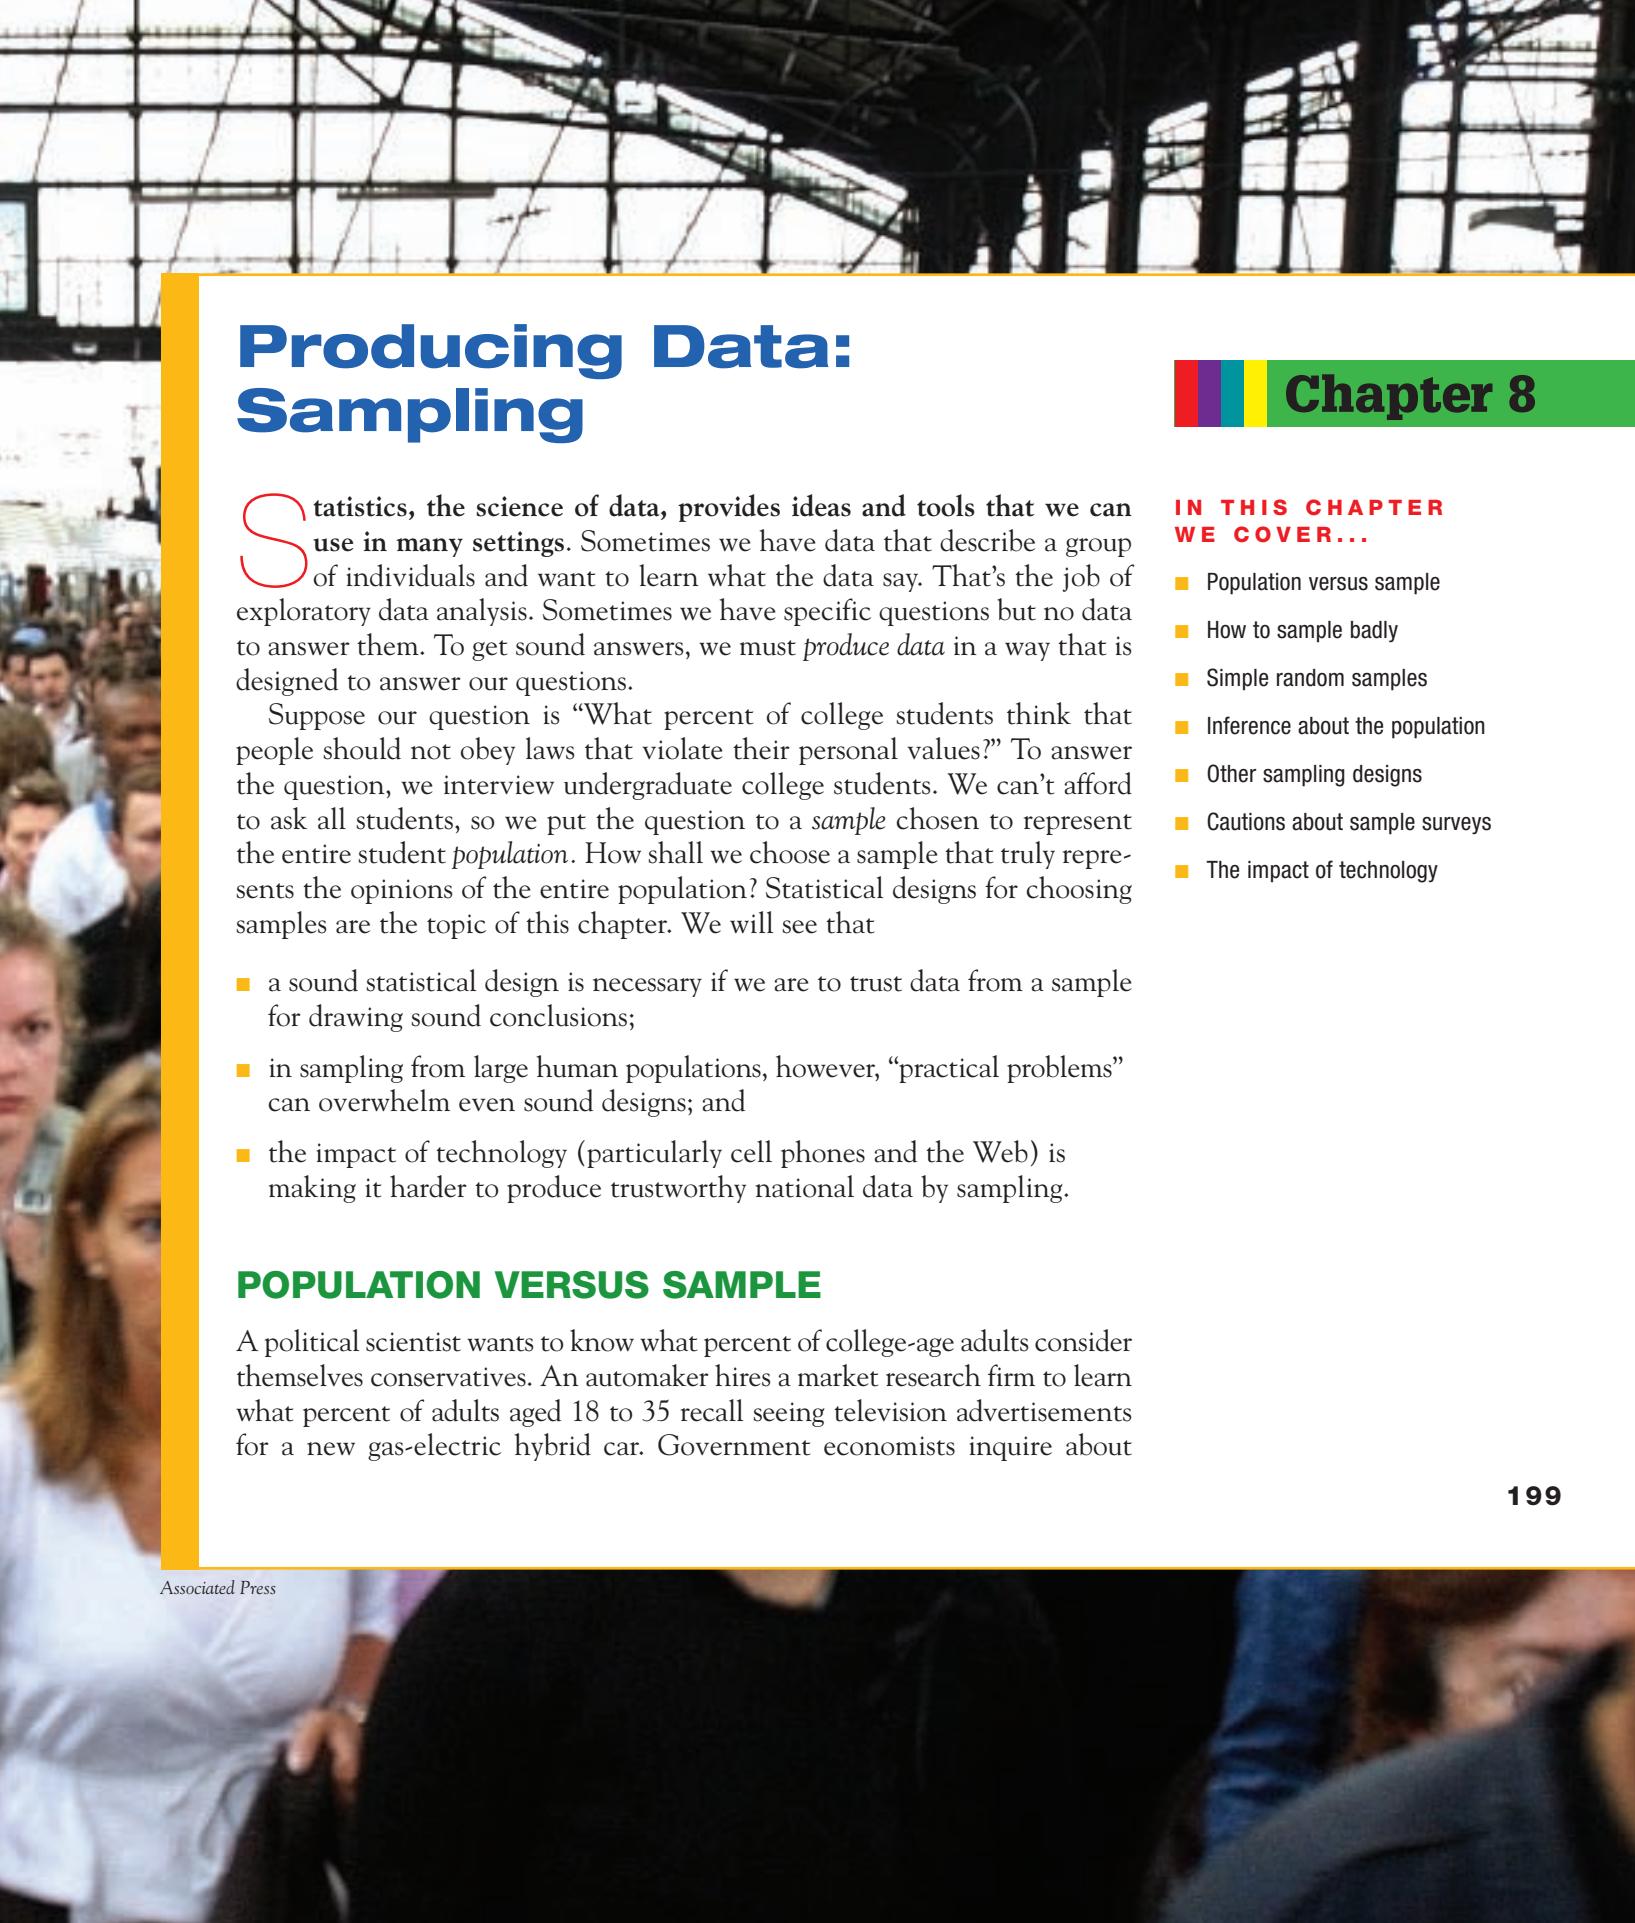
CHAPTER 14 Confidence Intervals: The Basics

CHAPTER 15 Tests of Significance: The Basics

CHAPTER 16 Inference in Practice

CHAPTER 17 From Exploration to Inference: Part II Review





Producing Data: Sampling



Chapter 8

Statistics, the science of data, provides ideas and tools that we can use in many settings. Sometimes we have data that describe a group of individuals and want to learn what the data say. That's the job of exploratory data analysis. Sometimes we have specific questions but no data to answer them. To get sound answers, we must *produce data* in a way that is designed to answer our questions.

Suppose our question is “What percent of college students think that people should not obey laws that violate their personal values?” To answer the question, we interview undergraduate college students. We can't afford to ask all students, so we put the question to a *sample* chosen to represent the entire student *population*. How shall we choose a sample that truly represents the opinions of the entire population? Statistical designs for choosing samples are the topic of this chapter. We will see that

- a sound statistical design is necessary if we are to trust data from a sample for drawing sound conclusions;
- in sampling from large human populations, however, “practical problems” can overwhelm even sound designs; and
- the impact of technology (particularly cell phones and the Web) is making it harder to produce trustworthy national data by sampling.

POPULATION VERSUS SAMPLE

A political scientist wants to know what percent of college-age adults consider themselves conservatives. An automaker hires a market research firm to learn what percent of adults aged 18 to 35 recall seeing television advertisements for a new gas-electric hybrid car. Government economists inquire about

IN THIS CHAPTER WE COVER...

- Population versus sample
- How to sample badly
- Simple random samples
- Inference about the population
- Other sampling designs
- Cautions about sample surveys
- The impact of technology

average household income. In all these cases, we want to gather information about a large group of individuals. Time, cost, and inconvenience forbid contacting every individual. So we gather information about only part of the group in order to draw conclusions about the whole.

POPULATION, SAMPLE, SAMPLING DESIGN

The **population** in a statistical study is the entire group of individuals about which we want information.

A **sample** is a part of the population from which we actually collect information. We use a sample to draw conclusions about the entire population.

A **sampling design** describes exactly how to choose a sample from the population.

Pay careful attention to the details of the definitions of “population” and “sample.” Look at Exercise 8.1 right now to check your understanding.

We often draw conclusions about a whole on the basis of a sample. Everyone has tasted a sample of ice cream and ordered a cone on the basis of that taste. But ice cream is uniform, so that the single taste represents the whole. Choosing a representative sample from a large and varied population is not so easy. The first step in planning a **sample survey** is to say exactly *what population* we want to describe. The second step is to say exactly *what we want to measure*, that is, to give exact definitions of our variables. These preliminary steps can be complicated, as the following example illustrates.

EXAMPLE 8.1 The Current Population Survey

The most important government sample survey in the United States is the monthly Current Population Survey (CPS) conducted by the Bureau of the Census for the Bureau of Labor Statistics. The CPS contacts about 60,000 households each month. It produces the monthly unemployment rate and much other economic and social information. (See Figure 8.1.) To measure unemployment, we must first specify the population we want to describe. Which age groups will we include? Will we include illegal immigrants or people in prisons? The CPS defines its population as all U.S. residents (legal or not) 16 years of age and over who are civilians and are not in an institution such as a prison. The unemployment rate announced in the news refers to this specific population.

The second question is harder: what does it mean to be “unemployed”? Someone who is not looking for work—for example, a full-time student—should not be called unemployed just because she is not working for pay. If you are chosen for the CPS sample, the interviewer first asks whether you are available to work and whether you actually looked for work in the past four weeks. If not, you are neither employed nor unemployed—you are not in the labor force. So discouraged workers who haven’t looked for a job in four weeks are excluded from the count.

If you are in the labor force, the interviewer goes on to ask about employment. If you did any work for pay or in your own business during the week of the survey, you

**FIGURE 8.1**

The home page of the Current Population Survey at the Bureau of Labor Statistics.

are employed. If you worked at least 15 hours in a family business without pay, you are employed. You are also employed if you have a job but didn't work because of vacation, being on strike, or other good reason. An unemployment rate of 6.7% means that 6.7% of the sample was unemployed, using the exact CPS definitions of both "labor force" and "unemployed." ■

The final step in planning a sample survey is the sampling design. We will now introduce basic statistical designs for sampling.

APPLY YOUR KNOWLEDGE

8.1 Sampling students. A political scientist wants to know how college students feel about the Social Security system. She obtains a list of the 3456 undergraduates at her college and mails a questionnaire to 250 students selected at random. Only 104 questionnaires are returned.

- What is the population in this study? Be careful: what group does she *want information about*?
- What is the sample? Be careful: from what group does she *actually obtain information*?

The important message in this problem is that the sample can redefine the population about which information is obtained.



John ELK III/Alamy

8.2 Student archaeologists. An archaeological dig turns up large numbers of pottery shards, broken stone implements, and other artifacts. Students working on the project classify each artifact and assign it a number. The counts in different categories are important for understanding the site, so the project director chooses 2% of the artifacts at random and checks the students' work. What are the population and the sample here?

8.3 Customer satisfaction. A department store mails a customer satisfaction survey to people who make credit card purchases at the store. This month, 45,000 people made credit card purchases. Surveys are mailed to 1000 of these people, chosen at random, and 137 people return the survey form.

- (a) What is the population of interest for this survey?
- (b) What is the sample? From what group is information actually obtained?

HOW TO SAMPLE BADLY

convenience sample

How can we choose a sample that we can trust to represent the population? A sampling design is a specific method for choosing a sample from the population. The easiest—but not the best—design just chooses individuals close at hand. If we are interested in finding out how many people have jobs, for example, we might go to a shopping mall and ask people passing by if they are employed. A sample selected by taking the members of the population that are easiest to reach is called a **convenience sample**. Convenience samples often produce unrepresentative data.

EXAMPLE 8.2 Sampling at the mall

A sample of mall shoppers is fast and cheap. But people at shopping malls tend to be more prosperous than typical Americans. They are also more likely to be teenagers or retired. Moreover, unless interviewers are carefully trained, they tend to question well-dressed, respectable-looking people and avoid poorly dressed or tough-looking individuals. The type of people at the mall will also vary by time of day and day of week. In short, mall interviews will not contact a sample that is representative of the entire population. ■

Interviews at shopping malls will almost surely overrepresent middle-class and retired people and underrepresent the poor. This will happen almost every time we take such a sample. That is, it is a systematic error caused by a bad sampling design, not just bad luck on one sample. This is *bias*: the outcomes of mall surveys will repeatedly miss the truth about the population in the same ways.

BIAS

The design of a statistical study is **biased** if it systematically favors certain outcomes.

EXAMPLE 8.3 Online polls

Former CNN evening commentator Lou Dobbs doesn't like illegal immigration. One of his broadcasts in 2007 was largely devoted to attacking a proposal by the governor of New York State to offer drivers' licenses to illegal immigrants as a public safety measure. During the show, Mr. Dobbs invited his viewers to go to loudobbs.com to vote on the question "Would you be more or less likely to vote for a presidential candidate who supports giving drivers' licenses to illegal aliens?" We aren't surprised that 97% of the 7350 people who voted by the end of the broadcast said, "Less likely." ■

The loudobbs.com poll was biased because people chose whether or not to participate. Most who voted were viewers of Lou Dobbs's program who had just heard him denounce the governor's idea. *People who take the trouble to respond to an open invitation are usually not representative of any clearly defined population.* That's true of the people who bother to respond to write-in, call-in, or online polls in general. Polls like these are examples of *voluntary response sampling*.



VOLUNTARY RESPONSE SAMPLE

A **voluntary response sample** consists of people who choose themselves by responding to a broad appeal. Voluntary response samples are biased because people with strong opinions are most likely to respond.

APPLY YOUR KNOWLEDGE

8.4 Sampling on campus. You see a woman student standing in front of the student center, now and then stopping other students to ask them questions. She says that she is collecting student opinions for a class assignment. Explain why this sampling method is almost certainly biased.

8.5 More sampling on campus. You would like to start a club for psychology majors on campus, and you are interested in finding out what proportion of psychology majors would join. The dues would be \$35 and used to pay for speakers to come to campus. You ask 5 psychology majors from your senior psychology honors seminar whether they would be interested in joining this club and find that 4 of the 5 students questioned are interested. Is this sampling method biased, and, if so, what is the likely direction of bias?

SIMPLE RANDOM SAMPLES

In a voluntary response sample, people choose whether to respond. In a convenience sample, the interviewer makes the choice. In both cases, personal choice produces bias. The statistician's remedy is to allow impersonal chance to choose the sample. A sample chosen by chance rules out both favoritism by the sampler

and self-selection by respondents. Choosing a sample by chance attacks bias by giving all individuals an equal chance to be chosen. Rich and poor, young and old, black and white, all have the same chance to be in the sample.

The simplest way to use chance to select a sample is to place names in a hat (the population) and draw out a handful (the sample). This is the idea of *simple random sampling*. Although the idea of drawing names from a hat is a good way to conceptualize a simple random sample, it is generally not a good method for obtaining a simple random sample. Writing names on slips of paper can lead to bias if the slips of paper are not well mixed or there is a tendency to select them from, say, the top or bottom. Drawing names from a hat would be particularly difficult to implement if the population size is large, possibly requiring thousands of slips of paper.

SIMPLE RANDOM SAMPLE

A **simple random sample (SRS)** of size n consists of n individuals from the population chosen in such a way that every set of n individuals has an equal chance to be the sample actually selected.

An SRS not only gives each individual an equal chance to be chosen but also gives every possible sample an equal chance to be chosen. There are other random sampling designs that give each individual, but not each sample, an equal chance. Exercise 8.43 (page 219) describes one such design.

When you think of an SRS, you can still picture the conceptual situation of drawing names from a hat to remind yourself that an SRS doesn't favor any part of the population. That's why an SRS is a better method of choosing samples than convenience or voluntary response sampling. But writing names on slips of paper, mixing them well, and drawing them from a hat is slow and inconvenient. That's especially true if, like the Current Population Survey, we must draw a sample of size 60,000. In practice, samplers use software. The *Simple Random Sample* applet makes choosing an SRS very fast. If you don't use the applet or other software, you can randomize by using a *table of random digits*. In fact, software for choosing samples starts by generating random digits, so using a table just does by hand what the software does more quickly.

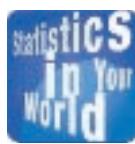


RANDOM DIGITS

A **table of random digits** is a long string of the digits 0, 1, 2, 3, 4, 5, 6, 7, 8, 9 with these two properties:

1. Each entry in the table is equally likely to be any of the 10 digits 0 through 9.
2. The entries are independent of each other. That is, knowledge of one part of the table gives no information about any other part.

Table B at the back of the book is a table of random digits. Table B begins with the digits 19223950340575628713. To make the table easier to read, the digits appear in groups of five and in numbered rows. The groups and rows have no meaning—the table is just a long list of randomly chosen digits. There are two steps in using the table to choose a simple random sample.



Are these random digits really random?

Not a chance.

The random

digits in Table B were produced by a computer program. Computer programs do exactly what you tell them to do. Give the program the same input and it will produce exactly the same “random” digits. Of course, clever people have devised computer programs that produce output that looks like random digits. These are called “pseudo-random numbers,” and that’s what Table B contains. Pseudo-random numbers work fine for statistical randomizing, but they have hidden nonrandom patterns that can mess up more refined uses.

USING TABLE B TO CHOOSE AN SRS

Label: Give each member of the population a numerical label of the *same length*.

Table: To choose an SRS, read from Table B successive groups of digits of the length you used as labels. Your sample contains the individuals whose labels you find in the table.

You can label up to 100 items with two digits: 01, 02, ..., 99, 00. Up to 1000 items can be labeled with three digits, and so on. Always use the shortest labels that will cover your population. As standard practice, we recommend that you begin with label 1 (or 01 or 001, as needed). Reading groups of digits from the table gives all individuals the same chance to be chosen because all labels of the same length have the same chance to be found in the table. For example, any pair of digits in the table is equally likely to be any of the 100 possible labels 01, 02, ..., 99, 00. Ignore any group of digits that was not used as a label or that duplicates a label already in the sample. You can read digits from Table B in any order—across a row, down a column, and so on—because the table has no order. As standard practice, we recommend reading across rows.

EXAMPLE 8.4 Sampling spring break resorts

A campus newspaper plans a major article on spring break destinations. The authors intend to call 4 randomly chosen resorts at each destination to ask about their attitudes toward groups of students as guests. Here are the resorts listed in one city:

01	Aloha Kai	08	Captiva	15	Palm Tree	22	Sea Shell
02	Anchor Down	09	Casa del Mar	16	Radisson	23	Silver Beach
03	Banana Bay	10	Coconuts	17	Ramada	24	Sunset Beach
04	Banyan Tree	11	Diplomat	18	Sandpiper	25	Tradewinds
05	Beach Castle	12	Holiday Inn	19	Sea Castle	26	Tropical Breeze
06	Best Western	13	Lime Tree	20	Sea Club	27	Tropical Shores
07	Cabana	14	Outrigger	21	Sea Grape	28	Veranda

Label: Because two digits are needed to label the 28 resorts, all labels will have two digits. We have added labels 01 to 28 in the list of resorts. Always say how you labeled the members of the population. To sample from the 1240 resorts in a major vacation area, you would label the resorts 0001, 0002, ..., 1239, 1240.

Table: To use the *Simple Random Sample* applet, just enter 28 in the “Population =” box and 4 in the “Select a sample” box, click “Reset,” and click “Sample.” Figure 8.2 shows the result of one sample.

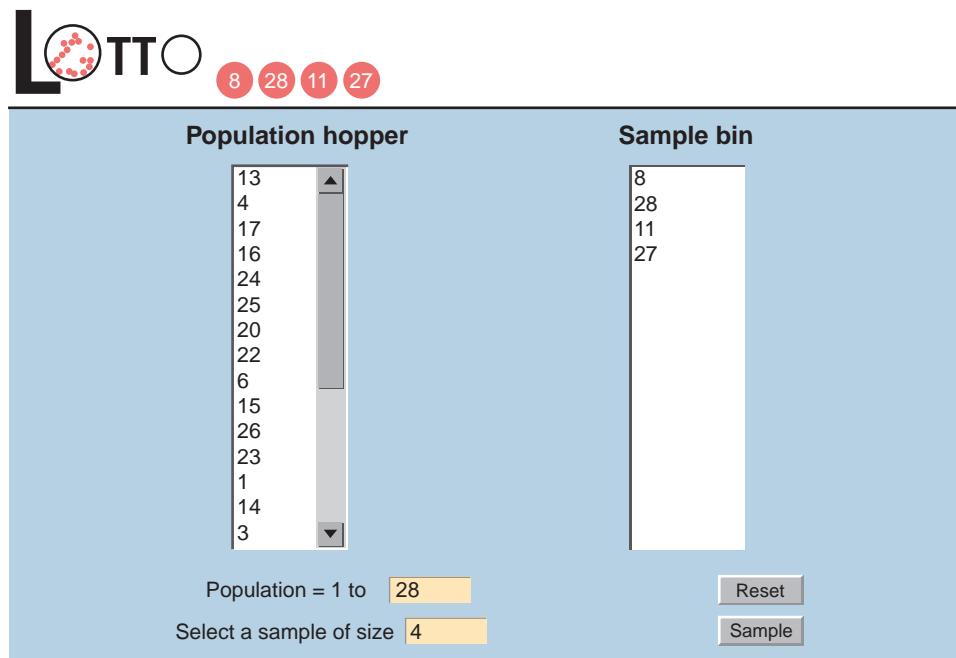


Robert Daly/Getty Images



FIGURE 8.2

The *Simple Random Sample* applet used to choose an SRS of size $n = 4$ from a population of size 28.



To use Table B, read two-digit groups until you have chosen four resorts. Starting at line 130 (any line will do), we find

69051 64817 87174 09517 84534 06489 87201 97245

Because the labels are two digits long, read successive two-digit groups from the table; so the first three two-digit groups here are 69, 05, and 16. Ignore groups not used as labels, like the initial 69. Also ignore any repeated labels, like the second and third 17s in this row, because you can't choose the same resort twice. Your sample contains the resorts labeled 05, 16, 17, and 20. These are Beach Castle, Radisson, Ramada, and Sea Club. ■

We can trust results from an SRS, as well as from other types of random samples that we will meet later, because the use of impersonal chance avoids bias. Online polls and mall interviews also produce samples. We can't trust results from these samples, because they are chosen in ways that invite bias. *The first question to ask about any sample is whether it was chosen at random.*



EXAMPLE 8.5 Texting while driving

"Do you think sending a text message while driving, either on a cell phone or other electronic device, should be legal or illegal?" When the *New York Times* and CBS News asked this question of 829 adults in October 2009, 97% said "illegal" and just 1% said "legal." Can we trust the opinions of this sample to fairly represent the opinions of all adults? Here's part of the statement by the *Times* on how the poll was conducted:



iStockphoto

The latest New York Times/CBS News poll is based on telephone interviews conducted October 5 through October 8 with 829 adults throughout the United States.

The sample of land line telephone exchanges called was randomly selected by a computer from a complete list of more than 69,000 active residential exchanges across the country. The exchanges were chosen so as to ensure that each region of the country was represented in proportion to its population.

Within each exchange, random digits were added to form a complete telephone number, thus permitting access to listed and unlisted numbers alike. Within each household, one adult was designated by a random procedure to be the respondent for the survey.¹

This is a good description of the most common method for choosing national samples, called **random digit dialing**. We'll come back to random digit dialing and its problems later (see Exercise 8.16), but this statement is a good start toward gaining our confidence. We know the size of the sample, when the poll was taken, and the comforting word "random" appears three times. ■

random digit dialing

APPLY YOUR KNOWLEDGE

- 8.6 Apartment living.** You are planning a report on apartment living in a college town. You decide to select four apartment complexes at random for in-depth interviews with residents. Use the *Simple Random Sample* applet, other software, or Table B to select a simple random sample of four of the following apartment complexes. If you use Table B, start at line 122.

Ashley Oaks	Country View	Mayfair Village
Bay Pointe	Country Villa	Nobb Hill
Beau Jardin	Crestview	Pemberly Courts
Bluffs	Del-Lynn	Peppermill
Brandon Place	Fairington	Pheasant Run
Briarwood	Fairway Knolls	River Walk
Brownstone	Fowler	Sagamore Ridge
Burberry Place	Franklin Park	Salem Courthouse
Cambridge	Georgetown	Village Square
Chauncey Village	Greenacres	Waterford Court

- 8.7 Minority managers.** A firm wants to understand the attitudes of its minority managers toward its system for assessing management performance. Below is a list of all the firm's managers who are members of minority groups. Use the *Simple Random Sample* applet, other software, or Table B at line 134 to choose five to be interviewed in detail about the performance appraisal system.

Adelaja	Draguljic	Huo	Modur
Ahmadiani	Fernandez	Ippolito	Rettiganti
Barnes	Fox	Jiang	Rodriguez
Bonds	Gao	Jung	Sanchez
Burke	Gemayel	Mani	Sgambellone
Deis	Gupta	Mazzeo	Yajima
Ding	Hernandez		

- 8.8 Sampling gravestones.** The local genealogical society in Coles County, Illinois, has compiled records on all 55,914 gravestones in cemeteries in the county for the years 1825 to 1985. Historians plan to use these records to learn about African



© The Photo Works

Americans in Coles County's history. They first choose an SRS of 395 records to check their accuracy by visiting the actual gravestones.²

- (a) How would you label the 55,914 records?
- (b) Use Table B, beginning at line 120, to choose the first 5 records for the SRS.

inference

INFERENCE ABOUT THE POPULATION

The purpose of a sample is to give us information about a larger population. The process of drawing conclusions about a population on the basis of sample data is called **inference** because we *infer* information about the population from what we know about the sample.

Inference from convenience samples or voluntary response samples would be misleading because these methods of choosing a sample are biased. We are almost certain that the sample does not fairly represent the population. *The first reason to rely on random sampling is to eliminate bias in selecting samples from the list of available individuals.*

Nonetheless, it is unlikely that results from a random sample are exactly the same as for the entire population. Sample results, like the unemployment rate obtained from the monthly Current Population Survey, are only estimates of the truth about the population. If we select two samples at random from the same population, we will almost certainly draw different individuals. So the sample results will differ somewhat, just by chance. Properly designed samples avoid systematic bias, but their results are rarely exactly correct and they vary from sample to sample.

Why can we trust random samples? The big idea is that the results of random sampling don't change haphazardly from sample to sample. Because we deliberately use chance, the results obey the laws of probability that govern chance behavior. These laws allow us to say how likely it is that sample results are close to the truth about the population. *The second reason to use random sampling is that the laws of probability allow trustworthy inference about the population.* Results from random samples come with a margin of error that sets bounds on the size of the likely error. How to do this is part of the technique of statistical inference. We will describe the reasoning in Chapter 14 and present details throughout the rest of the book.

One point is worth making now: *larger random samples give more accurate results than smaller random samples.* By taking a very large sample, you can be confident that the sample result is very close to the truth about the population. The Current Population Survey contacts about 60,000 households, so it estimates the national unemployment rate very accurately. Opinion polls that contact 1000 or 1500 people give less accurate results. Of course, only samples chosen by chance carry this guarantee. Lou Dobbs's online sample tells us little about overall American public opinion even though 7350 people clicked a response.



APPLY YOUR KNOWLEDGE

- 8.9 Ask more people.** Just before a presidential election, a national opinion-polling firm increases the size of its weekly sample from the usual 1500 people to 4000 people. Why do you think the firm does this?

8.10 Sampling Pentecostals. Pentecostals are among the fastest-growing Christian groups in many countries. The Pew Forum on Religion and Public Life surveyed Pentecostal Christians in 10 countries and compared their opinions with those of the general population. In South Korea, random samples by Gallup Korea had margins of error (we will give more detail in later chapters) of $\pm 4\%$ for the general public and $\pm 9\%$ for Pentecostals.³ What do you think explains the fact that estimates for Pentecostals were less precise?

OTHER SAMPLING DESIGNS

Random sampling, the use of chance to select the sample, is the essential principle of statistical sampling. Designs for random sampling from large populations spread out over a wide area are usually more complex than an SRS. For example, it is common to sample important groups within the population separately, then combine these samples. This is the idea of a *stratified random sample*.

STRATIFIED RANDOM SAMPLE

To select a **stratified random sample**, first classify the population into groups of similar individuals, called **strata**. Then choose a separate SRS in each stratum and combine these SRSs to form the full sample.

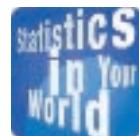
Choose the strata based on facts known before the sample is taken. For example, a population of election districts might be divided into urban, suburban, and rural strata. A stratified design can produce more precise information than an SRS of the same size by taking advantage of the fact that individuals in the same stratum are similar to one another.

EXAMPLE 8.6 Seat belt use in Hawaii

Each state conducts an annual survey of seat belt use by drivers, following guidelines set by the federal government. The guidelines require random sampling. Seat belt use is observed at randomly chosen road locations at random times during daylight hours. The locations are not an SRS of all locations in the state but rather a stratified sample using the state's counties as strata.

In Hawaii, the counties are the islands that make up the state's territory. The seat belt survey sample consists of 135 road locations in the four most populated islands: 66 in Oahu, 24 in Maui, 23 in Hawaii, and 22 in Kauai. The sample sizes on the islands are proportional to the amount of road traffic.⁴

Most large-scale sample surveys use **multistage samples**. For example, the opinion poll described in Example 8.5 has three stages: choose a random sample of telephone exchanges (stratified by region of the country), then an SRS of household telephone numbers within each exchange, then a random adult in each household.



Golfing at random

Random drawings give everyone the same chance to be chosen, so they offer a fair way to decide who gets a scarce good—like a round of golf. Lots of golfers want to play the famous Old Course at St. Andrews, Scotland. Some can reserve in advance, at considerable expense. Most must hope that chance favors them in the daily random drawing for tee times. At the height of the summer season, only 1 in 6 wins the right to pay \$250 for a round.



Ryan McVay/Photo Disc/Getty Images

multistage sample

Analysis of data from sampling designs more complex than an SRS takes us beyond basic statistics. But the SRS is the building block of more elaborate designs, and analysis of other designs differs more in complexity of detail than in fundamental concepts.

APPLY YOUR KNOWLEDGE

8.11 Sampling metro Chicago. Cook County, Illinois, has the second-largest population of any county in the United States (after Los Angeles County, California). Cook County has 30 suburban townships and an additional 8 townships that make up the city of Chicago. The suburban townships are

Barrington	Elk Grove	Maine	Orland	Riverside
Berwyn	Evanston	New Trier	Palatine	Schaumburg
Bloom	Hanover	Niles	Palos	Stickney
Bremen	Lemont	Northfield	Proviso	Thornton
Calumet	Leyden	Norwood Park	Rich	Wheeling
Cicero	Lyons	Oak Park	River Forest	Worth

The Chicago townships are

Hyde Park	Lake	North Chicago	South Chicago
Jefferson	Lake View	Rogers Park	West Chicago

Because city and suburban areas may differ, the first stage of a multistage sample chooses a stratified sample of 5 suburban townships and 3 of the more heavily populated Chicago townships. Use Table B or software to choose this sample. (If you use Table B, assign labels in alphabetical order and start at line 105 for the suburbs and at line 115 for Chicago.)

8.12 Academic dishonesty. A study of academic dishonesty among college students used a two-stage sampling design. The first stage chose a sample of 30 colleges and universities. Then the study authors mailed questionnaires to a stratified sample of 200 seniors, 100 juniors, and 100 sophomores at each school.⁵ One of the schools chosen has 898 sophomores, 943 juniors, and 895 seniors. You have alphabetical lists of the students in each class. Explain how you would assign labels for stratified sampling. Then use software or Table B, starting at line 122, to select the first 5 students in the sample from each stratum. After selecting 5 students for a stratum, continue to select the students for the next stratum from where you left off in the table.

CAUTIONS ABOUT SAMPLE SURVEYS

Random selection eliminates bias in the choice of a sample from a list of the population. When the population consists of human beings, however, accurate information from a sample requires more than a good sampling design.

To begin, we need an accurate and complete list of the population. Because such a list is rarely available, most samples suffer from some degree of *undercoverage*. A sample survey of households, for example, will miss not only homeless people but prison inmates and students in dormitories. An opinion poll conducted by calling

landline telephone numbers will miss households that have only cell phones as well as households without a phone. The results of national sample surveys therefore have some bias if the people not covered differ from the rest of the population.

A more serious source of bias in most sample surveys is *nonresponse*, which occurs when a selected individual cannot be contacted or refuses to cooperate. Nonresponse to sample surveys often exceeds 50%, even with careful planning and several callbacks. Because nonresponse is higher in urban areas, most sample surveys substitute other people in the same area to avoid favoring rural areas in the final sample. If the people contacted differ from those who are rarely at home or who refuse to answer questions, some bias remains.

UNDERCOVERAGE AND NONRESPONSE

Undercoverage occurs when some groups in the population are left out of the process of choosing the sample.

Nonresponse occurs when an individual chosen for the sample can't be contacted or refuses to participate.

EXAMPLE 8.7 How bad is nonresponse?

The U.S. Census Bureau's American Community Survey (ACS) has the lowest nonresponse rate of any poll we know: only about 1% of the households in the sample refuse to respond; the overall nonresponse rate, including "never at home" and other causes, is just 2.5%.⁶ This monthly survey of about 250,000 households replaces the "long form" that in the past was sent to some households in the every-ten-years national census. Participation in the ACS is mandatory, and the U.S. Census Bureau follows up by telephone and then in person if a household fails to return the mail questionnaire.

The University of Chicago's General Social Survey (GSS) is the nation's most important social science survey. (See Figure 8.3.) The GSS contacts its sample in person, and it is run by a university. Despite these advantages, a recent survey had a 30% rate of nonresponse.

What about opinion polls by news media and opinion-polling firms? We don't know their rates of nonresponse because they won't say. That itself is a bad sign. The Pew



FIGURE 8.3

The home page of the General Social Survey at the University of Chicago's National Opinion Research Center. The GSS has tracked opinions about a wide variety of issues since 1972.

Research Center for the People and the Press imitated a careful random digit dialing survey and published the results: over 5 days, the survey reached 76% of the households in its chosen sample, but “because of busy schedules, skepticism and outright refusals, interviews were completed in just 38% of households that were reached.” Combining households that could not be contacted with those who did not complete the interview gave a nonresponse rate of 73%.⁷ ■

In addition, the behavior of the respondent or of the interviewer can cause **response bias** in sample results. People know that they should take the trouble to vote, for example, so many who didn’t vote in the last election will tell an interviewer that they did. The race or sex of the interviewer can influence responses to questions about race relations or attitudes toward feminism. Answers to questions that ask respondents to recall past events are often inaccurate because of faulty memory. For example, many people “telescope” events in the past, bringing them forward in memory to more recent time periods. “Have you visited a dentist in the last 6 months?” will often draw a “Yes” from someone who last visited a dentist 8 months ago.⁸ Careful training of interviewers and careful supervision to avoid variation among the interviewers can reduce response bias. Good interviewing technique is another aspect of a well-done sample survey.

wording effects

The **wording of questions** is the most important influence on the answers given to a sample survey. Confusing or leading questions can introduce strong bias, and changes in wording can greatly change a survey’s outcome. Even the order in which questions are asked matters. Here are some examples.⁹

EXAMPLE 8.8 What was that question?

How do Americans feel about illegal immigrants? “Should illegal immigrants be prosecuted and deported for being in the U.S. illegally, or shouldn’t they?” Asked this question in an opinion poll, 69% favored deportation. But when the very same sample was asked whether illegal immigrants who have worked in the United States for two years “should be given a chance to keep their jobs and eventually apply for legal status,” 62% said that they should. Different questions give quite different impressions of attitudes toward illegal immigrants.

What about government help for the poor? Only 13% think we are spending too much on “assistance to the poor,” but 44% think we are spending too much on “welfare.” ■

EXAMPLE 8.9 Are you happy?

Ask a sample of college students these two questions:

- “How happy are you with your life in general?” (Answers on a scale of 1 to 5)
- “How many dates did you have last month?”

The correlation between answers is $r = -0.012$ when asked in this order. It appears that dating has little to do with happiness. Reverse the order of the questions, however, and $r = 0.66$. Asking a question that brings dating to mind makes dating success a big factor in happiness. ■

Don't trust the results of a sample survey until you have read the exact questions asked. The amount of nonresponse and the date of the survey are also important. Good statistical design is a part, but only a part, of a trustworthy survey.



APPLY YOUR KNOWLEDGE

8.13 Ring-no-answer. A common form of nonresponse in telephone surveys is “ring-no-answer.” That is, a call is made to an active number but no one answers. The Italian National Statistical Institute looked at nonresponse to a government survey of households in Italy during the periods January 1 to Easter and July 1 to August 31. All calls were made between 7 and 10 P.M., but 21.4% gave “ring-no-answer” in one period versus 41.5% “ring-no-answer” in the other period.¹⁰ Which period do you think had the higher rate of no answers? Why? Explain why a high rate of nonresponse makes sample results less reliable.

8.14 Gays in the military. In 2010, a Quinnipiac University Poll and a CNN Poll each asked a nationwide sample about their views on openly gay men and women serving in the military.¹¹ Here are the two questions:

Question A: *Federal law currently prohibits openly gay men and women from serving in the military. Do you think this law should be repealed or not?*

Question B: *Do you think people who are openly gay or homosexual should or should not be allowed to serve in the U.S. military?*

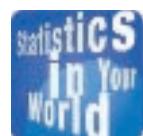
One of these questions had 78% responding “should,” and the other question had only 57% responding “should.” Which wording is slanted toward a more negative response on gays in the military? Why?

THE IMPACT OF TECHNOLOGY

A few national sample surveys, including the General Social Survey and the government’s American Community Survey and Current Population Survey, interview some or all of their subjects in person. This is expensive and time-consuming, so most national surveys contact subjects by telephone using the random digit dialing (RDD) method described in Example 8.5 (page 206). Technology, especially the spread of cell phones, is making traditional RDD methods outdated.

First, *call screening* is now common. A large majority of American households have answering machines, voice mail, or caller ID, and many use these methods to screen their calls. Calls from polling organizations are rarely returned.

More seriously, the number of *cell-phone-only households* is increasing rapidly. By mid-2007, 14% of American households had a cell phone but no landline phone, and by the end of 2009 that number had increased to almost 25%. Even if the United States and Canada don’t approach the 61% of households in Finland that have no landline phone, it’s clear that RDD reaching only landline numbers is in trouble. Can surveys just add cell phone numbers? Not easily. Federal regulations prohibit automated dialing to cell phones, which rules out computerized RDD sampling and requires hand dialing of cell phone numbers, which is expensive.



Do not call!

People who do sample surveys hate telemarketing.

We all get so many unwanted sales pitches by phone that many people hang up before learning that the caller is conducting a survey rather than selling vinyl siding. You can eliminate calls from commercial telemarketers by placing your phone number on the National Do Not Call Registry. Just go to www.donotcall.gov to sign up.

A cell phone can be anywhere, and many people keep their cell number despite moving, so stratifying by location becomes difficult. And a cell phone user may be driving or otherwise unable to talk safely.

People who screen calls and people who have only a cell phone tend to be younger than the general population. By the end of 2009 almost half of adults aged 25 to 29 years (48.6%) lived in households with no landline phone. So RDD surveys using only landlines may be biased (see Exercise 8.16). Careful surveys weight their responses to reduce bias. For example, if a sample contains too few young adults, the responses of the young adults who do respond are given extra weight. But with response rates steadily dropping and cell-phone-only use steadily growing, the future of RDD landline telephone surveys is not promising. Some polling organizations include a minimum quota of cell phone users in their samples to help adjust for bias¹² (see Exercise 8.45).

One alternative is to use *Web surveys* rather than telephone surveys. It's important to distinguish professional Web surveys from the overwhelming number of voluntary response online surveys that are intended to be entertaining rather than to give trustworthy information about a clearly defined population. Undercoverage is a serious problem for even careful Web surveys because about a quarter of Americans lack Internet access and only about half have broadband access. People without Internet access are more likely to be poor, elderly, minority, or rural than the overall population, so the potential for bias in a Web survey is clear. There is no easy way to choose a random sample even from people with Web access, because there is no technology that generates personal email addresses at random in the way that RDD generates residential telephone numbers. Even if such technology existed, etiquette and regulations aimed at spammers would prevent mass emailing. For the present, Web surveys work well only for restricted populations, for example, surveying students at your university using the school's list of student email addresses.¹³ Here is an example of a successful Web survey.

EXAMPLE 8.10 Doctors and placebos

A placebo is a dummy treatment such as a pill that has no direct effect on a patient but may bring about a response because patients expect it to. Do academic physicians who maintain private practices sometimes give their patients placebos? A Web survey of doctors in internal medicine departments at Chicago-area medical schools was possible because almost all the doctors had listed email addresses.

Send an email to each doctor explaining the purpose of the study, promising anonymity, and giving an individual Web link for response. In all, 231 of 443 doctors responded. The response rate was helped by the fact that the email came from a team at a medical school. Result: 45% said they sometimes used placebos in their clinical practice.¹⁴ ■

APPLY YOUR KNOWLEDGE

- 8.15 Let's go polling.** Use Google or your favorite search engine to search the Web for "Web polling software." Choose one of the sites that offer software that allows you to conduct your own online opinion polls. (At the time of writing,

www.twiigs.com and www.micropoll.com were two options, but things change quickly on the Web.)

- (a) Choose a site and give the name of the site that you are using.
- (b) Briefly describe two attractive features that the software offers. (For example, you would like to list the answer choices in random order, so that the same choice is not always in the first position.)
- (c) Despite these features in (b), all such polls share a fatal weakness. What is this?

8.16 More on random digit dialing. By the end of 2009, about 25% of adults lived in households with a cell phone and no landline phone, and among adults aged 25 to 29 this number was almost 50%.

- (a) Write a survey question for which the opinions of adults with landline phones only are likely to differ from the opinions of adults with cell phones only. Give the direction of the difference of opinion.
- (b) For the survey question in (a), suppose a survey was conducted using random digit dialing of landline phones only. Would the results be biased? What would be the direction of bias?
- (c) Most surveys now supplement the landline sample contacted by RDD with a second sample of respondents reached through random dialing of cell phone numbers. The landline respondents are weighted to take account of household size and number of telephone lines into the residence, while the cell phone respondents are weighted according to whether they were reachable only by cell phone or also by landline. Explain why it is important to include both a landline sample and a cell phone sample. Why is the number of telephone lines into the residence important? (*Hint:* How does the number of telephone lines into the residence affect the chance of the household being included in the RDD sample?)

CHAPTER 8 SUMMARY

CHAPTER SPECIFICS

- A **sample survey** selects a **sample** from the **population** of all individuals about which we desire information. We base conclusions about the population on data from the sample. It is important to specify exactly what population you are interested in and what variables you will measure.
- The **design** of a sample describes the method used to select the sample from the population. **Random sampling** designs use chance to select a sample.
- The basic random sampling design is a **simple random sample (SRS)**. An SRS gives every possible sample of a given size the same chance to be chosen.
- Choose an SRS by labeling the members of the population and using **random digits** to select the sample. Software can automate this process.
- To choose a **stratified random sample**, classify the population into **strata**, groups of individuals that are similar in some way that is important to the response. Then choose a separate SRS from each stratum.
- Failure to use random sampling often results in **bias**, or systematic errors in the way the sample represents the population. **Voluntary response samples**, in which the respondents choose themselves, are particularly prone to large bias.

- In human populations, even random samples can suffer from bias due to **undercoverage** or **nonresponse**, from **response bias**, or from misleading results due to **poorly worded questions**. Sample surveys must deal expertly with these potential problems in addition to using a random sampling design.
- Most national sample surveys are carried out by telephone, using **random digit dialing** to choose residential telephone numbers at random. Call screening is increasing nonresponse to such surveys, and the rise of cell-phone-only households is increasing undercoverage.

LINK IT

The methods of Chapters 1 to 6 can be used to describe data regardless of how the data were obtained. However, if we want to reason from data to give answers to specific questions or to draw conclusions about the larger population, then the method that was used to collect the data is important. Sampling is one way to collect data, but it does not guarantee that we can draw meaningful conclusions. Biased sampling methods, such as convenience sampling and voluntary response samples, produce data that can be misleading, resulting in incorrect conclusions. Simple random sampling avoids bias and produces data that can lead to valid conclusions regarding the population. Even with perfect sampling methods, there is still sample-to-sample variation; we will begin our study of the connection between sampling variation and drawing conclusions in Chapter 11.

Even when we take a simple random sample, our conclusions can be weakened by undercoverage, nonresponse, and poor wording of questions. Careful attention must be given to all aspects of the sampling process to ensure that the conclusions we make are valid. In some cases, more complex designs are required, such as stratified sampling or multistage sampling. But the use of impersonal chance to select the sample remains a key ingredient in the sampling process. And issues such as undercoverage and nonresponse still remain for these more complex designs.

CHECK YOUR SKILLS

8.17 An online store contacts 1000 customers from its list of customers who have purchased something from them in the last year. In all, 696 of the 1000 say that they are very satisfied with the store's Web site. The population in this setting is

- (a) all customers who have purchased something in the last year.
- (b) the 1000 customers contacted.
- (c) the 696 customers who were very satisfied with the store's Web site.

8.18 A state representative wants to know how voters in his district feel about enacting a statewide smoking ban in all enclosed public places, including bars and restaurants, as well as several other current statewide issues. He mails a questionnaire addressing these issues to an SRS of 800 voters in his

district. Of the 800 questionnaires mailed, 152 were returned. The sample is

- (a) the 800 voters receiving the questionnaire.
- (b) the 152 voters returning the questionnaire.
- (c) all voters in his district.

8.19 In the survey for the previous exercise, there are 8741 registered voters in his district. You label the voters 0001 to 8741 in alphabetical order. Using line 123 of Table B to select the sample, the first 5 voters in your sample would be

- (a) 5458, 0815, 7271, 2560, 2755.
- (b) 5458, 0815, 0727, 1025, 6027.
- (c) 5458, 8150, 7271, 2560, 2755.

8.20 The Web site www.twiigs.com allows you to vote on polls that interest you or to post one of your own. Once

you have found a poll of interest, you just click on “Vote,” and your response becomes part of the sample. One of the questions in July 2010 was “How many times have you been pulled over by the police?” Of the 780 people responding, 70% said “1–5 times.” You can conclude that

- (a) about 70% of Americans have been pulled over by the police “1–5 times.”
- (b) the poll uses voluntary response, so the results tell us little about the population of all adults.
- (c) more people still need to vote on the question, as a larger sample is required to reduce bias.

8.21 Archaeologists plan to examine a sample of two-meter-square plots near an ancient Greek city for artifacts visible in the ground. They choose separate samples of plots from floodplain, coast, foothills, and high hills. What kind of sample is this?

- (a) A simple random sample
- (b) A stratified random sample
- (c) A voluntary response sample

8.22 You must choose an SRS of 10 of the 440 retail outlets in New York that sell your company’s products. How would you label this population in order to use Table B?

- (a) 001, 002, 003, ..., 439, 440
- (b) 000, 001, 002, ..., 439, 440
- (c) 1, 2, ..., 439, 440

8.23 You are using the table of random digits to choose a simple random sample of 6 students from a class of 30 students.

You label the students 01 to 30 in alphabetical order. You are going to select the sample using Table B. Which of the following is a possible sample that could be obtained?

- (a) 45, 74, 04, 18, 07, 65
- (b) 04, 18, 07, 13, 02, 07
- (c) 04, 18, 07, 13, 02, 05

8.24 A sample of households in a community is selected at random from the telephone directory. In this community, 4% of households have no telephone, 10% have only cell phones, and another 25% have unlisted telephone numbers. The sample will certainly suffer from

- (a) nonresponse.
- (b) undercoverage.
- (c) false responses.

8.25 The Pew Research Center survey asked a random sample of 1500 adults, “Do you think the use of marijuana should be made legal, or not?” In the entire sample, 41% said, “Yes, legal.” But only 24% of the Republicans in the sample said, “Yes, legal.” Which of these two sample percents will be more accurate as an estimate of the truth about the population?

- (a) The result for Republicans is more accurate because it is easier to estimate a proportion for a smaller group.
- (b) The result for the entire sample is more accurate because it comes from a larger sample.
- (c) Both are equally accurate because both come from the same sample.

CHAPTER 8 EXERCISES

In all exercises asking for an SRS, you may use Table B, the Simple Random Sample applet, or other software.

8.26 Immigration reform priorities. A Gallup Poll asked, “If you had to choose, what should be the main focus of the U.S. government in dealing with the issue of illegal immigration—

developing a plan for halting the flow of illegal immigrants into the U.S. (or) developing a plan to deal with immigrants who are currently in the U.S. illegally?” Gallup’s report said, “Results are based on telephone interviews conducted June 11–13, 2010, with a random sample of 1,014 adults, aged 18 and older, living in the continental U.S.”¹⁵ What is the population for this sample survey? What is the sample?

8.27 Sampling stuffed envelopes. A large retailer prepares its customers’ monthly credit card bills using an automatic machine that folds the bills, stuffs them into envelopes, and seals the envelopes for mailing. Are the envelopes completely sealed? Inspectors choose 40 envelopes from the 1000 stuffed each hour for visual inspection. What is the population for this sample survey? What is the sample?

8.28 Do you trust the Internet? You want to ask a sample of college students the question “How much do you trust information about health that you find on the Internet—a



Paul J. Richards/AFP/Getty Images

great deal, somewhat, not much, or not at all?" You try out this and other questions on a pilot group of 8 students chosen from your class. The class members are

Adams	Devore	Guo	Newberg	Shoepf
Aeffner	Ding	Heaton	Paulsen	Spagnola
Barnes	Drake	Huling	Payton	Terry
Bower	Eckstein	Kahler	Prince	Vore
Burke	Fassnacht	Kessis	Pulak	Wallace
Cao	Fullmer	Lu	Rabin	Wanner
Cisse	Gandhi	Mattos	Roberts	Zhang

Choose an SRS of 8 students. If you use Table B, start at line 131.

8.29 Sampling telephone area codes. The United States currently has approximately 287 Numbering Plan Areas (NPAs) in service, corresponding to geographic regions. Each NPA is identified by a three-digit code, commonly called an area code. (More are created regularly.)¹⁶ You want to choose an SRS of 20 of these area codes for a study of available telephone numbers. Label the codes 001 to 287 and use the *Simple Random Sample* applet or other software to choose your sample. (If you use Table B, start at line 135 and choose only the first 5 area codes in the sample.)

8.30 Sampling the forest. To gather data on a 1200-acre pine forest in Louisiana, the U.S. Forest Service laid a grid of 1410 equally spaced circular plots over a map of the forest. A ground survey visited a sample of 10% of these plots.¹⁷

- (a) How would you label the plots?
- (b) Choose the first 5 plots in an SRS of 141 plots. (If you use Table B, start at line 105.)

8.31 Sampling students. The freshman class at The Ohio State University contains 6168 students. The Office of International Affairs is considering increasing its programming staff for its study abroad program and is going to sample the entering freshman class to see how many students are considering taking advantage of the opportunity to travel abroad while attending Ohio State.

- (a) How would you label the names in order to select an SRS?
- (b) Use software or Table B, starting at line 135, to select an SRS of 8 Ohio State freshmen.

8.32 Random digits. In using Table B repeatedly to choose random samples, you should not always begin at the same place, such as line 101. Why not?

8.33 Random digits. Which of the following statements are true of a table of random digits, and which are false? Briefly explain your answers.

- (a) There are exactly four 0s in each row of 40 digits.

- (b) Each pair of digits has chance 1/100 of being 00.
- (c) The digits 0000 can never appear as a group, because this pattern is not random.

8.34 Movie viewing. An opinion poll calls 2000 randomly chosen residential telephone numbers and asks to speak with an adult member of the household. The interviewer asks, "How many movies have you watched in a movie theater in the past 12 months?"

- (a) What population do you think the poll has in mind?
- (b) In all, 831 people respond. What is the rate (percent) of nonresponse?
- (c) What source of response error is likely for the question asked?

8.35 Online polls. Example 8.3 (page 203) reports an online poll in which 97% of the respondents opposed issuing driver's licenses to illegal immigrants. National random samples taken at the same time showed about 70% of the respondents opposed to such licenses. Explain briefly to someone who knows no statistics why the random samples report public opinion more reliably than the online poll.

8.36 Nonresponse. Academic sample surveys, unlike commercial polls, often discuss nonresponse. A survey of drivers began by randomly sampling all listed residential telephone numbers in the United States. Of 45,956 calls to these numbers, 5029 were completed.¹⁸ What was the rate of nonresponse for this sample? (Only one call was made to each number. Nonresponse would be lower if more calls were made.)

8.37 Running red lights. The sample described in the previous exercise produced a list of 5024 licensed drivers. The investigators then chose an SRS of 880 of these drivers to answer questions about their driving habits.

- (a) How would you assign labels to the 5024 drivers? Use Table B, starting at line 104, to choose the first 5 drivers in the sample.
- (b) One question asked was "Recalling the last ten traffic lights you drove through, how many of them were red when you entered the intersections?" Of the 880 respondents, 171 admitted that at least one light had been red. A practical problem with this survey is that people may not give truthful answers. What is the likely direction of the bias: do you think more or fewer than 171 of the 880 respondents really ran a red light? Why?

8.38 Seat belt use. A study in El Paso, Texas, looked at seat belt use by drivers. Drivers were observed at randomly chosen convenience stores. After they left their cars, they were invited to answer questions that included questions about seat belt use. In all, 75% said they always used seat

belts, yet only 61.5% were wearing seat belts when they pulled into the store parking lots.¹⁹ Explain the reason for the bias observed in responses to the survey. Do you expect bias in the same direction in most surveys about seat belt use?

8.39 Sampling at a party. At a large block party there are 40 men and 30 women. You want to ask opinions about how to improve the next party. You choose at random 4 of the men and separately choose at random 3 of the women to interview.

(a) What is the probability that any of the 40 men is in your random sample of 4 men to be interviewed? What is the probability that any of the 30 women is in your random sample of 3 women to be interviewed?

(b) If you have done the calculations correctly in part (a), the probability of any person at the party being interviewed is the same. Why is your sample of 7 men and women not an SRS of people from the party?

8.40 Sampling pharmacists. All pharmacists in the Canadian province of Ontario are required to be members of the Ontario College of Pharmacists. In 2009, there were 11,361 members of the college divided into 17 electoral districts, with each district having an elected member on the Council. The number of members in each district follow:²⁰

District	1	2	3	4	5	6	7	8	9
Membership	997	803	694	771	536	1126	1104	864	286
District	10	11	12	13	14	15	16	17	
Membership	414	458	572	537	303	263	743	890	

Suppose they are interested in obtaining members' views on the Minor Ailments Program, which proposes making pharmacists the primary source of care for patients with some 30 minor ailments. To be sure that the opinions of all districts are represented, you choose a stratified random sample of 5 pharmacists from each district. Explain how you will assign labels within each district, and then give the labels of the pharmacists from Districts 1 and 2 in your sample. If you use Table B, start at line 122 for District 1 and at line 131 for District 2. Why should you not start at the same line in Table B to obtain your samples for Districts 1 and 2?

8.41 Sampling Amazon forests. Stratified samples are widely used to study large areas of forest. Based on satellite images, a forest area in the Amazon basin is divided into 14 types. Foresters studied the four most commercially val-

able types: alluvial climax forests of quality levels 1, 2, and 3, and mature secondary forest. They divided the area of each type into large parcels, chose parcels of each type at random, and counted tree species in a 20- by 25-meter rectangle randomly placed within each parcel selected. Here is some detail:



© Age fotostock/SuperStock

Forest type	Total parcels	Sample size
Climax 1	36	4
Climax 2	72	7
Climax 3	31	3
Secondary	42	4

Choose the stratified sample of 18 parcels. Be sure to explain how you assigned labels to parcels. If you use Table B, start at line 102.

8.42 Canadian health care survey. The Tenth Annual Health Care in Canada Survey is a survey of the Canadian public's and health care providers' opinions on a variety of health care issues, including quality of health care, access to health care, health and the environment, and so forth. A description of the survey follows:

The 10th edition of the Health Care in Canada Survey was conducted by POLLARA Research between October 3rd and November 8th, 2007. Results for the survey are based on telephone interviews with nationally representative samples of 1,223 members of the Canadian public, 202 doctors, 201 nurses, 202 pharmacists and 201 health managers. Public results are considered to be accurate within $\pm 2.8\%$, while the margin of error for results for doctors, nurses, pharmacists and managers is $\pm 6.9\%$.²¹

- (a) Why is the accuracy greater for the public than for health care providers and managers?
- (b) Why do you think they sampled the public as well as health care providers and managers?

8.43 Systematic random samples. Systematic random samples go through a list of the population at fixed intervals from a randomly chosen starting point. For example, a study

of dating among college students chose a systematic sample of 200 single male students at a university as follows.²² Start with a list of all 9000 single male students. Because $9000/200 = 45$, choose one of the first 45 names on the list at random and then every 45th name after that. For example, if the first name chosen is at position 23, the systematic sample consists of the names at positions 23, 68, 113, 158, and so on up to 8978.

- (a) Use Table B to choose a systematic random sample of 5 names from a list of 200. Enter the table at line 120.
- (b) Like an SRS, a systematic sample gives all individuals the same chance to be chosen. Explain why this is true, then explain carefully why a systematic sample is nonetheless *not* an SRS.

8.44 Why random digit dialing is common. The list of individuals from which a sample is actually selected is called the *sampling frame*. Ideally, the frame should list every individual in the population, but in practice this is often difficult. A frame that leaves out part of the population is a common source of undercoverage.

- (a) Suppose that a sample of households in a community is selected at random from the telephone directory. What households are omitted from this frame? What types of people do you think are likely to live in these households? These people will probably be underrepresented in the sample.
- (b) It is usual in telephone surveys to use random digit dialing equipment that selects the last four digits of a telephone number at random after being given the exchange (the first three digits), as described in Example 8.5 (page 206). Which of the households you mentioned in your answer to (a) will be included in the sampling frame by random digit dialing?

8.45 Gulf oil spill. Two months after the Gulf oil spill began in April 2010, nearly half of Americans (49%) believed that at least some of the affected beaches will never recover, according to a Gallup Poll conducted June 11 to 13, 2010. Results are based on telephone interviews of a random sample of 1014 adults, aged 18 and older, selected using random digit dialing sampling. In the survey methods section, Gallup reports: “Interviews are conducted with respondents on landline tele-

phones (for respondents with a landline telephone) and cellular phones (for respondents who are cell phone-only). Each sample includes a minimum quota of 150 cell phone-only respondents and 850 landline respondents, with additional minimum quotas among landline respondents for gender within region.”²³

- (a) What is automated random digit dialing? Why is it a practical method for obtaining (almost) an SRS of households with landline phones?
- (b) The survey wants the opinion of an individual adult. Several adults may live in a household. In that case, the survey interviewed the adult with the most recent birthday. Why is this preferable to simply interviewing the person who answers the phone?
- (c) The survey included both landline telephones and cellular phones. Why do you think this may be important? Sampling landline telephones and cellular phones separately corresponds to what type of sampling design?

8.46 Wording survey questions. Comment on each of the following as a potential sample survey question. Is the question sufficiently clear? Is it slanted toward a desired response?

- (a) “Some cell phone users have developed brain cancer. Should all cell phones come with a warning label explaining the danger of using cell phones?”
- (b) “Do you agree that a national system of health insurance should be favored because it would provide health insurance for everyone and would reduce administrative costs?”
- (c) “In view of the negative externalities in parent labor force participation and pediatric evidence associating increased group size with morbidity of children in day care, do you support government subsidies for day care programs?”

8.47 Your own bad questions. Write your own examples of bad sample survey questions.

- (a) Write a biased question designed to get one answer rather than another.
- (b) Write a question to which many people may not give truthful answers.

8.48 The Canadian census. The Canadian government’s decision to eliminate the mandatory long-form version of the census and to move these questions to an optional survey has many concerned. Many members of the business community and economists stressed the importance of the census data for crafting public policy. The minister of industry was given the task of defending the government’s decision. In response to an argument that making the long form of the census voluntary would skew the data by eliminating the statistical randomness of the survey, the minister replied: “Wrong. Statisticians can ensure validity with a larger sample size.”²⁴ Is the minister correct? If not, explain in simple terms the error in his statement.



Derick E. Hingle/Bloomberg via Getty Images



EXPLORING THE WEB

8.49 Poor survey designs. The Web site for the American Association for Public Opinion Research discusses several issues about polls. This information can be found at www.aapor.org/Poll_andamp_Survey_FAQ.htm. Click on the link “Questions to Ask When Writing about Polls” for suggestions about how to determine if a poll is good or bad. Click on the link “What Is a Random Sample?” and then the link “Bad Samples” for some examples of flawed samples.

- (a) You are going to design a survey at your university. Give a question of interest and two examples of bad ways to collect your sample, along with the likely direction of bias that would result. Explain your answers.
- (b) How would you modify your examples in (a) to produce a better sample? What are some difficulties you might encounter when collecting your sample?

8.50 Find a survey. The Web site for the Pew Research Center for the People and the Press is www.people-press.org. Go to the Web site and read one of the featured surveys. Which questions listed under “Questions to Ask When Writing about Polls” from the previous exercise can be answered from the information in the featured survey you have read? You may find the concluding section entitled “About the Survey” and some of the links at the end of the featured survey helpful for finding answers to these questions.



Producing Data: Experiments

Chapter 9

IN THIS CHAPTER WE COVER...

- Observation versus experiment
- Subjects, factors, treatments
- How to experiment badly
- Randomized comparative experiments
- The logic of randomized comparative experiments
- Cautions about experimentation
- Matched pairs and other block designs

A sample survey aims to gather information about a population without disturbing the population in the process. Sample surveys are one kind of *observational study*. Other observational studies observe the behavior of animals in the wild or the interactions between teacher and students in the classroom. This chapter is about statistical designs for *experiments*, a quite different way to produce data.

OBSERVATION VERSUS EXPERIMENT

In contrast to observational studies, experiments don't just observe individuals or ask them questions. They actively impose some treatment in order to observe the response. Experiments can answer questions such as "Does aspirin reduce the chance of a heart attack?" and "Do a majority of college students prefer Pepsi to Coke when they taste both without knowing which they are drinking?"

OBSERVATION VERSUS EXPERIMENT

An **observational study** observes individuals and measures variables of interest but does not attempt to influence the responses. The purpose of an observational study is to describe some group or situation.

An **experiment**, on the other hand, deliberately imposes some treatment on individuals in order to observe their responses. The purpose of an experiment is to study whether the treatment causes a change in the response.



You just don't understand

A sample survey of journalists and scientists

found quite a communications gap. Journalists think that scientists are arrogant, while scientists think that journalists are ignorant. We won't take sides, but here is one interesting result from the survey: 82% of the scientists agree that the "media do not understand statistics well enough to explain new findings" in medicine and other fields.

An observational study, even one based on a statistical sample, is a poor way to gauge the effect of a treatment. To see the response to a change, we must actually impose the change. When our goal is to understand cause and effect, experiments are the only source for fully convincing data. For this reason, the distinction between observation and experiment is one of the most important in statistics.

EXAMPLE 9.1 Drink a little, but not a lot

Many observational studies show that people who drink a moderate amount of alcohol have less heart disease than people who drink no alcohol or who drink heavily.¹ ("Moderate" means one or two drinks a day for men and one drink a day for women.) Is this association good reason to think that moderate drinking actually causes less heart disease? People who choose to drink in moderation are, as a group, different from both heavy drinkers and abstainers. They are more likely to maintain a healthy weight, get enough sleep, and exercise regularly. Moderate drinkers may be healthier because of these healthy habits rather than because of the effect of alcohol on health.

It is easy to imagine an experiment that would settle the issue of whether moderate drinking really causes reduced heart disease. Choose half of a large group of adults at random to be the "treatment" group. The remaining half becomes the "control" group. Require the treatment group to have one alcoholic drink every day. Require the control group to abstain from alcohol. Follow both groups for a decade. This experiment isolates the effect of alcohol. Of course, it isn't practical or ethical to carry out such an experiment. ■

The point of Example 9.1 is the contrast between observing people who choose for themselves what to drink and an experiment that requires some people to drink and others to abstain. When we simply observe people's drinking choices, the effect of moderate drinking is confounded with (mixed up with) the characteristics of people who choose to drink in moderation. These characteristics are lurking variables (see page 143) that make it hard to see the true relationship between the explanatory and response variables. Figure 9.1 shows the confounding in picture form.

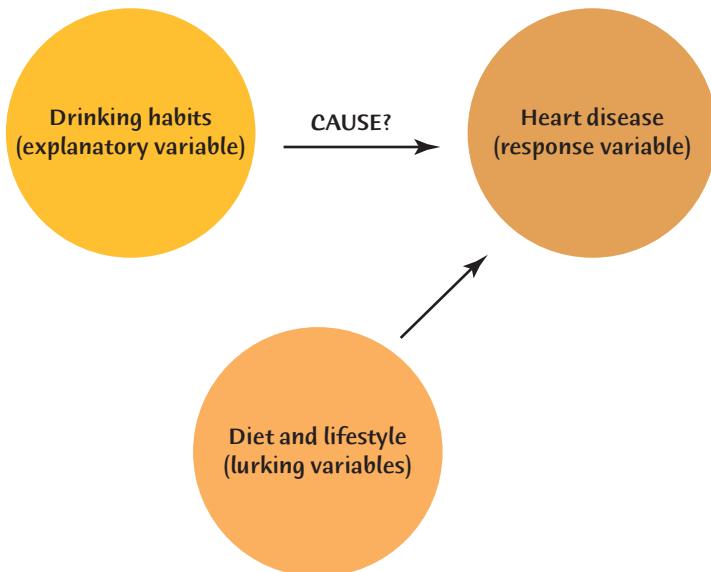


FIGURE 9.1

Confounding: we can't distinguish the effects of drinking habits from the effects of overall diet and lifestyle.

CONFOUNDING

Two variables (explanatory variables or lurking variables) are **confounded** when their effects on a response variable cannot be distinguished from each other.

Observational studies of the effect of one variable on another often fail because the explanatory variable is confounded with lurking variables. Well-designed experiments take steps to prevent confounding.



APPLY YOUR KNOWLEDGE

- 9.1 **Cell phones and brain cancer.** A study of cell phones and the risk of brain cancer looked at a group of 469 people who have brain cancer. The investigators matched each cancer patient with a person of the same sex, age, and race who did not have brain cancer, then asked about use of cell phones. Result: “Our data suggest that use of handheld cellular telephones is not associated with risk of brain cancer.”² Is this an observational study or an experiment? Why? What are the explanatory and response variables?
- 9.2 **The font matters!** In general, high perceived effort is an impediment to changes in behavior, whether it is modifying your diet or adopting an exercise routine. Yet little is known about how individuals estimate effort for a novel behavior. Researchers divide 40 students into two groups of 20. The first group reads instructions for an exercise program printed in an easy-to-read font (Arial, 12 point), and the second group reads identical instructions in a difficult-to-read font (Brush, 12 point). They estimated how many minutes the program would take (open-ended) and used a 7-point rating scale to report whether they were likely to make the exercise program part of their daily routine (7 = very likely).³ As hypothesized by the researchers, those reading about the exercise program in the more difficult-to-read font estimated that it would take longer and were less likely to make the exercise program part of their regular routine. Is this an experiment? Why or why not? What are the explanatory and response variables?
- 9.3 **Quitting smoking and risk for type 2 diabetes.** Researchers studied a group of 10,892 middle-aged adults over a period of nine years. They found that smokers who quit had a higher risk for diabetes within three years of quitting than either nonsmokers or continuing smokers.⁴ Does this show that stopping smoking causes the short-term risk for diabetes to increase? (Weight gain has been shown to be a major risk factor for developing type 2 diabetes and is often a side effect of quitting smoking.) Based on this research, should you tell a middle-aged adult who smokes that stopping smoking can cause diabetes and advise him or her to continue smoking? Carefully explain your answers to both questions.



Paula Solloway/Alamy

SUBJECTS, FACTORS, TREATMENTS

A study is an experiment when we actually do something to people, animals, or objects in order to observe the response. Because the purpose of an experiment is to reveal the response of one variable to changes in other variables, the distinction between explanatory and response variables is essential. Here is the basic vocabulary of experiments.

SUBJECTS, FACTORS, TREATMENTS

The individuals studied in an experiment are often called **subjects**, particularly when they are people.

The explanatory variables in an experiment are often called **factors**.

A **treatment** is any specific experimental condition applied to the subjects. If an experiment has more than one factor, a treatment is a combination of specific values of each factor.

EXAMPLE 9.2 Foster care versus orphans

Do abandoned children placed in foster homes do better than similar children placed in an institution? The Bucharest Early Intervention Project found that the answer is a clear “Yes.” The *subjects* were 136 young children abandoned at birth and living in orphanages in Bucharest, Romania. Half of the children, chosen at random, were placed in foster homes. The other half remained in the orphanages. The experiment compared these two *treatments*. There is a single *factor*, type of care, with two values, foster and institutional care. When there is only one factor, the levels or values of the factor correspond to the treatments. The *response variables* included measures of mental and physical development.⁵ (Foster care was not easily available in Romania at the time and so was paid for by the study. See Exercise 15 on page 256 in the Data Ethics essay for ethical questions concerning this experiment.) ■

EXAMPLE 9.3 Effects of TV advertising

What are the effects of repeated exposure to an advertising message? The answer may depend both on the length of the ad and on how often it is repeated. An experiment investigated this question using undergraduate students as *subjects*. All subjects viewed a 40-minute television program that included ads for a digital camera. Some subjects saw a 30-second commercial; others, a 90-second version. The same commercial was shown either 1, 3, or 5 times during the program.

This experiment has 2 *factors*: length of the commercial, with 2 values, and repetitions, with 3 values. The 6 combinations of 1 value of each factor form 6 *treatments*. Figure 9.2 shows the layout of the treatments. After viewing, all the subjects answered questions about their recall of the ad, their attitude toward the camera, and their intention to purchase it. These are the *response variables*. ■

		Factor B Repetitions			
		1 time	3 times	5 times	
Factor A Length	30 seconds	1	2	3	
	90 seconds	4	5	6	

Subjects assigned to Treatment 3 see a 30-second ad five times during the program.

FIGURE 9.2

The treatments in the experimental design of Example 9.3. Combinations of values of the two factors form six treatments.

Examples 9.2 and 9.3 illustrate the advantages of experiments over observational studies. In an experiment, we can study the effects of the specific treatments we are interested in. By assigning subjects to treatments, we can avoid confounding. For example, observational studies of the effects of foster homes versus institutions on the development of children have often been biased because healthier or more alert children tend to be placed in homes. The random assignment in Example 9.2 eliminates bias in placing the children. Moreover, we can control the environment of the subjects to hold constant factors that are of no interest to us, such as the specific product advertised in Example 9.3.

Another advantage of experiments is that we can study the combined effects of several factors simultaneously. The interaction of several factors can produce effects that could not be predicted from looking at the effect of each factor alone. Perhaps longer commercials increase interest in a product, and more commercials also increase interest, but if we both make a commercial longer and show it more often, viewers get annoyed and their interest in the product drops. The two-factor experiment in Example 9.3 will help us find out.

APPLY YOUR KNOWLEDGE

For each of the following experiments, identify the subjects, the factors, the treatments, and the response variables.

9.4 Ginkgo extract and the post-lunch dip. The post-lunch dip is the drop in mental alertness after a midday meal. Does an extract of the leaves of the ginkgo tree reduce the post-lunch dip? Assign healthy people aged 18 to 40 to take either ginkgo extract or a placebo pill. After lunch, ask them to read seven pages of random letters and place an X over every e. Count the number of misses.

9.5 Growing in the shade. Ability to grow in shade may help pines found in the dry forests of Arizona resist drought. How well do these pines grow in shade? Plant pine seedlings in a greenhouse in either full light, light reduced to 25% of normal by shade cloth, or light reduced to 5% of normal. At the end of the study, dry the young trees and weigh them.

9.6 Reactions to simulated news reports. A sample of University of Colorado students each viewed one of two simulated news reports about a terrorist bombing against the United States by a fictitious country. One report showed the bombing attack on a military target and the other on a cultural/educational site. Additionally, before viewing the news report, each student read one of two “primes.” The first was a prime for *forgiveness* based on the biblical saying “Love thy enemy,” while the second was a *retaliatory* prime based on the biblical saying “An eye for an eye, and a tooth for a tooth.” After viewing the news report, the students were asked to rate on a scale of 1 to 12 what the U.S. reaction should be, with the lowest score (1) corresponding to the United States sending a special ambassador to the country and the highest score (12) corresponding to an all-out nuclear attack against the country.⁶ (Use a diagram like Figure 9.2 to display the factors and treatments.)



Howard Bjornson/Getty Images

HOW TO EXPERIMENT BADLY

Experiments are the preferred method for examining the effect of one variable on another. By imposing the specific treatment of interest and controlling other influences, we can pin down cause and effect. Statistical designs are often essential for effective experiments. To see why, let's look at an example in which an experiment suffers from confounding just as observational studies do.

EXAMPLE 9.4 An uncontrolled experiment

A college regularly offers a review course to prepare candidates for the Graduate Management Admission Test (GMAT), which is required by most graduate business schools. This year, it offers only an online version of the course. The average GMAT score of students in the online course is 10% higher than the longtime average for those who took the classroom review course. Is the online course more effective?

This experiment has a very simple design. A group of subjects (the students) were exposed to a treatment (the online course), and the outcome (GMAT scores) was observed. Here is the design:

Subjects → Online course → GMAT scores

A closer look at the GMAT review course showed that the students in the online review course were quite different from the students who in past years took the classroom course. In particular, they were older and more likely to be employed. An online course appeals to these mature people, but we can't compare their performance with that of the undergraduates who previously dominated the course. The online course might even be less effective than the classroom version. The effect of online versus in-class instruction is confounded with the effect of lurking variables. As a result of confounding, the experiment is biased in favor of the online course.

Would the situation have been different if both the online and the classroom courses had been given this year? If students still chose the course they wanted, with older students tending to sign up for the online course and younger students tending to sign up for the classroom course, then the effect of type of course would still be confounded with the lurking variable age. The solution will be described in the next section. ■

Most laboratory experiments use a design like that in Example 9.4:

Subjects → Treatment → Measure response

In the controlled environment of the laboratory, simple designs often work well. Field experiments and experiments with living subjects are exposed to more variable conditions and deal with more variable subjects. *Outside the laboratory, uncontrolled experiments often yield worthless results because of confounding with lurking variables.*

APPLY YOUR KNOWLEDGE

- 9.7 Reducing unemployment.** Will cash bonuses speed the return to work of unemployed people? A state department of labor notes that last year 68% of people who

filed claims for unemployment insurance found a new job within 15 weeks. As an experiment, the state offers \$500 to people filing unemployment claims if they find a job within 15 weeks. The percent who do so increases to 77%. Suggest some conditions that might make it easier or harder to find a job this year as opposed to last year. Confounding with these lurking variables makes it impossible to say whether the bonus really caused the increase.

RANDOMIZED COMPARATIVE EXPERIMENTS

The remedy for the confounding in Example 9.4 is to be sure that we do a *comparative experiment* in which some students are taught in the classroom and other, similar students take the course online. The first group is called a **control group**. Most well-designed experiments compare two or more treatments. Part of the design of an experiment is a description of the factors (explanatory variables) and the layout of the treatments, with comparison as the leading principle.

control group

However, as discussed at the end of Example 9.4, comparison alone isn't enough to produce results we can trust. If the treatments are given to groups that differ markedly when the experiment begins, bias will result. If we allow students to elect online or classroom instruction, students who are older and employed are likely to sign up for the online course. Personal choice will bias our results in the same way that volunteers bias the results of online opinion polls. The solution to the problem of bias in sampling is random selection, and the same is true in experiments. The subjects assigned to any treatment should be chosen at random from the available subjects.

RANDOMIZED COMPARATIVE EXPERIMENT

An experiment that uses both comparison of two or more treatments and random assignment of subjects to treatments is a **randomized comparative experiment**.

EXAMPLE 9.5 Classroom versus online

The college decides to compare the progress of 25 on-campus students taught in the classroom with that of 25 students taught the same material online. Select the students who will be taught online by taking a simple random sample of size 25 from the 50 available subjects. The remaining 25 students form the control group. They will receive classroom instruction. The result is a randomized comparative experiment with two groups. Figure 9.3 outlines the design in graphical form.

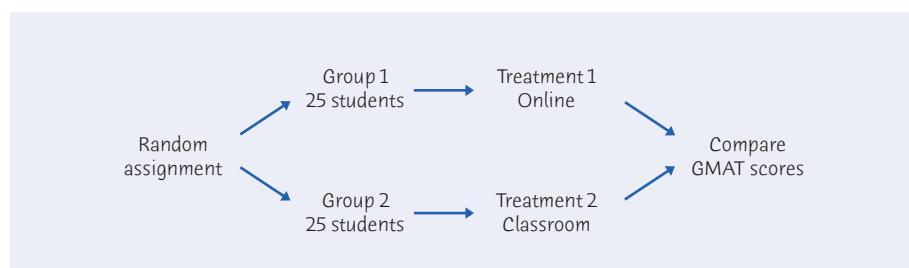


FIGURE 9.3

Outline of a randomized comparative experiment to compare online and classroom instruction, for Example 9.5.



The selection procedure is exactly the same as it is for sampling. **Label:** Label the 50 students 01 to 50. **Table:** Go to the table of random digits and read successive two-digit groups. The first 25 labels encountered select the online group. As usual, ignore repeated labels and groups of digits not used as labels. For example, if you begin at line 125 in Table B, the first 5 students chosen are those labeled 21, 49, 37, 18, and 44. Software such as the *Simple Random Sample* applet makes it particularly easy to choose treatment groups at random. ■

The design in Example 9.5 is *comparative* because it compares two treatments (the two instructional settings). It is *randomized* because the subjects are assigned to the treatments by chance. This “flowchart” outline in Figure 9.3 presents all the essentials: randomization, the sizes of the groups and which treatment they receive, and the response variable. There are, as we will see later, statistical reasons for generally using treatment groups that are about equal in size. We call designs like that in Figure 9.3 *completely randomized*.

COMPLETELY RANDOMIZED DESIGN

In a **completely randomized** experimental design, all the subjects are allocated at random among all the treatments.

Completely randomized designs can compare any number of treatments. Here is an example that compares three treatments.

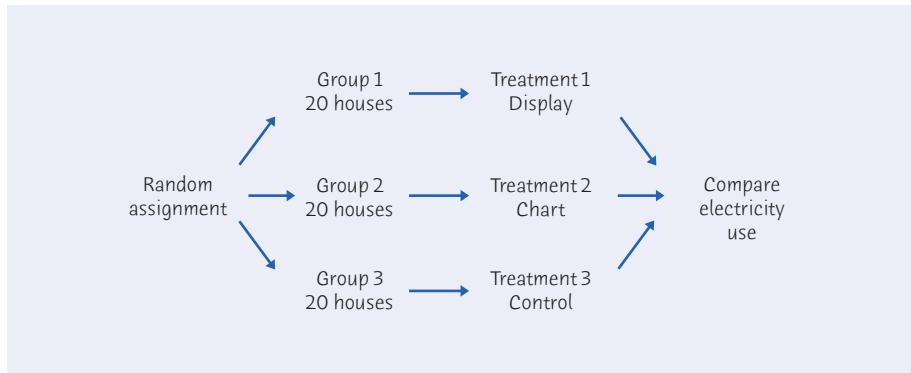
EXAMPLE 9.6 Conserving energy

Many utility companies have introduced programs to encourage energy conservation among their customers. An electric company considers placing small digital displays in households to show current electricity use and what the cost would be if this use continued for a month. Will the displays reduce electricity use? Would cheaper methods work almost as well? The company decides to conduct an experiment.

One cheaper approach is to give customers a chart and information about monitoring their electricity use from their outside meter. The experiment compares these two approaches (display, chart) and also a control. The control group of customers receives information about energy conservation but no help in monitoring electricity use. The response variable is total electricity used in a year. The company finds 60 single-family residences in the same city willing to participate, so it assigns 20 residences at random to each of the three treatments. Figure 9.4 outlines the design.

To use the *Simple Random Sample* applet, set the population labels as 1 to 60 and the sample size to 20 and click “Reset” and “Sample.” The 20 households chosen receive the displays. The “Population hopper” now contains the 40 remaining households, in scrambled order. Click “Sample” again to choose 20 of these to receive charts. The 20 households remaining in the “Population hopper” form the control group.



**FIGURE 9.4**

Outline of a completely randomized design comparing three energy-saving programs, for Example 9.6.

To use Table B, label the 60 households 01 to 60. Enter the table to select an SRS of 20 to receive the displays. Continue in Table B, selecting 20 more to receive charts. The remaining 20 form the control group. ■

Examples 9.5 and 9.6 describe completely randomized designs that compare values of a single factor. In Example 9.5, the factor is the type of instruction. In Example 9.6, it is the method used to encourage energy conservation. Completely randomized designs can have more than one factor. The advertising experiment of Example 9.3 (see page 226) has two factors: the length and the number of repetitions of a television commercial. Their combinations form the six treatments outlined in Figure 9.2. A completely randomized design assigns subjects at random to these six treatments. Once the layout of treatments is set, the randomization needed for a completely randomized design is tedious but straightforward.

APPLY YOUR KNOWLEDGE

9.8 Adolescent obesity. Adolescent obesity is a serious health risk affecting more than 5 million young people in the United States alone. Laparoscopic adjustable gastric banding has the potential to provide a safe and effective treatment. Fifty adolescents between 14 and 18 years old with a body mass index higher than 35 were recruited from the Melbourne, Australia, community for the study. Twenty-five were randomly selected to undergo gastric banding, and the remaining 25 were assigned to a supervised lifestyle intervention program involving diet, exercise, and behavior modification. All subjects were followed for two years and their weight loss was recorded.⁷

- Outline the design of this experiment, following the model of Figure 9.3 (page 229). What is the response variable?
- Carry out the random assignment of 25 adolescents to the gastric-banding group, using the *Simple Random Sample* applet, other software, or Table B, starting at line 130.

9.9 More rain for California? The changing climate will probably bring more rain to California, but we don't know whether the additional rain will come during the winter wet season or extend into the long dry season in spring and summer. Kenwyn Suttle of the University of California at Berkeley and his coworkers carried out a randomized controlled experiment to study the effects of more rain in either season. They randomly





Courtesy Blake Suttle

assigned plots of open grassland to 3 treatments: added water equal to 20% of annual rainfall either during January to March (winter) or during April to June (spring), and no added water (control). Thirty-six circular plots of area 70 square meters were available (see the photo), of which 18 were used for this study. One response variable was total plant biomass, in grams per square meter, produced in a plot over a year.⁸

- Outline the design of the experiment, following the model of Figure 9.4 (page 231).
- Number all 36 plots and choose 6 at random for each of the 3 treatments. Be sure to explain how you did the random selection.

9.10 Effects of TV advertising. Figure 9.2 (page 226) displays the 6 treatments for the two-factor experiment on TV advertising described in Example 9.3. The 24 students named below will serve as subjects. Outline the design and randomly assign the subjects to the 6 treatments, an equal number of subjects to each treatment. If you use Table B, start at line 132.

Abramson	Biry	Cohen	Greenberg	Linder	Stanley
Anthony	Blake	Cote	Kessis	Minor	Tory
Austen	Brower	Delp	Koster	Schwartz	Truitt
Baker	Carroll	Disbro	Kruger	Shi	Walsh

THE LOGIC OF RANDOMIZED COMPARATIVE EXPERIMENTS

Randomized comparative experiments are designed to give good evidence that differences in the treatments actually *cause* the differences we see in the response. The logic is as follows:

- Random assignment of subjects forms groups that should be similar in all respects before the treatments are applied. Exercise 9.50 uses the *Simple Random Sample* applet to demonstrate this.
- Comparative design ensures that influences other than the experimental treatments operate equally on all groups.
- Therefore, differences in average response must be due either to the treatments or to the play of chance in the random assignment of subjects to the treatments.

That “either-or” deserves more thought. In Example 9.5, we cannot say that *any* difference between the average GMAT scores of students enrolled online and in the classroom must be caused by a difference in the effectiveness of the two types of instruction. There would be some difference even if both groups received the same instruction, because of variation among students in background and study habits. Chance assigns students to one group or the other, and this creates a chance difference between the groups. We would not trust an experiment with just one student in each group, for example. The results would depend too much on which group got lucky and received the stronger student. If we assign many subjects to each group, however, the effects of chance will average out, and there will be little difference in the average responses in the two groups unless the treatments themselves cause a difference. “Use enough subjects to reduce chance variation” is the third big idea of statistical design of experiments.



PRINCIPLES OF EXPERIMENTAL DESIGN

The basic principles of statistical design of experiments are

1. **Control** the effects of lurking variables on the response, most simply by comparing two or more treatments.
2. **Randomize**—use chance to assign subjects to treatments.
3. **Use enough subjects** in each group to reduce chance variation in the results.



What's news?

Randomized comparative experiments provide the best evidence for medical advances. Do newspapers care? Maybe not. University researchers looked at 1192 articles in medical journals, of which 7% were turned into stories by the two newspapers examined. Of the journal articles, 37% concerned observational studies and 25% described randomized experiments. Among the articles publicized by the newspapers, 58% were observational studies and only 6% were randomized experiments. Conclusion: the newspapers want exciting stories, especially bad-news stories, whether or not the evidence is good.

We hope to see a difference in the responses so large that it is unlikely to happen just because of chance variation. We can use the laws of probability, which describe chance behavior, to learn if the treatment effects are larger than we would expect to see if only chance were operating. If they are, we call them *statistically significant*.

STATISTICAL SIGNIFICANCE

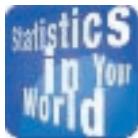
An observed effect so large that it would rarely occur by chance is called **statistically significant**.

If we observe statistically significant differences among the groups in a randomized comparative experiment, we have good evidence that the treatments actually caused these differences. You will often see the phrase “statistically significant” in reports of investigations in many fields of study. The great advantage of randomized comparative experiments is that they can produce data that give good evidence for a cause-and-effect relationship between the explanatory and response variables. We know that in general a strong association does not imply causation. A statistically significant association in data from a well-designed experiment *does* imply causation.

APPLY YOUR KNOWLEDGE

9.11 Prayer and meditation. You read in a magazine that “nonphysical treatments such as meditation and prayer have been shown to be effective in controlled scientific studies for such ailments as high blood pressure, insomnia, ulcers, and asthma.” Explain in simple language what the article means by “controlled scientific studies.” Why can such studies in principle provide good evidence that, for example, meditation is an effective treatment for high blood pressure?

9.12 Conserving energy. Example 9.6 (page 230) describes an experiment to learn whether providing households with digital displays or charts will reduce their electricity consumption. An executive of the electric company objects to including a control group. He says: “It would be simpler to just compare electricity use last year (before the display or chart was provided) with consumption in the same period this year. If households use less electricity this year, the display or chart must be working.” Explain clearly why this design is inferior to that in Example 9.6.



Scratch my furry ears

Rats and rabbits, specially bred to be uniform in their inherited characteristics, are the subjects in many experiments. Animals, like people, are quite sensitive to how they are treated. This can create opportunities for hidden bias. For example, human affection can change the cholesterol level of rabbits. Choose some rabbits at random and regularly remove them from their cages to have their heads scratched by friendly people. Leave other rabbits unloved. All the rabbits eat the same diet, but the rabbits that receive affection have lower cholesterol.

placebo

9.13 Healthy diet and cataracts. The relationship between healthy diet and prevalence of cataracts was assessed using a sample of 1808 participants from the Women's Health Initiative Observational Study. Having a high Healthy Eating Index score was the strongest predictor of a reduced risk of cataracts, among modifiable behaviors considered. The Healthy Eating Index score was created by the U.S. Department of Agriculture and measures how well a person's diet conforms to recommended healthy eating patterns. The report concludes: "These data add to the body of evidence suggesting that eating foods rich in a variety of vitamins and minerals may contribute to postponing the occurrence of the most common type of cataract in the United States."⁹

- (a) Explain why this is an observational study rather than an experiment.
- (b) Although the result was statistically significant, the authors did not use strong language in stating their conclusions, using words such as "suggesting" and "may." Do you think that their language is appropriate given the nature of the study? Why?

CAUTIONS ABOUT EXPERIMENTATION

The logic of a randomized comparative experiment depends on our ability to treat all the subjects identically in every way except for the actual treatments being compared. Good experiments therefore require careful attention to details to ensure that all subjects really are treated identically.

If some subjects in a medical experiment take a pill each day and a control group takes no pill, the subjects are not treated identically. Many medical experiments are therefore "placebo-controlled." A study of the effects of taking vitamin E on heart disease is typical. All of the subjects receive the same medical attention during the several years of the experiment. All of them take a pill every day, vitamin E in the treatment group and a placebo in the control group. A *placebo* is a dummy treatment. Many patients respond favorably to any treatment, even a placebo, perhaps because they trust the doctor. The response to a dummy treatment is called the *placebo effect*. If the control group did not take any pills, the effect of vitamin E in the treatment group would be confounded with the placebo effect, the effect of simply taking pills.

In addition, such studies are usually *double-blind*. The subjects don't know whether they are taking vitamin E or a placebo. Neither do the medical personnel who work with them. The double-blind method avoids unconscious bias by, for example, a doctor who is convinced that a vitamin must be better than a placebo. In many medical studies, only the statistician who does the randomization knows which treatment each patient is receiving.

DOUBLE-BLIND EXPERIMENTS

In a *double-blind* experiment, neither the subjects nor the people who interact with them know which treatment each subject is receiving.

Placebo controls and the double-blind method are more ways to eliminate possible confounding. But even well-designed experiments often face another problem: **lack of realism**. Practical constraints may mean that the subjects or treatments or setting of an experiment don't realistically duplicate the conditions we really want to study. Here are two examples.

Lack of realism

EXAMPLE 9.7 Response to advertising

The study of television advertising in Example 9.3 (page 226) showed a 40-minute video to students who knew an experiment was going on. We can't be sure that the results apply to everyday television viewers. Many behavioral science experiments use as subjects students or other volunteers who know they are subjects in an experiment. That's not a realistic setting. ■

EXAMPLE 9.8 Center brake lights

Do those high center brake lights, required on all cars sold in the United States since 1986, really reduce rear-end collisions? Randomized comparative experiments with fleets of rental and business cars, done before the lights were required, showed that the third brake light reduced rear-end collisions by as much as 50%. Alas, requiring the third light in all cars led to only a 5% drop.

What happened? Most cars did not have the extra brake light when the experiments were carried out, so it caught the eye of following drivers. Now that almost all cars have the third light, they no longer capture attention. ■



© Image 100/CORBIS

Lack of realism can limit our ability to apply the conclusions of an experiment to the settings of greatest interest. Most experimenters want to generalize their conclusions to some setting wider than that of the actual experiment. *Statistical analysis of an experiment cannot tell us how far the results will generalize.* Nonetheless, the randomized comparative experiment, because of its ability to give convincing evidence for causation, is one of the most important ideas in statistics.



APPLY YOUR KNOWLEDGE

9.14 Testosterone for older men. As men age, their testosterone levels gradually decrease. This may cause a reduction in lean body mass, an increase in fat, and other undesirable changes. Do testosterone supplements reverse some of these effects? A study in the Netherlands assigned 237 men aged 60 to 80 with low or low-normal testosterone levels to either a testosterone supplement or a placebo. The report in the *Journal of the American Medical Association* described the study as a “double-blind, randomized, placebo-controlled trial.”¹⁰ Explain each of these terms to someone who knows no statistics.

9.15 Does meditation reduce anxiety? An experiment that claimed to show that meditation reduces anxiety proceeded as follows. The experimenter interviewed the

subjects and rated their level of anxiety. Then the subjects were randomly assigned to two groups. The experimenter taught one group how to meditate and they meditated daily for a month. The other group was simply told to relax more. At the end of the month, the experimenter interviewed all the subjects again and rated their anxiety level. The meditation group now had less anxiety. Psychologists said that the results were suspect because the ratings were not blind. Explain what this means and how lack of blindness could bias the reported results.

MATCHED PAIRS AND OTHER BLOCK DESIGNS

Completely randomized designs are the simplest statistical designs for experiments. They illustrate clearly the principles of control, randomization, and adequate number of subjects. However, completely randomized designs are often inferior to more elaborate statistical designs. In particular, matching the subjects in various ways can produce more precise results than simple randomization.

matched pairs design

One common design that combines matching with randomization is the **matched pairs design**. A matched pairs design compares just two treatments. Choose pairs of subjects that are as closely matched as possible. Use chance to decide which subject in a pair gets the first treatment. The other subject in that pair gets the other treatment. That is, the random assignment of subjects to treatments is done within each matched pair, not for all subjects at once. Sometimes each “pair” in a matched pairs design consists of just one subject, who gets both treatments one after the other. Each subject serves as his or her own control. The order of the treatments can influence the subject’s response, so we randomize the order for each subject.



Royalty Free/CORBIS

EXAMPLE 9.9 Cell phones and driving

Does talking on a hands-free cell phone distract drivers? Undergraduate students “drove” in a high-fidelity driving simulator equipped with a hands-free cell phone. The car ahead brakes: how quickly does the subject react? Let’s compare two designs for this experiment. There are 40 student subjects available.

In a *completely randomized design*, all 40 subjects are assigned at random, 20 to simply drive and the other 20 to talk on the cell phone while driving. In the *matched pairs design* that was actually used, all subjects drive both with and without using the cell phone. The two drives are on separate days to reduce carryover effects. The order of the two treatments is assigned at random: 20 subjects are chosen to drive first with the phone, and the remaining 20 drive first without the phone.¹¹

Some subjects naturally react faster than others. The completely randomized design relies on chance to distribute the faster subjects roughly evenly between the two groups. The matched pairs design compares each subject’s reaction time with and without the cell phone. This makes it easier to see the effects of using the phone. ■

Matched pairs designs use the principles of comparison of treatments and randomization. However, the randomization is not complete—we do not randomly assign all the subjects at once to the two treatments. Instead, we randomize only

within each matched pair. This allows matching to reduce the effect of variation among the subjects. Matched pairs are one kind of *block design*, with each pair forming a *block*.

BLOCK DESIGN

A **block** is a group of individuals that are known before the experiment to be similar in some way that is expected to affect the response to the treatments.

In a **block design**, the random assignment of individuals to treatments is carried out separately within each block.

A block design combines the idea of creating equivalent treatment groups by matching with the principle of forming treatment groups at random. Blocks are another form of *control*. They control the effects of some outside variables by bringing those variables into the experiment to form the blocks. Here are some typical examples of block designs.

EXAMPLE 9.10 Men, women, and advertising

Women and men respond differently to advertising. An experiment to compare the effectiveness of three advertisements for the same product will want to look separately at the reactions of men and women, as well as assess the overall response to the ads.

A *completely randomized design* considers all subjects, both men and women, as a single pool. The randomization assigns subjects to three treatment groups without regard to their sex. This ignores the differences between men and women. A *block design* considers women and men separately. Randomly assign the women to three groups, one to view each advertisement. Then separately assign the men at random to three groups. Figure 9.5 outlines this improved design. ■

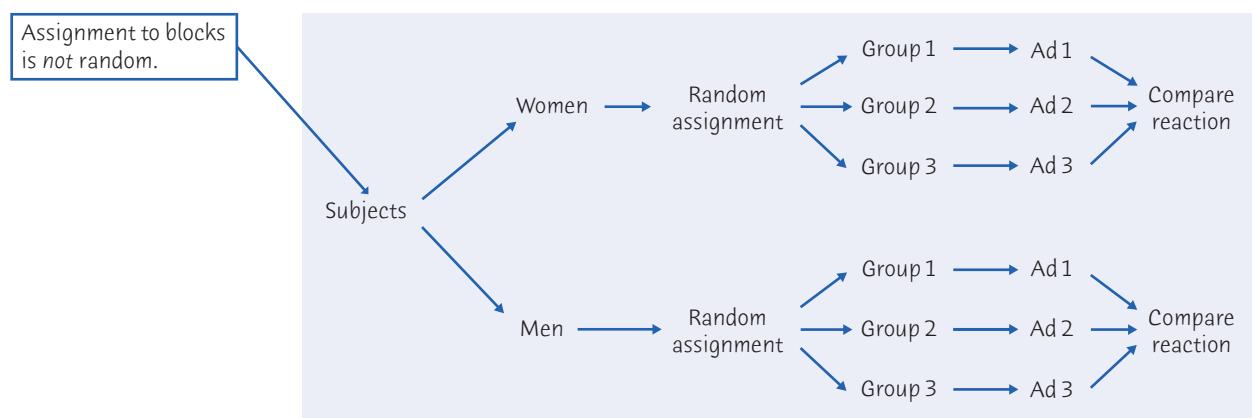


FIGURE 9.5

Outline of a block design, for Example 9.10. The blocks consist of male and female subjects. The treatments are three advertisements for the same product.



EXAMPLE 9.11 Comparing welfare policies

A social policy experiment will assess the effect on family income of several proposed new welfare systems and compare them with the present welfare system. Because the future income of a family is strongly related to its present income, the families who agree to participate are divided into blocks of similar income levels. The families in each block are then allocated at random among the welfare systems. ■

A block design allows us to draw separate conclusions about each block, for example, about men and women in Example 9.10. Blocking also allows more precise overall conclusions, because the systematic differences between men and women can be removed when we study the overall effects of the three advertisements. The idea of blocking is an important additional principle of statistical design of experiments. A wise experimenter will form blocks based on the most important unavoidable sources of variability among the subjects. Randomization will then average out the effects of the remaining variation and allow an unbiased comparison of the treatments.

Like the design of samples, the design of complex experiments is a job for experts. Now that we have seen a bit of what is involved, we will concentrate for the most part on completely randomized experiments.



APPLY YOUR KNOWLEDGE



Flirt/Superstock

9.16 Comparing breathing frequencies in swimming. Researchers from the United Kingdom studied the effect of two breathing frequencies on both performance times and several physiological parameters in front crawl swimming.¹² The breathing frequencies were one breath every second stroke (B2) and one breath every fourth stroke (B4). Subjects were 10 male collegiate swimmers. Each subject swam 200 meters, once with breathing frequency B2 and once on a different day with breathing frequency B4.

- Describe the design of this matched pairs experiment, including the randomization required by this design.
- Could this experiment be conducted using a completely randomized design? How would the design differ from the matched pairs experiment?
- Are there any problems with having swimmers choose their own breathing frequency and then swim 200 meters using their selected frequency?

9.17 How long did I work? A psychologist wants to know if the difficulty of a task influences our estimate of how long we spend working at it. She designs two sets of mazes that subjects can work through on a computer. One set has easy mazes and the other has hard mazes. Subjects work until told to stop (after six minutes, but subjects do not know this). They are then asked to estimate how long they worked. The psychologist has 30 students available to serve as subjects.

- Describe the design of a completely randomized experiment to learn the effect of difficulty on estimated time.
- Describe the design of a matched pairs experiment using the same 30 subjects.

9.18 Technology for teaching statistics. The Brigham Young University statistics department is performing randomized comparative experiments to compare teaching methods. Response variables include students' final-exam scores and a measure of their attitude toward statistics. One study compares two levels of technology for large lectures: standard (overhead projectors and chalk) and multimedia. The individuals in the study are the eight lectures in a basic statistics course. There are four instructors, each of whom teaches two lectures. Because the lecturers differ, their lectures form four blocks.¹³ Suppose the lectures and lecturers are as follows:

Lecture	Lecturer	Lecture	Lecturer
1	Hilton	5	Tolley
2	Christensen	6	Hilton
3	Hadfield	7	Tolley
4	Hadfield	8	Christensen

Outline a block design and do the randomization that your design requires.

CHAPTER 9 SUMMARY

CHAPTER SPECIFICS

- We can produce data intended to answer specific questions by **observational studies** or **experiments**. Sample surveys that select a part of a population of interest to represent the whole are one type of observational study. **Experiments**, unlike observational studies, actively impose some treatment on the subjects of the experiment.
- Variables are **confounded** when their effects on a response can't be distinguished from each other. Observational studies and uncontrolled experiments often fail to show that changes in an explanatory variable actually cause changes in a response variable because the explanatory variable is confounded with lurking variables.
- In an experiment, we impose one or more **treatments** on individuals, often called **subjects**. Each treatment is a combination of values of the explanatory variables, which we call **factors**.
- The **design** of an experiment describes the choice of treatments and the manner in which the subjects are assigned to the treatments. The basic principles of statistical design of experiments are **control** and **randomization** to combat bias and **using enough subjects** to reduce chance variation.
- The simplest form of control is **comparison**. Experiments should compare two or more treatments in order to avoid confounding of the effect of a treatment with other influences, such as lurking variables.
- **Randomization** uses chance to assign subjects to the treatments. Randomization creates treatment groups that are similar (except for chance variation) before the treatments are applied. Randomization and comparison together prevent bias, or systematic favoritism, in experiments.
- You can carry out randomization by using software or by giving numerical labels to the subjects and using a **table of random digits** to choose treatment groups.

- Applying each treatment to many subjects reduces the role of chance variation and makes the experiment more sensitive to differences among the treatments.
- Good experiments also require attention to detail as well as good statistical design. Many behavioral and medical experiments are **double-blind**. Some give a **placebo** to a control group. **Lack of realism** in an experiment can prevent us from generalizing its results.
- In addition to comparison, a second form of control is to restrict randomization by forming **blocks** of individuals that are similar in some way that is important to the response. Randomization is then carried out separately within each block.
- **Matched pairs** are a common form of blocking for comparing just two treatments. In some matched pairs designs, each subject receives both treatments in a random order. In others, the subjects are matched in pairs as closely as possible, and each subject in a pair receives one of the treatments.

LINK IT

Observational studies and experiments are two methods for producing data. Observational studies are useful when the conclusion involves describing a group or situation without disturbing the scene we observe. Sample surveys, discussed in Chapter 8, are an important type of observational study in which we draw conclusions about a population by observing only a part of the population (the sample). In contrast, experiments are used when the situation calls for a conclusion about whether a treatment *causes* a change in a response. The distinction between observational studies and experiments will be important when stating your conclusions in later chapters.

Only well-designed experiments provide a sound basis for concluding cause-and-effect relationships. In a simple comparative experiment, two treatments are imposed on two groups of individuals. Reaching the conclusion that the difference between the groups is caused by the treatments, rather than lurking variables, requires that the two groups of individuals be similar at the outset. A randomized comparative experiment is used to create groups that are similar. If there is a sufficiently large difference between the groups after imposing the treatments, we can say that the results are statistically significant and conclude that the differences in the response were *caused* by the treatments. In later chapters, the specific statistical procedures for reaching these conclusions will be described.

As with sampling, when conducting an experiment, attention to detail is important because our conclusions can be weakened by several factors. A lack of blinding can result in the expectations of the researcher influencing the results, while the placebo effect can confound the comparison between a treatment and a control group. In many instances, a more complex design is required to overcome difficulties and can produce more precise results.

CHECK YOUR SKILLS

9.19 The Nurses' Health Study has interviewed a sample of more than 100,000 female registered nurses every two years since 1976. The study finds that "light-to-moderate drinkers had a significantly lower risk of death" than either nondrinkers or heavy drinkers. The Nurses' Health Study is

- (a) an observational study.
- (b) an experiment.
- (c) Can't tell without more information.

9.20 What electrical changes occur in muscles as they get tired? Student subjects hold their arms above their shoulders until they have to drop them. Meanwhile, the electrical activity in their arm muscles is measured. This is

- (a) an observational study.
- (b) an uncontrolled experiment.
- (c) a randomized comparative experiment.

9.21 Do violence and sex in television programs help sell products in advertisements? Subjects were randomly assigned to watch one of four types of TV shows: (1) neither sex nor violence in the content code; (2) violence but no sex in the content code; (3) sex but no violence in the content code; and (4) both sex and violence in the content code. For each TV show, the original advertisements were replaced with the same set of twelve advertisements. Subjects were not told the purpose of the study but were instead told that the researchers were studying attitudes toward TV shows. After viewing the show, subjects received a surprise memory test to check their recall of the products advertised.¹⁴ This experiment has

- (a) four factors, the four TV shows being compared.
- (b) twelve factors, the advertisements being shown.
- (c) two factors, with/without violent content and with/without sexual content.

9.22 In the experiment of the previous exercise, the 336 subjects are labeled 001 to 336. Labels are selected at random by software, with the first 84 selected assigned to view TV show 1, the next 84 to view TV show 2, and the next 84 to view TV show 3. The 84 remaining subjects view TV show 4. This is a

- (a) matched pairs design because subjects are matched to the TV shows.
- (b) completely randomized design.
- (c) block design with TV shows representing the four blocks.

9.23 In the experiment described in Exercise 9.21,

- (a) it would have been better to have subjects choose the type of TV show they preferred to view in order to improve their recall and reduce confounding.
- (b) the score on the memory test of their recall of advertisements is the response.
- (c) the experimenters should have used different advertisements for each type of TV show in order to reduce confounding.

9.24 A medical experiment compares an antidepressant medicine with a placebo for relief of chronic headaches. There are 36 headache patients available to serve as subjects. To choose 18 patients to receive the medicine, you would

- (a) assign labels 01 to 36 and use Table B or a random number generator to choose 18.
- (b) assign labels 01 to 18, because only 18 need to be chosen.
- (c) assign the first 18 who signed up to get the medicine.

9.25 The Community Intervention Trial for Smoking Cessation asked whether a community-wide advertising campaign would reduce smoking. The researchers located 11 pairs of communities, each pair similar in location, size, economic status, and so on. One community in each pair participated in the advertising campaign and the other did not. This is

- (a) an observational study.
- (b) a matched pairs experiment.
- (c) a completely randomized experiment.

9.26 To decide which community in each pair in the previous exercise should get the advertising campaign, it is best to

- (a) toss a coin.
- (b) choose the community that will help pay for the campaign.
- (c) choose the community with a mayor who will participate.

9.27 A marketing class designs two videos advertising an expensive Mercedes sports car. They test the videos by asking fellow students to view both (in random order) and say which makes them more likely to buy the car. Mercedes should be reluctant to agree that the video favored in this study will sell more cars because

- (a) the study used a matched pairs design instead of a completely randomized design.
- (b) results from students may not generalize to the older and richer customers who might buy a Mercedes.
- (c) this is an observational study, not an experiment.

CHAPTER 9 EXERCISES

 In all exercises that require randomization, you may use Table B, the Simple Random Sample applet, or other software. See Example 9.6 (page 230) for directions on using the applet for more than two treatment groups.

9.28 Alcohol and heart attacks. Many studies have found that people who drink alcohol in moderation have lower risk of heart attacks than either nondrinkers or heavy drinkers. Does alcohol consumption also improve survival after a heart attack? One study followed 1913 people who were hospitalized after severe heart attacks. In the year before their heart

attacks, 47% of these people did not drink, 36% drank moderately, and 17% drank heavily. After four years, fewer of the moderate drinkers had died.¹⁵

- (a) Is this an observational study or an experiment? Why? What are the explanatory and response variables?
- (b) Suggest some lurking variables that may be confounded with the drinking habits of the subjects. The possible confounding makes it difficult to conclude that drinking habits explain death rates.

9.29 Reducing nonresponse. How can we reduce the rate of refusals in telephone surveys? Most people who answer at all listen to the interviewer's introductory remarks and then decide whether to continue. One study made telephone calls to randomly selected households to ask opinions about the next election. In some calls, the interviewer gave her name, in others she identified the university she was representing, and in still others she identified both herself and the university. The study recorded what percent of each group of interviews was completed. Is this an observational study or an experiment? Why? What are the explanatory and response variables?

9.30 Samples versus experiments. Give an example of a question about college students, their behavior, or their opinions that would best be answered by

- (a) a sample survey. (b) an experiment.

9.31 Observation versus experiment. Observational studies have suggested that vitamin E reduces the risk of heart disease. Careful experiments, however, showed that vitamin E has no effect. According to a commentary in the *Journal of the American Medical Association*:

Thus, vitamin E enters the category of therapies that were promising in epidemiologic and observational studies but failed to deliver in adequately powered randomized controlled trials. As in other studies, the “healthy user” bias must be considered; i.e., the healthy lifestyle behaviors that characterize individuals who care enough about their health to take various supplements are actually responsible for the better health, but this is minimized with the rigorous trial design.¹⁶

A friend who knows no statistics asks you to explain this.

- (a) What is the difference between observational studies and experiments?
 (b) What is a “randomized controlled trial”? (We’ll discuss “adequately powered” in Chapter 16.)
 (c) How does “healthy user bias” explain how people who take vitamin E supplements have better health in observational studies but not in controlled experiments?

9.32 Attitudes toward homeless people. Negative attitudes toward poor people are common. Are attitudes more negative when a person is homeless? To find out, a description of a poor person is read to subjects. There are two versions of this description. One begins

Jim is a 30-year-old single man. He is currently living in a small single-room apartment.

The other description begins

Jim is a 30-year-old single man. He is currently homeless and lives in a shelter for homeless people.

Otherwise, the descriptions are the same. After reading the description, you ask subjects what they believe about Jim and what they think should be done to help him. The subjects are 544 adults interviewed by telephone.¹⁷ Outline the design of this experiment.

9.33 Getting teachers to come to school. Elementary schools in rural India are usually small, with a single teacher. The teachers often fail to show up for work. Here is an idea for improving attendance: give the teacher a digital camera with a tamper-proof time and date stamp and ask a student to take a photo of the teacher and class at the beginning and end of the day. Offer the teacher better pay for good attendance—verified by the photos. Will this work? A randomized comparative experiment started with 120 rural schools in Rajasthan and assigned 60 to this treatment and 60 to a control group. Random checks for teacher attendance showed that 21% of teachers in the treatment group were absent, as opposed to 42% in the control group.¹⁸

- (a) Outline the design of this experiment.
 (b) Label the schools and choose the first 10 schools for the treatment group. If you use Table B, start at line 108.

9.34 Marijuana and work. How does smoking marijuana affect willingness to work? Canadian researchers persuaded young adult men who used marijuana to live for 98 days in a “planned environment.” The men earned money by weaving belts. They used their earnings to pay for meals and other consumption and could keep any money left over. One group smoked two potent marijuana cigarettes every evening. The other group smoked two weak marijuana cigarettes. All subjects could buy more cigarettes but were given strong or weak cigarettes depending on their group. Did the weak and strong groups differ in work output and earnings?¹⁹

- (a) Outline the design of this experiment.
 (b) Here are the names of the 30 subjects. Use software or Table B at line 120 to carry out the randomization your design requires.

Abel	DeVorce	Kennedy	Reichert	Stout
Aeffner	Fleming	Lamone	Riddle	Williams
Birkel	Fritz	Mani	Sawant	Wilson
Bower	Giriunas	Mattos	Scannell	Worbis
Burke	Glosup	Molnar	Sheldon	Zaccari
Deis	Heaton	Newlen	Simmons	Zelaski

- (c) Do you think this can be run as a double-blind experiment? Explain.

9.35 The benefits of red wine. Some people think that red wine protects moderate drinkers from heart disease better than other alcoholic beverages. This calls for a randomized

comparative experiment. The subjects were healthy men aged 35 to 65. They were randomly assigned to drink red wine (9 subjects), drink white wine (9 subjects), drink white wine and also take polyphenols from red wine (6 subjects), take polyphenols alone (9 subjects), or drink vodka and lemonade (6 subjects).²⁰ Outline the design of the experiment and randomly assign the 39 subjects to the 5 groups. If you use Table B, start at line 107.

9.36 Can low-fat food labels lead to obesity?

What are the effects of low-fat food labels on food consumption?



Marcia Baker

Do people eat more of a snack food when the food is labeled as low-fat? The answer may depend both on whether the snack food is labeled low-fat and whether the label includes serving-size information. An experiment investigated this question using

university staff, graduate students, and undergraduate students at a large university as subjects. Subjects were asked to evaluate a pilot episode for an upcoming TV show in a theater on campus and were given a cold 24-ounce bottle of water and a bag of granola from a respected campus restaurant called The Spice Box. They were told to enjoy as much or as little of the granola as they wanted. Depending on the condition randomly assigned to the subjects, the granola was labeled as either "Regular Rocky Mountain Granola" or "Low-Fat Rocky Mountain Granola." Below this, the label indicated "Contains 1 Serving," or "Contains 2 Servings," or it provided no serving-size information.²¹ Twenty subjects are assigned to each treatment, and their granola bags were weighed at the end of the session to determine how much granola was eaten.

- (a) What are the factors and the treatments? How many subjects does the experiment require?
- (b) Outline a completely randomized design for this experiment. (You need not actually do the randomization.)

9.37 Relieving headaches. Doctors identify "chronic tension-type headaches" as headaches that occur almost daily for at least six months. Can antidepressant medications or stress management training reduce the number and severity of these headaches? Are both together more effective than either alone?

- (a) Use a diagram like Figure 9.2 (page 226) to display the treatments in a design with two factors: "medication, yes or no" and "stress management, yes or no." Then outline the design of a completely randomized experiment to compare these treatments.
- (b) The following headache sufferers have agreed to participate in the study. Randomly assign the subjects to the treatments.

If you use the *Simple Random Sample* applet or other software, assign all the subjects. If you use Table B, start at line 125 and assign subjects to only the first treatment group.

Abbott	Decker	Herrera	Lucero	Richter
Abdalla	Devlin	Hersch	Masters	Riley
Alawi	Engel	Hurwitz	Morgan	Samuels
Broden	Fuentes	Irwin	Nelson	Smith
Chai	Garrett	Jiang	Nho	Suarez
Chuang	Gill	Kelley	Ortiz	Upasani
Cordoba	Glover	Kim	Ramdas	Wilson
Custer	Hammond	Landers	Reed	Xiang

Treating sinus infections. Sinus infections are common, and doctors commonly treat them with antibiotics. Another treatment is to spray a steroid solution into the nose. A well-designed clinical trial found that these treatments, alone or in combination, do not reduce the severity or the length of sinus infections.²² Exercises 9.38 to 9.40 concern this trial.

9.38 Experimental design. The clinical trial was a completely randomized experiment that assigned 240 patients at random among four treatments as follows:

	Antibiotic pill	Placebo pill
Steroid spray	53	64
Placebo spray	60	63

- (a) Outline the design of the experiment.
- (b) How will you label the 240 subjects?
- (c) Explain briefly how you would do the random assignment of patients to treatments. Assign the first 5 patients who will receive the first treatment.

9.39 Describing the design. The report of this study in the *Journal of the American Medical Association* describes it as a "double-blind, randomized, placebo-controlled factorial trial." "Factorial" means that the treatments are formed from more than one factor. What are the factors? What do "double-blind" and "placebo-controlled" mean?

9.40 Checking the randomization. If the random assignment of patients to treatments did a good job of eliminating bias, possible lurking variables such as smoking history, asthma, and hay fever should be similar in all four groups. After recording and comparing many such variables, the investigators said that "all showed no significant difference between groups." Explain to someone who knows no statistics what "no significant difference" means. Does it mean that the presence of all these variables was exactly the same in all four treatment groups?

9.41 Frappuccino light? Here's the opening of a Starbucks press release: "Starbucks Corp. on Monday said it would roll out a line of blended coffee drinks intended to tap into the growing popularity of reduced-calorie and reduced-fat menu choices for Americans." You wonder if Starbucks customers like the new "Mocha Frappuccino Light" as well as the regular Mocha Frappuccino coffee.

- (a) Describe a matched pairs design to answer this question. Be sure to include proper blinding of your subjects. What is your response variable going to be?
- (b) You have 20 regular Starbucks customers on hand. Use the *Simple Random Sample* applet or Table B at line 141 to do the randomization that your design requires.

9.42 Growing trees faster. The concentration of carbon dioxide (CO_2) in the atmosphere is increasing rapidly due to our use of fossil fuels. Because green plants use CO_2 to fuel photosynthesis, more CO_2 may cause trees to grow faster. An elaborate apparatus allows researchers to pipe extra CO_2 to a 30-meter circle of forest. We want to compare the growth in base area of trees in treated and untreated areas to see if extra CO_2 does in fact increase growth. We can afford to treat three circular areas.²³

- (a) Describe the design of a completely randomized experiment using six well-separated 30-meter circular areas in a pine forest. Sketch the circles and carry out the randomization your design calls for.
- (b) Areas within the forest may differ in soil fertility. Describe a matched pairs design using three pairs of circles that will reduce the extra variation due to different fertility. Sketch the circles and carry out the randomization your design calls for.

9.43 Athletes taking oxygen. We often see players on the sidelines of a football game inhaling oxygen. Their coaches think this will speed their recovery. We might measure recovery from intense exertion as follows: Have a

football player run 100 yards three times in quick succession. Then allow three minutes to rest before running 100 yards again. Time the final run. Because players vary greatly in speed, you plan a matched pairs experiment using 25 football players as subjects. Discuss the design of such an experiment to investigate the effect of inhaling oxygen during the rest period.



Wade Payne/AP Photos

9.44 Protecting ultramarathon runners. An ultramarathon, as you might guess, is a footrace longer than the 26.2 miles of a marathon. Runners commonly develop respiratory infections after an ultramarathon. Will taking 600 milligrams of vitamin C daily reduce these infections? Researchers randomly assigned ultramarathon runners to receive either vitamin C or a placebo. Separately, they also randomly assigned these treatments to a group of nonrunners the same age as the runners. All subjects were watched for 14 days after the big race to see if infections developed.²⁴

- (a) What is the name for this experimental design?
- (b) Use a diagram to outline the design.

9.45 Wine, beer, or spirits? There is good evidence that moderate alcohol use improves health. Some people think that red wine is better for your health than other alcoholic drinks. You have recruited 300 adults aged 45 to 65 who are willing to follow your orders about alcohol consumption over the next five years. You want to compare the effects on heart disease of moderate drinking of just red wine, just beer, or just spirits. Outline the design of a completely randomized experiment to do this. (No such experiment has been done because subjects aren't willing to have their drinking regulated for years.)

9.46 Wine, beer, or spirits? Women as a group develop heart disease much later than men. We can improve the completely randomized design of Exercise 9.45 by using women and men as blocks. Your 300 subjects include 120 women and 180 men. Outline a block design for comparing wine, beer, and spirits. Be sure to say how many subjects you will put in each group in your design.

9.47 Quick randomizing. Here's a quick and easy way to randomize. You have 100 subjects, 50 women and 50 men. Toss a coin. If it's heads, assign all the men to the treatment group and all the women to the control group. If the coin comes up tails, assign all the women to treatment and all the men to control. This gives every individual subject a 50-50 chance of being assigned to treatment or control. Why isn't this a good way to randomly assign subjects to treatment groups?

9.48 Do antioxidants prevent cancer? People who eat lots of fruits and vegetables have lower rates of colon cancer than those who eat little of these foods. Fruits and vegetables are rich in "antioxidants" such as vitamins A, C, and E. Will taking antioxidants help prevent colon cancer? A medical experiment studied this question with 864 people who were at risk of colon cancer. The subjects were divided into four groups: daily beta-carotene, daily vitamins C and E, all three vitamins every day, or daily placebo. After four years, the researchers were surprised to find no significant difference in colon cancer among the groups.²⁵

- (a) What are the explanatory and response variables in this experiment?

- (b) Outline the design of the experiment. Use your judgment in choosing the group sizes.
- (c) The study was double-blind. What does this mean?
- (d) What does “no significant difference” mean in describing the outcome of the study?
- (e) Suggest some lurking variables that could explain why people who eat lots of fruits and vegetables have lower rates of colon cancer. The experiment suggests that these variables, rather than the antioxidants, may be responsible for the observed benefits of fruits and vegetables.

9.49 An herb for depression? Does the herb Saint-John’s-wort relieve major depression? Here are some excerpts from the report of a study of this issue.²⁶ The study concluded that the herb is no more effective than a placebo.

- (a) “Design: Randomized, double-blind, placebo-controlled clinical trial....” A clinical trial is a medical experiment using actual patients as subjects. Explain the meaning of each of the other terms in this description.

Organic Image Library/Alamy



EXPLORING THE WEB

9.51 Smoking cessation. Go to the *New England Journal of Medicine* Web site, www.nejm.org, and find the article “A Randomized, Controlled Trial of Financial Incentives for Smoking Cessation” by Volpp et al. in the February 12, 2009, issue. Under the “ISSUES” link, you need to go to the “Browse full index” link and then to the February 12, 2009, issue. You can then download the pdf of the article for free. Was this a comparative study? Was randomization used? How many subjects took part? There were 22 subjects in the control group and 64 in the incentive group who were still not smoking six months after they stopped. What were the percents in each group? This difference is statistically significant. Explain in simple language what this means.

9.52 Find an experiment. You can find the latest medical research in the *Journal of the American Medical Association* at www.jama.ama-assn.org and the *New England Journal of Medicine* at www.nejm.org. Many of the articles describe randomized comparative experiments and use the language of statistical significance when giving conclusions. Look through the abstracts and find an experiment of interest to you. If your institution has a subscription to these journals, you should be able to view the entire article. Otherwise, use the information in the abstract to answer as many of these questions as you can. What was the purpose of the experiment? How many factors were in the experiment, and what were the levels of the factors? What response(s) were measured? How many subjects were assigned to each of the treatments, and was randomization used? Was it a double-blind experiment? What were the conclusions, and were the results statistically significant?

- (b) “Participants ... were randomly assigned to receive either Saint-John’s-wort extract ($n = 98$) or placebo ($n = 102$).... The primary outcome measure was the rate of change in the Hamilton Rating Scale for Depression over the treatment period.” Based on this information, use a diagram to outline the design of this clinical trial.

9.50 Randomization avoids bias. Suppose that the 25 even-numbered students among the 50 students available for the comparison of classroom and online instruction (Example 9.5) are older, employed students. We hope that randomization will distribute these students roughly equally between the classroom and online groups. Use the *Simple Random Sample* applet to take 20 samples of size 25 from the 50 students. These 25 students will be the classroom instruction group. (Be sure to click “Reset” after each sample.) Record the counts of even-numbered students in each of your 20 samples.

- (a) How many older students would you expect to see in the classroom instruction group?
- (b) You see that there is considerable chance variation in the number of older (even-numbered) students assigned to the classroom group. Draw a stem-and-leaf plot of the number of older students assigned to the classroom group. Do you see any systematic bias in favor of one or the other group being assigned the older students? Larger samples from a larger population will, on the average, do an even better job of creating two similar groups.



Commentary: Data Ethics*

The production and use of data, like all human endeavors, raise ethical questions. We won't discuss the telemarketer who begins a telephone sales pitch with "I'm conducting a survey." Such deception is clearly unethical. It enrages legitimate survey organizations, which find the public less willing to talk with them. Neither will we discuss those few researchers who, in the pursuit of professional advancement, publish fake data. There is no ethical question here—faking data to advance your career is just wrong. It will end your career when uncovered. But just how honest must researchers be about real, unfaked data? Here is an example that suggests the answer is "More honest than they often are."

IN THIS COMMENTARY WE COVER...

- Institutional review boards
- Informed consent
- Confidentiality
- Clinical trials
- Behavioral and social science experiments

EXAMPLE 1 The whole truth?

Papers reporting scientific research are supposed to be short, with no extra baggage. Brevity, however, can allow researchers to avoid complete honesty about their data. Did they choose their subjects in a biased way? Did they report data on only some of their subjects? Did they try several statistical analyses and report only the ones that looked best? The statistician John Bailar screened more than 4000 medical papers in more than a decade as consultant to the *New England Journal of Medicine*. He says, "When it came to the statistical review, it was often clear that critical information was lacking, and the gaps nearly always had the practical effect of making the authors' conclusions look stronger than they should have."¹ The situation is no doubt worse in fields that screen published work less carefully. ■

*This short essay concerns a very important topic, but the material is not needed to read the rest of the book.

The most complex issues of data ethics arise when we collect data from people. The ethical difficulties are more severe for experiments that impose some treatment on people than for sample surveys that simply gather information. Trials of new medical treatments, for example, can do harm as well as good to their subjects. Here are some basic standards of data ethics that must be obeyed by all studies that gather data from human subjects, both observational studies and experiments.

BASIC DATA ETHICS

All planned studies must be reviewed in advance by an **institutional review board** charged with protecting the safety and well-being of the subjects.

All individuals who are subjects in a study must give their **informed consent** before data are collected.

All individual data must be kept **confidential**. Only statistical summaries for groups of subjects may be made public.

The law requires that studies carried out or funded by the federal government obey these principles.² But neither the law nor the consensus of experts is completely clear about the details of their application.

INSTITUTIONAL REVIEW BOARDS

The purpose of an institutional review board is not to decide whether a proposed study will produce valuable information or whether it is statistically sound. The board's purpose is, in the words of one university's board, "to protect the rights and welfare of human subjects (including patients) recruited to participate in research activities." The board reviews the plan of the study and can require changes. It reviews the consent form to ensure that subjects are informed about the nature of the study and about any potential risks. Once research begins, the board monitors the study's progress at least once a year.

The most pressing issue concerning institutional review boards is whether their workload has become so large that their effectiveness in protecting subjects drops. When the government temporarily stopped human subject research at Duke University Medical Center in 1999 due to inadequate protection of subjects, more than 2000 studies were going on. That's a lot of review work. There are shorter review procedures for projects that involve only minimal risks to subjects, such as most sample surveys. When a board is overloaded, there is a temptation to put more proposals in the minimal-risk category to speed the work.

INFORMED CONSENT

Both words in the phrase "informed consent" are important, and both can be controversial. Subjects must be *informed* in advance about the nature of a study and any risk of harm it may bring. In the case of a sample survey, physical harm

Institutional Review Board (IRB)

Home
More Pages
Education and Training
Federalwide Assurance
Glossary of Terms
Policy Manual

Give Touch Transform
Support Mayo Now

Institutional Board (IRB)

Director, Mayo Clinic Office of Human Research Protection
William J Tramaine, M.D.

Overview
The Mayo Clinic Institutional Review Board (IRB) reviews all human subject research conducted at Mayo Clinic Florida (MCF), Mayo Clinic Rochester (MCR), or Mayo Clinic Arizona (MCA) and research conducted at other facilities under the direction of MCF, MCR, or MCA staff. A guarantee that all human subject research at Mayo will be reviewed by the IRB has been given to the U.S. Department of Health and Human Services (HHS) in a Federalwide Assurance (FWA00005001).

[Read More](#)

Mission
The primary mission of Mayo Clinic's IRB is to ensure the protection of rights, privacy and welfare of all human participants in research programs conducted by Mayo Clinic and associated faculty, professional staff, and students. Consistent with participant protection is the goal of providing quality service to enhance the conduct of research. To achieve this goal, the IRB has the authority to review, approve, modify or disapprove research protocols submitted by faculty, staff and student investigators. The IRB review process is guided by federal rules and regulations, and is based on the Protection of Human Subject Code of Federal Regulations, the Belmont Report and provisions of 45 CFR 46 – Protection of Human Subjects requiring institutions receiving federal funds to have all research involving human participants be approved by an IRB.

Related Resources
Food and Drug Administration (FDA)
Guidance for Institutional Review Boards and Clinical Investigators (FDA)
Office for Human Research Protections (OHRP)
National Institutes of Health (NIH)

The Web page of the Mayo Clinic's institutional review board. It begins by describing the job of such boards.

is not possible. The subjects should be told what kinds of questions the survey will ask and about how much of their time it will take. Experimenters must tell subjects the nature and purpose of the study and outline possible risks. Subjects must then consent in writing.

EXAMPLE 2 Who can consent?

Are there some subjects who can't give informed consent? It was once common, for example, to test new vaccines on prison inmates who gave their consent in return for



Bernardo Bucci/CORBIS

good-behavior credit. Now we worry that prisoners are not really free to refuse, and the law forbids almost all medical research in prisons.

Children can't give fully informed consent, so the usual procedure is to ask their parents. A study of new ways to teach reading is about to start at a local elementary school, so the study team sends consent forms home to parents. Many parents don't return the forms. Can their children take part in the study because the parents did not say "No," or should we allow only children whose parents returned the form and said "Yes"?

What about research into new medical treatments for people with mental disorders? What about studies of new ways to help emergency room patients who may be unconscious? In most cases, there is not time to get the consent of the family. Does the principle of informed consent bar realistic trials of new treatments for unconscious patients?

These are questions without clear answers. Reasonable people differ strongly on all of them. There is nothing simple about informed consent.³ ■

The difficulties of informed consent do not vanish even for capable subjects. Some researchers, especially in medical trials, regard consent as a barrier to getting patients to participate in research. They may not explain all possible risks; they may not point out that there are other therapies that might be better than those being studied; they may be too optimistic in talking with patients even when the consent form has all the right details. On the other hand, mentioning every possible risk leads to very long consent forms that really are barriers. "They are like rental car contracts," one lawyer said. Some subjects don't read forms that run five or six printed pages. Others are frightened by the large number of possible (but unlikely) disasters that might happen and so refuse to participate. Of course, unlikely disasters sometimes happen. When they do, lawsuits follow and the consent forms become yet longer and more detailed.

CONFIDENTIALITY

Ethical problems do not disappear once a study has been cleared by the review board, has obtained consent from its subjects, and has actually collected data about the subjects. It is important to protect the subjects' privacy by keeping all data about individuals confidential. The report of an opinion poll may say what percent of the 1200 respondents felt that legal immigration should be reduced. It may not report what you said about this or any other issue.

anonymity

Confidentiality is not the same as **anonymity**. Anonymity means that subjects are anonymous—their names are not known even to the director of the study. Anonymity is rare in statistical studies. Even where it is possible (mainly in surveys conducted by mail), anonymity prevents any follow-up to improve non-response or inform subjects of results.

Any breach of confidentiality is a serious violation of data ethics. The best practice is to separate the identity of the subjects from the rest of the data at once. Sample surveys, for example, use the identification only to check on who did or did not respond. In an era of advanced technology, however, it is no longer

The screenshot shows the "Internet Privacy Policy" page of the Social Security Administration's website. At the top, there's a navigation bar with links like "Home", "About SSA", "Contact Us", and "Feedback". Below the navigation, a large red banner features the text "Our Commitment to You". The main content area is titled "Social Notice" and contains several sections: "What We Automatically Collect Online", "Other Information We May Collect", "Why We Collect Personal Information", "Sharing Your Information", "How We Use Your Personal Information", and "DISCLAIMER". Each section has a detailed description of the data collected and how it is used.

The privacy policy of the government's Social Security Administration Web site.

enough to be sure that each individual set of data protects people's privacy. The government, for example, maintains a vast amount of information about citizens in many separate data bases—census responses, tax returns, Social Security information, data from surveys such as the Current Population Survey, and so on. Many of these data bases can be searched by computers for statistical studies. A clever computer search of several data bases might be able, by combining information, to identify you and learn a great deal about you even if your name and other identification have been removed from the data available for search. A colleague from Germany once remarked that “female full professor of statistics with a PhD from the United States” was enough to identify her among all the 83 million residents of Germany. Privacy and confidentiality of data are hot issues among statisticians in the computer age.

EXAMPLE 3 Uncle Sam knows

Citizens are required to give information to the government. Think of tax returns and Social Security contributions. The government needs these data for administrative purposes—to see if you paid the right amount of tax and how large a Social Security benefit you are owed when you retire. Some people feel that individuals should be able to forbid any other use of their data, even with all identification removed. This would prevent using government records to study, say, the ages, incomes, and household sizes of Social Security recipients. Such a study could well be vital to debates on reforming Social Security. ■

CLINICAL TRIALS

Clinical trials are experiments that study the effectiveness of medical treatments on actual patients. Medical treatments can harm as well as heal, so clinical trials spotlight the ethical problems of experiments with human subjects. Here are the starting points for a discussion:

- Randomized comparative experiments are the only way to see the true effects of new treatments. Without them, risky treatments that are no more effective than placebos will become common.
- Clinical trials produce great benefits, but most of these benefits go to future patients. The trials also pose risks, and these risks are borne by the subjects of the trial. So we must balance future benefits against present risks.
- Both medical ethics and international human rights standards say that “the interests of the subject must always prevail over the interests of science and society.”

The quoted words are from the 1964 Helsinki Declaration of the World Medical Association, the most respected international standard. The most outrageous examples of unethical experiments are those that ignore the interests of the subjects.

EXAMPLE 4 The Tuskegee study

In the 1930s, syphilis was common among black men in the rural South, a group that had almost no access to medical care. The Public Health Service Tuskegee study recruited 399 poor black sharecroppers with syphilis and 201 others without the disease in order to observe how syphilis progressed when no treatment was given. Beginning in 1943, penicillin became available to treat syphilis. The study subjects were not treated. In fact, the Public Health Service prevented any treatment until word leaked out and forced an end to the study in the 1970s.

The Tuskegee study is an extreme example of investigators following their own interests and ignoring the well-being of their subjects. A 1996 review said, “It has come to symbolize racism in medicine, ethical misconduct in human research, paternalism by physicians, and government abuse of vulnerable people.” In 1997, President Clinton formally apologized to the surviving participants in a White House ceremony.⁴

Because “the interests of the subject must always prevail,” medical treatments can be tested in clinical trials only when there is reason to hope that they will help the patients who are subjects in the trials. Future benefits aren’t enough to justify experiments with human subjects. Of course, if there is already strong evidence that a treatment works and is safe, it is unethical not to give it. Here are the words of Dr. Charles Hennekens of the Harvard Medical School, who

directed the large clinical trial that showed that aspirin reduces the risk of heart attacks:

There's a delicate balance between when to do or not do a randomized trial. On the one hand, there must be sufficient belief in the agent's potential to justify exposing half the subjects to it. On the other hand, there must be sufficient doubt about its efficacy to justify withholding it from the other half of subjects who might be assigned to placebos.⁵

Why is it ethical to give a control group of patients a placebo? Well, we know that placebos often work. Moreover, placebos have no harmful side effects. So in the state of balanced doubt described by Dr. Hennekens, the placebo group may be getting a better treatment than the drug group. If we knew which treatment was better, we would give it to everyone. When we don't know, it is ethical to try both and compare them.

BEHAVIORAL AND SOCIAL SCIENCE EXPERIMENTS

When we move from medicine to the behavioral and social sciences, the direct risks to experimental subjects are less acute, but so are the possible benefits to the subjects. Consider, for example, the experiments conducted by psychologists in their study of human behavior.

EXAMPLE 5 Psychologists in the men's room

Psychologists observe that people have a “personal space” and are uneasy if others come too close to them. We don’t like strangers to sit at our table in a coffee shop if other tables are available, and we see people move apart in elevators if there is room to do so. Americans tend to require more personal space than people in most other cultures. Can violations of personal space have physical, as well as emotional, effects?

Investigators set up shop in a men’s public restroom. They blocked off urinals to force men walking in to use either a urinal next to an experimenter (treatment group) or a urinal separated from the experimenter (control group). Another experimenter, using a periscope from a toilet stall, measured how long the subject took to start urinating and how long he continued.⁶ ■



David Pollack/CORBIS

This personal space experiment illustrates the difficulties facing those who plan and review behavioral studies.

- There is no risk of harm to the subjects, although they would certainly object to being watched through a periscope. What should we protect subjects from when physical harm is unlikely? Possible emotional harm? Undignified situations? Invasion of privacy?
- What about informed consent? The subjects did not even know they were participating in an experiment. Many behavioral experiments rely on hiding

the true purpose of the study. The subjects would change their behavior if told in advance what the investigators were looking for. Subjects are asked to consent on the basis of vague information. They receive full information only after the experiment.

The “Ethical Principles” of the American Psychological Association require consent unless a study merely observes behavior in a public place. They allow deception only when it is necessary to the study, does not hide information that might influence a subject’s willingness to participate, and is explained to subjects as soon as possible. The personal space study (from the 1970s) does not meet current ethical standards.

We see that the basic requirement for informed consent is understood differently in medicine and psychology. Here is an example of another setting with yet another interpretation of what is ethical. The subjects get no information and give no consent. They don’t even know that an experiment may be sending them to jail for the night.

EXAMPLE 6 Reducing domestic violence

How should police respond to domestic violence calls? In the past, the usual practice was to remove the offender and order him to stay out of the household overnight. Police were reluctant to make arrests because the victims rarely pressed charges. Women’s groups argued that arresting offenders would help prevent future violence even if no charges were filed. Is there evidence that arrest will reduce future offenses? That’s a question that experiments have tried to answer.

A typical domestic violence experiment compares two treatments: arrest the suspect and hold him overnight, or warn the suspect and release him. When police officers reach the scene of a domestic violence call, they calm the participants and investigate. Weapons or death threats require an arrest. If the facts permit an arrest but do not require it, an officer radios headquarters for instructions. The person on duty opens the next envelope in a file prepared in advance by a statistician. The envelopes contain the treatments in random order. The police either arrest the suspect or warn and release him, depending on the contents of the envelope. The researchers then watch police records and visit the victim to see if the domestic violence reoccurs.

Such experiments show that arresting domestic violence suspects does reduce their future violent behavior.⁷ As a result of this evidence, arrest has become the common police response to domestic violence. ■

The domestic violence experiments shed light on an important issue of public policy. Because there is no informed consent, the ethical rules that govern clinical trials and most social science studies would forbid these experiments. They were cleared by review boards because, in the words of one domestic violence researcher, “These people became subjects by committing acts that allow the police to arrest them. You don’t need consent to arrest someone.”

DISCUSSION EXERCISES

Most of these exercises pose issues for discussion. There are no right or wrong answers, but there are more and less thoughtful answers.

1. Minimal risk? You are a member of your college's institutional review board. You must decide whether several research proposals qualify for less rigorous review because they involve only minimal risk to subjects. Federal regulations say that "minimal risk" means the risks are no greater than "those ordinarily encountered in daily life or during the performance of routine physical or psychological examinations or tests." That's vague. Which of these do you think qualifies as "minimal risk"?

- (a) Take hair and nail clippings in a nondisfiguring manner.
- (b) Draw a drop of blood by pricking a finger in order to measure blood sugar.
- (c) Draw blood from the arm for a full set of blood tests.
- (d) Insert a tube that remains in the arm so that blood can be drawn regularly.
- (e) Take extra specimens from a subject who is undergoing an invasive clinical procedure such as a bronchoscopy (a procedure in which a physician views the inside of the airways for diagnostic and therapeutic purposes using an instrument that is inserted into the airways, usually through the nose or mouth).

2. Who reviews? Government regulations require that institutional review boards consist of at least five people, including at least one scientist, one nonscientist, and one person from outside the institution. Most boards are larger, but many contain just one outsider.

- (a) Why should review boards contain people who are not scientists?
- (b) Do you think that one outside member is enough? How would you choose that member? (For example, would you prefer a medical doctor? A member of the clergy? An activist for patients' rights?)

3. Informed consent. A researcher suspects that people with ultraliberal political beliefs tend to be more prone to depression. She prepares a questionnaire that measures depression and also asks many political questions. Write a description of the purpose of this research to be read by subjects in order to obtain their informed consent. You must balance the conflicting goals of not deceiving the subjects as to what the questionnaire will tell about them and of not biasing the sample by scaring off people with ultraliberal political views.

4. Is consent needed? In which of the circumstances below would you allow collecting personal information without the subjects' consent?

(a) A government agency takes a random sample of income tax returns to obtain information on the marital status and average income of people who identify themselves as clergy. Only the marital status and income are recorded from the returns, not the names.

(b) A social psychologist attends public meetings of a religious group to study the behavior patterns of members.

(c) A social psychologist pretends to be converted to membership in a religious group and attends private meetings to study the behavior patterns of members.

5. Studying your blood. Long ago, doctors drew a blood specimen from you as part of treating minor anemia. Unknown to you, the sample was stored. Now researchers plan to use stored samples from you and many other people to look for genetic factors that may influence anemia. It is no longer possible to ask your consent. Modern technology can read your entire genetic makeup from the blood sample.

(a) Do you think it violates the principle of informed consent to use your blood sample if your name is on it but you were not told that it might be saved and studied later?

(b) Suppose that your identity is not attached. The blood sample is known only to come from (say) "a 20-year-old white female being treated for anemia." Is it now OK to use the sample for research?

(c) Perhaps we should use biological materials such as blood samples only from patients who have agreed to allow the material to be stored for later use in research. It isn't possible to say in advance what kind of research, so this falls short of the usual standard for informed consent. Is it nonetheless acceptable, given complete confidentiality and the fact that using the sample can't physically harm the patient?

6. Anonymous? Confidential? One of the most important nongovernment surveys in the United States is the National Opinion Research Center's General Social Survey. The GSS regularly monitors public opinion on a wide variety of political and social issues. Interviews are conducted in person in the subject's home. Are a subject's responses to GSS questions anonymous, confidential, or both? Explain your answer.

7. Anonymous or confidential? The University of Wisconsin at Madison, like many universities, offers free screening for HIV, the virus that causes AIDS. The announcement at



Melba Melba/Photolibrary

the University Health Services Web site says that for persons who seek testing one option is the following. “A code is used instead of a name. The person tested receives a copy of the report for their own information, but only the code identifies the report as theirs. The report does not go into a medical record.” Does this practice offer anonymity or just confidentiality? Explain your answer.

8. Political polls. The presidential election campaign is in full swing, and the candidates have hired polling organizations to take sample surveys to find out what the voters think about the issues. What information should the pollsters be required to give out?

- (a) What does the standard of informed consent require the pollsters to tell potential respondents?
- (b) The standards accepted by polling organizations also require giving respondents the name and address of the organization that carries out the poll. Why do you think this is required?
- (c) The polling organization usually has a professional name such as “Samples Incorporated,” so respondents don’t know that the poll is being paid for by a political party or candidate. Would revealing the sponsor to respondents bias the poll? Should the sponsor always be announced whenever poll results are made public?

9. Making poll results public. Some people think that the law should require that all political poll results be made public. Otherwise, the possessors of poll results can use the information to their own advantage. They can act on the information, release only selected parts of it, or time the release for best effect. A candidate’s organization replies that they are paying for the poll in order to gain information for their own use, not to amuse the public. Do you favor requiring complete disclosure of political poll results? What about other private surveys, such as market research surveys of consumer tastes?

10. Student subjects. Students taking Psychology 001 are required to serve as experimental subjects. Students in Psychology 002 are not required to serve, but they are given extra credit if they do so. Students in Psychology 003 are required either to sign up as subjects or to write a term paper. Serving as an experimental subject may be educational, but current ethical standards frown on using “dependent subjects” such as prisoners or charity medical patients. Students are certainly somewhat dependent on their teachers. Do you object to any of these course policies? If so, which ones, and why?

11. The Willowbrook hepatitis studies. In the 1960s, children entering the Willowbrook State School, an institution for the mentally retarded, were deliberately infected with hepatitis. The researchers argued that almost all children in

the institution quickly became infected anyway. The studies showed for the first time that two strains of hepatitis existed. This finding contributed to the development of effective vaccines. Despite these valuable results, the Willowbrook studies are now considered an example of unethical research. Explain why, according to current ethical standards, useful results are not enough to allow a study.

12. Unequal benefits. Researchers on aging proposed to investigate the effect of supplemental health services on the quality of life of older people. Eligible patients on the rolls of a large medical clinic were to be randomly assigned to treatment and control groups. The treatment group would be offered hearing aids, dentures, transportation, and other services not available without charge to the control group. The review board felt that providing these services to some but not other persons in the same institution raised ethical questions. Do you agree?

13. How many have HIV? Researchers from Yale, working with medical teams in Tanzania, wanted to know how common infection with HIV, the virus that causes AIDS, is among pregnant women in that African country. To do this, they planned to test blood samples drawn from pregnant women.

Yale’s institutional review board insisted that the researchers get the informed consent of each woman and tell her the results of the test. This is the usual procedure in developed nations. The Tanzanian government did not want to tell the women why blood was drawn or tell them the test results. The government feared panic if many people turned out to have an incurable disease for which the country’s medical system could not provide care. The study was canceled. Do you think that Yale was right to apply its usual standards for protecting subjects?

14. AIDS trials in Africa. The drug programs that treat AIDS in rich countries are very expensive, so some African nations cannot afford to give them to large numbers of people. Yet AIDS is more common in parts of Africa than anywhere else. “Short-course” drug programs that are much less expensive might help, for example, in preventing infected pregnant women from passing the infection to their unborn children. Is it ethical to compare a short-course program with a placebo in a clinical trial? Some say “No”: this is a double standard, because in rich countries the full drug program would be the control treatment. Others say “Yes”: the intent is to find treatments that are practical in Africa, and the trial does not withhold any treatment that subjects would otherwise receive. What do you think?

15. Abandoned children in Romania. The study described in Example 9.2 randomly assigned abandoned children in Romanian orphanages to move to foster homes or to remain

in an orphanage. All of the children would otherwise have remained in an orphanage. The foster care was paid for by the study. There was no informed consent because the children had been abandoned and had no adult to speak for them. The experiment was considered ethical because “people who cannot consent can be protected by enrolling them only in minimal-risk research, whose risks do not exceed those of everyday life,” and because the study “aimed to produce results that would primarily benefit abandoned, institutionalized children.”⁸ Do you agree?

16. Asking teens about sex. The Centers for Disease Control and Prevention, in a survey of teenagers, asked the subjects if they had ever had sexual intercourse. Males who said “Yes” were then asked, “That very first time that you had sexual intercourse with a female, how old were you?” and “Please tell me the name or initials of your first sexual partner so that I can refer to her during the interview.” Should consent of parents be required to ask minors about sex, drugs, and other such issues, or is consent of the minors themselves enough? Give reasons for your opinion.

17. Deceiving subjects. Students sign up to be subjects in a psychology experiment. When they arrive, they are placed in a room and assigned a task. During the task, the subject hears a loud thud from an adjacent room and then a piercing cry for help. Some subjects are placed in a room by themselves. Others are placed in a room with “confederates” who have been instructed by the researcher to look up upon hearing the cry, then return to their task. The treatments being compared are whether the subject is alone in the room or in

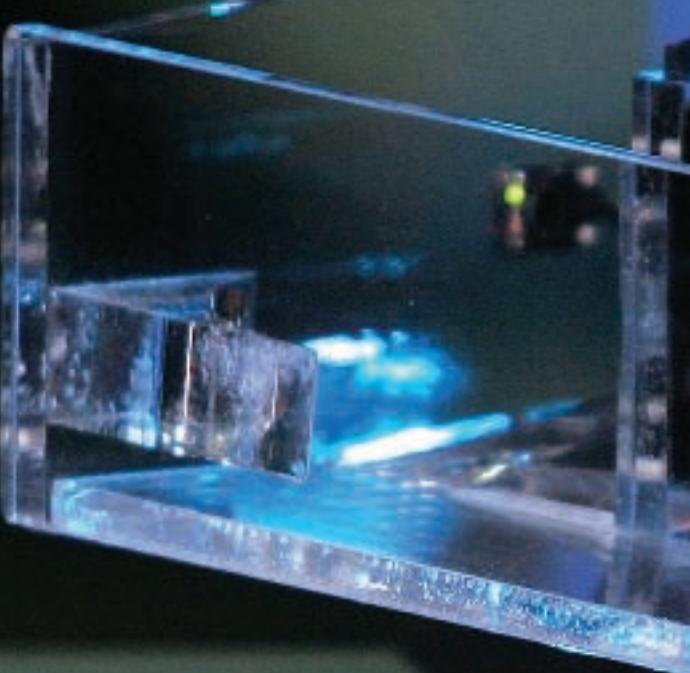
the room with confederates. Will the subject ignore the cry for help?

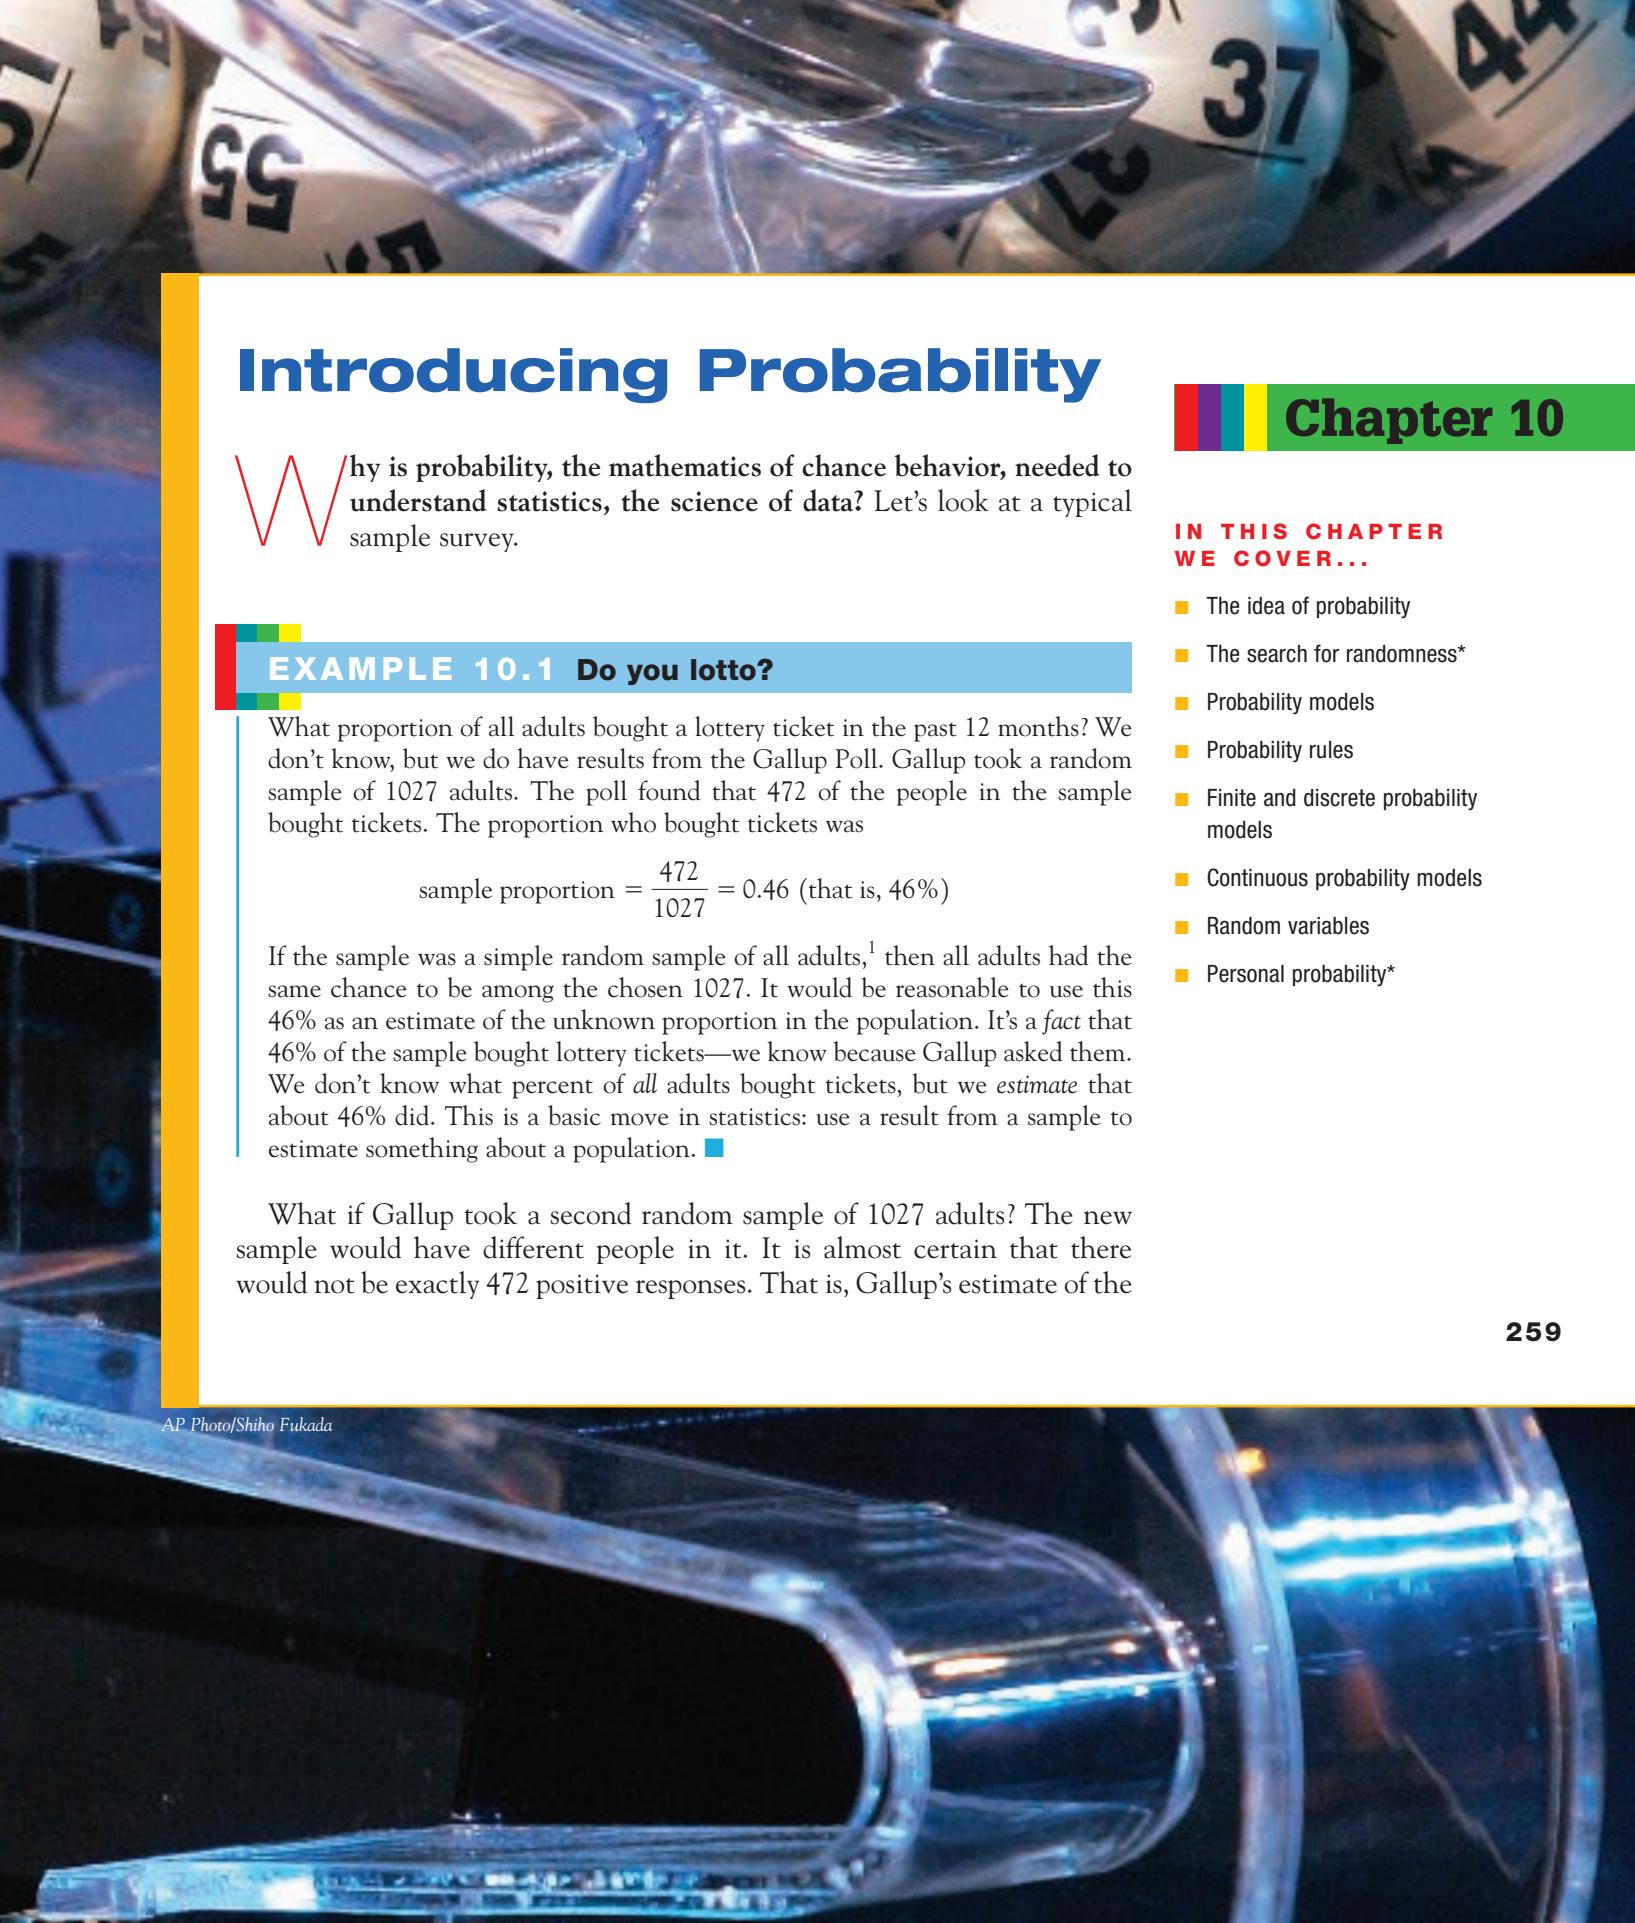
The students had agreed to take part in an unspecified study, and the true nature of the experiment is explained to them afterward. Do you think this study is ethically okay?

18. Deceiving subjects. A psychologist conducts the following experiment: he measures the attitude of subjects toward cheating, then has them take a mathematics skills exam in which the subjects are tempted to cheat. Subjects are told that high scores will receive a \$100.00 gift certificate and that the purpose of the experiment is to see if rewards affect performance. The exam is computer-based and multiple choice. Subjects are left alone in a room with a computer on which the exam is available and are told that they are to click on the answer they believe is correct. However, when subjects click on an answer, a small pop-up window appears with the correct answer indicated. When the pop-up window is closed, it is possible to change the answer selected. The computer records—unknown to the subjects—whether or not they change their answers after closing the pop-up window. After completing the exam, attitude toward cheating is retested.

Subjects who cheat tend to change their attitudes to find cheating more acceptable. Those who resist the temptation to cheat tend to condemn cheating more strongly on the second test of attitude. These results confirm the psychologist’s theory.

This experiment tempts subjects to cheat. The subjects are led to believe that they can cheat secretly when in fact they are observed. Is this experiment ethically objectionable? Explain your position.



A close-up photograph of several lottery balls in a machine. The balls are white with black numbers printed on them. One ball clearly visible has the number 37. Another ball partially visible behind it has the number 44. The lighting is dramatic, with bright highlights reflecting off the metallic surfaces of the balls and the machine's interior.

Chapter 10

Introducing Probability

Why is probability, the mathematics of chance behavior, needed to understand statistics, the science of data? Let's look at a typical sample survey.

EXAMPLE 10.1 Do you lotto?

What proportion of all adults bought a lottery ticket in the past 12 months? We don't know, but we do have results from the Gallup Poll. Gallup took a random sample of 1027 adults. The poll found that 472 of the people in the sample bought tickets. The proportion who bought tickets was

$$\text{sample proportion} = \frac{472}{1027} = 0.46 \text{ (that is, } 46\%)$$

If the sample was a simple random sample of all adults,¹ then all adults had the same chance to be among the chosen 1027. It would be reasonable to use this 46% as an estimate of the unknown proportion in the population. It's a fact that 46% of the sample bought lottery tickets—we know because Gallup asked them. We don't know what percent of *all* adults bought tickets, but we *estimate* that about 46% did. This is a basic move in statistics: use a result from a sample to estimate something about a population. ■

What if Gallup took a second random sample of 1027 adults? The new sample would have different people in it. It is almost certain that there would not be exactly 472 positive responses. That is, Gallup's estimate of the

IN THIS CHAPTER WE COVER...

- The idea of probability
- The search for randomness*
- Probability models
- Probability rules
- Finite and discrete probability models
- Continuous probability models
- Random variables
- Personal probability*

AP Photo/Shiro Fukada



proportion of adults who bought a lottery ticket will vary from sample to sample. Could it happen that one random sample finds that 46% of adults recently bought a lottery ticket and a second random sample finds that 66% had done so? *Random samples eliminate bias from the act of choosing a sample, but they can still be wrong because of the variability that results when we choose at random.* If the variation when we take repeat samples from the same population is too great, we can't trust the results of any one sample.

This is where we need facts about probability to make progress in statistics. Because Gallup uses chance to choose its samples, the laws of probability govern the behavior of the samples. Gallup says that the probability is 0.95 that an estimate from one of their samples comes within ± 3 percentage points of the truth about the population of all adults. The first step toward understanding this statement is to understand what “probability 0.95” means. Our purpose in this chapter is to understand the language of probability, but without going into the mathematics of probability theory.

THE IDEA OF PROBABILITY

To understand why we can trust random samples and randomized comparative experiments, we must look closely at chance behavior. The big fact that emerges is this: **chance behavior is unpredictable in the short run but has a regular and predictable pattern in the long run.**

Toss a coin, or choose a random sample. The result can't be predicted in advance, because the result will vary when you toss the coin or choose the sample repeatedly. But there is still a regular pattern in the results, a pattern that emerges clearly only after many repetitions. This remarkable fact is the basis for the idea of probability.

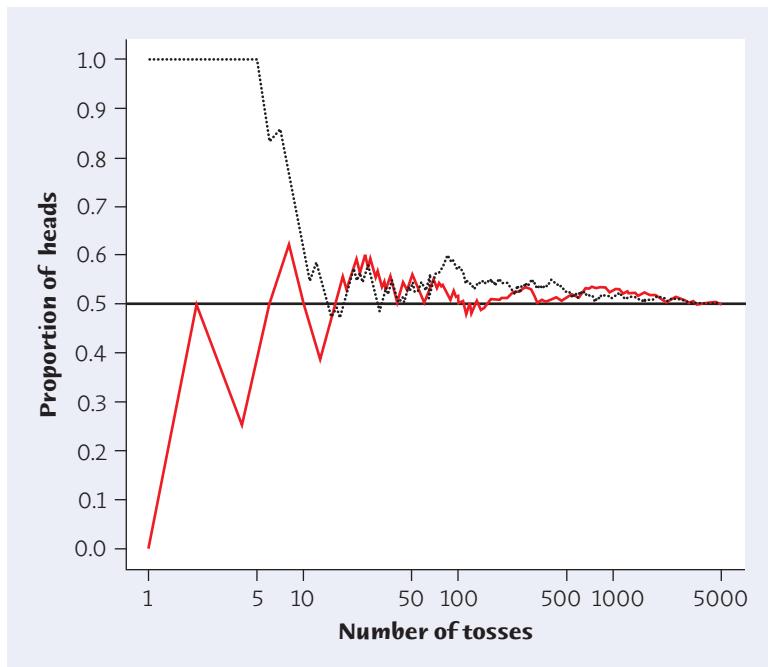
EXAMPLE 10.2 Coin tossing



SuperStock

When you toss a coin, there are only two possible outcomes, heads or tails. Figure 10.1 shows the results of tossing a coin 5000 times twice. For each number of tosses from 1 to 5000, we have plotted the proportion of those tosses that gave a head. Trial A (solid red line) begins tail, head, tail, tail. You can see that the proportion of heads for Trial A starts at 0 on the first toss, rises to 0.5 when the second toss gives a head, then falls to 0.33 and 0.25 as we get two more tails. Trial B (dashed gray line), on the other hand, starts with five straight heads, so the proportion of heads is 1 until the sixth toss.

The proportion of tosses that produce heads is quite variable at first. Trial A starts low and Trial B starts high. As we make more and more tosses, however, the proportion of heads for both trials gets close to 0.5 and stays there. If we made yet a third trial at tossing the coin a great many times, the proportion of heads would again settle down to 0.5 in the long run. This is the intuitive idea of probability. Probability 0.5 means “occurs half the time in a very large number of trials.” The probability 0.5 appears as a horizontal line on the graph. ■

**FIGURE 10.1**

The proportion of tosses of a coin that give a head changes as we make more tosses. Eventually, however, the proportion approaches 0.5, the probability of a head. This figure shows the results of two trials of 5000 tosses each.

We might suspect that a coin has probability 0.5 of coming up heads just because the coin has two sides. But we can't be sure. In fact, spinning a penny on a flat surface, rather than tossing the coin, gives heads probability about 0.45 rather than 0.5.² The idea of probability is empirical. That is, it is based on observation rather than theorizing. Probability describes what happens in very many trials, and we must actually observe many trials to pin down a probability. In the case of tossing a coin, some diligent people have in fact made thousands of tosses.

EXAMPLE 10.3 Some coin tossers

The French naturalist Count Buffon (1707–1788) tossed a coin 4040 times. Result: 2048 heads, or proportion $2048/4040 = 0.5069$ for heads.

Around 1900, the English statistician Karl Pearson heroically tossed a coin 24,000 times. Result: 12,012 heads, a proportion of 0.5005.

While imprisoned by the Germans during World War II, the South African mathematician John Kerrich tossed a coin 10,000 times. Result: 5067 heads, a proportion of 0.5067. ■

RANDOMNESS AND PROBABILITY

We call a phenomenon **random** if individual outcomes are uncertain but there is nonetheless a regular distribution of outcomes in a large number of repetitions.

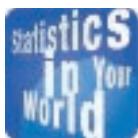
The **probability** of any outcome of a random phenomenon is the proportion of times the outcome would occur in a very long series of repetitions.



The best way to understand randomness is to observe random behavior, as in Figure 10.1. You can do this with physical devices like coins, but computer simulations (imitations) of random behavior allow faster exploration. The *Probability* applet is a computer simulation that animates Figure 10.1. It allows you to choose the probability of a head and simulate any number of tosses of a coin with that probability. Experience shows that the proportion of heads gradually settles down close to the probability. Equally important, it also shows that *the proportion in a small or moderate number of tosses can be far from the probability.*

Probability describes only what happens in the long run. Of course, we can never observe a probability exactly. We could always continue tossing the coin, for example. Mathematical probability is an idealization based on imagining what would happen in an indefinitely long series of trials.

THE SEARCH FOR RANDOMNESS*



Does God play dice?

Few things in the world are truly random in the sense that no amount of information will allow us to predict the outcome. But according to the branch of physics called quantum mechanics, randomness does rule events inside individual atoms. Although Albert Einstein helped quantum theory get started, he always insisted that nature must have some fixed reality, not just probabilities. “I shall never believe that God plays dice with the world,” said the great scientist. A century after Einstein’s first work on quantum theory, it appears that he was wrong.

Random numbers are valuable. They are used to choose random samples, to shuffle the cards in online poker games, to encrypt our credit card numbers when we buy online, and as part of simulations of the flow of traffic and the spread of epidemics. Where does randomness come from, and how can we get random numbers? We defined randomness by how it behaves: unpredictable in the short run, regular pattern in the long run. Probability describes the long-run regular pattern. That many things are random in this sense is an observed fact about the world. Not all these things are “really” random. Here’s a quick tour of how to find random behavior and get random numbers.

The easiest way to get random numbers is from a *computer program*. Of course, a computer program just does what it is told to do. Run the program again and you get exactly the same result. The random numbers in Table B, the outcomes of the *Probability* applet, and the random numbers that shuffle cards for online poker come from computer programs, so they aren’t “really” random. Clever computer programs produce outcomes that look random even though they really aren’t. These *pseudo-random numbers* are more than good enough for choosing samples and shuffling cards. But they may have hidden patterns that can distort scientific simulations.

You might think that *physical devices such as coins and dice* produce really random outcomes. But a tossed coin obeys the laws of physics. If we knew all the inputs of the toss (forces, angles, and so on), then we could say in advance whether the outcome will be heads or tails. The outcome of a toss is predictable rather than random. Why do the results of tossing a coin *look* random? The outcomes are extremely sensitive to the inputs, so that very small changes in the forces you apply when you toss a coin change the outcome from heads to tails and back again. In practice, the outcomes are not predictable. Probability is a lot more useful than physics for describing coin tosses.

We call a phenomenon with “small changes in, big changes out” behavior *chaotic*. If we can feed chaotic behavior into a computer, we can do better than pseudo-random numbers. Coins and dice are awkward, but you can go to the

*This short discussion is optional.

Web site www.random.org to get random numbers from radio noise in the atmosphere, a chaotic phenomenon that is easy to feed to a computer.

Is anything really random? As far as current science can say, behavior inside atoms really is random—that is, there isn’t any way to predict behavior in advance no matter how much information we have. It was this “really, truly random” idea that Einstein disliked as he watched the new science of quantum mechanics emerge. You can go to the HotBits Web site www.fourmilab.ch/hotbits to get really, truly random numbers generated from the radioactive decay of atoms.

APPLY YOUR KNOWLEDGE

- 10.1 Texas hold ’em.** In the popular Texas hold ’em variety of poker, players make their best five-card poker hand by combining the two cards they are dealt with three of five cards available to all players. You read in a book on poker that if you hold a pair (two cards of the same rank) in your hand, the probability of getting four of a kind is $2/245$. Explain carefully what this means. In particular, explain why it does *not* mean that if you play 245 such hands, exactly 2 will be four of a kind.
- 10.2 Probability says ...** Probability is a measure of how likely an event is to occur. Match one of the probabilities that follow with each statement of likelihood given. (The probability is usually a more exact measure of likelihood than is the verbal statement.)

0 0.01 0.45 0.50 0.55 0.99 1

- (a) This event is impossible. It can never occur.
- (b) This event is certain. It will occur on every trial.
- (c) This event is very likely, but it will not occur once in a while in a long sequence of trials.
- (d) This event will occur slightly less often than not.

- 10.3 Random digits.** The table of random digits (Table B) was produced by a random mechanism that gives each digit probability 0.1 of being a 0.
- (a) What proportion of the first 200 digits (those in the first five lines) in the table are 0s? This proportion is an estimate, based on 200 repetitions, of the true probability, which we know is 0.1.
 - (b) The *Probability* applet can imitate random digits. Set the probability of heads in the applet to 0.1. Check “Show true probability” to show this value on the graph. A head stands for a 0 in the random digit table and a tail stands for any other digit. Simulate 200 digits (keep clicking “Toss” to get 40 at a time—don’t click “Reset”). If you kept going forever, presumably you would get 10% heads. What was the result of your 200 tosses?

- 10.4 The long run but not the short run.** Our intuition about chance behavior is not very accurate. In particular, we tend to expect that the long-run pattern described by probability will show up in the short run as well. For example, we tend to think that tossing a coin 20 times will give close to 10 heads.
- (a) Set the probability of heads in the *Probability* applet to 0.5 and the number of tosses to 20. Click “Toss” to simulate 20 tosses of a balanced coin. What was the proportion of heads?



Cut and Deal Ltd./Alamy



- (b) Click “Reset” and toss again. The simulation is fast, so do it 25 times and keep a record of the proportion of heads in each set of 20 tosses. Make a stemplot of your results. You see that the result of tossing a coin 20 times is quite variable and need not be very close to the probability 0.5 of heads.

PROBABILITY MODELS

Gamblers have known for centuries that the fall of coins, cards, and dice displays clear patterns in the long run. The idea of probability rests on the observed fact that the average result of many thousands of chance outcomes can be known with near certainty. How can we give a mathematical description of long-run regularity?

To see how to proceed, think first about a very simple random phenomenon, tossing a coin once. When we toss a coin, we cannot know the outcome in advance. What *do* we know? We are willing to say that the outcome will be either heads or tails. We believe that each of these outcomes has probability 1/2. This description of coin tossing has two parts:

- a list of possible outcomes
- a probability for each outcome

Such a description is the basis for all *probability models*. Here is the basic vocabulary we use.

PROBABILITY MODELS

The **sample space** S of a random phenomenon is the set of all possible outcomes.

An **event** is an outcome or a set of outcomes of a random phenomenon. That is, an event is a subset of the sample space.

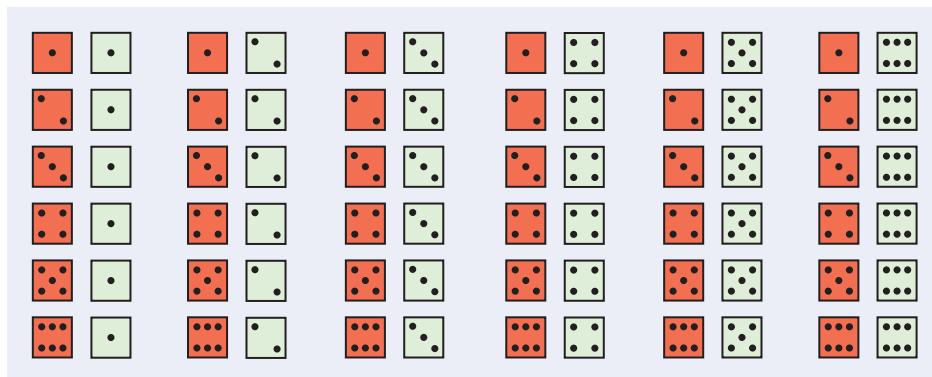
A **probability model** is a mathematical description of a random phenomenon consisting of two parts: a sample space S and a way of assigning probabilities to events.

A sample space S can be very simple or very complex. When we toss a coin once, there are only two outcomes, heads and tails. The sample space is $S = \{H, T\}$. When Gallup draws a random sample of 1523 adults, the sample space contains all possible choices of 1523 of the 235 million adults in the United States. This S is extremely large. Each member of S is a possible sample, so S is the collection or “space” of all possible samples. This explains the term *sample space*.

EXAMPLE 10.4 Rolling dice

Rolling two dice is a common way to lose money in casinos. There are 36 possible outcomes when we roll two dice and record the up-faces in order (first die, second die). Figure 10.2 displays these outcomes. They make up the sample space S . “Roll a 5” is an event, call it A , that contains four of these 36 outcomes:

$$A = \left\{ \begin{array}{c} \bullet \\ \square \end{array}, \begin{array}{c} \bullet \bullet \\ \square \square \end{array} \right\}$$

**FIGURE 10.2**

The 36 possible outcomes in rolling two dice. If the dice are carefully made, all of these outcomes have the same probability.

How can we assign probabilities to this sample space? We can find the actual probabilities for two specific dice only by actually tossing the dice many times, and even then only approximately. So we will give a probability model that assumes ideal, perfectly balanced dice. This model will be quite accurate for carefully made casino dice and less accurate for the cheap dice that come with a board game.

If the dice are perfectly balanced, all 36 outcomes in Figure 10.2 will be *equally likely*. That is, each of the 36 outcomes will come up on one thirty-sixth of all rolls in the long run. So each outcome has probability $1/36$. There are 4 outcomes in the event A (“roll a 5”), so this event has probability $4/36$. In this way we can assign a probability to any event. So we have a complete probability model. ■

EXAMPLE 10.5 Rolling dice and counting the spots

Gamblers care only about the total number of spots on the up-faces of the dice. The sample space for rolling two dice and counting the spots is

$$S = \{2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12\}$$

Comparing this S with Figure 10.2 reminds us that *we can change S by changing the detailed description of the random phenomenon we are describing.*



What are the probabilities for this new sample space? The 11 possible outcomes are not equally likely, because there are six ways to roll a 7 and only one way to roll a 2 or a 12. That’s the key: each outcome in Figure 10.2 has probability $1/36$. So “roll a 7” has probability $6/36$ because this event contains 6 of the 36 outcomes. Similarly, “roll a 2” has probability $1/36$, and “roll a 5” (4 outcomes from Figure 10.2) has probability $4/36$. Here is the complete probability model:

Spots	2	3	4	5	6	7	8	9	10	11	12
Probability	$1/36$	$2/36$	$3/36$	$4/36$	$5/36$	$6/36$	$5/36$	$4/36$	$3/36$	$2/36$	$1/36$

APPLY YOUR KNOWLEDGE

10.5 Sample space. Choose a student at random from a large statistics class. Describe a sample space S for each of the following. (In some cases you may have some freedom in specifying S .)

- Does the student live on campus or off campus?
- What is the student’s age in years?

- (c) Ask how much money in coins (not bills) the student is carrying.
- (d) Record the student's letter grade at the end of the course.



Slpix/Dreamstime.com

10.6 Role-playing games. Computer games in which the players take the roles of characters are very popular. They go back to earlier tabletop games such as Dungeons & Dragons. These games use many different types of dice. A four-sided die has faces with 1, 2, 3, and 4 spots.

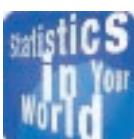
- (a) What is the sample space for rolling a four-sided die twice (spots on first and second rolls)? Follow the example of Figure 10.2.
- (b) What is the assignment of probabilities to outcomes in this sample space? Assume that the die is perfectly balanced, and follow the method of Example 10.4.

10.7 Role-playing games. The intelligence of a character in a game is determined by rolling the four-sided die twice and adding 1 to the sum of the spots. Start with your work in the previous exercise to give a probability model (sample space and probabilities of outcomes) for the character's intelligence. Follow the method of Example 10.5.

PROBABILITY RULES

In Examples 10.4 and 10.5 we found probabilities for tossing dice. As random phenomena go, dice are pretty simple. Even so, we had to assume idealized perfectly balanced dice. In most situations, it isn't easy to give a "correct" probability model. We can make progress by listing some facts that must be true for *any* assignment of probabilities. These facts follow from the idea of probability as "the long-run proportion of repetitions on which an event occurs."

- 1. Any probability is a number between 0 and 1.** Any proportion is a number between 0 and 1, so any probability is also a number between 0 and 1. An event with probability 0 never occurs, and an event with probability 1 occurs on every trial. An event with probability 0.5 occurs in half the trials in the long run.
- 2. All possible outcomes together must have probability 1.** Because some outcome must occur on every trial, the sum of the probabilities for all possible outcomes must be exactly 1.
- 3. If two events have no outcomes in common, the probability that one or the other occurs is the sum of their individual probabilities.** If one event occurs in 40% of all trials, a different event occurs in 25% of all trials, and the two can never occur together, then one or the other occurs on 65% of all trials because $40\% + 25\% = 65\%$.
- 4. The probability that an event does not occur is 1 minus the probability that the event does occur.** If an event occurs in (say) 70% of all trials, it fails to occur in the other 30%. The probability that an event occurs and the probability that it does not occur always add to 100%, or 1.



Equally likely?

A game of bridge begins by dealing all 52 cards in the deck to the four players, 13 to each. If the deck is well shuffled, all of the immense number of possible hands will be equally likely. But don't expect the hands that appear in newspaper bridge columns to reflect the equally likely probability model. Writers on bridge choose "interesting" hands, especially those that lead to high bids that are rare in actual play.

We can use mathematical notation to state Facts 1 to 4 more concisely. Capital letters near the beginning of the alphabet denote events. If A is any event, we write its probability as $P(A)$. Here are our probability facts in formal language. As

you apply these rules, remember that they are just another form of intuitively true facts about long-run proportions.

PROBABILITY RULES

Rule 1. The probability $P(A)$ of any event A satisfies $0 \leq P(A) \leq 1$.

Rule 2. If S is the sample space in a probability model, then $P(S) = 1$.

Rule 3. Two events A and B are **disjoint** if they have no outcomes in common and so can never occur together. If A and B are disjoint,

$$P(A \text{ or } B) = P(A) + P(B)$$

This is the **addition rule for disjoint events**.

Rule 4. For any event A ,

$$P(A \text{ does not occur}) = 1 - P(A)$$

The addition rule extends to more than two events that are disjoint in the sense that no two have any outcomes in common. If events A , B , and C are disjoint, the probability that one of these events occurs is $P(A) + P(B) + P(C)$.

EXAMPLE 10.6 Using the probability rules

We already used the addition rule, without calling it by that name, to find the probabilities in Example 10.5. The event “roll a 5” contains the four disjoint outcomes displayed in Example 10.4, so the addition rule (Rule 3) says that its probability is

$$\begin{aligned} P(\text{roll a 5}) &= P\left(\begin{array}{|c|c|}\hline \bullet & \cdot \\ \hline \end{array}\right) + P\left(\begin{array}{|c|c|}\hline \bullet & \cdot \\ \hline \cdot & \bullet \\ \hline \end{array}\right) + P\left(\begin{array}{|c|c|}\hline \bullet & \cdot \\ \hline \cdot & \bullet \\ \hline \end{array}\right) + P\left(\begin{array}{|c|c|}\hline \bullet & \bullet \\ \hline \cdot & \cdot \\ \hline \end{array}\right) \\ &= \frac{1}{36} + \frac{1}{36} + \frac{1}{36} + \frac{1}{36} \\ &= \frac{4}{36} = 0.111 \end{aligned}$$

Check that the probabilities in Example 10.5, found using the addition rule, are all between 0 and 1 and add to exactly 1. That is, this probability model obeys Rules 1 and 2.

What is the probability of rolling anything other than a 5? By Rule 4,

$$\begin{aligned} P(\text{roll does not give a 5}) &= 1 - P(\text{roll a 5}) \\ &= 1 - 0.111 = 0.889 \end{aligned}$$

Our model assigns probabilities to individual outcomes. To find the probability of an event, just add the probabilities of the outcomes that make up the event. For example:

$$\begin{aligned} P(\text{outcome is odd}) &= P(3) + P(5) + P(7) + P(9) + P(11) \\ &= \frac{2}{36} + \frac{4}{36} + \frac{6}{36} + \frac{4}{36} + \frac{2}{36} \\ &= \frac{18}{36} = \frac{1}{2} \blacksquare \end{aligned}$$



Image Source/Alamy



APPLY YOUR KNOWLEDGE

10.8 Who takes the GMAT? In many settings, the “rules of probability” are just basic facts about percents. The Graduate Management Admission Test (GMAT) Web site provides the following information about the undergraduate majors of those who took the test in 2009–2010: 53% majored in business or commerce; 17% majored in engineering; 16% majored in the social sciences; 6% majored in the sciences; 5% majored in the humanities; and 3% listed some major other than the preceding.³

- What percent of those who took the test in 2009–2010 majored in either engineering or science? Which rule of probability did you use to find the answer?
- What percent of those who took the test in 2009–2010 majored in something other than business or commerce? Which rule of probability did you use to find the answer?

10.9 Overweight? Although the rules of probability are just basic facts about percents or proportions, we need to be able to use the language of events and their probabilities. Choose an American adult at random. Define two events:

A = the person chosen is obese

B = the person chosen is overweight, but not obese

According to the National Center for Health Statistics, $P(A) = 0.34$ and $P(B) = 0.33$.

- Explain why events A and B are disjoint.
- Say in plain language what the event “ A or B ” is. What is $P(A \text{ or } B)$?
- If C is the event that the person chosen has normal weight or less, what is $P(C)$?

10.10 Languages in Canada. Canada has two official languages, English and French. Choose a Canadian at random and ask, “What is your mother tongue?” Here is the distribution of responses, combining many separate languages from the province of Quebec:⁴

Language	English	French	Italian	Other
Probability	0.08	0.80	0.02	?

- What probability should replace “?” in the distribution?
- What is the probability that a Canadian’s mother tongue is not English?
- What is the probability that a Canadian’s mother tongue is a language other than English or French?

FINITE AND DISCRETE PROBABILITY MODELS

Examples 10.4, 10.5, and 10.6 illustrate one way to assign probabilities to events: assign a probability to every individual outcome, then add these probabilities to find the probability of any event. This idea works well when there are only a finite (fixed and limited) number of outcomes.

FINITE PROBABILITY MODEL

A probability model with a finite sample space is called **finite**.

To assign probabilities in a finite model, list the probabilities of all the individual outcomes. These probabilities must be numbers between 0 and 1 that add to exactly 1. The probability of any event is the sum of the probabilities of the outcomes making up the event.

Finite probability models are sometimes called **discrete** probability models. However, discrete probability models include finite sample spaces as well as sample spaces that are infinite and equivalent to the set of all positive integers. An example of a discrete but not finite sample space would be the sample space for the number of free-throw attempts until a basketball player makes her first free throw. This could occur on her first attempt, her second attempt, her third attempt, etc. Assigning probabilities to individual outcomes in an infinite discrete sample space is more complicated than for a finite sample space. In this book we will often refer to finite probability models as discrete, and in practice statisticians often refer to finite probability models as discrete.

EXAMPLE 10.7 Benford's law

Faked numbers in tax returns, invoices, or expense account claims often display patterns that aren't present in legitimate records. Some patterns, such as too many round numbers, are obvious and easily avoided by a clever crook. Others are more subtle. It is a striking fact that the first digits of numbers in legitimate records often follow a model known as Benford's law.⁵ Call the first digit of a randomly chosen record X for short. Benford's law gives this probability model for X (note that a first digit can't be 0):

First digit X	1	2	3	4	5	6	7	8	9
Probability	0.301	0.176	0.125	0.097	0.079	0.067	0.058	0.051	0.046

Check that the probabilities of the outcomes sum to exactly 1. This is therefore a legitimate finite (or discrete) probability model. Investigators can detect fraud by comparing the first digits in records such as invoices paid by a business with these probabilities.

The probability that a first digit is equal to or greater than 6 is

$$\begin{aligned} P(X \geq 6) &= P(X = 6) + P(X = 7) + P(X = 8) + P(X = 9) \\ &= 0.067 + 0.058 + 0.051 + 0.046 = 0.222 \end{aligned}$$

This is less than the probability that a record has first digit 1,

$$P(X = 1) = 0.301$$

Fraudulent records tend to have too few 1s and too many higher first digits.

Note that the probability that a first digit is greater than or equal to 6 is not the same as the probability that a first digit is strictly greater than 6.



The latter probability is

$$P(X > 6) = 0.058 + 0.051 + 0.046 = 0.155$$

The outcome $X = 6$ is included in “greater than or equal to” and is not included in “strictly greater than.” ■

APPLY YOUR KNOWLEDGE

10.11 Rolling a die. Figure 10.3 displays several finite probability models for rolling a die. We can learn which model is actually *accurate* for a particular die only by rolling the die many times. However, some of the models are not *legitimate*. That is, they do not obey the rules. Which are legitimate and which are not? In the case of the illegitimate models, explain what is wrong.

FIGURE 10.3

Four assignments of probabilities to the six faces of a die, for Exercise 10.11.

Outcome	Probability			
	Model 1	Model 2	Model 3	Model 4
	1/7	1/3	1/3	1
	1/7	1/6	1/6	1
	1/7	1/6	1/6	2
	1/7	0	1/6	1
	1/7	1/6	1/6	1
	1/7	1/6	1/6	2

10.12 Benford’s law. The first digit of a randomly chosen expense account claim follows Benford’s law (Example 10.7). Consider the events

$$A = \{\text{first digit is 4 or greater}\}$$

$$P = \{\text{first digit is even}\}$$

- What outcomes make up the event A ? What is $P(A)$?
- What outcomes make up the event B ? What is $P(B)$?
- What outcomes make up the event “ A or B ”? What is $P(A \text{ or } B)$? Why is this probability not equal to $P(A) + P(B)$?

10.13 Weighty behavior. Choose an adult in the United States at random and ask, “How many days per week do you lift weights?” Call the response X for short. Based on a large sample survey, here is a probability model for the answer you will get:⁶

Days	0	1	2	3	4	5	6	7
Probability	0.73	0.06	0.06	0.06	0.04	0.02	0.01	0.02

- (a) Verify that this is a legitimate finite probability model.
- (b) Describe the event $X < 4$ in words. What is $P(X < 4)$?
- (c) Express the event “lifted weights at least once” in terms of X . What is the probability of this event?

CONTINUOUS PROBABILITY MODELS

When we use the table of random digits to select a digit between 0 and 9, the finite probability model assigns probability 1/10 to each of the 10 possible outcomes. Suppose that we want to choose a number at random between 0 and 1, allowing *any* number between 0 and 1 as the outcome. Software random number generators will do this. For example, here is the result of asking software to produce 5 random numbers between 0 and 1:

0.2893511 0.3213787 0.5816462 0.9787920 0.4475373

The sample space is now an entire interval of numbers:

$$S = \{\text{all numbers between 0 and 1}\}$$

Call the outcome of the random number generator Y for short. How can we assign probabilities to such events as $\{0.3 \leq Y \leq 0.7\}$? As in the case of selecting a random digit, we would like all possible outcomes to be equally likely. But we cannot assign probabilities to each individual value of Y and then add them, because there is an infinite interval of possible values. In fact, we cannot even make a list of the individual values of Y . For example, what is the next largest value of Y after 0?

We use a new way of assigning probabilities directly to events—as *areas under a density curve*. Any density curve has area exactly 1 underneath it, corresponding to total probability 1. We met density curves as models for data in Chapter 3 (page 71).

CONTINUOUS PROBABILITY MODEL

A continuous probability model assigns probabilities as areas under a density curve. The area under the curve and above any range of values is the probability of an outcome in that range.



Really random digits

For purists,

the RAND

Corporation long ago published a book titled *One Million Random Digits*. The book lists 1,000,000 digits that were produced by a very elaborate physical randomization and really are random. An employee of RAND once said that this is not the most boring book that RAND has ever published.

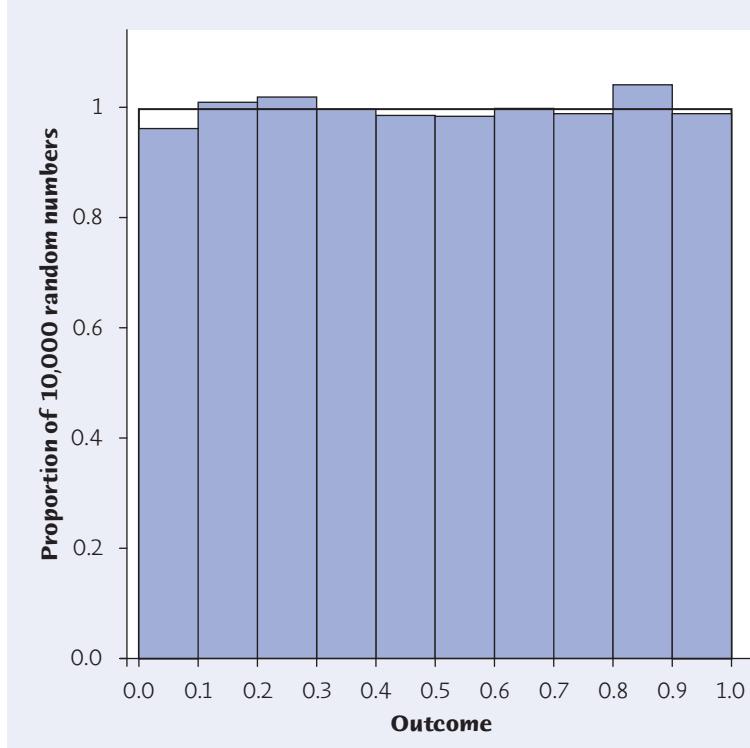
EXAMPLE 10.8 Random numbers

The random number generator will spread its output uniformly across the entire interval from 0 to 1 as we allow it to generate a long sequence of numbers. Figure 10.4 is a histogram of 10,000 random numbers. They are quite uniform, but not exactly so. The bar heights would all be exactly equal (1000 numbers for each bar) if the 10,000 numbers were exactly uniform. In fact, the counts vary from a low of 960 to a high of 1022.

As in Chapter 3, we have adjusted the histogram scale so that the total area of the bars is exactly 1. Now we can add the density curve that describes the distribution of

FIGURE 10.4

The probability model for the outcomes of a software random number generator, for Example 10.8. Compare the histogram of 10,000 actual outcomes with the uniform density curve that spreads probability evenly between 0 and 1.



uniform distribution

perfectly random numbers. This density curve also appears in Figure 10.4. It has height 1 over the interval from 0 to 1. This is the density curve of a **uniform distribution**. It is the continuous probability model for the results of generating very many random numbers. Like the probability models for perfectly balanced coins and dice, the density curve is an idealized description of the outcomes of a perfectly uniform random number generator. It is a good approximation for software outcomes, but even 10,000 tries isn't enough for actual outcomes to look exactly like the idealized model. ■

The uniform density curve has height 1 over the interval from 0 to 1. The area under the curve is 1, and the probability of any event is the area under the curve and above the event in question. Figure 10.5 illustrates finding probabilities as areas under the density curve. The probability that the random number generator produces a number between 0.3 and 0.7 is

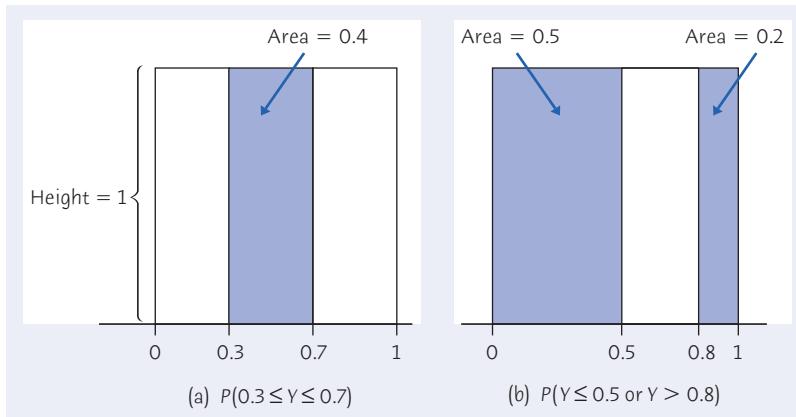
$$P(0.3 \leq Y \leq 0.7) = 0.4$$

because the area under the density curve and above the interval from 0.3 to 0.7 is 0.4. The height of the curve is 1 and the area of a rectangle is the product of height and length, so the probability of any interval of outcomes is just the length of the interval. Similarly,

$$P(Y \leq 0.5) = 0.5$$

$$P(Y > 0.8) = 0.2$$

$$P(Y \leq 0.5 \text{ or } Y > 0.8) = 0.7$$

**FIGURE 10.5**

Probability as area under a density curve. The uniform density curve spreads probability evenly between 0 and 1.

The last event consists of two nonoverlapping intervals, so the total area above the event is found by adding two areas, as illustrated by Figure 10.5(b). This assignment of probabilities obeys all of our rules for probability.

Continuous probability models assign probabilities to intervals of outcomes rather than to individual outcomes. In fact, *all continuous probability models assign probability 0 to every individual outcome*. Only intervals of values have positive probability. To see that this is true, consider a specific outcome such as $P(Y = 0.8)$. The probability of any interval is the same as its length. The point 0.8 has no length, so its probability is 0. Put another way, $P(Y > 0.8)$ and $P(Y \geq 0.8)$ are both 0.2 because that is the area in Figure 10.5(b) between 0.8 and 1.

We can use any density curve to assign probabilities. The density curves that are most familiar to us are the Normal curves. **Normal distributions are continuous probability models** as well as descriptions of data. There is a close connection between a Normal distribution as an idealized description for data and a Normal probability model. If we look at the heights of all young women, we find that they closely follow the Normal distribution with mean $\mu = 64.3$ inches and standard deviation $\sigma = 2.7$ inches. This is a distribution for a large set of data. Now choose one young woman at random. Call her height X . If we repeat the random choice very many times, the distribution of values of X is the same Normal distribution that describes the heights of all young women.

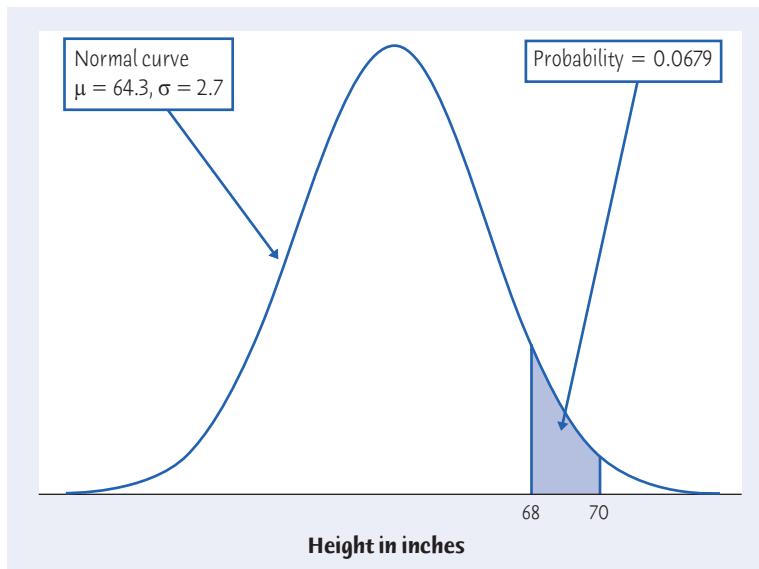
EXAMPLE 10.9 The heights of young women

What is the probability that a randomly chosen young woman has height between 68 and 70 inches? The height X of the woman we choose has the $N(64.3, 2.7)$ distribution. We want $P(68 \leq X \leq 70)$. This is the area under the Normal curve in Figure 10.6. Software or the *Normal Curve* applet will give us the answer at once: $P(68 \leq X \leq 70) = 0.0679$.



FIGURE 10.6

The probability in Example 10.9 as an area under a Normal curve.



We can also find the probability by standardizing and using Table A, the table of standard Normal probabilities. We will reserve capital Z for a standard Normal variable.

$$\begin{aligned} P(68 \leq X \leq 70) &= P\left(\frac{68 - 64.3}{2.7} \leq \frac{X - 64.3}{2.7} \leq \frac{70 - 64.3}{2.7}\right) \\ &= P(1.37 \leq Z \leq 2.11) \\ &= P(Z \leq 2.11) - P(Z \leq 1.37) \\ &= 0.9826 - 0.9147 = 0.0679 \end{aligned}$$

The calculation is the same as those we did in Chapter 3. Only the language of probability is new. ■



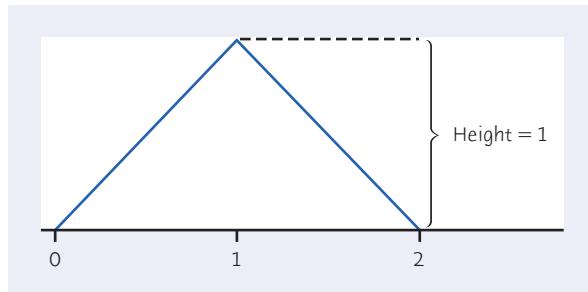
APPLY YOUR KNOWLEDGE

10.14 Random numbers. Let Y be a random number between 0 and 1 produced by the idealized random number generator described in Example 10.8 and Figure 10.4. Find the following probabilities:

- (a) $P(Y \leq 0.6)$
- (b) $P(Y < 0.6)$
- (c) $P(0.4 \leq Y \leq 0.8)$

10.15 Adding random numbers. Generate two random numbers between 0 and 1 and take X to be their sum. The sum X can take any value between 0 and 2. The density curve of X is the triangle shown in Figure 10.7.

- (a) Verify by geometry that the area under this curve is 1.
- (b) What is the probability that X is less than 1? (Sketch the density curve, shade the area that represents the probability, then find that area. Do this for (c) also.)
- (c) What is the probability that X is less than 0.5?

**FIGURE 10.7**

The density curve for the sum of two random numbers, for Exercise 10.15. This density curve spreads probability between 0 and 2.

10.16 The Medical College Admission Test. The Normal distribution with mean $\mu = 25.0$ and standard deviation $\sigma = 6.4$ is a good description of the total score on the Medical College Admission Test (MCAT). This is a continuous probability model for the score of a randomly chosen student. Call the score of a randomly chosen student X for short.

- Write the event “the student chosen has a score of 35 or higher” in terms of X .
- Find the probability of this event.

RANDOM VARIABLES

Examples 10.7 to 10.9 use a shorthand notation that is often convenient. In Example 10.9, we let X stand for the result of choosing a woman at random and measuring her height. We know that X would take a different value if we made another random choice. Because its value changes from one random choice to another, we call the height X a *random variable*.

RANDOM VARIABLE

A **random variable** is a variable whose value is a numerical outcome of a random phenomenon.

The **probability distribution** of a random variable X tells us what values X can take and how to assign probabilities to those values.

We usually denote random variables by capital letters near the end of the alphabet, such as X or Y . Of course, the random variables of greatest interest to us are outcomes such as the mean \bar{x} of a random sample, for which we will keep the familiar notation. There are two main types of random variables, corresponding to two types of probability models: *discrete* and *continuous*. Notice that neither a finite sample space nor a sample space consisting of all positive integers would be considered continuous, because neither sample space is an interval of all possible numbers (to an arbitrary number of decimal places) between two given values. Thus, we classify random variables as either discrete or continuous, rather than as either finite or continuous.

discrete random variable**continuous random variable**

EXAMPLE 10.10 Discrete and continuous random variables

The first digit X in Example 10.7 is a random variable whose possible values are the whole numbers $\{1, 2, 3, 4, 5, 6, 7, 8, 9\}$. The distribution of X assigns a probability to each of these outcomes. Random variables that have a finite list of possible outcomes are called **discrete**.

Compare the output Y of the random number generator in Example 10.8. The values of Y fill the entire interval of numbers between 0 and 1. The probability distribution of Y is given by its density curve, shown in Figure 10.4. Random variables that can take on any value in an interval, with probabilities given as areas under a density curve, are called **continuous**. ■

APPLY YOUR KNOWLEDGE

10.17 Grades in an economics course. Indiana University posts the grade distributions for its courses online.⁷ Students in Economics 201 in the fall 2009 semester received 9% A's, 8% A-'s, 10% B+'s, 14% B's, 13% B-'s, 10% C+'s, 12% C's, 4% C-'s, 4% D+'s, 8% D's, and 8% F's. Choose an Economics 201 student at random. To "choose at random" means to give every student the same chance to be chosen. The student's grade on a four-point scale (with $A = 4$, $A- = 3.7$, $B+ = 3.3$, $B = 3.0$, $B- = 2.7$, $C+ = 2.3$, $C = 2.0$, $C- = 1.7$, $D+ = 1.3$, $D = 1.0$, and $F = 0.0$) is a discrete random variable X with this probability distribution:

Value of X	0.0	1.0	1.3	1.7	2.0	2.3	2.7	3.0	3.3	3.7	4.0
Probability	0.08	0.08	0.04	0.04	0.12	0.10	0.13	0.14	0.10	0.08	0.09

- (a) Say in words what the meaning of $P(X \geq 3.0)$ is. What is this probability?
- (b) Write the event "the student got a grade poorer than B−" in terms of values of the random variable X . What is the probability of this event?

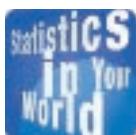
10.18 Running a mile. A study of 12,000 able-bodied male students at the University of Illinois found that their times for the mile run were approximately Normal with mean 7.11 minutes and standard deviation 0.74 minute.⁸ Choose a student at random from this group and call his time for the mile Y .

- (a) Say in words what the meaning of $P(Y \geq 8)$ is. What is this probability?
- (b) Write the event "the student could run a mile in less than 6 minutes" in terms of values of the random variable Y . What is the probability of this event?

PERSONAL PROBABILITY*

We began our discussion of probability with one idea: the probability of an outcome of a random phenomenon is the proportion of times that outcome would occur in a very long series of repetitions. This idea ties probability to actual outcomes. It

*This short section is optional.



What are the odds?

Gamblers often express chance in terms of *odds* rather than probability. Odds of A to B against an outcome means that the probability of that outcome is $B/(A + B)$. So "odds of 5 to 1" is another way of saying "probability 1/6." A probability is always between 0 and 1, but odds range from 0 to infinity. Although odds are mainly used in gambling, they give us a way to make very small probabilities clearer. "Odds of 999 to 1" may be easier to understand than "probability 0.001."

allows us, for example, to estimate probabilities by simulating random phenomena. Yet we often meet another, quite different, idea of probability.

EXAMPLE 10.11 Joe and the Chicago Cubs

Joe sits staring into his beer as his favorite baseball team, the Chicago Cubs, loses another game. The Cubbies have some good young players, so let's ask Joe, "What's the chance that the Cubs will go to the World Series next year?" Joe brightens up. "Oh, about 10%," he says.

Does Joe assign probability 0.10 to the Cubs' appearing in the World Series? The outcome of next year's pennant race is certainly unpredictable, but we can't reasonably ask what would happen in many repetitions. Next year's baseball season will happen only once and will differ from all other seasons in players, weather, and many other ways. If probability measures "what would happen if we did this many times," Joe's 0.10 is not a probability. Probability is based on data about many repetitions of the same random phenomenon. Joe is giving us something else, his personal judgment. ■

Although Joe's 0.10 isn't a probability in our usual sense, it gives useful information about Joe's opinion. More seriously, a company asking, "How likely is it that building this plant will pay off within five years?" can't employ an idea of probability based on many repetitions of the same thing. The opinions of company officers and advisers are nonetheless useful information, and these opinions can be expressed in the language of probability. These are *personal probabilities*.

PERSONAL PROBABILITY

A **personal probability** of an outcome is a number between 0 and 1 that expresses an individual's judgment of how likely the outcome is.

Rachel's opinion about the Cubs may differ from Joe's, and the opinions of several company officers about the new plant may differ. Personal probabilities are indeed personal: they vary from person to person. Moreover, if two people assign different personal probabilities to an event, it may be difficult or impossible to determine who is more correct. If we say, "In the long run, this coin will come up heads 60% of the time," we can find out if we are right by actually tossing the coin several thousand times. If Joe says, "I think the Cubs have a 10% chance of going to the World Series next year," that's just Joe's opinion. Why think of personal probabilities as probabilities? Because *any set of personal probabilities that makes sense obeys the same basic Rules 1 to 4 that describe any legitimate assignment of probabilities to events*. If Joe thinks there's a 10% chance that the Cubs will go to the World Series, he must also think that there's a 90% chance that they won't go. There is just one set of rules of probability, even though we now have two interpretations of what probability means.



APPLY YOUR KNOWLEDGE

10.19 Will you have an accident? The probability that a randomly chosen driver will be involved in an accident in the next year is about 0.2. This is based on the proportion of millions of drivers who have accidents. “Accident” includes things like crumpling a fender in your own driveway, not just highway accidents.

- What do you think is your own probability of being in an accident in the next year? This is a personal probability.
- Give some reasons why your personal probability might be a more accurate prediction of your “true chance” of having an accident than the probability for a random driver.
- Almost everyone says their personal probability is lower than the random driver probability. Why do you think this is true?

10.20 Winning the ACC tournament. The annual Atlantic Coast Conference men’s basketball tournament has temporarily taken Joe’s mind off the Chicago Cubs. He says to himself, “I think that Maryland has probability 0.1 of winning. Duke’s probability is twice Maryland’s, and North Carolina’s probability is three times Maryland’s.”

- What are Joe’s personal probabilities for Duke and North Carolina?
- What is Joe’s personal probability that one of the 9 teams other than Maryland, Duke, and North Carolina will win the tournament?

CHAPTER 10 SUMMARY

CHAPTER SPECIFICS

- A **random phenomenon** has outcomes that we cannot predict but that nonetheless have a regular distribution in very many repetitions.
- The **probability** of an event is the proportion of times the event occurs in many repeated trials of a random phenomenon.
- A **probability model** for a random phenomenon consists of a sample space S and an assignment of probabilities P .
- The **sample space S** is the set of all possible outcomes of the random phenomenon. Sets of outcomes are called **events**. P assigns a number $P(A)$ to an event A as its probability.
- Any assignment of probability must obey the rules that state the basic properties of probability:
 1. $0 \leq P(A) \leq 1$ for any event A .
 2. $P(S) = 1$.
 3. **Addition rule:** Events A and B are **disjoint** if they have no outcomes in common. If A and B are disjoint, then $P(A \text{ or } B) = P(A) + P(B)$.
 4. For any event A , $P(A \text{ does not occur}) = 1 - P(A)$.
- When a sample space S contains finitely many possible values, a **finite probability model** assigns each of these values a probability between 0 and 1 such that the sum of all the probabilities is exactly 1. The probability of any event is the sum of the probabilities of all the values that make up the event. Finite probability models are also referred to as discrete probability models.
- A sample space can contain all values in some interval of numbers. A **continuous probability model** assigns probabilities as areas under a **density curve**. The

probability of any event is the area under the curve above the values that make up the event.

- A **random variable** is a variable taking numerical values determined by the outcome of a random phenomenon. The **probability distribution** of a random variable X tells us what the possible values of X are and how probabilities are assigned to those values.
- A random variable X and its distribution can be **discrete** or **continuous**. The distribution of a **discrete random variable** with finitely many possible values gives the probability of each value. A **continuous random variable** takes all values in some interval of numbers. A density curve describes the probability distribution of a continuous random variable.

LINK IT

This chapter begins our study of probability. The important fact is that random phenomena are unpredictable in the short run but have a regular and predictable behavior in the long run. Probability rules and probability models provide the tools for describing and predicting the long-run behavior of random phenomena.

Probability helps us understand why we can trust random samples and randomized comparative experiments, the subjects of Chapters 8 and 9. It is the key to generalizing what we learn from data produced by random samples and randomized comparative experiments to some wider universe or population. How we use probability to do this will be the topic of the remainder of this book.

CHECK YOUR SKILLS

10.21 You read in a book on poker that the probability of being dealt two pairs in a five-card poker hand is $1/20$. This means that

- if you deal thousands of poker hands, the fraction of them that contain two pairs will be very close to $1/20$.
- if you deal 20 poker hands, exactly 1 of them will contain two pairs.
- if you deal 10,000 poker hands, exactly 500 of them will contain two pairs.

10.22 A basketball player shoots 5 free throws during a game. The sample space for counting the number she makes is

- $S =$ any number between 0 and 1.
- $S =$ whole numbers 0 to 5.
- $S =$ all sequences of 5 hits or misses, like HMMHH.

Here is the probability model for the political affiliation of a randomly chosen adult in the United States.⁹ Exercises 10.23 to 10.26 use this information.

Political affiliation	Republican	Independent	Democrat	Other
Probability	0.28	0.42	0.28	?

10.23 This probability model is

- continuous.
- finite.
- equally likely.

10.24 The probability that a randomly chosen American adult's political affiliation is "Other" must be

- any number between 0 and 1.
- 0.02.
- 0.2.

10.25 What is the probability that a randomly chosen American adult is a member of one of the two major political parties (Republicans and Democrats)?

- 0.42
- 0.44
- 0.56

10.26 What is the probability that a randomly chosen American adult is not a Republican?

- 0.28
- 0.72
- 0.02

10.27 In a table of random digits such as Table B, each digit is equally likely to be any of 0, 1, 2, 3, 4, 5, 6, 7, 8, or 9. What is the probability that a digit in the table is a 7?

- 1/9
- 1/10
- 9/10

10.28 In a table of random digits such as Table B, each digit is equally likely to be any of 0, 1, 2, 3, 4, 5, 6, 7, 8, or 9. What is the probability that a digit in the table is 7 or greater?

- 7/10
- 4/10
- 3/10

10.29 Choose an American household at random and let the random variable X be the number of cars (including SUVs and light trucks) they own. Here is the probability model if we ignore the few households that own more than 6 cars:

Number of cars X	0	1	2	3	4	5	6
Probability	0.09	0.29	0.38	0.16	0.05	0.02	0.01

A housing company builds houses with two-car garages. What percent of households have more cars than the garage can hold?

- (a) 16% (b) 24% (c) 62%

- 10.30** Choose a common fruit fly *Drosophila melanogaster* at random. Call the length of the thorax (where the wings and legs attach) Y. The random variable Y has the Normal distribution with mean $\mu = 0.800$ millimeter (mm) and standard deviation $\sigma = 0.078$ mm. The probability $P(Y > 1)$ that the fly you choose has a thorax more than 1 mm long is about
 (a) 0.995. (b) 0.5. (c) 0.005.

CHAPTER 10 EXERCISES

10.31 Sample space. In each of the following situations, describe a sample space S for the random phenomenon.

- (a) A basketball player shoots four free throws. You record the sequence of hits and misses.
 (b) A basketball player shoots four free throws. You record the number of baskets she makes.



Darrell Walker/HWMS/Icon SMI/Newscom

10.32 Probability models? In each of the following situations, state whether or not the given assignment of probabilities to individual outcomes is legitimate, that is, satisfies the rules of probability. Remember, a legitimate model need not be a practically reasonable model. If the assignment of probabilities is not legitimate, give specific reasons for your answer.

- (a) Roll a six-sided die and record the count of spots on the up-face:

$$\begin{aligned} P(1) &= 0 & P(2) &= 1/6 & P(3) &= 1/3 \\ P(4) &= 1/3 & P(5) &= 1/6 & P(6) &= 0 \end{aligned}$$

- (b) Deal a card from a shuffled deck:

$$\begin{aligned} P(\text{clubs}) &= 12/52 & P(\text{diamonds}) &= 12/52 \\ P(\text{hearts}) &= 12/52 & P(\text{spades}) &= 16/52 \end{aligned}$$

- (c) Choose a college student at random and record sex and enrollment status:

$$\begin{aligned} P(\text{female full-time}) &= 0.56 & P(\text{male full-time}) &= 0.44 \\ P(\text{female part-time}) &= 0.24 & P(\text{male part-time}) &= 0.17 \end{aligned}$$

10.33 Education among young adults. Choose a young adult (aged 25 to 29) at random. The probability is 0.13 that the person chosen did not

complete high school, 0.31 that the person has a high school diploma but no further education, and 0.29 that the person has at least a bachelor's degree.

- (a) What must be the probability that a randomly chosen young adult has some education beyond high school but does not have a bachelor's degree?
 (b) What is the probability that a randomly chosen young adult has at least a high school education?

10.34 Land in Canada. Canada's national statistics agency, Statistics Canada, says that the land area of Canada is 9,094,000 square kilometers. Of this land, 4,176,000 square kilometers are forested. Choose a square kilometer of land in Canada at random.

- (a) What is the probability that the area you choose is forested?
 (b) What is the probability that it is not forested?

10.35 Foreign-language study. Choose a student in a U.S. public high school at random and ask if he or she is studying a language other than English. Here is the distribution of results:

Language	Spanish	French	German	All others	None
Probability	0.30	0.08	0.02	0.03	0.57

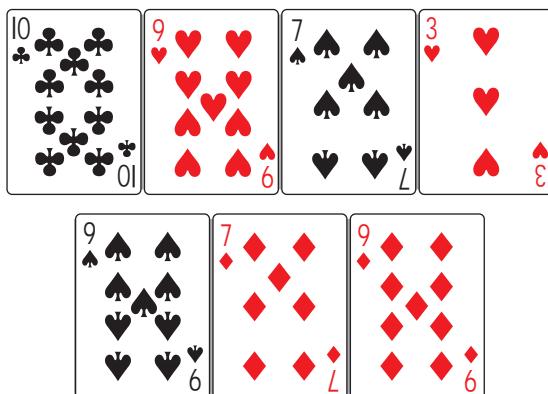
- (a) Explain why this is a legitimate probability model.
 (b) What is the probability that a randomly chosen student is studying a language other than English?
 (c) What is the probability that a randomly chosen student is studying French, German, or Spanish?

10.36 Car colors. Choose a new car or light truck at random and note its color. Here are the probabilities of the most popular colors for vehicles sold globally in 2010:¹⁰

Color	Silver	Black	White	Gray	Red	Blue	Beige, brown
Probability	0.26	0.24	0.16	0.16	0.06	0.05	0.03

- (a) What is the probability that the vehicle you choose has any color other than those listed?
 (b) What is the probability that a randomly chosen vehicle is neither silver nor white?

10.37 Drawing cards. You are about to draw a card at random (that is, all choices have the same probability) from a set of 7 cards. Although you can't see the cards, here they are:



- (a) What is the probability that you draw a 9?
 (b) What is the probability that you draw a red 9?
 (c) What is the probability that you do not draw a 7?

10.38 Loaded dice. There are many ways to produce crooked dice. To *load* a die so that 6 comes up too often and 1 (which is opposite 6) comes up too seldom, add a bit of lead to the filling of the spot on the 1 face. If a die is loaded so that 6 comes up with probability 0.2 and the probabilities of the 2, 3, 4, and 5 faces are not affected, what is the assignment of probabilities to the six faces?

10.39 A door prize. A party host gives a door prize to one guest chosen at random. There are 48 men and 42 women at the party. What is the probability that the prize goes to a woman? Explain how you arrived at your answer.

10.40 Race and ethnicity. The U.S. Census Bureau allows each person to choose from a long list of races. That is, in the eyes of the U.S. Census Bureau, you belong to whatever race you say you belong to. “Hispanic/Latino” is a separate category; Hispanics may be of any race. If we choose a resident of the United States at random, the U.S. Census Bureau gives these probabilities:¹¹

	Hispanic	Not Hispanic
Asian	0.001	0.044
Black	0.006	0.124
White	0.144	0.667
Other	0.005	0.009

- (a) Verify that this is a legitimate assignment of probabilities.
 (b) What is the probability that a randomly chosen American is Hispanic?

(c) Non-Hispanic whites are the historical majority in the United States. What is the probability that a randomly chosen American is not a member of this group?

Choose at random a person aged 15 to 44 years. Ask their age and who they live with (alone, with spouse, with other persons). Here is the probability model for 12 possible answers:¹²

	Age in Years			
	15–19	20–24	25–34	35–44
Alone	0.001	0.011	0.031	0.030
With spouse	0.001	0.023	0.155	0.216
With others	0.169	0.132	0.142	0.089

Exercises 10.41 to 10.43 use this probability model.

10.41 Living arrangements.

- (a) Why is this a legitimate finite probability model?
 (b) What is the probability that the person chosen is a 15- to 19-year-old who lives with others?
 (c) What is the probability that the person is 15 to 19 years old?
 (d) What is the probability that the person chosen lives with others?

10.42 Living arrangements, continued.

- (a) List the outcomes that make up the event

$$A = \{\text{The person chosen is either 15 to 19 years old or lives with others, or both}\}$$

- (b) What is $P(A)$? Explain carefully why $P(A)$ is not the sum of the probabilities you found in parts (b) and (c) of the previous exercise.

10.43 Living arrangements, continued.

- (a) What is the probability that the person chosen is 20 years old or older?
 (b) What is the probability that the person chosen does not live alone?

10.44 Spelling errors. Spell-checking software catches “nonword errors” that result in a string of letters that is not a word, as when “the” is typed as “teh.” When undergraduates are asked to type a 250-word essay (without spell-checking), the number X of nonword errors has the following distribution:

Value of X	0	1	2	3	4
Probability	0.1	0.2	0.3	0.3	0.1

- (a) Is the random variable X discrete or continuous? Why?
 (b) Write the event “at least one nonword error” in terms of X . What is the probability of this event?
 (c) Describe the event $X \leq 2$ in words. What is its probability? What is the probability that $X < 2$?

10.45 First digits again. A crook who never heard of Benford’s law might choose the first digits of his faked invoices so that all of 1, 2, 3, 4, 5, 6, 7, 8, and 9 are equally likely. Call the first digit of a randomly chosen fake invoice W for short.

- (a) Write the probability distribution for the random variable W .
 (b) Find $P(W \geq 6)$ and compare your result with the Benford’s law probability from Example 10.7.

10.46 Who gets interviewed? Abby, Deborah, Mei-Ling, Sam, and Roberto are students in a small seminar course. Their professor decides to choose two of them to interview about the course. To avoid unfairness, the choice will be made by drawing two names from a hat. (This is an SRS of size 2.)

- (a) Write down all possible choices of two of the five names. This is the sample space.
 (b) The random drawing makes all choices equally likely. What is the probability of each choice?
 (c) What is the probability that Mei-Ling is chosen?
 (d) Abby, Deborah, and Mei-Ling liked the course. Sam and Roberto did not like the course. What is the probability that both people selected liked the course?

10.47 Birth order. A couple plans to have three children. There are 8 possible arrangements of girls and boys. For example, GGB means the first two children are girls and the third child is a boy.

All 8 arrangements are (approximately) equally likely.

- (a) Write down all 8 arrangements of the sexes of three children. What is the probability of any one of these arrangements?
 (b) Let X be the number of girls the couple has. What is the probability that $X = 2$?
 (c) Starting from your work in (a), find the distribution of X . That is, what values can X take, and what are the probabilities for each value?



Picture Press/Alamy

10.48 Unusual dice. Nonstandard dice can produce interesting distributions of outcomes. You have two balanced, six-sided dice. One is a standard die, with faces having 1, 2, 3, 4, 5, and 6 spots. The other die has three faces with 0 spots and three faces with 6 spots. Find the probability distribution for the total number of spots Y on the up-faces when you roll these two dice. (*Hint:* Start with a picture like Figure 10.2 for the possible up-faces. Label the three 0 faces on the second die 0a, 0b, 0c in your picture, and similarly distinguish the three 6 faces.)

10.49 Random numbers. Many random number generators allow users to specify the range of the random numbers to be produced. Suppose that you specify that the random number T can take any value between 0 and 2. Then the density curve of the outcomes has constant height between 0 and 2, and height 0 elsewhere.

- (a) Is the random variable Y discrete or continuous? Why?
 (b) What is the height of the density curve between 0 and 2? Draw a graph of the density curve.
 (c) Use your graph from (b) and the fact that probability is area under the curve to find $P(Y \leq 1)$.

10.50 More random numbers. Find these probabilities as areas under the density curve you sketched in Exercise 10.49.

- (a) $P(0.5 < Y < 1.3)$
 (b) $P(Y \geq 0.8)$

10.51 Survey accuracy. A sample survey contacted an SRS of 3050 registered voters shortly before the 2008 presidential election and asked respondents whom they planned to vote for. Election results show that 53% of registered voters voted for Barack Obama. We will see later that in this situation the proportion of the sample who planned to vote for Barack Obama (call this proportion V) has approximately the Normal distribution with mean $\mu = 0.53$ and standard deviation $\sigma = 0.009$.

- (a) If the respondents answer truthfully, what is $P(0.51 \leq V \leq 0.55)$? This is the probability that the sample proportion V estimates the population proportion 0.53 within plus or minus 0.02.
 (b) In fact, 55% of the respondents said they planned to vote for Barack Obama ($V = 0.55$). If respondents answer truthfully, what is $P(V \geq 0.55)$?

10.52 Friends. How many close friends do you have? Suppose that the number of close friends adults claim to have varies from person to person with mean $\mu = 9$ and standard deviation $\sigma = 2.5$. An opinion poll asks this question of an SRS of 1100 adults. We will see later that in this situation the sample mean response \bar{x} has approximately the Normal distribution with mean 9 and standard deviation 0.075. What is $P(8.9 \leq \bar{x} \leq 9.1)$, the probability that the sample result \bar{x} estimates the population truth $\mu = 9$ to within ± 0.1 ?

10.53 Playing Pick 4. The Pick 4 games in many state lotteries announce a four-digit winning number each day. Each

of the 10,000 possible numbers 0000 to 9999 has the same chance of winning. You win if your choice matches the winning digits. Suppose your chosen number is 5974.

(a) What is the probability that the winning number matches your number exactly?

(b) What is the probability that the winning number has the same digits as your number *in any order*?

10.54 Nickels falling over. You may feel that it is obvious that the probability of a head in tossing a coin is about 1/2 because the coin has two faces. Such opinions are not always correct. Stand a nickel on edge on a hard, flat surface. Pound the surface with your hand so that the nickel falls over. What is the probability that it falls with heads upward? Make at least 50 trials to estimate the probability of a head.

10.55 What probability doesn't say. The idea of probability is that the *proportion* of heads in many tosses of a balanced coin eventually gets close to 0.5. But does the actual *count* of heads get close to one-half the number of tosses? Let's find out. Set the "Probability of heads" in the *Probability* applet to 0.5 and the number of tosses to 40. You can extend the number of tosses by clicking "Toss" again to get 40 more. Don't click "Reset" during this exercise.

(a) After 40 tosses, what is the proportion of heads? What is the count of heads? What is the difference between the count of heads and 20 (one-half the number of tosses)?

(b) Keep going to 120 tosses. Again record the proportion and count of heads and the difference between the count and 60 (half the number of tosses).

(c) Keep going. Stop at 240 tosses and again at 480 tosses to record the same facts. Although it may take a long time, the laws of probability say that the proportion of heads will always get close to 0.5 and also that the difference between the count of heads and half the number of tosses will always grow without limit.

10.56 LeBron's free throws. The basketball player LeBron James makes about three-quarters of his free throws over an entire season. Use the *Probability* applet or statistical software to simulate 100 free throws shot by a player who has probability 0.75 of making each shot. (In most software, the key phrase to look for is "Bernoulli trials." This is the technical term for independent trials with Yes/No outcomes. Our outcomes here are "Hit" and "Miss.")

(a) What percent of the 100 shots did he hit?

(b) Examine the sequence of hits and misses. How long was the longest run of shots made? Of shots missed? (Sequences of random outcomes often show runs longer than our intuition thinks likely.)

10.57 Simulating an opinion poll. A 2009 opinion poll showed that about 40% of the American public have very little or no confidence in big business. Suppose that this is exactly true. Choosing a person at random then has probability 0.40 of getting one who has very little or no confidence in big business. Use the *Probability* applet or statistical software to simulate choosing many people at random. (In most software, the key phrase to look for is "Bernoulli trials." This is the technical term for independent trials with Yes/No outcomes. Our outcomes here are "Favorable" or not.)

(a) Simulate drawing 20 people, then 80 people, then 320 people. What proportion have very little or no confidence in big business in each case? We expect (but because of chance variation we can't be sure) that the proportion will be closer to 0.40 in longer runs of trials.

(b) Simulate drawing 20 people 10 times and record the percents in each sample who have very little or no confidence in big business. Then simulate drawing 320 people 10 times and again record the 10 percents. Which set of 10 results is less variable? We expect the results of samples of size 320 to be more predictable (less variable) than the results of samples of size 20. That is "long-run regularity" showing itself.

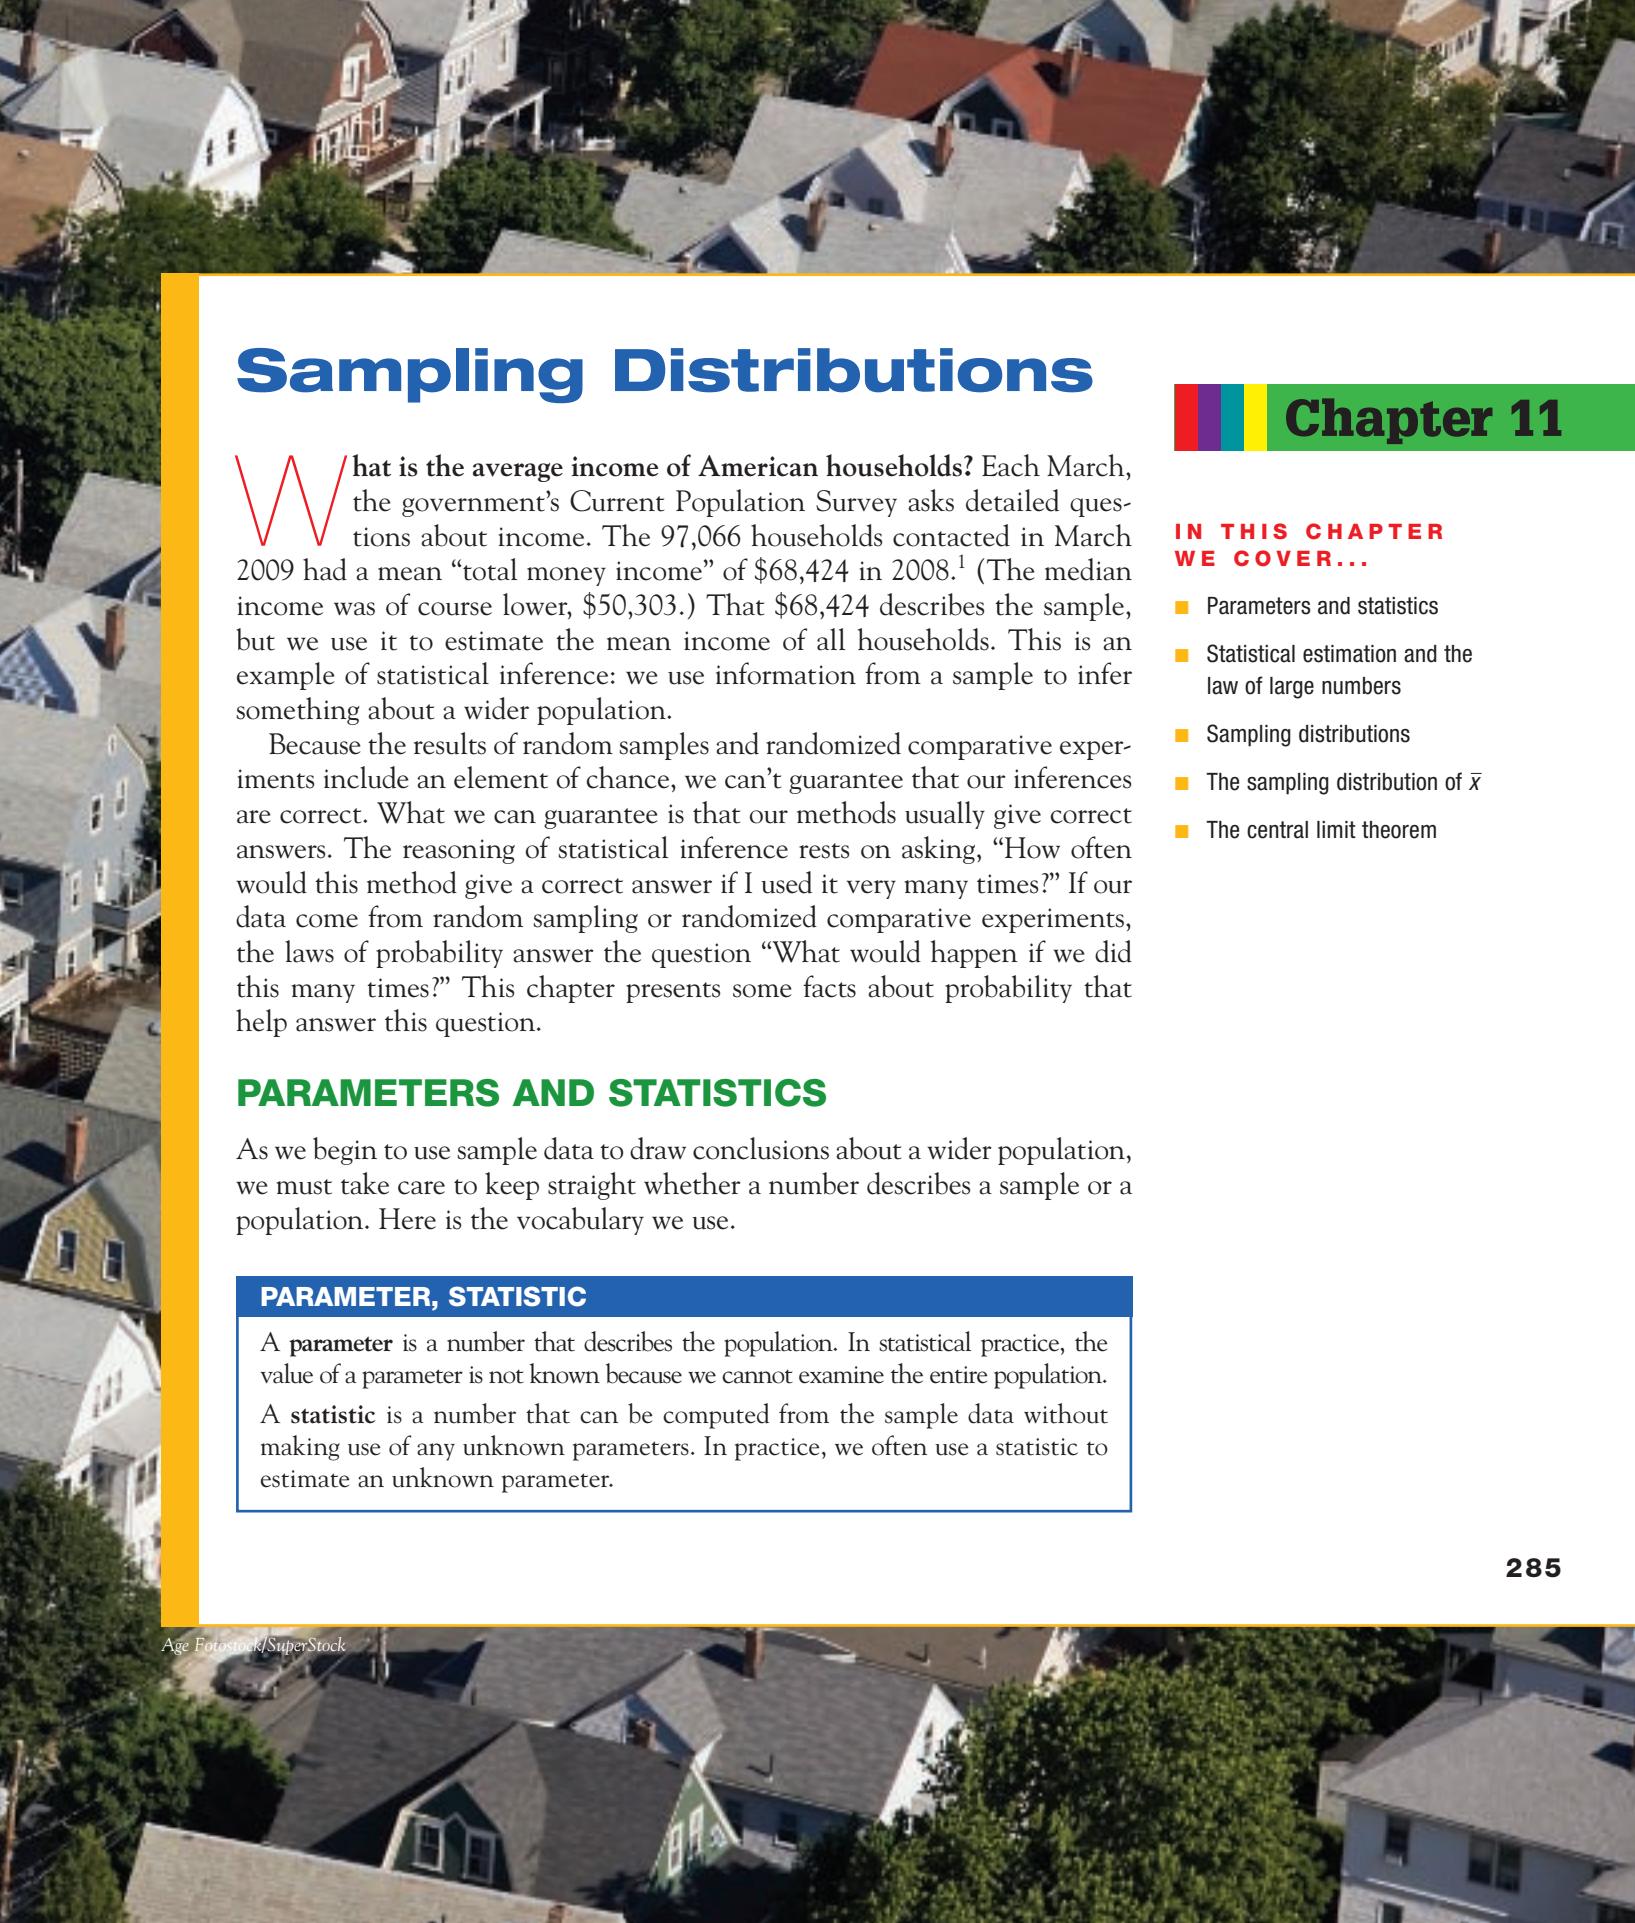


EXPLORING THE WEB

10.58 Super Bowl odds. Oddsmakers often list the odds for certain sporting events on the Web. For example, one can find the current odds of winning the next Super Bowl for each NFL team. We found a list of such odds at www.vegas.com/gaming/futures/superbowl.html. When an oddsmaker says the odds are A to B of winning, he or she means that the probability of winning is $B/(A + B)$. For example, when we checked the Web site listed above, the odds that the Indianapolis Colts would win Super Bowl XLIV were 13 to 2. This corresponds to a probability of winning of $2/(13 + 2) = 2/15$.

On the Web, find the current odds, according to an oddsmaker, of winning the Super Bowl for each NFL team. Convert these odds to probabilities. Do these probabilities satisfy Rules 1 and 2 given in this chapter? If they don't, can you think of a reason why?





Sampling Distributions

What is the average income of American households? Each March, the government's Current Population Survey asks detailed questions about income. The 97,066 households contacted in March 2009 had a mean "total money income" of \$68,424 in 2008.¹ (The median income was of course lower, \$50,303.) That \$68,424 describes the sample, but we use it to estimate the mean income of all households. This is an example of statistical inference: we use information from a sample to infer something about a wider population.

Because the results of random samples and randomized comparative experiments include an element of chance, we can't guarantee that our inferences are correct. What we can guarantee is that our methods usually give correct answers. The reasoning of statistical inference rests on asking, "How often would this method give a correct answer if I used it very many times?" If our data come from random sampling or randomized comparative experiments, the laws of probability answer the question "What would happen if we did this many times?" This chapter presents some facts about probability that help answer this question.

PARAMETERS AND STATISTICS

As we begin to use sample data to draw conclusions about a wider population, we must take care to keep straight whether a number describes a sample or a population. Here is the vocabulary we use.

PARAMETER, STATISTIC

A **parameter** is a number that describes the population. In statistical practice, the value of a parameter is not known because we cannot examine the entire population.

A **statistic** is a number that can be computed from the sample data without making use of any unknown parameters. In practice, we often use a statistic to estimate an unknown parameter.



Chapter 11

IN THIS CHAPTER WE COVER...

- Parameters and statistics
- Statistical estimation and the law of large numbers
- Sampling distributions
- The sampling distribution of \bar{x}
- The central limit theorem



EXAMPLE 11.1 Household earnings

The mean income of the sample of 97,066 households contacted by the Current Population Survey was $\bar{x} = \$68,424$. The number \$68,424 is a *statistic* because it describes this one Current Population Survey sample. The population that the poll wants to draw conclusions about is all 117 million U.S. households. The *parameter* of interest is the mean income of all these households. We don't know the value of this parameter. ■

population mean μ
Population standard deviation
sample mean \bar{x}
sample standard deviations

Remember s and p : statistics come from samples, and parameters come from populations. As long as we were just doing data analysis, searching for patterns or summarizing features of our data, the distinction between population and sample was not important. Now, as we begin to understand what our data (sample) tell us about a population, it is essential. The notation we use must reflect this distinction. We write μ (the Greek letter mu) for the **mean of a population** and σ (the Greek letter sigma) for the **standard deviation of a population**. These are fixed parameters that are unknown when we use a sample for inference. The **mean of the sample** is the familiar \bar{x} , the average of the observations in the population standard deviation σ sample. The **standard deviation of the sample** is denoted by s , the standard deviation of the observations in the sample. These are statistics that would almost certainly take different values if we chose another sample from the same population. The sample mean \bar{x} and sample standard deviation s from a sample or an experiment are estimates of the mean μ and standard deviation σ of the underlying population.



APPLY YOUR KNOWLEDGE

11.1 Genetic engineering. Here's a new idea for treating advanced melanoma, the most serious kind of skin cancer. Genetically engineer white blood cells to better recognize and destroy cancer cells, then infuse these cells into patients. The subjects in a small initial study of this approach were 11 patients whose melanoma had not responded to existing treatments. One outcome of this experiment was measured by a test for the presence of cells that trigger an immune response in the body and so may help fight cancer. The mean counts of active cells per 100,000 cells for the 11 subjects were **3.8** before infusion and **160.2** after infusion. Is each of the boldface numbers a parameter or a statistic?

11.2 Florida voters. Florida played a key role in the 2000 and 2004 presidential elections. Voter registration records in August 2010 show that **41%** of Florida voters are registered as Democrats and **36%** as Republicans. (Most of the others did not choose a party.) To test a random digit dialing device that you plan to use to poll voters for the 2010 Senate elections, you use it to call 250 randomly chosen residential telephones in Florida. Of the registered voters contacted, **34%** are registered Democrats. Is each of the boldface numbers a parameter or a statistic?

11.3 Human growth hormone. Researchers surveyed more than 230 American male weight lifters, ranging in age from 18 to 40, and found that **12%** of them had used

HGH, which has been banned in sports for more than 20 years now. The median usage time for those who reported HGH use was 23 weeks. Is each of the boldface numbers a parameter or a statistic?

STATISTICAL ESTIMATION AND THE LAW OF LARGE NUMBERS

Statistical inference uses sample data to draw conclusions about the entire population. Because good samples are chosen randomly, statistics such as \bar{x} computed from these samples are random variables. We can describe the behavior of a sample statistic by a probability model that answers the question “What would happen if we did this many times?” Here is an example that will lead us toward the probability ideas most important for statistical inference.

EXAMPLE 11.2 Does this wine smell bad?

Sulfur compounds such as dimethyl sulfide (DMS) are sometimes present in wine. DMS causes “off-odors” in wine, so winemakers want to know the odor threshold, the lowest concentration of DMS that the human nose can detect. Different people have different thresholds, so we start by asking about the mean threshold μ in the population of all adults. The number μ is a parameter that describes this population.

To estimate μ , we present tasters with both natural wine and the same wine spiked with DMS at different concentrations to find the lowest concentration at which they identify the spiked wine. Here are the odor thresholds (measured in micrograms of DMS per liter of wine) for 10 randomly chosen subjects:

28 40 28 33 20 31 29 27 17 21

The mean threshold for these subjects is $\bar{x} = 27.4$. It seems reasonable to use the sample result $\bar{x} = 27.4$ to estimate the unknown μ . An SRS should fairly represent the population, so the mean \bar{x} of the sample should be somewhere near the mean μ of the population. Of course, we don’t expect \bar{x} to be exactly equal to μ . We realize that if we choose another SRS, the luck of the draw will probably produce a different \bar{x} . ■

If \bar{x} is rarely exactly right and varies from sample to sample, why is it nonetheless a reasonable estimate of the population mean μ ? Here is one answer: *if we keep on taking larger and larger samples, the statistic \bar{x} is guaranteed to get closer and closer to the parameter μ .* We have the comfort of knowing that if we can afford to keep on measuring more subjects, eventually we will estimate the mean odor threshold of all adults very accurately. This remarkable fact is called the *law of large numbers*. It is remarkable because it holds for *any* population, not just for some special class such as Normal distributions.



Foodpix/Getty Images



High-tech gambling

There are twice as many slot machines as

bank ATMs in the United States. Once upon a time, you put in a coin and pulled the lever to spin three wheels, each with 20 symbols. No longer. Now the machines are video games with flashy graphics and outcomes produced by random number generators. Machines can accept many coins at once, can pay off on a bewildering variety of outcomes, and can be networked to allow common jackpots. Gamblers still search for systems, but in the long run the law of large numbers guarantees the house its 5% profit.

LAW OF LARGE NUMBERS

Draw observations at random from any population with finite mean μ . As the number of observations drawn increases, the mean \bar{x} of the observed values gets closer and closer to the mean μ of the population.

The law of large numbers can be proved mathematically starting from the basic laws of probability. The behavior of \bar{x} is similar to the idea of probability. In the long run, the proportion of outcomes taking any value gets close to the probability of that value, and the average outcome gets close to the population mean. Figure 10.1 (page 261) shows how proportions approach probability in one example. Here is an example of how sample means approach the population mean.

EXAMPLE 11.3 The law of large numbers in action

In fact, the distribution of odor thresholds among all adults has mean 25. The mean $\mu = 25$ is the true value of the parameter we seek to estimate. Figure 11.1 shows how the sample mean \bar{x} of an SRS drawn from this population changes as we add more subjects to our sample.

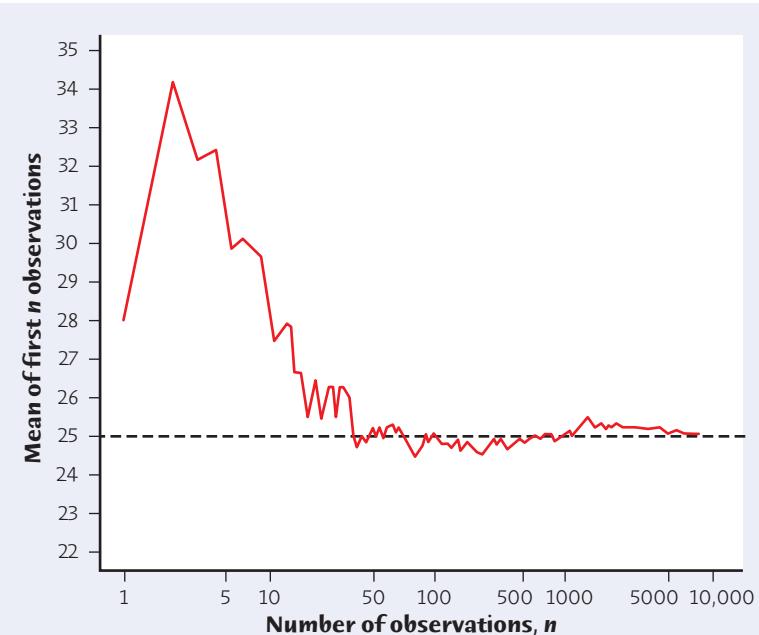


FIGURE 11.1

The law of large numbers in action: as we take more observations, the sample mean \bar{x} always approaches the mean μ of the population.

The first subject in Example 11.2 had threshold 28, so the line in Figure 11.1 starts there. The mean for the first two subjects is

$$\bar{x} = \frac{28 + 40}{2} = 34$$

This is the second point on the graph. At first, the graph shows that the mean of the sample changes as we take more observations. Eventually, however, the mean of the observations gets close to the population mean $\mu = 25$ and settles down at that value.

If we started over, again choosing people at random from the population, we would get a different path from left to right in Figure 11.1. The law of large numbers says that whatever path we get will always settle down at 25 as we draw more and more people. ■

The *Law of Large Numbers* applet animates Figure 11.1 in a different setting. You can use the applet to watch \bar{x} change as you average more observations until it eventually settles down at the mean μ .



The law of large numbers is the foundation of such business enterprises as gambling casinos and insurance companies. The winnings (or losses) of a gambler on a few plays are uncertain—that's why some people find gambling exciting. In Figure 11.1, the mean of even 100 observations is not yet very close to μ . It is only in the long run that the mean outcome is predictable. The house plays tens of thousands of times. So the house, unlike individual gamblers, can count on the long-run regularity described by the law of large numbers. The average winnings of the house on tens of thousands of plays will be very close to the mean of the distribution of winnings. Needless to say, this mean guarantees the house a profit. That's why gambling can be a business.

APPLY YOUR KNOWLEDGE

11.4 The law of large numbers made visible. Roll two balanced dice and count the total spots on the up-faces. The probability model appears in Example 10.5 (page 265). You can see that this distribution is symmetric with 7 as its center, so it's no surprise that the mean is $\mu = 7$. This is the population mean for the idealized population that contains the results of rolling two dice forever. The law of large numbers says that the average \bar{x} from a finite number of rolls gets closer and closer to 7 as we do more and more rolls.



- Click “More dice” once in the *Law of Large Numbers* applet to get two dice. Click “Show mean” to see the mean 7 on the graph. Leaving the number of rolls at 1, click “Roll dice” three times. How many spots did each roll produce? What is the average for the three rolls? You see that the graph displays at each point the average number of spots for all rolls up to the last one. This is exactly like Figure 11.1.
- Set the number of rolls to 100 and click “Roll dice.” The applet rolls the two dice 100 times. The graph shows how the average count of spots changes as we make more rolls. That is, the graph shows \bar{x} as we continue to roll the dice. Sketch (or print out) the final graph.

- (c) Repeat your work from (b). Click “Reset” to start over, then roll two dice 100 times. Make a sketch of the final graph of the mean \bar{x} against the number of rolls. Your two graphs will often look very different. What they have in common is that the average eventually gets close to the population mean $\mu = 7$. The law of large numbers says that this will *always* happen if you keep on rolling the dice.

11.5 Insurance. The idea of insurance is that we all face risks that are unlikely but carry high cost. Think of a fire or flood destroying your apartment. Insurance spreads the risk: we all pay a small amount, and the insurance policy pays a large amount to those few of us whose apartments are damaged. An insurance company looks at the records for millions of apartment owners and sees that the mean loss from apartment damage in a year is $\mu = \$75$ per person. (Most of us have no loss, but a few lose most of their possessions. The \$75 is the average loss.) The company plans to sell fire insurance for \$75 plus enough to cover its costs and profit. Explain clearly why it would be unwise to sell only 12 policies. Then explain why selling thousands of such policies is a safe business.

SAMPLING DISTRIBUTIONS

The law of large numbers assures us that if we measure enough subjects, the statistic \bar{x} will eventually get very close to the unknown parameter μ . But the odor threshold study in Example 11.2 had just 10 subjects. What can we say about estimating μ by \bar{x} from a sample of 10 subjects? Put this one sample in the context of all such samples by asking, “What would happen if we took many samples of 10 subjects from this population?” Here’s how to answer this question:

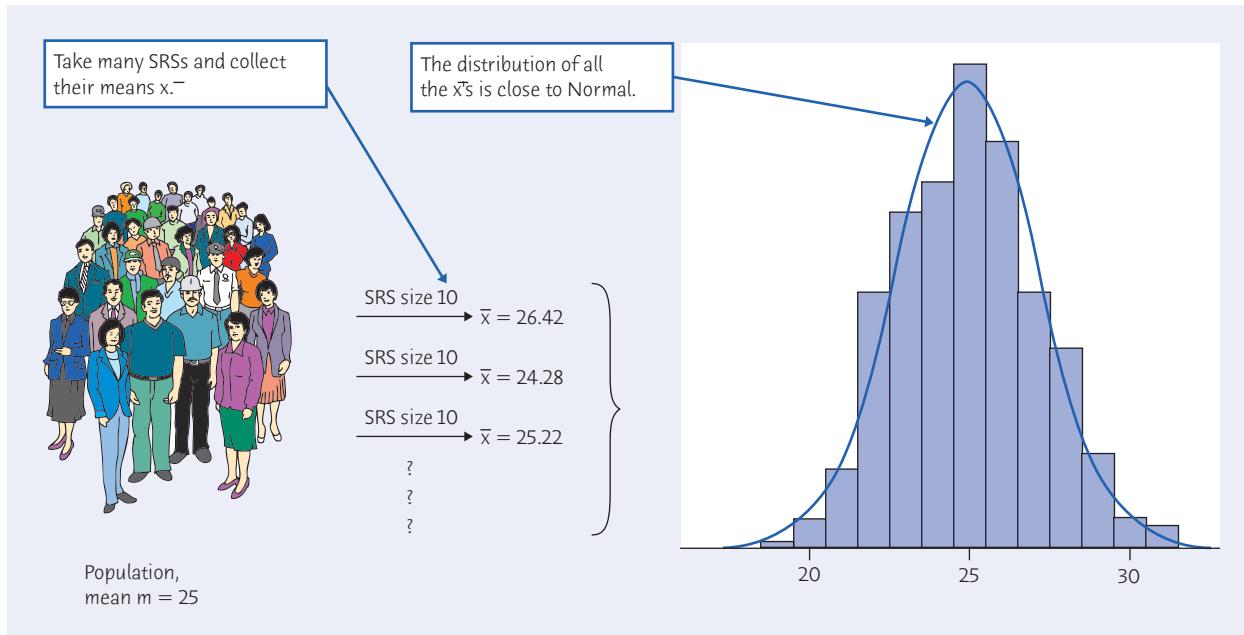
- Take a large number of samples of size 10 from the population.
- Calculate the sample mean \bar{x} for each sample.
- Make a histogram of the values of \bar{x} .
- Examine the shape, center, and spread of the distribution displayed in the histogram.

In practice it is too expensive to take many samples from a large population such as all adult U.S. residents. But we can imitate many samples by using software. Using software to imitate chance behavior is called **simulation**.

EXAMPLE 11.4 What would happen in many samples?

Extensive studies have found that the DMS odor threshold of adults follows roughly a Normal distribution with mean $\mu = 25$ micrograms per liter and standard deviation $\sigma = 7$ micrograms per liter. We call this the *population distribution* of odor threshold.

Figure 11.2 illustrates the process of choosing many samples and finding the sample mean threshold \bar{x} for each one. Follow the flow of the figure from the population at the left, to choosing an SRS and finding the \bar{x} for this sample, to collecting together the \bar{x} ’s from many samples. The first sample has $\bar{x} = 26.42$. The second sample

**FIGURE 11.2**

The idea of a sampling distribution: take many samples from the same population, collect the \bar{x} 's from all the samples, and display the distribution of the \bar{x} 's. The histogram shows the results of 1000 samples.

contains a different 10 people, with $\bar{x} = 24.28$, and so on. The histogram at the right of the figure shows the distribution of the values of \bar{x} from 1000 separate SRSs of size 10. This histogram displays the *sampling distribution* of the statistic \bar{x} . ■

POPULATION DISTRIBUTION, SAMPLING DISTRIBUTION

The **population distribution** of a variable is the distribution of values of the variable among all the individuals in the population.

The **sampling distribution** of a statistic is the distribution of values taken by the statistic in all possible samples of the same size from the same population.

Be careful: The population distribution describes the *individuals* that make up the population. A sampling distribution describes how a *statistic* varies in many samples from the population.

Strictly speaking, the sampling distribution is the ideal pattern that would emerge if we looked at all possible samples of size 10 from our population. A distribution obtained from a fixed number of trials, like the 1000 trials in Figure 11.2, is only an approximation to the sampling distribution. One of the uses of probability theory in statistics is to obtain sampling distributions without simulation. The interpretation of a sampling distribution is the same, however, whether we obtain it by simulation or by the mathematics of probability.

We can use the tools of data analysis to describe any distribution. Let's apply those tools to Figure 11.2. What can we say about the shape, center, and spread of this distribution?

- **Shape:** It looks Normal! Detailed examination confirms that the distribution of \bar{x} from many samples is very close to Normal.
- **Center:** The mean of the 1000 \bar{x} 's is 24.95. That is, the distribution is centered very close to the population mean $\mu = 25$.
- **Spread:** The standard deviation of the 1000 \bar{x} 's is 2.217, notably smaller than the standard deviation $\sigma = 7$ of the population of individual subjects.

Although these results describe just one simulation of a sampling distribution, they reflect facts that are true whenever we use random sampling.

APPLY YOUR KNOWLEDGE

11.6 Sampling distribution versus population distribution. During World War II, 12,000 able-bodied male undergraduates at the University of Illinois participated in required physical training. Each student ran a timed mile. Their times followed the Normal distribution with mean 7.11 minutes and standard deviation 0.74 minute. An SRS of 100 of these students has mean time $\bar{x} = 7.15$ minutes. A second SRS of size 100 has mean $\bar{x} = 6.97$ minutes. After many SRSs, the many values of the sample mean \bar{x} follow the Normal distribution with mean 7.11 minutes and standard deviation 0.074 minute.

- What is the population? What values does the population distribution describe? What is this distribution?
- What values does the sampling distribution of \bar{x} describe? What is the sampling distribution?

11.7 Generating a sampling distribution. Let's illustrate the idea of a sampling distribution in the case of a very small sample from a very small population. The population is the scores of 10 students on an exam:  SAMPLING

Student	0	1	2	3	4	5	6	7	8	9
Score	86	63	81	55	72	72	65	66	75	59

The parameter of interest is the mean score μ in this population. The sample is an SRS of size $n = 4$ drawn from the population. Because the students are labeled 0 to 9, a single random digit from Table B chooses one student for the sample.

- Find the mean of the 10 scores in the population. This is the population mean μ .
- Use the first digits in row 116 of Table B to draw an SRS of size 4 from this population. What are the four scores in your sample? What is their mean \bar{x} ? This statistic is an estimate of μ .
- Repeat this process 9 more times, using the first digits in rows 117 to 125 of Table B. Make a histogram of the 10 values of \bar{x} . You are constructing the sampling distribution of \bar{x} . Is the center of your histogram close to μ ?

THE SAMPLING DISTRIBUTION OF \bar{x}

Figure 11.2 suggests that when we choose many SRSs from a population, the sampling distribution of the sample means is centered at the mean of the original population and is less spread out than the distribution of individual observations. Here are the facts.

MEAN AND STANDARD DEVIATION OF A SAMPLE MEAN²

Suppose that \bar{x} is the mean of an SRS of size n drawn from a large population with mean μ and standard deviation σ . Then the sampling distribution of \bar{x} has mean μ and standard deviation σ/\sqrt{n} .

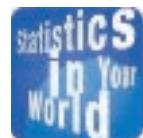
These facts about the mean and the standard deviation of the sampling distribution of \bar{x} are true for *any* population, not just for some special class such as Normal distributions. They have important implications for statistical inference:

- The mean of the statistic \bar{x} is always equal to the mean μ of the population. That is, the sampling distribution of \bar{x} is centered at μ . In repeated sampling, \bar{x} will sometimes fall above the true value of the parameter μ and sometimes below, but there is no systematic tendency to overestimate or underestimate the parameter. This makes the idea of lack of bias in the sense of “no favoritism” more precise. Because the mean of \bar{x} is equal to μ , we say that the statistic \bar{x} is an **unbiased estimator** of the parameter μ .
- An unbiased estimator is “correct on the average” in many samples. How close the estimator falls to the parameter in most samples is determined by the spread of the sampling distribution. If individual observations have standard deviation σ , then sample means \bar{x} from samples of size n have standard deviation σ/\sqrt{n} . That is, **averages are less variable than individual observations**.
- Not only is the standard deviation of the distribution of \bar{x} smaller than the standard deviation of individual observations, but it gets smaller as we take larger samples. **The results of large samples are less variable than the results of small samples.**

The upshot of all this is that we can trust the sample mean from a large random sample to estimate the population mean accurately. If the sample size n is large, the standard deviation of \bar{x} is small, and almost all samples will give values of \bar{x} that lie very close to the true parameter μ . However, the standard deviation of the sampling distribution gets smaller only at the rate \sqrt{n} . To cut the standard deviation of \bar{x} in half, we must take four times as many observations, not just twice as many. So very precise estimates (estimates with very small standard deviation) may be expensive.

We have described the center and spread of the sampling distribution of a sample mean \bar{x} , but not its shape. The shape of the sampling distribution depends on the shape of the population distribution. In one important case there is a

unbiased estimator



Sample size matters

The new thing in baseball is using statistics to evaluate players, with new measures of performance to help decide which players are worth the high salaries they demand. This challenges traditional subjective evaluation of young players and the usefulness of traditional measures such as batting average. But success has led many major league teams to hire statisticians. The statisticians say that sample size matters in baseball also: the 162-game regular season is long enough for the better teams to come out on top, but 5-game and 7-game play-off series are so short that luck has a lot to say about who wins.

simple relationship between the two distributions: if the population distribution is Normal, then so is the sampling distribution of the sample mean.

SAMPLING DISTRIBUTION OF A SAMPLE MEAN

If individual observations have the $N(\mu, \sigma)$ distribution, then the sample mean \bar{x} of an SRS of size n has the $N(\mu, \sigma/\sqrt{n})$ distribution.

Notice that if the population distribution is Normal, then the sampling distribution of the sample mean is Normal regardless of the sample size n .

EXAMPLE 11.5 Population distribution, sampling distribution

If we measure the DMS odor thresholds of individual adults, the values follow the Normal distribution with mean $\mu = 25$ micrograms per liter and standard deviation $\sigma = 7$ micrograms per liter. This is the population distribution of odor threshold.

Take many SRSs of size 10 from this population and find the sample mean \bar{x} for each sample, as in Figure 11.2. The sampling distribution describes how the values of \bar{x} vary among samples. That sampling distribution is also Normal, with mean $\mu = 25$ and standard deviation

$$\frac{\sigma}{\sqrt{n}} = \frac{7}{\sqrt{10}} = 2.2136$$

Figure 11.3 contrasts these two Normal distributions. Both are centered at the population mean, but sample means are much less variable than individual observations.

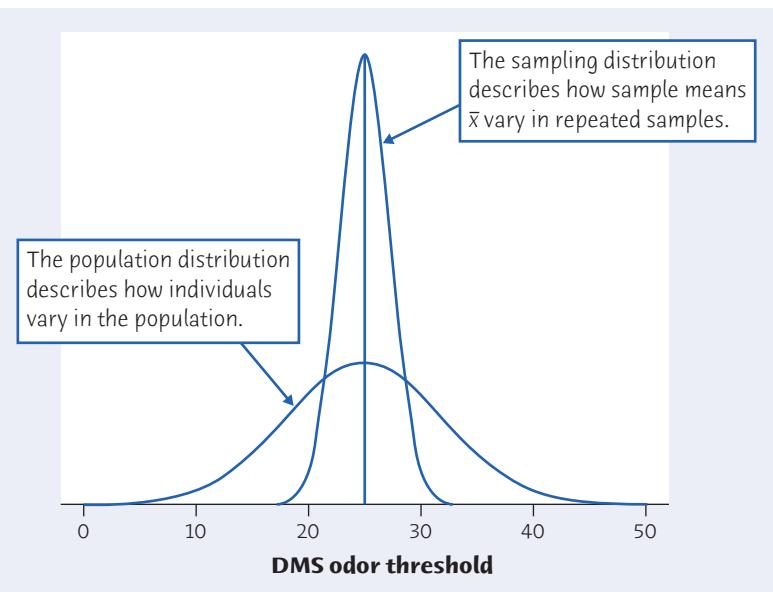


FIGURE 11.3

The distribution of single observations (the population distribution) compared with the sampling distribution of the means \bar{x} of 10 observations, for Example 11.5. Both have the same mean, but averages are less variable than individual observations.

The smaller variation of sample means shows up in probability calculations. You can show (using software or standardizing and using Table A) that about 52% of all adults have odor thresholds between 20 and 30. But almost 98% of means of samples of size 10 lie in this range. ■

APPLY YOUR KNOWLEDGE

- 11.8 A sample of young men.** A government sample survey plans to measure the blood cholesterol level of an SRS of men aged 20 to 34. The researchers will report the mean \bar{x} from their sample as an estimate of the mean cholesterol level μ in this population.
- Explain to someone who knows no statistics what it means to say that \bar{x} is an “unbiased” estimator of μ .
 - The sample result \bar{x} is an unbiased estimator of the population truth μ no matter what size SRS the study uses. Explain to someone who knows no statistics why a large sample gives more trustworthy results than a small sample.

- 11.9 Larger sample, more accurate estimate.** Suppose that in fact the blood cholesterol level of all men aged 20 to 34 follows the Normal distribution with mean $\mu = 186$ milligrams per deciliter (mg/dl) and standard deviation $\sigma = 41$ mg/dl.
- Choose an SRS of 100 men from this population. What is the sampling distribution of \bar{x} ? What is the probability that \bar{x} takes a value between 183 and 189 mg/dl? This is the probability that \bar{x} estimates μ within ± 3 mg/dl.
 - Choose an SRS of 1000 men from this population. Now what is the probability that \bar{x} falls within ± 3 mg/dl of μ ? The larger sample is much more likely to give an accurate estimate of μ .

- 11.10 Measurements in the lab.** Juan makes a measurement in a chemistry laboratory and records the result in his lab report. The standard deviation of students' lab measurements is $\sigma = 10$ milligrams. Juan repeats the measurement 4 times and records the mean \bar{x} of his 4 measurements.
- What is the standard deviation of Juan's mean result? (That is, if Juan kept on making 4 measurements and averaging them, what would be the standard deviation of all his \bar{x} 's?)
 - How many times must Juan repeat the measurement to reduce the standard deviation of \bar{x} to 2? Explain to someone who knows no statistics the advantage of reporting the average of several measurements rather than the result of a single measurement.

THE CENTRAL LIMIT THEOREM

The facts about the mean and standard deviation of \bar{x} are true no matter what the shape of the population distribution may be. But what is the shape of the sampling distribution when the population distribution is not Normal? *It is a remarkable fact that as the sample size increases, the distribution of \bar{x} changes shape: it looks less like that of the population and more like a Normal distribution.* When the sample is large enough, the distribution of \bar{x} is very close to Normal. This is true no matter what



What was that probability again?

Wall Street uses

fancy mathematics to predict the probabilities that fancy investments will go wrong. The probabilities are always too low—sometimes because something was assumed to be Normal but was not. Probability predictions in other areas also go wrong. In mid-September 2007, the New York Mets had probability 0.998 of making the National League play-offs, or so an elaborate calculation said. Then the Mets lost 12 of their final 17 games, the Phillies won 13 of their final 17, and the Mets were out.

shape the population distribution has, as long as the population has a finite standard deviation σ . This famous fact of probability theory is called the *central limit theorem*. It is much more useful than the fact that the distribution of \bar{x} is exactly Normal if the population is exactly Normal.

CENTRAL LIMIT THEOREM

Draw an SRS of size n from any population with mean μ and finite standard deviation σ . The **central limit theorem** says that when n is large, the sampling distribution of the sample mean \bar{x} is approximately Normal:

$$\bar{x} \text{ is approximately } N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$$

The central limit theorem allows us to use Normal probability calculations to answer questions about sample means from many observations even when the population distribution is not Normal.

More general versions of the central limit theorem say that the distribution of any sum or average of many small random quantities is close to Normal. This is true even if the quantities are correlated with each other (as long as they are not too highly correlated) and even if they have different distributions (as long as no one random quantity is so large that it dominates the others). The central limit theorem suggests why the Normal distributions are common models for observed data. Any variable that is a sum of many small influences will have approximately a Normal distribution.

How large a sample size n is needed for \bar{x} to be close to Normal depends on the population distribution. More observations are required if the shape of the population distribution is far from Normal. Here are two examples in which the population is far from Normal.

EXAMPLE 11.6 The central limit theorem in action

In March 2009, the Current Population Survey contacted 97,066 households. Figure 11.4(a) is a histogram of the earnings of the 75,310 households that had earned income greater than zero in 2008.³ As we expect, the distribution of earned incomes is strongly skewed to the right and very spread out. The right tail of the distribution is even longer than the histogram shows because there are too few high incomes for their bars to be visible on this scale. In fact, we cut off the earnings scale at \$400,000 to save space—a few households earned even more than \$400,000. The mean earnings for these 75,310 households was \$71,305.

Regard these 75,310 households as a population with mean $\mu = \$71,305$. Take an SRS of 100 households. The mean earnings in this sample is $\bar{x} = \$83,143$. That's higher than the mean of the population. Take another SRS of size 100. The mean for this sample is $\bar{x} = \$63,115$. That's less than the mean of the population. *What would happen if we did this many times?* Figure 11.4(b) is a histogram of the mean earnings for

500 samples, each of size 100. The scales in Figures 11.4(a) and 11.4(b) are the same, for easy comparison. Although the distribution of individual earnings is skewed and very spread out, the distribution of sample means is roughly symmetric and much less spread out.

Figure 11.4(c) zooms in on the center part of the histogram in Figure 11.4(b) to more clearly show its shape. Although $n = 100$ is not a very large sample size and the population distribution is extremely skewed, we can see that the distribution of sample means is close to Normal. ■

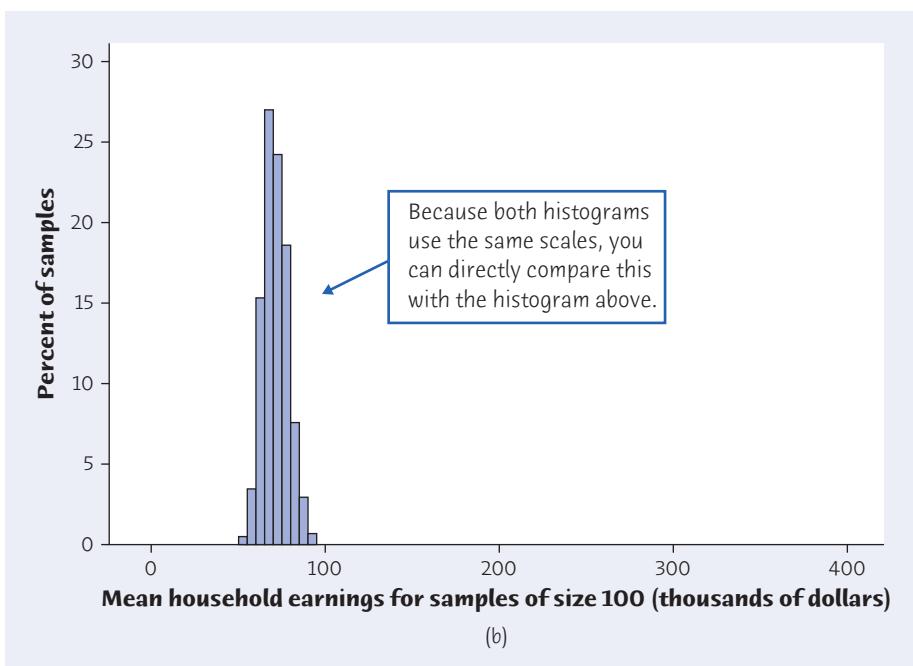
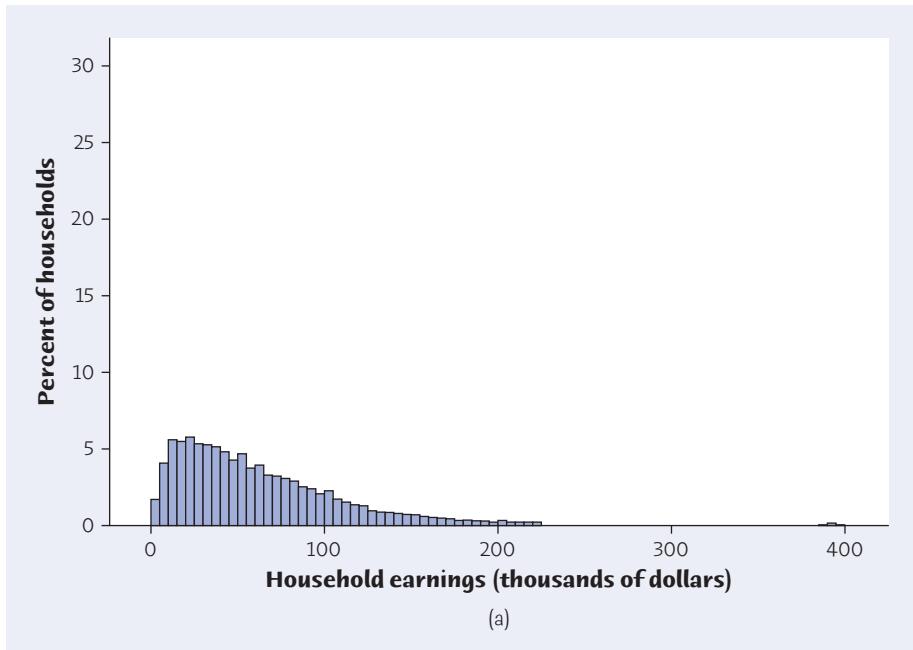
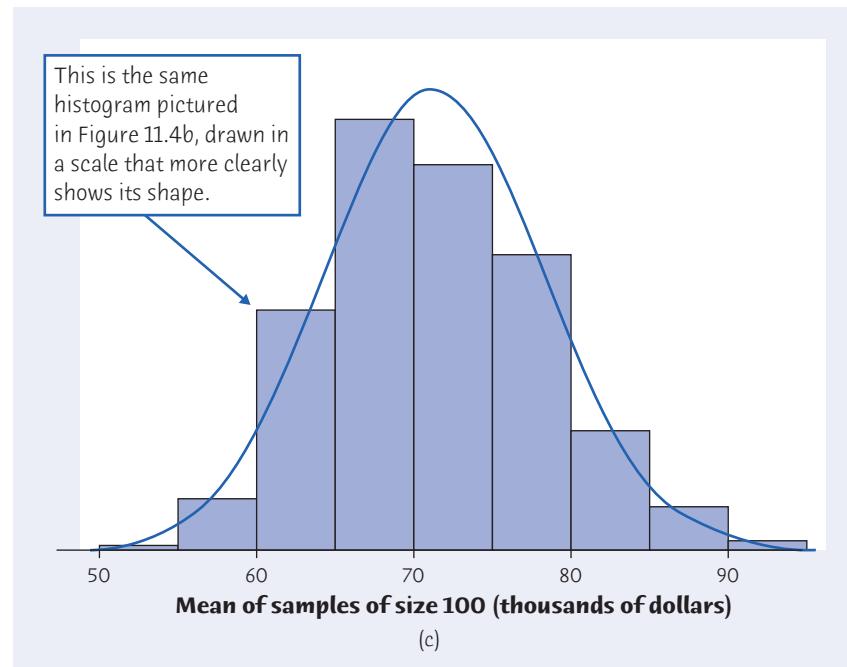


FIGURE 11.4

The central limit theorem in action, for Example 11.6. (a) The distribution of earned income in a population of 75,310 households. (b) The distribution of the mean earnings for 500 SRSs of 100 households each from this population.

FIGURE 11.4 (Continued)

(c) The distribution of the sample means in more detail: the shape is close to Normal.



Comparing Figure 11.4(a) with Figures 11.4(b) and 11.4(c) illustrates the two most important ideas of this chapter.

THINKING ABOUT SAMPLE MEANS

Means of random samples are **less variable** than individual observations.

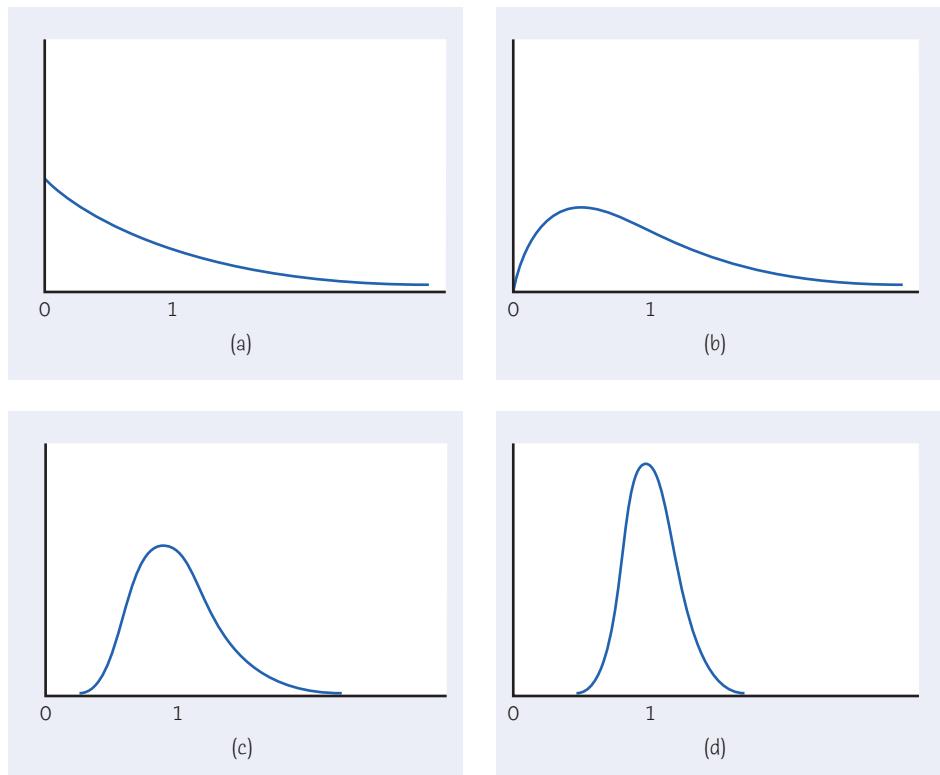
Means of random samples are **more Normal** than individual observations.



EXAMPLE 11.7 The central limit theorem in action

The *Central Limit Theorem* applet allows you to watch the central limit theorem in action. Figure 11.5 presents snapshots from the applet, drawn on the same scales for easy comparison. Figure 11.5(a) shows the population distribution, that is, the density curve of a single observation. This distribution is strongly right-skewed, and the most probable outcomes are near 0. The mean μ of this distribution is 1, and its standard deviation σ is also 1. This particular distribution is called an *exponential distribution*. Exponential distributions are used as models for the lifetime in service of electronic components and for the time required to serve a customer or repair a machine.

Figures 11.5(b), (c), and (d) are the density curves of the sample means of 2, 10, and 25 observations from this population. As n increases, the shape becomes more Normal. The mean remains at $\mu = 1$, and the standard deviation decreases, taking the value $1/\sqrt{n}$. The density curve for 10 observations is still somewhat skewed to the right but already resembles a Normal curve having $\mu = 1$ and $\sigma = 1/\sqrt{10} = 0.32$. The density curve for $n = 25$ is yet more Normal. The contrast between the shapes of the population distribution and of the distribution of the mean of 10 or 25 observations is striking. ■

**FIGURE 11.5**

The central limit theorem in action, for Example 11.7. The distribution of sample means \bar{x} from a strongly non-Normal population becomes more Normal as the sample size increases. (a) The distribution of 1 observation. (b) The distribution of \bar{x} for 2 observations. (c) The distribution of \bar{x} for 10 observations. (d) The distribution of \bar{x} for 25 observations.

Let's use Normal calculations based on the central limit theorem to answer a question about the very non-Normal distribution in Figure 11.5(a).

EXAMPLE 11.8 Maintaining air conditioners

STATE: The time (in hours) that a technician requires to perform preventive maintenance on an air-conditioning unit is governed by the exponential distribution whose density curve appears in Figure 11.5(a). The exponential distribution arises in many engineering and industrial problems, such as time until failure of a machine or time until a success. The mean time is $\mu = 1$ hour and the standard deviation is $\sigma = 1$ hour. Your company has a contract to maintain 70 of these units in an apartment building. You must schedule technicians' time for a visit to this building. Is it safe to budget an average of 1.1 hours for each unit? Or should you budget an average of 1.25 hours?



PLAN: We believe that the manufacturing and distribution process associated with this type of air-conditioning unit is such that variation from one unit to the next is random. Thus, we treat these 70 air conditioners as an SRS from all units of this type. What is the probability that the average maintenance time for 70 units exceeds 1.1 hours? That the average time exceeds 1.25 hours?

SOLVE: The central limit theorem says that the sample mean time \bar{x} spent working on 70 units has approximately the Normal distribution with mean equal to the population mean $\mu = 1$ hour and standard deviation

$$\frac{\sigma}{\sqrt{n}} = \frac{1}{\sqrt{70}} = 0.12 \text{ hour}$$

The distribution of \bar{x} is therefore approximately $N(1, 0.12)$. This Normal curve is the solid curve in Figure 11.6.

Using this Normal distribution, the probabilities we want are

$$P(\bar{x} > 1.10 \text{ hours}) = 0.2014$$

$$P(\bar{x} > 1.25 \text{ hours}) = 0.0182$$

Software gives these probabilities immediately, or you can standardize and use Table A. For example,

$$\begin{aligned} P(\bar{x} > 1.10) &= P\left(\frac{\bar{x} - 1}{0.12}\right) > P\left(\frac{1.10 - 1}{0.12}\right) \\ &= P(Z > 0.83) = 1 - 0.7967 = 0.2033 \end{aligned}$$

with the usual roundoff error. Don't forget to use standard deviation 0.12 in your software or when you standardize \bar{x} .

CONCLUDE: If you budget 1.1 hours per unit, there is a 20% chance that the technicians will not complete the work in the building within the budgeted time. This chance drops to 2% if you budget 1.25 hours. You should therefore budget 1.25 hours per unit. ■

Using more mathematics, we can start with the exponential distribution and find the actual density curve of \bar{x} for 70 observations. This is the dotted curve in Figure 11.6. You can see that the solid Normal curve is a good approximation. The exactly correct probability for 1.1 hours is an area to the right of 1.1 under the dotted density curve. It is 0.1977. The central limit theorem Normal approximation 0.2014 is off by only about 0.004.

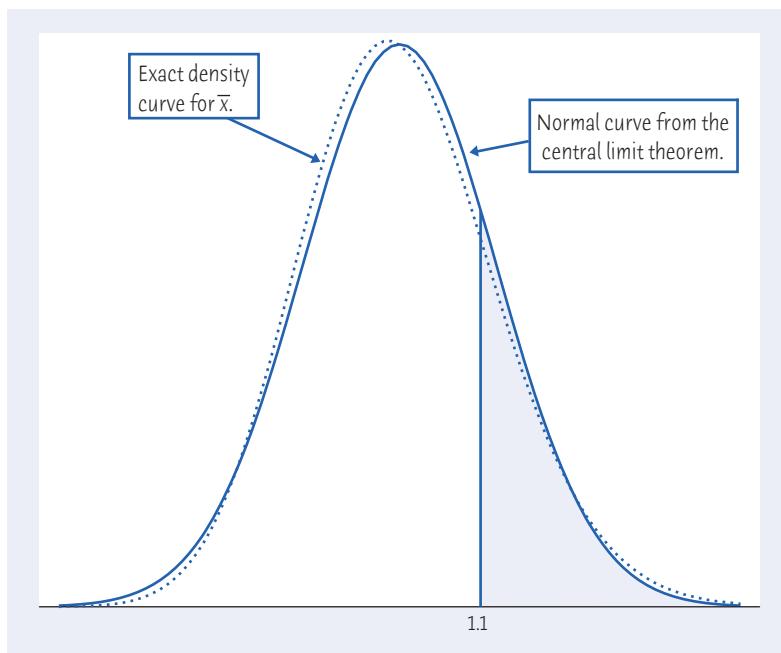


FIGURE 11.6

The exact distribution (dotted) and the Normal approximation from the central limit theorem (solid) for the average time needed to maintain an air conditioner, for Example 11.8. The probability we want is the area to the right of 1.1.

APPLY YOUR KNOWLEDGE

11.11 What does the central limit theorem say? Asked what the central limit theorem says, a student replies, “As you take larger and larger samples from a population, the histogram of the sample values looks more and more Normal.” Is the student right? Explain your answer.

11.12 Detecting gypsy moths. The gypsy moth is a serious threat to oak and aspen trees. A state agriculture department places traps throughout the state to detect the moths. When traps are checked periodically, the mean number of moths trapped is only 0.5, but some traps have several moths. The distribution of moth counts is discrete and strongly skewed, with standard deviation 0.7.

- What are the mean and standard deviation of the average number of moths \bar{x} in 50 traps?
- Use the central limit theorem to find the probability that the average number of moths in 50 traps is greater than 0.6.

11.13 More on insurance. An insurance company knows that in the entire population of millions of apartment owners, the mean annual loss from damage is $\mu = \$75$ and the standard deviation of the loss is $\sigma = \$300$. The distribution of losses is strongly right-skewed: most policies have \$0 loss, but a few have large losses. If the company sells 10,000 policies, can it safely base its rates on the assumption that its average loss will be no greater than \$85? Follow the four-step process as illustrated in Example 11.8.



Bruce Coleman/Alamy



CHAPTER 11 SUMMARY

CHAPTER SPECIFICS

- A **parameter** in a statistical problem is a number that describes a population, such as the **population mean μ** . To estimate an unknown parameter, use a **statistic** calculated from a sample, such as the **sample mean \bar{x}** .
- The **law of large numbers** states that the actual observed mean outcome \bar{x} must approach the mean μ of the population as the number of observations increases.
- The **population distribution** of a variable describes the values of the variable for all individuals in a population.
- The **sampling distribution** of a statistic describes the values of the statistic in all possible samples of the same size from the same population.
- When the sample is an SRS from the population, the **mean** of the sampling distribution of the sample mean \bar{x} is the same as the population mean μ . That is, \bar{x} is an **unbiased estimator** of μ .
- The **standard deviation** of the sampling distribution of \bar{x} is σ/\sqrt{n} for an SRS of size n if the population has standard deviation σ . That is, averages are less variable than individual observations.
- When the sample is an SRS from a population that has a Normal distribution, the sample mean \bar{x} also has a Normal distribution.
- Choose an SRS of size n from any population with mean μ and finite standard deviation σ . The **central limit theorem** states that when n is large, the sampling distribution of \bar{x} is approximately Normal. That is, averages are more Normal than individual observations. We can use the $N(\mu, \sigma/\sqrt{n})$ distribution to calculate approximate probabilities for events involving \bar{x} .

LINK IT

As we mentioned in Chapter 10, probability is the tool we will use to generalize from data produced by random samples and randomized comparative experiments to some wider population. In this chapter we begin to formalize this process. We use a statistic to estimate an unknown parameter. We use the sampling distribution to summarize the behavior of a statistic in all possible random samples of the same size from a population.

More specifically, in this chapter we begin to think about how a sample mean, \bar{x} , can provide information about a population mean, μ . When the sample mean is computed from an SRS drawn from a large population, its sampling distribution has properties that help us understand how the sample mean can be used to draw conclusions about a population mean. The law of large numbers tells us that a sample mean computed from a *random* sample from some population gets closer and closer to the population mean as the sample size increases. The central limit theorem describes the sampling distribution of the sample mean for “large” SRSs and allows us to make probability statements about possible values of the sample mean. Beginning with Chapter 14, we will develop specific methods for drawing conclusions about a population mean based on a sample mean computed from an SRS. These methods will use the tools developed in this chapter.

CHECK YOUR SKILLS

11.14 The Bureau of Labor Statistics announces that last month it interviewed all members of the labor force in a sample of 60,000 households; **9.5%** of the people interviewed were unemployed. The boldface number is a

- (a) sampling distribution.
- (b) statistic.
- (c) parameter.

11.15 A study of voting chose 663 registered Canadian voters at random shortly after the 2008 elections. Of these, 72% said they had voted in the election. Election records show that only **58.8%** of registered voters voted in the election (a record low). The boldface number is a

- (a) sampling distribution.
- (b) statistic.
- (c) parameter.

11.16 Annual returns on the more than 5000 common stocks available to investors vary a lot. In a recent year, the mean return was 8.3% and the standard deviation of returns was 28.5%. The law of large numbers says that

- (a) you can get an average return higher than the mean 8.3% by investing in a large number of stocks.
- (b) as you invest in more and more stocks chosen at random, your average return on these stocks gets ever closer to 8.3%.
- (c) if you invest in a large number of stocks chosen at random, your average return will have approximately a Normal distribution.

11.17 Scores on the Critical Reading part of the SAT exam in a recent year were roughly Normal with mean 501 and standard deviation 112. You choose an SRS of 100 students and average their SAT Critical Reading scores. If you do this many times, the mean of the average scores you get will be close to

- (a) 501.
- (b) $501/100 = 5.01$.
- (c) $501/\sqrt{100} = 50.1$.

11.18 Scores on the Critical Reading part of the SAT exam in a recent year were roughly Normal with mean 501 and standard deviation 112. You choose an SRS of 100 students and average their SAT Critical Reading scores. If you do this many times, the standard deviation of the average scores you get will be close to

- (a) 112.
- (b) $112/100 = 1.12$.
- (c) $112/\sqrt{100} = 11.2$.

11.19 A newborn baby has extremely low birth weight (ELBW) if it weighs less than 1000 grams. A study of the health of such children in later years examined a random sample of 219 children. Their mean weight at birth was $\bar{x} = 810$ grams. This sample mean is an *unbiased estimator* of the mean weight μ in the population of all ELBW babies. This means that

- (a) in many samples from this population, the mean of the many values of \bar{x} will be equal to μ .
- (b) as we take larger and larger samples from this population, \bar{x} will get closer and closer to μ .
- (c) in many samples from this population, the many values of \bar{x} will have a distribution that is close to Normal.

11.20 The number of hours a battery lasts before failing varies from battery to battery. The distribution of failure times follows an exponential distribution (see Example 11.8 page 299), which is strongly skewed to the right. The central limit theorem says that

- (a) as we look at more and more batteries, their average failure time gets ever closer to the mean μ for all batteries of this type.
- (b) the average failure time of a large number of batteries has a distribution of the same shape (strongly skewed) as the distribution for individual batteries.

(c) the average failure time of a large number of batteries has a distribution that is close to Normal.

11.21 The length of human pregnancies from conception to birth varies according to a distribution that is approximately Normal with mean 266 days and standard deviation 16 days. The probability that the average pregnancy length for 6 randomly chosen women exceeds 270 days is about

- (a) 0.40.
- (b) 0.27.
- (c) 0.07.

CHAPTER 11 EXERCISES

11.22 Testing glass. How well materials conduct heat matters when designing houses. As a test of a new measurement process, 10 measurements are made on pieces of glass known to have conductivity 1. The average of the 10 measurements is **1.07**. For each of the boldface numbers, indicate whether it is a parameter or a statistic. Explain your answer.

11.23 Statistics anxiety. What can teachers do to alleviate statistics anxiety in their students? To explore this question, statistics anxiety for students in two classes was compared. In one class, the instructor lectured in a formal manner, including dressing formally. In the other, the instructor was less formal, dressed informally, was more personal, used humor, and called on students by their first names. Anxiety was measured using a questionnaire. Higher scores indicate a greater level of anxiety. The mean anxiety score for students in the formal lecture class was **25.40**; in the informal class the mean was **20.41**. For each of the boldface numbers, indicate whether it is a parameter or a statistic. Explain your answer.

11.24 Roulette. A roulette wheel has 38 slots, of which 18 are black, 18 are red, and 2 are green. When the wheel is spun, the ball is equally likely to come to rest in any of the slots. One of the simplest wagers chooses red or black. A bet of \$1 on red returns \$2 if the ball lands in a red slot. Otherwise, the player loses his dollar. When gamblers bet on red or black, the two green slots belong to the house. Because the probability of winning \$2 is $18/38$, the mean payoff from a \$1 bet is twice $18/38$, or 94.7 cents. Explain what the law of large numbers tells us about what will happen if a gambler makes very many bets on red.

11.25 Monsoon rains. The summer monsoon rains in India follow approximately a Normal distribution with mean 852 millimeters (mm) of rainfall and standard deviation 82 mm. Rainfall is to be recorded each year for a decade and the mean rainfall \bar{x} computed. What are the mean and standard deviation of \bar{x} , the mean rainfall per year?

11.26 The Medical College Admission Test. Almost all medical schools in the United States require students to take the Medical College Admission Test (MCAT). To estimate the mean score μ of those who took the MCAT on your campus, you will obtain the scores of an SRS of students. The scores follow a Normal distribution, and from published information you know that the standard deviation is 6.4. Suppose that (unknown to you) the mean score of those taking the MCAT on your campus is 25.0.

- (a) If you choose one student at random, what is the probability that the student's score is between 20 and 30?
- (b) You sample 25 students. What is the sampling distribution of their average score \bar{x} ?
- (c) What is the probability that the mean score of your sample is between 20 and 30?

11.27 Glucose testing. Shelia's doctor is concerned that she may suffer from gestational diabetes (high blood glucose levels during pregnancy). There is variation both in the actual glucose level and in the blood test that measures the level. A patient is classified as having gestational diabetes if the glucose level is above 140 milligrams per deciliter (mg/dl) one hour after having a sugary drink. Shelia's measured glucose level one hour after the sugary drink varies according to the Normal distribution with $\mu = 122 \text{ mg/dl}$ and $\sigma = 12 \text{ mg/dl}$.

- (a) If a single glucose measurement is made, what is the probability that Shelia is diagnosed as having gestational diabetes?
- (b) If measurements are made on 4 separate days and the mean result is compared with the criterion 140 mg/dl, what is the probability that Shelia is diagnosed as having gestational diabetes?

11.28 Daily activity. It appears that people who are mildly obese are less active than leaner people. One study looked at the average number of minutes per day that people spend standing or walking.⁴ Among mildly obese people, the mean number of minutes of daily activity (standing or walking) is approximately Normally distributed with mean 373 minutes

and standard deviation 67 minutes. The mean number of minutes of daily activity for lean people is approximately Normally distributed with mean 526 minutes and standard deviation 107 minutes. A researcher records the minutes of activity for an SRS of 5 mildly obese people and an SRS of 5 lean people.

- What is the probability that the mean number of minutes of daily activity of the 5 mildly obese people exceeds 420 minutes?
- What is the probability that the mean number of minutes of daily activity of the 5 lean people exceeds 420 minutes?

11.29 Glucose testing, continued. Shelia's measured glucose level one hour after having a sugary drink varies according to the Normal distribution with $\mu = 122$ mg/dl and $\sigma = 12$ mg/dl. What is the level L such that there is probability only 0.05 that the mean glucose level of 4 test results falls above L ? (Hint: This requires a backward Normal calculation. See page 86 in Chapter 3 if you need to review.)

11.30 Pollutants in auto exhausts. Light vehicles sold in the United States must emit an average of no more than 0.07 grams per mile (g/mi) of nitrogen oxides (NOX). NOX emissions for one car model vary Normally with mean 0.05 g/mi and standard deviation 0.01 g/mi.

- What is the probability that a single car of this model emits more than 0.07 g/mi of NOX?
- A company has 25 cars of this model in its fleet. What is the probability that the average NOX level \bar{x} of these cars is above 0.07 g/mi?

11.31 Runners. In a study of exercise, a large group of male runners walk on a treadmill for 6 minutes. After this exercise, their heart rates vary with mean 8.8 beats per five seconds and standard deviation 1.0 beats per five seconds. This distribution takes only whole-number values, so it is certainly not Normal.

- Let \bar{x} be the mean number of beats per five seconds after measuring heart rate for 12 five-second intervals (a minute). What is the approximate distribution of \bar{x} according to the central limit theorem?
- What is the approximate probability that \bar{x} is less than 8?
- What is the approximate probability that the heart rate of a runner is less than 100 beats per minute? (Hint: Restate this event in terms of \bar{x} .)



Bruce Laurance/Getty Images

11.32 Pollutants in auto exhausts, continued. The level of nitrogen oxides (NOX) in the exhaust of cars of a particular model varies Normally with mean 0.05 g/mi and standard deviation 0.01 g/mi. A company has 25 cars of this model in its fleet. What is the level L such that the probability that the average NOX level \bar{x} for the fleet is greater than L is only

0.01? (Hint: This requires a backward Normal calculation. See page 86 in Chapter 3 if you need to review.)

11.33 Returns on stocks. Andrew plans to retire in 40 years. He plans to invest part of his retirement funds in stocks, so he seeks out information on past returns. He learns that from 1960 to 2009, the annual returns on U.S. common stocks had mean 10.8% and standard deviation 17.1%.⁵ The distribution of annual returns on common stocks is roughly symmetric, so the mean return over even a moderate number of years is close to Normal. What is the probability (assuming that the past pattern of variation continues) that the mean annual return on common stocks over the next 40 years will exceed 10%? What is the probability that the mean return will be less than 5%? Follow the four-step process as illustrated in Example 11.8.

11.34 Airline passengers get heavier. In response to the increasing weight of airline passengers, the Federal Aviation Administration (FAA) in 2003 told airlines to assume that passengers average 190 pounds in the summer, including clothing and carry-on baggage. But passengers vary, and the FAA did not specify a standard deviation. A reasonable standard deviation is 35 pounds. Weights are not Normally distributed, especially when the population includes both men and women, but they are not very non-Normal. A commuter plane carries 22 passengers. What is the approximate probability that the total weight of the passengers exceeds 4500 pounds? Use the four-step process to guide your work. (Hint: To apply the central limit theorem, restate the problem in terms of the mean weight.)



Jeff Greenberg/The Image Works

11.35 Sampling students. To estimate the mean score μ of those who took the Medical College Admission Test on your campus, you will obtain the scores of an SRS of students. From published information you know that the scores are approximately Normal with standard deviation about 6.4. How large an SRS must you take to reduce the standard deviation of the sample mean score to 1?

11.36 Sampling students, continued. To estimate the mean score μ of those who took the Medical College Admission Test on your campus, you will obtain the scores of an SRS of students. From published information you know that the scores are approximately Normal with standard deviation about 6.4. You want your sample mean \bar{x} to estimate μ with an error of no more than 1 point in either direction.

- What standard deviation must \bar{x} have so that 99.7% of all samples give an \bar{x} within 1 point of μ ? (Use the 68–95–99.7 rule.)
- How large an SRS do you need in order to reduce the standard deviation of \bar{x} to the value you found in part (a)?

11.37 Playing the numbers. The numbers racket is a well-entrenched illegal gambling operation in most large cities. One version works as follows: you choose one of the 1000 three-digit numbers 000 to 999 and pay your local numbers runner a dollar to enter your bet. Each day, one three-digit number is chosen at random and pays off \$600. The mean payoff for the population of thousands of bets is $\mu = 60$ cents. Joe makes one bet every day for many years. Explain what the law of large numbers says about Joe's results as he keeps on betting.

11.38 Playing the numbers: a gambler gets chance outcomes. The law of large numbers tells us what happens in the long run. Like many games of chance, the numbers racket has outcomes so variable—one three-digit number wins \$600 and all others win nothing—that gamblers never reach “the long run.” Even after many bets, their average winnings may not be close to the mean. For the numbers racket, the mean payout for single bets is \$0.60 (60 cents) and the standard deviation of payouts is about \$18.96. If Joe plays 350 days a year for 40 years, he makes 14,000 bets.

- (a) What are the mean and standard deviation of the average payout \bar{x} that Joe receives from his 14,000 bets?
- (b) The central limit theorem says that his average payout is approximately Normal with the mean and standard deviation you found in part (a). What is the approximate probability that Joe's average payout per bet is between \$0.50 and \$0.70? You see that Joe's average may not be very close to the mean \$0.60 even after 14,000 bets.

11.39 Playing the numbers: the house has a business. Unlike Joe (see the previous exercise) the operators of the numbers racket can rely on the law of large numbers. It is said that the New York City mobster Casper Holstein took as many as 25,000 bets per day in the Prohibition era. That's 150,000 bets in a week if he takes Sunday off. Casper's mean winnings per bet are \$0.40 (he pays out 60 cents of each dollar bet to people like Joe and keeps the other 40 cents.) His standard deviation for single bets is about \$18.96, the same as Joe's.

- (a) What are the mean and standard deviation of Casper's average winnings \bar{x} on his 150,000 bets?

(b) According to the central limit theorem, what is the approximate probability that Casper's average winnings per bet are between \$0.30 and \$0.50? After only a week, Casper can be pretty confident that his winnings will be quite close to \$0.40 per bet.

11.40 Can we trust the central limit theorem? The central limit theorem says that “when n is large” we can act as if the distribution of a sample mean \bar{x} is close to Normal. How large a sample we need depends on how far the population distribution is from being Normal. Example 11.8 shows that we can trust this Normal approximation for quite moderate sample sizes even when the population has a strongly skewed continuous distribution.

The central limit theorem requires much larger samples for Joe's bets with his local numbers racket. The population of individual bets has a finite distribution with only two possible outcomes: \$600 (probability 0.001) and \$0 (probability 0.999). This distribution has mean $\mu = 0.6$ and standard deviation about $\sigma = 18.96$. With more math and good software we can find exact probabilities for Joe's average winnings.

- (a) If Joe makes 14,000 bets, the exact probability $P(0.5 \leq \bar{x} \leq 0.7) = 0.4674$. How accurate was your Normal approximation from part (b) of Exercise 11.38?
- (b) If Joe makes only 3500 bets, $P(0.5 \leq \bar{x} \leq 0.7) = 0.2450$. How accurate is the Normal approximation for this probability?
- (c) If Joe and his buddies make 150,000 bets, $P(0.5 \leq \bar{x} \leq 0.7) = 0.9589$. How accurate is the Normal approximation?

11.41 What's the mean? Suppose that you roll three balanced dice. We wonder what the mean number of spots on the up-faces of the three dice is. The law of large numbers says that we can find out by experience: roll three dice many times, and the average number of spots will eventually approach the true mean. Set up the *Law of Large Numbers* applet to roll three dice. Don't click “Show mean” yet. Roll the dice until you are confident you know the mean quite closely, then click “Show mean” to verify your discovery. What is the mean? Make a rough sketch of the path the averages \bar{x} followed as you kept adding more rolls.

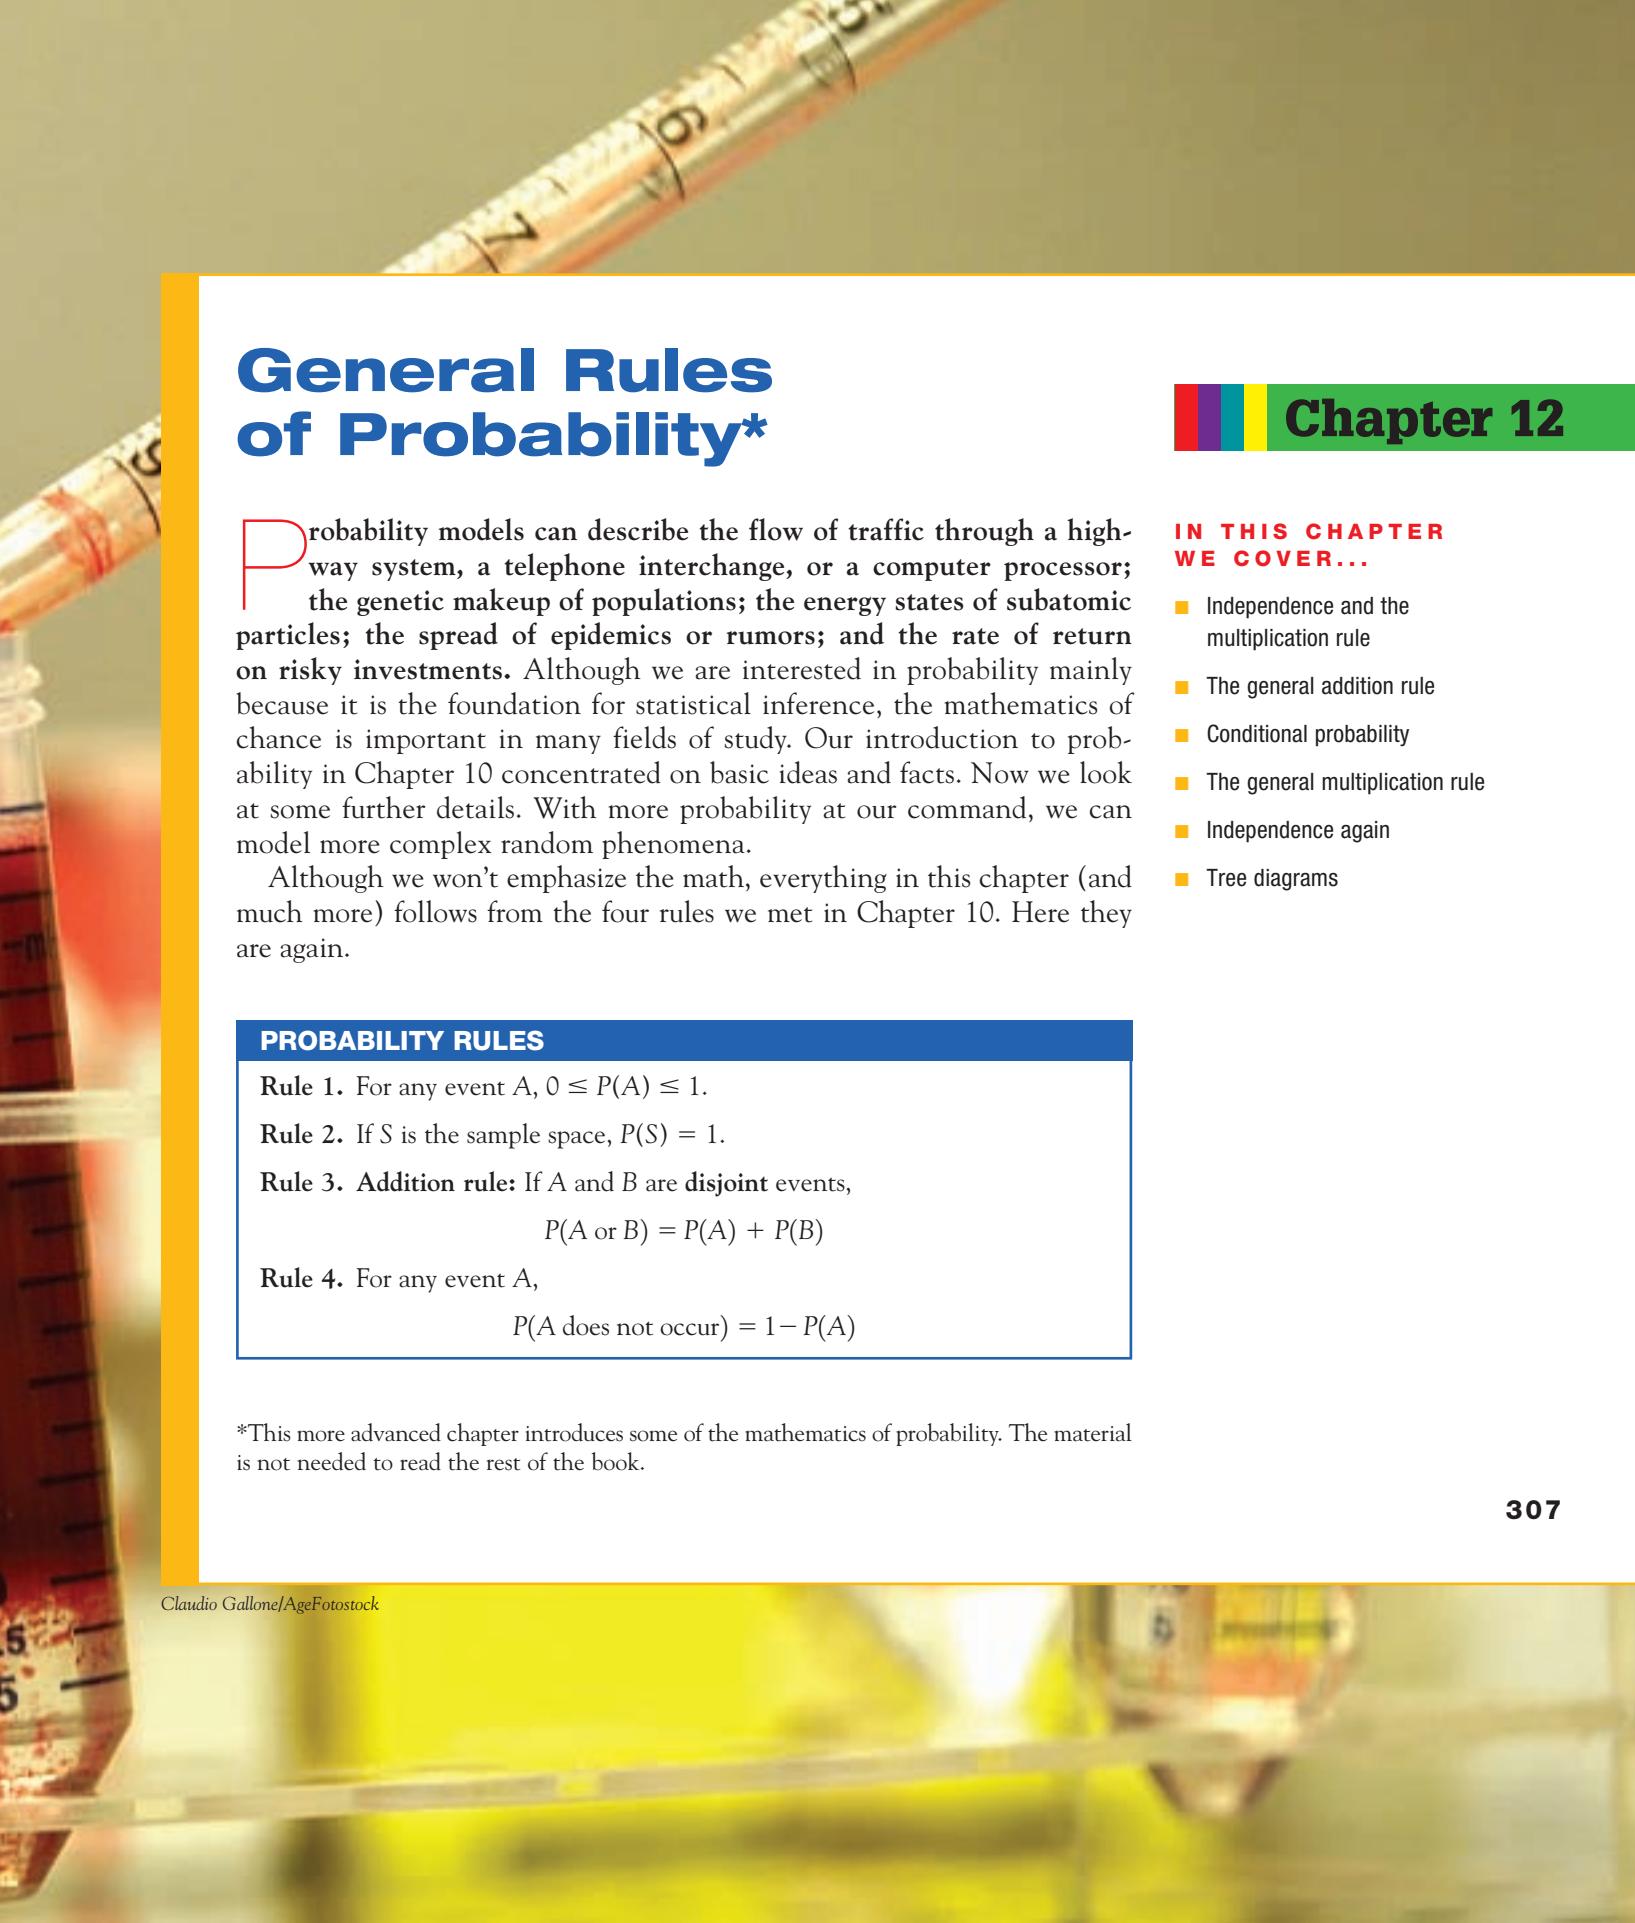


EXPLORING THE WEB

11.42 Online videos. There are several online videos of the law of large numbers and central limit theorem. Locate one such video, watch it, and write a brief summary of the video. We did a Google search of “videos for law of large numbers” and “videos for the central limit theorem” and found several links.

11.43 Work the law of large numbers. Read the online article at ezinearticles.com/?Work-The-Law-of-Large-Numbers-But-Remember-It-Only-Takes-One-to-Succeed!&id=932026. Does this article accurately describe the law of large numbers? Explain your answer.





General Rules of Probability*

Probability models can describe the flow of traffic through a highway system, a telephone interchange, or a computer processor; the genetic makeup of populations; the energy states of subatomic particles; the spread of epidemics or rumors; and the rate of return on risky investments. Although we are interested in probability mainly because it is the foundation for statistical inference, the mathematics of chance is important in many fields of study. Our introduction to probability in Chapter 10 concentrated on basic ideas and facts. Now we look at some further details. With more probability at our command, we can model more complex random phenomena.

Although we won't emphasize the math, everything in this chapter (and much more) follows from the four rules we met in Chapter 10. Here they are again.

PROBABILITY RULES

Rule 1. For any event A , $0 \leq P(A) \leq 1$.

Rule 2. If S is the sample space, $P(S) = 1$.

Rule 3. Addition rule: If A and B are disjoint events,

$$P(A \text{ or } B) = P(A) + P(B)$$

Rule 4. For any event A ,

$$P(A \text{ does not occur}) = 1 - P(A)$$



IN THIS CHAPTER WE COVER...

- Independence and the multiplication rule
- The general addition rule
- Conditional probability
- The general multiplication rule
- Independence again
- Tree diagrams

*This more advanced chapter introduces some of the mathematics of probability. The material is not needed to read the rest of the book.

INDEPENDENCE AND THE MULTIPLICATION RULE

Rule 3, the addition rule for disjoint events, describes the probability that *one or the other* of two events A and B occurs in the special situation when A and B cannot occur together. Now we will describe the probability that *both* events A and B occur, again only in a special situation.

Venn diagram

You may find it helpful to draw a picture to display relations among several events. A picture like Figure 12.1 that shows the sample space S as a rectangular area and events as areas within S is called a **Venn diagram**. The events A and B in Figure 12.1 are disjoint because they do not overlap. The Venn diagram in Figure 12.2 illustrates two events that are not disjoint. The event {A and B} appears as the overlapping area that is common to both A and B. Can we find the probability $P(A \text{ and } B)$ that both events occur if we know the individual probabilities $P(A)$ and $P(B)$?

FIGURE 12.1

Venn diagram showing disjoint events A and B.

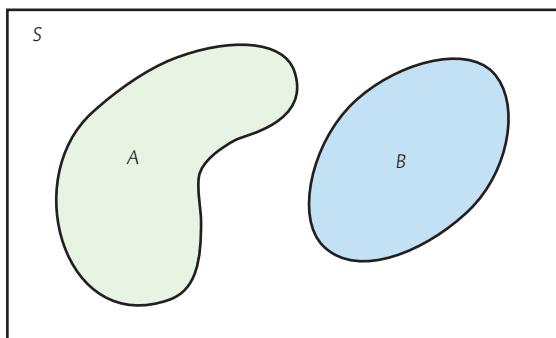
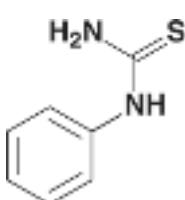
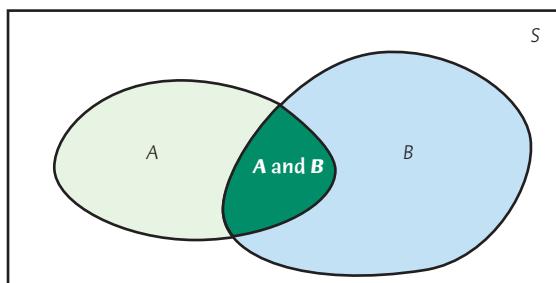


FIGURE 12.2

Venn diagram showing events A and B that are not disjoint. The event {A and B} consists of outcomes common to A and B.



EXAMPLE 12.1 Can you taste PTC?

That molecule in the diagram to the left is PTC, a substance with an unusual property: 70% of people find that it has a bitter taste and the other 30% can't taste it at all. The difference is genetic, depending on a single gene. Ask two people chosen at random to taste PTC. We are interested in the events

$$A = \{\text{first person can taste PTC}\}$$

$$B = \{\text{second person can taste PTC}\}$$

We know that $P(A) = 0.7$ and $P(B) = 0.7$. What is the probability $P(A \text{ and } B)$ that both can taste PTC?

We can think our way to the answer. The first person chosen can taste PTC in 70% of all samples and then the second person can taste it in 70% of those samples. We will get two tasters in 70% of 70% of all samples. That's $P(A \text{ and } B) = 0.7 \times 0.7 = 0.49$. ■

The argument in Example 12.1 works because knowing that the first person can taste PTC tells us nothing about the second person. The probability is still 0.7 that the second person can taste PTC whether or not the first person can. We say that the events "first person can taste PTC" and "second person can taste PTC" are **independent**. Now we have another rule of probability.

independent events

MULTIPLICATION RULE FOR INDEPENDENT EVENTS

Two events A and B are **independent** if knowing that one occurs does not change the probability that the other occurs. If A and B are independent,

$$P(A \text{ and } B) = P(A)P(B)$$

EXAMPLE 12.2 Independent or not?

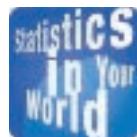
To use this multiplication rule, we must decide whether events are independent. Here are some examples to help you recognize when you can assume that events are independent.

In Example 12.1, we think that the ability of one randomly chosen person to taste PTC tells us nothing about whether or not a second person, also randomly chosen, can taste PTC. That's independence. But if the two people are members of the same family, the fact that ability to taste PTC is inherited warns us that they are not independent.

Independence is clearly recognized in artificial settings such as games of chance. Because a coin has no memory and most coin tossers cannot influence the fall of the coin, it is safe to assume that successive coin tosses are independent, so that the probability of three heads in succession is $0.5 \times 0.5 \times 0.5 = 0.125$.

On the other hand, the colors of successive cards dealt from the same deck are not independent. A standard 52-card deck contains 26 red and 26 black cards. For the first card dealt from a shuffled deck, the probability of a red card is $26/52 = 0.50$. Once we see that the first card is red, we know that there are only 25 reds among the remaining 51 cards. The probability that the second card is red is therefore only $25/51 = 0.49$. Knowing the outcome of the first deal changes the probabilities for the second. ■

The multiplication rule extends to collections of more than two events, provided that all are independent. Independence of events A, B, and C means that no information about any one or any two can change the probability of the remaining events. Independence is often assumed in setting up a probability model when the events we are describing seem to have no connection.



Condemned by independence

Assuming independence when it isn't true can

lead to disaster. Several mothers in England were convicted of murder simply because two of their children had died in their cribs with no visible cause. An "expert witness" for the prosecution said that the probability of an unexplained crib death in a nonsmoking middle-class family is $1/8500$. He then multiplied $1/8500$ by $1/8500$ to claim that there is only a 1 in 73 million chance that two children in the same family could have died naturally. This is nonsense: it assumes that crib deaths are independent, and data suggest that they are not. Some common genetic or environmental cause, not murder, probably explains the deaths.

If two events A and B are independent, the event that A does not occur is also independent of B, and so on. For example, choose two people at random and ask if they can taste PTC. Because 70% can taste PTC and 30% cannot, the probability that the first person is a taster and the second is not is $(0.7)(0.3) = 0.21$.

EXAMPLE 12.3 Surviving?

During World War II, the British found that the probability that a bomber is lost through enemy action on a mission over occupied Europe was 0.05. The probability that the bomber returns safely from a mission was therefore 0.95. It is reasonable to assume that missions are independent. Take A_i to be the event that a bomber survives its i th mission. The probability of surviving 2 missions is

$$\begin{aligned}P(A_1 \text{ and } A_2) &= P(A_1)P(A_2) \\&= (0.95)(0.95) = 0.9025\end{aligned}$$

The multiplication rule also applies to more than two independent events, so the probability of surviving 3 missions is

$$\begin{aligned}P(A_1 \text{ and } A_2 \text{ and } A_3) &= P(A_1)P(A_2)P(A_3) \\&= (0.95)(0.95)(0.95) = 0.8574\end{aligned}$$

In 1941, the tour of duty for an airman was established as 30 missions. The probability of surviving 30 missions is only

$$\begin{aligned}P(A_1 \text{ and } A_2 \text{ and } \dots \text{ and } A_{30}) &= P(A_1)P(A_2) \cdots P(A_{30}) \\&= (0.95)(0.95) \cdots (0.95) \\&= (0.95)^{30} = 0.2146\end{aligned}$$

The probability of surviving two tours of duty was much smaller. ■

Here is another example of using the multiplication rule for independent events to compute probabilities.

EXAMPLE 12.4 Rapid HIV testing



STATE: Many people who come to clinics to be tested for HIV, the virus that causes AIDS, don't come back to learn the test results. Clinics now use "rapid HIV tests" that give a result while the client waits. In a clinic in Malawi, for example, use of rapid tests increased the percent of clients who learned their test results from 69% to 99.7%.

The trade-off for fast results is that rapid tests are less accurate than slower laboratory tests. Applied to people who have no HIV antibodies, one rapid test has probability about 0.004 of producing a false-positive (that is, of falsely indicating that antibodies are present).¹ If a clinic tests 200 people who are free of HIV antibodies, what is the chance that at least one false-positive will occur?

PLAN: It is reasonable to assume that the test results for different individuals are independent. We have 200 independent events, each with probability 0.004. What is the probability that at least one of these events occurs?

SOLVE: “At least one” combines many outcomes. It is much easier to use the fact that

$$P(\text{at least one positive}) = 1 - P(\text{no positives})$$

and find $P(\text{no positives})$ first.

The probability of a negative result for any one person is $1 - 0.004 = 0.996$. To find the probability that all 200 people tested have negative results, use the multiplication rule:

$$\begin{aligned} P(\text{no positives}) &= P(\text{all 200 negative}) \\ &= (0.996)(0.996) \cdots (0.996) \\ &= 0.996^{200} = 0.4486 \end{aligned}$$

The probability we want is therefore

$$P(\text{at least one positive}) = 1 - 0.4486 = 0.5514$$

CONCLUDE: The probability is greater than 1/2 that at least one of the 200 people will test positive for HIV even though no one has the virus. ■

The multiplication rule $P(A \text{ and } B) = P(A)P(B)$ holds if A and B are independent but not otherwise. The addition rule $P(A \text{ or } B) = P(A) + P(B)$ holds if A and B are disjoint but not otherwise. Resist the temptation to use these simple rules when the circumstances that justify them are not present. You must also be careful not to confuse disjointness and independence. If A and B are disjoint, then the fact that A occurs tells us that B cannot occur—look again at Figure 12.1. So disjoint events are not independent. Unlike disjointness, we cannot picture independence in a Venn diagram, because it involves the probabilities of the events rather than just the outcomes that make up the events.



APPLY YOUR KNOWLEDGE

12.1 Older college students. Government data show that 8% of adults are full-time college students and that 30% of adults are age 55 or older. Nonetheless, we can't conclude that, because $(0.08)(0.30) = 0.024$, about 2.4% of adults are college students 55 or older. Why not?

12.2 Common names. The U.S. Census Bureau says that the 10 most common names in the United States are (in order) Smith, Johnson, Williams, Brown, Jones, Miller, Davis, Garcia, Rodriguez, and Wilson. These names account for 9.6% of all U.S. residents. Out of curiosity, you look at the authors of the textbooks for your current courses. There are 9 authors in all. Would you be surprised if none of the names of these authors were among the 10 most common? (Assume that authors' names are independent and follow the same probability distribution as the names of all residents.)

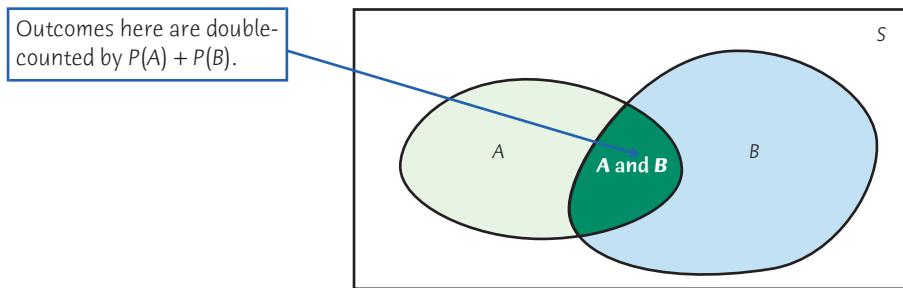
12.3 Lost Internet sites. Internet sites often vanish or move, so that references to them can't be followed. In fact, 13% of Internet sites referenced in major scientific journals are lost within two years after publication.² If a paper contains seven Internet references, what is the probability that all seven are still good two years later? What specific assumptions did you make in order to calculate this probability?

THE GENERAL ADDITION RULE

We know that if A and B are disjoint events, then $P(A \text{ or } B) = P(A) + P(B)$. If events A and B are *not* disjoint, they can occur together. The probability that one or the other occurs is then *less* than the sum of their probabilities. As Figure 12.3 illustrates, outcomes common to both are counted twice when we add probabilities, so we must subtract this probability once. Here is the addition rule for any two events, disjoint or not.

FIGURE 12.3

The general addition rule: for any events A and B, $P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$.



ADDITION RULE FOR ANY TWO EVENTS

For any two events A and B,

$$P(A \text{ or } B) = P(A) + P(B) - P(A \text{ and } B)$$

If A and B are disjoint, the event {A and B} that both occur contains no outcomes and therefore has probability 0. So the general addition rule includes Rule 3, the addition rule for disjoint events.



Justin Sullivan/Getty Images

EXAMPLE 12.5 Motor vehicle sales

Motor vehicles sold in the United States (ignoring heavy trucks) are classified as either cars or light trucks and as either domestic or imported. “Light trucks” include SUVs and minivans. “Domestic” means made in Canada, Mexico, or the United States, so that a Toyota made in Canada counts as domestic.

In 2010, 76% of the new vehicles sold to individuals were domestic, 50% were light trucks, and 43% were domestic light trucks.³ Choose a vehicle sale at random. Then

$$\begin{aligned} P(\text{domestic or light truck}) &= P(\text{domestic}) + P(\text{light truck}) - P(\text{domestic light truck}) \\ &= 0.76 + 0.50 - 0.43 = 0.83 \end{aligned}$$

That is, 83% of vehicles sold were either domestic or light trucks. A vehicle is an imported car if it is *neither* domestic *nor* a light truck. So

$$P(\text{imported car}) = 1 - 0.83 = 0.17 \blacksquare$$

Venn diagrams clarify events and their probabilities because you can just think of adding and subtracting areas. Figure 12.4 shows all the events formed from “domestic” and “truck” in Example 12.5. The four probabilities that appear in the figure add to 1 because they refer to four disjoint events that make up the entire sample space. All these probabilities come from the information in Example 12.5. For example, the probability that a randomly chosen vehicle sale is a domestic car (“D and not T” in the figure) is

$$\begin{aligned} P(\text{domestic car}) &= P(\text{domestic}) - P(\text{domestic light truck}) \\ &= 0.76 - 0.43 = 0.33 \end{aligned}$$

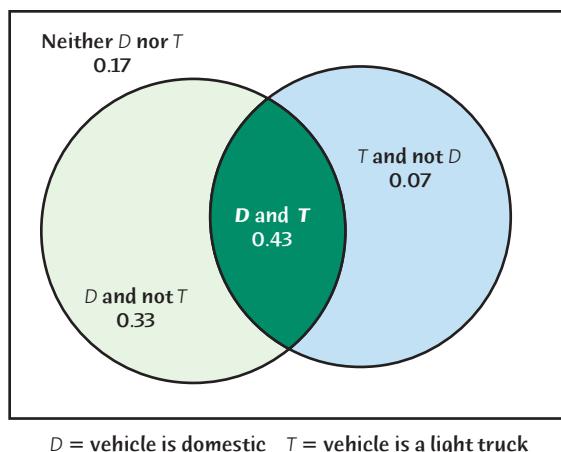


FIGURE 12.4

Venn diagram and probabilities for motor vehicle sales, for Example 12.5.

APPLY YOUR KNOWLEDGE

12.4 College degrees. Of all college degrees awarded in the United States, 50% are bachelor's degrees, 59% are earned by women, and 29% are bachelor's degrees earned by women. Make a Venn diagram and use it to answer these questions.

- What percent of all degrees are earned by men?
- What percent of all degrees are bachelor's degrees earned by men?
- Are the events earning a bachelor's degree and being a man independent? Why?

12.5 Distance learning. A study of the students taking distance-learning courses at a university finds that they are mostly older students not living in the university town. Choose a distance-learning student at random. Let A be the event that the student is 25 years old or older and B the event that the student is local. The study finds that $P(A) = 0.7$, $P(B) = 0.25$, and $P(A \text{ and } B) = 0.05$.

- Make a Venn diagram similar to Figure 12.4 showing the events $\{A \text{ and } B\}$, $\{A \text{ and not } B\}$, $\{B \text{ and not } A\}$, and $\{\text{neither } A \text{ nor } B\}$.
- Describe each of these events in words.
- Find the probabilities of all four events and add the probabilities to your Venn diagram.



Barry Austin Photography/Getty

CONDITIONAL PROBABILITY

The probability we assign to an event can change if we know that some other event has occurred. This idea is the key to many applications of probability.

EXAMPLE 12.6 Trucks among imported motor vehicles

Figure 12.4, based on the information in Example 12.5, gives the following probabilities for a randomly chosen light motor vehicle sold at retail in the United States:

	Domestic	Imported	Total
Light truck	0.43	0.07	0.50
Car	0.33	0.17	0.50
Total	0.76	0.24	1

The four probabilities in the body of the table add to 1 because they describe all vehicles sold. We obtain the “Total” row and column from these probabilities by the addition rule. For example, the probability that a randomly chosen vehicle is a light truck is

$$\begin{aligned} P(\text{truck}) &= P(\text{truck and domestic}) + P(\text{truck and imported}) \\ &= 0.43 + 0.07 = 0.50 \end{aligned}$$

Suppose we are told that the vehicle chosen is imported. That is, it is one of the 24% in the “Imported” column of the table. The probability that a vehicle is a light truck *given the information that it is imported* is the proportion of trucks in the “Imported” column,

$$P(\text{truck} | \text{imported}) = \frac{0.07}{0.24} = 0.29$$

conditional probability

This is a **conditional probability**. You can read the bar | as “given the information that.” ■

Although 50% of all vehicles sold are trucks, only 29% of imported vehicles are trucks. It’s common sense that knowing that one event (the vehicle is imported) occurs often changes the probability of another event (the vehicle is a truck). The example also shows how we should define conditional probability. The idea of a conditional probability $P(B | A)$ of one event B given that another event A occurs is the proportion of *all occurrences of A* for which B also occurs.

CONDITIONAL PROBABILITY

When $P(A) > 0$, the **conditional probability** of B given A is

$$P(B | A) = \frac{P(A \text{ and } B)}{P(A)}$$

The conditional probability $P(B | A)$ makes no sense if the event A can never occur, so we require that $P(A) > 0$ whenever we talk about $P(B | A)$. Be sure to keep in mind the distinct roles of the events A and B in $P(B | A)$. Event A represents the information we are given, and B is the event whose probability we are calculating. Here is an example that emphasizes this distinction.



EXAMPLE 12.7 Imports among trucks

What is the conditional probability that a randomly chosen vehicle is imported *given the information that it is a truck?* Using the definition of conditional probability,

$$\begin{aligned} P(\text{imported} | \text{truck}) &= \frac{P(\text{imported and truck})}{P(\text{truck})} \\ &= \frac{0.07}{0.50} = 0.14 \end{aligned}$$

Only 14% of trucks sold are imports. ■

Be careful not to confuse the two different conditional probabilities

$$P(\text{truck} | \text{imported}) = 0.29$$

$$P(\text{imported} | \text{truck}) = 0.14$$

The first answers the question “What proportion of imports are trucks?” The second answers “What proportion of trucks are imports?”

APPLY YOUR KNOWLEDGE

12.6 College degrees. In the setting of Exercise 12.4, what is the conditional probability that a degree is earned by a woman given that it is a bachelor’s degree?

12.7 Distance learning. In the setting of Exercise 12.5, what is the conditional probability that a student is local given that he or she is less than 25 years old?

12.8 Computer games. Here is the distribution of computer games sold by type of game:⁴

Game type	Probability
Strategy	0.354
Role playing	0.139
Family entertainment	0.127
Shooters	0.109
Children’s	0.057
Other	0.214

What is the conditional probability that a computer game is a role-playing game given that it is not a strategy game?



Winning the lottery twice

In 1986, Evelyn Marie Adams won the New Jersey lottery for the second time, adding \$1.5 million to her previous \$3.9 million jackpot. The *New York Times* claimed that the odds of one person winning the big prize twice were 1 in 17 trillion. Nonsense, said two statisticians in a letter to the *Times*. The chance that Adams would win twice is indeed tiny, but it is almost certain that *someone* among the millions of lottery players would win two jackpots. There have been many multiple winners since that time. Ernest Pullen of St. Louis won the Missouri lottery in June 2010 and then won again in September 2010 (\$3 million total). When commenting on the double win, Pullen said that he considers himself to be a “lucky guy.”

THE GENERAL MULTIPLICATION RULE

The definition of conditional probability reminds us that in principle all probabilities, including conditional probabilities, can be found from the assignment of probabilities to events that describe a random phenomenon. More often, however, conditional probabilities are part of the information given to us in a probability model. The definition of conditional probability then turns into a rule for finding the probability that both of two events occur.

MULTIPLICATION RULE FOR ANY TWO EVENTS

The probability that both of two events A and B happen together can be found by

$$P(A \text{ and } B) = P(A)P(B | A)$$

Here $P(B | A)$ is the conditional probability that B occurs given the information that A occurs.

In words, this rule says that for both of two events to occur, first one must occur and then, given that the first event has occurred, the second must occur. This is just common sense expressed in the language of probability, as the following example illustrates.

EXAMPLE 12.8 Teens with online profiles

The Pew Internet and American Life Project finds that 93% of teenagers (ages 12 to 17) use the Internet, and that 55% of online teens have posted a profile on a social-networking site.⁵ What percent of teens are online *and* have posted a profile?

Use the multiplication rule:

$$P(\text{online}) = 0.93$$

$$P(\text{profile} | \text{online}) = 0.55$$

$$\begin{aligned} P(\text{online and have profile}) &= P(\text{online}) \times P(\text{profile} | \text{online}) \\ &= (0.93)(0.55) = 0.5115 \end{aligned}$$

That is, about 51% of all teens use the Internet and have a profile on a social-networking site.

You should think your way through this: if 93% of teens are online and 55% of these have posted a profile, then 55% of 93% are both online and have a profile. ■

We can extend the multiplication rule to find the probability that all of several events occur. The key is to condition each event on the occurrence of *all* of the preceding events. So for any three events A, B, and C,

$$P(A \text{ and } B \text{ and } C) = P(A)P(B | A)P(C | \text{both } A \text{ and } B)$$

Here is an example of the extended multiplication rule.

EXAMPLE 12.9 Fundraising by telephone

STATE: A charity raises funds by calling a list of prospective donors to ask for pledges. It is able to talk with 40% of the names on its list. Of those the charity reaches, 30% make a pledge. But only half of those who pledge actually make a contribution. What percent of the donor list contributes?



PLAN: Express the information we are given in terms of events and their probabilities:

$$\begin{aligned} \text{If } A &= \{\text{the charity reaches a prospect}\} & \text{then } P(A) = 0.4 \\ \text{If } B &= \{\text{the prospect makes a pledge}\} & \text{then } P(B|A) = 0.3 \\ \text{If } C &= \{\text{the prospect makes a contribution}\} & \text{then } P(C|\text{both } A \text{ and } B) = 0.5 \end{aligned}$$

We want to find $P(A \text{ and } B \text{ and } C)$.

SOLVE: Use the multiplication rule:

$$\begin{aligned} P(A \text{ and } B \text{ and } C) &= P(A)P(B|A)P(C|\text{both } A \text{ and } B) \\ &= 0.4 \times 0.3 \times 0.5 = 0.06 \end{aligned}$$

CONCLUDE: Only 6% of the prospective donors make a contribution. ■

As Example 12.9 illustrates, formulating a problem in the language of probability is often the key to success in applying probability ideas.

APPLY YOUR KNOWLEDGE

- 12.9 At the gym.** Suppose that 10% of adults belong to health clubs, and 40% of these health club members go to the club at least twice a week. What percent of all adults go to a health club at least twice a week? Write the information given in terms of probabilities and use the general multiplication rule.

- 12.10 Teens online.** We saw in Example 12.8 that 93% of teenagers are online and that 55% of online teens have posted a profile on a social-networking site. Of online teens with a profile, 76% have placed comments on a friend's blog. What percent of all teens are online, have a profile, and comment on a friend's blog? Define events and probabilities and follow the pattern of Example 12.9.



- 12.11 The probability of a flush.** A poker player holds a flush when all 5 cards in the hand belong to the same suit (clubs, diamonds, hearts, or spades). We will find the probability of a flush when 5 cards are dealt. Remember that a deck contains 52 cards, 13 of each suit, and that when the deck is well shuffled, each card dealt is equally likely to be any of those that remain in the deck.

- Concentrate on spades. What is the probability that the first card dealt is a spade? What is the conditional probability that the second card is a spade given that the first is a spade? (Hint: How many cards remain? How many of these are spades?)
- Continue to count the remaining cards to find the conditional probabilities of a spade on the third, the fourth, and the fifth card given in each case that all previous cards are spades.



The Photo Works

- (c) The probability of being dealt 5 spades is the product of the 5 probabilities you have found. Why? What is this probability?
 - (d) The probability of being dealt 5 hearts or 5 diamonds or 5 clubs is the same as the probability of being dealt 5 spades. What is the probability of being dealt a flush?
-

INDEPENDENCE AGAIN

The conditional probability $P(B | A)$ is generally not equal to the unconditional probability $P(B)$. That's because the occurrence of event A generally gives us some additional information about whether or not event B occurs. If knowing that A occurs gives no additional information about B , then A and B are independent events. The precise definition of independence is expressed in terms of conditional probability.

INDEPENDENT EVENTS

Two events A and B that both have positive probability are **independent** if

$$P(B | A) = P(B)$$

We now see that the multiplication rule for independent events, $P(A \text{ and } B) = P(A)P(B)$, is a special case of the general multiplication rule, $P(A \text{ and } B) = P(A)P(B | A)$, just as the addition rule for disjoint events is a special case of the general addition rule. We rarely use the definition of independence because most often independence is part of the information given to us in a probability model.

APPLY YOUR KNOWLEDGE

12.12 Independent? The Clemson University Fact Book for 2007 shows that 123 of the university's 338 assistant professors were women, along with 76 of the 263 associate professors and 73 of the 375 full professors.

- (a) What is the probability that a randomly chosen Clemson professor is a woman?
 - (b) What is the conditional probability that a randomly chosen professor is a woman, given that the person chosen is a full professor?
 - (c) Are the rank and sex of Clemson professors independent? How do you know?
-

TREE DIAGRAMS

Probability models often have several stages, with probabilities at each stage conditional on the outcomes of earlier states. These models require us to combine several of the basic rules into a more elaborate calculation. Here is an example.

EXAMPLE 12.10 Who visits YouTube?

STATE: Video-sharing sites, led by YouTube, are popular destinations on the Internet. Let's look only at adult Internet users, aged 18 and over. About 27% of adult Internet users are 18 to 29 years old, another 45% are 30 to 49 years old, and the remaining 28% are 50 and over. The Pew Internet and American Life Project finds that 70% of Internet users aged 18 to 29 have visited a video-sharing site, along with 51% of those aged 30 to 49 and 26% of those 50 or older. What percent of all adult Internet users visit video-sharing sites?



PLAN: To use the tools of probability, restate all these percents as probabilities. If we choose an online adult at random,

$$P(\text{age 18 to 29}) = 0.27$$

$$P(\text{age 30 to 49}) = 0.45$$

$$P(\text{age 50 and older}) = 0.28$$

These three probabilities add to 1 because all adult Internet users are in one of the three age groups. The percents of each group who visit video-sharing sites are *conditional* probabilities:

$$P(\text{video yes} \mid \text{age 18 to 29}) = 0.70$$

$$P(\text{video yes} \mid \text{age 30 to 49}) = 0.51$$

$$P(\text{video yes} \mid \text{age 50 and older}) = 0.26$$

We want to find the unconditional probability $P(\text{video yes})$.

SOLVE: The **tree diagram** in Figure 12.5 organizes this information. Each segment in the tree is one stage of the problem. Each complete branch shows a path through the two stages. The probability written on each segment is the conditional probability of an Internet user following that segment given that he or she has reached the node from which it branches.

tree diagram

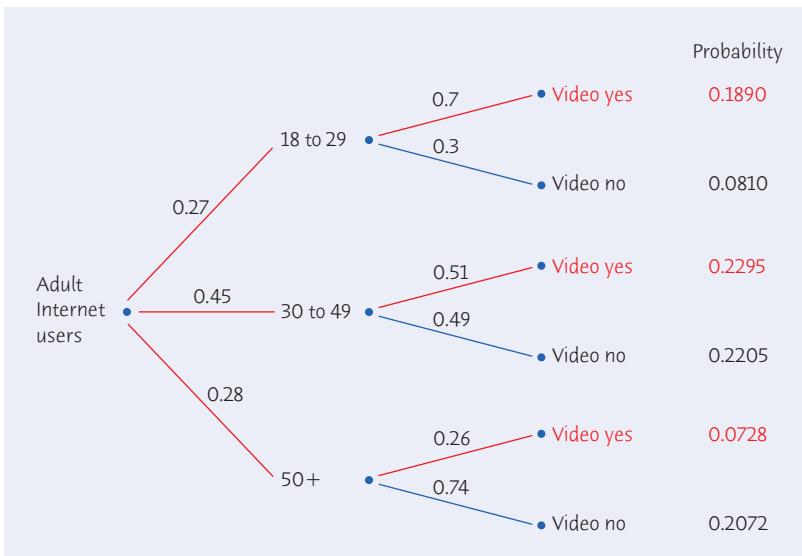


FIGURE 12.5

Tree diagram for use of the Internet and video-sharing sites such as YouTube, for Example 12.10. The three disjoint paths to the outcome that an adult Internet user visits video-sharing sites are colored red.

Starting at the left, an Internet user falls into one of the three age groups. The probabilities of these groups mark the leftmost segments in the tree. Look at age 18 to 29, the top branch. The two segments going out from the “18 to 29” branch point carry the conditional probabilities

$$P(\text{video yes} \mid \text{age 18 to 29}) = 0.70$$

$$P(\text{video no} \mid \text{age 18 to 29}) = 0.30$$

The full tree shows the probabilities for all three age groups.

Now use the multiplication rule. The probability that a randomly chosen Internet user is an 18- to 29-year-old who visits video-sharing sites is

$$\begin{aligned} P(18 \text{ to } 29 \text{ and video yes}) &= P(18 \text{ to } 29)P(\text{video yes} \mid 18 \text{ to } 29) \\ &= (0.27)(0.70) = 0.1890 \end{aligned}$$

This probability appears at the end of the topmost branch. The multiplication rule says that the probability of any complete branch in the tree is the product of the probabilities of the segments in that branch.

There are three disjoint paths to “video yes,” one for each of the three age groups. These paths are colored red in Figure 12.5. Because the three paths are disjoint, the probability that an adult Internet user visits video-sharing sites is the sum of their probabilities:

$$\begin{aligned} P(\text{video yes}) &= (0.27)(0.70) + (0.45)(0.51) + (0.28)(0.26) \\ &= 0.1890 + 0.2295 + 0.0728 = 0.4913 \end{aligned}$$

CONCLUDE: About 49% of all adult Internet users have visited a video-sharing site. ■

It takes longer to explain a tree diagram than it does to use it. Once you have understood a problem well enough to draw the tree, the rest is easy. Here is another question about video-sharing sites that the tree diagram helps us answer.



EXAMPLE 12.11 Young adults at video-sharing sites

STATE: What percent of adult Internet users who visit video-sharing sites are age 18 to 29?

PLAN: In probability language, we want the conditional probability $P(18 \text{ to } 29 \mid \text{video yes})$. Use the tree diagram and the definition of conditional probability:

$$P(18 \text{ to } 29 \mid \text{video yes}) = \frac{P(18 \text{ to } 29 \text{ and video yes})}{P(\text{video yes})}$$

SOLVE: Look again at the tree diagram in Figure 12.5. $P(\text{video yes})$ is the sum of the three red probabilities, as in Example 12.10. $P(18 \text{ to } 29 \text{ and video yes})$ is the result of following just the top branch in the tree diagram. So

$$\begin{aligned} P(18 \text{ to } 29 \mid \text{video yes}) &= \frac{P(18 \text{ to } 29 \text{ and video yes})}{P(\text{video yes})} \\ &= \frac{0.1890}{0.4913} = 0.3847 \end{aligned}$$

CONCLUDE: About 38% of adults who visit video-sharing sites are between 18 and 29 years old. Compare this conditional probability with the original information (unconditional) that 27% of adult Internet users are between 18 and 29 years old. Knowing that a person visits video-sharing sites increases the probability that he or she is young. ■

Examples 12.10 and 12.11 illustrate a common setting for tree diagrams. Some outcome (such as visiting video-sharing sites) has several sources (such as the three age groups). Starting from

- the probability of each source, and
- the conditional probability of the outcome given each source

the tree diagram leads to the overall probability of the outcome. Example 12.10 does this. You can then use the probability of the outcome and the definition of conditional probability to find the conditional probability of one of the sources given that the outcome occurred. Example 12.11 shows how.

APPLY YOUR KNOWLEDGE



12.13 Peanut and tree nut allergies. About 1% of the American population is allergic to peanuts or tree nuts.⁶ Choose 3 individuals at random and let the random variable X be the number in this sample who are allergic to peanuts or tree nuts. The possible values that X can take are 0, 1, 2, and 3. Make a three-stage tree diagram of the outcomes (allergic or not allergic) for the 3 individuals, and use it to find the probability distribution of X .

12.14 Testing for HIV. Enzyme immunoassay tests are used to screen blood specimens for the presence of antibodies to HIV, the virus that causes AIDS. Antibodies indicate the presence of the virus. The test is quite accurate but is not always correct. Here are approximate probabilities of positive and negative test results when the blood tested does and does not actually contain antibodies to HIV:⁷

		Test Result	
		Positive	Negative
Antibodies present		0.9985	0.0015
Antibodies absent		0.0060	0.9940

Suppose that 1% of a large population carries antibodies to HIV in their blood.

- Draw a tree diagram for selecting a person from this population (outcomes: antibodies present or absent) and testing his or her blood (outcomes: test positive or negative).
- What is the probability that the test is positive for a randomly chosen person from this population?

12.15 Peanut and tree nut allergies. Continue your work from Exercise 12.13. What is the conditional probability that exactly 2 of the people will be allergic to peanuts or tree nuts, given that at least 1 of the 3 people suffers from one of these allergies?



Politically correct

In 1950, the Soviet mathematician B. V. Gnedenko

(1912–1995) wrote *The Theory of Probability*, a text that was popular around the world. The introduction contains a mystifying paragraph that begins, “We note that the entire development of probability theory shows evidence of how its concepts and ideas were crystallized in a severe struggle between materialistic and idealistic conceptions.” It turns out that “materialistic” is jargon for “Marxist-Leninist.” It was good for the health of Soviet scientists in the Stalin era to add such statements to their books.

12.16 False HIV positives. Continue your work from Exercise 12.14. What is the probability that a person has the antibody, given that the test is positive? (Your result illustrates a fact that is important when considering proposals for widespread testing for HIV, illegal drugs, or agents of biological warfare: if the condition being tested is uncommon in the population, most positives will be false-positives.)

CHAPTER 12 SUMMARY

CHAPTER SPECIFICS

- Events A and B are **disjoint** if they have no outcomes in common. In that case, $P(A \text{ or } B) = P(A) + P(B)$.
- The **conditional probability** $P(B | A)$ of an event B given an event A is defined by

$$P(B | A) = \frac{P(\text{A and B})}{P(\text{A})}$$

when $P(A) > 0$. In practice, we most often find conditional probabilities from directly available information rather than from the definition.

- Events A and B are **independent** if knowing that one event occurs does not change the probability we would assign to the other event; that is, $P(B | A) = P(B)$. In that case, $P(\text{A and B}) = P(\text{A})P(\text{B})$.
- Any assignment of probability obeys these rules:

Addition rule for disjoint events: If events A, B, C, … are all disjoint in pairs, then

$$P(\text{at least one of these events occurs}) = P(\text{A}) + P(\text{B}) + P(\text{C}) + \dots$$

Multiplication rule for independent events: If events A, B, C, … are independent, then

$$P(\text{all of these events occur}) = P(\text{A})P(\text{B})P(\text{C}) \dots$$

General addition rule: For any two events A and B,

$$P(\text{A or B}) = P(\text{A}) + P(\text{B}) - P(\text{A and B})$$

General multiplication rule: For any two events A and B,

$$P(\text{A and B}) = P(\text{A})P(\text{B} | \text{A})$$

- Tree diagrams organize probability models that have several stages.

LINK IT

Probability models provide the important connection between the observed data and the process that generated the data. Probability is also the foundation of statistical inference and provides the language by which we answer questions and draw conclusions from our data. For these reasons, it is important that we have at least a basic understanding of what we mean by probability and of some of the rules and properties of probabilities.

Chapter 10 introduced basic ideas and facts about probability, and in this chapter we have considered some further details. The conditional distributions discussed in Chapter 6 are an example of the computation of conditional probabilities. Conditional probabilities play a central role when studying the relationship between two categorical variables, and we will return to them in Chapter 23. In terms of later chapters, the most important idea of this chapter is independence. This is an assumption that many statistical procedures make about the observations in a data set, and therefore, it is an important assumption to fully understand.

CHECK YOUR SKILLS

12.17 An instant lottery game gives you probability 0.02 of winning on any one play. Plays are independent of each other. If you play 3 times, the probability that you win on none of your plays is about
 (a) 0.98. (b) 0.94. (c) 0.000008.

12.18 The probability that you win on 1 or more of your 3 plays of the game in the previous exercise is about
 (a) 0.02. (b) 0.06. (c) 0.999992.

12.19 An athlete suspected of having used steroids is given two tests that operate independently of each other. Test A has probability 0.9 of being positive if steroids have been used. Test B has probability 0.8 of being positive if steroids have been used. What is the probability that *at least one* test is positive if steroids have been used?

- (a) 0.98 (b) 0.72 (c) 0.28

What is the distribution of doctorates conferred by field and sex? Here are the counts from the most popular fields in 2009.⁸ The physical sciences include mathematics and computer and information sciences; the life sciences include agricultural sciences/natural resources, biological/biomedical sciences, and health sciences; and the social sciences include psychology.

	Male	Female
Engineering	6,006	1,623
Physical sciences	5,868	2,450
Life sciences	5,180	6,212
Social sciences	3,259	4,575
Education	2,160	4,370
Other	3,865	3,960
Total	26,338	23,190

Exercises 12.20 to 12.23 are based on this table.

12.20 Choose a doctoral recipient at random from this group. The probability that the recipient is female is about

- (a) 0.42. (b) 0.47. (c) 0.58.

12.21 The conditional probability that the recipient is female given that the degree is in engineering is about

- (a) 0.03. (b) 0.07. (c) 0.21.

12.22 The conditional probability that the degree is in engineering given that the recipient is female is about

- (a) 0.03. (b) 0.07. (c) 0.21.

12.23 Let A be the event that the degree is in engineering and B the event that the recipient is female. The proportion of engineering doctorates conferred on females is expressed in probability notation as

- (a) $P(A \text{ and } B)$. (b) $P(A | B)$. (c) $P(B | A)$.

12.24 Choose an American adult at random. The probability that you choose a woman is 0.52. The probability that the person you choose has never married is 0.25. The probability that you choose a woman who has never married is 0.11. The probability that the person you choose is either a woman or never married (or both) is therefore about

- (a) 0.77. (b) 0.66. (c) 0.13.

12.25 Of people who died in the United States in recent years, 86% were white, 12% were black, and 2% were Asian. (This ignores a small number of deaths among other races.) Diabetes caused 2.8% of deaths among whites, 4.4% among blacks, and 3.5% among Asians. The probability that a randomly chosen death is a white who died of diabetes is about

- (a) 0.107. (b) 0.030. (c) 0.024.

12.26 Using the information in the previous exercise, the probability that a randomly chosen death was due to diabetes is about

- (a) 0.107. (b) 0.030. (c) 0.024.

CHAPTER 12 EXERCISES

12.27 Playing the lottery. New York State's "Quick Draw" lottery moves right along. Players choose between 1 and 10 numbers from the range 1 to 80; 20 winning numbers are displayed on a screen every four minutes. If you choose just 1 number, your probability of winning is $20/80$, or 0.25. Lester plays 1 number 8 times as he sits in a bar. What is the probability that all 8 bets lose?

12.28 Universal blood donors. People with type O-negative blood are universal donors. That is, any patient can receive a transfusion of O-negative blood. Only 7.2% of the American population have O-negative blood. If 10 people appear at random to give blood, what is the probability that at least 1 of them is a universal donor?

12.29 Playing the slots. Slot machines are now video games, with outcomes determined by random number generators. In the old days, slot machines were like this: you pull the lever to spin three wheels; each wheel has 20 symbols, all equally likely to show when the wheel stops spinning; the three wheels are independent of each other. Suppose that the middle wheel has 9 cherries among its 20 symbols, and the left and right wheels have 1 cherry each.

- (a) You win the jackpot if all three wheels show cherries. What is the probability of winning the jackpot?
- (b) There are three ways that the three wheels can show 2 cherries and 1 symbol other than a cherry. Find the probability of each of these ways.
- (c) What is the probability that the wheels stop with exactly 2 cherries showing among them?

12.30 A whale of a time. Hacksaw's Boats of St. Lucia takes tourists on a daily dolphin/whale watch cruise. Their brochure claims an 80% chance of sighting a dolphin or a whale, and you can assume that sightings from day to day are independent.

- (a) If you take the dolphin/whale watch cruise on two consecutive



Peter Dazeley/Getty



Mark Conlin/Alamy

days, what is the probability that you see a dolphin or a whale on both days?

- (b) If you take the dolphin/whale watch cruise on two consecutive days, what is the probability that you see a dolphin or a whale on at least one day? (*Hint:* First compute the probability that you don't see a dolphin or a whale on either day.)
- (c) If you want to have a 99% probability of seeing a dolphin or a whale at least once, what is the minimum number of days that you will need to take the cruise?

12.31 Tendon surgery. You have torn a tendon and are facing surgery to repair it. The surgeon explains the risks to you: infection occurs in 3% of such operations, the repair fails in 14%, and both infection and failure occur together in 1%. What percent of these operations succeed and are free from infection? Follow the four-step process in your answer.

12.32 A whale of a time, continued. Hacksaw's Boats of St. Lucia takes tourists on a daily dolphin/whale watch cruise. Their brochure claims an 80% chance of sighting a dolphin or a whale. Suppose that there is a 75% chance of seeing a dolphin and a 15% chance of seeing both a dolphin and a whale. Make a Venn diagram. Then answer these questions.

- (a) What is the probability of seeing a whale on the cruise?
- (b) What is the probability of seeing a whale but not a dolphin?
- (c) Are seeing a whale and seeing a dolphin independent events?

12.33 Tendon surgery, continued. You have torn a tendon and are facing surgery to repair it. The surgeon explains the risks to you: infection occurs in 3% of such operations, the repair fails in 14%, and both infection and failure occur together in 1%. What is the probability of infection given that the repair is successful? Follow the four-step process in your answer.

12.34 Screening job applicants. A company retains a psychologist to assess whether job applicants are suited for assembly-line work. The psychologist classifies applicants as one of A (well suited), B (marginal), or C (not suited). The company is concerned about the event D that an employee leaves the company within a year of being hired. Data on all people hired in the past five years give these probabilities:

$$\begin{array}{lll} P(A) = 0.4 & P(B) = 0.3 & P(C) = 0.3 \\ P(A \text{ and } D) = 0.1 & P(B \text{ and } D) = 0.1 & P(C \text{ and } D) = 0.2 \end{array}$$

Sketch a Venn diagram of the events A, B, C, and D and mark on your diagram the probabilities of all combinations

of psychological assessment and leaving (or not) within a year. What is $P(D)$, the probability that an employee leaves within a year?

12.35 Type of high school attended. Choose a college freshman at random and ask what type of high school they attended. Here is the distribution of results:⁹

Type	Regular public	Public charter	Public magnet	Private religious	Private independent	Home school
Probability	0.781	0.018	0.031	0.105	0.059	0.006

What is the conditional probability that a college freshman was home schooled given that he or she did not attend a regular public high school?

12.36 Income tax returns. Here is the distribution of the adjusted gross income (in thousands of dollars) reported on individual federal income tax returns in 2008:¹⁰

Income	<15	15–29	30–74	75–199	≥200
Probability	0.265	0.209	0.315	0.180	0.031

- (a) What is the probability that a randomly chosen return shows an adjusted gross income of \$30,000 or more?
- (b) Given that a return shows an income of at least \$30,000, what is the conditional probability that the income is at least \$75,000?

12.37 Thomas's pizza. You work at Thomas's pizza shop. You have the following information about the 7 pizzas in the oven: 3 of the 7 have thick crust, and of these 1 has only sausage and 2 have only mushrooms; the remaining 4 pizzas have regular crust, and of these 2 have only sausage and 2 have only mushrooms. Choose a pizza at random from the oven.

- (a) Are the events {getting a thick-crust pizza} and {getting a pizza with mushrooms} independent? Explain.
- (b) You add an eighth pizza to the oven. This pizza has thick crust with only cheese. Now are the events {getting a thick-crust pizza} and {getting a pizza with mushrooms} independent? Explain.

12.38 A probability teaser. Suppose (as is roughly correct) that each child born is equally likely to be a boy or a girl and that the sexes of successive children are independent. If we let BG mean that the older child is a boy and the younger child is a girl, then each of the combinations BB, BG, GB, GG has probability 0.25. Ashley and Brianna each have two children.

- (a) You know that at least one of Ashley's children is a boy. What is the conditional probability that she has two boys?

- (b) You know that Brianna's older child is a boy. What is the conditional probability that she has two boys?

12.39 College degrees. A striking trend in higher education is that more women than men reach each level of attainment. The National Center for Education Statistics provides projections for the number of degrees earned, classified by level and by the sex of the degree recipient. Here are the projected number of earned degrees (in thousands) in the United States for the 2015–2016 academic year:¹¹

	Associate's	Bachelor's	Master's	Professional	Doctorate	Total
Female	556	1034	450	54	45	2139
Male	311	737	282	53	38	1421
Total	867	1771	732	107	83	3560

- (a) If you choose a degree recipient at random, what is the probability that the person you choose is a man?
- (b) What is the conditional probability that you choose a man given that the person chosen received a master's?
- (c) Are the events "choose a man" and "choose a master's degree recipient" independent? How do you know?

12.40 College degrees. Exercise 12.39 gives the projected counts (in thousands) of earned degrees in the United States in the 2015–2016 academic year. Use these data to answer the following questions.

- (a) What is the probability that a randomly chosen degree recipient is a woman?
- (b) What is the conditional probability that the person chosen received an associate's degree given that she is a woman?
- (c) Use the multiplication rule to find the probability of choosing a female associate's degree recipient. Check your result by finding this probability directly from the table of counts.

12.41 Deer and pine seedlings. As suburban gardeners know, deer will eat almost anything green. In a study of pine seedlings at an environmental center in Ohio, researchers noted how deer damage varied with how much of the seedling was covered by thorny undergrowth:¹²

Thorny cover	Deer Damage	
	Yes	No
None	60	151
<1/3	76	158
1/3 to 2/3	44	177
>2/3	29	176

(a) What is the probability that a randomly selected seedling was damaged by deer?

(b) What are the conditional probabilities that a randomly selected seedling was damaged given each level of cover?

(c) Does knowing about the amount of thorny cover on a seedling change the probability of deer damage? If so, cover and damage are not independent.



Peter Skinner/Photo Researchers

12.42 Deer and pine seedlings. In the setting of Exercise 12.41, what percent of the trees that were not damaged by deer were more than two-thirds covered by thorny plants?

12.43 Deer and pine seedlings. In the setting of Exercise 12.41, what percent of the trees that were damaged by deer were less than one-third covered by thorny plants?

Julie graduates from college. Julie has studied biology, chemistry, and computing and hopes to use her science background in crime investigation. Late one night she thinks about some jobs for which she has applied. Let A, B, and C be the events that Julie is offered a job by

A = the Connecticut Office of the Chief Medical Examiner

B = the New Jersey Division of Criminal Justice

C = the federal Disaster Mortuary Operations Response Team

Julie writes down her personal probabilities for being offered these jobs:

$$\begin{aligned} P(A) &= 0.5 & P(B) &= 0.4 & P(C) &= 0.2 \\ P(A \text{ and } B) &= 0.1 & P(A \text{ and } C) &= 0.5 & P(B \text{ and } C) &= 0.05 \\ P(A \text{ and } B \text{ and } C) &= 0 \end{aligned}$$

Make a Venn diagram of the events A, B, and C. As in Figure 12.4 (see page 313), mark the probabilities of every intersection involving these events. Use this diagram for Exercises 12.44 to 12.46.

12.44 Will Julie get a job offer? What is the probability that Julie is not offered any of the three jobs?

12.45 Will Julie get just this offer? What is the probability that Julie is offered the Connecticut job but not the New Jersey or federal job?

12.46 Julie's conditional probabilities. If Julie is offered the federal job, what is the conditional probability that she is also offered the New Jersey job? If Julie is offered the New

Jersey job, what is the conditional probability that she is also offered the federal job?

12.47 The geometric distributions. You are rolling a pair of balanced dice in a board game. Rolls are independent. You land in a danger zone that requires you to roll doubles (both faces show the same number of spots) before you are allowed to play again. How long will you wait to play again?

(a) What is the probability of rolling doubles on a single toss of the dice? (If you need review, the possible outcomes appear in Figure 10.2 (page 265). All 36 outcomes are equally likely.)

(b) What is the probability that you do not roll doubles on the first toss, but you do on the second toss?

(c) What is the probability that the first two tosses are not doubles and the third toss is doubles? This is the probability that the first doubles occurs on the third toss.

(d) Now you see the pattern. What is the probability that the first doubles occurs on the fourth toss? On the fifth toss? Give the general result: what is the probability that the first doubles occurs on the k th toss?

(e) What is the probability that you get to go again within 3 turns?

(Comment: The distribution of the number of trials to the first success is called a *geometric distribution*. In this problem you have found geometric distribution probabilities when the probability of a success on each trial is $1/6$. The same idea works for any probability of success.)

12.48 Winning at tennis. A player serving in tennis has two chances to get a serve into play. If the first serve is out, the player serves again. If the second serve is also out, the player loses the point. Here are probabilities based on four years of the Wimbledon Championship:¹³

$$P(\text{1st serve in}) = 0.59$$

$$P(\text{win point} | \text{1st serve in}) = 0.73$$

$$P(\text{2nd serve in} | \text{1st serve out}) = 0.86$$

$$P(\text{win point} | \text{1st serve out and 2nd serve in}) = 0.59$$

Make a tree diagram for the results of the two serves and the outcome (win or lose) of the point. (The branches in your tree have different numbers of stages depending on the outcome of the first serve.) What is the probability that the serving player wins the point?

12.49 Urban voters. The voters in a large city are 40% white, 40% black, and 20% Hispanic. (Hispanics may be of any race in official statistics, but here we are speaking of political blocks.) A black mayoral candidate anticipates attracting 30% of the white vote, 90% of the black vote, and 50% of the Hispanic vote. Draw a tree

diagram with probabilities for the race (white, black, or Hispanic) and vote (for or against the candidate) of a randomly chosen voter. What percent of the overall vote does the candidate expect to get? Use the four-step process to guide your work.

12.50 Winning at tennis, continued. Based on your work in Exercise 12.48, in what percent of points won by the server was the first serve in? (Write this as a conditional probability and use the definition of conditional probability.)

12.51 Where do the votes come from? In the election described in Exercise 12.49, what percent of the candidate's votes come from black voters? (Write this as a conditional probability and use the definition of conditional probability.)

12.52 Teens and texting. The Pew Internet and American Life Project finds that 75% of teenagers (ages 12 to 17) now own cell phones, and of the teens who own cell phones, 87% use text messaging.¹⁴

- (a) What percent of teens own cell phones and are "texters?"
- (b) Among teens who own cell phones and are texters, 15% send more than 200 texts a day, or more than 6,000 texts a month. What percent of all teens own a cell phone, are texters and send more than 6,000 texts a month?

12.53 Lactose intolerance. Lactose intolerance causes difficulty digesting dairy products that contain lactose (milk sugar). It is particularly common among people of African and Asian ancestry. In the United States (ignoring other groups and people who consider themselves to belong to more than one race), 82% of the population is white, 14% is black,

and 4% is Asian. Moreover, 15% of whites, 70% of blacks, and 90% of Asians are lactose intolerant.¹⁵

- (a) What percent of the entire population is lactose intolerant?
- (b) What percent of people who are lactose intolerant are Asian?

12.54 Fundraising by telephone. Tree diagrams can organize problems having more than two stages. Figure 12.6 shows probabilities for a charity calling potential donors by telephone.¹⁶ Each person called is either a recent donor, a past donor, or a new prospect. At the next stage, the person called either does or does not pledge to contribute, with conditional probabilities that depend on the donor class the person belongs to. Finally, those who make a pledge either do or don't actually make a contribution.

- (a) What percent of calls result in a contribution?
- (b) What percent of those who contribute are recent donors?

DNA forensics. When a suspect's DNA is compared with a sample of DNA collected at a crime scene, the comparison is made between certain sections of the DNA called loci. Each locus has two alleles (gene forms), one inherited from the mother and the other from the father. Suppose that there are two alleles, called A and B, for a particular locus. These alleles can be present at the locus in three combinations. A person's alleles at the locus could both be A, one allele could be A and the other B, or both alleles could be B, giving the three combinations (A and A), (A and B), and (B and B). Here's how the math works. If the proportion of the population with allele A as one of their alleles at the locus is a , and the proportion of the population with allele

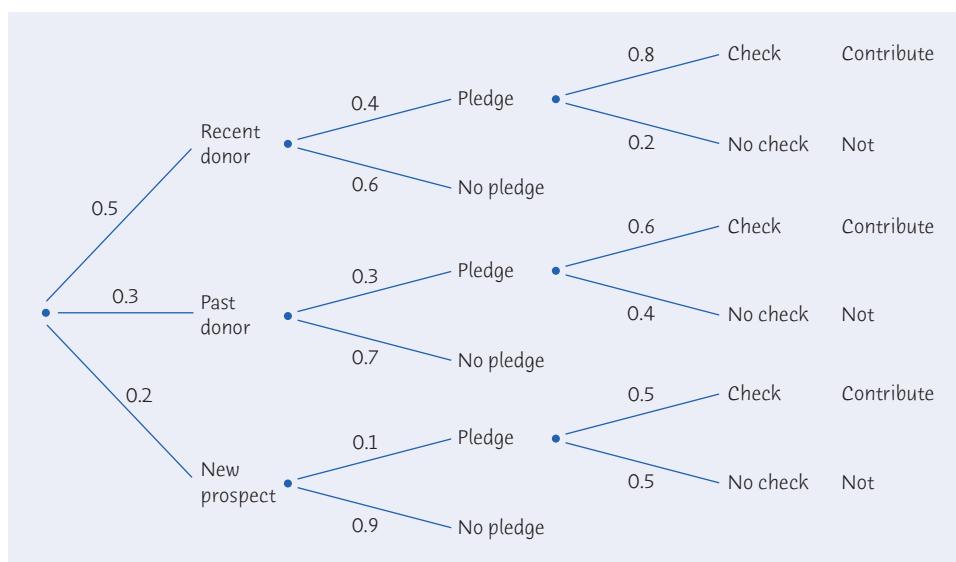


FIGURE 12.6

Tree diagram for fundraising by telephone, for Exercise 12.54. The three stages are the type of prospect called, whether or not the person makes a pledge, and whether or not a person who pledges actually makes a contribution.

B as one of their alleles at the locus is *b*, then the proportions of the population with the three combinations of these allele types at the locus are:

Alleles at the locus	Population proportion with allele combination
A and A	a^2
A and B	$2ab$
B and B	b^2

Use this information in Exercises 12.55 and 12.56. The numbers used in the exercises are from the FBI data base.¹⁷

12.55 Alleles at D21S11. Suppose that the locus D21S11 has two alleles called 29 and 31. The proportion of the Caucasian population with allele 29 is 0.181 and with allele 31 is 0.071. What proportion of the Caucasian population has the combination (29, 31) at the locus D21S11? What proportion has the combination (29, 29)?

12.56 Alleles at D3S1358. Suppose that the locus D3S1358 has two alleles called 16 and 17. The proportion of the Caucasian population with allele 16 is 0.232 and with allele 17 is 0.212. What proportion of the Caucasian population has the combination (16, 17) at the locus D3S1358.

12.57 Alleles at multiple loci. One important fact regarding the loci evaluated in such forensic tests is that the allele combinations at each locus have been shown to be independent. What proportion of the Caucasian population has the combination (29, 31) at the loci D21S11 and combination (16, 17) at the loci D3S1358? As we specify the alleles present at more loci, what will happen to the proportion of the Caucasian population that matches the allele combinations at all the loci?

12.58 How many match? A defendant in Ohio was indicted on December 17, 2009, on charges of aggravated burglary and assault. A hair found at the crime scene was tested at six loci and demonstrated a specific combination of alleles found in about 1 in 1.6 million individuals in the population. Comparison of the DNA profile found on the hair with a data base of convicted felons revealed a match between the allelic profile found on the hair and an individual in the data base (the defendant). Defense attorneys in the case requested that the State perform additional DNA testing since several previously untested loci were available to test. The results of this testing revealed that the defendant did not match at some of these newly tested markers, indicating that the DNA from the hair was not the defendant's. The charges were dropped. If the DNA profile (or combination of alleles) found on the hair is possessed by 1 in 1.6 million individuals, and the data base of convicted felons contains 4.5 million individuals, approximately how many individuals in the data base would demonstrate a match between their DNA and that found on the hair?



EXPLORING THE WEB

12.59 Mathematics SAT scores. The Web site <http://professionals.collegeboard.com/data-reports-research/sat> presents data for high school seniors who participated in the SAT Program during the current year as well as previous years. Under *SAT Data & Reports*, click on the link for *College Bound Seniors* for the most recent year given. In the window that opens, click on the link for *Total Group Report: College Bound Seniors* for this year. The Total Group Profile Report presents data for high school graduates who participated in the SAT Program that year. Students are counted only once, no matter how often they are tested, and only their latest scores are summarized. Suppose a high school graduate is selected at random. Use the information in the tables available to answer the following questions.

- What is the probability that the selected student is female?
- What is the probability that the selected student scores 600 or over on the Mathematics section of the SAT?
- What is the conditional probability that the selected student scores 600 or over on the Mathematics section given that the student is male?

- (d) What is the conditional probability that the selected student scores 600 or over on the Mathematics section given that the student is female?
- (e) Are scoring over 600 on the Mathematics section and sex independent? If not, explain in a simple sentence the nature of the dependence.

12.60 Let's make a deal. The Monty Hall Problem is an example of a simple probability problem with an answer that is counterintuitive. The problem was made popular when Marilyn vos Savant published the problem in her *Parade Magazine* column.¹⁸ Here is the question:

"Suppose you're on a game show, and you're given a choice of three doors: Behind one door is a car; behind the others, goats. You pick a door, say number 1, and the host, who knows what's behind the doors, opens another door, say number 3, which has a goat. He says to you, "Do you want to pick door number 2?" Is it to your advantage to switch your choice of doors?"

Go to the Web site www.letsmakeadeal.com/problem.htm. Then use the *Three Doors Simulation* link further down on the Web page to play the game 30 times.

- (a) What proportion of the time would you have won if you had switched doors? If you had not switched? Based on your simulation, does it seem better to switch or to stay with your initial choice?
- (b) The Web site provides an explanation of why it is better to switch doors. The explanation uses conditional probabilities, and in the notation of the Web site, $P(A \wedge B)$ is another way of writing $P(A \text{ and } B)$. Following the guidelines on the Web site and using the notation in our text, show why the probability of getting the car is $1/3$ if you do not switch doors and $2/3$ if you do switch doors. How well do these probabilities agree with the proportions you found in your simulation?



Binomial Distributions*

A basketball player shoots 5 free throws. How many does she make? A sample survey dials 1200 residential phone numbers at random. How many live people answer the phone? You plant 10 dogwood trees. How many live through the winter? In all these situations, we want a probability model for a *count* of successful outcomes.

THE BINOMIAL SETTING AND BINOMIAL DISTRIBUTIONS

The distribution of a count depends on how the data are produced. Here is a common situation.

THE BINOMIAL SETTING

1. There are a fixed number n of observations.
2. The n observations are all **independent**. That is, knowing the result of one observation does not change the probabilities we assign to other observations.
3. Each observation falls into one of just two categories, which for convenience we call “success” and “failure.”
4. The probability of a success, call it p , is the same for each observation.

Think of tossing a coin n times as an example of the binomial setting. Each toss gives either heads or tails. Knowing the outcome of one toss

*This more advanced chapter concerns a special topic in probability. The material is not needed to read the rest of the book.

Chapter 13

IN THIS CHAPTER WE COVER...

- The binomial setting and binomial distributions
- Binomial distributions in statistical sampling
- Binomial probabilities
- Using technology
- Binomial mean and standard deviation
- The Normal approximation to binomial distributions



Was he good or was he lucky?

When a baseball player hits .300, everyone applauds. A .300 hitter gets a hit in 30% of times at bat. Could a .300 year just be luck? Typical major leaguers bat about 500 times a season and hit about .260. A hitter's successive tries seem to be independent, so we have a binomial setting. From this model, we can calculate or simulate the probability of hitting .300. It is about 0.025. Out of 100 run-of-the-mill major league hitters, two or three each year will bat .300 because they were lucky.

doesn't change the probability of a head on any other toss, so the tosses are independent. If we call heads a success, then p is the probability of a head and remains the same as long as we toss the same coin. For tossing a coin, p is close to 0.5. If we spin the coin on a flat surface rather than toss it, p is not equal to 0.5. The number of heads we count is a discrete random variable X . The distribution of X is called a *binomial distribution*.

BINOMIAL DISTRIBUTION

The count X of successes in the binomial setting has the **binomial distribution** with parameters n and p . The parameter n is the number of observations, and p is the probability of a success on any one observation. The possible values of X are the whole numbers from 0 to n .



The binomial distributions are an important class of finite probability models. Pay attention to the binomial setting, because not all counts have binomial distributions.

EXAMPLE 13.1 Blood types

Genetics says that children receive genes from their parents independently. Each child of a certain pair of parents has probability 0.25 of having type O blood. If these parents have 5 children, the number who have type O blood is the count X of successes in 5 independent observations with probability 0.25 of a success on each observation. So X has the binomial distribution with $n = 5$ and $p = 0.25$.

EXAMPLE 13.2 Counting boys

Here is set of genetic examples that require more thought.

Choose two births at random from the last year's births at a large hospital and count the number of boys (0, 1, or 2). The sexes of children born to different mothers are surely independent. The probability that a randomly chosen birth in Canada and the United States is a boy is about 0.52. (Why it is not 0.5 is something of a mystery.) So the count of boys has a binomial distribution with $n = 2$ and $p = 0.52$.

Next, observe successive births at a large hospital and let X be the number of births until the first boy is born. Births are independent and each has probability 0.52 of being a boy. Yet X is not binomial, because there is no fixed number of observations. "Count observations until the first success" is a different setting than "count the number of successes in a fixed number of observations."

Finally, choose at random a family with exactly two children and count the number of boys. Careful study of such families shows that the count of boys is not binomial: the probability of exactly 1 boy is too high.¹ Families are less likely to have a third child if the first two are a boy and a girl, so when we look at families that stopped at two children, "one of each" is more common than if we look at randomly chosen births. The sexes of successive children in two-child families are not independent, because the parents' choices interfere with the genetics.

BINOMIAL DISTRIBUTIONS IN STATISTICAL SAMPLING

The binomial distributions are important in statistics when we wish to make inferences about the proportion p of “successes” in a population. Here is a typical example.

EXAMPLE 13.3 Choosing an SRS of CDs

A music distributor inspects an SRS of 10 CDs from a shipment of 10,000 music CDs. Suppose that (unknown to the distributor) 10% of the CDs in the shipment have defective copy-protection schemes that will harm personal computers. Count the number X of bad CDs in the sample.

This is not quite a binomial setting. Removing 1 CD changes the proportion of bad CDs remaining in the shipment. So the probability that the second CD chosen is bad changes when we know whether the first is good or bad. But removing 1 CD from a shipment of 10,000 changes the makeup of the remaining 9999 CDs very little. In practice, the distribution of X is very close to the binomial distribution with $n = 10$ and $p = 0.1$. ■

Example 13.3 shows how we can use the binomial distributions in the statistical setting of selecting an SRS. When the population is much larger than the sample, a count of successes in an SRS of size n has approximately the binomial distribution with n equal to the sample size and p equal to the proportion of successes in the population.

SAMPLING DISTRIBUTION OF A COUNT

Choose an SRS of size n from a population with proportion p of successes. When the population is much larger than the sample, the count X of successes in the sample has approximately the binomial distribution with parameters n and p .

APPLY YOUR KNOWLEDGE

In each of Exercises 13.1 to 13.3, X is a count. Does X have a binomial distribution? Give your reasons in each case.

- 13.1 **Random digit dialing.** When an opinion poll calls residential telephone numbers at random, only 20% of the calls reach a live person. You watch the random dialing machine make 15 calls. X is the number that reach a live person.
- 13.2 **Random digit dialing.** When an opinion poll calls residential telephone numbers at random, only 20% of the calls reach a live person. You watch the random dialing machine make calls. X is the number of calls until the first live person answers.
- 13.3 **Boxes of tiles.** Boxes of six-inch slate flooring tile contain 40 tiles per box. The count X is the number of cracked tiles in a box. You have noticed that most boxes contain no cracked tiles, but if there are cracked tiles in a box, then there are usually several.

13.4 Canadian Internet use. A survey finds that in 2009, 80% of Canadians aged 16 and older used the Internet for personal reasons.² In the survey, an “Internet user” is defined as someone who used the Internet for personal reasons from any location in the 12 months preceding the survey. If you take an SRS of 1500 Canadians aged 16 and over, what is the approximate distribution of the number in your sample who have used the Internet for personal reasons?

BINOMIAL PROBABILITIES

We can find a formula for the probability that a binomial random variable takes any value by adding probabilities for the different ways of getting exactly that many successes in n observations. Here is an example that illustrates the idea.

EXAMPLE 13.4 Inheriting blood type

The blood types of successive children born to the same parents are independent and have fixed probabilities that depend on the genetic makeup of the parents. Each child born to a certain set of parents has probability 0.25 of having blood type O. If these parents have 5 children, what is the probability that exactly 2 of them have type O blood?

The count of children with type O blood is a binomial random variable X with $n = 5$ tries and probability $p = 0.25$ of a success on each try. We want $P(X = 2)$. ■

Because the method doesn’t depend on the specific example, let’s use “S” for success and “F” for failure for short. Do the work in two steps.

Step 1. Find the probability that a specific 2 of the 5 tries, say the first and the third, give successes. This is the outcome SFSFF. Because tries are independent, the multiplication rule for independent events applies. The probability we want is

$$\begin{aligned} P(\text{SFSFF}) &= P(\text{S})P(\text{F})P(\text{S})P(\text{F})P(\text{F}) \\ &= (0.25)(0.75)(0.25)(0.75)(0.75) \\ &= (0.25)^2(0.75)^3 \end{aligned}$$

Step 2. Observe that *any one arrangement* of 2 S’s and 3 F’s has this same probability. This is true because we multiply together 0.25 twice and 0.75 three times whenever we have 2 S’s and 3 F’s. The probability that $X = 2$ is the probability of getting 2 S’s and 3 F’s in any arrangement whatsoever. Here are all the possible arrangements:

SSFFF	SFSFF	SFFSF	SFFFS	FSSFF
FSFSF	FSFFS	FFSSF	FFSFS	FFFSS

There are 10 of them, all with the same probability. The overall probability of 2 successes is therefore

$$P(X = 2) = 10(0.25)^2(0.75)^3 = 0.2637$$

The pattern of this calculation works for any binomial probability. To use it, we must count the number of arrangements of k successes in n observations. We use the following fact to do the counting without actually listing all the arrangements.

BINOMIAL COEFFICIENT

The number of ways of arranging k successes among n observations is given by the **binomial coefficient**

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

for $k = 0, 1, 2, \dots, n$.

The formula for binomial coefficients uses the **factorial** notation. For any positive whole number n , its factorial $n!$ is

$$n! = n \times (n-1) \times (n-2) \times \cdots \times 3 \times 2 \times 1$$

In addition, we define $0! = 1$.

The larger of the two factorials in the denominator of a binomial coefficient will cancel much of the $n!$ in the numerator. For example, the binomial coefficient we need for Example 13.4 is

$$\begin{aligned}\binom{5}{2} &= \frac{5!}{2!3!} \\ &= \frac{(5)(4)(3)(2)(1)}{(2)(1) \times (3)(2)(1)} \\ &= \frac{(5)(4)}{(2)(1)} = \frac{20}{2} = 10\end{aligned}$$

The binomial coefficient $\binom{5}{2}$ is not related to the fraction $\frac{5}{2}$. A helpful way to remember its meaning is to read it as “5 choose 2.” Binomial coefficients have many uses, but we are interested in them only as an aid to finding binomial probabilities. The binomial coefficient $\binom{n}{k}$ counts the number of different ways in which k successes can be arranged among n observations. The binomial probability $P(X = k)$ is this count multiplied by the probability of any one specific arrangement of the k successes. Here is the result we seek.



factorial

What looks random?

Toss a coin six times and record heads (H) or tails (T) on each toss. Which of these outcomes is more probable: HTHTTH or TTTHHH? Almost everyone says that HTHTTH is more probable, because TTTHHH does not “look random.” In fact, both are equally probable. That heads has probability 0.5 says that about half of a very long sequence of tosses will be heads. It doesn’t say that heads and tails must come close to alternating in the short run. The coin doesn’t know what past outcomes were, and it can’t try to create a balanced sequence.

BINOMIAL PROBABILITY

If X has the binomial distribution with n observations and probability p of success on each observation, the possible values of X are $0, 1, 2, \dots, n$. If k is any one of these values,

$$P(X = k) = \binom{n}{k} p^k (1-p)^{n-k}$$

EXAMPLE 13.5 Inspecting CDs

The number X of CDs with defective copy protection in Example 13.3 has approximately the binomial distribution with $n = 10$ and $p = 0.1$.

The probability that the sample contains no more than 1 defective CD is

$$\begin{aligned} P(X \leq 1) &= P(X = 1) + P(X = 0) \\ &= \binom{10}{1}(0.1)^1(0.9)^9 + \binom{10}{0}(0.1)^0(0.9)^{10} \\ &= \frac{10!}{1!9!}(0.1)(0.3874) + \frac{10!}{0!10!}(1)(0.3487) \\ &= (10)(0.1)(0.3874) + (1)(1)(0.3487) \\ &= 0.3874 + 0.3487 = 0.7361 \end{aligned}$$

This calculation uses the facts that $0! = 1$ and that $a^0 = 1$ for any number a other than 0. We see that about 74% of all samples will contain no more than 1 bad CD. In fact, 35% of the samples will contain no bad CDs. A sample of size 10 cannot be trusted to alert the distributor to the presence of unacceptable CDs in the shipment.

Rule 4 described in Chapter 10 can make the computation of certain binomial probabilities simpler. For example, the probability that the sample contains at least 1 defective CD is

$$\begin{aligned} P(X \geq 1) &= P(X = 1) + P(X = 2) + \cdots + P(X = 10) \\ &= 1 - P(X = 0) \\ &= 1 - 0.3487 = 0.6513 \end{aligned}$$

When computing binomial probabilities by hand, it is useful to keep this rule in mind. ■

USING TECHNOLOGY

The binomial probability formula is awkward to use unless the number of observations n is quite small. You can find tables of binomial probabilities $P(X = k)$ and cumulative probabilities $P(X \leq k)$ for selected values of n and p , but the most efficient way to do binomial calculations is to use technology.

Figure 13.1 shows output for the calculation in Example 13.5 from a graphing calculator, two statistical programs, and a spreadsheet program. We asked all four to give cumulative probabilities. The calculator, Minitab, and CrunchIt! have menu

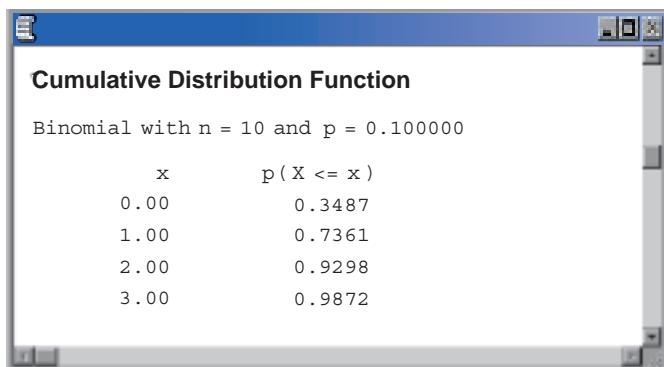
Texas Instruments Graphing Calculator

```
binomcdf(10,0.1,
1)
.7361
```

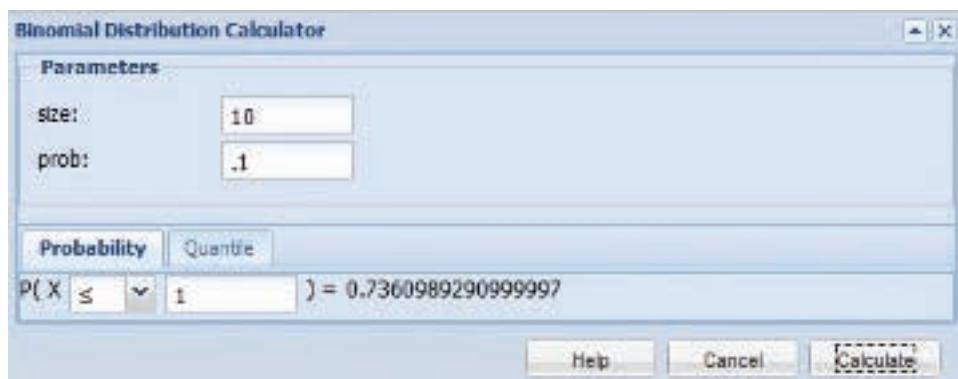
FIGURE 13.1

The binomial probability $P(X \leq 1)$ for Example 13.5: output from a graphing calculator, two statistical programs, and a spreadsheet program.

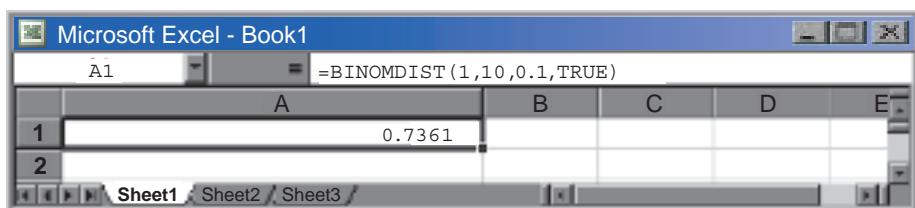
Minitab



CrunchIt!



Microsoft Excel

**FIGURE 13.1** (Continued)

entries for binomial cumulative probabilities. Excel has no menu entry, but the worksheet function BINOMDIST is available. All the outputs agree with the result 0.7361 of Example 13.5.

APPLY YOUR KNOWLEDGE

13.5 Proofreading. Typing errors in a text are either nonword errors (as when “the” is typed as “teh”) or word errors that result in a real but incorrect word. Spell-checking software will catch nonword errors but not word errors. Human proofreaders catch 70% of word errors. You ask a fellow student to proofread an essay in which you have deliberately made 10 word errors.

- (a) If the student matches the usual 70% rate, what is the distribution of the number of errors caught? What is the distribution of the number of errors missed?
- (b) Missing 3 or more out of 10 errors seems a poor performance. What is the probability that a proofreader who catches 70% of word errors misses exactly 3 out of 10? If you use software, also find the probability of missing 3 or more out of 10.

13.6 Random digit dialing. When an opinion poll calls residential telephone numbers at random, only 20% of the calls reach a live person. You watch the random digit dialing machine make 15 calls.

- (a) What is the probability that exactly 3 calls reach a person?
- (b) What is the probability that at most 3 calls reach a person?
- (c) What is the probability that at least 3 calls reach a person?
- (d) What is the probability that fewer than 3 calls reach a person?
- (e) What is the probability that more than 3 calls reach a person?

13.7 Google does binomial. Point your Web browser to www.google.com. Instead of searching the Web or looking for images, you can request a calculation in the Search box.

- (a) Enter `5 choose 2` and click Search. What does Google return?
- (b) You see that Google calculates the binomial coefficient “5 choose 2.” What are the values of the binomial coefficients for “500 choose 2” and “500 choose 100”? We expect that there are more ways to choose 100 than to choose 2, but how many more may be a surprise. That 10^{107} in Google’s answer means a 1 followed by 107 zeros.
- (c) Google also does binomial probabilities. Enter `(10 choose 1)*0.1*0.9^9` to find the first binomial probability in Example 13.5 (page 336). What is Google’s answer with all its decimal places?

BINOMIAL MEAN AND STANDARD DEVIATION

If a count X has the binomial distribution based on n observations with probability p of success, what is its mean μ ? That is, in very many repetitions of the binomial setting, what will be the average count of successes? We can guess the answer. If a basketball player makes 80% of her free throws, the mean number made in 10 tries should be 80% of 10, or 8. In general, the mean of a binomial distribution should be $\mu = np$. Here are the facts.

BINOMIAL MEAN AND STANDARD DEVIATION

If a count X has the binomial distribution with number of observations n and probability of success p , the **mean** and **standard deviation** of X are

$$\begin{aligned}\mu &= np \\ \sigma &= \sqrt{np(1 - p)}\end{aligned}$$



Remember that these short formulas are good only for binomial distributions. They can't be used for other distributions.

EXAMPLE 13.6 Inspecting CDs

Continuing Example 13.5, the count X of bad CDs is binomial with $n = 10$ and $p = 0.1$. The histogram in Figure 13.2 displays this probability distribution. (Because probabilities are long-run proportions, using probabilities as the heights of the bars shows what the distribution of X would be in very many repetitions.) The distribution is strongly right-skewed. Although X can take any whole-number value from 0 to 10, the probabilities of values larger than 5 are so small that they do not appear in the histogram.

The mean and standard deviation of the binomial distribution in Figure 13.2 are

$$\begin{aligned}\mu &= np \\ &= (10)(0.1) = 1 \\ \sigma &= \sqrt{np(1 - p)} \\ &= \sqrt{(10)(0.1)(0.9)} = \sqrt{0.9} = 0.9487\end{aligned}$$

The mean is marked on the probability histogram in Figure 13.2. ■

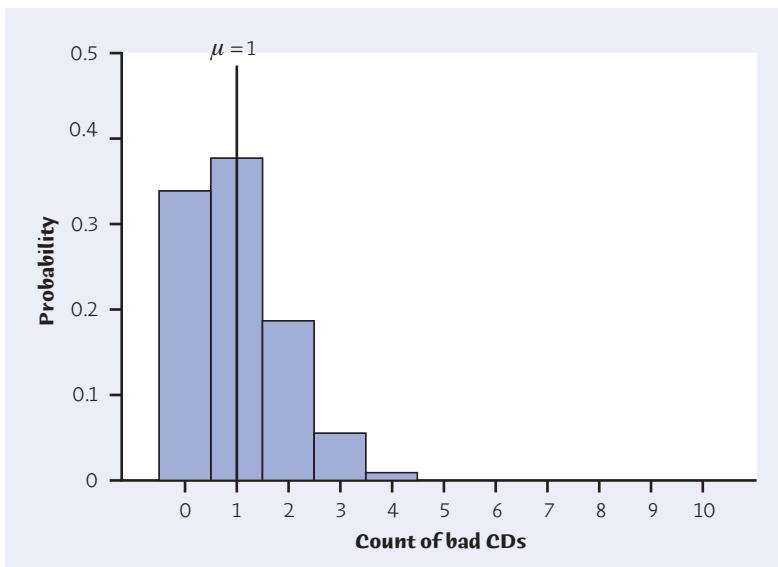
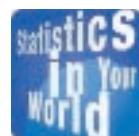


FIGURE 13.2

Probability histogram for the binomial distribution with $n = 10$ and $p = 0.1$, for Example 13.6.



Randomness turns silver to bronze

After many charges of favoritism by judges, the rules for scoring international figure skating competitions changed in 2004. The big change is that 12 judges score all performances, then scores from 3 judges chosen at random are dropped for each part of the program. So there are $\binom{12}{9} = 220$ possible panels of 9 judges for (say) the "Free Skate" and these panels will have slightly different scores. Result: at the 2006 World Figure Skating Championships, the Russian pair Maria Petrova and Alexei Tikhonov received the bronze medal when the consensus of all 12 judges would have given them the silver medal. Perhaps the system needs another change.

APPLY YOUR KNOWLEDGE

- 13.8 Random digit dialing.** When an opinion poll calls residential telephone numbers at random, only 20% of the calls reach a live person. You watch the random digit dialing machine make 15 calls.

- What is the mean number of calls that reach a person?
- What is the standard deviation σ of the count of calls that reach a person?
- If calls are made to New York City rather than nationally, the probability that a call reaches a person is only $p = 0.08$. How does this new p affect the standard deviation? What would be the standard deviation if $p = 0.01$? What does your work show about the behavior of the standard deviation of a binomial distribution as the probability of a success gets closer to 0?

13.9 Proofreading. Return to the proofreading setting of Exercise 13.5 (page 336).

- If X is the number of word errors missed, what is the distribution of X ? If Y is the number of word errors caught, what is the distribution of Y ?
- What is the mean number of errors caught? What is the mean number of errors missed? The mean counts of successes and of failures always add to n , the number of observations.
- What is the standard deviation of the number of errors caught? What is the standard deviation of the number of errors missed? The standard deviations of the count of successes and the count of failures are always the same.

THE NORMAL APPROXIMATION TO BINOMIAL DISTRIBUTIONS

It isn't practical to use the formula for binomial probabilities when the number of observations n is large. (Look at part (b) of Exercise 13.7 to see why.) Software or a graphing calculator will handle many problems that are beyond the reach of hand calculation. If technology does not rescue you, there is another alternative: *as the number of observations n gets larger, the binomial distribution gets close to a Normal distribution*. When n is large, we can use Normal probability calculations to approximate binomial probabilities. Here are the facts.

NORMAL APPROXIMATION FOR BINOMIAL DISTRIBUTIONS

Suppose that a count X has the binomial distribution with n observations and success probability p . When n is large, the distribution of X is approximately Normal, $N(np, \sqrt{np(1-p)})$.

As a rule of thumb, we will use the Normal approximation when n is so large that $np \geq 10$ and $n(1-p) \geq 10$.

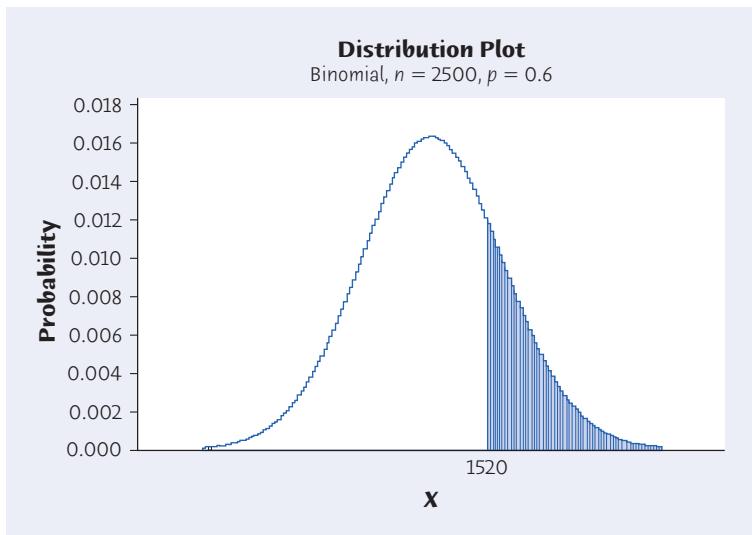
The Normal approximation is easy to remember because it says to act as if X is Normal with exactly the same mean and standard deviation as the binomial distribution. The accuracy of the Normal approximation improves as the sample size n increases. It is most accurate for any fixed n when p is close to 1/2 and least accurate when p is near 0 or 1. This is why the rule of thumb in the box depends on p as well as n .



Erica Shires/zefa/CORBIS

EXAMPLE 13.7 Attitudes toward shopping

How many people enjoy shopping? A survey asked a nationwide random sample of 2500 adults if they agreed or disagreed that "I like buying new clothes, but shopping is often frustrating and time-consuming."³ The population that the poll wants to draw conclusions about is all U.S. residents aged 18 and over. Suppose that in fact 60% of all adult U.S. residents would say "Agree" if asked the same question. What is the probability that 1520 or more of the sample agree? ■

**FIGURE 13.3**

Probability histogram for the binomial distribution with $n = 2500$ and $p = 0.6$. The bars at and above 1520 are shaded to highlight the probability of getting at least 1520 successes. The shape of this binomial probability distribution closely resembles a Normal curve.

Because there are about 235 million adults in the United States, the responses of 2500 randomly chosen adults are very close to independent. So the number in our sample who agree that shopping is frustrating is a random variable X having the binomial distribution with $n = 2500$ and $p = 0.6$. To find the probability $P(X \geq 1520)$ that at least 1520 of the people in the sample find shopping frustrating, we must add the binomial probabilities of all outcomes from $X = 1520$ to $X = 2500$. Figure 13.3 is a probability histogram of this binomial distribution, from Minitab. As the Normal approximation suggests, the shape of the distribution looks Normal. The probability we want is the sum of the heights of the shaded bars. Here are three ways to find this probability.

1. Use technology. Statistical software can find the exact binomial probability. In most cases, software finds cumulative probabilities $P(X \leq x)$. So start by writing

$$P(X \geq 1520) = 1 - P(X \leq 1519)$$

Here is Minitab's answer for $P(X \leq 1519)$:

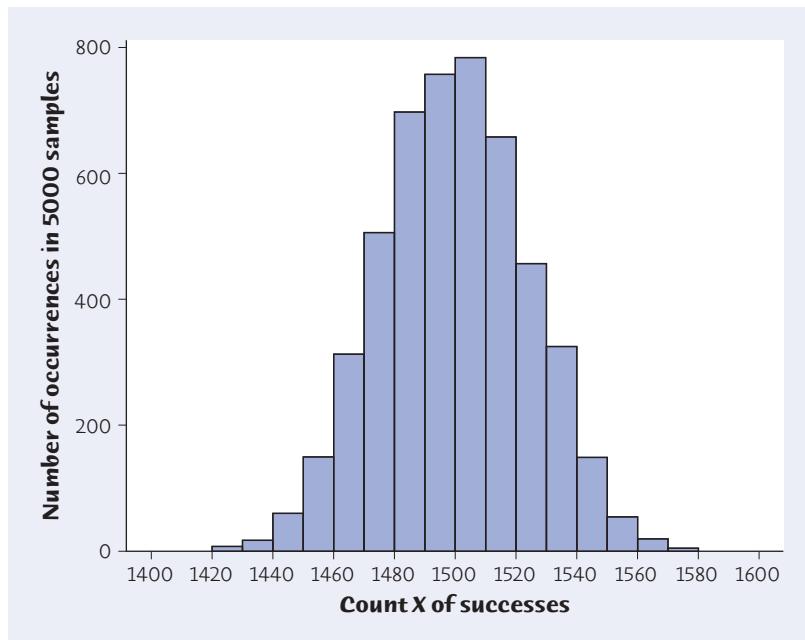
Binomial with $n = 2500$ and $p = 0.6$

X	$P(X \leq x)$
1519	0.786861

The probability we want is $1 - 0.786861 = 0.213139$, correct to 6 decimal places.

2. Simulate a large number of samples. Figure 13.4 displays a histogram of the counts X from 5000 samples of size 2500 when the truth about the population is $p = 0.6$. The simulated distribution, like the exact distribution in Figure 13.3, looks Normal. Because 1085 of these 5000 samples have X at least 1520, the probability estimated from the simulation is

$$P(X \geq 1520) = \frac{1085}{5000} = 0.2170$$

**FIGURE 13.4**

Histogram of 5000 simulated binomial counts ($n = 2500$ and $p = 0.6$).

This estimate misses the true probability by about 0.004. The law of large numbers says that the results of such simulations always get closer to the true probability as we simulate more and more samples.

3. Both of the previous methods require software. We can avoid the need for software by using the Normal approximation.

EXAMPLE 13.8 Normal calculation of a binomial probability

Act as though the count X in Example 13.7 has the Normal distribution with the same mean and standard deviation as the binomial distribution:

$$\begin{aligned}\mu &= np = (2500)(0.6) = 1500 \\ \sigma &= \sqrt{np(1-p)} = \sqrt{(2500)(0.6)(0.4)} = 24.49\end{aligned}$$

Standardizing X gives a standard Normal variable Z . The probability we want is

$$\begin{aligned}P(X \geq 1520) &= P\left(\frac{X - 1500}{24.49} \geq \frac{1520 - 1500}{24.49}\right) \\ &= P(Z \geq 0.82) \\ &= 1 - 0.7939 = 0.2061\end{aligned}$$

The Normal approximation 0.2061 misses the true probability calculated in Example 13.7 by about 0.007. ■

The *Normal Approximation to Binomial* applet shows in visual form how well the Normal approximation fits the binomial distribution for any n and p . You can slide n and watch the approximation get better. Whether or not the Normal approximation is satisfactory depends on how accurate your calculations need to be. For most statistical purposes, great accuracy is not required. Our rule of thumb for use of the Normal approximation reflects this judgment.



APPLY YOUR KNOWLEDGE

13.10 Using Benford's law. According to Benford's law (Example 10.7, page 269) the probability that the first digit of the amount of a randomly chosen invoice is a 1 or a 2 is 0.477. You examine 90 invoices from a vendor and find that 29 have first digits 1 or 2. If Benford's law holds, the count of 1s and 2s will have the binomial distribution with $n = 90$ and $p = 0.477$. Too few 1s and 2s suggests fraud. What is the approximate probability of 29 or fewer if the invoices follow Benford's law? Do you suspect that the invoice amounts are not genuine?

13.11 College admissions. A small liberal arts college in Ohio would like to have an entering class of 475 students next year. Past experience shows that about 31% of the students admitted will decide to attend. The college is planning to admit 1520 students. Suppose that students make their decisions independently and that the probability is 0.31 that a randomly chosen student will accept the offer of admission.

- What are the mean and standard deviation of the number of students who accept the admissions offer from this college?
- Use the Normal approximation: what is the approximate probability that the college gets more students than it wants?
- Use software to compute the exact probability that the college gets more students than it wants. How good is the approximation in part (b)?

13.12 Checking for survey errors. One way of checking the effect of undercoverage, nonresponse, and other sources of error in a sample survey is to compare the sample with known facts about the population. About 24% of the Canadian population over 15 years of age are first generation; that is, they were born outside Canada. The number X of first-generation Canadians in random samples of 1000 persons over 15 should therefore vary with the binomial ($n = 1000$, $p = 0.24$) distribution.

- What are the mean and standard deviation of X ?
- Use the Normal approximation to find the probability that the sample will contain between 210 and 270 first-generation Canadians. Be sure to check that you can safely use the approximation.

CHAPTER 13 SUMMARY

CHAPTER SPECIFICS

- A count X of successes has a **binomial distribution** in the **binomial setting**: there are n observations; the observations are independent of each other; each observation results in a success or a failure; each observation has the same probability p of a success.

- The binomial distribution with n observations and probability p of success gives a good approximation to the sampling distribution of the count of successes in an SRS of size n from a large population containing proportion p of successes.
- If X has the binomial distribution with parameters n and p , the possible values of X are the whole numbers $0, 1, 2, \dots, n$. The **binomial probability** that X takes any of these values is

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n - k}$$

In practice, binomial probabilities are best found using software.

- The **binomial coefficient**

$$\binom{n}{k} = \frac{n!}{k!(n - k)!}$$

counts the number of ways k successes can be arranged among n observations. Here the **factorial $n!$** is

$$n! = n \times (n - 1) \times (n - 2) \times \cdots \times 3 \times 2 \times 1$$

for positive whole numbers n , and $0! = 1$.

- The **mean** and **standard deviation** of a binomial count X are

$$\begin{aligned}\mu &= np \\ \sigma &= \sqrt{np(1 - p)}\end{aligned}$$

- The **Normal approximation** to the binomial distribution says that if X is a count having the binomial distribution with parameters n and p , then when n is large, X is approximately $N(np, \sqrt{np(1 - p)})$. Use this approximation only when $np \geq 10$ and $n(1 - p) \geq 10$.

LINK IT

The binomial distribution is used to compute probabilities for the count of successes among n observations that are produced under the binomial setting. An important situation for which the binomial setting can be used is when we choose an SRS from a population with a proportion p of successes. When the success probability p is known, probabilities associated with the number of successes among n observations can be computed using either the binomial formula or the Normal approximation when n and p are such that both the mean number of successes and failures are large enough.

An important application of the binomial distribution is in making inferences about the *proportion* of some outcome in a population. This is described in Chapter 20, where our data are collected under the binomial setting, but the proportion with the given outcome in the population is not known. For example, we may be interested in learning about the proportion of young adults (ages 19 to 25) who still live at home with their parents based on a sample from the population of young adults. When we want to make inferences about an unknown proportion in a population, we generally work with the *proportion of successes in the sample* rather than the count of successes in the sample. When using the proportion of successes in the sample to answer questions and draw conclusions about this unknown proportion in the population, the statistical methods are still related to the binomial distribution and the Normal approximation to the binomial.


CHECK YOUR SKILLS

13.13 James reads that 1 out of 4 eggs contains salmonella bacteria. So he never uses more than 3 eggs in cooking. If eggs do or don't contain salmonella independently of each other, the number of contaminated eggs when James uses 3 chosen at random has the distribution

- (a) binomial with $n = 4$ and $p = 1/4$.
- (b) binomial with $n = 3$ and $p = 1/4$.
- (c) binomial with $n = 3$ and $p = 1/3$.

13.14 In the previous exercise, the probability that at least 1 of James's 3 eggs contains salmonella is about

- (a) 0.68.
- (b) 0.58.
- (c) 0.30.

13.15 In a group of 10 college students, 4 are business majors. You choose 3 of the 10 students at random and ask their major. The distribution of the number of business majors you choose is

- (a) binomial with $n = 10$ and $p = 0.4$.
- (b) binomial with $n = 3$ and $p = 0.4$.
- (c) not binomial.

13.16 Virginia is the star player on her middle school basketball team. If she makes 3 free throws and misses 2 free throws during a game, in how many ways can you arrange the sequence of hits and misses?

$$(a) \binom{3}{2} = 3 \quad (b) \binom{5}{2} = 20 \quad (c) \binom{5}{3} = 10$$

13.17 Virginia makes 40% of her free throws. She takes 5 free throws in a game. If the shots are independent of each other, the probability that she misses the first and last shot but makes the other 3 is about

- (a) 0.230.
- (b) 0.115.
- (c) 0.023.

13.18 Virginia makes 40% of her free throws. She takes 5 free throws in a game. If the shots are independent of each other, the probability that she makes 3 out of the 5 shots is about

- (a) 0.230.
- (b) 0.115.
- (c) 0.023.

Each entry in a table of random digits like Table B has probability 0.1 of being any of the ten digits 0 to 9, and digits are independent of each other. Exercises 13.19 to 13.21 use this setting.

13.19 The probability of an entry being either an 8 or a 9 is

- (a) 0.1.
- (b) 0.2.
- (c) 0.4.

13.20 Each line in Table B has 40 digits. The number of times an 8 or a 9 occurs in two lines of the table has a

- (a) binomial distribution with $n = 80$ and $p = 0.2$.
- (b) binomial distribution with $n = 80$ and $p = 0.1$.
- (c) binomial distribution with $n = 40$ and $p = 0.2$.

13.21 The mean number of times an 8 or a 9 occurs in two lines of the table is

- (a) 16.
- (b) 12.8.
- (c) 8.


CHAPTER 13 EXERCISES

13.22 Binomial setting? In each situation below, is it reasonable to use a binomial distribution for the random variable X ? Give reasons for your answer in each case.

(a) An auto manufacturer chooses one car from each hour's production for a detailed quality inspection. One variable recorded is the count X of finish defects (dimples, ripples, etc.) in the car's paint.

(b) The pool of potential jurors for a murder case contains 100 persons chosen at random from the adult residents of a large city. Each person in the pool is asked whether he or she opposes the death penalty; X is the number who say "Yes."

(c) Joe buys a ticket in his state's Pick 3 lottery game every week; X is the number of times in a year that he wins a prize.

13.23 Binomial setting? A binomial distribution will be approximately correct as a model for one of these two sports settings and not for the other. Explain why by briefly discussing both settings.

(a) A National Football League kicker has made 80% of his field goal attempts in the past. This season he attempts 20 field goals. The attempts differ widely in distance, angle, wind, and so on.

(b) A National Basketball Association player has made 80% of his free-throw attempts in the past. This season he takes 150 free throws. Basketball free throws are always attempted from 15 feet away from the basket with no interference from other players.



David Bergman/CORBIS

13.24 Testing ESP. In a test for ESP (extrasensory perception), a subject is told that cards the experimenter can see but he cannot contain either a star, a circle, a wave, or a square. As the experimenter looks at each of 20 cards in turn, the subject

names the shape on the card. A subject who is just guessing has probability 0.25 of guessing correctly on each card.

- The count of correct guesses in 20 cards has a binomial distribution. What are n and p ?
- What is the mean number of correct guesses in many repetitions of the experiment?
- What is the probability of exactly 5 correct guesses?

13.25 Random stock prices. A believer in the random walk theory of stock markets thinks that an index of stock prices has probability 0.65 of increasing in any year. Moreover, the change in the index in any given year is not influenced by whether it rose or fell in earlier years. Let X be the number of years among the next 5 years in which the index rises.

- X has a binomial distribution. What are n and p ?
- What are the possible values that X can take?
- Find the probability of each value of X . Draw a probability histogram for the distribution of X . (See Figure 13.2 [page 339] for an example of a probability histogram.)
- What are the mean and standard deviation of this distribution? Mark the location of the mean on your histogram.

13.26 Betting on red. A roulette wheel has 38 slots, numbered 0, 00, and 1 to 36. The slots 0 and 00 are colored green, 18 of the others are red, and 18 are black. The dealer spins the wheel and at the same time rolls a small ball along the wheel in the opposite direction. The wheel is carefully balanced so that the ball is equally likely to land in any slot when the wheel slows. Gamblers can bet on various combinations of numbers and colors.

- If you bet on “red,” you win if the ball lands in a red slot. What is the probability of winning with a bet on red in a single play of roulette?
- You decide to play roulette 4 times, each time betting on red. What is the distribution of X , the number of times you win?
- If you bet the same amount on each play and win on exactly 2 of the 4 plays, you will “break even.” What is the probability that you will break even?
- If you win on fewer than 2 of the 4 plays, you will lose money. What is the probability that you will lose money?

13.27 The birth control shot

shot. The birth control shot is one of the most effective methods of birth control available, and it works best when you get the shot regularly, every 12 weeks. Under ideal conditions, only 1% of



Urbano Delvalle/Time Life Pictures/Getty Images

women getting the shot become pregnant within one year. In typical use, however, 3% become pregnant.⁴ Choose at random 20 women using the shot as their method of birth control. We count the number who become pregnant in the next year.

- Explain why this is a binomial setting.
- What is the probability that at least 1 of the women becomes pregnant under ideal conditions? What is the probability in typical use?

13.28 Betting on red, continued. You decide to play roulette 200 times, each time betting the same amount on red. You will lose money if you win on fewer than 100 of the plays. Based on the information in Exercise 13.26, what is the probability that you will lose money? (Check that the Normal approximation is permissible, and use it to find this probability. If your software allows, find the exact binomial probability and compare the two results.) In general, if you bet the same amount on red every time, you will lose money if you win on fewer than half of the plays. What do you think happens to the probability of making money the longer you continue to play?

13.29 The birth control shot, continued. A study of the effectiveness of the birth control shot interviews a random sample of 600 women who are using the shot as their method of birth control.

- Based on the information about typical use in Exercise 13.27, what is the probability that at least 20 of these women become pregnant in the next year? (Check that the Normal approximation is permissible, and use it to find this probability. If your software allows, find the exact binomial probability and compare the two results.)
- We can't use the Normal approximation to the binomial distribution to find this probability under ideal conditions as described in Exercise 13.27. Why not?

13.30 Hitting the fairway. One statistic used to assess professional golfers is driving accuracy, the percent of drives that land in the fairway. In 2009, driving accuracy for PGA Tour professionals ranged from about 50% to about 75%. Phil Mickelson, the third-highest money winner on the PGA Tour in 2009, hits the fairway only about 52% of the time.⁵ Phil is also one of the longest drivers on the tour, and increased distance is generally associated with decreased accuracy.

- Phil hits 14 drives in a round. What assumptions must you make in order to use a binomial distribution for the count X of fairways he hits? Which of these assumptions is least realistic?
- Assuming that a binomial distribution can be used, what is the expected number of fairways that Phil hits in a round in which he hits 14 drives?

13.31 Genetics. According to genetic theory, the blossom color in the second generation of a certain cross of sweet peas should be red or white in a 3:1 ratio. That is, each plant has probability $3/4$ of having red blossoms, and the blossom colors of separate plants are independent.

- (a) What is the probability that exactly 3 out of 4 of these plants have red blossoms?
- (b) What is the mean number of red-blossomed plants when 60 plants of this type are grown from seeds?
- (c) What is the probability of obtaining at least 45 red-blossomed plants when 60 plants are grown from seeds? Use the Normal approximation. If your software allows, find the exact binomial probability and compare the two results.



© blickwinkel/Alamy

13.32 False-positives in testing for HIV. A rapid test for the presence in the blood of antibodies to HIV, the virus that causes AIDS, gives a positive result with probability about 0.004 when a person who is free of HIV antibodies is tested. A clinic tests 1000 people who are all free of HIV antibodies.

- (a) What is the distribution of the number of positive tests?
- (b) What is the mean number of positive tests?
- (c) You cannot safely use the Normal approximation for this distribution. Explain why.

13.33 Chevrolet sales in 2010. Chevrolet sold 4.26 million vehicles globally in 2010, making it the only one among the top five global car companies to increase its market share. The five best-selling Chevrolet vehicles in 2010 were Silverado pickups, with approximately 435,000 sold; Cruze compact cars, with 335,000; Aveo compacts, with 322,000; Malibus, with 222,000; and Impalas, with 184,000.⁶ Chevrolet wants to undertake a survey of buyers of these five vehicle types to ask them about satisfaction with their purchases.

- (a) What proportion of the five best-selling vehicle types were Impalas?
- (b) If they plan to survey a total of 1000 customers, what is the expected number and standard deviation of the number of Impala buyers in the sample?
- (c) What is the probability that they will get more than 100 Impala buyers in their sample?

13.34 Retention rates in a weight-loss program. Americans spend over \$30 billion dollars on a variety of weight-loss

products and services. In a study of retention rates of those using the Platinum Program at Jenny Craig from May 2001 to May 2002, it was found that about 25% of those who began the program dropped out in the first four weeks.⁷ Assume that we have a random sample of 300 people who are beginning the program.

- (a) What is the mean number of people who will drop out of the Platinum Program in the first four weeks in a sample of this size? What is the standard deviation?
- (b) What is the approximate probability that at least 210 people in the sample will still be in the Platinum Program after the first four weeks?

13.35 Multiple-choice tests. Here is a simple probability model for multiple-choice tests. Suppose that each student has probability p of correctly answering a question chosen at random from a universe of possible questions. (A strong student has a higher p than a weak student.) Answers to different questions are independent.

- (a) Jodi is a good student for whom $p = 0.75$. Use the Normal approximation to find the probability that Jodi scores between 70% and 80% on a 100-question test.
- (b) If the test contains 250 questions, what is the probability that Jodi will score between 70% and 80%? You see that Jodi's score on the longer test is more likely to be close to her "true score."

13.36 Is this coin balanced? While he was a prisoner of war during World War II, John Kerrich tossed a coin 10,000 times. He got 5067 heads. If the coin is perfectly balanced, the probability of a head is 0.5. Is there reason to think that Kerrich's coin was not balanced? To answer this question, find the probability that tossing a balanced coin 10,000 times would give a count of heads at least this far from 5000 (that is, at least 5067 heads or no more than 4933 heads).

13.37 Binomial variation. Never forget that probability describes only what happens in the long run. Example 13.5 (page 336) concerns the count of bad CDs in inspection samples of size 10. The count has the binomial distribution with $n = 10$ and $p = 0.1$. The *Probability* applet simulates inspecting a lot of CDs if you set the probability of heads to 0.1, toss 10 times, and let each head stand for a bad CD.

- (a) The mean number of bad CDs in a sample is 1. Click "Toss" and "Reset" repeatedly to simulate 20 samples. How many bad CDs did you find in each sample? How close to the mean 1 is the average number of bad CDs in these samples?
- (b) Example 13.5 shows that the probability of exactly 1 bad CD is 0.3874. How close to the probability is the proportion of the 20 lots that have exactly 1 bad CD?

Whooping cough. Whooping cough (*pertussis*) is a highly contagious bacterial infection that was a major cause of childhood deaths before the development of vaccines. About 80% of unvaccinated children who are exposed to whooping cough will develop the infection, as opposed to only about 5% of vaccinated children. Exercises 13.38 to 13.41 are based on this information.

13.38 Vaccination at work. A group of 20 children at a nursery school are exposed to whooping cough by playing with an infected child.

- (a) If all 20 have been vaccinated, what is the mean number of new infections? What is the probability that no more than 2 of the 20 children develop infections?
- (b) If none of the 20 have been vaccinated, what is the mean number of new infections? What is the probability that 18 or more of the 20 children develop infections?

13.39 A whooping cough outbreak. In 2007, Bob Jones University in Greenville, South Carolina, ended its fall semester a week early because of a whooping cough outbreak; 158 students were isolated and another 1200 given antibiotics as a precaution.⁸ Authorities react strongly to whooping cough outbreaks because the disease is so contagious. Because the effect of childhood vaccination often wears off by late adolescence, treat the Bob Jones students as if they were unvaccinated. It appears that about 1400 students were exposed. What is the probability that at least 75% of these students develop infections if not treated? (Fortunately, whooping cough is much less serious after infancy.)

13.40 A mixed group: means. A group of 20 children at a nursery school are exposed to whooping cough by playing with an infected child. Of these children 17 have been vaccinated and 3 have not.

- (a) What is the distribution of the number of new infections among the 17 vaccinated children? What is the mean number of new infections?
- (b) What is the distribution of the number of new infections among the 3 unvaccinated children? What is the mean number of new infections?
- (c) Add your means from parts (a) and (b). This is the mean number of new infections among all 20 exposed children.

13.41 A mixed group: probabilities. We would like to find the probability that exactly 2 of the 20 exposed children in the previous exercise develop whooping cough.

- (a) One way to get 2 infections is to get 1 among the 17 vaccinated children and 1 among the 3 unvaccinated children. Find the probability of exactly 1 infection among the 17 vaccinated children. Find the probability of exactly 1 infection among the 3 unvaccinated children. These events are independent: what is the probability of exactly 1 infection in each group?

- (b) Write down all the ways in which 2 infections can be divided between the two groups of children. Follow the pattern of part (a) to find the probability of each of these possibilities. Add all of your results, including the result of part (a), to obtain the probability of exactly 2 infections among the 20 children.

13.42 Estimating π from random numbers. Kenyon College student Eric Newman used basic geometry to evaluate software random number generators as part of a summer research project. He generated 2000 independent random points (X, Y) in the unit square. That is, X and Y are independent random numbers between 0 and 1, each having the density function illustrated in Figure 10.4 (page 272). The probability that (X, Y) falls in any region within the unit square is the area of the region.⁹

- (a) Sketch the unit square, the region of possible values for the point (X, Y) .
- (b) The set of points (X, Y) where $X^2 + Y^2 < 1$ describes a circle of radius 1. Add this circle to your sketch in part (a), and label the intersection of the two regions A.
- (c) Let T be the total number of the 2000 points that fall into the region A. T follows a binomial distribution. Identify n and p . (Hint: Recall that the area of a circle is πr^2 .)
- (d) What are the mean and standard deviation of T ?
- (e) Explain how Eric used a random number generator and the facts above to estimate π .

13.43 The continuity correction. One reason why the Normal approximation may fail to give accurate estimates of binomial probabilities is that the binomial distributions are discrete and the Normal distributions are continuous. That is, counts take only whole-number values but Normal variables can take any value. We can improve the Normal approximation by treating each whole-number count as if it occupied the interval from 0.5 below the number to 0.5 above the number. For example, approximate a binomial probability $P(X \geq 10)$ by finding the Normal probability $P(X \geq 9.5)$. Be careful: binomial $P(X > 10)$ is approximated by Normal $P(X \geq 10.5)$.

We saw in Exercise 13.30 that Phil Mickelson hits the fairway in 52% of his drives. We will assume that his drives are independent and that each has probability 0.52 of hitting the fairway. Suppose Phil drives 24 times. The exact binomial probability that he hits 17 or more fairways is 0.0487.

- (a) Show that this setting satisfies the rule of thumb for use of the Normal approximation (just barely).
- (b) What is the Normal approximation to $P(X \geq 17)$?
- (c) What is the Normal approximation using the continuity correction? That's a lot closer to the true binomial probability.



EXPLORING THE WEB

13.44 MCAT writing sample. Go to the Web site aamc.org/students/applying/mcat/admissionsadvisors/mcat_stats/, which reports the distribution of the total scaled MCAT composite scores as well as the scores on the individual areas of assessment for recent years. The areas of assessment include four multiple-choice portions that are combined to give the overall composite score, plus a writing sample consisting of two essays. The essays are scored individually, with the results converted to an alphabetic score ranging from J (lowest) to T (highest). Most competitive medical schools look for a writing MCAT score of at least P or Q.

- (a) Open the pdf with the percentage and scaled score tables for the most recent year provided. Can the distribution of writing-sample scores be assumed to be approximately Normal? Explain.
- (b) A survey organization is planning on contacting an SRS of 1000 of the examinees from the most recent year to see how they prepared for the writing sample. What is the distribution of the number of examinees in the sample who had a score of at least P on the writing sample?
- (c) Use the Normal distribution to approximate the probability that at least half of the examinees in the sample scored a P or higher.
- (d) Use software or a calculator to compute the exact probability that at least half of the examinees in the sample scored a P or higher. How do the exact probability and the Normal approximation to this probability compare? Which of the two answers would you report to the survey organization? Why?

13.45 Binomial calculators. A number of Web sites will do exact binomial probability calculations for you.

- (a) Find a Web site with a binomial calculator and give its URL.
- (b) In Example 13.7, the number in the sample that agree that shopping is frustrating is a binomial random variable X with $n = 2500$ and $p = 0.6$. Use the binomial calculator to compute the probability that 1530 or more of the sample agree.





Confidence Intervals: The Basics

Chapter 14

After we have selected a sample, we know the responses of the individuals in the sample. The usual reason for taking a sample is not to learn about the individuals in the sample but to *infer* from the sample data some conclusion about the wider population that the sample represents.

STATISTICAL INFERENCE

Statistical inference provides methods for drawing conclusions about a population from sample data.

Because a different sample might lead to different conclusions, we can't be certain that our conclusions are correct. Statistical inference uses the language of probability to say how trustworthy our conclusions are. This chapter introduces one of the two most common types of inference, *confidence intervals* for estimating the value of a population parameter. The next chapter discusses the other common type of inference, *tests of significance* for assessing the evidence for a claim about a population. Both types of inference are based on the sampling distributions of statistics. That is, both use probability to say what would happen if we applied the inference method many times.

This chapter presents the basic reasoning of statistical inference. To make the reasoning as clear as possible, we start with a setting that is too simple to be realistic. Here is the setting for our work in this chapter.

IN THIS CHAPTER WE COVER...

- The reasoning of statistical estimation
- Margin of error and confidence level
- Confidence intervals for a population mean
- How confidence intervals behave



SIMPLE CONDITIONS FOR INFERENCE ABOUT A MEAN

1. We have an SRS from the population of interest. There is no nonresponse or other practical difficulty.
2. The variable we measure has an exactly Normal distribution $N(\mu, \sigma)$ in the population.
3. We don't know the population mean μ . But we do know the population standard deviation σ .

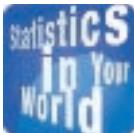


The conditions that we have a perfect SRS, that the population is exactly Normal, and that we know the population σ are all unrealistic. Chapter 16 begins to move from the “simple conditions” toward the reality of statistical practice. Later chapters deal with inference in fully realistic settings.

If these “simple conditions” are unrealistic, why study them? One reason is that under these simple conditions we can apply what we have learned in previous chapters about the Normal distribution and the sampling distribution of a sample mean to develop, step-by-step, methods for inference about a mean. The reasoning used under simple conditions applies to more realistic settings where the mathematics is more complicated.

Another reason for studying inference under these simple conditions is that we can carry out calculations using what we have already learned about the Normal distribution in previous chapters. This includes the calculations used to plan statistical studies, which we discuss in Chapter 16. Unfortunately, under more realistic conditions, such calculations are more complicated, and the connection to previous material about the Normal distribution is less clear.

Although we never know whether a population is exactly Normal and we never know the population σ , the methods we discuss in this and the next two chapters are approximately correct for sufficiently large sample sizes, provided we treat the sample standard deviation as though it were the population σ . Thus, there are situations (admittedly rare) where these methods can be used in practice.



Ranges are for statistics?

Many people like to think that statistical estimates are exact. The Nobel Prize-winning economist Daniel McFadden tells a story of his time on the Council of Economic Advisers. Presented with a range of forecasts for economic growth, President Lyndon Johnson replied: “Ranges are for cattle; give me one number.”



Thomas Northcut/Getty

THE REASONING OF STATISTICAL ESTIMATION

Body mass index (BMI) is used to screen for possible weight problems. It is calculated as weight divided by the square of height, measuring weight in kilograms and height in meters. Many online BMI calculators allow you to enter weight in pounds and height in inches. Adults with BMI less than 18.5 are considered underweight, and those with BMI greater than 25 may be overweight. For data about BMI, we turn to the National Health and Nutrition Examination Survey (NHANES), a continuing government sample survey that monitors the health of the American population.

EXAMPLE 14.1 Body mass index of young women

An NHANES report gives data for 654 women aged 20 to 29 years.¹ The mean BMI of these 654 women was $\bar{x} = 26.8$. On the basis of this sample, we want to estimate the mean BMI μ in the population of all 20.6 million women in this age group.

To match the “simple conditions,” we will treat the NHANES sample as an SRS from a Normal population with standard deviation $\sigma = 7.5$. ■

Here is the reasoning of statistical estimation in a nutshell:

1. To estimate the unknown population mean BMI μ , use the mean $\bar{x} = 26.8$ of the random sample. We don’t expect \bar{x} to be exactly equal to μ , so we want to say how accurate this estimate is.
2. We know the sampling distribution of \bar{x} . In repeated samples, \bar{x} has the Normal distribution with mean μ and standard deviation σ/\sqrt{n} . So the average BMI \bar{x} of an SRS of 654 young women has standard deviation

$$\frac{\sigma}{\sqrt{n}} = \frac{7.5}{\sqrt{654}} = 0.3 \text{ (rounded off)}$$

3. The 95 part of the 68–95–99.7 rule for Normal distributions says that \bar{x} is within 0.6 (that’s two standard deviations) of the mean μ in 95% of all samples. That is, for 95% of *all* samples of size 654, the distance between the sample mean \bar{x} and the population mean μ is less than 0.6. So if we estimate that μ lies somewhere in the interval from $\bar{x} - 0.6$ to $\bar{x} + 0.6$, we’ll be right for 95% of all possible samples. For this particular sample, this interval is

$$\bar{x} - 0.6 = 26.8 - 0.6 = 26.2$$

to

$$\bar{x} + 0.6 = 26.8 + 0.6 = 27.4$$

4. Because we got the interval 26.2 to 27.4 from a method that captures the population mean for 95% of all possible samples, we say that we are *95% confident* that the mean BMI μ of all young women is some value in that interval, no lower than 26.2 and no higher than 27.4.

The big idea is that the sampling distribution of \bar{x} tells us how close to μ the sample mean \bar{x} is likely to be. Statistical estimation just turns that information around to say how close to \bar{x} the unknown population mean μ is likely to be. We call the interval of numbers between the values $\bar{x} \pm 0.6$ a *95% confidence interval* for μ .

APPLY YOUR KNOWLEDGE

- 14.1 Number skills of high school seniors.** The National Assessment of Educational Progress (NAEP) includes a mathematics test for high school seniors.² Scores on the test range from 0 to 300. Demonstrating the ability to use the Pythagorean theorem to determine the length of a hypotenuse is an example of the skills and knowledge associated with performance at the Basic level. An example of the knowledge and skills associated with the Proficient level is using trigonometric ratios to determine length.

In 2009, 51,000 12th-graders were in the NAEP sample for the mathematics test. The mean mathematics score was $\bar{x} = 153$. We want to estimate the mean

score μ in the population of all 12th-graders. Consider the NAEP sample as an SRS from a Normal population with standard deviation $\sigma = 34$.

- If we take many samples, the sample mean \bar{x} varies from sample to sample according to a Normal distribution with mean equal to the unknown mean score μ in the population. What is the standard deviation of this sampling distribution?
- According to the 95 part of the 68–95–99.7 rule, 95% of all values of \bar{x} fall within _____ on either side of the unknown mean μ . What is the missing number?
- What is the 95% confidence interval for the population mean score μ based on this one sample?

14.2 Retaking the SAT. An SRS of 400 high school seniors gained an average of $\bar{x} = 22$ points in their second attempt at the SAT Mathematics exam. Assume that the change in score has a Normal distribution with standard deviation $\sigma = 50$. We want to estimate the mean change in score μ in the population of all high school seniors. Give a 95% confidence interval for μ based on this sample.

MARGIN OF ERROR AND CONFIDENCE LEVEL

The 95% confidence interval for the mean BMI of young women, based on the NHANES sample, is $\bar{x} \pm 0.6$. Once we have the sample results in hand, we know that for this sample $\bar{x} = 26.8$, so that our confidence interval is 26.8 ± 0.6 . Most confidence intervals have a form similar to this,

$$\text{estimate} \pm \text{margin of error}$$

The estimate ($\bar{x} = 26.8$ in our example) is our guess for the value of the unknown parameter. The **margin of error** ± 0.6 shows how accurate we believe our guess is, based on the variability of the estimate. We have a 95% confidence interval because the interval $\bar{x} \pm 0.6$ catches the unknown parameter in 95% of all possible samples.

CONFIDENCE INTERVAL

A level C confidence interval for a parameter has two parts:

- An interval calculated from the data, usually of the form

$$\text{estimate} \pm \text{margin of error}$$

- A confidence level C, which gives the probability that the interval will capture the true parameter value in repeated samples. That is, the confidence level is the success rate for the method.

Users can choose the confidence level, usually 90% or higher because we usually want to be quite sure of our conclusions. The most common confidence level is 95%.

INTERPRETING A CONFIDENCE LEVEL

The confidence level is the success rate of the method that produces the interval. We don't know whether the 95% confidence interval from a particular sample is one of the 95% that capture μ or one of the unlucky 5% that miss.

To say that we are **95% confident** that the unknown μ lies between 26.2 and 27.4 is shorthand for “We got these numbers using a method that gives correct results **95% of the time.**”

EXAMPLE 14.2 Statistical estimation in pictures

Figures 14.1 and 14.2 illustrate the behavior of confidence intervals. Study these figures carefully. If you understand what they say, you have mastered one of the big ideas of statistics.

Figure 14.1 illustrates the behavior of the interval $\bar{x} \pm 0.6$ for the mean BMI of young women. Starting with the population, imagine taking many SRSs of 654 young women. The first sample has $\bar{x} = 26.8$, the second has $\bar{x} = 27.0$, the third has $\bar{x} = 26.2$, and so on. The sample mean varies from sample to sample, but when we use the formula $\bar{x} \pm 0.6$ to get an interval based on each sample, 95% of these intervals capture the unknown population mean μ .

Figure 14.2 illustrates the idea of a 95% confidence interval in a different form. It shows the result of drawing many SRSs from the same population and calculating a 95% confidence interval from each sample. The center of each interval is at \bar{x} and therefore varies from sample to sample. The sampling distribution of \bar{x} appears at the top of the figure to show the long-term pattern of this variation. The population mean μ is at the center of the sampling distribution. The 95% confidence intervals from 25 SRSs appear underneath. The center \bar{x} of each interval is marked by a dot. The arrows on either side of the dot span the confidence interval. All except 1 of these 25 intervals capture the true value of μ . If we take a very large number of samples, 95% of the confidence intervals will contain μ . ■

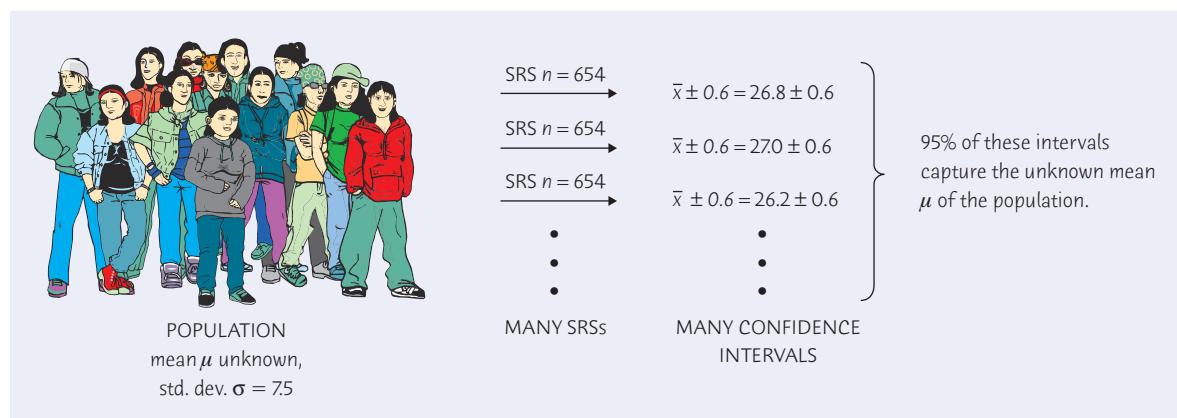
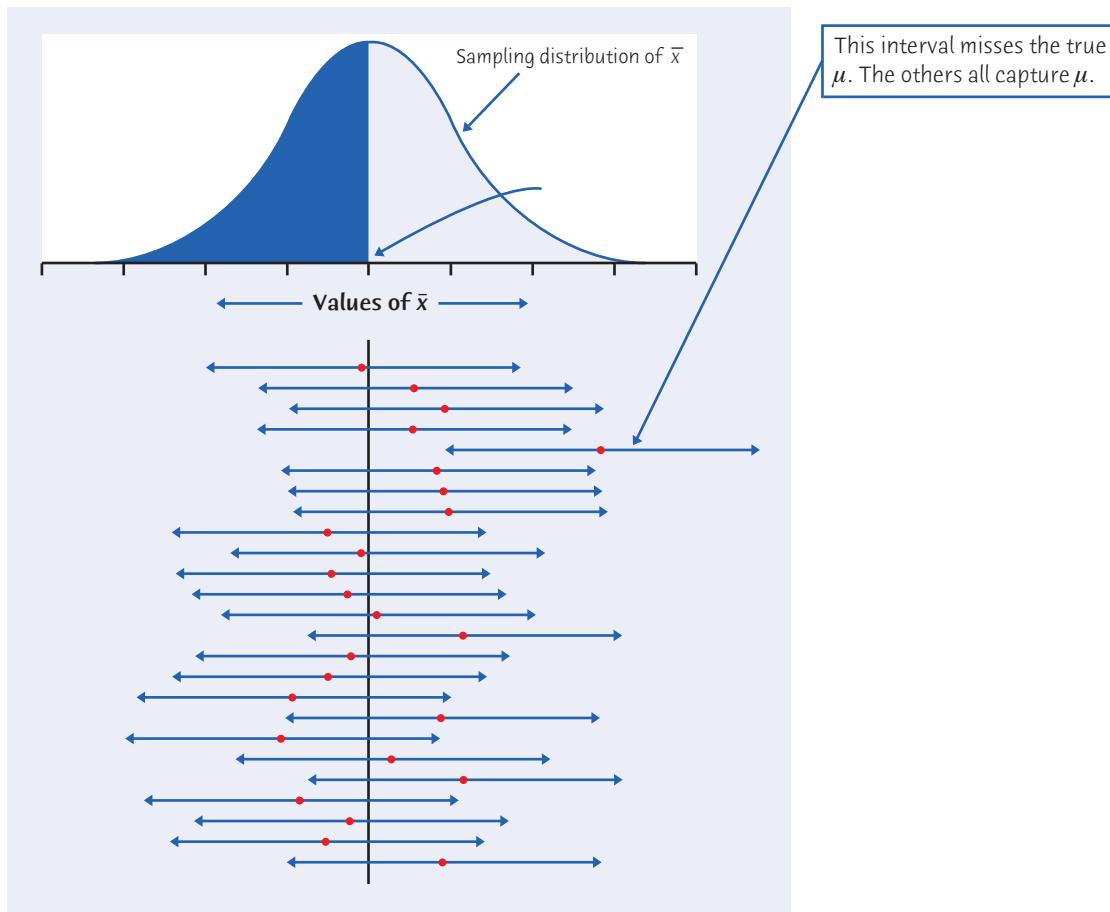


FIGURE 14.1

To say that $\bar{x} \pm 0.6$ is a 95% confidence interval for the population mean μ is to say that, in repeated samples, 95% of these intervals capture μ .

**FIGURE 14.2**

Twenty-five samples from the same population gave these 95% confidence intervals. In the long run, 95% of all samples give an interval that contains the population mean μ .



The *Confidence Interval* applet animates Figure 14.2. You can use the applet to watch confidence intervals from one sample after another capture or fail to capture the true parameter.

APPLY YOUR KNOWLEDGE



14.3 Confidence intervals in action. The idea of an 80% confidence interval is that in 80% of all samples the method produces an interval that captures the true parameter value. That's not high enough confidence for practical use, but 80% hits and 20% misses make it easy to see how a confidence interval behaves in repeated samples from the same population. Go to the *Confidence Interval* applet.

- Set the confidence level to 80%. Click “Sample” to choose an SRS and calculate the confidence interval. Do this 10 times to simulate 10 SRSs with their 10 confidence intervals. How many of the 10 intervals captured the true mean μ ? How many missed?

- (b) You see that we can't predict whether the next sample will hit or miss. The confidence level, however, tells us what percent will hit in the long run. Reset the applet and click "Sample 50" to get the confidence intervals from 50 SRSs. How many hit?
- (c) Click "Sample 50" repeatedly and write down the number of hits each time. What was the percent of hits among 100, 200, 300, 400, 500, 600, 700, 800, and 1000 SRSs? Even 1000 samples is not truly "the long run," but we expect the percent of hits in 1000 samples to be fairly close to the confidence level, 80%.

- 14.4 Losing weight.** A Gallup Poll in November 2010 found that 54% of the people in the sample said they want to lose weight. Gallup announced, "For results based on the total sample of national adults, one can say with 95% confidence that the maximum margin of sampling error is ± 4 percentage points."
- (a) What is the 95% confidence interval for the percent of all adults who want to lose weight?
 - (b) What does it mean to have 95% confidence in this interval?
-

CONFIDENCE INTERVALS FOR A POPULATION MEAN

In the setting of Example 14.1 we outlined the reasoning that leads to a 95% confidence interval for the unknown mean μ of a population. Now we will reduce the reasoning to a formula.

To find a 95% confidence interval for the mean BMI of young women, we first caught the central 95% of the Normal sampling distribution by going out two standard deviations in both directions from the mean. To find a level C confidence interval, we first catch the central area C under the Normal sampling distribution. Because all Normal distributions are the same in the standard scale, we can obtain everything we need from the standard Normal curve.

Figure 14.3 shows how the central area C under a standard Normal curve is marked off by two points z^* and $-z^*$. Numbers like z^* that mark off specified areas are called **critical values** of the standard Normal distribution. Values of z^*

critical value

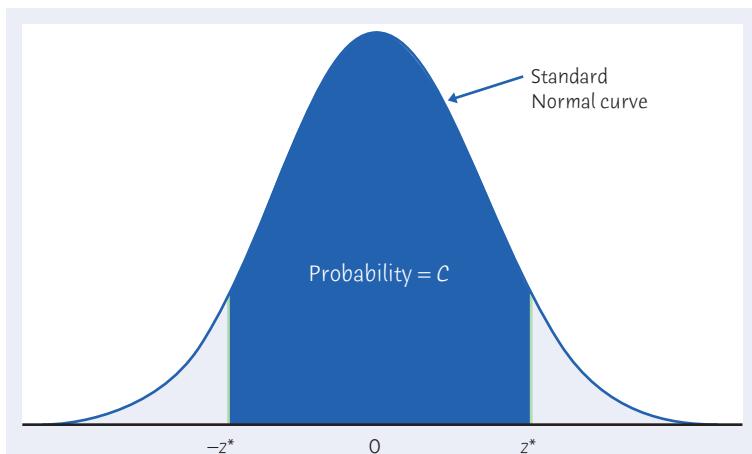


FIGURE 14.3

The critical value z^* is the number that catches central probability C under a standard Normal curve between $-z^*$ and z^* .

for many choices of C appear at the bottom of Table C in the back of the book, in the row labeled z^* . Here are the entries for the most common confidence levels:

Confidence level C	90%	95%	99%
Critical value z^*	1.645	1.960	2.576

You see that for $C = 95\%$ the table gives $z^* = 1.960$. This is a bit more precise than the approximate value $z^* = 2$ based on the 68–95–99.7 rule. You can of course use software to find critical values z^* , as well as the entire confidence interval.

Figure 14.3 shows that there is area C under the standard Normal curve between $-z^*$ and z^* . So any Normal curve has area C within z^* standard deviations on either side of its mean. The Normal sampling distribution of \bar{x} has area C within $z^* \sigma/\sqrt{n}$ on either side of the population mean μ because it has mean μ and standard deviation σ/\sqrt{n} . If we start at \bar{x} and go out $z^* \sigma/\sqrt{n}$ in both directions, we get an interval that contains the population mean μ in a proportion C of all samples. This interval is

$$\text{from } \bar{x} - z^* \frac{\sigma}{\sqrt{n}} \text{ to } \bar{x} + z^* \frac{\sigma}{\sqrt{n}}$$

or

$$\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}}$$

It is a level C confidence interval for μ .

CONFIDENCE INTERVAL FOR THE MEAN OF A NORMAL POPULATION

Draw an SRS of size n from a Normal population having unknown mean μ and known standard deviation σ . A level C confidence interval for μ is

$$\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}}$$

The critical value z^* is illustrated in Figure 14.3 and found at the bottom of Table C.

The steps in finding a confidence interval mirror the overall four-step process for organizing statistical problems.

CONFIDENCE INTERVALS: The Four-Step Process

STATE: What is the practical question that requires estimating a parameter?

PLAN: Identify the parameter, choose a level of confidence, and select the type of confidence interval that fits your situation.

SOLVE: Carry out the work in two phases:

1. Check the conditions for the interval you plan to use.
2. Calculate the confidence interval.

CONCLUDE: Return to the practical question to describe your results in this setting.

EXAMPLE 14.3 Good weather, good tips?

STATE: Does the expectation of good weather lead to more generous behavior? Psychologists studied the size of the tip in a restaurant when a message indicating that the next day's weather would be good was written on the bill. Here are tips from 20 patrons, measured in percent of the total bill:³

20.8	18.7	19.9	20.6	21.9	23.4	22.8	24.9	22.2	20.3
24.9	22.3	27.0	20.4	22.2	24.0	21.1	22.1	22.0	22.7

This is one of three sets of measurements made, the others being tips received when the message on the bill said that the next day's weather would not be good or there was no message on the bill. We want to estimate the mean tip for comparison with tips under the other conditions.

PLAN: We will estimate the mean percentage tip μ for all patrons of this restaurant when they receive a message on their bill indicating that the next day's weather will be good by giving a 95% confidence interval. The confidence interval just introduced fits this situation.

SOLVE: We should start by checking the conditions for inference. For this example, we will first find the interval and then discuss how statistical practice deals with conditions that are never perfectly satisfied.

The mean percentage tip of the sample is $\bar{x} = 22.21$. As part of the “simple conditions,” suppose that from past experience with patrons of this restaurant we know that the standard deviation of percentage tip is $\sigma = 2$. For 95% confidence, the critical value is $z^* = 1.960$. A 95% confidence interval for μ is therefore

$$\begin{aligned}\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}} &= 22.21 \pm 1.960 \frac{2}{\sqrt{20}} \\ &= 22.21 \pm 0.88 \\ &= 21.33 \text{ to } 23.09\end{aligned}$$

CONCLUDE: We are 95% confident that the mean percentage tip for all patrons of this restaurant when their bill contains a message that the next day's weather will be good is between 21.33 and 23.09. ■

In practice, the first part of the “Solve” step is to check the conditions for inference. The “simple conditions” are as follows:

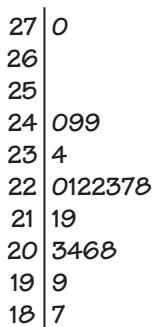
- 1. SRS:** We don't have an actual SRS from the population of all patrons of this restaurant. Scientists often act as if subjects are SRSs if there is nothing special about how the subjects were obtained. But it is always better to have an actual SRS because otherwise we can never be sure that hidden biases aren't present. This study was actually a randomized comparative experiment in which these 20 patrons were assigned at random from a larger group of patrons to get one of the treatments being compared.



TIPPING2



Ariel Skelley/AgeFotostock

**FIGURE 14.4**

Stemplot of the percentage tips in Example 14.3.

2. **Normal distribution:** The psychologists expect from past experience that measurements like this on patrons of the same restaurant under the same conditions will follow approximately a Normal distribution. We can't look at the population, but we can examine the sample. Figure 14.4 is a stemplot. The shape is roughly bell-shaped, with perhaps a modest outlier but no strong skewness. Shapes like this often occur in small samples from Normal populations, so we have no reason to doubt that the population distribution is Normal.
3. **Known σ :** It really is unrealistic to suppose that we know that $\sigma = 2$. We will see in Chapter 18 that it is easy to do away with the need to know σ .

As this discussion suggests, inference methods are often used when conditions like SRS and Normal population are not exactly satisfied. In this introductory chapter, we act as though the “simple conditions” are satisfied. In reality, wise use of inference requires judgment. Chapter 16 and the later chapters on each inference method will give you a better basis for judgment.

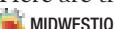
APPLY YOUR KNOWLEDGE

- 14.5 Find a critical value.** The critical value z^* for confidence level 85% is not in Table C. Use software or Table A of standard Normal probabilities to find z^* . Include in your answer a sketch like Figure 14.3 with $C = 0.85$ and your critical value z^* marked on the axis.

-  **14.6 Measuring conductivity.** The National Institute of Standards and Technology (NIST) supplies “standard materials” whose physical properties are supposed to be known. For example, you can buy from NIST an iron rod whose electrical conductivity is supposed to be 10.1 at 293 kelvin. (The units for conductivity are microsiemens per centimeter. Distilled water has conductivity 0.5.) Of course, no measurement is exactly correct. NIST knows the variability of its measurements very well, so it is quite realistic to assume that the population of all measurements of the same rod has the Normal distribution with mean μ equal to the true conductivity and standard deviation $\sigma = 0.1$. Here are 6 measurements on the same standard iron rod, which is supposed to have conductivity 10.1: 

$$10.08 \quad 9.89 \quad 10.05 \quad 10.16 \quad 10.21 \quad 10.11$$

NIST wants to give the buyer of this iron rod a 90% confidence interval for its true conductivity. What is this interval? Follow the four-step process as illustrated in Example 14.3.

- 14.7 IQ test scores.** Here are the IQ test scores of 31 seventh-grade girls in a Midwest school district.⁴ 

114	100	104	89	102	91	114	114	103	105
108	130	120	132	111	128	118	119	86	72
111	103	74	112	107	103	98	96	112	112



- (a) These 31 girls are an SRS of all seventh-grade girls in the school district. Suppose that the standard deviation of IQ scores in this population is known to be $\sigma = 15$. We expect the distribution of IQ scores to be close to Normal. Make a stemplot of the distribution of these 31 scores (split the stems) to verify that there are no major departures from Normality. You have now checked the “simple conditions” to the extent possible.
- (b) Estimate the mean IQ score for all seventh-grade girls in the school district, using a 99% confidence interval. Follow the four-step process as illustrated in Example 14.3.

HOW CONFIDENCE INTERVALS BEHAVE

The z confidence interval $\bar{x} \pm z^* \sigma / \sqrt{n}$ for the mean of a Normal population illustrates several important properties that are shared by all confidence intervals in common use. The user chooses the confidence level, and the margin of error follows from this choice. We would like high confidence and also a small margin of error. High confidence says that our method almost always gives correct answers. A small margin of error says that we have pinned down the parameter quite precisely. The factors that influence the margin of error of the z confidence interval are typical of most confidence intervals.

How do we get a small margin of error? The margin of error for the z confidence interval is

$$\text{margin of error} = z^* \frac{\sigma}{\sqrt{n}}$$

This expression has z^* and σ in the numerator and \sqrt{n} in the denominator. Therefore, the margin of error gets smaller when

- z^* gets smaller. Smaller z^* is the same as lower confidence level C (look again at Figure 14.3 on page 357). *There is a trade-off between the confidence level and the margin of error. To obtain a smaller margin of error from the same data, you must be willing to accept lower confidence.* 
- σ is smaller. The standard deviation σ measures the variation in the population. You can think of the variation among individuals in the population as noise that obscures the average value μ . It is easier to pin down μ when σ is small.
- n gets larger. Increasing the sample size n reduces the margin of error for any confidence level. Larger samples thus allow more precise estimates. *However, because n appears under a square root sign, we must take four times as many observations to cut the margin of error in half.* 

EXAMPLE 14.4 Changing the margin of error

In Example 14.3, psychologists recorded the size of the tip of 20 patrons in a restaurant when a message indicating that the next day's weather would be good was written on their bill. The data gave the mean size of the tip, as a percentage of the total bill, as

$\bar{x} = 22.21$, and we know that $\sigma = 2$. The 95% confidence interval for the mean percentage tip for all patrons of the restaurant when their bill contains a message that the next day's weather will be good is

$$\begin{aligned}\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}} &= 22.21 \pm 1.960 \frac{2}{\sqrt{20}} \\ &= 22.21 \pm 0.88\end{aligned}$$

The 90% confidence interval based on the same data replaces the 95% critical value $z^* = 1.960$ by the 90% critical value $z^* = 1.645$. This interval is

$$\begin{aligned}\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}} &= 22.21 \pm 1.645 \frac{2}{\sqrt{20}} \\ &= 22.21 \pm 0.74\end{aligned}$$

Lower confidence results in a smaller margin of error, ± 0.74 in place of ± 0.88 . You can calculate that the margin of error for 99% confidence is larger, ± 1.15 . Figure 14.5 compares these three confidence intervals.

If we had a sample of only 10 patrons, you can check that the margin of error for 95% confidence increases from ± 0.88 to ± 1.24 . Cutting the sample size in half does not double the margin of error, because the sample size n appears under a square root sign. ■

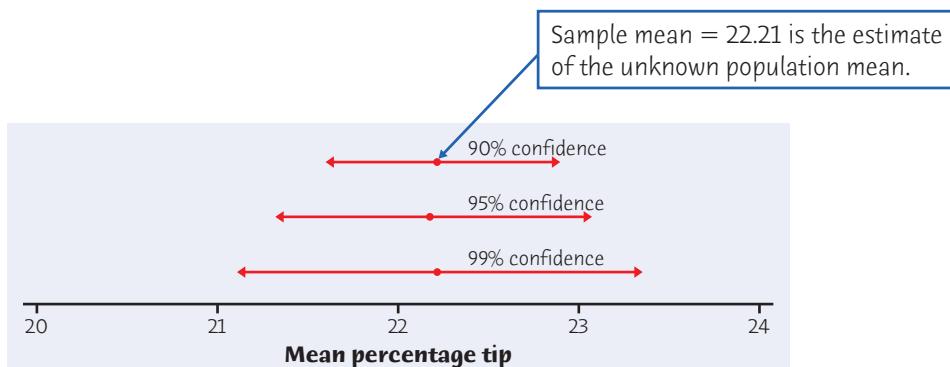


FIGURE 14.5

The lengths of three confidence intervals for Example 14.4. All three are centered at the estimate $\bar{x} = 22.21$. When the data and the sample size remain the same, higher confidence results in a larger margin of error.



APPLY YOUR KNOWLEDGE

14.8 Confidence level and margin of error. Example 14.1 described NHANES survey data on the body mass index (BMI) of 654 young women. The mean BMI in the sample was $\bar{x} = 26.8$. We treated these data as an SRS from a Normally distributed population with standard deviation $\sigma = 7.5$.

- (a) Give three confidence intervals for the mean BMI μ in this population, using 90%, 95%, and 99% confidence.

- (b) What are the margins of error for 90%, 95%, and 99% confidence? How does increasing the confidence level change the margin of error of a confidence interval when the sample size and population standard deviation remain the same?

14.9 Sample size and margin of error. Example 14.1 described NHANES survey data on the body mass index (BMI) of 654 young women. The mean BMI in the sample was $\bar{x} = 26.8$. We treated these data as an SRS from a Normally distributed population with standard deviation $\sigma = 7.5$.

- Suppose that we had an SRS of just 100 young women. What would be the margin of error for 95% confidence?
- Find the margins of error for 95% confidence based on SRSs of 400 young women and 1600 young women.
- Compare the three margins of error. How does increasing the sample size change the margin of error of a confidence interval when the confidence level and population standard deviation remain the same?

14.10 Retaking the SAT. In Exercise 14.2 we saw that an SRS of 400 high school seniors gained an average of $\bar{x} = 22$ points in their second attempt at the SAT Mathematics exam. Assuming that the change in score has a Normal distribution with standard deviation $\sigma = 50$, we computed a 95% confidence interval for the mean change in score μ in the population of all high school seniors.

- Find a 90% confidence interval for μ based on this sample.
- What is the margin of error for 90%? How does decreasing the confidence level change the margin of error of a confidence interval when the sample size and population standard deviation remain the same?
- Suppose we had an SRS of just 100 high school seniors. What would be the margin of error for 95% confidence?
- How does decreasing the sample size change the margin of error of a confidence interval when the confidence level and population standard deviation remain the same?

CHAPTER 14 SUMMARY

CHAPTER SPECIFICS

- A **confidence interval** uses sample data to estimate an unknown population parameter with an indication of how accurate the estimate is and of how confident we are that the result is correct.
- Any confidence interval has two parts: an interval calculated from the data and a confidence level C . The **confidence interval** often has the form

$$\text{estimate} \pm \text{margin of error}$$

- The **confidence level** is the success rate of the method that produces the interval. That is, C is the probability that the method will give a correct answer. If you use 95% confidence intervals often, in the long run 95% of your intervals will contain the true parameter value. You do not know whether or not a 95% confidence interval calculated from a particular set of data contains the true parameter value.

- A level C confidence interval for the mean μ of a Normal population with known standard deviation σ , based on an SRS of size n , is given by

$$\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}}$$

- The critical value z^* is chosen so that the standard Normal curve has area C between $-z^*$ and z^* .
- Other things being equal, the margin of error of a confidence interval gets smaller as
 - the confidence level C decreases,
 - the population standard deviation σ decreases,
 - the sample size n increases.

LINK IT

The reason we collect data is not to learn about the individuals that we observed but to infer from the data to some wider population that the individuals represent. Chapters 8 and 9 tell us that the way we produce the data (sampling, experimental design) affects whether we have a good basis for generalizing to some wider population. Chapters 10, 11, 12, and 13 discuss probability, the formal mathematical tool that determines the nature of the inferences we make. In particular, Chapter 11 discusses sampling distributions, which tell us how repeated SRSs behave and hence what a statistic (in particular, a sample mean) computed from our sample is likely to tell us about the corresponding parameter of the population (in particular, a population mean) from which the sample was selected.

In this chapter we discuss the basic reasoning of statistical estimation, with emphasis on estimating a population mean. To an estimate of the population mean we attach a margin of error and a confidence level. The result is a confidence interval. The sampling distribution of the sample mean, discussed in Chapter 11, provides the mathematical basis for constructing confidence intervals and understanding their properties. Although we apply the reasoning of statistical estimation in a simple and artificial setting (we assume we know the population standard deviation), we will use the same logic in future chapters to construct confidence intervals for population parameters in more realistic settings.

CHECK YOUR SKILLS

- 14.11** To give a 99.9% confidence interval for a population mean μ , you would use the critical value

- (a) $z^* = 1.960$. (b) $z^* = 2.576$. (c) $z^* = 3.291$.

Use the following information for Exercises 14.12 through 14.14. A laboratory scale is known to have a standard deviation of $\sigma = 0.001$ gram in repeated weighings. Scale readings in repeated weighings are Normally distributed, with mean equal to the true

weight of the specimen. Three weighings of a specimen on this scale give 3.412, 3.416, and 3.414 grams.



- 14.12** A 95% confidence interval for the true weight of this specimen is

- (a) 3.414 ± 0.00113 .
 (b) 3.414 ± 0.00065 .
 (c) 3.414 ± 0.00196 .

- 14.13** You want a 99% confidence interval for the true weight of this specimen. The margin of error for this interval will be
 (a) smaller than the margin of error for 95% confidence.
 (b) greater than the margin of error for 95% confidence.
 (c) about the same as the margin of error for 95% confidence.

14.14 Another specimen is weighed 8 times on this scale. The average weight is 4.1602 grams. A 99% confidence interval for the true weight of this specimen is

- (a) 4.1602 ± 0.00032 . (b) 4.1602 ± 0.00069 .
 (c) 4.1602 ± 0.00091 .

Use the following information for Exercises 14.15 through 14.18. The National Assessment of Educational Progress (NAEP) includes a mathematics test for high school seniors. Scores on the test range from 0 to 300. Suppose that you give the NAEP test to an SRS of 900 12th-graders from a large population in which the scores have mean $\mu = 150$ and standard deviation $\sigma = 35$. The mean \bar{x} will vary if you take repeated samples.

14.15 The sampling distribution of \bar{x} is approximately Normal. It has mean $\mu = 150$. What is its standard deviation?

- (a) 35. (b) 1.167. (c) 0.039.

14.16 Suppose that an SRS of 900 12th-graders has $\bar{x} = 148$. Based on this sample, a 95% confidence interval for μ is
 (a) 2.29. (b) 148 ± 2.29 . (c) 150 ± 2.29 .

14.17 In the previous exercise, suppose that we computed a 99% confidence interval for μ .

- (a) This 99% confidence interval would have a smaller margin of error than the 95% confidence interval.
 (b) This 99% confidence interval would have a larger margin of error than the 95% confidence interval.
 (c) This 99% confidence interval could have either a smaller or a larger margin of error than the 95% confidence interval. This varies from sample to sample.

14.18 Suppose that we took an SRS of 1600 12th-graders and found $\bar{x} = 148$. Compared with an SRS of 900 12th-graders, the margin of error for a 95% confidence interval for μ is

- (a) smaller. (b) larger.
 (c) either smaller or larger, but we can't say which.

CHAPTER 14 EXERCISES

14.19 Student study times. A class survey in a large class for first-year college students asked, “About how many minutes do you study on a typical weeknight?” The mean response of the 463 students was $\bar{x} = 118$ minutes. Suppose that we know that the study time follows a Normal distribution with standard deviation $\sigma = 65$ minutes in the population of all first-year students at this university.

- (a) Use the survey result to give a 99% confidence interval for the mean study time of all first-year students.
 (b) What condition not yet mentioned must be met for your confidence interval to be valid?

14.20 I want more muscle.

Young men in North America and Europe (but not in Asia) tend to think they need more muscle to be attractive. One study presented 200 young American men with 100 images of men with various levels of muscle.⁵ Researchers measure level of muscle in kilograms per square meter (kg/m^2) of fat-free body mass. Typical young men have about $20 \text{ kg}/\text{m}^2$. Each subject chose two images, one that



© Rubberball/Age fotostock

represented his own level of body muscle and one that he thought represented “what women prefer.” The mean gap between self-image and “what women prefer” was $2.35 \text{ kg}/\text{m}^2$.

Suppose that the “muscle gap” in the population of all young men has a Normal distribution with standard deviation $2.5 \text{ kg}/\text{m}^2$. Give a 90% confidence interval for the mean amount of muscle young men think they should add to be attractive to women. (They are wrong: women actually prefer a level close to that of typical men.)

14.21 An outlier strikes. There were actually 464 responses to the class survey in Exercise 14.19. One student claimed to study 60,000 minutes per night. We know he’s joking, so we left out this value. If we did a calculation without looking at the data, we would get $\bar{x} = 247$ minutes for all 464 students. Now what is the 99% confidence interval for the population mean? (Continue to use $\sigma = 65$.) Compare the new interval with that in Exercise 14.19. The message is clear: always look at your data, because outliers can greatly change your result.

14.22 Explaining confidence. A student reads that a 95% confidence interval for the mean ideal weight given by adult American women is 140 ± 1.4 pounds. Asked to explain the meaning of this interval, the student says, “95% of all adult American women would say that their ideal weight is between 138.6 and 141.4 pounds.” Is the student right? Explain your answer.

14.23 Explaining confidence. You ask another student to explain the confidence interval for mean ideal weight described in the previous exercise. The student answers, “We can be 95% confident that future samples of adult American women will say that their mean ideal weight is between 138.6 and 141.4 pounds.” Is this explanation correct? Explain your answer.

14.24 Explaining confidence. Here is an explanation from the Associated Press concerning one of its opinion polls. Explain briefly but clearly in what way this explanation is incorrect.

For a poll of 1,600 adults, the variation due to sampling error is no more than three percentage points either way. The error margin is said to be valid at the 95 percent confidence level. This means that, if the same questions were repeated in 20 polls, the results of at least 19 surveys would be within three percentage points of the results of this survey.

Exercises 14.25 to 14.27 ask you to answer questions from data. Assume that the “simple conditions” hold in each case. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan**, **Solve**, and **Conclude** steps, illustrated in Example 14.3 (page 359) for a confidence interval.

14.25 Pulling wood apart. How heavy a load (pounds) is needed to pull apart pieces of Douglas fir 4 inches long and 1.5 inches square? Here are data from students doing a laboratory exercise:

33,190	31,860	32,590	26,520	33,280
32,320	33,020	32,030	30,460	32,700
23,040	30,930	32,720	33,650	32,340
24,050	30,170	31,300	28,730	31,920

(a) We are willing to regard the wood pieces prepared for the lab session as an SRS of all similar pieces of Douglas fir. Engineers also commonly assume that characteristics of materials vary Normally. Make a graph to show the shape of the distribution for these data. Does it appear safe to assume that the Normality condition is satisfied? Suppose that the strength of pieces of wood like these follows a Normal distribution with standard deviation 3000 pounds.

(b) Give a 95% confidence interval for the mean load required to pull the wood apart. 

14.26 Bone loss by nursing mothers. Breast-feeding mothers secrete calcium into their milk. Some of the calcium may come from their bones, so mothers may lose bone mineral. Researchers measured the percent change in mineral content of the spines of 47 mothers during three months of breast-feeding.⁶ Here are the data:

-4.7	-2.5	-4.9	-2.7	-0.8	-5.3	-8.3	-2.1	-6.8	-4.3
2.2	-7.8	-3.1	-1.0	-6.5	-1.8	-5.2	-5.7	-7.0	-2.2
-6.5	-1.0	-3.0	-3.6	-5.2	-2.0	-2.1	-5.6	-4.4	-3.3
-4.0	-4.9	-4.7	-3.8	-5.9	-2.5	-0.3	-6.2	-6.8	1.7
0.3	-2.3	0.4	-5.3	0.2	-2.2	-5.1			



Blend Images/Superstock

(a) The researchers are willing to consider these 47 women as an SRS from the population of all nursing mothers. Suppose that the percent change in this population has standard deviation $\sigma = 2.5\%$. Make a stemplot of the data to verify that the data follow a Normal distribution quite closely. (Don’t forget that you need both a 0 and a -0 stem because there are both positive and negative values.)

(b) Use a 99% confidence interval to estimate the mean percent change in the population.  BONELOSS

14.27 This wine stinks. Sulfur compounds cause “off-odors” in wine, so winemakers want to know the odor threshold, the lowest concentration of a compound that the human nose can detect. The odor threshold for dimethyl sulfide (DMS) in trained wine tasters is about 25 micrograms per liter of wine ($\mu\text{g/l}$). The untrained noses of consumers may be less sensitive, however. Here are the DMS odor thresholds for 10 untrained students:

30	30	42	35	22	33	31	29	19	23
----	----	----	----	----	----	----	----	----	----

(a) Assume that the standard deviation of the odor threshold for untrained noses is known to be $\sigma = 7 \mu\text{g/l}$. Briefly discuss the other two “simple conditions,” using a stemplot to verify that the distribution is roughly symmetric with no outliers.

(b) Give a 95% confidence interval for the mean DMS odor threshold among all students.  WINE2

14.28 Why are larger samples better? Statisticians prefer large samples. Describe briefly the effect of increasing the size of a sample on the margin of error of a 95% confidence interval.



EXPLORING THE WEB

14.29 A statistics glossary. An editorial was published in the *Journal of the National Cancer Institute*, Vol. 101, No. 23 (December 2, 2009) that announced some online resources for journalists, including a statistics glossary. The glossary can be found at www.oxfordjournals.org/our_journals/jnc/resource/statistics%20glossary.pdf. Read the definition of a confidence interval. Is this an accurate definition? Explain your answer.

14.30 Getting around No Child Left Behind. The PBS Web site has an interesting article from 2007 discussing how school districts were getting around certain requirements of the No Child Left Behind law. You can find the article at www.pbs.org/newshour/bb/education/july-dec07/nclb_08-14.html. What does the article have to say about the use of confidence intervals in reporting results about the percent of students passing proficiency tests?



EXIT



Tests of Significance: The Basics

Chapter 15

Confidence intervals are one of the two most common types of statistical inference. Use a confidence interval when your goal is to estimate a population parameter. The second common type of inference, called *tests of significance*, has a different goal: to assess the evidence provided by data about some claim concerning a population. Here is the reasoning of statistical tests in a nutshell.

EXAMPLE 15.1 I'm a good free-throw shooter

I claim that I make 75% of my basketball free throws. To test my claim, you ask me to shoot 20 free throws. I make only 8 of the 20. "Aha!" you say. "Someone who makes 75% of his free throws would almost never make only 8 out of 20. So I don't believe your claim."

Your reasoning is based on asking what would happen if my claim were true and we repeated the sample of 20 free throws many times—I would almost never make as few as 8. This outcome is so unlikely that it gives strong evidence that my claim is not true.

You can say how strong the evidence against my claim is by giving the probability that I would make as few as 8 out of 20 free throws if I really make 75% in the long run. This probability is 0.0009. I would make as few as 8 of 20 only 9 times in 10,000 tries in the long run if my claim to make 75% were true. The small probability convinces you that my claim is false. ■

The *Reasoning of a Statistical Test* applet animates Example 15.1. You can ask a player to shoot free throws until the data do (or don't) convince you that he makes fewer than 75%. Significance tests use an elaborate vocabulary, but the basic idea is simple: *an outcome that would rarely happen if a claim were true is good evidence that the claim is not true*.

IN THIS CHAPTER WE COVER...

- The reasoning of tests of significance
- Stating hypotheses
- *P*-value and statistical significance
- Tests for a population mean
- Significance from a table*



THE REASONING OF TESTS OF SIGNIFICANCE

The reasoning of statistical tests, like that of confidence intervals, is based on asking what would happen if we repeated the sample or experiment many times. We will act as if the “simple conditions” listed on page 352 are true: we have a perfect SRS from an exactly Normal population with standard deviation σ known to us. Here is an example we will explore.



Ramin/Talaie/CORBIS

EXAMPLE 15.2 Sweetening colas

Diet colas use artificial sweeteners to avoid sugar. These sweeteners gradually lose their sweetness over time. Manufacturers therefore test new colas for loss of sweetness before marketing them. Trained tasters sip the cola along with drinks of standard sweetness and score the cola on a “sweetness score” of 1 to 10. The cola is then stored for a month at high temperature to imitate the effect of four months’ storage at room temperature. Each taster scores the cola again after storage. This is a matched pairs experiment. Our data are the differences (score before storage minus score after storage) in the tasters’ scores. The bigger these differences, the bigger the loss of sweetness.

Suppose we know that for any cola, the sweetness loss scores vary from taster to taster according to a Normal distribution with standard deviation $\sigma = 1$. The mean μ for all tasters measures loss of sweetness and is different for different colas.

Here are the sweetness losses for a new cola, as measured by 10 trained tasters:

2.0 0.4 0.7 2.0 -0.4 2.2 -1.3 1.2 1.1 2.3

Most are positive. That is, most tasters found a loss of sweetness. But the losses are small, and two tasters (the negative scores) thought the cola gained sweetness. The average sweetness loss is given by the sample mean $\bar{x} = 1.02$. Are these data good evidence that the cola lost sweetness in storage? ■

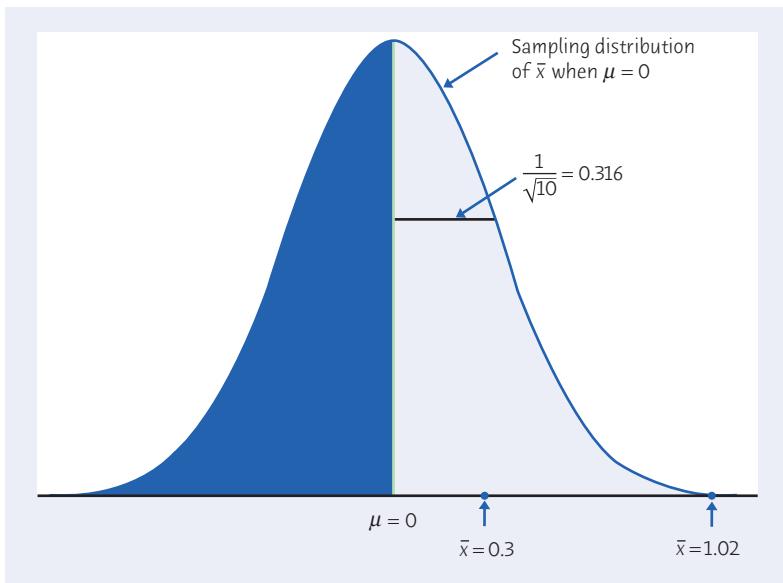
The reasoning is the same as in Example 15.1. We make a claim and ask if the data give evidence *against* it. We seek evidence that there *is* a sweetness loss, so the claim we test is that there *is not* a loss. In that case, the mean loss for the population of all trained testers would be $\mu = 0$.

- If the claim that $\mu = 0$ is true, the sampling distribution of \bar{x} from 10 tasters is Normal with mean $\mu = 0$ and standard deviation

$$\frac{\sigma}{\sqrt{n}} = \frac{1}{\sqrt{10}} = 0.316$$

This is just like the calculations we did in Chapter 11 (see Example 11.5 on page 294) and Chapter 14 (see Example 14.1 on page 352). Figure 15.1 shows this sampling distribution. We can judge whether any observed \bar{x} is surprising by locating it on this distribution.

- For a cola already on the market, 10 tasters had mean loss $\bar{x} = 0.3$. It is clear from Figure 15.1 that an \bar{x} this large could easily occur just by chance when

**FIGURE 15.1**

If the cola does not lose sweetness in storage, the mean score \bar{x} for 10 tasters will have this sampling distribution. The actual result for one cola was $\bar{x} = 0.3$. That could easily happen just by chance. Another cola had $\bar{x} = 1.02$. That's so far out on the Normal curve that it is good evidence that this cola did lose sweetness.

the population mean is $\mu = 0$. That 10 tasters found $\bar{x} = 0.3$ is not evidence that this cola loses sweetness.

- The taste test for the new cola produced $\bar{x} = 1.02$. That's way out on the Normal curve in Figure 15.1—so far out that *an observed value this large would rarely occur just by chance if the true μ were 0*. This observed value is good evidence that the true μ is in fact greater than 0, that is, that the cola lost sweetness. The manufacturer must reformulate the cola and try again.

APPLY YOUR KNOWLEDGE

- 15.1 Student attitudes.** The Survey of Study Habits and Attitudes (SSHA) is a psychological test that measures students' study habits and attitude toward school. Scores range from 0 to 200. The mean score for college students is about 115, and the standard deviation is about 30. A teacher suspects that the mean μ for older students is higher than 115. She gives the SSHA to an SRS of 25 students who are at least 30 years old. Suppose we know that scores in the population of older students are Normally distributed with standard deviation $\sigma = 30$.

- We seek evidence *against* the claim that $\mu = 115$. What is the sampling distribution of the mean score \bar{x} of a sample of 25 students if the claim is true? Draw the density curve of this distribution. (Sketch a Normal curve, then mark on the axis the values of the mean and 1, 2, and 3 standard deviations on either side of the mean.)
- Suppose that the sample data give $\bar{x} = 118.6$. Mark this point on the axis of your sketch. In fact, the result was $\bar{x} = 125.8$. Mark this point on your sketch. Using your sketch, explain in simple language why one result is good evidence that the mean score of all older students is greater than 115 and why the other outcome is not.

15.2 Measuring conductivity. The National Institute of Standards and Technology (NIST) supplies a “standard iron rod” whose electrical conductivity is supposed to be exactly 10.1. Is there reason to think that the true conductivity is not 10.1? To find out, NIST measures the conductivity of one rod 6 times. Repeated measurements of the same thing vary, which is why NIST makes 6 measurements. These measurements are an SRS from the population of all possible measurements. This population has a Normal distribution with mean μ equal to the true conductivity and standard deviation $\sigma = 0.1$.

- We seek evidence *against* the claim that $\mu = 10.1$. What is the sampling distribution of the mean \bar{x} in many samples of 6 measurements of one rod if the claim is true? Make a sketch of the Normal curve for this distribution. (Draw a Normal curve, then mark on the axis the values of the mean and 1, 2, and 3 standard deviations on either side of the mean.)
- Suppose that the sample mean is $\bar{x} = 10.09$. Mark this value on the axis of your sketch. Another rod has $\bar{x} = 9.95$ for 6 measurements. Mark this value on the axis as well. Explain in simple language why one result is good evidence that the true conductivity differs from 10.1 and why the other result gives no reason to doubt that 10.1 is correct.

STATING HYPOTHESES

A statistical test starts with a careful statement of the claims we want to compare. In Example 15.2, we saw that the taste test data are not plausible if the new cola loses no sweetness. Because the reasoning of tests looks for evidence *against* a claim, we start with the claim we seek evidence against, such as “no loss of sweetness.”

NULL AND ALTERNATIVE HYPOTHESES

The claim tested by a statistical test is called the **null hypothesis**. The test is designed to assess the strength of the evidence *against* the null hypothesis. Usually the null hypothesis is a statement of “no effect” or “no difference.”

The claim about the population that we are trying to find evidence *for* is the **alternative hypothesis**. The alternative hypothesis is **one-sided** if it states that a parameter is *larger than* or *smaller than* the null hypothesis value. It is **two-sided** if it states that the parameter is *different from* the null value (it could be either smaller or larger).



We abbreviate the null hypothesis as H_0 and the alternative hypothesis as H_a . Hypotheses always refer to a population, not to a particular outcome. Be sure to state H_0 and H_a in terms of population parameters. Because H_a expresses the effect that we hope to find evidence *for*, it is sometimes easier to begin by stating H_a and then set up H_0 as the statement that the hoped-for effect is not present.

In Example 15.2, we are seeking evidence *for* loss in sweetness. The null hypothesis says “no loss” on the average in a large population of tasters. The alternative hypothesis says “there is a loss.” So the hypotheses are

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

The alternative hypothesis is *one-sided* because we are interested only in whether the cola lost sweetness.

EXAMPLE 15.3 Studying job satisfaction

Does the job satisfaction of assembly workers differ when their work is machine-paced rather than self-paced? Assign workers either to an assembly line moving at a fixed pace or to a self-paced setting. All subjects work in both settings, in random order. This is a matched pairs design. After two weeks in each work setting, the workers take a test of job satisfaction. The response variable is the difference in satisfaction scores, self-paced minus machine-paced.

The parameter of interest is the mean μ of the differences in scores in the population of all assembly workers. The null hypothesis says that there is no difference between self-paced and machine-paced work, that is,

$$H_0: \mu = 0$$

The authors of the study wanted to know if the two work conditions have different levels of job satisfaction. They did not specify the direction of the difference. The alternative hypothesis is therefore *two-sided*:

$$H_a: \mu \neq 0$$

The hypotheses should express the hopes or suspicions we have before we see the data. It is cheating to first look at the data and then frame hypotheses to fit what the data show. For example, the data for the study in Example 15.3 showed that the workers were more satisfied with self-paced work, but this should not influence the choice of H_a . If you do not have a specific direction firmly in mind in advance, use a two-sided alternative.



APPLY YOUR KNOWLEDGE

- 15.3 Student attitudes.** State the null and alternative hypotheses for the study of older students' attitudes described in Exercise 15.1. (Is the alternative hypothesis one-sided or two-sided?)
- 15.4 Measuring conductivity.** State the null and alternative hypotheses for the study of electrical conductivity described in Exercise 15.2. (Is the alternative hypothesis one-sided or two-sided?)
- 15.5 Grading a teaching assistant.** The examinations in a large statistics class are scaled after grading so that the mean score is 75. The professor thinks that one teaching assistant is a poor teacher and suspects that his students have a lower mean score than the class as a whole. The TA's students this semester can be considered a sample from the population of all students in the course, so the professor compares their mean score with 75. State the hypotheses H_0 and H_a .
- 15.6 Women's incomes.** The average income of American women who work full-time and have only a high school degree is \$31,666. You wonder whether the mean income of female graduates from your local high school who work full-time but have only a high

Honest hypotheses?

Chinese and Japanese, for whom the number

4 is unlucky, die more often on the fourth day of the month than on other days. The authors of a study did a statistical test of the claim that the fourth day has more deaths than other days and found good evidence in favor of this claim. Can we trust this? Not if the authors looked at all days, picked the one with the most deaths, then made "this day is different" the claim to be tested. A critic raised that issue, and the authors replied: "No, we had day 4 in mind in advance, so our test was legitimate."

school degree is different from the national average. You obtain income information from an SRS of 62 female graduates who work full-time and have only a high school degree and find that $\bar{x} = \$30,052$. What are your null and alternative hypotheses?

- 15.7 Stating hypotheses.** In planning a study of the annual consumption of carbonated soft drinks by high school students, a researcher states the hypotheses as

$$\begin{aligned} H_0: \bar{x} &= 60 \text{ gallons per year} \\ H_a: \bar{x} &> 60 \text{ gallons per year} \end{aligned}$$

What's wrong with this?

P-VALUE AND STATISTICAL SIGNIFICANCE

The idea of stating a null hypothesis that we want to find evidence *against* seems odd at first. It may help to think of a criminal trial. The defendant is “innocent until proven guilty.” That is, the null hypothesis is innocence and the prosecution must try to provide convincing evidence against this hypothesis. That’s exactly how statistical tests work, though in statistics we deal with evidence provided by data and use a probability to say how strong the evidence is.

The probability that measures the strength of the evidence against a null hypothesis is called a *P-value*. Statistical tests generally work like this:

TEST STATISTIC AND P-VALUE

A **test statistic** calculated from the sample data measures how far the data diverge from what we would expect if the null hypothesis H_0 were true. Large values of the statistic show that the data are not consistent with H_0 .

The probability, computed assuming that H_0 is true, that the test statistic would take a value as extreme or more extreme than that actually observed is called the **P-value** of the test. The smaller the *P*-value, the stronger the evidence against H_0 provided by the data.

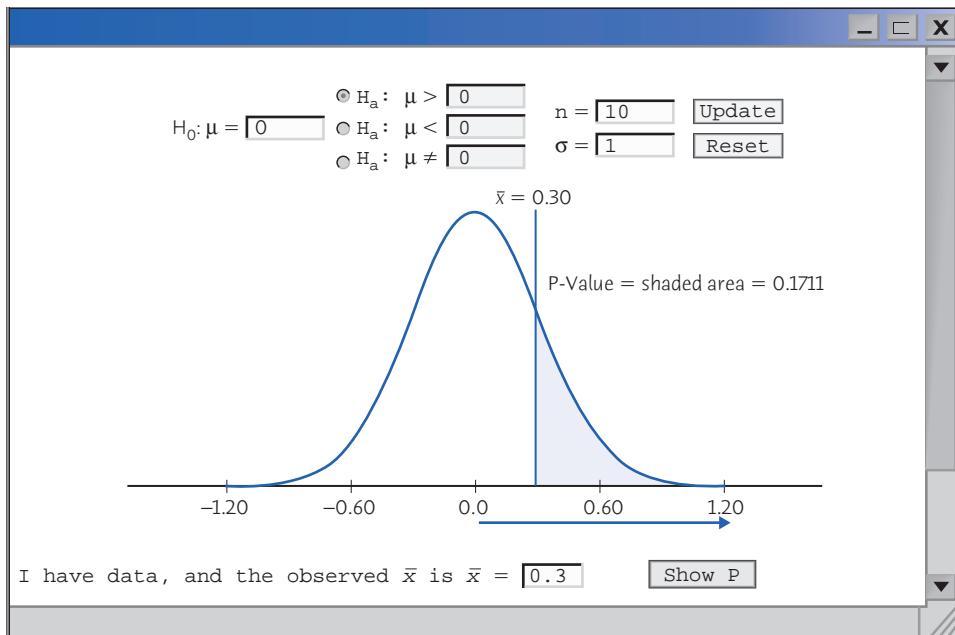
Small *P*-values are evidence against H_0 because they say that the observed result would be unlikely to occur if H_0 were true. Large *P*-values fail to give evidence against H_0 . Statistical software will give you the *P*-value of a test when you enter your null and alternative hypotheses and your data. So your most important task is to understand what a *P*-value says.

EXAMPLE 15.4 Sweetening colas: one-sided *P*-value

The study of sweetness loss in Example 15.2 tests the hypotheses

$$\begin{aligned} H_0: \mu &= 0 \\ H_a: \mu &> 0 \end{aligned}$$

Because the alternative hypothesis says that $\mu > 0$, values of \bar{x} greater than 0 favor H_a over H_0 . The test statistic compares the observed \bar{x} with the hypothesized value $\mu = 0$. For now, let’s concentrate on the *P*-value.

**FIGURE 15.2**

The one-sided *P*-value for the cola with mean sweetness loss $\bar{x} = 0.3$ in Example 15.4. The figure shows both the input and the output for the *P-Value of a Test of Significance* applet. Note that the *P*-value is the shaded area under the curve, not the unshaded area.

The experiment presented in Example 15.2 actually compared two colas, though Example 15.2 gives actual data only for one. For the first cola, the 10 tasters found mean sweetness loss $\bar{x} = 0.3$. For the second, the data gave $\bar{x} = 1.02$. *The P-value for each test is the probability of getting an \bar{x} this large when the mean sweetness loss is really $\mu = 0$.*

The shaded area in Figure 15.2 shows the *P*-value when $\bar{x} = 0.3$. The Normal curve is the sampling distribution of \bar{x} when the null hypothesis $H_0: \mu = 0$ is true. A Normal probability calculation (Exercise 15.8 page 377) shows that the *P*-value is $P(\bar{x} \geq 0.3) = 0.1711$.

A value as large as $\bar{x} = 0.3$ would occur just by chance in 17% of all samples when $H_0: \mu = 0$ is true. So observing $\bar{x} = 0.3$ is not strong evidence against H_0 . On the other hand, you can calculate that the probability that \bar{x} is 1.02 or larger when in fact $\mu = 0$ is only 0.0006. We would very rarely observe a mean sweetness loss of 1.02 or larger if H_0 were true. This small *P*-value provides strong evidence against H_0 and in favor of the alternative $H_a: \mu > 0$. ■

Figure 15.2 is actually the output of the *P-Value of a Test of Significance* applet, along with the information we entered into the applet. This applet automates the work of finding *P*-values for samples of size 50 or smaller under the “simple conditions” for inference about a mean.

The alternative hypothesis sets the direction that counts as evidence against H_0 . In Example 15.4, only large positive values count because the alternative is one-sided on the high side. If the alternative is two-sided, both directions count.



EXAMPLE 15.5 Job satisfaction: two-sided P -value

The study of job satisfaction in Example 15.3 requires that we test

$$H_0: \mu = 0$$

$$H_a: \mu \neq 0$$

Suppose we know that differences in job satisfaction scores (self-paced minus machine-paced) in the population of all workers follow a Normal distribution with standard deviation $\sigma = 60$.

Data from 18 workers give $\bar{x} = 17$. That is, these workers prefer the self-paced environment on the average. Because the alternative is two-sided, the P -value is the probability of getting an \bar{x} at least as far from $\mu = 0$ in either direction as the observed $\bar{x} = 17$.

Enter the information for this example into the *P-Value of a Test of Significance* applet and click “Show P.” Figure 15.3 shows the applet output as well as the information we entered. The P -value is the sum of the two shaded areas under the Normal curve. It is $P = 0.2302$. Values as far from 0 as $\bar{x} = 17$ (in either direction) would happen 23% of the time when the true population mean is $\mu = 0$. An outcome that would occur so often when H_0 is true is not good evidence against H_0 . ■

The conclusion of Example 15.5 is *not* that H_0 is true. The study looked for evidence against $H_0: \mu = 0$ and failed to find strong evidence. That is all we can say. No doubt the mean μ for the population of all assembly workers is not

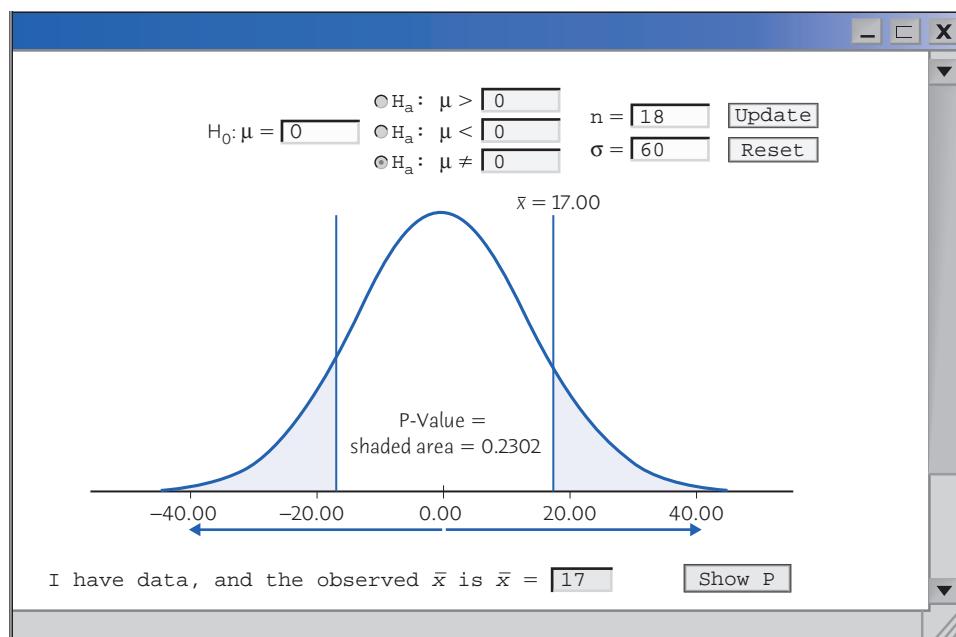


FIGURE 15.3

The two-sided P -value for Example 15.5. The figure shows both the input and the output for the *P-Value of a Test of Significance* applet. Note that the P -value is the shaded area under the curve, not the unshaded area.

exactly equal to 0. A large enough sample would give evidence of the difference, even if the difference is very small. Tests of significance assess the evidence against H_0 . If the evidence is strong, we can confidently reject H_0 in favor of the alternative. *Failing to find evidence against H_0 means only that the data are not inconsistent with H_0 , not that we have clear evidence that H_0 is true.* Only data that are inconsistent with H_0 provide evidence against H_0 .



In Examples 15.4 and 15.5, we decided that P -value $P = 0.0006$ was strong evidence against the null hypothesis and that P -values $P = 0.1711$ and $P = 0.2302$ did not give convincing evidence. There is no rule for how small a P -value we should require to reject H_0 —it's a matter of judgment and depends on the specific circumstances.

Nonetheless, we can compare a P -value with some fixed values that are in common use as standards for evidence against H_0 . The most common fixed values are 0.05 and 0.01. If $P \leq 0.05$, there is no more than 1 chance in 20 that a sample would give evidence this strong just by chance when H_0 is actually true. If $P \leq 0.01$, we have a result that in the long run would happen no more than once per 100 samples if H_0 were true. These fixed standards for P -values are called **significance levels**. We use α , the Greek letter alpha, to stand for a significance level.

significance level

STATISTICAL SIGNIFICANCE

If the P -value is as small or smaller than α , we say that the data are **statistically significant at level α** . The quantity α is called the **significance level** or the **level of significance**.

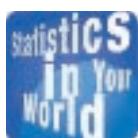
“Significant” in the statistical sense does not mean “important.” It means simply “not likely to happen just by chance.” The significance level α makes “not likely” more exact. Significance at level 0.01 is often expressed by the statement “The results were significant ($P < 0.01$).” Here P stands for the P -value. The actual P -value is more informative than a statement of significance because it allows us to assess significance at any level we choose. For example, a result with $P = 0.03$ is significant at the $\alpha = 0.05$ level but is not significant at the $\alpha = 0.01$ level.



APPLY YOUR KNOWLEDGE

15.8 Sweetening colas: find the P -value. The P -value for the first cola in Example 15.4 (page 375) is the probability (taking the null hypothesis $\mu = 0$ to be true) that \bar{x} takes a value at least as large as 0.3.

- What is the sampling distribution of \bar{x} when $\mu = 0$? This distribution is shown in Figure 15.2.
- Do a Normal probability calculation to find the P -value. Your result should agree with Example 15.4 up to roundoff error.



Significance strikes down a new drug

The pharmaceutical company Pfizer spent \$1 billion developing a new cholesterol-fighting drug. The final test for its effectiveness was a clinical trial with 15,000 subjects. To enforce double-blindness, only an independent group of experts saw the data during the trial. Three years into the trial, the monitors declared that there was a statistically significant excess of deaths and of heart problems in the group assigned to the new drug. Pfizer ended the trial. There went \$1 billion.

15.9 Job satisfaction: find the *P*-value. The *P*-value in Example 15.5 is the probability (taking the null hypothesis $\mu = 0$ to be true) that \bar{x} takes a value at least as far from 0 in either direction as 17.

- What is the sampling distribution of \bar{x} when $\mu = 0$? This distribution is shown in Figure 15.3.
- Do a Normal probability calculation to find the *P*-value. Your result should agree with Example 15.5 up to roundoff error.

15.10 Lorcaserin and weight loss. A double-blind, randomized comparative experiment compared the effect of the drug lorcaserin and a placebo on weight loss in overweight adults. All subjects also underwent diet and exercise counseling. The study reported that after one year, patients in the lorcaserin group had an average weight loss of 5.8 kilograms (kg), while those on the placebo had an average weight loss of 2.2 kg ($P < 0.001$).¹ Explain to someone who knows no statistics why these results mean that there is good reason to think that lorcaserin works. Include an explanation of what $P < 0.001$ means.



15.11 Student attitudes. Exercise 15.1 describes a study of the attitudes of older college students. You stated the null and alternative hypotheses in Exercise 15.3 (page 373).

- One sample of 25 students had mean SSHA score $\bar{x} = 118.6$. Enter this \bar{x} , along with the other required information, into the *P-Value of a Test of Significance* applet. What is the *P*-value? Is this outcome statistically significant at the $\alpha = 0.05$ level? At the $\alpha = 0.01$ level?
- Another sample of 25 students had $\bar{x} = 125.8$. Use the applet to find the *P*-value for this outcome. Is it statistically significant at the $\alpha = 0.05$ level? At the $\alpha = 0.01$ level?
- Explain briefly why these *P*-values tell us that one outcome is strong evidence against the null hypothesis and that the other outcome is not.



15.12 Measuring conductivity. Exercise 15.2 describes 6 measurements of the electrical conductivity of an iron rod. You stated the null and alternative hypotheses in Exercise 15.4 (page 373).

- One set of measurements has mean conductivity $\bar{x} = 10.09$. Enter this \bar{x} , along with the other required information, into the *P-Value of a Test of Significance* applet. What is the *P*-value? Is this outcome statistically significant at the $\alpha = 0.05$ level? At the $\alpha = 0.01$ level?
- Another set of measurements has $\bar{x} = 9.95$. Use the applet to find the *P*-value for this outcome. Is it statistically significant at the $\alpha = 0.05$ level? At the $\alpha = 0.01$ level?
- Explain briefly why these *P*-values tell us that one outcome is strong evidence against the null hypothesis and that the other outcome is not.

TESTS FOR A POPULATION MEAN

We have used tests for hypotheses about the mean μ of a population, under the “simple conditions,” to introduce tests of significance. The big idea is the reasoning of a test: *data that would rarely occur if the null hypothesis H_0 were true provide evidence that H_0 is not true*. The *P*-value gives us a probability to measure “would

rarely occur." In practice, the steps in carrying out a significance test mirror the overall four-step process for organizing realistic statistical problems.

TESTS OF SIGNIFICANCE: A FOUR-STEP PROCESS

STATE: What is the practical question that requires a statistical test?

PLAN: Identify the parameter, state null and alternative hypotheses, and choose the type of test that fits your situation.

SOLVE: Carry out the test in three phases:

1. Check the conditions for the test you plan to use.
2. Calculate the test statistic.
3. Find the *P*-value.

CONCLUDE: Return to the practical question to describe your results in this setting.

Once you have stated your question, formulated hypotheses, and checked the conditions for your test, you or your software can find the test statistic and *P*-value by following a rule. Here is the rule for the test we have used in our examples.

z TEST FOR A POPULATION MEAN

Draw an SRS of size n from a Normal population that has unknown mean μ and known standard deviation σ . To test the null hypothesis that μ has a specified value,

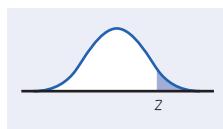
$$H_0: \mu = \mu_0$$

calculate the one-sample *z* test statistic

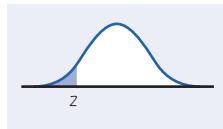
$$z = \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}}$$

In terms of a variable Z having the standard Normal distribution, the *P*-value for a test of H_0 against

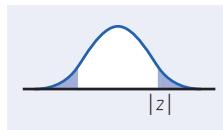
$$H_a: \mu > \mu_0 \text{ is } P(Z \geq z)$$



$$H_a: \mu < \mu_0 \text{ is } P(Z \leq z)$$



$$H_a: \mu \neq \mu_0 \text{ is } 2P(Z \geq |z|)$$



As promised, the test statistic z measures how far the observed sample mean \bar{x} deviates from the hypothesized population value μ_0 . The measurement is in the familiar standard scale obtained by dividing by the standard deviation of \bar{x} . So we have a common scale for all z tests, and the 68–95–99.7 rule helps us see at once if \bar{x} is far from μ_0 . The pictures that illustrate the P -value look just like the curves in Figures 15.2 (page 375) and 15.3 (page 376) except that they are in the standard scale.



ImageState/Alamy

EXAMPLE 15.6 Executives' blood pressures

STATE: The National Center for Health Statistics reports that the systolic blood pressure for males 35 to 44 years of age has mean 128 and standard deviation 15. The medical director of a large company looks at the medical records of 72 executives in this age group and finds that the mean systolic blood pressure in this sample is $\bar{x} = 126.07$. Is this evidence that the company's executives have a different mean systolic blood pressure from the general population?

PLAN: The null hypothesis is “no difference” from the national mean $\mu_0 = 128$. The alternative is two-sided because the medical director did not have a particular direction in mind before examining the data. So the hypotheses about the unknown mean μ of the executive population are

$$H_0: \mu = 128$$

$$H_a: \mu \neq 128$$

We know that the one-sample z test is appropriate for these hypotheses under the “simple conditions.”

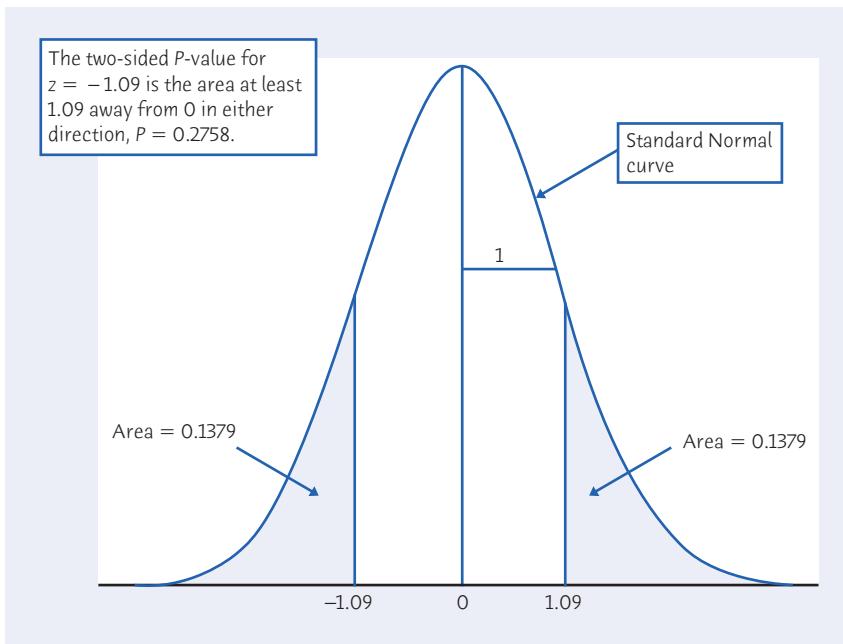
SOLVE: As part of the “simple conditions,” suppose we know that executives’ systolic blood pressures follow a Normal distribution with standard deviation $\sigma = 15$. Software can now calculate z and P for you. Going ahead by hand, the **test statistic** is

$$\begin{aligned} z &= \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}} = \frac{126.07 - 128}{15/\sqrt{72}} \\ &= -1.09 \end{aligned}$$

To help find the **P -value**, sketch the standard Normal curve and mark on it the observed value of z . Figure 15.4 shows that the P -value is the probability that a standard Normal variable Z takes a value at least 1.09 away from zero. From Table A or software, this probability is

$$P = 2P(Z < -1.09) = (2)(0.1379) = 0.2758$$

CONCLUDE: More than 27% of the time, an SRS of size 72 from the general male population would have a mean systolic blood pressure at least as far from 128 as that of the executive sample. The observed $\bar{x} = 126.07$ is therefore not good evidence that executives differ from other men. ■

**FIGURE 15.4**

The P -value for the two-sided test in Example 15.6. The observed value of the test statistic is $z = -1.09$.

In this chapter we are acting as if the “simple conditions” stated on page 352 are true. In practice, you must verify these conditions.



1. **SRS:** The most important condition is that the 72 executives in the sample are an SRS from the population of all middle-aged male executives in the company. We should check this requirement by asking how the data were produced. If medical records are available only for executives with recent medical problems, for example, the data are of little value for our purpose because of the obvious health bias. It turns out that all executives are given a free annual medical exam, and that the medical director selected 72 exam results at random.
2. **Normal distribution:** We should also examine the distribution of the 72 observations to look for signs that the population distribution is not Normal.
3. **Known σ :** It really is unrealistic to suppose that we know that $\sigma = 15$. We will see in Chapter 18 that it is easy to do away with the need to know σ .



APPLY YOUR KNOWLEDGE

15.13 The z statistic. Published reports of research work are terse. They often report just a test statistic and P -value. For example, the conclusion of Example 15.6 might be stated as “($z = -1.09, P = 0.2758$).” Find the values of the one-sample z statistic needed to complete these conclusions:

- (a) For the first cola in Example 15.4 (page 374), $z = ?, P = 0.1711$.

- (b) For the second cola in Example 15.4, $z = ?$, $P = 0.0006$.
 (c) For Example 15.5, $z = ?$, $P = 0.2302$.



15.14 Measuring conductivity. Here are 6 measurements of the electrical conductivity of an iron rod:

10.08 9.89 10.05 10.16 10.21 10.11

The iron rod is supposed to have conductivity 10.1. Do the measurements give good evidence that the true conductivity is not 10.1?

The 6 measurements are an SRS from the population of all results we would get if we kept measuring conductivity forever. This population has a Normal distribution with mean equal to the true conductivity of the rod and standard deviation 0.1. Use this information to carry out a test, following the four-step process as illustrated in Example 15.6. CONDUCTIVITY



15.15 Bad weather, bad tip? People tend to be more generous after receiving good news. Are they less generous after receiving bad news? The average tip left by adult Americans is 20%. Give 20 patrons of a restaurant a message on their bill warning them that tomorrow's weather will be bad and record the percentage tip they leave. Here are the tips as a percentage of the total bill:²

18.0 19.1 19.2 18.8 18.4 19.0 18.5 16.1 16.8 18.2
 14.0 17.0 13.6 17.5 20.0 20.2 18.8 18.0 23.2 19.4

Suppose that percentage tips are Normal with $\sigma = 2$. Is there good evidence that the mean percentage tip is less than 20? Follow the four-step process as illustrated in Example 15.6. TIPPING3

SIGNIFICANCE FROM A TABLE*

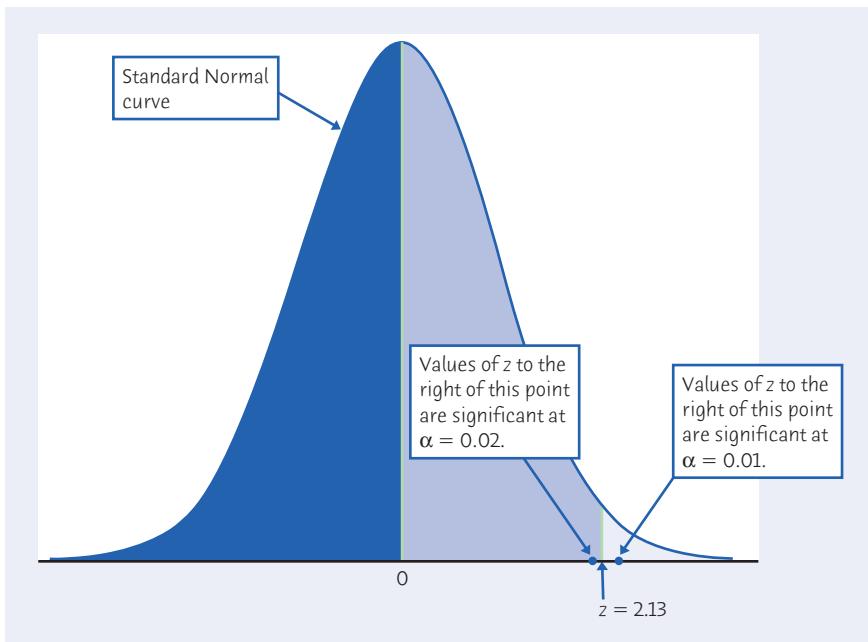
Statistics in practice uses technology (graphing calculator or software) to get P -values quickly and accurately. In the absence of suitable technology, you can get approximate P -values quickly by comparing the value of your test statistic with critical values from a table. For the z statistic, the table is Table C, the same table we used for confidence intervals.

Look at the bottom row of critical values in Table C, labeled z^* . At the top of the table, you see the confidence level C for each z^* . At the bottom of the table, you see both the one-sided and two-sided P -values for each z^* . Values of a test statistic z that are farther out than a z^* (in the direction given by the alternative hypothesis) are significant at the level that matches z^* .

SIGNIFICANCE FROM A TABLE OF CRITICAL VALUES

To find the approximate P -value for any z statistic, compare z (ignoring its sign) with the critical values z^* at the bottom of Table C. If z falls between two values of z^* , the P -value falls between the two corresponding values of P in the “One-sided P ” or the “Two-sided P ” row of Table C.

*This material can be skipped if you use software to compute P -values.

**FIGURE 15.5**

Is it significant? The test statistic value $z = 2.13$ falls between the critical values required for significance at the $\alpha = 0.02$ and $\alpha = 0.01$ levels. So the test is significant at $\alpha = 0.02$ and is not significant at $\alpha = 0.01$.

EXAMPLE 15.7 Is it significant?

The z statistic for a one-sided test is $z = 2.13$. How significant is this result? Compare $z = 2.13$ with the z^* row in Table C. It lies between $z^* = 2.054$ and $z^* = 2.326$. So the P -value lies between the corresponding entries in the “One-sided P ” row, which are $P = 0.02$ and $P = 0.01$. This z is significant at the $\alpha = 0.02$ level and is not significant at the $\alpha = 0.01$ level.

Figure 15.5 illustrates the situation. The shaded area under the Normal curve is the P -value for $z = 2.13$. You can see that P falls between the areas to the right of the two critical values, for $P = 0.02$ and $P = 0.01$.

The z statistic in Example 15.6 is $z = -1.09$. The alternative hypothesis is two-sided. Compare $z = -1.09$ (ignoring the minus sign) with the z^* row in Table C. It lies between $z^* = 1.036$ and $z^* = 1.282$. So the P -value lies between the matching entries in the “Two-sided P ” row, $P = 0.30$ and $P = 0.20$. This is enough to conclude that the data do not provide good evidence against the null hypothesis. ■

z^*	2.054	2.326
One-sided P	0.02	0.01

z^*	1.036	1.282
Two-sided P	0.30	0.20

APPLY YOUR KNOWLEDGE

- 15.16 Significance from a table.** A test of $H_0: \mu = 0$ against $H_a: \mu > 0$ has test statistic $z = 1.876$. Is this test significant at the 5% level ($\alpha = 0.05$)? Is it significant at the 1% level ($\alpha = 0.01$)?

15.17 Significance from a table. A test of $H_0: \mu = 0$ against $H_a: \mu \neq 0$ has test statistic $z = 1.876$. Is this test significant at the 5% level ($\alpha = 0.05$)? Is it significant at the 1% level ($\alpha = 0.01$)?

15.18 Testing a random number generator. A random number generator is supposed to produce random numbers that are uniformly distributed on the interval from 0 to 1. If this is true, the numbers generated come from a population with $\mu = 0.5$ and $\sigma = 0.2887$. A command to generate 100 random numbers gives outcomes with mean $\bar{x} = 0.4365$. Assume that the population σ remains fixed. We want to test

$$H_0: \mu = 0.5$$

$$H_a: \mu \neq 0.5$$

- (a) Calculate the value of the z test statistic.
- (b) Use Table C: is z significant at the 5% level ($\alpha = 0.05$)?
- (c) Use Table C: is z significant at the 1% level ($\alpha = 0.01$)?
- (d) Between which two Normal critical values z^* in the bottom row of Table C does z lie? Between what two numbers does the P -value lie? Does the test give good evidence against the null hypothesis?

CHAPTER 15 SUMMARY

CHAPTER SPECIFICS

- A **test of significance** assesses the evidence provided by data against a **null hypothesis** H_0 in favor of an **alternative hypothesis** H_a .
- Hypotheses are always stated in terms of population parameters. Usually H_0 is a statement that no effect is present, and H_a says that a parameter differs from its null value in a specific direction (**one-sided alternative**) or in either direction (**two-sided alternative**).
- The essential reasoning of a significance test is as follows. Suppose for the sake of argument that the null hypothesis is true. If we repeated our data production many times, would we often get data as inconsistent with H_0 as the data we actually have? Data that would rarely occur if H_0 were true provide evidence against H_0 .
- A test is based on a **test statistic** that measures how far the sample outcome is from the value stated by H_0 .
- The **P -value** of a test is the probability, computed supposing H_0 to be true, that the test statistic will take a value at least as extreme as that actually observed. Small P -values indicate strong evidence against H_0 . To calculate a P -value we must know the sampling distribution of the test statistic when H_0 is true.
- If the P -value is as small or smaller than a specified value α , the data are **statistically significant at significance level α** .
- **Significance tests for the null hypothesis $H_0: \mu = \mu_0$** concerning the unknown mean μ of a population are based on the **one-sample z test statistic**

$$z = \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}}$$

- The z test assumes an SRS of size n from a Normal population with known population standard deviation σ . P -values can be obtained either with computations from the standard Normal distribution or by using technology (applet or software).

LINK IT

In this chapter we discuss tests of significance, the second type of statistical inference. The mathematics of probability, in particular the sampling distributions discussed in Chapter 11, provides the formal basis for a test of significance. The sampling distribution allows us to assess “probabilistically” the strength of evidence against a null hypothesis, either through a level of significance or a P -value. The goal of hypothesis testing, which is used to assess the evidence provided by data about some claim concerning a population, is different from the goal of confidence interval estimation, which is used to estimate a population parameter.

Although we apply the reasoning of tests of significance for the mean of a population that has a Normal distribution in a simple and artificial setting (we assume that we know the population standard deviation), we will use the same logic in future chapters to construct tests of significance for population parameters in more realistic settings.

CHECK YOUR SKILLS

15.19 You use software to carry out a test of significance. The program tells you that the P -value is $P = 0.031$. You conclude

- that the probability, computed assuming that H_0 is true, that the test statistic would take a value as extreme or more extreme than that actually observed is 0.031.
- that the probability, computed assuming that H_0 is true, that the test statistic would take a value as extreme or less extreme than that actually observed is 0.031.
- that the probability, computed assuming that H_0 is false, that the test statistic would take a value as extreme or more extreme than that actually observed is 0.031.

15.20 You use software to carry out a test of significance. The program tells you that the P -value is $P = 0.031$. This result is

- not significant at either $\alpha = 0.05$ or $\alpha = 0.01$.
- significant at $\alpha = 0.05$ but not at $\alpha = 0.01$.
- significant at both $\alpha = 0.05$ and $\alpha = 0.01$.

15.21 The z statistic for a one-sided test is $z = 2.433$. This test is

- not significant at either $\alpha = 0.05$ or $\alpha = 0.01$.
- significant at $\alpha = 0.05$ but not at $\alpha = 0.01$.
- significant at both $\alpha = 0.05$ and $\alpha = 0.01$.

15.22 The gas mileage for a particular model car is known to have a standard deviation of $\sigma = 1.0$ miles per gallon in repeated tests in a controlled laboratory environment at a fixed speed. For a fixed speed, gas mileages in repeated tests are Normally distributed. Tests on three cars of this model at 35 miles per hour give gas mileages of 29.3, 29.9, and 29.8 miles per gallon. The z statistic for testing $H_0: \mu = 30$ miles per gallon based on these three measurements is

- $z = 0.286$.
- $z = 0.5$.
- $z = -0.286$.

15.23 Experiments on learning in animals sometimes measure how long it takes mice to find their way through a maze. The mean time is 18 seconds for one particular maze. A researcher thinks that a loud noise will cause the mice to complete the maze faster. She measures how long each of 10 mice takes with a noise as stimulus. The sample mean is $\bar{x} = 16.5$ seconds. The null hypothesis for the significance test is

- $H_0: \mu = 18$.
- $H_0: \mu = 16.5$.
- $H_0: \mu < 18$.



Garry Gay/Getty Images

15.24 The alternative hypothesis for the test in Exercise 15.23 is

- (a) $H_a: \mu \neq 18$. (b) $H_a: \mu < 18$. (c) $H_a: \mu = 16.5$.

15.25 You read an article about an experiment in which the researcher conducted a test of significance. The article tells you that the P -value is $P = 0.19$. This means that

- (a) the probability that the null hypothesis is true is 0.19.
 (b) the value of the test statistic is not particularly large.
 (c) neither of the above.

15.26 You are testing $H_0: \mu = 0$ against $H_a: \mu \neq 0$ based on an SRS of 20 observations from a Normal population. What values of the z statistic are statistically significant at the $\alpha = 0.005$ level?

- (a) All values for which $z > 2.576$
 (b) All values for which $z > 2.807$
 (c) All values for which $|z| > 2.807$

15.27 You are testing $H_0: \mu = 0$ against $H_a: \mu > 0$ based on an SRS of 20 observations from a Normal population. What values of the z statistic are statistically significant at the $\alpha = 0.005$ level?

- (a) All values for which $z > 2.576$
 (b) All values for which $z > 2.807$
 (c) All values for which $|z| > 2.807$

CHAPTER 15 EXERCISES

In all exercises that call for P -values, give the actual value if you use software or the P -Value applet. Otherwise, use Table C to give values between which P must fall.

15.28 Student study times. Exercise 14.19 (page 365) describes a class survey in which students claimed to study an average of $\bar{x} = 118$ minutes on a typical weeknight. Regard these students as an SRS from the population of all first-year students at this university. Does the study give good evidence that students claim to study less than 2 hours per night on the average?

- (a) State null and alternative hypotheses in terms of the mean study time in minutes for the population.
 (b) What is the value of the test statistic z ?
 (c) What is the P -value of the test? Can you conclude that students do claim to study less than 2 hours per weeknight on the average?

15.29 I want more muscle. If young men thought that their own level of muscle was about what women prefer, the mean “muscle gap” in the study described in Exercise 14.20 (page 365) would be 0. We suspect (before seeing the data) that young men think women prefer more muscle than they themselves have.

- (a) State null and alternative hypotheses for testing this suspicion.
 (b) What is the value of the test statistic z ?
 (c) You can tell just from the value of z that the evidence in favor of the alternative is very strong (that is, the P -value is very small). Explain why this is true.

15.30 Hotel managers’ personalities. Successful hotel managers must have personality characteristics often thought of as feminine (such as “compassionate”) as well as those often thought of as masculine (such as “forceful”). The Bem Sex-Role Inventory (BSRI) is a personality test that gives separate ratings for female and male stereotypes, both on a scale of 1 to 7. A sample of 148 male general managers of three-star and four-star hotels had mean BSRI femininity score $\bar{x} = 5.29$.³ The mean score for the general male population is $\mu = 5.19$. Do hotel managers, on the average, differ significantly in femininity score from men in general? Assume that the standard deviation of scores in the population of all male hotel managers is the same as the $\sigma = 0.78$ for the adult male population.

- (a) State null and alternative hypotheses in terms of the mean femininity score μ for male hotel managers.
 (b) Find the z test statistic.
 (c) What is the P -value for your z ? What do you conclude about male hotel managers?

15.31 Is this what P means? When asked to explain the meaning of “the P -value was $P = 0.03$,” a student says, “This means there is only probability 0.03 that the null hypothesis is true.” Explain what $P = 0.03$ really means in a way that makes it clear that the student’s explanation is wrong.

15.32 How to show that you are rich. Every society has its own marks of wealth and prestige. In ancient China, it appears that owning pigs was such a mark. Evidence comes

from examining burial sites. The skulls of sacrificed pigs tend to appear along with expensive ornaments, which suggests that the pigs, like the ornaments, signal the wealth and prestige of the person buried. A study of burials from around 3500 B.C. concluded that “there are striking differences in grave goods between burials with pig skulls and burials without them.... A test indicates that the two samples of total artifacts are significantly different at the 0.01 level.”⁴ Explain clearly why “significantly different at the 0.01 level” gives good reason to think that there really is a systematic difference between burials that contain pig skulls and those that lack them.

15.33 Alleviating test anxiety. Research suggests that pressure to perform well can reduce performance on exams. Are there effective strategies to deal with pressure? In an experiment, researchers had students take a test on mathematical skills. The same students were asked to take a second test on the same skills, but now each student was paired with a partner and only if both improved their scores would they receive a monetary reward for participating in the experiment. They were also told that their performance would be videotaped and watched by teachers and students. To help them cope with the pressure, ten minutes before the second exam they were asked to write as candidly as possible about their thoughts and feelings regarding the exam. “Students who expressed their thoughts before the high-pressure test showed a significant 5% math accuracy improvement from the pretest to posttest” ($P < 0.03$).⁵ A colleague who knows no statistics says that an increase of 5% isn’t a lot—maybe it’s just an accident due to natural variation among the students. Explain in simple language how “ $P < 0.03$ ” answers this objection.

15.34 Treating Parkinson’s disease. A randomized comparative experiment compared the effects of two types of deep-brain stimulation (pallidal stimulation and subthalamic stimulation) on change in motor function, as blindly assessed on the Unified Parkinson’s Disease Rating Scale, part III (UPDRS-III). The abstract of the study said: “Mean changes in the primary outcome did not differ significantly between the two study groups ($P = 0.50$).”⁶ The P -value refers to a null hypothesis of “no change” in measurements between pallidal stimulation and subthalamic stimulation. Explain clearly why this value provides no evidence of change.

15.35 5% versus 1%. Sketch the standard Normal curve for the z test statistic and mark off areas under the curve to show why a value of z that is significant at the 1% level in a one-sided test is always significant at the 5% level. If z is significant at the 5% level, what can you say about its significance at the 1% level?

15.36 The wrong alternative. A graduate student is comparing final-exam test scores of male and female students in an introductory physics class. She starts with no expectations as to which sex will score more highly. After seeing that men did better than women on the first quiz, she tests a one-sided alternative about the mean final-exam scores,

$$H_0: \mu_M = \mu_F$$

$$H_a: \mu_M > \mu_F$$

She finds $z = 1.9$ with one-sided P -value $P = 0.0287$.

- (a) Explain why she should have used the two-sided alternative hypothesis.
- (b) What is the correct P -value for $z = 1.9$?

15.37 The wrong P . The report of a study of seat belt use by drivers says, “Hispanic drivers were not significantly more likely than White/non-Hispanic drivers to overreport safety belt use (27.4 vs. 21.1%, respectively; $z = 1.33$, $P > 1.0$).”⁷ How do you know that the P -value given is incorrect? What is the correct one-sided P -value for test statistic $z = 1.33$?

Exercises 15.38 to 15.41 ask you to answer questions from data. Assume that the “simple conditions” hold in each case. The exercise statements give you the **State** step of the four-step process. In your work, follow the **Plan**, **Solve**, and **Conclude** steps, illustrated in Example 14.3 (page 359) for a confidence interval and in Example 15.6 (page 380) for a test of significance.

15.38 Pulling wood apart. How heavy a load (pounds) is needed to pull apart pieces of Douglas fir 4 inches long and 1.5 inches square? Here are data from students doing a laboratory exercise:

33,190	31,860	32,590	26,520	33,280
32,320	33,020	32,030	30,460	32,700
23,040	30,930	32,720	33,650	32,340
24,050	30,170	31,300	28,730	31,920

We are willing to regard the wood pieces prepared for the lab session as an SRS of all similar pieces of Douglas fir. Engineers also commonly assume that characteristics of materials vary Normally. Suppose that the strength of pieces of wood like these follows a Normal distribution with standard deviation 3000 pounds.

- (a) Is there significant evidence at the $\alpha = 0.10$ level against the hypothesis that the mean is 32,500 pounds for the two-sided alternative?
- (b) Is there significant evidence at the $\alpha = 0.10$ level against the hypothesis that the mean is 31,500 pounds for the two-sided alternative? 

15.39 Bone loss by nursing mothers. As discussed in Exercise 14.26 (page 366), breast-feeding mothers secrete calcium into their milk. Some of the calcium

may come from their bones, so mothers may lose bone mineral. Researchers measured the percent change in mineral content of the spines of 47 mothers during three months of breast-feeding.⁸ Here are the data:

-4.7	-2.5	-4.9	-2.7	-0.8	-5.3	-8.3	-2.1	-6.8	-4.3
2.2	-7.8	-3.1	-1.0	-6.5	-1.8	-5.2	-5.7	-7.0	-2.2
-6.5	-1.0	-3.0	-3.6	-5.2	-2.0	-2.1	-5.6	-4.4	-3.3
-4.0	-4.9	-4.7	-3.8	-5.9	-2.5	-0.3	-6.2	-6.8	1.7
0.3	-2.3	0.4	-5.3	0.2	-2.2	-5.1			

The researchers are willing to consider these 47 women as an SRS from the population of all nursing mothers. Suppose that the percent change in this population has a Normal distribution with standard deviation $\sigma = 2.5\%$. Do these data give good evidence that, on the average, nursing mothers lose bone mineral?  **BONELOSS**

 **15.40 This wine stinks.** Sulfur compounds cause “off odors” in wine, so winemakers want to know the odor threshold, the lowest concentration of a compound that the human nose can detect. The odor threshold for dimethyl sulfide (DMS) in trained wine tasters is about 25 micrograms per liter of wine ($\mu\text{g/l}$). The untrained noses of consumers may be less sensitive, however. Here are the DMS odor thresholds for 10 untrained students:

30 30 42 35 22 33 31 29 19 23

Assume that the odor threshold for untrained noses is Normally distributed with $\sigma = 7 \mu\text{g/l}$. Is there evidence that the mean threshold for untrained tasters is greater than 25 $\mu\text{g/l}$?  **WINE**

 **15.41 Eye grease.** Athletes performing in bright sunlight often smear black eye grease under their eyes to reduce glare. Does eye grease work? In one study, 16 student subjects took a test of sensitivity to contrast after 3 hours facing into bright sun, both with and without eye grease. This is a matched pairs design. Here are the differences in sensitivity, with eye grease minus without eye grease:⁹



AP Photo/Kathy Willens

0.07	0.64	-0.12	-0.05	-0.18	0.14	-0.16	0.03
0.05	0.02	0.43	0.24	-0.11	0.28	0.05	0.29

We want to know whether eye grease increases sensitivity on the average.

(a) What are the null and alternative hypotheses? Say in words what mean μ your hypotheses concern.

(b) Suppose that the subjects are an SRS of all young people with normal vision, that contrast differences follow a Normal distribution in this population, and that the standard deviation of differences is $\sigma = 0.22$. Carry out a test of significance.  **EYEGREASE**

15.42 Tests from confidence intervals. A confidence interval for the population mean μ tells us which values of μ are plausible (those inside the interval) and which values are not plausible (those outside the interval) at the chosen level of confidence. You can use this idea to carry out a test of any null hypothesis $H_0: \mu = \mu_0$ starting with a confidence interval: *reject H_0 if μ_0 is outside the interval and fail to reject if μ_0 is inside the interval.*

The alternative hypothesis is always two-sided, $H_a: \mu \neq \mu_0$, because the confidence interval extends in both directions from \bar{x} . A 95% confidence interval leads to a test at the 5% significance level because the interval is wrong 5% of the time. In general, confidence level C leads to a test at significance level $\alpha = 1 - C$.

(a) In Example 15.6 (page 380), a medical director found mean blood pressure $\bar{x} = 126.07$ for an SRS of 72 executives. The standard deviation of the blood pressures of all executives is $\sigma = 15$. Give a 90% confidence interval for the mean blood pressure μ of all executives.

(b) The hypothesized value $\mu_0 = 128$ falls *inside* this confidence interval. Carry out the z test for $H_0: \mu = 128$ against the two-sided alternative. Show that the test is *not significant* at the 10% level.

(c) The hypothesized value $\mu_0 = 129$ falls *outside* this confidence interval. Carry out the z test for $H_0: \mu = 129$ against the two-sided alternative. Show that the test is *significant* at the 10% level.

15.43 Tests from confidence intervals. A 95% confidence interval for a population mean is 30.7 ± 3.2 . Use the method described in the previous exercise to answer these questions.

(a) With a two-sided alternative, can you reject the null hypothesis that $\mu = 33$ at the 5% ($\alpha = 0.05$) significance level? Why?

(b) With a two-sided alternative, can you reject the null hypothesis that $\mu = 34$ at the 5% significance level? Why?



EXPLORING THE WEB

15.44 Significance in journals. Choose a major journal in your field of study. Use a Web search engine to find its Web site—just search on the journal’s name. Find a paper that uses a phrase like “significant ($P < 0.01$)” and summarize the findings in the paper.

15.45 A statistics glossary. An editorial was published in the *Journal of the National Cancer Institute*, Vol. 101, No. 23 (December 2, 2009) that announced some online resources for journalists, including a statistics glossary. The glossary can be found at www.oxfordjournals.org/our_journals/jnc/resource/statistics%20glossary.pdf. Read the definition of a P -value. Is this an accurate definition? Explain your answer.



Inference in Practice

To this point, we have met just two procedures for statistical inference. Both concern inference about the mean μ of a population when the “simple conditions” (page 352) are true: the data are an SRS, the population has a Normal distribution, and we know the standard deviation σ of the population. Under these conditions, a confidence interval for the mean μ is

$$\bar{x} \pm z^* \frac{\sigma}{\sqrt{n}}$$

To test a hypothesis $H_0: \mu = \mu_0$ we use the one-sample z statistic:

$$z = \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}}$$

We call these **z procedures** because they both start with the one-sample z statistic and use the standard Normal distribution.

In later chapters we will modify these procedures for inference about a population mean to make them useful in practice. We will also introduce procedures for confidence intervals and tests in most of the settings we met in learning to explore data. There are libraries—both of books and of software—full of more elaborate statistical techniques. The reasoning of confidence intervals and tests is the same, no matter how elaborate the details of the procedure are.

There is a saying among statisticians that “mathematical theorems are true; statistical methods are effective when used with judgment.” That the one-sample z statistic has the standard Normal distribution when the null hypothesis is true is a mathematical theorem. Effective use of statistical methods requires more than knowing such facts. It requires even more than understanding the underlying reasoning. This chapter begins the process of helping you develop the judgment needed to use statistics in practice. That process will continue in examples and exercises through the rest of this book.

Chapter 16

IN THIS CHAPTER WE COVER...

- Conditions for inference in practice
- Cautions about confidence intervals
- Cautions about significance tests
- Planning studies: sample size for confidence intervals
- Planning studies: the power of a statistical test*

procedures

CONDITIONS FOR INFERENCE IN PRACTICE



Any confidence interval or significance test can be trusted only under specific conditions. It's up to you to understand these conditions and judge whether they fit your problem. With that in mind, let's look back at the "simple conditions" for the z procedures.

The final "simple condition," that we know the standard deviation σ of the population, is rarely satisfied in practice. The z procedures are therefore of little practical use. Fortunately, it's easy to remove the "known σ " condition. Chapter 18 shows how. The first two "simple conditions" (SRS, Normal population) are harder to escape. In fact, they represent the kinds of conditions needed if we are to trust almost any statistical inference. As you plan inference, you should always ask, "Where did the data come from?" and you must often also ask, "What is the shape of the population distribution?" This is the point where knowing mathematical facts gives way to the need for judgment.

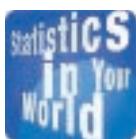
Where did the data come from? *The most important requirement for any inference procedure is that the data come from a process to which the laws of probability apply.* Inference is most reliable when the data come from a random sample or a randomized comparative experiment. Random samples use chance to choose respondents. Randomized comparative experiments use chance to assign subjects to treatments. The deliberate use of chance ensures that the laws of probability apply to the outcomes, and this in turn ensures that statistical inference makes sense.

WHERE THE DATA COME FROM MATTERS

When you use statistical inference, you are acting as if your data are a random sample or come from a randomized comparative experiment.



If your data don't come from a random sample or a randomized comparative experiment, your conclusions may be challenged. To answer the challenge, you must usually rely on subject-matter knowledge, not on statistics. It is common to apply statistical inference to data that are not produced by random selection. When you see such a study, ask whether the data can be trusted as a basis for the conclusions of the study.



Don't touch the plants

We know that confounding can distort inference.

We don't always recognize how easy it is to confound data. Consider the innocent scientist who visits plants in the field once a week to measure their size. To measure the plants, he has to touch them. A study of six plant species found that one touch a week significantly increased leaf damage by insects in two species and significantly decreased damage in another species.



EXAMPLE 16.1 The psychologist and the sociologist

A psychologist is interested in how our visual perception can be fooled by optical illusions. Her subjects are students in Psychology 101 at her university. Most psychologists would agree that it's safe to treat the students as an SRS of all people with normal vision. There is nothing special about being a student that changes visual perception.

A sociologist at the same university uses students in Sociology 101 to examine attitudes toward poor people and antipoverty programs. Students as a group are younger

than the adult population as a whole. Even among young people, students as a group come from more prosperous and better-educated homes. Even among students, this university isn't typical of all campuses. Even on this campus, students in a sociology course may have opinions that are quite different from those of engineering students. The sociologist can't reasonably act as if these students are a random sample from any interesting population. ■

Our first examples of inference, using the z procedures, act as if the data are an SRS from the population of interest. Let's look back at the examples in Chapters 14 and 15.

EXAMPLE 16.2 Is it really an SRS?

The NHANES survey that produced the BMI data for Example 14.1 used a complex multistage sample design, so it's a bit oversimplified to treat the BMI data as coming from an SRS from the population of young women.¹ Although the overall effect of the NHANES sample is close to an SRS, professional statisticians would use more complex inference procedures to match the more complex design of the sample.

The 20 patrons in the tipping study in Example 14.3 were chosen from those eating at a particular restaurant to receive one of several treatments being compared in a randomized comparative experiment. Recall that each treatment group in a completely randomized experiment is an SRS of the available subjects. Researchers sometimes act as if the available subjects are an SRS from some population if there is nothing special about where the subjects came from. In some cases, researchers collect demographic data on subjects to help justify the assumption that the subjects are a representative sample from some population. We are willing to regard the subjects as an SRS from the population of patrons of this particular restaurant, but perhaps this needs to be explored further.

The cola taste test in Example 15.2 uses scores from 10 tasters. All were examined to be sure that they have no medical condition that interferes with normal taste and then carefully trained to score sweetness using a set of standard drinks. We are willing to take their scores as an SRS from the population of trained tasters.

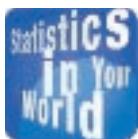
The medical director who examined executives' blood pressures in Example 15.6 actually chose an SRS from the medical records of all executives in this company. ■

These examples are typical. One is an actual SRS, two are situations in which common practice is to act as if the sample were an SRS, and in the remaining example procedures that assume an SRS are used for a quick analysis of data from a more complex random sample. *There is no simple rule for deciding when you can act as if a sample is an SRS. Pay attention to these cautions:*



- Practical problems such as nonresponse in samples or dropouts from an experiment can hinder inference even from a well-designed study. The NHANES survey has about an 80% response rate. This is much higher than opinion polls and most other national surveys, so by realistic standards NHANES data are quite trustworthy. (NHANES uses advanced methods to try to correct for nonresponse, but these methods work a lot better when response is high to start with.)

- Different methods are needed for different designs. The z procedures aren't correct for random sampling designs more complex than an SRS. Later chapters give methods for some other designs, but we won't discuss inference for really complex designs like that used by NHANES. Always be sure that you (or your statistical consultant) know how to carry out the inference your design calls for.
- There is no cure for fundamental flaws like voluntary response surveys or uncontrolled experiments. Look back at the bad examples in Chapters 8 and 9 and steel yourself to just ignore data from such studies.



Really wrong numbers

By now you know that "statistics" that don't come

from properly designed studies are often dubious and sometimes just made up. It's rare to find wrong numbers that anyone can see are wrong, but it does happen. A German physicist claimed that 2006 was the first year since 1441 with more than one Friday the 13th. Sorry: Friday the 13th occurred in February and August of 2004, which is a bit more recent than 1441.

What is the shape of the population distribution? Most statistical inference procedures require some conditions on the shape of the population distribution. Many of the most basic methods of inference are designed for Normal populations. That's the case for the z procedures and also for the more practical procedures for inference about means that we will meet in Chapters 18 and 19. Fortunately, this condition is less essential than where the data come from.

This is true because the z procedures and many other procedures designed for Normal distributions are based on Normality of the distribution of sample mean \bar{x} , not Normality of the distribution of individual observations. The central limit theorem tells us that \bar{x} is more Normal than the individual observations and that \bar{x} becomes more Normal as the size of the sample increases. In practice, the z procedures are reasonably accurate for any roughly symmetric distribution for samples of even moderate size. If the sample is large, \bar{x} will be close to Normal even if individual measurements are strongly skewed, as Figures 11.4 (page 297) and 11.5 (page 299) illustrate. Later chapters give practical guidelines for specific inference procedures.

There is one important exception to the principle that the shape of the population is less critical than how the data were produced. Outliers can distort the results of inference. Any inference procedure based on sample statistics like the sample mean \bar{x} that are not resistant to outliers can be strongly influenced by a few extreme observations.

We rarely know the shape of the population distribution. In practice we rely on previous studies and on data analysis. Sometimes long experience suggests that our data are likely to come from a roughly Normal distribution, or not. For example, heights of people of the same sex and similar ages are close to Normal, but weights are not. Always explore your data before doing inference. When the data are chosen at random from a population, the shape of the data distribution mirrors the shape of the population distribution. Make a stemplot or histogram of your data and look to see whether the shape is roughly Normal. Remember that small samples have a lot of chance variation, so that Normality is hard to judge from just a few observations. Always look for outliers and try to correct them or justify their removal before performing the z procedures or other inference based on statistics like \bar{x} that are not resistant.

When outliers are present or the data suggest that the population is strongly non-Normal, consider alternative methods that don't require Normality and are not sensitive to outliers. Some of these methods appear in Chapter 26 (available online and on the text CD).




APPLY YOUR KNOWLEDGE

- 16.1 Rate the lecture.** A professor is interested in how the 500 students in his class will rate today's lecture. He selects the first 20 students on his class list, reads the names at the beginning of the lecture, and asks them to go online to the course Web site and rate the lecture on a scale of 0 to 5. Which of the following is the most important reason why a confidence interval for the mean rating by all his students based on these data is of little use? Comment briefly on each reason to explain your answer.
- The number of students selected is small, so the margin of error will be large.
 - Most of the students selected will not respond.
 - The students selected can't be considered a random sample from the population of all students in the course.
- 16.2 Running red lights.** A survey of licensed drivers inquired about running red lights. One question asked, "Of every ten motorists who run a red light, about how many do you think will be caught?" The mean result for 880 respondents was $\bar{x} = 1.92$ and the standard deviation was $s = 1.83$.² For this large sample, s will be close to the population standard deviation σ , so suppose we know that $\sigma = 1.83$.
- Give a 95% confidence interval for the mean opinion in the population of all licensed drivers.
 - The distribution of responses is skewed to the right rather than Normal. This will not strongly affect the z confidence interval for this sample. Why not?
 - The 880 respondents are an SRS from completed calls among 45,956 calls to randomly chosen residential telephone numbers listed in telephone directories. Only 5029 of the calls were completed. This information gives two reasons to suspect that the sample may not represent all licensed drivers. What are these reasons?
- 16.3 Sampling shoppers.** A marketing consultant observes 50 consecutive shoppers at a department store the Friday after Thanksgiving, recording how much each shopper spends in the store. Suggest some reasons why it may be risky to act as if 50 consecutive shoppers at this particular time are an SRS of all shoppers at this store.



Ilene MacDonald/Alamy

CAUTIONS ABOUT CONFIDENCE INTERVALS

The most important caution about confidence intervals in general is a consequence of the use of a sampling distribution. A sampling distribution shows how a statistic such as \bar{x} varies in repeated random sampling. This variation causes *random sampling error* because the statistic misses the true parameter by a random amount. No other source of variation or bias in the sample data influences the sampling distribution. So the *margin of error in a confidence interval ignores everything except the sample-to-sample variation due to choosing the sample randomly*.



THE MARGIN OF ERROR DOESN'T COVER ALL ERRORS

The margin of error in a confidence interval covers only random sampling errors.

Practical difficulties such as undercoverage and nonresponse are often more serious than random sampling error. The margin of error does not take such difficulties into account.

Recall from Chapter 8 that national opinion polls often have response rates less than 50% and that even small changes in the wording of questions can strongly influence results. In such cases, the announced margin of error is probably unrealistically small. And of course there is no way to assign a meaningful margin of error to results from voluntary response or convenience samples, because there is no random selection. Look carefully at the details of a study before you trust a confidence interval.

APPLY YOUR KNOWLEDGE

16.4 What's your weight? A Gallup Poll asked a national random sample of 501 adult women to state their current weight. The mean weight in the sample was $\bar{x} = 159$. We will treat these data as an SRS from a Normally distributed population with standard deviation $\sigma = 35$.

- Give a 95% confidence interval for the mean weight of adult women based on these data.
- Do you trust the interval you computed in part (a) as a 95% confidence interval for the mean weight of all U.S. adult women? Why or why not?

16.5 Good weather, good tips? Example 14.3 (page 359) described an experiment exploring the size of the tip in a particular restaurant when a message indicating that the next day's weather would be good was written on the bill. You work part-time as a server in a restaurant. You read a newspaper article about the study that reports that with 95% confidence the mean percentage tip from restaurant patrons will be between 21.33 and 23.09 when the server writes a message on the bill stating that the next day's weather will be good. Can you conclude that if you begin writing a message on patron's bills that the next day's weather will be good, approximately 95% of the days you work your mean percentage tip will be between 21.33 and 23.09? Why or why not?

16.6 Sample size and margin of error. Example 14.1 (page 352) described NHANES data on the body mass index (BMI) of 654 young women. The mean BMI in the sample was $\bar{x} = 26.8$. We treated these data as an SRS from a Normally distributed population with standard deviation $\sigma = 7.5$.

- Suppose that we had an SRS of just 100 young women. What would be the margin of error for 95% confidence?
- Find the margins of error for 95% confidence based on SRSs of 400 young women and 1600 young women.
- Compare the three margins of error. How does increasing the sample size change the margin of error of a confidence interval when the confidence level and population standard deviation remain the same?

16.7 Is your food safe? “Do you feel confident or not confident that the food available at most grocery stores is safe to eat?” When a Gallup Poll asked this question, 82% of the sample said they were confident.³ Gallup announced the poll’s margin of error for 95% confidence as ± 3 percentage points. Which of the following sources of error are included in this margin of error?

- Gallup dialed landline telephone numbers at random and so missed all people without landline phones, including people whose only phone is a cell phone.



Creatas/Thinkstock

- (b) Some people whose numbers were chosen never answered the phone in several calls or answered but refused to participate in the poll.
 - (c) There is chance variation in the random selection of telephone numbers.
-

CAUTIONS ABOUT SIGNIFICANCE TESTS

Significance tests are widely used in most areas of statistical work. New pharmaceutical products require significant evidence of effectiveness and safety. Courts inquire about statistical significance in hearing class action discrimination cases. Marketers want to know whether a new package design will significantly increase sales. Medical researchers want to know whether a new therapy performs significantly better. In all these uses, statistical significance is valued because it points to an effect that is unlikely to occur simply by chance. Here are some points to keep in mind when you use or interpret significance tests.

How small a P is convincing? The purpose of a test of significance is to describe the degree of evidence provided by the sample against the null hypothesis. The P -value does this. But how small a P -value is convincing evidence against the null hypothesis? This depends mainly on two circumstances:

- *How plausible is H_0 ?* If H_0 represents an assumption that the people you must convince have believed for years, strong evidence (small P) will be needed to persuade them.
- *What are the consequences of rejecting H_0 ?* If rejecting H_0 in favor of H_a means making an expensive changeover from one type of product packaging to another, you need strong evidence that the new packaging will boost sales.

These criteria are a bit subjective. Different people will often insist on different levels of significance. Giving the P -value allows each of us to decide individually if the evidence is sufficiently strong.

Users of statistics have often emphasized standard levels of significance such as 10%, 5%, and 1%. For example, courts have tended to accept 5% as a standard in discrimination cases.⁴ This emphasis reflects the time when tables of critical values rather than software dominated statistical practice. The 5% level ($\alpha = 0.05$) is particularly common. *There is no sharp border between “significant” and “not significant,” only increasingly strong evidence as the P -value decreases. There is no practical distinction between the P -values 0.049 and 0.051. It makes no sense to treat $P \leq 0.05$ as a universal rule for what is significant.*



Significance depends on the alternative hypothesis. You may have noticed that the P -value for a one-sided test is one-half the P -value for the two-sided test of the same null hypothesis based on the same data. The two-sided P -value combines two equal areas, one in each tail of a Normal curve. The one-sided P -value is just one of these areas, in the direction specified by the alternative hypothesis. It makes sense that the evidence against H_0 is stronger when the alternative is one-sided, because it is based on the data *plus* information about

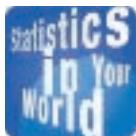
the direction of possible deviations from H_0 . If you lack this added information, always use a two-sided alternative hypothesis.

Significance depends on sample size. A sample survey shows that significantly fewer students are heavy drinkers at colleges that ban alcohol on campus. “Significantly fewer” is not enough information to decide whether there is an important difference in drinking behavior at schools that ban alcohol. *How important an effect is depends on the size of the effect as well as on its statistical significance.* If the number of heavy drinkers is only 1% less at colleges that ban alcohol than at other colleges, this is not an important effect even if it is statistically significant. In fact, the sample survey found that 38% of students at colleges that ban alcohol are “heavy episodic drinkers” compared with 48% at other colleges.⁵ That difference is large enough to be important. (Of course, this observational study doesn’t prove that an alcohol ban directly reduces drinking; it may be that colleges that ban alcohol attract more students who don’t want to drink heavily.)

Such examples remind us to always look at the size of an effect (like 38% versus 48%) as well as its significance. They also raise a question: can a tiny effect really be highly significant? Yes. The behavior of the z test statistic is typical. The statistic is

$$z = \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}}$$

The numerator measures how far the sample mean deviates from the hypothesized mean μ_0 . Larger values of the numerator give stronger evidence against $H_0: \mu = \mu_0$. The denominator is the standard deviation of \bar{x} . It measures how much random variation we expect. There is less variation when the number of observations n is large. So z gets larger (more significant) when the estimated effect $\bar{x} - \mu_0$ gets larger *or* when the number of observations n gets larger. Significance depends both on the size of the effect we observe *and* on the size of the sample. Understanding this fact is essential to understanding significance tests.



Should tests be banned?

Significance tests don't tell us how large or how

important an effect is. Research in psychology has emphasized tests, so much so that some think their weaknesses should ban them from use. The American Psychological Association asked a group of experts. They said: “Use anything that sheds light on your study. Use more data analysis and confidence intervals.” But: “The task force does not support any action that could be interpreted as banning the use of null hypothesis significance testing or P -values in psychological research and publication.”

SAMPLE SIZE AFFECTS STATISTICAL SIGNIFICANCE

Because large random samples have small chance variation, very small population effects can be highly significant if the sample is large.

Because small random samples have a lot of chance variation, even large population effects can fail to be significant if the sample is small.

Statistical significance does not tell us whether an effect is large enough to be important. That is, **statistical significance is not the same thing as practical significance.**

Keep in mind that “statistical significance” means “the sample showed an effect larger than would often occur just by chance.” The extent of chance variation changes with the size of the sample, so the size of the sample does matter. Exercise 16.9 demonstrates in detail how increasing the sample size drives down the P -value. Here is another example.

EXAMPLE 16.3 It's significant. Or not. So what?

We are testing the hypothesis of no correlation between two variables. With 1000 observations, an observed correlation of only $r = 0.08$ is significant evidence at the 1% level that the correlation in the population is not zero but positive. *The small P-value does not mean that there is a strong association, only that there is strong evidence of some association.* The true population correlation is probably quite close to the observed sample value, $r = 0.08$. We might well conclude that for practical purposes we can ignore the association between these variables, even though we are confident (at the 1% level) that the correlation is positive.



On the other hand, if we have only 10 observations, a correlation of $r = 0.5$ is not significantly greater than zero even at the 5% level. Small samples vary so much that a large r is needed if we are to be confident that we aren't just seeing chance variation at work. So a small sample will often fall short of significance even if the true population correlation is quite large. ■

Beware of multiple analyses. Statistical significance ought to mean that you have found an effect that you were looking for. The reasoning behind statistical significance works well if you decide what effect you are seeking, design a study to search for it, and use a test of significance to weigh the evidence you get. In other settings, significance may have little meaning.

EXAMPLE 16.4 Cell phones and brain cancer

Might the radiation from cell phones be harmful to users? Many studies have found little or no connection between using cell phones and various illnesses. Here is part of a news account of one study:

A hospital study that compared brain cancer patients and a similar group without brain cancer found no statistically significant association between cell phone use and a group of brain cancers known as gliomas. But when 20 types of glioma were considered separately an association was found between phone use and one rare form. Puzzlingly, however, this risk appeared to decrease rather than increase with greater mobile phone use.⁶

Think for a moment. Suppose that the 20 null hypotheses (no association) for these 20 significance tests are all true. Then each test has a 5% chance of being significant at the 5% level. That's what $\alpha = 0.05$ means: results this extreme occur 5% of the time just by chance when the null hypothesis is true. Because 5% is $1/20$, we expect about 1 of 20 tests to give a significant result just by chance. That's what the study observed. ■



iStockphoto

Running one test and reaching the 5% level of significance is reasonably good evidence that you have found something. Running 20 tests and reaching that level only once is not. The caution about multiple analyses applies to confidence intervals as well. A single 95% confidence interval has probability 0.95 of capturing the true parameter each time you use it. The probability that all of



20 confidence intervals will capture their parameters is much less than 95%. If you think that multiple tests or intervals may have discovered an important effect, you need to gather new data to do inference about that specific effect.

APPLY YOUR KNOWLEDGE



- 16.8 Is it significant?** In the absence of special preparation, SAT Mathematics (SATM) scores in 2009 varied Normally with mean $\mu = 515$ and $\sigma = 116$. Fifty students go through a rigorous training program designed to raise their SATM scores by improving their mathematics skills. Either by hand or by using the *P-Value of a Test of Significance* applet, carry out a test of

$$H_0: \mu = 515$$

$$H_a: \mu > 515$$

(with $\sigma = 116$) in each of the following situations:

- (a) The students' average score is $\bar{x} = 541$. Is this result significant at the 5% level?
- (b) The average score is $\bar{x} = 542$. Is this result significant at the 5% level?

The difference between the two outcomes in (a) and (b) is of no importance. Beware attempts to treat $\alpha = 0.05$ as sacred.

- 16.9 Detecting acid rain.** Emissions of sulfur dioxide by industry set off chemical changes in the atmosphere that result in “acid rain.” The acidity of liquids is measured by pH on a scale of 0 to 14. Distilled water has pH 7.0, and lower pH values indicate acidity. Normal rain is somewhat acidic, so acid rain is sometimes defined as rainfall with a pH below 5.0. Suppose that pH measurements of rainfall on different days in a Canadian forest follow a Normal distribution with standard deviation $\sigma = 0.5$. A sample of n days finds that the mean pH is $\bar{x} = 4.8$. Is this good evidence that the mean pH μ for all rainy days is less than 5.0? The answer depends on the size of the sample.



Either by hand or using the *P-Value of a Test of Significance* applet, carry out three tests of

$$H_0: \mu = 5.0$$

$$H_a: \mu < 5.0$$

Use $\sigma = 0.5$ and $\bar{x} = 4.8$ in all three tests. But use three different sample sizes, $n = 5$, $n = 15$, and $n = 40$.

- (a) What are the *P*-values for the three tests? *The P-value of the same result $\bar{x} = 4.8$ gets smaller (more significant) as the sample size increases.*
- (b) For each test, sketch the Normal curve for the sampling distribution of \bar{x} when H_0 is true. This curve has mean 5.0 and standard deviation $0.5/\sqrt{n}$. Mark the observed $\bar{x} = 4.8$ on each curve. (If you use the applet, you can just copy the curves displayed by the applet.) *The same result $\bar{x} = 4.8$ gets more extreme on the sampling distribution as the sample size increases.*

- 16.10 Confidence intervals help.** Give a 95% confidence interval for the mean pH μ for each sample size in the previous exercise. The intervals, unlike the *P*-values, give a clear picture of what mean pH values are plausible for each sample.

16.11 Searching for ESP. A researcher looking for evidence of extrasensory perception (ESP) tests 500 subjects. Four of these subjects do significantly better ($P < 0.01$) than random guessing.

- You can't conclude that these 4 people have ESP. Why not?
- What should the researcher now do to test whether any of these 4 subjects have ESP?

PLANNING STUDIES: SAMPLE SIZE FOR CONFIDENCE INTERVALS

A wise user of statistics never plans a sample or an experiment without at the same time planning the inference. The number of observations is a critical part of planning a study. Larger samples give smaller margins of error in confidence intervals and make significance tests better able to detect effects in the population. But taking observations costs both time and money. How many observations are enough? We will look at this question first for confidence intervals and then for tests. Planning a confidence interval is much simpler than planning a test. It is also more useful, because estimation is generally more informative than testing. The section on planning tests is therefore optional.

You can arrange to have both high confidence and a small margin of error by taking enough observations. The margin of error of the z confidence interval for the mean of a Normally distributed population is $m = z^* \sigma / \sqrt{n}$. To obtain a desired margin of error m , put in the value of z^* for your desired confidence level, and solve for the sample size n . Here is the result.

SAMPLE SIZE FOR DESIRED MARGIN OF ERROR

The z confidence interval for the mean of a Normal population will have a specified margin of error m when the sample size is

$$n = \left(\frac{z^* \sigma}{m} \right)^2$$

Notice that it is the size of the sample that determines the margin of error. The size of the population does not influence the sample size we need. (This is true as long as the population is much larger than the sample.)



EXAMPLE 16.5 How many observations?

In Example 14.3 (page 359), psychologists recorded the size of the tip of 20 patrons in a restaurant when a message indicating that the next day's weather would be good was written on their bill. We know that the population standard deviation is $\sigma = 2$. We want to estimate the mean percentage tip μ for patrons of this restaurant who receive this message on their bill within ± 0.5 with 90% confidence. How many patrons must we observe?

The desired margin of error is $m = 0.5$. For 90% confidence, Table C gives $z^* = 1.645$. Therefore,

$$n = \left(\frac{z^* \sigma}{m} \right)^2 = \left(\frac{1.645 \times 2}{0.5} \right)^2 = 43.3$$

Because 43 patrons will give a slightly larger margin of error than desired, and 44 patrons a slightly smaller margin of error, we must observe 44 patrons. *Always round up to the next higher whole number when finding n.* ■

APPLY YOUR KNOWLEDGE

16.12 Body mass index of young women. Example 14.1 (page 352) assumed that the body mass index (BMI) of all American young women follows a Normal distribution with standard deviation $\sigma = 7.5$. How large a sample would be needed to estimate the mean BMI μ in this population to within ± 1 with 95% confidence?

16.13 Number skills of young men. Suppose that scores on the mathematics part of the National Assessment of Educational Progress (NAEP) test for high school seniors follow a Normal distribution with standard deviation $\sigma = 30$. You want to estimate the mean score within ± 10 with 90% confidence. How large an SRS of scores must you choose?

PLANNING STUDIES: THE POWER OF A STATISTICAL TEST*

How large a sample should we take when we plan to carry out a test of significance? We know that if our sample is too small, even large effects in the population will often fail to give statistically significant results. Here are the questions we must answer to decide how many observations we need:

Significance level. How much protection do we want against getting a significant result from our sample when there really is no effect in the population?

Effect size. How large an effect in the population is important in practice?

Power. How confident do we want to be that our study will detect an effect of the size we think is important?

The three boldface terms are statistical shorthand for three pieces of information. Power is a new idea.

EXAMPLE 16.6 Sweetening colas: planning a study

Let's illustrate typical answers to these questions in the example of testing a new cola for loss of sweetness in storage (Example 15.2, page 370). Ten trained tasters rated the sweetness on a 10-point scale before and after storage, so that we have each taster's judgment of loss of sweetness. From experience, we know that sweetness loss scores

*Power calculations are important in planning studies, but this more advanced material is not needed to read the rest of the book.

vary from taster to taster according to a Normal distribution with standard deviation about $\sigma = 1$. To see if the taste test gives reason to think that the cola does lose sweetness, we will test

$$\begin{aligned} H_0: \mu &= 0 \\ H_a: \mu &> 0 \end{aligned}$$

Are 10 tasters enough, or should we use more?

Significance level. Requiring significance at the 5% level is enough protection against declaring there is a loss in sweetness when in fact there is no change if we could look at the entire population. This means that when there is no change in sweetness in the population, 1 out of 20 samples of tasters will wrongly find a significant loss.

Effect size. Researchers know that a mean sweetness loss of 0.8 point on the 10-point scale will be noticed by consumers and so is important in practice.

Power. We want a high level of confidence (90%) that our test will detect a mean loss of 0.8 point in the population of all tasters. We agreed to use significance at the 5% level as our standard for detecting an effect. So we want probability at least 0.9 that a test at the $\alpha = 0.05$ level will reject the null hypothesis $H_0: \mu = 0$ when the true population mean is $\mu = 0.8$. ■

The probability that the test successfully detects a sweetness loss of the specified size is the *power* of the test. You can think of tests with high power as being highly sensitive to deviations from the null hypothesis. In Example 16.6, we decided that we want power 90% when the truth about the population is that $\mu = 0.8$.

POWER

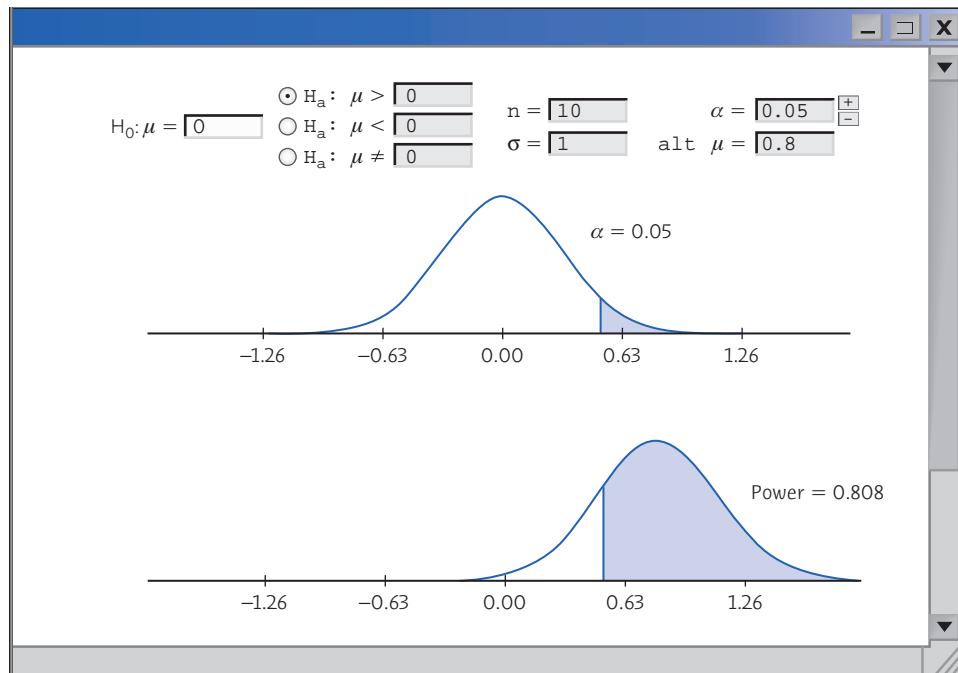
The **power** of a test against a specific alternative is the probability that the test will reject H_0 at a chosen significance level α when the specified alternative value of the parameter is true.

For most statistical tests, calculating power is a job for comprehensive statistical software. The z test is easier, but we will nonetheless skip the details. The two following examples illustrate two approaches: an applet that shows the meaning of power, and statistical software.

EXAMPLE 16.7 Finding power: use an applet

Finding the power of the z test is less challenging than most other power calculations because it requires only a Normal distribution probability calculation. The *Power of a Test* applet does this and illustrates the calculation with Normal curves. Enter the information from Example 16.6 into the applet: hypotheses, significance level $\alpha = 0.05$, alternative value $\mu = 0.8$, standard deviation $\sigma = 1$, and sample size $n = 10$. Click “Update.” The applet output appears in Figure 16.1.



**FIGURE 16.1**

Output from the *Power of a Test* applet for Example 16.7, along with the information entered into the applet. The top curve shows the behavior of \bar{x} when the null hypothesis is true ($\mu = 0$). The bottom curve shows the distribution of \bar{x} when $\mu = 0.8$.

The power of the test against the specific alternative $\mu = 0.8$ is 0.808. That is, the test will reject H_0 about 81% of the time when this alternative is true. So 10 observations are too few to give power 90%. ■

The two Normal curves in Figure 16.1 show the sampling distribution of \bar{x} under the null hypothesis $\mu = 0$ (top) and also under the specific alternative $\mu = 0.8$ (bottom). The curves have the same shape because σ does not change. The top curve is centered at $\mu = 0$ and the bottom curve at $\mu = 0.8$. The shaded region at the right of the top curve has area 0.05. It marks off values of \bar{x} that are statistically significant at the $\alpha = 0.05$ level. The lower curve shows the probability of these same values when $\mu = 0.8$. This area is the power, 0.808.

The applet will find the power for any given sample size. It's more helpful in practice to turn the process around and learn what sample size we need to achieve a given power. Statistical software will do this but usually doesn't show the helpful Normal curves that are part of the applet's output.

EXAMPLE 16.8 Finding power: use software

Some software packages (for example, SAS, JMP, Minitab, and R) will calculate power. We asked Minitab to find the number of observations needed for the one-sided z test to have power 0.9 against several specific alternatives at the 5% signifi-

cance level when the population standard deviation is $\sigma = 1$. Here is the table that results:

Difference	Sample Size	Target Power	Actual Power
0.1	857	0.9	0.900184
0.2	215	0.9	0.901079
0.3	96	0.9	0.902259
0.4	54	0.9	0.902259
0.5	35	0.9	0.905440
0.6	24	0.9	0.902259
0.7	18	0.9	0.907414
0.8	14	0.9	0.911247
0.9	11	0.9	0.909895
1.0	9	0.9	0.912315

In this output, “Difference” is the difference between the null hypothesis value $\mu = 0$ and the alternative we want to detect. This is the effect size. The “Sample Size” column shows the smallest number of observations needed for power 0.9 against each effect size.

We see again that our earlier sample of 10 tasters is not large enough to be 90% confident of detecting (at the 5% significance level) an effect of size 0.8. If we want power 90% against effect size 0.8, we need at least 14 tasters. The actual power with 14 tasters is 0.911247.

Statistical software, unlike the applet, will do power calculations for most of the tests in this book. ■

The table in Example 16.8 makes it clear that smaller effects require larger samples to reach power 90%. Here is an overview of influences on “How large a sample do I need?”

- If you insist on a smaller significance level (such as 1% rather than 5%), you will need a larger sample. A smaller significance level requires stronger evidence to reject the null hypothesis.
- If you insist on higher power (such as 99% rather than 90%), you will need a larger sample. Higher power gives a better chance of detecting an effect when it is really there.
- At any significance level and desired power, a two-sided alternative requires a larger sample than a one-sided alternative.
- At any significance level and desired power, detecting a small effect requires a larger sample than detecting a large effect.

Planning a serious statistical study always requires an answer to the question “How large a sample do I need?” If you intend to test the hypothesis $H_0: \mu = \mu_0$ about the mean μ of a population, you need at least a rough idea of the size of the population standard deviation σ and of how big a deviation $\mu - \mu_0$ of the population mean from its hypothesized value you want to be able to detect. More



Fish, fishermen, and power

Are the stocks of cod in the ocean off eastern Canada declining? Studies over many years failed to find significant evidence of a decline. These studies had low power—that is, they might fail to find a decline even if one were present. When it became clear that the cod were vanishing, quotas on fishing ravaged the economy in parts of Canada. If the earlier studies had had high power, they would likely have detected the decline. Quick action might have reduced the economic and environmental costs.

elaborate settings, such as comparing the mean effects of several treatments, require more elaborate advance information. You can leave the details to experts, but you should understand the idea of power and the factors that influence how large a sample you need.

To calculate the power of a test, we act as if we are interested in a fixed level of significance such as $\alpha = 0.05$. That's essential to do a power calculation, but remember that in practice we think in terms of *P*-values rather than a fixed level α . To effectively plan a statistical test we must find the power for several significance levels and for a range of sample sizes and effect sizes to get a full picture of how the test will behave.

Type I and Type II errors in significance tests. We can assess the performance of a test by giving two probabilities: the significance level α and the power for an alternative that we want to be able to detect. The significance level of a test is the probability of reaching the *wrong* conclusion when the null hypothesis is true. The power for a specific alternative is the probability of reaching the *right* conclusion when that alternative is true. We can just as well describe the test by giving the probabilities of being *wrong* under both conditions.

TYPE I AND TYPE II ERRORS

If we reject H_0 when in fact H_0 is true, this is a **Type I error**.

If we fail to reject H_0 when in fact H_a is true, this is a **Type II error**.

The **significance level α** of any fixed level test is the probability of a Type I error.

The **power** of a test against any alternative is 1 minus the probability of a Type II error for that alternative.

The possibilities are summed up in Figure 16.2. If H_0 is true, our conclusion is correct if we fail to reject H_0 and is a Type I error if we reject H_0 . If H_a is true, our conclusion is either correct or a Type II error. Only one error is possible at one time.

		Truth about the population	
		H_0 true	H_a true
Conclusion based on sample	Reject H_0	Type I error	Correct conclusion
	Fail to reject H_0	Correct conclusion	Type II error

FIGURE 16.2

The two types of error in testing hypotheses.

EXAMPLE 16.9 Calculating error probabilities

Because the probabilities of the two types of error are just a rewording of significance level and power, we can see from Figure 16.1 what the error probabilities are for the test in Example 16.6.

$$\begin{aligned} P(\text{Type I error}) &= P(\text{reject } H_0 \text{ when in fact } \mu = 0) \\ &= \text{significance level } \alpha = 0.05 \end{aligned}$$

$$\begin{aligned} P(\text{Type II error}) &= P(\text{fail to reject } H_0 \text{ when in fact } \mu = 0.8) \\ &= 1 - \text{power} = 1 - 0.808 = 0.192 \end{aligned}$$

The two Normal curves in Figure 16.1 are used to find the probabilities of a Type I error (top curve, $\mu = 0$) and of a Type II error (bottom curve, $\mu = 0.8$). ■



APPLY YOUR KNOWLEDGE

16.14 What is power? You manufacture and sell an iron rod whose electrical conductivity is supposed to be 10.1. You plan to make 6 measurements of the conductivity of a rod you plan to sell. You know that the standard deviation of your measurements is $\sigma = 0.1$. If the product meets specifications, the mean of many measurements will be 10.1. You will therefore test

$$H_0: \mu = 10.1$$

$$H_a: \mu \neq 10.1$$

If the true conductivity is 10.15, the rod is not suitable for its intended use. You learn that the power of your test at the 5% significance level against the alternative $\mu = 10.15$ is 0.24.

- (a) Explain in simple language what “power = 0.24” means.
- (b) Explain why the test you plan will not adequately protect you against selling a rod with conductivity 10.15.

16.15 Thinking about power. Answer these questions in the setting of the previous exercise about measuring the conductivity of an iron rod.

- (a) You could get higher power against the same alternative with the same α by changing the number of measurements you make. Should you make more measurements or fewer to increase power?
- (b) If you decide to use $\alpha = 0.10$ in place of $\alpha = 0.05$, with no other changes in the test, will the power increase or decrease?
- (c) If you shift your interest to the alternative $\mu = 10.2$ with no other changes, will the power increase or decrease?

16.16 How power behaves. In the setting of Exercise 16.14, use the *Power of a Test* applet to find the power in each of the following circumstances. Be sure to set the applet to the two-sided alternative.

- (a) Standard deviation $\sigma = 0.1$, significance level $\alpha = 0.05$, alternative $\mu = 10.15$, and sample sizes $n = 6$, $n = 12$, and $n = 24$. How does increasing the sample size with no other changes affect the power?



- (b) Standard deviation $\sigma = 0.1$, significance level $\alpha = 0.05$, sample size $n = 6$, and alternatives $\mu = 10.15$, $\mu = 10.2$, and $\mu = 10.25$. How do alternatives more distant from the hypothesis (larger effect sizes) affect the power?
- (c) Standard deviation $\sigma = 0.1$, sample size $n = 6$, alternative $\mu = 10.15$, and significance levels $\alpha = 0.05$, $\alpha = 0.10$, and $\alpha = 0.25$. (Click the + and – buttons to change α .) How does increasing the desired significance level affect the power?



16.17 How power behaves. Another approach to improving the unsatisfactory power of the test in Exercise 16.14 is to improve the measurement process. That is, use a measurement process that is less variable. Use the *Power of a Test* applet to find the power of the test in Exercise 16.14 in each of these circumstances: significance level $\alpha = 0.05$, alternative $\mu = 10.15$, sample size $n = 6$, and $\sigma = 0.1$, $\sigma = 0.05$, and $\sigma = 0.025$. How does decreasing the variability of the population of measurements affect the power?

16.18 Two types of error. Your company markets a computerized medical diagnostic program used to evaluate thousands of people. The program scans the results of routine medical tests (pulse rate, blood tests, etc.) and refers the case to a doctor if there is evidence of a medical problem. The program makes a decision about each person.

- (a) What are the two hypotheses and the two types of error that the program can make? Describe the two types of error in terms of “false-positive” and “false-negative” test results.
- (b) The program can be adjusted to decrease one error probability, at the cost of an increase in the other error probability. Which error probability would you choose to make smaller, and why? (This is a matter of judgment. There is no single correct answer.)

CHAPTER 16 SUMMARY

CHAPTER SPECIFICS

- A specific confidence interval or test is correct only under specific conditions. The most important conditions concern the method used to produce the data. Other factors such as the shape of the population distribution may also be important.
- Whenever you use statistical inference, you are acting as if your data are a random sample or come from a randomized comparative experiment.
- Always do data analysis before inference to detect outliers or other problems that would make inference untrustworthy.
- The margin of error in a confidence interval accounts for only the chance variation due to random sampling. In practice, errors due to nonresponse or undercoverage are often more serious.
- There is no universal rule for how small a P -value in a test of significance is convincing evidence against the null hypothesis. Beware of placing too much weight on traditional significance levels such as $\alpha = 0.05$.
- Very small effects can be highly significant (small P) when a test is based on a large sample. A statistically significant effect need not be practically important. Plot the

data to display the effect you are seeking, and use confidence intervals to estimate the actual values of parameters.

- On the other hand, lack of significance does not imply that H_0 is true. Even a large effect can fail to be significant when a test is based on a small sample.
- Many tests run at once will probably produce some significant results by chance alone, even if all the null hypotheses are true.
- When you plan a statistical study, plan the inference as well. In particular, ask what sample size you need for successful inference.
- The z confidence interval for a Normal mean has specified margin of error m when the sample size is

$$n = \left(\frac{z^* \sigma}{m} \right)^2$$

Here z^* is the critical value for the desired level of confidence. Always round n up when you use this formula.

- The **power** of a significance test measures its ability to detect an alternative hypothesis. The power against a specific alternative is the probability that the test will reject H_0 at a particular level α when that alternative is true.
- Increasing the size of the sample increases the power of a significance test. You can use statistical software to find the sample size needed to achieve a desired power.

LINK IT

In Chapters 14 and 15 we introduced the basic reasoning behind statistical estimation and tests of significance. We applied this reasoning to the problem of making inferences about a mean of a Normally distributed population in a simple and artificial setting (the population standard deviation is known). In this chapter we begin to move away from our simple and artificial conditions toward the reality of statistical practice. Later chapters deal with inference in fully realistic settings. Here we discuss more carefully conditions under which our procedures for inference do and do not hold. Where the data come from is crucial. Even if the population distribution is not Normal, our procedures are approximately correct if we have a moderate sample size and the shape of the population distribution is roughly symmetric. If the sample size is large, our procedures are approximately correct even if the population distribution is strongly skewed. However, any procedure based on sample statistics like the sample mean that are not resistant to outliers can be influenced by a few extreme observations. This is a caution that we first encountered in Chapter 4.

In this chapter we also provide some cautions about confidence intervals and tests of significance. Some of these echo statements made in Chapters 8 and 9. Some are based on the behavior of confidence intervals and significance tests. These cautions will help you evaluate studies that report confidence intervals or the results of a test of significance.

We conclude this chapter with material about planning a study that will use a confidence interval or a significance test. The fundamental question is “What sample size do I need for successful inference?”


CHECK YOUR SKILLS

16.19 The most important condition for sound conclusions from statistical inference is usually

- (a) that the data can be thought of as a random sample from the population of interest.
- (b) that the population distribution is exactly Normal.
- (c) that the data contain no outliers.

16.20 The coach of a college men's soccer team records the resting heart rates of the 27 team members. You should not trust a confidence interval for the mean resting heart rate of all male students at this college based on these data because

- (a) with only 27 observations, the margin of error will be large.
- (b) heart rates may not have a Normal distribution.
- (c) the members of the soccer team can't be considered a random sample of all students.

16.21 You turn your Web browser to the online Harris Interactive Poll. Based on 2163 responses, the poll reports that 43% of U.S. adults said they would like to be richer, 21% said thinner, 14% said smarter, and 12% said younger. Nearly 1 in 10 (9%) said they would not want to choose any of the given options.⁷ You should refuse to calculate a 95% confidence interval for the proportion of all U.S. adults who would like to be richer based on this sample because

- (a) the poll was taken a week ago.
- (b) inference from a voluntary response sample can't be trusted.
- (c) the sample is too large.

16.22 Many sample surveys use well-designed random samples but half or more of the original sample can't be contacted or refuse to take part. Any errors due to this nonresponse

- (a) have no effect on the accuracy of confidence intervals.
- (b) are included in the announced margin of error.
- (c) are in addition to the random variation accounted for by the announced margin of error.

16.23 A writer in a medical journal says: "An uncontrolled experiment in 37 women found a significantly improved mean clinical symptom score after treatment. Methodologic flaws make it difficult to interpret the results of this study." The writer is skeptical about the significant improvement because

- (a) there is no control group, so the improvement might be due to the placebo effect or to the fact that many medical conditions improve over time.

(b) the *P*-value given was $P = 0.048$, which is too large to be convincing.

- (c) the response variable might not have an exactly Normal distribution in the population.

16.24 Vigorous exercise helps people live several years longer (on the average). Whether mild activities like slow walking extend life is not clear. Suppose that the added life expectancy from regular slow walking is just 2 months. A statistical test is more likely to find a significant increase in mean life if

- (a) it is based on a very large random sample.
- (b) it is based on a very small random sample.
- (c) The size of the sample doesn't have any effect on the significance of the test.

16.25 A medical experiment compared the herb echinacea with a placebo for preventing colds. One response variable was "volume of nasal secretions" (if you have a cold, you blow your nose a lot). Take the average volume of nasal secretions in people without colds to be $\mu = 1$. An increase to $\mu = 3$ indicates a cold. The significance level of a test of $H_0: \mu = 1$ versus $H_a: \mu > 1$ is defined as

- (a) the probability that the test rejects H_0 when $\mu = 1$ is true.
- (b) the probability that the test rejects H_0 when $\mu = 3$ is true.
- (c) the probability that the test fails to reject H_0 when $\mu = 3$ is true.

16.26 (Optional) The power of the test in the previous exercise against the specific alternative $\mu = 3$ is defined as

- (a) the probability that the test rejects H_0 when $\mu = 1$ is true.
- (b) the probability that the test rejects H_0 when $\mu = 3$ is true.
- (c) the probability that the test fails to reject H_0 when $\mu = 3$ is true.

16.27 (Optional) The power of a test is important in practice because power

- (a) describes how well the test performs when the null hypothesis is actually true.
- (b) describes how sensitive the test is to violations of conditions such as Normal population distribution.
- (c) describes how well the test performs when the null hypothesis is actually not true.

CHAPTER 16 EXERCISES

16.28 Hotel managers. In Exercise 15.30 (page 386) you carried out a test of significance based on data from 148 general managers of three-star and four-star hotels. Before you trust your results, you would like more information about the data. What facts would you most like to know?

16.29 Color blindness in Africa. An anthropologist claims that color blindness is less common in societies that live by hunting and gathering than in settled agricultural societies. He tests a number of adults in two populations in Africa, one of each type. The proportion of color-blind people is significantly lower ($P < 0.05$) in the hunter-gatherer population. What additional information would you want to help you decide whether you believe the anthropologist's claim?

16.30 Sampling at the mall. A market researcher chooses at random from women entering a large suburban shopping mall. One outcome of the study is a 95% confidence interval for the mean of "the highest price you would pay for a pair of jeans."

- (a) Explain why this confidence interval does not give useful information about the population of all women.
- (b) Explain why it may give useful information about the population of women who shop at large suburban malls.

16.31 Sensitive questions. The National AIDS Behavioral Surveys found that 320 individuals in its random sample of 1334 urban heterosexuals aged 18 to 25 years said that they had multiple sexual partners in the past year. That's 24% of the sample. Why is this estimate likely to be biased? Do you think it is biased high or low? Does the margin of error of a 95% confidence interval for the proportion of all urban heterosexuals aged 18 to 25 with multiple partners allow for this bias?

16.32 College degrees. At the Statistics Canada Web site, www.statcan.gc.ca, you can find the percent of adults in each province or territory who have at least a university certificate, diploma, or degree at bachelor's level or above. It makes no sense to find \bar{x} for these data and use it to get a confidence interval for the mean percent μ in all 13 provinces or territories. Why not?

16.33 An outlier strikes. You have data on an SRS of recent graduates from your college that shows how long each student took to complete a bachelor's degree. The data contain one high outlier. Will this outlier have a greater effect on a confidence interval for mean completion time if your sample is small or if it is large? Why?

16.34 Can we trust this interval? Here are data on the percent change in the total mass (in tons) of wildlife in several West African game preserves in the years 1971 to 1999.⁸

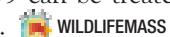


Adolfo López Pérez/Photolibrary

1971	1972	1973	1974	1975	1976	1977	1978
2.9	3.1	-1.2	-1.1	-3.3	3.7	1.9	-0.3
1979	1980	1981	1982	1983	1984	1985	1986
-5.9	-7.9	-5.5	-7.2	-4.1	-8.6	-5.5	-0.7
1987	1988	1989	1990	1991	1992	1993	1994
-5.1	-7.1	-4.2	0.9	-6.1	-4.1	-4.8	-11.3
1995	1996	1997	1998	1999			
-9.3	-10.7	-1.8	-7.4	-22.9			

Software gives the 95% confidence interval for the mean annual percent change as -6.66% to -2.55% . There are several reasons why we might not trust this interval.

- (a) Examine the distribution of the data. What feature of the distribution throws doubt on the validity of statistical inference?
- (b) Plot the percents against year. What trend do you see in this time series? Explain why a trend over time casts doubt on the condition that years 1971 to 1999 can be treated as an SRS from a larger population of years.



16.35 When to use pacemakers. A medical panel prepared guidelines for when cardiac pacemakers should be implanted in patients with heart problems. The panel reviewed a large number of medical studies to judge the strength of the evidence supporting each recommendation. For each recommendation, they ranked the evidence as level A (strongest), B, or C (weakest). Here, in



Layne Kennedy/CORBIS

scrambled order, are the panel's descriptions of the three levels of evidence.⁹ Which is A, which B, and which C? Explain your ranking.

Evidence was ranked as level _____ when data were derived from a limited number of trials involving comparatively small numbers of patients or from well-designed data analysis of nonrandomized studies or observational data registries.

Evidence was ranked as level _____ if the data were derived from multiple randomized clinical trials involving a large number of individuals.

Evidence was ranked as level _____ when consensus of expert opinion was the primary source of recommendation.

16.36 What is significance good for? Which of the following questions does a test of significance answer? Briefly explain your replies.

- Is the sample or experiment properly designed?
- Is the observed effect due to chance?
- Is the observed effect important?

16.37 Why are larger samples better? Statisticians prefer large samples. Describe briefly the effect of increasing the size of a sample (or the number of subjects in an experiment) on each of the following:

- The P -value of a test, when H_0 is false and all facts about the population remain unchanged as n increases.
- (Optional) The power of a fixed level α test, when α , the alternative hypothesis, and all facts about the population remain unchanged.

16.38 Divorce rates. Divorce rates vary from city to city in the United States. We have lots of data on many U.S. cities. Statistical software makes it easy to perform dozens of significance tests on dozens of variables to see which ones best predict divorce rate. One interesting finding is that those cities with major league ballparks tend to have significantly lower divorce rates than other cities. To improve your chances of a successful marriage, should you use this "significant" variable to decide where to live? Explain your answer.

16.39 A test goes wrong. Software can generate samples from (almost) exactly Normal distributions. Here is a random sample of size 5 from the Normal distribution with mean 8 and standard deviation 2:

4.47 5.51 8.10 11.63 7.91

These data match the conditions for a z test better than real data will: the population is very close to Normal and has

known standard deviation $\sigma = 2$, and the population mean is $\mu = 8$. Although we know the true value of μ , suppose we pretend that we do not and we test the hypotheses

$$H_0: \mu = 6$$

$$H_a: \mu \neq 6$$

- What are the z statistic and its P -value? Is the test significant at the 5% level?
- We know that the null hypothesis does not hold, but the test failed to give strong evidence against H_0 . Explain why this is not surprising.  RANDOMSAMPLE

16.40 Reducing the gender gap. In many science disciplines women are outperformed by men on test scores. Will "values affirmation training" improve self-confidence and hence performance of women relative to men in science courses? A study conducted at a large university compares the scores of men and women at the end of a large introductory physics course on a nationally normed standardized test of conceptual physics, the Force and Motion Conceptual Evaluation (FMCE). Half the women in the course were given values affirmation training during the course; the other half received no training. The study reports that there was a significant difference ($P < 0.01$) in the gap between men's and women's scores, although the gap for women who received the values affirmation training was much smaller than that for women who did not receive training. The study also reports that a 95% confidence interval for the mean difference in scores on the FMCE exam between women who received the training and those who didn't is 13 ± 8 points. You are a faculty member in the physics department, and the provost, who is interested in women in science, asks you about the study.

- Explain in simple language what "a significant difference ($P < 0.01$)" means.
- Explain clearly and briefly what "95% confidence" means.
- Is this study good evidence that requiring values affirmation training of all female students would greatly reduce the gender gap in scores on science tests in college courses?

16.41 How far do rich parents take us? How much education children get is strongly associated with the wealth and social status of their parents. In social science jargon, this is "socioeconomic status," or SES. But the SES of parents has little influence on whether children who have graduated from college go on to yet more education. One study looked at whether college graduates took the graduate admissions tests for business, law, and other graduate programs. The effects of

the parents' SES on taking the LSAT test for law school were "both statistically insignificant and small."

- What does "statistically insignificant" mean?
- Why is it important that the effects were small in size as well as insignificant?

16.42 This wine stinks. How sensitive are the untrained noses of students? Exercise 14.27 (page 366) gives the lowest levels of dimethyl sulfide (DMS) that 10 students could detect. You want to estimate the mean DMS odor threshold among all students, and you would be satisfied to estimate the mean to within ± 0.1 with 99% confidence. The standard deviation of the odor threshold for untrained noses is known to be $\sigma = 7$ micrograms per liter of wine. How large an SRS of untrained students do you need?

16.43 Pulling wood apart. You want to estimate the mean load needed to pull apart the pieces of wood in Exercise 14.25 (page 366) to within ± 600 pounds with 95% confidence. How large a sample is needed (Assume as in Exercise 14.25 that the load needed to pull apart pieces of wood follows a Normal distribution with standard deviation 3000 pounds)?

The following exercises concern the optional material on the power of a test.

16.44 The first child has higher IQ. Does the birth order of a family's children influence their IQ scores? A careful study of 241,310 Norwegian 18- and 19-year-olds found that firstborn children scored 2.3 points higher on the average than second children in the same family. This difference was highly significant ($P < 0.001$). A commentator said, "One puzzle highlighted by these latest findings is why certain other within-family studies have failed to show equally consistent results. Some of these previous null findings, which have all been obtained in much smaller samples, may be explained by inadequate statistical power."¹⁰ Explain in simple language why tests having low power often fail to give evidence against a null hypothesis even when the hypothesis is really false.

16.45 How valium works. Valium is a common antidepressant and sedative. A study investigated how valium works by comparing its effect on sleep in 7 genetically modified mice and 8 normal control mice. There was no significant difference between the two groups. The authors say that this lack of significance "is related to the large inter-individual variability that is also reflected in the low power (20%) of the test."¹¹

- Explain exactly what power 20% against a specific alternative means.
- Explain in simple language why tests having low power often fail to give evidence against a null hypothesis even when the null hypothesis is really false.

- What fact about this experiment most likely explains the low power?

16.46 Dialysis. An article in the *New England Journal of Medicine* describes a randomized controlled trial that compared early versus late initiation of dialysis on the survival of adults with progressive chronic kidney disease. The experiment found no significant difference between early and late initiation of dialysis. According to the article, the study was designed to have power 80%, with a two-sided Type I error of 0.05, to detect a clinically important difference of approximately 10 percentage points in the absolute risk of death.¹²



Phanie/Photo Researchers

- What fixed significance level was used in calculating the power?
- Explain to someone who knows no statistics why power 80% means that the experiment would probably have been significant if there was a difference between early and late initiation of dialysis.

16.47 Power. In Exercise 16.39, a sample from a Normal population with mean $\mu = 8$ and standard deviation $\sigma = 2$ failed to reject the null hypothesis $H_0: \mu = 6$ at the $\alpha = 0.05$ significance level. Enter the information from this example into the *Power of a Test* applet. (Don't forget that the alternative hypothesis is two-sided.) What is the power of the test against the alternative $\mu = 10$? Because the power is not high, it isn't surprising that the sample in Exercise 16.39 failed to reject H_0 .

16.48 Finding power by hand. Even though software is used in practice to calculate power, doing the work by hand builds your understanding. Return to the test in Example 16.6 (page 402). There are $n = 10$ observations from a population with standard deviation $\sigma = 1$ and unknown mean μ . We will test

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

with fixed significance level $\alpha = 0.05$. Find the power against the alternative $\mu = 0.8$ by following these steps.

- The z test statistic is

$$z = \frac{\bar{x} - \mu_0}{\sigma/\sqrt{n}} = \frac{\bar{x} - 0}{1/\sqrt{10}} = 3.162\bar{x}$$

(Remember that you won't know the numerical value of \bar{x} until you have data.) What values of z lead to rejecting H_0 at the 5% significance level?

(b) Starting from your result in (a), what values of \bar{x} lead to rejecting H_0 ? The area above these values is shaded under the top curve in Figure 16.1.

(c) The power is the probability that you observe any of these values of \bar{x} when $\mu = 0.8$. This is the shaded area under the bottom curve in Figure 16.1. What is this probability?

16.49 Finding power by hand: two-sided test. The previous exercise shows how to calculate the power of a one-sided z test. Power calculations for two-sided tests follow the same outline. We will find the power of a test based on 6 measurements of the conductivity of an iron rod, reported in Exercise 16.14. The hypotheses are

$$H_0: \mu = 10.1$$

$$H_a: \mu \neq 10.1$$

The population of all measurements is Normal with standard deviation $\sigma = 0.1$, and the alternative we hope to be able to detect is $\mu = 10.15$. (If you used the *Power of a Test* applet for Exercise 16.16, the two Normal curves for $n = 6$ illustrate parts (a) and (b) below.)

(a) Write the z test statistic in terms of the sample mean \bar{x} . For what values of z does this two-sided test reject H_0 at the 5% significance level?

(b) Restate your result from part (a): what values of \bar{x} lead to rejection of H_0 ?

(c) Now suppose that $\mu = 10.15$. What is the probability of observing an \bar{x} that leads to rejection of H_0 ? This is the power of the test.

16.50 Error probabilities. You read that a statistical test at significance level $\alpha = 0.01$ has power 0.78. What are the probabilities of Type I and Type II errors for this test?

16.51 Power. You read that a statistical test at the $\alpha = 0.05$ level has probability 0.14 of making a Type II error when a specific alternative is true. What is the power of the test against this alternative?

16.52 Find the error probabilities. You have an SRS of size $n = 16$ from a Normal distribution with $\sigma = 1$. You wish to test

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

You decide to reject H_0 if $\bar{x} > 0$ and to accept H_0 otherwise.

(a) Find the probability of a Type I error. That is, find the probability that the test rejects H_0 when in fact $\mu = 0$.

(b) Find the probability of a Type II error when $\mu = 0.2$. This is the probability that the test accepts H_0 when in fact $\mu = 0.2$.

(c) Find the probability of a Type II error when $\mu = 0.5$.

16.53 Two types of error. Go to the *Statistical Significance* applet. This applet carries out tests at a fixed significance level. When you arrive, the applet is set for the cola-tasting test of Example 16.6. That is, the hypotheses are

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

We have an SRS of size 10 from a Normal population with standard deviation $\sigma = 1$, and we will do a test at level $\alpha = 0.05$. At the bottom of the screen, a button allows you to choose a value of the mean μ and then to generate samples from a population with that mean.

(a) Set $\mu = 0$, so that the null hypothesis is true. Each time you click the button, a new sample appears. If the sample \bar{x} lands in the colored region, that sample rejects H_0 at the 5% level. Click 100 times rapidly, keeping track of how many samples reject H_0 . Use your results to estimate the probability of a Type I error. If you kept clicking forever, what probability would you get?

(b) Now set $\mu = 0.8$. Example 16.6 shows that the test has power 0.808 against this alternative. Click 100 times rapidly, keeping track of how many samples fail to reject H_0 . Use your results to estimate the probability of a Type II error. If you kept clicking forever, what probability would you get?



EXPLORING THE WEB

16.54 Statistically significant but not practically important. Find an example of a study in which a statistically significant result may not be practically important. Summarize the study and its conclusions in your own words. The CHANCE Web site at www.causeweb.org/wiki/chance//index.php/Main_Page is a good place to look for examples.

16.55 The American Psychological Association and testing. The report of the American Psychological Association's Task Force on Statistical Inference is an excellent brief introduction to wise use of inference. The report appeared in the journal *American Psychologist* in 1999. You can find a copy of the initial report on the Web in the list of "TFSI Publications/Links" from this journal at www.apa.org/science/leadership/bsa/statistical/index.aspx. Read the initial report. Are the authors of the report opposed to the use of hypothesis testing? Describe one abuse of hypothesis testing that is cited in this report.



From Exploration to Inference: Part II Review

Chapter 17

IN THIS CHAPTER
WE COVER...

- Part II Summary
- Test Yourself
- Supplementary Exercises

In Part I of this book, you mastered data analysis, the use of graphs and numerical summaries to organize and explore any set of data. Part II has introduced designs for data production, probability, and the reasoning of statistical inference. Parts III and IV will deal in detail with practical inference.

Designs for producing data are essential if the data are intended to represent some wider population or process. Figures 17.1 and 17.2 display the big ideas visually. Random sampling and randomized comparative experiments are perhaps the most important statistical inventions of the 20th century. Both were slow to gain acceptance, and you will still see many voluntary response samples and uncontrolled experiments. You should

Figure 17.1 STATISTICS IN SUMMARY

Simple Random Sample

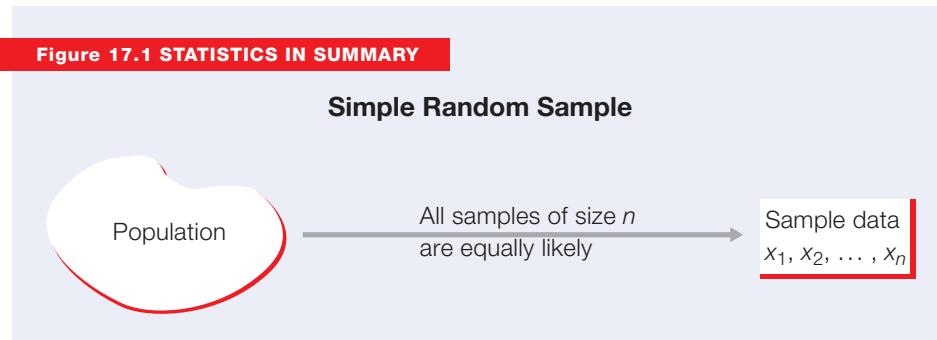
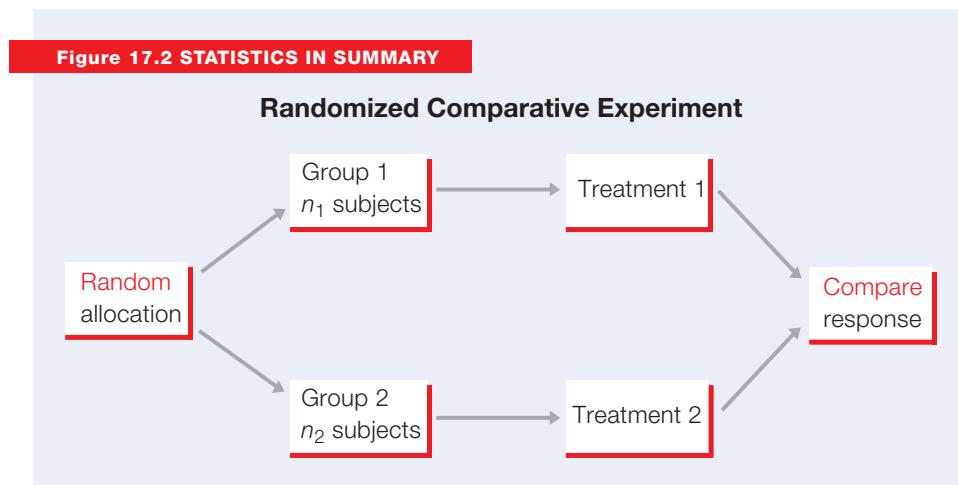


FIGURE 17.1

**FIGURE 17.2**

now understand good designs for producing data and also why bad designs often produce data that are worthless for inference. The deliberate use of chance in producing data is a central idea in statistics. It not only reduces bias but allows us to use **probability**, the mathematics of chance, as the basis for inference. Fortunately, we need only some basic facts about probability in order to understand statistical inference.

Statistical inference draws conclusions about a population on the basis of sample data and uses probability to indicate how reliable the conclusions are. A confidence interval estimates an unknown parameter. A significance test shows how strong the evidence is for some claim about a parameter.

The probabilities in both confidence intervals and tests tell us what would happen if we used the method for the interval or test very many times.

- A confidence level is the success rate of the method for a confidence interval. This is the probability that the method actually produces an interval that captures the unknown parameter. A 95% confidence interval gives a correct result 95% of the time when we use it repeatedly.
- A *P*-value tells us how surprising the observed outcome would be if the null hypothesis were true. That is, *P* is the probability that the test would produce a result at least as extreme as the observed result if the null hypothesis really were true. Very surprising outcomes (small *P*-values) are good evidence that the null hypothesis is not true.

Figures 17.3 and 17.4 use the *z* procedures introduced in Chapters 14 and 15 to present in picture form the big ideas of confidence intervals and significance tests. These ideas are the foundation for the rest of this book. We will have much to say about many statistical methods and their use in practice. In every case, the basic reasoning of confidence intervals and significance tests remains the same.

Figure 17.3 STATISTICS IN SUMMARY

The Idea of a Confidence Interval

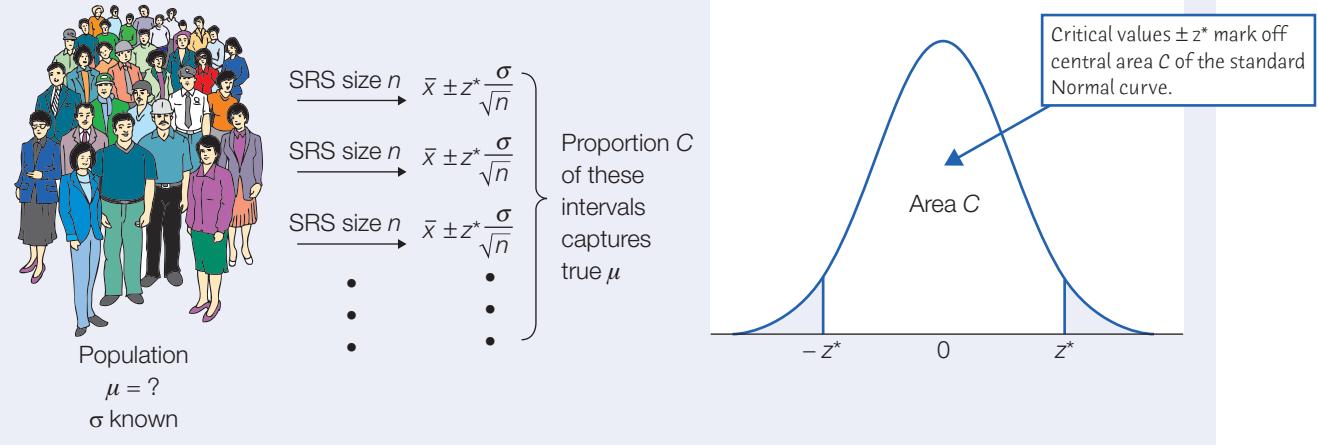


FIGURE 17.3

Figure 17.4 STATISTICS IN SUMMARY

The Idea of a Significance Test

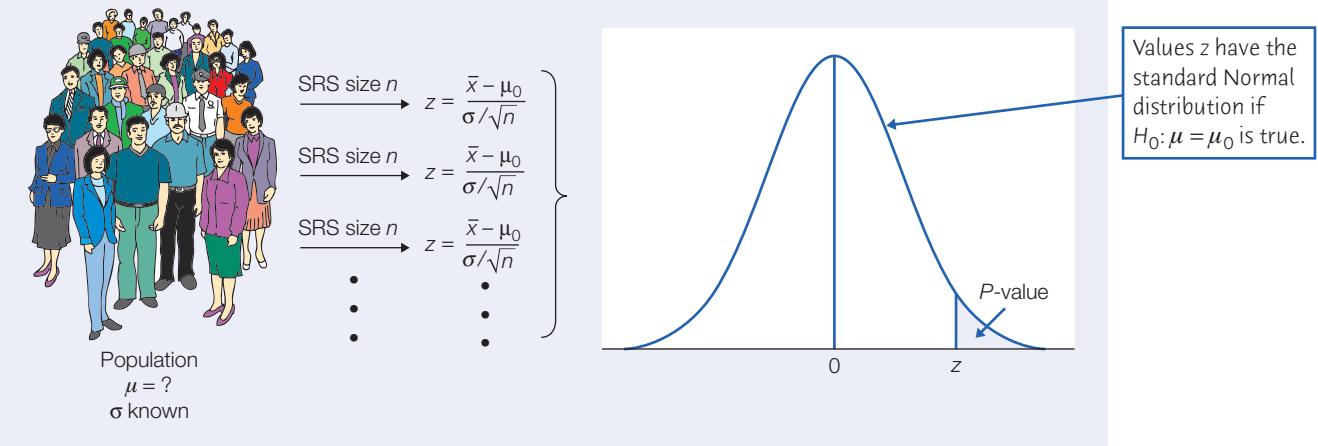


FIGURE 17.4

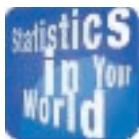
PART II SUMMARY

Here are the most important skills you should have acquired from reading Chapters 8 to 16.

A. Sampling

1. Identify the population in a sampling situation.
2. Recognize bias due to voluntary response samples and other inferior sampling methods.

3. Use software or Table B of random digits to select a simple random sample (SRS) from a population.
4. Recognize the presence of undercoverage and nonresponse as sources of error in a sample survey. Recognize the effect of the wording of questions on the responses.
5. Use random digits to select a stratified random sample from a population when the strata are identified.



Icing the kicker

The football team lines up for what they hope will be the winning field goal ... and the other team calls time out. "Make the kicker think about it" is their motto. Does "icing the kicker" really work? That is, does the probability of making a field goal go down when the kicker must wait around during the time out? This isn't a simple question. A detailed statistical study considered the distance, the weather, the kicker's skill, and so on. The conclusion is cheering to coaches: yes, icing the kicker does reduce the probability of success.

B. Experiments

1. Recognize whether a study is an observational study or an experiment.
2. Recognize bias due to confounding of explanatory variables with lurking variables in either an observational study or an experiment.
3. Identify the factors (explanatory variables), treatments, response variables, and individuals or subjects in an experiment.
4. Outline the design of a completely randomized experiment using a diagram like that in Figure 9.3 (page 229). The diagram in a specific case should show the sizes of the groups, the specific treatments, and the response variable.
5. Use software or Table B of random digits to carry out the random assignment of subjects to groups in a completely randomized experiment.
6. Recognize the placebo effect. Recognize when the double-blind technique should be used.
7. Explain why randomized comparative experiments can give good evidence for cause-and-effect relationships.

C. Probability

1. Recognize that some phenomena are random. Probability describes the long-run regularity of random phenomena.
2. Understand that the probability of an event is the proportion of times the event occurs in very many repetitions of a random phenomenon. Use the idea of probability as long-run proportion to think about probability.
3. Use basic probability rules to detect illegitimate assignments of probability: any probability must be a number between 0 and 1, and the total probability assigned to all possible outcomes must be 1.
4. Use basic probability rules to find the probabilities of events that are formed from other events. The probability that an event does not occur is 1 minus its probability. If two events are disjoint, the probability that one or the other occurs is the sum of their individual probabilities.
5. Find probabilities in a finite probability model by adding the probabilities of their outcomes. Find probabilities in a continuous probability model as areas under a density curve.
6. Use the notation of random variables to make compact statements about random outcomes, such as $P(\bar{x} \leq 4) = 0.3$. Be able to interpret such statements.

D. Sampling Distributions

1. Identify parameters and statistics in a statistical study.
2. Recognize the fact of sampling variability: a statistic will take different values when you repeat a sample or experiment.
3. Interpret a sampling distribution as describing the values taken by a statistic in all possible repetitions of a sample or experiment under the same conditions.
4. Interpret the sampling distribution of a statistic as describing the probabilities of its possible values.

E. The Sampling Distribution of a Sample Mean

1. Recognize when a problem involves the mean \bar{x} of a sample. Understand that \bar{x} estimates the mean μ of the population from which the sample is drawn.
2. Use the law of large numbers to describe the behavior of \bar{x} as the size of the sample increases.
3. Find the mean and standard deviation of a sample mean \bar{x} from an SRS of size n when the mean μ and standard deviation σ of the population are known.
4. Understand that \bar{x} is an unbiased estimator of μ and that the variability of \bar{x} about its mean μ gets smaller as the sample size increases.
5. Understand that \bar{x} has approximately a Normal distribution when the sample is large (central limit theorem). Use this Normal distribution to calculate probabilities that concern \bar{x} .

F. General Rules of Probability (not required for later chapters)

1. Use Venn diagrams to picture relationships among several events.
2. Use the general addition rule to find probabilities that involve overlapping events.
3. Understand the idea of independence. Judge when it is reasonable to assume independence as part of a probability model.
4. Use the multiplication rule for independent events to find the probability that all of several independent events occur.
5. Use the multiplication rule for independent events in combination with other probability rules to find the probabilities of complex events.
6. Understand the idea of conditional probability. Find conditional probabilities for individuals chosen at random from a table of counts of possible outcomes.
7. Use the general multiplication rule to find $P(A \text{ and } B)$ from $P(A)$ and the conditional probability $P(B | A)$.
8. Use tree diagrams to organize several-stage probability models.

G. Binomial Distributions (not required for later chapters)

1. Recognize the binomial setting: a fixed number n of independent success-failure trials with the same probability p of success on each trial.
2. Recognize and use the binomial distribution of the count of successes in a binomial setting.
3. Use the binomial probability formula to find probabilities of events involving the count X of successes in a binomial setting for small values of n .
4. Find the mean and standard deviation of a binomial count X .
5. Recognize when you can use the Normal approximation to a binomial distribution. Use the Normal approximation to calculate probabilities that concern a binomial count X .

H. Confidence Intervals



1. State in nontechnical language what is meant by “95% confidence” or other statements of confidence in statistical reports.
2. Know the four-step process (page 359) for any confidence interval. This process will be used more extensively in later chapters.
3. Calculate a confidence interval for the mean μ of a Normal population with known standard deviation σ , using the formula $\bar{x} \pm z^*\sigma/\sqrt{n}$.
4. Understand how the margin of error of a confidence interval changes with the sample size and the level of confidence C .
5. Find the sample size required to obtain a confidence interval of specified margin of error m when the confidence level and other information are given.
6. Identify sources of error in a study that are *not* included in the margin of error of a confidence interval, such as undercoverage or nonresponse.

I. Significance Tests



1. State the null and alternative hypotheses in a testing situation when the parameter in question is a population mean μ .
2. Explain in nontechnical language the meaning of the P -value when you are given the numerical value of P for a test.
3. Know the four-step process (page 379) for any significance test. This process will be used more extensively in later chapters.
4. Calculate the one-sample z test statistic and the P -value for both one-sided and two-sided tests about the mean μ of a Normal population.
5. Assess statistical significance at standard levels α , either by comparing P with α or by comparing z with standard Normal critical values.
6. Recognize that significance testing does not measure the size or importance of an effect. Explain why a small effect can be significant in a large sample and why a large effect can fail to be significant in a small sample.
7. Recognize that any inference procedure acts as if the data were properly produced. The z confidence interval and test require that the data be an SRS from the population.
8. Explain in nontechnical language the meaning of power, Type I, and Type II errors.

TEST YOURSELF

The questions below include both multiple-choice and short-answer questions and calculations. They will help you review the basic ideas and skills presented in Chapters 8 to 16.

Elephants and bees. Elephants sometimes damage crops in Africa. It turns out that elephants dislike bees. They recognize beehives in areas where they are common and avoid them. Can this be used to keep elephants away from trees? A group in Kenya placed active beehives in some trees, empty beehives in others, while others received no beehives.¹ Will elephant damage be less in trees with hives? Will even empty hives keep elephants away? Use this information to answer Questions 17.1 and 17.2.



© David Paynter/Age fotostock

17.1 This experiment has

- (a) two factors, beehives present or absent.
- (b) matched pairs.
- (c) three treatments.
- (d) stratification by beehive.

17.2 The response in this experiment is

- (a) the type of crop.
- (b) the presence or absence of bees.
- (c) the presence or absence of hives.
- (d) elephant damage.

American Community Survey. Each month the U.S. Census Bureau's American Community Survey mails survey forms to 250,000 households asking questions about demographic, social, economic, and housing characteristics such as mortgage and utility costs. Telephone calls are made to households that don't return the form. In one month, responses were obtained from 240,000 of the households contacted. Use this information to answer Questions 17.3 and 17.4.

17.3 The sample is

- (a) the 250,000 households initially contacted.
- (b) the 240,000 households that responded.
- (c) the 10,000 households that did not respond.
- (d) all U.S. households.

17.4 The population of interest is

- (a) all households with mortgages.
- (b) the 250,000 households contacted.
- (c) only U.S. households with phones.
- (d) all U.S. households.

17.5 At a local health club, a researcher samples 75 people whose primary exercise is cardiovascular and 75 people whose primary exercise is strength training. The researchers' objective is to assess the effect of type of exercise on cholesterol. Each subject reported to a clinic to have his or her cholesterol measured. The subjects were unaware of the purpose of the study, and the technician measuring the cholesterol was not aware of the subjects' type of exercise. This is

- (a) an observational study.
- (b) an experiment, but not a double-blind experiment.
- (c) a double-blind experiment.
- (d) a matched pairs experiment.

17.6 A university's financial aid office wants to know how much it can expect students to earn from summer employment. This information will be used to set the level of financial aid. The population contains 3478 students who have completed at least one year of study but have not yet graduated. The university

will send a questionnaire to an SRS of 100 of these students, drawn from an alphabetized list.

- (a) Describe how you will label the students in order to select the sample.
- (b) Use Table B, beginning at line 105, to select the first 5 students in the sample.
- (c) What is the response variable in this study?

17.7 A common definition of “binge drinking” is 5 or more drinks at one setting for men, and 4 or more for women. An observational study finds that students who binge have lower average GPA than those who don’t. Suggest two lurking variables that may be confounded with binge drinking, and be sure to give a reason why you have chosen each of these variables. The possibility of confounding means that we can’t conclude that binge drinking *causes* lower GPA.

17.8 The evidence linking chocolate to chronic headaches is inconsistent. In one study, 64 women with chronic headaches ate a restricted diet for two weeks. They then ate candy bars containing either chocolate or carob, prepared to taste the same, and reported whether they had a headache in the next 12 hours.²

- (a) Outline the design of this experiment.
- (b) Use Table B, beginning at line 110, to choose the first 5 members of the chocolate group.

17.9 In 2000, when the federal budget showed a large surplus, the Pew Research Center asked random samples of adults two questions about using the remaining surplus. Both questions stated that Social Security would be “fixed.”

Question A: *Should the money be used for a tax cut, or should it be used to fund new government programs?*

Question B: *Should the money be used for a tax cut, or should it be spent on programs for education, the environment, health care, crime-fighting and military defense?*

One of these questions drew 60% favoring a tax cut. The other drew only 22%. Which wording pulls respondents toward a tax cut? Why?

Snacking and movies. In a study of human development, investigators showed two movies that were different types to a group of children. Crackers were available in a bowl at each movie, and the investigators compared the number of crackers eaten by children watching each movie. One movie was shown at 8 A.M. (right after the children had breakfast) and the other at 11 A.M. (right before the children had lunch). It was found that during the movie shown at 11 A.M., more crackers were eaten than during the movie shown at 8 A.M. The investigators concluded that the different types of movies had different effects on appetite. Use this information to answer Questions 17.10 and 17.11.

17.10 The results cannot be trusted because

- (a) the study was not double-blind. Neither the investigators nor the children should have been aware of which movie was being shown.
- (b) the investigators were biased. They knew beforehand what the study would show.
- (c) the investigators should have used several bowls of crackers randomly placed in the room.
- (d) the time each movie was shown is a confounding variable.

17.11 The treatment in this experiment is

- (a) the number of crackers eaten.
- (b) the different types of movies.
- (c) the time each movie was shown.
- (d) the type of cracker.

17.12 The Web site of the PBS television program NOVA Science Now invites viewers to vote on issues such as re-creating the virus responsible for the deadly flu epidemic of 1918. This online poll is unusual in offering detailed arguments for both sides. Of the 790 viewers who read the arguments and voted, 64% said that re-creating the virus was justified.³ Explain to someone who knows no statistics why these 790 responses probably don't represent the opinions of all American adults.

17.13 A study attempts to determine whether a football filled with helium travels farther when kicked than one filled with air. Each subject kicks twice, once with a football filled with helium and once with a football filled with air. The order of the type of football kicked is randomized. This is an example of

- (a) a matched pairs experiment.
- (b) a randomized controlled experiment.
- (c) a stratified experiment.
- (d) the placebo effect.

A student survey. To assess the opinion of students about campus safety at the Ohio State University, a reporter for the student newspaper interviews 15 students she meets walking on the campus late at night who are willing to give their opinion. Use this information to answer Questions 17.14 and 17.15.

17.14 The sample is

- (a) all those students walking on campus late at night.
- (b) all students at universities with safety issues.
- (c) the 15 students interviewed.
- (d) all students approached by the reporter.

17.15 The sample obtained is

- (a) a simple random sample of students who feel safe.
- (b) a stratified random sample of students who feel safe.
- (c) a probability sample of students with night classes.
- (d) probably biased.

17.16 A randomly chosen subject arrives for a study of exercise and fitness. Describe a sample space for each of the following. (In some cases, you may have some freedom in your choice of S .)

- (a) The subject is either female or male.
- (b) After 10 minutes on an exercise bicycle, you ask the subject to rate his or her effort on the Rate of Perceived Exertion (RPE) scale. RPE ranges in whole-number steps from 6 (no exertion at all) to 20 (maximal exertion).
- (c) You measure VO₂, the maximum volume of oxygen consumed per minute during exercise. VO₂ is generally between 2.5 and 6.1 liters per minute.
- (d) You measure the maximum heart rate (beats per minute).

Internet search engines. Internet search sites compete for users because they sell advertising space on their sites and can charge more if they are heavily used. Choose an Internet search attempt at random. Here is the probability distribution for the site the search uses:⁴

Site	Google	Yahoo	MSN	Ask.com	Others
Probability	0.66	0.21	0.07	0.04	?

Use this information to answer Questions 17.17 and 17.18



Justin Sullivan/Getty Images

17.17 What is the probability that a search attempt is made at a site other than the leading four?

- (a) 0.02 (b) 0.34 (c) 0.98

(d) Cannot be determined from the information given.

17.18 What is the probability that a search attempt is directed to a site other than Google?

- (a) 0.02 (b) 0.34 (c) 0.98

(d) Cannot be determined from the information given.

How many in the house? In government data, a household consists of all occupants of a dwelling unit. Here is the distribution of household size in the United States:

Number of persons	1	2	3	4	5	6	7
Probability	0.26	0.33	0.16	0.15	0.07	0.02	0.01

Choose an American household at random and let the random variable Y be the number of persons living in the household. Use this information to answer Questions 17.19 to 17.21.

17.19 Express “more than one person lives in this household” in terms of Y . What is the probability of this event?

17.20 What is $P(2 < Y \leq 4)$?

17.21 What is $P(Y \neq 2)$?

- (a) 0.26 (b) 0.33 (c) 0.41 (d) 0.67

How many children? How many children do women give birth to during their childbearing years? Choose at random an American woman who is past childbearing years:⁵

Number of children	0	1	2	3	4	5
Probability	0.193	0.174	0.344	0.181	0.074	0.034

(The few women with 6 or more children are included in the “5 children” group.) Use this information to answer Questions 17.22 to 17.25.

17.22 Check that this distribution satisfies the two requirements for a legitimate finite probability model.

17.23 Describe in words the event $P(X \leq 2)$. What is the probability of this event?

17.24 What is $P(X < 2)$?

- (a) 0.289 (b) 0.344 (c) 0.367 (d) 0.711

17.25 Write the event “a woman gives birth to 3 or more children” in terms of values of X . What is the probability of this event?

Random number generators. Many random number generators allow users to specify the range of the random numbers to be produced. Suppose that you specify that the random number Y can take any value between 0 and 5. The density curve of the outcome has a constant height between 0 and 5, and height 0 elsewhere. Use this information to answer Questions 17.26 to 17.28.

17.26 The random variable Y is

- (a) discrete.
- (b) continuous, but not Normal.
- (c) continuous and Normal.
- (d) none of the above.

17.27 The height of the density curve between 0 and 5 is

- (a) 0.2.
- (b) 1.
- (c) 5.
- (d) none of the above.

17.28 Draw a graph of the density curve and find $P(1 \leq Y \leq 3)$.

17.29 (Optional). Byron claims that the probability that Florida and Alabama will play in the Southeastern Conference (SEC) championship football game this year is 78%. The number 78% is

- (a) the proportion of times Florida and Alabama have played in the championship game in the past.
- (b) Byron's personal probability that Florida and Alabama will play in the SEC championship football game this year.
- (c) the area under a Normal density curve.
- (d) all of the above.

An IQ test. The Wechsler Adult Intelligence Scale (WAIS) is a common "IQ test" for adults. The distribution of WAIS scores for persons over 16 years of age is approximately Normal with mean 100 and standard deviation 15. Use this information to answer Questions 17.30 to 17.33.

17.30 What is the probability that a randomly chosen individual has a WAIS score of 105 or higher?

- (a) 0.0005
- (b) 0.3607
- (c) 0.4400
- (d) 0.6293

17.31 What are the mean and standard deviation of the average WAIS score \bar{x} for an SRS of 60 people?

- (a) mean = 13.56, standard deviation = 15.
- (b) mean = 100, standard deviation = 15.
- (c) mean = 100, standard deviation = 1.94.
- (d) mean = 100, standard deviation = 0.25.

17.32 What is the probability that the average WAIS score of an SRS of 60 people is 105 or higher?

- (a) 0.0049
- (b) 0.3607
- (c) 0.9738
- (d) none of the above.

17.33 Would your answers to any of Questions 17.30, 17.31, or 17.32 be affected if the distribution of WAIS scores in the adult population were distinctly non-Normal? Explain.

Reaction times. The time that people require to react to a stimulus usually has a right-skewed distribution, as lack of attention or tiredness causes some lengthy reaction times. Reaction times for children with attention-deficit/hyperactivity disorder (ADHD) are more skewed, as their condition causes more frequent lack of attention. In one study, children with ADHD were asked to press the spacebar on a computer keyboard when any letter other than X appeared on the screen. With 2 seconds between letters, the mean reaction time was 445 milliseconds (ms) and the standard deviation was 82 ms.⁶ Take these values to be the population μ and σ for ADHD children. Use this information to answer Questions 17.34 to 17.36.

17.34 What are the mean and standard deviation of the mean reaction time \bar{x} for a randomly chosen group of 15 ADHD children? For a group of 150 such children?

17.35 The distribution of reaction time is strongly skewed. Explain briefly why we hesitate to regard \bar{x} as Normally distributed for 15 children but are willing to use a Normal distribution for the mean reaction time of 150 children.

17.36 What is the approximate probability that the mean reaction time in a group of 150 ADHD children is greater than 450 ms?

17.37 (Optional). Accidents, suicide, and murder are the leading causes of death for young adults. Here are the counts of violent deaths in a recent year among people 20 to 24 years of age:

	Female	Male
Accidents	1818	6457
Homicide	457	2870
Suicide	345	2152

- (a) Choose a violent death in this age group at random. What is the probability that the victim was male?
- (b) Find the conditional probability that the victim was male given that the death was accidental.

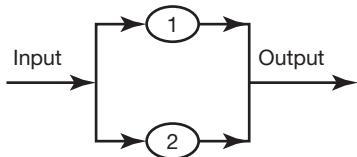


FIGURE 17.5

Parallel systems for Exercise 17.38.

Parallel systems (optional). A system has two components that operate in parallel, as shown in Figure 17.5. Because the components operate in parallel, at least one of the components must function properly if the system is to function properly. The probabilities of failure for Components 1 and 2 during one period of operation are 0.20 and 0.03, respectively. Let F denote the event that Component 1 fails during one period of operation and G denote the event that Component 2 fails during one period of operation. The component failures are independent. Use this information to answer Questions 17.38 and 17.39.

17.38 The event corresponding to the system failing during one period of operation is

- (a) F and G. (b) F or G. (c) not F or not G. (d) not F and not G.

17.39 The probability that the system functions properly during one period of operation is closest to

- (a) 0.994. (b) 0.970. (c) 0.940. (d) 0.776.

17.40 (Optional). A survey of college students finds that 35% like country music, 25% like gospel music, and 15% like both. The proportion of students that like country music but not gospel music is

- (a) 15%. (b) 20%. (c) 25%. (d) 40%.

17.41 (Optional). Opinion polls find that 63% of American teens say that their parents put at least some pressure on them to get into a good college.⁷ If you take an SRS of 1000 teens, the approximate distribution of the number in your sample who say that they feel at least some pressure from their parents to get into a good college is

- (a) $N(0.63, 15.27)$. (b) $N(0.63, 233.1)$.
- (c) $N(630, 15.27)$. (d) $N(630, 233.1)$.

17.42 (Optional). What kinds of Web sites do males aged 18 to 34 visit? About 50% of male Internet users in this age group visit an auction site such as eBay at least once a month.⁸

- (a) If we interview a random sample of 12 male Internet users aged 18 to 34, what is the probability that exactly 8 of the 12 have visited an auction site in the past month?

- (b) Suppose that we had interviewed a random sample of 500 men aged 18 to 34. What is the probability that at least 235 of the men in the sample visit an online auction site at least once a month? (Check that the Normal approximation is permissible and use it to find this probability.)

Pesticides in whale blubber: estimation. The level of pesticides found in the blubber of whales is a measure of pollution of the oceans by runoff from land and can also be used to identify different populations of whales. A sample of 8 male minke whales in the West Greenland area of the North Atlantic found the mean concentration of the insecticide dieldrin to be $\bar{x} = 357$ nanograms per gram of blubber (ng/g).⁹ Suppose that the concentration in all such whales varies Normally with standard deviation $\sigma = 50$ ng/g. Use this information to answer Questions 17.44 to 17.47.

17.43 A 95% confidence interval to estimate the mean level of dieldrin is

- (a) 344.75 to 369.25. (b) 339.32 to 374.68.
 (c) 322.35 to 391.65. (d) 259.00 to 455.00.

17.44 A 90% confidence interval to estimate the mean level of dieldrin is

- (a) 346.72 to 367.28. (b) 327.92 to 386.08.
 (c) 311.36 to 402.54. (d) 274.75 to 439.25.

17.45 Find an 80% confidence interval for the mean concentration of dieldrin in the whale population.

17.46 What general fact about confidence intervals do the margins of error of your three intervals in the previous problems illustrate?

Estimating blood cholesterol. The distribution of blood cholesterol level in the population of young men aged 20 to 34 years is close to Normal with standard deviation $\sigma = 41$ milligrams per deciliter (mg/dl). You measure the blood cholesterol of 14 cross-country runners. The mean level is $\bar{x} = 172$ mg/dl. Assume that σ is the same as in the general population. Use this information to answer Questions 17.47 to 17.49.

17.47 A 90% confidence interval for the mean level μ among cross-country runners is

- (a) 172 ± 4.82 mg/dl. (b) 172 ± 18.03 mg/dl.
 (c) 172 ± 21.48 mg/dl. (d) none of the above.

17.48 How large a sample is needed to cut the margin of error in the previous exercise in half?

- (a) 2 (b) 4 (c) 28 (d) 56

17.49 How large a sample is needed to cut the margin of error to ± 5 mg/dl?

- (a) 14 (b) 68 (c) 182 (d) 259

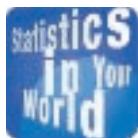
17.50 The Environmental Protection Agency (EPA) fuel economy ratings say that the Toyota Prius hybrid car gets 48 miles per gallon (mpg) on the highway. Deborah wonders whether the actual long-term average highway mileage μ of her new Prius is less than 48 mpg. She keeps careful records of gas mileage for 3000 miles of highway driving. Her result is $\bar{x} = 47.2$ mpg. What are her null and alternative hypotheses?

- (a) $H_0: \mu = 48$, $H_a: \mu < 48$. (b) $H_0: \mu = 48$, $H_a: \mu > 48$.
 (c) $H_0: \bar{x} = 48$, $H_a: \bar{x} < 48$. (d) $H_0: \bar{x} = 48$, $H_a: \bar{x} > 48$.

17.51 The average amount of time that high school students spend on homework is about 5 hours per week. Only 25% of college freshmen say they spent at least



Flip Nicklin/Getty



How many miles per gallon?

As gasoline prices rise, more people pay attention to the government's gas mileage ratings of their vehicles. Until recently these ratings overstated the miles per gallon we can expect in real-world driving. The ratings assumed a top speed of 60 miles per hour, slow acceleration, and no air-conditioning. That doesn't resemble what we see around us on the highway. Maybe it doesn't resemble the way we ourselves drive. Starting with 2008 models, the ratings assume higher speeds (80 miles per hour tops), faster acceleration, and air-conditioning in warm weather. Mileage ratings of the same vehicle dropped by about 12% in the city and 8% on the highway.

6 hours per week on homework in high school. Your college wonders if the average μ for its freshmen differs from the national average. A random sample of 500 freshmen claims to have spent an average of $\bar{x} = 6.2$ hours per week on homework in high school. What are the null and alternative hypotheses for a comparison of freshmen at your college with national freshmen?

- (a) $H_0: \bar{x} = 5, H_a: \bar{x} \neq 5.$ (b) $H_0: \bar{x} = 6, H_a: \bar{x} > 6.$
 (c) $H_0: \mu = 5, H_a: \mu \neq 5.$ (d) $H_0: \mu = 6, H_a: \mu > 6.$

Testing blood cholesterol. The distribution of blood cholesterol level in the population of young men aged 20 to 34 years is close to Normal with mean 188 milligrams per deciliter (mg/dl) and standard deviation 41 mg/dl. You measure the blood cholesterol of 14 cross-country runners. The mean level is $\bar{x} = 172$ mg/dl. Assume that σ is the same as in the general population. Use this information to answer Questions 17.52 to 17.54.

17.52 We suspect that the mean μ for all cross-country runners is lower than that for the population of young men aged 20 to 34 years. Thus, we decide to test the hypotheses $H_0: \mu = 188, H_a: \mu < 188.$ The z test statistic for testing these hypotheses is

- (a) 5.46 (b) -5.46 (c) 1.46 (d) -1.46

17.53 The result is significant at

- (a) $\alpha = 0.01.$ (b) $\alpha = 0.05$ but not at $\alpha = 0.01.$
 (c) $\alpha = 0.10$ but not at $\alpha = 0.05.$ (d) $\alpha = 0.25$ but not at $\alpha = 0.10.$

17.54 You increase the sample of cross-country runners from 14 to 56. Suppose that this larger sample gives the same mean level, $\bar{x} = 172$ mg/dl. Redo the test in the previous exercises. The result is significant at

- (a) $\alpha = 0.01.$ (b) $\alpha = 0.05$ but not at $\alpha = 0.01.$
 (c) $\alpha = 0.10$ but not at $\alpha = 0.05.$ (d) $\alpha = 0.25$ but not at $\alpha = 0.10$

17.55 The Food and Drug Administration regulates the amount of dieldrin in raw food. For some foods, no more than 100 nanograms per gram (ng/g) is allowed. Using the information in Questions 17.43 to 17.46, is there good evidence that the mean concentration μ in whale blubber is above 100 ng/g? Carry out a test of the hypotheses $H_0: \mu = 100, H_a: \mu > 100$ assuming that the “simple conditions” (page 352) hold. The P -value of your test is

- (a) above 0.10.
 (b) less than or equal to 0.10 but greater than 0.05.
 (c) less than or equal to 0.05 but greater than 0.01.
 (d) no more than 0.01.

17.56 Infants weighing less than 1500 grams at birth are classed as “very low birth weight.” Low birth weight carries many risks. One study followed 113 male infants with very low birth weight to adulthood. At age 20, the mean IQ score for these men was $\bar{x} = 87.6.$ ¹⁰ IQ scores vary Normally with standard deviation $\sigma = 15.$ Give a 95% confidence interval for the mean IQ score at age 20 for all very-low-birth-weight males.

17.57 IQ tests are scaled so that the mean score in a large population should be $\mu = 100.$ We suspect that the very-low-birth-weight population has mean score less than 100. Does the study described in the previous exercise give good evidence that this is true? State hypotheses, carry out a test assuming that the “simple conditions” (page 352) hold, compute the P -value, and give your conclusion in plain language.

17.58 Very-low-birth-weight babies are more likely to be born to unmarried mothers and to mothers who did not complete high school. Is the study of the previous examples an experiment? Explain. Also explain clearly why confounding prevents us from concluding that very low birth weight in itself reduces adult IQ.

17.59 When our brains store information, complicated chemical changes take place. In trying to understand these changes, researchers blocked some processes in brain cells taken from rats and compared these cells with a control group of normal cells. They say that “no differences were seen” between the two groups in four response variables. They give P -values 0.45, 0.83, 0.26, and 0.84 for these four comparisons.¹¹ Which of the following statements is correct?

- (a) It is literally true that “no differences were seen.” That is, the mean responses were exactly alike in the two groups.
- (b) The mean responses were exactly alike in the two groups for at least one of the four response variables measured, but not for all of them.
- (c) The statement “no differences were seen” means that the observed differences were not statistically significant at the significance level used by the researchers.
- (d) The statement “no differences were seen” means that the observed differences were all less than 1 (and were actually 0.45, 0.83, 0.26, and 0.84 for these four comparisons).

17.60 Here are some of the results of the experiment described in Question 17.8. There was no significant difference in headaches between the chocolate and carob groups ($P = 0.68$). But subjects who said they had a mild headache before eating the candy bar were more likely to report a headache afterward ($P < 0.001$). Explain carefully why $P = 0.68$ means that there is no evidence that chocolate and carob differ in their effects and why $P < 0.001$ is evidence that having a headache before eating the candy bar does increase reports of a headache after eating.

17.61 We often see televised reports of brushfires threatening homes in California. Some people argue that the modern practice of quickly putting out small fires allows fuel to accumulate and so increases the damage done by large fires. A detailed study of historical data suggests that this is wrong—the damage has risen simply because there are more houses in risky areas. As usual, the study report gives statistical information tersely. Here is the summary of a regression of number of fires on decade (9 data points, for the 1910s to the 1990s): “Collectively, since 1910, there has been a highly significant increase ($r^2 = 0.61$, $P < 0.01$) in the number of fires per decade.”¹² How would you explain this statement to someone who knows no statistics? Include an explanation of both the description given by r^2 and its statistical significance.



CORBIS

SUPPLEMENTARY EXERCISES

Supplementary exercises apply the skills you have learned in ways that require more thought or more elaborate use of technology.

17.62 Sampling students. You want to investigate the attitudes of students at your school toward the school’s policy on sexual harassment. You have a grant that will pay the costs of contacting about 500 students.

- (a) Specify the exact population for your study. For example, will you include part-time students?
- (b) Describe your sample design. Will you use a stratified sample?
- (c) Briefly discuss the practical difficulties that you anticipate. For example, how will you contact the students in your sample?

17.63 The placebo effect. A survey of physicians found that some doctors give a placebo to a patient who complains of pain for which the physician can find no cause. If the patient's pain improves, these doctors conclude that it had no physical basis. The medical school researchers who conducted the survey claimed that these doctors do not understand the placebo effect. Why?

17.64 Informed consent. The requirement that human subjects give their informed consent to participate in an experiment can greatly reduce the number of available subjects. For example, a study of new teaching methods asks the consent of parents for their children to be taught by either a new method or the standard method. Many parents do not return the forms, so their children must continue to follow the standard curriculum. Why is it not correct to consider these children as part of the control group along with children who are randomly assigned to the standard method?

17.65 Fixing health care. The cost of health care and health insurance is the biggest health concern among Americans, even ahead of cancer and other diseases. Changing to a national government health insurance system is controversial. An opinion poll will give different results depending on the wording of the question asked. For each of the following claims, say whether including it in the question would *increase* or *decrease* the percent of a poll sample who support a government health insurance system.

- (a) A national system would mean that everybody has health insurance.
- (b) A national system would probably require an increase in taxes.
- (c) Eliminating private insurance companies and their profits would reduce insurance costs.
- (d) A national system would limit the medical treatments available in order to contain costs.

17.66 Market research. Stores advertise price reductions to attract customers. What type of price cut is most attractive? Market researchers prepared ads for athletic shoes announcing different levels of discounts (20%, 40%, or 60%). The student subjects who read the ads were also given "inside information" about the fraction of shoes on sale (50% or 100%). Each subject then rated the attractiveness of the sale on a scale of 1 to 7.¹³

- (a) There are two factors. Make a sketch like Figure 9.2 (page 226) that displays the treatments formed by all combinations of levels of the factors.
- (b) Outline a completely randomized design using 60 student subjects. Use software or Table B at line 111 to choose the subjects for the first treatment.

17.67 Making french fries. Few people want to eat discolored french fries. Potatoes are kept refrigerated before being cut for french fries to prevent spoiling and preserve flavor. But immediate processing of cold potatoes causes discoloring due to complex chemical reactions. The potatoes must therefore be brought to room temperature before processing. Design an experiment in which tasters will rate the color and flavor of french fries prepared from several groups of potatoes. The potatoes will be freshly picked or stored for a month at room temperature or stored for a month refrigerated. They will then be sliced and cooked either immediately or after an hour at room temperature.

- (a) What are the factors and their levels, the treatments, and the response variables?
- (b) Describe and outline the design of this experiment.
- (c) It is efficient to have each taster rate fries from all treatments. How will you use randomization in presenting fries to the tasters?

17.68 The addition rule. The addition rule for probabilities, $P(A \text{ or } B) = P(A) + P(B)$, is not always true. Give (in words) an example of real-world events A and B for which this rule is not true.

17.69 Comparing wine tasters. Two wine tasters rate each wine they taste on a scale of 1 to 5. From data on their ratings of a large number of wines, we obtain the following probabilities for both tasters' ratings of a randomly chosen wine:

		Taster 2				
		1	2	3	4	5
Taster 1	1	0.03	0.02	0.01	0.00	0.00
1	2	0.02	0.08	0.05	0.02	0.01
2	3	0.01	0.05	0.25	0.05	0.01
3	4	0.00	0.02	0.05	0.20	0.02
4	5	0.00	0.01	0.01	0.02	0.06

- (a) Why is this a legitimate finite probability model?
- (b) What is the probability that the tasters agree when rating a wine?
- (c) What is the probability that Taster 1 rates a wine higher than Taster 2? What is the probability that Taster 2 rates a wine higher than Taster 1?

17.70 A 14-sided die. An ancient Korean drinking game involves a 14-sided die. The players roll the die in turn and must submit to whatever humiliation is written on the

up-face: something like “Keep still when tickled on face.” Six of the 14 faces are squares. Let’s call them A, B, C, D, E, and F for short. The other eight faces are triangles, which we will call 1, 2, 3, 4, 5, 6, 7, and 8. Each of the squares is equally likely. Each of the triangles is also equally likely, but the triangle probability differs from the square probability. The probability of getting a square is 0.72. Give the probability model for the 14 possible outcomes.



David Moore

17.71 Distributions: means versus individuals. The z confidence interval and test are based on the sampling distribution of the sample mean \bar{x} . Suppose that the distribution of body mass index (BMI) among young women is Normal with mean $\mu = 27$ and standard deviation $\sigma = 7.5$.

- You take an SRS of 100 young women. According to the 99.7 part of the 68–95–99.7 rule, about what range of BMI values do you expect to see in your sample?
- You look at many SRSs of size 100. About what range of sample mean BMIs \bar{x} do you expect to see?

17.72 Distributions: larger samples. In the setting of the previous exercise, how many women must you sample to cut the range of values of \bar{x} in half? This will also cut the margin of error of a confidence interval for μ in half. Do you expect the range of individual scores in the new sample to also be much less than in a sample of size 100? Why?

17.73 Alcohol and mortality. It appears that people who drink alcohol in moderation have lower death rates than either people who drink heavily or people who do not drink at all. The protection offered by moderate drinking is concentrated among people over 50 and on deaths from heart disease. The Nurses’ Health Study played an essential role in establishing these facts for women. This part of the study followed 85,709 female nurses for 12 years, during which time 2658 of the subjects died. The nurses completed a questionnaire that described their diet, including their use of alcohol. They were reexamined every two years. Conclusion: “As compared with nondrinkers and heavy drinkers, light-to-moderate drinkers had a significantly lower risk of death.”¹⁴

- Was this study an experiment? Explain your answer.
- What does “significantly lower risk of death” mean in simple language?
- Suggest some lurking variables that might be confounded with how much a person drinks. The investigators used advanced statistical methods to adjust for many such variables before concluding that moderate drinkers really do have a lower risk of death.

17.74 Time in a restaurant. The owner of a pizza restaurant in France knows that the time customers spend in the restaurant on Saturday evening has mean 90 minutes and standard deviation 15 minutes. He has read that pleasant odors can influence customers, so he spreads a lavender odor throughout the restaurant. Here are the times (minutes) for customers on the next Saturday evening:¹⁵

RESTAURANT

92	126	114	106	89	137	93	76	98	108
124	105	129	103	107	109	94	105	102	108
95	121	109	104	116	88	109	97	101	106

- Make a stemplot of the times. The distribution is roughly symmetric and single-peaked, so the distribution of \bar{x} should be close to Normal.
- Suppose that the standard deviation $\sigma = 15$ minutes is not changed by the odor. Is there reason to think that the lavender odor has changed the mean time customers spend in the restaurant? Follow the four-step process for significance tests (page 379).

17.75 Normal body temperature? Here are the daily average body temperatures (degrees Fahrenheit) for 20 healthy adults:¹⁶

BODYTEMP

98.74	98.83	96.80	98.12	97.89	98.09	97.87
97.42	97.30	97.84	100.27	97.90	99.64	97.88
98.54	98.33	97.87	97.48	98.92	98.33	

- Make a stemplot of the data. The distribution is roughly symmetric and single-peaked. There is one mild outlier. We expect the distribution of the sample mean \bar{x} to be close to Normal.
- Do these data give evidence that the mean body temperature for all healthy adults is not equal to the traditional 98.6 degrees? Follow the four-step process for significance tests (page 379). (Suppose that body temperature varies Normally with standard deviation 0.7 degree.)

17.76 Time in a restaurant. Use the data in Exercise 17.74 to estimate the mean time customers spend in this restaurant on Saturday evenings with 95% confidence. Follow the four-step process for confidence intervals (page xxx). RESTAURANT

17.77 Normal body temperature. Use the data in Exercise 17.75 to estimate mean body temperature with 90% confidence. Follow the four-step process for confidence intervals (page 359).  BODYTEMP

17.78 Tests from confidence intervals. You read in a U.S. Census Bureau report that a 99% confidence interval for the mean income in 2005 of American households headed by a college-educated person at least 25 years old was $\$100,272 \pm \1651 . (The median income of these households was lower, \$77,179.) Based on this interval, can you reject the null hypothesis that the mean income in this group is \$95,000? What is the alternative hypothesis of the test? What is its significance level?

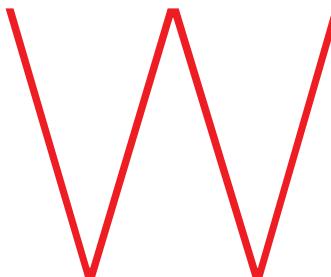
17.79 Low power? (Optional) It appears that eating oat bran lowers cholesterol slightly. At a time when oat

bran was something of a fad, a paper in the *New England Journal of Medicine* found that it had no significant effect on cholesterol.¹⁷ The paper reported a study with just 20 subjects. Letters to the journal denounced publication of a negative finding from a study with very low power. Explain why lack of significance in a study with low power gives no reason to accept the null hypothesis that oat bran has no effect.

17.80 Type I and Type II errors (Optional). Exercise 17.57 asks for a significance test of the null hypothesis that the mean IQ of very-low-birth-weight male babies is 100 against the alternative hypothesis that the mean is less than 100. State in words what it means to make a Type I error and a Type II error in this setting.

Inference about Variables

Part III



With the principles in hand, we proceed to practice, that is, to inference in fully realistic settings. In the remaining chapters of this book, you will meet many of the most commonly used statistical procedures. We have grouped these procedures into two classes, corresponding to our division of data analysis into exploring variables and distributions and exploring relationships. The five chapters of Part III concern inference about the distribution of a single variable and inference for comparing the distributions of two variables. Part IV deals with inference for relationships among variables. In Chapters 18 and 19, we analyze data on quantitative variables. We begin with the familiar Normal distribution for a quantitative variable. Chapters 20 and 21 concern categorical variables, so that inference begins with counts and proportions of outcomes. Chapter 22 reviews this part of the text.

The four-step process for approaching a statistical problem can guide much of your work in these chapters. You should review the outlines of the four-step process for a confidence interval (page 359) and for a test of significance (page 379). The statement of an exercise usually does the “State” step for you, leaving the “Plan,” “Solve,” and “Conclude” steps for you to complete. It is helpful to first summarize the “State” step in your own words to organize your thinking. Many examples and exercises in these chapters involve both carrying out inference and thinking about inference in practice. Remember that any inference method is useful only under certain conditions, and that you must judge these conditions before rushing to inference.

QUANTITATIVE RESPONSE VARIABLE

CHAPTER 18 Inference about a Population Mean

CHAPTER 19 Two-Sample Problems

CATEGORICAL RESPONSE VARIABLE

CHAPTER 20 Inference about a Population Proportion

CHAPTER 21 Comparing Two Proportions

CHAPTER 22 Inference about Variables: Part III Review



Inference about a Population Mean

Chapter 18

This chapter describes confidence intervals and significance tests for the mean μ of a population. We used the z procedures in this same setting to introduce the ideas of confidence intervals and tests. Now we discard the unrealistic condition that we know the population standard deviation σ and present procedures for practical use. We also pay more attention to the real-data setting of our work. The details of confidence intervals and tests change only slightly when you don't know σ . More important, you can interpret your results exactly as before. To illustrate this, Example 18.2 repeats an example from Chapter 14.

CONDITIONS FOR INFERENCE ABOUT A MEAN

Confidence intervals and tests of significance for the mean μ of a Normal population are based on the sample mean \bar{x} . Confidence intervals and P -values involve probabilities calculated from the sampling distribution of \bar{x} . Here are the conditions needed for realistic inference about a population mean.

CONDITIONS FOR INFERENCE ABOUT A MEAN

- We can regard our data as a **simple random sample** (SRS) from the population. This condition is very important.
- Observations from the population have a **Normal distribution** with mean μ and standard deviation σ . In practice, it is enough that the distribution be symmetric and single-peaked unless the sample is very small. Both μ and σ are unknown parameters.

IN THIS CHAPTER WE COVER...

- Conditions for inference about a mean
- The t distributions
- The one-sample t confidence interval
- The one-sample t test
- Using technology
- Matched pairs t procedures
- Robustness of t procedures

There is another condition that applies to all the inference methods in this book: *the population must be much larger than the sample, say at least 20 times as large.*¹ All our examples and exercises satisfy this condition. Practical settings in which the sample is a large part of the population are rather special, and we will not discuss them.

When the conditions for inference are satisfied, the sample mean \bar{x} has the Normal distribution with mean μ and standard deviation σ/\sqrt{n} . Because we don't know σ , we estimate it by the sample standard deviation s . We then estimate the standard deviation of \bar{x} by s/\sqrt{n} . This quantity is called the *standard error* of the sample mean \bar{x} .

STANDARD ERROR

When the standard deviation of a statistic is estimated from data, the result is called the **standard error** of the statistic. The standard error of the sample mean \bar{x} is s/\sqrt{n} .

APPLY YOUR KNOWLEDGE



© Dote Boe Photography/Alamy

18.1 Travel time to work. A study of commuting times reports the travel times to work of a random sample of 1000 employed adults. The mean is $\bar{x} = 49.2$ minutes and the standard deviation is $s = 63.9$ minutes. What is the standard error of the mean?

18.2 Comparing breathing frequencies in swimming. Researchers from the United Kingdom studied the effect of two breathing frequencies on performance times and on several physiological parameters in front crawl swimming. The breathing frequencies were one breath every second stroke (B2) and one breath every fourth stroke (B4). Subjects were 10 male collegiate swimmers. Each subject swam 200 meters using each breathing frequency: once with breathing frequency B2 and once on a different day with breathing frequency B4. A paper states that the results are expressed as mean plus or minus the standard deviation.² One result reported in the paper states that the immediate postexercise heart rate for subjects when using breathing frequency B2 was 163 ± 15 beats per minute. What are \bar{x} and the standard error of the mean for these subjects? (This exercise is also a warning to read carefully: that 163 ± 15 is not a confidence interval, yet summaries in this form are common in scientific reports.)

THE *t* DISTRIBUTIONS

If we knew the value of σ , we would base confidence intervals and tests for μ on the one-sample z statistic

$$z = \frac{\bar{x} - \mu}{\sigma/\sqrt{n}}$$

This z statistic has the standard Normal distribution $N(0, 1)$. In practice, we don't know σ , so we substitute the standard error s/\sqrt{n} of \bar{x} for its standard deviation σ/\sqrt{n} . The statistic that results does not have a Normal distribution. It has a distribution that is new to us, called a *t distribution*.

THE ONE-SAMPLE *t* STATISTIC AND THE *t* DISTRIBUTIONS

Draw an SRS of size n from a large population that has the Normal distribution with mean μ and standard deviation σ . The **one-sample *t* statistic**

$$t = \frac{\bar{x} - \mu}{s/\sqrt{n}}$$

has the ***t* distribution** with $n - 1$ degrees of freedom.

The *t* statistic has the same interpretation as any standardized statistic: it says how far \bar{x} is from its mean μ in standard deviation units. There is a different *t* distribution for each sample size. We specify a particular *t* distribution by giving its **degrees of freedom**. The degrees of freedom for the one-sample *t* statistic come from the sample standard deviation s in the denominator of t . We saw in Chapter 2 (page 51) that s has $n - 1$ degrees of freedom. There are other *t* statistics with different degrees of freedom, some of which we will meet later. We will write the *t* distribution with $n - 1$ degrees of freedom as $t(n - 1)$ for short.

degrees of freedom

Figure 18.1 compares the density curves of the standard Normal distribution and the *t* distributions with 2 and 9 degrees of freedom. The figure illustrates these facts about the *t* distributions:

- The density curves of the *t* distributions are similar in shape to the standard Normal curve. They are symmetric about 0, single-peaked, and bell-shaped.
- The spread of the *t* distributions is a bit greater than that of the standard Normal distribution. The *t* distributions in Figure 18.1 have more probability in the tails and less in the center than does the standard Normal. This is true because substituting the estimate s for the fixed parameter σ introduces more variation into the statistic.

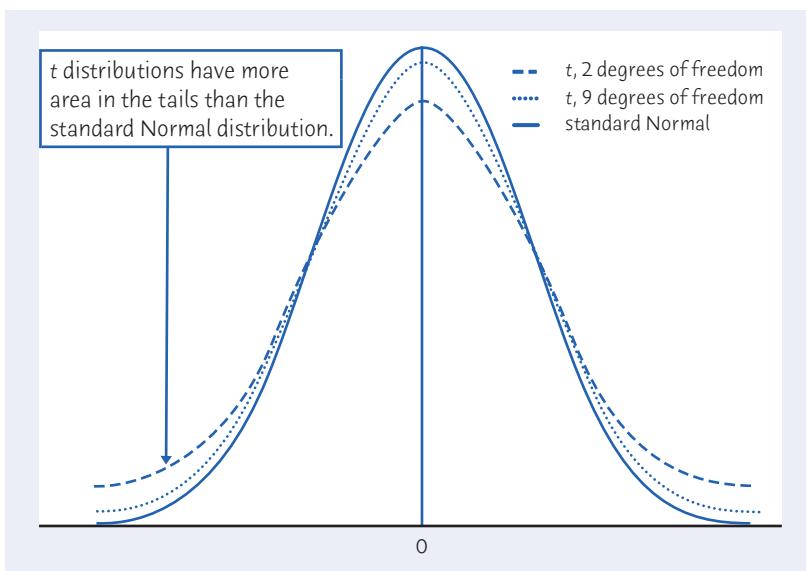


FIGURE 18.1

Density curves for the *t* distributions with 2 and 9 degrees of freedom and for the standard Normal distribution. All are symmetric with center 0. The *t* distributions are somewhat more spread out.

- As the degrees of freedom increase, the t density curve approaches the $N(0, 1)$ curve ever more closely. This happens because s estimates σ more accurately as the sample size increases. So using s in place of σ causes little extra variation when the sample is large.

Table C in the back of the book gives critical values for the t distributions. Each row in the table contains critical values for the t distribution whose degrees of freedom appear at the left of the row. For convenience, we label the table entries both by the confidence level C (in percent) required for confidence intervals and by the one-sided and two-sided P -values for each critical value. You have already used the standard Normal critical values in the z^* row at the bottom of Table C. By looking down any column, you can check that the t critical values approach the Normal values as the degrees of freedom increase. If you use statistical software, you don't need Table C.

EXAMPLE 18.1 t critical values

Figure 18.1 shows the density curve for the t distribution with 9 degrees of freedom. What point on this distribution has probability 0.05 to its right? In Table C, look in the $df = 9$ row above one-sided P -value .05 and you will find that this critical value is $t^* = 1.833$. To use software, enter the degrees of freedom and the probability you want to the left, 0.95 in this case. Here is Minitab's output:

```
Student's t distribution with 9 DF
P( X <= x )      x
    0.95  1.83311
```

APPLY YOUR KNOWLEDGE

18.3 Critical values.

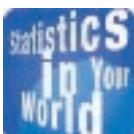
Use Table C or software to find

- the critical value for a one-sided test with level $\alpha = 0.05$ based on the $t(4)$ distribution.
- the critical value for a 98% confidence interval based on the $t(26)$ distribution.

18.4 More critical values.

You have an SRS of size 30 and calculate the one-sample t statistic. What is the critical value t^* such that

- t has probability 0.025 to the right of t^* ?
- t has probability 0.75 to the left of t^* ?



Better statistics, better beer

The t distribution and the t inference

procedures were invented by William S. Gosset (1876–1937). Gosset worked for the Guinness brewery, and his goal in life was to make better beer. He used his new t procedures to find the best varieties of barley and hops. Gosset's statistical work helped him become head brewer, a more interesting title than professor of statistics. Because Gosset published under the pen name "Student" you will often see the t distribution called "Student's t " in his honor.

THE ONE-SAMPLE t CONFIDENCE INTERVAL

To analyze samples from Normal populations with unknown σ , just replace the standard deviation σ/\sqrt{n} of \bar{x} by its standard error s/\sqrt{n} in the z procedures of Chapters 14, 15, and 16. The confidence interval and test that result are *one-sample t procedures*. Critical values and P -values come from the t distribution with $n - 1$ degrees of freedom. The one-sample t procedures are similar in both reasoning and computational detail to the z procedures.

THE ONE-SAMPLE t CONFIDENCE INTERVAL

Draw an SRS of size n from a large population having unknown mean μ . A level C confidence interval for μ is

$$\bar{x} \pm t^* \frac{s}{\sqrt{n}}$$

where t^* is the critical value for the $t(n - 1)$ density curve with area C between $-t^*$ and t^* . This interval is exact when the population distribution is Normal and is approximately correct for large n in other cases.

EXAMPLE 18.2 Good weather, good tips?

Let's look again at the study of tipping in a restaurant that we met in Example 14.3. We follow the four-step process for a confidence interval, outlined on page 359.



STATE: Does the expectation of good weather lead to more generous behavior? Psychologists studied the size of the tip in a restaurant when a message indicating that the next day's weather would be good was written on the bill. Here are tips from 20 patrons, measured in percent of the total bill:³



20.8	18.7	19.9	20.6	21.9	23.4	22.8	24.9	22.2	20.3
24.9	22.3	27.0	20.4	22.2	24.0	21.1	22.1	22.0	22.7

This is one of three sets of measurements made, the others being tips received when the message on the bill said that the next day's weather would not be good or there was no message on the bill. We want to estimate the mean tip for comparison with tips under the other conditions.

PLAN: We will estimate the mean percentage tip μ for all patrons of this restaurant when they receive a message on their bill indicating that the next day's weather will be good by giving a 95% confidence interval.

SOLVE: We must first check the conditions for inference.

- As in Chapter 14 (page 359), we are willing to regard these patrons as an SRS from all patrons of this restaurant.
- The stemplot in Figure 18.2 does not suggest any strong departures from Normality.

We can proceed to calculation. For these data,

$$\bar{x} = 22.21 \text{ and } s = 1.963$$

The degrees of freedom are $n - 1 = 19$. From Table C we find that for 95% confidence $t^* = 2.093$. The confidence interval is

$$\begin{aligned} \bar{x} \pm t^* \frac{s}{\sqrt{n}} &= 22.21 \pm 2.093 \frac{1.963}{\sqrt{20}} \\ &= 22.21 \pm 0.92 \\ &= 21.29 \text{ to } 23.13 \text{ percent} \end{aligned}$$

CONCLUDE: We are 95% confident that the mean percentage tip for all patrons of this restaurant when their bill contains a message that the next day's weather will be good is between 21.29 and 23.13. ■

18	7
19	9
20	3 4 6 8
21	1 9
22	0 1 2 2 3 7 8
23	4
24	0 9 9
25	
26	
27	0

FIGURE 18.2

Stemplot of the percentage tips, for Example 18.2.

Our work in Example 18.2 is very similar to what we did in Example 14.3 (page 359). To make the inference realistic we replaced the assumed $\sigma = 2$ by $s = 1.963$ calculated from the data and replaced the standard Normal critical value $z^* = 1.960$ by the t critical value $t^* = 2.093$.

The one-sample t confidence interval has the form

$$\text{estimate} \pm t^* \text{SE}_{\text{estimate}}$$

where “SE” stands for “standard error.” We will meet a number of confidence intervals that have this common form. In Example 18.2, the estimate is the sample mean \bar{x} , and its standard error is

$$\begin{aligned}\text{SE}_{\bar{x}} &= \frac{s}{\sqrt{n}} \\ &= \frac{1.963}{\sqrt{20}} = 0.439\end{aligned}$$

Software will find \bar{x} , s , $\text{SE}_{\bar{x}}$, and the confidence interval from the data. Figure 18.5 (page 447) displays typical software output for Example 18.2.

APPLY YOUR KNOWLEDGE

18.5 Critical values. What critical value t^* from Table C would you use for a confidence interval for the mean of the population in each of the following situations?

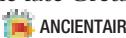
- (a) A 95% confidence interval based on $n = 12$ observations.
- (b) A 99% confidence interval from an SRS of 18 observations.
- (c) A 90% confidence interval from a sample of size 6.

18.6 How much will I bet? Our decisions depend on how the options are presented to us. Here's an experiment that illustrates this phenomenon. Tell 20 subjects that they have been given \$50 but can't keep it all. Then present them with a long series of choices between bets they can make with the \$50. Scattered among these choices in random order are 64 choices that ask the subject to choose between betting a fixed amount and an all-or-nothing gamble. The odds for all the bets are the same, but in 32 of the choices, the fixed option reads “Keep \$20” and in the other 32 choices the fixed option reads “Lose \$30.” These two fixed options lead to exactly the same outcome, but people are more likely to choose the fixed option that says they lose money. Here are the percent differences (“Number of times chose ‘Lose \$30’ minus ‘Number of times chose ‘Keep \$20’” divided by the number of trials on which the 20 subjects chose the fixed-option gamble rather than the all-or-nothing bet).⁴  GAMBLING1

37.5	30.8	6.2	17.6	14.3	8.3	16.7	20.0	10.5	21.7
30.8	27.3	22.7	38.5	8.3	10.5	8.3	10.5	25.0	7.7

- (a) Make a stemplot. Is there any sign of a major deviation from Normality?
- (b) All 20 subjects gambled a fixed amount more often when faced with a sure loss than when faced with a sure win. Give a 95% confidence interval for the mean percent increase in gambling a fixed amount when faced with a sure loss.

18.7 Ancient air. The composition of the earth's atmosphere may have changed over time. To try to discover the nature of the atmosphere long ago, we can examine the gas in bubbles inside ancient amber. Amber is tree resin that has hardened and been trapped in rocks. The gas in bubbles within amber should be a sample of the atmosphere at the time the amber was formed. Measurements on specimens of amber from the late Cretaceous era (75 to 95 million years ago) give these percents of nitrogen:⁵



63.4 65.0 64.4 63.3 54.8 64.5 60.8 49.1 51.0

Assume (this is not yet agreed on by experts) that these observations are an SRS from the late Cretaceous atmosphere. Use a 90% confidence interval to estimate the mean percent of nitrogen in ancient air. Follow the four-step process as illustrated in Example 18.2. (Our present-day atmosphere is about 78.1% nitrogen.)



David Sanger Photography/Alamy

THE ONE-SAMPLE t TEST

Like the confidence interval, the t test is very similar to the z test we met earlier.

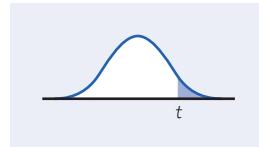
THE ONE-SAMPLE t TEST

Draw an SRS of size n from a large population having unknown mean μ . To **test the hypothesis $H_0: \mu = \mu_0$** , compute the **one-sample t statistic**

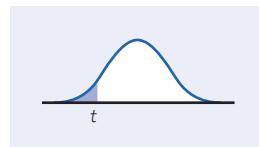
$$t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}}$$

In terms of a variable T having the $t(n - 1)$ distribution, the P -value for a test of H_0 against

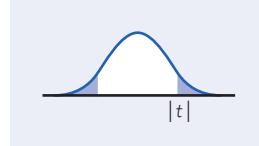
$$H_a: \mu > \mu_0 \text{ is } P(T \geq t)$$



$$H_a: \mu < \mu_0 \text{ is } P(T \leq t)$$



$$H_a: \mu \neq \mu_0 \text{ is } 2P(T \geq |t|)$$



These P -values are exact if the population distribution is Normal and are approximately correct for large n in other cases.



EXAMPLE 18.3 Water quality

We follow the four-step process for a significance test, outlined on page 379.

STATE: To investigate water quality, on August 8, 2010, the *Columbus Dispatch* took water samples at 20 Ohio State Park swimming areas. Those samples were taken to laboratories and tested for fecal coliform, which are bacteria found in human and animal feces. An unsafe level of fecal coliform means there's a higher chance that disease-causing bacteria are present and more risk that a swimmer will become ill. Ohio considers it unsafe if a 100-milliliter sample (about 3.3 ounces) of water contains more than 400 coliform bacteria. Here are the fecal coliform levels found by the laboratories:⁶

160	40	2800	80	2000	2000	1500	400	150	500
3000	2200	15	80	2000	2000	2600	600	1000	1500

Are these data good evidence that, on average, the fecal coliform levels in these swimming areas were unsafe?

PLAN: Experts caution that the tests are a snapshot of the quality of the water at the time they were taken. Fecal coliform levels can change as weather and other conditions change. So we ask the question in terms of the mean fecal coliform level μ for all these swimming areas. The null hypothesis is “level is not unsafe,” and the alternative hypothesis is “level is unsafe.”

$$H_0: \mu = 400$$

$$H_a: \mu > 400$$

SOLVE: First check the conditions for inference. We are willing to regard these particular 20 samples as an SRS from a large population of possible samples. Figure 18.3 is a histogram of the data. We can't accurately judge Normality from 20 observations; there are no outliers but the data are somewhat skewed. *P-values for the t test may be only approximately accurate.*

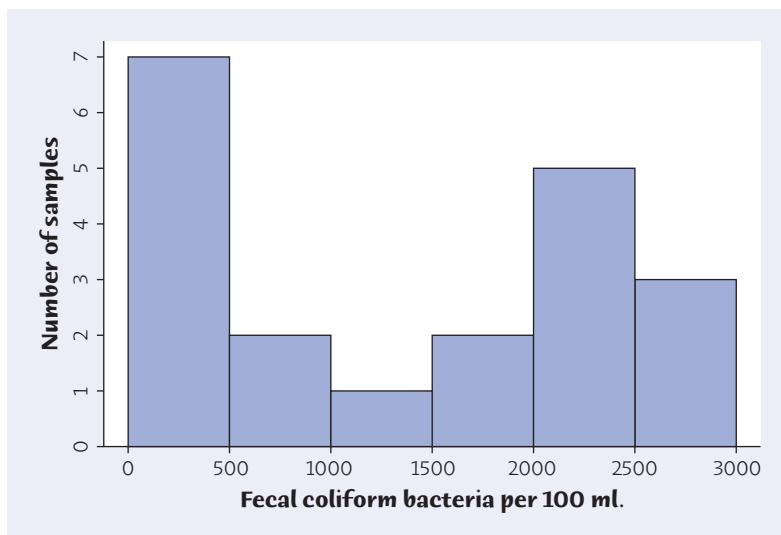
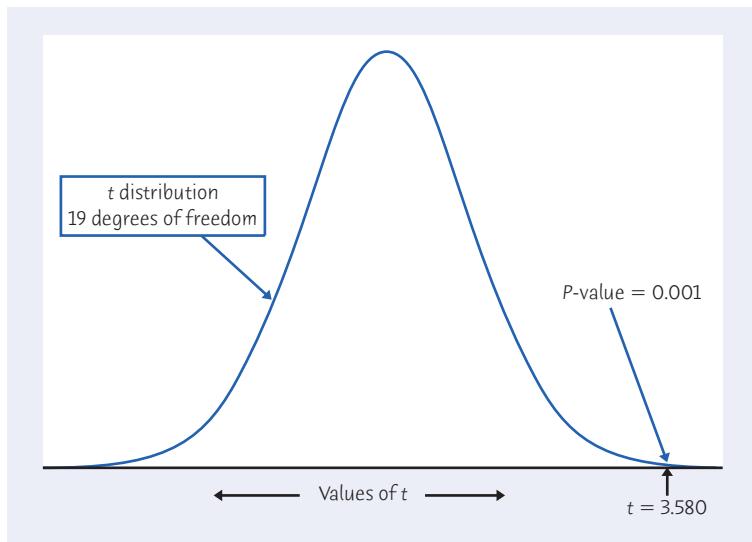


FIGURE 18.3

Histogram of the fecal coliform levels, for Example 18.3.

**FIGURE 18.4**

The P -value for the one-sided t test in Example 18.3.

The basic statistics are

$$\bar{x} = 1231 \text{ and } s = 1038$$

The one-sample t statistic is

$$\begin{aligned} t &= \frac{\bar{x} - \mu_0}{s/\sqrt{n}} = \frac{1231 - 400}{1038/\sqrt{20}} \\ &= 3.580 \end{aligned}$$

The P -value for $t = 3.580$ is the area to the right of 3.580 under the t distribution curve with degrees of freedom $n - 1 = 19$. Figure 18.4 shows this area. Software (see Figure 18.6) tells us that $P = 0.001$.

Without software, we can pin P between two values by using Table C. Search the $df = 19$ row of Table C for entries that bracket $t = 3.580$. The observed t lies between the critical values for one-sided P -values 0.001 and 0.0005.

CONCLUDE: There is quite strong evidence ($P < 0.001$) that, on average, fecal coliform levels in these Ohio State Park swimming areas are unsafe. ■

df = 19		
t^*	3.579	3.883
One-sided P	.001	.0005

APPLY YOUR KNOWLEDGE

18.8 Is it significant? The one-sample t statistic for testing

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

from a sample of $n = 20$ observations has the value $t = 1.84$.

- What are the degrees of freedom for this statistic?
- Give the two critical values t^* from Table C that bracket t . What are the one-sided P -values for these two entries?

- (c) Is the value $t = 1.84$ significant at the 5% level? Is it significant at the 1% level?
- (d) (Optional) If you have access to suitable technology, give the exact one-sided P -value for $t = 1.84$.

18.9 Is it significant? The one-sample t statistic from a sample of $n = 15$ observations for the two-sided test of

$$\begin{aligned} H_0: \mu &= 64 \\ H_a: \mu &\neq 64 \end{aligned}$$

has the value $t = 2.12$.

- (a) What are the degrees of freedom for t ?
- (b) Locate the two critical values t^* from Table C that bracket t . What are the two-sided P -values for these two entries?
- (c) Is the value $t = 2.12$ statistically significant at the 10% level? At the 5% level?
- (d) (Optional) If you have access to suitable technology, give the exact two-sided P -value for $t = 2.12$.

18.10 Ancient air, continued. Do the data of Exercise 18.7 give good reason to think that the percent of nitrogen in the air during the Cretaceous era was different from the present 78.1%? Carry out a test of significance, following the four-step process as illustrated in Example 18.3.  ANCIENTAIR



USING TECHNOLOGY

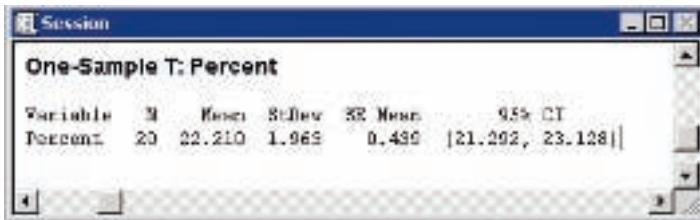
Any technology suitable for statistics will implement the one-sample t procedures. As usual, you can read and use almost any output now that you know what to look for. Figure 18.5 displays output for the 95% confidence interval of Example 18.2 from a graphing calculator, a statistical program, a spreadsheet program, and the CrunchIt! software package. The calculator, Minitab, and CrunchIt! outputs are straightforward. All three give the estimate \bar{x} and the confidence interval plus a clearly labeled selection of other information. The confidence interval agrees with our hand calculation in Example 18.2. In general, software results are more accurate because of the rounding in hand calculations. Excel gives several descriptive measures but does not give the confidence interval. The entry labeled “Confidence Level (95.0%)” is the margin of error. You can use this together with \bar{x} to get the interval using either a calculator or the spreadsheet’s formula capability.

Figure 18.6 displays output for the t test in Example 18.3. The graphing calculator, Minitab, and CrunchIt! give the sample mean \bar{x} , the t statistic, and its P -value. Accurate P -values are the biggest advantage of software for the t procedures. Excel is, as usual, more awkward than software designed for statistics. It lacks a one-sample t test menu selection but does have a function named TDIST for tail areas under t density curves. The Excel output shows functions for the t statistic and its P -value to the right of the main display, along with their values $t = 3.582967678$ and $P = 0.00099192$.

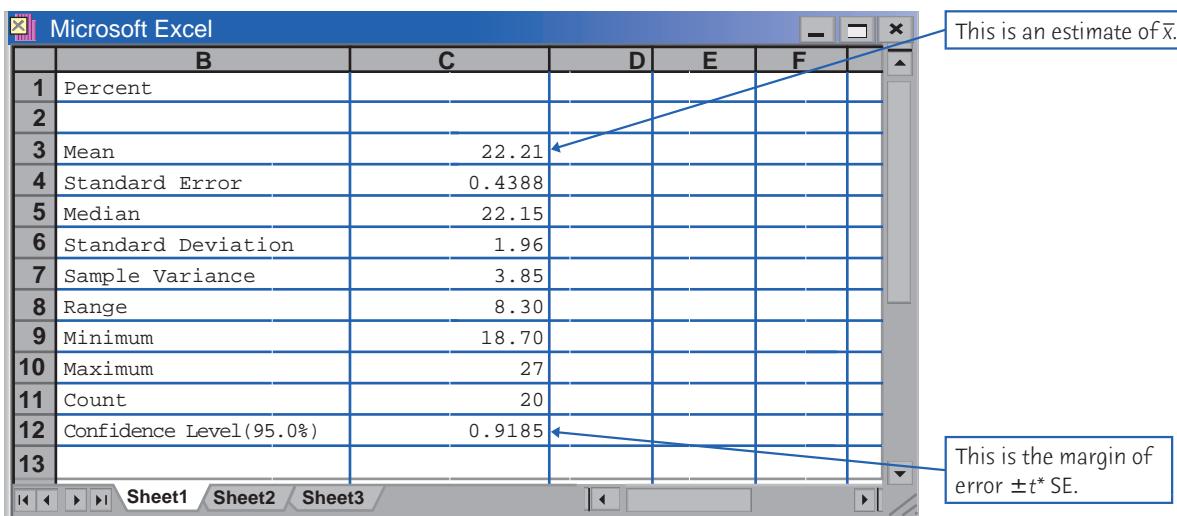
Texas Instruments Graphing Calculator

```
TInterval
(21.292, 23.128)
x=22.2100
Sx=1.9625
n=20.0000
```

Minitab



Excel



CrunchIt!

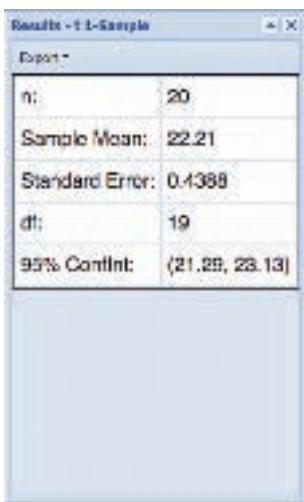


FIGURE 18.5

The t confidence interval for Example 18.2: output from a graphing calculator, two statistical programs, and a spreadsheet program.

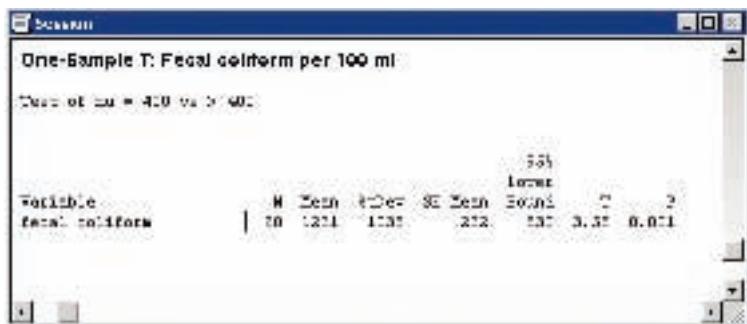
Texas Instruments Graphing Calculator

```

T-Test
μ>400.0000
t=3.5803
P=.0010
x=1231.2500
Sx=1037.5374
n=20.0000

```

Minitab



Excel

The Microsoft Excel output table shows various statistical measures for 'fecal coliform per 100 ml'. The t statistic is calculated as (C3.400)/C4 (3.582967678) and the P-value is given as TDIST (E2,19,1) (0.00099192).

	B	C	D	E	F
1	fecal coliform per 100 ml				
2			(C3.400)/C4	3.582967678	This is the t statistic.
3	Mean	1231.25	TDIST (E2,19,1)	0.00099192	This is the P-value.
4	Standard Error	232.0004			
5	Median	1250			
6	Standard Deviation	1037.5374			
7	Sample Variance	1076483.88			
8	Range	2985			
9	Minimum	15			
10	Maximum	3000			
11	Count	20			
12	Confidence Level(95%)	485.5825			
13					

CrunchIt!

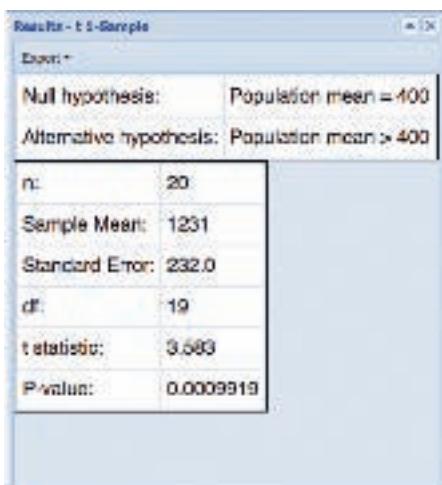


FIGURE 18.6

The t test for Example 18.3: output from a graphing calculator, two statistical programs, and a spreadsheet program.

MATCHED PAIRS *t* PROCEDURES

Often the goal of an investigation is to demonstrate that a treatment causes an observed effect. In Chapter 9 we learned that randomized comparative studies are more convincing than single-sample investigations for demonstrating causation. For that reason, one-sample inference is less common than comparative inference. One common design to compare two treatments makes use of one-sample procedures. Matched pairs designs were discussed in Chapter 9. In a **matched pairs design**, subjects are matched in pairs and each treatment is given to one subject in each pair. Another situation calling for matched pairs is before-and-after observations on the same subjects.

matched pairs design

MATCHED PAIRS *t* PROCEDURES

To compare the responses to the two treatments in a matched pairs design, find the difference between the responses within each pair. Then apply the one-sample *t* procedures to these differences.

The parameter μ in a matched pairs *t* procedure is the mean difference in the responses to the two treatments within matched pairs of subjects in the entire population.

EXAMPLE 18.4 Do chimpanzees collaborate?

STATE: Humans often collaborate to solve problems. Will chimpanzees recruit another chimp when solving a problem requires collaboration? Researchers presented chimpanzee subjects with food outside their cage that they could bring within reach by pulling two ropes, one attached to each end of the food tray. If a chimp pulled only one rope, the rope came loose and the food was lost. Another chimp was available as a partner, but only if the subject unlocked a door joining two cages. (Chimpanzees learn these things quickly.) The same 8 chimpanzee subjects faced this problem in two versions: the two ropes were close enough together that one chimp could pull both (no collaboration needed) or the two ropes were too far apart for one chimp to pull both (collaboration needed). Table 18.1 shows how often in 24 trials for each version each subject opened the door to recruit another chimp as partner.⁷ Is there evidence that chimpanzees recruit partners more often when a problem requires collaboration?

PLAN: Take μ to be the mean difference (collaboration required minus not) in the number of times a subject recruited a partner. The null hypothesis says that the need for collaboration has no effect, and H_a says that partners are recruited more often when the problem requires collaboration. So we test the hypotheses

$$H_0: \mu = 0$$

$$H_a: \mu > 0$$

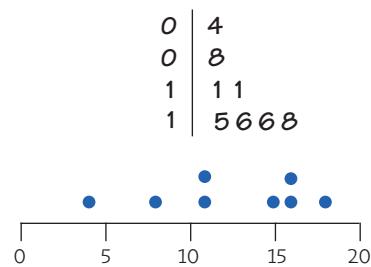
SOLVE: The subjects are “semi-free-ranging chimpanzees at Ngamba Island Chimpanzee Sanctuary in Uganda.” We are willing to regard them as an SRS from their species. To analyze the data, we examine the difference in the number of times a chimp



CHIMPS



Manoj Shah/Getty

**FIGURE 18.7**

Stemplot and dotplot of the differences, for Example 18.4.

$df = 7$		
t^*	4.785	5.408
One-sided P	.001	.0005

recruited a partner, so subtract the “no collaboration needed” count from the “collaboration needed” count for each subject. The 8 differences form a single sample from a population with unknown mean μ . They appear in the “Difference” column in Table 18.1. All the chimpanzees recruited a partner more often when the ropes were too far apart to be pulled by one chimp.

The stemplot in Figure 18.7 creates the impression of a left-skew. This is a bit misleading, as the dotplot in the bottom part of Figure 18.7 shows. A dotplot simply places the observations on an axis, stacking observations that have the same value. It gives a good picture of distributions with only whole-number values. We know that observations that can take only whole-number values cannot come from a Normal population. In practice, researchers are willing to treat such observations as coming from a Normal population if there are more than just a few possible values and the distribution appears approximately Normal. Of course, we can't assess approximate Normality from just 8 observations, but there are no signs of major departures from Normality. The researchers used the matched pairs t test.

The 8 differences have

$$\bar{x} = 12.375 \text{ and } s = 4.749$$

The one-sample t statistic is therefore

$$t = \frac{\bar{x} - 0}{s/\sqrt{n}} = \frac{12.375 - 0}{4.749/\sqrt{8}} \\ = 7.37$$

Find the P -value from the $t(7)$ distribution. (Remember that the degrees of freedom are 1 less than the sample size.) Table C shows that 7.37 is greater than the critical value for one-sided $P = 0.0005$. The P -value is therefore less than 0.0005. Software says that $P = 0.000077$.

CONCLUDE: The data give very strong evidence ($P < 0.0005$) that chimpanzees recruit a collaborator more often when faced with a problem that requires a collaborator to solve. That is, chimpanzees recognize when collaboration is necessary, a skill that they share with humans. ■

TABLE 18.1 Trials (out of 24) on which chimpanzees recruited a partner

CHIMPANZEE	COLLABORATION NEEDED		
	YES	NO	DIFFERENCE
Namuiska	16	0	16
Kalema	16	1	15
Okech	23	5	18
Baluku	19	3	16
Umugenzi	15	4	11
Indi	20	9	11
Bili	24	16	8
Asega	24	20	4

Example 18.4 illustrates how to turn matched pairs data into single-sample data by taking differences within each pair. We are making inferences about a single population, the population of all differences within matched pairs. *It is incorrect to ignore the matching and analyze the data as if we had two samples of chimpanzees, one facing ropes close together and the other facing ropes far apart.* Inference procedures for comparing two samples assume that the samples are selected independently of each other. This condition does not hold when the same subjects are measured twice. The proper analysis depends on the design used to produce the data.



APPLY YOUR KNOWLEDGE

Many exercises from this point on ask you to give the *P*-value of a *t* test. If you have suitable technology, give the exact *P*-value. Otherwise, use Table C to give two values between which *P* lies.

18.11 The brain responds to sound. The usual way to study the brain's response to sounds is to have subjects listen to "pure tones." The response to recognizable sounds may differ. To compare responses, researchers anesthetized macaque monkeys. They fed pure tones and also monkey calls directly to their brains by inserting electrodes. Response to the stimulus was measured by the firing rate (electrical spikes per second) of neurons in various areas of the brain. Table 18.2 contains the responses for 37 neurons.⁸ Researchers suspected that the response to monkey calls would be stronger than the response to a pure tone. Do the data support this idea? Complete the "Plan," "Solve," and "Conclude" steps of the four-step process, following the model of Example 18.4.

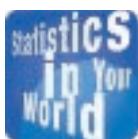


18.12 The brain responds, continued. How much more strongly do monkey brains respond to monkey calls than to pure tones? Give a 99% confidence interval to answer this question.



TABLE 18.2 Neuron response to tones and monkey calls

TONE	CALL	TONE	CALL	TONE	CALL	TONE	CALL
474	500	145	42	71	134	35	103
256	138	141	241	68	65	31	70
241	485	129	194	59	182	28	192
226	338	113	123	59	97	26	203
185	194	112	182	57	318	26	135
174	159	102	141	56	201	21	129
176	341	100	118	47	279	20	193
168	85	74	62	46	62	20	54
161	303	72	112	41	84	19	66
150	208						



Catching cheaters

A certification test for surgeons asks 277 multiple-choice questions. Smith and Jones have 193 common right answers and 53 identical wrong choices. The computer flags their 246 identical answers as evidence of possible cheating. They sue. The court wants to know how unlikely it is that exams this similar would occur just by chance. That is, the court wants a P -value. Statisticians offer several P -values based on different models for the exam-taking process. They all say that results this similar would almost never happen just by chance. Smith and Jones fail the exam.

ROBUSTNESS OF t PROCEDURES

The t confidence interval and test are exactly correct when the distribution of the population is exactly Normal. No real data are exactly Normal. The usefulness of the t procedures in practice therefore depends on how strongly they are affected by lack of Normality.

ROBUST PROCEDURES

A confidence interval or significance test is called **robust** if the confidence level or P -value does not change very much when the conditions for use of the procedure are violated.

The condition that the population is Normal rules out outliers, so the presence of outliers shows that this condition is not fulfilled. The t procedures are not robust against outliers unless the sample is large, because \bar{x} and s are not resistant to outliers.

Fortunately, the t procedures are quite robust against non-Normality of the population except when outliers or strong skewness are present. (Skewness is more serious than other kinds of non-Normality.) As the size of the sample increases, the central limit theorem ensures that the distribution of the sample mean \bar{x} becomes more nearly Normal and that the t distribution becomes more accurate for critical values and P -values of the t procedures.

Always make a plot to check for skewness and outliers before you use the t procedures for small samples. For most purposes, you can safely use the one-sample t procedures when $n \geq 15$ unless an outlier or quite strong skewness is present. Here are practical guidelines for inference on a single mean.⁹

USING THE t PROCEDURES

- Except in the case of small samples, the condition that the data are an SRS from the population of interest is more important than the condition that the population distribution is Normal.
- *Sample size less than 15:* Use t procedures if the data appear close to Normal (roughly symmetric, single peak, no outliers). If the data are clearly skewed or if outliers are present, do not use t .
- *Sample size at least 15:* The t procedures can be used except in the presence of outliers or strong skewness.
- *Large samples:* The t procedures can be used even for clearly skewed distributions when the sample is large, roughly $n \geq 40$.

EXAMPLE 18.5 Can we use t ?

Figure 18.8 shows plots of several data sets. For which of these can we safely use the t procedures?¹⁰

- Figure 18.8(a) is a histogram of the percent of each state's adult residents who are college graduates. *We have data on the entire population of 50 states, so inference is not needed.* We can calculate the exact mean for the population. There is no uncertainty



due to having only a sample from the population, and no need for a confidence interval or test. If these data were an SRS from a larger population, t inference would be safe despite the mild skewness because $n = 50$.

- Figure 18.8(b) is a stemplot of the force required to pull apart 20 pieces of Douglas fir. The data are strongly skewed to the left with possible low outliers, so we cannot trust the t procedures for $n = 20$.

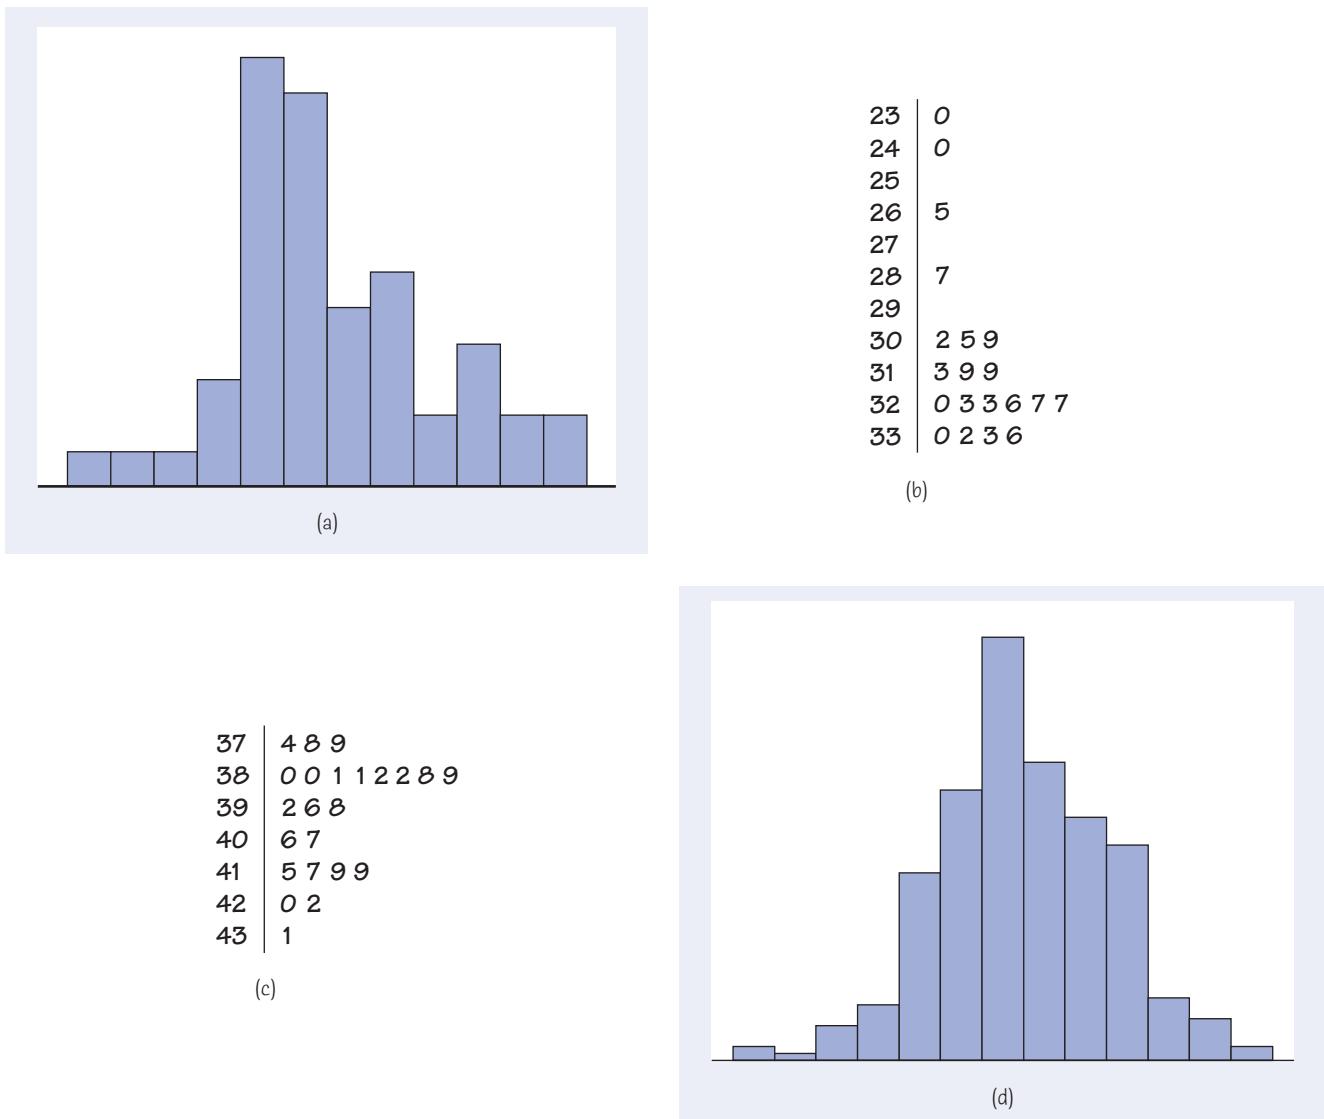


FIGURE 18.8

Can we use t procedures for these data? (a) Percent of adult college graduates in the 50 states. No, this is an entire population, not a sample. (b) Force required to pull apart 20 pieces of Douglas fir. No, there are just 20 observations and strong skewness. (c) Lengths of 23 tropical flowers of the same variety. Yes, the sample is large enough to overcome the mild skewness. (d) Heights of college students in a college class. Yes, for any size sample, because the distribution is close to Normal.

■ Figure 18.8(c) is a stemplot of the lengths of 23 specimens of the red variety of the tropical flower *Heliconia*. The data contain no outliers and so the skewness is relatively mild. We can use the *t* distributions for such data.

■ Figure 18.8(d) is a histogram of the heights of the students in a college class. This distribution is quite symmetric and appears close to Normal. We can use the *t* procedures for any sample size. ■

APPLY YOUR KNOWLEDGE



Eric Nathan/Alamy

18.13 Diamonds. A group of earth scientists studied the small diamonds found in a nodule of rock carried up to the earth's surface in surrounding rock. This is an opportunity to examine a sample from a single population of diamonds formed in a single event deep in the earth.¹¹ Table 18.3 presents data on the nitrogen content (parts per million) and the abundance of carbon 13 in these diamonds. (Carbon has several isotopes, forms with different numbers of neutrons in the nuclei of their atoms. Carbon 12 makes up almost 99% of natural carbon. The abundance of carbon 13 is measured by the ratio of carbon 13 to carbon 12, in parts per thousand more or less than a standard. The minus signs in the data mean that the ratio is smaller in these diamonds than in standard carbon.) DIAMONDS

We would like to estimate the mean abundance of both nitrogen and carbon 13 in the population of diamonds represented by this sample. Examine the data for nitrogen. Can we use a *t* confidence interval for mean nitrogen? Explain your answer. Give a 90% confidence interval if you think the result can be trusted.

18.14 Diamonds, continued. Examine the data in Table 18.3 on abundance of carbon 13. Can we use a *t* confidence interval for mean carbon 13? Explain your answer. Give a 90% confidence interval if you think the result can be trusted. DIAMONDS

TABLE 18.3 Nitrogen and carbon 13 in a sample of diamonds

DIAMOND	NITROGEN (PPM)	CARBON 13 RATIO	DIAMOND	NITROGEN (PPM)	CARBON 13 RATIO
1	487	-2.78	13	273	-2.73
2	1430	-1.39	14	94	-2.33
3	60	-4.26	15	69	-3.83
4	244	-1.19	16	262	-2.04
5	196	-2.12	17	120	-2.82
6	274	-2.87	18	302	-0.84
7	41	-3.68	19	75	-3.57
8	54	-3.29	20	242	-2.42
9	473	-3.79	21	115	-3.89
10	30	-4.06	22	65	-3.87
11	98	-1.83	23	311	-1.58
12	41	-4.03	24	61	-3.97

CHAPTER 18 SUMMARY

CHAPTER SPECIFICS

- Tests and confidence intervals for the mean μ of a Normal population are based on the sample mean \bar{x} of an SRS. Because of the central limit theorem, the resulting procedures are approximately correct for other population distributions when the sample is large.
- The standardized sample mean is the **one-sample z statistic**

$$z = \frac{\bar{x} - \mu}{\sigma/\sqrt{n}}$$

If we knew σ , we would use the z statistic and the standard Normal distribution.

- In practice, we do not know σ . Replace the standard deviation σ/\sqrt{n} of \bar{x} by the **standard error** s/\sqrt{n} to get the **one-sample t statistic**

$$t = \frac{\bar{x} - \mu}{s/\sqrt{n}}$$

The t statistic has the **t distribution** with $n - 1$ degrees of freedom.

- There is a t distribution for every positive **degrees of freedom**. All are symmetric distributions similar in shape to the standard Normal distribution. The t distribution approaches the $N(0, 1)$ distribution as the degrees of freedom increase.
- A level C **confidence interval for the mean μ** of a Normal population is

$$\bar{x} \pm t^* \frac{s}{\sqrt{n}}$$

The **critical value t^*** is chosen so that the t curve with $n - 1$ degrees of freedom has area C between $-t^*$ and t^* .

- **Significance tests** for $H_0: \mu = \mu_0$ are based on the t statistic. Use P -values or fixed significance levels from the $t(n - 1)$ distribution.
- Use these one-sample procedures to analyze **matched pairs** data by first taking the difference within each matched pair to produce a single sample.
- The t procedures are quite **robust** when the population is non-Normal, especially for larger sample sizes. The t procedures are useful for non-Normal data when $n \geq 15$ unless the data show outliers or strong skewness. When $n \geq 40$, the t procedures can be used even for clearly skewed distributions.

LINK IT

In Chapters 14 to 16 we began our study of inference. We focused on inference for a population mean based on a sample from a Normal population. We included the unrealistic assumption that we knew the population standard deviation σ . The reason was to make the underlying mathematics simpler. We were able to use what we learned about the Normal distribution in Chapters 3 and 10 and what we learned about the sampling distribution of the sample mean in Chapter 11 to construct confidence intervals for and conduct hypothesis tests about the population mean. Computing P -values and sample sizes was possible using what we had learned about Normal probability calculations.

In this chapter we continue our study of inference about a population mean based on a sample from a Normal population in the more realistic setting that we do not know σ . The basic ideas of Chapters 14 to 16 still apply, but now we use the t distribution rather than the standard Normal distribution. Unfortunately, the mathematics associated with the t distribution is more complicated than that associated with the standard Normal distribution. We must rely on approximations from tables or statistical software to calculate P -values (and determining sample sizes is also more complicated when we use the t distribution). As we continue our study of statistical inference in realistic settings, statistical software will be invaluable.

As we saw in Chapter 9, statistical studies comparing two or more groups are preferable to one-sample procedures if we want to demonstrate that a treatment causes an observed response. As a first step toward developing methods for comparing means from two populations, we considered matched pairs studies (also discussed in Chapter 9) and saw that by looking at differences, we could use our one-sample t procedures to compare two means. In the next chapter, we consider inference for comparing the means of two populations when we have independent samples from the two populations. This will expand our tools for doing statistical inference in settings that we encounter in practice.

CHECK YOUR SKILLS

18.15 We prefer the t procedures to the z procedures for inference about a population mean because

- (a) z can be used only for large samples.
- (b) z requires that you know the population standard deviation σ .
- (c) z requires that you can regard your data as an SRS from the population.

18.16 You are testing $H_0: \mu = 10$ against $H_a: \mu < 10$ based on an SRS of 16 observations from a Normal population. The data give $\bar{x} = 8$ and $s = 4$. The value of the t statistic is

- (a) -0.5 .
- (b) -2 .
- (c) -8 .

18.17 You are testing $H_0: \mu = 100$ against $H_a: \mu < 100$ based on an SRS of 25 observations from a Normal population. The t statistic is $t = -2.5$. The degrees of freedom for the t statistic are

- (a) 26.
- (b) 25.
- (c) 24.

18.18 The P -value for the statistic in the previous exercise

- (a) falls between 0.02 and 0.04.
- (b) falls between 0.01 and 0.02.
- (c) is less than 0.01.

18.19 You have an SRS of 12 observations from a Normally distributed population. What critical value would you use to obtain a 98% confidence interval for the mean μ of the population?

- (a) 2.718
- (b) 2.681
- (c) 2.650

18.20 You are testing $H_0: \mu = 0$ against $H_a: \mu \neq 0$ based on an SRS of 12 observations from a Normal population. What values of the t statistic are statistically significant at the $\alpha = 0.005$ level?

- (a) $t > 3.497$
- (b) $t < -3.497$ or $t > 3.497$
- (c) $t < -3.428$ or $t > 3.428$

18.21 Data on the blood cholesterol levels of 10 rats (milligrams per deciliter of blood) give $\bar{x} = 85$ and $s = 12$. A 99% confidence interval for the mean blood cholesterol of rats is

- (a) 76.4 to 93.6.
- (b) 73.0 to 97.0.
- (c) 72.7 to 97.3.

18.22 Which of the following would cause the most worry about the validity of the confidence interval you calculated in the previous exercise?

- (a) There is a clear outlier in the data.
- (b) A stemplot of the data shows a mild right-skew.
- (c) You do not know the population standard deviation σ .

18.23 Which of these settings does not allow use of a matched pairs t procedure?

- (a) You interview both the husband and the wife in 64 married couples and ask each about their ideal number of children.
- (b) You interview a sample of 64 unmarried male students and another sample of 64 unmarried female students and ask each about their ideal number of children.
- (c) You interview 64 female students in their freshman year and again in their senior year and ask each about their ideal number of children.

18.24 Because the t procedures are robust, the most important condition for their safe use is that

- (a) the population standard deviation σ is known.
- (b) the population distribution is exactly Normal.
- (c) the data can be regarded as an SRS from the population.

CHAPTER 18 EXERCISES

18.25 Read carefully. You read in the report of a psychology experiment: “Separate analyses for our two groups of 12 participants revealed no overall placebo effect for our student group (mean = 0.08, SD = 0.37, $t(11) = 0.49$) and a significant effect for our non-student group (mean = 0.35, SD = 0.37, $t(11) = 3.25, p < 0.01$).”¹² The null hypothesis is that the mean effect is zero. What are the correct values of the two t statistics based on the means and standard deviations? Compare each correct t -value with the critical values in Table C. What can you say about the two-sided P -value in each case?

18.26 Body mass index of young women. In Example 14.1 (page 352) we developed a 95% z confidence interval for the mean body mass index (BMI) of women aged 20 to 29 years, based on a national random sample of 654 such women. We assumed there that the population standard deviation was known to be $\sigma = 7.5$. In fact, the sample data had mean BMI $\bar{x} = 26.8$ and standard deviation $s = 7.42$. What is the 95% t confidence interval for the mean BMI of all young women?

18.27 Reading scores in Atlanta. The Trial Urban District Assessment (TUDA) is a government-sponsored study of student achievement in large urban school districts. TUDA gives a reading test scored from 0 to 500. A score of 243 is a “basic” reading level and a score of 281 is “proficient.” Scores for a random sample of 3000 eighth-graders in Atlanta had $\bar{x} = 250$ with standard error 1.0.¹³

- (a) We don’t have the 3000 individual scores, but use of the t procedures is surely safe. Why?
- (b) Give a 99% confidence interval for the mean score of all Atlanta eighth-graders. (Be careful: the report gives the standard error of \bar{x} , not the standard deviation s .)
- (c) Urban children often perform below the basic level. Is there good evidence that the mean for all Atlanta eighth-graders is less than the basic level?

18.28 Color and cognition. In a randomized comparative experiment on the effect of color on the performance of a cognitive task, researchers randomly divided 69 subjects (27 males and 42 females ranging in age from 17 to 25 years) into three groups. Participants were asked to solve a series of 6 anagrams. One group was presented with the anagrams on a blue screen; one group saw them on a red screen; and one group had a neutral screen. The time, in seconds, taken to solve the anagrams was recorded. The paper reporting the study gives $\bar{x} = 11.58$ and $s = 4.37$ for the times of the 23 members of the neutral group.¹⁴

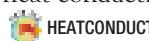
- (a) Give a 95% confidence interval for the mean time in the population from which the subjects were recruited.

(b) What conditions for the population and the study design are required by the procedure you used in (a)? Which of these conditions are important for the validity of the procedure in this case?

18.29 The placebo effect. The placebo effect is particularly strong in patients with Parkinson’s disease. To understand the workings of the placebo effect, scientists measure activity at a key point in the brain when patients receive a placebo that they think is an active drug and also when no treatment is given.¹⁵ The same 6 patients are measured both with and without the placebo, at different times.

- (a) Explain why the proper procedure to compare the mean response to placebo with control (no treatment) is a matched pairs t test.
- (b) The 6 differences (treatment minus control) had $\bar{x} = -0.326$ and $s = 0.181$. Is there significant evidence of a difference between treatment and control?

18.30 The conductivity of fibrous-glass board. How well materials conduct heat matters when designing houses, for example. Conductivity is measured in terms of watts of heat power transmitted per square meter of surface per degree Celsius of temperature difference on the two sides of the material. The National Institute of Standards and Technology (NIST) provides data on properties of materials. Here are 9 NIST measurements of the heat conductivity of a particular type of fibrous-glass board:¹⁶



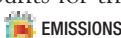
0.0339	0.0337	0.0334	0.0334	0.0333
0.0333	0.0333	0.0332	0.0330	

- (a) We can consider this an SRS of all specimens of fibrous-glass board of this type. Make a stemplot. Is there any sign of major deviation from Normality?
- (b) Give a 95% confidence interval for the mean conductivity.
- (c) Is there significant evidence at the 5% level that the mean conductivity of this type of fibrous-glass board is not 0.0330?

18.31 Exhaust from school buses. In a study of exhaust emissions from school buses, the pollution intake by passengers was determined for a sample of 9 school buses used in the Southern California Air Basin. The pollution intake is the amount of exhaust emissions, in grams per person, that would be breathed in while traveling on the bus during its usual 18-mile trip on congested freeways from South Central LA to a magnet school in West LA. (As a reference, the average intake of motor emissions of carbon monoxide in the LA area is estimated to be about 0.000046 grams per person.)

Here are the amounts for the 9 buses when driven with the windows open:¹⁷

1.15 0.33 0.40 0.33 1.35 0.38 0.25 0.40 0.35



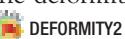
(a) Make a stemplot. Are there outliers or strong skewness that would forbid use of the *t* procedures?

(b) A good way to judge the effect of outliers is to do your analysis twice, once with the outliers and a second time without them. Give two 90% confidence intervals, one with all the data and one with the outliers removed, for the mean pollution intake among all school buses used in the Southern California Air Basin that travel the route investigated in the study.

(c) Compare the two intervals in part (b). What is the most important effect of removing the outliers?

18.32 A big-toe problem.

Hallux abducto valgus (call it HAV) is a deformation of the big toe that often requires surgery. Doctors used X-rays to measure the angle (in degrees) of deformity in 38 consecutive patients under the age of 21 who came to a medical center for surgery to correct HAV. The angle is a measure of the seriousness of the deformity. Here are the data:¹⁸



Wellcome Trust Medical Library/
Custom Medical Stock Photo

28 32 25 34 38 26 25 18 30 26 28 13 20
21 17 16 21 23 14 32 25 21 22 20 18 26
16 30 30 20 50 25 26 28 31 38 32 21

It is reasonable to regard these patients as a random sample of young patients who require HAV surgery. Carry out the “Solve” and “Conclude” steps for a 95% confidence interval for the mean HAV angle in the population of all such patients.

18.33 An outlier’s effect. Our bodies have a natural electrical field that is known to help wounds heal. Does changing the field strength slow healing? A series of experiments with newts investigated this question. In one experiment, the two hind limbs of 12 newts were assigned at random to either experimental or control groups. This is a matched pairs design. The electrical field in the experimental limbs was reduced to zero by applying a voltage. The control limbs were left alone. Here are the rates at which new cells closed a razor cut in each limb, in micrometers per hour:¹⁹



Newt	1	2	3	4	5	6	7	8	9	10	11	12
Control limb	36	41	39	42	44	39	39	56	33	20	49	30
Experimental limb	28	31	27	33	33	38	45	25	28	33	47	23

(a) Make a stemplot of the differences between limbs of the same newt (control limb minus experimental limb). There is a high outlier.

(b) A good way to judge the effect of an outlier is to do your analysis twice, once with the outlier and a second time without it. Carry out two *t* tests to see if the mean healing rate is significantly lower in the experimental limbs, one test including all 12 newts and another that omits the outlier. What are the test statistics and their *P*-values? Does the outlier have a strong influence on your conclusion?

18.34 An outlier’s effect. A good way to judge the effect of an outlier is to do your analysis twice, once with the outlier and a second time without it. The data in Exercise 18.32 follow a Normal distribution quite closely except for one patient with HAV angle 50 degrees, a high outlier.



(a) Find the 95% confidence interval for the population mean based on the 37 patients who remain after you drop the outlier.
(b) Compare your interval in (a) with your interval from Exercise 18.32. What is the most important effect of removing the outlier?

18.35 Men of few words? Researchers claim that women speak significantly more words per day than men. One estimate is that a woman uses about 20,000 words per day while a man uses about 7,000. To investigate such claims, one study used a special device to record the conversations of male and female university students over a four-day period. From these recordings, the daily word count of the 20 men in the study was determined. Here are their daily word counts:²⁰



28,408	10,084	15,931	21,688	37,786
10,575	12,880	11,071	17,799	13,182
8,918	6,495	8,153	7,015	4,429
10,054	3,998	12,639	10,974	5,255

(a) Examine the data. Is it reasonable to use the *t* procedures (assume these men are an SRS of all male students at this university)?

(b) If your conclusion in part (a) is “Yes,” do the data give convincing evidence that the mean number of words per day of men at this university differs from 7,000?

18.36 Genetic engineering for cancer treatment. Here’s a new idea for treating advanced melanoma, the most serious kind of skin cancer. Genetically engineer white blood cells to

better recognize and destroy cancer cells, then infuse these cells into patients. The subjects in a small initial study were 11 patients whose melanoma had not responded to existing treatments. One question was how rapidly the new cells would multiply after infusion, as measured by the doubling time in days. Here are the doubling times.²¹ 

1.4 1.0 1.3 1.0 1.3 2.0 0.6 0.8 0.7 0.9 1.9

- (a) Examine the data. Is it reasonable to use the t procedures?
- (b) Give a 90% confidence interval for the mean doubling time. Are you willing to use this interval to make an inference about the mean doubling time in a population of similar patients?

18.37 Genetic engineering for cancer treatment, continued.

Another outcome in the cancer experiment described in Exercise 18.36 is measured by a test for the presence of cells that trigger an immune response in the body and so may help fight cancer. Here are data for the 11 subjects: counts of active cells per 100,000 cells before and after infusion of the modified cells. The difference (after minus before) is the response variable. 

Before	14	0	1	0	0	0	0	20	1	6	0
After	41	7	1	215	20	700	13	530	35	92	108
Difference	27	7	0	215	20	700	13	510	34	86	108

- (a) Examine the data. Is it reasonable to use the t procedures?
- (b) If your conclusion in part (a) is “Yes,” do the data give convincing evidence that the count of active cells is higher after treatment?

18.38 Kicking a helium-filled football. Does a football filled with helium travel farther than one filled with ordinary air? To test this, the *Columbus Dispatch* conducted a study. Two identical footballs, one filled with helium and one filled with ordinary air, were used. A casual observer was unable to detect a difference in the two footballs. A novice kicker was used to punt the footballs. A trial consisted of kicking both footballs in a random order. The kicker did not know which football (the helium-filled or the air-filled football) he was kicking. The distance of each punt was recorded. Then another trial was conducted. A total of 39 trials were run. Here are the data for the 39 trials, in yards that the footballs traveled. The difference (helium minus air) is the response variable.²² 

Helium	25	16	25	14	23	29	25	26	22	26
Air	25	23	18	16	35	15	26	24	24	28
Difference	0	-7	7	-2	-12	14	-1	2	-2	-2
Helium	12	28	28	31	22	29	23	26	35	24
Air	25	19	27	25	34	26	20	22	33	29
Difference	-13	9	1	6	-12	3	3	4	2	-5
Helium	31	34	39	32	14	28	30	27	33	11
Air	31	27	22	29	28	29	22	31	25	20
Difference	0	7	17	3	-14	-1	8	-4	8	-9
Helium	26	32	30	29	30	29	29	30	26	
Air	27	26	28	32	28	25	31	28	28	
Difference	-1	6	2	-3	2	4	-2	2	-2	

- (a) Examine the data. Is it reasonable to use the t procedures?
- (b) If your conclusion in part (a) is “Yes,” do the data give convincing evidence that the helium-filled football travels farther than the air-filled football?

18.39 Growing trees faster. The concentration of carbon dioxide (CO_2) in the atmosphere is increasing rapidly due to our use of fossil fuels. Because plants use CO_2 to fuel photosynthesis, more CO_2 may cause trees and other plants to grow faster. An elaborate apparatus allows researchers to pipe extra CO_2 to a 30-meter circle of forest. They selected two nearby circles in each of three parts of a pine forest and randomly chose one of each pair to receive extra CO_2 . The response variable is the mean increase in base area for 30 to 40 trees in a circle during a growing season. We measure this in percent increase per year. Here are one year’s data.²³ 

fastest growing tree in each pair of plots. The data are as follows:

Pair	Control plot	Treated plot
1	9.752	10.587
2	7.263	9.244
3	5.742	8.675

- (a) State the null and alternative hypotheses. Explain clearly why the investigators used a one-sided alternative.
- (b) Carry out a test and report your conclusion in simple language.
- (c) The investigators used the test you just carried out. Any use of the t procedures with samples this size is risky. Why?

18.40 Fungus in the air. The air in poultry-processing plants often contains fungus spores. Inadequate ventilation

can affect the health of the workers. The problem is most serious during the summer. To measure the presence of spores, air samples are pumped to an agar plate and “colony-forming units (CFUs)” are counted after an incubation period. Here are data from two locations in a plant that processes 37,000 turkeys per day, taken on four days in the summer. The units are CFUs per cubic meter of air.²⁴



FUNGUS

	Day 1	Day 2	Day 3	Day 4
Kill room	3175	2526	1763	1090
Processing	529	141	362	224

- (a) Explain carefully why these are matched pairs data.
- (b) The spore count is clearly higher in the kill room. Give sample means and a 90% confidence interval to estimate how much higher. Be sure to state your conclusion in plain English.
- (c) You will often see the *t* procedures used for data like these. You should regard the results as only rough approximations. Why?

18.41 Weeds among the corn.

Velvetleaf is a particularly annoying weed in corn fields. It produces lots of seeds, and the seeds wait in the soil for years until conditions are right. How many seeds do velvetleaf plants produce? Here are counts from 28 plants that came up in a corn field when no herbicide was used.²⁵



WEEDS



Cuboimages art/Alamy

2450	2504	2114	1110	2137	8015	1623	1531	2008	1716
721	863	1136	2819	1911	2101	1051	218	1711	164
2228	363	5973	1050	1961	1809	130	880		

We would like to give a confidence interval for the mean number of seeds produced by velvetleaf plants. Alas, the *t* interval can't be safely used for these data. Why not?

18.42 Sweetening colas. Cola makers test new recipes for loss of sweetness during storage. Trained tasters rate the sweetness before and after storage. Here are the sweetness losses (sweetness before storage minus sweetness after storage) found by 10 tasters for one new cola recipe:



COLA

2.0	0.4	0.7	2.0	-0.4	2.2	-1.3	1.2	1.1	2.3
-----	-----	-----	-----	------	-----	------	-----	-----	-----

Take the data from these 10 carefully trained tasters as an SRS from a large population of all trained tasters.

- (a) Use these data to see if there is good evidence that the cola lost sweetness.

- (b) It is not uncommon to see the *t* procedures used for data like these. However, you should regard the results as only rough approximations. Why?

18.43 How much oil? How much oil wells in a given field will ultimately produce is key information in deciding whether to drill more wells. Here are the estimated total amounts of oil recovered from 64 wells in the Devonian Richmond Dolomite area of the Michigan basin, in thousands of barrels:²⁶



21.7	53.2	46.4	42.7	50.4	97.7	103.1	51.9
43.4	69.5	156.5	34.6	37.9	12.9	2.5	31.4
79.5	26.9	18.5	14.7	32.9	196.0	24.9	118.2
82.2	35.1	47.6	54.2	63.1	69.8	57.4	65.6
56.4	49.4	44.9	34.6	92.2	37.0	58.8	21.3
36.6	64.9	14.8	17.6	29.1	61.4	38.6	32.5
12.0	28.3	204.9	44.5	10.3	37.7	33.7	81.1
12.1	20.1	30.5	7.1	10.1	18.0	3.0	2.0

Take these wells to be an SRS of wells in this area.

- (a) Give a 95% *t* confidence interval for the mean amount of oil recovered from all wells in this area.
- (b) Make a graph of the data. The distribution is very skewed, with several high outliers. A computer-intensive method that gives accurate confidence intervals without assuming any specific shape for the distribution gives a 95% confidence interval of 40.28 to 60.32. How does the *t* interval compare with this? Should the *t* procedures be used with these data?

18.44 *E. coli* in swimming areas. To investigate water quality, the Columbus Dispatch took water samples at 16 Ohio State Park swimming areas in central Ohio. Those samples were taken to laboratories and tested for *E. coli*, which are bacteria that can cause serious gastrointestinal problems. If a 100-milliliter sample (about 3.3 ounces) of water contains more than 130 *E. coli* bacteria, it is considered unsafe. Here are the *E. coli* levels found by the laboratories:²⁷



ECOLI

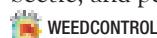
291.0	190.4	47.0	86.0	44.0	18.9	1.0	50.0
10.9	45.7	28.5	8.6	9.6	16.0	34.0	18.9

Take these water samples to be an SRS of the water in all swimming areas in central Ohio.

- (a) Are these data good evidence that, on average, the *E. coli* levels in these swimming areas were unsafe?
- (b) Make a graph of the data. The distribution is very skewed. Another method that gives *P*-values without assuming any specific shape for the distribution gives a *P*-value of 0.9997 for the question in part (a). How does the one-sample *t* test compare with this? Should the *t* procedures be used with these data?

The following exercises ask you to answer questions from data without having the details outlined for you. The four-step process is illustrated in Examples 18.2, 18.3, and 18.4. The exercise statements give you the **State** step. Follow the **Plan**, **Solve**, and **Conclude** steps in your work.

18.45 Natural weed control? Fortunately, we aren't really interested in the number of seeds velvetleaf plants produce (see Exercise 18.41). The velvetleaf seed beetle feeds on the seeds and might be a natural weed control. Here are the total seeds, seeds infected by the beetle, and percent of seeds infected for 28 velvetleaf plants:



Seeds	2450	2504	2114	1110	2137	8015	1623	1531	2008	1716
Infected	135	101	76	24	121	189	31	44	73	12
Percent	5.5	4.0	3.6	2.2	5.7	2.4	1.9	2.9	3.6	0.7
Seeds	721	863	1136	2819	1911	2101	1051	218	1711	164
Infected	27	40	41	79	82	85	42	0	64	7
Percent	3.7	4.6	3.6	2.8	4.3	4.0	4.0	0.0	3.7	4.3
Seeds	2228	363	5973	1050	1961	1809	130	880		
Infected	156	31	240	91	137	92	5	23		
Percent	7.0	8.5	4.0	8.7	7.0	5.1	3.8	2.6		

Do a complete analysis of the percent of seeds infected by the beetle. Include a 90% confidence interval for the mean percent infected in the population of all velvetleaf plants. Do you think that the beetle is very helpful in controlling the weed?

18.46 Recruiting T cells. There is evidence that cytotoxic T lymphocytes (T cells) participate in controlling tumor growth and that they can be harnessed to use the body's immune system to treat cancer. One study investigated the use of a T cell-engaging antibody, blinatumomab, to recruit T cells to control tumor growth. The data below are T cell counts (1000 per microliter) at baseline (beginning of the study) and after 20 days on blinatumomab for 6 subjects in the study.²⁸ The difference (after 20 days minus baseline) is the response variable.



Baseline	0.04	0.02	0.00	0.02	0.38	0.33
After 20 days	0.28	0.47	1.30	0.25	1.22	0.44
Difference	0.24	0.45	1.30	0.23	0.84	0.11

Do the data give convincing evidence that the mean count of T cells is higher after 20 days on blinatumomab?

18.47 Recruiting T cells, continued. Give a 95% confidence interval for the mean difference in T cell counts (after 20 days minus baseline) in the previous exercise.



18.48 Mutual funds performance. Mutual funds often compare their performance with a benchmark provided by an "index" that describes the performance of the class of assets in which the fund invests. For example, the Vanguard International Growth Fund benchmarks its performance against the EAFE (Europe, Australasia, Far East) index. Table 18.4 gives annual returns (percent) for the fund and the index. Does the fund's performance differ significantly from that of its benchmark?



- Explain clearly why the matched pairs *t* test is the proper choice to answer this question.
- Do a complete analysis that answers the question posed.

TABLE 18.4 A mutual fund versus its benchmark index

YEAR	FUND RETURN (%)	INDEX RETURN (%)	YEAR	FUND RETURN (%)	INDEX RETURN (%)
1984	-1.02	7.38	1998	16.93	20.00
1985	56.94	56.16	1999	26.34	26.96
1986	56.71	69.44	2000	-8.60	-14.17
1987	12.48	24.63	2001	-18.92	-21.44
1988	11.61	28.27	2002	-17.79	-15.94
1989	24.76	10.54	2003	34.45	38.59
1990	-12.05	-23.45	2004	18.95	20.25
1991	4.74	12.13	2005	15.00	13.54
1992	-5.79	-12.17	2006	25.92	26.34
1993	44.74	32.56	2007	15.98	11.17
1994	0.76	7.78	2008	-44.94	-43.38
1995	14.89	11.21	2009	41.63	31.78
1996	14.65	6.05	2010	15.66	8.13
1997	4.12	1.78			

18.49 Right versus left. The design of controls and instruments affects how easily people can use them. Timothy Sturm investigated this effect in a course project, asking 25 right-handed students to turn a knob (with their right hands) that moved an indicator by screw action. There were two identical instruments, one with a right-hand thread (the knob turns clockwise) and the other with a left-hand thread (the knob turns counterclockwise). Table 18.5 gives the times in seconds each subject took to move the indicator a fixed distance.²⁹  

- (a) Each of the 25 students used both instruments. Explain briefly how you would use randomization in arranging the experiment.
- (b) The project hoped to show that right-handed people find right-hand threads easier to use. Do an analysis that leads to a conclusion about this issue.

TABLE 18.5 Performance times (seconds) using right-hand and left-hand threads

SUBJECT	RIGHT		LEFT		SUBJECT	RIGHT		LEFT	
	THREAD	THREAD	THREAD	THREAD		THREAD	THREAD	THREAD	THREAD
1	113	137	14	107	87				
2	105	105	15	118	166				
3	130	133	16	103	146				
4	101	108	17	111	123				
5	138	115	18	104	135				
6	118	170	19	111	112				
7	87	103	20	89	93				
8	116	145	21	78	76				
9	75	78	22	100	116				
10	96	107	23	89	78				
11	122	84	24	85	101				
12	103	148	25	88	123				
13	116	147							

18.50 Comparing two drugs. Makers of generic drugs must show that they do not differ significantly from the “reference” drugs that they imitate. One aspect in which drugs might differ is their extent of absorption in the blood. Table 18.6 gives data taken from 20 healthy nonsmoking male subjects for one pair of drugs.³⁰ This is a matched pairs design. Numbers 1 to 20 were assigned at random to the subjects. Subjects 1 to 10 received the generic drug first, followed by the reference drug. Subjects 11 to 20 received the reference drug first, followed by the generic drug. In all

TABLE 18.6 Absorption extent for two versions of a drug

SUBJECT	REFERENCE DRUG	GENERIC DRUG
15	4108	1755
3	2526	1138
9	2779	1613
13	3852	2254
12	1833	1310
8	2463	2120
18	2059	1851
20	1709	1878
17	1829	1682
2	2594	2613
4	2344	2738
16	1864	2302
6	1022	1284
10	2256	3052
5	938	1287
7	1339	1930
14	1262	1964
11	1438	2549
1	1735	3340
19	1020	3050

cases, a washout period separated the two drugs so that the first had disappeared from the blood before the subject took the second. By randomizing the order, we eliminate the order in which the drugs were administered from being confounded with the difference in the absorption in the blood. Do the drugs differ significantly in the amount absorbed in the blood?  

18.51 Practical significance? Give a 90% confidence interval for the mean time advantage of right-hand over left-hand threads in the setting of Exercise 18.49. Do you think that the time saved would be of practical importance if the task were performed many times—for example, by an assembly-line worker? To help answer this question, find the mean time for right-hand threads as a percent of the mean time for left-hand threads.  

18.52 Bad weather, bad tips? As part of the study of tipping in a restaurant that we met in Example 14.3 (page 359), the psychologists also studied the size of the tip in a restaurant when a message indicating that the next day’s weather would be bad was written on the bill.

Here are tips from 20 patrons, measured in percent of the total bill:³¹

18.0 19.1 19.2 18.8 18.4 19.0 18.5 16.1 16.8 14.0
17.0 13.6 17.5 20.0 20.2 18.8 18.0 23.2 18.2 19.4

Do the data give convincing evidence that the mean percentage tip for all patrons of this restaurant when their bill contains a message that the next day's weather will be bad is less than 20%? (20% is an often-recommended size for restaurant tips.)



EXPLORING THE WEB

18.53 A matched pairs study. Find an example of a matched pairs study on the Web. The *Journal of the American Medical Association* (jama.ama-assn.org), *Science Magazine* (www.sciencemag.org), the *Canadian Medical Association Journal* (www.cmaj.ca), the *Journal of Statistics Education* (www.amstat.org/publications/jse), or perhaps the *Journal of Quantitative Analysis in Sports* (www.bepress.com/jqas) are possible sources. To help locate an article, look through the abstracts of articles. Once you find a suitable article, read it, and then briefly describe the study (including why it is a matched pairs study) and its conclusions. If *P*-values, means, standard deviations, etc. are reported, be sure to include them in your summary. Also, be sure to give the reference (either the Web link or the journal, issue, year, title of the paper, authors, and page numbers).

18.54 An improper use of a *t* procedure. Search the Web for an example of an improper use of a *t* procedure. You might try using Google to do a search on “improper use of a *t*-test.” Summarize the study in which the *t* procedure was used and discuss how it was used improperly. Be sure to provide a link to your example or an appropriate reference.

18.55 How big a sample size do you need? If you examine Table C you will notice that critical values of the *t* distribution get closer and closer to the corresponding critical values of the Normal distribution as the number of degrees of freedom increase. You can see this by comparing the *z* critical values at the bottom of Table C with the *t* critical values in the corresponding column. This suggests that for very large sample sizes, inference based on the Normal probability calculations in Chapters 14 and 15 (pretending σ is known) and inference based on the *t* distribution as discussed in this chapter (σ is not known) may give essentially the same answer if sample sizes are large and we pretend that our estimate of σ is the true value of σ . Professor R. Webster West, Department of Statistics, Texas A&M University, has created an applet that allows one to compute *t* probabilities. The link to the applet is www.stat.tamu.edu/~west/applets/tdemo.html.

- Use this applet to determine how large a sample size (or how many degrees of freedom) is needed for the critical value of the *t* distribution to be within 0.01 of the corresponding critical value of the Normal distribution for a 90%, 95%, and 99% confidence interval for a population mean.
- Based on your findings, how large a sample size do you think is needed for inference using the Normal distribution and inference using the *t* distribution to give very similar results if σ (both the true value and its estimate) is 1? If σ (both the true value and its estimate) is 100?



Two-Sample Problems

Comparing two populations or two treatments is one of the most common situations encountered in statistical practice. We call such situations *two-sample problems*.

TWO-SAMPLE PROBLEMS

- The goal of inference is to compare the responses to two treatments or to compare the characteristics of two populations.
- We have a separate sample from each treatment or each population.

TWO-SAMPLE PROBLEMS

A two-sample problem can arise from a randomized comparative experiment that randomly divides subjects into two groups and exposes each group to a different treatment. Comparing random samples separately selected from two populations is also a two-sample problem. Unlike the matched pairs designs studied earlier, there is no matching of the individuals in the two samples, and the two samples can be of different sizes. Inference procedures for two-sample data differ from those for matched pairs. Here are some typical two-sample problems:

EXAMPLE 19.1 Two-sample problems

- Does regular physical therapy help lower-back pain? A randomized experiment assigned patients with lower-back pain to two groups: 142 received an examination and advice from a physical therapist; another 144 received regular physical therapy for up to five weeks. After a year, the change in their level of disability (0% to 100%) was assessed by a doctor who did not know which treatment the patients had received.

Chapter 19

IN THIS CHAPTER WE COVER...

- Two-sample problems
- Comparing two population means
- Two-sample *t* procedures
- Using technology
- Robustness again
- Details of the *t* approximation*
- Avoid the pooled two-sample *t* procedures*
- Avoid inference about standard deviations*

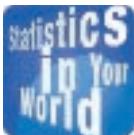
- A psychologist develops a test that measures social insight. He compares the social insight of female college students with that of male college students by giving the test to a sample of female students and a separate sample of male students.
- A bank wants to know which of two incentive plans will most increase the use of its credit cards. It offers each incentive to a random sample of credit card customers and compares the amounts charged during the following six months. ■

APPLY YOUR KNOWLEDGE

Which data design? Each situation described in Exercises 19.1 to 19.4 requires inference about a mean or means. Identify each as involving (1) a single sample, (2) matched pairs, or (3) two independent samples. The procedures of Chapter 18 apply to designs (1) and (2). We are about to learn procedures for (3).



Betsie Van Der Meer/Getty



Driving while fasting

Muslims fast from sunrise to sunset during the month of Ramadan. Does this affect the rate of traffic accidents? Fasting can improve alertness, reducing accidents. Or it can cause dehydration, increasing accidents. Data from Turkey show a statistically significant increase, starting two weeks into Ramadan. Ah, but because Ramadan follows a lunar calendar, it cycles through the year. Perhaps accidents go down during a winter Ramadan (alertness) but go up during a summer Ramadan (longer fast and dehydration). Ask the statisticians this question and get their favorite answer: we need more data.

19.1 Managing the finances. Choose 50 engaged couples who have not been previously married. Interview the man and woman separately about how their joint finances will be handled after marriage. Compare the views of men and women.

19.2 Does peer discussion promote learning? Undergraduate students in a biology class were randomly divided into two groups. In the first group, students worked alone on a multiple-choice exam on material recently covered in the class. In the second group, students worked in teams of four on the same multiple-choice exam. The teams were encouraged to discuss the questions among themselves before answering questions. All exams were graded and returned to the students. Then all students worked alone on another multiple-choice exam on the same material as the first exam. Compare the mean scores of the two groups on the second exam.

19.3 Chemical analysis. To check a new analytical method, a chemist obtains a reference specimen of known concentration from the National Institute of Standards and Technology. She then makes 20 measurements of the concentration of this specimen with the new method and checks for bias by comparing the mean result with the known concentration.

19.4 Chemical analysis, continued. Another chemist is checking the same new method. He has no reference specimen, but a familiar analytic method is available. He wants to know if the new and old methods agree. He takes a specimen of unknown concentration and measures the concentration 10 times with the new method and 10 times with the old method.

COMPARING TWO POPULATION MEANS

Comparing two populations or the responses to two treatments starts with data analysis: make boxplots, stemplots (for small samples), or histograms (for larger samples) and compare the shapes, centers, and spreads of the two samples. The most common goal of inference is to compare the average or typical responses in the two populations. When data analysis suggests that both population distributions are symmetric, and especially when they are at least approximately Normal, we want to compare the population means. Here are the conditions for inference about means.

CONDITIONS FOR INFERENCE COMPARING TWO MEANS

- We have two SRSs, from two distinct populations. The samples are **independent**. That is, one sample has no influence on the other. Matching violates independence, for example. We measure the same response variable for both samples.
- Both populations are **Normally distributed**. The means and standard deviations of the populations are unknown. In practice, it is enough that the distributions have similar shapes and that the data have no strong outliers.

Call the variable we measure x_1 in the first population and x_2 in the second because the variable may have different distributions in the two populations. Here is how we describe the two populations:

Population	Variable	Mean	Standard Deviation
1	x_1	μ_1	σ_1
2	x_2	μ_2	σ_2

There are four unknown parameters, the two means and the two standard deviations. The subscripts remind us which population a parameter describes. We want to compare the two population means, either by giving a confidence interval for their difference $\mu_1 - \mu_2$ or by testing the hypothesis of no difference, $H_0: \mu_1 = \mu_2$.

We use the sample means and standard deviations to estimate the unknown parameters. Again, subscripts remind us which sample a statistic comes from. Here is how we describe the samples:

Population	Sample Size	Sample Mean	Sample Standard Deviation
1	n_1	\bar{x}_1	s_1
2	n_2	\bar{x}_2	s_2

To do inference about the difference $\mu_1 - \mu_2$ between the means of the two populations, we start from the difference $\bar{x}_1 - \bar{x}_2$ between the means of the two samples.

EXAMPLE 19.2 Daily activity and obesity

STATE: People gain weight when they take in more energy from food than they expend. James Levine and his collaborators at the Mayo Clinic investigated the link between obesity and energy spent on daily activity.¹

Choose 20 healthy volunteers who don't exercise. Deliberately choose 10 who are lean and 10 who are mildly obese but still healthy. Attach sensors that monitor the subjects' every move for 10 days. Table 19.1 presents data on the time (in minutes per day) that the subjects spent standing or walking, sitting, and lying down. Do lean and obese people differ in the average time they spend standing and walking?



ACTIVITY



AP Photo/Toby Talbot

TABLE 19.1 Time (minutes per day) spent in three different postures by lean and obese subjects

GROUP	SUBJECT	STAND/WALK	SIT	LIE
Lean	1	511.100	370.300	555.500
Lean	2	607.925	374.512	450.650
Lean	3	319.212	582.138	537.362
Lean	4	584.644	357.144	489.269
Lean	5	578.869	348.994	514.081
Lean	6	543.388	385.312	506.500
Lean	7	677.188	268.188	467.700
Lean	8	555.656	322.219	567.006
Lean	9	374.831	537.031	531.431
Lean	10	504.700	528.838	396.962
Obese	11	260.244	646.281	521.044
Obese	12	464.756	456.644	514.931
Obese	13	367.138	578.662	563.300
Obese	14	413.667	463.333	532.208
Obese	15	347.375	567.556	504.931
Obese	16	416.531	567.556	448.856
Obese	17	358.650	621.262	460.550
Obese	18	267.344	646.181	509.981
Obese	19	410.631	572.769	448.706
Obese	20	426.356	591.369	412.919

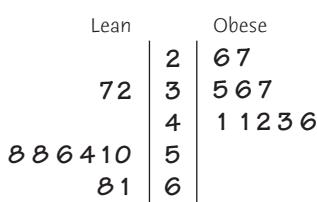
PLAN: Examine the data and carry out a test of hypotheses. We suspect in advance that lean subjects (Group 1) are more active than obese subjects (Group 2), so we test the hypotheses

$$H_0: \mu_1 = \mu_2$$

$$H_a: \mu_1 > \mu_2$$

SOLVE (first steps): Are the conditions for inference met? The subjects are volunteers, so they are not SRSs from all lean and mildly obese adults. The study tried to recruit comparable groups: all worked in sedentary jobs, none smoked or were taking medication, and so on. Setting clear standards like these helps make up for the fact that we can't reasonably get SRSs for so invasive a study. The subjects were not told that they were chosen from a larger group of volunteers because they did not exercise and were either lean or mildly obese. Because their willingness to volunteer isn't related to the purpose of the experiment, we will treat them as two independent SRSs.

A back-to-back stemplot (Figure 19.1) displays the data in detail. To make the plot, we rounded the data to the nearest 10 minutes and used 100s as stems and 10s as leaves. The distributions are a bit irregular, as we expect with just 10 observations. There are no clear departures from Normality such as extreme outliers or skewness. The lean

**FIGURE 19.1**

Back-to-back stemplot of the times spent walking or standing, for Example 19.2.

subjects as a group spend much more time standing and walking than do the obese subjects. Calculating the group means confirms this:

Group	<i>n</i>	Mean \bar{x}	Std. dev. <i>s</i>
Group 1 (lean)	10	525.751	107.121
Group 2 (obese)	10	373.269	67.498

The observed difference in mean time per day spent standing or walking is

$$\bar{x}_1 - \bar{x}_2 = 525.751 - 373.269 = 152.482 \text{ minutes}$$

To complete the “Solve” step, we must learn the details of inference for comparing two means. ■

TWO-SAMPLE *t* PROCEDURES

To assess the significance of the observed difference between the means of our two samples, we follow a familiar path. Whether an observed difference is surprising depends on the spread of the observations as well as on the two means. Widely different means can arise just by chance if the individual observations vary a great deal. To take variation into account, we would like to standardize the observed difference $\bar{x}_1 - \bar{x}_2$ by dividing by its standard deviation. This standard deviation of the difference in sample means is

$$\sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}}$$

This standard deviation gets larger as either population gets more variable, that is, as σ_1 or σ_2 increases. It gets smaller as the sample sizes n_1 and n_2 increase.

Because we don’t know the population standard deviations, we estimate them by the sample standard deviations from our two samples. The result is the **standard error**, or estimated standard deviation, of the difference in sample means:

$$\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

When we standardize the estimate by dividing it by its standard error, the result is the **two-sample *t* statistic**:

$$t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

The statistic *t* has the same interpretation as any *z* or *t* statistic: it says how far $\bar{x}_1 - \bar{x}_2$ is from 0 in standard deviation units.

The two-sample *t* statistic has approximately a *t* distribution. It does not have exactly a *t* distribution even if the populations are both exactly Normal. In practice,

standard error

two-sample *t* statistic

however, the approximation is very accurate. There are two practical options for using the two-sample t procedures:

Option 1. With software, use the statistic t with accurate critical values from the approximating t distribution. The degrees of freedom are calculated from the data by a somewhat messy formula. Moreover, the degrees of freedom may not be a whole number.

Option 2. Without software, use the statistic t with critical values from the t distribution with *degrees of freedom equal to the smaller of $n_1 - 1$ and $n_2 - 1$* . These procedures are always conservative for any two Normal populations. The confidence interval has a margin of error as large as or larger than is needed for the desired confidence level. The significance test gives a P -value *equal to or greater than* the true P -value.

The two options are exactly the same except for the degrees of freedom used for t critical values and P -values. As the sample sizes increase, confidence levels and P -values from Option 2 become more accurate. The gap between what Option 2 reports and the truth is quite small unless the sample sizes are both small and unequal.²

THE TWO-SAMPLE t PROCEDURES

Draw an SRS of size n_1 from a large Normal population with unknown mean μ_1 , and draw an independent SRS of size n_2 from another large Normal population with unknown mean μ_2 . A level C confidence interval for $\mu_1 - \mu_2$ is given by

$$(\bar{x}_1 - \bar{x}_2) \pm t^* \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

Here t^* is the critical value for confidence level C for the t distribution with degrees of freedom from either Option 1 (software) or Option 2 (the smaller of $n_1 - 1$ and $n_2 - 1$).

To test the hypothesis $H_0: \mu_1 = \mu_2$, calculate the two-sample t statistic

$$t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

Find P -values from the t distribution with degrees of freedom from either Option 1 (software) or Option 2 (the smaller of $n_1 - 1$ and $n_2 - 1$).



EXAMPLE 19.3 Daily activity and obesity

We can now complete Example 19.2.

SOLVE (inference): The two-sample t statistic comparing the average minutes spent standing and walking in Group 1 (lean) and Group 2 (obese) is

$$\begin{aligned}
 t &= \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}} \\
 &= \frac{525.751 - 373.269}{\sqrt{\frac{107.121^2}{10} + \frac{67.498^2}{10}}} \\
 &= \frac{152.482}{40.039} = 3.808
 \end{aligned}$$

Software (Option 1) gives one-sided P -value $P = 0.0008$ based on $df = 15.174$.

Without software, use the conservative Option 2. Because $n_1 - 1 = 9$ and $n_2 - 1 = 9$, there are 9 degrees of freedom. Because H_a is one-sided, the P -value is the area to the right of $t = 3.808$ under the $t(9)$ curve. Figure 19.2 illustrates this P -value. Table C shows that $t = 3.808$ lies between the critical values t^* for 0.0025 and 0.001. So $0.001 < P < 0.0025$. Option 2 gives a larger (more conservative) P -value than Option 1. As usual, the practical conclusion is the same for both versions of the test.

CONCLUDE: There is very strong evidence ($P = 0.0008$) that, on average, lean people spend more time walking and standing than do moderately obese people. ■

df = 9		
t^*	3.690	4.297
One-sided P	.0025	.001

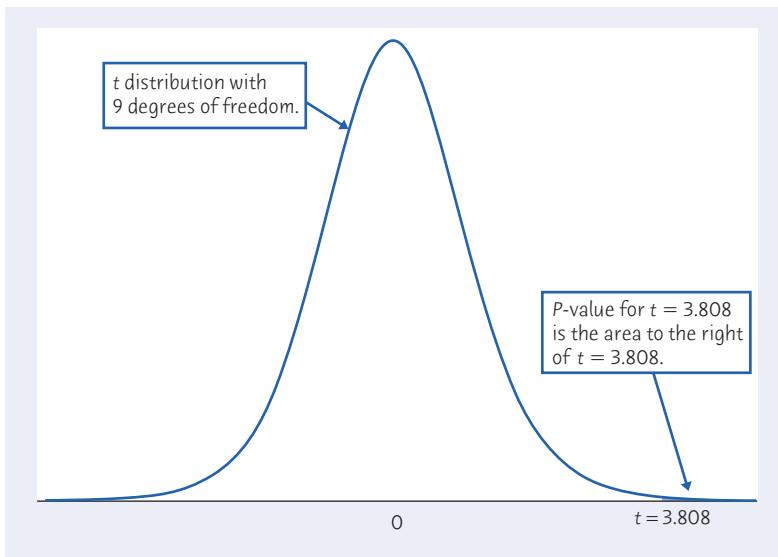


FIGURE 19.2

Using the conservative Option 2, the P -value in Example 19.3 comes from the t distribution with 9 degrees of freedom.

Does lack of daily activity *cause* obesity? This is an observational study, and that affects our ability to draw cause-and-effect conclusions. It may be that some people are naturally more active and are therefore less likely to gain weight. Or it may be that people who gain weight reduce their activity level. The study went on to enroll most of the obese subjects in a weight-reduction program and most of the lean subjects in a supervised program of overeating. After 8 weeks, the obese subjects had lost weight (mean 8 kilograms) and the lean subjects had gained weight (mean 4 kg). But both groups kept their original allocation of time

to the different postures. This suggests that time allocation may be biological and influences weight, rather than the other way around. The authors remark: “It should be emphasized that this was a pilot study and that the results need to be confirmed in larger studies.”



EXAMPLE 19.4 How much more active are lean people?

PLAN: Give a 90% confidence interval for $\mu_1 - \mu_2$, the difference in average daily minutes spent standing and walking between lean and mildly obese adults.

SOLVE AND CONCLUDE: As in Example 19.3, the conservative Option 2 uses 9 degrees of freedom. Table C shows that the $t(9)$ critical value is $t^* = 1.833$. We are 90% confident that $\mu_1 - \mu_2$ lies in the interval

$$\begin{aligned} (\bar{x}_1 - \bar{x}_2) &\pm t^* \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} \\ &= (525.751 - 373.269) \pm 1.833 \sqrt{\frac{107.121^2}{10} + \frac{67.498^2}{10}} \\ &= 152.482 \pm 73.390 \\ &= 79.09 \text{ to } 225.87 \text{ minutes} \end{aligned}$$

Software using Option 1 gives the 90% interval as 82.35 to 222.62 minutes, based on t with 15.174 degrees of freedom. The Option 2 interval is wider because this method is conservative. Both intervals are quite wide because the samples are small and the variation among individuals, as measured by the two sample standard deviations, is large. Whichever interval we report, we are (at least) 90% confident that the mean difference in average daily minutes spent standing and walking between lean and mildly obese adults lies in this interval. ■



EXAMPLE 19.5 Community service and attachment to friends

STATE: Do college students who have volunteered for community service work differ from those who have not? A study obtained data from 57 students who had done service work and 17 who had not. One of the response variables was a measure of attachment to friends (roughly, secure relationships), measured by the Inventory of Parent and Peer Attachment. In particular, the response is a score based on the responses to 25 questions. Here are the results:³

Group	Condition	n	\bar{x}	s
1	Service	57	105.32	14.68
2	No service	17	96.82	14.26

PLAN: The investigator had no specific direction for the difference in mind before looking at the data, so the alternative is two-sided. We will test the hypotheses

$$H_0: \mu_1 = \mu_2$$

$$H_a: \mu_1 \neq \mu_2$$

SOLVE: The investigator says that the individual scores, examined separately in the two samples, appear roughly Normal. There is a serious problem with the more important condition that the two samples can be regarded as SRSs from two student populations. We will discuss that after we illustrate the calculations.

The two-sample *t* statistic is

$$\begin{aligned} t &= \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}} \\ &= \frac{105.32 - 96.82}{\sqrt{\frac{14.68^2}{57} + \frac{14.26^2}{17}}} \\ &= \frac{8.5}{3.9677} = 2.142 \end{aligned}$$

Software (Option 1) says that the two-sided *P*-value is $P = 0.0414$.

Without software, use Option 2 to find a conservative *P*-value. There are 16 degrees of freedom, the smaller of

$$n_1 - 1 = 57 - 1 = 56 \quad \text{and} \quad n_2 - 1 = 17 - 1 = 16$$

Figure 19.3 illustrates the *P*-value. Find it by comparing $t = 2.142$ with the two-sided critical values for the $t(16)$ distribution. Table C shows that the *P*-value is between 0.05 and 0.04.

CONCLUDE: The data give moderately strong evidence ($P < 0.05$) that students who have engaged in community service are, on the average, more attached to their friends. ■

df = 16		
t^*	2.120	2.235
Two-sided <i>P</i>	.05	.04

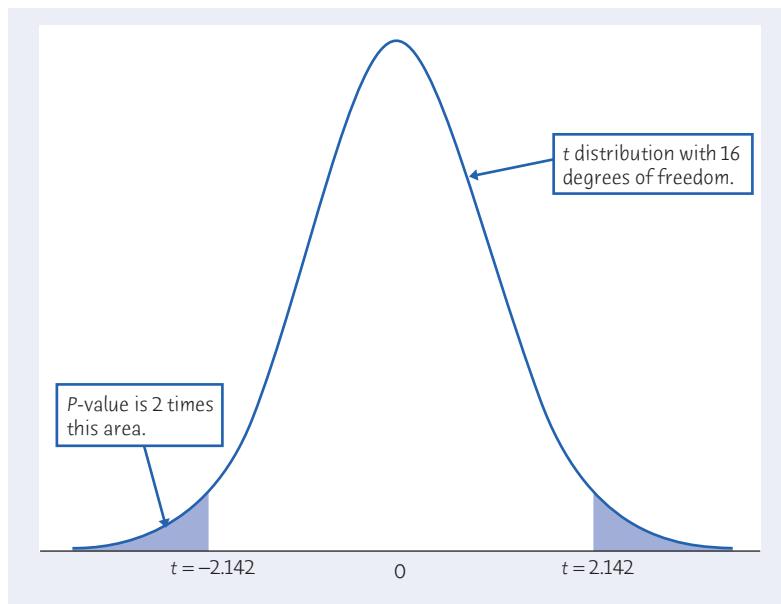


FIGURE 19.3

The *P*-value in Example 19.5. Because the alternative is two-sided, the *P*-value is double the area to the left of $t = -2.142$.

Is the *t* test in Example 19.5 justified? The student subjects were “enrolled in a course on U.S. Diversity at a large mid-western university.” Unless this course is

required of all students, the subjects cannot be considered a random sample even from this campus. Students were placed in the two groups on the basis of a questionnaire, 39 in the “no service” group and 71 in the “service” group. The data were gathered from a follow-up survey two years later; 17 of the 39 “no service” students responded (44%), compared with 80% response (57 of 71) in the “service” group. Nonresponse is confounded with group: students who had done community service were much more likely to respond. Finally, 75% of the “service” respondents were women, compared with 47% of the “no service” respondents. Sex, which can strongly affect attachment, is badly confounded with the presence or absence of community service. The data are so far from meeting the SRS condition for inference that the t test is meaningless. Difficulties like these are common in social science research, where confounding variables have stronger effects than is usual when biological or physical variables are measured. This researcher honestly disclosed the weaknesses in data production but left it to readers to decide whether to trust her inferences.

APPLY YOUR KNOWLEDGE

In exercises that call for two-sample t procedures, use Option 1 if you have technology that implements that method. Otherwise, use Option 2 (degrees of freedom the smaller of $n_1 - 1$ and $n_2 - 1$).



Digital Vision/Getty Images

19.5 Logging in the rain forest. “Conservationists have despaired over destruction of tropical rain forest by logging, clearing, and burning.” These words begin a report on a statistical study of the effects of logging in Borneo.⁴ Here are data on the number of tree species in 12 unlogged forest plots and 9 similar plots logged 8 years earlier:



Unlogged	22	18	22	20	15	21	13	13	19	13	19	15
Logged	17	4	18	14	18	15	15	10	12			

- (a) The study report says, “Loggers were unaware that the effects of logging would be assessed.” Why is this important? The study report also explains why the plots can be considered to be randomly assigned.
- (b) Does logging significantly reduce the mean number of species in a plot after 8 years? Follow the four-step process as illustrated in Examples 19.2 and 19.3.

19.6 Daily activity and obesity. We can conclude from Examples 19.2 and 19.3 that mildly obese people spend less time standing and walking (on the average) than lean people. Is there a significant difference between the mean times the two groups spend lying down? Use the four-step process to answer this question from the data in Table 19.1. Follow the model of Examples 19.2 and 19.3.

19.7 Logging in the rain forest, continued. Use the data in Exercise 19.5 to give a 99% confidence interval for the difference in mean number of species between unlogged and logged plots.



USING TECHNOLOGY

Software should use Option 1 for the degrees of freedom to give accurate confidence intervals and P -values. Unfortunately, there is variation in how well software implements Option 1. Figure 19.4 displays output from a graphing

Texas Instruments Graphing Calculator

```
2-SampTTest
μ1 > μ2
t = 3.808375604
P = 8.4101904e-4
df = 15.17355038
x̄1 = 525.7513
x̄2 = 373.2692
```

Minitab

Session

Two-sample T for stand

group	N	Mean	StDev	SE Mean
1	10	526	107	34
2	10	373.3	67.5	21

Difference = mu (1) - mu (2)
Estimate for difference: 152.5
T-Test of difference = 0 (vs >): T-Value = 3.81 P-Value = 0.001 DF = 15

Microsoft Excel

Excel

A	B	C
1	t-Test: Two-Sample Assuming Unequal Variances	
2		
3	Lean	Obese
4	Mean	525.7513
5	Variance	11474.8903
6	Observations	10
7	Hypothesized Mean	0
8	df	15
9	t Stat	3.808375604
10	P(T<=t) one-tail	0.000856818
11	t Critical one-tail	1.753051038
12	P(T<=t) two-tail	0.001713635
13	P(T>t) two-tail	2.131450856
14		

CrunchIt!

Results - t-2-Sample

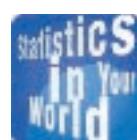
Null hypothesis: Difference of means = 0
Alternative hypothesis: Difference of means > 0

Group	Std Dev	Sample Mean	n
1	107.1	525.8	10
2	67.50	373.3	10

df: 15.17
Difference of means: 152.5
T statistic: 3.808
P-value: 0.0008410

FIGURE 19.4

The two-sample t procedures applied to the data on activity and obesity: output from a graphing calculator, two statistical programs, and a spreadsheet program.



Meta-analysis

Small samples have large margins of error. Large

samples are expensive. Often we can find several studies of the same issue; if we could combine their results, we would have a large sample with a small margin of error. That is the idea of "meta-analysis." Of course, we can't just lump the studies together, because of differences in design and quality. Statisticians have more sophisticated ways of combining the results. Meta-analysis has been applied to issues ranging from the effect of secondhand smoke to whether coaching improves SAT scores.

calculator, two statistical programs, and a spreadsheet program for the test of Example 19.3 (page 470). All four claim to use Option 1. The two-sample t statistic is exactly as in Example 19.3, $t = 3.808$. You can find this in all four outputs (Minitab rounds to 3.81; Excel and the graphing calculator give additional decimal places). The different technologies use different methods to find the P -value for $t = 3.808$, however.

- CrunchIt! and the calculator get Option 1 completely right. The accurate approximation uses the t distribution with approximately 15.174 (CrunchIt! rounds this to 15.17) degrees of freedom. The P -value is $P = 0.0008$.
- Minitab uses Option 1, but it *truncates* the exact degrees of freedom to the next smaller whole number to get critical values and P -values. In this example, the exact $df = 15.174$ is truncated to $df = 15$, so that Minitab's results are slightly conservative. That is, Minitab's P -value (rounded to $P = 0.001$ in the output) is slightly larger than the full Option 1 P -value.
- Excel *rounds* the exact degrees of freedom to the nearest whole number, so that $df = 15.174$ becomes $df = 15$. Excel's method agrees with Minitab's in this example. But when rounding moves the degrees of freedom up to the next higher whole number, Excel's P -values are slightly smaller than is correct. This is misleading, another illustration of the fact that Excel is substandard as statistical software.



Excel's label for the test, "Two-Sample Assuming Unequal Variances," is seriously misleading. *The two-sample t procedures we have described work whether or not the two populations have the same variance.* There is an old-fashioned special procedure that works only when the two variances are equal. We discuss this method in an optional section on page 481, but you should never use it.

Although different calculators and software give slightly different P -values, in practice you can just accept what your technology says. The small differences in P don't affect the conclusion. Even "between 0.001 and 0.0025" from Option 2 (Example 19.3) is close enough for practical purposes.



APPLY YOUR KNOWLEDGE

- 19.8 Perception of life expectancy.** Do women and men differ in how they perceive their life expectancy? A researcher asked a sample of men and women to indicate their life expectancy. This was compared with values from actuarial tables, and the relative percent difference was computed (perceived life expectancy minus life expectancy from actuarial tables was divided by life expectancy from actuarial tables and converted to a percent). Here are the relative percent differences for all men and women over the age of 70 in the sample:⁵ LIFEEXPECTANCY

Men	-28	-23	-20	-19	-14	-13
Women	-20	-19	-15	-12	-10	-8

Figure 19.5 shows output for the two-sample t test using Option 1. (This output is from CrunchIt! software, which does Option 1 without rounding or truncating the degrees of freedom.) Do men and women over 70 years old differ in their perceptions of life expectancy? Using the output in Figure 19.5, write a summary in a sentence or two, including t , df, P , and a conclusion.

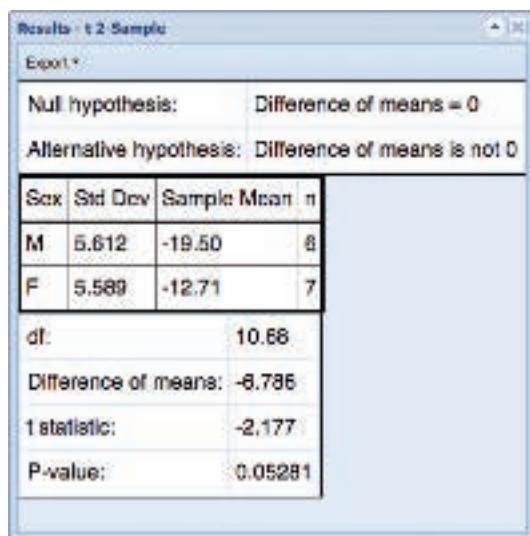


FIGURE 19.5

Two-sample t output from CrunchIt! for Exercise 19.8.

ROBUSTNESS AGAIN

The two-sample t procedures are more robust than the one-sample t methods, particularly when the distributions are not symmetric. When the sizes of the two samples are equal and the two populations being compared have distributions with similar shapes, probability values from the t table are quite accurate for a broad range of distributions when the sample sizes are as small as $n_1 = n_2 = 5$.⁶ When the two population distributions have different shapes, larger samples are needed.

As a guide to practice, adapt the guidelines given on page 452 for the use of one-sample t procedures to two-sample procedures by replacing “sample size” with the “sum of the sample sizes,” $n_1 + n_2$. These guidelines err on the side of safety, especially when the two samples are of equal size. *In planning a two-sample study, choose equal sample sizes whenever possible. The two-sample t procedures are most robust against non-Normality in this case, and the conservative Option 2 probability values are most accurate.*





APPLY YOUR KNOWLEDGE

19.9 Do good smells bring good business? Businesses know that customers often respond to background music. Do they also respond to odors? One study of this question took place in a small pizza restaurant in France on two Saturday evenings in May. On one of these evenings, a relaxing lavender odor was spread through the restaurant. On the other evening, no scent was used. Table 19.2 gives the

TABLE 19.2 Time (minutes) and spending (euros) by restaurant customers

NO ODOR		LAVENDER	
MINUTES	EUROS SPENT	MINUTES	EUROS SPENT
103	15.9	92	21.9
68	18.5	126	18.5
79	15.9	114	22.3
106	18.5	106	21.9
72	18.5	89	18.5
121	21.9	137	24.9
92	15.9	93	18.5
84	15.9	76	22.5
72	15.9	98	21.5
92	15.9	108	21.9
85	15.9	124	21.5
69	18.5	105	18.5
73	18.5	129	25.5
87	18.5	103	18.5
109	20.5	107	18.5
115	18.5	109	21.9
91	18.5	94	18.5
84	15.9	105	18.5
76	15.9	102	24.9
96	15.9	108	21.9
107	18.5	95	25.9
98	18.5	121	21.9
92	15.9	109	18.5
107	18.5	104	18.5
93	15.9	116	22.8
118	18.5	88	18.5
87	15.9	109	21.9
101	25.5	97	20.7
75	12.9	101	21.9
86	15.9	106	22.5

time (minutes) that two samples of 30 customers spent in the restaurant and the amount they spent (in euros).⁷ The two evenings were comparable in many ways (weather, customer count, and so on), so we are willing to regard the data as independent SRSs from spring Saturday evenings at this restaurant. The authors say, “Therefore at this stage it would be impossible to generalize the results to other restaurants.”  ODORS2

- (a) Does a lavender odor encourage customers to stay longer in the restaurant? Examine the time data and explain why they are suitable for two-sample t procedures. Use the two-sample t test to answer the question posed.
- (b) Does a lavender odor encourage customers to spend more while in the restaurant? Examine the spending data. In what ways do these data deviate from Normality? With 30 observations, the t procedures are nonetheless reasonably accurate. Use the two-sample t test to answer the question posed.

19.10 Compressing soil. Farmers know that driving heavy equipment on wet soil compresses the soil and injures future crops. Here are data on the “penetrability” of the same type of soil at two levels of compression.⁸ Penetrability is a measure of how much resistance plant roots will meet when they try to grow through the soil.  SOILCOMPRESS

Compressed Soil									
2.86	2.68	2.92	2.82	2.76	2.81	2.78	3.08	2.94	2.86
3.08	2.82	2.78	2.98	3.00	2.78	2.96	2.90	3.18	3.16
Intermediate Soil									
3.14	3.38	3.10	3.40	3.38	3.14	3.18	3.26	2.96	3.02
3.54	3.36	3.18	3.12	3.86	2.92	3.46	3.44	3.62	4.26

- (a) Make stemplots to investigate the shape of the distributions. The penetrabilities for intermediate soil are skewed to the right and have a high outlier. Returning to the source of the data shows that the outlying sample had unusually low soil density, so that it belongs in the “loose soil” class. We are justified in removing the outlier.
- (b) We suspect that the penetrability of compressed soil is less than that of intermediate soil. Do the data (with the outlier removed) support this suspicion?

19.11 Weeds among the corn. Lamb’s-quarter is a common weed that interferes with the growth of corn. An agriculture researcher planted corn at the same rate in 16 small plots of ground, then weeded the plots by hand to allow a fixed number of lamb’s-quarter plants to grow in each meter of corn row. No other weeds were allowed to grow. Here are the yields of corn (bushels per acre) for only the experimental plots controlled to have 1 weed per meter of row and 9 weeds per meter of row:⁹  WEEDSANDCORN

1 weed/meter	166.2	157.3	166.7	161.1
9 weeds/meter	162.8	142.4	162.8	162.4

Explain carefully why a two-sample t confidence interval for the difference in mean yields may not be accurate.

19.12 Compressing soil, continued. Use the data in Exercise 19.10, omitting the outlier, to give a 90% confidence interval for the decrease in penetrability of compressed soil relative to intermediate soil.  SOILCOMPRESS2

DETAILS OF THE t APPROXIMATION*

The exact distribution of the two-sample t statistic is not a t distribution. Moreover, the distribution changes as the unknown population standard deviations σ_1 and σ_2 change. However, an excellent approximation is available. We call this Option 1 for t procedures.

APPROXIMATE DISTRIBUTION OF THE TWO-SAMPLE t STATISTIC

The distribution of the two-sample t statistic is very close to the t distribution with degrees of freedom df given by

$$df = \frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\frac{1}{n_1 - 1} \left(\frac{s_1^2}{n_1}\right)^2 + \frac{1}{n_2 - 1} \left(\frac{s_2^2}{n_2}\right)^2}$$

This approximation is accurate when both sample sizes n_1 and n_2 are 5 or larger.



ACTIVITY

EXAMPLE 19.6 Daily activity and obesity

In the experiment of Examples 19.2 and 19.3, the data on minutes per day spent standing and walking give

Group	n	\bar{x}	s
Group 1 (lean)	10	525.751	107.121
Group 2 (obese)	10	373.269	67.498

The two-sample t test statistic calculated from these values is $t = 3.808$.

The one-sided P -value is the area to the right of 3.808 under a t density curve, as in Figure 19.2 (page 471). The conservative Option 2 uses the t distribution with 9 degrees of freedom. Option 1 finds a very accurate P -value by using the t distribution with degrees of freedom df given by

$$\begin{aligned} df &= \frac{\left(\frac{107.121^2}{10} + \frac{67.498^2}{10}\right)^2}{\frac{1}{9}\left(\frac{107.121^2}{10}\right)^2 + \frac{1}{9}\left(\frac{67.498^2}{10}\right)^2} \\ &= \frac{2,569,894}{169,367.2} = 15.1735 \end{aligned}$$

These degrees of freedom appear in the graphing calculator output in Figure 19.4. Because the formula is messy and roundoff errors are likely, we don't recommend calculating df by hand. ■

*This section can be omitted unless you are using software and wish to understand what the software does.

The degrees of freedom df is generally not a whole number. It is always at least as large as the smaller of $n_1 - 1$ and $n_2 - 1$. The larger degrees of freedom that result from Option 1 give slightly shorter confidence intervals and slightly smaller P -values than the conservative Option 2 produces. There is a t distribution for any positive degrees of freedom, even though Table C contains entries only for whole-number degrees of freedom.

The difference between the t procedures using Options 1 and 2 is rarely of practical importance. That is why we recommend the simpler, conservative Option 2 for inference without software. With software, the more accurate Option 1 procedures are painless.

APPLY YOUR KNOWLEDGE

19.13 Students' self-concept. A study of the self-concept of seventh-grade students asked if male and female students differ in mean score on the Piers-Harris Children's Self-Concept Scale.¹⁰ Software that uses Option 1 gives these summary results:

Gender	n	Mean	Std dev	Std err	t	df	P
F	31	55.5161	12.6961	2.2803	-0.8276	62.8	0.4110
M	47	57.9149	12.2649	1.7890			

Starting from the sample means and standard deviations, verify each of these entries: the standard errors of the means; the degrees of freedom for two-sample t ; the value of t .

19.14 Perception of life expectancy. Figure 19.5 (page 477) gives output for the life expectancy data in Exercise 19.8 from software that does Option 1 with the correct degrees of freedom. What are \bar{x}_i and s_i for the two samples? Starting from these values, find the t test statistic and its degrees of freedom. Your work should agree with Figure 19.5.

19.15 Students' self-concept, continued. Write a sentence or two summarizing the comparison of female and male students in Exercise 19.13, as if you were preparing a report for publication. Use the output in Exercise 19.13.

AVOID THE POOLED TWO-SAMPLE t PROCEDURES*

Most software and graphing calculators, including all four illustrated in Figure 19.4, offer a choice of two-sample t statistics. One is often labeled for “unequal” variances; the other for “equal” variances. The “unequal” variance procedure is our two-sample t . This test is valid whether or not the population variances are equal. The other choice is a special version of the two-sample t statistic that assumes that the two populations have the same variance. This procedure averages (the statistical term is “pools”) the two sample variances to estimate the common population variance. The resulting statistic is called the *pooled two-sample t statistic*. It is equal to our t statistic if the two sample sizes are the same, but not otherwise. We could choose to use the pooled t for tests and confidence intervals.

*This short section offers advice on what not to do. This material is not needed to read the rest of the book.

The pooled t statistic has exactly the t distribution with $n_1 + n_2 - 2$ degrees of freedom if the two population variances really are equal and the population distributions are exactly Normal. The pooled t was in common use before software made it easy to use Option 1 for our two-sample t statistic. Of course, in the real world distributions are not exactly Normal and population variances are not exactly equal. In practice, the Option 1 two-sample t procedures are almost always more accurate than the pooled procedures. Our advice: *Never use the pooled t procedures if you have software that will implement Option 1.*



AVOID INFERENCE ABOUT STANDARD DEVIATIONS*

Two basic features of a distribution are its center and spread. In a Normal population, we measure center by the mean and spread by the standard deviation. We use the t procedures for inference about population means for Normal populations, and we know that t procedures are widely useful for non-Normal populations as well. It is natural to turn next to inference about the standard deviations of Normal populations. Our advice here is short and clear: don't do it without expert advice.

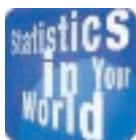
There are methods for inference about the standard deviations of Normal populations. The most common such method is the “F test” for comparing the standard deviations of two Normal populations. You will find this test in the menus of most statistical software. *Unlike the t procedures for means, the F test for standard deviations is extremely sensitive to non-Normal distributions.* This lack of robustness does not improve in large samples. It is difficult in practice to tell whether a significant test result is evidence of unequal population spreads or simply a sign that the populations are not Normal. Because this test is of little use in practice, we don't give its details.

The deeper difficulty underlying the very poor robustness of Normal population procedures for inference about spread already appeared in our work on describing data. The standard deviation is a natural measure of spread for Normal distributions but not for distributions in general. In fact, because skewed distributions have unequally spread tails, no single numerical measure does a good job of describing the spread of a skewed distribution. In summary, the standard deviation is not always a useful parameter, and even when it is (for symmetric distributions), the results of inference about the standard deviation are not trustworthy. Consequently, *we do not recommend trying to do inference about population standard deviations in basic statistical practice.*¹¹

CHAPTER 19 SUMMARY

CHAPTER SPECIFICS

- The data in a **two-sample problem** are two independent SRSs, each drawn from a separate population.
- Tests and confidence intervals for the difference between the means μ_1 and μ_2 of two Normal populations start from the difference $\bar{x}_1 - \bar{x}_2$ between the two sample means. Because of the central limit theorem, the resulting procedures are approximately correct for other population distributions when the sample sizes are large.



Is money the root of all evil?

That may go too far, but Kathleen

Vohs and her coworkers show that even thinking about money has strong effects. Exercise 19.46 describes a small part of their work. What does Professor Vohs say about the consequences of having money? “Money makes people feel self-sufficient and behave accordingly.” With money, you can achieve your goals with less help from others. You feel less dependent on others and more willing to work toward your own goals. Maybe that's good. You also prefer to be less involved with others, so that self-sufficiency is a barrier to close relationships with others. Maybe that's not good. Scientists don't tell us what's good or not good, just that money increases our sense of self-sufficiency.



- Draw independent SRSs of sizes n_1 and n_2 from two Normal populations with parameters μ_1 , σ_1 , and μ_2 , σ_2 . The two-sample **t statistic** is

$$t = \frac{(\bar{x}_1 - \bar{x}_2) - (\mu_1 - \mu_2)}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

The statistic t has approximately a t distribution.

- There are two choices for the **degrees of freedom** of the two-sample t statistic. Option 1: software produces accurate probability values using degrees of freedom calculated from the data. Option 2: for conservative inference procedures, use degrees of freedom equal to the smaller of $n_1 - 1$ and $n_2 - 1$.
- The confidence interval for $\mu_1 - \mu_2$ is

$$(\bar{x}_1 - \bar{x}_2) \pm t^* \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$$

The critical value t^* from Option 1 gives a confidence level very close to the desired level C . Option 2 produces a margin of error at least as wide as is needed for the desired level C .

- Significance tests for $H_0: \mu_1 = \mu_2$ are based on

$$t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}}$$

P -values calculated from Option 1 are very accurate. Option 2 P -values are always at least as large as the true P .

- The two-sample t procedures are quite **robust** against departures from Normality. Guidelines for practical use are similar to those for one-sample t procedures. Equal sample sizes are recommended.
- Procedures for inference about the standard deviations of Normal populations are very sensitive to departures from Normality. Avoid inference about standard deviations unless you have expert advice.

LINK IT

In Chapter 18 we studied inference for the mean of a Normal population using procedures based on the t distribution. These procedures are more realistic than those studied in Chapters 14 to 16 because t procedures do not require that we know the population variance. In practice, the most common use of t procedures for a single population mean is with matched pairs data because most research studies make comparisons between two or more populations.

In this chapter, we discuss t procedures for comparing the means of two Normal populations when we have independent samples from these two populations. These t procedures are quite common in practice, and one encounters them in many research papers. Researchers typically report the means and standard deviations for the two samples and, in the case of tests of hypotheses, the value of the t statistic and the corresponding P -value.

In the next two chapters we extend our procedures for inference to additional settings that occur frequently in practice—in particular, to inference for population proportions. And we will see that the basic ideas about confidence intervals and tests of hypotheses that we learned in Chapters 14 to 16 still apply.

CHECK YOUR SKILLS

19.16 The 2009 National Assessment of Educational Progress (NAEP) gave a mathematics test to a random sample of eighth-graders in Texas. The mean score was 287 out of 500. To give a confidence interval for the mean score of all Texas eighth-graders, you would use

- (a) the two-sample t interval.
- (b) the matched pairs t interval.
- (c) the one-sample t interval.

19.17 In the 2009 NAEP sample of Texas eighth-graders, the mean mathematics scores were 286 for female students and 287 for male students. To see if this difference is statistically significant, you would use

- (a) the two-sample t test.
- (b) the matched pairs t test.
- (c) the one-sample t test.

19.18 A new method for measuring blood pressure in laboratory mice using radiotelemetry has been proposed. How does this compare with the existing method? You apply both methods to a sample of 16 mice and use

- (a) the two-sample t test.
- (b) the matched pairs t test.
- (c) the one-sample t test.

19.19 One major reason that the two-sample t procedures are widely used is that they are quite *robust*. This means that

- (a) t procedures do not require that we know the standard deviations of the populations.
- (b) confidence levels and P -values from the t procedures are quite accurate even if the population distribution is not exactly Normal.
- (c) t procedures compare population means, a comparison that answers many practical questions.

19.20 A study of the effect of exposure to color (red or blue) on the ability to solve puzzles used 42 subjects. Half the subjects (21) were asked to solve a series of puzzles while in a red-colored environment. The other half were asked to solve the same series of puzzles while in a blue-colored environment. The time taken to solve the puzzles was recorded for each subject. To compare the mean times for the two groups of subjects using the two-sample t proce-

dures with the conservative Option 2, the correct degrees of freedom is

- (a) 41. (b) 40. (c) 20.

19.21 The 21 subjects in the red-colored environment had a mean time for solving the puzzles of 9.64 seconds with standard deviation 3.43; the 21 subjects in the blue-colored environment had a mean time of 15.84 seconds with standard deviation 8.65. The two-sample t statistic for comparing the population means has value

- (a) 1.50. (b) 3.05. (c) 6.2.

19.22 A study of the use of social media asked a sample of 488 American adults under the age of 40 and a sample of 421 American adults aged 40 or over about their use of social media. Based on their answers, each subject was assigned a social media usage score on a scale of 0 to 25. Higher scores indicate greater usage. The subjects were chosen by random digit dialing of telephone numbers. Are the conditions for two-sample t inference satisfied?

- (a) Maybe: the SRS condition is OK but we need to look at the data to check Normality.
- (b) No: scores in a range between 0 and 25 can't be Normal.
- (c) Yes: the SRS condition is OK and large sample sizes make the Normality condition unnecessary.

19.23 We suspect that younger adults use social media more than adults aged 40 or over. To see if this is true, test these hypotheses for the mean social media usage scores of all adults under 40 and all adults 40 and over:

- (a) $H_0: \mu_{<40} = \mu_{\geq 40}$ versus $H_a: \mu_{<40} > \mu_{\geq 40}$
- (b) $H_0: \mu_{<40} = \mu_{\geq 40}$ versus $H_a: \mu_{<40} \neq \mu_{\geq 40}$
- (c) $H_0: \mu_{<40} = \mu_{\geq 40}$ versus $H_a: \mu_{<40} < \mu_{\geq 40}$

19.24 The two-sample t statistic for the social media use study (“under 40” mean minus “40 and over” mean) is $t = 3.18$. The P -value for testing the hypotheses from the previous exercise satisfies

- (a) $0.001 < P < 0.005$.
- (b) $0.0005 < P < 0.001$.
- (c) $0.001 < P < 0.002$.

CHAPTER 19 EXERCISES

Exercises 19.25 to 19.34 are based on summary statistics rather than raw data. This information is typically all that is presented in published reports. You can perform inference procedures by hand from the summaries. Use the conservative Option 2 (degrees of freedom the smaller of $n_1 - 1$ and $n_2 - 1$) for two-sample t confidence intervals and P -values. You must trust that the authors understood the conditions for inference and verified that they apply. This isn't always true.

19.25 Do women talk more than men? Equip male and female students with a small device that secretly records sound for a random 30 seconds during each 12.5-minute period over two days. Count the words each subject speaks during each recording period, and from this, estimate how many words per day each subject speaks. The published report includes a table summarizing six such studies.¹² Here are two of the six:

Study	Sample Size		Estimated Average Number (SD) of Words Spoken per Day	
	Women	Men	Women	Men
1	56	56	16,177 (7520)	16,569 (9108)
2	27	20	16,496 (7914)	12,867 (8343)

Readers are supposed to understand that, for example, the 56 women in the first study had $\bar{x} = 16,177$ and $s = 7520$. It is commonly thought that women talk more than men. Does either of the two samples support this idea? For each study:

- (a) State hypotheses in terms of the population means for men (μ_M) and women (μ_F).
- (b) Find the two-sample t statistic.
- (c) What degrees of freedom does Option 2 use to get a conservative P -value?
- (d) Compare your value of t with the critical values in Table C. What can you say about the P -value of the test?
- (e) What do you conclude from the results of these two studies?

19.26 Alcohol and zoning out. Healthy men aged 21 to 35 were randomly assigned to one of two groups: half received 0.82 grams of alcohol per kilogram of body weight; half received a placebo. Participants were then given 30 minutes to read up to 34 pages of Tolstoy's *War and Peace* (beginning at chapter 1, with each page containing approximately 22 lines of text). Every two to four minutes participants were prompted to indicate whether they were "zoning out." The proportion of times participants indicated they were zoning out was recorded for each subject. The table below summarizes data on the proportion of episodes of zoning out.¹³

(The study report gave the standard error of the mean s/\sqrt{n} , abbreviated as SEM, rather than the standard deviation s .)

Group	<i>n</i>	\bar{x}	SEM
Alcohol	25	0.25	0.05
Placebo	25	0.12	0.03

- (a) What are the two sample standard deviations?
- (b) What degrees of freedom does the conservative Option 2 use for two-sample t procedures for these samples?
- (c) Using Option 2, give a 90% confidence interval for the mean difference between the two groups.

19.27 Stress and weight in rats

In a study of the effects of stress on behavior in rats, 71 rats were randomly assigned to either a stressful environment or a control (nonstressful) environment. After 21 days, the change in weight (in grams) was determined for each rat. The table below summarizes data on weight gain.¹⁴ (The study report gave the standard error of the mean s/\sqrt{n} , abbreviated as SEM, rather than the standard deviation s .)



R.L. Brinster/Photolibrary

Group	<i>n</i>	\bar{x}	SEM
Stress	20	26	3
No stress	51	32	2

- (a) What are the standard deviations for the two groups?
- (b) What degrees of freedom does the conservative Option 2 use for two-sample t procedures for these data?
- (c) Test the null hypothesis of no difference between the two group means against the two-sided alternative. Use the degrees of freedom from (b).

19.28 Is Montessori preschool beneficial? Do education programs for preschool children that follow the Montessori method perform better than other programs? A study compared 5-year-old children in Milwaukee, Wisconsin, who had been enrolled in preschool programs from the age of 3.¹⁵

- (a) Explain why comparing children whose parents chose a Montessori school with children of other parents would not show whether Montessori schools perform better than other programs. (In fact, all the children in the study applied to the

Montessori school. The school district assigned students to Montessori or other preschools by a random lottery.)

(b) In all, 54 children were assigned to the Montessori school and 112 to other schools at age 3. When the children were 5, parents of 30 of the Montessori children and 25 of the others could be located. Those parents who were located agreed to and subsequently participated in testing. This information reveals a possible source of bias in the comparison of outcomes. Explain why.

(c) One of the many response variables was score on a test of ability to apply basic mathematics to solve problems. Here are summaries for the children who took this test:

Group	<i>n</i>	\bar{x}	<i>s</i>
Montessori	30	19	3.11
Control	25	17	4.19

Is there evidence of a difference in the population mean scores? (The researchers used two-sided alternative hypotheses.)

19.29 Ginkgo extract and the post-lunch dip.

The post-lunch dip is the drop in mental alertness after a midday meal. Does an extract of the leaves of the ginkgo tree reduce the post-lunch dip? Assign healthy people aged 18 to 40 to take either ginkgo extract or a placebo pill. After lunch, ask them to read seven pages of random letters and place an X over every e. Count the number of misses per line read.¹⁶



Emilio Ereza/Alamy

- (a) What is a placebo and why was one group given a placebo?
- (b) What is the double-blind method and why should it be used in this experiment?
- (c) Here are summaries of performance after 13 weeks of either ginkgo extract or placebo:

Group	Group size	Mean	Std. dev.
Ginkgo	21	0.06383	0.01462
Placebo	18	0.05342	0.01549

Is there a significant difference between the two groups? What do these data show about the effect of ginkgo extract?

19.30 Illusory pattern perceptions. When one experiences a lack of control in one's life, does one compensate by seeking structure elsewhere? Assign 36 undergraduate

students to one of two conditions. All are asked to identify a concept associated with a series of "grainy" pictures that contain an embedded image and are presented to them sequentially. During the task they can ask questions to help them determine the associated concept. Half of the subjects (the lack-of-control group) receive feedback that is random and noncontingent on their questions. The other half receive useful feedback (the in-control group). After attempting to complete the task, all subjects are presented with 12 grainy pictures that are similar to those used in the task but lack embedded images. All are asked whether they perceive an image in the pictures, and the number of pictures identified as containing an image is counted for each subject. Here are the summary statistics:¹⁷

Group	Group size	Mean	Std. dev.
Lack-of-control	18	5.16	3.5
In-control	18	3.47	2.0

- (a) What degrees of freedom would you use in the conservative two-sample *t* procedures to compare the lack-of-control and in-control groups?
- (b) What is the two-sample *t* test statistic for comparing the mean number of pictures identified as having an image for the two groups?
- (c) Test the null hypothesis of no difference between the two population means against the two-sided alternative. Use your statistic from part (b) with degrees of freedom from part (a).

19.31 Asperger's syndrome and the ability to mentalize.

Asperger's syndrome is a form of autism. People with this syndrome have difficulty interacting and communicating socially. Some researchers believe that behind these difficulties is an inability to "mentalize," that is, to automatically attribute mental states to the self and others. To test this, researchers had 36 subjects (19 with Asperger's syndrome and 17 without Asperger's syndrome) watch a play in which an actor is known to have a false belief (mental state) and his actions are different than would be the case if he did not have this false belief. In particular, the eye movements of subjects were tracked to see if subjects anticipated the direction the actor would move based on his false belief rather than the direction he would move if he did not have the false belief. A score reflecting the bias toward looking in the correct direction (the direction the actor moved based on his false belief) was recorded for each subject. Higher scores indicate the subject looked in this "correct" direction more often. (The study report gave the standard error of the mean s/\sqrt{n} , abbreviated as SEM, rather than the standard deviation *s*.) Here are the summary statistics:¹⁸

Group	Group size	Mean	SEM
Asperger's	19	-0.001	0.15
Non-Asperger's	17	0.42	0.17

Is there evidence of a difference in mean scores between those with Asperger's syndrome and those without it?

19.32 Coaching and SAT scores. Coaching companies claim that their courses can raise the SAT scores of high school students. Of course, students who retake the SAT without paying for coaching generally raise their scores, too. A random sample of students who took the SAT twice found 427 who were coached and 2733 who were uncoached.¹⁹ Starting with their Verbal scores on the first and second tries, we have these summary statistics:

	Try 1			Try 2			Gain	
	n	\bar{x}	s	\bar{x}	s	\bar{x}	\bar{x}	s
Coached	427	500	92	529	97	29	59	
Uncoached	2733	506	101	527	101	21	52	

Let's first ask if students who are coached increased their scores significantly.

- (a) You could use the information on the Coached line to carry out either a two-sample *t* test comparing Try 1 with Try 2 for coached students or a matched pairs *t* test using Gain. Which is the correct test? Why?
- (b) Carry out the proper test. What do you conclude?
- (c) Give a 99% confidence interval for the mean gain of all students who are coached.

19.33 Coaching and SAT scores, continued. What we really want to know is whether coached students improve more than uncoached students, and whether any advantage is large enough to be worth paying for. Use the information in the previous exercise to answer these questions:

- (a) Is there good evidence that coached students gained more on the average than uncoached students?
- (b) How much more do coached students gain on the average? Give a 99% confidence interval.
- (c) Based on your work, what is your opinion: do you think coaching courses are worth paying for?

19.34 Coaching and SAT scores: critique. The data you used in the previous two problems came from a random sample of students who took the SAT twice. The response rate was 63%, which is pretty good for nongovernment surveys, so let's accept that the respondents do represent all students who took the exam twice. Nonetheless, we can't be

sure that coaching actually *caused* the coached students to gain more than the uncoached students. Explain briefly but clearly why this is so.

Exercises 19.35 to 19.42 include the actual data. To apply the two-sample *t* procedures, use Option 1 if you have technology that implements that method. Otherwise, use Option 2.

19.35 Improving your tips. Researchers gave 40 index cards to a waitress at an Italian restaurant in New Jersey. Before delivering the bill to each customer, the waitress randomly selected a card and wrote on the bill the same message that was printed on the index card. Twenty of the cards had the message "The weather is supposed to be really good tomorrow. I hope you enjoy the day!" Another 20 cards contained the message "The weather is supposed to be not so good tomorrow. I hope you enjoy the day anyway!" After the customers left, the waitress recorded the amount of the tip (percent of bill) before taxes. Here are the tips for those receiving the good-weather message:²⁰



20.8	18.7	19.9	20.6	21.9	23.4	22.8	24.9	22.2	20.3
24.9	22.3	27.0	20.5	22.2	24.0	21.2	22.1	22.0	22.7

The tips for the 20 customers who received the bad-weather message are

18.0	19.1	19.2	18.8	18.4	19.0	18.5	16.1	16.8	14.0
17.0	13.6	17.5	20.0	20.2	18.8	18.0	23.2	18.2	19.4

- (a) Make stemplots or histograms of both sets of data. Because the distributions are reasonably symmetric with no extreme outliers, the *t* procedures will work well.
- (b) Is there good evidence that the two different messages produce different percent tips? State hypotheses, carry out a two-sample *t* test, and report your conclusions.

19.36 Do good smells bring good business? In Exercise 19.9 (page 478) you examined the effects of a lavender odor on customer behavior in a small restaurant. Lavender is a relaxing odor. The researchers also looked at the effects of lemon, a stimulating odor. The design of the study is described in Exercise 19.9. Here are the times in minutes that customers spent in the restaurant when no odor was present:



103	68	79	106	72	121	92	84	72	92
85	69	73	87	109	115	91	84	76	96
107	98	92	107	93	118	87	101	75	86

When a lemon odor was present, customers lingered for these times:

78	104	74	75	112	88	105	97	101	89
88	73	94	63	83	108	91	88	83	106
108	60	96	94	56	90	113	97		

- (a) Examine both samples. Does it appear that use of two-sample *t* procedures is justified? Do the sample means suggest that a lemon odor changes the average length of stay?
 (b) Does a lemon odor influence the length of time customers stay in the restaurant? State hypotheses, carry out a *t* test, and report your conclusions.

19.37 Improving your tips, continued. Use the data in Exercise 19.35 to give a 95% confidence interval for the difference between the mean percent tips for the two different messages.  TIPPING4

19.38 How strong are durable press fabrics? “Durable press” cotton fabrics are treated to improve their recovery from wrinkles after washing. Unfortunately, the treatment also reduces the strength of the fabric. A study compared the breaking strength of fabrics treated by two commercial durable press processes. Five swatches of the same fabric were assigned at random to each process. Here are the data, in pounds of pull needed to tear the fabric:²¹  FABRICS

Permafresh	29.9	30.7	30.0	29.5	27.6
Hylite	28.8	23.9	27.0	22.1	24.2

Is there good evidence that the two processes result in different mean breaking strengths?

- (a) Do the sample means suggest that one of the processes is superior in breaking strength?
 (b) Make stemplots for both samples. The Permafresh sample contains a mild outlier. With just 5 observations per group, we worry that this outlier will affect our conclusions.
 (c) Test the hypothesis $H_0: \mu_1 = \mu_2$ against the two-sided alternative twice: once using all the data and again without the outlier in the Permafresh sample. Do the two tests lead to similar conclusions? Can we safely conclude that one treatment has significantly higher mean breaking strength than the other?

19.39 Reducing wrinkles. Of course, the reason for durable press treatment is to reduce wrinkling. “Wrinkle recovery angle” measures how well a fabric recovers from wrinkles. Higher is better. Here are data on the wrinkle recovery angle (in degrees) for the same fabric swatches discussed in the previous exercise:  WRINKLES

Permafresh	136	135	132	137	134
Hylite	143	141	146	141	145

Is there a significant difference in wrinkle resistance?

- (a) Do the sample means suggest that one process has better wrinkle resistance?

- (b) Make stemplots for both samples. There are no obvious deviations from Normality.

- (c) Test the hypothesis $H_0: \mu_1 = \mu_2$ against the two-sided alternative. What do you conclude from part (a) and from the result of your test?

19.40 How much stronger? Continue your work from Exercise 19.38. A fabric manufacturer wants to know how large an advantage in strength fabrics treated by the Permafresh method have over fabrics treated by the Hylite process. Give a 90% confidence interval for the difference in mean breaking strengths. (Use all 5 fabric swatches.)  FABRICS

19.41 How much less wrinkling? In Exercise 19.39, you found that the Hylite process results in significantly greater wrinkle resistance than the Permafresh process. How large is the difference in mean wrinkle recovery angle? Give a 90% confidence interval.  WRINKLES

19.42 Do women talk more than men? Another study. Exercise 19.25 described a series of six studies investigating the number of words women and men speak per day. Exercise 19.25 gives results from two of these studies. Here are the results from another of these studies. The estimated numbers of words spoken per day for 27 women are  TALKING2

15,357	13,618	9,783	26,451	12,151	8,391	19,763
25,246	8,427	6,998	24,876	6,272	10,047	15,569
39,681	23,079	24,814	19,287	10,351	8,866	10,827
12,584	12,764	19,086	26,852	17,639	16,616	

The estimated numbers of words spoken per day for 20 men are

28,408	10,084	15,931	21,688	37,786	10,575	12,880
11,071	17,799	13,182	8,918	6,495	8,153	7,015
4,429	10,054	3,998	12,639	10,974	5,255	

Does this study provide good evidence that women talk more than men, on average?

- (a) Make stemplots for both samples. Are there any obvious deviations from Normality? In spite of these deviations from Normality, it is safe to use the *t* procedures. Explain.
 (b) Test the hypothesis $H_0: \mu_1 = \mu_2$ against the one-sided alternative that the mean number of words per day for women (μ_1) is greater than the mean number of words per day for men (μ_2). What do you conclude?

Do birds learn to time their breeding? Blue titmice eat caterpillars. The birds would like lots of caterpillars around when they have young to feed, but they breed earlier than peak caterpillar season. Do the birds time when they breed based on the previous year's caterpillar supply? Researchers randomly assigned

7 pairs of birds to have the natural caterpillar supply supplemented while feeding their young and another 6 pairs to serve as a control group relying on natural food supply. The next year, they measured how many days after the caterpillar peak the birds produced their nestlings.²² Exercises 19.43 to 19.45 are based on this experiment.



Hugh Clark/Frank Lane Picture Agency/CORBIS

19.43 Did the randomization produce similar groups?

The first thing to do is to compare the two groups in the first year. The only difference should be the chance effect of the random assignment. The study report says: "In the experimental year, the degree of synchronization did not differ between food-supplemented and control females." For this comparison, the report gives $t = -1.05$. What type of t statistic (paired or two-sample) is this? What are the degrees of freedom for this statistic? Show that this t leads to the quoted conclusion.

19.44 Did the treatment have an effect? The investigators expected the control group to adjust their breeding date the next year, whereas the well-fed, supplemented group had no reason to change. The report continues: "but in the following year food-supplemented females were more out of synchrony with the caterpillar peak than the controls." Here are the data (days behind the caterpillar peak):



Control	4.6	2.3	7.7	6.0	4.6	-1.2
Supplemented	15.5	11.3	5.4	16.5	11.3	11.4

Carry out a t test and show that it leads to the quoted conclusion.

19.45 Year-to-year comparison. Rather than comparing the two groups in each year, we could compare the behavior of each group in the first and second years. The study report says: "Our main prediction was that females receiving additional food in the nestling period should not change laying date the next year, whereas controls, which (in our area) breed too late in their first year, were expected to advance their laying date in the second year."

Comparing days behind the caterpillar peak in Years 1 and 2 gave $t = 0.63$ for the control group and $t = -2.63$ for the supplemented group. Are these paired or two-sample t statistics? What are the degrees of freedom for each t ? Show that these t -values do not agree with the prediction.

The remaining exercises ask you to answer questions from data without having the details outlined for you. The exercise statements give you the **State** step of the four-step process. Follow the **Plan**, **Solve**, and **Conclude** steps as illustrated in Examples 19.2

and 19.3 for tests and Example 19.4 for confidence intervals. Remember that examining the data and discussing the conditions for inference are part of the **Solve** step.

19.46 Thinking about money changes behavior.

Kathleen Vohs of the University of Minnesota and her coworkers carried out several randomized comparative experiments on the effects of thinking about money. Here's part of one such experiment.²³ Ask student subjects to unscramble 30 sets of five words to make a meaningful phrase from four of the five words. The control group unscrambled phrases like "cold it desk outside is" into "it is cold outside." The treatment group unscrambled phrases that lead to thinking about money, turning "high a salary desk paying" into "a high-paying salary." Then each subject worked a hard puzzle, knowing that he or she could ask for help. Here are the times in seconds until subjects asked for help. For the treatment group:



MONEYTHINK

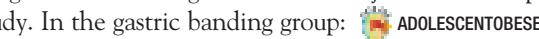
609	444	242	199	174	55	251	466	443
531	135	241	476	482	362	69	160	

For the control group:

118	272	413	291	140	104	55	189	126
400	92	64	88	142	141	373	156	

The researchers suspected that money is connected with self-sufficiency, so that the treatment group will ask for help less quickly on the average. Do the data support this idea?

19.47 Adolescent obesity. Adolescent obesity is a serious health risk affecting more than 5 million young people in the United States alone. Laparoscopic adjustable gastric banding has the potential to provide a safe and effective treatment. Fifty adolescents between 14 and 18 years old with a body mass index higher than 35 were recruited from the Melbourne, Australia, community for the study.²⁴ Twenty-five were randomly selected to undergo gastric banding, and the remaining twenty-five were assigned to a supervised lifestyle intervention program involving diet, exercise, and behavior modification. All subjects were followed for two years. Here are the weight losses in kilograms for the subjects who completed the study. In the gastric banding group:



35.6	81.4	57.6	32.8	31.0	37.6	36.5	-5.4
27.9	49.0	64.8	39.0	43.0	33.9	29.7	20.2
15.2	41.7	53.4	13.4	24.8	19.4	32.3	22.0

In the lifestyle intervention group

6.0	2.0	-3.0	20.6	11.6	15.5	-17.0	1.4	4.0
-4.6	15.8	34.6	6.0	-3.1	-4.3	-16.7	-1.8	-12.8

Is there good evidence that gastric banding is superior to the lifestyle intervention program?

19.48 Active versus traditional learning. Can active learning improve knowledge retention? Two undergraduate calculus-based engineering statistics courses were taught in different academic quarters, with one employing active-learning methods and another using traditional learning methods. The traditional class was taught lecture-style with relatively little in-class interaction between peers and with the instructor. The active-learning course integrated four group projects into the curriculum, with in-class time devoted to group work on the projects and fewer homework assignments. To assess knowledge retention, two five-question versions of a test were created. They had similar but not identical questions covering core statistics topics, worth a total of 18 possible points. All students in both sections were randomly given one version of the test as part of their final exam. Then, eight months later, a volunteer subset of the original students were given the version that they had not taken previously. To encourage students to take the second version of the exam, a ten-dollar gift card to the university bookstore was given to each participant. The change in the score from the first version to the second is used to measure a student's long-term ability to retain the course material. 

Here are the changes in exam scores for the 15 students in the Active group:²⁵

0	5	7	8	0	3	6	2	5	1
3	2	4	3	5					

The changes in exam scores for the 23 students in the Traditional group are

7	0	8	2	4	3	1	2	5	8
5	6	3	12	1	6	3	6	7	7
5	6	2							

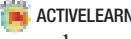
Is there good evidence that active learning is superior to traditional lecturing?

19.49 Each day I am getting better in math. A “subliminal” message is below our threshold of awareness but may nonetheless influence us. Can subliminal messages help students learn math? A group of students who had failed the mathematics part of the City University of New York Skills Assessment Test agreed to participate in a study to find out. 

All received a daily subliminal message, flashed on a screen too rapidly to be consciously read. The treatment group of 10 students (chosen at random) was exposed to “Each day I am getting better in math.” The control group of 8 students was exposed to a neutral message, “People are walking on the street.” All students participated in a summer program designed to raise their math skills, and all took the assessment test again at the end of the program. Table 19.3 gives data on the subjects’ scores

before and after the program.²⁶ Is there good evidence that the treatment brought about a greater improvement in math scores than the neutral message? How large is the mean difference in gains between treatment and control? (Use 90% confidence.)

19.50 Active versus traditional learning, continued.

- Use the data in Exercise 19.48 to give a 90% confidence interval for the difference in the mean change in score for students in the active and traditional classes. 
- Give a 90% confidence interval for the mean change in score of students in the active-learning class.

19.51 Tropical flowers.

Different varieties of the tropical flower *Heliconia* are fertilized by different species of hummingbirds. Over time, the lengths of the flowers and the forms of the hummingbirds' beaks have evolved to match each other. Here are data on the lengths in millimeters of two color varieties of the same species of flower on the island of Dominica:²⁷ 



Art Wolfe/Getty Images

H. caribaea red

41.90	42.01	41.93	43.09	41.47	41.69	39.78	40.57
39.63	42.18	40.66	37.87	39.16	37.40	38.20	38.07
38.10	37.97	38.79	38.23	38.87	37.78	38.01	

H. caribaea yellow

36.78	37.02	36.52	36.11	36.03	35.45	38.13	37.10
35.17	36.82	36.66	35.68	36.03	34.57	34.63	

Is there good evidence that the mean lengths of the two varieties differ? Estimate the difference between the population means. (Use 95% confidence.)

19.52 Student drinking.

A professor asked her sophomore students, “How many drinks do you typically have per session? (A drink is defined as one 12 oz beer, one 4 oz glass of wine, or one 1 oz shot of liquor.)” Some of the students didn’t drink. Table 19.4 gives the responses of the female and male students who did drink.²⁸ It is likely that some of the students exaggerated a bit. The sample is all students in one large sophomore-level class. The class is popular, so we are tentatively willing to regard its members as an SRS of sophomore students at this college. Do a complete analysis that reports on 

- the drinking behavior claimed by sophomore women.
- the drinking behavior claimed by sophomore men.
- a comparison of the behavior of women and men.

TABLE 19.3 Mathematics skills scores before and after a subliminal message

TREATMENT GROUP		CONTROL GROUP	
BEFORE	AFTER	BEFORE	AFTER
18	24	18	29
18	25	24	29
21	33	20	24
18	29	18	26
18	33	24	38
20	36	22	27
23	34	15	22
23	36	19	31
21	34		
17	27		

TABLE 19.4 Drinks per session claimed by female and male students

FEMALE STUDENTS													
2.5	9	1	3.5	2.5	3	1	3	3	3	3	2.5	2.5	
5	3.5	5	1	2	1	7	3	7	4	4	6.5	4	
3	6	5	3	8	6	6	3	6	8	3	4	7	
4	5	3.5	4	2	1	5	5	3	3	6	4	2	
7	7	7	3.5	3	2.5	10	5	4	9	8	1	6	
2	5	2.5	3	4.5	9	5	4	4	3	4	6	7	
4	5	1	5	3	4	10	7	3	4	4	4	4	
2	1	2.5	2.5										
MALE STUDENTS													
7	7.5	8	15	3	4	1	5	11	4.5	6	4	10	
16	4	8	5	9	7	7	3	5	6.5	1	12	4	
6	8	8	4.5	10.5	8	6	10	1	9	8	7	8	
15	3	10	7	4	6	5	2	10	7	9	5	8	
7	3	7	6	4	5	2	5	5.5	9	10	10	4	
8	4	2	4	12.5	3	15	2	6	3	4	3	10	
6	4.5	5											



EXPLORING THE WEB

19.53 A two-sample t test example. Find an example of a two-sample t test on the Web. The *Journal of the American Medical Association* (jama.ama-assn.org), *Science Magazine* (www.sciencemag.org), the *Canadian Medical Association Journal* (www.cmaj.ca), the *Journal of Statistics Education* (www.amstat.org/publications/jse), or perhaps the *Journal of Quantitative Analysis in Sports* (www.bepress.com/jqas) are possible sources. To help locate an article, look through the abstracts of articles. Once you find a suitable article, read it and then briefly describe the study (including why it is a two-sample study) and its conclusions. If P -values, means, standard deviations, t statistics, etc. are reported, be sure to include them in your summary. Also, be sure to give the reference (either the Web link or the journal, issue, year, title of the paper, authors, and page numbers).

19.54 Antibiotics after surgery. If your college has online access to the *Archives of Otolaryngology—Head and Neck Surgery*, read the article “Duration-Related Efficacy of Postoperative Antibiotics Following Pediatric Tonsillectomy: A Prospective, Randomized, Placebo-Controlled Trial” (available online at archotol.ama-assn.org/cgi/content/full/135/10/984). The authors appear to use two-sample t procedures in the paper. After reading the article, read the (brief) discussion on the *Chance* Web site, www.causeweb.org/wiki/chance/index.php/Chance_News_57 (You can find this under the heading “Some recent studies of potential interest.”). What are some criticisms or concerns expressed about the study in the *Chance* article?

5:53 P

Eric Ho

as
tom

Hello Eric, we had a great bbq yester
Do u get

Q W E R T Y

A S D F G H

Z X C V B

pace

123

Inference about a Population Proportion

Our discussion of statistical inference to this point has concerned making inferences about population *means*. Now we turn to questions about the *proportion* of some outcome in a population. Here are some examples that call for inference about population proportions.

EXAMPLE 20.1 Risky behavior in the age of AIDS

How common is behavior that puts people at risk of AIDS? In the early 1990s, the landmark National AIDS Behavioral Surveys interviewed a random sample of 2673 adult heterosexuals. Of these, 170 had more than one sexual partner in the past year. That's 6.36% of the sample.¹ Based on these data, what can we say about the percent of all adult heterosexuals who have multiple partners? We want to *estimate a single population proportion*. This chapter concerns inference about one proportion. ■

EXAMPLE 20.2 Young adults living at home

A surprising number of young adults (ages 19 to 25) still live at home with their parents. A random sample of 2253 men and 2629 women in this age group found that 44% of the men but only 35% of the women lived at home. Is this significant evidence that the proportions living at home differ in the populations of all young men and all young women? We want to *compare two population proportions*. This is the topic of Chapter 21. ■

To do inference about a population mean μ , we use the mean \bar{x} of a random sample from the population. The reasoning of inference starts with the sampling distribution of \bar{x} . Now we follow the same pattern, replacing means by proportions.

IN THIS CHAPTER WE COVER...

- The sample proportion \hat{p}
- Large-sample confidence intervals for a proportion
- Accurate confidence intervals for a proportion
- Choosing the sample size
- Significance tests for a proportion

THE SAMPLE PROPORTION \hat{p}

We are interested in the unknown proportion p of a population that has some outcome. For convenience, call the outcome we are looking for a “success.” In Example 20.1, the population is adult heterosexuals, and the parameter p is the proportion who have had more than one sexual partner in the past year. To estimate p , the National AIDS Behavioral Surveys used random dialing of telephone numbers to contact a sample of 2673 people. Of these, 170 said they had multiple sexual partners. The statistic that estimates the parameter p is the **sample proportion**

sample proportion

$$\begin{aligned}\hat{p} &= \frac{\text{number of successes in the sample}}{\text{total number of individuals in the sample}} \\ &= \frac{170}{2673} = 0.0636\end{aligned}$$

Read the sample proportion \hat{p} as “p-hat.”

How good is the statistic \hat{p} as an estimate of the parameter p ? To find out, we ask, “What would happen if we took many samples?” The sampling distribution of \hat{p} answers this question. Here are the facts.²

SAMPLING DISTRIBUTION OF A SAMPLE PROPORTION

Draw an SRS of size n from a large population that contains proportion p of successes. Let \hat{p} be the **sample proportion** of successes,

$$\hat{p} = \frac{\text{number of successes in the sample}}{n}$$

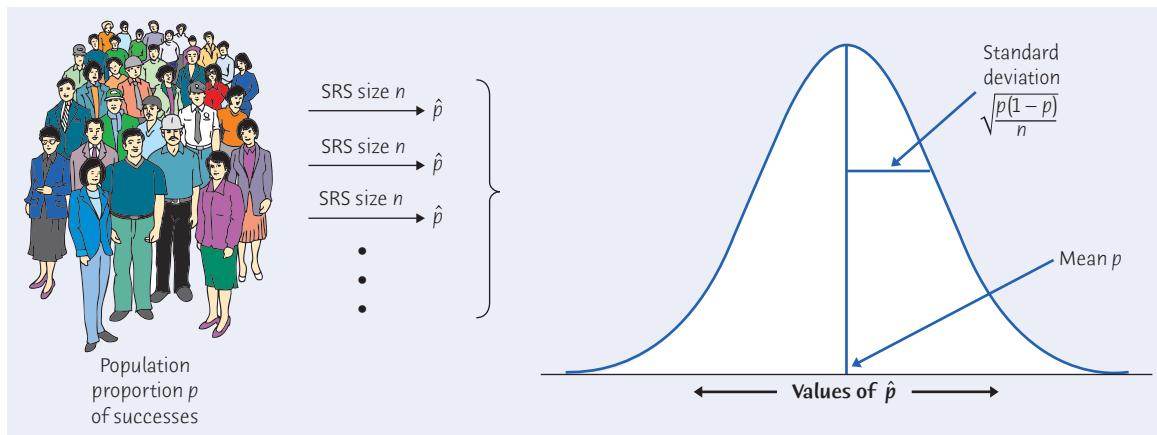
Then:

- The **mean** of the sampling distribution is p .
- The **standard deviation** of the sampling distribution is

$$\sqrt{\frac{p(1-p)}{n}}$$

- As the sample size increases, the sampling distribution of \hat{p} becomes **approximately Normal**. That is, for large n , \hat{p} has approximately the $N(p, \sqrt{p(1-p)/n})$ distribution.

Figure 20.1 summarizes these facts in a form that helps you recall the big idea of a sampling distribution. The behavior of sample proportions \hat{p} is similar to the behavior of sample means \bar{x} , except that the distribution of \hat{p} is only approximately Normal. The mean of the sampling distribution of \hat{p} is the true value of the population proportion p . That is, \hat{p} is an unbiased estimator of p . The standard deviation of \hat{p} gets smaller as the sample size n gets larger, so that estimation is likely to be more accurate when the sample is larger. As is the case for \bar{x} , the standard deviation gets smaller only at the rate \sqrt{n} . We need four times as many observations to cut the standard deviation in half.

**FIGURE 20.1**

Select a large SRS from a population in which the proportion p are successes. The sampling distribution of the proportion \hat{p} of successes in the sample is approximately Normal. The mean is p and the standard deviation is $\sqrt{p(1 - p)/n}$.

EXAMPLE 20.3 Asking about risky behavior

Suppose that in fact 6% of all adult heterosexuals had more than one sexual partner in the past year (and would admit it when asked). The National AIDS Behavioral Surveys interviewed a random sample of 2673 people from this population. In many such samples, the proportion \hat{p} of the 2673 people in the sample who had more than one partner would vary according to (approximately) the Normal distribution with mean 0.06 and standard deviation

$$\begin{aligned}\sqrt{\frac{p(1-p)}{n}} &= \sqrt{\frac{(0.06)(0.94)}{2673}} \\ &= \sqrt{0.0000211} = 0.00459\end{aligned}$$

APPLY YOUR KNOWLEDGE

20.1 Prayer among the Millennials. The Millennial generation (so called because they were born after 1980 and began to come of age around the year 2000) are less religiously active than older Americans. One of the questions in the General Social Survey in 2008 was “How often does the respondent pray?” Among the 385 respondents in the survey between 18 and 30 years of age, 247 prayed at least once a week.³

- (a) Describe the population and explain in words what the parameter p is.
- (b) Give the numerical value of the statistic \hat{p} that estimates p .

20.2 Texting. Consumers turned to their mobile devices in growing numbers and increasing frequency throughout 2010. Text messaging led as the top mobile activity, with 68% of American mobile subscribers texting in 2010.⁴ A polling firm contacts an SRS of 1200 people chosen from the population of American mobile subscribers. If the sample were repeated many times, what would be the range of the sample proportions of mobile subscribers who have texted, according to the 95 part of the 68–95–99.7 rule?

20.3 Social-networking sites. About 70% of young adult Internet users (ages 18 to 29) use social-networking sites. Suppose that a sample survey contacts an SRS of 1500 young adult Internet users and calculates the proportion \hat{p} in this sample who use social-networking sites.

- What is the approximate distribution of \hat{p} ?
- If the sample size were 6000 rather than 1500, what would be the approximate distribution of \hat{p} ?

LARGE-SAMPLE CONFIDENCE INTERVALS FOR A PROPORTION

We can follow the same path from sampling distribution to confidence interval as we did for \bar{x} in Chapter 14. To obtain a level C confidence interval for p , we start by capturing the central probability C in the distribution of \hat{p} . To do this, go out z^* standard deviations from the mean p , where z^* is the critical value that captures the central area C under the standard Normal curve. Figure 20.2 shows the result. The confidence interval is

$$\hat{p} \pm z^* \sqrt{\frac{p(1-p)}{n}}$$

This won't do, because we don't know the value of p . So we replace the standard deviation by the **standard error of \hat{p}**

$$SE = \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

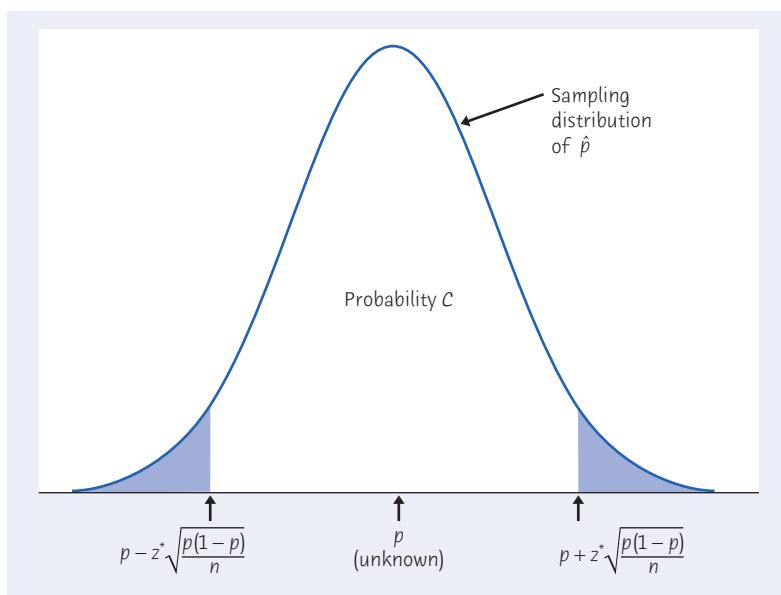


FIGURE 20.2

With probability C , \hat{p} lies within $\pm z^* \sqrt{p(1-p)/n}$ of the unknown population proportion p . That is to say that in these samples p lies within $\pm z^* \sqrt{p(1-p)/n}$ of \hat{p} .

to get the confidence interval

$$\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

This interval has the form

$$\text{estimate} \pm z^* \text{SE}_{\text{estimate}}$$

We can trust this confidence interval only for large samples. Because the number of successes must be a whole number, using a continuous Normal distribution to describe the behavior of \hat{p} may not be accurate unless n is large. Because the approximation is least accurate for populations that are almost all successes or almost all failures, we require that the sample have both enough successes and enough failures rather than that the overall sample size be large. *Pay attention to both conditions for inference in the box below that summarizes the confidence interval: we must as usual be willing to regard the sample as an SRS from the population, and the sample must have both enough successes and enough failures.*



LARGE-SAMPLE CONFIDENCE INTERVAL FOR A POPULATION PROPORTION

Draw an SRS of size n from a large population that contains an unknown proportion p of successes. An approximate level C confidence interval for p is

$$\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

where z^* is the critical value for the standard Normal density curve with area C between $-z^*$ and z^* .

Use this interval only when the numbers of successes and failures in the sample are both at least 15.⁵

Why not t ? Notice that we don't change z^* to t^* when we replace the standard deviation by the standard error. When the sample mean \bar{x} estimates the population mean μ , a separate parameter σ describes the spread of the distribution of \bar{x} . We separately estimate σ , and this leads to a t distribution. When the sample proportion \hat{p} estimates the population proportion p , the spread depends on p , not on a separate parameter. There is no t distribution—we just make the Normal approximation a bit less accurate when we replace p in the standard deviation by \hat{p} .

EXAMPLE 20.4 Estimating risky behavior

The four-step process for any confidence interval is outlined on pages 358–359.

STATE: The National AIDS Behavioral Surveys found that 170 of a sample of 2673 adult heterosexuals had multiple partners. That is,

$$\hat{p} = \frac{170}{2673} = 0.0636$$



What can we say about the population of all adult heterosexuals?

PLAN: We will give a 99% confidence interval to estimate the proportion p of all adult heterosexuals who have multiple partners.

SOLVE: First verify the conditions for inference:

- The sampling design was a complex stratified sample, and the survey used inference procedures for that design. The overall effect is close to an SRS, however.
- The sample is large enough: the numbers of successes (170) and failures (2503) in the sample are both much larger than 15.

The sample size condition is easily satisfied. The condition that the sample be an SRS is only approximately met.

A 99% confidence interval for the proportion p of all adult heterosexuals with multiple partners uses the standard Normal critical value $z^* = 2.576$. The confidence interval is

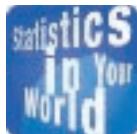
$$\begin{aligned}\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} &= 0.0636 \pm 2.576 \sqrt{\frac{(0.0636)(0.9364)}{2673}} \\ &= 0.0636 \pm 0.0122 \\ &= 0.0514 \text{ to } 0.0758\end{aligned}$$

CONCLUDE: We are 99% confident that the percent of adult heterosexuals who have had more than one sexual partner in the past year lies between about 5.1% and 7.6%. ■



As usual, the practical problems of a large sample survey weaken our confidence in the AIDS survey's conclusions. Only people in households with landline telephones could be reached. Although at the time of the survey about 89% of American households had landline telephones, as the number of cell-phone-only users increases, using a sample of households with landline phones is becoming less acceptable for surveys of the general population (see pages 213–214). Additionally, some groups at high risk for AIDS, such as people who inject illegal drugs, often don't live in settled households and were therefore underrepresented in the sample. About 30% of the people reached refused to cooperate. A nonresponse rate of 30% is not unusual in large sample surveys, but it may cause some bias if those who refuse differ systematically from those who cooperate. The survey used statistical methods that adjust for unequal response rates in different groups. Finally, some respondents may not have told the truth when asked about their sexual behavior. The survey team tried to make respondents feel comfortable. For example, Hispanic women were interviewed only by Hispanic women, and Spanish speakers were interviewed by Spanish speakers with the same regional accent (Cuban, Mexican, or Puerto Rican). Nonetheless, the survey report says that some bias is probably present:

It is more likely that the present figures are underestimates; some respondents may underreport their numbers of sexual partners and intravenous drug use because of embarrassment and fear of reprisal, or they may forget or not know details of their own or of their partner's HIV risk and their antibody testing history.⁶



Who is a smoker?

When estimating a proportion p , be sure you know

what counts as a "success." The news says that 20% of adolescents smoke. Shocking. It turns out that this is the percent who smoked at least once in the past month. If we say that a smoker is someone who smoked on at least 20 of the past 30 days and smoked at least half a pack on those days, fewer than 4% of adolescents qualify.

Reading the report of a large study like the National AIDS Behavioral Surveys reminds us that statistics in practice involves much more than formulas for inference.


APPLY YOUR KNOWLEDGE

20.4 No confidence interval. In the National AIDS Behavioral Surveys sample of 2673 adult heterosexuals, 0.2% (that's 0.002 as a decimal fraction) had both received a blood transfusion and had a sexual partner from a group at high risk of AIDS. Explain why we can't use the large-sample confidence interval to estimate the proportion p in the population who share these two risk factors.

20.5 Canadian attitudes toward guns. Canada has much stronger gun control laws than the United States, and Canadians support gun control more strongly than do Americans. A sample survey asked a random sample of 1505 adult Canadians, "Do you agree or disagree that all firearms should be registered?" Of the 1505 people in the sample, 1288 answered either "Agree strongly" or "Agree somewhat."⁷

- The survey dialed residential telephone numbers at random in ten Canadian provinces (omitting the sparsely populated northern territories). Based on what you know about sample surveys, what is likely to be the biggest weakness in this survey?
- Nonetheless, act as if we have an SRS from adults in the Canadian provinces. Give a 95% confidence interval for the proportion who support registration of all firearms.

20.6 Weight-lifting injuries. Resistance training is a popular form of conditioning aimed at enhancing sports performance and is widely used among high school, college, and professional athletes, although its use for younger athletes is controversial. Researchers obtained a random sample of 4111 patients between the ages of 8 and 30 who were admitted to U.S. emergency rooms with injuries classified by the Consumer Product Safety Commission code "weightlifting." These injuries were further classified as "accidental" if caused by dropped weight or improper equipment use. Of the 4111 weight-lifting injuries, 1552 were classified as accidental.⁸ Give a 90% confidence interval for the proportion of weight-lifting injuries in this age group that were accidental. Follow the four-step process as illustrated in Example 20.4.



ACCURATE CONFIDENCE INTERVALS FOR A PROPORTION

The confidence interval $\hat{p} \pm z^* \sqrt{\hat{p}(1 - \hat{p})/n}$ for a sample proportion p is easy to calculate. It is also easy to understand because it is based directly on the approximately Normal distribution of \hat{p} . Unfortunately, confidence levels from this interval are often quite inaccurate unless the sample is very large. The actual confidence level is usually *less* than the confidence level you asked for in choosing the critical value z^* . That's bad. What is worse, accuracy does not consistently get better as the sample size n increases. There are "lucky" and "unlucky" combinations of the sample size n and the true population proportion p .

Fortunately, there is a simple modification that is almost magically effective in improving the accuracy of the confidence interval. We call it the "plus four"

plus four estimate

method because all you need to do is *add four imaginary observations, two successes and two failures*. With the added observations, the **plus four estimate** of p is

$$\tilde{p} = \frac{\text{number of successes in the sample} + 2}{n + 4}$$

The formula for the confidence interval is exactly as before, with the new sample size and number of successes.⁹ You do not need software that offers the plus four interval—just enter the new sample size (actual size + 4) and number of successes (actual number + 2) into the large-sample procedure.

PLUS FOUR CONFIDENCE INTERVAL FOR A PROPORTION

Draw an SRS of size n from a large population that contains an unknown proportion p of successes. To get the **plus four confidence interval for p** , add four imaginary observations, two successes and two failures. Then use the large-sample confidence interval with the new sample size ($n + 4$) and number of successes (actual number + 2).

Use this interval when the confidence level is at least 90% and the sample size n is at least 10, with any counts of successes and failures.



Kolvenbach/Alamy

EXAMPLE 20.5 Cocaine traces in Spanish currency

STATE: Cocaine users commonly snort the powder up the nose through a rolled-up paper currency bill. Spain has a high rate of cocaine use, so it's not surprising that euro paper currency in Spain often contains traces of cocaine. Researchers collected 20 euro bills in each of several Spanish cities. In Madrid, 17 out of 20 contained traces of cocaine.¹⁰ The researchers note that we can't tell whether the bills had been used to snort cocaine or had been contaminated in currency-sorting machines. Estimate the proportion of all euro bills in Madrid that have traces of cocaine.

PLAN: Take p to be the proportion of bills that contain cocaine traces. Give a 95% confidence interval for p .

SOLVE: The conditions for use of the large-sample interval are not met because there are only 3 failures. To apply the plus four method, add two successes and two failures to the original data. The plus four estimate of p is

$$\tilde{p} = \frac{17 + 2}{20 + 4} = \frac{19}{24} = 0.7917$$

We calculate the plus four confidence interval in the same way as we do the large-sample interval, but we base it on 19 successes in 24 observations. Here it is:

$$\begin{aligned}\tilde{p} \pm z^* \sqrt{\frac{\tilde{p}(1 - \tilde{p})}{n + 4}} &= 0.7917 \pm 1.960 \sqrt{\frac{(0.7917)(0.2083)}{24}} \\ &= 0.7917 \pm 0.1625 \\ &= 0.6292 \text{ to } 0.9542\end{aligned}$$

CONCLUDE: We estimate with 95% confidence that between about 63% and 95% of all euro bills in Madrid contain traces of cocaine. ■

For comparison, the ordinary sample proportion is

$$\hat{p} = \frac{17}{20} = 0.85$$

The plus four estimate $\tilde{p} = 0.7917$ in Example 20.5 is farther away from 1 than $\hat{p} = 0.85$. The plus four estimate gains its added accuracy by always moving toward 0.5 and away from 1 or 0, whichever is closer. This is particularly helpful when the sample contains only a few successes or a few failures. The numerical difference between a large-sample interval and the corresponding plus four interval is often small. Remember that the confidence level is the probability that the interval will catch the true population proportion *in very many uses*. Small differences every time add up to accurate confidence levels from plus four versus inaccurate levels from the large-sample interval.

How much more accurate is the plus four interval? Computer studies have asked how large n must be to guarantee that the actual probability that a 95% confidence interval covers the true parameter value is at least 0.94 for all samples of size n or larger. If $p = 0.1$, for example, the answer is $n = 646$ for the large-sample interval and $n = 11$ for the plus four interval.¹¹ The consensus of computational and theoretical studies is that plus four is very much better than the large-sample interval for many combinations of n and p . (If you use software such as Minitab, you may find an “exact method” on the menu. Despite the appealing name, this method is often less accurate than plus four.) **We recommend that you always use the plus four interval for estimating a proportion.**

APPLY YOUR KNOWLEDGE

20.7 Black raspberries and cancer. Sample surveys usually contact large samples, so we can use the large-sample confidence interval if the sample design is close to an SRS. Scientific studies often use smaller samples that require the plus four method. For example, familial adenomatous polyposis (FAP) is a rare inherited disease characterized by the development of an extreme number of polyps early in life and by colon cancer in virtually 100% of patients before the age of 40. A group of 14 people suffering from FAP and being treated at the Cleveland Clinic drank black raspberry powder in a slurry of water every day for nine months. The number of polyps was reduced in 11 out of 14 of these patients.¹²

- Why can't we use the large-sample confidence interval for the proportion p of patients suffering from FAP who will have the number of polyps reduced after 9 months of treatment?
- The plus four method adds four observations, two successes and two failures. What are the sample size and the number of successes after you do this? What is the plus four estimate \tilde{p} of p ?
- Give the plus four 90% confidence interval for the proportion of patients suffering from FAP who will have the number of polyps reduced after nine months of treatment.

20.8 Computer/Internet-based crime. With over 50% of adults spending more than an hour a day on the Internet, the number experiencing computer- or Internet-based



Jennifer Shields/Getty Images

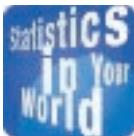
crime continues to rise. A survey in 2010 of a random sample of 1025 adults, aged 18 and older, reached by random digit dialing found 113 adults in the sample who said that they or a household member was a victim of a computer or Internet crime on their home computer in the past year.¹³

- Give the 95% large-sample confidence interval for the proportion p of all households that have experienced computer or Internet crime during the year before the survey was conducted.
- Give the plus four 95% confidence interval for p . If you express the two intervals in percents, rounded to the nearest tenth of a percent, how do they differ? (The plus four interval always pulls the results away from 0% or 100%, whichever is closer. Even though the condition for using the large-sample interval is met, the plus four interval is more trustworthy.)

20.9 Cocaine traces in Spanish currency, continued. The plus four method is particularly useful when there are no successes or no failures in the data. The study of Spanish currency described in Example 20.5 found that in Seville, all 20 of a sample of 20 euro bills had cocaine traces.

- What is the sample proportion \hat{p} of contaminated bills? What is the large-sample 95% confidence interval for p ? It's not plausible that *every* bill in Seville has cocaine traces, as this interval says.
- Find the plus four estimate \tilde{p} and the plus four 95% confidence interval for p . These results are more reasonable.

CHOOSING THE SAMPLE SIZE



New York, New York

New York City, they say, is bigger, richer, faster,

ruder. Maybe there's something to that. The sample survey firm Zogby International says that as a national average it takes 5 telephone calls to reach a live person. When calling to New York, it takes 12 calls. Survey firms assign their best interviewers to make calls to New York and often pay them bonuses to cope with the stress.

In planning a study, we may want to choose a sample size that will allow us to estimate the parameter within a given margin of error. We saw earlier (page 401) how to do this for a population mean. The method is similar for estimating a population proportion.

The margin of error in the large-sample confidence interval for p is

$$m = z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

Here z^* is the standard Normal critical value for the level of confidence we want. Because the margin of error involves the sample proportion of successes \hat{p} , we need to guess this value when choosing n . Call our guess p^* . Here are two ways to get p^* :

- Use a guess p^* based on a pilot study or on past experience with similar studies. You can do several calculations to cover the range of values of \hat{p} you might get.
- Use $p^* = 0.5$ as the guess. The margin of error m is largest when $\hat{p} = 0.5$, so this guess is conservative in the sense that if we get any other \hat{p} when we do our study, we will get a margin of error smaller than planned.

Once you have a guess p^* , the recipe for the margin of error can be solved to give the sample size n needed. Here is the result for the large-sample confidence interval. For simplicity, use this result even if you plan to use the plus four interval.

SAMPLE SIZE FOR DESIRED MARGIN OF ERROR

The level C confidence interval for a population proportion p will have margin of error approximately equal to a specified value m when the sample size is

$$n = \left(\frac{z^*}{m} \right)^2 p^*(1 - p^*)$$

where p^* is a guessed value for the sample proportion. The margin of error will always be less than or equal to m if you take the guess p^* to be 0.5.

Which method for finding the guess p^* should you use? The n you get doesn't change much when you change p^* as long as p^* is not too far from 0.5. You can use the conservative guess $p^* = 0.5$ if you expect the true \hat{p} to be roughly between 0.3 and 0.7. If the true \hat{p} is close to 0 or 1, using $p^* = 0.5$ as your guess will give a sample much larger than you need. Try to use a better guess from a pilot study when you suspect that \hat{p} will be less than 0.3 or greater than 0.7.

EXAMPLE 20.6 Planning a poll

STATE: Gloria Chavez and Ronald Flynn are the candidates for mayor in a large city. You are planning a sample survey to determine what percent of the voters intend to vote for Chavez. You will contact an SRS of registered voters in the city. You want to estimate the proportion p of Chavez voters with 95% confidence and a margin of error no greater than 3%, or 0.03. How large a sample do you need?



PLAN: Find the sample size n needed for margin of error $m = 0.03$ and 95% confidence. The winner's share in all but the most lopsided elections is between 30% and 70% of the vote. You can use the guess $p^* = 0.5$.

SOLVE: The sample size you need is

$$n = \left(\frac{1.96}{0.03} \right)^2 (0.5)(1 - 0.5) = 1067.1$$

Round the result up to $n = 1068$. (Rounding down would give a margin of error slightly greater than 0.03.)

CONCLUDE: An SRS of 1068 registered voters is adequate for margin of error $\pm 3\%$. ■



Colin Anderson/BrandX/Age fotostock

If you want a 2.5% margin of error rather than 3%, then (after rounding up)

$$n = \left(\frac{1.96}{0.025} \right)^2 (0.5)(1 - 0.5) = 1537$$

For a 2% margin of error the sample size you need is

$$n = \left(\frac{1.96}{0.02} \right)^2 (0.5)(1 - 0.5) = 2401$$

As usual, smaller margins of error call for larger samples.


APPLY YOUR KNOWLEDGE

20.10 Canadians and doctor-assisted suicide. A Gallup Poll asked a sample of Canadian adults if they thought the law should allow doctors to end the life of a patient who is in great pain and near death if the patient makes a request in writing. The poll included 270 people in Québec, 221 of whom agreed that doctor-assisted suicide should be allowed.¹⁴

- What is the margin of error of the large-sample 95% confidence interval for the proportion of all Québec adults who would allow doctor-assisted suicide?
- How large a sample is needed to get the common ± 3 percentage point margin of error? Use the previous sample as a pilot study to get p^* .

20.11 Can you taste PTC? PTC is a substance that has a strong bitter taste for some people and is tasteless for others. The ability to taste PTC is inherited. About 75% of Italians can taste PTC, for example. You want to estimate the proportion of Americans with at least one Italian grandparent who can taste PTC. Starting with the 75% estimate for Italians, how large a sample must you collect in order to estimate the proportion of PTC tasters within ± 0.04 with 90% confidence?

SIGNIFICANCE TESTS FOR A PROPORTION

The test statistic for the null hypothesis $H_0: p = p_0$ is the sample proportion \hat{p} standardized using the value p_0 specified by H_0 ,

$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}}$$

This z statistic has approximately the standard Normal distribution when H_0 is true. P -values therefore come from the standard Normal distribution. Because H_0 fixes a value of p , the inaccuracy that plagues the large-sample confidence interval does not affect tests. Here is the procedure for tests and the conditions for inference.

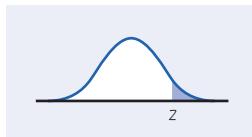
SIGNIFICANCE TESTS FOR A PROPORTION

Draw an SRS of size n from a large population that contains an unknown proportion p of successes. To **test the hypothesis $H_0: p = p_0$** , compute the z statistic

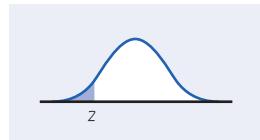
$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}}$$

In terms of a variable Z having the standard Normal distribution, the approximate P -value for a test of H_0 against

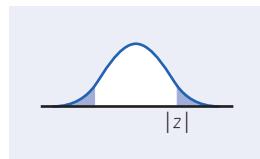
$$H_a: p > p_0 \quad \text{is} \quad P(Z \geq z)$$



$$H_a: p < p_0 \text{ is } P(Z \leq z)$$



$$H_a: p \neq p_0 \text{ is } 2P(|Z| \geq |z|)$$



Use this test when the sample size n is so large that both the expected number of successes and failures are at least 10 when H_0 is true. Specifically, check that both np_0 and $n(1 - p_0)$ are 10 or more.¹⁵

EXAMPLE 20.7 Are boys more likely?

The four-step process for any significance test is outlined on page 379.

STATE: We hear that newborn babies are more likely to be boys than girls, presumably to compensate for higher mortality among boys in early life. Is this true? A random sample found 13,173 boys among 25,468 firstborn children.¹⁶ The sample proportion of boys was

$$\hat{p} = \frac{13,173}{25,468} = 0.5172$$

Boys do make up more than half of the sample, but of course we don't expect a perfect 50-50 split in a random sample. Is this sample evidence that boys are more common than girls in the entire population?

PLAN: Take p to be the proportion of boys among all firstborn children of American mothers. (Biology says that this should be the same as the proportion among all children, but the survey data concern first births.) We want to test the hypotheses

$$H_0: p = 0.5$$

$$H_a: p > 0.5$$

SOLVE: The conditions for inference require that we have a random sample and that $np_0 = (25,468)(0.5) = 12,734$ and $n(1 - p_0) = (25,468)(0.5) = 12,734$ are both greater than 10. Since the conditions for inference are met, we can go on to find the z test statistic:

$$\begin{aligned} z &= \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}} \\ &= \frac{0.5172 - 0.5}{\sqrt{\frac{(0.5)(0.5)}{25,468}}} = 5.49 \end{aligned}$$



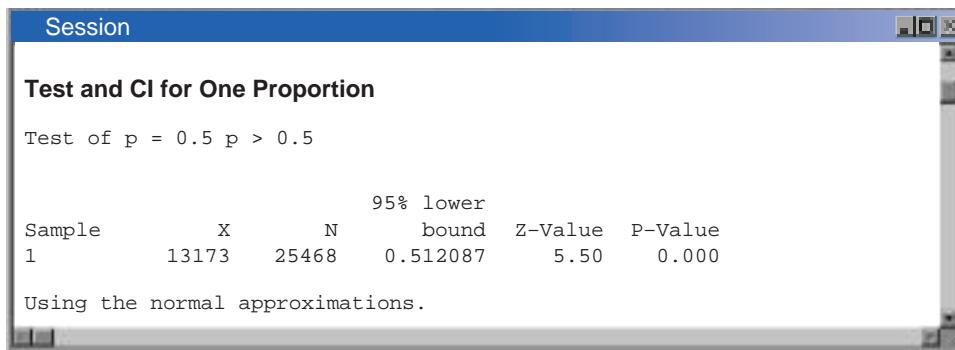
Blaine Harrington III/CORBIS

The P -value is the area under the standard Normal curve to the right of $z = 5.49$. We know that this is very small; Table C shows that $P < 0.0005$. Minitab (Figure 20.3) says that P is 0 to three decimal places.

CONCLUDE: There is very strong evidence that more than half of firstborns are boys ($P < 0.001$). ■

FIGURE 20.3

Minitab output for the significance test of Example 20.7. Roundoff error in Example 20.7 explains the small difference (5.49 versus 5.50) in the values of the z statistic.



EXAMPLE 20.8 Estimating the chance of a boy

With 13,173 successes in 25,468 trials, we have at least 15 successes and 15 failures in the sample. The conditions for the large-sample confidence interval, as well as for the plus four confidence interval, are easily met. Because of the large sample size, the estimates of p are almost identical. Both are 0.5172 to four decimal places. So both methods give the 99% confidence interval

$$\begin{aligned}\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} &= 0.5172 \pm 2.576 \sqrt{\frac{(0.5172)(0.4828)}{25,468}} \\ &= 0.5172 \pm 0.0081 \\ &= 0.5091 \text{ to } 0.5253\end{aligned}$$

We are 99% confident that between about 51% and 52.5% of first children are boys.

The confidence interval is more informative than the test in Example 20.7, which tells us only that more than half are boys. ■

APPLY YOUR KNOWLEDGE



20.12 Spinning euros. All euros have a national image on the “heads” side and a common design on the “tails” side. Spinning a coin, unlike tossing it, may not give heads and tails equal probabilities. Polish students spun the Belgium euro 250 times, with its portly king, Albert, displayed on the heads side. The result was 140 heads.¹⁷ How significant is this evidence against equal probabilities? Follow the four-step process as illustrated in Example 20.7.



20.13 Vote for the best face? We often judge other people by their faces. It appears that some people judge candidates for elected office by their faces. Psychologists showed head-and-shoulders photos of the two main candidates in 32 races for the U.S. Senate to many subjects (dropping subjects who recognized one of the

candidates) to see which candidate was rated “more competent” based on nothing but the photos. On election day, the candidates whose faces looked more competent won 22 of the 32 contests.¹⁸ If faces don’t influence voting, half of all races in the long run should be won by the candidate with the better face. Is there evidence that the candidate with the better face wins more than half the time? Follow the four-step process as illustrated in Example 20.7.

20.14 No test. Explain whether we can use the z test for a proportion in these situations:

- You toss a coin 10 times to test the hypothesis $H_0: p = 0.5$ that the coin is balanced.
- A local candidate contacts an SRS of 900 of the registered voters in his district to see if there is evidence that more than half support the bill he is sponsoring.
- A college president says, “99% of the alumni support my firing of Coach Boggs.” You contact an SRS of 200 of the college’s 15,000 living alumni to test the hypothesis $H_0: p = 0.99$.

CHAPTER 20 SUMMARY

CHAPTER SPECIFICS

- Tests and confidence intervals for a population proportion p when the data are an SRS of size n are based on the **sample proportion** \hat{p} .
- When n is large, \hat{p} has approximately the Normal distribution with mean p and standard deviation $\sqrt{p(1 - p)/n}$.
- The level C **large-sample confidence interval for p** is

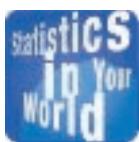
$$\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

where z^* is the critical value for the standard Normal curve with area C between $-z^*$ and z^* .

- The true confidence level of the large-sample interval can be substantially less than the planned level C unless the sample is very large. We recommend using the plus four interval instead.
- To get a more accurate confidence interval, add four imaginary observations, two successes and two failures, to your sample. Then use the same formula for the confidence interval. This is the **plus four confidence interval**. Use this interval in practice for confidence level 90% or higher and sample size n at least 10.
- The **sample size** needed to obtain a confidence interval with approximate margin of error m for a population proportion is

$$n = \left(\frac{z^*}{m}\right)^2 p^*(1 - p^*)$$

where p^* is a guessed value for the sample proportion \hat{p} , and z^* is the standard Normal critical point for the level of confidence you want. If you use $p^* = 0.5$ in this formula, the margin of error of the interval will be less than or equal to m no matter what the value of \hat{p} is.



Kids on bikes

In the most recent year for which data are available, 77% of children

killed in bicycle accidents were boys. You might take these data as a sample and start from $\hat{p} = 0.77$ to do inference about bicycle deaths in the near future. What you should not do is conclude that boys on bikes are in greater danger than girls. We don't know how many boys and girls ride bikes—it may be that most fatalities are boys because most riders are boys.

- Significance tests for $H_0: p = p_0$ are based on the z statistic

$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}}$$

with P -values calculated from the standard Normal distribution. Use this test in practice when $np_0 \geq 10$ and $n(1 - p_0) \geq 10$.

LINK IT

The methods of this chapter can be used to compute confidence intervals and test hypotheses about a population proportion. This may be the proportion in a population with some attribute of interest such as the population proportion of young adults who use social-networking sites. The proportion could also correspond to the probability of an outcome in an experiment such as the probability that a spinning coin will land on heads. Since the methods of this chapter are approximations, it is important to always check the conditions required for the approximations to work well, whether you are using the large-sample confidence interval, the plus four method, or the z test.

When making inferences about a proportion in a population, an important assumption for the methods of this chapter is that the data are an SRS from the population. This is the most difficult assumption to guarantee because of the many difficulties associated with obtaining an SRS. These difficulties were described in Chapter 8. With nonresponse rates in many surveys that are over 80%, the final sample may not be representative of the population even when researchers initially selected an SRS. Because of this, most surveys of large populations need to use more complicated sampling schemes as well as modify the estimates to adjust for problems such as nonresponse.

In Chapters 18 and 19, we first considered inference about a single mean and then turned our attention to situations that required the comparison of two population means. In many instances, we want to compare two population proportions rather than making inferences about a single proportion. Methods to compare two population proportions will be described in the next chapter.

CHECK YOUR SKILLS

20.15 The proportion of drivers who use seat belts depends on things like age, gender, and ethnicity. As part of a broader study, investigators observed a random sample of 117 female Hispanic drivers in Boston. Suppose that in fact 60% of all female Hispanic drivers in the Boston area wear seat belts. In repeated samples, the sample proportion \hat{p} would follow approximately a Normal distribution with mean

- (a) 70.2. (b) 0.6. (c) 0.002.

20.16 The standard deviation of the distribution of \hat{p} in the previous exercise is about

- (a) 0.002. (b) 0.045. (c) 0.24.

20.17 A 2010 study finds that in a random sample of 3000 American adults aged 18 and over, 1410 owned an MP3 player such as an iPod. The sample proportion \hat{p} who own an MP3 player is

- (a) 47. (b) 53. (c) 0.47.

20.18 Based on the sample in the previous exercise, the 95% large-sample confidence interval for the proportion of all American adults aged 18 and over who own an MP3 player is

- (a) 0.47 ± 0.009 . (b) 0.47 ± 0.015 .
(c) 0.47 ± 0.018 .

20.19 How many American adults aged 18 and over must be interviewed to estimate the proportion who own MP3 players within ± 0.02 with 99% confidence? Use 0.5 as the conservative guess for p .

- (a) $n = 1692$ (b) $n = 2401$ (c) $n = 4148$

20.20 An opinion poll asks an SRS of 100 college seniors how they view their job prospects. In all, 53 say “Good.” The plus four 95% confidence interval for estimating the proportion of all college seniors who think their job prospects are good is

- (a) 0.529 ± 0.096 . (b) 0.529 ± 0.098 . (c) 0.529 ± 0.049 .

20.21 The sample survey in Exercise 20.20 actually called 130 seniors, but 30 of the seniors refused to answer. This nonresponse could cause the survey result to be in error. The error due to nonresponse

- (a) is in addition to the margin of error found in Exercise 20.20.
 (b) is included in the margin of error found in Exercise 20.20.
 (c) can be ignored because it isn’t random.

20.22 Does the poll in Exercise 20.20 give reason to conclude that more than half of all seniors think their job prospects are good? The hypotheses for a test to answer this question are

- (a) $H_0: p = 0.5$, $H_a: p > 0.5$.
 (b) $H_0: p > 0.5$, $H_a: p = 0.5$.
 (c) $H_0: p = 0.5$, $H_a: p \neq 0.5$.

20.23 The value of the z statistic for the test of the previous exercise is about

- (a) $z = 12$. (b) $z = 6$. (c) $z = 0.6$.

20.24 A Gallup Poll found that only 28% of American adults expect to inherit money or valuable possessions from a relative. The poll’s margin of error was 3%. This means that

- (a) the poll used a method that gets an answer within 3% of the truth about the population 95% of the time.
 (b) we can be sure that the percent of all adults who expect an inheritance is between 25% and 31%.
 (c) if Gallup takes another poll using the same method, the results of the second poll will lie between 25% and 31%.

CHAPTER 20 EXERCISES

We recommend using the plus four method for all confidence intervals for a proportion. However, the large-sample method is acceptable when the guidelines for its use are met.

20.25 Do smokers know that smoking is bad for them? The Harris Poll asked a sample of smokers, “Do you believe that smoking will probably shorten your life, or not?” Of the 1010 people in the sample, 848 said “Yes.”

- (a) Harris called residential telephone numbers at random in an attempt to contact an SRS of smokers. Based on what you know about national sample surveys, what is likely to be the biggest weakness in the survey?
 (b) We will nonetheless act as if the people interviewed are an SRS of smokers. Give a 95% confidence interval for the percent of smokers who agree that smoking will probably shorten their lives.



20.26 Reporting cheating. Students are reluctant to report cheating by other students. A student project put this question to an SRS of 172 undergraduates at a large university: “You witness two students cheating on a quiz. Do you go to

the professor?” Only 19 answered “Yes.”¹⁹ Give a 95% confidence interval for the proportion of all undergraduates at this university who would report cheating.

20.27 Harris announces a margin of error. Exercise 20.25 describes a Harris Poll survey of smokers in which 848 of a sample of 1010 smokers agreed that smoking would probably shorten their lives. Harris announces a margin of error of ± 3 percentage points for all samples of about this size. Opinion polls announce the margin of error for 95% confidence.

- (a) What is the actual margin of error (in percent) for the large-sample confidence interval from this sample?
 (b) The margin of error is largest when $\hat{p} = 0.5$. What would the margin of error (in percent) be if the sample had resulted in $\hat{p} = 0.5$?
 (c) Why do you think that Harris announces a $\pm 3\%$ margin of error for all samples of about this size?

20.28 Prayer among the Millennials, continued. The Millennial generation (so called because they were born after 1980 and began to come of age around the year 2000) are less religiously active than older Americans. One of the questions in the General Social Survey in 2008 was “How often does the respondent pray?” Among the 385 respondents in the survey between 18 and 30 years of age, 247 prayed at least once a week.²⁰

- (a) What is the large-sample 99% confidence interval for the proportion p of all adults between 18 and 30 years of age who pray at least once a week?
- (b) Give the plus four 99% confidence interval for p . If you express the two intervals in percents and round to the nearest tenth of a percent, how do they differ? (As always, the plus four method pulls results away from 0% or 100%, whichever is closer. Although the condition for the large-sample interval is met, the plus four interval is more trustworthy.)

20.29 Internet searches and cell phones. Pew Internet and American Life Project asked a random sample of 2485 cell phone users whether they had used their cell phone to look up health or medical information. Of these, 422 said “Yes.”²¹

- (a) Pew dialed cell phone telephone numbers at random in the continental United States in an attempt to contact a random sample of adults. Based on what you know about national sample surveys, what is likely to be the biggest weakness in the survey?
- (b) Act as if the sample is an SRS. Give a large-sample 90% confidence interval for the proportion p of all cell phone users who have used their cell phone to look up health or medical information.
- (c) Three out of the five most popular health-related searches on cell phones have to do with sex: “pregnancy,” “herpes,” and “STD” (sexually transmitted diseases). Sex-related queries don’t even show up on Google and Yahoo’s lists of the top five health searches on computers. What do you think explains the difference in the topics of health-related searches on cell phones versus computers? When drawing conclusions from a sample, you must always be careful to take into account the relevant population.

20.30 Which font? Plain type fonts such as Times New Roman are easier to read than fancy fonts such as Gigi. A group of 25 volunteer subjects read the same text in both fonts. (This is a matched pairs design. One-sample procedures for proportions, like those for means, are used to analyze data from matched pairs designs.) Of the 25 subjects, 17 said that they preferred Times New Roman for Web use. But 20 said that Gigi was more attractive.²²

- (a) Because the subjects were volunteers, conclusions from this sample can be challenged. Show that the sample size condition for the large-sample confidence interval is not met, but that the condition for the plus four interval is met.
- (b) Give a 95% confidence interval for the proportion of all adults who prefer Times New Roman for Web use. Give a 90% confidence interval for the proportion of all adults who think Gigi is more attractive.

20.31 Testing the waters. In August 2010, the *Columbus Dispatch* took water samples at 20 Ohio State Park swimming areas and tested for fecal coliform, which are bacteria found in human and animal feces. Experts warn that the tests are a snapshot of the quality of the water at the time they were taken, and levels can change as weather and other conditions vary. An unsafe level of fecal coliform means that there’s a higher chance that disease-causing bacteria are present and more risk that a swimmer will become ill. Of the 20 swimming areas tested, 13 were found to have unsafe levels of fecal coliform according to state standards. Assume that the swimming areas tested represent a random sample of swimming areas throughout the state.²³

- (a) Show that the conditions for the large-sample confidence interval are not met. Show that the conditions for the plus four interval are met.
- (b) Use the plus four method to give a 90% confidence interval for the percent of Ohio State Park swimming areas that have unsafe levels of fecal coliform.

20.32 Running red lights. A random digit dialing telephone survey of 880 drivers asked, “Recalling the last ten traffic lights you drove through, how many of them were red when you entered the intersections?” Of the 880 respondents, 171 admitted that at least one light had been red.²⁴

- (a) Give a 95% confidence interval for the proportion of all drivers who ran one or more of the last ten red lights they encountered.
- (b) Nonresponse is a practical problem for this survey—only 21.6% of calls that reached a live person were completed. Another practical problem is that people may not give truthful answers. What is the likely direction of the bias: do you think more or fewer than 171 of the 880 respondents really ran a red light? Why?

20.33 The IRS plans an SRS. The Internal Revenue Service plans to examine an SRS of individual federal income tax returns from each state. One variable of interest is the proportion of returns claiming itemized deductions. The total number of tax returns in a state varies from more than 15 million in California to fewer than 250,000 in Wyoming.

- (a) Will the margin of error for estimating the population proportion change from state to state if an SRS of 2000 tax returns is selected in each state? Explain your answer.
- (b) Will the margin of error change from state to state if an SRS of 1% of all tax returns is selected in each state? Explain your answer.

20.34 Customer satisfaction. An automobile manufacturer would like to know what proportion of its customers are not satisfied with the service provided by the local dealer.

The customer relations department will survey a random sample of customers and compute a 99% confidence interval for the proportion who are not satisfied.

- Past studies suggest that this proportion will be about 0.2. Find the sample size needed if the margin of error of the confidence interval is to be about 0.015.
- When the sample of the size found in (a) is actually contacted, 10% of the sample say they are not satisfied. What is the margin of error of the 99% confidence interval?

20.35 Surveying students. You are planning a survey of students at a large university to determine what proportion favor an increase in student fees to support an expansion of the student newspaper. Using records provided by the registrar, you can select a random sample of students. You will ask each student in the sample whether he or she is in favor of the proposed increase. Your budget will allow a sample of 100 students.

- For a sample of size 100, construct a table of the margins of error for 95% confidence intervals when \hat{p} takes the values 0.1, 0.3, 0.5, 0.7, and 0.9.
- A former editor of the student newspaper offers to provide funds for a sample of size 500. Repeat the margin of error calculations in (a) for the larger sample size. Then write a short thank-you note to the former editor describing how the larger sample size will improve the results of the survey.

In responding to Exercises 20.36 to 20.44, follow the **Plan, Solve, and Conclude** steps of the four-step process.

20.36 College-educated parents. The National Assessment of Educational Progress (NAEP) includes a “long-term trend” study that tracks reading and mathematics skills over time and obtains demographic information. In the 2008 study, a random sample of 9600 17-year-old students was selected.²⁵ The NAEP sample used a multistage design, but the overall effect is quite similar to an SRS of 17-year-olds who are still in school.

- In the sample, 46% of students had at least one parent who was a college graduate. Estimate with 99% confidence the proportion of all 17-year-old students in 2008 who had at least one parent who was a college graduate.
- The sample does not include 17-year-olds who dropped out of school, so your estimate is valid only for students. Do you think that the proportion of all 17-year-olds with at least one parent who was a college graduate would be higher or lower than 46%? Explain.

20.37 Shrubs that survive fires. Some shrubs have the useful ability to resprout from their roots after their tops are destroyed. Fire is a particular threat to shrubs

in dry climates, as it can injure the roots as well as destroy the aboveground material. One study of resprouting took place in a dry area of Mexico.²⁶ The investigators clipped the tops of samples of several species of shrubs. In some cases, they also applied a propane torch to the stumps to simulate a fire. Of 12 specimens of the shrub *Krameria cistoïdes*, 5 resprouted after fire. Estimate with 90% confidence the proportion of all shrubs of this species that will resprout after fire.

20.38 Downloading music. A husband and wife, Ted and Suzanne, share a digital music player that has a feature that randomly selects which song to play. A total of 3476 songs have been loaded into the player, some by Ted and the rest by Suzanne. They are interested in determining whether they have each loaded a different proportion of songs into the player. Suppose that when the player was in the random-selection mode, 22 of the first 30 songs selected were songs loaded by Suzanne. Let p denote the proportion of songs that were loaded by Suzanne. State the null and alternative hypotheses to be tested. How strong is the evidence that Ted and Suzanne have each loaded a different proportion of songs into the player?

20.39 Opinions about evolution. A sample survey funded by the National Science Foundation asked a random sample of American adults about biological evolution.²⁷ One question asked subjects to answer “True,” “False,” or “Not sure” to the statement “Human beings, as we know them today, developed from earlier species of animals.” Of the 1484 respondents, 594 said “True.” What can you say with 95% confidence about the percent of all American adults who think that humans developed from earlier species of animals?

20.40 Order in choice. Does the order in which wine is presented make a difference? Several choices of wine are presented one at a time, and the subject is then asked to choose his or her preferred wine at the end of the sequence. In one study, subjects were asked to taste two wine samples in sequence. Both samples given to a subject were the same wine, although the subjects were expecting to taste two different samples of a particular variety. Of the 32 subjects in the study, 22 selected the wine presented first when presented with two identical wine samples.²⁸

- Give a 95% confidence interval for the proportion of subjects who would select the first choice presented.
- The subjects were recruited in Ontario, Canada, via advertisements to participate in a study of “attitudes and values towards wine.” What assumption are you making about these subjects?

20.41 Opinions about evolution, continued. Does the sample in Exercise 20.39 give good evidence to support the claim “Fewer than half of American adults think that humans developed from earlier species of animals”?

20.42 Order in choice, continued. Do the data in Exercise 20.40 give good reason to conclude that the subjects are not equally likely to choose either of the two wines when presented with two identical wine samples in sequence? Can we generalize our conclusions to all wine tasters? Explain.

20.43 Chick-fil-A gets it right. Which fast-food chain fills orders most accurately at the drive-thru window? The Quick Service Restaurant (QSR) magazine drive-thru study involved a total of 7594 visits to restaurants in the



© Flo Minton

25 largest fast-food chains in all 50 states. All visits occurred during the lunch hours of 11:00 A.M. to 2:30 P.M. or during the dinner hours of 4:00 to 7:00 P.M. During each visit, the researcher ordered a main item, a side item, and a drink. One item was left off of each order; for example, a field researcher could order a burger with no pickles. After receiving the order, all food and drink items were checked for complete accuracy. Any food or drink item received that was not exactly as ordered resulted in the order being classified as inaccurate. Also included in the measurement of accuracy were condiments asked for, napkins, straws, and correct change. Any errors in these resulted in the order being classified as inaccurate. Chick-fil-A had the fewest inaccuracies, with only 14 of 196 orders classified as inaccurate.²⁹ What proportion of orders are filled *accurately* by Chick-fil-A? (Use 95% confidence.)

20.44 Order in choice: planning a study. How large a sample would be needed to obtain margin of error ± 0.05 in the study of choice order for tasting wine? Use the \hat{p} from Exercise 20.40 as your guess for the unknown p .



EXPLORING THE WEB

20.45 Health care access/coverage. The Behavioral Risk Factor Surveillance System (BRFSS) is an ongoing data collection program designed to measure behavioral risk factors for the adult population (18 years of age or older) living in households. Data are collected from a random sample of adults (one per household) through a telephone survey. Go to the Web site apps.nccd.cdc.gov/BRFSS/ and under “Category” go to “Health Care Access/Coverage.” Under the topic “Adults aged 18–64 who have any kind of health care coverage,” you will find the percent with coverage in each state.

- Which state has the highest percent of coverage, and what is the reported value? Which state has the lowest percent, and what is its value? Are the reported percents statistics or parameters?
- Choose a state of interest to you and click on the link. In the table that opens, there is a line for n , and the entries are the numbers who answered “Yes” and “No.” Find the percent in the sample who answered “Yes.” Notice that it is different from the percent reported in the table. The table estimates are weighted to try to reduce bias. If it is determined that certain portions of the population are underrepresented in the sample, then

that portion of the sample receives more weight when computing the estimate of the percent. The assumptions for an SRS are rarely met in practice, and more complicated methods are often necessary to estimate proportions and compute confidence intervals.

20.46 Find a poll. Search the Web for a recent poll in which the sample statistic is a proportion, for example, the proportion in the sample responding “Yes” to a question. Calculate a 95% confidence interval for the population proportion (assume that the sample is a random sample). State the question asked, how the sample was collected, the sample size, and the population of interest. Possible Web sites are www.gallup.com and www.cbsnews.com/sections/opinion/polls/main500160.shtml.



Comparing Two Proportions

In a two-sample problem, we want to compare two populations or the responses to two treatments based on two independent samples. When the comparison involves the *means* of two populations, we use the two-sample *t* methods of Chapter 19. Now we turn to methods to compare the *proportions* of successes in two populations.

TWO-SAMPLE PROBLEMS: PROPORTIONS

We will use notation similar to that used in our study of two-sample *t* statistics. The groups we want to compare are Population 1 and Population 2. We have a separate SRS from each population or responses from two treatments in a randomized comparative experiment. A subscript shows which group a parameter or statistic describes. Here is our notation:

Population	Population proportion	Sample size	Sample proportion
1	p_1	n_1	\hat{p}_1
2	p_2	n_2	\hat{p}_2

We compare the populations by doing inference about the difference $p_1 - p_2$ between the population proportions. The statistic that estimates this difference is the difference between the two sample proportions, $\hat{p}_1 - \hat{p}_2$.

EXAMPLE 21.1 Young adults living with their parents

STATE: A surprising number of young adults (ages 19 to 25) still live in their parents' home. A random sample by the National Institutes of Health included 2253 men and 2629 women in this age group.¹ The survey found that 986 of the

IN THIS CHAPTER WE COVER...

- Two-sample problems: proportions
- The sampling distribution of a difference between proportions
- Large-sample confidence intervals for comparing proportions
- Using technology
- Accurate confidence intervals for comparing proportions
- Significance tests for comparing proportions



men and 923 of the women lived with their parents. Is this good evidence that different proportions of young men and young women live with their parents? How large is the difference between the proportions of young men and young women who live with their parents?

PLAN: Take young men to be Population 1 and young women to be Population 2. The population proportions who live in their parents' home are p_1 for men and p_2 for women. We want to test the hypotheses

$$H_0: p_1 = p_2 \quad (\text{the same as } H_0: p_1 - p_2 = 0)$$

$$H_a: p_1 \neq p_2 \quad (\text{the same as } H_a: p_1 - p_2 \neq 0)$$

We also want to give a confidence interval for the difference $p_1 - p_2$.

SOLVE: Inference about population proportions is based on the sample proportions

$$\hat{p}_1 = \frac{986}{2253} = 0.4376 \quad (\text{men})$$

$$\hat{p}_2 = \frac{923}{2629} = 0.3511 \quad (\text{women})$$

We see that about 44% of the men but only about 35% of the women lived with their parents. Because the samples are large and the sample proportions are quite different, we expect that a test will be highly significant (in fact, $P < 0.0001$). So we concentrate on the confidence interval. To estimate $p_1 - p_2$, start from the difference between sample proportions

$$\hat{p}_1 - \hat{p}_2 = 0.4376 - 0.3511 = 0.0865$$

To complete the “Solve” step, we must know how this difference behaves. ■

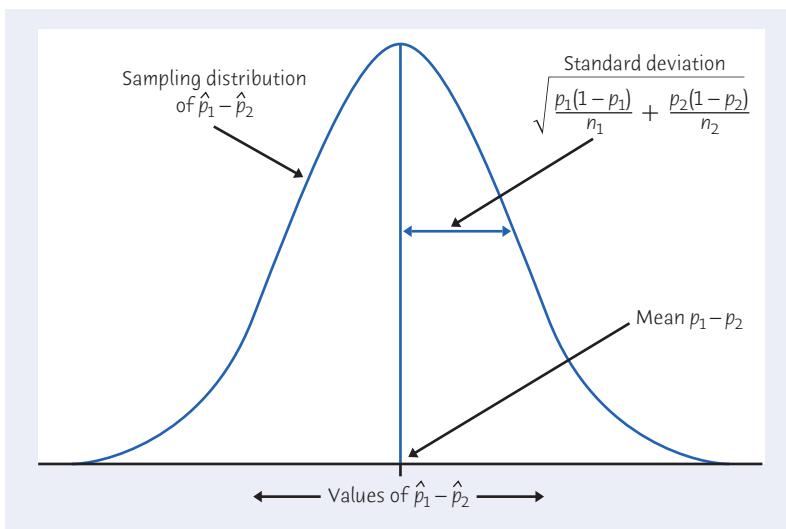
THE SAMPLING DISTRIBUTION OF A DIFFERENCE BETWEEN PROPORTIONS

To use $\hat{p}_1 - \hat{p}_2$ for inference, we must know its sampling distribution. Here are the facts we need:

- When the samples are large, the distribution of $\hat{p}_1 - \hat{p}_2$ is **approximately Normal**.
- The **mean** of the sampling distribution is $p_1 - p_2$. That is, the difference between sample proportions is an unbiased estimator of the difference between population proportions.
- The **standard deviation** of the distribution is

$$\sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}$$

Figure 21.1 displays the distribution of $\hat{p}_1 - \hat{p}_2$. The standard deviation of $\hat{p}_1 - \hat{p}_2$ involves the unknown parameters p_1 and p_2 . Just as in the previous chapter, we must replace these by estimates in order to do inference. And just as in the previous chapter, we do this a bit differently for confidence intervals and for tests.

**FIGURE 21.1**

Select independent SRSs from two populations having proportions of successes p_1 and p_2 . The proportions of successes in the two samples are \hat{p}_1 and \hat{p}_2 . When the samples are large, the sampling distribution of the difference $\hat{p}_1 - \hat{p}_2$ is approximately Normal.

LARGE-SAMPLE CONFIDENCE INTERVALS FOR COMPARING PROPORTIONS

To obtain a confidence interval, replace the population proportions p_1 and p_2 in the standard deviation by the sample proportions. The result is the **standard error** of the statistic $\hat{p}_1 - \hat{p}_2$:

$$SE = \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

The confidence interval has the same form we met in the previous chapter:

$$\text{estimate} \pm z^*SE_{\text{estimate}}$$

LARGE-SAMPLE CONFIDENCE INTERVAL FOR COMPARING TWO PROPORTIONS

Draw an SRS of size n_1 from a large population having proportion p_1 of successes and draw an independent SRS of size n_2 from another large population having proportion p_2 of successes. When n_1 and n_2 are large, an approximate level C **confidence interval for $p_1 - p_2$** is

$$(\hat{p}_1 - \hat{p}_2) \pm z^*SE$$

In this formula the standard error SE of $\hat{p}_1 - \hat{p}_2$ is

$$SE = \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

and z^* is the critical value for the standard Normal density curve with area C between $-z^*$ and z^* .

Use this interval only when the numbers of successes and failures are each 10 or more in both samples.

standard error



EXAMPLE 21.2 Living with parents: men versus women

We can now complete Example 21.1. Here is a summary of the basic information:

Population	Population description	Sample size	Number of successes	Sample proportion
1	men	$n_1 = 2253$	986	$\hat{p}_1 = 986/2253 = 0.4376$
2	women	$n_2 = 2629$	923	$\hat{p}_2 = 923/2629 = 0.3511$

SOLVE: We will give a 95% confidence interval for $p_1 - p_2$, the difference between the proportions of young men and young women who live with their parents. To check that the large-sample confidence interval is safe to use, look at the counts of successes and failures in the two samples. All these four counts are much larger than 10, so the large-sample method will be accurate. The standard error is

$$\begin{aligned} \text{SE} &= \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}} \\ &= \sqrt{\frac{(0.4376)(0.5624)}{2253} + \frac{(0.3511)(0.6489)}{2629}} \\ &= \sqrt{0.0001959} = 0.01400 \end{aligned}$$

The 95% confidence interval is

$$\begin{aligned} (\hat{p}_1 - \hat{p}_2) \pm z^* \text{SE} &= (0.4376 - 0.3511) \pm (1.960)(0.01400) \\ &= 0.0865 \pm 0.0274 \\ &= 0.059 \text{ to } 0.114 \end{aligned}$$

CONCLUDE: We are 95% confident that the percent of young men living with their parents is between 5.9 and 11.4 percentage points higher than the percent of young women who live with their parents. ■

The sample survey in this example selected a single random sample of young adults, not two separate random samples of young men and young women. To get two samples, we divided the single sample by sex. This means that we did not know the two sample sizes n_1 and n_2 until after the data were in hand. The two-sample z procedures for comparing proportions are valid in such situations. This is an important fact about these methods.

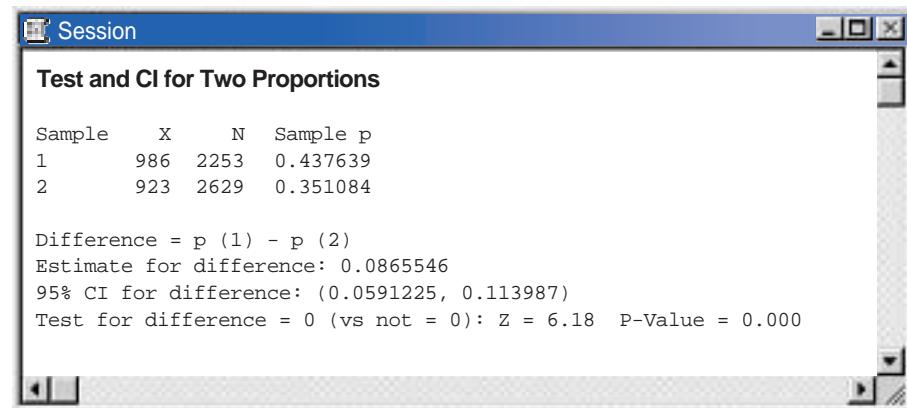
USING TECHNOLOGY

Figure 21.2 displays software output for Example 21.2 from a graphing calculator and two statistical software programs. As usual, you can understand the output even without knowledge of the program that produced it. Minitab gives the test as well as the confidence interval, confirming that the difference between men and women is highly significant. In CrunchIt!, the test and the confidence interval must be requested using separate commands, resulting in the two outputs in the figure. Excel spreadsheet output is not shown because Excel lacks menu items for inference about proportions. You must use the spreadsheet's formula capability to program the confidence interval or test statistic and then to find the P -value of a test.

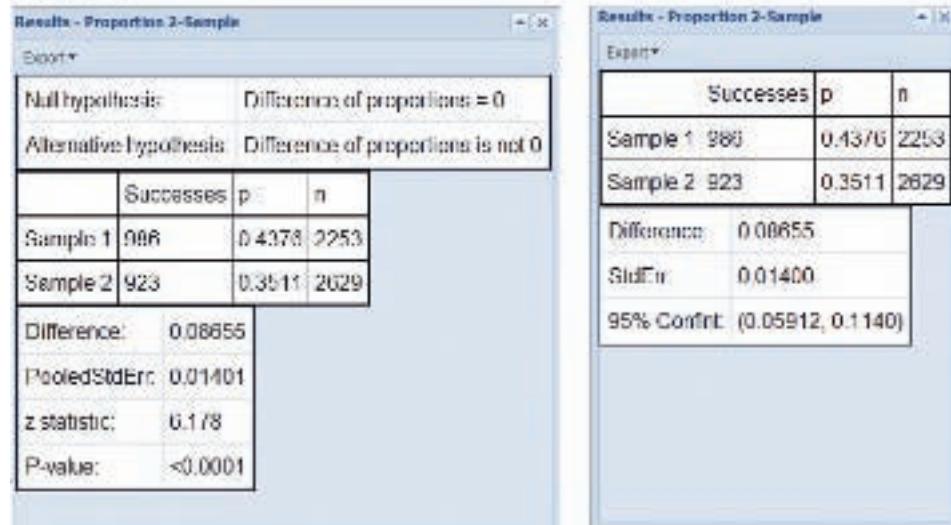
Texas Instruments Graphing Calculator

```
2-PropZInt
(.05912, .11399)
P1=.437638704
P2=.3510840624
n1=2253
n2=2629
```

Minitab



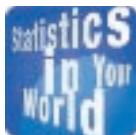
CrunchIt!

**FIGURE 21.2**

Output from a graphing calculator, Minitab, and CrunchIt! for the 95% confidence interval of Example 21.2.



Patrik Giardino/Photolibrary



Computer-assisted interviewing

The days of the interviewer with a clipboard are past. Interviewers now read questions from a computer screen and use the keyboard to enter responses. The computer skips irrelevant items—once a woman says that she has no children, further questions about her children never appear. The computer can even present questions in random order to avoid bias due to always following the same order. Software keeps records of who has responded and prepares a file of data from the responses. The tedious process of transferring responses from paper to computer, once a source of errors, has disappeared.

APPLY YOUR KNOWLEDGE

21.1 Who texts? Younger people use their cell phones to text more often than older people do. A random sample of 625 teens aged 12 to 17 who use their cell phones to text found that 475 sent more than 10 text messages in a typical day. In a random sample of 1917 adults aged 18 and over who use their cell phones to text, 786 sent more than 10 text messages in a single day.² Give a 95% confidence interval for the difference between the proportions of cell phone users who send more than 10 texts a day for these two age groups. Follow the four-step process as illustrated in Examples 21.1 and 21.2.

21.2 An issue of free speech. In 2008, respondents to the General Social Survey were asked: “There are always some people whose ideas are considered bad or dangerous by other people. For instance, somebody who is against churches and religion. If such a person wanted to make a speech in your (city/town/community) against churches and religion, should he be allowed to speak, or not?”³ Among the 464 respondents who considered themselves Democrats, 362 said, “Allow,” while among the 340 respondents who considered themselves Republicans, 254 said, “Allow.” Give a 95% confidence interval for the difference between the proportions of respondents from the two political parties who would allow such a person to speak. Follow the four-step process as illustrated in Examples 21.1 and 21.2.

21.3 High school students in action. A government survey randomly selected 8164 female high school students and 7881 male high school students.⁴ Of these students, 2261 females and 3594 males met recommended levels of physical activity. (These levels are quite high: at least 60 minutes of activity that makes you breathe hard on at least 5 of the past 7 days.) Give a 99% confidence interval for the difference between the proportions of all female and male high school students who meet the recommended levels of activity.

ACCURATE CONFIDENCE INTERVALS FOR COMPARING PROPORTIONS

 Like the large-sample confidence interval for a single proportion p , the large-sample interval for $p_1 - p_2$ generally has a true confidence level less than the level you asked for. The inaccuracy is not as serious as in the one-sample case, at least if our guidelines for use are followed. Once again, adding imaginary observations greatly improves the accuracy.⁵

PLUS FOUR CONFIDENCE INTERVAL FOR COMPARING TWO PROPORTIONS

Draw independent SRSs from two large populations with population proportions of successes p_1 and p_2 . To get the plus four confidence interval for the difference $p_1 - p_2$, add four imaginary observations, one success and one failure in each of the two samples. Then use the large-sample confidence interval with the new sample sizes (actual sample sizes + 2) and counts of successes (actual counts + 1).

Use this interval when the sample size is at least 5 in each group, with any counts of successes and failures.

If your software does not offer the plus four method, just enter the new plus four sample sizes and success counts into the large-sample procedure.

EXAMPLE 21.3 Shrubs that withstand fire



STATE: Fire is a serious threat to shrubs in dry climates. Some shrubs can resprout from their roots after their tops are destroyed. One study of resprouting took place in a dry area of Mexico.⁶ The investigators randomly assigned shrubs to treatment and control groups. They clipped the tops of all the shrubs. They then applied a propane torch to the stumps of the treatment group to simulate a fire. A shrub is a success if it resprouts. Here are the data for the shrub *Xerospirea hartwegiana*:

Population	Population description	Sample size	Number of successes	Sample proportion
1	control	$n_1 = 12$	12	$\hat{p}_1 = 12/12 = 1.000$
2	treatment	$n_2 = 12$	8	$\hat{p}_2 = 8/12 = 0.667$

How much does burning reduce the proportion of shrubs of this species that resprout?

PLAN: Give a 90% confidence interval for the difference between population proportions, $p_1 - p_2$.

SOLVE: The conditions for the large-sample interval are not met. In fact, there are no failures in the control group. We will use the plus four method. Add four imaginary observations. The new data summary is

Population	Population description	Sample size	Number of successes	Plus four sample proportion
1	control	$n_1 + 2 = 14$	$12 + 1 = 13$	$\tilde{p}_1 = 13/14 = 0.9286$
2	treatment	$n_2 + 2 = 14$	$8 + 1 = 9$	$\tilde{p}_2 = 9/14 = 0.6429$

The standard error based on the new facts is

$$\begin{aligned} \text{SE} &= \sqrt{\frac{\tilde{p}_1(1 - \tilde{p}_1)}{n_1 + 2} + \frac{\tilde{p}_2(1 - \tilde{p}_2)}{n_2 + 2}} \\ &= \sqrt{\frac{(0.9286)(0.0714)}{14} + \frac{(0.6429)(0.3571)}{14}} \\ &= \sqrt{0.02113} = 0.1454 \end{aligned}$$

The plus four 90% confidence interval is

$$\begin{aligned} (\tilde{p}_1 - \tilde{p}_2) \pm z^* \text{SE} &= (0.9286 - 0.6429) \pm (1.645)(0.1454) \\ &= 0.2857 \pm 0.2392 \\ &= 0.047 \text{ to } 0.525 \end{aligned}$$

CONCLUDE: We are 90% confident that burning reduces the percent of these shrubs that resprout by between 4.7% and 52.5%. ■

The plus four interval may be conservative (that is, the true confidence level may be *higher* than you asked for) for very small samples and population p 's close to 0 or 1,

as in this example. It is generally much more accurate than the large-sample interval when the samples are small. Nevertheless, the plus four interval in Example 21.3 cannot save us from the fact that small samples produce wide confidence intervals.



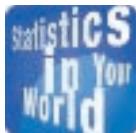
Jupiterimages/Comstock Images/Alamy

APPLY YOUR KNOWLEDGE

21.4 In-line skaters. A study of injuries to in-line skaters used data from the National Electronic Injury Surveillance System, which collects data from a random sample of hospital emergency rooms. The researchers interviewed 161 people who came to emergency rooms with injuries from in-line skating. Wrist injuries (mostly fractures) were the most common.⁷

- The interviews found that 53 people were wearing wrist guards and 6 of these had wrist injuries. Of the 108 who did not wear wrist guards, 45 had wrist injuries. Why should we not use the large-sample confidence interval for these data?
- The plus four method adds one success and one failure in each sample. What are the sample sizes and counts of successes after you do this?
- Give the plus four 95% confidence interval for the difference between the two population proportions of wrist injuries. State carefully what populations your inference compares.

21.5 Broken crackers. We don't like to find broken crackers when we open the package. How can makers reduce breaking? One idea is to microwave the crackers for 30 seconds right after baking them. Breaks start as hairline cracks called "checking." Assign 65 newly baked crackers to the microwave and another 65 to a control group that is not microwaved. After one day, none of the microwave group and 16 of the control group show checking.⁸ Give the 95% plus four confidence interval for the amount by which microwaving reduces the proportion of checking. The plus four method is particularly helpful when, as here, a count of successes is zero. Follow the four-step process as illustrated in Example 21.3.



The cookie strikes

How many different people clicked on your

business Web site last month? Technology tries to help: when someone visits your site, a little piece of code called a cookie is left on their computer. When the same person clicks again, the cookie says not to count them as a "unique visitor" because this isn't their first visit. But lots of Web users delete cookies, either by hand or automatically with software. These people get counted again when they visit your site again. That's bias: your counts of unique visitors are systematically too high. One study found that unique-visitor counts were as much as 50% too high.

SIGNIFICANCE TESTS FOR COMPARING PROPORTIONS

An observed difference between two sample proportions can reflect an actual difference between the populations, or it may just be due to chance variation in random sampling. Significance tests help us decide if the effect we see in the samples is really there in the populations. The null hypothesis says that there is no difference between the two populations:

$$H_0: p_1 = p_2$$

The alternative hypothesis says what kind of difference we expect.

EXAMPLE 21.4 Interracial dating

STATE: "Would you date a person of a different race?" Researchers answered this question for black males and females by collecting data from the Internet dating site Match.com. When people post profiles on the site, they indicate which races they

are willing to date. A random sample of 100 black males and a random sample of 100 black females were selected from the dating site, with 75 of the black males indicating their willingness to date white females and 56 of the black females indicating their willingness to date white males.⁹ Is there reason to think that different proportions of black males and females on this Internet dating site would be willing to date whites?



PLAN: Call the population proportions p_1 for men and p_2 for women. We had no direction for the difference in mind before looking at the data, so we have a two-sided alternative:

$$\begin{aligned} H_0: p_1 &= p_2 \\ H_a: p_1 &\neq p_2 \end{aligned}$$

SOLVE: The men and women can be considered separate SRSs of black men and women from the Internet dating site Match.com. The sample proportions who would be willing to date whites are

$$\begin{aligned} \hat{p}_1 &= \frac{75}{100} = 0.75 \quad (\text{men}) \\ \hat{p}_2 &= \frac{56}{100} = 0.56 \quad (\text{women}) \end{aligned}$$

That is, 75% of the men but only 56% of the women would be willing to date whites. Is this apparent difference statistically significant? To continue the solution, we must learn the proper test. ■

To do a test, standardize the difference between the sample proportions $\hat{p}_1 - \hat{p}_2$ to get a z statistic. If H_0 is true, both samples come from populations in which the same unknown proportion p would be willing to date whites. We take advantage of this by combining the two samples to estimate this single p instead of estimating p_1 and p_2 separately. Call this the **pooled sample proportion**. It is

pooled sample proportion

$$\hat{p} = \frac{\text{number of successes in both samples combined}}{\text{number of individuals in both samples combined}}$$

Use \hat{p} in place of both \hat{p}_1 and \hat{p}_2 in the expression for the standard error SE of $\hat{p}_1 - \hat{p}_2$ to get a z statistic that has the standard Normal distribution when H_0 is true. Here is the test.

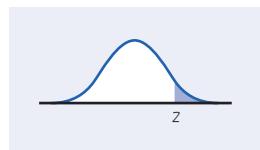
SIGNIFICANCE TEST FOR COMPARING TWO PROPORTIONS

Draw an SRS of size n_1 from a large population having proportion p_1 of successes and draw an independent SRS of size n_2 from another large population having proportion p_2 of successes. To **test the hypothesis $H_0: p_1 = p_2$** , first find the pooled proportion \hat{p} of successes in both samples combined. Then compute the z statistic

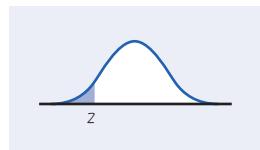
$$z = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

In terms of a variable Z having the standard Normal distribution, the P -value for a test of H_0 against

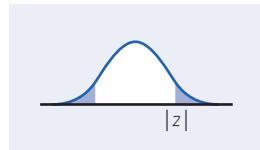
$$H_a: p_1 > p_2 \text{ is } P(Z \geq z)$$



$$H_a: p_1 < p_2 \text{ is } P(Z \leq z)$$



$$H_a: p_1 \neq p_2 \text{ is } 2P(|Z| \geq |z|)$$



Use this test when the counts of successes and failures are each 5 or more in both samples.¹⁰



EXAMPLE 21.5 Interracial dating, continued

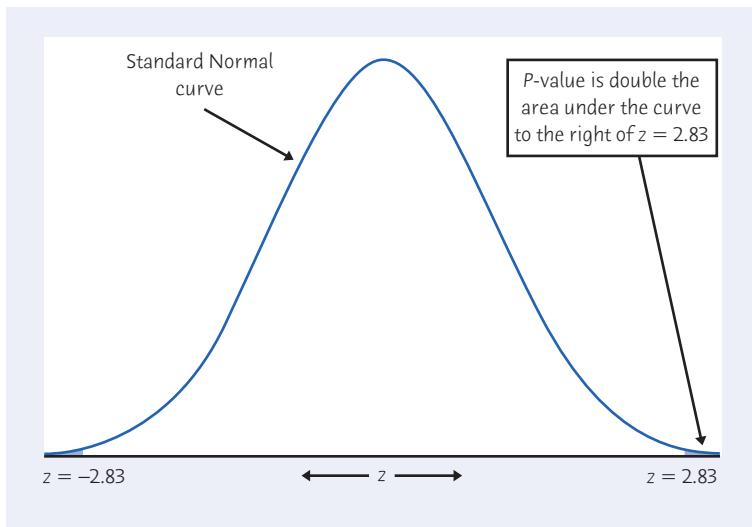
SOLVE: The data come from an SRS and the counts of successes and failures are all much larger than 5. The pooled proportion of blacks who would date whites is

$$\begin{aligned}\hat{p} &= \frac{\text{number "willing to date whites" among men and women combined}}{\text{number of men and women combined}} \\ &= \frac{75 + 56}{100 + 100} \\ &= \frac{131}{200} = 0.655\end{aligned}$$

The z test statistic is

$$\begin{aligned}z &= \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}} \\ &= \frac{0.75 - 0.56}{\sqrt{(0.655)(0.345)\left(\frac{1}{100} + \frac{1}{100}\right)}} \\ &= \frac{0.19}{0.06723} = 2.83\end{aligned}$$

The two-sided P -value is the area under the standard Normal curve more than 2.83 distant from 0. Figure 21.3 shows this area. Software tells us that $P = 0.0046$.

**FIGURE 21.3**

The P -value for the two-side test of Example 21.5.

Without software, you can compare $z = 2.83$ with the bottom row of Table C (standard Normal critical values) to approximate P . It lies between the critical values 2.807 and 3.091 for two-sided P -values 0.005 and 0.002.

CONCLUDE: There is strong evidence ($P < 0.005$) that black men are more likely than black women to be willing to date whites on comparable Internet dating sites. In a similar study by the authors, it was found that white men were more willing than white women to date blacks. ■

z^*	2.807	3.091
Two-sided P	.005	.002

APPLY YOUR KNOWLEDGE

- 21.6 Seat belt use.** The proportion of drivers who use seat belts depends on things like age (young people are more likely to go unbelted) and sex (women are more likely to use belts). It also depends on local law. In New York City, police can stop a driver who is not belted. In Boston at the time of the survey, police could cite a driver for not wearing a seat belt only if the driver had been stopped for some other violation. Here are data from observing random samples of female Hispanic drivers in these two cities in 2002:¹¹



City	Drivers	Belted
New York	220	183
Boston	117	68

- (a) Is this an experiment or an observational study? Why?
- (b) Comparing local laws suggests the hypothesis that a smaller proportion of drivers wear seat belts in Boston than in New York. Do the data give good evidence that this is true for female Hispanic drivers? Follow the four-step process as illustrated in Examples 21.4 and 21.5.
- 21.7 Protecting skiers and snowboarders.** Most alpine skiers and snowboarders do not use helmets. Do helmets reduce the risk of head injuries? A study in Norway





Jupiterimages/Age fotostock



compared skiers and snowboarders who suffered head injuries with a control group who were not injured. Of 578 injured subjects, 96 had worn a helmet. Of the 2992 in the control group, 656 wore helmets.¹² Is helmet use less common among skiers and snowboarders who have head injuries? Follow the four-step process as illustrated in Examples 21.4 and 21.5. (Note that this is an observational study that compares injured and uninjured subjects. An experiment that assigned subjects to helmet and no-helmet groups would be more convincing.)

21.8 Breast cancer treatment. In sentinel lymph node dissection (SLND), surgeons remove two or three lymph nodes close to the breast that are most likely to contain cancer cells. If these “sentinel” lymph nodes are free of tumors, SLND alone is the accepted management for patients. When the sentinel lymph nodes contain metastases, axillary lymph node dissection (ALND; that is, the removal of further nodes) remains the standard of care, although its contribution to survival is controversial. ALND carries the additional risk of complications such as seroma, infection, and lymphedema. In one study, patients with sentinel metastases identified by SLND were randomized to undergo ALND or no further treatment (SLND alone). Here are the five-year disease-free survival numbers for the two groups:¹³

Group	Sample size	Disease-free after five years
ALND	420	345
SLND alone	436	366

How strong is the evidence that the proportions of disease-free patients after five years differ for patients who underwent ALND or only SLND? Follow the four-step process as illustrated in Examples 21.4 and 21.5.

CHAPTER 21 SUMMARY

CHAPTER SPECIFICS

- The data in a **two-sample** problem are two independent SRSs, each drawn from a separate population.
- Tests and confidence intervals to compare the proportions p_1 and p_2 of successes in the two populations are based on the difference $\hat{p}_1 - \hat{p}_2$ between the sample proportions of successes in the two SRSs.
- When the sample sizes n_1 and n_2 are large, the sampling distribution of $\hat{p}_1 - \hat{p}_2$ is close to Normal with mean $p_1 - p_2$.
- The level C **large-sample confidence interval** for $p_1 - p_2$ is

$$(\hat{p}_1 - \hat{p}_2) \pm z^* \text{SE}$$

where the **standard error** of $\hat{p}_1 - \hat{p}_2$ is

$$\text{SE} = \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}$$

and z^* is a standard Normal critical value.

- The true confidence level of the large-sample interval can be substantially less than the planned level C . Use this interval only if the counts of successes and failures in both samples are 10 or greater.
- To get a more accurate confidence interval, add four imaginary observations, one success and one failure in each sample. Then use the same formula for the confidence interval. This is the **plus four confidence interval**. You can use it whenever both samples have 5 or more observations.
- Significance tests for $H_0: p_1 = p_2$ use the pooled sample proportion

$$\hat{p} = \frac{\text{number of successes in both samples combined}}{\text{number of individuals in both samples combined}}$$

and the z statistic

$$z = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1 - \hat{p}) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}}$$

P -values come from the standard Normal distribution. Use this test when there are 5 or more successes and 5 or more failures in both samples.

LINK IT

Most studies compare two or more treatments rather than investigating a single treatment. These can be observational studies or comparative experiments, and the differences in the type of conclusions that can be reached were described in Chapter 9. When there are two treatments and the response is a continuous measurement, the comparison between the treatments is often based on a comparison of the treatment means using the methods of Chapter 19.

This chapter considers the case where the response classifies individuals into two categories such as young adults living with their parents or not. The comparison is then a comparison of this proportion for two groups, such as the proportion of males living with their parents versus the proportion of females living with their parents. Since the methods of this chapter are approximate, it is important to always check the conditions required for the approximations to work well, whether you are using the large-sample confidence interval, the plus four method, or the z test.

After the appropriate statistical method has been applied, care must be taken when stating the conclusion. The issues described in Chapter 9 are still important. For example, a lack of blinding can result in the expectations of the researcher influencing the results, while confounding can mix up the comparison of the two groups with other factors.

CHECK YOUR SKILLS

In the past decade there have been intensive antismoking campaigns sponsored by both federal and private agencies. The Behavioral Risk Factor Surveillance System (BRFSS) is an ongoing data collection program that monitors behaviors such as smoking on a statewide level using data collected on a random sample of adults through a telephone survey.¹⁴ The first sample, taken in

1999 in Alaska, involved 2045 adults, of which 592 were current smokers. The second sample, taken in 2009 in Alaska, involved 2411 adults, of which 511 were smokers. The samples are to be compared to determine whether the proportion of U.S. adults in Alaska that smoke declined during the 10-year period between the samples. Exercises 21.9 to 21.13 are based on these surveys.

21.9 Take p_{1999} and p_{2009} to be the proportions of all adults in Alaska over 18 who were current smokers in these years. The hypotheses to be tested are

- (a) $H_0: p_{1999} = p_{2009}$ versus $H_a: p_{1999} \neq p_{2009}$.
- (b) $H_0: p_{1999} = p_{2009}$ versus $H_a: p_{1999} > p_{2009}$.
- (c) $H_0: p_{1999} = p_{2009}$ versus $H_a: p_{1999} < p_{2009}$.

21.10 The sample proportions of adults who were current smokers in 1999 and 2009 are about

- (a) $\hat{p}_{1999} = 0.21$ and $\hat{p}_{2009} = 0.29$.
- (b) $\hat{p}_{1999} = 0.29$ and $\hat{p}_{2009} = 0.21$.
- (c) $\hat{p}_{1999} = 0.21$ and $\hat{p}_{2009} = 0.25$.

21.11 The pooled sample proportion of adult smokers in Alaska is about

- (a) $\hat{p} = 0.23$.
- (b) $\hat{p} = 0.25$.
- (c) $\hat{p} = 0.28$.

21.12 The z test for comparing the proportions of smokers in Alaska in 1999 and 2009 has

- (a) $P < 0.01$.
- (b) $0.01 < P < 0.05$.
- (c) $P > 0.05$.

21.13 The 95% large-sample confidence interval for the difference $p_{1999} - p_{2009}$ in the proportions of smokers in Alaska in 1999 and 2009 is about

- (a) 0.08 ± 0.013 .
- (b) 0.08 ± 0.021 .
- (c) 0.08 ± 0.026 .

21.14 In an experiment to learn if substance M can help restore memory, the brains of 20 rats were treated to damage their memories. The rats were trained to run a maze. After a day, 10 rats were given M and 7 of them succeeded in the maze; only 2 of the 10 control rats were successful. The z test for “no difference” against “a higher proportion of the M group succeeds” has

- (a) $z = 2.25, P < 0.02$.
- (b) $z = 2.60, P < 0.005$.
- (c) $z = 2.25, 0.02 < P < 0.04$.

21.15 The z test in the previous exercise

- (a) may be inaccurate because the populations are too small.
- (b) may be inaccurate because some counts of successes and failures are too small.
- (c) is reasonably accurate because the conditions for inference are met.

21.16 The plus four 90% confidence interval for the difference between the proportion of rats that succeed when given M and the proportion that succeed without it is

- (a) 0.455 ± 0.312 .
- (b) 0.417 ± 0.304 .
- (c) 0.417 ± 0.185 .

CHAPTER 21 EXERCISES

We recommend using the plus four method for all confidence intervals for proportions. However, the large-sample method is acceptable when the guidelines for its use are met.

21.17 Truthfulness in online profiles. Many teens have posted profiles on a social-networking Web site. A sample survey asked random samples of teens with online profiles if they included false information in their profiles. Of 170 younger teens (ages 12 to 14), 117 said “Yes.” Of 317 older teens (ages 15 to 17), 152 said “Yes.”¹⁵

- (a) Do these samples satisfy the guidelines for the large-sample confidence interval?
- (b) Give a 95% confidence interval for the difference between the proportions of younger and older teens who include false information in their online profiles.

21.18 Effects of an appetite suppressant. Subjects with preexisting cardiovascular symptoms who were receiving subitramine, an appetite suppressant, were found to be at increased risk of cardiovascular events while taking the drug. The study included 9804 overweight or obese subjects with preexisting cardiovascular disease and/or type 2 diabetes. The subjects were randomly assigned to subitramine (4906 subjects)

or a placebo (4898 subjects) in a double-blind fashion. The primary outcome measured was the occurrence of any of the following events: nonfatal myocardial infarction or stroke, resuscitation after cardiac arrest, or cardiovascular death. The primary outcome was observed in 561 subjects in the subitramine group and 490 subjects in the placebo group.¹⁶

- (a) Find the proportion of subjects experiencing the primary outcome for both the subitramine and placebo groups.
- (b) Can we safely use the large-sample confidence interval for comparing the proportions of subitramine and placebo subjects who experienced the primary outcome? Explain.
- (c) Give a 95% confidence interval for the difference between the proportions of subitramine and placebo subjects who experienced the primary outcome.

21.19 Genetically altered mice. Genetic influences on cancer can be studied by manipulating the genetic makeup of mice. One of the processes that turn genes on or off (so to speak) in



Mark Harmel/Getty Images

particular locations is called “DNA methylation.” Do low levels of this process help cause tumors? Compare mice altered to have low levels with normal mice. Of 33 mice with lowered levels of DNA methylation, 23 developed tumors. None of the control group of 18 normal mice developed tumors in the same time period.¹⁷

- (a) Explain why we cannot safely use either the large-sample confidence interval or the test for comparing the proportions of normal and altered mice that develop tumors.
- (b) The plus four method adds two observations, a success and a failure, to each sample. What are the sample sizes and the numbers of mice with tumors after you do this?
- (c) Give a 99% confidence interval for the difference in the proportions of the two populations that develop tumors.

21.20 Effects of an appetite suppressant, continued.

Exercise 21.18 describes a study to determine if subjects with preexisting cardiovascular symptoms were at an increased risk of cardiovascular events while taking subitramine. Do the data give good reason to think that there is a difference between the proportions of treatment and placebo subjects who experienced the primary outcome? (Note that subitramine is no longer available in the United States due to the manufacturer’s concerns over increased risk of heart attack or stroke.)

- (a) State hypotheses, find the test statistic, and use either software or the bottom row of Table C for the P -value. Be sure to state your conclusion.
- (b) Explain simply why it was important to have a placebo group in this study.

Adolescence, music, and algebra. Research has suggested that musicians process music in the same cortical regions in which adolescents process algebra. When taking introductory algebra, will students who were enrolled in formal instrumental or choral music instruction during middle school outperform those who experienced neither of these modes of musical instruction? The sample consisted of 6026 ninth-grade students in Maryland who had completed introductory algebra. Of these, 3239 students had received formal instrumental or choral instruction during all three years of middle school, while the remaining students had not. Of those receiving formal musical instruction, 2818 received a passing grade on the Maryland Algebra/Data Analysis High School Assessment (HSA). In contrast, 2091 of the 2787 students not receiving musical instruction received a passing grade.¹⁸ Exercises 21.21 to 21.23 are based on this study.

21.21 Does music make a difference?

- (a) Is there a significant difference in the proportions of students with and without musical instruction who receive a passing grade on the Maryland HSA? State hypotheses, find

the test statistic, and use software or the bottom row of Table C to get a P -value.

- (b) Is this an observational study or an experiment? Why?
- (c) In view of your answer in (b), carefully state your conclusions about the relationship between music instruction and success in algebra.

21.22 How many students pass? Give a 95% confidence interval for the proportion of ninth-grade students who receive a passing grade on the HSA.

21.23 How big a difference? Give a 95% confidence interval for the difference between the proportions of students passing the HSA who have received or not received formal musical instruction in middle school.

21.24 The design of the study matters. How accurate are the tests that grain-handling facilities make to detect the presence of genetically modified (GM) soybeans in shipments to countries that do not allow GM beans? Batches of soybeans containing some genetically modified (GM) beans were submitted to 23 grain-handling facilities. When batches contained 1% of GM beans, 18 of the facilities detected the presence of GM beans. Only 7 of the facilities detected GM beans when they made up one-tenth of 1% of the beans in the batches.¹⁹ Explain why we cannot use the methods of this chapter to compare the proportions of facilities that will detect the two levels of GM soybeans.

21.25 Significant does not mean important. Never forget that even small effects can be statistically significant if the samples are large. To illustrate this fact, consider a sample of 148 small businesses. During a three-year period, 15 of the 106 headed by men and 7 of the 42 headed by women failed.²⁰

(a) Find the proportions of failures for businesses headed by women and businesses headed by men. These sample proportions are quite close to each other. Give the P -value for the z test of the hypothesis that the same proportion of women’s and men’s businesses fail. (Use the two-sided alternative.) The test is very far from being significant.

(b) Now suppose that the same sample proportions came from a sample 30 times as large. That is, 210 out of 1260 businesses headed by women and 450 out of 3180 businesses headed by men fail. Verify that the proportions of failures are exactly the same as in (a). Repeat the z test for the new data, and show that it is now significant at the $\alpha = 0.05$ level.

(c) It is wise to use a confidence interval to estimate the size of an effect rather than just giving a P -value. Give 95% confidence intervals for the difference between the proportions of women’s and men’s businesses that fail for the settings of

both (a) and (b). What is the effect of larger samples on the confidence interval?

In responding to Exercises 21.26 to 21.35, follow the **Plan**, **Solve**, and **Conclude** steps of the four-step process.

21.26 Are urban students more successful? North Carolina State University looked at the factors that affect the success of students in a required chemical engineering course. Students must get a C or better in the course to continue as chemical engineering majors, so a “success” is a grade of C or better. There were 65 students from urban or suburban backgrounds, and 52 of these students succeeded. Another 55 students were from rural or small-town backgrounds; 30 of these students succeeded in the course.²¹ Is there good evidence that the proportion of students who succeed is different for urban/suburban versus rural/small-town backgrounds?

21.27 Female and male students. The North Carolina State University study in the previous exercise also looked at possible differences in the proportions of female and male students who succeeded in the course. They found that 23 of the 34 women and 60 of the 89 men succeeded. Is there evidence of a difference between the proportions of women and men who succeed?

21.28 More on urban and rural students. Continue your work from Exercise 21.26. Estimate the difference between the success rates for all urban/suburban and rural/small-town students who plan to study chemical engineering at North Carolina State. (Use 90% confidence.)

21.29 Smoking cessation. Chantix is different from most other quit-smoking products in that it targets nicotine receptors in the brain, attaches to them, and blocks nicotine from reaching them. As part of a larger randomized controlled trial, generally healthy smokers who smoked at least 10 cigarettes per day were assigned at random to take Chantix or a placebo. The study was double-blind, with the response measure being continuous absence from smoking for Weeks 9 through 12 of the study. Of the 352 subjects taking Chantix, 155 abstained from smoking during Weeks 9 through 12, while 61 of the 344 subjects taking the placebo abstained during this same time period.²² Give a 99% confidence interval for the difference (treatment minus placebo) in the proportions of smokers who abstained from smoking during Weeks 9 through 12.

21.30 The Gold Coast. A historian examining British colonial records for the Gold Coast in Africa suspects that the death rate was higher among African miners than among European miners. In the year 1936, there were 223 deaths among 33,809 African miners and 7 deaths among 1541 European miners on the Gold Coast.²³ (The Gold Coast became the independent nation of Ghana in 1957.) Consider this year as a random sample from the colonial era in West Africa. Is there good evidence that the proportion of African

miners who died was higher than the proportion of European miners who died?

21.31 I refuse! Do our emotions influence economic decisions?

One way to examine the issue is to have subjects play an “ultimatum game” against other people and against a computer. Your partner (person or computer) gets \$10, on the condition that it be shared with you. The partner makes you an offer. If you refuse, neither of you gets anything. So it’s to your advantage to accept even the unfair offer of \$2 out of the \$10. Some people get mad and refuse unfair offers. Here are data on the responses of 76 subjects randomly assigned to receive an offer of \$2 from either a person they were introduced to or a computer:²⁴

	Accept	Reject
Human offers	20	18
Computer offers	32	6

We suspect that emotion will lead to offers from another person being rejected more often than offers from an impersonal computer. Do a test to assess the evidence for this conjecture.

21.32 Did the random assignment work? A large clinical trial of the effect of diet on breast cancer assigned women at random to either a normal diet or a low-fat diet. To check that the random assignment did produce comparable groups, we can compare the two groups at the start of the study. Ask if there is a family history of breast cancer: 3396 of the 19,541 women in the low-fat group and 4929 of the 29,294 women in the control group said “Yes.”²⁵ If the random assignment worked well, there should not be a significant difference in the proportions with a family history of breast cancer. How significant is the observed difference?

21.33 Lyme disease. Lyme disease is spread in the northeastern United States by infected ticks. The ticks are infected mainly by feeding on mice, so more mice result in more infected ticks. The mouse population in turn rises and falls with the abundance of acorns, their favored food. Experimenters studied two similar forest areas in a year when the acorn crop failed. They added hundreds of thousands of acorns to one area to imitate an abundant acorn crop, while leaving the other area untouched. The next spring, 54 of the 72 mice trapped in the first area were in breeding condition, versus 10 of the 17 mice trapped in the second area.²⁶



Scott Camazine/Photo Researchers

Estimate the difference between the proportions of mice ready to breed in good acorn years and bad acorn years. (Use 90% confidence. Be sure to justify your choice of confidence interval.)

21.34 Does preschool help? To study the long-term effects of preschool programs for poor children, the High/Scope Educational Research Foundation has followed two groups of Michigan children since early childhood.²⁷ One group of 62 attended preschool as three- and four-year-olds. A control group of 61 children from the same area and similar backgrounds did not attend preschool. Over a 10-year period as adults, 38 of the preschool sample and 49 of the control sample needed social services (mainly welfare). Does the study provide significant evidence that children who attend preschool have less need for social services as adults? How large is the difference between the proportions of the preschool and no-preschool populations that require social services? Do inference to answer both questions. Be sure to explain exactly what inference you choose to do.

21.35 Hand sanitizers. Hand disinfection is frequently recommended for prevention of transmission of the rhinovirus that causes the common cold. In particular,

hand lotion containing 2% citric acid and 2% malic acid in 70% ethanol (HL+) has been found to have both immediate and persistent ability to inactivate rhinovirus (RV) on the hands in an experimental setting. Is hand disinfection effective in reducing the risk of infection in a natural setting? A total of 212 volunteers were assigned at random to either the HL+ group, which used the hand lotion every three hours or after hand washing, and a control group, which was asked to use routine hand washing but to avoid the use of alcohol-based hand sanitizers. Here are the data on the numbers of subjects with and without RV infection in the two groups over the 10-week study period.²⁸

RV Infection		
	Yes	No
HL+	49	67
Control group	49	47

- (a) Is this an experiment or an observational study? Why?
- (b) Do the data give good evidence that hand sanitizers reduce the chance of an RV infection?



EXPLORING THE WEB

21.36 Hearing loss in adolescents. Go to the *Journal of the American Medical Association* Web site, <http://jama.ama-assn.org/content/by/year>, and find the article “Change in Prevalence of Hearing Loss in US Adolescents” by Shargorodsky et al. in the August 18, 2010, issue. If you cannot get the full text of the article, use the information in the abstract plus the information given below to answer the questions. NHANES III is the earlier sample and NHANES 2005–2006 is the more recent sample.

- (a) Is this an observational study or an experiment?
- (b) How many people were in the earlier sample and how many were in the later sample?
- (c) If you do not have access to Table 2 of the full article, here are the facts you will need: in the earlier study 480 people experienced some hearing loss, while in the later study 333 people experienced some hearing loss. Is there evidence of an increase in hearing loss for children aged 12 to 19 in the later study? State hypotheses, find the test statistic, and use either software or Table A to compute the P-value. Although the article used a more sophisticated analysis, your P-value should be quite close to the P-value of 0.02 reported in the abstract.

21.37 Compare two surveys. Go to the Web site www.pollingreport.com, which contains the results of surveys conducted by several survey organizations. Choose a topic of interest to you, and then, to see if attitudes have changed over time, find two surveys that were conducted at two different times but that ask the same question. For example, you might choose the topic of abortion and compare the percents of people who feel abortion should always be legal at points in time separated by several years. State hypotheses to check for a difference over time, find the test statistic, and use either software or Table A to compute the P-value. What is your conclusion in context?



Inference about Variables: Part III Review

Chapter 22

IN THIS CHAPTER
WE COVER...

- Part III Summary
- Test Yourself
- Supplementary Exercises

The procedures of Chapters 18 to 21 are among the most common of all statistical inference methods. Now that you have mastered important ideas and practical methods for inference, it's time to review the big ideas of statistics in outline form. Here is a summary of Parts I and II of this book, leading up to Part III. The outline contains some important warnings: look for the Caution icon.

1. Data Production

- Data basics:
 - Individuals (subjects).
 - Variables: categorical versus quantitative, units of measurement, explanatory versus response.
 - Purpose of study.
- Data production basics:
 - Observation versus experiment.
 - Simple random samples.
 - Completely randomized experiments.
- Beware: really bad data production (voluntary response, confounding) can make interpretation impossible.
- Beware: weaknesses in data production (for example, sampling students at only one campus) can make generalizing conclusions difficult.



2. Data Analysis



- Plot your data. Look for an overall pattern and striking deviations.
- Add numerical descriptions based on what you see.
- Beware: averages and other simple descriptions can miss the real story.
- One quantitative variable:

Graphs: stemplot, histogram, boxplot.

Pattern: distribution shape, center, spread. Outliers?

Density curves (such as Normal curves) to describe overall pattern.

Numerical descriptions: five-number summary or \bar{x} and s .

- Relationships between two quantitative variables:

Graph: scatterplot.

Pattern: relationship form, direction, strength. Outliers? Influential observations?

Numerical description for linear relationships: correlation, regression line.

Beware the lurking variable: correlation does not imply causation.

- Beware the effects of outliers and influential observations.

3. The Reasoning of Inference



- Inference uses data to infer conclusions about a wider population.
- When you do inference, you are acting as if your data come from random samples or randomized comparative experiments. Beware: if they don't, you may have "garbage in, garbage out."
- Always examine your data before doing inference. Inference often requires a regular pattern, such as roughly Normal with no strong outliers.
- Key idea: "What would happen if we did this many times?"
- Confidence intervals: estimate a population parameter.

95% confidence: I used a method that captures the true parameter 95% of the time in repeated use.

Beware: the margin of error of a confidence interval does not include the effects of practical errors such as undercoverage and nonresponse.

- Significance tests: assess evidence against H_0 in favor of H_a .

P-value: If H_0 were true, how often would I get an outcome favoring the alternative this strongly? Smaller *P* = stronger evidence against H_0 .

Statistical significance at the 5% level, $P < 0.05$, means that an outcome this extreme would occur less than 5% of the time if H_0 were true.

Beware: $P < 0.05$ is not sacred.

Beware: statistical significance is not the same as practical significance.

Large samples can make small effects significant. Small samples can fail to declare large effects significant.

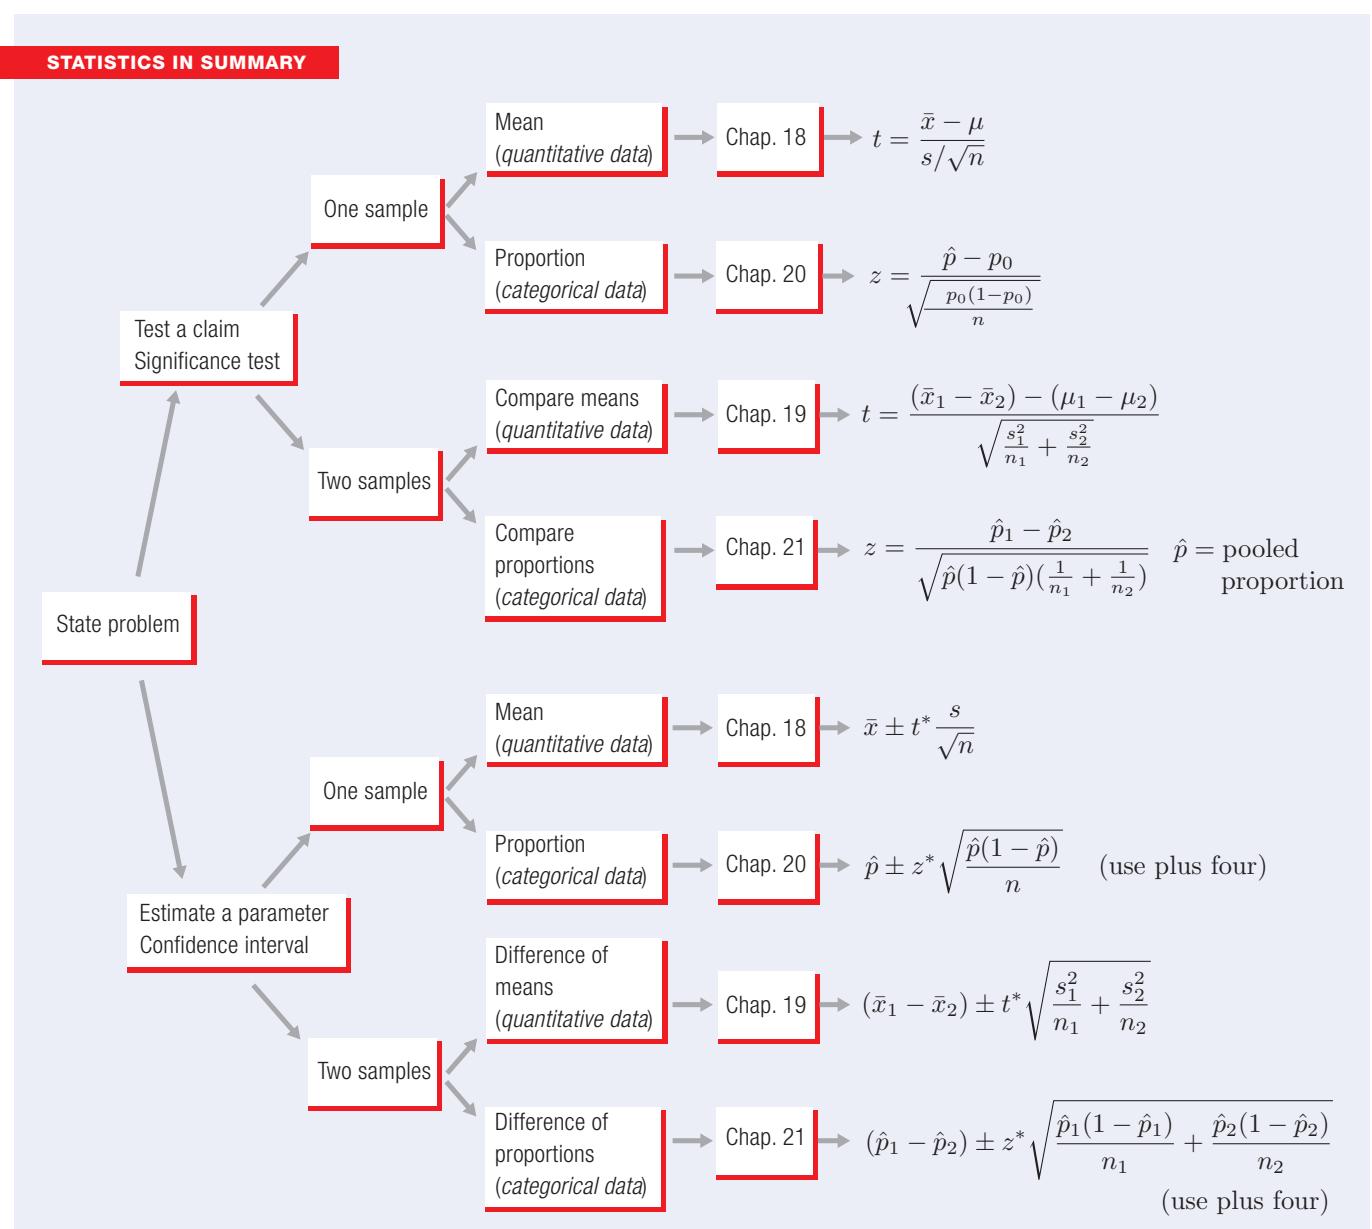
Always try to estimate the size of an effect (for example, with a confidence interval), not just its significance.

4. Methods of Inference



- Choose the right inference procedure.
- Carry out the details.
- State your conclusion.

Part III of this book introduces the fourth and last part of this outline. To actually do inference, you must choose the right procedure and carry out the details. The Statistics in Summary flowchart below offers a brief guide. It is important to do some of the review exercises because now, for the first time, you must decide which of several inference procedures to use. Learning to recognize problem settings in order to choose the right type of inference is a key step in advancing your mastery of statistics. This is the “Plan” step in the four-step process, in which you translate the real-world problem from the “State” step into a specific inference procedure.



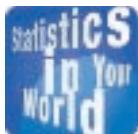
The flowchart organizes one way of planning inference problems. Let's go through it from left to right.

1. *Do you want to test a claim or estimate an unknown quantity?* That is, will you need a test of significance or a confidence interval?
2. *Are your data a single sample representing one population or two samples chosen to compare two populations or responses to two treatments in an experiment?* Remember that to work with *matched pairs* data you form one sample from the differences within pairs.
3. *Is the response variable quantitative or categorical?* Quantitative variables take numerical values with some unit of measurement such as inches or grams. The most common inference questions about quantitative variables concern *mean* responses. If the response variable is categorical, inference most often concerns the *proportion* of some category (call it a “success”) among the responses.

The flowchart leads you to a specific test or confidence interval, indicated by a formula at the end of each path. The formula is just an aid to guide you toward the “Solve” and “Conclude” steps. You (or your technology) will use the formula as part of the “Solve” step, but don’t forget that you must do more.

- *Are the conditions for this procedure met?* Can you act as if the data come from a random sample or randomized comparative experiment? Does data analysis show extreme outliers or strong skewness that forbid use of inference based on Normality? Do you have enough observations for your intended procedure?
- *Do your data come from an experiment or from an observational study?* The details of inference methods are the same for both. But the design of the study determines what conclusions you can reach, because experiments give much better evidence that an effect uncovered by inference can be explained by direct causation.

You may ask, as you study the Statistics in Summary flowchart, “What if I have an experiment comparing four treatments, or samples from three populations?” The flowchart allows only one or two, not three or four or more. Be patient: methods for comparing more than two means or proportions, as well as some other settings for inference, appear in Part IV.



How many was that?

Good causes often breed bad statistics. An

advocacy group claims, without much evidence, that 150,000 Americans suffer from the eating disorder anorexia nervosa. Soon someone misunderstands and says that 150,000 people die from anorexia nervosa each year. This wild number gets repeated in countless books and articles. It really is a wild number: only about 55,000 women aged 15 to 44 (the main group affected) die of all causes each year.

PART III SUMMARY

Here are the most important skills you should have acquired from reading Chapters 18 to 21.

A. Recognition

1. Recognize when a problem requires inference about population means (quantitative response variable) or population proportions (usually categorical response variable).

2. Recognize from the design of a study whether one-sample, matched pairs, or two-sample procedures are needed.
3. Based on recognizing the problem setting, choose among the one- and two-sample t procedures for means and the one- and two-sample z procedures for proportions.

B. Inference about One Mean

1. Verify that the t procedures are appropriate in a particular setting. Check the study design and the distribution of the data and take advantage of robustness against lack of Normality.
2. Recognize when poor study design, outliers, or a small sample from a skewed distribution make the t procedures risky.
3. Use the one-sample t procedure to obtain a confidence interval at a stated level of confidence for the mean μ of a population.
4. Carry out a one-sample t test for the hypothesis that a population mean μ has a specified value against either a one-sided or a two-sided alternative. Use software to find the P -value or Table C to get an approximate value.
5. Recognize matched pairs data and use the t procedures to obtain confidence intervals and to perform tests of significance for such data.

C. Comparing Two Means

1. Verify that the two-sample t procedures are appropriate in a particular setting. Check the study design and the distribution of the data and take advantage of robustness against lack of Normality.
2. Give a confidence interval for the difference between two means. Use software if you have it. Use the two-sample t statistic with conservative degrees of freedom and Table C if you do not have statistical software.
3. Test the hypothesis that two populations have equal means against either a one-sided or a two-sided alternative. Use software if you have it. Use the two-sample t test with conservative degrees of freedom and Table C if you do not have statistical software.
4. Know that procedures for comparing the standard deviations of two Normal populations are available, but that these procedures are risky because they are not at all robust against non-Normal distributions.

D. Inference about One Proportion

1. Verify that you can safely use either the large-sample or the plus four z procedures in a particular setting. Check the study design and the guidelines for sample size.
2. Use the large-sample z procedure to give a confidence interval for a population proportion p . Understand that the true confidence level may be substantially less than you ask for unless the sample is very large and the true p is not close to 0 or 1.

3. Use the plus four modification of the z procedure to give a confidence interval for p that is accurate even for small samples and for any value of p .
4. Use the z statistic to carry out a test of significance for the hypothesis $H_0: p = p_0$ about a population proportion p against either a one-sided or a two-sided alternative. Use software or Table A to find the P -value, or Table C to get an approximate value.

E. Comparing Two Proportions

1. Verify that you can safely use either the large-sample or the plus four z procedures in a particular setting. Check the study design and the guidelines for sample sizes.
2. Use the large-sample z procedure to give a confidence interval for the difference $p_1 - p_2$ between proportions in two populations based on independent samples from the populations. Understand that the true confidence level may be less than you ask for unless the samples are quite large.
3. Use the plus four modification of the z procedure to give a confidence interval for $p_1 - p_2$ that is accurate even for very small samples and for any values of p_1 and p_2 .
4. Use a z statistic to test the hypothesis $H_0: p_1 = p_2$ that proportions in two distinct populations are equal. Use software or Table A to find the P -value, or Table C to get an approximate value.

TEST YOURSELF

The questions below include both multiple-choice and short-answer questions and calculations. They will help you review the basic ideas and skills presented in Chapters 18 to 21.



Fotosearch Premium/Photolibrary

Calcium and blood pressure. In a randomized comparative experiment on the effect of dietary calcium on blood pressure, researchers divided 54 healthy white males at random into two groups. One group received calcium; the other, a placebo. At the beginning of the study, the researchers measured many variables on the subjects. The paper reporting the study gives $\bar{x} = 114.9$ and $s = 9.3$ for the seated systolic blood pressure of the 27 members of the placebo group. Use this information to answer Questions 22.1 and 22.2.

22.1 A 95% confidence interval for the mean blood pressure in the population from which the subjects were recruited is

- (a) 113.1 to 116.7.
- (b) 111.8 to 118.0.
- (c) 111.2 to 118.6.
- (d) 109.9 to 119.9.

22.2 What conditions for the population and the study design are required by the procedure you used to construct your confidence interval? Which of these conditions are important for the validity of the procedure in this case?

Does nature heal better? Our bodies have a natural electrical field that is known to help wounds heal. Does changing the field strength slow healing? A series of experiments with newts investigated this question. The data below are the healing rates of cuts (micrometers per hour) in a matched pairs experiment. The pairs are the two hind limbs of the same newt, with the body's natural field in one limb (control) and half the natural value in the other limb (experimental).¹

Newt	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Control	25	13	44	45	57	42	50	36	35	38	43	31	26	48
Experimental	24	23	47	42	26	46	38	33	28	28	21	27	25	45
Difference (control – experimental)	1	-10	-3	3	31	-4	12	3	7	10	22	4	1	3

The mean and standard deviation of the differences are 5.71 and 10.56 micrometers per hour, respectively. Use this information to answer Questions 22.3 and 22.4.

22.3 Is there good evidence that changing the electrical field from its natural level slows healing? The *P*-value for your test is

- (a) less than 0.01.
- (b) between 0.01 and 0.05.
- (c) between 0.05 and 0.10.
- (d) greater than 0.10.

22.4 Give a 99% confidence interval for the difference in healing rates (control minus experimental).

Game players. A government survey randomly selected 6889 female high school students and 7028 male high school students.² Of these students, 1020 females and 1926 males played video or computer games for three or more hours a day. Use this information to answer Questions 22.5 to 22.8.

22.5 The estimate of the proportion of males who played video or computer games for three or more hours a day is

- (a) 0.148.
- (b) 0.212.
- (c) 0.231.
- (d) 0.274.

22.6 The estimate of the proportion of females who played video or computer games for three or more hours a day is

- (a) 0.148.
- (b) 0.212.
- (c) 0.231.
- (d) 0.274.

22.7 The sampling distribution for the difference in the sample proportions has standard error

- (a) 0.0026.
- (b) 0.0043.
- (c) 0.0053.
- (d) 0.0068.

22.8 A 99% confidence interval for the difference in proportions of male and female high school students who play video or computer games for at least three hours a day is

- (a) 0.115 to 0.137.
- (b) 0.113 to 0.139.
- (c) 0.108 to 0.144.
- (d) 0.106 to 0.146.

22.9 Wikipedia. A sample survey of 1497 adult Internet users found that 36% consult the online collaborative encyclopedia Wikipedia.³

- (a) Give the standard error SE of \hat{p} , the proportion of all adult Internet users who refer to Wikipedia.
- (b) Give a 95% confidence interval for the proportion of all adult Internet users who refer to Wikipedia.

Men and muscle. Ask young men to estimate their own degree of body muscle by choosing from a set of 100 photos. Then ask them to choose what they think women prefer. The researchers know the actual degree of muscle, measured as kilograms per square meter of fat-free mass, for each of the photos. They can therefore measure the difference between what a subject thinks women prefer and the subject's own self-image. Call this difference the "muscle gap." Here are

summary statistics for the muscle gap from two samples, one of American and European young men and the other of Chinese young men from Taiwan:⁴

Group	<i>n</i>	\bar{x}	<i>s</i>
American/European	200	2.35	2.5
Chinese	55	1.20	3.2

Use this information to answer Questions 22.10 to 22.13.

- 22.10** A 95% confidence interval for the mean size of the muscle gap for all American and European young men is

- (a) 2.35 ± 0.18 .
- (b) 2.35 ± 0.35 .
- (c) 2.35 ± 4.95 .
- (d) 1.15 ± 0.93 .

- 22.11** A 95% confidence interval for the mean size of the muscle gap for all Chinese young men is

- (a) 1.20 ± 0.43 .
- (b) 1.20 ± 0.86 .
- (c) 1.20 ± 6.40 .
- (d) 1.15 ± 0.93 .

- 22.12** Is there a significant difference between the mean sizes of the muscle gap for American/European men and Chinese men? The value of the *t* statistic for testing the null hypothesis of no difference in the mean sizes of the muscle gap is

- (a) 0.47.
- (b) 1.15.
- (c) 2.13.
- (d) 2.47.

- 22.13** Is there a significant difference between the mean sizes of the muscle gap for American/European men and Chinese men? The degrees of freedom using the conservative Option 2 for the *t* statistic for testing the null hypothesis of no difference in the mean sizes of the muscle gap is

- (a) 54.
- (b) 126.5.
- (c) 199.
- (d) 253.

- 22.14 Butterflies mating.** Here's how butterflies mate: a male passes to a female a packet of sperm called a spermatophore. Females may mate several times. Will they remate sooner if the first spermatophore they receive is small? Among 20 females who received a large spermatophore (greater than 25 milligrams), the mean time to the next mating was 5.15 days, with standard deviation 0.18 day. For 21 females who received a small spermatophore (about 7 milligrams), the mean was 4.33 days and the standard deviation was 0.31 day.⁵ Is the observed difference in means statistically significant? Test using the conservative Option 2 for the degrees of freedom. The *P*-value is

- (a) less than 0.01.
- (b) between 0.01 and 0.05.
- (c) between 0.05 and 0.10.
- (d) greater than 0.10.

- Mouse endurance.** A study of the inheritance of speed and endurance in mice found a trade-off between these two characteristics, both of which help mice survive. To test endurance, mice were made to swim in a bucket with a weight attached to their tails. (The mice were rescued when exhausted.) Here are data on endurance in minutes for female and male mice:⁶

Group	<i>n</i>	Mean	Standard deviation
Female	162	11.4	26.09
Male	135	6.7	6.69

Use this information to answer Questions 22.15 to 22.18.



Alamy

22.15 Both sets of endurance data are skewed to the right. Why are t procedures nonetheless reasonably accurate for these data?

22.16 A 90% confidence interval for the mean endurance of female mice swimming is

- (a) 9.35 to 13.45.
- (b) 8.00 to 14.80.
- (c) 7.34 to 15.46.
- (d) 7.14 to 15.66.

22.17 A 90% confidence interval for the mean difference (female minus male) in endurance times is

- (a) 3.35 to 6.45.
- (b) 2.00 to 7.80.
- (c) 1.34 to 8.46.
- (d) 1.18 to 8.22.

22.18 Do the data show that female mice have significantly higher endurance on the average than male mice?

22.19 Pre-readers in kindergarten. A school has two kindergarten classes. There are 21 children in Ms. Toodle's kindergarten class. Of these, 17 are "pre-readers"—children on the verge of reading. There are 19 children in Mr. Grimace's kindergarten class. Of these, 13 are pre-readers. Using the plus four confidence interval method, a 90% confidence interval for the difference in proportions of children in these classes who are pre-readers is 0.104 to 0.336. Which of the following statements is correct?

- (a) This confidence interval is not reliable because the samples are so small.
- (b) This confidence interval is of no use because it contains 0, the value of no difference between classes.
- (c) This confidence interval is reasonable because the sample sizes are both at least 5.
- (d) This confidence interval is not reliable because these samples cannot be viewed as simple random samples taken from a larger population.

22.20 State of the economy. If we want to estimate p , the population proportion of likely voters who believe the state of the economy is the most urgent national concern, with 99% confidence and a margin of error no greater than 2%, how many likely voters need to be surveyed? Assume that you have no idea of the value of p .

- (a) 2401
- (b) 3484
- (c) 4148
- (d) 8256

Favoritism for college athletes? Sports Illustrated surveyed a random sample of 757 Division I college athletes in 36 sports. One question asked was "Have you ever received preferential treatment from a professor because of your status as an athlete?" Of the athletes polled, 225 said "Yes." Use this information to answer Questions 22.21 to 22.23.

22.21 The sample proportion of athletes who have received preferential treatment from a professor is

- (a) 0.160.
- (b) 0.297.
- (c) 0.703.
- (d) 0.840.

22.22 The standard error SE of \hat{p} , the proportion of athletes who have received preferential treatment from a professor, is

- (a) 0.003.
- (b) 0.017.
- (c) 0.209.
- (d) 0.457.

22.23 A 90% confidence interval for the proportion of athletes who have received preferential treatment from a professor is

- (a) 0.276 to 0.320.
- (b) 0.270 to 0.326.
- (c) 0.265 to 0.331.
- (d) 0.255 to 0.342.



Knut Mueller/Peter Arnold

Very-low-birth-weight babies. Starting in the 1970s, medical technology allowed babies with very low birth weight (VLBW, less than 1500 grams, about 3.3 pounds) to survive without major handicaps. It was noticed that these children nonetheless had difficulties in school and as adults. A long-term study has followed 242 VLBW babies to age 20 years, along with a control group of 233 babies from the same population who had normal birth weight.⁷ At age 20, 179 of the VLBW group and 193 of the control group had graduated from high school. Use this information to answer Questions 22.24 to 22.29.

22.24 This is an example of

- (a) an observational study.
- (b) a nonrandomized experiment.
- (c) a randomized controlled study.
- (d) a matched pairs experiment.

22.25 Take p_{VLBW} and $p_{control}$ to be the proportions of all VLBW and normal-birth-weight (control) babies who would graduate from high school. The hypotheses to be tested are

- (a) $H_0: p_{VLBW} = p_{control}$ versus $H_a: p_{VLBW} \neq p_{control}$.
- (b) $H_0: p_{VLBW} = p_{control}$ versus $H_a: p_{VLBW} > p_{control}$.
- (c) $H_0: p_{VLBW} = p_{control}$ versus $H_a: p_{VLBW} < p_{control}$.
- (d) $H_0: p_{VLBW} > p_{control}$ versus $H_a: p_{VLBW} = p_{control}$.

22.26 The pooled sample proportion of babies who would graduate from high school is

- (a) $\hat{p} = 0.74$.
- (b) $\hat{p} = 0.78$.
- (c) $\hat{p} = 0.81$.
- (d) $\hat{p} = 0.83$.

22.27 The numerical value of the z test for comparing the proportions of all VLBW and normal-birth-weight (control) babies who would graduate from high school is

- (a) $z = -1.65$.
- (b) $z = -2.34$.
- (c) $z = -2.77$.
- (d) $z = -3.14$.

22.28 IQ scores were available for 113 men in the VLBW group and for 106 men in the control group. The mean IQ for the 113 men in the VLBW group was 87.6, and the standard deviation was 15.1. The 106 men in the control group had mean IQ 94.7, with standard deviation 14.9. Is there good evidence that mean IQ is lower among VLBW men than among controls from similar backgrounds? To test this with a two-sample t test, the test statistic would be

- (a) $t = -1.72$
- (b) $t = -3.50$
- (c) $t = -5.00$
- (d) $t = -7.10$

22.29 Of the 126 women in the VLBW group, 38 said they had used illegal drugs; 54 of the 124 control group women had done so. The IQ scores for the VLBW women had mean 86.2 (standard deviation 13.4), and the normal-birth-weight controls had mean IQ 89.8 (standard deviation 14.0). Is there a statistically significant difference between the two groups in mean IQ? The P -value for this test is

- (a) less than 0.01.
- (b) between 0.01 and 0.05.
- (c) between 0.05 and 0.10.
- (d) greater than 0.10.

Binge drinking. According to the National Institute on Alcohol Abuse and Alcoholism (NIAAA) and the National Institutes of Health (NIH), 41% of college students nationwide engage in “binge-drinking” behavior: having 5 or more drinks on one occasion during the past two weeks. A college president wonders if the proportion of students enrolled at her college who binge drink is actually lower than the national proportion. In a commissioned study, 348 students are selected randomly from a list of all students enrolled at the college. Of these, 132 admit to having engaged in binge drinking. Use this information to answer Questions 22.30 to 22.32.

22.30 Based on the results of the commissioned study, a 95% confidence interval for the proportion of all students at this college who engage in binge drinking is

- (a) 0.330 to 0.428.
- (b) 0.338 to 0.420.
- (c) 0.341 to 0.417.
- (d) 0.343 to 0.415.

22.31 The college president is more interested in testing her belief that the proportion of students at her college who engage in binge drinking is lower than the national proportion of 0.41. Her staff tests the hypotheses $H_0: p = 0.41$ versus $H_a: p < 0.41$. The P -value is

- (a) between 0.15 and 0.20.
- (b) between 0.10 and 0.15.
- (c) between 0.05 and 0.10.
- (d) below 0.05.

22.32 Which of the following conclusions is reasonable, based on the P -value computed in the previous exercise?

- (a) There is little evidence to support a conclusion that the proportion of students at this particular college who binge drink is lower than the national proportion of 0.41.
- (b) There is moderate but not strong evidence that the proportion of binge-drinking students at this college is lower than the national proportion of 0.41.
- (c) There is strong evidence that the proportion of students at this college who binge drink is lower than the national proportion of 0.41.
- (d) We can't reach any reasonable conclusion, because the assumptions necessary for a significance test for a proportion are not met in this case.

Listening to rap. The Black Youth Project of the University of Chicago interviewed random samples of black, Hispanic, and white young people aged 15 to 25. We can consider this a stratified sample or three separate random samples of 634 blacks, 314 Hispanics, and 567 whites. The survey found that 58% of black youth listen to rap music every day, compared with 45% of Hispanics and 23% of whites. But attitudes were quite similar in the three groups. For example, 72% of blacks, 72% of Hispanics, and 68% of whites agreed that “rap music videos contain too many references to sex.”⁸ Questions 22.33 to 22.35 are based on this study.

22.33 Give a 90% confidence interval for the proportion of all black young people who listen to rap every day.

22.34 Give a 90% confidence interval for the difference between the proportions of all Hispanic and all white young people who listen to rap every day.

22.35 Is there a significant difference between the proportions of black and white young people who think that rap videos contain too much sex? State hypotheses, find the test statistic, and use either software or the bottom row of Table C for the P -value. Be sure to state your conclusion.

Cholesterol in dogs. High levels of cholesterol in the blood are not healthy in either humans or dogs. Because a diet rich in saturated fats raises the cholesterol level, it is plausible that dogs owned as pets have higher cholesterol levels than dogs owned by a veterinary research clinic. “Normal” levels of cholesterol based on the clinic’s dogs would then be misleading. A clinic compared healthy dogs it owned with healthy pets brought to the clinic to be neutered. The summary statistics for blood cholesterol levels (milligrams per deciliter of blood) appear below.⁹

Group	<i>n</i>	\bar{x}	<i>s</i>
Pets	26	193	68
Clinic	23	174	44

Questions 22.36 to 22.40 are based on this study.

22.36 A 95% confidence interval for the mean cholesterol level in pets is

- (a) 179.7 to 206.3.
- (b) 176.8 to 209.2.
- (c) 165.5 to 220.5.
- (d) 159.6 to 226.48.

22.37 Is there strong evidence that pets have a higher mean cholesterol level than clinic dogs? To test this with a two-sample t test, the values of the t statistic and its degrees of freedom using conservative Option 2 are

- (a) $t = 1.17$, $df = 22$.
- (b) $t = 1.17$, $df = 47$.
- (c) $t = 8.92$, $df = 22$.
- (d) $t = 8.92$, $df = 47$.

22.38 A 95% confidence interval for the difference in mean cholesterol levels between pets and clinic dogs is (use conservative Option 2 for the degrees of freedom)

- (a) -26.1 to 64.1.
- (b) -14.6 to 52.6.
- (c) -8.7 to 46.7.
- (d) 2.8 to 35.2.

22.39 What conditions must be satisfied to justify the procedures you used in Exercise 22.36? In Exercise 22.37? In Exercise 22.38?

22.40 Assuming that the cholesterol measurements have no outliers and are not strongly skewed, what is the chief threat to the validity of the results of this study?

Choosing an inference procedure. In each of Questions 22.41 to 22.46, say which type of inference procedure from the Statistics in Summary flowchart (page 535) you would use, or explain why none of these procedures fits the problem. You do not need to carry out any procedures.

22.41 Driving too fast. How seriously do people view speeding in comparison with other annoying behaviors? A large random sample of adults was asked to rate a number of behaviors on a scale of 1 (no problem at all) to 5 (very severe problem). Do speeding drivers get a higher average rating than noisy neighbors?

22.42 Preventing drowning. Drowning in bathtubs is a major cause of death in children less than 5 years old. A random sample of parents was asked many questions related to bathtub safety. Overall, 85% of the sample said they used baby bathtubs for infants. Estimate the percent of all parents of young children who use baby bathtubs.

22.43 Acid rain? You have data on rainwater collected at 16 locations in the Adirondack Mountains of New York State. One measurement is the acidity of the water, measured by pH on a scale of 0 to 14 (the pH of distilled water is 7.0). Estimate the average acidity of rainwater in the Adirondacks.

22.44 Athletes' salaries. Looking online, you find the salaries of the 25 players with guaranteed salaries on the roster of the Chicago Cubs as of opening day of the 2011 baseball season. The team total was \$125.0 million, seventh highest in the major leagues. Estimate the average salary of the Cubs players.

22.45 Looking back on love. How do young adults look back on adolescent romance? Investigators interviewed 40 couples in their midtwenties. The female and male partners were interviewed separately. Each was asked about his or her current relationship and also about a romantic relationship that lasted at least two months when they were aged 15 or 16. One response variable was a measure on a numerical scale of how much the attractiveness of the adolescent partner mattered. You want to compare the men and women on this measure.

22.46 Preventing AIDS through education. The Multisite HIV Prevention Trial was a randomized comparative experiment to compare the effects of twice-weekly

small-group AIDS discussion sessions (the treatment) with a single one-hour session (the control). Compare the effects of treatment and control on each of the following response variables:

- A subject does or does not use condoms six months after the education sessions.
- The number of unprotected intercourse acts by a subject between four and eight months after the sessions.
- A subject is or is not infected with a sexually transmitted disease six months after the sessions.

SUPPLEMENTARY EXERCISES

Supplementary exercises apply the skills you have learned in ways that require more thought or more use of technology. Some of these exercises start from actual data rather than from data summaries. Many of these exercises ask you to follow the **Plan, Solve, and Conclude** steps of the four-step process. Remember that the **Solve** step includes checking the conditions for the inference you plan.

22.47 Do you have confidence? A report of a survey distributed to randomly selected email addresses at a large university says: “We have collected 427 responses from our sample of 2,100 as of April 30, 2004. This number of responses is large enough to achieve a 95% confidence interval with $\pm 5\%$ margin of sampling error in generalizing the results to our study population.”¹⁰ Why would you be reluctant to trust a confidence interval based on these data?

22.48 Pain from a rubber hand. People who have had limbs amputated sometimes feel sensations from the limb that is no longer there. To study this effect, psychologists asked subjects to place their right arm on a table. They then put a rubber arm and hand next to the real arm, with a high partition arranged so that the subject could see only the rubber arm. After a few minutes during which the real and fake hand were both tapped by an experimenter, the subjects felt the taps coming from the location of the rubber hand they could see, not from the real hand they couldn’t see. Now the experiment begins: bend back a finger of the fake hand in a way that would cause pain, while merely lifting a real finger. Do electrical measurements show a response to pain? Because there would be some response from the surprise of being touched, a control treatment delayed the touch to the real hand to separate surprise from “pain.” Here are summary data for 16 undergraduate students who were subject to both stimuli:¹¹

Stimulus	\bar{x}	s
Treatment	0.39	0.28
Control	0.18	0.20

- Which t procedures are correct for comparing the mean response to treatment and control: one-sample, matched pairs, or two-sample?

- The data summary given is not enough information to carry out the correct t procedures. Explain why not.

22.49 Monkeys and music. Humans generally prefer music to silence. What about monkeys? Allow a tamarin monkey to enter a V-shaped cage with food in both arms of the V. After the monkey eats the food, which arm will it prefer? The monkey’s location determines what it hears, a lullaby played by a flute in one arm and silence in the other. Each of 4 monkeys was tested 6 times, on different days and with the music arm alternating between left and right (in case a monkey prefers one direction). The monkeys chose silence for about 65% of their time in the cage. The researchers reported a one-sample t test for the mean percent of time spent in the music arm, $H_0: \mu = 50\%$ against the two-sided alternative, $t = -5.26$, $df = 23$, $P < 0.0001$.¹²

Although the result is interesting, the statistical analysis is not correct. The degrees of freedom $df = 23$ show that the researchers assumed that they had 24 independent observations. Explain why the results of the 24 trials are not independent.

22.50 Drug-detecting rats? Dogs are big and expensive. Rats are small and cheap. Might rats be trained to replace dogs in sniffing out illegal drugs? A first study of this idea trained rats to rear up on their hind legs when they smelled simulated cocaine. To see how well rats performed after training, they were let loose on a surface with many cups sunk in it, one of which contained simulated cocaine. Four out of six trained rats succeeded in 80 out of 80 trials.¹³ How should we estimate the long-term success rate p of a rat that succeeds in every one of 80 trials?

- What is the rat’s sample proportion \hat{p} ? What is the large-sample 95% confidence interval for p ? It’s not plausible that the rat will always be successful, as this interval says.

- Find the plus four estimate \hat{p} and the plus four 95% confidence interval for p . These results are more reasonable.

22.51 A new vaccine. In 2006, the pharmaceutical company Merck released a vaccine named Gardasil for human papilloma virus, the most common cause of cervical cancer in young women. The Merck Web site gives results from “four placebo-controlled, double-blind, randomized clinical studies” with women 16 to 26 years of age, as follows:¹⁴

	Cervical cancer		Genital warts
	n		n
Gardasil	8487	0	7897
Placebo	8460	32	7899

- (a) Give a 99% confidence interval for the difference in the proportions of young women who develop cervical cancer with and without the vaccine.
 (b) Do the same for the proportions who develop genital warts.
 (c) What do you conclude about the overall effectiveness of the vaccine?

22.52 Starting to talk. At what age do infants speak their first word of English? Here are data on 20 children (ages in months):¹⁵



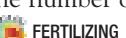
15	26	10	9	15	20	18	11	8	20
7	9	10	11	11	10	12	17	11	10

(In fact, the sample contained one more child, who began to speak at 42 months. Child development experts consider this abnormally late, so the investigators dropped the outlier to get a sample of “normal” children. We are willing to treat these data as an SRS.) Is there good evidence that the mean age at first word among all normal children is greater than one year?

22.53 Fertilizing a tropical plant. Bromeliads are tropical flowering plants. Many are epiphytes that attach to trees and obtain moisture and nutrients from air and rain. Their leaf bases form cups that collect water and are home to the larvae of many insects. In an experiment in Costa Rica, Jacqueline Ngai and Diane Srivastava studied whether added nitrogen increases the productivity of bromeliad plants. Bromeliads were randomly assigned to nitrogen or control groups. Here are data on the number of new leaves produced over a 7-month period:¹⁶



Kelly Kalhoefer/Jupiter Images



FERTILIZING

Control	11	13	16	15	15	11	12
Nitrogen	15	14	15	16	17	18	17

Is there evidence that adding nitrogen increases the mean number of new leaves formed?

22.54 Starting to talk, continued. Use the data in Exercise 22.52 to give a 90% confidence interval for the mean age at which children speak their first word.



22.55 Dyeing fabrics. Different fabrics respond differently when dyed. This matters to clothing manufacturers, who want the color of the fabric to be just right. A researcher dyed fabrics made of cotton and of ramie with the same “procion blue” dye applied in the same way. Then she used a colorimeter to measure the lightness of the color on a scale in which black is 0 and white is 100. Here are the data for 8 pieces of each fabric:¹⁷



Cotton	48.82	48.88	48.98	49.04	48.68	49.34	48.75	49.12
Ramie	41.72	41.83	42.05	41.44	41.27	42.27	41.12	41.49

Is there a significant difference between the fabrics? Which fabric is darker when dyed in this way?

22.56 More on dyeing fabrics. The color of a fabric depends on the dye used and also on how the dye is applied. This matters to clothing manufacturers, who want the color of the fabric to be just right. The study discussed in the previous exercise went on to dye fabric made of ramie with the same procion blue dye applied in two different ways. Here are the lightness scores for 8 pieces of identical fabric dyed in each way:



Method B	40.98	40.88	41.30	41.28	41.66	41.50	41.39	41.27
Method C	42.30	42.20	42.65	42.43	42.50	42.28	43.13	42.45

- (a) This is a randomized comparative experiment. Outline the design.
 (b) A clothing manufacturer wants to know which method gives the darker color (lower lightness score). Use sample means to answer this question. Is the difference between the two sample means statistically significant? Can you tell from just the P-value whether the difference is large enough to be important in practice?

22.57 Do parents matter? A professor asked her sophomore students, “Does either of your parents allow you to drink alcohol around him or her?” and “How many drinks do you typically have per session? (A drink is defined as one 12 oz beer, one 4 oz glass of wine, or one 1 oz shot of liquor.)” Table 22.1 contains the responses of the female

TABLE 22.1 Drinks per session by female students

PARENT ALLOWS STUDENT TO DRINK														
2.5	1	2.5	3	1	3	3	3	2.5	2.5	3.5	5	2		
7	7	6.5	4	8	6	6	3	6	3	4	7	5		
3.5	2	1	5	3	3	6	4	2	7	5	8	1		
6	5	2.5	3	4.5	9	5	4	4	3	4	6	4		
5	1	5	3	10	7	4	4	4	4	2	2.5	2.5		
PARENT DOES NOT ALLOW STUDENT TO DRINK														
9	3.5	3	5	1	1	3	4	4	3	6	5	3		
8	4	4	5	7	7	3.5	3	10	4	9	2	7		
4	3	1												

students who are not abstainers.¹⁸ The sample is all students in one large sophomore-level class. The class is popular, so we are tentatively willing to regard its members as an SRS of sophomore students at this college. Does the behavior of parents make a significant difference in how many drinks students have on the average? 

22.58 Parents' behavior. We wonder what proportion  of female students have at least one parent who allows them to drink around him or her. Table 22.1 contains information about a sample of 94 students. Use this sample to give a 95% confidence interval for this proportion. 

22.59 Diabetic mice. The body's natural electrical field  helps wounds heal. If diabetes changes this field, that might explain why people with diabetes heal slowly.

A study of this idea compared normal mice and mice bred to spontaneously develop diabetes. The investigators attached sensors to the right hip and front feet of the mice and measured the difference in electrical potential (millivolts) between these locations. The data are shown in the table below:¹⁹ 

- (a) Make a stemplot of each sample of potentials. There is a low outlier in the diabetic group. Does it appear that potentials in the two groups differ in a systematic way?
- (b) Is there significant evidence of a difference in mean potentials between the two groups?
- (c) Repeat your inference without the outlier. Does the outlier affect your conclusion?

Diabetic Mice						Normal Mice				
14.70	13.60	7.40	1.05	10.55	16.40	13.80	9.10	4.95	7.70	9.40
10.00	22.60	15.20	19.60	17.25	18.40	7.20	10.00	14.55	13.30	6.65
9.80	11.70	14.85	14.45	18.25	10.15	9.50	10.40	7.75	8.70	8.85
10.85	10.30	10.45	8.55	8.85	19.20	8.40	8.55	12.60		

22.60 Keeping crackers from breaking. We don't like to find broken crackers when we open the package. How can makers reduce breaking? One idea is to microwave the crackers for 30 seconds right after baking them. Analyze the following results from two experiments intended to examine this idea.²⁰ Does microwaving significantly improve indicators of future breaking? How large is the improvement? What do you conclude about the idea of microwaving crackers?

(a) The experimenter randomly assigned 65 newly baked crackers to be microwaved and another 65 to a control group that is not microwaved. Fourteen days after baking, 3 of the 65 microwaved crackers and 57 of the 65 crackers in the control group showed visible checking, which is the starting point for breaks.

(b) The experimenter randomly assigned 20 crackers to be microwaved and another 20 to a control group. After 14 days, he broke the crackers. Here are summaries of the pressure needed to break them, in pounds per square inch:

	Microwave	Control
Mean	139.6	77.0
Standard deviation	33.6	22.6

22.61 Falling through the ice. Table 7.3 (page 192) gives the dates on which a wooden tripod fell through the ice of the Tanana River in Alaska, thus deciding the winner of the Nenana Ice Classic contest, for the years 1917



2006 Bill Watkins/AlaskaStock.com

to 2010. Give a 95% confidence interval for the mean date on which the tripod falls through the ice. After calculating the interval in the scale used in the table (days from April 20, which is Day 1), translate your result into calendar dates and hours within the dates. (Each hour is $1/24$, or 0.042, of a day.)  TANANA

22.62 A case for the Supreme Court. In 1986, a Texas jury found a black man guilty of murder. The prosecutors had used “peremptory challenges” to remove 10 of the 11 blacks and 4 of the 31 whites in the pool from which the jury was chosen.²¹ The law says that there must be a plausible reason (that is, a reason other than race) for different treatment of blacks and whites in the jury pool. When the case reached the Supreme Court 17 years later, the Court said that “happenstance is unlikely to produce this disparity.” Explain why the methods we know can’t be safely used to do the inference that lies behind the Court’s finding that chance is unlikely to produce so large a black-white difference.

22.63 Mouse genes. A study of genetic influences on diabetes compared normal mice with similar mice genetically altered to remove a gene called *aP2*. Mice of both types were allowed to become obese by eating a high-fat diet. The researchers then measured the levels of insulin and glucose in their blood plasma. Here are some excerpts from their findings.²² The normal mice are called “wild-type” and the altered mice are called “*aP2*^{-/-}.”

*Each value is the mean \pm SEM of measurements on at least 10 mice. Mean values of each plasma component are compared between *aP2*^{-/-} mice and wild-type controls by Student’s t test (*P < 0.05 and **P < 0.005).*

Parameter	Wild type	<i>aP2</i> ^{-/-}
Insulin (ng/ml)	5.9 ± 0.9	$0.75 \pm 0.2^{**}$
Glucose (mg/dl)	230 ± 25	$150 \pm 17^*$

*Despite much greater circulating amounts of insulin, the wild-type mice had higher blood glucose than the *aP2*^{-/-} animals. These results indicate that the absence of *aP2* interferes with the development of dietary obesity-induced insulin resistance.*

Other biologists are supposed to understand the statistics reported so tersely.

(a) What does “SEM” mean? What is the expression for SEM based on n , \bar{x} , and s from a sample?

(b) Which of the tests we have studied did the researchers apply?

(c) Explain to a biologist who knows no statistics what $P < 0.05$ and $P < 0.005$ mean. Which is stronger evidence of a difference between the two types of mice?

22.64 Mouse genes, continued. The report quoted in the previous exercise says only that the sample sizes were “at least

10.” Suppose that the results are based on exactly 10 mice of each type. Use the values in the table to find \bar{x} and s for the insulin concentrations in the two types of mice. Carry out a test to assess the significance of the difference in mean insulin concentration. Does your P -value confirm the claim in the report that $P < 0.005$?

this page left intentionally blank

Inference about Relationships

Part IV

S

tatistical inference offers more methods than anyone can know well, as a glance at the offerings of any large statistical software package demonstrates. In an introductory text, we must be selective. Parts I to III have laid a foundation for understanding statistics:

- The nature and purpose of data analysis.
- The central ideas of designs for data production.
- The reasoning behind confidence intervals and significance tests.
- Experience applying these ideas in practice.

Each of the three chapters of Part IV offers an introduction to a more advanced topic in statistical inference. You may choose to read any or all of them, in any order.

What makes a statistical method “more advanced”? More complex data, for one thing. In Part III, we looked only at methods for inference about a single population parameter and for comparing two parameters. All the chapters in Part IV present methods for studying relationships between two variables. In Chapter 23, both variables are categorical, with data given as a two-way table of counts of outcomes. Chapter 24 considers inference in the setting of regressing a response variable on an explanatory variable. This is an important type of relationship between two quantitative variables. In Chapter 25 we meet methods for comparing the mean response in more than two groups. Here, the explanatory variable (group) is categorical and the response variable is quantitative. These chapters together bring our knowledge of inference to the same point that our study of data analysis reached in Chapters 1 to 7.

With greater complexity comes greater reliance on technology. In these final three chapters you will more often be interpreting the output of statistical software or using software yourself. With effort, you can do the calculations needed in Chapter 23 with a basic calculator. In Chapters 24 and 25, the pain is too great and the contribution to learning too small. Fortunately, you can grasp the ideas without step-by-step arithmetic.

Another aspect of “more advanced” methods is new concepts and ideas. This is where we draw the line in deciding what statistical topics we can master in a first course. Part IV builds elaborate methods on the foundation we have laid without introducing fundamentally new concepts. You can see that statistical practice does need additional big ideas by reading the sections on “the problem of multiple comparisons” in Chapters 23 and 25. But the ideas you already know place you among the world’s statistical sophisticates.

INFERENCE ABOUT RELATIONSHIPS

CHAPTER 23 Two Categorical Variables: The Chi-Square Test

CHAPTER 24 Inference for Regression

CHAPTER 25 One-Way Analysis of Variance: Comparing Several Means



Two Categorical Variables: The Chi-Square Test

The two-sample z procedures of Chapter 21 allow us to compare the proportions of successes in two groups, either two populations or two treatment groups in an experiment. In the first example in Chapter 21 (page 515), we compared young men and young women by looking at whether or not they lived with their parents. That is, we looked at a relationship between two categorical variables, sex (female or male) and “Where do you live?” (with parents or not). In fact, the data include four more outcomes for “Where do you live?”, in another person’s home, in your own place, in group quarters such as a dormitory, or “other.” When there are more than two outcomes, or when we want to compare more than two groups, we need a new statistical test. The new test addresses a general question: *is there a relationship between two categorical variables?*

TWO-WAY TABLES

We saw in Chapter 6 that we can present data on two categorical variables in a **two-way table** of counts. That’s our starting point. Let’s continue our exploration of where college-age young people live.

EXAMPLE 23.1 Where do young people live?

A sample survey asked a random sample of young adults, “Where do you live now? That is, where do you stay most often?” Table 23.1 is a two-way table of all 2984 people in the sample (both men and women) classified by their age and by where they lived.¹ Living arrangement is a categorical variable. Even though age is quantitative, the two-way table treats age as dividing young adults into

IN THIS CHAPTER WE COVER...

- Two-way tables
- The problem of multiple comparisons
- Expected counts in two-way tables
- The chi-square test statistic
- Cell counts required for the chi-square test
- Using technology
- Uses of the chi-square test
- The chi-square distributions
- The chi-square test for goodness of fit*

two-way table



WHERE LIVE

cell

four categories. Table 23.1 gives the counts for all 20 combinations of age and living arrangement. Each of the 20 counts occupies a **cell** of the table. ■

TABLE 23.1 Young adults by age and living arrangement

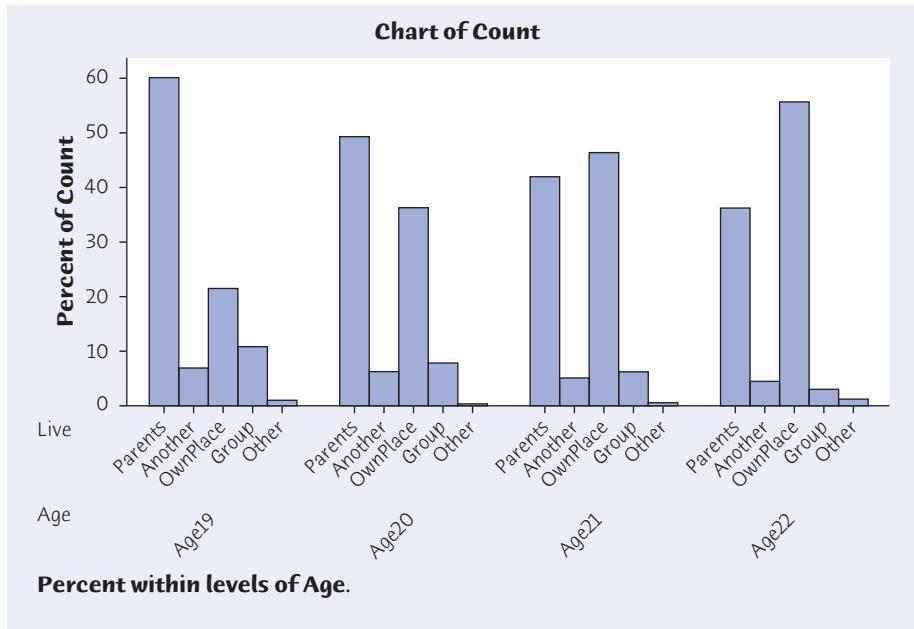
LIVING ARRANGEMENT	AGE (YEARS)				TOTAL
	19	20	21	22	
Parents' home	324	378	337	318	1357
Another person's home	37	47	40	38	162
Your own place	116	279	372	487	1254
Group quarters	58	60	49	25	192
Other	5	2	3	9	19
Total	540	766	801	877	2984

As usual, we prepare for inference by first doing data analysis. Because we think that age helps explain where young people live, find the percents of people in each age group who have each living arrangement. The percents appear in Table 23.2. Each column adds to 100% (up to roundoff error) because we are looking at each age group separately. In the language of Chapter 6 (page 162), Table 23.2 shows the four *conditional distributions* of living arrangements given a specific age.

Figure 23.1 is Minitab's bar graph comparing the four conditional distributions. The graph shows a strong relationship between age and living arrangement. As young adults age from 19 to 22, the percent living with their parents drops and the percent living in their own place rises. The percent living in group quarters also declines with age as college students move out of dormitories. Are these differences among the four age groups large enough to be statistically significant?

TABLE 23.2 Percents of each age group who have each living arrangement (read down columns)

LIVING ARRANGEMENT	AGE (YEARS)			
	19	20	21	22
Parents' home	60.0	49.3	42.1	36.3
Another person's home	6.9	6.1	5.0	4.3
Your own place	21.5	36.4	46.4	55.5
Group quarters	10.7	7.8	6.1	2.9
Other	0.9	0.3	0.4	1.0
Total	100.0	99.9	100.0	100.0

**FIGURE 23.1**

Minitab bar graph comparing the four conditional distributions of living arrangements given age, for Example 23.1.



APPLY YOUR KNOWLEDGE

- 23.1 Facebook at Penn State.** The Pennsylvania State University has its main campus in University Park and more than 20 smaller “commonwealth campuses” around the state. The Penn State Division of Student Affairs polled a random sample of undergraduates about their use of online social networking. (The response rate was only about 20%, which casts some doubt on the usefulness of the data.) Facebook was the most popular site, with more than 80% of students having an account. Here is a comparison of Facebook use by undergraduates at the University Park and commonwealth campuses:²  FACEBOOKUSE



Courtesy Pennsylvania State University

Use Facebook	University Park	Commonwealth
Do not use Facebook	68	248
Several times a month or less	55	76
At least once a week	215	157
At least once a day	640	394

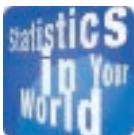
- What percent of University Park students fall in each Facebook category? What percent of commonwealth campus students fall in each category? Each column should add to 100% (up to roundoff error). These are the conditional distributions of Facebook use given campus setting.
- Make a bar graph that compares the two conditional distributions. What are the most important differences in Facebook use between the two campus settings?

23.2 Video-gaming and grades. The popularity of computer, video, online, and virtual reality games has raised concerns about their ability to negatively impact youth. The data in this exercise are based on a recent survey of 14- to 18-year-olds in Connecticut high schools. Here are the grade distributions of boys who have and have not played video games:³



	Grade Average		
	A's and B's	C's	D's and F's
Played games	736	450	193
Never played games	205	144	80

- (a) It appears that boys who have played video games have better grades than those who have never played video games. Give percents to back up this claim. Make a bar graph that compares your percents for boys who have and have not played video games.
- (b) Association does not prove causation. Explain why you can't conclude from this study that playing video games improves grades for boys.



He started it!

A study of deaths in bar fights showed that in 90% of the cases, the person who died started the fight. You shouldn't believe this. If you killed someone in a fight, what would you say when the police ask you who started the fight? After all, dead men tell no tales.

THE PROBLEM OF MULTIPLE COMPARISONS

The null hypothesis in Example 23.1 is that in the population of all American young adults there is *no difference* among the four distributions of living arrangements for people aged 19, 20, 21, and 22. If the null hypothesis is true, the differences in the sample are just accidents due to random selection of the sample. Put more generally, the null hypothesis is that there is *no relationship* between two categorical variables,

H_0 : there is no relationship between age and where young people live

The alternative hypothesis says that there is a relationship but does not specify any particular kind of relationship,

H_a : there is some relationship between age and living arrangement

Any difference among the four distributions of living arrangements in the population of all young adults means that the null hypothesis is false and the alternative hypothesis is true. The alternative hypothesis is not one-sided or two-sided. We might call it “many-sided” because it allows any kind of difference.

With only the methods we already know, we might start by comparing the proportions of people aged 19 and 22 who live with their parents. We could similarly compare other pairs of proportions, ending up with many tests and many P -values. This is a bad idea. The P -values belong to each test separately, not to the collection of all the tests together. Think of the distinction between the probability that a basketball player makes a free throw and the probability that she makes all her free throws in a game. When we do many individual tests or confidence intervals, the individual P -values and confidence levels don't tell us how confident we can be in all the inferences taken together.



Because of this, it's cheating to pick out one large difference from Table 23.2 and then test its significance as if it were the only comparison we had in mind. For example, the percents of people aged 19 and 22 who live with their parents are significantly different ($z = 8.72, P < 0.001$) if we make just this one comparison. But we could also pick a comparison that is not significant; for example, the proportions of people aged 21 and 22 who live in another person's home do not differ significantly ($z = 0.64, P = 0.522$). Individual comparisons can't tell us whether the four distributions, each with five outcomes, are significantly different.

The problem of how to do many comparisons at once with an overall measure of confidence in all our conclusions is common in statistics. This is the problem of **multiple comparisons**. Statistical methods for dealing with multiple comparisons usually have two steps:

1. An *overall test* to see if there is good evidence of *any* differences among the parameters that we want to compare.
2. A detailed *follow-up analysis* to decide which of the parameters differ and to estimate how large the differences are.

multiple comparisons

The overall test, though more complex than the tests we met earlier, is reasonably straightforward. The follow-up analysis can be quite elaborate. We will concentrate on the overall test and use data analysis to describe in detail the nature of the differences.

APPLY YOUR KNOWLEDGE

23.3 Facebook at Penn State. In the setting of Exercise 23.1, we might do several significance tests to compare University Park with the commonwealth campuses.

- (a) Is there a significant difference between the proportions of students in the two locations who do not use Facebook? Give the *P*-value.
- (b) Is there a significant difference between the proportions of students in the two locations who are in the "at least once a week" category? Give the *P*-value.
- (c) Explain clearly why *P*-values for individual outcomes like these can't tell us whether the two distributions for all four outcomes in the two locations differ significantly.  FACEBOOKUSE

23.4 Is astrology scientific? The University of Chicago's General Social Survey (GSS) is the nation's most important social science sample survey. The GSS asked a random sample of adults their opinion about whether astrology is very scientific, sort of scientific, or not at all scientific. Here is a two-way table of counts for people in the sample who had three levels of higher education degrees:⁴  ASTROLOGY

Opinion	Degree Held		
	Junior college	Bachelor's	Graduate
Not at all scientific	87	198	111
Very or sort of scientific	43	57	28

- (a) Give three 95% confidence intervals, for the percents of people with each degree who think that astrology is not at all scientific.
- (b) Explain clearly why we are *not* 95% confident that *all three* of these intervals capture their respective population proportions.

EXPECTED COUNTS IN TWO-WAY TABLES

Our general null hypothesis H_0 is that there is *no relationship* between the two categorical variables that label the rows and columns of a two-way table. To test H_0 , we compare the observed counts in the table with the *expected counts*, the counts we would expect—except for random variation—if H_0 were true. If the observed counts are far from the expected counts, that is evidence against H_0 . It is easy to find the expected counts.

EXPECTED COUNTS

The **expected count** in any cell of a two-way table when H_0 is true is

$$\text{expected count} = \frac{\text{row total} \times \text{column total}}{\text{table total}}$$

EXAMPLE 23.2 Where young people live: expected counts

Let's find the expected counts for the study of where young people live. Look back at the two-way table of counts, Table 23.1. That table includes the row and column totals. The expected count of 19-year-olds who live in their parents' home is

$$\frac{\text{row 1 total} \times \text{column 1 total}}{\text{table total}} = \frac{(1357)(540)}{2984} = 245.57$$

The expected count of 22-year-olds who live with their parents is

$$\frac{\text{row 1 total} \times \text{column 4 total}}{\text{table total}} = \frac{(1357)(877)}{2984} = 398.82$$

The actual counts are 324 and 318. More younger people and fewer older people live with their parents than we would expect if there were no relationship between age and living arrangement. Table 23.3 shows all 20 expected counts.

As this table shows, the *expected counts have exactly the same row and column totals (up to roundoff error) as the observed counts*. That's a good way to check your work. Comparing the actual counts (Table 23.1) and the expected counts (Table 23.3) shows in what ways the data diverge from the null hypothesis. ■

Why the formula works Where does the formula for an expected count come from? Think of a basketball player who makes 70% of her free throws in the long run. If she shoots 10 free throws in a game, we expect her to make 70% of them, or 7 of the 10. Of course, she won't make exactly 7 every time she shoots 10 free

TABLE 23.3 Young adults by age and living arrangement: expected counts

LIVING ARRANGEMENT	AGE (YEARS)				TOTAL
	19	20	21	22	
Parents' home	245.57	348.35	364.26	398.82	1357
Another person's home	29.32	41.59	43.49	47.61	162
Your own place	226.93	321.90	336.61	368.55	1254
Group quarters	34.75	49.29	51.54	56.43	192
Other	3.44	4.88	5.10	5.58	19
Total	540	766	801	877	2984

throws in a game. There is chance variation from game to game. But in the long run, 7 of 10 is what we expect. In more formal language, if we have n independent tries and the probability of a success on each try is p , we expect np successes.

Now go back to the count of 19-year-olds living in their parents' home. The proportion of all 2984 subjects who live with their parents is

$$\frac{\text{count of successes}}{\text{table total}} = \frac{\text{row 1 total}}{\text{table total}} = \frac{1357}{2984}$$

Think of this as p , the overall proportion of successes. If H_0 is true, we expect (except for random variation) this same proportion of successes in all four age groups. So the expected count of successes among the 540 19-year-olds is

$$np = (540)\left(\frac{1357}{2984}\right) = 245.57$$

That's the formula in the Expected Counts box.

APPLY YOUR KNOWLEDGE

23.5 Facebook at Penn State. The two-way table in Exercise 23.1 displays data on use of Facebook by two groups of Penn State students. It's clear that nonusers are much more frequent at the commonwealth campuses. Let's look just at students who have Facebook accounts:  FACEBOOKUSERS

Use Facebook	University Park	Commonwealth
Several times a month or less	55	76
At least once a week	215	157
At least once a day	640	394
Total Facebook users	910	627

The null hypothesis is that there is no relationship between campus and Facebook use.

- (a) If this hypothesis is true, what are the expected counts for Facebook use among commonwealth campus students? This is one column of the two-way table of expected counts. Find the column total and verify that it agrees with the column total for the observed counts.
- (b) Commonwealth campus students as a group are older and more likely to be married and employed than University Park students. What does comparing the observed and expected counts in this column show about Facebook use by these students?

23.6 Video-gaming and grades. Exercise 23.2 describes a comparison of the grade distribution of a sample of 14- to 18-year-old boys in Connecticut who do and don't play video games. The null hypothesis "no relationship" says that in the population of all 14- to 18-year-old boys in Connecticut, the proportions who have each grade average are the same for those who play and don't play video games.  GAMING

- (a) Find the expected counts if this hypothesis is true and display them in a two-way table. Add the row and column totals to your table and check that they agree with the totals for the observed counts.
- (b) Are there any large deviations between the observed counts and the expected counts? What kind of relationship between the two variables do these deviations point to?

THE CHI-SQUARE TEST STATISTIC

To test whether the observed differences among the four distributions of living arrangements given age are statistically significant, we compare the observed and expected counts. The test statistic that makes the comparison is the *chi-square statistic*.

CHI-SQUARE STATISTIC

The **chi-square statistic** is a measure of how far the observed counts in a two-way table are from the expected counts. The formula for the statistic is

$$\chi^2 = \sum \frac{(\text{observed count} - \text{expected count})^2}{\text{expected count}}$$

The sum is over all cells in the table.

As you might guess, the symbol χ in the box is the Greek letter chi. The chi-square statistic is a sum of terms, one for each cell in the table.

EXAMPLE 23.3 Where young people live: the test statistic

In the study of where young people live, 324 19-year-olds lived with their parents. The expected count for this cell is 245.57. So the term of the chi-square statistic from this cell is

$$\frac{(\text{observed count} - \text{expected count})^2}{\text{expected count}} = \frac{(324 - 245.57)^2}{245.57}$$

$$= \frac{6151.26}{245.57} = 25.05$$

The chi-square statistic χ^2 is the sum of 20 terms like this one. Here they are, arranged to match the layout of the two-way table:

$$\begin{aligned}\chi^2 &= 25.05 + 2.53 + 2.04 + 16.38 \\ &\quad + 2.01 + 0.71 + 0.28 + 1.94 \\ &\quad + 54.23 + 5.72 + 3.72 + 38.07 \\ &\quad + 15.56 + 2.33 + 0.13 + 17.51 \\ &\quad + 0.71 + 1.70 + 0.87 + 2.09 \\ &= 193.58\end{aligned}$$

To find the value $\chi^2 = 193.58$, we had to calculate the 20 expected cell counts in Table 23.3 and then the 20 terms of the sum. Moreover, rounding each term to two decimal places creates roundoff error in the sum. Software is very handy in finding χ^2 . ■

Think of χ^2 as a measure of the distance of the observed counts from the expected counts. Like any distance, it is always zero or positive, and it is zero only when the observed counts are exactly equal to the expected counts. Large values of χ^2 are evidence against H_0 because they say that the observed counts are far from what we would expect if H_0 were true. Although the alternative hypothesis H_a is many-sided, the chi-square test is one-sided because any violation of H_0 tends to produce a large value of χ^2 . Small values of χ^2 are not evidence against H_0 .



CELL COUNTS REQUIRED FOR THE CHI-SQUARE TEST

The chi-square test, like the z procedures for comparing two proportions, is an approximate method that becomes more accurate as the counts in the cells of the table get larger. We must therefore check that the counts are large enough to allow us to trust the P -value. Fortunately, the chi-square approximation is accurate for quite modest counts. Here is a practical guideline.⁵

CELL COUNTS REQUIRED FOR THE CHI-SQUARE TEST

You can safely use the chi-square test with critical values from the chi-square distribution when no more than 20% of the expected counts are less than 5 and all individual expected counts are 1 or greater. In particular, all four expected counts in a 2×2 table should be 5 or greater.

Note that the guideline uses *expected* cell counts. The expected counts for the living arrangements study of Example 23.1 appear in Table 23.3. Only 2 of the 20 expected counts (that's 10%) are less than 5 and all are greater than 1, so the data meet the guideline for safe use of chi-square.

USING TECHNOLOGY

Calculating the expected counts and then the chi-square statistic by hand is time-consuming. As usual, software saves time and always gets the arithmetic right. Figure 23.2 shows output for the chi-square test for the living arrangements data from a graphing calculator and two statistical programs.

Texas Instruments Graphing Calculator

```
X2-test
X2=193.5482790
P=6.981157e-35
df=12
round<[B1,2>
[[242.37 348.35...
[29.32 41.59 ...
[226.93 321.9 ...
[34.75 49.29 ...
[5.44 4.88 ...
]]
```

Minitab

The Minitab session window displays the following output:

Session

Expected counts are printed below observed counts
Chi-Square contributions are printed below expected counts

	Age 19	Age 29	Age 21	Age 22	Total
Parents	324	378	337	318	1357
	245.57	348.35	364.26	398.82	
	25.049	2.525	2.040	16.379	
Another	37	47	40	38	162
	29.32	41.59	43.49	47.61	
	2.014	0.705	0.279	1.940	
OwnPlace	116	279	372	437	1254
	226.93	321.90	336.61	368.55	
	54.226	5.719	3.720	38.068	
Group	58	60	49	25	192
	34.75	49.29	51.54	56.43	
	15.564	2.329	0.125	17.505	
Other	5	2	3	9	19
	3.44	4.88	5.10	5.58	
	0.709	1.697	0.865	2.090	
Total	540	766	801	877	2984

Chi-Sq = 193.548, DF = 12, P-Value = 0.000
2 cells with expected counts less than 5.

This key identifies the output for each cell in the table

FIGURE 23.2

Output from a graphing calculator, Minitab, and CrunchIt! for the two-way table for the study of where young people live, for Example 23.4.

CrunchIt!

		Age19	Age20	Age21	Age22	All						
		324	378	337	318	1357						
Parents	23.88	27.98	24.83	23.43	100							
	60	49.35	42.07	36.28	45.48							
	10.86	12.67	11.29	10.86	45.48							
	37	47	40	38	162							
Another	22.84	29.01	24.89	23.48	100							
	6.852	6.138	4.984	4.303	5.429							
	1.240	1.575	1.340	1.273	5.429							
	116	279	372	487	1254							
OwnPlace	9.250	22.25	29.87	38.84	100							
	21.48	36.42	46.44	55.53	42.02							
	3.887	9.350	12.47	16.32	42.02							
	58	80	49	25	192							
Group	30.21	31.25	25.52	13.02	100							
	10.74	7.833	6.117	2.851	6.434							
	1.944	2.011	1.642	0.8378	6.434							
	5	2	3	9	19							
Other	26.32	10.53	15.79	47.37	100							
	0.9259	0.2611	0.3745	1.026	0.6367							
	0.1676	0.06702	0.1005	0.3016	0.6367							
	540	766	801	877	2984							
All	18.10	25.67	26.84	29.39	100							
	100	100	100	100	100							
	18.10	25.67	26.84	29.39	100							
	Count											
% of Row												
% of Col												
% of Total												
<table border="1"> <tr> <td>Chi-squared statistic:</td> <td>193.5</td> </tr> <tr> <td>df:</td> <td>12</td> </tr> <tr> <td>P-value:</td> <td><0.0001</td> </tr> </table>							Chi-squared statistic:	193.5	df:	12	P-value:	<0.0001
Chi-squared statistic:	193.5											
df:	12											
P-value:	<0.0001											

FIGURE 23.2

(Continued)

EXAMPLE 23.4 Where young people live: chi-square output

All three outputs tell us that the chi-square statistic is $\chi^2 = 193.5$, with very small P -value. Minitab reports $P = 0.000$, rounded to three decimal places, while CrunchIt! reports $P < 0.0001$. The graphing calculator says that $P = 6.98 \times 10^{-35}$, a very small number indeed. This P -value comes from an approximation to the sampling distribution of χ^2 . The approximation is less accurate far out in the tail than near the center of

the distribution, so you should not take 6.98×10^{-35} literally. Just read it as “P is very small.” The sample gives very strong evidence that living arrangements differ among the four age groups.

Statistical software generally offers additional information on request. We asked Minitab to show the observed counts and expected counts and also the term in the chi-square statistic for each cell, called the “chi-square contribution.” The top-left cell has expected count 245.57 and contributes 25.049 to the chi-square statistic, as we calculated earlier. (Roundoff errors are smaller with software than in hand calculation.) The graphing calculator also displays the observed and expected cell counts on request. We told the calculator to display the expected counts rounded to two decimal places. To see the remaining columns of expected counts on the calculator’s small screen, you must scroll to the right. CrunchIt! displays the counts, the row and column percents associated with each count, and the overall percent of the total associated with the count. Depending on the problem, only one of these percents will generally be of interest.

What about the Excel spreadsheet program? Excel is, in general, a poor choice for statistics. It is particularly awkward for chi-square because its Data Analysis tool pack omits this common test. You will need add-in modules to use Excel effectively for chi-square. ■

The chi-square test is an overall test for detecting relationships between two categorical variables. If the test is significant, it is important to look at the data to learn the nature of the relationship. We have three ways to look at the living arrangements data:

- **Compare selected percents:** which living arrangements occur in quite different percents of the four age groups? This is the method we learned in Chapter 6.
- **Compare observed and expected cell counts:** which cells have more or fewer observations than we would expect if H_0 were true?
- **Look at the terms of the chi-square statistic:** which cells contribute the most to the value of χ^2 ?

EXAMPLE 23.5 Where young people live: conclusion

There is very strong evidence ($\chi^2 = 193.55$, $P < 0.001$) that living arrangements of young people are not the same for ages 19, 20, 21, and 22. Comparing selected percents—specifically, the four conditional distributions of living arrangements for each age in Table 23.2 and Figure 23.1—shows how young people become more independent as they grow older.

The additional information provided by programs like Minitab shows what differences among the age groups explain the large value of the chi-square statistic. Look at the 20 terms in the chi-square statistic in the Minitab output and compare the observed and expected counts in the cells that contribute most to chi-square. Just 6 of the 20 cells contribute 166.79 of the total chi-square $\chi^2 = 193.55$. These 6 cells occur in pairs:

- 54.226 and 38.068: fewer 19-year-olds than expected and more 22-year-olds than expected live in their own place.

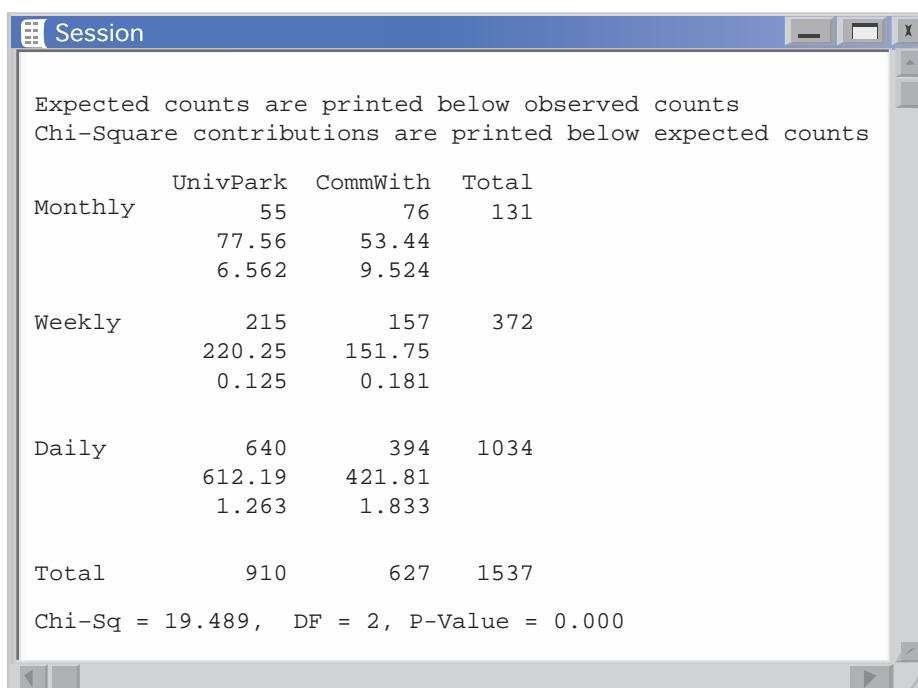
- 25.049 and 16.379: more 19-year-olds than expected and fewer 22-year-olds than expected live in their parents' home.
- 15.564 and 17.505: more 19-year-olds than expected and fewer 22-year-olds than expected live in group quarters.

These three trends display the increase in independent living between age 19 and age 22. ■

APPLY YOUR KNOWLEDGE

23.7 Facebook at Penn State. Figure 23.3 displays Minitab output for how frequently students at the University Park and commonwealth campuses of Penn State University who have Facebook accounts make use of their accounts. The output includes the two-way table of observed counts, the expected counts, and each cell's contribution to the chi-square statistic.  FACEBOOKUSERS

- Verify from the output that the data meet the cell count requirement for use of chi-square.
- What hypotheses does chi-square test? What are the test statistic and its P -value?
- Which cells contribute the most to χ^2 ? Compare the observed and expected counts in these cells and comment on the most important differences in Facebook use between students at the two locations.



Session

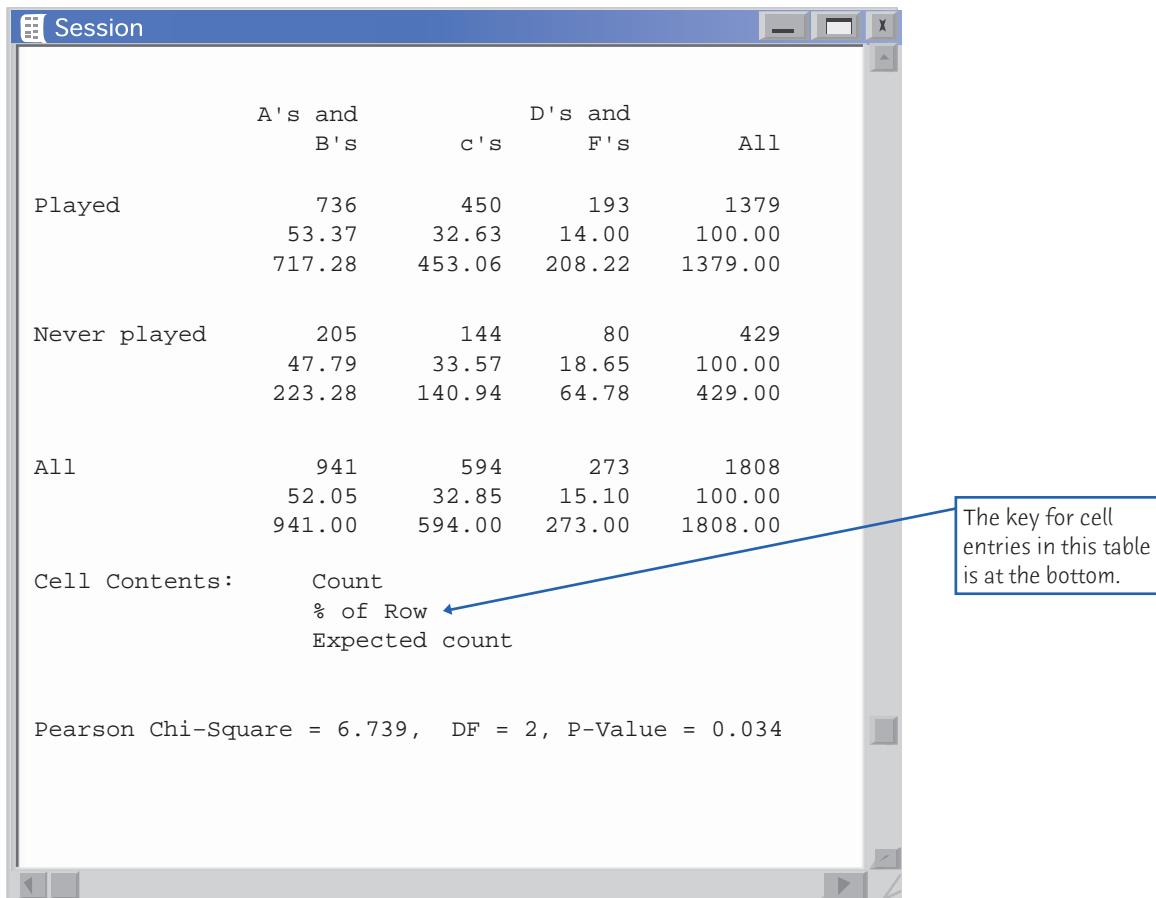
Expected counts are printed below observed counts
Chi-Square contributions are printed below expected counts

	UnivPark	CommWith	Total
Monthly	55	76	131
	77.56	53.44	
	6.562	9.524	
Weekly	215	157	372
	220.25	151.75	
	0.125	0.181	
Daily	640	394	1034
	612.19	421.81	
	1.263	1.833	
Total	910	627	1537
Chi-Sq = 19.489, DF = 2, P-Value = 0.000			

FIGURE 23.3

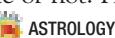
Minitab output for the two-way table of Facebook use and Penn State campus, for Exercise 23.7.

23.8 Video-gaming and grades. Your data analysis in Exercise 23.2 found that boys who have played video games tend to have higher grades than those who have not. Figure 23.4 gives Minitab output for the two-way table in Exercise 23.2.  GAMING

**FIGURE 23.4**

Minitab output for the study of video-gaming and grades, for Exercise 23.8.

- Verify from the output that the data meet the cell count requirement for use of chi-square.
- What are the chi-square statistic and its *P*-value? Explain in simple language what it means to reject H_0 in this setting.
- Give an overall conclusion that refers to row percents to describe the nature of the relationship between playing video games and grades.

23.9 Is astrology scientific? The General Social Survey asked a random sample of adults about their education and about their view of astrology as scientific or not. Here are the data for people with three levels of higher education degrees: 



Opinion	Degree Held		
	Junior college	Bachelor's	Graduate
Not at all scientific	87	198	111
Very or sort of scientific	43	57	28

Figure 23.5 gives Minitab chi-square output for these data. Follow the “Plan,” “Solve,” and “Conclude” steps of the four-step process in using the information in the output to describe how people with these levels of education differ in their opinions about astrology. Be sure that your “Solve” step includes data analysis and checking conditions for inference as well as a formal test.

	Junior			
	College	Bachelor	Graduate	All
NotScience	87 66.92 98.2 1.2869	198 77.65 192.7 0.1452	111 79.86 105.0 0.3375	396 75.57 396.0 *
Science	43 33.08 31.8 3.9814	57 22.35 62.3 0.4493	28 20.14 34.0 1.0441	128 24.43 128.0 *
All	130 100.00 130.0 *	255 100.00 255.0 *	139 100.00 139.0 *	524 100.00 524.0 *
Cell Contents:	Count % of Column Expected count Contribution to Chi-square			
Pearson Chi-Square = 7.244, DF = 2, P-Value = 0.027				

FIGURE 23.5

Minitab output for the two-way table of opinion about astrology and degree held, for Exercise 23.9.

USES OF THE CHI-SQUARE TEST

Two-way tables can arise in several ways. Most commonly, the subjects in a single sample are classified by two categorical variables. For example, we classified young adults by their age group and where they lived. The next example illustrates a different setting, in which we compare two separate samples. “Which sample” is now one of the variables for a two-way table.

EXAMPLE 23.6 Are cell-only telephone users different?

STATE: Random digit dialing telephone surveys do not call cell phone numbers. If the opinions of people who have only cell phones differ from those of people who still have



CELLPOLITICS



AB/Getty Images

landline service, the poll results may not represent the entire adult population. The Pew Research Center interviewed separate random samples of cell-only and landline telephone users. We will compare the 96 cell-only users and the 104 landline users who were less than 30 years old. Here's what the Pew survey found about how these people describe their political party affiliation.⁶

Party affiliation	Cell-only sample	Landline sample
Democrat or lean Democratic	49	47
Refuse to lean either way	15	27
Republican or lean Republican	32	30
Total	96	104

PLAN: Carry out a chi-square test for

H_0 : no relationship; that is, the distribution of party affiliation is the same in both populations

H_a : there is some relationship; that is, the party distribution in the cell-only population differs from that of landline users

Compare column percents or observed versus expected cell counts or terms of chi-square to see the nature of the relationship.

SOLVE: The Minitab output in Figure 23.6 includes the column percents. These give the conditional distributions of party given telephone use. Cell-only users are less likely to have no party affiliation (15.63% versus 25.96% of landline users). The party affiliations among the people who prefer one party are nearly the same in both groups, 60% Democrat for cell-only, 61% Democrat for landline.

To see if the differences are significant, first check the guideline for use of chi-square. The samples are reasonably close to SRSs, though nonresponse was higher for the cell phone calls. The Minitab output shows that all expected cell counts are greater than 5. The chi-square test shows that there is no significant difference between the party affiliations of the two groups of young adults ($\chi^2 = 3.22$, $P = 0.200$). Comparing observed and expected cell counts again shows that cell-only young adults are less likely to have no party preference than would be expected if there were no relationship, and that landline users are more likely to have no preference. The two “refuse to lean” cells contribute 2.54 of the total chi-square $\chi^2 = 3.22$. But the overall comparison is not significant.

CONCLUDE: There is no significant difference between the political party affiliations of young people who have a landline telephone and those who rely entirely on cell phones. The data do suggest that cell-only users are more likely to have some affiliation, so that a larger sample might find a significant difference. The Pew study found “little difference” to be true for all adults and for a variety of political questions. Traditional telephone sample surveys will live on, at least for a while. ■

	Cell-only	Landline	All
Democrat	49 51.04 46.08 0.1850	47 45.19 49.92 0.1708	96 48.00 96.00 *
RefuseToLean	15 15.63 20.16 1.3207	27 25.96 21.84 1.2191	42 21.00 42.00 *
Republican	32 33.33 29.76 0.1686	30 28.85 32.24 0.1556	62 31.00 62.00 *
All	96 100.00 96.00 *	104 100.00 104.00 *	200 100.00 200.00 *
Cell Contents:	Count % of Column Expected count Contribution to Chi-square		
Pearson Chi-Square = 3.220, DF = 2, P-Value = 0.200			

FIGURE 23.6

Minitab output for the two-way table of political party affiliation and telephone use, for Example 23.6.



More chi-square tests

There are other chi-square tests for hypotheses more specific than “no relationship.” A sociologist places people in classes by social status, waits ten years, then classifies the same people again.

The row and column variables are the classes at the two times. She might test the hypothesis that there has been no change in the overall distribution of social status in the group. Or she might ask if moves up in status are balanced by matching moves down. These and other null hypotheses can be tested by variations of the chi-square test.

One of the most useful properties of chi-square is that it tests the null hypothesis “the row and column variables are not related to each other” whenever this hypothesis makes sense for a two-way table. It makes sense when we are comparing a categorical response in two or more samples, as when we compared people who have only a cell phone with people who have a landline phone. The hypothesis also makes sense when we have data on two categorical variables for the individuals in a single sample, as when we examined age group and living arrangement for a sample of young adults. Statistical significance has the same meaning in both settings: “A relationship this strong is not likely to happen just by chance.”

USES OF THE CHI-SQUARE TEST

Use the chi-square test to test the null hypothesis

H_0 : there is no relationship between two categorical variables

when you have a two-way table from one of these situations:

- Independent SRSs from two or more populations, with each individual classified according to one categorical variable. (The other variable says which sample the individual comes from.)
- A single SRS, with each individual classified according to both of two categorical variables.



APPLY YOUR KNOWLEDGE

23.10 Cell-only versus landline users. We suspect that people who rely entirely on cell phones will as a group be younger than those who have a landline telephone. Do data confirm this guess? Here is a two-way table that breaks down both of Pew's samples (see Example 23.6) by age group: CELLAGE

Age (years)	Landline sample	Cell-only sample
18–29	104	96
30–49	265	70
50–64	204	26
65 or older	179	8
Total	752	200

Do a complete analysis of these data, following the four-step process as illustrated in Example 23.6.

THE CHI-SQUARE DISTRIBUTIONS

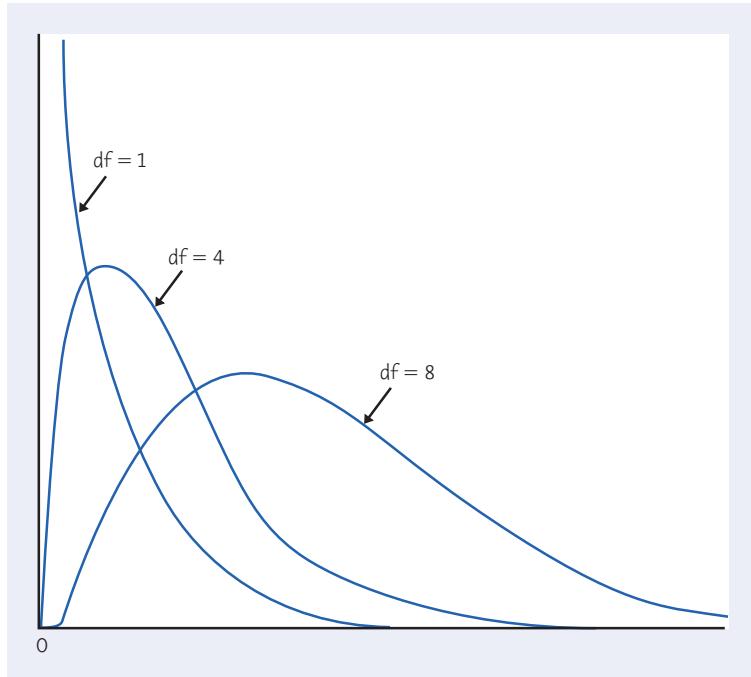
Software usually finds P -values for us. The P -value for a chi-square test comes from comparing the value of the chi-square statistic with critical values for a *chi-square distribution*.

THE CHI-SQUARE DISTRIBUTIONS

The **chi-square distributions** are a family of distributions that take only positive values and are skewed to the right. A specific chi-square distribution is specified by giving its **degrees of freedom**.

The chi-square test for a two-way table with r rows and c columns uses critical values from the chi-square distribution with $(r - 1)(c - 1)$ degrees of freedom. The P -value is the area under the density curve of this chi-square distribution to the right of the value of test statistic.

Figure 23.7 shows the density curves for three members of the chi-square family of distributions. As the degrees of freedom increase, the density curves become less skewed and larger values become more probable. Table D in the back of the book gives critical values for chi-square distributions. You can use Table D if you do not have software that gives you P -values for a chi-square test.

**FIGURE 23.7**

Density curves for the chi-square distributions with 1, 4, and 8 degrees of freedom. Chi-square distributions take only positive values and are right-skewed.

EXAMPLE 23.7 Using the chi-square table

The two-way table of 5 outcomes by 4 age groups for the living arrangements study (Table 23.1) has 5 rows and 4 columns. That is, $r = 5$ and $c = 4$. The chi-square statistic therefore has degrees of freedom

$$(r - 1)(c - 1) = (5 - 1)(4 - 1) = (4)(3) = 12$$

All three outputs in Figure 23.2 give 12 as the degrees of freedom.

The observed value of the chi-square statistic is $\chi^2 = 193.55$. Look in the $df = 12$ row of Table D. The value $\chi^2 = 193.55$ falls above the largest critical value in the table, for $P = 0.0005$. Remember that the chi-square test is always one-sided. So the P -value of $\chi^2 = 193.55$ is less than 0.0005. ■

$df = 12$

p	.001	.0005
χ^*	32.91	34.82

We know that all z and t statistics measure the size of an effect in the standard scale centered at zero. We can roughly assess the size of any z or t statistic by the 68–95–99.7 rule, though this is exact only for z . The chi-square statistic does not have any such natural interpretation. But here is a helpful fact: *the mean of any chi-square distribution is equal to its degrees of freedom*. In Example 23.7, χ^2 would have mean 12 if the null hypothesis were true. The observed value $\chi^2 = 193.55$ is so much larger than 12 that we suspect it is significant even before we look at Table D.



APPLY YOUR KNOWLEDGE

23.11 Facebook at Penn State. The Minitab output in Figure 23.3 (see page 565) gives the degrees of freedom for a table of Facebook use by students at two campus locations as $DF = 2$.

- Show that this is correct for a table with 3 rows and 2 columns.
- Minitab gives the chi-square statistic as $\text{Chi-Sq} = 19.489$. Where does this value fall when compared with critical values of the chi-square distribution with 2 degrees of freedom in Table D? How does Minitab's result $P\text{-value} = 0.000$ compare with the P -value from the table?
- The two-way table included only students who have Facebook accounts. The original table in Exercise 23.1 had 4 rows and 2 columns. What is the proper degrees of freedom for that table?

23.12 Video-gaming and grades. The Minitab output in Figure 23.4 gives 2 degrees of freedom for the table in Exercise 23.2.

- Verify that this is correct.
- The computer gives the value of the chi-square statistic as $\chi^2 = 6.739$. Between what two entries in Table D does this value lie? Verify that Minitab's P -value does fall between the tail probabilities p for these two entries.
- What is the mean value of the statistic χ^2 if the null hypothesis is true? How does the observed value of χ^2 compare with this mean?

THE CHI-SQUARE TEST FOR GOODNESS OF FIT*

The most common and most important use of the chi-square statistic is to test the hypothesis that there is *no relationship between two categorical variables*. A variation of the statistic can be used to test a different kind of null hypothesis: that *a categorical variable has a specified distribution*. Here is an example that illustrates this use of chi-square.



BIRTHDAY140

EXAMPLE 23.8 Never on Sunday?

Births are not evenly distributed across the days of the week. Fewer babies are born on Saturday and Sunday than on other days, probably because doctors find weekend births inconvenient.

A random sample of 140 births from local records shows this distribution across the days of the week:

Day	Sun.	Mon.	Tue.	Wed.	Thu.	Fri.	Sat.
Births	13	23	24	20	27	18	15

Sure enough, the two smallest counts of births are on Saturday and Sunday. Do these data give significant evidence that local births are not equally likely on all days of the week? ■

*This special topic is optional.

The chi-square test answers the question of Example 23.8 by comparing observed counts with expected counts under the null hypothesis. The null hypothesis for births says that they *are* evenly distributed. To state the hypotheses carefully, write the discrete probability distribution for days of birth:

Day	Sun.	Mon.	Tue.	Wed.	Thu.	Fri.	Sat.
Probability	p_1	p_2	p_3	p_4	p_5	p_6	p_7

The null hypothesis says that the probabilities are the same on all days. In that case, all 7 probabilities must be $1/7$. So the null hypothesis is

$$H_0: p_1 = p_2 = p_3 = p_4 = p_5 = p_6 = p_7 = \frac{1}{7}$$

The alternative hypothesis says that days are *not* all equally probable:

$$H_a: \text{not all } p_i = \frac{1}{7}$$

As usual in chi-square tests, H_a is a “many-sided” hypothesis that simply says that H_0 is not true. The chi-square statistic is also as usual:

$$\chi^2 = \sum \frac{(\text{observed count} - \text{expected count})^2}{\text{expected count}}$$

The expected count for an outcome with probability p is np , as we saw in the discussion following Example 23.2. Under the null hypothesis, all the probabilities p_i are the same, so all 7 expected counts are equal to

$$np_i = 140 \times \frac{1}{7} = 20$$

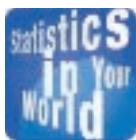
These expected counts easily satisfy our guideline for using chi-square. The chi-square statistic is

$$\begin{aligned} \chi^2 &= \sum \frac{(\text{observed count} - 20)^2}{20} \\ &= \frac{(13 - 20)^2}{20} + \frac{(23 - 20)^2}{20} + \dots + \frac{(15 - 20)^2}{20} \\ &= 7.6 \end{aligned}$$

This new use of χ^2 requires a different degrees of freedom. To find the P -value, compare χ^2 with critical values from the chi-square distribution with degrees of freedom 1 less than the number of values the birth day can take. That's $7 - 1 = 6$ degrees of freedom. From Table D, we see that $\chi^2 = 7.6$ is smaller than the smallest entry in the $df = 6$ row, which is the critical value for tail area 0.25. The P -value is therefore greater than 0.25 (software gives the more exact value $P = 0.269$). These 140 births don't give convincing evidence that births are not equally likely on all days of the week.

df = 6

p	.25	.20
χ^*	7.84	8.56



Chi-square in the casino

Gambling devices such as slot machines and

roulette wheels are supposed to have a fixed and known distribution of outcomes. Here's a job for the chi-square test of goodness of fit: state gambling regulators use it to verify that casino devices are honest. How much deviation a casino can get away with depends on the state. Nevada cracks down if chi-square is significant at the 5% level.

Mississippi gives more leeway, acting only when the 1% level is reached.

The chi-square test applied to the hypothesis that a categorical variable has a specified distribution is called the test for *goodness of fit*. The idea is that the test assesses whether the observed counts “fit” the distribution. The chi-square statistic is the same as for the two-way table test, but the expected counts and degrees of freedom are different. Here are the details.

THE CHI-SQUARE TEST FOR GOODNESS OF FIT

A categorical variable has k possible outcomes, with probabilities $p_1, p_2, p_3, \dots, p_k$. That is, p_i is the probability of the i th outcome. We have n independent observations from this categorical variable.

To test the null hypothesis that the probabilities have specified values

$$H_0: p_1 = p_{10}, p_2 = p_{20}, \dots, p_k = p_{k0}$$

find the **expected count** for the i th possible outcome as np_{i0} and use the **chi-square statistic**

$$\chi^2 = \sum \frac{(\text{observed count} - \text{expected count})^2}{\text{expected count}}$$

The sum is over all the possible outcomes.

The P -value is the area to the right of χ^2 under the density curve of the chi-square distribution with $k - 1$ degrees of freedom.

In Example 23.8, the outcomes are days of the week, with $k = 7$. The null hypothesis says that the probability of a birth on the i th day is $p_{i0} = 1/7$ for all days. We observe $n = 140$ births and count how many fall on each day. These are the counts used in the chi-square statistic.

APPLY YOUR KNOWLEDGE

23.13 Saving birds from windows. Many birds are injured or killed by flying into windows. It appears that birds don't see windows. Can tilting windows down so that they reflect earth rather than sky reduce bird strikes? Place six windows at the edge of a woods: two vertical, two tilted 20 degrees, and two tilted 40 degrees. During the next four months, there were 53 bird strikes, 31 on the vertical windows, 14 on the 20-degree windows, and 8 on the 40-degree windows.⁷ If the tilt has no effect, we expect strikes on windows with all three tilts to have equal probability. Test this null hypothesis. What do you conclude?

23.14 More on birth days. Births really are not evenly distributed across the days of the week. The data in Example 23.8 failed to reject this null hypothesis because of random variation in a quite small number of births. Here are data on 700 births in the same locale: **BIRTHDAYS700**

Day	Sun.	Mon.	Tue.	Wed.	Thu.	Fri.	Sat.
Births	84	110	124	104	94	112	72



Randy Duchaine/CORBIS

- The null hypothesis is that all days are equally probable. What are the probabilities specified by this null hypothesis? What are the expected counts for each day in 700 births?
- Calculate the chi-square statistic for goodness of fit.
- What are the degrees of freedom for this statistic? Do these 700 births give significant evidence that births are not equally probable on all days of the week?

23.15 Police harassment. Police may use minor violations such as not wearing a seat belt to stop motorists for other reasons. A large study in Michigan first studied the population of drivers not wearing seat belts during daylight hours by observation at more than 400 locations around the state. Here is the population distribution of seat belt violators by age group:⁸

Age group	16 to 29	30 to 59	60 or older
Proportion	0.328	0.594	0.078

The researchers then looked at court records and called a random sample of 803 drivers who had actually been cited by police for not wearing a seat belt. Here are the counts:

Age group	16 to 29	30 to 59	60 or older
Count	401	382	20

Does the age distribution of people cited differ significantly from the distribution of ages of all seat belt violators? Which age groups have the largest contributions to chi-square? Are these age groups cited more or less frequently than is justified? (The study found that males, blacks, and younger drivers were all overcited. This is an example in which the probabilities for the categories are not equal under the null hypothesis. You must use the given population probabilities and the sample size to compute the expected counts.)

23.16 Order in choice. Does the order in which wine is presented make a difference? Several choices of wine are presented one at a time, and subjects are asked to choose their preferred wine at the end of the sequence. In this study, subjects were asked to taste four wine samples in sequence. All four samples given to a subject were the *same* wine, although subjects were expecting to taste four different samples of a particular variety.⁹ There were 33 subjects in the study, and the positions in the sequence they selected for their preferred wine were



Position	1	2	3	4
Count	15	5	2	11

- What percent of the subjects chose each position?
- If the subjects are equally likely to select each position, what are the expected counts for each position?

- (c) Does the chi-square test for goodness of fit give good evidence that the subjects were not equally likely to choose each position? (State hypotheses, check the guideline for using chi-square, give the test statistic and its P -value, and state your conclusion.)
- (d) The *primacy* effect is a tendency for subjects to choose the first wine tasted, while the *recency* effect is a tendency for subjects to choose the most recent wine tasted. Are either of these effects present in these data?

23.17 What's your sign? For reasons known only to social scientists, the General Social Survey (GSS) regularly asks its subjects about their astrological sign. Here are the counts of responses for the 2008 GSS:



Sign	Aries	Taurus	Gemini	Cancer	Leo	Virgo
Count	164	152	159	167	157	201
Sign	Libra	Scorpio	Sagittarius	Capricorn	Aquarius	Pisces
Count	175	147	147	151	163	177

If births are spread uniformly across the year, we expect all 12 signs to be equally likely. Are they? Follow the four-step process in your answer.

CHAPTER 23 SUMMARY

CHAPTER SPECIFICS

- The **chi-square test** for a two-way table tests the null hypothesis H_0 that there is no relationship between the row variable and the column variable. The alternative hypothesis H_a says that there is some relationship but does not say what kind.
- The test compares the observed counts of observations in the cells of the table with the counts that would be expected if H_0 were true. The **expected count** in any cell is

$$\text{expected count} = \frac{\text{row total} \times \text{column total}}{\text{table total}}$$

- The **chi-square statistic** is

$$\chi^2 = \sum \frac{(\text{observed count} - \text{expected count})^2}{\text{expected count}}$$

- The chi-square test compares the value of the statistic χ^2 with critical values from the **chi-square distribution** with $(r - 1)(c - 1)$ degrees of freedom. Large values of χ^2 are evidence against H_0 , so the P -value is the area under the chi-square density curve to the right of χ^2 .
- The chi-square distribution is an approximation to the distribution of the statistic χ^2 . You can safely use this approximation when all expected cell counts are at least 1 and no more than 20% are less than 5.
- If the chi-square test finds a statistically significant relationship between the row and column variables in a two-way table, do data analysis to describe the nature of the relationship. You can do this by comparing well-chosen percents, comparing the observed counts with the expected counts, and looking for the largest **terms of the chi-square statistic**.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Two-Way Tables

1. Understand that the data for a chi-square test must be presented as a two-way table of counts of outcomes.
2. Use percents to describe the relationship between any two categorical variables, starting from the counts in a two-way table.

B. Interpreting Chi-Square Tests

1. Locate the chi-square statistic, its P -value, and other useful facts (row or column percents, expected counts, terms of chi-square) in output from your software or calculator.
2. Use the expected counts to check whether you can safely use the chi-square test.
3. Explain what null hypothesis the chi-square statistic tests in a specific two-way table.
4. If the test is significant, compare percents, compare observed with expected cell counts, or look for the largest terms of the chi-square statistic to see what deviations from the null hypothesis are most important.

C. Doing Chi-Square Tests by Hand

1. Calculate the expected count for any cell from the observed counts in a two-way table. Check whether you can safely use the chi-square test.
2. Calculate the term of the chi-square statistic for any cell, as well as the overall statistic.
3. Give the degrees of freedom of a chi-square statistic. Make a quick assessment of the significance of the statistic by comparing the observed value with the degrees of freedom.
4. Use the chi-square critical values in Table D to approximate the P -value of a chi-square test.

LINK IT

Part IV of the text studies relationships between variables. Relationships between two quantitative variables were introduced in Chapters 4 and 5, and these will be described in greater detail in the next chapter. In this chapter, the case of two categorical variables is considered, and a formal test for answering the question “Is there a relationship between the two categorical variables?” is developed. As with procedures described in earlier chapters, we must first consider how the data were produced, as this plays an important role in the conclusions we can reach. Were the data produced by an experiment or an observational study? If it is an observational study, are there lurking variables that can explain the observed relationship? In addition, we should begin with data analysis rather than a formal test. In the case of two-way tables, this typically involves looking at conditional distributions, both numerically and graphically, in order to first understand the nature of the relationship. When considering the relationship between

the age of young adults and their living arrangements in Example 23.1, we can see from our data analysis that as young adults age from 19 to 22, the percent living with their parents drops as the percent living in their own place rises.

Even though there appears to be a clear relationship between age and living arrangement in Example 23.1, we must still determine whether the observed differences are large enough to be statistically significant. The chi-square test can be used for this, but it is an approximate procedure, and the conditions for cell sizes need to be checked before applying the test. If the differences are statistically significant, the chi-square test, unlike some of the simpler procedures in earlier chapters, tells us only that there is evidence of a relationship, not the nature of the relationship. Although there are formal statistical procedures to further investigate the nature of the relationship, at this point we need to be satisfied with describing the relationship between the two categorical variables using our data analysis tools.

CHECK YOUR SKILLS

Resistance training is a popular form of conditioning aimed at enhancing sports performance and is widely used among high school, college, and professional athletes, although its use for younger athletes is controversial. A random sample of 4111 patients between the ages of 8 and 30 admitted to U.S. emergency rooms with the injury code “weightlifting” was obtained. These injuries were classified as “accidental” if caused by dropped weight or improper equipment use. The patients were also classified into the four age categories “8–13,” “14–18,” “19–22,” and “23–30.” Here is a two-way table of the results:¹⁰



CORBIS/Superstock

Age	Accidental	Not accidental
8–13	295	102
14–18	655	916
19–22	239	533
23–30	363	1008

- 23.18** The number of “accidental” injuries in the sample is
 (a) 1552. (b) 2559. (c) 4111.

23.19 The percent of the 14- to 18-year-olds in the sample whose injuries were classified as “accidental” is about WEIGHTLIFTING

- (a) 42.2%. (b) 41.7%. (c) 74.3%.

23.20 The percent of the 14- to 18-year-olds in the sample whose injuries were classified as “accidental” is

- (a) higher than the percent for 23- to 30-year-olds.
 (b) about the same as the percent for 23- to 30-year-olds.
 (c) lower than the percent for 23- to 30-year-olds.

23.21 The expected count of 14- to 18-year-olds whose injuries were classified as “accidental” is about

- (a) 593.09. (b) 655. (c) 977.91.

23.22 The term in the chi-square statistic for the cell of 14- to 18-year-olds whose injuries were classified as “accidental” is about

- (a) 593.09. (b) 3.919. (c) 6.463.

23.23 The degrees of freedom for the chi-square test for this two-way table are

- (a) 3. (b) 4. (c) 8.

23.24 The null hypothesis for the chi-square test for this two-way table is

- (a) The proportions of “Accidental” and “Not accidental” injuries are the same.
 (b) There is no difference in the probabilities of an “accidental” injury for each of the four age groups.
 (c) “Accidental” injuries are more likely for the younger age groups.

23.25 The alternative hypothesis for the chi-square test for this two-way table is

- (a) The proportions of “Accidental” and “Not accidental” injuries are different.
- (b) The probabilities of an “accidental” injury for each of the four age groups are not the same.
- (c) “Accidental” injuries are more likely for the younger age groups.

23.26 Software gives chi-square statistic $\chi^2 = 325.459$ for this table. From the table of critical values, we can say that the P -value is

- (a) between 0.0025 and 0.001.
- (b) between 0.001 and 0.0005.
- (c) less than 0.0005.

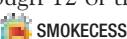
23.27 The most important fact that allows us to trust the results of the chi-square test is that

- (a) the sample is large, 4111 weight-lifting injuries in all.
- (b) the sample is close to an SRS of all weight-lifting injuries.
- (c) all the cell counts are greater than 100.

CHAPTER 23 EXERCISES

If you have access to software or a graphing calculator, use it to speed your analysis of the data in these exercises. Exercises 23.28 to 23.33 are suitable for hand calculation if necessary.

23.28 Smoking cessation. A large randomized trial was conducted to assess the efficacy of Chantix for smoking cessation compared with bupropion (more commonly known as Wellbutrin or Zyban) and a placebo. Chantix is different from most other quit-smoking products in that it targets nicotine receptors in the brain, attaches to them, and blocks nicotine from reaching them, while bupropion is an antidepressant often used to help people stop smoking. Generally healthy smokers who smoked at least 10 cigarettes per day were assigned at random to take Chantix ($n = 352$), bupropion ($n = 329$), or a placebo ($n = 344$). The study was double-blind, with the response measure being continuous cessation from smoking for Weeks 9 through 12 of the study. Here is a two-way table of the results:¹¹



	Treatment		
	Chantix	Bupropion	Placebo
No smoking in Weeks 9–12	155	97	61
Smoked in Weeks 9–12	197	232	283

(a) Give a 95% confidence interval for the difference between the proportions of smokers in the bupropion and placebo groups who did not smoke in Weeks 9 through 12 of the study.

(b) What proportion of each of the three groups in the sample did not smoke in Weeks 9 through 12 of the study? Are there statistically significant differences among these proportions? State hypotheses and give a test statistic and its P -value.

(c) Is this an observational study or an experiment? Why does this make a difference in the type of conclusion we can draw?

23.29 Attitudes toward recycled products. Some people think recycled products are lower in quality than other products, a fact that makes recycling less practical. Here are data on attitudes toward coffee filters made of recycled paper.¹²



RECYCLING

	Think the quality of the recycled product is		
	Higher	Same	Lower
Buyers	20	7	9
Nonbuyers	29	25	43

(a) Find the conditional distributions of opinions on the quality of recycled products for buyers and nonbuyers. Make a graph that compares the two conditional distributions. Use your work to describe the overall relationship between people who have and haven't bought recycled filters and their opinions on the quality of recycled products.

(b) Do buyers and nonbuyers of recycled filters differ significantly in their opinions on the quality of recycled products? State hypotheses, give the chi-square statistic and its P -value, and state your conclusion.

(c) Association does not prove causation. Explain how buying recycled filters might improve a person's opinion of their quality. Then explain how the opinion a person holds might influence his or her decision to buy or not. You see that the cause-and-effect relationship might go in either direction.

23.30 Do you use cocaine? Sample surveys on sensitive issues can give different results depending on how the question is asked. A University of Wisconsin study divided 2400 respondents into three groups at random. All were asked if they had ever used cocaine. One group of 800 was interviewed by phone; 21% said they had used cocaine. Another 800 people were asked the question in a one-on-one personal interview; 25% said “Yes.”

The remaining 800 were allowed to make an anonymous written response; 28% said "Yes."¹³ Are there statistically significant differences among these proportions? State the hypotheses, convert the information given into a two-way table of counts, give the test statistic and its P -value, and state your conclusions.

23.31 Did the randomization work? After randomly assigning subjects to treatments in a randomized comparative experiment, we can compare the treatment groups to see how well the randomization worked. We hope to find no significant differences among the groups. A study of how to provide premature infants with a substance essential to their development assigned infants at random to receive one of four types of supplement, called PBM, NLCP, PL-LCP, and TG-LCP.¹⁴

(a) The subjects were 77 premature infants. Outline the design of the experiment if 20 are assigned to the PBM group and 19 to each of the other treatments.

(b) The random assignment resulted in 9 females in the TG-LCP group and 11 females in each of the other groups. Make a two-way table of group by gender and do a chi-square test to see if there are significant differences among the groups. What do you find?

23.32 More on video-gaming. The data for comparing two sample proportions can be presented in a two-way table containing the counts of successes and failures in both samples, with two rows and two columns. In Exercise 23.2, a survey of the consequences of video-gaming on 14- to 18-year-olds is described. Another question from the survey was about aggressive behavior as evidenced by getting into serious fights, and the comparison was between girls that have and have not played video games. Here are the data:

		Serious Fights	
		Yes	No
Played games	36	55	
Never played games	578	1436	

(a) Is there evidence that the proportions of all 14- to 18-year-old girls who played or have never played video games and have gotten into serious fights differ? Find the two sample proportions, the z statistic, and its P -value.

(b) Is there evidence that the proportions of 14- to 18-year-old girls who have or have not gotten into serious fights differ between those who have played or have never played video games? Find the chi-square statistic χ^2 and its P -value.

(c) Show that (up to roundoff error) your χ^2 is the same as z^2 . The two P -values are also the same. These facts are always true, so you will often see chi-square for 2×2 tables used to compare two proportions.

(d) Suppose that we are interested in finding out if the data give good evidence that video-gaming is associated with increased aggression in girls as evidenced by getting into serious fights. Can we use the z test for this hypothesis? What about the χ^2 test? What is the important difference between these two procedures?

23.33 Unhappy rats and tumors. Some people think that the attitude of cancer patients can influence the progress of their disease. We can't experiment with humans, but here is a rat experiment on this theme. Inject 60 rats with tumor cells and then divide them at random into two groups of 30. All the rats receive electric shocks, but rats in Group 1 can end the shock by pressing a lever. (Rats learn this sort of thing quickly.) The rats in Group 2 cannot control the shocks, which presumably makes them feel helpless and unhappy. We suspect that the rats in Group 1 will develop fewer tumors. The results: 11 of the Group 1 rats and 22 of the Group 2 rats developed tumors.¹⁵

(a) Make a two-way table of tumors by group. State the null and alternative hypotheses for this investigation.

(b) Although we have a two-way table, the chi-square test can't test a one-sided alternative. Carry out the z test and report your conclusion.

23.34 I think I'll be rich by age 30. A sample survey asked  young adults (aged 19 to 25), "What do you think are the chances you will have much more than a middle-class income at age 30?" The Minitab output in Figure 23.8 shows the two-way table and related information, omitting a few subjects who refused to respond or who said they were already rich.¹⁶ Use the output as the basis for a discussion of the differences between young men and young women in assessing their chances of being rich by age 30.  RICHBY30

23.35 Sexy magazine ads? Look at full-page ads in magazines with a young adult readership. Classify ads that show a model as "not sexual" or "sexual" depending on how the model is dressed (or not dressed). Here are data on 1509 ads in magazines aimed at young men only, at young women only, or at young adults in general:¹⁷  SEXYADS

Ad type	Readers		
	Men	Women	General
Sexual	105	225	66
Not sexual	514	351	248

Figure 23.9 displays Minitab chi-square output. Use the information in the output to describe the relationship between the target audience and the sexual content of ads in magazines for young adults.

	Female	Male
A: Almost no chance	96 95.2 0.0076	98 98.8 0.0073
B: Some chance but probably not	426 349.2 16.8842	286 362.8 16.2525
C: A 50 50 chance	696 694.5 0.0032	720 721.5 0.0031
D: A good chance	663 697.0 1.6543	758 724.0 1.5924
E: Almost certain	486 531.2 3.8424	597 551.8 3.6986
Cell Contents:	Count Expected count Contribution to Chi-square	
Pearson Chi-Square =	43.946	DF = 4, P-Value = 0.000

FIGURE 23.8

Minitab output for the sample survey responses of Exercise 23.34.

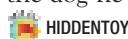
	Men	Women	General	All
Sexual	105 16.96 162.4 20.312	225 39.06 151.2 36.074	66 21.02 82.4 3.265	396 26.24 396.0 *
Not sexual	514 83.04 456.6 7.227	351 60.94 424.8 12.835	248 78.98 231.6 1.162	1113 73.76 1113.0 *
All	619 100.00 619.0 *	576 100.00 576.0 *	314 100.00 314.0 *	1509 100.00 1509.0 *
Cell Contents:	Count % of Column Expected count Contribution to Chi-square			
Pearson Chi-Square =	80.874	DF = 2, P-Value = 0.00		

FIGURE 23.9

Minitab output for a study of ads in magazines, for Exercise 23.35.

Mistakes in using the chi-square test are unusually common. Exercises 23.36 to 23.39 illustrate several kinds of mistake.

23.36 Sorry, no chi-square. An experimenter hid a toy from a dog behind either Screen A or Screen B. In the first phase the toy was always hidden behind Screen A, while in the second phase the toy was always hidden behind Screen B. Will the dog continue to look behind Screen A in the second phase? This was tried under three conditions. In the Social-Communicative condition the experimenter communicated with the dog by establishing eye contact and addressing the dog while hiding the toy; in the Noncommunicative condition the toy was hidden without communication; and in the Nonsocial condition the toy was dragged by a string so that it could be hidden without any interaction from the experimenter. There were 12 dogs assigned at random to each condition, and each dog had up to three trials to find the toy hidden behind Screen B in Phase 2. An error occurred if the dog continued to search behind Screen A. The number of errors ranged from 0 if the dog found the toy behind Screen B on the initial trial up to 3 if the dog never correctly chose Screen B. Here are the data:¹⁸



Condition	Number of Errors			
	0	1	2	3
Social-Communicative	0	3	3	6
Noncommunicative	5	3	1	3
Nonsocial	8	2	2	0

- (a) The data do show a difference in the number of errors for the different conditions. Show this by comparing suitable percents.
- (b) The researchers used a more complicated but exact procedure rather than chi-square to assess significance for these data. Why can't the chi-square test be trusted in this case?
- (c) If you use software, does the chi-square output for these data warn you against using the test?

23.37 Sorry, no chi-square. How do U.S. residents who travel overseas for leisure differ from those who travel for business? Here is the breakdown by occupation:¹⁹

Occupation	Leisure travelers	Business travelers
Professional/technical	36%	39%
Manager/executive	23%	48%
Retired	14%	3%
Student	7%	3%
Other	20%	7%
Total	100%	100%

Explain why we don't have enough information to use the chi-square test to learn whether these two distributions differ significantly.

23.38 Sorry, no chi-square. Here is more information about Internet use by students at Penn State, based on a random sample of 1852 undergraduates. Explain why it is not correct to use a chi-square test on this table to compare the University Park and commonwealth campuses. Note that in order to use the chi-square test in a two-way table, each individual must fall into one cell of the table.

Internet use	University Park	Commonwealth
Viewed a video on YouTube or similar site	875	700
Legally purchased music or videos online	514	348
Downloaded a podcast	235	145
Participated in Internet gambling	114	93

23.39 Sorry, no chi-square. Does eating chocolate trigger headaches? To find out, women with chronic headaches followed the same diet except for eating chocolate bars and carob bars that looked and tasted the same. Each subject ate both chocolate and carob bars in random order with at least three days between. Each woman then reported whether or not she had a headache within 12 hours of eating the bar. Here is a two-way table of the results for the 64 subjects:²⁰

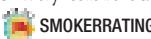
Bar	No headache	Headache
Chocolate	53	11
Carob (placebo)	38	26

The researchers carried out a chi-square test on this table to see if the two types of bar differ in triggering headaches. Explain why this test is incorrect. (Hint: There are 64 subjects. How many observations appear in the two-way table?)

The remaining exercises concern larger tables that require software for easy analysis. In many cases, you should follow the Plan, Solve, and Conclude steps of the four-step process in your answers.

23.40 Smokers rate their health. The University of Michigan Health and Retirement Study (HRS) surveys more than 22,000 Americans over the age of 50 every two years. A subsample of the HRS participated in the 2009 Internet-based survey that collected information on a number of topical areas, including health (physical and mental, health behaviors), psychosocial items, economics (income, assets, expectations, and consumption), and retirement.²¹ Two of

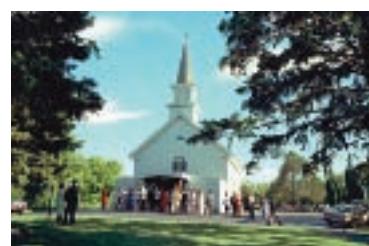
the questions asked on the Internet survey were “Would you say your health is excellent, very good, good, fair or poor?” and “Do you smoke cigarettes now?” The two-way table summarizes the answers to these two questions.



Health	Current Smoker	
	Yes	No
Excellent	25	484
Very good	115	1557
Good	145	1309
Fair	90	545
Poor	29	11

- (a) Regard the HRS Internet sample as approximately an SRS of Americans over the age of 50, and give a 99% confidence interval for the proportion of Americans over the age of 50 who are current smokers.
- (b) Compare the conditional distributions of self-evaluation of health for current smokers and nonsmokers using both a table and a graph. What are the most important differences?
- (c) Carry out the chi-square test for the hypothesis of no difference between the self-evaluation of health for current smokers and nonsmokers. What would be the mean of the test statistic if the null hypothesis were true? The value of the statistic is so far above this mean that you can see at once that it must be highly significant. What is the approximate P -value?
- (d) Look at the terms of the chi-square statistic and compare observed and expected counts in the cells that contribute the most to chi-square. Based on this and your findings in part (b), write a short comparison of the differences in self-evaluation of health for current smokers and nonsmokers.

23.41 Who goes to religious services? The General Social Survey (GSS) asked this question: “Have you attended religious services in the last week?” Here are the responses for those whose highest degree was high school or above:



Fotosearch/Superstock

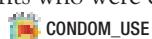
	Highest Degree Held			
	High school	Junior college	Bachelor's	Graduate
Attended services	400	62	146	76
Did not attend services	880	101	232	105

- (a) Carry out the chi-square test for the hypothesis of no relationship between the highest degree attained and attendance at religious services in the last week. What do you conclude?
- (b) Make a 2×3 table by omitting the column corresponding to those whose highest degree was high school. Carry out the chi-square test for the hypothesis of no relationship between the type of advanced degree attained and attendance at religious services in the last week. What do you conclude?
- (c) Make a 2×2 table by combining the counts in the three columns that have a highest degree beyond high school, so that you are comparing adults whose highest degree was high school with those whose highest degree was beyond high school. Carry out the chi-square test for the hypothesis of no relationship between attaining a degree beyond high school and attendance at religious services for this 2×2 table. What do you conclude?

- (d) Using the results from these three chi-square tests, write a short report explaining the relationship between attendance at religious services in the last week and the highest degree attained. As part of your report, you should give the percents who attended religious services for each of the four degrees.

23.42 Condom usage among high school students.

The Centers for Disease Control developed the Youth Risk Behavior Surveillance System (YRBSS) to monitor six categories of priority health risk behaviors among youth: behaviors that contribute to unintentional injuries and violence; tobacco use; alcohol and other drug use; sexual behaviors that contribute to unintended pregnancy and sexually transmitted diseases; unhealthy dietary behaviors; and physical inactivity. A multistage sample design is used to produce representative samples of students in grades 9 to 12, who then fill out a questionnaire on these behaviors. The data below are for the question “Did Not Use a Condom during Last Sexual Intercourse?” The two-way table of grade and condom usage includes only students who were currently sexually active. Here are the results:²²



Grade	Condom Used	
	Yes	No
9th	300	532
10th	350	736
11th	601	956
12th	873	1068

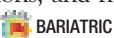
Describe the most important differences between condom usage and grade. Is there a significant overall difference between the proportions who used condoms in the different grades?

23.43 How are schools doing? The nonprofit group Public Agenda conducted telephone interviews with a stratified sample of parents of high school children. There were 202 black parents, 202 Hispanic parents, and 201 white parents. One question asked was “Are the high schools in your state doing an excellent, good, fair or poor job, or don’t you know enough to say?” Here are the survey results:²³

Opinion	Black parents	Hispanic parents	White parents
Excellent	12	34	22
Good	69	55	81
Fair	75	61	60
Poor	24	24	24
Don’t know	22	28	14
Total	202	202	201

Are the differences in the distributions of responses for the three groups of parents statistically significant? What departures from the null hypothesis “no relationship between group and response” contribute most to the value of the chi-square statistic? Write a brief conclusion based on your analysis.

23.44 Complications of bariatric surgery. Bariatric surgery, or weight-loss surgery, includes a variety of procedures performed on people who are obese. Weight loss is achieved by reducing the size of the stomach with an implanted medical device (gastric banding), through removal of a portion of the stomach (sleeve gastrectomy), or by resecting and rerouting the small intestines to a small stomach pouch (gastric bypass surgery). Because there can be complications using any of these methods, the National Institutes of Health recommends bariatric surgery for obese people with a body mass index (BMI) of at least 40 and for people with a BMI of at least 35 and serious coexisting medical conditions such as diabetes. Serious complications include potentially life-threatening, permanently disabling, and fatal outcomes. Here is a two-way table for data collected in Michigan over several years giving counts of non-life-threatening complications, serious complications, and no complications for these three types of surgeries:²⁴



Type of surgery	Type of Complication			Total
	Non-life-threatening	Serious	None	
Gastric banding	81	46	5253	5380
Sleeve gastrectomy	31	19	804	854
Gastric bypass	606	325	8110	9041

- (a) Is this study an experiment? Explain your answer.
 (b) Is there a significant difference in the distributions of type of complication for the three types of surgery? Which surgeries have the greatest chance of complications? Can we conclude that it is the surgery that is more dangerous, or could there be other factors associated with the increased risk?

23.45 Market research. Before bringing a new product to market, firms carry out extensive studies to learn how consumers react to the product and how best to advertise its advantages. Here are data from a study of a new laundry detergent.²⁵

The subjects are people who don’t currently use the established brand that the new product will compete with. Give subjects free samples of both detergents. After they have tried both for a while, ask which they prefer. The answers may depend on other facts about how people do laundry.



Preference	Laundry Practices			
	Soft water, warm wash	Soft water, hot wash	Hard water, warm wash	Hard water, hot wash
Prefer standard product	53	27	42	30
Prefer new product	63	29	68	42

How do laundry practices (water hardness and wash temperature) influence the choice of detergent? In which settings does the new detergent do best? Are the differences between the detergents statistically significant?

Support for political parties. Political parties want to know what groups of people support them. The General Social Survey (GSS) asked its 2008 sample, “Generally speaking, do you usually think of yourself as a Republican, Democrat, Independent, or what?” The GSS is essentially an SRS of American adults. Here is a large two-way table breaking down the responses by the highest degree the subject held:

Party support	Highest Degree Held				
	None	High school	Jr. college	Bachelor's	Graduate
Strong Democrat	63	185	32	56	54
Not strong Democrat	45	183	30	44	29
Independent, near Democrat	34	132	19	57	20
Independent	87	156	22	31	26
Independent, near Republican	19	78	20	35	10
Not strong Republican	20	147	33	76	27
Strong Republican	25	98	13	43	22
Other party	2	16	4	11	5

Exercises 23.46 to 23.48 are based on this table.

23.46 Other parties. Give a 95% confidence interval for the proportion of adults who are “Independent.”

23.47 Party support in brief. Make a 2×5 table by combining the counts in the three rows that mention “Democrat” and in the three rows that mention “Republican” and ignoring strict independents and supporters of other parties. We might think of this table as comparing all adults who lean Democrat and all adults who lean Republican. How

does support for the two major parties differ among adults with different levels of education?  POLPARTYCOMBINE

23.48 Party support in full. Use the full table to analyze the differences in political party support among levels of education. The sample is so large that the differences are bound to be highly significant, but give the chi-square statistic and its P -value nonetheless. The main challenge is in seeing what the data say. Does the full table yield any insights not found in the compressed table you analyzed in the previous exercise?  POLPARTYFULL



EXPLORING THE WEB

23.49 Make your own table. The Behavioral Risk Factor Surveillance System (BRFSS) is an ongoing data collection program designed to measure behavioral risk factors for the adult population (18 years of age or older) living in households. Data are collected from a random sample of adults (one per household) through a telephone survey. Go to the Web site apps.nccd.cdc.gov/BRFSS/ and under BRFSS Contents click on Web Enabled Analysis Tool (WEAT) and then click on Cross Tabulation Analysis. After selecting a year, a window will open that will allow you to produce two-way tables.

(a) Choose a state of interest to you and two variables for the two-way table. For example, you could choose Connecticut and look at the relationship between a demographic variable such as education level and a variable such as health care coverage. Once you have chosen your state and two variables, click on *run report* at the bottom of the page. A two-way table will appear in a new window.

(b) Is there a relationship between the two variables you selected? If the relationship is statistically significant, describe the relationship in a brief report using percents from the table and an appropriate graph.

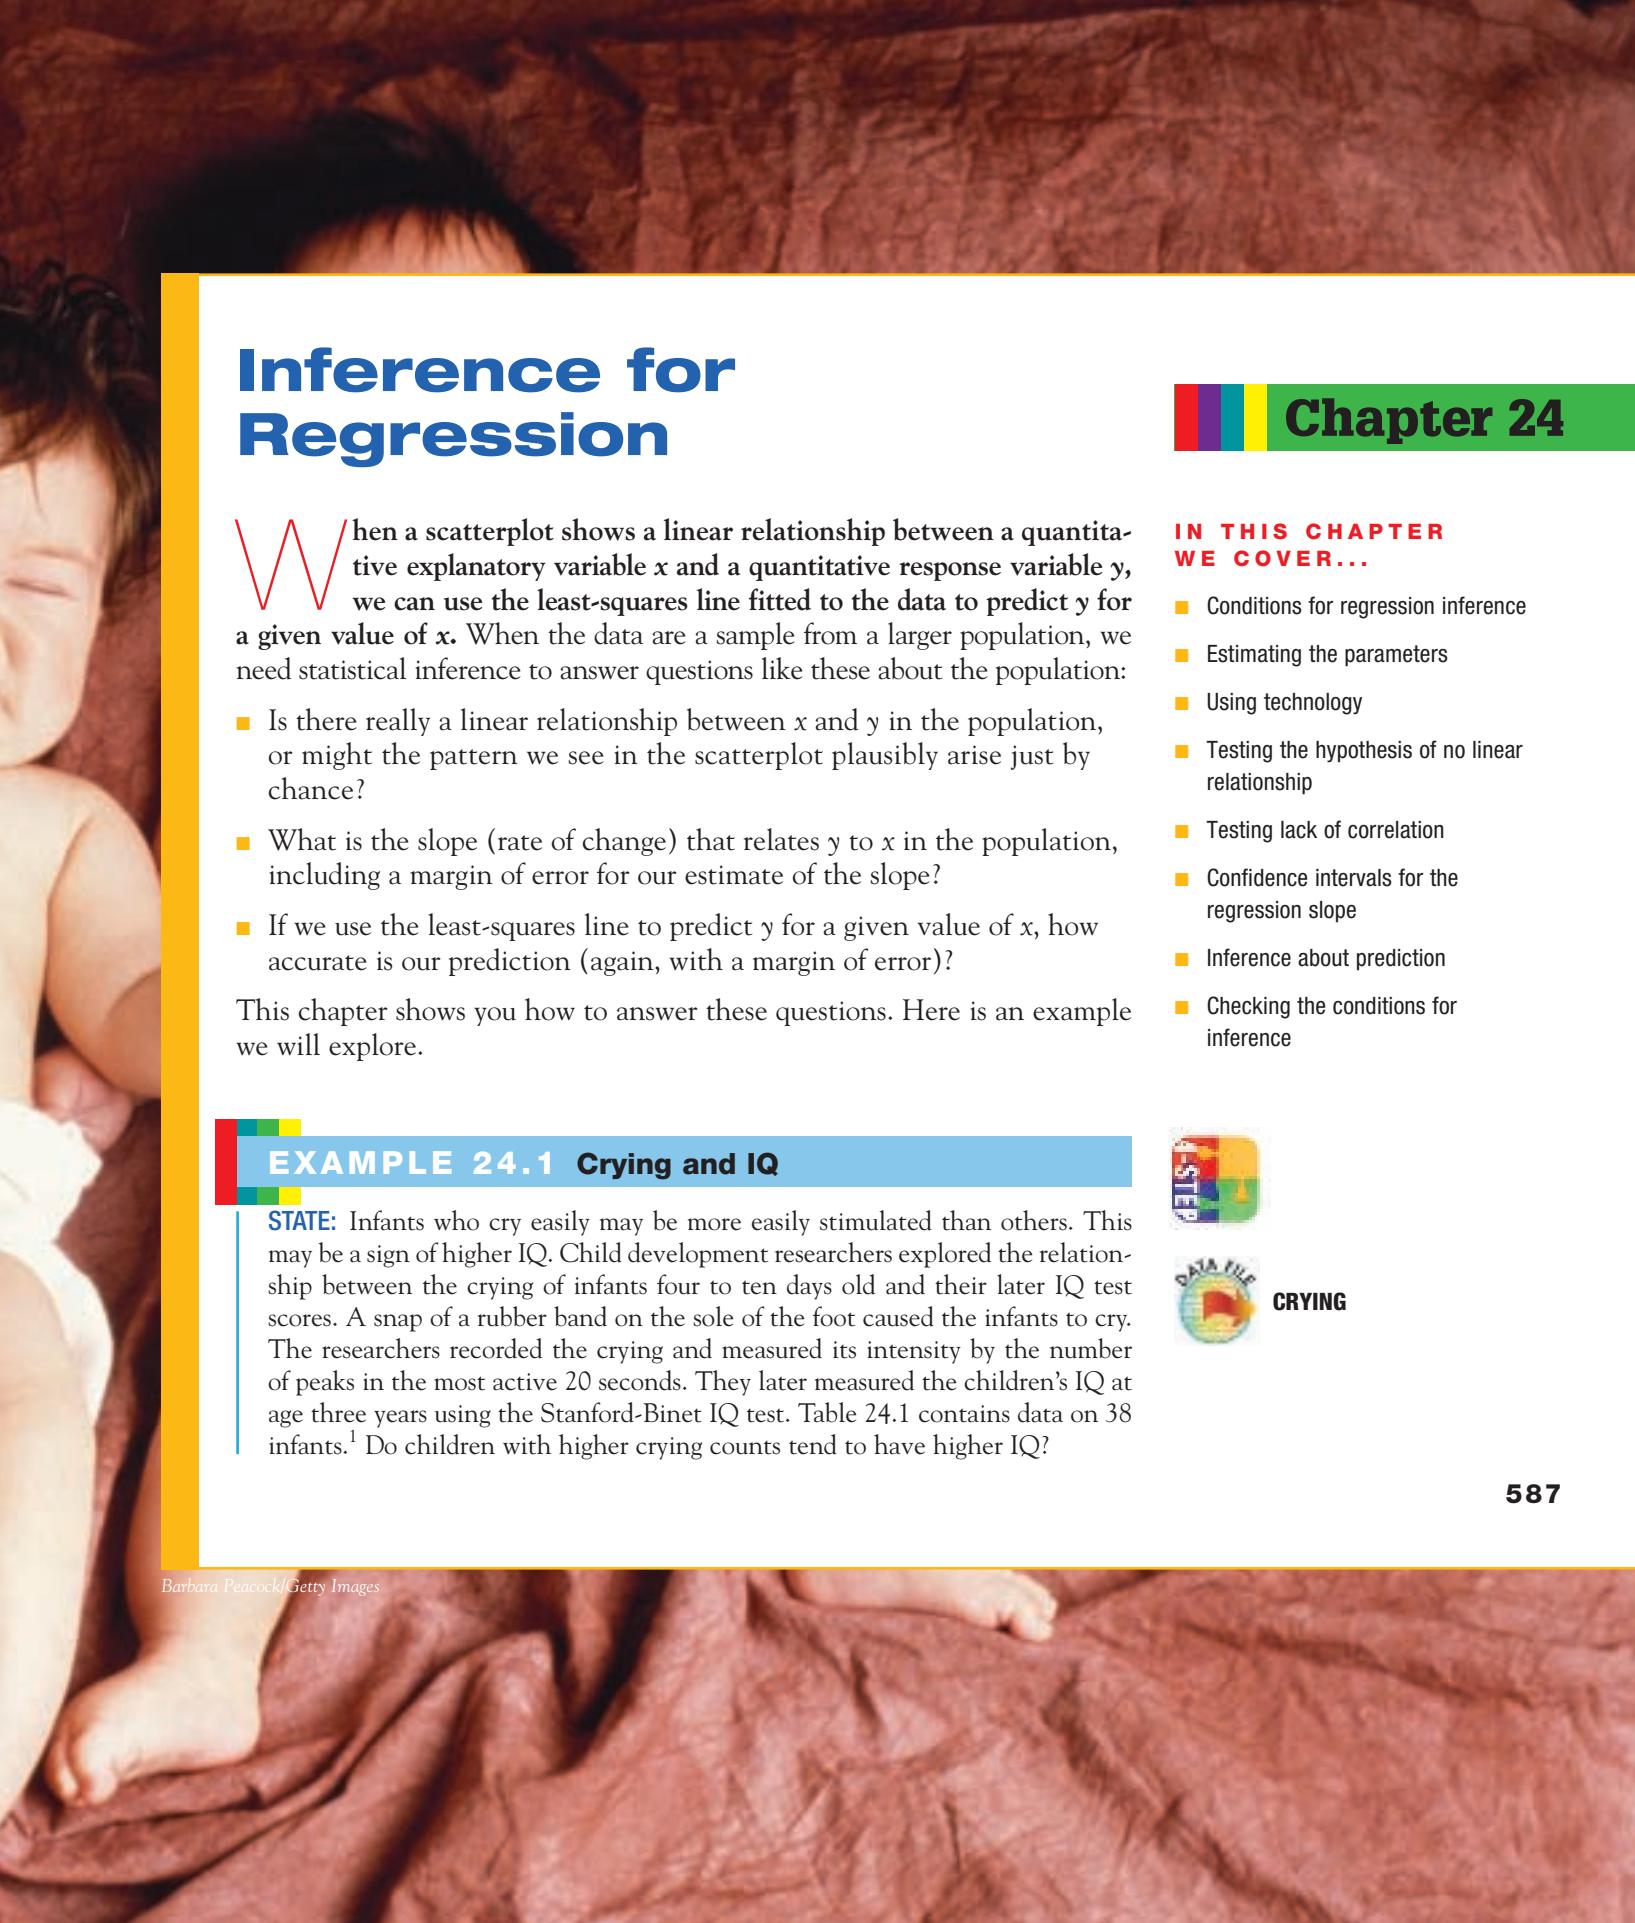
23.50 What do the voters think? The American National Election Studies (ANES) is the leading academically run national survey of voters in the United States and is conducted before and after every presidential election. SDA (Survey Documentation and Analysis) is a set of programs that allows you to analyze survey data and includes the ANES survey as part of its archive. Go to the Web site sda.berkeley.edu/ and click on Archive. Go to the 2008 ANES survey.

(a) Open the pre-election survey data. Under Liberal/Conservative, choose the variable “liberal/conservative self-placement on a 7 point scale.” Use this as your row variable. Under Issues, choose the variable “Iraq war increased or decreased the threat of terrorism.” Use this as your column variable. In the details for the table, set Weight to none, and for N of Cases to Display, make sure the unweighted box is checked. For Percentaging, choose row percents. Now click on “run the table.”

(b) To analyze the data, make a 3×3 table by combining the rows for extremely liberal and liberal; slightly liberal, middle of the road, and slightly conservative; and conservative and extremely conservative. Carry out a formal test to determine if there is a relationship between these two variables, and then describe the relationship in a brief report using percents from the table or an appropriate graph.

(c) Select two other variables of interest to you and analyze the relationship between them. If there is a more recent survey than 2008, you should use it.





Inference for Regression

Chapter 24

When a scatterplot shows a linear relationship between a quantitative explanatory variable x and a quantitative response variable y , we can use the least-squares line fitted to the data to predict y for a given value of x . When the data are a sample from a larger population, we need statistical inference to answer questions like these about the population:

- Is there really a linear relationship between x and y in the population, or might the pattern we see in the scatterplot plausibly arise just by chance?
- What is the slope (rate of change) that relates y to x in the population, including a margin of error for our estimate of the slope?
- If we use the least-squares line to predict y for a given value of x , how accurate is our prediction (again, with a margin of error)?

This chapter shows you how to answer these questions. Here is an example we will explore.

EXAMPLE 24.1 Crying and IQ

STATE: Infants who cry easily may be more easily stimulated than others. This may be a sign of higher IQ. Child development researchers explored the relationship between the crying of infants four to ten days old and their later IQ test scores. A snap of a rubber band on the sole of the foot caused the infants to cry. The researchers recorded the crying and measured its intensity by the number of peaks in the most active 20 seconds. They later measured the children's IQ at age three years using the Stanford-Binet IQ test. Table 24.1 contains data on 38 infants.¹ Do children with higher crying counts tend to have higher IQ?

IN THIS CHAPTER WE COVER...

- Conditions for regression inference
- Estimating the parameters
- Using technology
- Testing the hypothesis of no linear relationship
- Testing lack of correlation
- Confidence intervals for the regression slope
- Inference about prediction
- Checking the conditions for inference



CRYING

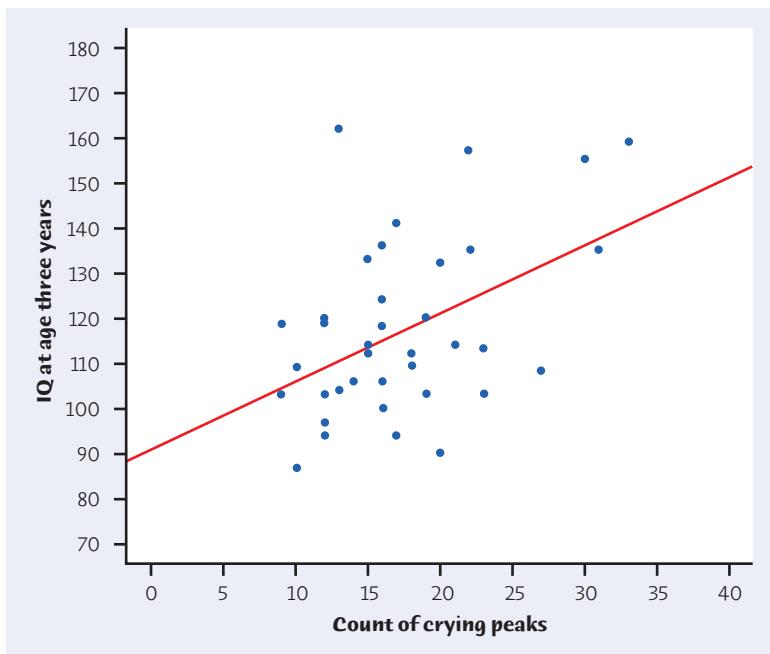
TABLE 24.1 Infants' crying (number of peaks) and IQ scores

CRYING	IQ	CRYING	IQ	CRYING	IQ	CRYING	IQ
10	87	20	90	17	94	12	94
12	97	16	100	19	103	12	103
9	103	23	103	13	104	14	106
16	106	27	108	18	109	10	109
18	109	15	112	18	112	23	113
15	114	21	114	16	118	9	119
12	119	12	120	19	120	16	124
20	132	15	133	22	135	31	135
16	136	17	141	30	155	22	157
33	159	13	162				

PLAN: Make a scatterplot. If the relationship appears linear, use correlation and regression to describe it. Finally, ask whether there is a *statistically significant* linear relationship between crying and IQ.

scatterplot

SOLVE (first steps): Chapters 4 and 5 introduced the data analysis that must come before inference. The first steps we take are a review of this data analysis. Figure 24.1 is a **scatterplot** of the crying data. Plot the explanatory variable (count of crying peaks) horizontally and the response variable (IQ) vertically. Look for the form, direction, and strength of the relationship as well as for outliers or other deviations. There is a moderately strong positive linear relationship, with no extreme outliers or potentially influential observations.

**FIGURE 24.1**

Scatterplot of the IQ score of infants at age three years against the intensity of their crying soon after birth, with the least-squares regression line, for Example 24.1.

Because the scatterplot shows a roughly linear (straight-line) pattern, the **correlation** describes the direction and strength of the relationship. The correlation between crying and IQ is $r = 0.455$. We are interested in predicting the response from information about the explanatory variable. So we find the **least-squares regression line** for predicting IQ from crying. The equation of the regression line is

$$\begin{aligned}\hat{y} &= a + bx \\ &= 91.27 + 1.493x\end{aligned}$$

CONCLUDE (first steps): Children who cry more vigorously do tend to have higher IQs. Because $r^2 = 0.207$, only about 21% of the variation in IQ scores is explained by crying intensity. Prediction of IQ will not be very accurate. It is nonetheless impressive that behavior soon after birth can even partly predict IQ three years later. Is this observed relationship statistically significant? We must now develop tools for inference in the regression setting. ■

correlation**least-squares line**

CONDITIONS FOR REGRESSION INFERENCE

We can fit a regression line to *any* data relating two quantitative variables, though the results are useful only if the scatterplot shows a linear pattern. Statistical inference requires more detailed conditions. Because the conclusions of inference always concern some *population*, the conditions describe the population and how the data are produced from it. The slope b and intercept a of the least-squares line are *statistics*. That is, we calculated them from the sample data. These statistics would take somewhat different values if we repeated the study with different infants. To do inference, think of a and b as estimates of unknown *parameters* that describe the population of all infants.

CONDITIONS FOR REGRESSION INFERENCE

We have n observations on an explanatory variable x and a response variable y . Our goal is to study or predict the behavior of y for given values of x .

- For any fixed value of x , the response y varies according to a **Normal distribution**. Repeated responses y are **independent** of each other.
- The mean response μ_y has a **straight-line relationship** with x given by a **population regression line**

$$\mu_y = \alpha + \beta x$$

- The slope β and intercept α are unknown parameters.
- The **standard deviation** of y (call it σ) is the same for all values of x . The value of σ is unknown.

There are thus three population parameters that we must estimate from the data: α , β , and σ .

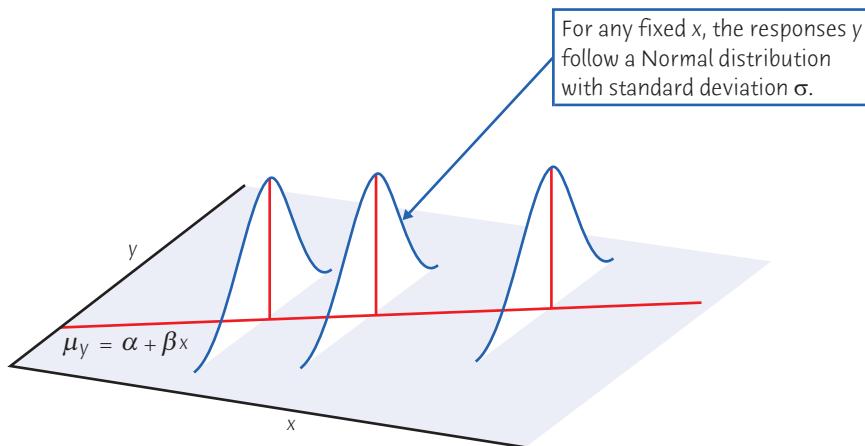
These conditions say that in the population there is an “on the average” straight-line relationship between y and x . The population regression line $\mu_y = \alpha + \beta x$ says that the *mean* response μ_y moves along a straight line as the explanatory variable x changes. We can’t observe the population regression line. The values of y that we

do observe vary about their means according to a Normal distribution. If we hold x fixed and take many observations on y , the Normal pattern will eventually appear in a stemplot or histogram. In practice, we observe y for many different values of x , so that we see an overall linear pattern formed by points scattered about the population line. The standard deviation σ determines whether the points fall close to the population regression line (small σ) or are widely scattered (large σ).

Figure 24.2 shows the conditions for regression inference in picture form. The line in the figure is the population regression line. The mean of the response y moves along this line as the explanatory variable x takes different values. The Normal curves show how y will vary when x is held fixed at different values. All the curves have the same σ , so the variability of y is the same for all values of x . You should check the conditions for inference when you do inference about regression. We will see later how to do that.

FIGURE 24.2

The nature of regression data when the conditions for inference are met. The line is the population regression line, which shows how the mean response μ_y changes as the explanatory variable x changes. For any fixed value of x , the observed response y varies according to a Normal distribution having mean μ_y and standard deviation σ .



ESTIMATING THE PARAMETERS

The first step in inference is to estimate the unknown parameters α , β , and σ .

ESTIMATING THE POPULATION REGRESSION LINE

When the conditions for regression are met and we calculate the least-squares line $\hat{y} = a + bx$, the slope b of the least-squares line is an unbiased estimator of the population slope β , and the intercept a of the least-squares line is an unbiased estimator of the population intercept α .

EXAMPLE 24.2 Crying and IQ: slope and intercept

The data in Figure 24.1 satisfy the condition of scatter about an invisible population regression line reasonably well. The least-squares line is $\hat{y} = 91.27 + 1.493x$. The slope is particularly important. A *slope is a rate of change*. The population slope β says how much higher average IQ is for children with one more peak in their crying

measurement. Because $b = 1.493$ estimates the unknown β , we estimate that, on the average, IQ is about 1.5 points higher for each added crying peak.

We need the intercept $a = 91.27$ to draw the line, but it has no statistical meaning in this example. No child had fewer than 9 crying peaks, so we have no data near $x = 0$. We suspect that all normal children would cry when snapped with a rubber band, so that we will never observe $x = 0$. ■

The remaining parameter is the standard deviation σ , which describes the variability of the response y about the population regression line. The least-squares line estimates the population regression line. So the **residuals** estimate how much y varies about the population line. Recall that the residuals are the vertical deviations of the data points from the least-squares line:

$$\begin{aligned}\text{residual} &= \text{observed } y - \text{predicted } y \\ &= y - \hat{y}\end{aligned}$$

There are n residuals, one for each data point. Because σ is the standard deviation of responses about the population regression line, we estimate it by a sample standard deviation of the residuals. We call this sample standard deviation the *regression standard error* to emphasize that it is estimated from data. The residuals from a least-squares line always have mean zero. That simplifies their standard error.

REGRESSION STANDARD ERROR

The **regression standard error** is

$$\begin{aligned}s &= \sqrt{\frac{1}{n-2} \sum \text{residual}^2} \\ &= \sqrt{\frac{1}{n-2} \sum (y - \hat{y})^2}\end{aligned}$$

Use s to estimate the standard deviation σ of responses about the mean given by the population regression line.

Because we use the regression standard error so often, we just call it s . The quantity $\sum(y - \hat{y})^2$ is the sum of the squared deviations of the data points from the line. We average the squared deviations by dividing by $n - 2$, the number of data points less 2. It turns out that if we know $n - 2$ of the n residuals, the other two are determined. That is, $n - 2$ are the **degrees of freedom** of s . We first met the idea of degrees of freedom in the case of the ordinary sample standard deviation of n observations, which has $n - 1$ degrees of freedom. Now we observe two variables rather than one, and the proper degrees of freedom are $n - 2$ rather than $n - 1$.

Calculating s is unpleasant. You must find the predicted response for each x in your data set, then the residuals, and then s . In practice you will use software that does this arithmetic instantly. Nonetheless, here is an example to help you understand the standard error s .

residuals

degrees of freedom



CRYINGRES

EXAMPLE 24.3 Crying and IQ: residuals and standard error

Table 24.1 shows that the first infant studied had 10 crying peaks and a later IQ of 87. The predicted IQ for $x = 10$ is

$$\begin{aligned}\hat{y} &= 91.27 + 1.493x \\ &= 91.27 + 1.493(10) = 106.2\end{aligned}$$

The residual for this observation is

$$\begin{aligned}\text{residual} &= y - \hat{y} \\ &= 87 - 106.2 = -19.2\end{aligned}$$

That is, the observed IQ for this infant lies 19.2 points below the least-squares line on the scatterplot.

Repeat this calculation 37 more times, once for each subject. The 38 residuals are

-19.20	-31.13	-22.65	-15.18	-12.18	-15.15	-16.63	-6.18
-1.70	-22.60	-6.68	-6.17	-9.15	-23.58	-9.14	2.80
-9.14	-1.66	-6.14	-12.60	0.34	-8.62	2.85	14.30
9.82	10.82	0.37	8.85	10.87	19.34	10.89	-2.55
20.85	24.35	18.94	32.89	18.47	51.32		

Check the calculations by verifying that the sum of the residuals is zero. It is 0.04, not quite zero, because of roundoff error. Another reason to use software in regression is that roundoff errors in hand calculation can accumulate to make the results inaccurate.

The variance about the line is

$$\begin{aligned}s^2 &= \frac{1}{n-2} \sum \text{residual}^2 \\ &= \frac{1}{38-2} [(-19.20)^2 + (-31.13)^2 + \dots + (51.32)^2] \\ &= \frac{1}{36}(11,023.3) = 306.20\end{aligned}$$

Finally, the regression standard error is

$$s = \sqrt{306.20} = 17.50 \blacksquare$$

We will study several kinds of inference in the regression setting. The regression standard error s is the key measure of the variability of the responses in regression. It is part of the standard error of all the statistics we will use for inference.



Glowimages/Age Fotostock

APPLY YOUR KNOWLEDGE

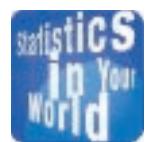
- 24.1 Wine and cancer in women.** Some studies have suggested that a nightly glass of wine may not only take the edge off a day but also improve health. Is wine good for your health? A study of nearly 1.3 million middle-aged British women examined wine consumption and the relative risk of breast cancer. The relative risk is the proportion of those in the study who drank a given amount of wine and who developed breast cancer divided by the proportion of nondrinkers in the study who developed

breast cancer. For example, if 10% of the women in the study who drank 10 grams of wine developed breast cancer and 9% of nondrinkers in the study developed breast cancer, the relative risk of breast cancer for women drinking 10 grams of wine per day would be $10\%/9\% = 1.11$. A relative risk greater than 1 indicates a greater proportion of drinkers in the study developed breast cancer than nondrinkers. Wine intake is the mean wine intake, in grams per day, of all women in the study who drank some wine but less than or equal to 2 drinks per week; who drank between 3 and 6 drinks per week; who drank between 7 and 14 drinks per week; and who drank 15 or more drinks per week. Here are the data (for drinkers only):²



Wine intake (grams per day) (x)	2.5	8.5	15.5	26.5
Relative risk (y)	1.00	1.08	1.15	1.22

- Examine the data. Make a scatterplot with wine intake as the explanatory variable and find the correlation. There is a strong linear relationship.
- Explain in words what the slope β of the population regression line would tell us if we knew it (note that these data represent averages over large numbers of women and are an example of an ecological correlation (see page 142), and one must be careful not to interpret the data as applying to individuals). Based on the data, what are the estimates of β and the intercept α of the population regression line?
- Calculate by hand the residuals for the four data points. Check that their sum is 0 (up to roundoff error). Use the residuals to estimate the standard deviation σ that measures variation in the responses (relative risk) about the means given by the population regression line. You have now estimated all three parameters.



The jinx!

Athletes are often jinxed. We read of “the rookie of the year jinx,” the “cover of Sports Illustrated jinx,” and many others. That is, athletes who are recognized for an outstanding performance often fail to do as well in the future. No, nature isn’t retaliating against them. It’s just random variation about their long-term mean performance. They were recognized because they randomly varied above their typical performance, and in the future they return to the mean or randomly vary down from it. If they randomly vary down, they can hope for a “comeback” award the next year.

USING TECHNOLOGY

Basic “two-variable statistics” calculators will find the slope b and intercept a of the least-squares line from keyed-in data. Inference about regression requires in addition the regression standard error s . At this point, software or a graphing calculator that includes procedures for regression inference becomes almost essential for practical work.

Figure 24.3 shows regression output for the data of Table 24.1 from a graphing calculator, two statistical programs, and a spreadsheet program. When we entered the data into the programs, we called the explanatory variable “Crycount.” The software outputs use that label. The graphing calculator just uses “x” and “y” to

Texas Instruments Graphing Calculator

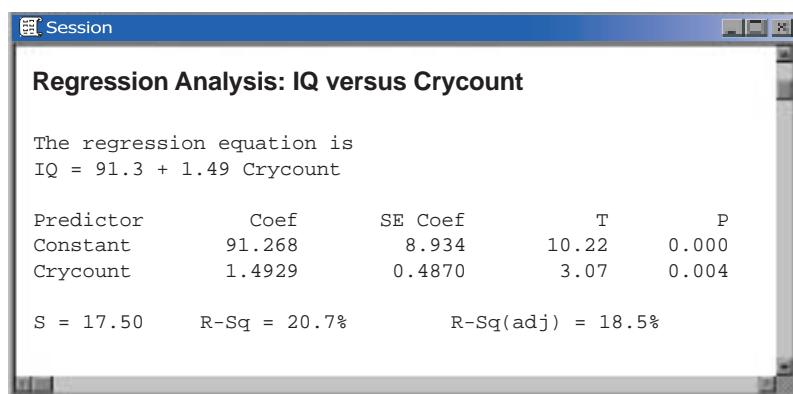
LinRegTTest	LinRegTTest
$y=a+bx$	$y=a+bx$
$B \neq 0$ and $P \neq 0$	$B \neq 0$ and $P \neq 0$
$t=3.0655$	$t=1.4928$
$P=.0041$	$s=17.4387$
$df=36.0000$	$r=.2070$
$t_a=91.2683$	$r=.4550$

FIGURE 24.3

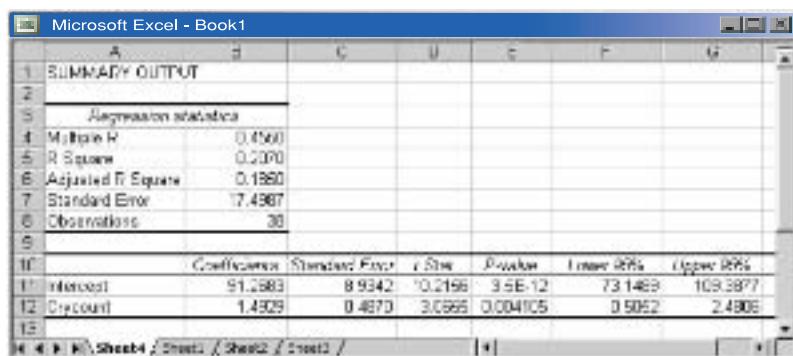
Regression of IQ on crying peaks: output from a graphing calculator, two statistical programs, and a spreadsheet program.

FIGURE 24.3 (Continued)

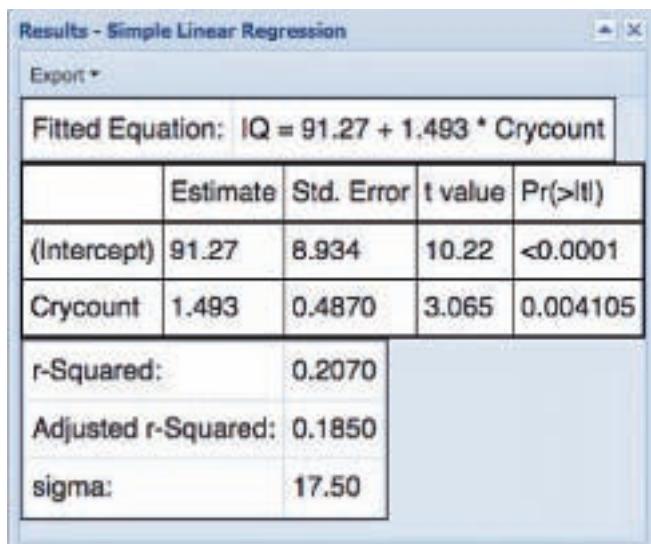
Minitab



Excel



CrunchIt!



label the explanatory and response variables. You can locate the basic information in all of the outputs. The regression slope is $b = 1.4929$ and the regression intercept is $a = 91.268$. The equation of the least-squares line is therefore (after rounding) just as given in Example 24.1. The regression standard error is $s = 17.4987$ and the squared correlation is $r^2 = 0.207$. Both of these results reflect the rather wide scatter of the points in Figure 24.1 about the least-squares line.

Each output contains other information, some of which we will need shortly and some of which we don't need. In fact, we left out some output to save space. Once you know what to look for, you can find what you want in almost any output and ignore what doesn't interest you.

APPLY YOUR KNOWLEDGE

24.2 Introspection and gray matter. The ability to introspect about self-performance is key to human subjective experience. Accurate introspection requires discriminating correct decisions from incorrect ones, a capacity that varies substantially across individuals. Are individual differences in introspective ability reflected in the anatomy of brain regions responsible for this function? The data below are a measure of introspective ability (labeled Aroc and based on the performance of subjects on a task) and a measure of gray-matter volume (Brodmann area) in the anterior prefrontal cortex of the brain of 29 subjects.³  GRAYMATTER

Volume	0.55	0.58	0.59	0.59	0.59	0.61	0.62	0.63	0.63	0.63
Aroc	59	62	43	63	83	61	55	57	57	67
Volume	0.63	0.64	0.65	0.65	0.65	0.65	0.65	0.66	0.66	0.67
Aroc	72	62	58	62	65	70	75	60	63	71
Volume	0.67	0.67	0.68	0.69	0.70	0.70	0.71	0.72	0.75	
Aroc	71	80	68	72	66	73	61	80	75	

We want to predict Aroc from volume. Figure 24.4 shows Minitab regression output for these data.

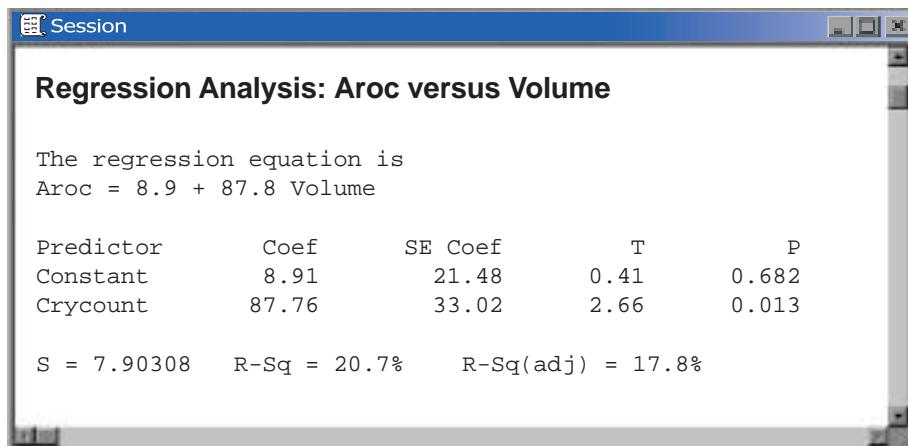


FIGURE 24.4

Minitab output for the introspective-ability data, for Exercise 24.2.

- (a) Make a scatterplot suitable for predicting Aroc from volume. What is the squared correlation r^2 ?
- (b) For regression inference, we must estimate the three parameters α , β , and σ . From the output, what are the estimates of these parameters?
- (c) What is the equation of the least-squares regression line of Aroc on volume? Add this line to your plot. We will continue the analysis of these data in later exercises.

24.3 Great Arctic rivers. One effect of global warming is to increase the flow of water into the Arctic Ocean from rivers. Such an increase may have major effects on the world's climate. Six rivers (Yenisey, Lena, Ob, Pechora, Kolyma, and Severnaya Dvina) drain two-thirds of the Arctic in Europe and Asia. Several of these are among the largest rivers on earth. Table 24.2 presents the total discharge (amount of water flowing from these rivers) each year from 1936 to 2008.⁴ Discharge is measured in cubic kilometers of water. Use software to analyze these data. 

- (a) Make a scatterplot of river discharge against time. Is there a clear increasing trend? Calculate r^2 and briefly interpret its value. There is considerable year-to-year variation, so we wonder if the trend is statistically significant.
- (b) As a first step, find the least-squares line and draw it on your plot. Then find the regression standard error s , which measures scatter about this line. We will continue the analysis in later exercises.

TABLE 24.2 Arctic river discharge (cubic kilometers), 1936 to 2008

YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE	YEAR	DISCHARGE
1936	1721	1955	1656	1974	2000	1993	1845
1937	1713	1956	1721	1975	1928	1994	1902
1938	1860	1957	1762	1976	1653	1995	1842
1939	1739	1958	1936	1977	1698	1996	1849
1940	1615	1959	1906	1978	2008	1997	2007
1941	1838	1960	1736	1979	1970	1998	1903
1942	1762	1961	1970	1980	1758	1999	1970
1943	1709	1962	1849	1981	1774	2000	1905
1944	1921	1963	1774	1982	1728	2001	1890
1945	1581	1964	1606	1983	1920	2002	2085
1946	1834	1965	1735	1984	1823	2003	1780
1947	1890	1966	1883	1985	1822	2004	1900
1948	1898	1967	1642	1986	1860	2005	1930
1949	1958	1968	1713	1987	1732	2006	1910
1950	1830	1969	1742	1988	1906	2007	2270
1951	1864	1970	1751	1989	1932	2008	2078
1952	1829	1971	1879	1990	1861		
1953	1652	1972	1736	1991	1801		
1954	1589	1973	1861	1992	1793		

TESTING THE HYPOTHESIS OF NO LINEAR RELATIONSHIP

Example 24.1 asked, “Do children with higher crying counts tend to have higher IQ?” Data analysis supports this conjecture. But is the positive association statistically significant? That is, is it too strong to often occur just by chance? To answer this question, test hypotheses about the slope β of the population regression line:

$$H_0: \beta = 0$$

$$H_a: \beta > 0$$

A regression line with slope 0 is horizontal. That is, the mean of y does not change at all when x changes. So H_0 says that there is *no linear relationship* between x and y in the population. Put another way, H_0 says that *linear regression of y on x is of no value for predicting y* .

The test statistic is just the standardized version of the least-squares slope b , using the hypothesized value $\beta = 0$ for the mean of b . It is another t statistic. Here are the details.

SIGNIFICANCE TEST FOR REGRESSION SLOPE

To test the hypothesis $H_0: \beta = 0$, compute the t statistic

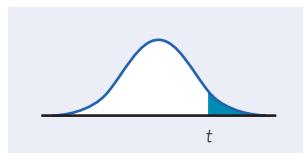
$$t = \frac{b}{\text{SE}_b}$$

In this formula, the standard error of the least-squares slope b is

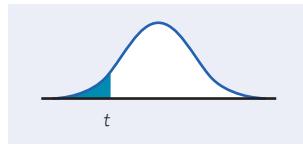
$$\text{SE}_b = \frac{s}{\sqrt{\sum(x - \bar{x})^2}}$$

The sum runs over all observations on the explanatory variable x . In terms of a random variable T having the $t(n - 2)$ distribution, the P -value for a test of H_0 against

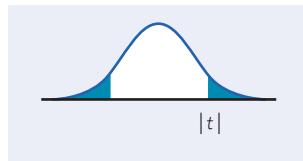
$$H_a: \beta > 0 \text{ is } P(T \geq t)$$



$$H_a: \beta < 0 \text{ is } P(T \leq t)$$



$$H_a: \beta \neq 0 \text{ is } 2P(T \geq |t|)$$



As advertised, the standard error of b is a multiple of the regression standard error s . The degrees of freedom $n - 2$ are the degrees of freedom of s . Although we give the formula for this standard error, you should not try to calculate it by hand. Regression software gives the standard error SE_b along with b itself.

EXAMPLE 24.4 Crying and IQ: is the relationship significant?

The hypothesis $H_0: \beta = 0$ says that crying has no straight-line relationship with IQ. We conjecture that there is a positive relationship, so we use the one-sided alternative $H_a: \beta > 0$.

Figure 24.1 shows that there is a positive relationship, and from Figure 24.3 we see $b = 1.4929$ and $\text{SE}_b = 0.4870$. Thus,

$$t = \frac{b}{\text{SE}_b} = \frac{1.4929}{0.4870} = 3.07$$

so it is not surprising that all the outputs in Figure 24.3 give $t = 3.07$ with two-sided P -value 0.004. The P -value for the one-sided test is half of this, $P = 0.002$. There is very strong evidence that IQ increases as the intensity of crying increases. ■

APPLY YOUR KNOWLEDGE

24.4 Wine and cancer in women. Exercise 24.1 gives data on daily wine consumption and the relative risk of breast cancer in women. Software tells us that the least-squares slope is $b = 0.009012$ with standard error $\text{SE}_b = 0.001112$.

- (a) What is the t statistic for testing $H_0: \beta = 0$?
- (b) How many degrees of freedom does t have? Use Table C to approximate the P -value of t against the one-sided alternative $H_a: \beta > 0$. What do you conclude?

24.5 Great Arctic rivers: testing. The most important question we ask of the data in Table 24.2 is this: is the increasing trend visible in your plot (Exercise 24.3) statistically significant? If so, changes in the Arctic may already be affecting the earth's climate. Use software to answer this question. Give a test statistic, its P -value, and the conclusion you draw from the test. 

24.6 Does fast driving waste fuel? Exercise 4.8 (page 104) gives data on the fuel consumption of a small car at various speeds from 10 to 150 kilometers per hour. Is there significant evidence of straight-line dependence between speed and fuel use? Make a scatterplot and use it to explain the result of your test. 

TESTING LACK OF CORRELATION

The least-squares slope b is closely related to the correlation r between the explanatory and response variables x and y . In the same way, the slope β of the population regression line is closely related to the correlation between x and y in the population. In particular, the slope is 0 exactly when the correlation is 0.

Testing the null hypothesis $H_0: \beta = 0$ is therefore exactly the same as testing that there is no correlation between x and y in the population from which we drew our data. You can use the test for zero slope to test the hypothesis of zero correlation between any two quantitative variables. That's a useful trick.

Because correlation also makes sense when there is no explanatory-response distinction, it is handy to be able to test correlation without doing regression. Table E in the back of the book gives critical values of the sample correlation r under the null hypothesis that the correlation is 0 in the population. Use this table when both variables have at least approximately Normal distributions or when the sample size is large.

EXAMPLE 24.5 Testing lack of correlation

Figure 24.5 displays two scatterplots that we will use to illustrate testing lack of correlation and also to illustrate once again the need for formal statistical tests. On the left are data from an experiment on the healing of cuts in the limbs of newts. The data are the healing rates (micrometers per hour) for the two front limbs of 18 newts. The right-hand scatterplot shows the first- and second-round scores for the 95 golfers in the 2010 Masters Tournament. (There are fewer than 95 points because of duplicate scores.)

We will test the hypotheses

$$H_0: \text{population correlation} = 0$$

$$H_a: \text{population correlation} \neq 0$$



AP Photo/Elise Amendola

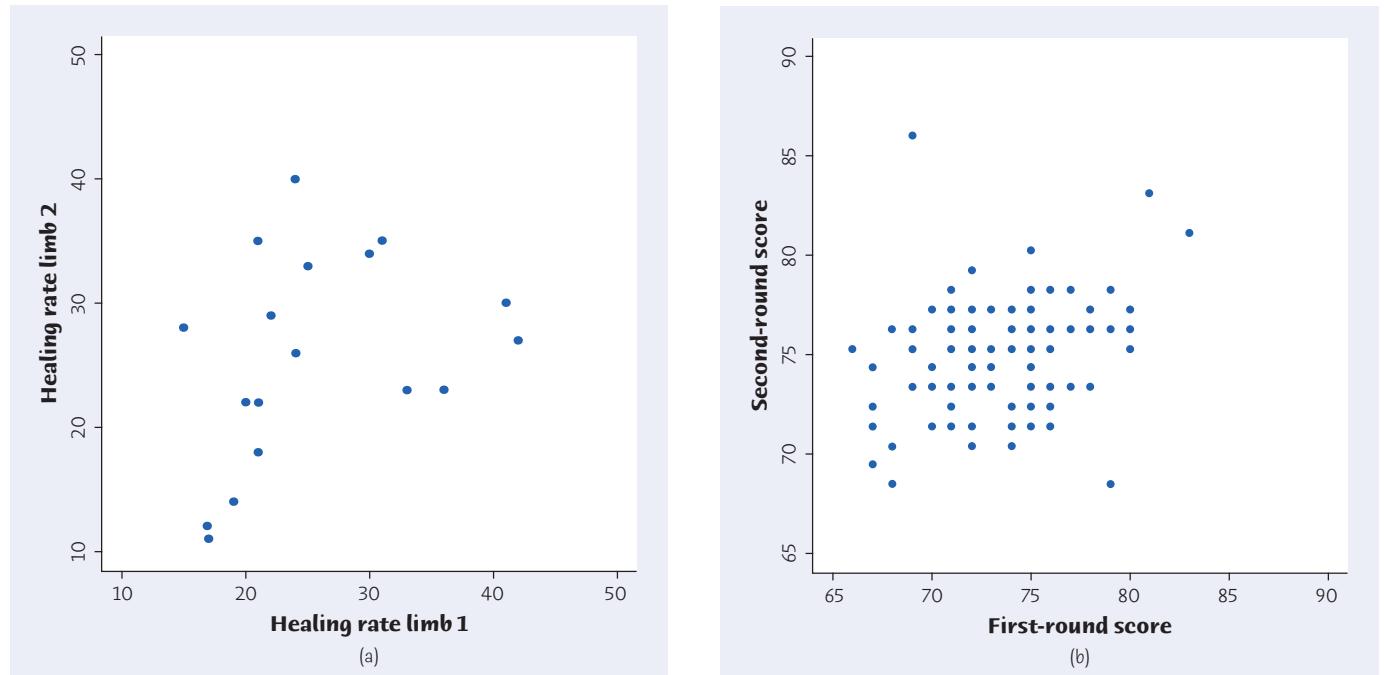


FIGURE 24.5

Two scatterplots for inference about the population correlation, for Example 24.5. (a) Healing rates for the two front limbs of 18 newts. (b) Scores on the first two rounds of the 2010 Masters Tournament.

for both sets of data. (The Masters scores are all whole numbers, but with $n = 95$ the robustness of t procedures allows their use.) Software gives

newts	$r=0.3581$	$t=1.5342$	$P=0.1445$
masters	$r=0.3465$	$t=3.56$	$P=0.001$

The two-sided P -values for the t statistic for testing slope 0 are also the two-sided P -values for testing correlation 0.

Without software, compare the correlation $r = 0.3581$ for newts with the critical values in the $n = 18$ row of Table E. It falls between the table entries for one-tail probabilities 0.05 and 0.10, so the two-sided P -value lies between 0.10 and 0.20. For the Masters data, use the $n = 80$ row of Table E. (There is no table entry for sample size $n = 95$, so we use the next smaller sample size.) The two-sided P -value lies between 0.001 and 0.002. ■

The evidence for nonzero correlation is strong for the Masters scores ($t = 3.56$, $P = 0.001$) but not for newts ($t = 1.5$, $P = 0.14$). Yet the correlation for the newts is slightly larger than that for the Masters, and the scatterplots suggest similar linear relationships for both. What happened? The larger sample size for the Masters data is largely responsible. The same r will have a smaller P -value for $n = 95$ than for $n = 18$. Our eyeball impression, even aided by calculating r , can't assess significance. We need the P -value from a formal test to guide us.

APPLY YOUR KNOWLEDGE

24.7 Wine and cancer in women: testing correlation. Exercise 24.1 gives data showing that the risk of breast cancer increases linearly with daily wine consumption. There are only 4 observations, so we worry that the apparent relationship may be just chance. Is the correlation significantly greater than 0? Answer this question in two ways.

- (a) Return to your t statistic from Exercise 24.4. What is the one-sided P -value for this t ? Apply your result to test the correlation.
- (b) Find the correlation r and use Table E to approximate the P -value of the one-sided test.  WINECANCER

24.8 Does social rejection hurt? Exercise 4.45 (page 121) gives data from a study of whether social rejection causes activity in areas of the brain that are known to be activated by physical pain. The explanatory variable is a subject's score on a test of "social distress" after being excluded from an activity. The response variable is activity in an area of the brain that responds to physical pain. Your scatterplot (Exercise 4.45) shows a positive linear relationship. The research report gives the correlation r and the P -value for a test that r is greater than 0. What are r and the P -value? (You can use Table E or you can get more accurate P -values for the correlation from regression software.) What do you conclude about the relationship?  REJECTION

CONFIDENCE INTERVALS FOR THE REGRESSION SLOPE

The slope β of the population regression line is usually the most important parameter in a regression problem. The slope is the rate of change of the mean response as the explanatory variable increases. We often want to estimate β . The slope b of

the least-squares line is an unbiased estimator of β . A confidence interval is more useful because it shows how accurate the estimate b is likely to be. The confidence interval for β has the familiar form

$$\text{estimate} \pm t^*SE_{\text{estimate}}$$

Because b is our estimate, the confidence interval is $b \pm t^*SE_b$. Here are the details.

CONFIDENCE INTERVAL FOR REGRESSION SLOPE

A level C confidence interval for the slope β of the population regression line is

$$b \pm t^*SE_b$$

Here t^* is the critical value for the $t(n - 2)$ density curve with area C between $-t^*$ and t^* . The formula for SE_b appears in the box on page 597.

EXAMPLE 24.6 Crying and IQ: estimating the slope

All the outputs in Figure 24.3 (pages 593 and 594) give the slope $b = 1.4929$ (or $b = 1.493$), and three also give the standard error $SE_b = 0.4870$. The outputs giving both use a similar arrangement, a table in which each regression coefficient is followed by its standard error. Excel also gives the lower and upper endpoints of the 95% confidence interval for the population slope β , 0.505 and 2.481.

Once we know b and SE_b , it is easy to find the confidence interval. There are 38 data points, so the degrees of freedom are $n - 2 = 36$. Because Table C does not have a row for $df = 36$, we must use either software or the next smaller degrees of freedom in the table, $df = 30$. To use software, enter 36 degrees of freedom. For 95% confidence, enter the cumulative proportion 0.975 that corresponds to upper-tail area 0.025. Minitab gives

Student's t distribution with 36 DF	
$P(X \leq x)$	x
0.975	2.02809

The 95% confidence interval for the population slope β is

$$\begin{aligned} b \pm t^*SE_b &= 1.4929 \pm (2.02809)(0.4870) \\ &= 1.4929 \pm 0.9877 \\ &= 0.505 \text{ to } 2.481 \end{aligned}$$

This agrees with Excel's result. We are 95% confident that mean IQ increases by between about 0.5 and 2.5 points for each additional peak in crying. ■

You can find a confidence interval for the intercept α of the population regression line in the same way, using a and SE_a from the "Constant" line of the Minitab output or the "Intercept" line in Excel and CrunchIt! We rarely need to estimate α .



APPLY YOUR KNOWLEDGE

24.9 Wine and cancer in women: estimating slope. Exercise 24.1 gives data on wine consumption and the risk of breast cancer. Software tells us that the least-squares slope is $b = 0.009012$ with standard error $SE_b = 0.001112$. Because there are only 4 observations, the observed slope b may not be an accurate estimate of the population slope β . Give a 90% confidence interval for β .

24.10 Introspection and gray matter: estimating slope. Exercise 24.2 gives data on introspective ability and gray-matter volume of the brains of subjects. We want a 95% confidence interval for the slope of the population regression line. Starting from the information in the Minitab output in Figure 24.4 (page 595), find this interval. Say in words what the slope of the population regression line tells us about the relationship between Aroc and gray-matter volume.

24.11 Great Arctic rivers: estimating slope. Use the data in Table 24.2 to give a 90% confidence interval for the slope of the population regression of Arctic river discharge on year. Does this interval convince you that discharge is actually increasing over time? Explain your answer.  ARCTIC

INFERENCE ABOUT PREDICTION

One of the most common reasons to fit a line to data is to predict the response to a particular value of the explanatory variable. This is another setting for regression inference: we want, not simply a prediction, but a prediction with a margin of error that describes how accurate the prediction is likely to be.



BEERS



Jame Shaffer/The Image Works

EXAMPLE 24.7 Beer and blood alcohol

STATE: The EESEE story “Blood Alcohol Content” describes a study in which 16 student volunteers at the Ohio State University drank a randomly assigned number of cans of beer. Thirty minutes later, a police officer measured their blood alcohol content (BAC) in grams of alcohol per deciliter of blood. Here are the data:⁵

Student	1	2	3	4	5	6	7	8
Beers	5	2	9	8	3	7	3	5
BAC	0.10	0.03	0.19	0.12	0.04	0.095	0.07	0.06
Student	9	10	11	12	13	14	15	16
Beers	3	5	4	6	5	7	1	4
BAC	0.02	0.05	0.07	0.10	0.085	0.09	0.01	0.05

The students were equally divided between men and women and differed in weight and usual drinking habits. Because of this variation, many students don’t believe that number of drinks predicts blood alcohol well. Steve thinks he can drive legally

30 minutes after he finishes drinking 5 beers. The legal limit for driving is BAC 0.08 in all states. We want to predict Steve's blood alcohol content, using no information except that he drinks 5 beers.

PLAN: Regress BAC on number of beers. Use the regression line to predict Steve's BAC. Give a margin of error that allows us to have 95% confidence in our prediction.

SOLVE: The scatterplot in Figure 24.6 and the regression output in Figure 24.7 show that student opinion is wrong: number of beers predicts blood alcohol content quite well. In fact, $r^2 = 0.80$, so that number of beers explains 80% of the observed

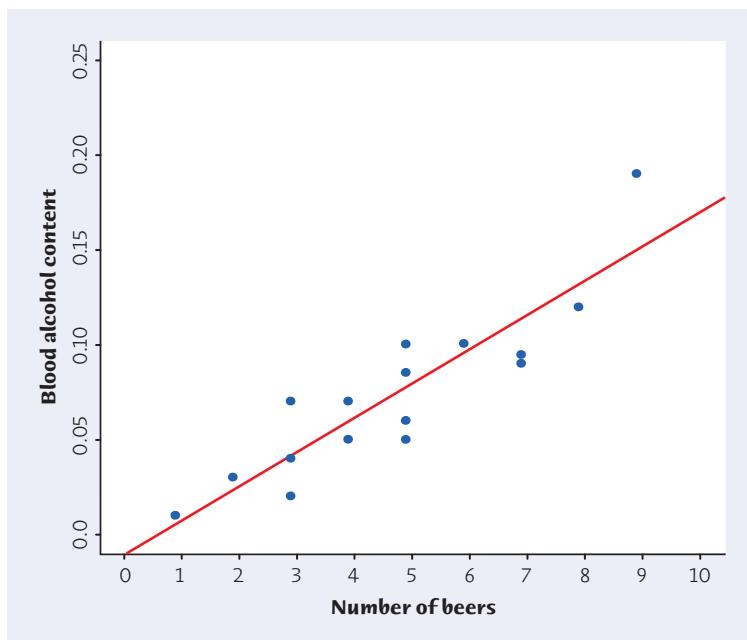


FIGURE 24.6

Scatterplot of students' blood alcohol content against the number of cans of beer consumed, with the least-squares regression line, for Example 24.7.

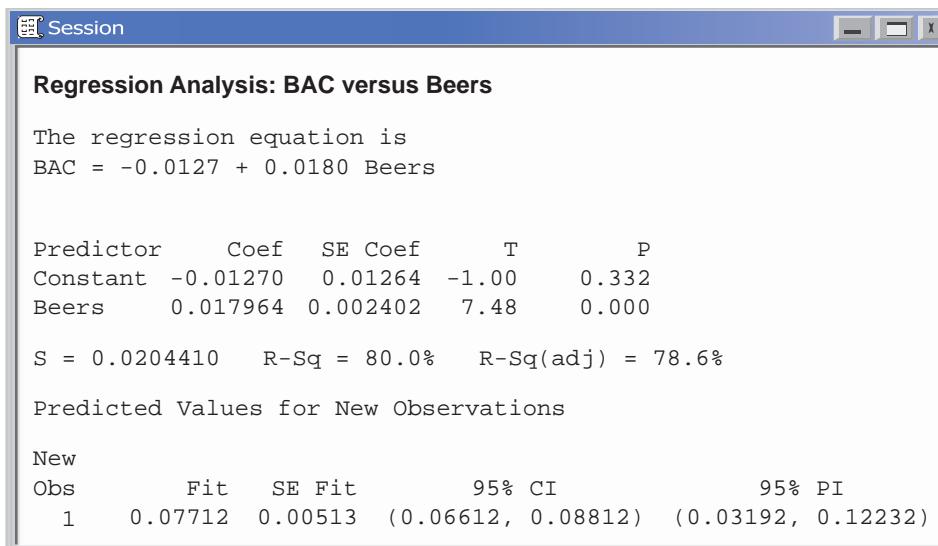


FIGURE 24.7

Minitab regression output for the blood alcohol content data, for Example 24.7.

variation in BAC. To predict Steve's BAC after 5 beers, use the equation of the regression line:

$$\begin{aligned}\hat{y} &= -0.0127 + 0.0180x \\ &= -0.0127 + 0.0180(5) = 0.077\end{aligned}$$

That's dangerously close to the legal limit 0.08. What about 95% confidence? The "Predicted Values" part of the output in Figure 24.7 shows two 95% intervals. Which should we use?

To decide which interval to use, you must answer this question: do you want to predict the *mean* BAC for all students who drink 5 beers, or do you want to predict the BAC of one *individual* student who drinks 5 beers? Both of these predictions may be interesting, but they are two different problems. The actual prediction is the same, $\hat{y} = 0.077$. But the margin of error is different for the two kinds of prediction. Individual students who drink 5 beers don't all have the same BAC. So we need a larger margin of error to pin down Steve's result than to estimate the mean BAC for all students who have 5 beers.

Write the given value of the explanatory variable x as x^* . In Example 24.7, $x^* = 5$. The distinction between predicting a single outcome and predicting the mean of all outcomes when $x = x^*$ determines what margin of error is correct. To emphasize the distinction, we use different terms for the two intervals.

- To estimate the *mean* response, we use a **confidence interval**. It is an ordinary confidence interval for the mean response when x has the value x^* , which is $\mu_y = \alpha + \beta x^*$. This is a parameter, a fixed number whose value we don't know.
- To estimate an *individual* response y , we use a **prediction interval**. A prediction interval estimates a single random response y rather than a parameter like μ_y . The response y is not a fixed number. If we took more observations with $x = x^*$, we would get different responses.

EXAMPLE 24.8 Beer and blood alcohol: conclusion

Steve is one individual, so we must use the prediction interval. The output in Figure 24.7 helpfully labels the confidence interval as "95% CI" and the prediction interval as "95% PI." We are 95% confident that Steve's BAC after 5 beers will lie between 0.032 and 0.122. The upper part of that range will get him arrested if he drives. The 95% confidence interval for the mean BAC of all students who drink 5 beers is much narrower, 0.066 to 0.088.

The meaning of a prediction interval is very much like the meaning of a confidence interval. A 95% prediction interval, like a 95% confidence interval, is right 95% of the time in repeated use. "Repeated use" now means that we take an observation on y for each of the n values of x in the original data and then take one more observation y with $x = x^*$. Form the prediction interval from the n observations, then see if it covers the one more y . It will in 95% of all repetitions.



Is regression garbage?

No—but garbage can be the setting for regression. The U.S. Census Bureau once asked if weighing a neighborhood's garbage would help count its people. So 63 households had their garbage sorted and weighed. It turned out that pounds of plastic in the trash gave the best garbage prediction of the number of people in a neighborhood. The margin of error for a 95% prediction interval in a neighborhood of about 100 households, based on five weeks' worth of garbage, was about ± 2.5 people. Alas, that is not accurate enough to help the U.S. Census Bureau.

prediction interval



The interpretation of prediction intervals is a minor point. The main point is that it is harder to predict one response than to predict a mean response. Both intervals have the usual form

$$\hat{y} \pm t^*SE$$

but the prediction interval is wider than the confidence interval because individuals are more variable than averages. As a further illustration, consider a professional athlete such as the basketball player Kobe Bryant. You can find his career statistics online. During his career as a starter, his season scoring *average* ranges from about 20 to 35 points per game, but individual game performances range from just a few points to a career high of 81 points. Individual game performances of professional athletes are more variable than season average performances. So the margin of error for predictions of individual game performances will be larger than for predictions of season averages. You will rarely need to know the details, because software automates the calculation, but here they are.



May the longer name win!

Regression is far from perfect, but it beats most other ways of predicting. A writer in the early 1960s noted a simple method for predicting presidential elections: just choose the candidate with the longer name. In the 22 elections from 1876 to 1960, this method failed only once. Let's hope that the writer didn't bet the family silver on this idea. The 12 elections from 1964 to 2008 presented 10 tests of the "long name wins" method (the 1980 candidates and the 2000 candidates had names of the same length). The longer name lost 6 of the 10.

CONFIDENCE AND PREDICTION INTERVALS FOR REGRESSION RESPONSE

A level C confidence interval for the mean response μ_y when x takes the value x^* is

$$\hat{y} \pm t^*SE_{\hat{\mu}}$$

The standard error $SE_{\hat{\mu}}$ is

$$SE_{\hat{\mu}} = s \sqrt{\frac{1}{n} + \frac{(x^* - \bar{x})^2}{\sum(x - \bar{x})^2}}$$

A level C prediction interval for a single observation y when x takes the value x^* is

$$\hat{y} \pm t^*SE_{\hat{y}}$$

The standard error for prediction $SE_{\hat{y}}$ is

$$SE_{\hat{y}} = s \sqrt{1 + \frac{1}{n} + \frac{(x^* - \bar{x})^2}{\sum(x - \bar{x})^2}}$$

In both intervals, t^* is the critical value for the $t(n - 2)$ density curve with area C between $-t^*$ and t^* .

There are two standard errors: $SE_{\hat{\mu}}$ for estimating the mean response μ_y and $SE_{\hat{y}}$ for predicting an individual response y . The only difference between the two standard errors is the extra 1 under the square root sign in the standard error for prediction. The extra 1 makes the prediction interval wider. Both standard errors are multiples of the regression standard error s . The degrees of freedom are again $n - 2$, the degrees of freedom of s .



APPLY YOUR KNOWLEDGE

24.12 Wine and cancer in women: prediction. Exercise 24.1 gives data on wine consumption and the risk of breast cancer. For a new group of women who drink an average of 10 grams of wine per day, predict their relative risk of breast cancer.

- Figure 24.8 is part of the output from Minitab for prediction when $x^* = 10.0$. Which interval in the output is the proper 95% interval for predicting the relative risk?
- Minitab gives only one of the two standard errors used in prediction. It is $SE_{\hat{\mu}}$, the standard error for estimating the mean response. Use this fact along with the output to give a 90% confidence interval for the mean relative risk of breast cancer in all women who drink an average of 10 grams of wine per day.

FIGURE 24.8

Partial Minitab output for regressing relative risk of breast cancer on mean daily intake of wine, for Exercise 24.12.

Predictor	Coef	SE Coef	T	P
Constant	0.99309	0.01777	55.88	0.000
Intake	0.009012	0.001112	8.10	0.015

Predicted Values for New Observations					
New	Obs	Fit	SE Fit	95% CI	95% PI
	1	1.08321	0.01057	(1.03775, 1.12868)	(0.98643, 1.18000)

24.13 Introspection and gray matter: prediction. Analysis of the data in Exercise 24.2 shows that the relationship between gray-matter volume and introspective ability, as measured by Aroc, is roughly linear. We might want to predict the mean introspective ability of a person with a gray-matter volume of 0.60. Here is the Minitab output for prediction when $x^* = 0.60$:

Predicted Values for New Observations					
New	Obs	Fit	SE Fit	95% CI	95% PI
	1	61.56	2.18	(57.08, 66.05)	(44.74, 78.39)

- (a) Use the regression line from Figure 24.4 (page 595) to verify that “Fit” is the predicted value for $x^* = 0.60$. (Start with the results in the “Coef” column of Figure 24.4 to reduce roundoff error.)
- (b) What is the 95% interval we want?

CHECKING THE CONDITIONS FOR INFERENCE

You can fit a least-squares line to any set of explanatory-response data when both variables are quantitative. If the scatterplot doesn’t show a roughly linear pattern, the fitted line may be almost useless. But it is still the line that fits the data best in the least-squares sense. To use regression inference, however, the data must satisfy additional conditions. *Before you can trust the results of inference, you must check the conditions for inference one by one.* There are ways to deal with violations of any of the conditions. If you see a clear violation, get expert advice.



Although the conditions for regression inference are a bit elaborate, it is not hard to check for major violations. The conditions involve the population regression line and the deviations of responses from this line. We can’t observe the population line, but the least-squares line estimates it and the residuals estimate the deviations from the population line. *You can check all the conditions for regression inference by looking at graphs of the residuals.* This is what we recommend in practice (and a failure to do so can sometimes lead to unjustified conclusions), and most regression software will calculate and save the residuals for you. Start by making a stemplot or histogram of the residuals and also a **residual plot**, a plot of the residuals against the explanatory variable x , with a horizontal line at the “residual = 0” position. The “residual = 0” line represents the position of the least-squares line in the scatterplot of y against x . Let’s look at each condition in turn.

residual plot

- **The relationship is linear in the population.** Look for curved patterns or other departures from a straight-line overall pattern in the residual plot. You can also use the original scatterplot, but the residual plot magnifies any effects.
- **The response varies Normally about the population regression line.** Because different y -values usually come from different x -values, the responses themselves need not be Normal. It is the deviations from the population line—estimated by the residuals—that must be Normal. Check for clear skewness or other major departures from Normality in your stemplot or histogram of the residuals.
- **Observations are independent.** In particular, repeated observations on the same individual are not allowed. You should not use ordinary regression to make inferences about the growth of a single child over time, for example. Signs of dependence in the residual plot are a bit subtle, so we usually rely on common sense.
- **The standard deviation of the responses is the same for all values of x .** Look at the scatter of the residuals above and below the “residual = 0” line in the residual plot. The scatter should be roughly the same from one end to the other. You will sometimes find that, as the response y gets larger, so does the scatter of the residuals. Rather than remaining fixed, the standard deviation σ about the

line changes with x as the mean response changes with x . There is no fixed σ for s to estimate. You cannot trust the results of inference when this happens.

You will always see some irregularity when you look for Normality and fixed standard deviation in the residuals, especially when you have few observations. Don't overreact to minor violations of the conditions. Like other t procedures, inference for regression is (with one exception) not very sensitive to lack of Normality, especially when we have many observations. Do beware of influential observations, which can greatly affect the results of inference.

The exception is the prediction interval for a single response y . This interval relies on Normality of individual observations, not just on the approximate Normality of statistics like the slope a and intercept b of the least-squares line. The statistics a and b become more Normal as we take more observations. This contributes to the robustness of regression inference, but it isn't enough for the prediction interval. We will not study methods that carefully check Normality of the residuals, so you should regard prediction intervals as rough approximations.



ANGLERFISH



Dave Harasti

EXAMPLE 24.9 Climate change chases fish north

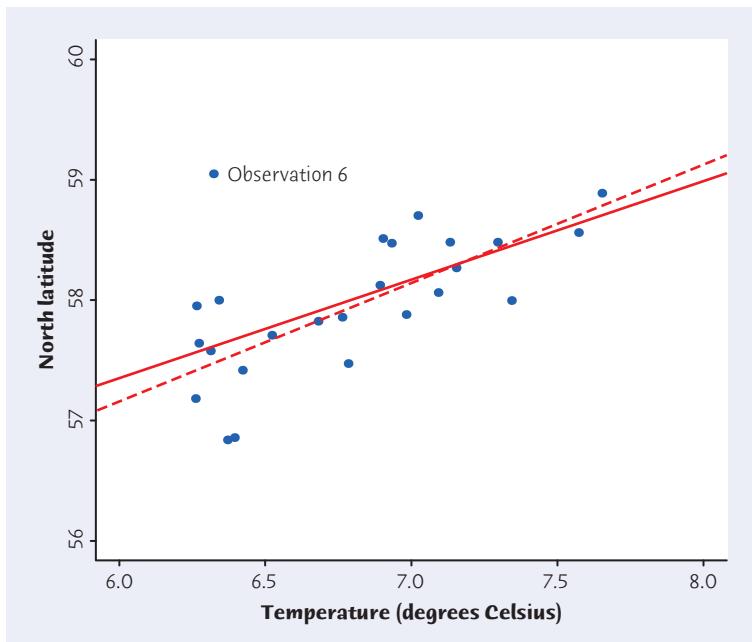
STATE: As the climate grows warmer, we expect many animal species to move toward the poles in an attempt to maintain their preferred temperature range. Do data on fish in the North Sea confirm this expectation? Table 24.3 gives data for 25 years on mean winter temperatures at the bottom of the North Sea (degrees Celsius) and the center of the distribution of anglerfish in degrees of north latitude.⁶

PLAN: Regress latitude on temperature. Look for a positive linear relationship and assess its significance. Be sure to check the conditions for regression inference.

SOLVE: The scatterplot in Figure 24.9 shows a clear positive linear relationship. The solid line in the plot is the least-squares regression line of the center of the fish distribution (north latitude) on winter ocean temperature. Software shows that the slope is $b = 0.818$. That is, each degree of ocean warming moves the fish about 0.8 degree of latitude farther north. The t statistic for testing $H_0: \beta = 0$ is $t = 3.6287$ with one-sided

TABLE 24.3 Winter temperature (°C) and anglerfish latitude, 1977 to 2001

YEAR	TEMP.	LATITUDE	YEAR	TEMP.	LATITUDE	YEAR	TEMP.	LATITUDE
1977	6.26	57.20	1986	6.52	57.72	1994	7.02	58.71
1978	6.26	57.96	1987	6.68	57.83	1995	7.09	58.07
1979	6.27	57.65	1988	6.76	57.87	1996	7.13	58.49
1980	6.31	57.59	1989	6.78	57.48	1997	7.15	58.28
1981	6.34	58.01	1990	6.89	58.13	1998	7.29	58.49
1982	6.32	59.06	1991	6.90	58.52	1999	7.34	58.01
1983	6.37	56.85	1992	6.93	58.48	2000	7.57	58.57
1984	6.39	56.87	1993	6.98	57.89	2001	7.65	58.90
1985	6.42	57.43						

**FIGURE 24.9**

Scatterplot of the latitude of the center of the distribution of anglerfish in the North Sea against mean winter temperature at the bottom of the sea, for Example 24.9. The two regression lines are for the data with (solid) and without (dashed) Observation 6.

P-value $P = 0.0007$. There is very strong evidence that the population slope is positive, $\beta > 0$.

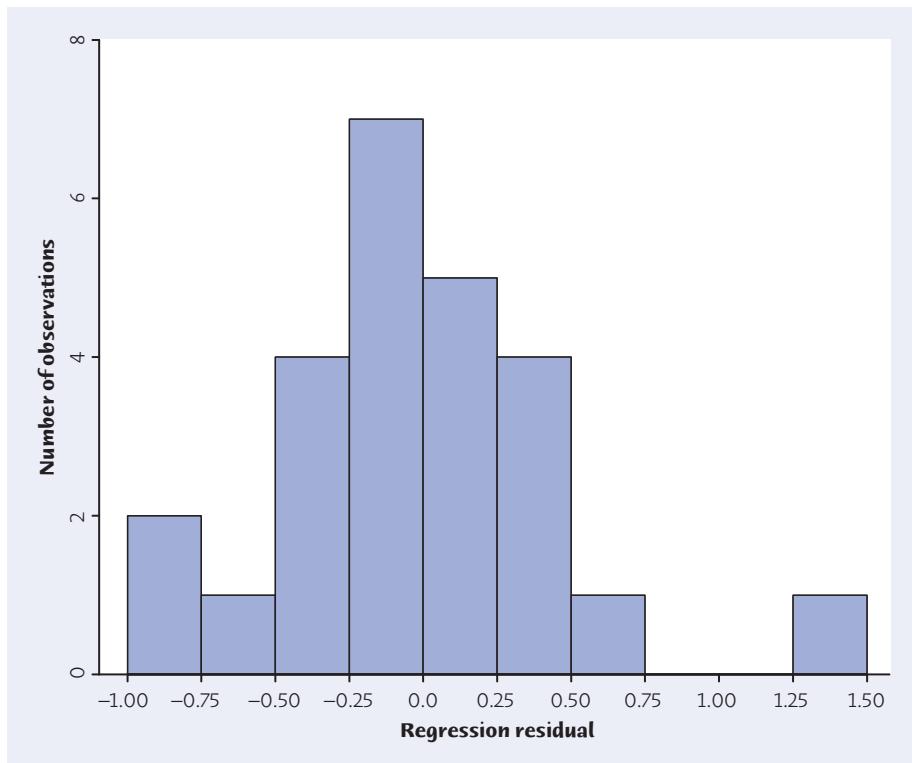
CONCLUDE: The data give highly significant evidence that anglerfish have moved north as the ocean has grown warmer. Before relying on this conclusion, we must check the conditions for inference. ■

The software that did the regression calculations also finds the 25 residuals. In the same order as the observations in Example 24.9, they are

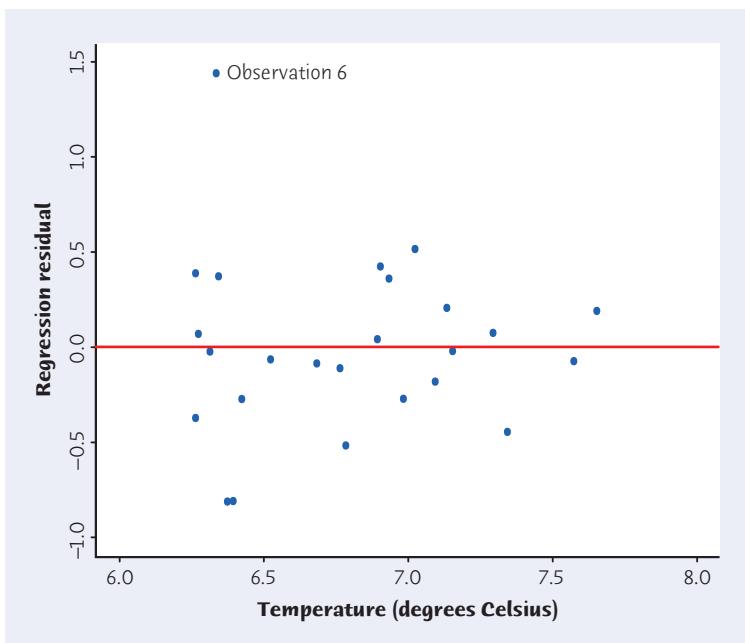
-0.3731	0.3869	0.0687	-0.0240	0.3714	1.4378	-0.8131
-0.8095	-0.2740	-0.0658	-0.0867	-0.1121	-0.5185	0.0415
0.4234	0.3588	-0.2721	0.5152	-0.1821	0.2052	-0.0211
0.0743	-0.4466	-0.0747	0.1899			

Begin by making two graphs of the residuals. Figure 24.10 is a histogram of the residuals. Figure 24.11 is the residual plot, a plot of the residuals against the explanatory variable, sea-bottom temperature. The “residual = 0” line marks the position of the regression line. Patterns in residual plots are often easier to see if you use a wider vertical scale than your software’s default plot and we suggest you do so if possible. Both graphs show that Observation 6 is a high outlier. Let’s check the conditions for regression inference.

- **Linear relationship.** The scatterplot in Figure 24.9 and the residual plot in Figure 24.11 both show a linear relationship except for the outlier.
- **Normal residuals.** The histogram in Figure 24.10 is roughly symmetric and single-peaked. There are no important departures from Normality except for the outlier.

**FIGURE 24.10**

Histogram of the residuals from the regression of latitude on temperature in Example 24.9.

**FIGURE 24.11**

Residual plot for the regression of latitude on temperature in Example 24.9.

- **Independent observations.** The observations were taken a year apart, so we are willing to regard them as close to independent. The residual plot shows no obvious pattern of dependence, such as runs of points all above or all below the line.
- **Constant standard deviation.** Again excepting the outlier, the residual plot shows no unusual variation in the scatter of the residuals above and below the line as x varies.

The outlier is the only serious violation of the conditions for inference. How influential is the outlier? The dashed line in Figure 24.9 is the regression line without Observation 6. Because there are several other observations with similar values of temperature, dropping Observation 6 does not move the regression line very much. Even though the outlier is not very influential for the regression line, it influences regression inference because of its effect on the regression standard error. The standard error is $s = 0.4734$ with Observation 6 and $s = 0.3622$ without it. When we omit the outlier, the t statistic changes from $t = 3.6287$ to $t = 5.5599$, and the one-sided P -value changes from $P = 0.0007$ to $P < 0.00001$. Fortunately, the outlier does not affect the conclusion we drew from the data. Dropping Observation 6 makes the test for the population slope more significant and increases the percent of variation in fish location explained by ocean temperature.



One more caution about inference in this example: as usual in an observational study, the possibility of lurking variables makes us hesitant to conclude that rising temperature is causing anglerfish to move north. Ocean temperature was steadily rising during these years. The effect on fish latitude of any lurking variable that increased over time—perhaps increased commercial fishing—is confounded with the effect of temperature.

APPLY YOUR KNOWLEDGE

24.14 Crying and IQ: residuals. The residuals for the study of crying and IQ appear in Example 24.3.  CRYINGRES

- Make a stemplot to display the distribution of the residuals (round to the nearest whole number). Are there strong outliers or other signs of departures from Normality?
- Make a residual plot, residuals against crying peaks. Try a vertical scale of -60 to 60 to show patterns more clearly. Draw the “residual = 0” line. Does the residual plot show clear deviations from a linear pattern or clearly unequal spread about the line?
- Using the information given in Example 24.1, explain why the 38 observations are independent.

24.15 Introspection and gray matter: residuals. Figure 24.4 (page 595) gives part of the Minitab output for the data on introspective ability and gray-matter volume in Exercise 24.2. Figure 24.12 comes from another part of the output. It gives x , y , the predicted response \hat{y} , the residual $y - \hat{y}$, and related quantities for each of the 29 observations. Most statistical software provides similar output. Examine the conditions for regression inference one by one. This example

FIGURE 24.12

Residuals from Minitab for Exercise 24.15. The table gives the predicted value (“Fit”) and the residual for each observation.

Obs	Volume	Aroc	Fit	SE Fit	Residual	St Resid
1	0.550	58.00	57.18	3.58	0.82	0.12
2	0.580	62.00	59.81	2.71	2.19	0.30
3	0.590	43.00	60.69	2.44	-17.69	-2.35R
4	0.590	63.00	60.69	2.44	2.31	0.31
5	0.590	83.00	60.69	2.44	22.31	2.97R
6	0.610	61.00	62.44	1.95	-1.44	-0.19
7	0.620	55.00	63.32	1.75	-8.32	-1.08
8	0.630	57.00	64.20	1.60	-7.20	-0.93
9	0.630	57.00	64.20	1.60	-7.20	-0.93
10	0.630	67.00	64.20	1.60	2.80	0.36
11	0.630	72.00	64.20	1.60	7.80	1.01
12	0.640	62.00	65.08	1.50	-3.08	-0.40
13	0.650	58.00	65.95	1.47	-7.95	-1.02
14	0.650	62.00	65.95	1.47	-3.95	-0.51
15	0.650	65.00	65.95	1.47	-0.95	-0.12
16	0.650	70.00	65.95	1.47	4.05	0.52
17	0.650	75.00	65.95	1.47	9.05	1.17
18	0.660	60.00	66.83	1.51	-6.83	-0.88
19	0.660	63.00	66.83	1.51	-3.83	-0.49
20	0.670	71.00	67.71	1.62	3.29	0.43
21	0.670	71.00	67.71	1.62	3.29	0.43
22	0.670	80.00	67.71	1.62	12.29	1.59
23	0.680	68.00	68.59	1.79	-0.59	-0.08
24	0.690	72.00	69.46	2.00	2.54	0.33
25	0.700	66.00	70.34	2.23	-4.34	-0.57
26	0.700	73.00	70.34	2.23	2.66	0.35
27	0.710	61.00	71.22	2.49	-10.22	-1.36
28	0.720	80.00	72.10	2.77	7.90	1.07
29	0.750	75.00	74.73	3.65	0.27	0.04

R denotes an observation with a large standardized residual.

illustrates mild violations of the conditions that did not prevent the researchers from doing inference.  CRAYMATTERES

- (a) **Linear relationship.** Your scatterplot and r^2 from Exercise 24.2 show that the relationship is roughly linear. Plot the residuals against volume. Are any deviations from a straight line apparent?
- (b) **Normal variation about the line.** Make a stemplot of the residuals (round to the nearest integer, use split stems, and don't forget that -0 and 0 are separate stems). With only 29 observations, a small amount of skew is not disturbing. Minitab suggests that Observations 3 and 5 may be outliers. Does your plot confirm or refute this suggestion?
- (c) **Independent observations.** The data come from 29 different subjects who were each measured separately.
- (d) **Spread about the line stays the same.** Is there any evidence that the spread may be larger at one end?

CHAPTER 24 SUMMARY

CHAPTER SPECIFICS

- Least-squares regression fits a straight line to data in order to predict a response variable y from an explanatory variable x . Inference about regression requires more conditions.
- The conditions for regression inference say that there is a population regression line $\mu_y = \alpha + \beta x$ that describes how the mean response varies as x changes. The observed response y for any x has a Normal distribution with mean given by the population regression line and with the same standard deviation σ for any value of x . Observations on y are independent.
- The parameters to be estimated are the intercept α and the slope β of the population regression line and also the standard deviation σ . The slope a and intercept b of the least-squares line estimate α and β . Use the regression standard error s to estimate σ .
- The regression standard error s has $n - 2$ degrees of freedom. All t procedures in regression inference have $n - 2$ degrees of freedom.
- To test the hypothesis that the slope is zero in the population, use the t statistic $t = b/\text{SE}_b$. This null hypothesis says that straight-line dependence on x has no value for predicting y . In practice, use software to find the slope b of the least-squares line, its standard error SE_b , and the t statistic.
- The t test for regression slope is also a test for the hypothesis that the population correlation between x and y is zero. To do this test without software, use the sample correlation r and Table E.
- Confidence intervals for the slope of the population regression line have the form $b \pm t^* \text{SE}_b$.
- Confidence intervals for the mean response when x has value x^* have the form $\hat{y} \pm t^* \text{SE}_{\hat{y}}$. Prediction intervals for an individual future response y have a similar form with a larger standard error, $\hat{y} \pm t^* \text{SE}_{\hat{y}}$. Software often gives these intervals.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Preliminaries

1. Make a scatterplot to show the relationship between an explanatory and a response variable.
2. Use a calculator or software to find the correlation and the equation of the least-squares regression line.
3. Recognize which type of inference you need in a particular regression setting.

B. Inference Using Software Output

1. Explain in any specific regression setting the meaning of the slope β of the population regression line.
2. Understand software output for regression. Find in the output the slope and intercept of the least-squares line, their standard errors, and the regression standard error.
3. Use that information to carry out tests of $H_0: \beta = 0$ and calculate confidence intervals for β .
4. Explain the distinction between a confidence interval for the mean response and a prediction interval for an individual response.
5. If software gives output for prediction, use that output to give either confidence or prediction intervals.

C. Checking the Conditions for Regression Inference

1. Make a stemplot or histogram of the residuals and look for strong departures from Normality.
2. Make a residual plot and look for departures from a linear pattern or unequal spread about the “residual = 0” line.
3. Ask whether the study design suggests that observations are independent.

LINK IT

In Chapters 4 and 5 we studied scatterplots, correlation, and the least-squares regression line as methods for exploring data. In Chapter 5 we mentioned that we must exercise caution in how we interpret any relationships we observe through such exploratory analyses. We also mentioned that such interpretations rest on the assumption that the relationship is valid in some broader sense. And we promised to explore this more carefully later in this book. In this chapter we did so by considering inference for regression. Inference for regression allows us to determine whether the relationship we observe in a scatterplot is valid for some larger population. It also allows us to attach margins of error to our estimates of the slope and intercept, as well as to predictions based on the least-squares regression line.

In Chapters 4 and 5 we discussed several cautions about correlation and the least-squares regression line. These same cautions apply to inference for regression. We discussed a systematic approach using the residuals and what we know of the study design to determine if the assumptions behind inference for regression are reasonable.

This chapter considers inference for relationships between two quantitative variables. The previous chapter considered inference for relationships between two categorical variables. In the next chapter, we consider inference for relationships between a quantitative and a categorical variable—in particular, for deciding whether the mean of a response is the same for more than two categories.

CHECK YOUR SKILLS

Florida reappraises real estate every year, so the county appraiser’s Web site lists the current “fair market value” of each piece of property. Property usually sells for somewhat more than the appraised market value. Here are the appraised market values and actual selling prices (in thousands of dollars) of condominium units sold in a beachfront building in a 93-month period between 2003 and 2010:⁷



Franz Marc Freil/CORBIS

Selling price	Appraised value	Month									
825	626	0	1325	1032	19	850	715	47	1510	1241	70
590	492	1	700	556	21	1100	997	54	1375	813	70
1075	930	3	1322	879	26	1164	953	59	560	496	73
890	790	9	1900	1016	26	1425	922	64	1050	774	79
845	648	13	1600	1040	28	1865	1190	64	605	470	86
1100	942	14	980	442	34	1450	610	64	675	545	88
715	345	15	940	771	37	875	806	64	693	690	93

Here is part of the Minitab output for regressing selling price on appraised value, along with prediction for a unit with appraised value \$800,000:

Predictor	Coef	SE Coef	T	P
Constant	86.0	156.8	0.55	0.588
Appraisal	1.2699	0.1938	6.55	0.000
<i>S</i> = 235.410 R-Sq = 62.3% R-Sq(adj) = 60.8%				
Predicted Values for New Observations				
New				
Obs	Fit	SE Fit	95% CI	95% PI
1	1101.9	44.7	(1010.0, 1193.9)	(609.4, 1594.5)

Exercises 24.16 to 24.24 are based on this information.

24.16 The equation of the least-squares regression line for predicting selling price from appraised value is

- (a) price = 86.0 + 1.2699 \times appraised value.
- (b) price = 1.2699 + 86.0 \times appraised value.
- (c) price = 156.8 + 0.1938 \times appraised value.

24.17 What is the correlation between selling price and appraised value?

- (a) 0.789
- (b) 0.623
- (c) 0.388

24.18 The slope β of the population regression line describes

- (a) the average selling price in a population of units when a unit's appraised value is 0.
- (b) the average increase in selling price in a population of units when appraised value increases by \$1000.

- (c) the exact increase in the selling price of an individual unit when its appraised value increases by \$1000.

24.19 Is there significant evidence that selling price increases as appraised value increases? To answer this question, test the hypotheses

- (a) $H_0: \beta = 0$ versus $H_a: \beta > 0$.
- (b) $H_0: \beta = 0$ versus $H_a: \beta \neq 0$.
- (c) $H_0: \alpha = 0$ versus $H_a: \alpha > 0$.

24.20 Minitab shows that the P-value for this test is

- (a) 0.588.
- (b) 0.1938.
- (c) less than 0.001.

24.21 The regression standard error for these data is

- (a) 0.1938.
- (b) 156.8.
- (c) 235.41.

24.22 Confidence intervals and tests for these data use the t distribution with degrees of freedom

- (a) 28.
- (b) 27.
- (c) 26.

24.23 A 95% confidence interval for the population slope β is

- (a) 1.2699 ± 0.3306 .
- (b) 1.2699 ± 322.3808 .
- (c) 1.2699 ± 0.3985 .

24.24 Louisa owns a unit in this building appraised at \$800,000. The Minitab output includes prediction for this appraised value. She can be 95% confident that her unit would sell for between

- (a) \$609,400 and \$1,594,500.
- (b) \$1,010,000 and \$1,193,900.
- (c) \$1,057,200 and \$1,146,6900.

CHAPTER 24 EXERCISES

24.25 Genetically engineered cotton

cotton. A strain of genetically engineered cotton, known as Bt cotton, is resistant to certain insects, which results in larger yields of cotton. Farmers in northern China have increased the number of acres planted in Bt cotton. Because Bt cotton is resistant to certain pests, farmers have also reduced their use of insecticide. Scientists in China were interested in the long-term effects of Bt cotton cultivation and decreased insecticide use on insect populations that are not affected by Bt cotton. One such insect is the mirid bug. Scientists measured the number of mirid bugs per 100 plants and the proportion of Bt cotton planted at



Softdreams/Dreamstime.com

38 locations in northern China for the 12-year period from 1997 and 2008. The scientists reported a regression analysis as follows:⁸

number of mirid bugs per 100 plants

$$= 0.54 + 6.81 \times \text{Bt cotton planting proportion}$$

$$r^2 = 0.90 \quad P < 0.0001$$

- (a) What does the slope $b = 6.81$ say about the relation between Bt cotton planting proportion and number of mirid bugs per 100 plants?
- (b) What does $r^2 = 0.90$ add to the information given by the equation of the least-squares line?
- (c) What null and alternative hypotheses do you think the P -value refers to? What does this P -value tell you?

- (d) Does the large value of r^2 and the small P -value indicate that increasing the proportion of acres planted in Bt cotton causes an increase in mirid bugs?

Exercise 7.51 (page 195) gives data from a study of the “gate velocity” of molten metal that experienced foundry workers choose based on the thickness of the aluminum piston being cast. Gate velocity is measured in feet per second, and the piston wall thickness is in inches. A scatterplot (you need not make one) shows a moderately strong positive linear relationship. Figure 24.13 displays part of the Minitab regression output. Exercises 24.26 to 24.28 analyze these data.

24.26 Casting aluminum: is there a relationship? Figure 24.13 leaves out the t statistics and their P -values. Based on the information in the output, test the hypothesis that there is no straight-line relationship between thickness and gate velocity. State hypotheses, give a test statistic and its approximate P -value, and state your conclusion.

24.27 Casting aluminum: intervals. The output in Figure 24.13 includes prediction for piston wall thickness $x^* = 0.5$ inch. Use the output to give 90% intervals for

- (a) the slope of the population regression line of gate velocity on piston thickness.

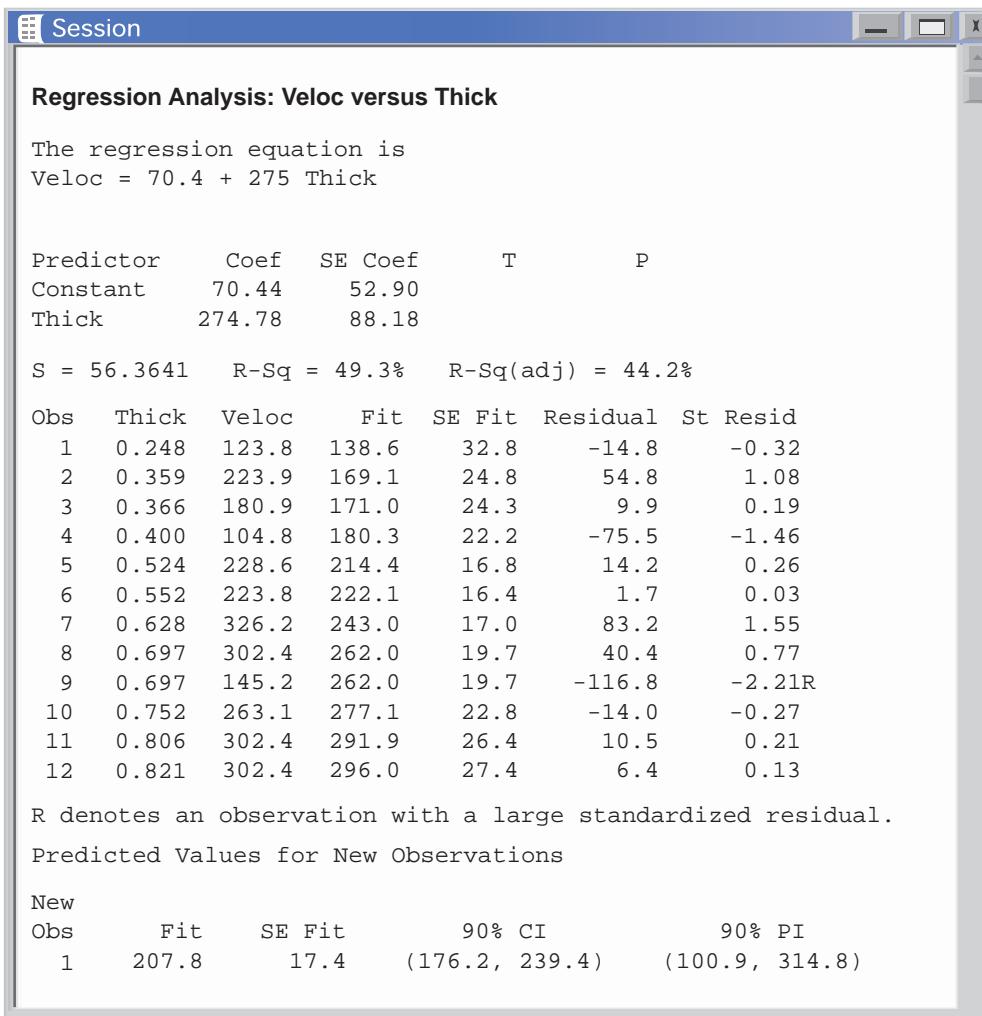


FIGURE 24.13

Minitab output for the regression of gate velocity on piston thickness in casting aluminum parts, for Exercises 24.26 to 24.28.

- (b) the average gate velocity for a type of piston with thickness 0.5 inch.

24.28 Casting aluminum: residuals. The output in Figure 24.13 includes a table of the x and y variables, the fitted values \hat{y} for each x , the residuals, and some related quantities.  ALUMINUMRES

- (a) Plot the residuals against thickness (the explanatory variable). Use vertical scale -200 to 200 so that the pattern is clearer. Add the “residual = 0” line. Does your plot show a systematically nonlinear relationship? Does it show systematic change in the spread about the regression line?
- (b) Make a histogram of the residuals. Minitab identifies the residual for Observation 9 as a suspected outlier. Does your histogram agree?
- (c) Redoing the regression without Observation 9 gives regression standard error $s = 42.4725$ and predicted mean velocity 216 feet per second (90% confidence interval 191.4 to 240.6) for piston walls 0.5 inch thick. Compare these values with those in Figure 24.13. Is Observation 9 influential for inference?

Table 4.1(page 103) gives 33 years’ data on boats registered in Florida and manatees killed by boats. Figure 4.2 (page 103) shows a strong linear relationship. The correlation is $r = 0.951$. Figure 24.14 shows part of the Minitab regression output. Exercises 24.29 to 24.31 analyze the manatee data.

24.29 Manatees: conditions for inference. We know that there is a strong linear relationship. Let’s check the other conditions for inference. Figure 24.14 includes a table of the two variables, the predicted values \hat{y} for each x in the data, the residuals, and related quantities.  MANATEESRES

- (a) Round the residuals to the nearest whole number and make a stemplot. The distribution is single-peaked and symmetric and appears close to Normal.
- (b) Make a residual plot, residuals against boats registered. Use a vertical scale from -25 to 25 to show the pattern more clearly. Add the “residual = 0” line. There is no clearly nonlinear pattern. The spread about the line may be a bit greater for larger values of the explanatory variable, but the effect is not large.
- (c) It is reasonable to regard the number of manatees killed by boats in successive years as independent. The number of boats grew over time. Someone says that pollution also grew over time and may explain more manatee deaths. How would you respond to this idea?

24.30 Manatees: do more boats bring more kills? The output in Figure 24.14 omits the t statistics and their P -values. Based on the information in the output, is there good evidence that the number of manatees killed increases as the

number of boats registered increases? State hypotheses and give a test statistic and its approximate P -value. What do you conclude?

24.31 Manatees: estimation. The output in Figure 24.14 includes prediction of the number of manatees killed when there are 1,050,000 boats registered in Florida. Give 95% intervals for

- (a) the increase in the number of manatees killed for each additional 1000 boats registered.
- (b) the number of manatees that will be killed next year if there are 1,050,000 boats registered next year.

24.32 Fidgeting keeps you slim: inference. Our first example of regression (Example 5.1, page 125) presented data showing that people who increased their nonexercise activity (NEA) when they were deliberately overfed gained less fat than other people. Use software to add formal inference to the data analysis for these data.  FATGAIN

- (a) Based on 16 subjects, the correlation between NEA increase and fat gain was $r = -0.7786$. Is this significant evidence that people with higher NEA increase gain less fat? (Report a t statistic from regression output and give the one-sided P -value.)
- (b) The slope of the least-squares regression line was $b = -0.00344$, so that fat gain decreased by 0.00344 kilogram for each added calorie of NEA. Give a 90% confidence interval for the slope of the population regression line. This rate of change is the most important parameter to be estimated.
- (c) Sam’s NEA increases by 400 calories. His predicted fat gain is 2.13 kilograms. Give a 95% interval for predicting Sam’s fat gain.

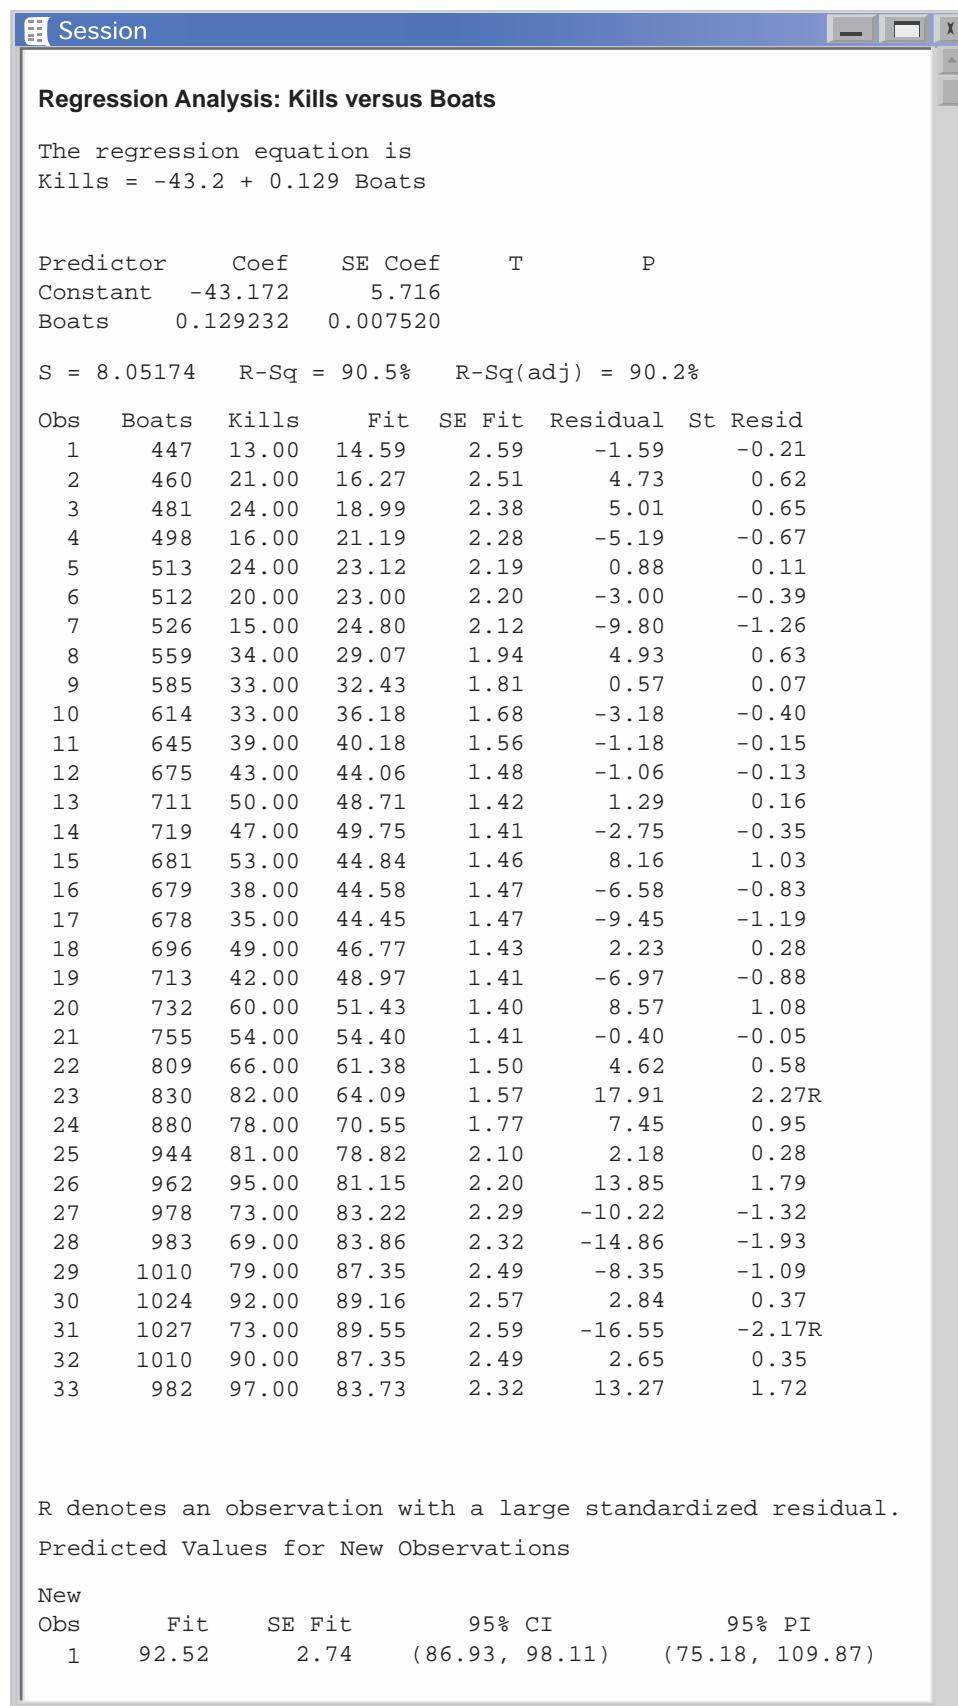
24.33 Predicting tropical storms. Exercise 5.55 (page 155) gives data on William Gray’s predictions of the number of named tropical storms in Atlantic hurricane seasons from 1984 to 2010. Use these data for regression inference as follows.  STORMS2

- (a) Does Professor Gray do better than random guessing? That is, is there a significantly positive correlation between his forecasts and the actual number of storms? (Report a t statistic from regression output and give the one-sided P -value.)
- (b) Give a 95% confidence interval for the mean number of storms in years when Professor Gray forecasts 16 storms.

24.34 Coral growth. Sea surface temperatures across much of the tropics have been increasing since the mid-1970s. At the same time, the growth of coral has been decreasing. Scientists examined data on mean sea surface temperatures

FIGURE 24.14

Minitab output for the regression of number of manatees killed by boats on the number of boats (in thousands) registered in Florida, for Exercises 24.29 to 24.31.



(SST) in degrees Celsius and mean coral growth in millimeters (mm) per year over a several-year period at locations in the Red Sea. Here are the data:⁹



SST	29.68	29.87	30.16	30.22	30.48	30.65	30.90
Growth	2.63	2.58	2.60	2.48	2.26	2.38	2.26

- (a) Do the data indicate that coral growth decreases linearly as SST increases? Is this change statistically significant?
- (b) Use the data to predict with 95% confidence the mean coral growth (mm per year) when SST is 30.0 degrees Celsius.

24.35 Predicting tropical storms: residuals. Make a stemplot of the residuals (round to the nearest tenth) from your regression in Exercise 24.33. Explain why your plot suggests that we should not use these data to get a prediction interval for the number of storms in a single year.



24.36 Coral growth: residuals. Do the data in Exercise 24.34 on mean sea surface temperatures and coral growth in the Red Sea satisfy the conditions for regression inference? To examine this, here are the residuals:



SST	29.68	29.87	30.16	30.22	30.48	30.65	30.90
Residual	-0.067	-0.060	0.128	0.066	0.025	-0.024	-0.068

- (a) **Linear relationship.** A plot of the residuals against the explanatory variable x magnifies the deviations from the least-squares line. Does the plot show any systematic deviation from a roughly linear pattern?
- (b) **Normal variation about the line.** Make a histogram of the residuals. With only 7 observations, no clear shape emerges. Do strong skewness or outliers suggest lack of Normality?
- (c) **Independent observations.** Why are the 7 observations independent?
- (d) **Spread about the line stays the same.** Does your plot in (a) show any systematic change in spread as x changes?

24.37 Our brains don't like losses. Exercise 4.29 (page 117) describes an experiment that showed a linear relationship between how sensitive people are to monetary losses ("behavioral loss aversion") and activity in one part of their brains ("neural loss aversion").



- (a) Make a scatterplot with neural loss aversion as x and behavioral loss aversion as y . One point is a high outlier in both the x and y directions. In Exercise 5.38 (page 152) you found that this outlier is not influential for the least-squares line.

- (b) The research report says that $r = 0.85$ and that the test for regression slope has $P < 0.001$. Verify these results, using all the observations.

(c) The report recognizes the outlier and says, "However, this regression also remained highly significant ($P = 0.004$) when the extreme data point (top right corner) was removed from the analysis." Repeat your analysis omitting the outlier. Show that the outlier influences regression inference by comparing the t statistic for testing slope with and without the outlier. Then verify the report's claim about the P -value of this test.

24.38 Time at the table. Does how long young children remain at the lunch table help predict how much they eat? Here are data on 20 toddlers observed over several months at a nursery school.¹⁰ "Time" is the average number of minutes a child spent at the table when lunch was served. "Calories" is the average number of calories the child consumed during lunch, calculated from careful observation of what the child ate each day.



TIMEATTABLE

Time	21.4	30.8	37.7	33.5	32.8	39.5	22.8	34.1	33.9	43.8
Calories	472	498	465	456	423	437	508	431	479	454
Time	42.4	43.1	29.2	31.3	28.6	32.9	30.6	35.1	33.0	43.7
Calories	450	410	504	437	489	436	480	439	444	408

- (a) Make a scatterplot. Find the correlation and the least-squares regression line. (Be sure to save the regression residuals.) Based on your work, describe the direction, form, and strength of the relationship.

(b) Check the conditions for regression inference. Parts (a) to (d) of Exercise 24.36 provide a handy outline. Use vertical limits -100 to 100 in your plot of the residuals against time to help you see the pattern. What do you conclude?

(c) Is there significant evidence that more time at the table is associated with more calories consumed? Give a 95% confidence interval to estimate how rapidly calories consumed changes as time at the table increases.

24.39 DNA on the ocean floor. We

think of DNA as the stuff that stores the genetic code. It turns out that DNA occurs, mainly outside living cells, on the ocean floor. It is important in nourishing seafloor life. Scientists think that this DNA comes from organic matter that settles to the bottom from the top layers of the ocean. "Phytopigments," which come mainly from algae, are a measure of the amount of organic matter that has settled to the bottom. The data contains



Minoru Toi/Getty Images

data on concentrations of DNA and phytopigments (both in grams per square meter) in 116 ocean locations around the world.¹¹ Look first at DNA alone. Describe the distribution of DNA concentration and give a confidence interval for the mean concentration. Be sure to explain why your confidence interval is trustworthy in the light of the shape of the distribution. The data show surprisingly high DNA concentrations, and this by itself was an important finding.  DNA

24.40 Time at the table: prediction. Rachel attends the nursery school of Exercise 24.38. Over several months, Rachel averages 40 minutes at the lunch table. Give a 95% interval to predict Rachel's average calorie consumption at lunch.  TIME TABLE

Exercises 24.41 to 24.45 ask practical questions involving regression inference without step-by-step instructions. Do complete regression analyses, using the **Plan**, **Solve**, and **Conclude** steps of the four-step process to organize your answers. Follow the model of Example 24.9 (page 608) and the following discussion, and check the conditions as part of the **Solve** step.

24.41 Squirrels and their food supply. The introduction to Exercises 7.24 to 7.26 (pages 185 and 186) gives data on the abundance of the pine cones that red squirrels feed on and the mean number of offspring per female squirrel over 16 years. The strength of the relationship is remarkable because females produce young before the food is available. How significant is the evidence that more cones leads to more offspring? (Use a vertical scale from -2 to 2 in your residual plot to show the pattern more clearly.)  SQUIRRELS

24.42 A big-toe problem. Table 7.4 (page 194) and Exercises 7.47 and 7.49 describe the relationship between two deformities of the feet in young patients. Metatarsus adductus (MA) may help predict the severity of hallux abducto valgus (HAV). The paper that reports this study says, "Linear regression analysis, using the hallux abducto angle as the response variable, demonstrated a significant correlation between the metatarsus adductus and hallux abducto angles."¹² Do a suitable analysis to verify this finding. The study authors note that the scatterplot suggests that the variation in y may change as x changes, so they offer a more elaborate analysis as well.  DEFORMITY

24.43 Beavers and beetles. Exercise 5.53 (page 155) describes a study that found that the number of stumps from trees felled by beavers predicts the abundance of beetle larvae. Is there good evidence that more beetle larvae clusters are present when beavers have left more tree stumps? Estimate how many more clusters accompany each additional stump, with 95% confidence.  BEAVERS

24.44 Sulfur, the ocean, and the sun. Sulfur in the atmosphere affects climate by influencing formation of clouds. The main natural source of sulfur is dimethylsulfide (DMS) produced by small organisms in the upper layers of the oceans. DMS production is in turn influenced by the amount of energy the upper ocean receives from sunlight. Exercise 4.30 (page 117) gives monthly data on solar radiation dose (SRD, in watts per square meter) and surface DMS concentration (in nanomolars) for a region in the Mediterranean. Do the data provide convincing evidence that DMS increases as SRD increases? We also want to estimate the rate of increase, with 90% confidence.  SULFUR

24.45 DNA on the ocean floor. Another conclusion of the study introduced in Exercise 24.39 was that organic matter settling down from the top layers of the ocean is the main source of DNA on the seafloor. An important piece of evidence is the relationship between DNA and phytopigments. Do the data give good reason to think that phytopigment concentration helps explain DNA concentration? (Try vertical limits -1 to 1 to make the pattern of your residual plot clearer.)  DNA

24.46 A lurking variable (optional). Return to the data on selling price versus appraised value for beachfront condominiums that are the basis for the Check Your Skills Exercises 24.16 to 24.24. The data are in order by date of the sale, and the data table includes the number of months from the start of the data period. Here are the residuals from the regression of selling price on appraised value (rounded):  CONDORES

-55.90	-120.78	-192.01	-199.21	-63.88	-182.24
190.90	-71.54	-92.05	119.76	523.78	193.30
332.72	-125.09	-143.97	-252.09	-132.21	168.15
267.81	589.37	-234.53	-151.95	256.58	-155.85
-18.90	-77.84	-103.08	-269.22		

(a) Plot the residuals against the explanatory variable (appraised value). To make the pattern clearer, use vertical limits -600 to 600 . Does the pattern you see agree with the conditions of linear relationship and constant standard deviation needed for regression inference?

(b) Make a stemplot of the residuals. Are there strong deviations from Normality that would prevent regression inference?

(c) Next, plot the residuals against month. Are the positive and negative residuals randomly scattered, as would be the case if the conditions for regression inference are satisfied?

(Comment: Prices for beachfront property were rising rapidly during the first 36 months of this period. Because property is reassessed just once a year, selling prices might pull away from appraised values over time in this period, creating a pattern of many negative residuals followed by several positive

residuals. As this example illustrates, it is often wise to plot residuals against important lurking variables as well as against the explanatory variable.)

24.47 Standardized residuals (optional). Software often calculates **standardized residuals** as well as the actual residuals from regression. Because the standardized residuals have the standard z -score scale, it is easier to judge whether any are extreme. Figure 24.13 (page 616) and the associated data include the standardized residuals for the regression of gate velocity on piston wall thickness.

- (a) Find the mean and standard deviation of the standardized residuals. Why do you expect values close to those you obtain?
- (b) Make a stemplot of the standardized residuals. Are there any striking deviations from Normality?
- (c) The most extreme standardized residual is $z = -2.21$. Minitab flags this as “large.” What is the probability that a standard Normal variable takes a value this extreme (that is, less than -2.21 or greater than 2.21)? Your result suggests that a residual this extreme would be a bit unusual when there are only 12 observations. That’s why we examined Observation 9 in Exercise 24.28.

24.48 Tests for the intercept (optional). Figure 24.7 (page xxx) gives Minitab output for the regression of blood alcohol

content (BAC) on number of beers consumed. The t test for the hypothesis that the population regression line has slope $\beta = 0$ has $P < 0.001$. The data show a positive linear relationship between BAC and beers. We might expect the intercept α of the population regression line to be 0, because no beers ($x = 0$) should produce no alcohol in the blood ($y = 0$). To test

$$\begin{aligned} H_0: \alpha &= 0 \\ H_a: \alpha &\neq 0 \end{aligned}$$

we use a t statistic formed by dividing the least-squares intercept a by its standard error SE_a . Locate this statistic in the output of Figure 24.7 and verify that it is in fact a divided by its standard error. What is the P -value? Do the data suggest that the intercept is not 0?

24.49 Confidence intervals for the intercept (optional). The output in Figure 24.7 (page 603) allows you to calculate confidence intervals for both the slope β and the intercept α of the population regression line of BAC on beers in the population of all students. Confidence intervals for the intercept α have the familiar form $a \pm t^*SE_a$ with degrees of freedom $n - 2$. What is the 95% confidence interval for the intercept? Does it contain 0, the value we might guess for α ?



EXPLORING THE WEB

24.50 Predicting batting averages. As you did in Exercise 5.59, go to www.mlb.com/ and find the batting averages for a diverse set of 30 players for both the 2009 and 2010 seasons. You can click on the “Stats” tab to find the results for the current season as well as historical data. You should select only players who played in at least 50 games both seasons. Find the least-squares regression line for predicting batting average in 2010 from that in 2009 based on your sample of 30 players. In 2009, the major league leader in batting was Joe Mauer, who had a batting average of .365. Find a 95% prediction interval for the 2010 batting average of someone who hit .365 in 2009. How does this prediction compare with Joe Mauer’s 2010 batting average?

24.51 Olympic medal counts. In Exercise 4.48 you made a scatterplot of the Winter Olympics medal counts for 2002 and 2006. We investigate these medal counts further. Go to the *Chance News* Web site at www.causeweb.org/wiki/chance/index.php/Chance_News_61#Predicting_medal_counts and read the article “Predicting Medal Counts.” Next, search the Web (as you did in Chapter 4) and locate the Winter Olympics medal counts for 2002 and 2006 (I found Winter Olympics medal counts on Wikipedia). Find the equation of the least-squares regression line for predicting the 2006 medal counts from the 2002 counts. Compute 95% confidence intervals for the slope and intercept of your regression line. Are your results consistent with the comment in the *Chance News* article that states “we would have done well simply predicting that the Vancouver totals would match the Torino totals”?



One-Way Analysis of Variance: Comparing Several Means

The two-sample t procedures of Chapter 19 compare the means of two populations or the mean responses to two treatments in an experiment. Of course, studies don't always compare just two groups. We need a method for comparing any number of means.

EXAMPLE 25.1 Comparing tropical flowers

STATE: Ethan Temeles and W. John Kress of Amherst College studied the relationship between varieties of the tropical flower *Heliconia* on the island of Dominica and the different species of hummingbirds that fertilize the flowers.¹ Over time, the researchers believe, the lengths of the flowers and the forms of the hummingbirds' beaks have evolved to match each other. If that is true, flower varieties fertilized by different hummingbird species should have distinct distributions of length.

Table 25.1 gives length measurements (in millimeters) for samples of three varieties of *Heliconia*, each fertilized by a different species of hummingbird. Do the three varieties display distinct distributions of length? In particular, are the mean lengths of their flowers different?

PLAN: Use graphs and numerical descriptions to describe and compare the three distributions of flower length. Finally, ask whether the differences among the mean lengths of the three varieties are *statistically significant*.

SOLVE (first steps): We first met these data in Chapter 2 (page 56), where we compared the distributions. Figure 25.1 repeats a side-by-side stemplot from

IN THIS CHAPTER WE COVER...

- Comparing several means
- The analysis of variance F test
- Using technology
- The idea of analysis of variance
- Conditions for ANOVA
- F distributions and degrees of freedom
- Some details of ANOVA*



FLOWERLENGTH



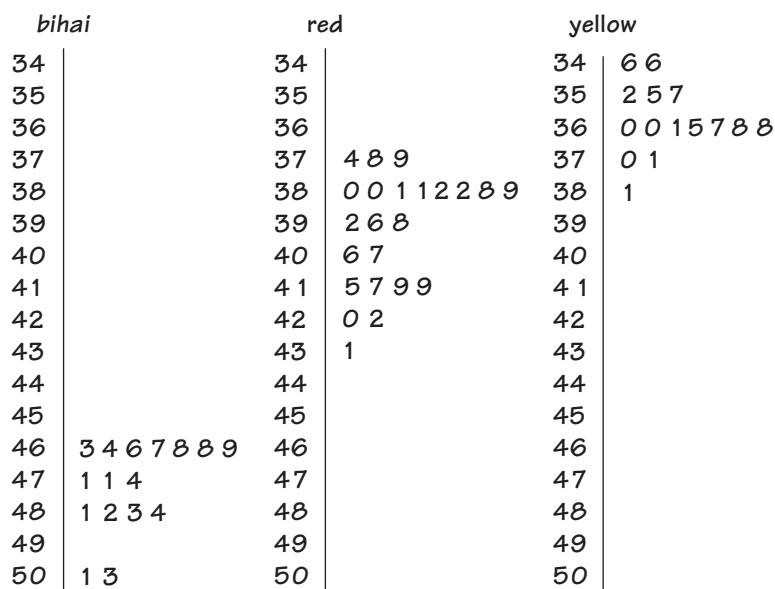
Kevin Shafer/Alamy

Chapter 2. The lengths have been rounded to the nearest tenth of a millimeter. Here are the summary measures we will use in further analysis:

Sample	Variety	Sample size	Mean length	Standard deviation
1	<i>bihai</i>	16	47.60	1.213
2	red	23	39.71	1.799
3	yellow	15	36.18	0.975

TABLE 25.1 Flower lengths (millimeters) for three *Heliconia* varieties

<i>H. bihai</i>								
47.12	46.75	46.81	47.12	46.67	47.43	46.44	46.64	
48.07	48.34	48.15	50.26	50.12	46.34	46.94	48.36	
<i>H. caribaea</i> red								
41.90	42.01	41.93	43.09	41.47	41.69	39.78	40.57	
39.63	42.18	40.66	37.87	39.16	37.40	38.20	38.07	
38.10	37.97	38.79	38.23	38.87	37.78	38.01		
<i>H. caribaea</i> yellow								
36.78	37.02	36.52	36.11	36.03	35.45	38.13	37.10	
35.17	36.82	36.66	35.68	36.03	34.57	34.63		

**FIGURE 25.1**

Side-by-side stemplots comparing the lengths in millimeters of samples of flowers from three varieties of *Heliconia*, from Table 25.1.

CONCLUDE (first steps): The three varieties differ so much in flower length that there is little overlap among them. In particular, the flowers of *bihai* are longer than either red or yellow. The mean lengths are 47.6 mm for *H. bihai*, 39.7 mm for *H. caribaea* red, and 36.2 mm for *H. caribaea* yellow. Are these observed differences in sample means statistically significant? We must develop a test for comparing more than two population means. ■

COMPARING SEVERAL MEANS

Call the mean lengths for the three populations of flowers μ_1 for *bihai*, μ_2 for red, and μ_3 for yellow. The subscript reminds us which group a parameter or statistic describes. To compare these three population means, we might use the two-sample *t* test several times:

- Test $H_0: \mu_1 = \mu_2$ to see if the mean length for *bihai* differs from the mean for red.
- Test $H_0: \mu_1 = \mu_3$ to see if *bihai* differs from yellow.
- Test $H_0: \mu_2 = \mu_3$ to see if red differs from yellow.

The weakness of doing three tests is that we get three *P*-values, one for each test alone. That doesn't tell us how likely it is that *three* sample means are spread apart as far as these are. It may be that $\bar{x}_1 = 47.60$ and $\bar{x}_3 = 36.18$ are significantly different if we look at just two groups but not significantly different if we know that they are the largest and the smallest means in three groups. As we look at more groups, we expect the gap between the largest and smallest sample mean to get larger. (Think of comparing the tallest and shortest person in larger and larger groups of people.) We can't safely compare many parameters by doing tests or confidence intervals for two parameters at a time.



The problem of how to do many comparisons at once with an overall measure of confidence in all our conclusions is common in statistics. This is the problem of **multiple comparisons**. Statistical methods for dealing with multiple comparisons usually have two steps:

1. An *overall test* to see if there is good evidence of *any* differences among the parameters that we want to compare.
2. A detailed *follow-up analysis* to decide which of the parameters differ and to estimate how large the differences are.

multiple comparisons

The overall test, though more complex than the tests we met in Chapters 18 to 21, is reasonably straightforward. Formal follow-up analysis can be quite elaborate. We will concentrate on the overall test and use data analysis to describe in detail the nature of the differences. Companion Chapter 29, on the text CD and Web site, presents some details of follow-up inference.

THE ANALYSIS OF VARIANCE F TEST

We want to test the null hypothesis that there are *no differences* among the mean lengths for the three populations of flowers:

$$H_0: \mu_1 = \mu_2 = \mu_3$$

The basic *conditions for inference* (more detail on page 634) are that we have random samples from the three populations and that flower lengths are Normally distributed in each population.

The alternative hypothesis is that there is *some difference*. That is, not all three population means are equal:

$$H_a: \text{not all of } \mu_1, \mu_2, \text{ and } \mu_3 \text{ are equal}$$

analysis of variance F test

The alternative hypothesis is no longer one-sided or two-sided. It is “many-sided,” because it allows any relationship other than “all three equal.” For example, H_a includes the case in which $\mu_2 = \mu_3$ but μ_1 has a different value. The test of H_0 against H_a is called the **analysis of variance F test**. Analysis of variance is usually abbreviated as ANOVA. The ANOVA F test is almost always carried out by software that reports the test statistic and its P -value.



EXAMPLE 25.2 Comparing tropical flowers: ANOVA

SOLVE (inference): Software tells us that for the flower length data in Table 25.1, the test statistic is $F = 259.12$ with P -value $P < 0.0001$. There is very strong evidence that the three varieties of flowers do not all have the same mean length.

The F test does not say *which* of the three means are significantly different. It appears from our preliminary data analysis that *bihai* flowers are distinctly longer than either red or yellow. Red and yellow are closer together, but the red flowers tend to be longer.

CONCLUDE: There is strong evidence ($P < 0.0001$) that the population means are not all equal. The most important difference among the means is that the *bihai* variety has longer flowers than the red and yellow varieties. ■

Example 25.2 illustrates our approach to comparing means. The ANOVA F test (done by software) assesses the evidence for *some difference* among the population means. Formal follow-up analysis would allow us to say which means differ and by how much, with (say) 95% confidence that *all* our conclusions are correct. We rely instead on examination of the data to show what differences are present and whether they are large enough to be interesting.

APPLY YOUR KNOWLEDGE

25.1 Angry women, sad men. What are the relationships among the portrayal of anger or sadness, sex, and the degree of status conferred? Sixty-eight subjects were randomly assigned to view a videotaped interview in which either a male or a female professional described feeling either anger or sadness. The people being interviewed (we'll call them the “targets”) wore professional attire and were ostensibly being interviewed for a job. The targets described an incident in which they and a colleague lost an account and, when asked by the interviewer how it made them feel, responded either that the incident made them feel angry or that it made them feel sad. The subjects were divided into four groups; each group evaluated one of the following four types of interviews:²

	Male target	Female target
Expressed anger	Group A	Group C
Expressed sadness	Group B	Group D

After watching the interview, subjects evaluated the target on a composite measure of status conferral that included items assessing how much status, power, and independence the target deserved in his or her future job. The measure of status ranged from 1 = none to 11 = a great deal.

- (a) What are the null and alternative hypotheses for the ANOVA F test? Be sure to explain what means the test compares.
- (b) Figure 25.2 is a graph displaying the means for the four groups. What is the approximate size of the mean difference in status conferred on angry men versus angry women? Which of these two groups has the higher mean status?

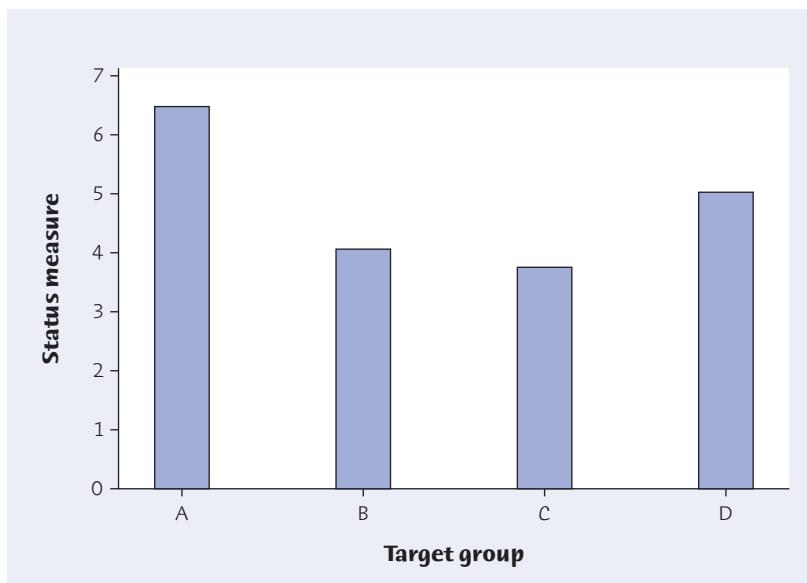


FIGURE 25.2

Bar graph comparing the mean status conferred for the four types of targets, for Exercise 25.1.

- 25.2 Road rage.** “The phenomenon of road rage has been frequently discussed but infrequently examined.” So begins a report based on interviews with 1382 randomly selected drivers.³ The respondents’ answers to interview questions produced scores on an “angry/threatening driving scale” with values between 0 and 19, larger values indicating more angry/threatening behaviors. What driver characteristics go with road rage? There were no significant differences among races or levels of education. What about the effect of the driver’s age? Here are the mean responses for three age groups:

<30 yr	30 to 55 yr	>55 yr
2.22	1.33	0.66

The report says that $F = 34.96$, with $P < 0.01$.

- (a) What are the null and alternative hypotheses for the ANOVA F test? Be sure to explain what means the test compares.
- (b) Based on the sample means and the F test, what do you conclude?

USING TECHNOLOGY

Any technology used for statistics should perform analysis of variance. Figure 25.3 displays ANOVA output for the data of Table 25.1 from a graphing calculator, two statistical programs, and a spreadsheet program.

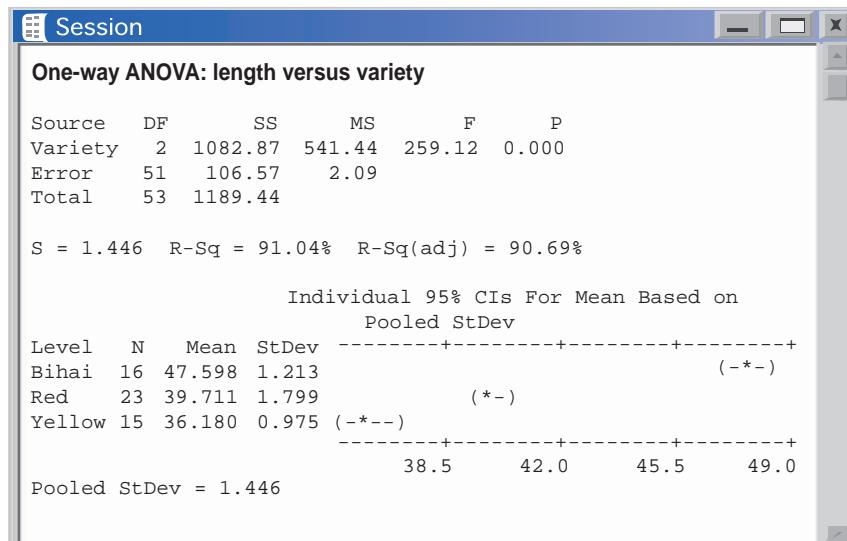
Minitab and Excel give the sizes of the three samples and their means. These agree with those in Example 25.1. Minitab also gives the standard deviations and Excel gives the variances. All four outputs report the F test statistic, $F = 259.1$, and its P -value. Minitab sensibly reports the P -value as 0 to three decimal places, while CrunchIt! reports that $P < 0.0001$. This is all we need to know about the P -value in practice. Excel and the graphing calculator offer the specific value 1.92×10^{-27} . (This would be correct if the population distributions were exactly Normal. In practice, read such values simply as “ P is very small.”) There is very strong evidence that the three varieties of flowers do not all have the same mean length.

All four outputs report degrees of freedom (df), sums of squares (SS), and mean squares (MS). We don’t need this information now. Minitab also gives confidence intervals for all three means that help us see which means differ and by how much. None of the intervals overlap, and *bihai* is much above the other two. These are 95% confidence intervals for each mean separately. We are *not* 95% confident that *all three* intervals cover the three means. This is another example of the peril of multiple comparisons.

Texas Instruments Graphing Calculator

One-way ANOVA $F=259.1192995$ $p=1.918818e-27$ Factor $df=2$ $SS=1082.87237$ $\downarrow MS=541.436183$	One-way ANOVA $F=541.436183$ Error $df=51$ $SS=106.565761$ $MS=2.08952472$ $S\times P=1.44551884$
---	---

Minitab



The Minitab session window displays the following output:

```

Session

One-way ANOVA: length versus variety

Source      DF        SS         MS          F       P
Variety      2    1082.87   541.44    259.12  0.000
Error       51    106.57    2.09
Total       53   1189.44

S = 1.446  R-Sq = 91.04%  R-Sq(adj) = 90.69%

Individual 95% CIs For Mean Based on
Pooled StDev
Level      N      Mean     StDev
Bihai      16    47.598   1.213      (*)-
Red        23    39.711   1.799      (*-)
Yellow     15    36.180   0.975      (-*-)
                                         38.5      42.0      45.5      49.0
Pooled StDev = 1.446

```

FIGURE 25.3

ANOVA for the flower length data: output from a graphing calculator, two statistical programs, and a spreadsheet program.

CrunchIt!

Source	Sum of Squares	df	Mean Square	F-value	P-value
Variety	1083	2	541.4	259.1	<0.0001
Error	106.6	51	2.090		
Total	1189	53			

FIGURE 25.3

(Continued)

Excel

Microsoft Excel - ta25-01.dat						
1	A	B	C	D	E	G
2	Anova: Single Factor					
3	SUMMARY					
4	Groups	Count	Sum	Average	Variance	
5	bihai	16	761.56	47.5975	1.471073	
6	red	23	913.36	39.7113	3.235548	
7	yellow	15	542.7	36.18	0.951257	
8						
9						
10	ANOVA					
11	Source of variation	SS	df	MS	F	P-value F crit
12	Between Groups	1082.872	2	541.4362	259.1193	1.92E-27 3.178799
13	Within Groups	106.5658	51	2.089525		
14						
15	Total	1189.438	53			


APPLY YOUR KNOWLEDGE

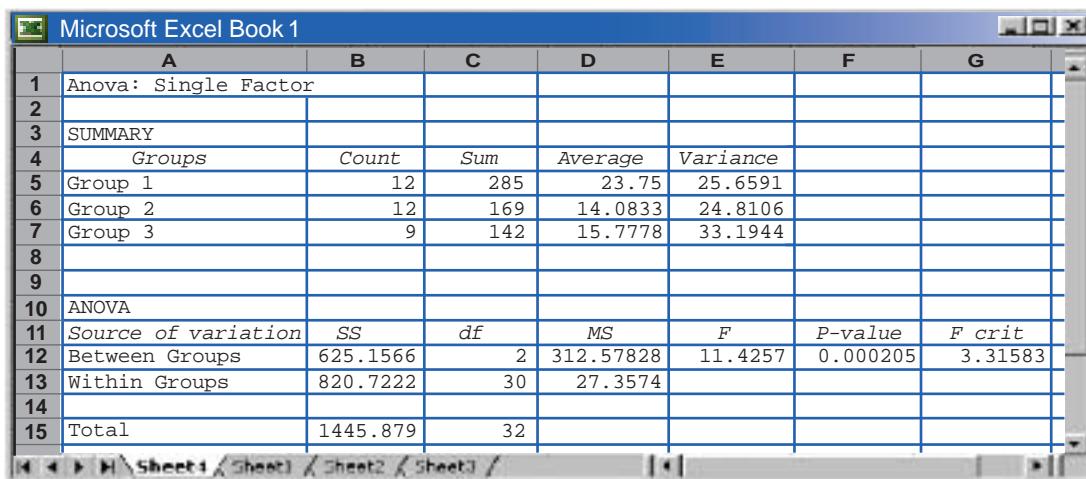
25.3 Logging in the rain forest. How does logging in a tropical rain forest affect the forest in later years? Researchers compared forest plots in Borneo that had never been logged (Group 1) with similar plots nearby that had been logged 1 year earlier (Group 2) and 8 years earlier (Group 3). Although the study was not an experiment, the authors explain why we can consider the plots to be randomly selected. The data appear in Table 25.2. The variable Trees is the count of trees in a plot; Species is the count of tree species in a plot. The variable Richness is Species/Trees, the number of species divided by the number of individual trees.⁴ 

- (a) Make side-by-side stemplots of Trees for the three groups. Use stems 0, 1, 2, and 3 and split the stems (see page 22). What effects of logging are visible?

TABLE 25.2 Data from a study of logging in Borneo

GROUP	TREES	SPECIES	RICHNESS	GROUP	TREES	SPECIES	RICHNESS
1	27	22	0.81481	2	18	15	0.83333
1	22	18	0.81818	2	17	15	0.88235
1	29	22	0.75862	2	14	12	0.85714
1	21	20	0.95238	2	14	13	0.92857
1	19	15	0.78947	2	2	2	1.00000
1	33	21	0.63636	2	17	15	0.88235
1	16	13	0.81250	2	19	8	0.42105
1	20	13	0.65000	3	18	17	0.94444
1	24	19	0.79167	3	4	4	1.00000
1	27	13	0.48148	3	22	18	0.81818
1	28	19	0.67857	3	15	14	0.93333
1	19	15	0.78947	3	18	18	1.00000
2	12	11	0.91667	3	19	15	0.78947
2	12	11	0.91667	3	22	15	0.68182
2	15	14	0.93333	3	12	10	0.83333
2	9	7	0.77778	3	12	12	1.00000
2	20	18	0.90000				

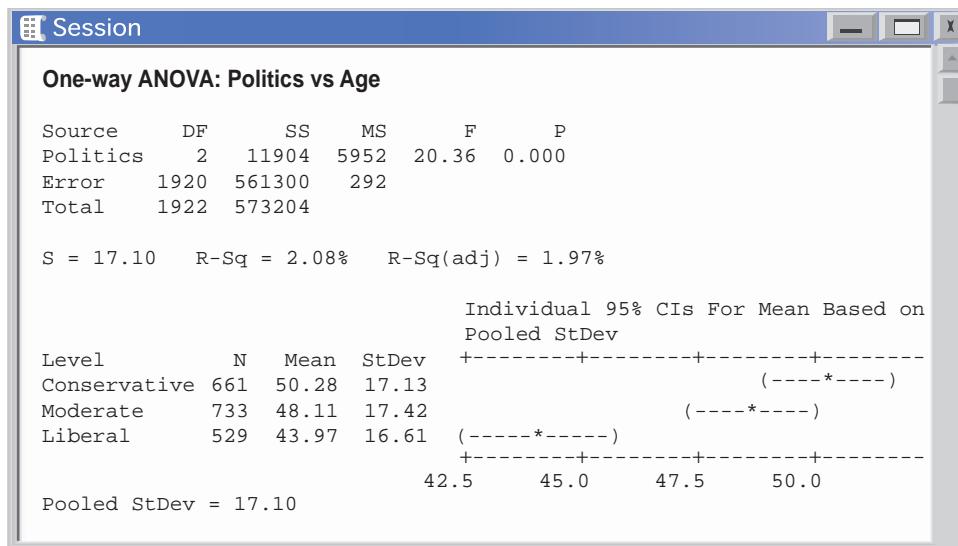
- (b) Figure 25.4 shows Excel ANOVA output for Trees. What do the group means show about the effects of logging?
- (c) What are the ANOVA F statistic and its P-value? What hypotheses does F test? What conclusions about the effects of logging on number of trees do the data lead to?

**FIGURE 25.4**

Excel output for analysis of variance on the number of trees in forest plots, for Exercise 25.3.

25.4 Political views and age. The University of Chicago's General Social Survey (GSS) is the nation's most important social science sample survey. The GSS asked a random sample of adults in 2008 both their age and where they placed themselves on the political spectrum from extremely liberal to extremely conservative. Is there a relationship between age and place on the political spectrum? The political spectrum categories in the original survey included slightly liberal, liberal, and extremely liberal, but these have been combined into the single category liberal, and similarly with conservative.⁵

- Figure 25.5 gives the Minitab ANOVA output for these data. What do the mean ages say about the relationship between age and political views?
- What are the ANOVA F statistic and its P-value? What hypotheses does F test? Briefly describe the conclusions you draw from these data.

**FIGURE 25.5**

Minitab output for the data on respondent's ages for three different political viewpoints, for Exercise 25.4.

THE IDEA OF ANALYSIS OF VARIANCE

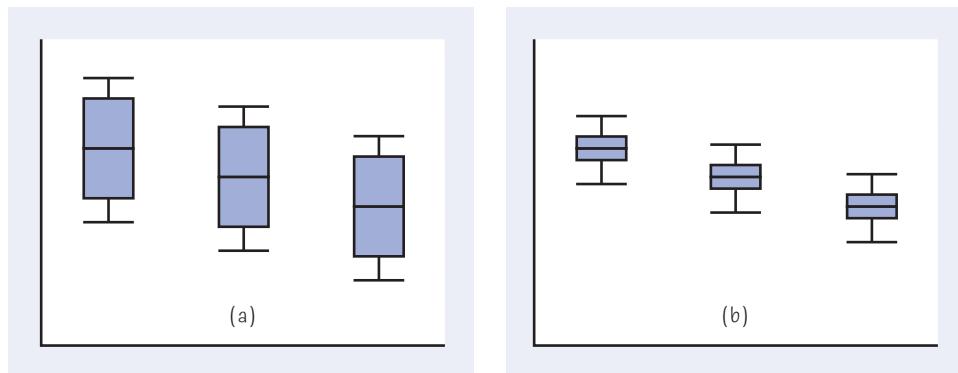
The details of ANOVA are a bit daunting (they appear in an optional section at the end of this chapter). The main idea of ANOVA is more accessible and much more important. Here it is: when we ask if a set of sample means gives evidence for differences among the population means, what matters is not how far apart the sample means are but how far apart they are *relative to the variability of individual observations*.

Look at the two sets of boxplots in Figure 25.6. For simplicity, these distributions are all symmetric, so that the mean and median are the same. The center line in each boxplot is therefore the sample mean. Both sets of boxplots compare three samples with the same three means. Could differences this large easily arise just due to chance, or are they statistically significant?

- The boxplots in Figure 25.6(a) have tall boxes, which indicates lots of variation among the individuals in each group. With this much variation among individuals, we would not be surprised if another set of samples gave

FIGURE 25.6

Boxplots for two sets of three samples each. The sample means are the same in (a) and (b). Analysis of variance will find a more significant difference among the means in (b) because there is less variation among the individuals within those samples.



quite different sample means. The observed differences among the sample means could easily happen just by chance.

- The boxplots in Figure 25.6(b) have the same centers as those in Figure 25.6(a), but the boxes are much shorter. That is, there is much less variation among the individuals in each group. It is unlikely that any sample from the first group would have a mean as small as the mean of the second group. Because means as far apart as those observed would rarely arise just by chance in repeated sampling, they are good evidence of real differences among the means of the three populations we are sampling from.

You can use the *One-Way ANOVA* applet to demonstrate the analysis of variance idea for yourself. The applet allows you to change both the group means and the spread within groups. You can watch the ANOVA F statistic and its P-value change as you work.

This comparison of the two parts of Figure 25.6 is too simple in one way. It ignores the effect of the sample sizes, an effect that boxplots do not show. Small differences among sample means can be significant if the samples are large. Large differences among sample means can fail to be significant if the samples are small. All we can be sure of is that for the same sample size, Figure 25.6(b) will give a much smaller P-value than Figure 25.6(a). Despite this qualification, the big idea remains: if sample means are far apart relative to the variation among individuals in the same groups, that's evidence that something other than chance is at work.



THE ANALYSIS OF VARIANCE IDEA

Analysis of variance compares the variation due to specific sources with the variation among individuals who should be similar. In particular, ANOVA tests whether several populations have the same mean by comparing how far apart the sample means are with how much variation there is within the samples.

It is one of the oddities of statistical language that methods for comparing means are named after the variance. The reason is that the test works by comparing two kinds of variation. Analysis of variance is a general method for studying sources of variation in responses. Comparing several means is the simplest form of ANOVA, called **one-way ANOVA**.

THE ANOVA F STATISTIC

The analysis of variance F statistic for testing the equality of several means has this form:

$$F = \frac{\text{variation among the sample means}}{\text{variation among individuals in the same sample}}$$

If you want more detail, read the optional section at the end of this chapter. The F statistic can take only values that are zero or positive. It is zero only when all the sample means are identical and gets larger as they move farther apart. Large values of F are evidence against the null hypothesis H_0 that all population means are the same. Although the alternative hypothesis H_a is many-sided, the ANOVA F test is one-sided because any violation of H_0 tends to produce a large value of F.

APPLY YOUR KNOWLEDGE

25.5 ANOVA compares several means. The One-Way ANOVA applet displays the observations in three groups, with the group means highlighted by black dots. When you open or reset the applet, the scale at the bottom of the display shows that for these groups the ANOVA F statistic is $F = 31.74$, with $P < 0.001$. (The P-value is marked by a red dot that moves along the scale.)



- The middle group has a larger mean than the other two. Grab its mean point with the mouse. How small can you make F? What did you do to the mean to make F small? Roughly how significant is your small F?
- Starting with the three means aligned from your configuration at the end of (a), drag any one of the group means either up or down. What happens to F? What happens to the P-value? Convince yourself that the same thing happens if you move any one of the means, or if you move one slightly and then another slightly in the opposite direction.

25.6 ANOVA uses within-group variation. Reset the One-Way ANOVA applet to its original state. As in Figure 25.6(b), the differences among the three means are highly significant (large F, small P-value) because the observations in each group cluster tightly about the group mean.



- Use the mouse to slide the Pooled Standard Error at the top of the display to the right. You see that the group means do not change, but the spread of the observations in each group increases. What happens to F and P as the spread among the observations in each group increases? What are the values of F and P when the slider is all the way to the right? This is similar to Figure 25.6(a): variation within groups hides the differences among the group means.
- Leave the Pooled Standard Error slider at the extreme right of its scale, so that spread within groups stays fixed. Use the mouse to move the group means apart. What happens to F and P as you do this?

CONDITIONS FOR ANOVA

Like all inference procedures, ANOVA is valid only in some circumstances. Here are the conditions under which we can use ANOVA to compare population means.

CONDITIONS FOR ANOVA INFERENCE

- We have **I independent SRSs**, one from each of I populations. We measure the same response variable for each sample.
- The i th population has a **Normal distribution** with unknown mean μ_i . One-way ANOVA tests the null hypothesis that all the population means are the same.
- All the populations have the **same standard deviation** σ , whose value is unknown.

The first two conditions are familiar from our study of the two-sample t procedures for comparing two means. As usual, the design of the data production is the most important condition for inference. Biased sampling or confounding can make any inference meaningless. *If we do not actually draw separate SRSs from each population or carry out a randomized comparative experiment, it may be unclear to what population the conclusions of inference apply.* ANOVA, like other inference procedures, is often used when random samples are not available. You must judge each use on its merits, a judgment that usually requires some knowledge of the subject of the study in addition to some knowledge of statistics.

Because no real population has exactly a Normal distribution, the usefulness of inference procedures that assume Normality depends on how sensitive they are to departures from Normality. Fortunately, procedures for comparing means are not very sensitive to lack of Normality. The ANOVA F test, like the t procedures, is **robust**. What matters is Normality of the sample means, so ANOVA becomes safer as the sample sizes get larger, because of the central limit theorem effect. Remember to check for outliers that change the value of sample means and for extreme skewness. When there are no outliers and the distributions are roughly symmetric, you can safely use ANOVA for sample sizes as small as 4 or 5.

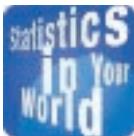
The third condition is annoying: ANOVA assumes that the variability of observations, measured by the standard deviation, is the same in all populations. The t test for comparing two means (Chapter 19) does not require equal standard deviations. Unfortunately, the ANOVA F for comparing more than two means is less broadly valid. It is not easy to check the condition that the populations have equal standard deviations. Statistical tests for equality of standard deviations are very sensitive to lack of Normality, so much so that they are of little practical value. You must either seek expert advice or rely on the robustness of ANOVA.

How serious are unequal standard deviations? ANOVA is not too sensitive to violations of the condition, especially when all samples have the same or similar sizes and no sample is very small. When designing a study, try to take samples of about the same size from all the groups you want to compare. The sample standard deviations estimate the population standard deviations, so check before doing ANOVA that the sample standard deviations are similar to each other. We expect some variation among them due to chance. Here is a rule of thumb that is safe in almost all situations.

CHECKING STANDARD DEVIATIONS IN ANOVA

The results of the ANOVA F test are approximately correct when the largest sample standard deviation is no more than twice as large as the smallest sample standard deviation.

robustness



We weren't working anyway

A “consultant” estimated that the annual NCAA men’s basketball tournament costs employers \$3.8 billion in time wasted by workers participating in office pools, checking game scores, and so on. That’s unlikely. Most of the games are played outside work hours, at night and on weekends. More to the point, economists note that workers waste lots of time every workday, by talking to other employees, chatting on the phone, shopping online, and so on. The dollar value of time spent on the basketball tournament probably comes in large part from time we were wasting anyway.



EXAMPLE 25.3 Comparing tropical flowers: conditions for ANOVA

The study of *Heliconia* blossoms is based on three independent samples that the researchers consider to be random samples from all flowers of these varieties in Dominica. The stemplots in Figure 25.1 show that the *bihai* and red varieties have slightly skewed distributions, but the sample means of samples of sizes 16 and 23 will have distributions that are close to Normal. The sample standard deviations for the three varieties are

$$s_1 = 1.213 \quad s_2 = 1.799 \quad s_3 = 0.975$$

These standard deviations satisfy our rule of thumb:

$$\frac{\text{largest } s}{\text{smallest } s} = \frac{1.799}{0.975} = 1.85 \quad (\text{less than } 2)$$

We can safely use ANOVA to compare the mean lengths for the three populations. ■

EXAMPLE 25.4 Thinking about money changes behavior

STATE: Kathleen Vohs of the University of Minnesota and her coworkers carried out several randomized comparative experiments on the effects of thinking about money. Here's an outline of one of the experiments. Ask student subjects to unscramble 30 sets of five words to make a meaningful phrase from four of the five. The control group unscrambled phrases like "cold it desk outside is" into "it is cold outside." The "play money" group unscrambled similar sets of words, but a stack of Monopoly money was placed nearby. The "money prime" group unscrambled phrases that lead to thinking about money, turning "high a salary desk paying" into "a high-paying salary." Then each subject worked a hard puzzle, knowing that he or she could ask for help. Table 25.3 shows the time in seconds that each subject worked on the puzzle before asking for help.⁶ Psychologists think that money tends to make people self-sufficient. If so, the two groups that were encouraged in different ways to think about money should take longer on the average to ask for help. Do the data support this idea?



THINKMONEY

PLAN: Examine the data to compare the effect of the treatments and check that we can safely use ANOVA. If the data allow ANOVA, assess the significance of observed differences in mean times to ask for help.

SOLVE: Figure 25.7 shows side-by-side stemplots of the data in the three groups. We expect some irregularity in small samples, but there are no outliers or strong skewness that would hinder use of ANOVA. The Minitab ANOVA output in Figure 25.8 shows that the group standard deviations easily satisfy our rule of thumb. The control group subjects asked for help much sooner (mean 186.1 seconds) than did subjects in the two money groups (means 305.2 seconds and 314.1 seconds). The three means are significantly different ($F = 3.73$, $P = 0.031$).

CONCLUDE: The experiment gives good evidence that reminding people of money in either of two ways does make them less willing to ask others for help. This is consistent with the idea that money makes people feel more self-sufficient. ■

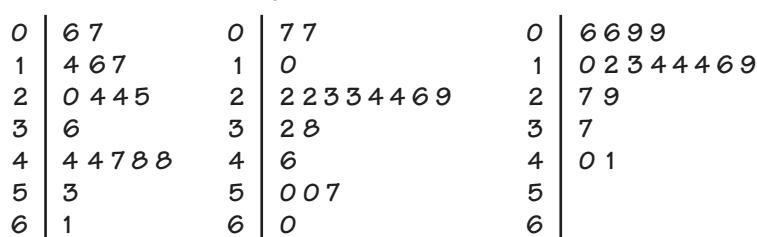
TABLE 25.3 Time (seconds) until subjects ask for help with a puzzle

GROUP	TIME	GROUP	TIME	GROUP	TIME
Prime	609	Play	455	Control	118
Prime	444	Play	100	Control	272
Prime	242	Play	238	Control	413
Prime	199	Play	243	Control	291
Prime	174	Play	500	Control	140
Prime	55	Play	570	Control	104
Prime	251	Play	231	Control	55
Prime	466	Play	380	Control	189
Prime	443	Play	222	Control	126
Prime	531	Play	71	Control	400
Prime	135	Play	232	Control	92
Prime	241	Play	219	Control	64
Prime	476	Play	320	Control	88
Prime	482	Play	261	Control	142
Prime	362	Play	290	Control	141
Prime	69	Play	495	Control	373
Prime	160	Play	600	Control	156
		Play	67		

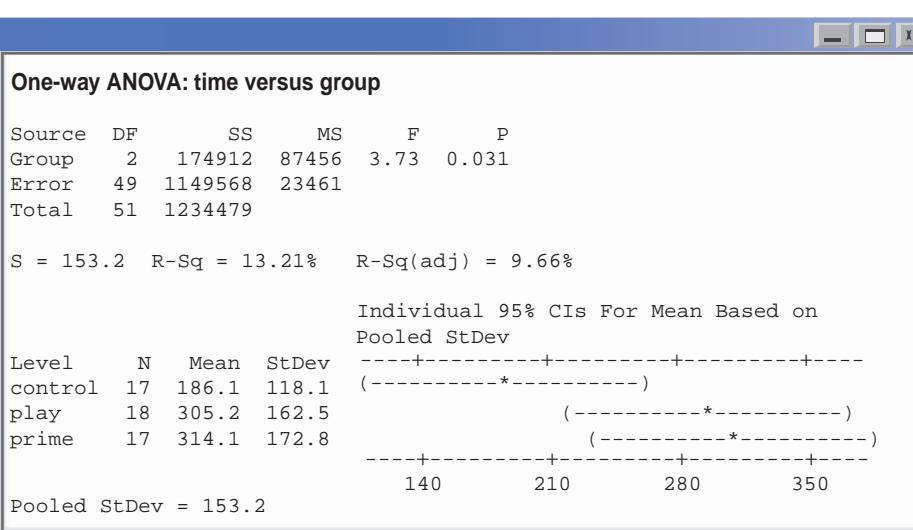
FIGURE 25.7

Side-by-side stemplots comparing the time until subjects asked for help with a puzzle, for Example 25.4.

Prime Play Control

**FIGURE 25.8**

Minitab ANOVA output for comparing the three treatments in Example 25.4.




APPLY YOUR KNOWLEDGE

25.7 Checking standard deviations. Verify that the sample standard deviations for these sets of data do allow use of ANOVA to compare the population means.

- The counts of trees in Exercise 25.3 and Figure 25.4.
- The ages of Exercise 25.4 and Figure 25.5.

25.8 Species richness after logging. Table 25.2 gives data on the species richness in rain forest plots, defined as the number of tree species in a plot divided by the number of trees in the plot. ANOVA may not be trustworthy for the richness data. Do data analysis: make side-by-side stemplots to examine the distributions of the response variable in the three groups, and also compare the standard deviations. What characteristic of the data makes ANOVA risky?  BORNEOLOGGING



25.9 Fertilizing bromeliads. Bromeliads are tropical flowering plants. Many are epiphytes that attach to trees and obtain moisture and nutrients from air and rain. Their leaf bases form cups that collect water and are home to the larvae of many insects. As a preliminary to a study of changes in the nutrient cycle, Jacqueline Ngai and Diane Srivastava examined the effects of adding nitrogen, phosphorus, or both to the cups. They randomly assigned 8 bromeliads growing in Costa Rica to each of four treatment groups, including an unfertilized control group. A monkey destroyed one of the plants in the control group, leaving 7 bromeliads in that group. Here are the numbers of new leaves on each plant over the 7 months following fertilization:⁷  BROMELIADS

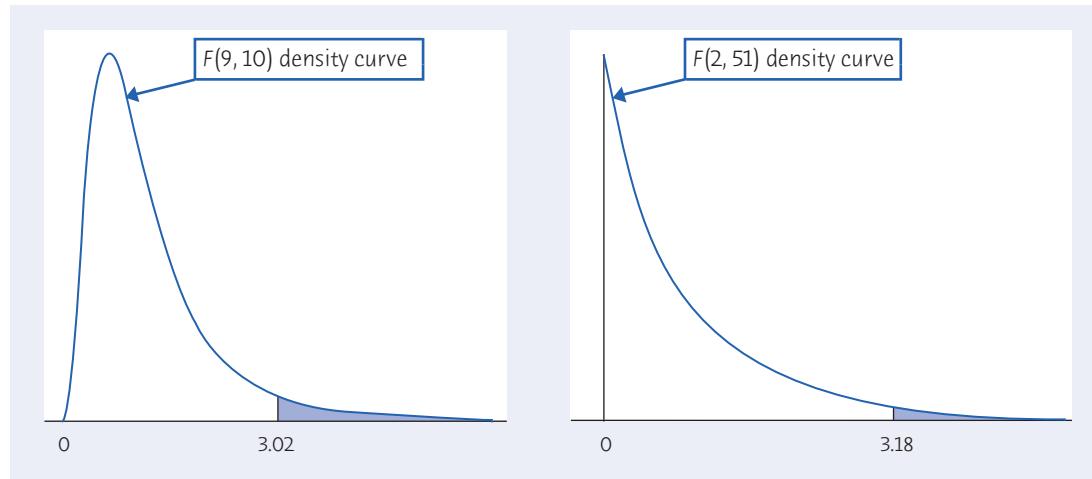
Nitrogen	Phosphorus	Both	Neither
15	14	14	11
14	14	16	13
15	14	15	16
16	11	14	15
17	13	14	15
18	12	13	11
17	15	17	12
13	15	14	

Analyze these data and discuss the results. Does nitrogen or phosphorus have a greater effect on the growth of bromeliads? Follow the four-step process as illustrated in Example 25.4.

F DISTRIBUTIONS AND DEGREES OF FREEDOM

The ANOVA F statistic is

$$F = \frac{\text{variation among the sample means}}{\text{variation among individuals in the same sample}}$$

**FIGURE 25.9**

Density curves for two F distributions. Both are right-skewed and take only positive values. The upper 5% critical values are marked under the curves.

F distribution

To find the P -value for this statistic, we must know the sampling distribution of F when the null hypothesis (all population means equal) is true. This sampling distribution is an **F distribution**.

The F distributions are a family of right-skewed distributions that take only values greater than 0. The density curves in Figure 25.9 illustrate their shapes. A specific F distribution is determined by the *degrees of freedom* of the numerator and denominator of the F statistic. You may have noticed that all our software outputs include degrees of freedom, labeled either “df” or “DF.” The optional section “Some Details of ANOVA” shows where the degrees of freedom come from. When describing an F distribution, always give the numerator degrees of freedom first. Our brief notation will be $F(df_1, df_2)$ for the F distribution with df_1 degrees of freedom in the numerator and df_2 in the denominator. *Interchanging the degrees of freedom changes the distribution, so the order is important.*

Tables of F critical points are awkward, because we need a separate table for every pair of degrees of freedom df_1 and df_2 . Fortunately, software gives you P -values for the ANOVA F test without the need for a table.

EXAMPLE 25.5 Comparing flowers: the F distribution

Look again at the software output in Figure 25.3 for the flower length data. All four outputs give the degrees of freedom for the F test, labeled “df” or “DF.” There are 2 degrees of freedom in the numerator and 51 in the denominator. P -values for the F test therefore come from the F distribution with 2 and 51 degrees of freedom: $F(2, 51)$. The right-hand curve in Figure 25.9 is the density curve of this distribution. The 5% critical value marked on that curve is 3.18, and the 1% critical value is 5.05. The observed value $F = 259.12$ of the ANOVA F statistic lies far to the right of these values, so the P -value is extremely small. ■

The degrees of freedom of the ANOVA F statistic depend on the number of means we are comparing and the number of observations in each sample. That is, the F test takes into account the number of observations. Here are the details.

DEGREES OF FREEDOM FOR THE F TEST

We want to compare the means of I populations. We have an SRS of size n_i from the i th population, so that the total number of observations in all samples combined is

$$N = n_1 + n_2 + \cdots + n_I$$

If the null hypothesis that all population means are equal is true, the ANOVA F statistic has the F distribution with $I - 1$ degrees of freedom in the numerator and $N - I$ degrees of freedom in the denominator.

EXAMPLE 25.6 Degrees of freedom for F

In Examples 25.1 and 25.2, we compared the mean lengths for three varieties of flowers, so $I = 3$. The three sample sizes are

$$n_1 = 16 \quad n_2 = 23 \quad n_3 = 15$$

The total number of observations is therefore

$$N = 16 + 23 + 15 = 54$$

The ANOVA F test has numerator degrees of freedom

$$I - 1 = 3 - 1 = 2$$

and denominator degrees of freedom

$$N - I = 54 - 3 = 51$$

These are the degrees of freedom given in the outputs in Figure 25.3. ■

APPLY YOUR KNOWLEDGE

25.10 Logging in the rain forest, continued. Exercise 25.3 (page 629) compares the number of tree species in rain forest plots that had never been logged (Group 1) with similar plots nearby that had been logged 1 year earlier (Group 2) and 8 years earlier (Group 3).

- What are I , the n_i , and N for these data? Identify these quantities in words and give their numerical values.
- Find the degrees of freedom for the ANOVA F statistic. Check your work against the Excel output in Figure 25.4.

25.11 What music will you play? People often match their behavior to their social environment. One study of this idea first established that the type of music most preferred by black college students is R&B and that whites' most preferred music is rock. Will students hosting a small group of other students choose music that



PhotoAlto/Alamy

matches the makeup of the people attending? Two studies were done, using either black or white students as subjects. In the first study, 90 black business students were assigned at random to three equal-sized groups, and in the second study the same was done for 96 white students. In both studies, each subject sees a picture of the people he or she will host. Group 1 sees 6 blacks, Group 2 sees 3 whites and 3 blacks, and Group 3 sees 6 whites. Ask how likely the host is to play the type of music preferred by the other race. Use ANOVA to compare the three groups to see whether the racial mix of the gathering affects the choice of music.⁸

- (a) For the white subjects, $F = 16.48$. What are the degrees of freedom?
- (b) For the black subjects, $F = 2.47$. What are the degrees of freedom?

SOME DETAILS OF ANOVA*

Now we will give the actual formula for the ANOVA F statistic. We have SRSs from each of I populations. Subscripts from 1 to I tell us which sample a statistic refers to:

Population	Sample size	Sample mean	Sample std. dev.
1	n_1	\bar{x}_1	s_1
2	n_2	\bar{x}_2	s_2
\vdots	\vdots	\vdots	\vdots
I	n_I	\bar{x}_I	s_I

You can find the F statistic from just the sample sizes n_i , the sample means \bar{x}_i , and the sample standard deviations s_i . You don't need to go back to the individual observations.

The ANOVA F statistic has the form

$$F = \frac{\text{variation among the sample means}}{\text{variation among individuals in the same sample}}$$

mean squares

The measures of variation in the numerator and denominator of F are called **mean squares**. A mean square is a more general form of a sample variance. An ordinary sample variance s^2 is an average (or mean) of the squared deviations of observations from their mean, so it qualifies as a "mean square."

Call the overall mean response \bar{x} . That is, \bar{x} is the mean of all N observations together. You can find \bar{x} from the I sample means by

$$\bar{x} = \frac{\text{sum of all observations}}{N} = \frac{n_1\bar{x}_1 + n_2\bar{x}_2 + \cdots + n_I\bar{x}_I}{N}$$

(This expression works because multiplying a group mean \bar{x}_i by the number of observations n_i it represents gives the sum of the observations in that group.)

*This more advanced section is optional if you are using software to find the F statistic.

The numerator of F is a mean square that measures variation among the I sample means $\bar{x}_1, \bar{x}_2, \dots, \bar{x}_I$. To measure this variation, look at the I deviations of the sample means from \bar{x} ,

$$\bar{x}_1 - \bar{x}, \bar{x}_2 - \bar{x}, \dots, \bar{x}_I - \bar{x}$$

The mean square in the numerator of F is an average of the squares of these deviations. We call it the **mean square for groups**, abbreviated as **MSG**:

$$\text{MSG} = \frac{n_1(\bar{x}_1 - \bar{x})^2 + n_2(\bar{x}_2 - \bar{x})^2 + \dots + n_I(\bar{x}_I - \bar{x})^2}{I - 1}$$

Each squared deviation is weighted by n_i , the number of observations it represents.

The mean square in the denominator of F measures variation among individual observations in the same sample. For any one sample, the sample variance s_i^2 does this job. For all I samples together, we use an average of the individual sample variances. It is another weighted average, in which each s_i^2 is weighted by its degrees of freedom $n_i - 1$. The resulting mean square is called the **mean square for error**, **MSE**:

$$\text{MSE} = \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2 + \dots + (n_I - 1)s_I^2}{N - I}$$

“Error” doesn’t mean a mistake has been made. It’s a traditional term for chance variation. Here is a summary of the ANOVA test.

THE ANOVA F TEST

Draw an independent SRS from each of I Normal populations that have a common standard deviation but may have different means. The sample from the i th population has size n_i , sample mean \bar{x}_i , and sample standard deviation s_i .

To test the null hypothesis that all I populations have the same mean against the alternative hypothesis that not all the means are equal, calculate the **ANOVA F statistic**

$$F = \frac{\text{MSG}}{\text{MSE}}$$

The numerator of F is the **mean square for groups**

$$\text{MSG} = \frac{n_1(\bar{x}_1 - \bar{x})^2 + n_2(\bar{x}_2 - \bar{x})^2 + \dots + n_I(\bar{x}_I - \bar{x})^2}{I - 1}$$

The denominator of F is the **mean square for error**

$$\text{MSE} = \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2 + \dots + (n_I - 1)s_I^2}{N - I}$$

When H_0 is true, F has the **F distribution** with $I - 1$ and $N - I$ degrees of freedom.

sums of squares**ANOVA table**

The denominators in the formulas for MSG and MSE are the two degrees of freedom $I - 1$ and $N - I$ of the F test. The numerators are called **sums of squares**, from their algebraic form. It is usual to present the results of ANOVA in an **ANOVA table**. Output from software usually includes an ANOVA table.

EXAMPLE 25.7 ANOVA calculations: software

Look again at the four outputs in Figure 25.3. The three software outputs give the ANOVA table. The calculator, with its small screen, gives the degrees of freedom, sums of squares, and mean squares separately. Each output uses slightly different language to identify the two sources of variation. The basic ANOVA table is

Source of variation	df	SS	MS	F statistic
Variation among samples	2	1082.87	MSG = 541.44	259.12
Variation within samples	51	106.57	MSE = 2.09	

You can check that each mean square MS is the corresponding sum of squares SS divided by its degrees of freedom df. The F statistic is MSG divided by MSE. ■

pooled standard deviation

Because MSE is an average of the individual sample variances, it is also called the *pooled sample variance*, written as s_p^2 . When all I populations have the same population variance σ^2 , as ANOVA assumes that they do, s_p^2 estimates the common variance σ^2 . The square root of MSE is the **pooled standard deviation** s_p . It estimates the common standard deviation σ of observations in each group. The Minitab and calculator outputs in Figure 25.3 give the value $s_p = 1.446$.

The pooled standard deviation s_p is a better estimator of the common σ than any individual sample standard deviation s_i because it combines (pools) the information in all I samples. We can get a confidence interval for any one of the means μ_i from the usual form

$$\text{estimate} \pm t^* \text{SE}_{\text{estimate}}$$

using s_p to estimate σ . The confidence interval for μ_i is

$$\bar{x}_i \pm t^* \frac{s_p}{\sqrt{n_i}}$$

Use the critical value t^* from the t distribution with $N - I$ degrees of freedom because s_p has $N - I$ degrees of freedom. These are the confidence intervals that appear in Minitab ANOVA output.

EXAMPLE 25.8 ANOVA calculations: without software

We can do the ANOVA test comparing the mean lengths of *bihai*, red, and yellow flower varieties using only the sample sizes, sample means, and sample standard deviations. These appear in Example 25.1, but it is easy to find them with a calculator. There are $I = 3$ groups with a total of $N = 54$ flowers.

The overall mean of the 54 lengths in Table 25.1 is

$$\begin{aligned}\bar{x} &= \frac{n_1\bar{x}_1 + n_2\bar{x}_2 + n_3\bar{x}_3}{N} \\ &= \frac{(16)(47.598) + (23)(39.711) + (15)(36.180)}{54} \\ &= \frac{2217.621}{54} = 41.067\end{aligned}$$

The mean square for groups is

$$\begin{aligned}\text{MSG} &= \frac{n_1(\bar{x}_1 - \bar{x})^2 + n_2(\bar{x}_2 - \bar{x})^2 + n_3(\bar{x}_3 - \bar{x})^2}{I - 1} \\ &= \frac{1}{3 - 1} [(16)(47.598 - 41.067)^2 + (23)(39.711 - 41.067)^2 \\ &\quad + (15)(36.180 - 41.067)^2] \\ &= \frac{1082.996}{2} = 541.50\end{aligned}$$

The mean square for error is

$$\begin{aligned}\text{MSE} &= \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2 + (n_3 - 1)s_3^2}{N - I} \\ &= \frac{(15)(1.213^2) + (22)(1.799^2) + (14)(0.975^2)}{51} \\ &= \frac{106.580}{51} = 2.09\end{aligned}$$

Finally, the ANOVA test statistic is

$$F = \frac{\text{MSG}}{\text{MSE}} = \frac{541.50}{2.09} = 259.09$$

Our work differs slightly from the output in Figure 25.3 because of roundoff error. We don't recommend doing these calculations, because tedium and roundoff errors cause frequent mistakes. ■

APPLY YOUR KNOWLEDGE

The calculations of ANOVA use only the sample sizes n_i , the sample means \bar{x}_i , and the sample standard deviations s_i . You can therefore re-create the ANOVA calculations when a report gives these summaries but does not give the actual data. These optional exercises ask you to do the ANOVA calculations starting with the summary statistics. P-values require either a table or software for the F distributions.

25.12 Road rage. Exercise 25.2 describes a study of road rage. Here are the means and standard deviations for a measure of “angry/threatening driving” for random samples of drivers in three age groups:

Age group	<i>n</i>	\bar{x}	<i>s</i>
Less than 30 yr	244	2.22	3.11
30 to 55 yr	734	1.33	2.21
Over 55 yr	364	0.66	1.60

- (a) The distributions of responses are somewhat right-skewed. ANOVA is nonetheless safe for these data. Why?
- (b) Check that the standard deviations satisfy the guideline for ANOVA inference.
- (c) Calculate the overall mean response \bar{x} , the mean squares MSG and MSE, and the ANOVA F statistic.
- (d) Which F distribution would you use to find the P-value of the ANOVA F test? Software gives $P < 0.001$. Write a brief conclusion based on the sample means and the ANOVA F test.

25.13 Angry women, sad men. Exercise 25.1 describes a study in which subjects conferred status on men and women who were displaying either anger or sadness. Subjects were randomly assigned to the four treatments, with 17 subjects assigned to each treatment. The study report contains the following information about status conferred for each of the four groups:

Treatment	n	\bar{x}	s
Males expressing anger	17	6.47	2.25
Females expressing anger	17	3.75	1.77
Males expressing sadness	17	4.05	1.61
Females expressing sadness	17	5.02	1.80

- (a) Do the standard deviations satisfy the rule of thumb for safe use of ANOVA?
- (b) Calculate the overall mean response \bar{x} , the mean squares MSG and MSE, and the F statistic.
- (c) Which F distribution would you use to find the P-value of the ANOVA F test? Write a brief conclusion based on the sample means and the ANOVA F test.

25.14 Attitudes toward math. Do high school students from different racial/ethnic groups have different attitudes toward mathematics? Measure the level of interest in mathematics on a 5-point scale for a national random sample of students. Here are summaries for students who were taking math at the time of the survey:⁹

Racial/ethnic group	n	\bar{x}	s
African American	809	2.57	1.40
White	1860	2.32	1.36
Asian/Pacific Islander	654	2.63	1.32
Hispanic	883	2.51	1.31
Native American	207	2.51	1.28

- (a) The conditions for ANOVA are clearly satisfied. Explain why.
- (b) Calculate the ANOVA table and the F statistic.
- (c) Software gives $P < 0.001$. What explains the small P-value? Do you think the differences are large enough to be important?

CHAPTER 25 SUMMARY

CHAPTER SPECIFICS

- One-way analysis of variance (ANOVA) compares the means of several populations. The **ANOVA F test** tests the null hypothesis that all the populations have the same mean. If the F test shows significant differences, examine the data to see where the differences lie and whether they are large enough to be important.
- The **conditions for ANOVA** state that we have an **independent SRS** from each population; that each population has a **Normal distribution**; and that all populations have the **same standard deviation**.
- In practice, ANOVA inference is relatively **robust** when the populations are non-Normal, especially when the samples are large. Before doing the F test, check the observations in each sample for outliers or strong skewness. Also verify that the largest sample standard deviation is no more than twice as large as the smallest standard deviation.
- When the null hypothesis is true, the **ANOVA F statistic** for comparing I means from a total of N observations in all samples combined has the **F distribution** with $I - 1$ and $N - I$ degrees of freedom.
- ANOVA calculations are reported in an **ANOVA table** that gives sums of squares, mean squares, and degrees of freedom for variation among groups and for variation within groups. In practice, we use software to do the calculations.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Recognition

1. Recognize when testing the equality of several means is helpful in understanding data.
2. Recognize that the statistical significance of differences among sample means depends on the sizes of the samples and on how much variation there is within the samples.
3. Recognize when you can safely use ANOVA to compare means. Check the data production, the presence of outliers, and the sample standard deviations for the groups you want to compare.

B. Interpreting ANOVA

1. Explain what null hypothesis F tests in a specific setting.
2. Locate the F statistic and its P -value on the output of analysis of variance software.
3. Find the degrees of freedom for the F statistic from the number and sizes of the samples.
4. If the test is significant, use graphs and descriptive statistics to see what differences among the means are most important.

LINK IT

Analysis of variance is a general statistical method for studying sources of variation in a response. In this chapter, we have studied one-way ANOVA, which is a specific statistical technique designed to test the null hypothesis of the equality of the means of several populations. As such, it is an extension of the two-sample *t* test of Chapter 19, which tested the null hypothesis of the equality of two population means.

Rejection of the null hypothesis of equality of two means with the two-sample *t* test means that we have evidence that the two means are different. Rejection of the null hypothesis with a one-way ANOVA is a more ambiguous conclusion: it is evidence of a difference in the means of the populations, but the result does not tell us *which* differences between the means are statistically significant. This is similar to the chi-square test in Chapter 23, where rejection of the null hypothesis indicates that there is a relationship between the two categorical variables but says nothing about the nature of that relationship. Typically, after rejection of the null hypothesis in a one-way ANOVA, we perform a more detailed *follow-up analysis* to decide which of the means are different and to estimate how large these differences are. The companion Chapter 29, on the text CD and Web site, presents some details of this follow-up inference.

CHECK YOUR SKILLS

25.15 The purpose of analysis of variance is to compare

- (a) the variances of several populations.
- (b) the proportions of successes in several populations.
- (c) the means of several populations.

25.16 The *F* distributions are

- (a) a family of distributions with bell-shaped density curves centered at 0.
- (b) a family of distributions that are right-skewed and take only values greater than 0.
- (c) a family of distributions that are left-skewed and take values between 0 and 1.

An experiment to help determine if insects sleep gave caffeine to fruit flies to see if it affected their rest. The three treatments were a control, a low caffeine dose of 1 milligram of caffeine per milliliter of blood (mg/ml) and a higher dose of 5 mg/ml. Nine fruit flies were assigned at random to the three treatments, three to each treatment, and the minutes of rest were measured over a 12-hour period. Here are the minutes of rest for the three groups:



MShields Photos/Alamy

Control	Low dose	High dose
450	466	265
413	420	330
418	435	389

Here is partial Minitab output for the ANOVA table (several numbers have been omitted), along with the means and standard deviations of the rest times for the three groups:

Source	DF	SS	MS	F	P
Caffeine		22598			0.027
Error				1600	
Total					
Level	N	Mean	StDev		
Control	3	427.00	20.07		
Low	3	440.33	23.46		
High	3	328.00	62.02		

Exercises 25.17 to 25.22 are based on this study.  CAFFINE

25.17 The degrees of freedom for the ANOVA *F* statistic comparing mean minutes of rest are

- (a) 2 and 7. (b) 2 and 6. (c) 3 and 7.

25.18 The null hypothesis for the ANOVA F test is

- (a) that the population mean rest time is the same for all three levels of caffeine.
- (b) that the population mean rest time decreases as the caffeine level increases.
- (c) that the population mean rest time is lowest for the high level of caffeine.

25.19 The value of the ANOVA F statistic for testing equality of the population means of the three caffeine levels is

- (a) 4.73.
- (b) 4.82.
- (c) 7.06.

25.20 The conclusion of the ANOVA test is that

- (a) there is strong evidence ($P = 0.027$) that the mean rest time is not the same for all three groups.
- (b) there is strong evidence ($P = 0.027$) that the mean rest time is lower in the high-caffeine group than in the other two.
- (c) the data give no evidence ($P = 0.027$) to suggest that mean rest time differs among the three treatments.

25.21 For this study, we notice that

- (a) ANOVA can be used on these data because ANOVA requires equal sample sizes.
- (b) there is an extreme outlier in the data.
- (c) the data show evidence of a violation of the assumption that the three populations have the same standard deviation.

25.22 To compare the treatments we might use three 90% two-sample t confidence intervals to compare each pair of treatments: the control versus low dose, the control versus high dose, and the low dose versus high dose. The weakness of doing this is that

- (a) we don't know how confident we can be that all three intervals cover the true differences in means.

CHAPTER 25 EXERCISES

Exercises 25.24 to 25.27 describe situations in which we want to compare the mean responses in several populations. For each setting, identify the populations and the response variable. Then give I , the n_i , and N . Finally, state the hypotheses to be tested and give the degrees of freedom of the ANOVA F statistic.

25.24 Morning or evening? Are you a morning person, an evening person, or neither? Does this personality trait affect how well you perform? A sample of 100 students took a psychological test that found 16 morning people, 30 evening people, and 54 who were neither. All the students then took a test of their ability to memorize at 8 A.M. and again at 9 P.M. The response variable is the score at 8 A.M. minus the score at 9 P.M.

(b) 90% confidence is OK for one comparison, but it isn't high enough for three comparisons done at once.

- (c) we can't compare two treatments that use different doses of caffeine.

25.23 A company runs a three-day workshop on strategies for working effectively in teams. On each day, a different strategy is presented. Forty-eight employees of the company attend the workshop. At the outset, all 48 are divided into 12 teams of 4. The teams remain the same for the entire workshop. Strategies are presented in the morning. In the afternoon, the teams are presented with a series of small tasks, and the number of these completed successfully using the strategy taught that morning is recorded for each team. The mean number of tasks completed successfully by all teams each day and the standard deviation follow:

Day	n	\bar{x}	s
1	12	17.25	7.10
2	12	17.64	14.14
3	12	17.21	14.03

In this study, we notice that

- (a) the data show very strong evidence of a violation of the assumption that the three populations have the same standard deviation.
- (b) ANOVA cannot be used on these data, because the sample sizes are less than 20.
- (c) the assumption that the data are independent for the three days is unreasonable because the same teams were observed each day.

25.25 Does art sell products? How does visual art affect the perception and evaluation of consumer products? Subjects were asked to evaluate an advertisement for bathroom fittings that contained an art image, a nonart image, or no image. The art image was Vermeer's painting *Girl with a Pearl Earring*, while the nonart image was a photograph of the actress Scarlett Johansson, in the same pose and wearing the same garments as the girl in the painting, that was taken from the motion picture *Girl with a Pearl Earring*. Thus, the art and nonart image were a match on content. College students were divided at random into three groups of 39 each, with each group assigned to one of the three types of advertisements. Students evaluated the product in the advertisement on a scale of 1 to 7, with 1 being the most unfavorable rating and 7

being the most favorable. The paper reported that a one-way ANOVA on the product evaluation index had $F = 6.29$ with $P < 0.05$.¹⁰

25.26 Test accommodations. Many states require schoolchildren to take regular statewide tests to assess their progress. Children with learning disabilities who read poorly may not do well on mathematics tests because they can't read the problems. Most states allow "accommodations" for learning-disabled children. Randomly assign 100 learning-disabled children in equal numbers to three types of accommodation and a control group: math problems are read by a teacher; by a computer; by a computer that also shows a video; and standard test conditions. Compare the mean scores on the state mathematics assessment.

25.27 Exercise and type 2 diabetes. It is generally accepted that regular exercise provides health benefits to individuals with type 2 diabetes, although it is unclear which exercise regimen (aerobic, resistance, or both) is the best. The subjects in this study were sedentary 30- to 75-year-old adults with type 2 diabetes and elevated hemoglobin A1c levels above 6.5%. The level of hemoglobin A1c correlates very well with a person's recent overall blood sugar levels. If the blood sugars have generally been running high during the previous few months, the level of hemoglobin A1c will be high. In a randomized controlled study, 41 subjects were assigned to a nonexercise control group, 73 to resistance training only, 72 to aerobic exercise only, and 76 to combined aerobic and resistance training. The weekly duration of exercise was similar for all three exercise groups, and subjects remained on the exercise regimens for 9 months. At the end of 9 months, the hemoglobin A1c levels of subjects were measured.¹¹

25.28 Don't handle the merchandise? Although consumers often want to touch products before purchasing them, they generally prefer that others have not touched products they would like to buy. Can another person touching a product create a positive reaction? Subjects were given instructions to contact a sales associate at a university bookstore who would provide them with a shirt to try on. When meeting the sales associate, subjects were told that there was only one shirt left and it was being tried on by another "customer." The other customer trying on the shirt was a confederate of the experimenter and was either an attractive, well-dressed professional female model or an average-looking female college student wearing jeans and a tee shirt. Subjects, who were either males or females, saw the confederate leaving the dressing room, where the shirt was left for them to try on. There was also a control group of subjects who were handed the shirt directly off the rack by the sales associate. Thus, there were five treatments: male subjects seeing a model, female subjects seeing a

model, male subjects seeing a college student, female subjects seeing a college student, and the control group. Subjects evaluated the product on five dimensions, each dimension on a 7-point scale, with the five scores then averaged to give the subject's evaluation measure, with higher numbers indicating a more positive evaluation. Here are the sample sizes, means, and standard deviations for the five groups:¹²

Treatment group	<i>n</i>	\bar{x}	<i>s</i>
Males seeing a model	22	5.34	0.87
Males seeing a student	23	3.32	1.21
Females seeing a model	24	4.10	1.32
Females seeing a student	23	3.50	1.43
Controls	27	4.17	1.50

(a) Verify that the sample standard deviations allow the use of ANOVA to compare the population means. What do the means suggest about the effect of the subject's sex and the attractiveness of the confederate on the evaluation of the product?

(b) The paper reports an ANOVA F statistic of $F = 8.30$. What are the degrees of freedom for the ANOVA F statistic and the P -value? State your conclusions.

25.29 Plants defend themselves. When some plants are attacked by leaf-eating insects, they release chemical compounds that attract other insects that prey on the leaf-eaters. A study carried out on plants growing naturally in the Utah desert demonstrated both the release of the compounds and that they not only repel the leaf-eaters but attract predators that act as the plants' bodyguards.¹³ The investigators chose 8 plants attacked by each of three leaf-eaters and 8 more that were undamaged, 32 plants of the same species in all. They then measured emissions of several compounds during seven hours. Here are data (mean \pm standard error of the mean for eight plants) for one compound. The emission rate is measured in nanograms (ng) per hour.

Group	Emission rate (ng/hr)
Control	9.22 ± 5.93
Hornworm	31.03 ± 8.75
Leaf bug	18.97 ± 6.64
Flea beetle	27.12 ± 8.62

(a) Make a graph that compares the mean emission rates for the four groups. Does it appear that emissions increase when the plant is attacked?

- (b) What hypotheses does ANOVA test in this setting?
 (c) We do not have all the data. What would you look for in deciding whether you can safely use ANOVA?
 (d) What is the relationship between the standard error of the mean (SEM) and the standard deviation for a sample? What are the four sample standard deviations? Do they satisfy our rule of thumb for safe use of ANOVA?

25.30 Can you hear these words?

To test whether a hearing aid is right for a patient, audiologists play a tape on which words are pronounced at low volume. The patient tries to repeat the words. There are several different lists of words that are supposed to be equally difficult. Are the lists equally difficult when there is background noise? To find out, an experimenter had subjects with normal hearing listen to four lists with a noisy background. The response variable was the percent of the 50 words in a list that the subject repeated correctly. The data set contains 96 responses.¹⁴ Here are two study designs that could produce these data:

Design A. The experimenter assigns 96 subjects to 4 groups at random. Each group of 24 subjects listens to one of the lists. All individuals listen and respond separately.

Design B. The experimenter has 24 subjects. Each subject listens to all four lists in random order. All individuals listen and respond separately.



Phanie/Photo Researchers

Does Design A allow use of one-way ANOVA to compare the lists? Does Design B allow use of one-way ANOVA to compare the lists? Briefly explain your answers.

25.31 More rain for California? The changing climate will probably bring more rain to California, but we don't know whether the additional rain will come during the winter wet season or extend into the long dry season in spring and summer. Kenwyn Suttle of the University of California at Berkeley and his coworkers randomly assigned plots of open grassland to three treatments: added water equal to 20% of annual rainfall either during January to March (winter) or during April to June (spring), and no added water (control). Here are some of the data, for plant biomass (in grams per square meter) produced by each plot in a single year:¹⁵



Winter	Spring	Control
264.1514	318.4182	129.0538
187.7312	281.6830	144.6578
291.1431	288.8433	172.7772
176.2879	382.6673	113.2813
141.7525	326.8877	142.1562
169.9737	293.8502	117.9808

Figure 25.10 shows Minitab ANOVA output for these data. (a) Make side-by-side stemplots of plant biomass for the three treatments, as well as a table of the sample means and standard deviations. What do the data appear to show about the effect of extra water in winter and in spring on biomass? Do these data satisfy the conditions for ANOVA?

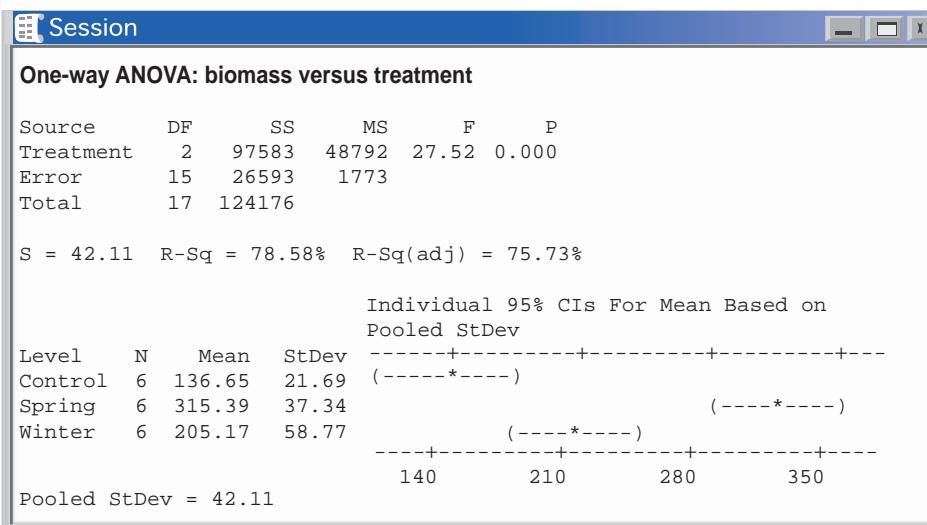
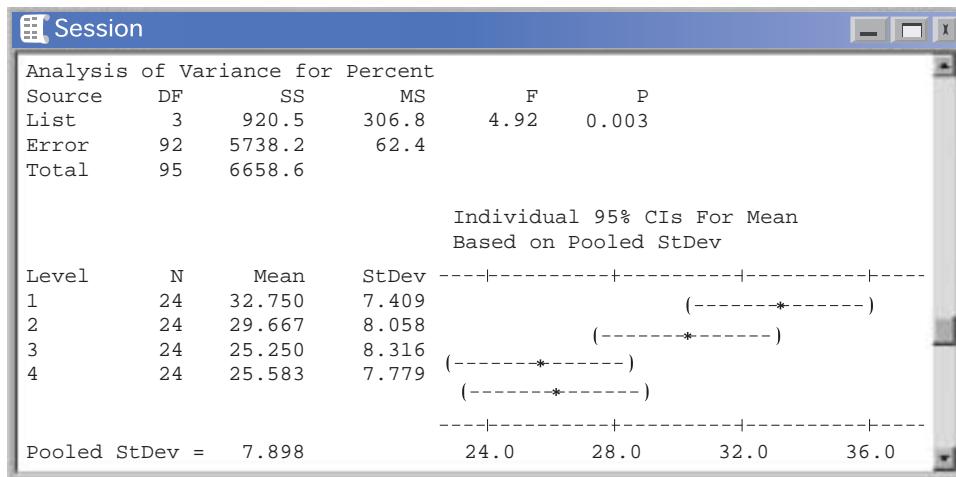


FIGURE 25.10

Minitab ANOVA output for comparing the total plant biomass of grassland plots under different water conditions, for Exercise 25.31.

FIGURE 25.11

Minitab ANOVA output for comparing the percents heard correctly in four lists of words, for Exercise 25.32.



- (b) State H_0 and H_a for the ANOVA F test, and explain in words what ANOVA tests in this setting.
 (c) Report your overall conclusions about the effect of added water on plant growth in California.

25.32 Can you hear these words? Figure 25.11 displays the Minitab output for one-way ANOVA applied to the hearing data described in Exercise 25.30. The response variable is “Percent,” and “List” identifies the four lists of words. Based on this analysis, is there good reason to think that the four lists are not all equally difficult? Write a brief summary of the study findings.

25.33 Which blue is most blue? The color of a fabric depends on the dye used and also on how the dye is applied. This matters to clothing manufacturers, who want the color of the fabric to be just right. A manufacturer dyes fabric made of ramie with the same “procion blue” dye applied in four different ways. She uses a colorimeter to measure the lightness of the color on a scale in which black is 0 and white is 100. Here are the data for 8 pieces of fabric dyed in each way:¹⁶ 

Method A	41.72	41.83	42.05	41.44	41.27	42.27	41.12	41.49
Method B	40.98	40.88	41.30	41.28	41.66	41.50	41.39	41.27
Method C	42.30	42.20	42.65	42.43	42.50	42.28	43.13	42.45
Method D	41.68	41.65	42.30	42.04	42.25	41.99	41.72	41.97

- (a) This is a randomized comparative experiment. Outline the design.
 (b) The clothing manufacturer wants to know which method gives the darkest color. Follow the four-step process in answering this question.

25.34 Do good smells bring good business? Businesses know that customers often respond to background music. Do they also respond to odors? Nicolas Guéguen and his colleagues studied this question in a small pizza restaurant in France on Saturday evenings in May. On one evening, a relaxing lavender odor was spread through the restaurant; on another evening, a stimulating lemon odor; a third evening served as a control, with no odor. The three evenings were comparable in many ways (weather, customer count, and so on), so we are willing to regard the data as independent SRSs from spring Saturday evenings at this restaurant. Table 25.4 contains data on how long (in minutes) customers stayed in the restaurant on each of the three evenings.¹⁷ 

- (a) Make stemplots of the customer times for each evening. Do any of the distributions show outliers, strong skewness, or other clear deviations from Normality?
 (b) Do a complete analysis to see whether the groups differ in the average amount of time spent in the restaurant. Follow the four-step process in your work. Did you find anything surprising?

25.35 Good weather and tipping. Favorable weather has been shown to be associated with increased tipping. Will just the belief that future weather will be favorable lead to higher tips? The researchers gave 60 index cards to a waitress at an Italian restaurant in New Jersey. Before delivering the bill to each customer, the waitress randomly selected a card and wrote on the bill the same message that was printed on the index card. Twenty of the cards had the message “The weather is supposed to be really good tomorrow. I hope you enjoy the day!” Another 20 cards contained the message “The

TABLE 25.4 Time (minutes) that customers remained in a restaurant when exposed to odors

LAVENDER ODOR									
92	126	114	106	89	137	93	76	98	108
124	105	129	103	107	109	94	105	102	108
95	121	109	104	116	88	109	97	101	106
LEMON ODOR									
78	104	74	75	112	88	105	97	101	89
88	73	94	63	83	108	91	88	83	106
108	60	96	94	56	90	113	97		
NO ODOR									
103	68	79	106	72	121	92	84	72	92
85	69	73	87	109	115	91	84	76	96
107	98	92	107	93	118	87	101	75	86

weather is supposed to be not so good tomorrow. I hope you enjoy the day anyway!" The remaining 20 cards were blank, indicating that the waitress was not supposed to write any message. Choosing a card at random ensured that there was a random assignment of the diners to the three experimental conditions. Here are the tips as a percent of the total bill for the three messages:¹⁸



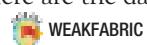
TIPPING

Good weather report	20.8	18.7	19.9	20.6	22.0	23.4	22.8	24.9	22.2	20.3
	24.9	22.3	27.0	20.4	22.2	24.0	21.2	22.1	22.0	22.7
Bad weather report	18.0	19.0	19.2	18.8	18.4	19.0	18.5	16.1	16.8	14.0
	17.0	13.6	17.5	19.9	20.2	18.8	18.0	23.2	18.2	19.4
No weather report	19.9	16.0	15.0	20.1	19.3	19.2	18.0	19.2	21.2	18.8
	18.5	19.3	19.3	19.4	10.8	19.1	19.7	19.8	21.3	20.6

Do the data support the hypothesis that there are differences among the tipping percents for the three experimental conditions? Does a prediction of good weather seem to increase the tip percent? Follow the four-step process in data analysis and ANOVA. Be sure to check the conditions for ANOVA and to include an appropriate graph that compares the tipping percents for the three conditions.

25.36 Durable press fabrics are weaker. "Durable press" cotton fabrics are treated to improve their recovery from wrinkles after washing. Unfortunately, the treatment also reduces the strength of the fabric. A study compared the breaking strength of untreated fabric with that of fabrics treated by three commercial durable press processes.

Five specimens of the same fabric were assigned at random to each group. Here are the data, in pounds of pull needed to tear the fabric:¹⁹



WEAKFABRIC

Untreated	60.1	56.7	61.5	55.1	59.4
Permafresh 55	29.9	30.7	30.0	29.5	27.6
Permafresh 48	24.8	24.6	27.3	28.1	30.3
Hylite LF	28.8	23.9	27.0	22.1	24.2

The untreated fabric is clearly much stronger than any of the treated fabrics. We want to know if there is a significant difference in breaking strength among the three durable press treatments. Analyze the data for the three processes and write a clear summary of your findings. Which process do you recommend if breaking strength is a main concern? Use the four-step process to guide your discussion. (Although the standard deviations do not quite satisfy our rule of thumb, that rule is conservative, and many statisticians would use ANOVA for these data.)

25.37 Durable press fabrics wrinkle less. The data in Exercise 25.36 show that durable press treatment greatly reduces the breaking strength of cotton fabric. Of course, durable press treatment also reduces wrinkling. How much? "Wrinkle recovery angle" measures how well a fabric recovers from wrinkles. Higher is better. Here are data on the wrinkle

recovery angle (in degrees) for the same fabric specimens discussed in the previous exercise:

Untreated	79	80	78	80	78
Permafresh 55	136	135	132	137	134
Permafresh 48	125	131	125	145	145
Hylite LF	143	141	146	141	145

The untreated fabric once again stands out, this time as inferior to the treated fabrics in wrinkle resistance. Examine the data for the three durable press processes and summarize your findings. How does the ranking of the three processes by wrinkle resistance compare with their ranking by breaking strength in Exercise 25.36? Explain why we can't trust the ANOVA F test. 

25.38 Logging in the rain forest: species counts. Table 25.2 gives data on the number of trees per forest plot, the number of species per plot, and species richness. Exercise 25.3 analyzed the effect of logging on number of trees. Exercise 25.8 concludes that it would be risky to use ANOVA to analyze richness. Use software to analyze the effect of logging on the number of species. 

- (a) Make a table of the group means and standard deviations. Do the standard deviations satisfy our rule of thumb for safe use of ANOVA? What do the means suggest about the effect of logging on the number of species?
- (b) Carry out the ANOVA. Report the F statistic and its P-value and state your conclusion.

More rain for California?

Exercise 25.31 describes a randomized experiment carried out by Kenwyn Suttle and his coworkers to examine the effects of additional water on California grassland. The experimental units are 18 plots of grassland, assigned at random among three treatments: added water in the winter wet season, added water in the spring dry season, and no added water (control group). Field experiments, unlike laboratory experiments, are exposed to variations in the natural environment. The experiment therefore continued over five years, from 2001 to 2005. Table 25.5 gives data on the total plant biomass (grams per square meter) that grew on each plot during each



Courtesy Blake Suttle

year.²⁰ The “Plot” column shows how the random assignment of 18 of the 36 available plots worked. Exercises 25.39 to 25.41 are based on this information.

25.39 Plot the means. Starting from the data in Table 25.5, you can calculate the mean plant biomass for each treatment in each year as follows: 

Treatment	Year				
	2001	2002	2003	2004	2005
Winter	132.58	203.33	205.17	223.58	332.84
Spring	257.69	388.85	315.39	299.54	289.66
Control	81.67	180.31	136.65	201.07	257.37

Plot the means for each of the three treatments against year, connecting the yearly means for each treatment by lines to show the pattern over time. Use the same plot for all three treatments, with a different color for each treatment. From this plot, you can get an overall picture of the experiment's results.

- (a) Across all five years, does more water in the wet season increase plant growth? What about more water in the dry season? Which seasonal addition of water has the larger effect?
- (b) One-way ANOVAs comparing the mean plant biomass separately in each year find significant differences in three years and no significant difference in two years. Based on your plot, in which three years do you think the treatment means differ significantly?
- (c) In 2005, there were unusually late rains during the spring. How does the effect of this natural rainfall show up in your plot? (You see that it would not be wise to do an experiment like this in just one year.)

25.40 The results for 2001. Your work in Exercise 25.31 shows that there were significant differences in mean plant biomass among the three treatments in 2003. Do a complete analysis of the data for 2001 and report your conclusions. 

25.41 Conditions for ANOVA. Examine the data for the year 2004. The conditions for ANOVA inference are not met. In what way do these data fail to meet the conditions? (It is not very surprising that in five ANOVAs one will fail to satisfy our quite conservative conditions.) 

25.42 Which test? Example 25.4 describes one of the experiments done by Kathleen Vohs and her coworkers to demonstrate that even being reminded of money makes people more self-sufficient and less involved with other

TABLE 25.5 Plant biomass (grams per square meter) for three water conditions over five years

TREATMENT	PLOT	YEAR				
		2001	2002	2003	2004	2005
Winter	3	136.8358	228.0717	264.1514	254.6453	344.3933
Winter	8	151.4154	189.9505	187.7312	233.8155	203.3908
Winter	14	136.1536	209.0485	291.1431	253.4506	331.9724
Winter	20	121.6323	189.6755	176.2879	228.5882	388.1056
Winter	27	124.1459	188.0090	141.7525	158.6675	382.8617
Winter	32	125.2986	215.2174	169.9737	212.3232	346.3042
Spring	4	338.1301	422.7411	318.4182	517.6650	344.0489
Spring	11	291.8597	339.8243	281.6830	342.2825	261.8016
Spring	18	244.8727	398.7296	288.8433	270.5785	262.7238
Spring	22	234.6599	400.6878	382.6673	212.5324	316.9683
Spring	25	197.5830	326.9497	326.8877	213.9879	224.1109
Spring	35	239.0122	444.1556	293.8502	240.1927	328.2783
Control	6	73.4288	148.8907	129.0538	178.9988	237.6596
Control	7	110.6306	182.6762	144.6578	205.5165	281.1442
Control	17	95.3405	196.8303	172.7772	242.6795	313.7242
Control	24	83.0584	186.1953	113.2813	231.7639	258.3631
Control	28	30.5886	154.0401	142.1562	134.9847	235.8320
Control	33	96.9709	213.2537	117.9808	212.4862	217.5060

people. Here are three more of these experiments. For each experiment, which statistical test from Chapters 18 to 25 would you use, and why?

(a) Randomly assign student subjects to money and control groups. The control group unscrambles neutral phrases, and the money group unscrambles money-oriented phrases, as described in Example 25.4. Then ask the subjects to volunteer to help the experimenter by coding data sheets, which takes about five minutes per sheet. Subjects said how many sheets they would volunteer to code. “Participants in the money condition volunteered to help code fewer data sheets than did participants in the control condition.”

(b) Randomly assign student subjects to high-money, low-money, and control groups. After playing Monopoly for a short time, the high-money group is left with \$4000 in Monopoly money, the low-money group with \$200, and the

control group with no money. Each subject is asked to imagine a future with lots of money (high-money group), a little money (low-money group), or just their future plans (control group). Another student walks in and spills a box of 27 pencils. How many pencils does the subject pick up? “Participants in the high-money condition gathered fewer pencils” than subjects in the other two groups.

(c) Randomly assign student subjects to three groups. All do paperwork while a computer on the desk shows a screensaver of currency floating underwater (Group 1), a screensaver of fish swimming underwater (Group 2), or a blank screen (Group 3). Each subject must now develop an advertisement and can choose whether to work alone or with a partner. Count how many in each group make each choice. “Choosing to perform the task with a coworker was reduced among money condition participants.”



EXPLORING THE WEB

25.43 Confidence in the banking system. The General Social Survey (GSS) is a sociological survey used to collect data on demographic characteristics and attitudes of residents of the United States. The survey is conducted by the National Opinion Research Center of the University of Chicago, which interviews face-to-face a randomly selected sample of adults (18 and older). SDA (Survey Documentation and Analysis) is a set of programs that allows you to analyze survey data and includes the GSS survey as part of its archive. Go to the Web site sda.berkeley.edu/ and click on Archive. Unless there is a more recent file, open the 1972–2010 cumulative data file (without the quick tables option).

- (a) In the “Analysis” tab at the top of the page, click on “Comparison of means”. Do an ANOVA that examines how the mean age of the respondents varies with their confidence in the banking and financial systems. To do this, type in the dependent variable as “Age” and the row (treatment) variable as “Confinan.” For the selection filter, type in “Year (2010)” or the most recent year available. For weight, change it to “noweight.” Finally, in the table options, the *only* boxes that should be checked are “Std dev,” “N,” and “ANOVA stats.” Make sure that the checks are removed from the other boxes. Now click on “Run the table.”
- (b) How many respondents are included in the analysis? What are the three means and standard deviations? Explain how the degrees of freedom were obtained. What are the *F*- and *P*-values? Write a brief report explaining the relationship between the average respondent age and confidence in the banking system.

25.44 Confidence in the banking system, continued. This exercise is a continuation of the previous Web exercise. You will download the data file and reproduce the analysis, as well as provide some additional plots. First, open the 1972–2010 cumulative data file following the instructions in the previous exercise.

- (a) In the “Download” tab at the top of the page, click on “Customized Subset.” For the data file, if you highlight the CSV bubble, an Excel spreadsheet will be downloaded (unclick “Codebook”). For the selection filter, again type in “year (2010)” or the year used in the previous exercise. In the box for entering the names of individual variables, enter “Age” and “Confinan.” Click “continue” at the bottom of the page. In the new window, click on “Create the Files,” and in the next window click on “data files.” You can now either open or save the data file to your computer.
- (b) Import the data into your statistical software package. You first need to “clean” the data a little because there are observations for which either the “Age” or the “Confinan” variable is missing. For the “Confinan” variable, any value other than a 1, 2, or 3 is a missing-value code. Delete these observations. For the “Age” variable, the missing-value codes are 0, 98, and 99. Eliminate any observations with these values for “Age.” You should now have the same number of observations as in the previous exercise.
- (c) Draw comparative boxplots of the age distribution for the three values of “Confinan.” Describe the shapes of the three distributions. What information can you obtain from the boxplots that was not included in the output for the previous exercise?
- (d) Reproduce the one-way ANOVA table using your software. Your results should agree with those of the previous exercise.



Notes and Data Sources

"About This Book" Notes

1. Lawrence D. Brown, Tony Cai, and Anirban DasGupta, "Interval estimation for a binomial proportion," *Statistical Science*, 16 (2001), pp. 101–133.

"Statistical Thinking" Notes

1. Parts of this essay are shared with David S. Moore, "Introduction: learning from data," in Roxy Peck et al. (eds.), *Statistics: A Guide to the Unknown*, 4th ed., Thomson, 2006.
2. See, for example, Martin Enserink, "The vanishing promises of hormone replacement," *Science*, 297 (2002), pp. 325–326; and Brian Vastag, "Hormone replacement therapy falls out of favor with expert committee," *Journal of the American Medical Association*, 287 (2002), pp. 1923–1924. A National Institutes of Health panel's comprehensive report is *International Position Paper on Women's Health and Menopause*, NIH Publication 02-3284, 2002.
3. A. C. Nielsen, Jr., "Statistics in marketing," in *Making Statistics More Effective in Schools of Business*, Graduate School of Business, University of Chicago, 1986.
4. The data in Figure 2 are based on a component of the Consumer Price Index, from the Bureau of Labor Statistics Web site: www.bls.gov. We converted the index number into cents per gallon using retail price information from the Automobile Association of America, www.fuelgaugereport.com.
5. FUTURE II Study Group, "Quadrivalent vaccine against human papillomavirus to prevent high-grade cervical lesions," *New England Journal of Medicine*, 356 (2007), pp. 1915–1927. We have simplified the conclusions so that students with as yet no statistics background can better follow the essay.
4. Higher Education Research Institute 2009 Freshman Survey, at www.heri.ucla.edu.
5. Centers for Disease Control and Prevention, National Center for Health Statistics, *Births: Final data for 2008*, National Vital Statistics Reports, 59, No. 1, December 2010, at www.cdc.gov. These are the most recent data available at the end of 2010, but the numbers change only slightly from year to year.
6. From the 2006 American Community Survey, at factfinder.census.gov.
7. Our eyes do respond to area, but not quite linearly. It appears that we perceive the ratio of two bars to be about the 0.7 power of the ratio of their actual areas. See W. S. Cleveland, *The Elements of Graphing Data*, Wadsworth, 1985, pp. 278–284.
8. From the 2008 population estimates and projections at factfinder.census.gov.
9. From the Gary Community School Corporation, courtesy of Celeste Foster, Purdue University.
10. From the College Board Web site, www.collegeboard.com.
11. The 35 countries with the highest GDP in 2007 were obtained at geohive.com/chartsec_gdp2.aspx/. The health care expenditures per capita in 2007 were obtained from the World Health Organization at who.int/whosis/whostat/en/. All amounts are in International dollars at purchasing power parity. That is, the exchange rate between each currency and the dollar is set not at the fluctuating market rate but at the rate that gives a dollar the same buying power in each country.
12. The U.S. Geological Survey maintains data for various water parameters at monitoring sites throughout the United States at waterdata.usgs.gov/nwis. The data can be graphed or downloaded. The data in Figure 1.12 are for USGS 254754080344300 SHARK RIVER SLOUGH NO. 1.
13. College Entrance Examination Board, *Trends in College Pricing*, 2010, at www.trends.collegeboard.org. The averages are "enrollment weighted," so that they give average tuition over students rather than over colleges. The reported averages have been adjusted to constant 2010 dollars.
14. See Note 6.

Chapter 1 Notes

1. Data for 2008 from the *Statistical Abstract of the United States* at the Census Bureau Web site, www.census.gov.
2. *The Infinite Dial 2010: Digital Platforms and the Future of Radio*, at www.arbitron.com.
3. *Arbitron Internet and Media 2006*, at www.arbitron.com.

15. DuPont 2010 Color Popularity Report, at www2.dupont.com.
16. The 2009 U.S. Digital Year in Review, at www.comscore.com.
17. Centers for Disease Control and Prevention, National Center for Health Statistics, *Deaths: Preliminary Data for 2008*, 59, No. 2, December 2010, at www.cdc.gov/nchs.
18. U.S. Hispanic Population 2006, at www.census.gov, based on the March 2006 Current Population Survey, Annual Social and Economic Supplement.
19. “2008 Student surveys: complete results,” *Macleans.ca*, February 19, 2008, at oncampus.macleans.ca.
20. Tom Lloyd et al., “Fruit consumption, fitness, and cardiovascular health in female adolescents: the Penn State Young Women’s Health Study,” *American Journal of Clinical Nutrition*, 67 (1998), pp. 624–630.
21. Data provided by Darlene Gordon from her PhD thesis, “Relationships among academic self-concept, academic achievement, and persistence with self-attribution, study habits, and perceived school environment,” Purdue University, 1997.
22. Monthly stock returns from the Web site of Professor Kenneth French of Dartmouth, mba.tuck.dartmouth.edu/pages/faculty/ken.french. A fine point: the data are actually the “excess returns” on stocks, the actual returns less the small monthly returns on Treasury bills. The data are in the file Fama/French Benchmark Factors.
23. National Institutes of Health, Essential Fatty Acids Education site, efaeducation.nih.gov.
24. 2010 Statistical Abstract of the United States, Table 159, at www.census.gov.
25. As of the end of 2010, yearly data were available through 2007 at the United Nations Website unstats.un.org/unsd/mdg/SeriesDetail.aspx?srid=751&crid=.
26. National Oceanic and Atmospheric Administration, at www.beringclimate.noaa.gov/.
27. David M. Fergusson and L. John Horwood, “Cannabis use and traffic accidents in a birth cohort of young adults,” *Accident Analysis and Prevention*, 33 (2001), pp. 703–711.
28. From a plot in K. Krishna Kumar et al., “Unraveling the mystery of Indian monsoon failure during El Niño,” *Science*, 314 (2006), pp. 115–119.
29. See Note 13.
30. Census Bureau, New Residential Construction page, at www.census.gov/const/startsua.pdf. These are monthly data that are not seasonally adjusted.
31. Ozone Hole Watch, at ozonewatch.gsfc.nasa.gov/index.html.

Chapter 2 Notes

1. From the 2003 American Community Survey, at the Census Bureau Web site, www.census.gov. The data are a subsample of the 13,194 individuals in the ACS North Carolina sample who had travel times greater than zero.
2. This isn’t a mathematical theorem. The mean can be less than the median in right-skewed distributions that take only a few values, many of which lie exactly at the median. The rule almost never fails for distributions taking many values, and most counterexamples don’t appear clearly skewed in graphs even though they may be slightly skewed according to technical measures of skewness. See Paul T. von Hippel, “Mean, median, and skew: correcting a textbook rule,” *Journal of Statistics Education*, 13, No. 2 (2005), online journal.
3. National Association of College and University Business Officers, 2009 Endowment Study, at www.nacubo.org.
4. From the U.S. Census Bureau, at www.census.gov/const/uspricemon.pdf.
5. U.S. Census Bureau, *Income, Poverty, and Health Insurance Coverage in the United States: 2009*, September 2010, at www.census.gov.
6. The U.S. Department of Energy, at www.fueleconomy.gov/feg/download.shtml.
7. We would like to thank Patricia Humphrey for supplying the test scores for students at Georgia Southern University.
8. From the Environmental Protection Agency, at www.epa.gov/radon/pubs/consguid.html.
9. Ethan J. Temeles and W. John Kress, “Adaptation in a plant-hummingbird association,” *Science*, 300 (2003), pp. 630–633. We thank Ethan J. Temeles for providing the data.
10. C. H. Cannon, D. R. Peart, and M. Leighton, “Tree species diversity in commercially logged Bornean rainforest,” *Science*, 281 (1998), pp. 1366–1367. We thank Charles Cannon for providing the data.
11. Raymond Fisman and Edward Miguel, “Cultures of corruption: evidence from diplomatic parking tickets,” National Bureau of Economic Research Working Paper 12312, June 2006, at www.nber.org.
12. D. G. Jakovljevic and A. K. McConnell, “Influence of different breathing frequencies on the severity of inspiratory muscle fatigue induced by high-intensity front crawl

- swimming," *Journal of Strength and Conditioning Research*, 23, No. 4 (2009), pp. 1169–1174.
13. Patrick J. Purcell, *Retirement Savings and Household Wealth in 2007*, Congressional Research Service, April 2009.
 14. T. Bjerkedal, "Acquisition of resistance in guinea pigs infected with different doses of virulent tubercle bacilli," *American Journal of Hygiene*, 72 (1960), pp. 130–148.
 15. See Note 5 for Chapter 1.
 16. Data for 1986 from David Brillinger, University of California, Berkeley. See David R. Brillinger, "Mapping aggregate birth data," in A. C. Singh and P. Whitridge (eds.), *Analysis of Data in Time*, Statistics Canada, 1990, pp. 77–83. A boxplot similar to Figure 2.6 appears in David R. Brillinger, "Some examples of random process environmental data analysis," in P. K. Sen and C. R. Rao (eds.), *Handbook of Statistics*, Vol. 18, *Bioenvironmental and Public Health Statistics*, North Holland, 2000.
 17. Paul E. O'Brien et al., "Laparoscopic adjustable gastric banding in severely obese adolescents," *Journal of the American Medical Association*, 303 (2010), pp. 519–526. We thank the authors for providing the data.
 18. The current roster as of January 2011 was obtained from canadiens.nhl.com, and their salaries were obtained from forecaster.thehockeynews.com.
 19. Nicolas Guéguen and Christine Petr, "Odors and consumer behavior in a restaurant," *Journal of Hospitality Management*, 25 (2006), pp. 335–339. We thank Nicolas Guéguen for providing the data.
 20. James A. Levine et al., "Inter-individual variation in posture allocation: possible role in human obesity," *Science*, 307 (2005), pp. 584–586. We thank James Levine for providing the data.
 21. Bruce Rind and David Strohmetz, "Effects of beliefs about future weather conditions on restaurant tipping," *Journal of Applied Social Psychology*, 31 (2001), pp. 2160–2164. We thank the authors for supplying the original data.
 22. A sample of responses to the 1901 Census of Canada is available at the Canadian Families Project of the University of Victoria Web site, web.uvic.ca/hrd/cfp/data. The sample and the data are described in Canadian Families Project, *The National Sample of the 1901 Census of Canada*, 2002, on this Web site. Table 2.6 is a random sample of the 47,417 positive incomes in the census sample. The information on bread and beef prices (for 1900) comes from James Powell, *A History of the Canadian Dollar*, Bank of Canada, no date, at www.bank-canada.ca/en/d.

Chapter 3 Notes

1. See Note 9 for Chapter 1.
2. Monsoon rainfall from B. Parthasarathy, Indian Institute of Tropical Meterology, at www.iges.org. The data cover the years 1871 to 2000.
3. Margaret A. McDowell et al., "Anthropometric reference data for children and adults: United States, 2003–2006," *National Health Statistics Reports*, No. 10 (October, 2008), at www.cdc.gov/nchs. This report provides the means of various anthropometric measurements. Standard deviations were computed from the first and third quartiles assuming Normality.
4. All SAT facts are from the College Board Web site, www.collegeboard.com, and all ACT facts are from the ACT Web site, www.act.org.
5. See Note 3.
6. From the 2009–2010 Guide for the College-Bound Student-Athlete, at www.ncaastudent.org/NCAA_Guide.pdf.
7. All MCAT facts are from the Medical College Admissions Test Web site, www.aamc.org/students/mcat/.
8. Detailed data appear in P. S. Levy et al., *Total Serum Cholesterol Values for Youths 12–17 Years*, Vital and Health Statistics, Series 11, No. 155, National Center for Health Statistics, 1976.
9. See Note 20 for Chapter 2.
10. See Note 3.
11. See Note 21 for Chapter 1.
12. The data were provided by Nicolas Fisher.
13. See Note 7 for Chapter 2.
14. See Note 2.

Chapter 4 Notes

1. Neal E. Cantin et al., "Ocean warming slows coral growth in the Central Red Sea," *Science*, 329 (2010), pp. 322–325.
2. Data for 2007 graduates from the College Board Web site, www.collegeboard.com.
3. Initial concerns were based on government data for 2005, presented in "An accident waiting to happen?" *Consumer Reports*, March 2007, pp. 16–19. Data for 2009 were found online at web.mit.edu/airlinedata/www/default.html (go to the Employee Data and Analysis link for each airline) and at the Department of Transportation Web site at airconsumer.dot.gov/reports/2011/February/2011FebruaryATCR.PDF.
4. The Florida Department of Highway Safety and Motor Vehicles (at www.flhsmv.gov/dmv/vslfacts.html)

- gives the number of registered vessels. The Florida Wildlife Commission maintains a manatee death data base at research.myfwc.com/manatees.
5. Based on T. N. Lam, "Estimating fuel consumption from engine size," *Journal of Transportation Engineering*, 111 (1985), pp. 339–357. The data for 10 to 50 km/h are measured; those for 60 and higher are calculated from a model given in the paper and are therefore smoothed.
 6. A careful study of this phenomenon is W. S. Cleveland, P. Diaconis, and R. McGill, "Variables on scatterplots look more highly correlated when the scales are increased," *Science*, 216 (1982), pp. 1138–1141.
 7. Neal E. Cantin et al., "Ocean warming slows coral growth in the Central Red Sea," *Science*, 329 (2010), pp. 322–325.
 8. Data for Figure 4.6(b) come from William Gray's Web site, at hurricane.atmos.colostate.edu. Data for Figure 4.6(c) were provided by Drina Iglesia, Purdue University, from a study reported in D. D. S. Iglesia, E. J. Cragoe, Jr., and J. W. Venable, "Electric field strength and epithelization in the newt (*Notophthalmus viridescens*)," *Journal of Experimental Zoology*, 274 (1996), pp. 56–62. Data for Figure 4.6(d) are for the Wilshire 5000 stock index. A fine point: plots (b), (c), and (d) are square with the same scales on both axes because both variables measure similar quantities in the same units.
 9. This exercise is motivated by Scott Berry, "Statistical fallacies in sports," *Chance*, 19, No. 4 (2006), pp. 50–56, where scores from the 2006 Masters are analyzed.
 10. Andrew J. Oswald et al., "Objective confirmation of subjective measures of human well-being: evidence from the U.S.A.," *Science*, 327 (2010), pp. 576–579.
 11. From a graph in Naomi E. Allen et al., "Moderate alcohol intake and cancer incidence in women," *Journal of the National Cancer Institute*, 101 (2009), pp. 296–305.
 12. From a graph in Magdalena Bermejo et al., "Ebola outbreak killed 5000 gorillas," *Science*, 314 (2006), p. 1564.
 13. From a graph in Bernt-Erik Saether, Steiner Engen, and Erik Mattysen, "Demographic characteristics and population dynamical patterns of solitary birds," *Science*, 295 (2002), pp. 2070–2073.
 14. From a graph in Sabrina M. Tom et al., "The neural basis of loss aversion in decision-making under risk," *Science*, 315 (2007), pp. 515–518.
 15. From a graph in Sergio M. Vallina and Rafel Simó, "Strong relationship between DMS and the solar radiation dose over the global surface ocean," *Science*, 315 (2007), pp. 506–508.
 16. From a graph in Camilla A. Hinde et al., "Parent-offspring conflict and coadaptation," *Science*, 327 (2010), pp. 1373–1376.
 17. Bruce Rind and David Strohmetz, "Effects of beliefs about future weather conditions on restaurant tipping," *Journal of Applied Social Psychology*, 31 (2001), pp. 2160–2164. We would like to thank the authors for supplying the original data.
 18. From a graph in Martin Wild et al., "From dimming to brightening: decadal changes in solar radiation at Earth's surface," *Science*, 308 (2005), pp. 847–850.
 19. Brian J. Whipp and Susan A. Ward, "Will women soon outrun men?" *Nature*, 355 (1992), p. 25.
 20. From a graph in Glenn J. Tattersall et al., "Heat exchange from the toucan bill reveals a controllable vascular thermal radiator" *Science*, 325 (2009), pp. 468–470.
 21. From a graph in Naomi I. Eisenberger, Matthew D. Lieberman, and Kipling D. Williams, "Does rejection hurt? An fMRI study of social exclusion," *Science*, 302 (2003), pp. 290–292.
 22. Justin S. Brashares et al., "Bushmeat hunting, wildlife declines, and fish supply in West Africa," *Science*, 306 (2004), pp. 1180–1183. The data used here are found in the online supplementary material. The published analysis omits data for 1999, an extreme low outlier, without explanation.

Chapter 5 Notes

1. From a graph in James A. Levine, Norman L. Eberhardt, and Michael D. Jensen, "Role of nonexercise activity thermogenesis in resistance to fat gain in humans," *Science*, 283 (1999), pp. 212–214.
2. See Note 1 for Chapter 4.
3. From a graph in Tania Singer et al., "Empathy for pain involves the affective but not sensory components of pain," *Science*, 303 (2004), pp. 1157–1162. Data for other brain regions showed a stronger correlation and no outliers.
4. Contributed by Marigene Arnold, Kalamazoo College.
5. Gannett News Service article appearing in the *Lafayette (Ind.) Journal and Courier*, April 23, 1994.
6. P. Goldblatt (ed.), *Longitudinal Study: Mortality and Social Organisation*, Her Majesty's Stationery Office, 1990. At least, so claims Richard Conniff, *The Natural History of the Rich*, Norton, 2002, p. 45. The Goldblatt report is not available to me.
7. Laura L. Calderon et al., "Risk factors for obesity in Mexican-American girls: dietary factors, anthropometric

- factors, physical activity, and hours of television viewing," *Journal of the American Dietetic Association*, 96 (1996), pp. 1177–1179.
8. *The Health Consequences of Smoking: 1983*, Public Health Service, Washington, D.C., 1983.
 9. Data provided by Robert Dale, Purdue University.
 10. G. L. Kooyman et al., "Diving behavior and energetics during foraging cycles in king penguins," *Ecological Monographs*, 62 (1992), pp. 143–163.
 11. Chu, S., "Diamond ring pricing using simple linear regression," *Journal of Statistics Education*, 4 (1996), available online at www.amstat.org/publications/jse/v4n3/datasets.chu.html.
 12. The last data pair are the heights of one of the authors and his sister. The first 11 data pairs are from Karl Pearson and A. Lee, "On the laws of inheritance in man," *Biometrika*, 2 (1902), p. 357. These first 11 data also appear in D. J. Hand et al., *A Handbook of Small Data Sets*, Chapman & Hall, 1994. This book offers more than 500 data sets that can be used in statistical exercises.
 13. From a presentation by Charles Knauf, Monroe County (N.Y.) Environmental Health Laboratory.
 14. See Note 11 for Chapter 4.
 15. Frank J. Anscombe, "Graphs in statistical analysis," *American Statistician*, 27 (1973), pp. 17–21.
 16. Debora L. Arsenau, "Comparison of diet management instruction for patients with non-insulin-dependent diabetes mellitus: learning activity package vs. group instruction," MS thesis, Purdue University, 1993.
 17. Gary Smith, "Do statistics test scores regress toward the mean?" *Chance*, 10, No. 4 (1997), pp. 42–45.
 18. From a graph in G. D. Martinsen, E. M. Driebe, and T. G. Whitham, "Indirect interactions mediated by changing plant chemistry: beaver browsing benefits beetles," *Ecology*, 79 (1998), pp. 192–200.
 19. P. Velleman, *ActivStats 2.0*, Addison Wesley Interactive, 1997.
 20. From William Gray's Web site, hurricane.atmos.colostate.edu. Forecasts are those made each June.
 21. Data for 1936–1999 are from a graph in Bruce J. Peterson et al., "Increasing river discharge to the Arctic Ocean," *Science*, 298 (2002), pp. 2171–2173. Data for 2000–2008 are from a graph in I. Ashik et al., "Arctic report card: update for 2010," available online at www.arctic.noaa.gov/reportcard/ArcticReportCard_full_report.pdf. The graph is on page 41 of the report.
 22. See Note 18 for Chapter 4.
 23. See Note 11 for Chapter 1.

Chapter 6 Notes

1. The National Longitudinal Study of Adolescent Health interviewed a stratified random sample of 27,000 adolescents, then reinterviewed many of the subjects six years later, when most were age 19 to 25. These data are from the Wave III reinterviews in 2000 and 2001, found at the Web site of the Carolina Population Center, www.cpc.unc.edu.
2. Rani A. Desai et al., "Video-gaming among high school students: health correlates, gender differences, and problematic gaming," *Pediatrics*, 126 (2010), pp. 1416–1424.
3. From the October 2008 Current Population Survey, at www.census.gov.
4. Siem Oppe and Frank De Charro, "The effect of medical care by a helicopter trauma team on the probability of survival and the quality of life of hospitalized victims," *Accident Analysis and Prevention*, 33 (2001), pp. 129–138. The authors give the data in Example 6.4 as a "theoretical example" to illustrate the need for their more elaborate analysis of actual data using severity scores for each victim.
5. Found online at www.math.kent.edu/darci/simpson/bballexamples.html, a Web site maintained by Darci L. Kracht at Kent State University. We thank Patricia Humphrey at Georgia Southern University for bringing this example to our attention.
6. I. Westbrooke, "Simpson's paradox: An example in a New Zealand survey of jury composition," *Chance*, 11 (1998), pp. 40–42.
7. These data are from an April 20, 2010 report, "Teens and mobile phones," by Amanda Lenhart, Rich Ling, Scott Campbell, Kristen Purcell of the Pew Internet and American Life Project. Found online at pewinternet.org/Reports/2010/Teens-and-Mobile-Phones.aspx.
8. This General Social Survey exercise presents a table constructed using the search function at the GSS archive, sda.berkeley.edu/archive.htm. These data are from the 2008 GSS.
9. Gregory D. Myer et al., "Youth versus adult weightlifting injuries presenting to United States emergency rooms: accidental versus nonaccidental injury mechanisms," *Journal of Strength and Conditioning Research*, 23 (2009), pp. 2054–2060.
10. Sanders Korenman and David Neumark, "Does marriage really make men more productive?" *Journal of Human Resources*, 26 (1991), pp. 282–307.
11. M. Radelet, "Racial characteristics and imposition of the death penalty," *American Sociological Review*, 46 (1981), pp. 918–927.

12. D. Gonzales et al., "Varenicline, an $\alpha 4\beta 2$ nicotinic acetylcholine receptor partial agonist, vs sustained-release bupropion and placebo for smoking cessation," *Journal of the American Medical Association*, 296 (2006), pp. 47–55.
13. Michael Gurian, "Where have the men gone? No place good," *Washington Post*, December 4, 2005, at www.washingtonpost.com. The data are from the 2009 *Digest of Education Statistics* at the Web site of the National Center for Education Statistics, nces.ed.gov.
14. Nancy J. O. Birkmeyer, "Hospital complication rates with bariatric surgery in Michigan," *Journal of the American Medical Association*, 304 (2010), pp. 435–442.
15. The data for the University of Michigan Health and Retirement Study (HRS) can be downloaded from the Web site ssl.isr.umich.edu/hrs/start.php.
16. R. Shine, T. R. L. Madsen, M. J. Elphick, and P. S. Harlow, "The influence of nest temperatures and maternal brooding on hatchling phenotypes in water pythons," *Ecology*, 78 (1997), pp. 1713–1721.
9. From a plot in Jon J. Ramsey et al., "Energy expenditure, body composition, and glucose metabolism in lean and obese rhesus monkeys treated with ephedrine and caffeine," *American Journal of Clinical Nutrition*, 68 (1998), pp. 42–51.
10. Centers for Disease Control and Prevention, *Cigarette Smoking among Adults—United States, 2006*, and related publications at www.cdc.gov.
11. Janice E. Williams et al., "Anger proneness predicts coronary heart disease risk," *Circulation*, 101 (2000), pp. 2034–2039.
12. From a graph in Peter A. Raymond and Jonathan J. Cole, "Increase in the export of alkalinity from North America's largest river," *Science*, 301 (2003), pp. 88–91.
13. From the Nenana Ice Classic Web site, www.nenanaakiceclassic.com. See Raphael Sagarin and Fiorenza Micheli, "Climate change in nontraditional data sets," *Science*, 294 (2001), p. 811, for a careful discussion.
14. Data for 2004 from Alan Heston, Robert Summers, and Bettina Aten, *Penn World Table Version 6.2*, Center for International Comparisons of Production, Income, and Prices at the University of Pennsylvania, September 2006, at pwt.econ.upenn.edu.
15. Louie H. Yang, "Periodical cicadas as resource pulses in North American forests," *Science*, 306 (2004), pp. 1565–1567. The data are simulated Normal values that match the means and standard deviations reported in this article.
16. Alan S. Banks et al., "Juvenile hallux abducto valgus association with metatarsus adductus," *Journal of the American Podiatric Medical Association*, 84 (1994), pp. 219–224.
17. Todd W. Anderson, "Predator responses, prey refuges, and density-dependent mortality of a marine fish," *Ecology*, 81 (2001), pp. 245–257.
18. From a graph in Craig Packer et al., "Ecological change, group territoriality, and population dynamics in Serengeti lions," *Science*, 307 (2005), pp. 390–393.
19. Peter H. Chen, Neftali Herrera, and Darren Christiansen, "Relationships between gate velocity and casting features among aluminum round castings," no date. Provided by Darren Christiansen.
20. Data compiled from a table of percents in "Americans view higher education as key to the American dream," press release from the National Center for Public Policy and Higher Education, at www.highereducation.org, May 3, 2000.

Chapter 7 Notes

1. Data for Cohort 2 in Richard A. Morgan et al., "Cancer regression in patients after transfer of genetically engineered lymphocytes," *Science*, 314 (2006), pp. 126–129. The doubling time data are given in the paper and the immune response data appear in the supplementary online material.
2. DuPont 2007 Color Popularity Report, at www2.dupont.com.
3. Data provided by Brigitte Baldi, University of California at Irvine.
4. From a graph in Stan Boutin et al., "Anticipatory reproduction and population growth in seed predators," *Science*, 314 (2006), pp. 1928–1930.
5. J. T. Dwyer et al., "Memory of food intake in the distant past," *American Journal of Epidemiology*, 130 (1989), pp. 1033–1046.
6. Data from a plot in Josef P. Rauschecker, Biao Tian, and Marc Hauser, "Processing of complex sounds in the macaque nonprimary auditory cortex," *Science*, 268 (1995), pp. 111–114. The paper states that there are $n = 41$ observations, but only $n = 37$ can be read accurately from the plot.
7. Mei-Hui Chen, "An exploratory comparison of American and Asian consumers' catalog patronage behavior," MS thesis, Purdue University, 1994.
8. "Dancing in step," *Economist*, March 22, 2001.
1. Data for Cohort 2 in Richard A. Morgan et al., "Cancer regression in patients after transfer of genetically engineered lymphocytes," *Science*, 314 (2006), pp. 126–129. The doubling time data are given in the paper and the immune response data appear in the supplementary online material.
2. DuPont 2007 Color Popularity Report, at www2.dupont.com.
3. Data provided by Brigitte Baldi, University of California at Irvine.
4. From a graph in Stan Boutin et al., "Anticipatory reproduction and population growth in seed predators," *Science*, 314 (2006), pp. 1928–1930.
5. J. T. Dwyer et al., "Memory of food intake in the distant past," *American Journal of Epidemiology*, 130 (1989), pp. 1033–1046.
6. Data from a plot in Josef P. Rauschecker, Biao Tian, and Marc Hauser, "Processing of complex sounds in the macaque nonprimary auditory cortex," *Science*, 268 (1995), pp. 111–114. The paper states that there are $n = 41$ observations, but only $n = 37$ can be read accurately from the plot.
7. Mei-Hui Chen, "An exploratory comparison of American and Asian consumers' catalog patronage behavior," MS thesis, Purdue University, 1994.
8. "Dancing in step," *Economist*, March 22, 2001.

Chapter 8 Notes

1. From the New York Times/CBS News poll at www.nytimes.com. The methodological statement is similar for most polls listed.
2. Gary S. Foster and Craig M. Eckert, "Up from the grave: a sociohistorical reconstruction of an African American community from cemetery data in the rural Midwest," *Journal of Black Studies*, 33 (2003), pp. 468–489.
3. Pew Forum on Religion and Public Life, *Spirit and Power: A 10-Country Survey of Pentecostals*, October 2006, at www.pewforum.org.
4. The regulations that govern seat belt survey design can be found at www-nrd.nhtsa.dot.gov. Details on the Hawaii survey are in Karl Kim et al., *Results of the 2002 Highway Seat Belt Use Survey*, at www.state.hi.us/dot.
5. Donald L. McCabe, Linda Klebe Trevino, and Kenneth D. Butterfield, "Dishonesty in academic environments," *Journal of Higher Education*, 72 (2001), pp. 29–45.
6. For information for the 2009 American Community Survey of households (there is a separate sample of group quarters), go to www.census.gov/acs.
7. The Pew press release and the full report, "Polls face growing resistance, but still representative" (April 20, 2004), are at people-press.org/reports.
8. For more detail on the limits of memory in surveys, see N. M. Bradburn, L. J. Rips, and S. K. Shevell, "Answering autobiographical questions: the impact of memory and inference on surveys," *Science*, 236 (1987), pp. 157–161.
9. The immigration questions are from the New York Times/CBS News Poll taken May 18 to 23, 2007, found at www.pollingreport.com. The responses on welfare are from a New York Times/CBS News Poll reported in the *New York Times*, July 5, 1992. Many other examples appear in T. W. Smith, "That which we call welfare by any other name would smell sweeter," *Public Opinion Quarterly*, 51 (1987), pp. 75–83. The example on the effect of question order is cited in Daniel Kahnemann et al., "Would you be happier if you were richer? A focusing illusion," *Science*, 312 (2006), pp. 1908–1910.
10. Giuliana Coccia, "An overview of non-response in Italian telephone surveys," *Proceedings of the 99th Session of the International Statistical Institute*, 1993, Book 3, pp. 271–272.
11. Go to www.pollingreport.com to see the results of many polling agencies compiled on a variety of issues.
12. Information from various articles in the special issue on cell phone surveys, *Public Opinion Quarterly*, 71, No. 5 (2007). See also the Pew study cited in Note 7 for a comparison of a standard RDD survey with a rigorous survey that reduced nonresponse from 73% to 49%. The 2009 cell phone use numbers were obtained from the CDC Web site, www.cdc.gov.
13. See Mick P. Couper, "Web surveys: A review of issues and approaches," *Public Opinion Quarterly*, 64 (2000), pp. 464–494.
14. Rachel Sherman and John Hickner, "Academic physicians use placebos in clinical practice and believe in the mind-body connection," *Journal of General Internal Medicine*, 23 (2008), pp. 7–10.
15. From the Web site of the Gallup Organization, www.gallup.com. Individual poll reports remain on this site for only a limited time.
16. Information about area codes can be found on the North American Numbering Plan Administration Web site, www.nanpa.com.
17. Robert C. Parker and Patrick A. Glass, "Preliminary results of double-sample forest inventory of pine and mixed stands with high- and low-density LiDAR," in Kristina F. Connoe (ed.), *Proceedings of the 12th Biennial Southern Silvicultural Research Conference*, U.S. Department of Agriculture, Forest Service, Southern Research Station, 2004. The researchers actually sampled every 10th plot. This is a systematic sample; see Exercise 8.43.
18. Bryan E. Porter and Thomas D. Berry, "A nationwide survey of self-reported red light running: measuring prevalence, predictors, and perceived consequences," *Accident Analysis and Prevention*, 33 (2001), pp. 735–741.
19. Mario A. Parada et al., "The validity of self-reported seatbelt use: Hispanic and non-Hispanic drivers in El Paso," *Accident Analysis and Prevention*, 33 (2001), pp. 139–143.
20. Information about the Ontario College of Pharmacists was obtained from its Web site at www.ocpinfo.com.
21. Information about the Health Care in Canada Survey can be found at its Web site, www.hcic-sssc.ca.
22. Clyde O. McDaniel, Jr., "Dating roles and reasons for dating," *Journal of Marriage and the Family*, 31 (1969), pp. 97–107.
23. Lydia Saad, "Gallup Poll: Many Americans say Gulf beaches, wildlife will never recover," at www.gallup.com/poll/140762/Americans-Say-Gulf-Beaches-Wildlife-Recover.aspx.

24. The article can be found at www2.macleans.ca/2010/07/16/sometimes-a-gaffe-is-more-than-a-gaffe.

Chapter 9 Notes

1. I. J. Goldberg et al., "Wine and your heart: a science advisory for healthcare professionals from the Nutrition Committee, Council on Epidemiology and Prevention, and Council on Cardiovascular Nursing of the American Heart Association," *Circulation*, 103 (2001), pp. 472–475.
2. J. E. Muscat et al., "Handheld cellular telephone use and risk of brain cancer," *Journal of the American Medical Association*, 284 (2000), pp. 3001–3007.
3. Hyunjin Song and Norbert Schwarz, "If it's hard to read, it's hard to do: processing fluency affects effort prediction and motivation," *Psychological Science*, 19 (2008), pp. 986–988.
4. Hsin-Chieh Yeh et al., "Smoking, smoking cessation, and risk for type 2 diabetes mellitus: a cohort study," *Annals of Internal Medicine*, 152 (2010), pp. 10–17.
5. Charles A. Nelson III et al., "Cognitive recovery in socially deprived young children: the Bucharest Early Intervention Project," *Science*, 318 (2007), pp. 1937–1940.
6. The description of the factors and the response is based on a portion of the study by Alice Healy et al., "Terrorism after 9/11: reactions to simulated news reports," *American Journal of Psychology*, 122 (2009), pp. 153–165.
7. See Note 17 for Chapter 2.
8. K. B. Suttle, Meredith A. Thomsen, and Mary E. Power, "Species interactions reverse grassland responses to changing climate," *Science*, 315 (2007), pp. 640–642. See Chapter 25 for an analysis of some data from this experiment.
9. Julie Mares et al., "Healthy diets and the subsequent prevalence of nuclear cataract in women," *Archives of Ophthalmology*, 128 (2010), pp. 738–749.
10. Marielle H. Emmelot-Vonk et al., "Effect of testosterone supplementation on functional mobility, cognition, and other parameters in older men," *Journal of the American Medical Association*, 299 (2008), pp. 39–52.
11. David L. Strayer, Frank A. Drews, and William A. Johnston, "Cell phone-induced failures of visual attention during simulated driving," *Journal of Experimental Psychology: Applied*, 9 (2003), pp. 23–32.
12. See Note 12 for Chapter 2.
13. Sterling C. Hilton et al., "A randomized controlled experiment to assess technological innovations in the classroom on student outcomes: an overview of a clinical trial in education," manuscript, no date. A brief report is Sterling C. Hilton and Howard B. Christensen, "Evaluating the impact of multimedia lectures on student learning and attitudes," *Proceedings of the 6th International Conference on the Teaching of Statistics*, at www.stat.auckland.ac.nz.
14. Brad J. Bushman, "Violence and sex in television programs do not sell products in advertisements," *Psychological Science*, 16 (2005), pp. 702–707.
15. K. J. Mukamal et al., "Prior alcohol consumption and mortality following acute myocardial infarction," *Journal of the American Medical Association*, 285 (2001), pp. 1965–1970.
16. Rita F. Redburg, "Vitamin E and cardiovascular health," *Journal of the American Medical Association*, 294 (2005), pp. 107–109.
17. Jo Phelan et al., "The stigma of homelessness: the impact of the label 'homeless' on attitudes towards poor persons," *Social Psychology Quarterly*, 60 (1997), pp. 323–337.
18. Esther Duflo, Rema Hanna, and Stephan Ryan, "Monitoring works: getting teachers to come to school," report dated November 21, 2007, at econ-mit.edu/files/2066.
19. John H. Kagel, Raymond C. Battalio, and C. G. Miles, "Marijuana and work performance: results from an experiment," *Journal of Human Resources*, 15 (1980), pp. 373–395.
20. Shailja V. Nigdikar et al., "Consumption of red wine polyphenols reduces the susceptibility of low-density lipoproteins to oxidation in vivo," *American Journal of Clinical Nutrition*, 68 (1998), pp. 258–265. (There were in fact only 30 subjects, some of whom received more than one treatment with a four-week period intervening.)
21. The description of the factors and the response is based on a portion of the study by Brian Wnasik and Perre Chandon, "Can 'low fat' nutrition labels lead to obesity?" *Journal of Marketing Research*, 43 (2006), pp. 605–617.
22. Ian G. Williamson et al., "Antibiotics and topical nasal steroid for treatment of acute maxillary sinusitis," *Journal of the American Medical Association*, 298 (2007), pp. 2487–2496.
23. Based on Evan H. DeLucia et al., "Net primary production of a forest ecosystem with experimental CO₂ enhancement," *Science*, 284 (1999), pp. 1177–1179. The investigators used the block design.
24. E. M. Peters et al., "Vitamin C supplementation reduces the incidence of postrace symptoms of upper-respiratory tract infection in ultramarathon runners,"

- American Journal of Clinical Nutrition*, 57 (1993), pp. 170–174.
25. The study is described in Gina Kolata, “New study finds vitamins are not cancer preventers,” *New York Times*, July 21, 1994. *Journal of the American Medical Association* of the same date reports the details.
 26. R. C. Shelton et al., “Effectiveness of St. John’s wort in major depression,” *Journal of the American Medical Association*, 285 (2001), pp. 1978–1986.

Data Ethics Notes

1. John C. Bailar III, “The real threats to the integrity of science,” *Chronicle of Higher Education*, April 21, 1995, pp. B1–B2.
 2. See the details on the Web site of the Office for Human Research Protections of the Department of Health and Human Services, www.hhs.gov/ohrp.
 3. The difficulties of interpreting guidelines for informed consent and for the work of institutional review boards in medical research are a main theme of Beverly Woodward, “Challenges to human subject protections in U.S. medical research,” *Journal of the American Medical Association*, 282 (1999), pp. 1947–1952. The references in this paper point to other discussions. Updated regulations and guidelines appear on the OHRP Web site (see Note 2).
 4. Quotation from the *Report of the Tuskegee Syphilis Study Legacy Committee*, May 20, 1996. A detailed history is James H. Jones, *Bad Blood: The Tuskegee Syphilis Experiment*, Free Press, 1993.
 5. Dr. Hennekens’s words are from an interview in the Annenberg/Corporation for Public Broadcasting video series *Against All Odds: Inside Statistics*. The lack of certainty that Dr. Hennekens refers to is now called “clinical equipoise” in discussions of ethics.
 6. R. D. Middlemist, E. S. Knowles, and C. F. Matter, “Personal space invasions in the lavatory: suggestive evidence for arousal,” *Journal of Personality and Social Psychology*, 33 (1976), pp. 541–546.
 7. For a review of domestic violence experiments, see C. D. Maxwell et al., *The Effects of Arrest on Intimate Partner Violence: New Evidence from the Spouse Assault Replication Program*, U.S. Department of Justice, NCH188199, 2001. Available online at www.ojp.usdoj.gov/nij/pubs-sum/188199.htm.
 8. Joseph Millum and Ezekial J. Emanuel, “The ethics of international research with abandoned children,” *Science*, 318 (2007), pp. 1874–1875. This paper has some useful comments on international research in general.
1. The Gallup Poll is based on telephone interviews. Each adult interviewed by Gallup had a known chance of being among those selected, but this chance depended on characteristics such as gender, age, type of phone (cell or landline), and geographic location. Gallup used special weights to adjust for differences in the probability of being selected to obtain an estimate of the proportion of all adults who bought a lottery ticket in the population of all U.S. adults. The actual estimate used by Gallup was close to 46% and uses the result from the sample to estimate what is true for the population.
 2. Note that pennies have rims that make spinning more stable. The probability of a head in spinning a coin depends on the type of coin and also on the surface. See Exercise 21.3 for an account of 56% of heads in spinning a Belgian 1-euro coin. *Chance News* 11.02 at www.dartmouth.edu/~chance reports about 45% heads in more than 20,000 spins of American pennies by Robin Lock’s students at Saint Lawrence University.
 3. The percentages were found at the GMAT Web site, www.gmac.com/gmac/ResearchandTrends/GMATStats/ProfileofCandidates.htm
 4. Data for 2006 from the Web site of Statistics Canada, www.statcan.gc.ca.
 5. A mathematical explanation of Benford’s law is in Ted Hill, “The first-digit phenomenon,” *American Scientist*, 86 (1996), pp. 358–363; and Ted Hill, “The difficulty of faking data,” *Chance*, 12, No. 3 (1999), pp. 27–31. Applications in fraud detection are discussed in the second paper by Hill and in Mark A. Nigrini, “I’ve got your number,” *Journal of Accountancy*, May 1999, available online at www.aicpa.org/pubs/jofa/joaiss.htm.
 6. Based on a November 2007 Gallup Poll, at www.gallup.com/poll/1648/Personal-Health-Issues.2.
 7. Information from www.indiana.edu/~registra/gradedist/.
 8. Thomas K. Cureton et al., *Endurance of Young Men*, Monographs of the Society for Research in Child Development, Vol. 10, No. 1, 1945.
 9. Based on a January 2007 Gallup Poll, at www.gallup.com/poll/15370/Party-Affiliation.aspx.
 10. See Note 15 for Chapter 1.
 11. National population estimates for July 1, 2008, at the U.S. Census Bureau Web site, www.census.gov. The table omits people who consider themselves as belonging to more than one race.

Chapter 10 Notes

12. Based on data from the 2010 *Statistical Abstract of the United States*, Table 58, at www.census.gov.

Chapter 11 Notes

1. U.S. Census Bureau, *Income, Poverty, and Health Insurance in the United States: 2009*, Current Population Reports P60-238. Available online at www.census.gov/prod/2010pubs/p60-238.pdf.
2. Strictly speaking, the formula σ/\sqrt{n} for the standard deviation of \bar{x} assumes that we draw an SRS of size n from an *infinite* population. If the population has finite size N , this standard deviation is multiplied by $\sqrt{1 - (n-1)/(N-1)}$. This “finite population correction” approaches 1 as N increases. When the population is at least 20 times as large as the sample, the correction factor is between about 0.97 and 1. It is reasonable to use the simpler form σ/\sqrt{n} in these settings.
3. Earnings for all 97,263 households were downloaded using the Census Bureau’s Data Ferret software. The histograms in Figure 11.4 were produced from the downloaded data.
4. See Note 9 for Chapter 3.
5. Found online at pages.stern.nyu.edu/adamodar/New_Home_Page/datafile/histret.html. Sophisticates will note that for compounding over several years we want the geometric mean return, which was 9.38%.
7. Probabilities from trials with 2897 people known to be free of HIV antibodies and 673 people known to be infected, reported in J. Richard George, “Alternative specimen sources: methods for confirming positives,” 1998 Conference on the Laboratory Science of HIV, found online at the Centers for Disease Control and Prevention Web site, www.cdc.gov.
8. From the statistics page of the National Science Foundation Web site, www.nsf.gov/statistics.
9. See Note 4 for Chapter 1.
10. From the Internal Revenue Service Web site, at www.irs.gov/taxstats.
11. Projections from U.S. Department of Education, *Projections of Education Statistics to 2016*, December 2007, at nces.ed.gov.
12. Data provided by Patricia Heithaus and the Department of Biology at Kenyon College.
13. F. J. G. M. Klaassen and J. R. Magnus, “How to reduce the service dominance in tennis? Empirical results from four years at Wimbledon,” in S. J. Haake and A. O. Coe (eds.), *Tennis Science and Technology*, Blackwell, 2000, pp. 277–284.
14. Amanda Lenhart et al., “Teens and mobile phones,” April 20, 2010, Pew Internet and American Life Project, at www.pewinternet.org.
15. From the National Institutes of Health’s National Digestive Diseases Information Clearinghouse, found at wrongdiagnosis.com.
16. The probabilities given are realistic, according to the fundraising firm SCM Associates, scmassoc.com.
17. B. Budowle et al. “Population data on the thirteen CODIS core short tandem repeat loci in African Americans, U.S. Caucasians, Hispanics, Bahamians, Jamaicans, and Trinidadians,” *Journal of Forensic Sciences*, 1999, pp. 1277–1286.
18. Marilyn vos Savant, “Ask Marilyn,” Parade Magazine, p. 16, September 9, 1990.

Chapter 12 Notes

1. This is one of several tests discussed in Bernard M. Branson, “Rapid HIV testing: 2005 update,” a presentation by the Centers for Disease Control and Prevention at www.cdc.gov. The Malawi clinic result is reported by Bernard M. Branson, “Point-of-care rapid tests for HIV antibody,” *Journal of Laboratory Medicine*, 27 (2003), pp. 288–295.
2. Robert P. Dellavalle et al., “Going, going, gone: lost Internet references,” *Science*, 302 (2003), pp. 787–788.
3. From the United States Department of Commerce Bureau of Economic Analysis at www.bea.gov, March 2011. Motor vehicle sales information is included with the National Economic Accounts.
4. Sales in 2006 from the Web site of the Entertainment Software Association, at www.theesa.com.
5. Information about Internet users comes from sample surveys carried out by the Pew Internet and American Life Project, at www.pewinternet.org.
6. S. H. Sicherer et al., “Prevalence of peanut and tree nut allergy in the US determined by random digit dial telephone survey,” *Journal of Allergy and Clinical Immunology*, 103 (1999), pp. 559–562.

Chapter 13 Notes

1. Matthew A. Carlton and William D. Stansfield, “Making babies by the flip of a coin?” *American Statistician*, 59 (2005), pp. 180–182.
2. From the Canadian Internet Use Survey, at www.statcan.gc.ca/daily-quotidien/100510/dq100510a-eng.htm.
3. The survey question is reported in Trish Hall, “Shop? Many say ‘Only if I must,’” *New York Times*, November 28, 1990. In fact, 66% (1650 of 2500) in the sample said “Agree.”

4. Information obtained from the Planned Parenthood Web site, www.plannedparenthood.org.
5. Results for the full 2009 season, at www.pgatour.com.
6. From the General Motors Web site at www.gm.com/news-article.jsp?brand=gm&id=/content/Pages/news/us/en/2011/Jan/0117_chev_global.html.
7. C. E. Finley et al., "Retention rates and weight loss in a commercial weight loss program," *International Journal of Obesity*, 31 (2006), pp. 292–298.
8. Associated Press news item dated December 9, 2007, found at www.msnbc.msn.com.
9. See demonstrations.wolfram.com/MonteCarloEstimateForPi/ for an online demonstration of this idea.
5. Gerardo Ramirez and Sian L. Bellock, "Writing about testing worries boosts exam performance in the classroom," *Science*, 331 (2011), pp. 211–213.
6. Kenneth A. Follett et al., "Pallidal versus subthalamic deep-brain stimulation for Parkinson's disease," *New England Journal of Medicine*, 362, No. 22 (2010), pp. 2077–2091.
7. Mario A. Parada et al., "The validity of self-reported seatbelt use: Hispanic and non-Hispanic drivers in El Paso," *Accident Analysis and Prevention*, 33 (2001), pp. 139–143.
8. See Note 6 for Chapter 14.
9. Data simulated from a Normal distribution based on information in Brian M. DeBroff and Patricia J. Pahk, "The ability of periorbitally applied antiglare products to improve contrast sensitivity in conditions of sunlight exposure," *Archives of Ophthalmology*, 121 (2003), pp. 997–1001.

Chapter 14 Notes

1. Margaret A. McDowell et al., "Anthropometric reference data for children and adults: U.S. population, 1999–2002," National Center for Health Statistics, Advance Data from Vital and Health Statistics, No. 361, 2005, at www.cdc.gov/nchs.
2. Information about the NAEP test can be found online at nationsreportcard.gov/math_2009/.
3. B. Rind and D. Strohmetz, "Effect of beliefs about future weather conditions on restaurant tipping," *Journal of Applied Social Psychology*, 31 (2001), pp. 2160–2164.
4. See Note 20 for Chapter 1.
5. Chi-Fu Jeffrey Yang, Peter Gray, Harrison G. Pope, Jr., "Male body image in Taiwan versus the West," *American Journal of Psychiatry*, 162 (2005), pp. 263–269.
6. M. Ann Laskey et al., "Bone changes after three months of lactation: influence of calcium intake, breast-milk output, and vitamin D-receptor genotype," *American Journal of Clinical Nutrition*, 67 (1998), pp. 685–692.

Chapter 15 Notes

1. Steven R. Smith et al., "Multicenter, placebo-controlled trial of lorcaserin for weight management," *New England Journal of Medicine*, 363, No. 3 (2010), pp. 245–256.
2. See Note 3 for Chapter 14.
3. Ajay Ghei, "An empirical analysis of psychological androgeny in the personality profile of the successful hotel manager," MS thesis, Purdue University, 1992.
4. Seung-Ok Kim, "Burials, pigs, and political prestige in Neolithic China," *Current Anthropology*, 35 (1994), pp. 119–141.

Chapter 16 Notes

1. See www.cdc.gov/nchs/tutorials/NHANES/Survey-Design/intro_iii.htm.
2. See Note 18 for Chapter 8.
3. From the Gallup Web site, www.gallup.com. The poll was taken in July 2008.
4. For a discussion of statistical significance in the legal setting, see D. H. Kaye, "Is proof of statistical significance relevant?" *Washington Law Review*, 61 (1986), pp. 1333–1365. Kaye argues: "Presenting the *P*-value without characterizing the evidence by a significance test is a step in the right direction. Interval estimation, in turn, is an improvement over *P*-values."
5. From a press release from the Harvard School of Public Health College Alcohol Study, April 12, 2001, at [www.hsph.harvard.edu/cas/](http://hsph.harvard.edu/cas/).
6. Warren E. Leary, "Cell phones: questions but no answers," *New York Times*, October 26, 1999.
7. Poll published August 26, 2010, at www.harrisinteractive.com/NewsRoom/HarrisPolls/tabid/447/mid/articleId/555/ct1/ReadCustom%20Default/Default.aspx. A note at the bottom of the page states: "Because the sample is based on those who agreed to participate in the Harris Interactive panel, no estimates of theoretical sampling error can be calculated."
8. Justin S. Brashares et al., "Bushmeat hunting, wildlife declines, and fish supply in West Africa," *Science*, 306 (2004), pp. 1180–1183. The data used here (and in Figure 1B of the article) are found in the online supplementary material.

9. Gabriel Gregoratos et al., "ACC/AHA guidelines for implantation of cardiac pacemakers and antiarrhythmia devices: executive summary," *Circulation*, 97 (1998), pp. 1325–1335.
10. From the commentary by Frank J. Sulloway, "Birth order and intelligence," *Science*, 316 (2007), pp. 1711–1712. The study report appears in the same issue, Petter Kristensen and Tor Bjerkedal, "Explaining the relation between birth order and intelligence," *Science*, 316 (2007), p. 1717.
11. C. Kopp et al., "Modulation of rhythmic brain activity by diazepam: GABA receptor subtype and state specificity," *Proceedings of the National Academy of Sciences*, 101 (2004), pp. 3674–3679.
12. Bruce A. Cooper et al., "A randomized, controlled trial of early versus late initiation of dialysis," *New England Journal of Medicine*, 363, No. 7 (2010), pp. 609–619.
13. Simplified from Sanjay K. Dhar, Claudia González-Vallejo, and Dilip Soman, "Modeling the effects of advertised price claims: tensile versus precise pricing," *Marketing Science*, 18 (1999), pp. 154–177.
14. Charles S. Fuchs et al., "Alcohol consumption and mortality among women," *New England Journal of Medicine*, 332 (1995), pp. 1245–1250.
15. See Note 3 for Chapter 14.
16. Data simulated from a Normal distribution with $\mu = 98.2$ and $\sigma = 0.7$. These values are based on P. A. Mackowiak, S. S. Wasserman, and M. M. Levine, "A critical appraisal of 98.6°F, the upper limit of the normal body temperature, and other legacies of Carl Reinhold August Wunderlich," *Journal of the American Medical Association*, 268 (1992), pp. 1578–1580.
17. J. F. Swain et al., "Comparison of the effects of oat bran and low-fiber wheat on serum lipoprotein levels and blood pressure," *New England Journal of Medicine*, 322 (1990), pp. 147–152.

Chapter 17 Notes

1. Based on a news item "Bee off with you," *Economist*, November 2, 2002, p. 78.
2. Simplified from D. A. Marcus et al., "A double-blind provocative study of chocolate as a trigger of headache," *Cephalgia*, 17 (1997), pp. 855–862.
3. Votes as of June 27, 2007, at www.pbs.org/wgbh/nova/sciencenow.
4. Data for U.S. searches in February 2008 from Hitwise, at www.hitwise.com.
5. U.S. Census Bureau, *Fertility of American Women: June 2004*, at www.census.gov.
6. Aaron S. Hervey et al., "Reaction time distribution analysis of neuropsychological performance in an ADHD sample," *Child Neuropsychology*, 12 (2006), pp. 125–140.
7. From a Gallup Poll taken in 2003, www.gallup.com.
8. John Schwartz, "Leisure pursuits of today's young men," *New York Times*, March 29, 2004. The source cited is comScore Media Matrix.
9. K. E. Hobbs et al., "Levels and patterns of persistent organochlorines in minke whale (*Balaenoptera acutorostrata*) stocks from the North Atlantic and European Arctic," *Environmental Pollution*, 121 (2003), pp. 239–252.
10. Maureen Hack et al., "Outcomes in young adulthood for very-low-birth-weight infants," *New England Journal of Medicine*, 346 (2002), pp. 149–157.
11. Mikyoung Park et al., "Recycling endosomes supply AMPA receptors for LTP," *Science*, 305 (2004), pp. 1972–1975.
12. Jon E. Keeley, C. J. Fotheringham, and Marco Morais, "Reexamining fire suppression impacts on brushland fire regimes," *Science*, 284 (1999), pp. 1829–1831.

Chapter 18 Notes

1. Note 2 for Chapter 11 explains the reason for this condition in the case of inference about a population mean.
2. D. G. Jakovljevic and A. K. McConnell, "Influence of different breathing frequencies on the severity of inspiratory muscle fatigue induced by high-intensity front crawl swimming," *Journal of Strength and Conditioning Research*, 23(4) (2009), pp. 1169–1174.
3. See Note 3 for Chapter 14.
4. From a graph in Benedetto De Martino et al., "Frames, biases, and rational decision-making in the human brain," *Science*, 313 (2006), pp. 684–687. I simplified the design a bit for easier comprehension: the starting amounts and gambles offered differed from trial to trial, though still matched in pairs; 32 very unbalanced "catch trials" were mixed with the 64 experimental trials to be sure subjects were paying attention; and all money amounts were in British pounds, not dollars.
5. R. A. Berner and G. P. Landis, "Gas bubbles in fossil amber as possible indicators of the major gas composition of ancient air," *Science*, 239 (1988), pp. 1406–1409. The 95% t confidence interval is 54.78 to 64.40. A bootstrap BCa interval is 55.03 to 62.63. So t is reasonably accurate despite the skew and the small sample.
6. This study is available online at www.dispatch.com/live/content/databases/index.html.
7. Alice P. Melis, Brian Hare, and Michael Tomasello, "Chimpanzees recruit the best collaborators," *Science*,

- 311 (2006), pp. 1297–1300. A Normal quantile plot does not show major lack of Normality and a saddlepoint approximation that allows for skew gives $P = 0.0039$. So the t test is reasonably accurate despite the skew and small sample size.
8. Josef P. Rauschecker, Biao Tian, and Marc Hauser, “Processing of complex sounds in the macaque nonprimary auditory cortex,” *Science*, 268 (1995), pp. 111–114.
 9. For a qualitative discussion explaining why skewness is the most serious violation of the Normal shape condition, see Dennis D. Boos and Jacqueline M. Hughes-Oliver, “How large does n have to be for the Z and t intervals?” *American Statistician*, 54 (2000), pp. 121–128. Our recommendations are based on extensive computer work. See, for example, Harry O. Posten, “The robustness of the one-sample t -test over the Pearson system,” *Journal of Statistical Computation and Simulation*, 9 (1979), pp. 133–149; and E. S. Pearson and N. W. Please, “Relation between the shape of population distribution and the robustness of four simple test statistics,” *Biometrika*, 62 (1975), pp. 223–241.
 10. For more advanced users, a good way to ascertain if the t procedures are safe is to compare the 95% confidence interval produced by t with the BCa interval from a bootstrap with at least 1000 resamples. For (b) the t interval is 29,428 to 32,254 and a BCa interval is 29,106 to 31,894. For (c), on the other hand, t gives 38.93 to 40.49 and BCa gives 38.97 to 40.44. These results confirm the judgment that t is safe for (c) but not for (b).
 11. Table 1 in E. Thomassot et al., “Methane-related diamond crystallization in the earth’s mantle: stable isotopes evidence from a single diamond-bearing xenolith,” *Earth and Planetary Science Letters*, 257 (2007), pp. 362–371.
 12. From the online supplement to Tor D. Wager et al., “Placebo-induced changes in fMRI in the anticipation and experience of pain,” *Science*, 303 (2004), pp. 1162–1167.
 13. TUDA results for 2009 from the National Center for Education Statistics, at nationsreportcard.gov/tuda.asp.
 14. Ravi Mehta and Rui Zhu, “Blue or red? Exploring the effect of color on cognitive task performances,” *Science*, 323 (2009), pp. 1226–1229.
 15. Raul de la Fuente-Fernandez et al., “Expectation and dopamine release: mechanism of the placebo effect in Parkinson’s disease,” *Science*, 293 (2001), pp. 1164–1166.
 16. Robert R. Zarr and Dennis D. Leber, “Evaluation and Selection of Candidate Thermal Insulation Materials for NIST SRM 1450d, Fibrous-Glass Board,” available online from the National Institute of Standards and Technology Web site, www.nist.gov/manuscript-publication-search.cfm?pub_id=902936.
 17. J. D. Marshall et al., “Vehicle self-pollution intake fraction: children’s exposure to school bus emissions,” *Environmental Science and Technology*, 39 (2005), pp. 2559–2563.
 18. See Note 16 for Chapter 7.
 19. Data provided by Drina Iglesia, Purdue University. The data are part of a larger study reported in D. D. S. Iglesia, E. J. Cragoe, Jr., and J. W. Vanable, “Electric field strength and epithelialization in the newt (*Notophthalmus viridescens*),” *Journal of Experimental Zoology*, 274 (1996), pp. 56–62.
 20. Matthias R. Mehl et al., “Are women really more talkative than men?” *Science*, 317 (2007), p. 82.
 21. See Note 1 for Chapter 7.
 22. M. B. Laferty “OSU scientist gets a kick of out sports controversy,” *Columbus Dispatch*, November 21, 1993.
 23. We thank Jason Hamilton, University of Illinois, for providing the data. The study is reported in Evan H. DeLucia et al., “Net primary production of a forest ecosystem with experimental CO₂ enhancement,” *Science*, 284 (1999), pp. 1177–1179. No method for inference can be trusted with $n = 3$. In this study, each observation is very costly, so the small n is inevitable.
 24. Michael W. Peugh, “Field investigation of ventilation and air quality in duck and turkey slaughter plants,” MS thesis, Purdue University, 1996.
 25. Harry B. Meyers, “Investigations of the life history of the velvetleaf seed beetle, *Althaeus folkertsi Kingsolver*,” MS thesis, Purdue University, 1996. The 95% t interval is 1227.9 to 2507.6. A 95% bootstrap BCa interval is 1444 to 2718, confirming that t inference is inaccurate for these data.
 26. J. Marcus Jobe and Hutch Jobe, “A statistical approach for additional infill development,” *Energy Exploration and Exploitation*, 18 (2000), pp. 89–103. The comparison interval is the BCa interval based on 1000 bootstrap resamples.
 27. This study is available online at www.dispatch.com/live/content/databases/index.html.
 28. Ralf Bargou et al., “Tumor regression in cancer patients by very low doses of a T cell engaging antibody,” *Science*, 321 (2008), pp. 974–977.
 29. Data provided by Timothy Sturm.
 30. Lianng Yuh, “A biopharmaceutical example for undergraduate students,” manuscript, no date.
 31. See Note 3 for Chapter 14.

Chapter 19 Notes

1. See Note 20 for Chapter 2.
2. Detailed information about the conservative t procedures can be found in Paul Leaverton and John J. Birch, "Small sample power curves for the two sample location problem," *Technometrics*, 11 (1969), pp. 299–307; Henry Scheffé, "Practical solutions of the Behrens-Fisher problem," *Journal of the American Statistical Association*, 65 (1970), pp. 1501–1508; and D. J. Best and J. C. W. Rayner, "Welch's approximate solution for the Behrens-Fisher problem," *Technometrics*, 29 (1987), pp. 205–210.
3. Kathleen G. McKinney, "Engagement in community service among college students: is it affected by significant attachment relationships?" *Journal of Adolescence*, 25 (2002), pp. 139–154. To see the questions in the Inventory of Parent and Peer Attachments, go to chipts.cch.ucla.edu/assessment/IB>List_Scales/inventory%20parent%20and%20peer%20attachment.htm.
4. See Note 10 for Chapter 2.
5. P. A. Handcock, "The effect of age and sex on the perception of time in life," *American Journal of Psychology*, 123 (2010), pp. 1–13.
6. See the extensive simulation studies in Harry O. Posten, "The robustness of the two-sample t -test over the Pearson system," *Journal of Statistical Computation and Simulation*, 6 (1978), pp. 295–311; and Harry O. Posten, H. Yeh, and Donald B. Owen, "Robustness of the two-sample t -test under violations of the homogeneity assumption," *Communications in Statistics*, 11 (1982), pp. 109–126.
7. See Note 15 for Chapter 2. Although the spending data are discrete, a bootstrap BCa 95% confidence interval for the difference in means based on 1000 resamples is 2.394 to 4.826, close to the Option 1 95% interval 2.209 to 4.736. So the sample means are sufficiently Normal to allow use of t procedures.
8. Parmeshwar S. Gupta, "Reaction of plants to the density of soil," *Journal of Ecology*, 21 (1933), pp. 452–474.
9. Data provided by Samuel Phillips, Purdue University.
10. See Note 21 for Chapter 1.
11. The problem of comparing spreads is difficult even with advanced methods. Common distribution-free procedures do not offer a satisfactory alternative to the F test because they are sensitive to unequal shapes when comparing two distributions. A survey of possible approaches is Dennis D. Boos and Cavell Brownie, "Comparing variances and other measures of dispersion," *Statistical Science*, 19 (2005), pp. 571–578.
12. Matthias R. Mehl et al., "Are women really more talkative than men?" *Science*, 317 (2007), p. 82.
13. Michael A. Sayette et al., "Lost in the sauce, the effects of alcohol on mind wandering," *Psychological Science*, 20 (2009), pp. 747–752.
14. Eduardo Dias-Ferreira et al., "Chronic stress causes frontostriatal reorganization and affects decision-making," *Science*, 325 (2009), pp. 621–625. Many of the details appear in the supporting online material.
15. Angeline Lillard and Nicole Else-Quest, "Evaluating Montessori education," *Science*, 313 (2006), pp. 1893–1894. Many of the details appear in the supporting online material.
16. Mary K. Pawlik, "The effect of ginkgo biloba on the post-lunch dip and chemosensory function," MS thesis, Purdue University, 2002.
17. Jennifer A. Whitson and Adam D. Galinsky, "Lacking control increases illusory pattern perception," *Science*, 322 (2008), pp. 115–117.
18. Atsushi Senju et al., "Mindblind eyes: an absence of spontaneous theory of mind in Asperger syndrome," *Science*, 325 (2009), pp. 883–885.
19. Wayne J. Camera and Donald Powers, "Coaching and the SAT I," *TIP* (online journal at www.siop.org/tip), July 1999.
20. Bruce Rind and David B. Strohmetz, "Effect of beliefs about future weather conditions on restaurant tipping," *Journal of Applied Social Psychology*, 31 (2001), pp. 2160–2164.
21. Sherri A. Buzinski, "The effect of position of methylation on the performance properties of durable press treated fabrics," CSR490 honors paper, Purdue University, 1985.
22. Fabrizio Grieco, Arie J. van Noordwijk, and Marcel E. Visser, "Evidence for the effect of learning on timing of reproduction in blue tits," *Science*, 296 (2002), pp. 136–138. The data in Exercise 18.48 are from a graph in this paper.
23. Kathleen D. Vohs, Nicole L. Mead, and Miranda R. Goode, "The psychological consequences of money," *Science*, 314 (2006), pp. 1154–1156. We thank Kathleen Vohs for supplying the data.
24. Paul E. O'Brien et al., "Laparoscopic adjustable gastric banding in severely obese adolescents," *Journal of the American Medical Association*, 303 (2010), pp. 519–526. I thank the authors for providing the data.
25. Paul Kvam, "The effect of active learning methods on student retention in engineering statistics," *American Statistician*, 54 (2000), pp. 136–140.
26. Data provided by Warren Page, New York City Technical College, from a study done by John Hudesman.
27. See Note 9 for Chapter 2.
28. Data provided by Marigene Arnold, Kalamazoo College.

Chapter 20 Notes

1. Joseph H. Catania et al., "Prevalence of AIDS-related risk factors and condom use in the United States," *Science*, 258 (1992), pp. 1101–1106.
2. Strictly speaking, the formula $\sqrt{p(1-p)/n}$ for the standard deviation of \hat{p} assumes that we draw an SRS of size n from an *infinite* population. If the population has finite size N , this standard deviation is multiplied by $\sqrt{1 - (n-1)/(N-1)}$. This "finite population correction" approaches 1 as N increases. When the population is at least 20 times as large as the sample, the correction factor is between about 0.97 and 1. It is reasonable to use the simpler form $\sqrt{p(1-p)/n}$ in these settings. See also Note 2 for Chapter 11.
3. The data were obtained from the GSS Cumulative Datafile 1972-2008-Quick Tables at sda.berkeley.edu/archive.htm. The data were restricted to 2008 and the proportion in the problem is consistent with the proportion obtained from the GSS.
4. The 2010 U.S. Digital Year in Review, at www.comscore.com.
5. This rule of thumb is based on study of computational results in the papers cited in Note 7 and discussion with Alan Agresti. We recommend using the plus four interval.
6. The quotation is from page 1104 of the article cited in Note 1.
7. G. A. Mauser and H. Taylor Buckner, "Canadian attitudes toward gun control: the real story," The Mackenzie Institute, 1997, at teapot.usask.ca/cdn-firearms/Mauser/gunstory.html.
8. D. Gregory Myer et al., "Youth versus adult weightlifting injuries presenting to United States emergency rooms: accidental versus nonaccidental injury mechanisms," *Journal of Strength and Conditioning Research*, 23 (2009), pp. 2054–2060.
9. This interval is proposed by Alan Agresti and Brent A. Coull, "Approximate is better than 'exact' for interval estimation of binomial proportions," *The American Statistician*, 52 (1998), pp. 119–126. Note in particular that the plus four interval is often more accurate than the Clopper-Pearson "exact interval" based on the binomial distribution of the sample count and implemented by, for example, Minitab.
- There are several even more accurate but considerably more complex intervals for p that might be used in professional practice. See Lawrence D. Brown, Tony Cai, and Anirban DasGupta, "Interval estimation for a binomial proportion," *Statistical Science*, 16 (2001), pp. 101–133. A detailed theoretical study that uncovers the reason the large-sample interval is inaccurate is Lawrence D. Brown, Tony Cai, and Anirban DasGupta, "Confidence intervals for a binomial proportion and asymptotic expansions," *Annals of Statistics*, 30 (2002), pp. 160–201.
10. BBC News, December 25, 2006, at news.bbc.co.uk.
11. From Alan Agresti and Brian Caffo, "Simple and effective confidence intervals for proportions and differences of proportions result from adding two successes and two failures," *American Statistician*, 45 (2000), pp. 280–288. When can the plus four interval be safely used? The answer depends on just how much accuracy you insist on. Brown and coauthors (see Note 7) recommend $n \geq 40$. Agresti and Coull (see Note 7) demonstrate that performance is almost always satisfactory in their eyes when $n \geq 5$. Our rule of thumb $n \geq 10$ allows for confidence levels C other than 95% and fits our philosophy of not insisting on more exact results than practice requires. The big point is that plus four is very much more accurate than the standard interval for most values of p and all but very large n .
12. Gary Stoner et al., "Regression of rectal polyps in familial adenomatous polyposis patients with freeze-dried black raspberries," abstract and paper presented at the AACR meeting in 2008.
13. Lydia Saad, "In U.S., 11% of households report computer crimes, a new high," December, 2010, at www.gallup.com. The sampling scheme was more complex than a SRS, so the computation of the number in the sample reporting crimes and acting as if it were a SRS is oversimplified.
14. Gary Edwards and Josephine Mazzuca, "Three quarters of Canadians support doctor-assisted suicide," Gallup Poll press release, March 24, 1999, at www.gallup.com.
15. In fact, P -values for two-sided tests are more accurate than those for one-sided tests. Our rule of thumb is a compromise to avoid the confusion of too many rules.
16. See Note 1 for Chapter 13.
17. Data found on the New Scientist Web site, at www.newscientist.com/article/dn1748-euro-coin-accused.
18. Alexander Todorov et al., "Inferences of competence from faces predict election outcomes," *Science*, 308 (2005), pp. 1623–1626.
19. Michele L. Head, "Examining college students' ethical values," Consumer Science and Retailing honors project, Purdue University, 2003.
20. See Note 3.
21. Elizabeth Cohen, "Your top health searches, asked and answered," Pew Internet and American Life Project,

- 2010, at pewinternet.org. The cell phone sample used random digit dialing drawn through a systematic sampling from dedicated wireless 100-blocks and shared service 100-blocks with no directory-listed landline numbers, so acting as if we have an SRS is oversimplified.
22. Data simulated from a Normal distribution with the mean and standard deviation reported by Sarah Morrison and Jan Noyes, "A comparison of two computer fonts: serif versus ornate sans serif," *Usability News*, 5.2 (2003), at psychology.wichita.edu/surl/usability_news.html.
 23. Data obtained from the Community Data Section of the Dispatch Data Center, at www.dispatch.com/live/content/databases/index.html.
 24. See Note 18 for Chapter 8.
 25. Bobby D. Rampey et al., *The Nation's Report Card: Trends in Academic Progress in Reading and Mathematics 2008*, can be found on the Web site nces.ed.gov/nationsreportcard/ under "Long term trends."
 26. Francisco Lloret et al., "Fire and resprouting in Mediterranean ecosystems: insights from an external biogeographical region, the Mexican shrubland," *American Journal of Botany*, 88 (1999), pp. 1655–1661.
 27. Jon D. Miller, Eugenie C. Scott, and Shinji Okamoto, "Public acceptance of evolution," *Science*, 313 (2006), pp. 765–766. The information in the exercise appears in the supplementary online material.
 28. A. Mantonakis, et al., "Order in choice: effects of serial position on preferences," *Psychological Science*, 20 (2009), pp. 1309–1312.
 29. Laura Tutor, "Navigating the Loop: The best drive-thru in America '02," *QSR*, October 2002, pp. 41–59.
 7. Modified from Richard A. Schieber et al., "Risk factors for injuries from in-line skating and the effectiveness of safety gear," *New England Journal of Medicine*, 335 (1996), Internet summary at content.nejm.org.
 8. Saiyad S. Ahmed, "Effects of microwave drying on checking and mechanical strength of low-moisture baked products," MS thesis, Purdue University, 1994.
 9. Shauna B. Wilson, et. al., "Dating across race: an examination of African American Internet personal advertisements," *Journal of Black Studies*, 37 (2007), pp. 964–982.
 10. This rule of thumb is quite conservative. It is in fact safe to arrange the data as a 2×2 table and apply the rule of thumb from Chapter 22 that all four *expected* counts must be 5 or greater. I give the conservative rule here because expected counts are messy to explain in the present context.
 11. JoAnn K. Wells, Allan F. Williams, and Charles M. Farmer, "Seat belt use among African Americans, Hispanics, and whites," *Accident Analysis and Prevention*, 34 (2002), pp. 523–529.
 12. Steiner Sulheim et al., "Helmet use and risk of head injuries in alpine skiers and snowboarders," *Journal of the American Medical Association*, 295 (2006), pp. 919–924.
 13. Armando E. Giuliano, M.D. et al., "Axillary dissection vs no axillary dissection in women with invasive breast cancer and sentinel node metastasis," *Journal of the American Medical Association*, 305 (2011), pp. 569–575. The sample sizes for the two groups and the proportions of patients in each group that are disease-free after 5 years have been chosen to match those in the paper.
 14. From the Prevalence and Trends Data of the Behavioral Risk Factor Surveillance System (BRFSS), at www.cdc.gov/BRFSS/.
 15. See Amanda Lenhart and Mary Madden, "Teens, privacy and online social networks," Pew Internet and American Life Project, 2007, at www.pewinternet.org.
 16. W. P. T. James et al., "Effect of sibutramine on cardiovascular outcomes in overweight and Obese Subjects," *New England Journal of Medicine*, 363 (2010), pp. 905–917.
 17. François Gaudet et al., "Induction of tumors in mice by genomic hypomethylation," *Science*, 300 (2003), pp. 489–492.
 18. Barbara Helmrich, "Window of opportunity? Adolescence, music and algebra," *Journal of Adolescent Research*, 25 (2010), pp. 557–577.
 19. John Fagan et al., "Performance assessment under field conditions of a rapid immunological test for transgenic soybeans," *International Journal of Food Science and Technology*, 36 (2001), pp. 357–367.

Chapter 21 Notes

1. See Note 1 for Chapter 6.
2. Based on data in Amanda Lenhart, "Cell phones and American adults," Pew Internet and American Life Project, September 2010, at pewinternet.org.
3. The data were obtained from the GSS Cumulative Datafile 1972–2008 at <http://sda.berkeley.edu/archive.htm>.
4. From the 2009 Youth Risk Behavior Surveillance System at <http://apps.nccd.cdc.gov/youthonline/App/Default.aspx?SID=HS>. The data are from a complex multistage sample, so that acting as if we have SRSs is oversimplified.
5. The plus four method is due to Alan Agresti and Brian Caffo. See Note 9 for Chapter 19.
6. See Note 26 for Chapter 20.

20. Arne L. Kalleberg and Kevin T. Leicht, "Gender and organizational performance: determinants of small business survival and success," *Academy of Management Journal*, 34 (1991), pp. 136–161.
21. Richard M. Felder et al., "Who gets it and who doesn't: a study of student performance in an introductory chemical engineering course," 1992 ASEE Annual Conference Proceedings, American Society for Engineering Education, Washington, D.C., 1992, pp. 1516–1519.
22. D. Gonzales et al., "Varenicline, an $\alpha 4\beta 2$ nicotinic acetylcholine receptor partial agonist, vs sustained-release bupropion and placebo for smoking cessation," *Journal of the American Medical Association*, 296 (2006), pp. 47–55.
23. Data courtesy of Raymond Dumett, Purdue University.
24. Based on Alan G. Sanfey et al., "The neural basis of economic decision-making in the ultimatum game," *Science*, 300 (2003), pp. 1755–1758. The paper reports a chi-square test (equivalent to a two-sided z test). This analysis is incorrect for the paper's data, as there were in fact only 19 participants, each appearing twice in each row of the table given in the exercise. Exercise 20.32 therefore amends the data, assuming 76 participants, so that the elementary analysis is correct.
25. Ross L. Prentice et al., "Low-fat dietary pattern and risk of invasive breast cancer," *Journal of the American Medical Association*, 295 (2006), pp. 629–642.
26. Clive G. Jones et al., "Chain reactions linking acorns to gypsy moth outbreaks and Lyme disease risk," *Science*, 279 (1998), pp. 1023–1026.
27. The study is reported in William Celis III, "Study suggests Head Start helps beyond school," *New York Times*, April 20, 1993. See www.highscope.org.
28. R. B. Turner et al., "Hand disinfection for the prevention of viral respiratory illness" ICAAC abstract 101, 2010.
5. K. S. Oberhauser, "Fecundity, lifespan and egg mass in butterflies: effects of male-derived nutrients and female size," *Functional Ecology*, 11 (1997), pp. 166–175.
6. Michael R. Dohm, Jack P. Hayes, and Theodore Garland, Jr., "Quantitative genetics of sprint running speed and swimming endurance in laboratory house mice (*Mus domesticus*)," *Evolution*, 50 (1996), pp. 1688–1701.
7. See Note 10 for Chapter 17. The exercises are simplified in that the measures reported in this paper have been statistically adjusted for "sociodemographic status."
8. See Note 3 for Chapter 21.
9. V. D. Bass, W. E. Hoffmann, and J. L. Dorner, "Normal canine lipid profiles and effects of experimentally induced pancreatitis and hepatic necrosis on lipids," *American Journal of Veterinary Research*, 37 (1976), pp. 1355–1357.
10. Jin Ha Lee and J. Stephen Downie, "Survey of music information needs, uses, and seeking behaviors: preliminary findings," online *Proceedings of the 5th International Conference on Music Information Retrieval*, 2004, at ismir2004.ismir.net.
11. K. Carrie Armel and V. S. Ramachandran, "Projecting sensations to external objects: evidence from skin conductance response," *Proceedings of the Royal Society of London, Series B*, 270 (2003), pp. 1499–1506.
12. Josh McDermott and Marc D. Hauser, "Nonhuman primates prefer slow tempos but dislike music overall," *Cognition*, 104 (2007), pp. 654–668. Failure to take account of repeated measures on the same subjects is one of the most common errors observed in statistical analysis.
13. James Otto, Michael F. Brown, and William Long III, "Training rats to search and alert on contraband odors," *Applied Animal Behaviour Science*, 77 (2002), pp. 217–232.
14. From the Merck Web site, www.merckvaccines.com/gardasilProductPage frmst.html.
15. These data were originally collected by L. M. Linde of UCLA but were first published by M. R. Mickey, O. J. Dunn, and V. Clark, "Note on the use of stepwise regression in detecting outliers," *Computers and Biomedical Research*, 1 (1967), pp. 105–111. The data have been used by several authors. I found them in N. R. Draper and J. A. John, "Influential observations and outliers in regression," *Technometrics*, 23 (1981), pp. 21–26.
16. Jacqueline T. Ngai and Diane S. Srivastava, "Predators accelerate nutrient cycling in a bromeliad ecosystem," *Science*, 314 (2006), p. 963. We thank Jacqueline Ngai for providing the data.

Chapter 22 Notes

1. Data provided by Drina Iglesia, Purdue University. The data are part of a larger study reported in D. D. S. Iglesia, E. J. Cragoe, Jr., and J. W. Vanable, "Electric field strength and epithelialization in the newt (*Notophthalmus viridescens*)," *Journal of Experimental Zoology*, 274 (1996), pp. 56–62.
2. See Note 4 for Chapter 21.
3. Lee Rainie and Bill Tancer, "36% of online American adults consult Wikipedia," Pew Internet and American Life Project, 2007, at www.pewinternet.org.
4. See Note 5 for Chapter 14.

17. Yvan R. Germain, "The dyeing of ramie with fiber reactive dyes using the cold pad-batch method," MS thesis, Purdue University, 1988.
18. Data provided by Marigene Arnold, Kalamazoo College.
19. Data provided by Corinne Lim, Purdue University, from a student project supervised by Professor Joseph Vanable.
20. See Note 8 for Chapter 21.
21. Michael O. Finkelstein and Bruce Levin, "Statistical proof of discrimination in peremptory challenges," *Chance*, 17, No. 1 (2004), pp. 35–38.
22. G. S. Hotamisligil et al., "Uncoupling of obesity from insulin resistance through a targeted mutation in *aP2*, the adipocyte fatty acid binding protein," *Science*, 274 (1996), pp. 1377–1379.
8. David W. Eby et al., "The effect of changing from secondary to primary safety belt enforcement on police harassment," *Accident Analysis and Prevention*, 36 (2000), pp. 819–828.
9. See Note 28 from Chapter 20.
10. See Note 9 for Chapter 6.
11. See Note 12 for Chapter 6.
12. Lien-Ti Bei, "Consumers' purchase behavior toward recycled products: an acquisition-transaction utility theory perspective," MS thesis, Purdue University, 1993.
13. Modified from Felicity Barringer, "Measuring sexuality through polls can be shaky," *New York Times*, April 25, 1993.
14. Virgilio P. Carnielli et al., "Intestinal absorption of long-chain polyunsaturated fatty acids in preterm infants fed breast milk or formula," *American Journal of Clinical Nutrition*, 67 (1998), pp. 97–103.
15. Adapted from M. A. Visintainer, J. R. Volpicelli, and M. E. P. Seligman, "Tumor rejection in rats after inescapable or escapable shock," *Science*, 216 (1982), pp. 437–439.
16. See Note 1 for Chapter 6.
17. Tom Reichert, "The prevalence of sexual imagery in ads targeted to young adults," *Journal of Consumer Affairs*, 37 (2003), pp. 403–412.
18. József Topál et al., "Differential sensitivity to human communication in dogs, wolves and human infants," *Science*, 325 (2009), pp. 1269–1272. Many statistical software packages offer "exact tests" that are valid even when there are small expected counts.
19. U.S. Department of Commerce, Office of Travel and Tourism Industries, in-flight survey, 2007, at tinet.ita.doc.gov.
20. See Note 2 for Chapter 16. I have simplified slightly: the table in the paper is exactly as in the exercise but contains data for 63 subjects plus data from one type of bar for 3 subjects who dropped out. Although the authors say that their chi-square refers to this table, they give a nonsignificant value that contradicts what the table shows.
21. See Note 15 for Chapter 6.
22. All General Social Survey exercises in this chapter present tables constructed using the search function at the GSS archive, sda.berkeley.edu/archive.htm. Most concern data from the 2008 GSS.
23. Two-way tables from the Youth Risk Behavior Surveillance System can be constructed from the Web site nccd.cdc.gov/youthonline/App/Default.aspx.
24. Data compiled from a table of percents in "Americans view higher education as key to the American dream,"

Chapter 23 Notes

1. See Note 1 for Chapter 6.
2. Pennsylvania State University Division of Student Affairs, "Net behaviors November 2006," *Penn State Pulse*, at www.sa.psu.edu.
3. See Note 2 for Chapter 6.
4. All General Social Survey exercises in this chapter present tables constructed using the search function at the GSS archive, sda.berkeley.edu/archive.htm. Most concern data from the 2008 GSS.
5. There are many computer studies of the accuracy of chi-square critical values for χ^2 . Our guideline goes back to W. G. Cochran (1954). Later work has shown that it is often conservative in the sense that, if the expected cell counts are all similar and the degrees of freedom exceed 1, the chi-square approximation works well for an average expected count as small as 1 or 2. Our guideline protects against dissimilar expected counts. It has the added advantage that it is safe in the 2×2 case, where the chi-square approximation is least good. So our guideline is helpful for beginners—there is no single condition that is not conservative and applies to 2×2 and larger tables with similar and dissimilar expected cell counts. There are exact procedures that (with software) should be used for tables that do not satisfy our guideline. For a survey, see Alan Agresti, "A survey of exact inference for contingency tables," *Statistical Science*, 7 (1992), pp. 131–177.
6. Pew Research Center for the People and the Press, "The cell phone challenge to survey research," news release for May 15, 2006, at www.people-press.org.
7. Based on a news item in *Science*, 305 (2004), p. 1560. The study, by Daniel Klem, appeared in *Wilson Journal*.
8. See Note 15 for Chapter 6.
9. All General Social Survey exercises in this chapter present tables constructed using the search function at the GSS archive, sda.berkeley.edu/archive.htm. Most concern data from the 2008 GSS.
10. Two-way tables from the Youth Risk Behavior Surveillance System can be constructed from the Web site nccd.cdc.gov/youthonline/App/Default.aspx.
11. Data compiled from a table of percents in "Americans view higher education as key to the American dream,"

- press release from the National Center for Public Policy and Higher Education, May 3, 2000, at www.highereducation.org.
25. See Note 14 for Chapter 6.
 26. Data produced by P. Ries and H. Smith, found in William D. Johnson and Gary G. Koch, "A note on the weighted least squares analysis of the Ries-Smith contingency table data," *Technometrics*, 13 (1971), pp. 438–447.

Chapter 24 Notes

1. Samuel Karelitz et al., "Relation of crying activity in early infancy to speech and intellectual development at age three years," *Child Development*, 35 (1964), pp. 769–777.
2. From a graph in Naomi E. Allen et al., "Moderate alcohol intake and cancer incidence in women," *Journal of the National Cancer Institute*, 101 (2009), pp. 296–305. These data represent averages over large numbers of women and are an example of an ecological correlation (see page 142 in Chapter 5), and one must be careful not to interpret the data as applying to individuals.
3. From a graph in Stephen M. Fleming et al., "Relating introspective accuracy to individual differences in brain structure," *Science*, 329 (2010), pp. 1541–1543.
4. Data for 1936–1999 are from a graph in Bruce J. Peterson et al., "Increasing river discharge to the Arctic Ocean," *Science*, 298 (2002), pp. 2171–2173. Data for 2000–2008 are from a graph in I. Ashik et al., "Arctic report card: update for 2010," available online at www.arctic.noaa.gov/reportcard/ArcticReportCard_full_report.pdf. The graph is on page 41 of the report.
5. Electronic Encyclopedia of Statistical Examples and Exercises (EESEE) at the text Web site, www.whfreeman.com/bps.
6. From a graph in Allison L. Perry et al., "Climate change and distribution shifts in marine fishes," *Science*, 308 (2005), pp. 1912–1915. The explanatory variable is the five-year running mean of winter (December to March) sea-bottom temperature.
7. Data for the building at 1800 Ben Franklin Drive, Sarasota, Florida, starting in February 2003. From the Web site of the Sarasota County Property Appraiser, www.sarasotaproperty.net.
8. Yanhui Lu et al., "Mirid bug outbreaks in multiple crops correlated with wide-scale adoption of Bt cotton in China," *Science*, 328 (2010), pp. 1151–1154.
9. Neal E. Cantin et al., "Ocean warming slows coral growth in the Central Red Sea," *Science*, 329 (2010), pp. 322–325.

10. Based on Marion E. Dunshee, "A study of factors affecting the amount and kind of food eaten by nursery school children," *Child Development*, 2 (1931), pp. 163–183. This article gives the means, standard deviations, and correlation for 37 children, from which the data in the exercise are simulated.
11. From Table S2 in the online supplement to Antonio Dell'Anno and Roberto Danovaro, "Extracellular DNA plays a key role in deep-sea ecosystem functioning," *Science*, 309 (2005), p. 2179.
12. See Note 21 for Chapter 7.

Chapter 25 Notes

1. See Note 9 for Chapter 2.
2. Victoria L. Brescoll and Eric L. Uhlmann, "Can an angry woman get ahead? Status conferral, gender and expression of emotion in the workplace," *Psychological Science*, 19 (2008), pp. 268–273. The description and data are based on study 1 in this article.
3. Elisabeth Wells-Parker et al., "An exploratory study of the relationship between road rage and crash experience in a representative sample of US drivers," *Accident Analysis and Prevention*, 34 (2002), pp. 271–278.
4. See Note 10 for Chapter 2.
5. The data from the General Social Survey for this exercise were constructed using the search function and download capabilities at the GSS archive, sda.berkeley.edu/archive.htm.
6. See Note 23 for Chapter 19.
7. See Note 16 for Chapter 22.
8. David B. Wooten, "One-of-a-kind in a full house: some consequences of ethnic and gender distinctiveness," *Journal of Consumer Psychology*, 4 (1995), 205–224.
9. John P. Thomas, "Influences on mathematics learning and attitudes among African American high school students," *Journal of Negro Education*, 69 (2000), pp. 165–183.
10. Henrik Hagvedt and Vanessa M. Patrick, "Art infusion: the influence of visual art on the perception and evaluation of consumer products," *Journal of Marketing Research*, XLV (2008), pp. 379–389.
11. Timothy Church et al., "Effects of aerobic and resistance training on hemoglobin A1c levels in patients with type 2 diabetes: a randomized controlled trial," *Journal of the American Medical Association*, 304 (2010), pp. 2253–2262.
12. Jennifer J. Argo et al., "Positive consumer contagion: Responses to attractive others in a retail context," *Journal of Marketing Research*, XLV (2008), pp. 690–701.

13. Data from the online supplement to André Kessler and Ian T. Baldwin, “Defensive function of herbivore-induced plant volatile emissions in nature,” *Science*, 291 (2001), pp. 2141–2144.
14. The data and the full story can be found in the Data and Story Library at lib.stat.cmu.edu. The original study is by Faith Loven, “A study of interlist equivalency of the CID W-22 word list presented in quiet and in noise,” MS thesis, University of Iowa, 1981.
15. See Note 8 for Chapter 9. We thank Kenwyn Suttle for providing these data for the year 2003.
16. See Note 17 for Chapter 22.
17. See Note 19 for Chapter 2.
18. See Note 17 for Chapter 4.
19. Sherri A. Buzinski, “The effect of position of methylation on the performance properties of durable press treated fabrics,” CSR490 honors paper, Purdue University, 1985.
20. See Note 8 for Chapter 9.



Tables

Table A Standard Normal Probabilities

Table B Random Digits

Table C t Distribution Critical Values

Table D Chi-square Distribution Critical Values

Table E Critical Values of the Correlation r

Table entry for z is the area under the standard Normal curve to the left of z .

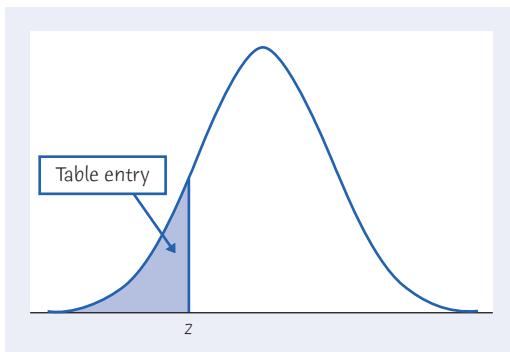


TABLE A Standard Normal cumulative proportions

<i>z</i>	.00	.01	.02	.03	.04	.05	.06	.07	.08	.09
-3.4	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0003	.0002
-3.3	.0005	.0005	.0005	.0004	.0004	.0004	.0004	.0004	.0004	.0003
-3.2	.0007	.0007	.0006	.0006	.0006	.0006	.0006	.0005	.0005	.0005
-3.1	.0010	.0009	.0009	.0009	.0008	.0008	.0008	.0008	.0007	.0007
-3.0	.0013	.0013	.0012	.0012	.0012	.0011	.0011	.0011	.0010	.0010
-2.9	.0019	.0018	.0018	.0017	.0016	.0016	.0015	.0015	.0014	.0014
-2.8	.0026	.0025	.0024	.0023	.0023	.0022	.0021	.0021	.0020	.0019
-2.7	.0035	.0034	.0033	.0032	.0031	.0030	.0029	.0028	.0027	.0026
-2.6	.0047	.0045	.0044	.0043	.0041	.0040	.0039	.0038	.0037	.0036
-2.5	.0062	.0060	.0059	.0057	.0055	.0054	.0052	.0051	.0049	.0048
-2.4	.0082	.0080	.0078	.0075	.0073	.0071	.0069	.0068	.0066	.0064
-2.3	.0107	.0104	.0102	.0099	.0096	.0094	.0091	.0089	.0087	.0084
-2.2	.0139	.0136	.0132	.0129	.0125	.0122	.0119	.0116	.0113	.0110
-2.1	.0179	.0174	.0170	.0166	.0162	.0158	.0154	.0150	.0146	.0143
-2.0	.0228	.0222	.0217	.0212	.0207	.0202	.0197	.0192	.0188	.0183
-1.9	.0287	.0281	.0274	.0268	.0262	.0256	.0250	.0244	.0239	.0233
-1.8	.0359	.0351	.0344	.0336	.0329	.0322	.0314	.0307	.0301	.0294
-1.7	.0446	.0436	.0427	.0418	.0409	.0401	.0392	.0384	.0375	.0367
-1.6	.0548	.0537	.0526	.0516	.0505	.0495	.0485	.0475	.0465	.0455
-1.5	.0668	.0655	.0643	.0630	.0618	.0606	.0594	.0582	.0571	.0559
-1.4	.0808	.0793	.0778	.0764	.0749	.0735	.0721	.0708	.0694	.0681
-1.3	.0968	.0951	.0934	.0918	.0901	.0885	.0869	.0853	.0838	.0823
-1.2	.1151	.1131	.1112	.1093	.1075	.1056	.1038	.1020	.1003	.0985
-1.1	.1357	.1335	.1314	.1292	.1271	.1251	.1230	.1210	.1190	.1170
-1.0	.1587	.1562	.1539	.1515	.1492	.1469	.1446	.1423	.1401	.1379
-0.9	.1841	.1814	.1788	.1762	.1736	.1711	.1685	.1660	.1635	.1611
-0.8	.2119	.2090	.2061	.2033	.2005	.1977	.1949	.1922	.1894	.1867
-0.7	.2420	.2389	.2358	.2327	.2296	.2266	.2236	.2206	.2177	.2148
-0.6	.2743	.2709	.2676	.2643	.2611	.2578	.2546	.2514	.2483	.2451
-0.5	.3085	.3050	.3015	.2981	.2946	.2912	.2877	.2843	.2810	.2776
-0.4	.3446	.3409	.3372	.3336	.3300	.3264	.3228	.3192	.3156	.3121
-0.3	.3821	.3783	.3745	.3707	.3669	.3632	.3594	.3557	.3520	.3483
-0.2	.4207	.4168	.4129	.4090	.4052	.4013	.3974	.3936	.3897	.3859
-0.1	.4602	.4562	.4522	.4483	.4443	.4404	.4364	.4325	.4286	.4247
-0.0	.5000	.4960	.4920	.4880	.4840	.4801	.4761	.4721	.4681	.4641

Table entry for z is the area under the standard Normal curve to the left of z .

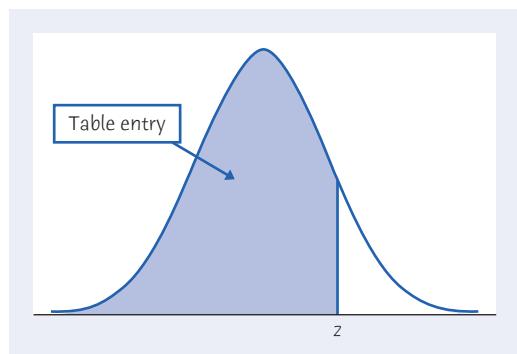


TABLE A Standard Normal cumulative proportions (continued)

TABLE B Random digits

LINE							
101	19223	95034	05756	28713	96409	12531	42544
102	73676	47150	99400	01927	27754	42648	82425
103	45467	71709	77558	00095	32863	29485	82226
104	52711	38889	93074	60227	40011	85848	48767
105	95592	94007	69971	91481	60779	53791	17297
106	68417	35013	15529	72765	85089	57067	50211
107	82739	57890	20807	47511	81676	55300	94383
108	60940	72024	17868	24943	61790	90656	87964
109	36009	19365	15412	39638	85453	46816	83485
110	38448	48789	18338	24697	39364	42006	76688
111	81486	69487	60513	09297	00412	71238	27649
112	59636	88804	04634	71197	19352	73089	84898
113	62568	70206	40325	03699	71080	22553	11486
114	45149	32992	75730	66280	03819	56202	02938
115	61041	77684	94322	24709	73698	14526	31893
116	14459	26056	31424	80371	65103	62253	50490
117	38167	98532	62183	70632	23417	26185	41448
118	73190	32533	04470	29669	84407	90785	65956
119	95857	07118	87664	92099	58806	66979	98624
120	35476	55972	39421	65850	04266	35435	43742
121	71487	09984	29077	14863	61683	47052	62224
122	13873	81598	95052	90908	73592	75186	87136
123	54580	81507	27102	56027	55892	33063	41842
124	71035	09001	43367	49497	72719	96758	27611
125	96746	12149	37823	71868	18442	35119	62103
126	96927	19931	36809	74192	77567	88741	48409
127	43909	99477	25330	64359	40085	16925	85117
128	15689	14227	06565	14374	13352	49367	81982
129	36759	58984	68288	22913	18638	54303	00795
130	69051	64817	87174	09517	84534	06489	87201
131	05007	16632	81194	14873	04197	85576	45195
132	68732	55259	84292	08796	43165	93739	31685
133	45740	41807	65561	33302	07051	93623	18132
134	27816	78416	18329	21337	35213	37741	04312
135	66925	55658	39100	78458	11206	19876	87151
136	08421	44753	77377	28744	75592	08563	79140
137	53645	66812	61421	47836	12609	15373	98481
138	66831	68908	40772	21558	47781	33586	79177
139	55588	99404	70708	41098	43563	56934	48394
140	12975	13258	13048	45144	72321	81940	00360
141	96767	35964	23822	96012	94591	65194	50842
142	72829	50232	97892	63408	77919	44575	24870
143	88565	42628	17797	49376	61762	16953	88604
144	62964	88145	83083	69453	46109	59505	69680
145	19687	12633	57857	95806	09931	02150	43163
146	37609	59057	66967	83401	60705	02384	90597
147	54973	86278	88737	74351	47500	84552	19909
148	00694	05977	19664	65441	20903	62371	22725
149	71546	05233	53946	68743	72460	27601	45403
150	07511	88915	41267	16853	84569	79367	32337

Table entry for C is the critical value t^* required for confidence level C. To approximate one- and two-sided P-values, compare the value of the t statistic with the critical values of t^* that match the P-values given at the bottom of the table.

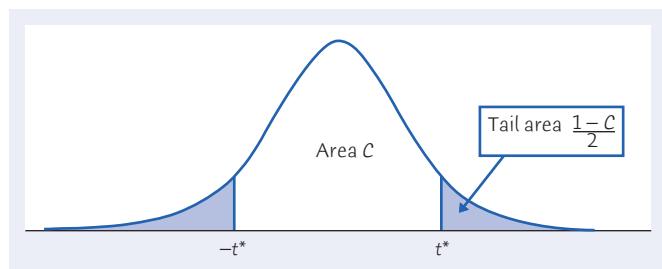


TABLE C *t* distribution critical values

DEGREES OF FREEDOM	CONFIDENCE LEVEL C											
	50%	60%	70%	80%	90%	95%	96%	98%	99%	99.5%	99.8%	99.9%
1	1.000	1.376	1.963	3.078	6.314	12.71	15.89	31.82	63.66	127.3	318.3	636.6
2	0.816	1.061	1.386	1.886	2.920	4.303	4.849	6.965	9.925	14.09	22.33	31.60
3	0.765	0.978	1.250	1.638	2.353	3.182	3.482	4.541	5.841	7.453	10.21	12.92
4	0.741	0.941	1.190	1.533	2.132	2.776	2.999	3.747	4.604	5.598	7.173	8.610
5	0.727	0.920	1.156	1.476	2.015	2.571	2.757	3.365	4.032	4.773	5.893	6.869
6	0.718	0.906	1.134	1.440	1.943	2.447	2.612	3.143	3.707	4.317	5.208	5.959
7	0.711	0.896	1.119	1.415	1.895	2.365	2.517	2.998	3.499	4.029	4.785	5.408
8	0.706	0.889	1.108	1.397	1.860	2.306	2.449	2.896	3.355	3.833	4.501	5.041
9	0.703	0.883	1.100	1.383	1.833	2.262	2.398	2.821	3.250	3.690	4.297	4.781
10	0.700	0.879	1.093	1.372	1.812	2.228	2.359	2.764	3.169	3.581	4.144	4.587
11	0.697	0.876	1.088	1.363	1.796	2.201	2.328	2.718	3.106	3.497	4.025	4.437
12	0.695	0.873	1.083	1.356	1.782	2.179	2.303	2.681	3.055	3.428	3.930	4.318
13	0.694	0.870	1.079	1.350	1.771	2.160	2.282	2.650	3.012	3.372	3.852	4.221
14	0.692	0.868	1.076	1.345	1.761	2.145	2.264	2.624	2.977	3.326	3.787	4.140
15	0.691	0.866	1.074	1.341	1.753	2.131	2.249	2.602	2.947	3.286	3.733	4.073
16	0.690	0.865	1.071	1.337	1.746	2.120	2.235	2.583	2.921	3.252	3.686	4.015
17	0.689	0.863	1.069	1.333	1.740	2.110	2.224	2.567	2.898	3.222	3.646	3.965
18	0.688	0.862	1.067	1.330	1.734	2.101	2.214	2.552	2.878	3.197	3.611	3.922
19	0.688	0.861	1.066	1.328	1.729	2.093	2.205	2.539	2.861	3.174	3.579	3.883
20	0.687	0.860	1.064	1.325	1.725	2.086	2.197	2.528	2.845	3.153	3.552	3.850
21	0.686	0.859	1.063	1.323	1.721	2.080	2.189	2.518	2.831	3.135	3.527	3.819
22	0.686	0.858	1.061	1.321	1.717	2.074	2.183	2.508	2.819	3.119	3.505	3.792
23	0.685	0.858	1.060	1.319	1.714	2.069	2.177	2.500	2.807	3.104	3.485	3.768
24	0.685	0.857	1.059	1.318	1.711	2.064	2.172	2.492	2.797	3.091	3.467	3.745
25	0.684	0.856	1.058	1.316	1.708	2.060	2.167	2.485	2.787	3.078	3.450	3.725
26	0.684	0.856	1.058	1.315	1.706	2.056	2.162	2.479	2.779	3.067	3.435	3.707
27	0.684	0.855	1.057	1.314	1.703	2.052	2.158	2.473	2.771	3.057	3.421	3.690
28	0.683	0.855	1.056	1.313	1.701	2.048	2.154	2.467	2.763	3.047	3.408	3.674
29	0.683	0.854	1.055	1.311	1.699	2.045	2.150	2.462	2.756	3.038	3.396	3.659
30	0.683	0.854	1.055	1.310	1.697	2.042	2.147	2.457	2.750	3.030	3.385	3.646
40	0.681	0.851	1.050	1.303	1.684	2.021	2.123	2.423	2.704	2.971	3.307	3.551
50	0.679	0.849	1.047	1.299	1.676	2.009	2.109	2.403	2.678	2.937	3.261	3.496
60	0.679	0.848	1.045	1.296	1.671	2.000	2.099	2.390	2.660	2.915	3.232	3.460
80	0.678	0.846	1.043	1.292	1.664	1.990	2.088	2.374	2.639	2.887	3.195	3.416
100	0.677	0.845	1.042	1.290	1.660	1.984	2.081	2.364	2.626	2.871	3.174	3.390
1000	0.675	0.842	1.037	1.282	1.646	1.962	2.056	2.330	2.581	2.813	3.098	3.300
t^*	0.674	0.841	1.036	1.282	1.645	1.960	2.054	2.326	2.576	2.807	3.091	3.291
One-sided P	.25	.20	.15	.10	.05	.025	.02	.01	.005	.0025	.001	.0005
Two-sided P	.50	.40	.30	.20	.10	.05	.04	.02	.01	.005	.002	.001

Table entry for p is the critical value χ^* with probability p lying to its right.

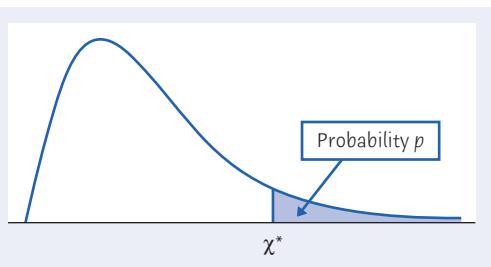


TABLE D Chi-square distribution critical values

df	<i>p</i>											
	.25	.20	.15	.10	.05	.025	.02	.01	.005	.0025	.001	.0005
1	1.32	1.64	2.07	2.71	3.84	5.02	5.41	6.63	7.88	9.14	10.83	12.12
2	2.77	3.22	3.79	4.61	5.99	7.38	7.82	9.21	10.60	11.98	13.82	15.20
3	4.11	4.64	5.32	6.25	7.81	9.35	9.84	11.34	12.84	14.32	16.27	17.73
4	5.39	5.99	6.74	7.78	9.49	11.14	11.67	13.28	14.86	16.42	18.47	20.00
5	6.63	7.29	8.12	9.24	11.07	12.83	13.39	15.09	16.75	18.39	20.51	22.11
6	7.84	8.56	9.45	10.64	12.59	14.45	15.03	16.81	18.55	20.25	22.46	24.10
7	9.04	9.80	10.75	12.02	14.07	16.01	16.62	18.48	20.28	22.04	24.32	26.02
8	10.22	11.03	12.03	13.36	15.51	17.53	18.17	20.09	21.95	23.77	26.12	27.87
9	11.39	12.24	13.29	14.68	16.92	19.02	19.68	21.67	23.59	25.46	27.88	29.67
10	12.55	13.44	14.53	15.99	18.31	20.48	21.16	23.21	25.19	27.11	29.59	31.42
11	13.70	14.63	15.77	17.28	19.68	21.92	22.62	24.72	26.76	28.73	31.26	33.14
12	14.85	15.81	16.99	18.55	21.03	23.34	24.05	26.22	28.30	30.32	32.91	34.82
13	15.98	16.98	18.20	19.81	22.36	24.74	25.47	27.69	29.82	31.88	34.53	36.48
14	17.12	18.15	19.41	21.06	23.68	26.12	26.87	29.14	31.32	33.43	36.12	38.11
15	18.25	19.31	20.60	22.31	25.00	27.49	28.26	30.58	32.80	34.95	37.70	39.72
16	19.37	20.47	21.79	23.54	26.30	28.85	29.63	32.00	34.27	36.46	39.25	41.31
17	20.49	21.61	22.98	24.77	27.59	30.19	31.00	33.41	35.72	37.95	40.79	42.88
18	21.60	22.76	24.16	25.99	28.87	31.53	32.35	34.81	37.16	39.42	42.31	44.43
19	22.72	23.90	25.33	27.20	30.14	32.85	33.69	36.19	38.58	40.88	43.82	45.97
20	23.83	25.04	26.50	28.41	31.41	34.17	35.02	37.57	40.00	42.34	45.31	47.50
21	24.93	26.17	27.66	29.62	32.67	35.48	36.34	38.93	41.40	43.78	46.80	49.01
22	26.04	27.30	28.82	30.81	33.92	36.78	37.66	40.29	42.80	45.20	48.27	50.51
23	27.14	28.43	29.98	32.01	35.17	38.08	38.97	41.64	44.18	46.62	49.73	52.00
24	28.24	29.55	31.13	33.20	36.42	39.36	40.27	42.98	45.56	48.03	51.18	53.48
25	29.34	30.68	32.28	34.38	37.65	40.65	41.57	44.31	46.93	49.44	52.62	54.95
26	30.43	31.79	33.43	35.56	38.89	41.92	42.86	45.64	48.29	50.83	54.05	56.41
27	31.53	32.91	34.57	36.74	40.11	43.19	44.14	46.96	49.64	52.22	55.48	57.86
28	32.62	34.03	35.71	37.92	41.34	44.46	45.42	48.28	50.99	53.59	56.89	59.30
29	33.71	35.14	36.85	39.09	42.56	45.72	46.69	49.59	52.34	54.97	58.30	60.73
30	34.80	36.25	37.99	40.26	43.77	46.98	47.96	50.89	53.67	56.33	59.70	62.16
40	45.62	47.27	49.24	51.81	55.76	59.34	60.44	63.69	66.77	69.70	73.40	76.09
50	56.33	58.16	60.35	63.17	67.50	71.42	72.61	76.15	79.49	82.66	86.66	89.56
60	66.98	68.97	71.34	74.40	79.08	83.30	84.58	88.38	91.95	95.34	99.61	102.7
80	88.13	90.41	93.11	96.58	101.9	106.6	108.1	112.3	116.3	120.1	124.8	128.3
100	109.1	111.7	114.7	118.5	124.3	129.6	131.1	135.8	140.2	144.3	149.4	153.2

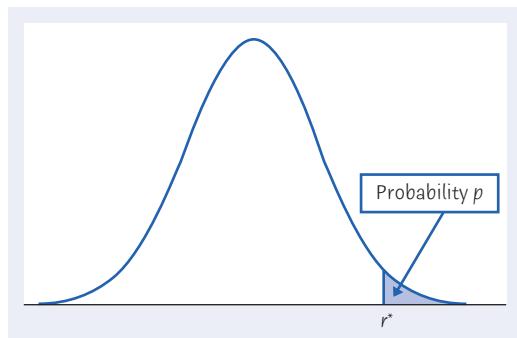


Table entry for p is the critical value r^* of the correlation coefficient r with probability p lying to its right.

TABLE E Critical values of the correlation r

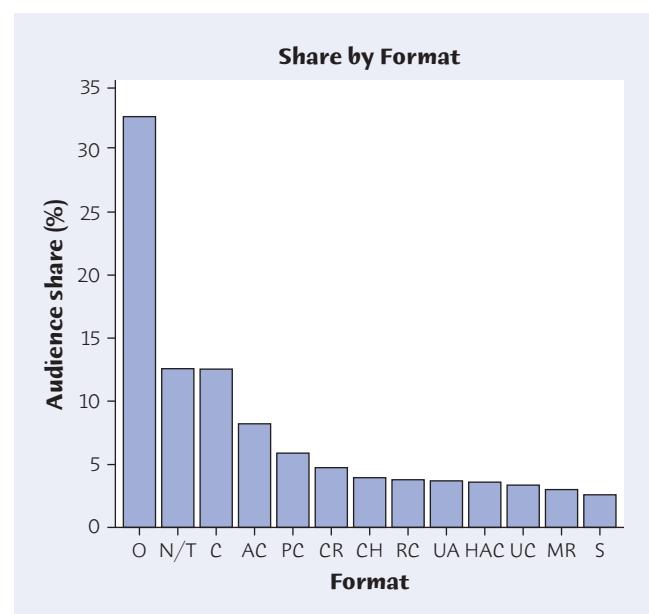
n	UPPER TAIL PROBABILITY p									
	.20	.10	.05	.025	.02	.01	.005	.0025	.001	.0005
3	0.8090	0.9511	0.9877	0.9969	0.9980	0.9995	0.9999	1.0000	1.0000	1.0000
4	0.6000	0.8000	0.9000	0.9500	0.9600	0.9800	0.9900	0.9950	0.9980	0.9990
5	0.4919	0.6870	0.8054	0.8783	0.8953	0.9343	0.9587	0.9740	0.9859	0.9911
6	0.4257	0.6084	0.7293	0.8114	0.8319	0.8822	0.9172	0.9417	0.9633	0.9741
7	0.3803	0.5509	0.6694	0.7545	0.7766	0.8329	0.8745	0.9056	0.9350	0.9509
8	0.3468	0.5067	0.6215	0.7067	0.7295	0.7887	0.8343	0.8697	0.9049	0.9249
9	0.3208	0.4716	0.5822	0.6664	0.6892	0.7498	0.7977	0.8359	0.8751	0.8983
10	0.2998	0.4428	0.5494	0.6319	0.6546	0.7155	0.7646	0.8046	0.8467	0.8721
11	0.2825	0.4187	0.5214	0.6021	0.6244	0.6851	0.7348	0.7759	0.8199	0.8470
12	0.2678	0.3981	0.4973	0.5760	0.5980	0.6581	0.7079	0.7496	0.7950	0.8233
13	0.2552	0.3802	0.4762	0.5529	0.5745	0.6339	0.6835	0.7255	0.7717	0.8010
14	0.2443	0.3646	0.4575	0.5324	0.5536	0.6120	0.6614	0.7034	0.7501	0.7800
15	0.2346	0.3507	0.4409	0.5140	0.5347	0.5923	0.6411	0.6831	0.7301	0.7604
16	0.2260	0.3383	0.4259	0.4973	0.5177	0.5742	0.6226	0.6643	0.7114	0.7419
17	0.2183	0.3271	0.4124	0.4821	0.5021	0.5577	0.6055	0.6470	0.6940	0.7247
18	0.2113	0.3170	0.4000	0.4683	0.4878	0.5425	0.5897	0.6308	0.6777	0.7084
19	0.2049	0.3077	0.3887	0.4555	0.4747	0.5285	0.5751	0.6158	0.6624	0.6932
20	0.1991	0.2992	0.3783	0.4438	0.4626	0.5155	0.5614	0.6018	0.6481	0.6788
21	0.1938	0.2914	0.3687	0.4329	0.4513	0.5034	0.5487	0.5886	0.6346	0.6652
22	0.1888	0.2841	0.3598	0.4227	0.4409	0.4921	0.5368	0.5763	0.6219	0.6524
23	0.1843	0.2774	0.3515	0.4132	0.4311	0.4815	0.5256	0.5647	0.6099	0.6402
24	0.1800	0.2711	0.3438	0.4044	0.4219	0.4716	0.5151	0.5537	0.5986	0.6287
25	0.1760	0.2653	0.3365	0.3961	0.4133	0.4622	0.5052	0.5434	0.5879	0.6178
26	0.1723	0.2598	0.3297	0.3882	0.4052	0.4534	0.4958	0.5336	0.5776	0.6074
27	0.1688	0.2546	0.3233	0.3809	0.3976	0.4451	0.4869	0.5243	0.5679	0.5974
28	0.1655	0.2497	0.3172	0.3739	0.3904	0.4372	0.4785	0.5154	0.5587	0.5880
29	0.1624	0.2451	0.3115	0.3673	0.3835	0.4297	0.4705	0.5070	0.5499	0.5790
30	0.1594	0.2407	0.3061	0.3610	0.3770	0.4226	0.4629	0.4990	0.5415	0.5703
40	0.1368	0.2070	0.2638	0.3120	0.3261	0.3665	0.4026	0.4353	0.4741	0.5007
50	0.1217	0.1843	0.2353	0.2787	0.2915	0.3281	0.3610	0.3909	0.4267	0.4514
60	0.1106	0.1678	0.2144	0.2542	0.2659	0.2997	0.3301	0.3578	0.3912	0.4143
80	0.0954	0.1448	0.1852	0.2199	0.2301	0.2597	0.2864	0.3109	0.3405	0.3611
100	0.0851	0.1292	0.1654	0.1966	0.2058	0.2324	0.2565	0.2786	0.3054	0.3242
1000	0.0266	0.0406	0.0520	0.0620	0.0650	0.0736	0.0814	0.0887	0.0976	0.1039

Answers to Selected Exercises

Chapter 1 Picturing Distributions with Graphs

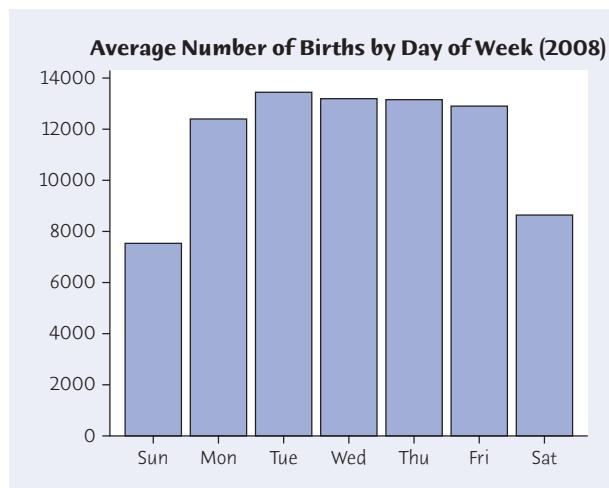
1.1: (a) Car makes and models. (b) Vehicle type (categorical), Transmission type (categorical), Number of cylinders (usually quantitative), City mpg (quantitative), Highway mpg (quantitative), Carbon footprint (tons, quantitative).

1.3: (a) 67.3%; $100\% - 67.3\% = 32.7\%$ of the radio audience listens to stations with other formats. (b)



(c) No, the shares do not sum to 100%. If you include a wedge for “other,” a pie chart would be reasonable.

1.5: A pie chart would make it more difficult to distinguish between the weekend days and the weekdays. Some births are scheduled (induced labor, for example), and probably most are scheduled for weekdays.



1.7: Use the applet to answer these questions.

1.9: (a) D.C. is the center of federal government and has many, many young professionals, many of whom may not be married. (b) Between 26 and 28. The spread is between 20 and 54, but virtually all are between 20 and 34. Again, D.C. is an outlier.

1.11: Data are rounded to units of hundreds. Stems are thousands and are split.

0	1	1	2	3
0	7	7	8	8
1	0	3		
1	7			
2	3			
2	7	7	7	8
3	0	3	4	4
3	5	5	6	7
4	4			
4	8			
5				
5				
6				
6				
7	3			

Distribution is somewhat right-skewed, with a single high outlier (United States) and two clusters of countries. Center is around 25 (\$2500 spent per capita), ignoring the outlier. Spread is from 1 (\$100 spent per capita) to 73 (\$7300 spent per capita).

1.13: (a)

1.15: (b)

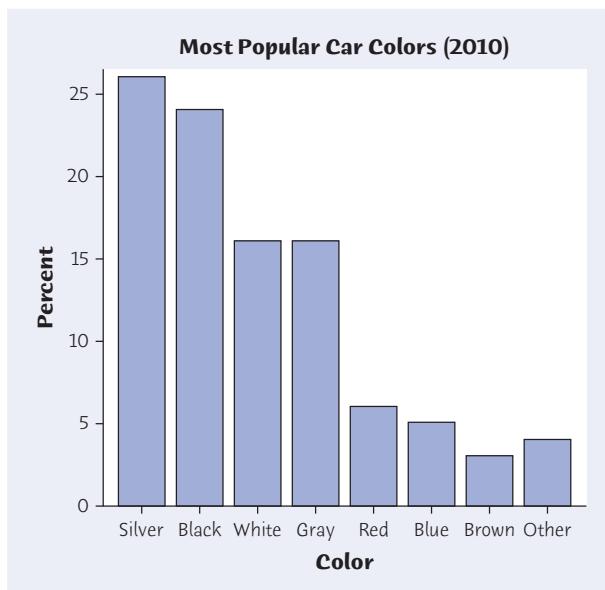
1.17: (b)

1.19: (c)

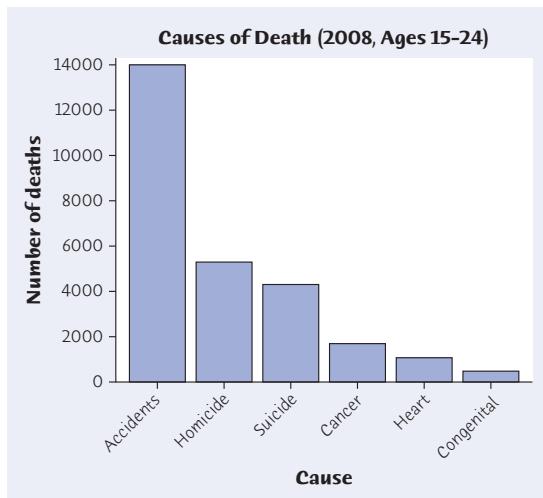
1.21: (b) Take the 26th ordered value.

1.23: (a) Students who have finished medical school. (b) 6; “Name,” “Age,” and “USMLE” are quantitative. The others are categorical.

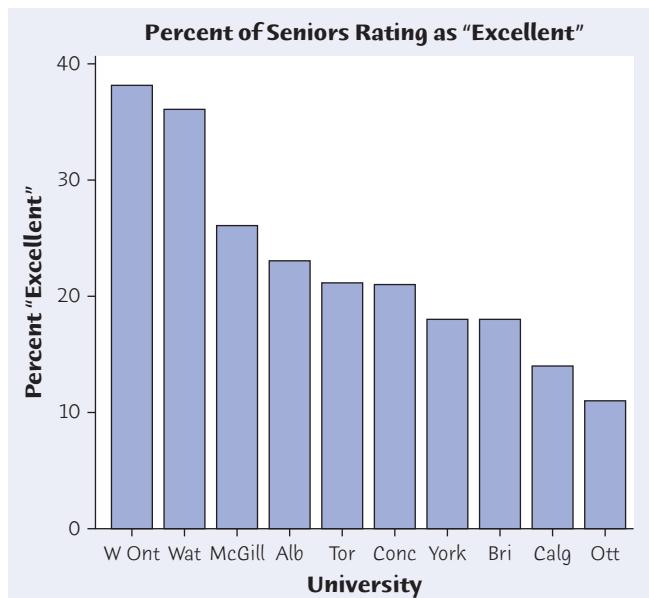
1.25: “Other colors” should account for 4%.



1.27: (a) See below. (b) You would need to know the total number of deaths in this age group or the number of deaths due to “other” causes.



1.29: (a) See below. (b) The percentages don't sum to 100%.

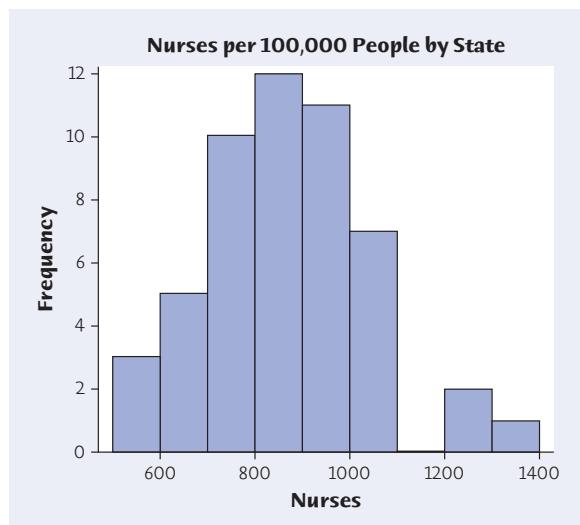


1.31: (a) Ignoring the outliers, the distribution is roughly symmetric, centered at about 110, with spread 86 to 136.
(b) $64/78 = 82.1\%$.

1.33:

1. Are you male or female → Histogram (c).
2. Are you right-handed or left-handed → Histogram (b).
3. Heights → Histogram (d).
4. Time spent studying → Histogram (a).

1.35: (a) States vary in population, so you would expect more nurses in California than in New Hampshire, for

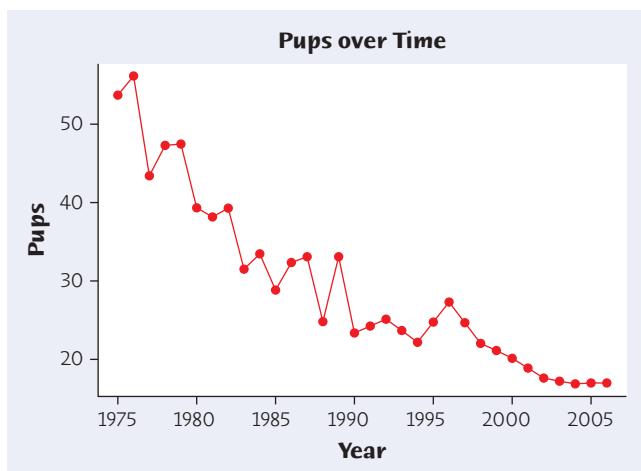


example. Nurses per 100,000 provides a better measure of how many nurses are available to serve a state's population. (b) See below. The District of Columbia, South Dakota, and Massachusetts are different from the others. Perhaps they could be considered outliers.

1.37: This is a right-skewed distribution, with center around 25 pups and spread of 17 pups to 56 pups. There were several extremely good years for pups, resulting in more than 45 births.

1	777789
2	0122344
2	555579
3	12333
3	899
4	3
4	77
5	4
5	6

1.39: The decline in population is not described by the stemplot made in Exercise 1.37.

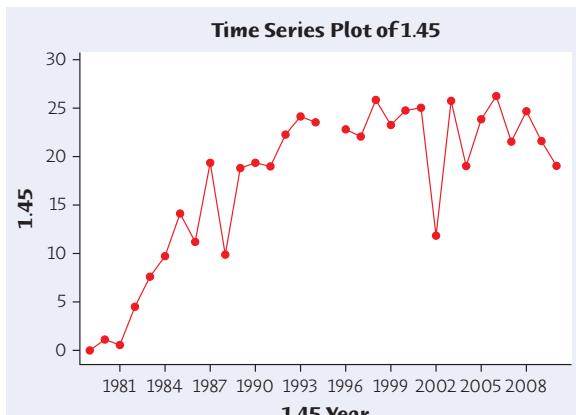


1.41: Coins with earlier (lower) dates are older and rarer. There are more coins with larger dates (newer coins) than with smaller dates (older coins).

1.43: (a) Graph (a). Vertical scaling can impact one's perception of the data. (b) In both graphs, tuition starts around \$2000 and rises to \$7700. Both plots describe the same data.

1.45: (a) See top right column. There is a trend, as well as year-to-year variability.

(b) See top right column. The midpoint is 19.3 million km². A stemplot fails to capture the relationship between size of hole and year.



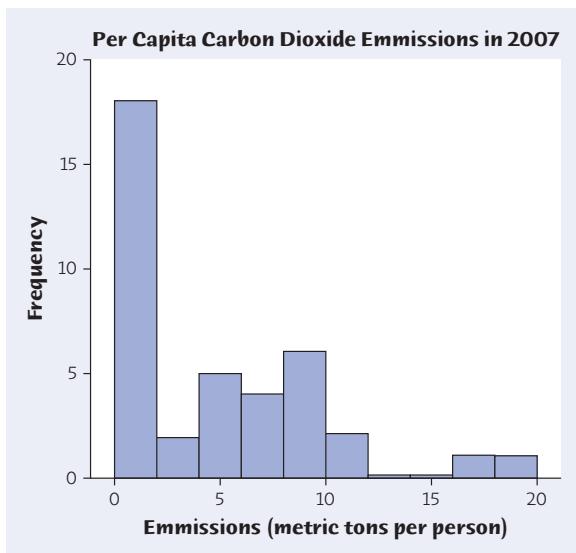
0	0004
0	79
1	0114
1	899999
2	11222333444
2	5556

Chapter 2 Describing Distributions with Numbers

2.1: Mean breaking strength = 308.4 pounds. The mean is so small relative to the data because of the sharp left skew.

2.3: Mean is 31.25 minutes. Median is 22.5 minutes. The mean is significantly larger than the median due to the right skew in the distribution of times.

2.5: The histogram shows a right skew. Hence, the mean is larger than the median. Here, the mean is 4.61 and the median is 3.95 tons per person.



2.7: (a) Minimum = 9, Q1 = 16, Median = 18, Q3 = 22, Maximum = 51. (b) The boxplot shows right skew in the distribution of MPG values.

2.9: $IQR = 22 - 16 = 6$, so $Q3 + 1.5 \times IQR = 22 + 1.5 \times 6 = 31$. Five values greater than 31 would be identified as potential outliers (33, 35, 41, 41, 51). Since $Q1 - 1.5 \times IQR = 16 - 1.5 \times 6 = 7$, there are no potential outliers below 7.

2.11: Both data sets have the same mean and standard deviation (about 7.5 and 2.0, respectively). Stemplots reveal that Data A have a very left-skewed distribution, while Data B have a slightly right-skewed distribution.

2.13: Group 1: $\bar{x} = 23.7500$, $s = 5.06548$. Group 2: $\bar{x} = 14.0833$, $s = 4.98102$. Group 3: $\bar{x} = 15.7778$, $s = 5.76146$.

2.15: (b)

2.17: (b)

2.19: (b)

2.21: (c)

2.23: (b)

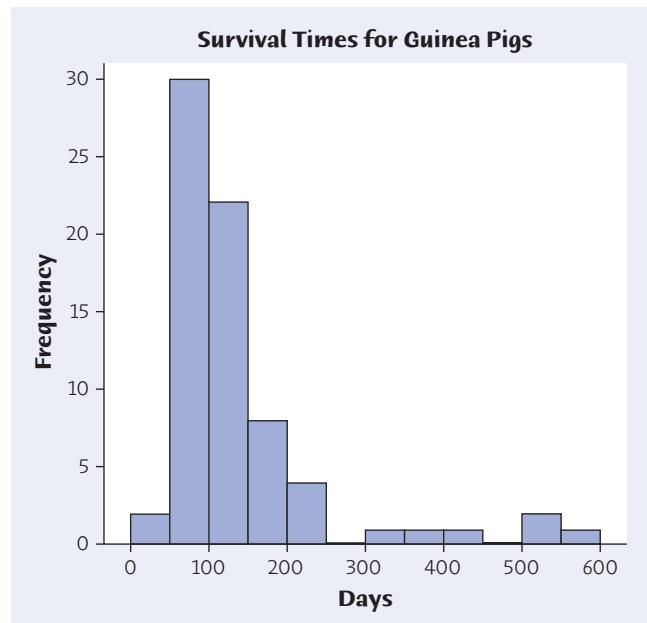
2.25: The distribution is almost certainly right-skewed, so the mean is \$58,762 and the median is \$46,931.

2.27: With 842 colleges (an even number), the median location is $(842 + 1)/2 = 421.5$, so the median is computed by averaging the 421st and 422nd endowments sizes. The 1st quartile, Q1, is found by taking the median of the first 421 endowments (when sorted). This would be the $(421 + 1)/2 = 211$ th endowment. Similarly, Q3 is found as the 632nd endowment (211 endowments above the median).

2.29: Box plots don't add much information not already present in the stemplots.

	Minimum	Q1	Median	Q3	Maximum
Bhai	46.34	46.71	47.12	48.25	50.26
Red	37.4	38.07	39.16	41.69	43.09
Yellow	34.57	35.45	36.11	36.82	38.13

2.31: (a) See top right column. The distribution is strongly right-skewed, with center around 100 days and spread 0 to 600 days. (b) Because of the extreme right skew, we should use the 5-number summary: 43, 82.5, 102.5, 151.5, 598 days. Notice that the median is closer to Q1 than to Q3.



2.33: (a) Symmetric distributions. (b) Removing the outliers reduces both means and both standard deviations.

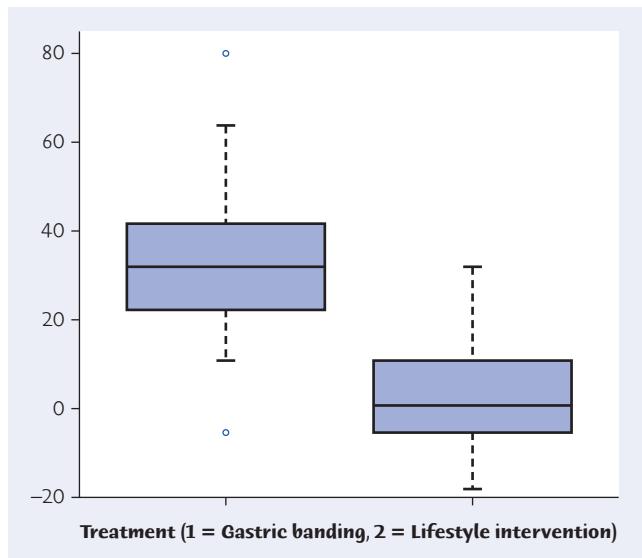
2.35: (a) The 6th observation must be placed at median for the original 5 observations. (b) No matter where you put the 7th observation, the median is one of the two repeated values above.

2.37: The mean is 8.4%, far from the national percentage of 12.5%. You can't average averages. Some states, like California and Florida, are larger and should carry more weight in the national percentage.

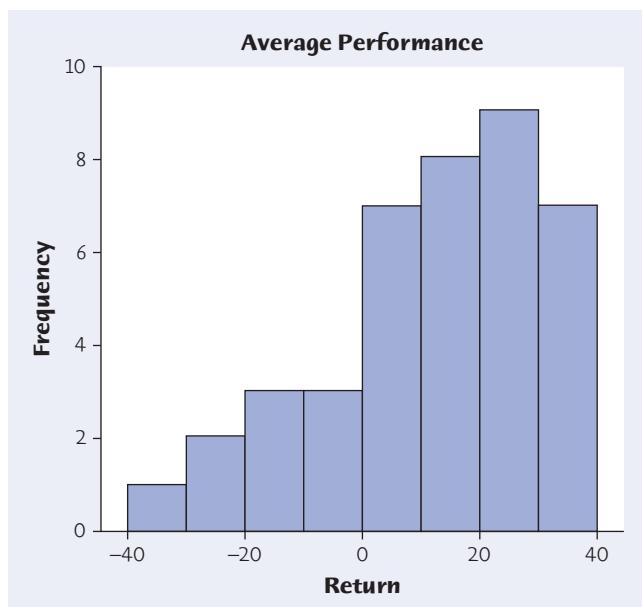
2.39: (a) Pick any four numbers all the same: e.g., (4,4,4,4) or (6,6,6,6). (b) (0,0,10,10). (c) There is more than one possible answer for (a) but not for (b).

2.41: Answers will vary. Start by insuring that the median is 7, by "locking" 7 as the 3rd smallest value. Then, adjust the minimum or maximum accordingly to acquire a mean of 10 (so they sum to 50). One solution: 5 6 7 8 24.

2.43: (a) Negative weight losses are weight gains. (b) See top page 686. Gastric banding seems to produce higher weight losses, typically. (c) It's better to measure weight loss relative to initial weight. (d) If the subjects that dropped out had continued, the difference between these groups would be as great or greater because many of the "lifestyle" dropouts had negative weight losses (i.e., weight gains), which would pull that group down.



2.45: The distribution of average returns is left-skewed. Most years, average return is positive. Returns range from about -40% to 40% , with the median return about 16% .



2.47: Based on side-by-side boxplots, lean people spend relatively more time active, but there is little difference in the time these groups spend lying down.

2.49: The distribution is right-skewed. The median salary was \$300, and the middle half of salaries were between \$167.50 and \$450. A handful of Canadians made \$1000 or more. One earned \$2200.

2.51: (a) Min = 0.0272. Q1 = 0.6449. Median = 3.954. Q3 = 8.1555. Max = 18.9144. Notice that the maximum is farther from Q3 than the minimum is from Q1. This suggests right skew. (b) IQR = $8.1555 - 0.6449 = 7.5106$. Hence, $1.5 \times \text{IQR} = 11.2659$. Now $\text{Q1} - 1.5 \times \text{IQR} = 0.6449 - 11.2659 < 0$, so no values are more than $1.5 \times \text{IQR}$ s below Q1. Also, $\text{Q3} + 1.5 \times \text{IQR} = 8.1555 + 11.2659 = 19.4214$, so there are no high outliers. This rule is rather conservative—most people would easily call the United States's value (18.9144) a far outlier, and perhaps Canada would be considered an outlier too.

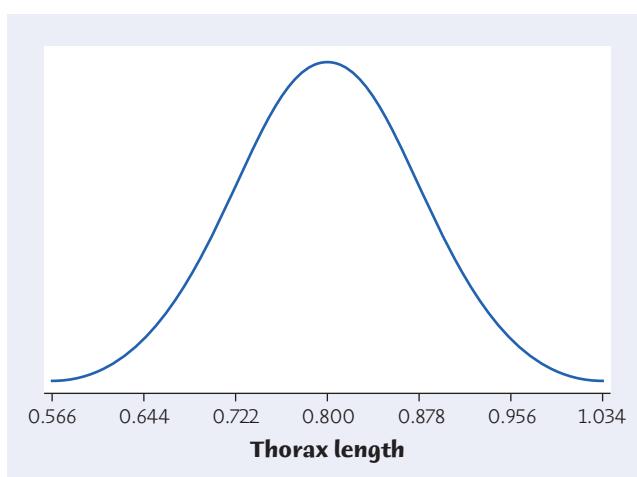
2.53: Any of the 11 incomes more than \$873.75 would be considered an outlier by the $1.5 \times \text{IQR}$ rule.

Chapter 3 The Normal Distributions

3.1: Sketches will vary.

3.3: $\mu = 2.5$. The median is also 2.5 because the distribution is symmetric.

3.5:



3.7: (a) In 95% of all years, monsoon rain levels are between two standard deviations above and below the mean: $852 \pm 2(82) = 688$ to 1016 mm. (b) The driest 2.5% of monsoon rainfalls are less than 688 mm.

3.9: A woman 6 feet tall has standardized score $z = \frac{72 - 64.3}{2.7} = 2.85$ (quite tall, relatively). A man 6 feet tall has standardized score $z = \frac{72 - 69.9}{3.1} = 0.68$.

3.11: Let x be the monsoon rainfall in a given year. (a) $x \leq 697$ mm corresponds to $z \leq \frac{697 - 852}{82} = -1.89$, for which Table A gives $0.0294 = 2.94\%$. (b) $683 < x < 1022$ corresponds to $\frac{683 - 852}{82} < z < \frac{1022 - 852}{82}$, or $-2.06 < z < 2.07$. This proportion is $0.9808 - 0.0197 = 0.9611 = 96.11\%$.

3.13: (a) Using Table A, looking for an area as close as possible to 0.1500, we find $z = -1.04$ (software gives $z = -1.0364$). (b) Now we want the value such that the proportion above is 0.70. This means that we want a proportion of 0.30 below. Using Table A, looking for an area as close to 0.3000 as possible, we find this value has $z = -0.52$ (software gives $z = -0.5244$).

3.15: (b) Income distributions are typically skewed to the right. Also, in a forest, there are likely to be many more relatively short trees than there are relatively tall trees. Although the distribution of home prices in a very large metropolitan area tends to be right-skewed, perhaps in a suburb, where the houses tend to be similar, the distribution is more symmetric.

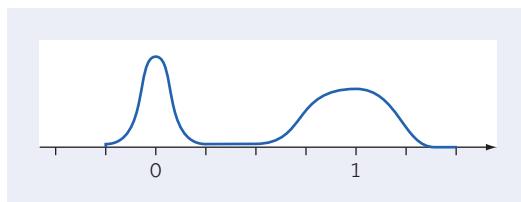
3.17: (b)

3.19: (b) $266 \pm 2(16) = 234$ to 298 days.

3.21: (b) $z = \frac{127 - 100}{15} = 1.80$.

3.23: (a)

3.25: Sketches will vary but should be some variation on the one shown here: the peak at 0 should be “tall and skinny,” while near 1, the curve should be “short and fat.”



3.27: 70 is two standard deviations below the mean (that is, it has standard score $z = -2$), so about 2.5% (half of the outer 5%) of adults would have WAIS scores below 70.

3.29: (a) We want the proportion less than z to be 0.60, so $z = 0.25$. (Software gives $z = 0.2533$.) (b) If 15% are more

than z , then 85% are less than or equal to z . Hence, $z = 1.04$. (Software gives $z = 1.0364$.)

3.31: $x < 5.0$ corresponds to $z < \frac{5.0 - 5.43}{0.54} = -0.80$, for which Table A gives 0.2119.

3.33: $0.8720 < x < 0.8780$ corresponds to $\frac{0.8720 - 0.8750}{0.0012} < z < \frac{0.8780 - 0.8750}{0.0012}$, or $-2.50 < z < 2.50$, for which Table A gives $0.9938 - 0.0062 = 0.9876$.

For problems 3.35 and 3.37, let x denote the gas mileage of a randomly selected vehicle type from the population of 2010 model vehicles (excluding the high-mileage outliers, as mentioned).

3.35: Cars with better mileage than the Camaro correspond to $x > 19$, which corresponds to $z > \frac{19 - 20.3}{4.3} = -0.30$. This proportion is $1 - 0.3821 = 0.6179$, or 61.79%.

3.37: The first and third quartiles have $z = -0.67$ and $z = 0.67$, respectively. The first quartile is $20.3 - (0.67)(4.3) = 17.42$ mpg, and the third quartile is $20.3 + (0.67)(4.3) = 23.18$ mpg.

3.39: Let x be the MCAT score for a randomly selected student that took it. The event $x < 32$ corresponds to $z < \frac{32 - 25.0}{6.4} = 1.09$. Hence, 0.8621 is the corresponding proportion, or 86.21%. William’s MCAT score is the 86.21 percentile.

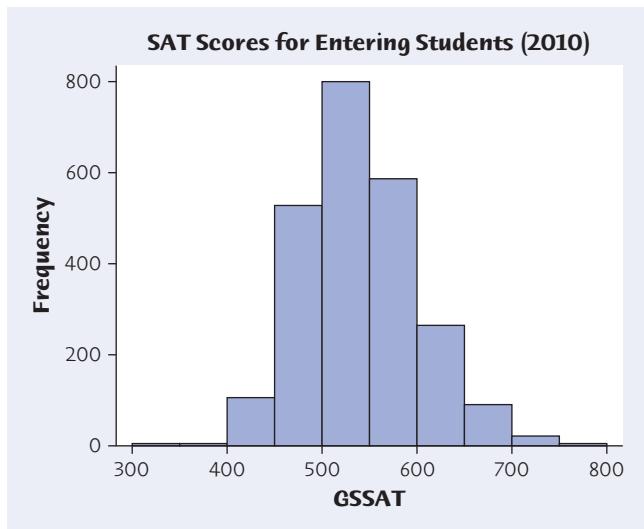
3.41: If x is the height of a randomly selected woman in this age group, we want the proportion corresponding to $x > 69.9$ inches. This corresponds to $z > \frac{69.9 - 64.3}{2.7} = 2.07$, which has proportion $1 - 0.9808 = 0.0192$, or 1.92%.

3.43: (a) Let x be a randomly selected man’s SAT math score. $x > 750$ corresponds to $z > \frac{750 - 534}{118} = 1.83$. The proportion is $1 - 0.9664 = 0.0336$. (b) Let x be a randomly selected woman’s SAT math score. $x > 750$ corresponds to $z > \frac{750 - 500}{112} = 2.23$. The proportion is $1 - 0.9871 = 0.0129$.

3.45: (a) About 0.6% of healthy young adults have osteoporosis ($z = -2.5$ gives 0.0062). (b) About 31% of this population of older women have osteoporosis ($z = -0.5$ gives 0.3085).

3.47: (a) $145,000/1,568,835 = 0.0924$, or 9.24%. (b) $50,860 + 145,000 = 195,860$ students have ACT score 28 or higher. This is $195,860/1,568,835 = 0.1248$, or 12.48%. (c) $z > \frac{28 - 21.0}{5.2} = 1.35$, so the corresponding proportion is $1 - 0.9115 = 0.0885$, or 8.85%.

3.49: (a) A histogram (page 688) appears to be roughly symmetric with no outliers. (b) Mean = 544.42. Median = 540. Standard deviation = 61.24. Q1 = 500. Q3 = 580. The mean and median are close, and the distances of each quartile to the median are equal, consistent with a Normal distribution. (c) $z > \frac{501 - 544.42}{61.24} = -0.71$, or $1 - 0.2389 = 0.7611$, or 76.11%. (d) In fact, 1776 entering GSU



students scored higher than 501, which represents $1776/2417 = 0.7348$, or 73.48%.

3.51: (a) The 65 Canadians with earnings greater than \$375 represent $65/200 = 0.325$, or 32.5%. $z > \frac{375 - 350.30}{292.20} = 0.08$, which has proportion $1 - 0.5319 = 0.4681$, or 46.81% above. (b) $z < \frac{0 - 350.30}{292.20} = -1.20$, or 0.1151, or 11.51%. (c) The Normal distribution model predicts 11.5% of Canadians to earn less than \$0, while (of course) none do. This is a substantial error, since the Normal model predicts 11.5% of values more than 375, where we actually observed 32.5% more than 375. The standard deviation (\$292.20) is large relative to the average (\$350.30), which suggests a strong right skew in the distribution, given that no values can be negative. In this application, the data seem to be far from Normal.

3.53: Because the quartiles of any distribution have 50% of observations between them, we seek to place the flags so that the reported area is 0.5. The closest the applet gets is an area of 0.5034, between -0.680 and 0.680 . Thus the quartiles of any Normal distribution are about 0.68 standard deviations above and below the mean.

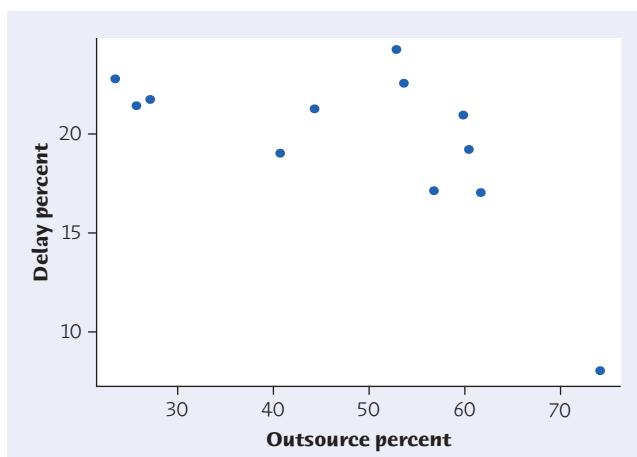
Note: Table A places the quartiles at about 0.67; other statistical software gives ± 0.6745 .

Chapter 4 Scatterplots and Correlation

4.1: (a) Explanatory: time spent studying; response: grade. (b) Explore the relationship. (c) Explanatory: time spent online using Facebook; response: GPA. (d) Explore the relationship.

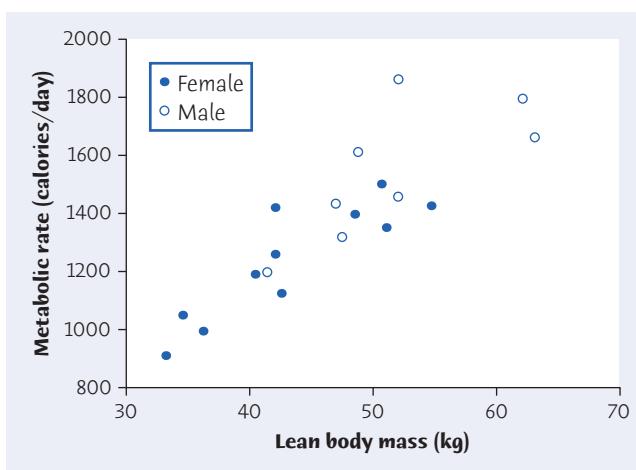
4.3: For example: weight, sex, other food eaten by the students, type of beer (light, imported, . . .).

4.5: Outsource percent is the explanatory variable and should be on the horizontal axis. Delay percent is the response and should be on the vertical axis.



4.7: There is an outlier (Hawaiian Airlines). Removing it, we would see no association between these variables. Without removing it, there is a very weak, negative association between the variables (which contradicts the suspicions described in Exercise 4.5).

4.9: (a) Women are marked with filled circles, men with open circles. (b) For both men and women, the association is linear and positive. The women's points show a stronger association. As a group, males typically have larger values for

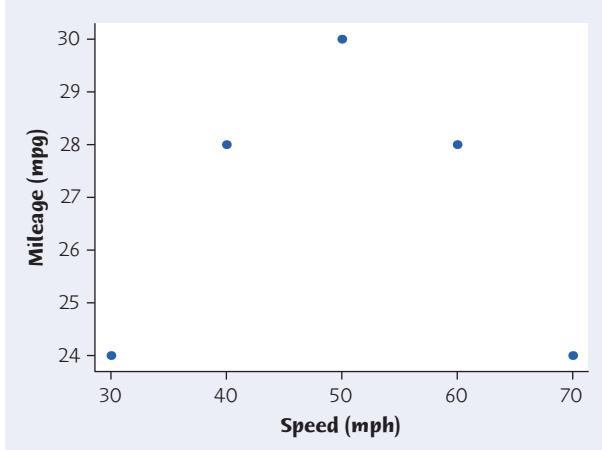


both variables (they tend to have more mass and tend to burn more calories per day).

4.11: r would not change; units do not affect correlation.

4.13: $\bar{x} = 50$ mph, $s_x = 15.8114$ mph, $\bar{y} = 26.8$ mpg, and $s_y = 2.6833$ mpg. Refer to the table of standardized scores below, then note that $r = 0/4 = 0$. The correlation is zero because these variables do not have a straight-line relationship; the association is neither positive nor negative.

z_x	z_y	$z_x z_y$
-1.2649	-1.0435	1.3199
-0.6325	0.4472	-0.2828
0	1.1926	0
0.6325	0.4472	0.2828
1.2649	-1.0435	-1.3199
0		



4.15: (a) The association should be positive (e.g., if oil prices rise, so do gas prices).

4.17: (a) 0.9. Without the outlier, there is a strong positive linear relationship.

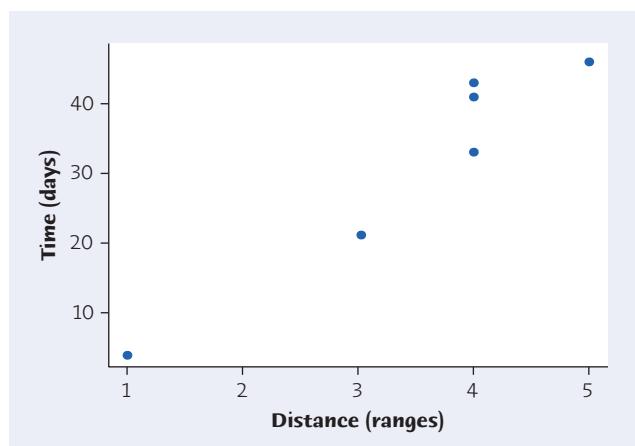
4.19: (c) A correlation close to 0 might arise from a scatterplot with no visible pattern, but there could be a nonlinear pattern.

4.21: (a)

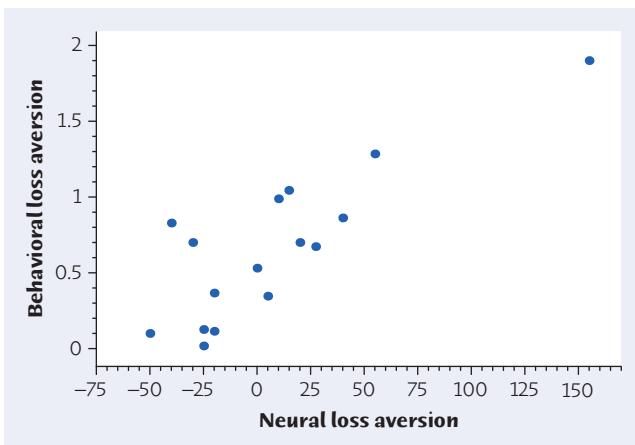
4.23: (b)

4.25: (a) There is a slightly negative association between these variables. (b) There is general disagreement. (c) It does not appear that any of the values are obviously outside the general pattern. Perhaps one value (Rank = 8, BRFSS = 0.30) is an outlier, but this is hard to say.

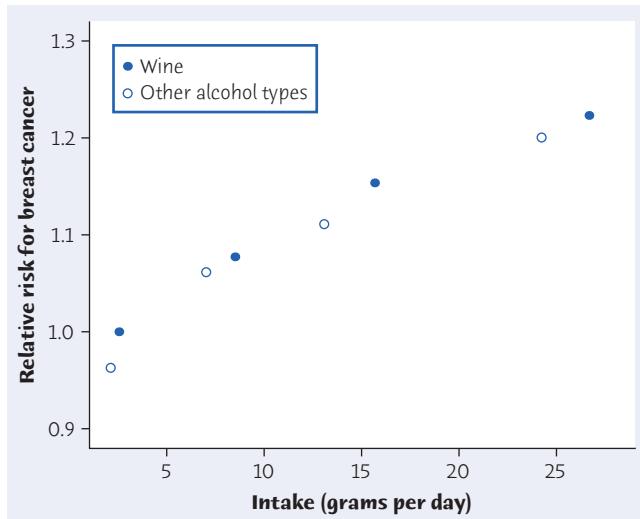
4.27: (a) The scatterplot suggests a strong positive linear association between distance and time with respect to the spread of Ebola. (b) $r = 0.9623$. (c) Correlation would not change, since it does not depend on units.



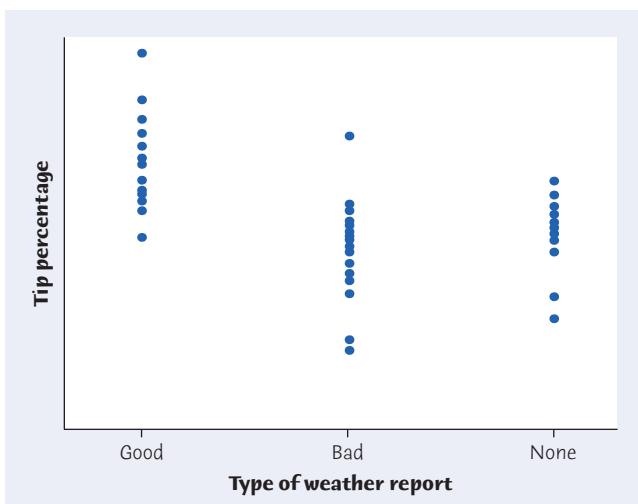
4.29: (a) See below. Neural activity is explanatory (and so should be on the horizontal axis). (b) The association is moderately strong, positive, and linear. The outlier is in the upper right corner. (c) For all points, $r = 0.8486$. Without the outlier, $r = 0.7015$. The correlation is greater with the outlier because it fits the pattern of the other points.



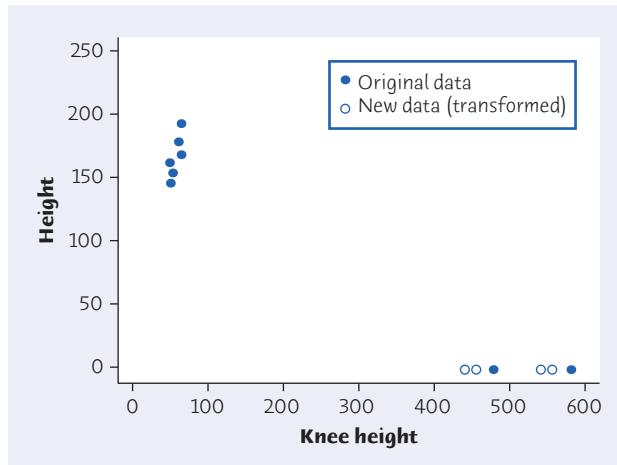
4.31: (a) See page 690. (b) There is a strong relationship between alcohol intake and relative risk of breast cancer (again, this is an observational study, so no causal relationship is established here). It seems that type of alcohol has nothing to do with the increase, since the same pattern and rate of increase is seen for both groups.



4.33: (a) A plot follows that suggests that “Good” weather reports tend to yield higher tips. (b) The explanatory variable is categorical, not quantitative, so r cannot be used. Note that we can arrange the categories any way, and these different arrangements would suggest different associations.



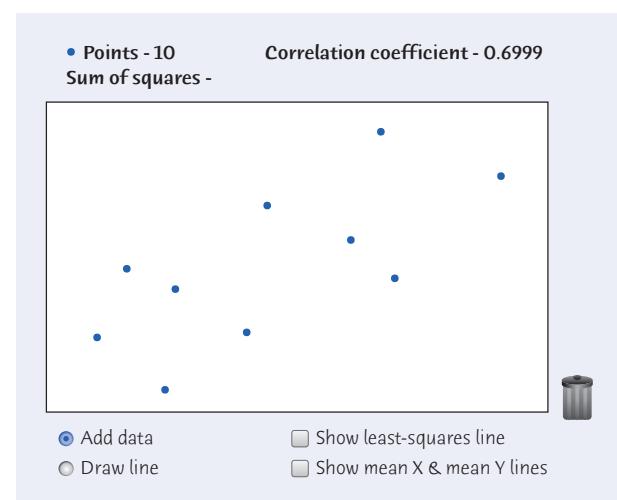
4.35: (a) See top right column. Changing the units has a dramatic impact on the plot. (b) Nevertheless, units do not impact correlation. For both data sets, $r = 0.8900$.

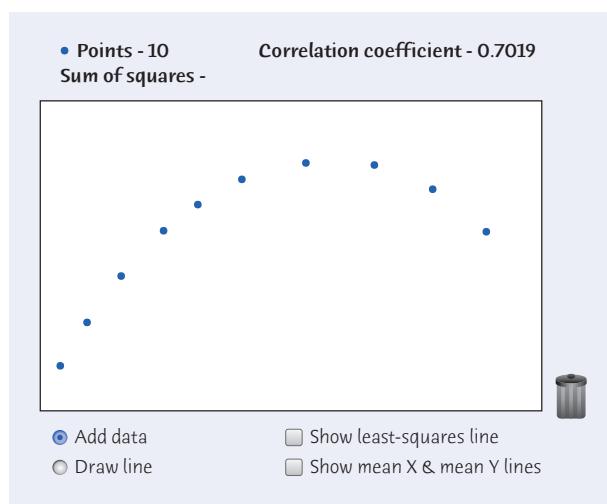
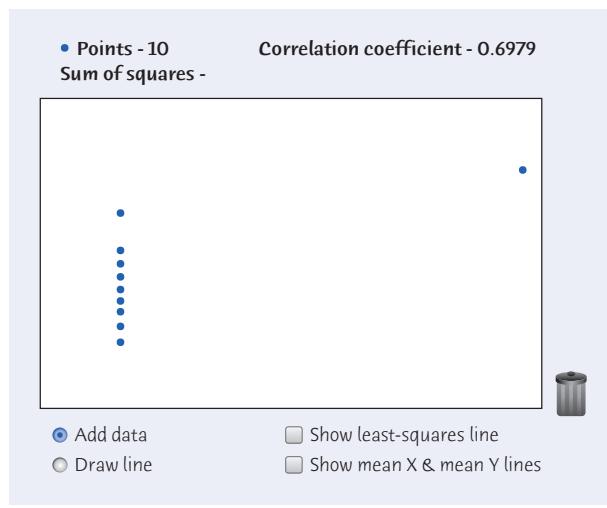


4.37: (a) Small-cap stocks have a lower correlation with municipal bonds, so the relationship is weaker. (b) She should look for a negative correlation.

4.39: (a) Because sex has a nominal scale, we cannot compute the correlation between sex and any other variable. Some writers and speakers use “correlation” as a synonym for “association,” but this is not correct. (b) $r = 1.09$ is impossible, because r is restricted to be between -1 and 1 . (c) Correlation has no units.

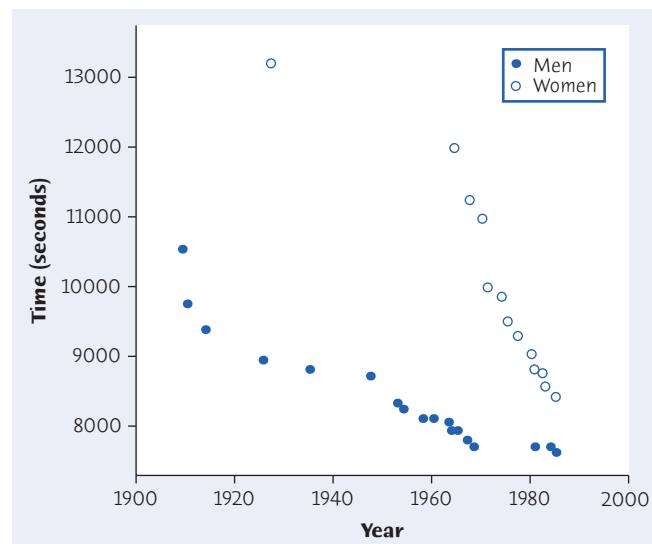
4.41: (a) Because two points determine a line, the correlation is always 1. (b) Sketches will vary; an example is shown in the graph below. Note that the scatterplot must be positively sloped, but r is affected only by the scatter about the line, not by the steepness of the slope of that line. (c) The first nine points cannot be spread from the top to the



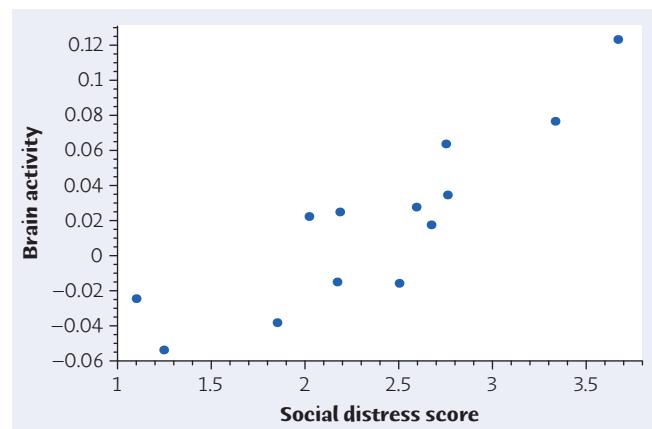


bottom of the graph because in such a case the correlation cannot exceed about 0.66 (this is based on experience – lots of playing around with the applet). One possibility is shown above, first-plot. (d) To have $r = 0.7$, the curve must be higher at the right than at the left. One possibility is shown above, second-plot.

4.43: We will not use correlation, but we will examine the plot to see if women are beginning to outrun men. By inspection, one might guess that the “lines” that fit these data sets will meet around 1998. This is how the researchers made this leap. Men’s and women’s times have, indeed, grown closer over time. Both sexes have improved their record marathon times over the years, but women’s times have improved at a faster rate. However, plots like this are designed to lure the reader into extrapolating—and extrapolation is never a good idea.



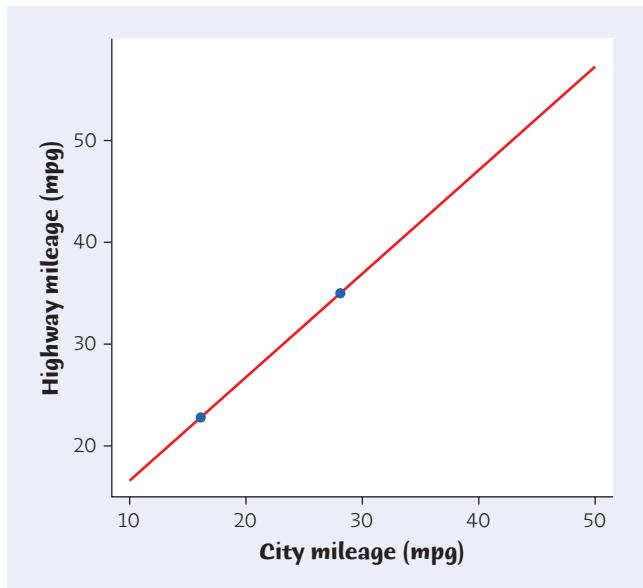
4.45: A scatterplot shows a fairly strong, positive, linear association. There are no particular outliers; each variable has low and high values, but those points do not deviate from the pattern of the rest. $r = 0.8782$. Social exclusion does appear to trigger a pain response: higher social distress measurements are associated with increased activity in the pain-sensing area of the brain. However, no cause-and-effect conclusion is possible since this was not a designed experiment.



Chapter 5 Regression

5.1: (a) Slope = 1.016. On average, highway mileage increases by 1.016 mpg for each additional 1 mpg change in city mileage. (b) Intercept = 6.554 mpg. This is the highway mileage for a nonexistent car that gets 0 mpg in the city.

(c) For a car that gets 16 mpg in the city, $6.554 + (1.016)(16) = 22.81$ mpg. For a car that gets 28 mpg in the city, $6.554 + (1.016)(28) = 35.002$ mpg. (d)



5.3: (a) $\bar{x} = 30.280$, $s_x = 0.4296$, $\bar{y} = 2.4557$, $s_y = 0.1579$, and $r = -0.8914$. Hence, $b = r \frac{s_y}{s_x} = (-0.8914) \frac{0.1579}{0.4296} = -0.3276$; $a = \bar{y} - b\bar{x} = 2.4557 - (-0.3275)(30.280) = 12.3754$. (b) Software agrees with these values to 3 decimal places, since we rounded to the 4th decimal place.

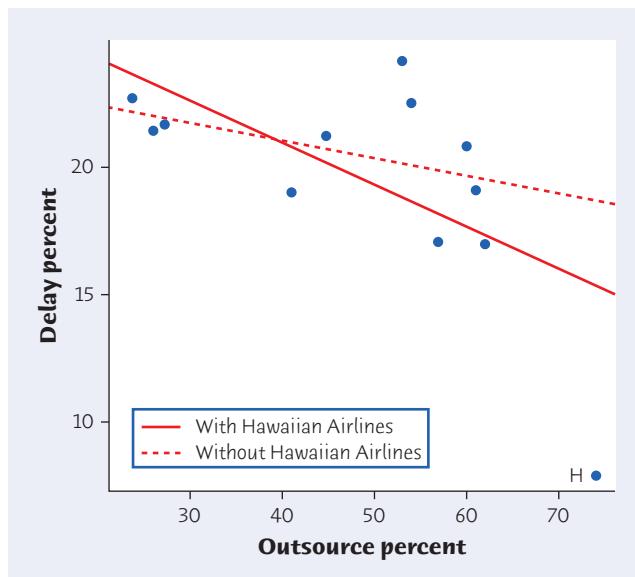
5.5: The farther r is from 0 (in either direction), the stronger the linear relationship between two variables.

5.7: (a) The residuals are computed in the table below using $\hat{y} = 12.3754 - 0.3276x$. (b) They sum to zero, except for rounding error. (c) From software, the correlation between x and $y - \hat{y}$ is 0.000025, which is zero except for rounding.

x	y	\hat{y}	$y - \hat{y}$
29.68	2.63	2.652	-0.022
29.87	2.58	2.590	-0.010
30.16	2.60	2.495	0.105
30.22	2.48	2.475	0.005
30.48	2.26	2.390	-0.130
30.65	2.38	2.335	0.045
30.90	2.26	2.253	0.007
			0

5.9: (a) Any point that falls exactly on the regression line will not increase the sum of squared vertical distances (which the regression line minimizes). Thus the regression line does not change. (b) Influential points are those whose x coordinates are outliers.

5.11: (a) In the plot, the outlier (Hawaiian Airlines) is the point identified with "H." This point falls outside the linear trend suggested by the other data points, so it is influential and will affect the regression line by "pulling" it. (b) With the outlier, $r = -0.624$. Without the outlier, $r = -0.441$. (c) The two regression lines (one including the outlier, the other without) are plotted. The regression line based on the complete (original) data set, including the outlier, is $\hat{y} = 27.486 - 0.164x$. Using this, when $x = 74.1$, we predict 15.33% delays. The other regression line (fit without the outlier) is $\hat{y} = 23.804 - 0.069x$, so our prediction would be 18.69% delays. The outlier impacts predictions because it impacts the regression line.



5.13: (a) $\hat{y} = -43.172 + 0.129x$. (b) If 975,000 boats are registered, then $x = 975$, and $\hat{y} = -43.172 + (0.129)(975) = 82.6$ manatees killed. This is not extrapolation. (c) If $x = 0$ (corresponding to no registered boats), then we would "predict" -43.172 manatees to be killed by boats. This illustrates the folly of extrapolation. $x = 0$ is well outside the range of observed values of x on which the regression line was based.

5.15: Possible lurking variables include the IQ and socio-economic status of the mother, as well as the mother's other habits (drinking, diet, etc.). These variables are associated with smoking in various ways, and are also predictive of a child's IQ.

5.17: Age is probably the most important lurking variable: married men would generally be older than single men, so they would have been in the workforce longer and therefore had more time to advance in their careers.

5.19: (b) Consider two points on the regression line, say $(90, 4)$ and $(130, 11)$. The slope of the line segment connecting these points is $\frac{11 - 4}{130 - 90} = 7/40$, which is close to 0.2.

5.21: (a)

5.23: (c)

5.25: (a) The slope of the line is positive.

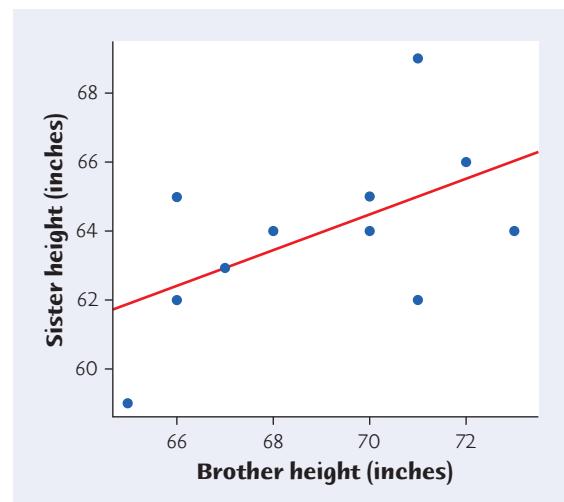
5.27: (a)

5.29: (a) Increasing the size of a diamond by an additional carat increases its price by 3721.02 Singapore dollars. (b) A diamond of size 0 carats would have a predicted price of 259.63 Singapore dollars. This is probably an extrapolation, since the data set on which the line was constructed almost certainly had no rings with diamonds of size 0 carats.

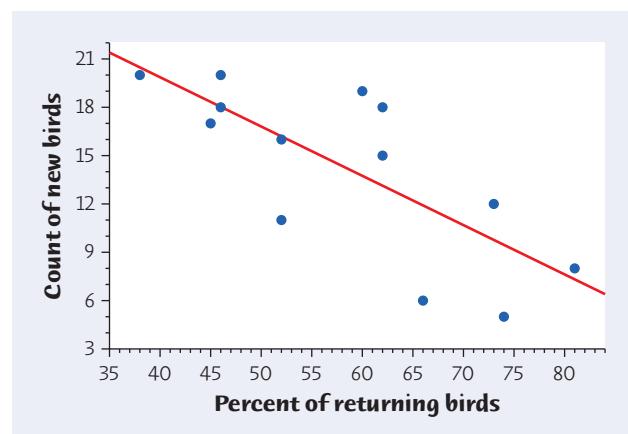
5.31: (a) $\hat{y} = 0.919 + 2.0647x$. At 25 degrees Celsius, we predict beak heat loss of $\hat{y} = 0.919 + (2.0647)(25) = 52.34$ percent. (b) Since $r^2 = 0.836$, 83.6% of the total variation in beak heat loss is explained by the straight-line relationship with temperature. (c) $r = \sqrt{r^2} = \sqrt{0.836} = 0.914$. Correlation is positive here, since the least-squares regression line has a positive slope.

5.33: The x -values will be pre-exam scores, and the y -values will be final exam scores. This is probably not the Princeton University, Nobel prize-winning economist Paul Krugman. (a) $b = 0.5 \times \frac{8}{40} = 0.1$, and $a = 75 - (0.1)(280) = 47$. The regression equation is $\hat{y} = 47 + 0.1x$. (b) $\hat{y} = 47 + (0.1)(300) = 77$. (c) Julie is right. With a correlation of $r = 0.5$, $r^2 = (0.5)^2 = 0.25$, so the regression line accounts for only 25% of the variability in student final exam scores.

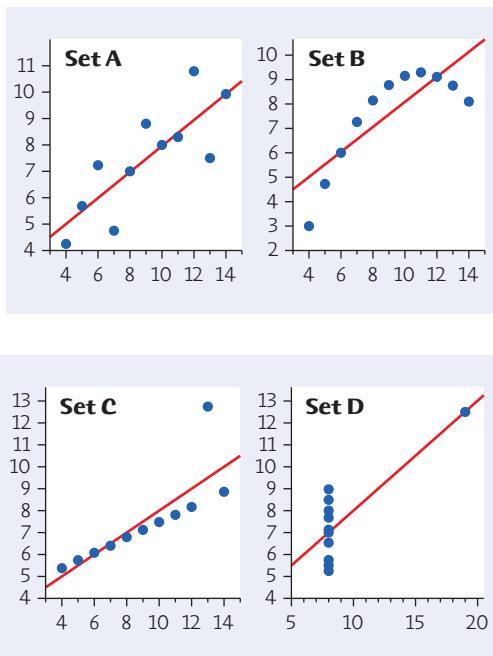
5.35: (a) $\hat{y} = 28.037 + 0.521x$. $r = 0.555$. (b) We predict Tonya to have height $\hat{y} = 28.037 + (0.521)(70) = 64.5$ inches (rounded). This prediction isn't expected to be very accurate because the correlation isn't very large. So $r^2 = (0.555)^2 = 0.308$. The regression line explains only 30.8% of the variation in sister heights.



5.37: (a) $\hat{y} = 31.9 - 0.304x$. (b) The slope (-0.304) tells us that, on the average, for each additional 1% increase in returning birds, the number of new birds joining the colony decreases by 0.304. (c) When $x = 60$, we predict $\hat{y} = 13.66$ new birds will join the colony.



5.39: (a) To three decimal places, the correlations are all approximately 0.816, and the regression lines are all approximately $\hat{y} = 3.000 + 0.500x$. For all four sets, we predict $\hat{y} = 8$ when $x = 10$. (b) Plots top of page 694. (c) For Set A, the use of the regression line seems to be reasonable. For Set B, there is an obvious nonlinear relationship; we should fit a parabola or other curve. For Set C, the point $(13, 12.74)$ deviates from the (highly linear) pattern of the other points. For Set D, the data point with $x = 19$ is a very influential point—the other points alone give no indication of slope for the line.



5.41: (a) $\hat{y} = 42.933 + 2.453x$. (b) $\hat{y} = 0.42933 + 0.002453x$. (c) Use the fact that $50 \text{ cm} = 500 \text{ mm}$. When $x = 50 \text{ cm}$, the first regression equation gives $\hat{y} = 165.583 \text{ cm}$. Using the second equation, with $x = 500 \text{ mm}$, $\hat{y} = 1.65583 \text{ m}$. These are the same.

5.43: The correlation would be much lower, because there is much greater variation in individuals than in the averages.

5.45: Responses will vary. For example, students who choose the online course might have more self-motivation, or have better computer skills (which might be helpful in doing well in the class).

5.47: Here is a (relatively) simple example to show how this can happen: suppose that most workers are currently 30 to 50 years old; of course, some are older or younger than that, but this age group dominates. Suppose further that each worker's current salary is his/her age (in thousands of dollars); for example, a 30-year-old worker is currently making \$30,000. Over the next 10 years, all workers age, and their salaries increase. Suppose every worker's salary increases by between \$4000 and \$8000. Then every worker will be making *more* money than he/she did 10 years before, but *less* money than a worker of that same age 10 years before. During that time, a few workers will retire, and others will enter the workforce, but that large cluster that had been between the ages of 30 and 50 (now between 40 and 60) will bring up the overall median salary despite the changes in older and younger workers.

5.49: For a player who shot 80 in the first round: $\hat{y} = 52.74 + (0.297)(80) = 76.5$. For a player who shot 70 in

the first round: $\hat{y} = 52.74 + (0.297)(70) = 73.53$. Notice that the player who shot 80 the first round (worse than average) is predicted to have a worse-than-average score the second round, but better than the first round. Similarly, the player who shot 70 the first round (better than average) is predicted to do better than average in the second round, but not as well (relatively) as in the first round. Both players are predicted to "regress" to the mean.

5.51: See Exercise 4.41 for the three sample scatterplots. A regression line is appropriate only for the scatterplot of part (b).

5.53: The scatterplot shows a positive linear association. The regression line is $-1.286 + 11.89x$. The straight-line relationship explains that $r^2 = 83.9\%$ of the variation in beetle larvae. The strong positive association supports the idea that beavers benefit beetles.

5.55: There is a reasonable but not very strong linear relationship between forecasted and actual hurricanes. In the plot, the 2005 season is noted with an open circle. It is an outlier and influential, pulling the regression line somewhat. We might consider deleting this point and fitting the line again. Deleting the point, we obtain the dotted regression line, $\hat{y} = 2.6725 + 0.7884x$. If the forecasts were perfect, the intercept of this line would be 0 and the slope would be 1, for reference. Deleting the 2005 season, $r = 0.621$ and $r^2 = 38.5\%$. Hence, even after deleting the outlier, the regression line explains only 38.5% of variation in number of hurricanes.

5.57: The regression lines are:

$$\text{For men: } \hat{y} = 67,825.3 - 30.44x$$

$$\text{For women: } \hat{y} = 182,976.15 - 87.73x$$

Although the lines appear to fit the data reasonably well (and the regression line for women would fit better if we omitted the outlier associated with year 1926), this analysis is inviting you to extrapolate, which is never advisable. Using the regression lines plotted, we might expect women to outrun men by the year 2010. Omitting the outlier, the line for women would decrease more steeply, and the intersection would occur sooner, by 1995.

Chapter 6 Two-Way Tables

6.1: (a) This table describes 1808 people. $736 + 450 + 193 = 1379$ played video games. (b) The percent of boys earning A's and B's is $(736 + 205)/1808 = 0.5205 = 52.05\%$. We do this for all three grade levels. The complete marginal distribution for grades is

Grade	Percent
A's and B's	52.05%
C's	32.85%
D's and F's	15.10%

Of all boys, $32.85\% + 15.10\% = 47.95\%$ received a grade of C or lower.

6.3: There are 1379 players. Of these, $736/1379 = 53.37\%$ earned A's or B's. Similarly, there are 429 nonplayers. Of these, $205/429 = 47.79\%$ earned A's or B's. Continuing in like manner, the conditional distribution of grades for players follows:

Grades	Players	Nonplayers
A's and B's	53.37%	47.79%
C's	32.63%	33.56%
D's and F's	14.00%	18.65%

It doesn't look as if there's a big difference between these conditional distributions.

6.5: Two examples are shown. In general, choose a to be any number from 10 to 50, and then all the other entries can be determined.

30	20
30	20

50	0
10	40

6.7: (a) For Rotura district, $79/8889 = 0.0089$ or 0.9%, of Maori are in the jury pool, while $258/24,009 = 0.0107$ or 1.07%, of the non-Maori are in the jury pool. For Nelson district, the corresponding percents are 0.08% for Maori and 0.17% for non-Maori. Hence, in each district, the percent of non-Maori in the jury pool exceeds the percent of Maori in the jury pool. (b) Combining the regions into one table:

	Maori	Non-Maori
In jury pool	80	314
Not in jury pool	10,138	56,353
Total	10,218	56,667

For the Maori, overall the percent in the jury pool is $80/10,218 = 0.0078$, or 0.78%, while for the non-Maori, the overall percent in the jury pool is $314/56,667 = 0.0055$, or 0.55%. Hence, overall the Maori have a larger percent in the jury pool, but in each region they have a lower percent in the jury pool. (c) The reason for Simpson's paradox occurring with this example is that the Maori constitute a large proportion of Rotura's population, while in Nelson they are small minority community.

6.9: (b)

6.11: (a)

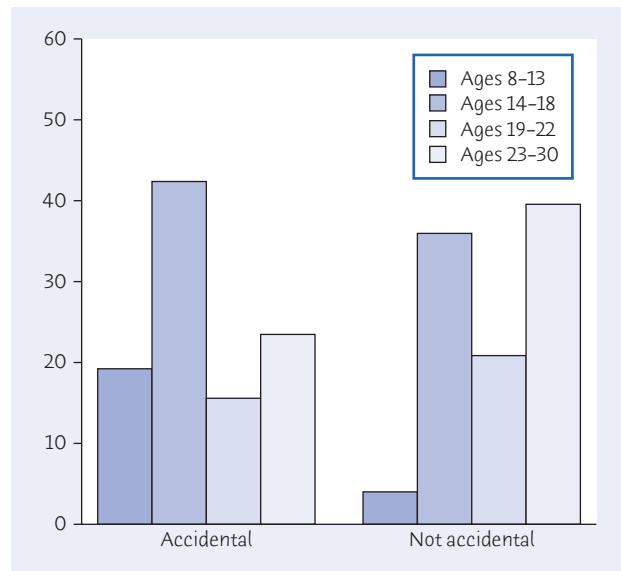
6.13: (c)

6.15: (b)

6.17: (b)

6.19: For each type of injury (accidental, not accidental), the distribution of ages is shown below.

	Accidental	Not accidental
8–13	19.0%	4.0%
14–18	42.2%	35.8%
19–22	15.4%	20.8%
23–30	23.4%	39.4%



Among accidental weight-lifting injuries, the percentage of relatively younger lifters is larger, while among the injuries that are not accidental, the percentage of relatively older lifters is larger.

6.21: The percent of single men in Grade 1 jobs is $58/337 = 0.172$, or 17.2%. The percent of Grade 1 jobs held by single men is $58/955 = 0.0607$, or 6.07%.

6.23: (a) We need to compute percents to account for the fact that the study included many more married men than single men, so we would expect their numbers to be higher in every job grade (even if marital status had no relationship with job level). (b) A table of percents is provided. Single and widowed men had higher percents of Grade 1 jobs; single men had the lowest (and widowed men the highest) percents of Grade 4 jobs.

	Single	Married	Divorced	Widowed
Grade 1	17.2%	11.3%	11.9%	19.0%
Grade 4	2.1%	6.9%	5.6%	9.5%

6.25: (a) The two-way table of race (White, Black) versus death penalty (Death penalty, no death penalty) follows.

	White defendant	Black defendant
Death penalty	19	17
No death penalty	141	149

(b) For black victims: percentage of white defendants given the death penalty is $0/9 = 0$, or 0%; percentage of black defendants given the death penalty is $6/103 = 0.058$, or 5.8%. For white victims: percentage of white defendants given the death penalty is $19/151 = 0.126$, or 12.6%; percentage of black defendants given the death penalty is $11/63 = 0.175$, or 17.5%. Hence, for both victim races, black defendants are given the death penalty relatively more often than white defendants. Overall, referring to the table in (a), $19/160 = 0.119$ or 11.9% of white defendants got the death penalty, while $17/166 = 0.102$ or 10.2% of black defendants got the death penalty. This illustrates Simpson's paradox. (c) For white defendants, $151/160 = 0.9438 = 94.4\%$ of victims were white. For black defendants, only $63/166 = 0.3795$ or 37.95% of victims were white. The death penalty was predominantly assigned to cases involving white victims: 14.0% of all cases with a white victim, while only 5.5% of all cases with a black victim had a death penalty assigned to the defendant. Hence, because most white defendants' victims are white and cases with white victims carry additional risk of a death penalty, white defendants are being assigned the death penalty more often overall.

6.27: The percentages for each column are provided in the table. For example, for Chantix, the percentage of successes (no smoking in weeks 9–12) is $155/(155 + 197) = 0.4403$, or 44.0%. Since we're comparing success rates, we'll leave off the row for "% smoking in weeks 9–12," since this is just $100\% - \%$ No smoking in weeks 9–12.

	Chantix	Bupropion	Placebo
% No smoking in weeks 9–12	44.0%	29.5%	7.7%

A larger percentage of subjects using Chantix were not smoking during weeks 9–12, compared with results for either of the other treatments.

6.29: We compute, for example, the percentage of women earning associate's degrees: $519/823 = 0.631$, or 63.1%. The table shows the percent of women at each degree level, which is all we need for comparison. Women constitute a substantial majority of associate's, bachelor's, and master's degrees, a scant majority of doctor's degrees, and slightly less than 50% of professional degrees.

Degree	% Female
Associate's	63.1
Bachelor's	57.5
Master's	61.1
Professional	49.5
Doctor's	53.3

6.31: The table provides the percent of subjects with various health outlooks for each group. The outlooks of current smokers are generally bleaker than that of current nonsmokers. Much larger percentages of nonsmokers reported being in "excellent" or "very good" health, while much larger percentages of smokers reported being in "fair" or "poor" health.

Health Outlook					
	Excellent	Very good	Good	Fair	Poor
Current smoker	6.2%	28.5%	35.9%	22.3%	7.2%
Current nonsmoker	12.4%	39.9%	33.5%	14.0%	0.3%

Chapter 7 Exploring Data: Part I Review

7.1: (c)

7.3: (c)

7.5: (b)

7.7: (a)

7.9: (d)

7.11: (b)

7.13: (b)

7.15: (a)

7.17: $P(X > 90) = P(Z > 1.81) = 1 - 0.9649 = 0.0351$, or 3.51%. (b) The middle 50% of all observations lie between the first and third quartiles, so the IQR is the range in which these observations lie. The first and third quartiles are 0.67 standard deviations above and below average. Hence, these values are $75 - 0.67(8.3) = 69.44$ and $75 + 0.67(8.3) = 80.56$ ksi. The range (IQR) in which the middle values lie is therefore $80.56 - 69.44 = 11.12$ ksi.

7.19: (a) Minimum = 7.2, $Q_1 = 8.5$, $M = 9.3$, $Q_3 = 10.9$, Maximum = 12.8. (b) $M = 27$. (c) 25% of values exceed $Q_3 = 30$. (d) Yes. Virtually all Torrey pine needles are longer than virtually all Aleppo pine needles. There is no overlap in the distributions, as seen by comparing, say, Minimum for Torrey pine needles (21) to Maximum for Aleppo pine needles (12.8).

7.21: (b)

7.23: (c)

7.25: (c)

7.27: (d)

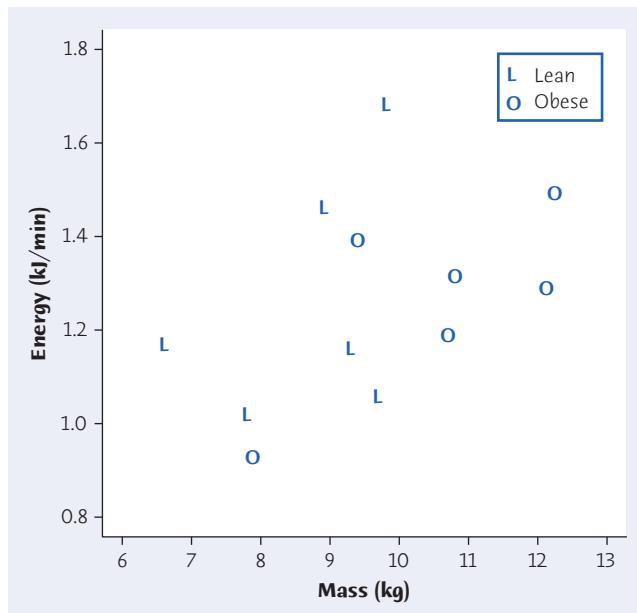
7.29: (d)

7.31: (c)

7.33: (a)

7.35: (a) No. (b) $r^2 = 0.64$, or 64%.

7.37: (a) 8.683 kg. (b) 10.517 kg. (c) Such a comparison would be unreasonable because the lean group is less massive and therefore would be expected to burn less energy on average. (d) See below. (e) It appears that the rate of increase in energy burned per kilogram of mass is about the same for both groups. Of course, the obese monkeys are more massive and therefore, on average, burn more energy, as computed in (a) and (b).



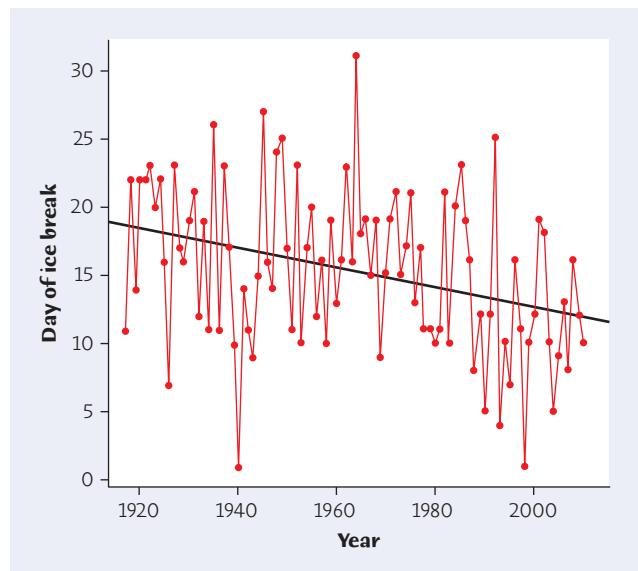
7.39: (a) $190/8474 = 0.0224$, or 2.24%. (b) $633/8474 = 0.0747$, or 7.47%. (c) $27/633 = 0.0427$, or 4.27%. (d) $4621/8284 = 0.5578$, or 55.78%. (e) The conditional distribution of CHD for each level of anger is tabulated below. Angrier people are at greater risk of CHD.

Low anger	Moderate anger	High anger
1.70%	2.33%	4.27%

7.41: The time plot shows a lot of fluctuation from year to year, but it also shows a recent increase: prior to 1972, the discharge rarely rose above 600 km^3 , but since then, it has

exceeded that level more than half the time. A histogram or stemplot cannot show this change over time.

7.43: (a) See below. (b) $\hat{y} = 160.79 - 0.07410x$. The slope is negative, suggesting that the ice breakup day is decreasing (by 0.07410 day per year). (c) The regression line is not very useful for prediction, as it accounts for only about 11.7% ($r^2 = 0.117$) of the variation in ice breakup time.



7.45: (a) and (b) Two stemplots are provided. The first shows all the data points, and the second omits the three highest countries, which are identified as outliers by use of the $1.5 \times \text{IQR}$ criterion. In the absence of those outliers, the distribution is roughly symmetric. (c) $\bar{x} = 18.2447\%$ and $s = 7.0451\%$. The 5-number summary is Min = 3.14%, $Q_1 = 13.64\%$, $M = 18.27\%$, $Q_3 = 24.08\%$, Max = 34.83%. (d) The U.S. share of G.D.P. is small compared to the other countries in this list: more than one standard deviation below the mean and below the first quartile of the distribution.

0	3456889
1	0000111122233334444444455566667788888999999
2	000111334444445666899999
3	044
4	8
5	
6	3
7	
8	
9	
10	6

0	3
0	45
0	6
0	889
1	00001111
1	2223333
1	44444444555
1	666677
1	888888999999
2	000111
2	33
2	4444445
2	666
2	899999
3	0
3	3
3	44

7.47: A stemplot is shown. The distribution seems to be fairly Normal apart from a high outlier of 50° . The five-number summary is preferred because of the outlier: $\text{Min} = 13^\circ$, $Q_1 = 20^\circ$, $M = 25^\circ$, $Q_3 = 30^\circ$, $\text{Max} = 50^\circ$. (The mean and standard deviation are $\bar{x} = 25.4211^\circ$ and $s = 7.4748^\circ$.) Most patients have a deformity angle in the range of 15° to 35° .

1	34
1	66788
2	000111123
2	55556666888
3	00012224
3	88
4	
4	
5	0

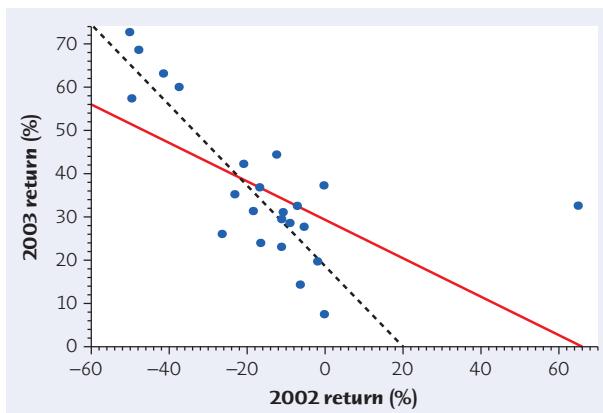
7.49: The scatterplot shows a moderate to weak positive linear association, with one clear outlier (the patient with HAV angle 50°). $r = 0.3021$, $\hat{y} = 19.723 + 0.3388x$. MA angle can be used to give (very rough, imprecise) estimates of HAV angle, but the spread is so wide that the estimates would not be very reliable. The linear relationship explains only $r^2 = 9.1\%$ of the variation in HAV angle.

7.51: Software gives us the regression line ($\hat{y} = 70.44 + 274.78x$), and a scatterplot shows a moderate positive linear relationship. The linear relationship explains about $r^2 = 49.3\%$ of the variation in gate velocity.

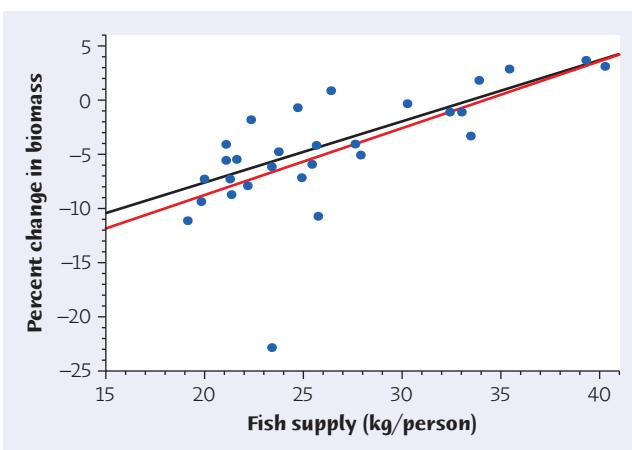
7.53: (a) The scatterplot (top of right column) of 2003 returns against 2002 returns shows (ignoring the outlier) a strong negative association. (b) The correlation for all 23 points is $r = -0.6230$; with the outlier removed, the correlation is

$r = -0.8722$. (c) Regression formulas are given in the table below. The first line is solid in the plot, the second is the dashed line. The line for the 22 other funds is so far below Fidelity Gold that the squared deviation is very large. The line must pivot up toward Fidelity Gold in order to minimize the sum of squares for all 23 deviations. Fidelity Gold is very influential.

	<i>r</i>	Equation
All 23 funds	-0.6230	$\hat{y} = 29.2512 - 0.4501x$
Without Fidelity Gold	-0.8722	$\hat{y} = 18.1106 - 0.9429x$



7.55: (a) Fish catch (on the horizontal axis) is the explanatory variable. The point for 1999 is at the bottom of the plot. (b) The correlations are given in the table below. The outlier decreases r because it weakens the strength of the association. (c) The two regression lines are given in the table; the solid line in the plot uses all points, while the dashed line omits the outlier. The effect of the outlier on the line is small: it pulls the line down on the left side (and increases the slope) very slightly, but for making predictions, both lines would give similar results.



	<i>r</i>	Equation
All points	0.6724	$\hat{y} = -21.09 + 0.6345x$
Without 1999	0.8042	$\hat{y} = -19.05 + 0.5788x$

Chapter 8 Producing Data: Sampling

8.1: (a) (All) college students. (b) The 104 students at the researcher's college who returned the questionnaire.

8.3: (a) All 45,000 people who made credit card purchases. (b) The 137 people who returned the survey form.

8.5: Since all the students surveyed are enrolled in a special senior honors class, these students may be more likely to be interested in joining the club (and more willing to pay \$35 to do so). The direction of bias is likely to overestimate the proportion of all psychology majors willing to pay to join this club. This is a convenience sample.

8.7: Number from 01 to 26 alphabetically (down the columns). With Table B, enter at line 134 and choose 16 = Ippolito, 18 = Jung, 13 = Gupta, 21 = Modur, and 04 = Bonds.

8.9: With the election close at hand, the polling organization wants to increase the accuracy of its results. Larger samples provide better information about the population.

8.11: Label the suburban townships from 01 to 30, down the columns. With Table B, enter at line 105 and choose 29 = Wheeling, 07 = Elk Grove, 19 = Orland, 14 = New Trier, and 17 = Norwood Park. Next, label the Chicago townships from 1 to 8, down the columns. With Table B, enter at line 115 and choose 6 = Rogers Park, 1 = Hyde Park, and 4 = Lake View.

8.13: The higher no-answer was probably the second period—more families are likely to be gone for vacations or to be outside enjoying the warmer weather, and so on.

8.15: (a) and (b) Features will vary depending on the website chosen. (c) The weakness of any online poll is that it relies on voluntary response. Most online poll samples are not representative of any larger population of use or interest to the researcher.

8.17: (a)

8.19: (b)

8.21: (b)

8.23: (c) Notice that in (b) "07" appears in the sample twice.

8.25: (b)

8.27: The population is the 1000 envelopes stuffed during a given hour. The sample is the 40 envelopes selected.

8.29: Using Table B, number the area codes 001 to 287. Then, enter at line 135, and pay attention to the instructions that if we use the table, we'll pick only 5 numbers. The selected area codes are 255, 100, 120, 126, 008.

8.31: (a) Alphabetize the 6168 names (using middle initials or a student ID to distinguish between two people with the same name). Label these students with an ID 0001 to 6168. (b) Using Table B, entering at line 135, the sample is 5556, 5839, 1007, 1120, 1513, 1260, 0842, and 1447.

8.33: (a) False. Such regularity holds only in the long run. (b) True. All pairs of digits (there are 100, from 00 to 99) are equally likely. (c) False. Four random digits have chance 1/10,000 to be 0000, so this sequence will occasionally occur.

8.35: Online polls, call-in polls, and voluntary response polls in general tend to attract responses from those who have strong opinions on the subject, and therefore are often not representative of the population as a whole.

8.37: (a) Assign labels 0001 through 5024, enter the table at line 104, and select: 1388, 0746, 0227, 4001, and 1858. (b) More than 171 respondents have run red lights. We would not expect very many people to claim they *have* run red lights when they have not, but some people will deny running red lights when they have.

8.39: (a) Each person has a 10% chance: 4 of 40 men, and 3 of 30 women. (b) This is not an SRS because not every group of 7 people can be chosen; the only possible samples are those with 4 men and 3 women.

8.41: Sample separately in each stratum; that is, assign separate labels, then choose the first sample, then continue on in the table to choose the next sample, etc. Beginning with line 102 in Table B, we choose:

Forest type	Labels	Parcels selected
Climax 1	01 to 36	19, 27, 26, 17
Climax 2	01 to 72	09, 55, 32, 22, 69, 56, 52
Climax 3	01 to 31	13, 07, 02
Secondary	01 to 42	27, 40, 01, 18

8.43: (a) Since $200/5 = 40$, we will choose one of the first 40 names at random. Beginning on line 120, the addresses selected are 35, 75, 115, 155, and 195. (Only the first number is chosen from the table.) (b) All addresses are equally likely; each has chance 1/40 of being selected. This is not an SRS because the only possible samples have exactly one address from the first 40, one address from the second 40, and so on. An SRS could contain any 5 of the 200 addresses in the population.

8.45: (a) Automated random digit dialing is a fast, economical way to randomly dial landline telephone numbers. (b) In some families, the adult that answers the phone regularly may be systematically different from an adult that does not. (c) There could be (and probably are) big differences between landline phone users and cellular phone users. The design is a stratified sample.

8.47: Answers will vary considerably.

Chapter 9 Producing Data: Experiments

9.1: This is an observational study: no treatment was assigned to the subjects; we merely observed cell phone usage (and presence/absence of cancer). Explanatory variable: cell phone usage; response variable: whether or not a subject has brain cancer.

9.3: This is an observational study, so it is not reasonable to conclude any cause-and-effect relationship.

9.5: Individuals: pine seedlings. Factor: amount of light. Treatments: full light, 25% light, or 5% light. Response variable: dry weight at the end of the study.

9.7: Making a comparison between the treatment group and the percent finding work *last year* is not helpful. Over a year, many things can change. (In order to draw conclusions, we would need to make the \$500 bonus offer to some people and not to others, and compare the two groups.)

9.9: (a) See below. (b) If using Table B, label 01 to 36 and take two digits at a time.



9.11: In a controlled scientific study, the effects of factors other than the nonphysical treatment (e.g., the placebo effect, differences in the prior health of the subjects) can be eliminated or accounted for so that the differences in improvement observed between the subjects can be attributed to the differences in treatments.

9.13: (a) The researchers simply observed the diets of subjects; they did not alter them. (No treatments were assigned.) (b) Such language is reasonable because with observational studies no “cause and effect” conclusion would be reasonable.

9.15: In this case, “lack of blindness” means that the experimenter knows which subjects were taught to meditate. He or she may have some expectation about whether or not meditation will lower anxiety; this could unconsciously influence the diagnosis.

9.17: (a) *Completely randomized design:* Randomly assign 15 students to Group 1 (easy mazes) and the other 15 to Group 2 (hard mazes). Compare the time estimates of Group 1 with those of Group 2. (b) *Matched-pairs design:* Each student does the activity twice, once with the easy mazes and once with the hard mazes. Randomly decide (for each student) which set of mazes is used

first. Compare each student’s “easy” and “hard” time estimate (for example, by looking at each “hard” minus “easy” difference).

9.19: (a) Behavior (alcohol consumption) is observed, but no treatment is imposed.

9.21: (c)

9.23: (b)

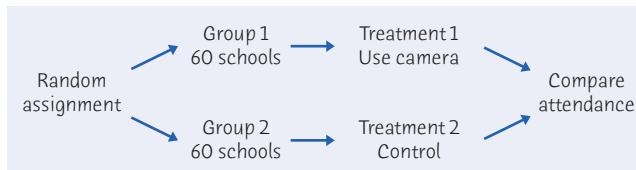
9.25: (b) The communities are paired up, then one is chosen to have the advertising campaign.

9.27: (b) This was a (matched-pairs) experiment, but in order to give useful information, the subjects should be chosen from those who might be expected to buy this car.

9.29: This is an experiment because the treatment is selected (randomly, we assume) by the interviewer. Explanatory variable: level of identification; response variable: whether or not the interview is completed.

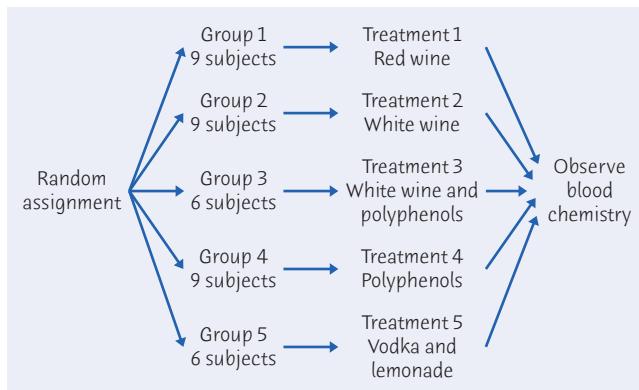
9.31: (a) In an observational study, we simply observe subjects who have chosen to take supplements and compare them with others who do not take supplements. In an experiment, we *assign* some subjects to take supplements and assign the others to take no supplements (or better yet, assign the others to take a placebo). (b) “Randomized” means that the assignment to treatments is made randomly, rather than by some other method (e.g., asking for volunteers). “Controlled” means that some subjects were used as a “control” group—probably meaning that they received placebos—which gives a basis for comparison to observe the effects of the treatment. (c) Subjects who choose to take supplements have other characteristics that are confounded with the effect of the supplements.

9.33: (a) Diagram below. (b) Assign labels 001 to 120. If using Table B, line 108 gives 090, 009, 067, 092, 041, 059, 040, 080, 029, 091.



9.35: Use a completely randomized design. Labeling the men from 01 through 39, and starting on line 107 of Table B, we make the assignments shown in the table below.

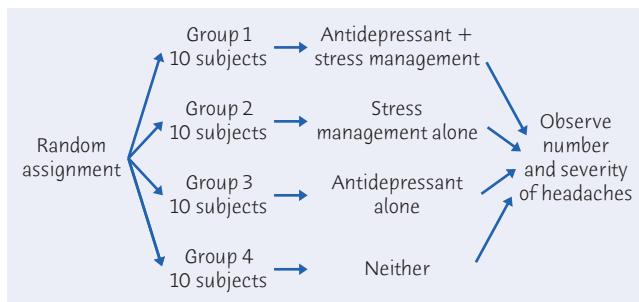
- | | |
|----------|------------------------------------|
| Group 1: | 20, 11, 38, 31, 07, 24, 17, 09, 06 |
| Group 2: | 36, 15, 23, 34, 16, 19, 18, 33, 39 |
| Group 3: | 08, 30, 27, 12, 04, 35 |
| Group 4: | 02, 32, 25, 14, 29, 03, 22, 26, 10 |
| Group 5: | Everyone else |



9.37: (a) There are 40 subjects, so we assign 10 subjects to each of the four treatments.

	Antidepressant	No drug
Stress management	1	2
None	3	4

(b) Assign labels 01 through 40 (in alphabetical order). Line 125 of Table B gives the following subjects for Group 1: 21 Jiang, 37 Suarez, 18 Hersch, 23 Kim, 19 Hurwitz, 10 Devlin, 33 Richter, 31 Ramdas, 36 Smith, and 40 Xiang.



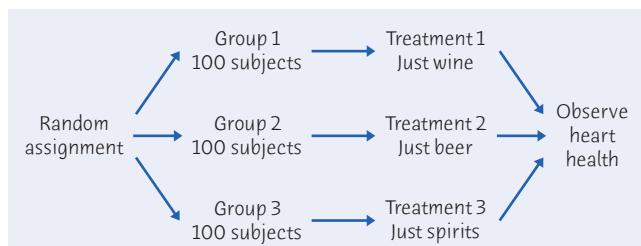
9.39: The factors are pill type and spray type. “Double-blind” means that the treatment assigned to a patient was unknown to both the patient and those responsible for assessing the effectiveness of that treatment. “Placebo-controlled” means that some of the subjects were given placebos.

9.41: (a) Subjects: randomly chosen Starbucks customers. Each subject tastes two cups of coffee, in identical unlabeled cups. One contains regular mocha Frappuccino, the other the new light version. The cups are presented in random order, half the subjects get regular then light, the other half light then regular. Each subject says which cup he or she prefers.

(b) We must assign 10 customers to get regular coffee first. Label the subjects 01 to 20. Starting at line 141, the “regular first” group is: 12, 16, 02, 08, 17, 10, 05, 09, 19, 06.

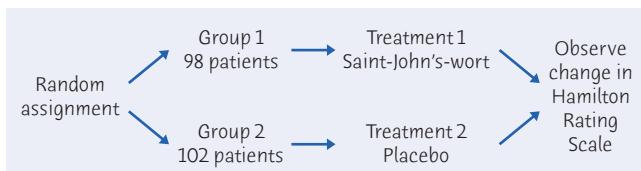
9.43: Each player will be put through the sequence (100 yards, four times) twice—once with oxygen and once without. For each player, randomly determine whether to use oxygen on the first or second trial. Allow ample time (perhaps a day or two) between trials for full recovery.

9.45:



9.47: Any experiment randomized in this way assigns all the women to one treatment and all the men to the other. That is, sex is completely confounded with treatment.

9.49: (a) “Randomized” means that patients were randomly assigned to receive either Saint-John’s-wort or a placebo. “Double-blind” means that the treatment assigned to a patient was unknown to both the patient and those responsible for assessing the effectiveness of that treatment. “Placebo-controlled” means that some of the subjects were given placebos. (b)



Data Ethics Solutions

As the text states, “Most of these exercises pose issues for discussion. There are no right or wrong answers, but there are more and less thoughtful answers.” We have not tried to supply answers for exercises that are largely matters of opinion. For that reason, only four solutions are provided here.

- These five proposals are clearly in increasing order of risk. Most students will consider that (a) qualifies as minimal risk, and most will agree that (e) goes beyond minimal risk. Opinions will vary on where to “draw the line,” of course.

3. It is good to state the purpose of the research plainly (“To study how people’s political beliefs are related to risk of depression”). Stating the research *thesis* (that people with ultraliberal beliefs are at greater risk of depression) would cause bias.

7. This offers anonymity because names are never revealed.

11. For example, informed consent is lacking. (Although most students might express shock or dismay at this research approach, they may struggle with identifying exactly *what* is wrong with these studies.)

Chapter 10 Introducing Probability

10.1: In the long run, of a large number of Texas Hold’em games in which you hold a pair, the fraction in which you can make four of a kind will be about $2/245$. It does not mean that exactly 2 out of 245 such hands would yield four of a kind.

10.3: (a) There are 21 zeros among the first 200 digits of the table (rows 101–105) for a proportion of 0.105. (b) Answers will vary.

10.5: (a) $S = \{\text{lives on campus, lives off campus}\}$. (b) $S = \{\text{All numbers between } \underline{\hspace{1cm}} \text{ and } \underline{\hspace{1cm}} \text{ years}\}$. (Choices of upper and lower limits will vary.) (c) $S = \{\text{all amounts greater than or equal to } 0\}$, or $S = \{0, 0.01, 0.02, 0.03, \dots\}$. (d) $S = \{A, B, C, D, F\}$.

10.7: $S = \{3, 4, 5, 6, 7, 8, 9\}$. As all faces are equally likely and the dice are independent, each of the 16 possible pairings is equally likely, so (for example) the probability of a total of 5 is $3/16$ because 3 pairings add to 4 (and then we add 1). The complete set of probabilities is shown in the table.

Total	Probability
3	1/16
4	2/16
5	3/16
6	4/16
7	3/16
8	2/16
9	1/16

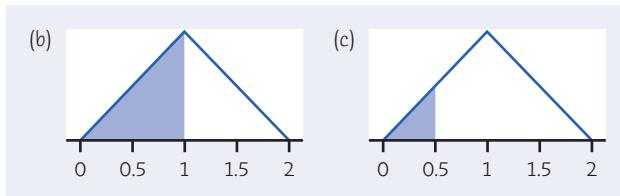
10.9: (a) Event B specifically rules out obese subjects, so there is no overlap with event A. (b) A or B is the event “The person chosen is overweight or obese.” $P(A \text{ or } B) = P(A) + P(B) = 0.34 + 0.33 = 0.67$. (c) $P(C) = 1 - P(A \text{ or } B) = 1 - 0.67 = 0.33$.

10.11: Model 1: Not legitimate (probabilities have sum $6/7$). Model 2: Legitimate. Model 3: Not legitimate (probabilities

have sum $7/6$). Model 4: Not legitimate (probabilities cannot be more than 1).

10.13: (a) This is a legitimate probability model because the probabilities sum to 1. (b) $\{X < 4\}$ is the event that somebody lifts weights 3 or fewer days per week. $P(X < 4) = 0.91$. (c) $\{X \geq 1\}$. $P(X \geq 1) = 0.27$.

10.15: (a) $\frac{1}{2}bh = \frac{1}{2}(2)(1) = 1$. (b) $P(X < 1) = 0.5$. (c) $P(X < 0.5) = 0.125$.



10.17: (a) $X \geq 3$ means the student’s grade is B or higher (B, B+, A– or A). $P(X \geq 3) = 0.41$. (b) “Poorer than B–” means any grade **lower** than B–. $P(X < 2.7) = P(X \leq 2.3) = 0.46$.

10.19: (a) Answers will vary. (b) A personal probability might take into account specific information about one’s own driving habits or about the kind of traffic one usually drives in. (c) Most people believe that they are better-than-average drivers (whether or not they have any evidence to support that belief).

10.21: (a) Probabilities express the *approximate* fraction of occurrences out of many trials.

10.23: (b) This is a discrete (but not equally likely) model.

10.25: (c) $P(\text{Republican or Democrat}) = P(\text{Republican}) + P(\text{Democrat}) = 0.28 + 0.28 = 0.56$.

10.27: (b) There are 10 equally likely possibilities, so $P(\text{seven}) = 1/10$.

10.29: (b) 24% ($0.16 + 0.05 + 0.02 + 0.01 = 0.24$, or 24%) have 3 or more cars.

10.31: (a) There are 16 possible outcomes: {HHHH, HHHM, HHMH, HMHH, MHHH, HHMM, HMHM, HMMH, MHMH, MMHH, MMHH, HMMM, MHMM, MMHM, MMMH, MMMM}. (b) $S = \{0, 1, 2, 3, 4\}$.

10.33: (a) $1 - 0.73 = 0.27$. (b) $P(\text{at least a high school education}) = 1 - P(\text{has not finished HS}) = 1 - 0.13 = 0.87$.

10.35: (a) All probabilities are between 0 and 1, and they add to 1. (We must assume that no one takes more than one language.) (b) $0.43 = 1 - 0.57$. (c) $0.40 = 0.30 + 0.08 + 0.02$.

10.37: Of the seven cards, there are three 9’s, two red 9’s, and two 7’s. (a) $P(\text{draw a 9}) = 3/7$. (b) $P(\text{draw a red 9}) = 2/7$. (c) $P(\text{don’t draw a 7}) = 1 - P(\text{draw a 7}) = 1 - 2/7 = 5/7$.

10.39: Each of the 90 guests has probability $1/90$ of winning the prize. Since there are 42 women, the probability is $42/90 = 0.467$.

10.41: (a) It is legitimate because every person must fall into exactly one category, the probabilities are all between 0 and 1, and they add up to 1. (b) 0.169. (c) 0.171—the sum of the numbers in the first column. (d) 0.532—the sum of the numbers in the third row.

10.43: (a) $1 - 0.171 = 0.829$. (b) $1 - 0.073 = 0.927$.

10.45: (a) All 9 digits are equally likely, so each has probability $1/9$:

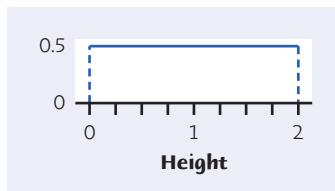
Value of W	1	2	3	4	5	6	7	8	9
Probability	$\frac{1}{9}$								

(b) $P(W \geq 6) = 4/9 = 0.444$, or twice as big as the Benford's law probability.

10.47: (a) BBB, BBG, BGB, GBB, GGB, GBG, BGG, GGG. Each has probability $1/8$. (b) $P(X = 2) = 3/8 = 0.375$. (c) See table.

Value of X	0	1	2	3
Probability	$1/8$	$3/8$	$3/8$	$1/8$

10.49: (a) This is a continuous random variable because the set of possible values is an interval. (b) The height should be $1/2$ because the area under the curve must be 1. The density curve is illustrated. (c) $1/2$.



10.51: (a) $P(0.51 \leq V \leq 0.55) = P(\frac{0.51 - 0.53}{0.009} \leq Z \leq \frac{0.55 - 0.53}{0.009}) = P(-2.22 \leq Z \leq 2.22) = 0.9868 - 0.0132 = 0.9736$. (b) $P(V \geq 0.55) = P(Z \geq \frac{0.55 - 0.53}{0.009}) = P(Z \geq 2.22) = 1 - 0.9868 = 0.0132$.

10.53: (a) Because there are 10,000 equally likely four-digit numbers (0000 through 9999), the probability of an exact match is $1/10,000$. (b) There is a total of $24 = 4 \cdot 3 \cdot 2 \cdot 1$ arrangements of the four digits 5, 9, 7, and 4, so the probability of a match in any order is $24/10,000$.

10.55: (a)–(c) Results will vary, but after n tosses, the distribution of the proportion \hat{p} is approximately Normal with mean 0.5 and standard deviation $1/(2\sqrt{n})$, while the

distribution of the count of heads is approximately Normal with mean $0.5n$ and standard deviation $\sqrt{n}/2$, so using the 68–95–99.7 rule, we have the results shown in the table on the right. Note that the range for \hat{p} gets narrower while the range for the count gets wider.

n	99.7% Range for \hat{p}	99.7% Range for count
40	0.5 ± 0.237	20 ± 9.5
120	0.5 ± 0.137	60 ± 16.4
240	0.5 ± 0.097	120 ± 23.2
480	0.5 ± 0.068	240 ± 32.9

10.57. (a) With $n = 20$, the variability in \hat{p} is larger. With $n = 80$, nearly all answers will be between 0.24 and 0.56. With $n = 320$, nearly all answers will be between 0.32 and 0.48.

Chapter 11 Sampling Distributions

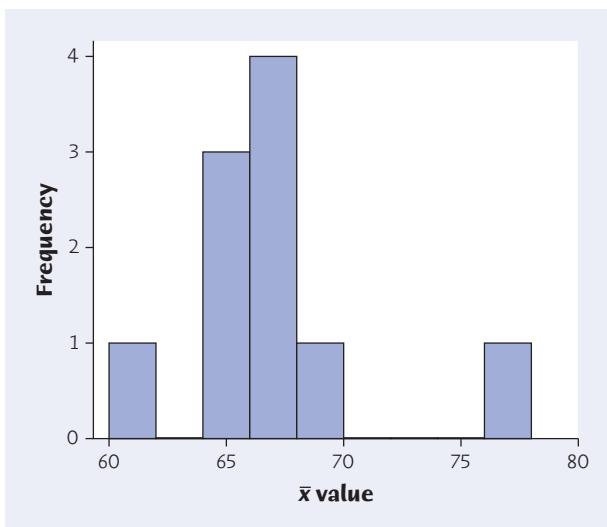
11.1: Both are statistics (related to one sample—the subjects before infusion and the same subjects after infusion).

11.3: Both 12% and 23 are statistics, as they describe the sample of 230 American male weight lifters.

11.5: Although the probability of having to pay for a total loss for 1 or more of the 12 policies is very small, if this were to happen, it would be financially disastrous. On the other hand, for thousands of policies, the law of large numbers says that the average claim on many policies will be close to the mean, so the insurance company can be assured that the premiums they collect will (almost certainly) cover the claims.

11.7: (a) $\mu = 694/10 = 69.4$. (b) The table below shows the results for line 116. Note that we need to choose 5 digits because the digit 4 appears twice. (c) The results for the other lines are in the table; the histogram is shown after the table.

Line	Digits	Scores	\bar{x}
116	14459	$63 + 72 + 72 + 59 = 266$	66.5
117	3816	$55 + 75 + 63 + 65 = 258$	64.5
118	7319	$66 + 55 + 63 + 59 = 243$	60.75
119	95857	$59 + 72 + 75 + 66 = 272$	68
120	3547	$55 + 72 + 72 + 66 = 265$	66.25
121	7148	$66 + 63 + 72 + 75 = 276$	69
122	1387	$63 + 55 + 75 + 66 = 259$	64.75
123	54580	$72 + 72 + 75 + 86 = 305$	76.25
124	7103	$66 + 63 + 86 + 55 = 270$	67.5
125	9674	$59 + 65 + 66 + 72 = 262$	65.5



11.9: (a) The sampling distribution of \bar{x} is $N(186 \text{ mg/dl}, 4.1 \text{ mg/dl})$. $P(183 < \bar{x} < 189) = P(-0.73 < Z < 0.73) = 0.5346$. (b) With $n = 1000$, the sample mean has the $N(186 \text{ mg/dl}, 1.2965 \text{ mg/dl})$ distribution, so $P(183 < \bar{x} < 189) = P(-2.31 < Z < 2.31) = 0.9792$.

11.11: No: the histogram of the sample values will look like the population distribution, whatever it might happen to be. The central limit theorem says that the histogram of *sample means* (from many large samples) will look more and more Normal.

11.13: The central limit theorem says that, in spite of the skewness of the population distribution, the average loss among 10,000 policies will be approximately $N(\$75, \$3)$. $P(\bar{x} > \$85) = P(Z > \frac{85 - 75}{3}) = P(Z > 3.33) = 1 - 0.9996 = 0.0004$.

11.15: (c) 58.8% is a proportion of all registered voters (the population).

11.17: (a) The mean of the sample means (\bar{x} 's) is the same as the population mean (μ).

11.19: (a) “Unbiased” means that the estimator is right “on the average.”

11.21: (b) For $n = 6$ women, \bar{x} has a $N(266, 16/\sqrt{6}) = N(266, 6.5320)$ distribution, so $P(\bar{x} > 270) = P(Z > 0.61) = 0.2709$.

11.23: Both 25.40 and 20.41 are statistics.

11.25: \bar{x} has mean $\mu = 852 \text{ mm}$, and standard deviation $\sigma/\sqrt{n} = 82/\sqrt{10} = 25.93 \text{ mm}$.

11.27: Let X be Shelia’s measured glucose level. (a) $P(X > 140) = P(Z > 1.5) = 0.0668$. (b) If \bar{x} is the mean of four measurements (assumed to be independent), then \bar{x} has a $N(122, 12/\sqrt{4}) = N(122 \text{ mg/dl}, 6 \text{ mg/dl})$ distribution, and $P(\bar{x} > 140) = P(Z > 3) = 0.0013$.

11.29: The mean of four measurements has a $N(122 \text{ mg/dl}, 6 \text{ mg/dl})$ distribution, and $P(Z > 1.645) = 0.05$ if Z is $N(0,1)$, so $L = 122 + 1.645 \cdot 6 = 131.87 \text{ mg/dl}$.

11.31: (a) The central limit theorem gives that \bar{x} will have a Normal distribution with mean 8.8 beats per five seconds and standard deviation $1/\sqrt{12} = 0.288675$ beats per five seconds. (b) $P(\bar{x} < 8) = P(Z < -2.77) = 0.0028$. (c) If the total number of beats in one minute is less than 100, then the average over 12 five-second intervals needs to be less than $100/12 = 8.333$ beats per five seconds. $P(\bar{x} < 8.333) = P(Z < -1.62) = 0.0526$.

11.33: The central limit theorem says that over 40 years, \bar{x} (the mean return) is approximately Normal with mean $\mu = 10.8\%$ and standard deviation $17.1\%/\sqrt{40} = 2.704\%$. Therefore, $P(\bar{x} > 10\%) = P(Z > -0.30) = 0.6179$, and $P(\bar{x} < 5\%) = P(Z < -2.14) = 0.0162$. Note: We have to assume that returns in separate years are independent.

11.35: We need to choose n so that $6.4/\sqrt{n} = 1$. That means $\sqrt{n} = 6.4$, so $n = 40.96$. Because n must be a whole number, take $n = 41$.

11.37: On the average, Joe loses 40 cents each time he plays (that is, he spends \$1 and gets back 60 cents).

11.39: (a) With $n = 150,000$, $\mu_{\bar{x}} = \$0.40$ and $\sigma_{\bar{x}} = \frac{\$18.96}{\sqrt{150,000}} = \0.0490 . (b) $P(\$0.30 < \bar{x} < \$0.50) = P(-2.04 < Z < 2.04) = 0.9586$.

11.41: The mean is $10.5 = (3)(3.5)$ because a single die has a mean of 3.5). Sketches will vary, as will the number of rolls.

Chapter 12 General Rules of Probability

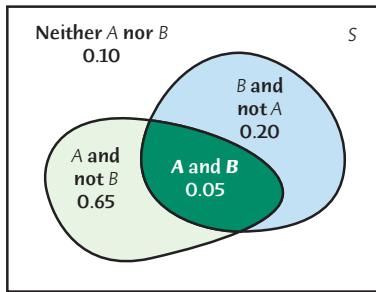
12.1: It is unlikely that these events are independent. In particular, it is reasonable to expect that younger adults are more likely than older adults to be college students.

12.3: If we assume that each site is independent of the others (and that they can be considered as a random sample from the collection of sites referenced in scientific journals), then $P(\text{all seven are still good}) = (0.87)^7 = 0.3773$.

12.5: (a) A Venn diagram is provided. (b) The events are

{A and B}	= {student is at least 25 and local}
{A and not B}	= {student is at least 25 and not local}
{B and not A}	= {student is less than 25 and local}
{neither A nor B}	= {student is less than 25 and not local}

(c) $P(A \text{ and } B)$ is given. Subtracting this from the given probabilities for A and B gives $P(A \text{ and not } B)$ and $P(B \text{ and not } A)$. Those probabilities add to 0.90, so $P(\text{neither } B \text{ nor } W) = 0.10$.



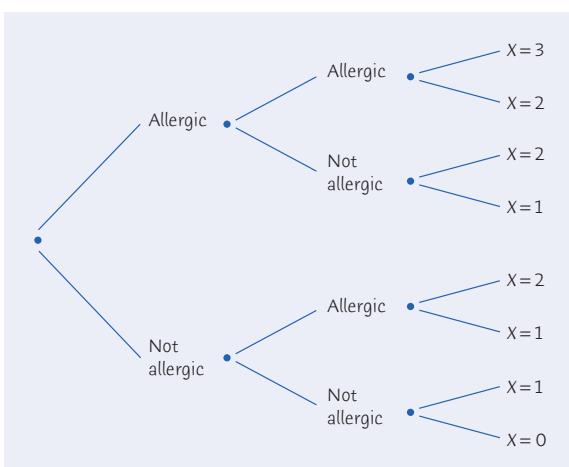
$$12.7: P(B \mid \text{not } A) = \frac{P(B \text{ and not } A)}{P(\text{not } A)} = \frac{P(B) - P(B \text{ and } A)}{P(\text{not } A)} = \frac{0.2}{0.3} = 0.667.$$

12.9: Let H be the event that an adult belongs to a club and T be the event that he/she goes at least twice a week. We have been given $P(H) = 0.15$ and $P(T \mid H) = 0.50$. Note also that $P(T \text{ and } H) = P(T)$, since one has to be a member of the club in order to attend. So $P(T) = P(H)P(T \mid H) = (0.15)(0.50) = 0.075$. About 7.5% of all adults go to health clubs at least twice a week.

12.11: (a) and (b) These probabilities are provided in the table. (c) The product of these conditional probabilities gives the probability of a flush in spades by the general multiplication rule: we must draw a spade, and then another, and then a third, a fourth, and a fifth. The product of these probabilities is about 0.0004952. (d) Because there are four possible suits in which to have a flush, the probability of a flush is four times that found in (c), or about 0.001981.

$$\begin{aligned} P(1\text{st card } \spadesuit) &= \frac{13}{52} = \frac{1}{4} = 0.25 \\ P(2\text{nd card } \spadesuit \mid 1 \spadesuit \text{ picked}) &= \frac{12}{51} = \frac{4}{17} \doteq 0.2353 \\ P(3\text{rd card } \spadesuit \mid 2 \spadesuit \text{s picked}) &= \frac{11}{50} = 0.22 \\ P(4\text{th card } \spadesuit \mid 3 \spadesuit \text{s picked}) &= \frac{10}{49} \doteq 0.2041 \\ P(5\text{th card } \spadesuit \mid 4 \spadesuit \text{s picked}) &= \frac{9}{48} = \frac{3}{16} = 0.1875 \end{aligned}$$

12.13: In the tree diagram, each “up-step” represents an allergic individual (and has probability 0.01), and each



“down-step” is a nonallergic individual (and has probability 0.99). At the end of each of the 8 complete branches is the value of X .

There are 3 branches each corresponding to $X = 2$ and $X = 1$, and only one branch each for $X = 3$ and $X = 0$. Because $X = 0$ and $X = 3$ appear on one branch each, $P(X = 0) = 0.99^3 = 0.970299$ and $P(X = 3) = 0.01^3 = 0.000001$. Meanwhile, $P(X = 1) = (3)(0.01)^1(0.99)^2 = 0.029403$, and $P(X = 2) = (3)(0.01)^2(0.99)^1 = 0.000297$.

$$12.15: P(X = 2 \mid X \geq 1) = \frac{P(X = 2 \text{ and } X \geq 1)}{P(X \geq 1)} = \frac{P(X = 2)}{P(X \geq 1)} = \frac{0.000297}{1 - 0.970299} = 0.010.$$

$$12.17: (b) (0.98)^3 = 0.9412.$$

$$12.19: (a) P(\text{at least one positive}) = 1 - P(\text{both negative}) = 1 - P(\text{first negative})P(\text{second negative}) = 1 - (0.1)(0.2) = 0.98.$$

$$12.21: (c) 2,349/3,294 = 0.7131.$$

12.23: (c) We want the fraction of engineering doctorates conferred to women. Hence, A (engineering degree) is what has been given. Hence, $P(B \mid A)$.

$$12.25: (c) P(W \text{ and } D) = P(W)P(D \mid W) = (0.86)(0.028) = 0.024.$$

$$12.27: (0.75)^8 = 0.1001.$$

12.29: (a) $(\frac{1}{20})(\frac{9}{20})(\frac{1}{20}) = 0.001125$. (b) The other (noncherry) symbol can show up on the middle wheel, with probability $(\frac{1}{20})(\frac{11}{20})(\frac{1}{20}) = 0.001375$, or on either of the outside wheels, with probability $= (\frac{19}{20})(\frac{9}{20})(\frac{1}{20})$ (each). (c) Combining all three cases from part (b), we have $P(\text{exactly two cherries}) = 0.001375 + 2 \cdot 0.021375 = 0.044125$.

12.31: Let I be the event “infection occurs” and let F be “the repair fails.” $P(I) = 0.03$, $P(F) = 0.14$, and $P(I \text{ and } F) = 0.01$. We want to find $P(\text{not } I \text{ and not } F)$. $P(I \text{ or } F) = P(I) + P(F) - P(I \text{ and } F) = 0.03 + 0.14 - 0.01 = 0.16$. Now observe that the desired probability is the complement of “ $I \text{ or } F$ ”: $P(\text{not } I \text{ and not } F) = 1 - P(I \text{ or } F) = 0.84$.

12.33: Let I be the event “infection occurs” and let F be “the repair fails.” $P(I \mid \text{not } F) = \frac{P(I \text{ and not } F)}{P(\text{not } F)} = \frac{0.02}{0.86} = 0.0233$.

12.35: Let H be the event “student was home schooled.” Let R be the event “student attended a regular public school.” Note that the event “ H and not R ” = “ H ”, since the events are disjoint. Then $P(H \mid \text{not } R) = \frac{P(H)}{P(\text{not } R)} = \frac{0.006}{(1 - 0.781)} = 0.0274$.

12.37: (a) These events are not independent because $P(\text{pizza with mushrooms}) = 4/7$, but $P(\text{mushrooms} \mid \text{thick crust}) = 2/3$. Alternatively, note that $P(\text{thick crust with mushrooms}) = 2/7$, which is not equal to the product of $P(\text{mushrooms}) = 4/7$ and $P(\text{thick crust pizza}) = 3/7$. (b) With the eighth pizza,

$P(\text{mushrooms}) = 4/8 = 1/2$ and $P(\text{mushrooms} \mid \text{thick crust}) = 2/4 = 1/2$, so these events are independent.

12.39: Let W be the event “the person is a woman” and M be “the person earned a master’s degree.” (a) $P(\text{not } W) = 1421/3560 = 0.3992$. (b) $P(\text{not } W \mid M) = 282/732 = 0.3852$. (c) The events “choose a man” and “choose a master’s degree recipient” are not independent. If they were, the two probabilities in (a) and (b) would be equal.

12.41: Let D be the event “a seedling was damaged by a deer.” (a) $P(D) = 209/871 = 0.2400$. (b) The conditional probabilities are

$$\begin{aligned} P(D \mid \text{no cover}) &= 60/211 = 0.2844 \\ P(D \mid \text{cover} < 1/3) &= 76/234 = 0.3248 \\ P(D \mid 1/3 \text{ to } 2/3 \text{ cover}) &= 44/221 = 0.1991 \\ P(D \mid \text{cover} > 2/3) &= 29/205 = 0.1415 \end{aligned}$$

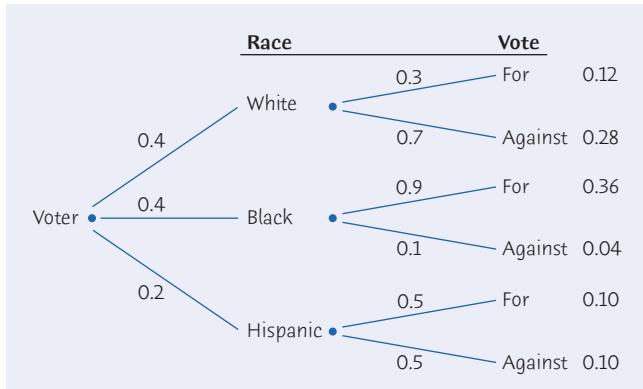
(c) Cover and damage are not independent; $P(D)$ decreases noticeably when thorny cover is $1/3$ or more.

12.43: $P(\text{cover} < 1/3 \mid D) = 76/209 = 0.3636$, or 36.36%.

12.45: $P(A \text{ and not } B \text{ and not } C) = 0.35$.

12.47: (a) $P(\text{doubles on first toss}) = 1/6$, since 6 of the 36 equally likely outcomes enumerated in Figure 10.2 involve rolling doubles. (b) We need no doubles on the first roll (which happens with probability $5/6$), then doubles on the second toss. $P(\text{first doubles appears on toss 2}) = (5/6)(1/6) = 5/36$. (c) Similarly, $P(\text{first doubles appears on toss 3}) = (5/6)^2(1/6) = 25/216$. (d) $P(\text{first doubles appears on toss 4}) = (5/6)^3(1/6)$, etc. In general, $P(\text{first doubles appears on toss } k) = (5/6)^{k-1}(1/6)$. (e) $P(\text{go again within 3 turns}) = P(\text{roll doubles in 3 or fewer rolls}) = P(\text{roll doubles on 1st, 2nd or 3rd try}) = (1/6) + (5/6)(1/6) + (5/6)^2(1/6) = 0.4213$.

12.49: Let W , B , and H be the events that a randomly selected voter is (respectively) white, black, and Hispanic.



We have been given $P(W) = 0.4$, $P(B) = 0.4$, $P(H) = 0.2$. If F = “a voter votes for the candidate,” then $P(F \mid W) = 0.3$, $P(F \mid B) = 0.9$, $P(F \mid H) = 0.5$. We find $P(F)$ by adding all the numbers next to the branches ending in “for”: $P(F) = 0.12 + 0.36 + 0.10 = 0.58$.

12.51: $P(B \mid F) = \frac{P(B \text{ and } F)}{P(F)} = \frac{0.36}{0.58} = 0.6207$, or about 62%.

12.53: Let W , B , A , and L be (respectively) the events that this person is white, black, Asian, and lactose intolerant. We have been given

$$\begin{array}{lll} P(W) = 0.82 & P(B) = 0.14 & P(A) = 0.04 \\ P(L \mid W) = 0.15 & P(L \mid B) = 0.70 & P(L \mid A) = 0.90 \end{array}$$

(a) $P(L) = (0.82)(0.15) + (0.14)(0.70) + (0.04)(0.90) = 0.257$, or 25.7%. (b) $P(A \mid L) = P(A \text{ and } L)/P(L) = (0.04)(0.90)/0.257 = 0.1401$, or 14%.

12.55: In this problem, allele 29 is playing the role of A and 0.181 is the proportion with this allele ($a = 0.181$). Similarly, allele 31 is playing the role of B and the proportion having this allele is $b = 0.071$. The proportion of the population with combination (29,31) is therefore $2(0.181)(0.071) = 0.025702$. The proportion with combination (29,29) is $(0.181)(0.181) = 0.032761$.

12.57: In Exercise 12.55, we found that the proportion of the population with allele (29,31) at loci D21S11 is 0.025702. In Exercise 12.56, we found that the proportion with allele (16,17) at loci D3S1358 is 0.098368. Assuming independence between loci, the proportion with allele (29,31) at D21S11 and (16,17) at D3S1358 is $(0.098368)(0.025702) = 0.002529$.

Chapter 13 Binomial Distributions

13.1: Binomial. (1) We have a fixed number of observations ($n = 15$). (2) It is reasonable to believe that each call is independent of the others. (3) “Success” means reaching a live person, “failure” is any other outcome. (4) Each randomly dialed number has chance $p = 0.2$ of reaching a live person.

13.3: Not binomial. The trials aren’t independent. If one tile in a box is cracked, there are likely more tiles cracked.

13.5: (a) C , the number caught, is binomial with $n = 10$ and $p = 0.7$. M , the number missed, is binomial with $n = 10$ and $p = 0.3$. (b) $P(M = 3) = \binom{10}{3}(0.3)^3(0.7)^7 = (120)(0.027)(0.08235) = 0.2668$. $P(M \geq 3) = 0.6172$.

13.7: (a) $5 \text{ choose } 2$ returns 10. (b) $500 \text{ choose } 2$ returns 124,750, and $500 \text{ choose } 100$ returns $2.04169424 \times 10^{107}$. (c) $(10 \text{ choose } 1) * 0.1 * 0.9^9$ returns 0.387420489.

13.9: (a) X is binomial with $n = 10$ and $p = 0.3$; Y is binomial with $n = 10$ and $p = 0.7$ (b) The mean of Y is $(10)(0.7) = 7$ errors caught, and for X the mean is $(10)(0.3) = 3$ errors

missed. (c) The standard deviation of Y (or X) is $\sigma = \sqrt{10(0.7)(0.3)} = 1.4491$ errors.

13.11: (a) $\mu = (1520)(0.31) = 471.2$ and $\sigma = \sqrt{1520(0.31)(1 - 0.31)} = 18.0313$ students. (b) $np = (1520)(0.31) = 471.2 \geq 10$ and $n(1 - p) = (1520)(0.69) = 1048.8 \geq 10$, so n is large enough for the Normal approximation to be reasonable. The college wants 475 students, so $P(X \geq 476) = P(Z \geq \frac{476 - 471.2}{18.0313}) = P(Z \geq 0.27) = 0.3936$. (c) The exact probability is 0.4045 (obtained from software), so the Normal approximation is 0.0109 too low.

13.13: (b) He has 3 independent eggs, each with probability 1/4 of containing salmonella.

13.15: (c) The selections are not independent; once we choose one student, it changes the probability that the next student is a business major.

13.17: (a) $(0.60)^2(0.40)^3 = 0.02304$.

13.19: (b) This is the event that a single digit is 8 or 9, so the probability is 0.20.

13.21: (a) $np = (80)(0.20) = 16$.

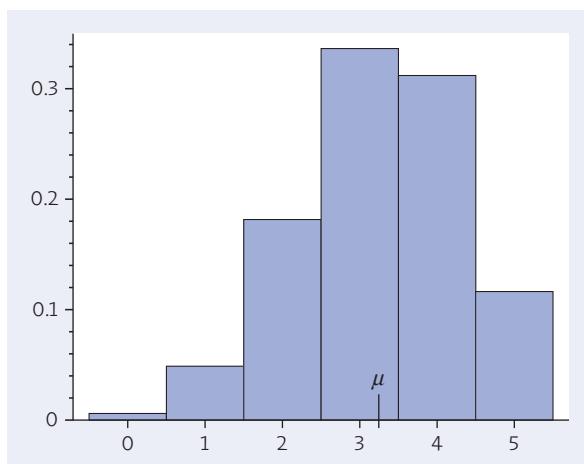
13.23: (a) A binomial distribution is not an appropriate choice for field goals made, because given the different situations the kicker faces, his probability of success is likely to change from one attempt to another. (b) It would be reasonable to use a binomial distribution for free throws made, because we have $n = 150$ attempts, presumably independent (or at least approximately so), with chance of success $p = 0.8$ each time.

13.25: (a) $n = 5$ and $p = 0.65$. (b) The possible values of X are the integers 0, 1, 2, 3, 4, 5. (c)

$$P(X = 0) = \binom{5}{0}(0.65)^0(0.35)^5 = 0.00525$$

$$P(X = 1) = \binom{5}{1}(0.65)^1(0.35)^4 = 0.04877$$

$$P(X = 2) = \binom{5}{2}(0.65)^2(0.35)^3 = 0.18115$$



$$P(X = 3) = \binom{5}{3}(0.65)^3(0.35)^2 = 0.33642$$

$$P(X = 4) = \binom{5}{4}(0.65)^4(0.35)^1 = 0.31239$$

$$P(X = 5) = \binom{5}{5}(0.65)^5(0.35)^0 = 0.11603$$

$$(d) \mu = np = 5(0.65) = 3.25 \text{ and}$$

$$\sigma = \sqrt{5(0.65)(1 - 0.65)} = 1.0665 \text{ years.}$$

13.27: (a) All women are independent, and each has the same probability of getting pregnant. (b) Under ideal conditions, the number who get pregnant is binomial with $n = 20$ and $p = 0.01$; $P(N \geq 1) = 1 - P(N = 0) = 1 - 0.8179 = 0.1821$. In typical use, $p = 0.03$, and $P(N \geq 1) = 1 - 0.5438 = 0.4562$.

13.29: (a) X , the number of women who get pregnant in typical use, is binomial with $n = 600$ and $p = 0.03$. $np = 18$ and $n(1 - p) = 582$ are both larger than 10. The mean is 18 and the standard deviation is 4.1785, so $P(X \geq 20) = P(Z \geq 0.48) = 0.3156$. The exact binomial probability is 0.3477. (b) Under ideal conditions, $p = 0.01$, so $np = 6$ is too small.

13.31: (a) If R is the number of red-blossomed plants out of a sample of 4, then $P(R = 3) = \binom{4}{3}(0.75)^3(0.25)^1 = 0.4219$. (b) With $n = 60$, the mean number of red-blossomed plants is $np = 45$. (c) If R is the number of red-blossomed plants out of a sample of 60, then $P(R \geq 45) = P(Z \geq 0) = 0.5000$ (software gives 0.5688 using the binomial distribution).

13.33: (a) $184,000/1,498,000 = 0.12283$. (b) If I is the number of Impala buyers in the 1000 surveyed buyers, then I has the binomial distribution with $n = 1000$, and $p = 0.12283$. $\mu = np = (1000)(0.12283) = 122.83$ and $\sigma = \sqrt{1000(0.12283)(1 - 0.12283)} = 10.38$. (c) $P(I > 100) = P(I \geq 101) = P(Z \geq -2.10) = 0.9821$.

13.35: (a) With $n = 100$, the mean and standard deviation are $\mu = 75$ and $\sigma = 4.3301$, so $P(70 \leq X \leq 80) = P(-1.15 \leq Z \leq 1.15) = 0.7498$ (software gives 0.7967). (b) With $n = 250$, we have $\mu = 187.5$ and $\sigma = 6.8465$, and a score between 70% and 80% means 175 to 200 correct answers, so $P(175 \leq X \leq 200) = P(-1.83 \leq Z \leq 1.83) = 0.9328$ (software gives 0.9428).

13.37: (a) Answers will vary. (b) Each time we choose a sample of size 10, the probability that we have exactly 1 bad CD is 0.3874; therefore, out of 20 samples, the number of times that we have exactly 1 bad CD has a binomial distribution with parameters $n = 20$ and $p = 0.3874$.

13.39: The number N of infections among untreated BJU students is binomial with $n = 1400$ and $p = 0.80$, so the mean is 1120 and the standard deviation is 14.9666 students. 75% of that group is 1050, and the Normal approximation is safe: $P(N \geq 1050) = P(Z \geq \frac{1050 - 1120}{14.9666}) = P(Z \geq -4.68)$, which is very near 1.

13.41: Let V and U be (respectively) the number of new infections among the vaccinated and unvaccinated children. (a) $P(V = 1) = 0.3741$ and $P(U = 1) = 0.0960$. Because these events are independent, $P(V = 1 \text{ and } U = 1) = P(V = 1)P(U = 1) = 0.0359$. (b) $P(2 \text{ infections}) = P(V = 0 \text{ and } U = 2) + P(V = 1 \text{ and } U = 1) + P(V = 2 \text{ and } U = 0) = P(V = 0)P(U = 2) + P(V = 1)P(U = 1) + P(V = 2)P(U = 0) = 0.1977$.

13.43: The number X of fairways Phil hits is binomial with $n = 24$ and $p = 0.52$. (a) $np = 12.48$ and $n(1 - p) = 11.52$, so the Normal approximation is (barely) safe. (b) The mean is $np = 12.48$ and the standard deviation is $\sqrt{24(0.52)(0.48)} = 2.447529$. Using the Normal approximation, $P(X \geq 17) = P(Z \geq 1.85) = 0.0322$. (c) With the continuity correction, $P(X \geq 17) = P(X \geq 16.5) = P(Z \geq 1.64) = 0.0505$. The answer using the continuity correction is closer to the exact answer (0.0487).

Chapter 14 Confidence Intervals: The Basics

14.1: (a) $\frac{\sigma}{\sqrt{n}} = \frac{34}{\sqrt{51,000}} = 0.1506$. (b) 95% of all values of \bar{x} fall within 2 standard deviations of the sampling distribution of μ , that is, within $2(0.1506) = 0.3012$. (c) 153 ± 0.3012 , or between 152.7 and 153.3.

14.3: In 99.4% of all repetitions of part (a), you should see between 5 and 10 hits (that is, at least 5 of the 10 SRS's capture the true mean μ). Out of 1000 80% confidence intervals, nearly all students will observe between 76% and 84% capturing the mean.

14.5: Search Table A for 0.075 (half of the 15% that is not included in the middle, shaded area corresponding to 85% confidence). This area corresponds to $-z^* = -1.44$, or $z^* = 1.44$.

14.7: (a) A stemplot is provided. The two low scores (72 and 74) are both possible outliers, but there are no other apparent deviations from Normality.

7	24
7	
8	
8	69
9	13
9	68
10	023334
10	578
11	11222444
11	89
12	0
12	8
13	02

(b) The problem states that these girls are an SRS of the population, which is very large, so conditions for inference are met. In part (a), we saw that the scores are consistent with having come from a Normal population. Our 99% confidence interval for μ is given by $105.84 \pm 2.576 \frac{15}{\sqrt{31}} = 98.90$ to 112.78. We are 99% confident that the mean IQ of seventh-grade girls in this district is between 98.90 and 112.78 points.

14.9: With $z^* = 1.96$ and $\sigma = 7.5$, the margin of error is $z^* \frac{\sigma}{\sqrt{n}} = \frac{14.7}{\sqrt{n}}$. (a) and (b) The margins of error are given in the table. (c) Margin of error decreases as n increases. (Specifically, every time the sample size n is quadrupled, the margin of error is halved.)

<i>n</i>	Margin of error
100	1.47
400	0.735
1600	0.3675

14.11: (c) $z = 3.291$. Using Table A, search for 0.9995.

14.13: (b) As the confidence level increases, z^* increases. This makes the margin of error larger.

14.15: (b) The standard deviation of \bar{x} is $\frac{\sigma}{\sqrt{n}} = \frac{35}{\sqrt{900}} = 1.167$.

14.17: (b) As the confidence level increases, z^* increases. This makes the margin of error larger.

14.19: (a) $118 \pm 2.576 \frac{65}{\sqrt{463}} = 110.22$ to 125.78 minutes. (b) The 463 students in this class must be a random sample of all of the first-year students at this university to satisfy conditions for inference.

14.21: The margin of error is now $2.576 \frac{65}{\sqrt{464}} = 7.77$, so the extra observation has minimal impact on the margin of error (the sample was large to begin with). If $\bar{x} = 247$, then the 99% confidence interval for average amount of time spent studying becomes $247 \pm 7.77 = 239.23$ to 254.77 minutes. The outlier had a huge impact on \bar{x} , which shifts the interval a lot.

14.23: This student is also confused. If we repeated the sample over and over, 95% of all future sample means would be within 1.96 standard deviations of μ (that is, within $1.96 \frac{\sigma}{\sqrt{n}}$) of the true, unknown value of μ . Future samples will have no memory of our sample.

14.25: (a) Notice that the distribution is noticeably skewed to the left. The data do not appear to follow a Normal distribution.

23	0
24	0
25	5
26	5
27	7
28	7
29	
30	149
31	389
32	033577
33	0126

(b) The problem states that we are willing to take this sample to be an SRS of the population. In spite of the shape of the stemplot, we are told to assume that this distribution is Normal with standard deviation $\sigma = 3000$ lb. The 95% confidence interval for μ is given by $30,841 \pm 1.96 \frac{3000}{\sqrt{20}} = 29,526.19$ to $32,155.81$. With 95% confidence, the mean load μ required to break apart pieces of Douglas fir is between 29,526.2 and 32,155.8 pounds.

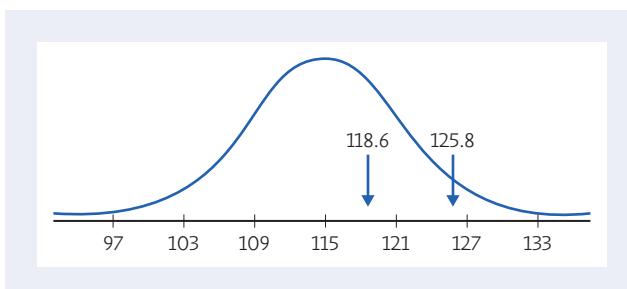
14.27: (a) A stemplot is given. There is little evidence that the sample does not come from a Normal distribution. For inference, we must assume that the 10 untrained students were selected randomly from the population of all untrained people.

1	9
2	239
3	00135
4	2

(b) We have assumed that we have a random sample and that the population we're sampling from is Normal. Our 95% confidence interval for μ is given by $29.4 \pm 1.96 \frac{7}{\sqrt{10}} = 25.06$ to $33.74 \mu\text{g/l}$. With 95% confidence, the mean sensitivity for all untrained people is between 25.06 and $33.74 \mu\text{g/l}$.

Chapter 15 Tests of Significance: The Basics

15.1: (a) The distribution is approximately Normal with mean $\mu = 115$ and standard deviation $\frac{\sigma}{\sqrt{n}} = 6$. (b) The actual result



lies out toward the high tail of the curve, while 118.6 is fairly close to the middle. If $\mu = 115$, observing a value similar to 118.6 would not be too surprising, but 125.8 is less likely, and it therefore provides some evidence that $\mu > 115$.

15.3: $H_0: \mu = 115$ vs. $H_a: \mu > 115$. Because the teacher suspects that older students have a higher mean, we have a one-sided alternative.

15.5: $H_0: \mu = 75$ vs. $H_a: \mu < 75$. The professor suspects this TA's students perform worse than the population of all students in the class on average.

15.7: Hypotheses are statements about parameters, not statistics. The research question is not about the sample mean (\bar{x}) but should be about the population mean (μ).

15.9: The standard deviation is $\frac{\sigma}{\sqrt{18}} = 14.1421$, so when $\mu = 0$, the distribution of \bar{x} is $N(0, 14.1421)$. (b) The P-value is $P = 2P(\bar{x} \geq 17) = 2P(Z \geq \frac{17 - 0}{14.1421}) = 0.2302$.

15.11: (a) P-value = 0.2743. This is not significant at either $\alpha = 0.05$ or $\alpha = 0.01$. (b) P-value = 0.0359. This is significant at $\alpha = 0.05$ but not at $\alpha = 0.01$. (c) If $\mu = 115$ (that is, if H_0 were true), observing a value similar to 118.6 would not be too surprising, but 125.8 is less likely, and it therefore provides some evidence that $\mu > 115$.

15.13: (a) $z = \frac{0.3 - 0}{1/\sqrt{10}} = 0.9488$. (b) $z = \frac{1.02 - 0}{1/\sqrt{10}} = 3.226$. (c) $z = \frac{17 - 0}{60/\sqrt{18}} = 1.2021$.

15.15: Let μ be the average percentage tip for all customers receiving bad news. $H_0: \mu = 20$ vs. $H_a: \mu < 20$. The standard deviation of \bar{x} is $\frac{2}{\sqrt{20}} = 0.4472$, so the test statistic is $z = \frac{18.19 - 20}{0.4472} = -4.05$. P-value is $P(Z \leq -4.05) \approx 0$. There is overwhelming evidence that the average tip percentage when bad news is delivered is lower than the average tip percentage overall.

15.17: This is not significant at the $\alpha = 0.05$ level because z is not larger than 1.96 or less than -1.96 . It is not significant at the $\alpha = 0.01$ level, since z is smaller than 2.576.

15.19: (a)

15.21: (c) P-value = 0.0075 (assuming that the difference is in the correct direction; that is, assuming that the alternative hypothesis was $H_a: \mu > \mu_0$).

15.23: (a) The null hypothesis states that μ takes on the “default” value, 18 seconds.

15.25: (c) The P-value refers to the probability of getting a sample as contrary to the null hypothesis as the sample observed, assuming H_0 is true.

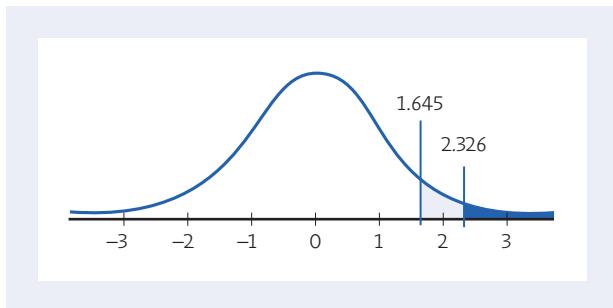
15.27: (b) This is a one-sided alternative, so we have 0.005 in the right tail of the Normal distribution, leading to $z > 2.807$.

15.29: (a) $H_0: \mu = 0$ vs. $H_a: \mu > 0$. (b) $z = \frac{2.35 - 0}{2.5/\sqrt{200}} = 13.29$. (c) This value of z is far outside the range we would expect from the $N(0,1)$ distribution. Under H_0 , it would be virtually impossible to observe a sample mean as large as 2.35 based on a sample of 200 men. Hence we would easily reject H_0 .

15.31: “ $P = 0.03$ ” means that if H_0 is true, a sample as contrary to H_0 as our sample would occur by chance alone only 3% of the time if the experiment was repeated over and over. However, it does not mean that there is a 3% chance that H_0 is true.

15.33: The person making the objection is confusing practical significance with statistical significance. In fact, a 5% increase isn't a lot in a pragmatic sense. $P = 0.03$ means that random chance does not easily explain the difference observed. That is, there does seem to be an increase in mean improvement for those that expressed their anxieties, but the significance test does not address whether the difference is large enough to matter.

15.35: In the sketch, the “significant at 1%” region includes only the dark shading ($z > 2.326$). The “significant at 5%” region of the sketch includes both the light and dark shading ($z > 1.645$). Significance at the 1% level implies significance at the 5% level (or at any level higher than 1%). The converse is false: something that occurs “less than 5 times in 100 repetitions” is not necessarily as rare as something that happens “less than once in 100 repetitions,” so a test that is significant at the 5% level is not necessarily significant at the 1% level.



15.37: (a) Because a P -value is a probability, it can never be greater than 1. (b) P -value = $P(Z \geq 1.33) = 0.0918$.

15.39: $H_0: \mu = 0\%$ vs. $H_a: \mu < 0\%$. $z = \frac{-3.857 - 0}{2.5/\sqrt{47}} = -9.84$; P -value = $P(Z \leq -9.84) \approx 0$. There is overwhelming evidence that, on average, nursing mothers lose bone mineral.

15.41: (a) $H_0: \mu = 0$ vs. $H_a: \mu > 0$, where μ is the mean sensitivity difference in the population. (b) $z = \frac{0.10125 - 0}{0.22/\sqrt{16}} = 1.84$. P -value = 0.0329. The sample gives signifi-

cant evidence (at the $\alpha = 0.05$ level) that eye grease increases sensitivity.

15.43: (a) No, because 33 falls in the 95% confidence interval, which is $(27.5, 33.9)$. (b) Yes, because 34 does not fall in the 95% confidence interval.

Chapter 16 Inference in Practice

16.1: The most important reason is (c); this is a convenience sample consisting of the first 20 students on a list. This is not an SRS. Anything we learn from this sample will not extend to the larger population.

16.3: Any number of things could go wrong with this convenience sample. The day after Thanksgiving is widely regarded (rightly or wrongly) as a day on which retailers offer great deals—and the kinds of shoppers found that day probably don't represent shoppers generally. Also, the sample isn't random.

16.5: No. The confidence interval does not describe the range of future values of \bar{x} .

16.7: The margin of error addresses only chance variation in the random selection of a sample. Hence, the answer is (c).

16.9: The reported P -values from the applet agree with the “hand-computed” values $z = \frac{4.8 - 5}{0.5/\sqrt{n}}$ and $P = P(Z \leq z)$, given in the table.

<i>n</i>	<i>z</i>	<i>P</i>
5	-0.89	0.1867
15	-1.55	0.0606
40	-2.53	0.0057

16.11: (a) In a sample of size $n = 500$, we expect to see about 5 people who have a P -value of 0.01 or less [$5 = (500)(0.01)$]. These 4 might have ESP, or they may simply be among the “lucky” ones that occurred by random chance, as expected. (b) The researcher should repeat the procedure on these 4 subjects to see whether they again perform well.

16.13: $n = \left(\frac{(1.645)(30)}{10}\right)^2 = 24.354$. Take $n = 25$.

16.15: (a) Increase power by taking more measurements. (b) If you increase α , you make it easier to reject H_0 , and increase power. (c) A value of $\mu = 10.2$ is even further from the stated value of $\mu = 10.1$ under H_0 , so power increases.

16.17: The table top of page 711 summarizes power as σ changes. As σ decreases, power increases. More precise measurements increase the researcher's ability to recognize a false null hypothesis.

σ	0.10	0.05	0.025
Power	0.232	0.688	0.998

16.19: (a) All statistical methods are based on probability samples. We must have a random sample in order to apply them.

16.21: (b) Inference from a voluntary response sample is never reasonable. Online web surveys are voluntary response surveys.

16.23: (a) There is no control group. Any observed improvement may be due to the treatment or may be due to another cause.

16.25: (a) The significance level (α) is the probability of rejecting H_0 when H_0 is true.

16.27: (c) Power describes the test's ability to reject a false H_0 .

16.29: We need to know that the samples taken from both populations (hunter-gatherers, agricultural) are random. Are the samples large? Recall that if the samples are very large, then even a small, practically insignificant difference in prevalence of color blindness in the two samples will be deemed statistically significant.

16.31: Many people might be reluctant to relate details of their sex lives, or perhaps some will be inclined to exaggerate. It would not be surprising that such an estimate would be biased.

16.33: The effect is greater if the sample is small. With a larger sample, the impact of any one value is small.

16.35: Opinion—even expert opinion—unsupported by data is the weakest type of evidence, so the third description is level C. The second description refers to experiments (clinical trials) and large samples; that is the strongest evidence (level A). The first description is level B: stronger than opinion, but not as strong as experiments with large numbers of subjects.

16.37: (a) The P -value decreases (the evidence against H_0 becomes stronger). (b) The power increases (the test becomes better at distinguishing between the null and alternative hypotheses).

16.39: (a) $z = \frac{7.524 - 6}{2/\sqrt{5}} = 1.704$. P -value = $2P(Z \geq 1.704) = 0.0884$ (using software). This is not significant at the 5% level of significance. (b) We would not reject 6 as a plausible value of μ , even though (unknown to the researcher) $\mu = 8$. This isn't surprising, since $\bar{x} = 7.524$.

16.41: (a) “Statistically insignificant” means that the differences observed were no more than might have been expected to occur by chance even if SES had no effect on LSAT results.

(b) If the results are based on a small sample, then even if the null hypothesis were not true, the test might not be sensitive enough to detect the effect. Knowing the effects were small tells us that the test was not insignificant merely because of a small sample size.

16.43: $n = \left(\frac{(1.96)(3000)}{600}\right)^2 = 96.04$. Take $n = 97$.

16.45: (a) This test has a 20% chance of rejecting H_0 when the alternative is true. (b) If the test has 20% power, then when the alternative is true, it will fail to reject H_0 80% of the time. (c) The sample sizes are very small, which typically leads to low-power tests.

16.47: From the applet, against the alternative $\mu = 8$, power = 0.609. Against the alternative $\mu = 10$, power = 0.994.

16.49: (a) We reject H_0 at the 5% level when $z \geq 1.96$ or $z \leq -1.96$. (b) Here, $z = \frac{\bar{x} - 10.1}{0.1/\sqrt{6}} = 24.4949(\bar{x} - 10.1)$. Hence, we reject H_0 if $24.4949(\bar{x} - 10.1) \leq -1.96$ or if $24.4949(\bar{x} - 10.1) \geq 1.96$. Equivalently (solving for \bar{x}), we reject H_0 if $\bar{x} \leq 10.2$ or $\bar{x} \geq 10.18$. (c) When $\mu = 10.15$, the power is $P(\bar{x} \leq 10.02) + P(\bar{x} \geq 10.18) = P(Z \leq \frac{10.02 - 10.15}{0.1/\sqrt{6}}) + P(Z \geq \frac{10.18 - 10.15}{0.1/\sqrt{6}}) = P(Z \leq -3.18) + P(Z \geq 0.74) = 0.2304$.

16.51: Power = $1 - P(\text{Type II error}) = 1 - 0.14 = 0.86$.

16.53: (a) In the long run, this probability should be 0.05. (b) If the power is 0.808, the probability of a Type II error is 0.192, so in the long run, this probability should be 0.192.

Chapter 17 From Exploration to Inference: Part II Review

17.1: (c) Hives with bees; Hives with no bees; No hives

17.3: (b)

17.5: (a) The subjects were not assigned to exercise type.

17.7: Many answers are possible. One possible lurking variable is “student attitude about purpose of college” (students with a view that college is about partying rather than studying may be more likely to binge drink and more likely to have lower grades).

17.9: Question A had 60% favoring a tax cut, while Question B had 22% favoring a tax cut.

17.11: (b)

17.13: (a)

17.15: (d) It's not a random sample, and those walking at night probably have a different view of campus safety than those that the campus community broadly defined.

17.17: (a) $1 - 0.66 - 0.21 - 0.07 - 0.04 = 0.02$.

17.19: $Y > 1$, or $Y \geq 2$. $P(Y \geq 2) = 1 - 0.26 = 0.74$.

17.21: (d) $1 - 0.33 = 0.67$.

17.23: $P(X \leq 2)$ is the probability of women giving birth to 2 or fewer children during their childbearing years. $P(X \leq 2) = 0.193 + 0.174 + 0.344 = 0.711$.

17.25: $P(X \geq 3) = 0.181 + 0.074 + 0.034 = 0.289$.

17.27: (a) The height of the density curve is $1/5 = 0.2$, since the area under the density function must be 1.

17.29: (b) This is a personal probability.

17.31: (c) mean = 100, standard deviation = $15/\sqrt{60} = 1.94$ (rounded).

17.33: The answer in 17.30 would change, since this refers to the population distribution, which is now non-Normal. The answer in 17.31 would not change—the mean of \bar{x} is 100, and the standard deviation of \bar{x} is 1.94, regardless of the population distribution. The answer in 17.32 would essentially not change. The central limit theorem tells us that the sampling distribution of \bar{x} is approximately Normal when n is large enough (and 60 should be large enough), no matter what the population distribution.

17.35: If the population we're sampling from is heavily skewed, then a larger sample is required for the central limit theorem to apply. Hence, if $n = 15$, the sampling distribution of \bar{x} may not be approximately Normal, but if $n = 150$, it will surely be approximately Normal.

17.37: (a) $11,479/14,099 = 0.8142$. (b) $6457/(6457 + 1818) = 0.7803$.

17.39: (a) $1 - P(\text{failure}) = 1 - P(\text{both components fail}) = 1 - (0.20)(0.03) = 0.994$.

17.41: (c) The mean is $1000(0.63) = 630$; standard deviation is $\sqrt{1000(0.63)(1 - 0.63)} = 15.27$.

17.43: (c)

17.45: 334.37 to 379.63.

17.47: (b)

17.49: (c)

17.51: (c)

17.53: (c) $P\text{-value} = 0.0721$.

17.55: (d)

17.57: We test $H_0: \mu = 100$ vs. $H_a: \mu < 100$; $z = \frac{87.6 - 100}{15/\sqrt{113}} = -8.79$; $P\text{-value} \approx 0$; Overwhelming evidence that the mean IQ for the very low-birth-weight population is less than 100.

17.59: (c)

17.61: Here, $r^2 = 0.61$ means that 61% of the total variability in number of wildfires is explained by our model (by know-

ing the year). If there is really no relationship between number of fires and year (a surrogate for population here), then an observed linear relationship in our data as strong as that observed ($r^2 = 0.61$) would have been very unlikely to occur by chance alone. It seems reasonable to conclude that “year” and “wildfires” are positively associated. However, a cause-and-effect conclusion is not possible.

Supplementary Exercises

17.63: Placebos do work with real pain, so the placebo response tells nothing about physical basis of the pain.

17.65: (a) Increase. (b) Decrease. (c) Increase. (d) Decrease.

Note: *The first and third statements make an argument in favor of a national health insurance system, while the second and fourth suggest reasons to oppose it.*

17.67: (a) Factors: storage method (three levels: fresh, room temperature for one month, refrigerated for one month) and preparation method (two levels: cooked immediately, or after one hour). There are six treatments. Response variables: the tasters' color and flavor ratings. (b) Randomly allocate n potatoes to each of the six groups, then compare ratings. (c) For each taster, randomly choose the order in which the fries are tasted.

	Cooked immediately	Wait one hour
Fresh	1	2
Stored	3	4
Refrigerated	5	6

17.69: (a) All probabilities are between 0 and 1, and their sum is 1. (b) Let R_1 be Taster 1's rating and R_2 be Taster 2's rating. $P(R_1 = R_2) = 0.03 + 0.08 + 0.25 + 0.20 + 0.06 = 0.62$. (c) $P(R_1 > R_2) = 0.19$. $P(R_2 > R_1) = 0.19$.

17.71: (a) Out of 100 BMIs, nearly all should be in the range $\mu \pm 3\sigma = 27 \pm 22.5 = 4.5$ to 49.5. (b) The sample mean \bar{x} has a $N(\mu, \sigma/\sqrt{100}) = N(27, 0.75)$ distribution, so nearly all such means should be in the range $27 \pm 3(0.75) = 27 \pm 2.25$, or 24.75 to 29.25.

17.73: (a) This is an observational study: behavior is observed, but no treatment is imposed. (b) “Significant” means unlikely to happen by chance. (c) Answers will vary. For example, some nondrinkers might avoid drinking because of other health concerns.

17.75: (a) The stemplot confirms the description given in the text. (b) Let μ be the mean body temperature. $H_0: \mu = 98.6^\circ$ vs. $H_a: \mu \neq 98.6^\circ$. Assume we have a Normal

distribution and an SRS. $z = \frac{98.203 - 98.6}{0.7/\sqrt{20}} = -2.54$. $P\text{-value} = 2P(Z < -2.54) = 0.0110$. We have fairly strong evidence—significant at $\alpha = 0.05$, but not at $\alpha = 0.01$ —that mean body temperature is not equal to 98.6° .

96	8
97	344
97	888889
98	0133
98	5789
99	
99	6
100	2

17.77: Assume we have a Normal distribution and an SRS. Our 90% confidence interval for μ is $98.203 \pm 1.645(\frac{0.7}{\sqrt{20}}) = 98.203 \pm 0.257$, or 97.95° to 98.46° . We are 90% confident that the mean body temperature for healthy adults is between 97.95° and 98.46° .

17.79: A low-power test has a small probability of rejecting the null hypothesis, at least for some alternatives. That is, we run a fairly high risk of making a Type II error (failing to reject H_0 when it is false) for such alternatives. Knowing that this can happen, we should not conclude that H_0 is “true” simply because we failed to reject it.

Chapter 18 Inference about a Population Mean

18.1: $s/\sqrt{n} = 63.9/\sqrt{1000} = 2.0207$ minutes.

18.3: (a) $t^* = 2.132$. (b) $t^* = 2.479$.

18.5: (a) $df = 12 - 1 = 11$, so $t^* = 2.201$. (b) $df = 18 - 1 = 17$, so $t^* = 2.898$. (c) $df = 6 - 1 = 5$, so $t^* = 2.015$.

18.7: We are told to view the observations as an SRS. A stemplot shows some left-skewness; however, for such a small sample, the data are not unreasonably skewed. There are no outliers. $t^* = 1.860$ ($df = 8$); $59.5889 \pm 1.860\frac{6.2553}{\sqrt{9}} = 55.71\%$ to 63.47% . We are 90% confident that the mean percent of nitrogen in ancient air is between 55.71% and 63.47% .

18.9: (a) $df = 15 - 1 = 14$. (b) $t = 2.12$ is bracketed by $t^* = 1.761$ (with two-tail probability 0.10) and $t^* = 2.145$ (with two-tail probability 0.05). Since this is a two-sided significance test, $0.05 < P < 0.10$. (c) This test is significant at the 10% level since the $P < 0.10$. It is not significant at the 5% level since the $P > 0.05$. (d) From software, $P = 0.0524$.

18.11: Let μ be the mean difference (monkey call minus pure tone) in firing rate. $H_0: \mu = 0$ vs. $H_0: \mu > 0$. We must assume that the monkeys can be regarded as an SRS. For each monkey, we compute the call minus pure tone differences; a

stemplot of these differences shows no outliers or deviations from Normality. $t = \frac{70.378 - 0}{88.447/\sqrt{37}} = 4.84$ with $df = 36$. This has a very small $P\text{-value}$: $P < 0.0001$. We have very strong evidence that macaque neural response to monkey calls is stronger than the response to pure tones.

18.13: A stemplot suggests that the distribution of nitrogen contents is heavily skewed. Although t procedures are robust, they should not be used if the population being sampled from is this heavily skewed. In this case, t procedures are not reliable.

18.15: (b) We virtually never know the value of σ .

18.17: (c) $df = 25 - 1 = 24$.

18.19: (a) 2.718. Here, $df = 11$.

18.21: (c) $85 \pm 3.250\frac{12}{\sqrt{10}}$.

18.23: (b) If you sample 64 unmarried male students and then sample 64 unmarried female students, no matching is present.

18.25: For the student group: $t = \frac{0.08 - 0}{0.37/\sqrt{12}} = 0.749$ (not 0.49, as stated). For the nonstudent group: $t = \frac{0.35 - 0}{0.37/\sqrt{12}} = 3.277$ (rather than 3.25, a difference that might be due to rounding error). From Table C, the first $P\text{-value}$ is between 0.4 and 0.5 (software gives 0.47), and the second $P\text{-value}$ is between 0.005 and 0.01 (software gives 0.007).

18.27: (a) The sample size is very large, so the only potential hazard is extreme skewness. Since scores range only from 0 to 500, there is a limit to how skewed the distribution could be.

(b) From Table C, we take $t^* = 2.581$ ($df = 1000$), or using software take $t^* = 2.5775$. For either value of t^* , the 99% confidence interval is $250 \pm 2.581 = 247.4$ to 252.6 . (c) Because the 99% confidence interval for μ does not contain 243 and is entirely above 243, we would fail to reject $H_0: \mu = 243$ against the one-sided alternative hypothesis $H_a: \mu < 243$ at the 1% significance level.

18.29: (a) A subject's responses to the two treatments would not be independent. (b) $t = \frac{-0.326 - 0}{0.181/\sqrt{6}} = -4.41$. With $df = 5$, $P = 0.0069$, significant evidence of a difference.

18.31: (a) A stemplot is provided and suggests the presence of outliers. The sample is small and the stemplot is skewed, so use of t procedures is not appropriate. (b) We will compute two confidence intervals, as called for. In the first interval, using all 9 observations, we have $df = 8$ and $t^* = 1.860$. For the second interval, removing the two outliers (1.15 and 1.35), $df = 6$ and $t^* = 1.943$. The two 90% confidence intervals are:

$$0.549 \pm 1.860\left(\frac{0.403}{\sqrt{9}}\right) = 0.299 \text{ to } 0.799 \text{ grams}$$

and

$$0.349 \pm 1.943\left(\frac{0.053}{\sqrt{7}}\right) = 0.310 \text{ to } 0.388 \text{ grams.}$$

(c) The confidence interval computed without the two outliers is much narrower. Using fewer data values reduces degrees

of freedom (yielding a larger value of t^*). Also, smaller sample sizes yield larger margins of error. However both of these effects are offset by removing two values far from the others — s reduces from 0.403 to 0.053 by removing them.

2	5
3	3 3 5 8
4	0 0
5	
6	
7	
8	
9	
10	
11	5
12	
13	5

18.33: (a) The stemplot shows the high outlier mentioned in the text. (b) Let μ be the mean difference (control minus experimental) in healing rates. $H_0: \mu = 0$ vs. $H_a: \mu > 0$. $t = \frac{6.417 - 0}{10.7065/\sqrt{12}} = 2.08$. With $df = 11$, $P = 0.0311$ (using software). Omitting the outlier: $\bar{x} = 4.182$ and $s = 7.7565$, so $t = \frac{4.182 - 0}{7.7565/\sqrt{11}} = 1.79$. With $df = 10$, $P = 0.052$. Hence, with all 12 differences there is greater evidence that the mean healing time is greater for the control limb. When we omit the outlier, the evidence is weaker.

-1	3
-0	6
-0	
0	12
0	5789
1	012
1	
2	
2	
3	1

18.35: (a) The sample has a significant outlier, and indicates skew. We might consider applying t procedures to the sample after removing the most extreme observation (37,786). (b) If we remove the largest observation, the remaining sample is not heavily skewed and has no outliers. $H_0: \mu = 7000$ vs. $H_a: \mu \neq 7000$. With the outlier removed, $\bar{x} = 11,555.16$ and $s = 6,095.015$. $t = \frac{11,555.16 - 7000}{6,095.015/\sqrt{19}} = 3.258$. With $df = 18$ with software, $P = 0.0044$ (this is a two-sided test). There is overwhelming evidence that the mean number of words per day of men at this university differs from 7000.

18.37: (a) A stemplot of differences shows an extreme right skew and one or two high outliers. The t procedures should not

be used. (b) Some students might perform the test ($H_0: \mu = 0$ vs. $H_a: \mu > 0$) using t procedures, despite the presence of strong skew and outliers in the sample. If so, they should find $\bar{x} = 156.36$, $s = 234.2952$ and $t = 2.213$, yielding $P = 0.0256$.

18.39: (a) $H_0: \mu = 0$ vs. $H_a: \mu > 0$, where μ is the mean difference (treated minus control). (b) $t = \frac{1.916 - 0}{1.050/\sqrt{3}} = 3.16$ with $df = 2$. Hence, $P = 0.044$. This is significant at the 5% significance level. (c) For very small samples, t procedures should only be used when we can assume that the population is Normal. We have no way to assess the Normality of the population based on these four observations. Hence, the validity of the analysis in (b) is dubious.

18.41: A stemplot reveals that these data contain two extreme high outliers (5973 and 8015). Hence, t procedures are not appropriate.

18.43: (a) From Table C, $t^* = 2.000$ ($df = 60$). Using software, with $df = 63$, $t^* = 1.998$. The 95% confidence interval for μ is $48.25 \pm 2.000(\frac{40.24}{\sqrt{64}}) = 38.19$ to 58.31 thousand barrels. (Using the software version of t^* , the confidence interval is almost identical: 38.20 to 58.30 thousand barrels.) (b) The stemplot confirms the skewness and outliers described in the exercise. The two intervals have similar widths, but the new interval is shifted higher by about 2000 barrels. Although t procedures are fairly robust, we should be cautious about trusting the result in (a) because of the strong skew and outliers. The computer-intensive method may produce a more reliable interval.

0	00001111111111
0	22222223333333333333
0	44444445555555
0	6666667
0	8899
1	01
1	
1	5
1	
1	9
2	0

18.45: Let μ be the mean percent of beetle-infected seeds. A stemplot shows a single-peak and roughly symmetric distribution. We assume that the 28 plants can be viewed as an SRS of the population, so t procedures are appropriate. Using $df = 27$, the 90% confidence interval for μ is $4.0786 \pm 1.703(\frac{2.0135}{\sqrt{28}}) = 3.43\%$ to 4.73% . The beetle infects less than 5% of seeds, so it is unlikely to be effective in controlling velvetleaf.

18.47: A 95% confidence interval for the mean difference in T cell counts after 20 days on blinatumomab is $0.5283 \pm 2.517(\frac{0.4574}{\sqrt{6}}) = 0.5283 \pm 0.4801 = 0.0482$ to 1.0084 thousand cells.

18.49: (a) For each subject, randomly select which knob (right or left) that subject should use first. (b) $H_0: \mu = 0$ vs. $H_a: \mu < 0$, where μ denotes the mean difference in time (right-thread time–left-thread time), so that $\mu < 0$ means “right-hand time is less than left-hand time on average.” A stemplot of the differences gives no reason that t procedures are not appropriate. We assume our sample can be viewed as an SRS. $t = \frac{-13.32 - 0}{22.936/\sqrt{25}} = -2.90$. With $df = 24$ we find $P = 0.0039$. We have good evidence (significant at the 1% level) that the mean difference really is negative—that is, the mean time for right-hand-thread knobs is less than the mean time for left-hand-thread knobs.

18.51: With $df = 24$, $t^* = 1.711$, so the confidence interval for μ is given by $-13.32 \pm 1.711(\frac{22.936}{\sqrt{25}}) = -13.32 \pm 7.85 = -21.2$ to -5.5 seconds. Now $\bar{x}_{RH}/\bar{x}_{LH} = 104.12/117.44 = 0.887$. Hence, right-handers working with right-handed knobs can accomplish the task in about 89% of the time needed by those working with left-handed knobs.

Chapter 19 Two-Sample Problems

19.1: This is a matched pairs design. Each couple is a matched pair.

19.3: This involves a single sample.

19.5: (a) If the loggers had known that a study would be done, they might have (consciously or subconsciously) cut down fewer trees than they typically would, in order to reduce the impact of logging. (b) $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 > \mu_2$, where μ_1 is the mean number of species in unlogged plots and μ_2 is the mean number of species in plots logged 8 years earlier. We assume that the data come from SRSs of the two populations. Stemplots suggest some deviation from Normality and a possible low outlier for the logged-plot counts, but there is not strong evidence of non-Normality in either sample. With $\bar{x}_1 = 17.50$, $\bar{x}_2 = 13.67$, $s_1 = 3.53$, $s_2 = 4.50$, $n_1 = 12$, and $n_2 = 9$: $SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 1.813$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{SE} = 2.11$. Using df as the smaller of $9 - 1$ and $12 - 1$, we have $df = 8$ and $0.025 < P < 0.05$. Using software, $df = 14.8$ and $P = 0.026$. There is strong evidence that the mean number of species in unlogged plots is greater than that for logged plots 8 years after logging.

19.7: $\bar{x}_1 = 17.50$, $\bar{x}_2 = 13.67$, and $SE = 1.813$. Using $df = 8$, $t^* = 3.355$. A 99% confidence interval for the mean difference in number of species in unlogged and logged plots is $\bar{x}_1 - \bar{x}_2 \pm t^*SE = -2.253$ to 9.913 species.

19.9: (a) Back-to-back stemplots of the time data are shown below. They appear to be reasonably Normal, and the discussion in the exercise justifies our treating the data as independent SRSs, so we can use the t procedures. $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 < \mu_2$, where μ_1 is the population mean time in the restaurant with no

scent, and μ_2 is the mean time in the restaurant with a lavender odor. With $\bar{x}_1 = 91.27$, $\bar{x}_2 = 105.700$, $s_1 = 14.930$, $s_2 = 13.105$, $n_1 = 30$, and $n_2 = 30$: $SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 3.627$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{SE} = -3.98$. Using software, $df = 57.041$ and $P = 0.0001$. Using the more conservative $df = 29$ (lesser of $30 - 1$ and $30 - 1$) and Table C, $P < 0.0005$. There is very strong evidence that customers spend more time on average in the restaurant when the lavender scent is present. (b) Back-to-back stemplots of the spending data are below. The distributions are skewed and have many gaps. $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 < \mu_2$, where μ_1 is the population mean amount spent in the restaurant with no scent and μ_2 is the mean amount spent in the restaurant with lavender odor. With $\bar{x}_1 = \$17.5133$, $\bar{x}_2 = \$21.1233$, $s_1 = \$2.3588$, $s_2 = \$2.3450$, $n_1 = 30$, and $n_2 = 30$: $SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = \0.6073 and $t = \frac{\bar{x}_1 - \bar{x}_2}{SE} = -5.95$. Using software, $df = 57.041$ and $P < 0.0001$. Using the more conservative $df = 29$ and Table C, $P < 0.0005$. There is very strong evidence that customers spend more money on average when the lavender scent is present.

No scent	Lavender
98	6
322	7
965	7 6
44	8
7765	8 89
32221	9 234
86	9 578
31	10 1234
9776	10 5566788999
	11 4
85	11 6
1	12 14
	12 69
	13
	13 7
No scent	Lavender
9	12
	13
	14
99999999999999	15
	16
	17
555555555555	18 555555555555
	19
5	20 7
9	21 5599999999
	22 3558
	23
	24 99
5	25 59

19.11: We have two small samples ($n_1 = n_2 = 4$), so the t procedures are not reliable unless both distributions are Normal.

19.13: Here are the details of the calculations:

$$\text{SE}_F = \frac{12.6961}{\sqrt{31}} \doteq 2.2803$$

$$\text{SE}_M = \frac{12.2649}{\sqrt{47}} \doteq 1.7890$$

$$\text{SE} = \sqrt{\text{SE}_F^2 + \text{SE}_M^2} \doteq 2.8983$$

$$\text{df} = \frac{\text{SE}^4}{\frac{1}{30}\left(\frac{12.6961^2}{31}\right)^2 + \frac{1}{46}\left(\frac{12.2649^2}{47}\right)^2} = \frac{70.565}{1.1239} \doteq 62.8$$

$$t = \frac{55.5161 - 57.9149}{\text{SE}} \doteq -0.8276.$$

19.15: Reading from the software output shown in the statement of Exercise 19.13, we find that there was no significant difference in mean Self-Concept Scale scores for men and women ($t = -0.8276$, $\text{df} = 62.8$, $P = 0.4110$).

19.17: (a) We have two independent populations: females and males.

19.19: (b)

19.21: (b)

19.23: (a) We suspect that younger people use social networks more than older people, so this is a one-sided alternative.

19.25: (a) $H_0: \mu_M = \mu_F$ vs. $H_a: \mu_M < \mu_F$. (b)–(d) The small table below provides a summary of t statistics, degrees of freedom, and P -values for both studies. The two-sample t statistic is computed as $t = \frac{\bar{x}_M - \bar{x}_F}{\sqrt{\frac{s_M^2}{n_M} + \frac{s_F^2}{n_F}}}$, and we take the conservative approach for computing df as the smaller sample size minus 1.

Study	t	df	Table C values	P-value
1	-0.248	55	$ t < 0.679$	$P > 0.25$
2	1.507	19	$1.328 < t < 1.729$	$0.05 < P < 0.10$

Note that for Study 1 we reference $\text{df} = 50$ in Table C. (e) The first study gives no support to the belief that women talk more than men; the second study gives weak support, significant only at a relatively high significance level (say $\alpha = 0.10$).

19.27: (a) Call group 1 the Stress group and group 2 the No stress group. Then, since $\text{SEM} = s/\sqrt{n}$, we have $s = \text{SEM}\sqrt{n}$. Hence, $s_1 = 3\sqrt{20} = 13.416$ and $s_2 = 2\sqrt{51} = 14.283$. (b) Using the conservative Option 2, $\text{df} = 19$ (the lesser of 20 and 51, minus 1). (c) $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$.

With $n_1 = 20$ and $n_2 = 51$, $\text{SE} = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 3.605$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{\text{SE}} = \frac{26 - 32}{3.605} = -1.664$. With $\text{df} = 19$, using Table C, $0.10 < P < 0.20$. There is little evidence in support of a conclusion that mean weights of rats in stressful environments differ from those of rats without stress.

19.29: (a) A placebo is an inert pill that allows researchers to account for any psychological benefit (or detriment) the subject might get from taking a pill. (b) Neither the subjects nor the researchers who worked with them knew who was getting ginkgo extract; this prevents expectations or prejudices from affecting the evaluation of the effectiveness of the treatment. (c) $\text{SE} = \sqrt{\frac{0.01462^2}{21} + \frac{0.01549^2}{18}} = 0.0048$; $t = \frac{0.06383 - 0.05342}{\text{SE}} = 2.147$. This is significant at the 5% level: $P = 0.0387$ ($\text{df} = 35.35$) or $0.04 < P < 0.05$ ($\text{df} = 17$). There is strong evidence that those who take gingko extract average more misses per line.

19.31: Let μ_1 be mean for people with Asperger Syndrome and let μ_2 be the mean for people without Asperger Syndrome. $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$. Here $\bar{x}_1 = -0.001$, $\bar{x}_2 = 0.42$, $n_1 = 19$, and $n_2 = 17$. Since $\text{SEM} = s/\sqrt{n}$, we have $s = \text{SEM}\sqrt{n}$. Hence, $s_1 = 0.15\sqrt{19} = 0.6538$ and $0.17\sqrt{17} = 0.7009$. $\text{SE} = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 0.2267$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{\text{SE}} = -1.857$. Using the conservative version for df (Option 2), $\text{df} = 16$ and $0.05 < P < 0.10$. Using software, $\text{df} = 32.89$ and $P = 0.0723$. There is strong evidence that the mean score for Asperger Syndrome population is different from that of the non-Asperger population.

19.33: (a) $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 > \mu_2$, where μ_1 is the mean gain among all coached students and μ_2 is the mean gain among uncoached students. $\text{SE} = \sqrt{\frac{59^2}{427} + \frac{52^2}{2733}} = 3.0235$ and $t = \frac{29 - 21}{3.0235} = 2.646$. Using the conservative approach, $\text{df} = 426$ is rounded down to $\text{df} = 100$ in Table C and we obtain $0.0025 < P < 0.005$. Using software, $\text{df} = 534.45$ and $P = 0.004$. There is evidence that coached students had a greater average increase. (b) $8 \pm t^*(3.0235)$ where t^* equals 2.626 (using $\text{df} = 100$ with Table C) or 2.585 ($\text{df} = 534.45$ with software). This gives either 0.06 to 15.94 points, or 0.184 to 15.816 points, respectively. (c) Increasing one's score by 0 to 16 points is not likely to make a difference in being granted admission or scholarships from any colleges.

19.35: (a) Neither sample histogram suggests strong skew or presence of far outliers. t procedures are reasonable here. (b) Let μ_1 be the mean tip percentage when the forecast is good and μ_2 be the mean tip percentage when the forecast is bad. $\bar{x}_1 = 22.22$, $\bar{x}_2 = 18.19$, $s_1 = 1.955$, $s_2 = 2.105$, $n_1 = 20$, and $n_2 = 20$. $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$. $\text{SE} = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 0.642$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{\text{SE}} = 6.274$. Using $\text{df} = 19$ (the conservative

Option 2) and Table C, we have $P < 0.001$. Using software, $df = 37.8$, and $P < 0.00001$. There is overwhelming evidence that the mean tip percentage differs between the two types of forecasts presented to patrons.

19.37: $df = 19$, $t^* = 2.093$. The 95% confidence interval for the difference in mean tip percentages between these two populations is $22.22 - 18.19 \pm 2.093(0.642) = 2.69$ to 5.37 percent. Using $df = 37.8$ with software, the corresponding 95% confidence interval is 2.73% to 5.33%.

19.39: (a) The Hylite mean is greater than the Permafresh mean. (b) Shown are back-to-back stemplots for the two processes, which confirm that there are no extreme outliers. (c) $SE = 1.334$ and $t = -6.296$. $0.002 < P < 0.005$ (using $df = 4$) or P -value = 0.0003 (using software, with $df = 7.779$). There is very strong evidence of a difference between the population means. As we might expect, the stronger process (Permafresh) is less resistant to wrinkles.

	<i>n</i>	\bar{x}	<i>s</i>
Permafresh	5	134.8	1.9235
Hylite	5	143.2	2.2804

Permafresh		Hylite
2	13	
54	13	
76	13	
	13	
	14	11
	14	3
	14	5
	14	6

19.41: The 90% confidence interval is $\bar{x}_1 - \bar{x}_2 \pm t^* SE$, where $t^* = 2.132$ ($df = 4$) or $t^* = 1.867$ ($df = 7.779$). This gives either

$$-8.4 \pm 2.844 = -11.244 \text{ to } -5.556 \text{ degrees} \quad (\text{with } df = 4) \text{ or}$$

$$-8.4 \pm 2.491 = -10.891 \text{ to } -5.909 \text{ degrees} \quad (\text{with } df = 7.779).$$

19.43: This is a two-sample *t* statistic, comparing two independent groups (supplemented and control). Using the conservative $df = 5$, $t = -1.05$ would have a *P*-value between 0.30 and 0.40, which (as the report said) is not significant. The test statistic $t = -1.05$ would not be significant for any value of *df*.

19.45: These are paired *t* statistics: for each bird, the number of days behind the caterpillar peak was observed, and the *t* values were computed based on the pairwise differences between the first and second years. For the control group, $df = 5$, and for the supplemented group, $df =$

6. The control *t* is not significant (so the birds in that group did not “advance their laying date in the second year”), while the supplemented group *t* is significant with one-sided $P = 0.0195$ (so those birds did change their laying date).

19.47: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 > \mu_2$, where μ_1 is the mean weight loss for adolescents in the gastric banding group and μ_2 is the mean time for the lifestyle intervention group. We must assume that the data come from an SRS of the intended population; we cannot check this with the data. Stemplots for each sample show no heavy skew and no outliers. With $\bar{x}_1 = 34.87$, $\bar{x}_2 = 3.01$, $s_1 = 18.12$, $s_2 = 13.22$, $n_1 = 24$, and $n_2 = 18$ (note that not all subjects completed the study), $SE = \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}} = 4.84$ and $t = \frac{314.0588 - 186.1176}{SE} = 6.59$, for which $P < 0.0005$ ($df = 17$) or $P < 0.00001$ ($df = 39.98$ using software). There is strong evidence that adolescents using gastric banding lose more weight on average than those that use lifestyle modification.

19.49: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 > \mu_2$, and find a 90% confidence interval for $\mu_1 - \mu_2$, where μ_1 is the mean for the treatment population and μ_2 is the mean for the control population. We must assume that we have two SRSs and that the distributions of score improvements are Normal. Back-to-back stemplots of the differences (“after” minus “before”) for the two groups; the samples are too small to assess Normality, but there are no outliers. With $\bar{x}_1 = 11.4$, $\bar{x}_2 = 8.25$, $s_1 = 3.1693$, $s_2 = 3.6936$, $n_1 = 10$, and $n_2 = 8$, we find $SE = 1.646$ and $t = 1.914$. With $df = 7$, $0.025 < P < 0.05$. With $df = 13.92$ (software), $P = 0.0382$. The 90% confidence interval is $(11.4 - 8.25) \pm t^* SE$, where $t^* = 1.895$ ($df = 7$) or $t^* = 1.762$ ($df = 13.92$): either 0.03 to 6.27 points or 0.25 to 6.05 points. We have fairly strong evidence that the encouraging subliminal message led to a greater improvement in math scores, on average. We are 90% confident that this increase is between 0.03 and 6.27 points (or 0.25 and 6.05 points).

19.51: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$, and find a 95% confidence interval for $\mu_1 - \mu_2$, where μ_1 is the mean for the Red population and μ_2 is the mean for the Yellow population. We must assume that the data come from an SRS. We also assume that the data are close to Normal. Back-to-back stemplots show some skewness in the red lengths, but the *t* procedures should be reasonably safe. With $\bar{x}_1 = 39.7113$, $\bar{x}_2 = 36.1800$, $s_1 = 1.7988$, $s_2 = 0.9753$, $n_1 = 23$, and $n_2 = 15$, we find $SE = 0.4518$ and $t = 7.817$. With either $df = 14$ or $df = 35.10$, $P < 0.0001$. The 95% confidence interval is $(39.711 - 36.180) \pm t^* SE$, where $t^* = 2.145$ ($df = 14$) or $t^* = 2.030$ ($df = 35.1$): either 2.562 to 4.500 mm or 2.614 to 4.448 mm. We have very strong evidence that the two

varieties differ in mean length. We are 95% confident that the mean red length minus yellow length is between 2.562 and 4.500 mm (or 2.614 and 4.448 mm).

Chapter 20 Inference for a Population Proportion

20.1: (a) The population consists of all persons between the ages of 18 and 30 living in the United States. The parameter p is the proportion of this population that prays at least once a week. (b) $\hat{p} = \frac{247}{385} = 0.6416$.

20.3: (a) Approximately Normal with mean $p = 0.70$ and standard deviation $\sqrt{\frac{0.70(1 - 0.70)}{1500}} = 0.0118$. (b) Approximately Normal with mean $p = 0.70$ and standard deviation $\sqrt{\frac{0.70(1 - 0.70)}{6000}} = 0.0059$. Notice that quadrupling the sample size (from 1500 to 6000) results in halving the standard deviation of \hat{p} (0.0059 is one-half of 0.0118).

20.5: (a) The survey excludes residents of the northern territories, as well as those who have no phones or have only cell phone service. (b) $\hat{p} = \frac{1288}{1505} = 0.8558$ so $SE_{\hat{p}} = \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} = 0.009055$. The 95% confidence interval is $0.8558 \pm (1.96)(0.009055) = 0.8381$ to 0.8736.

20.7: (a) Among the 14 observations, we have 11 successes and 3 failures. The number of successes and failures should both be at least 15 for the Normal approximation to be valid. (b) We add 4 observations: 2 successes and 2 failures. We now have 18 observations: 13 successes and 5 failures. Now $\tilde{p} = \frac{11 + 2}{14 + 4} = \frac{13}{18} = 0.7222$. (c) Using the plus four method, $SE_{\tilde{p}} = \sqrt{\frac{\tilde{p}(1 - \tilde{p})}{n + 4}} = \sqrt{\frac{0.7222(1 - 0.7222)}{18}} = 0.1056$. A 90% confidence interval for p is $0.7222 \pm 1.645(0.1056) = 0.5485$ to 0.8959.

20.9: (a) $\hat{p} = \frac{20}{20} = 1$, so $SE_{\hat{p}} = \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} = 0$. The margin of error would therefore be 0 (regardless of the confidence level), so large-sample methods give the useless interval 1 to 1. (b) The plus four estimate is $\tilde{p} = \frac{22}{24} = 0.9167$, and $SE_{\tilde{p}} = \sqrt{\frac{\tilde{p}(1 - \tilde{p})}{24}} = 0.0564$. A 95% confidence interval for p is then $0.9167 \pm 1.96(0.0564) = 0.8062$ to 1.0272. Since proportions can't exceed 1, we say that a 95% confidence interval for p is 0.8061 to 1.

20.11: $n = (\frac{z^*}{m})^2 p^*(1 - p^*) = (\frac{1.645}{0.04})^2 (0.75)(1 - 0.75) = 317.1$, so use $n = 318$.

20.13: Let p be the proportion of times the “best face” wins. $H_0: p = 0.50$ vs. $H_a: p > 0.50$. Since the sample consists of 32 trials, we expect 16 “successes” (best face wins) and 16 “failures” (best face does not win). The sample is large enough to use the Normal approximation to describe the sampling distribution of \hat{p} . We assume the sample is an SRS. $\hat{p} = \frac{22}{32} = 0.6875$,

and $SE_{\hat{p}} = \sqrt{\frac{0.50(1 - 0.50)}{32}} = 0.0884$. $z = \frac{\hat{p} - p_0}{SE_{\hat{p}}} = \frac{0.6875 - 0.50}{0.0884} = 2.12$, and $P = 0.0170$. There is strong evidence that the proportion of times the “best face” wins is more than 0.50.

20.15: (b)

20.17: (c) $\hat{p} = \frac{1410}{3000} = 0.47$.

20.19: (c) $n = (\frac{z^*}{m})^2 p^*(1 - p^*) = (\frac{2.58}{0.02})^2 (0.50)(1 - 0.50) = 4147.36$; round up to $n = 4148$.

20.21: (a) Sources of bias are not accounted for in a margin of error.

20.23: (c) $z = \frac{0.53 - 0.50}{\sqrt{\frac{0.50(1 - 0.50)}{100}}} = 0.60$.

20.25: (a) The survey excludes those who have no phones or have only cell phone service. (b) Note that we have plenty of successes and plenty of failures, so conditions for large-sample confidence interval are met. $\hat{p} = \frac{848}{1010} = 0.8396$; the large-sample 95% confidence interval is $\hat{p} \pm z^* \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} = 0.8396 \pm 1.96 \sqrt{\frac{0.8396(1 - 0.8396)}{1010}} = 0.8170$ to 0.8622. If we instead use the plus four method, $\tilde{p} = \frac{848 + 2}{1010 + 4} = 0.8383$, $SE_{\tilde{p}} = 0.01156$, the margin of error is $1.96(0.01156) = 0.02266$, and the 95% confidence interval is 0.8156 to 0.8610.

20.27: (a) $\hat{p} = \frac{848}{1010} = 0.8396$, $SE_{\hat{p}} = 0.01155$, so the margin of error is $1.96SE_{\hat{p}} = 0.02263 = 2.26\%$. (b) If instead $\hat{p} = 0.50$, then $SE_{\hat{p}} = 0.01573$ and the margin of error for 95% confidence would be $1.96 SE_{\hat{p}} = 0.03084 = 3.08\%$. (c) For samples of about this size, the margin of error is no more than about $\pm 3\%$ no matter what \hat{p} is.

20.29: (a) The survey excludes residents of Alaska and Hawaii and those who do not have cell phone service. (b) We have 422 successes and 2063 failures, so the sample is large enough to use large-sample inference procedure. We have $\hat{p} = \frac{422}{2485} = 0.1698$, and $SE_{\hat{p}} = 0.0075$. For 90% confidence, the margin of error is $1.645SE_{\hat{p}} = 0.0124$ and the confidence interval is 0.1574 to 0.1822, or 15.7% to 18.2%. Using the plus four method, $\tilde{p} = \frac{422 + 2}{2485 + 4} = 0.1703$, $SE_{\tilde{p}} = 0.0075$, the margin of error is $1.645SE_{\tilde{p}} = 0.0124$, and the 90% confidence interval is 0.1579 to 0.1827, or 15.8% to 18.3%. (c) Perhaps people that use the cell phone to search for information online are younger and interested in more sexually related topics.

20.31: (a) In order to construct a large-sample confidence interval, we require at least 15 successes (swimming areas with unsafe levels of fecal coliform) and at least 15 failures (swimming areas with safe levels of fecal coliform). Here we have 13 successes and 7 failures. In order to use the plus four confidence intervals, we require at least 90% confidence and at least 10 trials. Hence, conditions for using the plus

four method are satisfied. (b) $\tilde{p} = \frac{13+2}{20+4} = 0.625$, and $SE_{\tilde{p}} = \sqrt{\frac{\tilde{p}(1-\tilde{p})}{24}} = 0.0988$. The margin of error for 90% confidence is $1.645(0.0988) = 0.1626$, and the 90% confidence interval for the proportion of swimming areas with unsafe coliform levels is 0.4624 to 0.7879, or 46.2% to 78.8%.

20.33: (a) Because the smallest number of total tax returns (i.e., the smallest population) is still more than 100 times the sample size, the margin of error will be (approximately) same for all states. (b) Yes, it will change—the sample taken from Wyoming will be about the same size, but the sample from, for example, California will be considerably larger, and therefore the margin of error will decrease.

20.35: (a) The margins of error are $1.96\sqrt{\frac{\hat{p}(1-\hat{p})}{100}} = 0.196\sqrt{\hat{p}(1-\hat{p})}$ (below). (b) With $n = 500$, the margins of error are $1.96\sqrt{\frac{\hat{p}(1-\hat{p})}{500}} = 0.088\sqrt{\hat{p}(1-\hat{p})}$. The new margins of error are less than half their former size.

p	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9
(a) Margin of error	0.0588	0.0784	0.0898	0.0960	0.0980	0.0960	0.0898	0.0784	0.0588
(b) Margin of error	0.0263	0.0351	0.0402	0.0429	0.0438	0.0429	0.0402	0.0351	0.0263

20.37: We assume that the 12 shrubs in the sample can be treated as an SRS. Because the number of resprouting shrubs is just 5, the conditions for a large sample interval are not met. Using the plus four method: $\tilde{p} = \frac{5+2}{12+4} = 0.4375$, $SE_{\tilde{p}} = 0.1240$, the margin of error is $1.645SE_{\tilde{p}} = 0.2040$, and the 90% confidence interval is 0.2335 to 0.6415. We are 90% confident that the proportion of *Krameria cytisoides* shrubs that will resprout after fire is between about 0.23 and 0.64.

20.39: Let p be the proportion of American adults who think that humans developed from earlier species of animals. We have an SRS with a very large sample size, so both large-sample and plus four methods can be used. $\tilde{p} = \frac{594}{1484} = 0.4003$, $SE_{\tilde{p}} = 0.01272$, margin of error $1.96SE_{\tilde{p}} = 0.02493$, and the 95% confidence interval is 0.3754 to 0.4252. We are 95% confident that the percent of American adults thinking that humans developed from earlier species of animals is between about 37.5% and 42.5%. Using the plus four method, we have $\tilde{p} = \frac{594+2}{1484+4} = 0.4005$, $SE_{\tilde{p}} = 0.01270$, margin of error $1.96SE_{\tilde{p}} = 0.02489$. Hence the plus four 95% confidence interval is 0.3756 to 0.4254.

20.41: Let p represent the proportion of American adults who think that humans developed from earlier species of animals. $H_0: p = 0.50$ vs. $H_a: p < 0.50$. We have an SRS with a very large sample size, so expected counts (successes and failures) are easily large enough to apply the large-sample z test.

$\hat{p} = \frac{594}{1484} = 0.4003$, so $z = \frac{0.4003 - 0.50}{\sqrt{\frac{0.50(1-0.50)}{1484}}} = -7.68$, for which

$P < 0.0001$. We have very strong evidence that fewer than half of adults believe that humans developed from earlier species of animals.

20.43: Let p be the proportion of Chick-fil-A orders correctly filled. We will assume that the 196 visits constitute a random sample of all possible visits. In our sample, we have 182 successes (correctly filled orders) and 14 failures (incorrectly filled orders). We will use the plus four method, since we do not have at least 15 failures: $\tilde{p} = \frac{182+2}{196+4} = 0.92$, $SE_{\tilde{p}} = 0.0192$, margin of error $1.96SE_{\tilde{p}} = 0.0376$, and the 95% confidence interval is 0.8824 to 0.9576. We are 95% confident that the proportion of orders filled correctly by Chick-fil-A is between 0.882 and 0.958, or 88.2% to 95.8%.

Chapter 21 Comparing Two Proportions

21.1: Let p_1 denote the proportion of younger people that text often and p_2 denote the proportion of older people. We have two large samples: 625 younger people and 1,917 older people. The number of successes in each sample (475 and 786, respectively) and the number of failures in each sample (150 and 1131) are large enough to use large-sample methods. $\hat{p}_1 = \frac{475}{625} = 0.76$, and $\hat{p}_2 = \frac{786}{1917} = 0.41$. $SE = \sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{625}} + \sqrt{\frac{\hat{p}_2(1-\hat{p}_2)}{1917}} = 0.0204$, so the margin of error for 95% confidence is $1.96(0.0204) = 0.0400$ and a 95% confidence interval for the difference in proportions is 0.3100 to 0.3900, or 31% to 39%. With 95% confidence, the proportion of teenagers that text exceeds that of persons 18 and over by somewhere between 0.31 and 0.39.

21.3: Let p_1 denote the proportion of males that meet recommended levels, and let p_2 denote the proportion for females. We have many successes and failures in both samples, so large-sample methods are reasonable. $\hat{p}_1 = \frac{3594}{7881} = 0.4560$, and $\hat{p}_2 = \frac{2261}{8164} = 0.2769$. $SE = 0.0075$, and the margin of error is $2.576SE = 0.0193$. A 99% confidence interval for the difference in proportions between males and females meeting recommended levels of physical activity is 0.1598 to 0.1984, or 16.0% to 19.8%.

21.5: Let p_1 denote the proportion of checking in the control group, and let p_2 denote the proportion of checking in the microwaved group. To use plus four methods, we want samples of at least size 5; this condition is easily met here. $\tilde{p}_1 = \frac{16+1}{65+2} = 0.2537$ and $\tilde{p}_2 = \frac{0+1}{65+2} = 0.0149$. A plus four 95% confidence interval for $p_1 - p_2$ is $\tilde{p}_1 - \tilde{p}_2 \pm 1.96\sqrt{\frac{\tilde{p}_1(1-\tilde{p}_1)}{67} + \frac{\tilde{p}_2(1-\tilde{p}_2)}{67}} = 0.1306$ to 0.3470. We are 95% confident that microwaving reduces checking by between about 13% and 35%.

21.7: Let p_1 and p_2 be (respectively) the proportions of injured skiers and injured snowboarders who wear helmets. $H_0: p_1 = p_2$ vs. $H_a: p_1 < p_2$. The smallest count is 96, so the significance testing procedure is safe. $\hat{p}_1 = \frac{96}{578} = 0.1661$ and $\hat{p}_2 = \frac{656}{2992} = 0.2193$. $\hat{p} = \frac{96 + 656}{578 + 2992} = 0.2106$. $SE = \sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{578} + \frac{1}{2992}\right)} = 0.01853$. $z = \frac{0.1661 - 0.2193}{0.01853} = -2.87$, and $P = 0.0021$. We have strong evidence (significant at $\alpha = 0.01$) that skiers and snowboarders with head injuries are less likely to use helmets than skiers and snowboarders without head injuries.

21.9: (b) We look for evidence that the proportion for 2009 is lower than 1999.

21.11: (b) $\hat{p} = \frac{511 + 592}{2411 + 2045} = 0.2475$, which rounds to 0.25.

21.13: (c) For a 95% confidence interval, the margin of error is $1.96\sqrt{\frac{\hat{p}_{1999}(1 - \hat{p}_{1999})}{2411} + \frac{\hat{p}_{2009}(1 - \hat{p}_{2009})}{2045}} = 0.026$.

21.15: (b) We have only 3 failures in the treatment group and only 2 successes in the control group.

21.17: (a) The four counts are 117, 53, 152, and 165, so all counts are large enough. (b) Using the plus four method, $\tilde{p}_1 = \frac{117 + 1}{170 + 2} = 0.6860$ and $\tilde{p}_2 = \frac{152 + 1}{317 + 2} = 0.4796$, and the 95% confidence interval is $\tilde{p}_1 - \tilde{p}_2 \pm 1.96\sqrt{\frac{\tilde{p}_1(1 - \tilde{p}_1)}{172} + \frac{\tilde{p}_2(1 - \tilde{p}_2)}{317}} = 0.2064 \pm 0.08841 = 0.1180$ to 0.2948. Using the large-sample method, $\tilde{p}_1 = \frac{117}{170} = 0.6882$, and $\tilde{p}_2 = \frac{152}{317} = 0.4795$, and the 95% confidence interval is $\tilde{p}_1 - \tilde{p}_2 \pm 1.96\sqrt{\frac{\tilde{p}_1(1 - \tilde{p}_1)}{170} + \frac{\tilde{p}_2(1 - \tilde{p}_2)}{317}} = 0.2087 \pm 0.08873 = 0.1200$ to 0.2974.

21.19: (a) One of the counts is 0; for large-sample intervals, we want all counts to be at least 10, and for significance testing, we want all counts to be at least 5. (b) The sample size for the treatment group is 35, 24 of which have tumors; the sample size for the control group is 20, 1 of which has a tumor. (c) $\tilde{p}_1 = \frac{23 + 1}{33 + 2} = 0.6857$ and $\tilde{p}_2 = \frac{0 + 1}{18 + 2} = 0.05$. The plus four 99% confidence interval is $\tilde{p}_1 - \tilde{p}_2 \pm 2.576\sqrt{\frac{\tilde{p}_1(1 - \tilde{p}_1)}{35} + \frac{\tilde{p}_2(1 - \tilde{p}_2)}{20}} = 0.3977$ to 0.8737. We are 99% confidence that lowering DNA methylation increases the incidence of tumors by between about 40% and 87%.

21.21: (a) Let p_1 and p_2 be (respectively) the proportions of subjects in the music and no music groups that receive a passing grade on the Maryland HSA. $H_0: p_1 = p_2$ vs. $H_a: p_1 \neq p_2$. $\hat{p}_1 = \frac{2818}{3239} = 0.870$. $\hat{p}_2 = \frac{2091}{2787} = 0.750$. $\hat{p} = \frac{2818 + 2091}{3239 + 2787} = 0.815$. $z = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{3239} + \frac{1}{2787}\right)}} = 11.94$. An observed difference of $0.87 - 0.75 = 0.12$ in group proportions is much too large to be explained by chance alone, and $P < 0.001$. We have overwhelming evidence that the proportion of music students passing the Maryland HSA is greater than that for the no music group. (b) and (c) This is an observational study—people that choose to (or can afford to) take music lessons differ in many

ways from those that do not. We cannot conclude that music causes an improvement in Maryland HSA achievement.

21.23: The samples are so large, either confidence interval procedure is appropriate. Using the plus four method, we have $\tilde{p}_1 = \frac{2818 + 1}{3239 + 2} = 0.870$, and $\tilde{p}_2 = \frac{2091 + 1}{2787 + 2} = 0.750$. A 95% confidence interval for $p_1 - p_2$ is then $\tilde{p}_1 - \tilde{p}_2 \pm 1.96\sqrt{\frac{\tilde{p}_1(1 - \tilde{p}_1)}{3241} + \frac{\tilde{p}_2(1 - \tilde{p}_2)}{2789}} = 0.100$ to 0.140, or 10.0% to 14.0%. With such large samples, the large sample methods are appropriate also, but will yield virtually identical results: With $\hat{p}_1 = \frac{2818}{3239} = 0.870$. For the no music group, $\hat{p}_2 = \frac{2091}{2787} = 0.750$. $\hat{p}_1 - \hat{p}_2 \pm 1.96\sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{3289} + \frac{\hat{p}_2(1 - \hat{p}_2)}{2787}} = 0.100$ to 0.140, or 10.0% to 14.0%.

21.25: (a) To test $H_0: p_M = p_F$ vs. $H_a: p_M \neq p_F$, we find $\hat{p}_M = \frac{15}{106} = 0.1415$, $\hat{p}_F = \frac{7}{42} = 0.1667$, and $\hat{p} = 0.1486$. Then $SE = \sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{106} + \frac{1}{42}\right)} = 0.06485$, so $z = \frac{\hat{p}_M - \hat{p}_F}{SE} = -0.39$. This gives $P = 0.6966$, which provides virtually no evidence of a difference in failure rates. (b) We have $\hat{p}_M = \frac{450}{3180} = 0.1415$, $\hat{p}_F = \frac{210}{1260} = 0.1667$, and $\hat{p} = 0.1486$, but now $SE = \sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{3180} + \frac{1}{1260}\right)} = 0.01184$, so $z = \frac{\hat{p}_M - \hat{p}_F}{SE} = -2.13$ and $P = 0.0332$. (c) We are asked to construct two confidence intervals—one based on the smaller samples of part (a) and one based on the larger samples of part (b). In each case, we provide both large sample and plus four intervals, which are both appropriate here. First, for case (a), $\hat{p}_M = 0.1415$ and $\hat{p}_F = 0.1667$, so a 95% confidence interval for the difference is $\hat{p}_M - \hat{p}_F \pm 1.96\sqrt{\frac{\hat{p}_M(1 - \hat{p}_M)}{106} + \frac{\hat{p}_F(1 - \hat{p}_F)}{42}} = -0.1560$ to 0.1056. Using the plus four method, $\tilde{p}_M = \frac{15 + 1}{106 + 2} = 0.1481$ and $\tilde{p}_F = \frac{7 + 1}{42 + 2} = 0.1818$, so a 95% confidence interval for the difference is $\tilde{p}_M - \tilde{p}_F \pm 1.96\sqrt{\frac{\tilde{p}_M(1 - \tilde{p}_M)}{108} + \frac{\tilde{p}_F(1 - \tilde{p}_F)}{44}} = -0.0337 \pm 0.1322 = -0.1659$ to 0.0985. For case (b), $\hat{p}_M = 0.1415$ and $\hat{p}_F = 0.1667$, but now $\hat{p}_M = \frac{450 + 1}{3180 + 2} = 0.1417$ and $\hat{p}_F = \frac{210 + 1}{1260 + 2} = 0.1672$. The resulting confidence intervals are then $\hat{p}_M - \hat{p}_F \pm 1.96\sqrt{\frac{\hat{p}_M(1 - \hat{p}_M)}{3180} + \frac{\hat{p}_F(1 - \hat{p}_F)}{1260}} = -0.0491$ to -0.0013 . The plus four interval is $\tilde{p}_M - \tilde{p}_F \pm 1.96\sqrt{\frac{\tilde{p}_M(1 - \tilde{p}_M)}{3182} + \frac{\tilde{p}_F(1 - \tilde{p}_F)}{1262}} = -0.0494$ to -0.0016 .

21.27: Let p_1 be the proportion of women who succeed, and let p_2 be the proportion of men who succeed. $H_0: p_1 = p_2$ vs. $H_a: p_1 \neq p_2$. The smallest count is 11, so the significance test should be safe. $\hat{p}_1 = \frac{23}{34} = 0.6765$ and $\hat{p}_2 = \frac{60}{89} = 0.6742$. $\hat{p} = \frac{23 + 60}{34 + 89} = 0.6748$, and $SE = \sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{34} + \frac{1}{89}\right)} = 0.09445$. $z = \frac{0.6765 - 0.6742}{0.09445} = 0.02$, $P = 0.9840$. We have no evidence to support a conclusion that women's and men's success rates differ.

21.29: Let p_1 denote the proportion of people on Chantix who abstained from smoking, and let p_2 be the corresponding proportion for the placebo population. The sample counts are 155 and 61 (successes for treatment and control, respectively) and 197 and 283 (failures for the groups), so the large-sample procedures are safe. We will apply the plus four procedure:

$\tilde{p}_1 = \frac{155+1}{352+2} = 0.4407$, $\tilde{p}_2 = \frac{61+1}{344+2} = 0.1792$, and the 99% confidence interval is $\tilde{p}_1 - \tilde{p}_2 \pm 2.576\sqrt{\frac{\tilde{p}_1(1-\tilde{p}_1)}{354} + \frac{\tilde{p}_2(1-\tilde{p}_2)}{346}} = 0.2615 \pm 0.0863 = 0.1752$ to 0.3478 . We are 99% confident that the success rate for abstaining from smoking is between 17.5 and 34.8 percentage points higher for smokers using Chantix than for smokers on a placebo. Using the large-sample method, $\hat{p}_1 = \frac{155}{352} = 0.4403$, and $\hat{p}_2 = \frac{61}{344} = 0.1773$, and the 99% confidence interval is $\hat{p}_1 - \hat{p}_2 \pm 2.576\sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{352} + \frac{\hat{p}_2(1-\hat{p}_2)}{344}} = 0.2630 \pm 0.0864 = 0.1766$ to 0.3494 .

21.31: Let p_1 be the proportion of people that will reject an unfair offer from another person, and let p_2 be the proportion for offers from a computer. $H_0: p_1 = p_2$ vs. $H_a: p_1 > p_2$. All counts are greater than 5, so the conditions for a significance test are met. $\hat{p}_1 = \frac{18}{38} = 0.4737$ and $\hat{p}_2 = \frac{6}{38} = 0.1579$. $\hat{p} = \frac{18+6}{38+38} = 0.3158$, and $SE = \sqrt{\hat{p}(1-\hat{p})(\frac{1}{38} + \frac{1}{38})} = 0.1066$. $z = \frac{0.4737 - 0.1579}{0.1066} = 2.96$, $P = 0.0015$. There is very strong evidence that people are more likely to reject an unfair offer from another person than from a computer.

21.33: Let p_1 and p_2 be (respectively) the proportions of mice ready to breed in good acorn years and bad acorn years. One count is only 7, and the guidelines for using the large-sample method call for all counts to be at least 10, so we use the plus four method. $\tilde{p}_1 = \frac{54+1}{72+2} = 0.7432$, and $\tilde{p}_2 = \frac{10+1}{17+2} = 0.5789$. The plus four 90% confidence interval is $\tilde{p}_1 - \tilde{p}_2 \pm 1.645\sqrt{\frac{\tilde{p}_1(1-\tilde{p}_1)}{74} + \frac{\tilde{p}_2(1-\tilde{p}_2)}{19}} = -0.0399$ to 0.3685 . We are 90% confident that the proportion of mice ready to breed in good acorn years is between 0.04 lower than and 0.37 higher than the proportion in bad acorn years.

21.35: (a) This is an experiment because the researchers assigned subjects to the groups being compared. (b) Let p_1 and p_2 be (respectively) the proportions that have an RV infection for the HL+ group and control group. $H_0: p_1 = p_2$ vs. $H_a: p_1 < p_2$. We have large enough counts to use a large-sample significance testing procedure safely. $\hat{p}_1 = \frac{49}{49+67} = 0.4224$, $\hat{p}_2 = \frac{49}{49+47} = 0.5104$, and $\hat{p} = \frac{49+49}{116+96} = 0.4623$. $SE = \sqrt{\hat{p}(1-\hat{p})(\frac{1}{116} + \frac{1}{96})} = 0.0688$. $z = \frac{0.4224 - 0.5104}{0.0688} = -1.28$, $P = 0.1003$. We do not have enough evidence to reject the null hypothesis; there is little evidence to conclude that the proportion of HL+ users with a rhinovirus infection is less than that for non-HL+ users.

Chapter 22 Inference about Variables: Part III Review

22.1: (c) The margin of error is $2.056(9.3)/\sqrt{27} = 3.7$.

22.3: (b) $t = 2.023$, $df = 13$.

22.5: (d) $\hat{p} = 1926/7028 = 0.274$.

22.7: (d) The standard error is 0.0068.

22.9: (a) The standard error is 0.0124. (b) A 95% confidence interval is 0.336 to 0.384.

22.11: (b) The margin of error is $2.005(3.2)/\sqrt{55} = 0.865$.

22.13: (a) df is the lesser of $(55 - 1)$ and $(200 - 1)$.

22.15: With such large samples, t procedures are reasonable.

22.17: (d) The margin of error is 3.52. The point estimate is $11.4 - 6.7 = 4.7$.

22.19: (c) Plus four confidence intervals are reliable for samples of 5 or more in each group.

22.21: (b) $\hat{p} = 225/757 = 0.297$.

22.23: (b) $0.297 \pm 1.645(0.017)$

22.25: (c) It seems reasonable that the researchers suspect that VLBW babies are less likely to graduate from high school.

22.27: (b) $z = \frac{0.7397 - 0.8283}{\sqrt{\hat{p}(1-\hat{p})(\frac{1}{242} + \frac{1}{233})}} = -2.34$.

22.29: (d) $t = \frac{86.2 - 89.8}{\sqrt{\frac{13.4^2}{38} + \frac{14^2}{54}}} = -1.25$, and the test is two-sided.

22.31: (b) $z = \frac{0.379 - 0.41}{\sqrt{0.41(1-0.41)/348}} = -1.18$, so $P = 0.1190$.

22.33: $0.58 \pm 1.645\sqrt{\frac{0.58(1-0.58)}{634}} = 0.55$ to 0.61 . A plus four interval will agree to three decimal places and would also be appropriate.

22.35: $H_0: p_b = p_w$ vs. $H_a: p_b \neq p_w$. $\hat{p} = [0.72(634) + 0.68(567)]/(634 + 567) = 0.701$,

$z = \frac{0.72 - 0.68}{\sqrt{0.70(1-0.70)(\frac{1}{634} + \frac{1}{567})}} = 1.51$ and $P = 2P(Z > 1.51) = 0.131$.

There is little evidence of a difference between black and white young people in the proportion believing that rap music videos contain too many references to sex.

22.37: $t = \frac{193 - 174}{\sqrt{\frac{68^2}{26} + \frac{44^2}{23}}} = 1.174$, and $df = 22$, the lesser of $23 - 1 = 22$ and $26 - 1 = 25$.

22.39: We must assume that each sample is an SRS taken from its respective populations (clinic dogs and pet dogs). We must also assume that the populations (cholesterol levels of pet dogs and cholesterol levels of clinic dogs) are Normal.

22.41: Two-sample test for difference in means.

22.43: If the sample can be viewed as an SRS, a t confidence interval for a population mean.

22.45: Matched pairs t test or confidence interval.

22.47: The response rate for the survey was only about 20% ($427/2100 = 0.203$), which might make the conclusions unreliable.

22.49: Each of a monkey's six trials are not independent. If a monkey prefers silence, it will almost certainly spend more time in the silent arm of the cage each time it is tested.

22.51: (a) Let p_1 be the proportion of subjects on Gardasil that get cancer, and let p_2 be the corresponding proportion for the control group. We assume that we have SRSs from each population. Because there were no cases of cervical cancer in the Gardasil group, we should use the plus four procedure. $\tilde{p}_1 = \frac{0+1}{8487+2} = 0.000118$, and $\tilde{p}_2 = \frac{32+1}{8460+2} = 0.003900$. A 99% confidence interval for $p_2 - p_1$ is then given by $\tilde{p}_2 - \tilde{p}_1 \pm 2.576\sqrt{\frac{\tilde{p}_1(1-\tilde{p}_1)}{8489} + \frac{\tilde{p}_2(1-\tilde{p}_2)}{8462}} = 0.0020$ to 0.0056. (b) Let p_1 denote the proportion in the Gardasil group with genital warts, and let p_2 be the corresponding proportion for the control group. Because we have fewer than 10 "successes" in the Gardasil group, conditions for using the large-sample interval are not met. However, we can use the plus four interval. $\tilde{p}_1 = 0.000253$, and $\tilde{p}_2 = 0.011644$. A 99% confidence interval for $p_2 - p_1$ is then 0.0082 to 0.0145. (c) Gardasil is seen to be effective in reducing the risk of both cervical cancer (by between 0.0020 and 0.0056, with 99% confidence) and genital warts (by between 0.0082 and 0.0145, with 99% confidence).

22.53: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 < \mu_2$, where μ_1 is the mean number of new leaves on plants from the control population, and μ_2 is the mean for the nitrogen population. $\bar{x}_1 = 13.2857$, $\bar{x}_2 = 15.6250$, $s_1 = 2.0587$, $s_2 = 1.6850$, $n_1 = 7$, $n_2 = 8$, $SE = \sqrt{\frac{2.0587^2}{7} + \frac{1.6850^2}{8}} = 0.9800$, $t = \frac{13.2857 - 15.6250}{SE} = -2.387$. With Option 2, $df = 6$ and $P = 0.0271$. Or, using Option 1, $df = 11.66$ and $P = 0.0175$. We have strong evidence that nitrogen increases the formation of new leaves.

22.55: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$. We view the data as coming from two SRSs; the distributions show no strong departures from Normality. $\bar{x}_1 = 48.9513$, $s_1 = 0.2154$ (cotton), $\bar{x}_2 = 41.6488$, and $s_2 = 0.3922$ (ramie). $SE = 0.1582$ and $t = 46.16$. With either $df = 7$ or $df = 10.87$ (software), $P \approx 0$. There is overwhelming evidence that cotton is lighter than ramie.

22.57: $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$. We are told that the samples may be regarded as SRSs from their respective populations. Back-to-back stemplots show that t procedures are reasonably safe, since both distributions are only slightly skewed, with no outliers and with fairly large sample sizes $\bar{x}_1 = 4.1769$, $s_1 = 2.0261$, and $n_1 = 65$ (parent allows drinking); $\bar{x}_2 = 4.5517$, $s_2 = 2.4251$, and $n_2 = 29$ (parent does not allow drinking). $SE = 0.5157$ and $t = \frac{\bar{x}_1 - \bar{x}_2}{SE} = -0.727$. This is very close to zero, so we will certainly not reject the null hypothesis. Indeed, with $df = 46.19$ (software), $P = 0.47$. There is no significant differ-

ence in the mean number of drinks between female students with a parent that allows drinking and those whose parents do not allow drinking.

22.59: (a) Stemplots are provided. The diabetic potentials appear to be larger. (b) $H_0: \mu_1 = \mu_2$ vs. $H_a: \mu_1 \neq \mu_2$, where μ_1 is the mean potential for diabetics and μ_2 is the mean for the normal population. We assume we have two SRSs; the distributions appear to be safe for t procedures. $\bar{x}_1 = 13.0896$, $\bar{x}_2 = 9.5222$, $s_1 = 4.8391$, $s_2 = 2.5765$, $n_1 = 24$, $n_2 = 18$, $SE = 1.1595$, and $t = 3.077$. With Option 2, $df = 17$ and $0.005 < P < 0.01$. Or, using Option 1, $df = 36.6$, and $P = 0.0040$. We have strong evidence that the electric potential in diabetic mice is greater than the potential in normal mice. (c) If we remove the outlier, the diabetic mouse statistics change: $\bar{x}_1 = 13.6130$, $s_1 = 4.1959$, $n_1 = 23$. Now $SE = 1.065$ and $t = 3.841$. With $df = 16$, $0.001 < P < 0.002$. With $df = 37.15$, $P = 0.0005$. With the outlier removed, the evidence that diabetic mice have higher mean electric potential is even stronger.

Diabetic	Normal
1	0
	0
	0 4
7	0 6777
988	0 8888999
1000000	1 00
3	1 233
5444	1 4
76	1
9988	1
	2
2	2

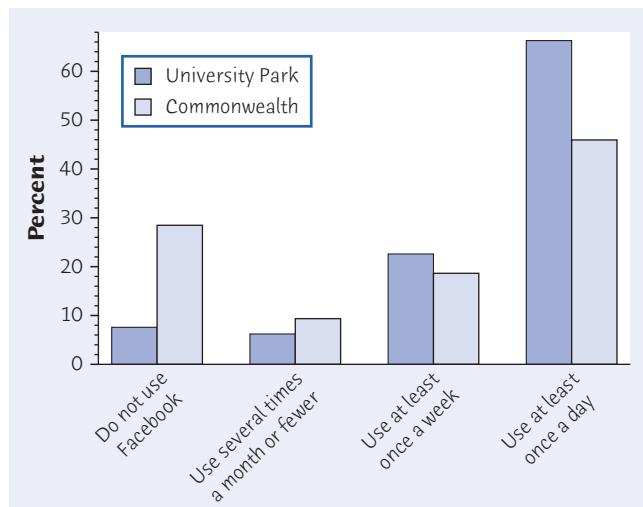
22.61: Let μ be the mean date on which the tripod falls through the ice. We assume that the data can be viewed as an SRS of fall-through times and that the distribution is roughly Normal. $n = 91$, $\bar{x} = 15.3736$, and $s = 5.9789$ days. $df = 90$. A 95% confidence interval is given by 14.13 to 16.62 days. We are 95% confident that the mean number of days for the tripod to or through the ice is 14.13 days to 16.62 days from April 20 or between May 3 and May 5.

22.63: (a) "SEM" stands for "standard error of the mean"; $SEM = s/\sqrt{n}$. (b) Two-sample t tests were done because there are two separate, independent groups of mice. (c) The observed differences between the two groups of mice were so large that it would be unlikely to occur by chance alone if the two groups were the same in average.

Chapter 23 Two Categorical Variables: The Chi-Square Test

23.1: (a) The proportion of University Park campus students that do not use Facebook is $68/978 = 0.0695$, which rounds to 0.070 and is represented as 7% in the table. (b) The bar graph reveals that students on the main campus are much more likely to use Facebook at least daily, while Commonwealth campus students are more likely not to use it at all.

	University Park	Commonwealth
Do not use Facebook	7.0%	28.3%
Use several times a month or fewer	5.6%	8.7%
Use at least once a week	22.0%	17.9%
Use at least once a day	65.4%	45.0%



23.3: (a) To test $H_0: p_1 = p_2$ vs. $H_a: p_1 \neq p_2$ for the proportions not using Facebook, we have $\hat{p}_1 = 0.0695$ and $\hat{p}_2 = 0.2834$. $\hat{p} = \frac{68 + 248}{978 + 875} = 0.1705$, $SE = 0.01750$, and $z = -12.22$, for which P is close to zero. (b) $H_0: p_1 = p_2$ vs. $H_a: p_1 \neq p_2$ for the proportions that use Facebook at least weekly, $\hat{p}_1 = 0.2198$ and $\hat{p}_2 = 0.1794$. $\hat{p} = 0.2008$, $SE = 0.01864$, $z = 2.17$, and $P = 0.0300$. (c) If we did four individual tests, we would not know how confident we could be in all four results when taken together.

23.5: (a) Expected observed counts are in the table provided. For example, $\frac{(131)(627)}{1537} = 53.44$. (b) Commonwealth students actually use Facebook less than once weekly more often than we would expect. Also, Commonwealth students use Facebook daily less often than we would expect.

	Expected counts	University Park	Commonwealth
Monthly	77.56	53.44	
Weekly	220.25	151.75	
Daily	612.19	421.81	

23.7: (a) All expected counts are well above 5 (the smallest is 53.44). (b) H_0 : there is no relationship between setting and Facebook use vs. H_a : there is a relationship between campus and Facebook use. Using software, $\chi^2 = 19.489$ and $P < 0.0005$. (c) The largest contributions come from the first row, reflecting the fact that monthly use is lower among University Park students and higher among commonwealth students.

23.9 H_0 : there is no relationship between education level and astrology opinion vs. H_a : there is some relationship between education level and astrology opinion. Examining the output provided in Figure 23.5, we see that all expected cell counts are greater than 5 and all observed cell counts are at least 1, so conditions for use of the chi-square test are satisfied. $\chi^2 = 7.244$ and $P = 0.027$. There is strong evidence of an association between education level and opinion of astrology.

23.11: (a) $df = (r - 1)(c - 1) = (3 - 1)(2 - 1) = 2$. (b) The largest critical value shown for $df = 2$ is 15.20; since the computed value (19.489) is greater than this, we conclude that $P < 0.0005$. (c) With $r = 4$ and $c = 2$, the appropriate degrees of freedom would be $df = 3$.

23.13: $H_0: p_1 = p_2 = p_3 = \frac{1}{3}$ vs. H_a : not all three are equally likely. The expected counts are each $53 \times \frac{1}{3} = 17.67$. $\chi^2 = \sum \frac{(\text{observed count} - 17.67)^2}{17.67} = \frac{(31 - 17.67)^2}{17.67} + \frac{(14 - 17.67)^2}{17.67} + \frac{(8 - 17.67)^2}{17.67} = 10.06 + 0.76 + 5.29 = 16.11$. $df = 2$. From Table D, $\chi^2 = 16.11$ falls beyond the 0.005 critical value, so $P < 0.005$. There is very strong evidence that the three tilts differ.

23.15: The details of the computation are shown below. The expected counts are found by multiplying the expected frequencies by 803 (the total number of observations).

	Expected frequency	Observed count	Expected count	$O - E$	$\frac{(O - E)^2}{E}$
16 to 29	0.328	401	263.384	137.616	71.9032
30 to 59	0.594	382	476.982	-94.982	18.9139
60 or older	0.078	20	62.634	-42.634	29.0203
		803			119.8374

The difference is significant: $\chi^2 = 119.84$, $df = 2$, and $P < 0.0005$ (using software, $P = 0.000$ to three decimal places).

23.17: $H_0: p_1 = p_2 = \dots = p_{12} = \frac{1}{12}$ vs. H_a : the 12 astrological sign birth probabilities are not equally likely. Under H_0 we expect $1960/12 = 163.33$ subjects per sign. All cells have expected counts greater than 5, and all cells have at

least one observation. A chi-square test is appropriate. $\chi^2 = \frac{(164 - 163.33)^2}{163.33} + \frac{(152 - 163.33)^2}{163.33} + \dots + \frac{(177 - 163.33)^2}{163.33} = 16.09$. $df = 12 - 1 = 11$, using Table D, $0.10 < P < 0.15$. There is little evidence that some astrological signs are more likely in birth than others.

23.19: (b) $655/(655 + 916) = 655/1571 = 0.4169$.

23.21: (a) $(1571)(1552)/4111 = 593.09$.

23.23: (a) $df = (r - 1)(c - 1) = (4 - 1)(2 - 1) = 3$.

23.25: (b) This is the hypothesis of association between “age” and “type of injury.”

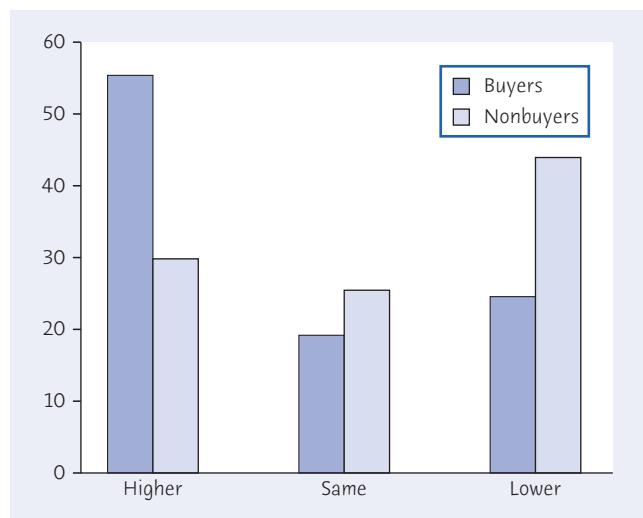
23.27: (b) We assume that the sample is an SRS, or essentially an SRS from all weightlifting injuries.

23.29: (a) The table below summarizes conditional distributions of opinion for each type of consumer. For example, there are $20 + 7 + 9 = 36$ buyers, so the proportion of buyers that think the quality of the recycled product is higher is $20/36 = 0.556$, or 55.6%.

Think the Quality of Recycled Product Is			
	Higher	Same	Lower
Buyers	55.6%	19.4%	25.0%
Nonbuyers	29.9%	25.8%	44.3%

It seems that buyers of recycled products are more likely to feel that recycled products are of higher quality, while nonbuyers are more likely to feel that recycled products are of lower quality.

(b) H_0 : No association between “opinion of quality” and “buyer status” vs. H_a : there is some association between buyer status and opinion of quality. All expected cell counts are more than 5, so the guidelines for the chi-square test are satisfied. $\chi^2 =$



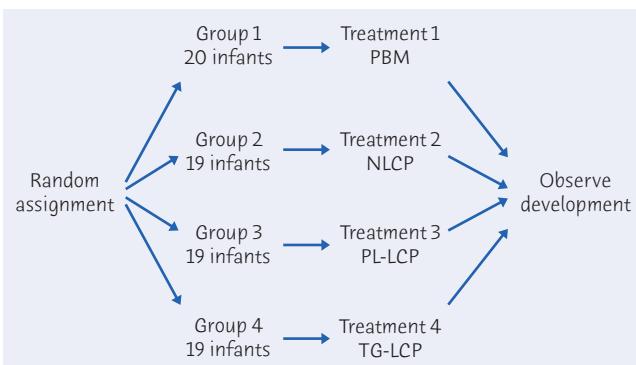
7.64, $df = 2$, and $0.02 < P < 0.025$. There is strong evidence of an association between buyer status and opinion of quality.

Expected counts	Better	Same	Lower
Buyers	13.26	8.66	14.08
Nonbuyers	35.74	23.34	37.92

(c) We see that there is a relationship between opinion of quality and whether somebody buys the recycled product. However, it is impossible to determine whether (i) prior opinion on quality drives the decision to buy or not to buy or (ii) perhaps quality of both types of products are excellent, and whichever product you happen to buy drives your opinion of that product.

23.31: (a) The diagram is shown below. To perform the randomization, label the infants 01 to 77, and choose pairs of random digits. (b) See the table for the expected counts. $\chi^2 = 0.568$, $df = 3$, and $P = 0.904$. There is no reason to doubt that the randomization “worked.”

Expected counts	Female	Male
PBM	10.91	9.09
NLCP	10.36	8.64
PL-LCP	10.36	8.64
TG-LCP	10.36	8.64



23.33: (a) $H_0: p_1 = p_2$ vs. $H_a: p_1 < p_2$. (b) The z -test must be used because the chi-square procedure measures evidence in support of evidence of any association, and is implicitly two-sided. $\hat{p}_1 = 0.3667$ and $\hat{p}_2 = 0.7333$. $\hat{p} = (11 + 22)/(30 + 30) = 0.55$, and $SE = 0.12845$, so $z = -2.85$ and $P = 0.0022$. We have strong evidence that rats that can stop the shock (and

therefore presumably have better attitudes) develop tumors less often than rats that cannot (and therefore are presumably depressed).

23.35: H_0 : there is no relationship between sexual content of ads and magazine audience vs. H_a : there is some relationship between sexual content of ads and magazine audience. Examining the Minitab output, we see that conditions for use of the chi-square test are satisfied, since all expected cell counts exceed 5. $\chi^2 = 80.874$ with $df = 2$, leading to $P < 0.0005$. Magazines aimed at women are much more likely to have sexual depictions of models than are the other two types of magazines.

23.37: We need cell counts, not just percents. If we had been given the number of travelers in each group—leisure and business—we could have estimated the counts.

23.39: In order to do a chi-square test, each subject can be counted only once.

23.41: (a) H_0 : there is no relationship between degree held and service attendance vs. H_a : there is some relationship between degree held and service attendance. Expected counts are shown in the table below. $\chi^2 = 14.19$ with $df = 3$, yielding $P\text{-value} = 0.0027$. There is strong evidence of an association between degree held and service attendance.

Expected counts	High school			
	Junior college	Bachelor's	Graduate	
Attended services	437.3	55.7	129.2	61.8
Did not attend services	842.7	107.3	248.9	119.1

(b) Expected counts are shown in the table below. $\chi^2 = 0.73$ on $df = 2$. Hence, $P > 0.25$ (0.694 by software). In this table, we find no evidence of association between religious service attendance and degree held.

Expected counts	Junior college		
	Bachelor's	Graduate	
Attended services	64.1	148.7	71.2
Did not attend services	98.9	229.3	109.8

(c) Expected counts are shown in the table below. $\chi^2 = 13.50$ on $df = 1$. Hence, $P < 0.0005$ (0.0002 by software). There is overwhelming evidence of association between level of education (High School versus Beyond High School) and religious service attendance.

Expected counts	High school	Beyond HS
Attended services	437.3	246.7
Did not attend services	842.7	475.3

(d) In general, we find that people with degrees beyond high school attend service more often than expected, while people with high school degrees attend services less often than expected. Of those with high school degrees, 31.3% attended services, while the percentages are 38.0%, 38.6% and 42.0%, respectively, for people with junior college, bachelor's, and graduate degrees.

23.43: H_0 : there is no relationship between race and opinion about schools vs. H_a : there is some relationship between race and opinion about schools. All expected cell counts exceed 5, so use of a chi-square test is appropriate. $\chi^2 = 22.426$, $df = 8$, and $P = 0.004$. We have strong evidence of a relationship between race and opinion of schools.

23.45: H_0 : there is no relationship between laundry habits and preference vs. H_a : there is some relationship between laundry habits and preference. To compare people with different laundry habits, we compare the percent in each class who prefer the new product.

	Soft water, warm wash	Soft water, hot wash	Hard water, warm wash	Hard water, hot wash
Prefer new product	54.3%	51.8%	61.8%	58.3%

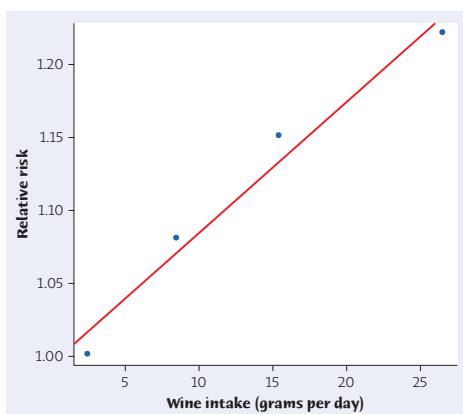
The differences are not large, but the “Hard water, warm wash” group is most likely to prefer the new detergent. With expected cell counts exceeding 5, a chi-square test is appropriate. $\chi^2 = 2.058$, $df = 3$, $P = 0.560$. The data provide no evidence to conclude that laundry habits and brand preference are related.

23.47: We compare the percentages leaning toward each party within each education group. At each education level, we compute the percentage leaning each party. For example, among bachelor's degree holders, $157/(157 + 154) = 50.5\%$ lean Democrat, while the other 49.5% lean Republican.

At every education level, people leaning Democrat outweigh people leaning Republican. The difference is greatest at the “None” level of education, then decreases until the party support is nearly equal for bachelor's holders. Among graduate degree holders, Democrats strongly outnumber Republicans.

Chapter 24 Inference for Regression

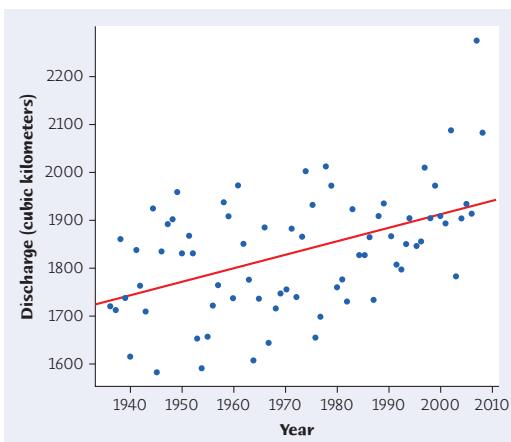
24.1: (a) A scatterplot of the data is provided.



(b) We estimate that an increase in intake of 1 gram per day increases relative risk of breast cancer by 0.0009. According to our estimate, wine intake of 0 grams per day is associated with a relative risk of breast cancer of 0.9931 (about 1). (c) $\hat{y} = 0.9931 + 0.0009x$. See table below. $s^2 = 0.00079/2 = 0.000395$. We estimate σ by $s = \sqrt{0.000395} = 0.01987$.

Residual				
x	y	\hat{y}	$y - \hat{y}$	$(y - \hat{y})^2$
2.5	1.00	1.0156	-0.0156	0.00024
8.5	1.08	1.0697	0.0103	0.00011
15.5	1.15	1.1328	0.0172	0.00030
26.5	1.22	1.2319	-0.0119	0.00014
		0		0.00079

24.3: (a) A scatterplot of discharge by year is provided. Discharge seems to be increasing over time, but there is also



a lot of variation in this trend, and our impression is easily influenced by the most recent years' data. $r^2 = 0.225$. (b) $\hat{y} = -3690.08 + 2.80x$; $s = 111$.

24.5: $H_0: \beta = 0$ vs. $H_a: \beta > 0$. $t = \frac{b}{SE_b} = \frac{2.800}{0.6168} = 4.539$. $df = 73 - 2 = 71$. Round df down to df = 60; $P < 0.0005$ (software: $P = 0.0002$). There is strong evidence of an increase in arctic discharge over time.

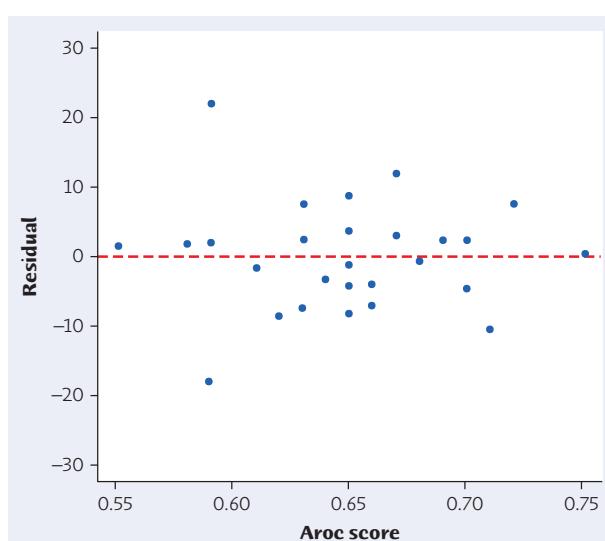
24.7: (a) $H_0: \beta = 0$ vs. $H_a: \beta > 0$. $t = 8.104$ with $df = 2$. $0.005 < P < 0.01$. This test is equivalent to testing $H_0: \text{population correlation} = 0$ vs. $H_a: \text{population correlation} > 0$. (b) $r = 0.985$. $0.005 < P < 0.01$ (using Table E with $n = 4$).

24.9: $t^* = 2.920$ ($df = 4 - 2 = 2$, with 90% confidence). $0.009012 \pm 2.920(0.001112) = 0.009012 \pm 0.003247 = 0.00577$ to 0.01226. With 90% confidence, the increase in relative risk of breast cancer associated with an increase in alcohol consumption by 1 gram per day is between 0.00577 and 0.01226.

24.11: $b = 2.800$ and $SE_b = 0.6168$. $df = 71$. With Table C, use $df = 60$ (the nearest smaller value of df in the table): $t^* = 1.671$. $2.8000 \pm 1.671(0.6168) = 2.8000 \pm 1.0307 = 1.7693$ to 3.8307 cubic kilometers per year. With 90% confidence, the yearly increase in arctic discharge is between 1.7693 and 3.8307 cubic kilometers. This confidence interval excludes "0," so there is evidence arctic discharge is increasing over time.

24.13: (a) $\hat{\mu} = 8.91 + 87.76(0.60) = 61.57$. (b) $SE_{\hat{\mu}} = 2.184$, $df = 29 - 2 = 27$, $t^* = 2.052$. $61.57 \pm 2.052(2.184) = 57.088$ to 66.052.

24.15: (a) The residual plot provided does not suggest any deviation from a straight-line relationship between volume



and Aroc score, although there are two residuals of larger magnitude present, both for Aroc scores slightly lower than 0.60.

(b) A stemplot of residuals, provided below, does not suggest that the distribution of residuals departs strongly from Normality. There are two possible outliers, which agrees with the output provided by Minitab referenced in the problem statement.

-1	8
-1	0
-0	88777
-0	4443211
0	0222333334
0	889
1	2
1	
2	2

- (c) It is reasonable to assume that observations are independent, since we have 29 different subjects, measured separately.
 (d) It may be the case that spread is larger for smaller values of Aroc, but these happen to be the two outliers. It is difficult to make a definitive argument either way.

24.17: (a) $r = +\sqrt{r^2} = +\sqrt{0.623} = 0.789$.

24.19: (a) This is a one-sided alternative, because we wonder if larger appraisal values are associated with larger selling prices.

24.21: (c)

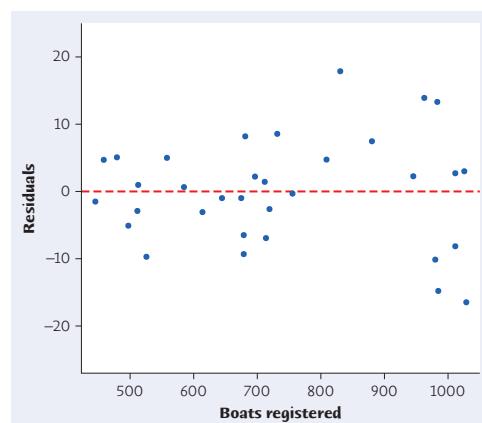
24.23: (c) $t^* = 2.056$, so the margin of error is $2.056(0.1938) = 0.3985$.

24.25: (a) Scientists estimate that each additional 1% increase in the percentage of Bt cotton plants results in an increase of 6.81 mirid bugs per 100 plants. (b) The regression model explains 90% of the variability in mirid bug density. (c) $H_0: \beta = 0$ vs. $H_a: \beta > 0$ (H_0 : population correlation = 0 vs. H_a : population correlation > 0). $P < 0.0001$; there is strong evidence of a positive linear relationship between the proportion of Bt cotton plants and the density of mirid bugs. (d) We cannot conclude a causal relationship.

24.27: $df = 10$, $t^* = 1.812$. (a) $b \pm t^*SE_b = 274.78 \pm 1.812(88.18) = 274.78 \pm 159.78 = 115.0$ to 434.6 fps/inch.
 (b) $\hat{y} \pm t^*SE_{\hat{y}} = 207.8 \pm 1.812(17.4) = 176.3$ to 239.3 fps.

24.29: (a) There is little evidence of non-Normality in the residuals, and there don't appear to be any strong outliers.
 (b) The scatterplot confirms the comments made in the text.
 (c) Presumably, close inspection of a manatee's corpse will reveal nonsubtle clues when cause of death is from collision with a boat rotor. Hence, it seems reasonable that the

-1	75
-1	00
-0	98775
-0	3332110
0	1112233
0	5555789
1	34
1	8



number of kills listed in the table are mostly not caused by pollution.

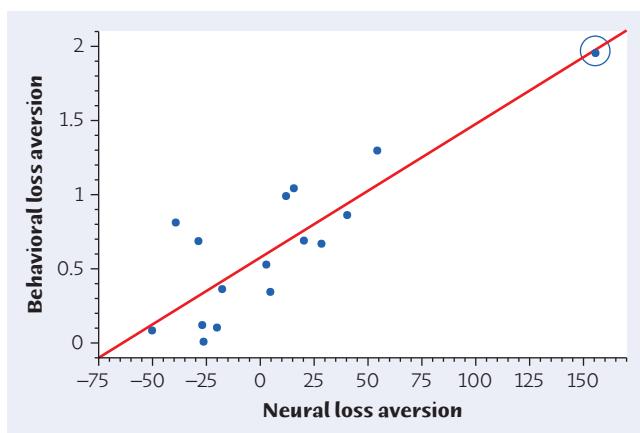
24.31: (a) This is a confidence interval for β . With $df = 31$, $t^* = 2.042$. $b \pm t^*SE_b = 0.129232 \pm 2.042(0.00752) = 0.129232 \pm 0.01536 = 0.11387$ to 0.14459 additional killed manatees per 1000 additional boats. (b) $\hat{y} = -43.17195 + 0.129232(1050) = 92.5217$ killed manatees. A 95% prediction interval for the number of killed manatees if 1,050,000 boats are registered is 75.18 to 109.87 kills.

24.33: (a) We test H_0 : population correlation = 0 against H_a : population correlation is positive. $t = 3.88$; $df = 27 - 2 = 25$; $P < 0.0005$. There is very strong evidence of a positive correlation between Gray's forecasted number of storms and the number of storms that actually occur. (b) $\hat{\mu} = 1.022 + 0.9696(16) = 16.535$, and $SE_{\hat{\mu}} = 1.3070$. $df = 25$, $t^* = 2.060$. The 95% confidence interval is given by $16.535 \pm 2.060(1.3070) = 13.843$ to 19.227 storms.

24.35: The stemplot is provided, where residuals are rounded to the nearest tenth. The plot suggests that the residuals do not follow a Normal distribution. Specifically, there are a number of rather extreme outliers. This makes regression inference and interval procedures unreliable.

-7	5
-6	
-5	
-4	7
-3	75
-2	88776
-1	8
-0	876
0	3345
1	334
2	333
3	23
4	
5	
6	3
7	
8	
9	
10	
11	4

24.37: (a) Shown is the scatterplot with two (nearly identical) regression lines, one using all points and one with the outlier omitted. (b) The correlation for all points is $r = 0.8486$. For testing the slope, $t = 6.00$, for which $P < 0.001$. (c) Without the outlier, $r = 0.7014$, the test statistic for the slope is $t = 3.55$, and $P = 0.004$. In both cases there is strong evidence of a linear relationship between neural loss aversion and behavioral loss aversion. However, omitting the outlier weakens this evidence somewhat.



24.39: The distribution is skewed right but the sample is large, so t procedures should be safe. $\bar{x} = 0.2781 \text{ g/m}^2$, $s = 0.1803 \text{ g/m}^2$. $t^* = 1.984$ for $df = 100$ (rounded down from 115). The 95% confidence interval for μ is 0.2449 to 0.3113 g/m^2 .

24.41: $\hat{y} = 1.4146 + 0.4399x$. The slope is significantly different from zero ($t = 4.33$, $P = 0.001$). To assess the evidence that more cones leads to more offspring, we should use the one-sided alternative, $H_a: \beta > 0$, for which P is half as large (so $P < 0.001$). The conditions for inference seem to be satisfied. One might also choose to find a confidence interval for β : $df = 14$, $t^* = 2.145$. A 95% confidence interval for β is $0.4399 \pm 2.145(0.1016) = 0.2220$ to 0.6578 offspring per cone.

24.43: $\hat{y} = -1.286 + 11.894x$. An examination of the residuals does not suggest any severe violations of the conditions for regression inference. To test $H_0: \beta = 0$ vs. $H_a: \beta > 0$, $t = 10.47$ ($df = 21$), $P < 0.0005$. For $df = 21$, $t^* = 2.080$ for 95% confidence, so with b and $SE_b = 1.136$. We are 95% confident that β is between 9.531 and 14.257.

24.45: $\hat{y} = 0.1523 + 8.1676x$. A stemplot of the residuals looks reasonably Normal, but the scatterplot suggests that the spread about the line is greater when phytopigment concentration is greater. This may make regression inference unreliable, but we will proceed. The slope is significantly different from 0 ($t = 13.25$, $df = 114$, $P < 0.001$). A 95% confidence interval for β is $8.1676 \pm 1.984(0.6163) = 6.95$ to 9.39 .

24.47: (a) $\bar{x} = -0.00333$, $s = 1.0233$. For a standardized set of values, we expect the mean and standard deviation to be (up to rounding error) 0 and 1, respectively. (b) The stemplot does not look particularly symmetric, but it is not strikingly non-Normal for such a small sample. (c) The probability that a standard Normal variable is as extreme as this is about 0.0272.

-2	2
-1	4
-0	
-0	32
0	01122
0	7
1	0
1	5

24.49: $df = 14$, $t^* = 2.145$. $-0.01270 \pm 2.145(0.01264) = -0.0398$ to 0.0144 . This interval does contain 0.

Chapter 25 One-Way Analysis of Variance: Comparing Several Means

25.1: (a) $H_0: \mu_A = \mu_B = \mu_C = \mu_D$ vs. H_a : not all means agree. (b) Referring to Figure 25.2, comparing Groups A and C, we see that the mean status for men expressing anger is about 6.3, while the mean status for women expressing anger is about 4.

With both groups expressing anger, men receive higher mean status scores than women, and the mean difference is about 2.3. Notice that comparing Groups B and D, we see that women expressing sadness receive higher status scores than men expressing sadness, but the difference is relatively small.

25.3: (a) The stemplots appear to suggest that logging reduces the number of trees per plot and that recovery is slow (the 1-year-after and 8-years-after stemplots are similar). (b) The means lead one to the same conclusion as in (a): the first mean is much larger than the other two. (c) $H_0: \mu_1 = \mu_2 = \mu_3$ vs. H_a : not all means are the same, $F = 11.43$, $df = 2$ and 30, $P = 0.000205$, so we conclude that these differences are significant: the number of trees per plot really is lower in logged areas.

Never logged	1 year ago	8 years ago
0	0 2	0 4
0	0 9	0
1	1 2244	1 22
1 699	1 57789	1 5889
2 0124	2 0	2 22
2 7789	2	2
3 3	3	3

25.5: (a) By moving the middle mean to the same level as the other two, it is possible to reduce F to about 0.02, which has a P -value very close to the left end of the scale (near 1). (b) By moving any mean about 1 centimeter up or down (or any two means about 0.5 cm in opposite directions), the value of F increases (and P decreases) until it moves to the right end of the scale.

25.7: (a) $s_1^2 = 25.6591$, $s_2^2 = 24.8106$, and $s_3^2 = 33.1944$. $s_1 = 5.065$, $s_2 = 4.981$, and $s_3 = 5.761$. The largest standard deviation (5.761) is not more than twice the size of the smallest standard deviation (4.981). Conditions are satisfied. (b) The three standard deviations are $s_L = 16.61$, $s_M = 17.42$ and $s_C = 17.13$. The ratio of largest to smallest standard deviation is $17.42/16.61 = 1.05$, which is less than 2. Conditions are satisfied.

25.9: Side-by-side stemplots show some irregularity but no outliers or strong skewness. ANOVA output shows that the group standard deviations easily satisfy our rule of thumb ($2.059/1.302 = 1.58 < 2$). The differences among the groups were significant at $\alpha = 0.05$: $F = 3.44$, $df = 3$ and 27, $P = 0.031$. Nitrogen had a positive effect, the phosphorus and control groups were similar, and the plants that got both nutrients fell between the others.

25.11: (a) $I = 3$ and $N = 96$, so $df = 2$ and 93. (b) $I = 3$ and $N = 90$, so $df = 2$ and 87.

25.13: (a) No sample standard deviation is larger than twice any other. Specifically, the ratio of largest to smallest standard deviation is $2.25/1.61 = 1.40$, which is less than 2. Conditions are safe for use of ANOVA. (b)

$$\bar{x} = 4.8225$$

$$MSG = 25.502$$

$$MSE = 3.507$$

$$F = \frac{MSG}{MSE} = 7.272$$

(c) We have $df = 4 - 1 = 3$ and $68 - 4 = 64$, so we refer to the F distribution with 3 and 64 degrees of freedom. The P -value is 0.000 rounded to three decimal places. In fact, $P = 0.0003$ (obtained using software). There is strong evidence that the mean status scores between the four groups studied are not equal—a conclusion consistent with the solution to Exercise 25.1.

25.15: (c)

25.17: (b) $I - 1 = 3 - 1 = 2$, and $N - I = 9 - 3 = 6$.

25.19: (c) Since $MSG = 22,598/(3 - 1) = 22,598/2 = 11,299$, $F = MSG/MSE = 11,299/1600 = 7.06$.

25.21: (c) The largest standard deviation is 62.02 and the smallest is 20.07. Hence, the largest standard deviation is more than twice the smallest.

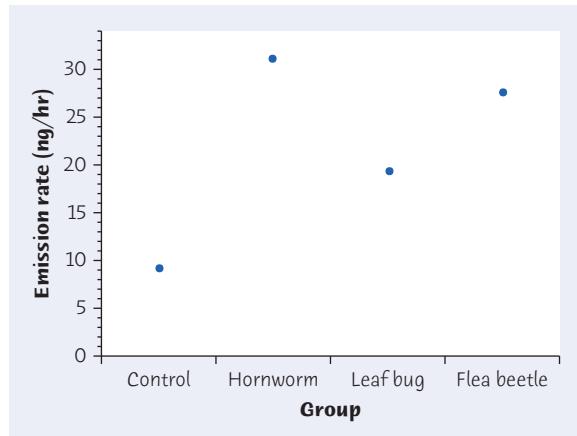
25.23: (c) We do not have three independent samples from three populations.

25.25: The populations are college students that might view the advertisement with art image, college students that might view the advertisement with a nonart image, and college students that might view the advertisement with no image. The response variable is student evaluation of the advertisement on the 1–7 scale. $H_0: \mu_1 = \mu_2 = \mu_3$ (all three groups have equal mean advertisement evaluation) vs. H_a : not all means are equal. There are $I = 3$ populations; the sample sizes are $n_1 = n_2 = n_3 = 39$, so there are $N = 39 + 39 + 39 = 117$ individuals in the total sample. There are then $I - 1 = 3 - 1 = 2$ and $N - I = 117 - 3 = 114$ df.

25.27: The response variable is hemoglobin A1c level. We have $I = 4$ populations: a control (sedentary) population, an aerobic exercise population, a resistance training population, and a combined aerobic and resistance training population. $H_0: \mu_1 = \mu_2 = \mu_3 = \mu_4$ (all four groups have equal mean hemoglobin A1c levels) vs. H_a : not all means are equal. Sample sizes are $n_1 = 41$, $n_2 = 73$, $n_3 = 72$, and $n_4 = 76$. The total sample size is $N = 41 + 73 + 72 + 76 = 262$. We have $I - 1 = 4 - 1 = 3$ and $N - I = 262 - 4 = 258$ df.

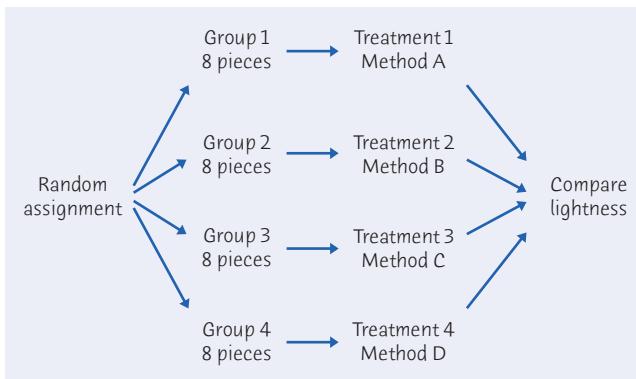
25.29: (a) The graph suggests that emissions rise when a plant is attacked because the mean control emission rate is half the smallest of the other rates. (b) The null hypothesis

is “all groups have the same mean emission rate.” The alternative is “at least one group has a different mean emission rate.” (c) The most important piece of additional information would be whether the data are sufficiently close to Normally distributed. (From the description, it seems reasonably safe to assume that these are more or less random samples.) (d) The SEM equals $s/\sqrt{8}$, so we can find the standard deviations by multiplying by $\sqrt{8}$; they are 16.77, 24.75, 18.78, and 24.38. However, this factor of $\sqrt{8}$ would cancel out in the process of finding the ratio of the largest and smallest standard deviations, so we can simply find this ratio directly from the SEMs: $\frac{8.75}{5.93} = \frac{24.75}{16.77} = 1.48$, which satisfies our rule of thumb.



25.31: (a) The means suggest that extra water in the spring has the greatest effect on biomass, with a lesser effect from added water in the winter. ANOVA is risky with these data; the standard deviation ratio is nearly 3, and the winter and spring distributions may have skewness or outliers (although it is difficult to judge with such small samples). (b) $H_0: \mu_w = \mu_s = \mu_c$ vs. H_a : at least one mean is different. (c) ANOVA gives a statistically significant result ($F = 27.52$, df 2 and 15, $P < 0.0005$), but as noted in (a), the conditions for ANOVA are not satisfied. Based on the stemplots and the means, however, we should still be safe in concluding that added water increases biomass.

25.33: (a) The design, with four treatments, is shown. (b) ANOVA should be safe: it is reasonable to view the samples as SRSs from the four populations, the distributions do not



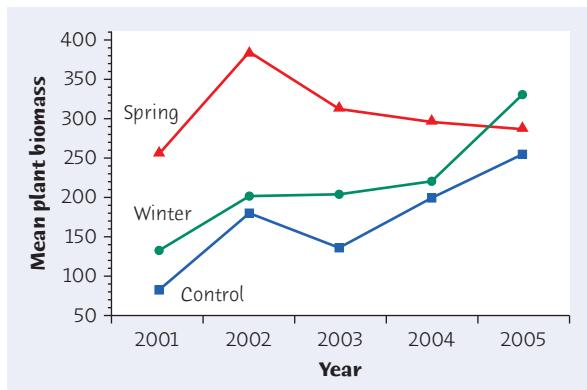
show drastic deviations from Normality, and the standard deviations satisfy our rule of thumb ($0.392/0.250 = 1.568$). The means have rather small differences in lightness score; Method C is lightest and Method B is darkest. The differences in mean lightness are nonetheless highly significant ($F = 22.77$, $P < 0.001$). The manufacturer will prefer Method B. Whether these differences are large enough to be important in practice requires more information about the scale of lightness scores.

25.35: First, we see that the ratio of largest standard deviation to smallest standard deviation is $2.388/1.959 = 1.219$, which is less than 2. There is some evidence of non-Normality, and perhaps one outlier in the “No Weather Report” group. We proceed, as the samples are reasonably large. $F = 20.679$, $df = 3 - 1 = 2$ and $60 - 3 = 57$. $P = 0.000$, to three decimal places. There is overwhelming evidence that the mean tip percentages are not the same for all three groups.

25.37: First, we note that the mean angle for untreated fabric is 79 degrees, showing much less wrinkle resistance than any of the treated fabrics. ANOVA on four groups gives $F = 153.76$ and $P < 0.001$. A comparison of wrinkle recovery angle for the three durable press treatments is more interesting.

The ANOVA F -test cannot be trusted because the standard deviations violate our rule of thumb: $10.16/1.92 = 5.29$. This is much larger than 2. In particular, Permafresh 48 shows much more variability from piece to piece than either of the other treatments. Large variability in performance is a serious defect in a commercial product, so it appears that Permafresh 48 is unsuited for use on these grounds. The data are very helpful to a maker of durable press fabrics despite the fact that the formal test is not valid.

25.39: (a) There is a slight increase in growth when water is added in the wet season and a much greater increase when it



is added during the dry season. (b) The means differ significantly during the first three years. (c) The year 2005 is the only one for which the winter biomass was higher than the spring biomass.

25.41: In addition to a high standard deviation ratio ($117.18/35.57 = 3.29$), the spring biomass distribution has a high outlier.

this page left intentionally blank



Index

- Absolute values, 26–20
Addition rule for disjoint events, 267, 278
 defined, 308
 Venn diagram, 308
Alternative hypothesis. *See also* Hypotheses
 defined, 372, 384
 one-sided, 372, 373, 384
 relationships, 556
 two-sided, 372, 375, 384
American Community Survey, 213
Analysis of variance (ANOVA), 623–654,
 29–3–29–45
 boxplots, 631–632
 calculation examples, 642–643
 conditions for, 633–637, 645, 29–3
 defined, 632
 details of, 631, 640–644
 F distributions, 637–640, 645
 F statistic, 633, 637, 645, 28–18, 28–61, 29–4
 F statistic formula, 637, 640
 F test, 625–627, 641, 645, 28–54
 idea of, 631–633
 independent SRSs, 634, 645
 interpretation skills, 645
 Normal distribution, 634, 645
 one-way, 632, 645, 29–3
 robustness of, 634, 645
 squared multiple correlation coefficient, 28–15
 standard deviations in, 634, 645
 technology use in, 628–631
 two-way, 29–6, 29–16–29–37
Analysis of variance (ANOVA) tables, 642, 645
 general form, 28–53
 one-way ANOVA, 29–33
 software provision, 28–53
 sum of squares row, 28–61
Anonymity, 250
Applets
 Central Limit Theorem, 298
 Confidence Interval, 356
 Correlation and Regression, 141
 Law of Large Numbers, 289, 305
 Normal Approximation to Binomial, 343
 Normal Curve, 83
 One-Variable Statistical Calculator, 22
 One-Way ANOVA, 632, 633
 Power of a Test, 403–404
 Probability, 262
 P-Value of a Test of Significance, 375
 Reasoning of a Statistical Test, 369
 Simple Random Sample, 204, 205–206, 231
 Two-Variable Statistical Calculator, 101, 112
Association
 causation and, 144–147
 consistent, 146
 examples, 144, 145, 146
 strong, 144, 146
Automated algorithms, 28–46
Average
 correlation based on, 142
 as less variable than individual observations, 293
 ranks, 26–14
Balanced designs, 29–17
Bar graphs. *See also* Graphs
 in categorical data representation, 169
 defined, 8, 10, 26
 distribution display, 8–10
 flexibility, 9
 histograms versus, 13–14
 marginal distributions, 161
 Pareto charts, 27–7
 percents, 165
 in quantity comparison, 10
Behavior chaotic, 262
Behavioral experiments, 253–254
Bell curve, 88
Benford's law example, 269–270
Bias
 defined, 202, 215
 eliminating with random sampling, 208
 random sample elimination of, 260
RDD, 214
 response, 212, 216
 sample choice, 210
 wording of questions, 212–213
Biased statistical studies, 202
Binomial coefficients
 defined, 335, 344
 formula, 335
Binomial distributions, 331–367
 caution, 338
 defined, 332, 343
 examples, 332
 in finite probability models, 332
 mean, 338–340, 344
 Normal approximation to, 340–343
 probability histogram, 339
 sampling distribution approximation, 344
 skills, 422
 standard deviation, 338–340, 344
 in statistical sampling, 333–334
Binomial probabilities
 defined, 335, 344
example, 336
Normal calculation, 340
with technology, 336–337
Binomial setting, 331, 343
Block designs
 benefits of, 238
 defined, 237
 illustrated, 237
 matched pairs, 240
 matched pairs design, 236–237
Blocks
 conclusions about, 238
 defined, 237, 240
Body mass index (BMI), 352–353
Bootstrap methods, 26–3–26–4
Boxplots
 analysis of variance (ANOVA), 631–632
 defined, 46, 58
 illustrated, 46
 for side-by-side comparison, 46
Call screening, 213
Capability
 control versus, 27–34–27–36
 defined, 27–35, 27–42
 example, 27–35
 skills, 27–43
Categorical variables, 6–11. *See also* Variables
 adding to scatterplots, 104–106
 defined, 4, 25
 distribution, 7
 lurking variables and, 166
 marginal distributions for, 161
 possible outcomes, 574
 relationships, 553
 skills, 179–180
 values of, 7
Causation
 association and, 144–147
 basic meaning of, 144
 correlation and, 148
 criteria for establishing, 146
 experiments and, 145
Cause-and-effect diagrams
 defined, 27–4, 27–42
 illustrated, 27–7
 outline, 27–5
Cell counts, for chi-square test, 561
Cell-phone-only households, 213
Center, 18, 26, 58
 defined, 15
 example, 16

Center (*continued*)
 measures, choosing, 51–53
 measuring, 41–42
 numerical measures of, 52
 resistant measure of, 40
 sampling distribution, 292
 Central limit theorem, 295–301
 defined, 296, 301
 exact distribution and normal approximation from, 300
 examples, 296–299
 general versions, 296
Central Limit Theorem applet, 298
 Chance behavior, in short and long term, 260
 Chance outcomes, 76
 Chi-square distributions
 approximation, 576
 defined, 570, 576
 degrees of freedom, 570, 576
 density curve, 570–571
 example, 571
 Kruskal-Wallis statistic, 26–35
 mean, 571
 Chi-square statistic
 defined, 560, 576
 example, 560–561
 in goodness of fit, 574
 terms of, 564, 576
 Chi-square test
 in casinos, 574
 cell counts requirement, 561
 defined, 576
 example, 564–565
 for goodness of fit, 572–576
 “many-sided” hypothesis, 573
 null hypothesis, 569
 as one-sided, 561
 relationships detection with, 564
 skills, 577
 technology use in, 562–567
 uses of, 567–570
 Clinical trials, 252–253
 Clusters, 102
 Coefficients
 binomial distributions, 335, 344
 regression, 28–39–28–41, 28–55
t tests for, 28–19
 Column totals, 160, 168
 Column variables, 5, 159, 168
 Common cause variation, 27–9
 Conditional distributions
 comparing, 168
 defined, 163, 168
 example, 163
 explanatory-response relationship, 165
 finding, 169
 sets of, 165
 software output, 163–164

Conditional probability
 caution, 315
 defined, 314, 322
 examples, 314, 315
 in finding sources, 321
Confidence Interval applet, 356
 Confidence intervals, 351–367
 behavior, 361–363
 cautions about, 395–397
 for contrasts, 29–15
 critical values, 357
 defined, 351, 363, 418
 example, 359
 form, 363
 four-step process, 358
 illustrated, 356
 large-scale, for proportions, 496–499, 507
 lengths of, 362
 level, 354–357
 margin of error of, 364, 395, 408
 for mean response, 605, 613, 28–61
 more accurate, obtaining, 527
 multiple analyses applied to, 399–400
 one-sample *t*, 440–443
 planning, 401
 plus four, 499–501, 507, 527
 pooled *t* for, 481
 for population mean, 357–361, 364, 437, 455
 for proportion comparison, 520–522
 for proportions, 499–502, 507
 regression coefficients, 28–55
 regression models, 28–55
 for regression response, 605
 for regression slope, 600–602
 robust, 452
 sample proportion, 499
 sample size for, 401–402
 skills, 422
 for slope, 613
 standard error, 496, 517
 trust of, 392
 Tukey simultaneous, 29–9
 two-sample *t* procedures, 470
 unknown population mean, 355
 z , 409
 Confidence levels
 defined, 354, 363
 example, 355–356
 large-sample interval, 527
 margin of error and, 354–357
 margin of error trade-off, 361
 overall, 29–9
 Confidentiality. *See also* Data ethics
 anonymity versus, 250
 breach of, 250
 defined, 248
 example, 251
 Confounding
 defined, 225, 239
 example, 224
 prevention, 225
 Continuity correction, 26–9
 Continuous probability models
 defined, 271, 278
 example, 271–272
 Normal distributions as, 273
 probability assignment, 273
 Continuous random variables, 276, 279
 Contrasts
 confidence intervals for, 29–15
 defined, 29–13
 estimating, 29–13
 examples, 29–13, 29–14–29–15
 follow-up analysis, 29–12–29–16
 inference about, 29–14
 sample, 29–13–29–14
 Control
 block designs, 237
 comparison, 239
 experiment design, 233, 239
 groups, 229
 placebo, 235
 Control charts. *See also* Statistical process control
 center line, 27–11
 chart setup, 27–10, 27–26, 27–42
 constants, 27–18
 control limits, 27–11, 27–42
 defined, 27–9, 27–42
 functioning of, 27–10
 out-of-control signals, 27–24–27–25, 27–42
 p charts, 27–36–27–41, 27–42
 with past data, 27–28, 27–42
 process mean estimation, 27–26–27–27
 for process monitoring, 27–10–27–23, 27–25–27–26
 R charts, 27–23–27–24, 27–42
 s charts, 27–16–27–23, 27–42
 for sample proportions, 27–36–27–37
 sample range, 27–23
 setting up, 27–25–27–32
 skills, 27–43
 standard deviation estimation, 27–27
 three-sigma, 27–17, 27–42
 using, 27–23–27–25
 \bar{x} charts, 27–10–27–16, 27–42
 Control limits, 27–11, 27–37–27–41, 27–42
 Convenience samples. *See also* Samples
 defined, 202
 inference from, 208
 Correlation
 based on averages, 142
 causation and, 148
 cautions, 142–144, 148, 179
 defined, 106, 113, 589
 ecological, 142, 148
 example, 109–111
 explanatory-response distinction and, 599
 explanatory/response variables and, 108

- facts about, 108–112
 formula, 107, 108, 111
 high, 148
 lack of, testing, 598–600
 lurking variables and, 143
 in measuring linear association, 106–108
 negative, 108
 negative association, 108
 nonzero, evidence of, 600
 outlying observations and, 111
 population, hypothesis of, 613
 positive, 108
 positive association, 108
 quantitative variable requirement, 111
 regression and, 147
 scatterplots for inference, 599
 skills, 178–179
 squared, 133–134, 147
 strength measurement, 111
 units of measure and, 108
 values, 108–109
- Correlation and Regression* applet, 141
- Critical values
 defined, 357, 364
 t curve, 455
 table, significance from, 382–384
- Crossed designs, 29–17
- Cross-sectional data, 25
- CrunchIt!
 analysis of variance, 629
 binomial probabilities, 337
 chi-square test, 563
 descriptive measures, 53
 Kruskal-Wallis test, 26–32
 least-squares regression, 131
 multiple regression model, 28–45
 output, 54
 parallel regression lines, 28–14
 parameter estimates and t statistics, 28–47
 proportion comparison, 519
 rank tests, 26–10–26–12
 regression inference, 594
 signed rank test ties, 26–26
 simple linear regression model, 28–46
 t confidence interval, 447
 t test, 448
 two-sample problems, 475, 476
 two-way tables, 164
 Wilcoxon signed rank test, 26–23
- Cumulative proportions
 defined, 82, 89
 for Normal curves, 83
- Cycles, 24, 26
- Data
 categorical, 169, 179–180
 consistency, checking, 7
 cross-sectional, 25
 designs for producing, 417
- deviations, 15
 overall pattern, 15
 past, control charts using, 27–28, 27–42
 past, p chart using, 27–38
 plotting, 26, 52, 58, 148
 representing with bar graphs/pie charts, 9
 rounding, 21
 skills, 177
 source, 392
 time series, 25
 two-sample problems, 482
 unrepresentative, 202
- Data analysis
 defined, 1, 175
 of experiments, 235
 exploratory, 6, 25, 175
 importance of, 1
 before inference, 408
 purpose of, 175
 from sampling design, 210
 several-variable data, 98
 skills, 534
 symmetric population distributions, 466
- Data ethics, 247–257
 basics, 248
 behavioral and social science experiments, 253–254
 clinical trials, 252–253
 complex issues of, 248
 confidentiality, 248, 250–251
 institutional review boards, 248
 questions, 247
- Data production
 experiments, 223–245
 sampling, 199–221
 skills, 533
- Degrees of freedom
 chi-square distributions, 570, 576
 defined, 50, 439, 455
 F test, 639
 multiple regression, 28–8
 regression standard error, 591, 613
 t distribution, 439, 455
 two-sample t statistic, 483
 two-way ANOVA, 29–35
- Density curves
 area under, 278–279
 areas under, 271
 in assigning probability, 273
 chi-square distribution, 570–571
 defined, 71, 89
 describing, 73–75
 equal-areas point, 73
 F distributions, 638
 histogram example, 70–71
 illustrated, 70, 71, 73, 75
 mean, 74, 89
 median, 73, 74, 89
 outliers and, 72
 as overall pattern description, 72
- right-skewed distribution, 72
 skills, 178
 standard deviation, 74, 89
 sum of random numbers, 275
 symmetric, 73, 89
 t distributions, 439
 total area, 89
 uniform, 272
- Dependent variables. *See* Response variables
- Deviations
 defined, 15, 26
 outliers, 15, 16, 26
 in scatterplot interpretation, 101
- squared, 51
 standard, 49–51
- Direction, 101, 113
- Discrete probability models. *See* Finite probability models
- Discrete random variables, 276, 279
- Disjoint events, 267, 278, 322
- Distributions
 bar graph display, 8–10, 26
 binomial, 331–367
 boxplots, 46–47
 categorical variables, 7
 center, 18
 chi-square, 570–572
 conditional, 162–166, 168
 continuous probability models, 26–4
 defined, 7, 26
 describing with numbers, 39
 description skills, 177
 display skills, 177
 displaying, 7–10
 F , 637–640
 five-number summary, 45–46
 histogram display, 11–15, 26
 irregular shapes, 19
 marginal, 160–162, 168
 midpoint, 15
 Normal, 58, 75–81
 numerical summaries of, 177–178
 outliers, 15, 16
 pie chart display, 8, 26
 population, 301
 of real data, 76
 as resistant, 45
 sampling, 285–305
 shape, 15, 16, 18
 single-peaked, 17
 skewed, 20, 26
 skewed to left, 16
 skewed to right, 16
 spread, 15, 16, 18
 stemplot display, 20–23, 26
 symmetric, 16, 17, 26
 t , 438–440
 two-peaked, 18
- Double-blind experiments, 234, 240

- Ecological correlation
caution, 142
defined, 142, 148
- Economic variables, 43
- Errors. *See also* Mean square error
regression standard, 591, 592, 613, 28-8,
28-61
- roundoff, 7, 161
- standard, 438, 455, 496, 517
- testing, 406-407
- Type I, 406-407
- Type II, 406-407
- Estimated regression model, 28-61
- Events
defined, 278
- disjoint, 267, 278, 322
- independent, 318, 322
- Examining, 101
- Excel
analysis of variance, 629
- binomial probabilities, 337
- least-squares regression, 131
- output, 54
- quartile function, 53
- regression inference, 594
- t* confidence interval, 447
- t* test, 448
- two-sample problems, 475, 476
- Expected counts
defined, 558, 576
- examples, 558
- formula, 558-559
- in two-way tables, 558-560
- Experimental design
control, 233, 239
- defined, 239
- enough subjects, 233, 239
- principles of, 233
- randomization, 233, 239
- Experiments, 223-245
advantages of, 227
- behavioral and social science,
253-254
- causation and, 145
- cautions about, 234-236
- clinical trials, 252-253
- comparative, 229-232
- defined, 148, 223, 239
- domestic violence, 254
- double-blind, 234, 240
- factors, 225-227
- good, requirements of, 240
- lack of realism, 235, 240
- lurking variables and, 233
- observational studies versus, 223-225
- personal space, 253-254
- placebos, 234, 240
- randomized comparative, 229-232
- skills, 420
- statistical analysis of, 235
- statistical design in, 228
- statistical significance, 233
- subjects, 225-227
- treatments, 225-227
- two-factor, 226, 227
- uncontrolled, 228
- vocabulary, 225-227
- Explanatory variables. *See also* Variables
in conditional distributions, 165
- correlation and, 108
- defined, 97, 113
- examples, 98
- identifying, 98
- names for, 98
- plotting on scatterplot, 113
- plotting residuals against, 28-58
- plotting response variables against, 28-58
- response variable relationship, 125
- Exploratory data analysis
defined, 6, 26
- principles, 6
- purpose of, 175
- statistical inference versus, 197
- Extrapolation
caution, 143
- defined, 143, 148
- F distributions, 29-4
ANOVA, 637-640
- defined, 638, 645
- degrees of freedom, 638, 639
- density curves, 638
- example, 638
- in *F* test, 641
- F statistic
analysis of variance, 28-18
- ANOVA, 28-61, 29-4
- calculating, 641
- defined, 633, 645
- elements of, 640
- formula, 637, 640
- mean squares, 640
- null hypothesis, 28-19
- numerator and denominator formulas,
28-54
- one-way ANOVA, 29-33
- F tests
analysis of variance, 625-627, 641, 645
- degrees of freedom, 639
- example, 28-20-28-23
- follow-up analysis, 29-6, 29-8-29-16
- one-way ANOVA, 29-33
- parallel lines, 28-18-28-19
- two-way ANOVA, 29-35-29-36, 29-37
- Factorial notation, 335
- Factors. *See also* Experiments
defined, 226, 239
- interaction of, 227
- Finite probability models
assigning probabilities in, 269
- binomial distributions in, 332
- defined, 269, 278
- example, 269-270
- using, 268-271
- First quartile
defined, 44, 58
- finding, 88
- Five-number summary
boxplots, 46-47
- defined, 45
- illustrated, 45
- for skewed distributions, 52
- Flowcharts
defined, 27-4, 27-42
- illustrated, 27-6
- Follow-up analysis
contrasts, 29-12-29-16
- defined, 29-6, 29-37
- skills, 29-38
- Tukey pairwise multiple comparisons, 29-8-29-12
- Form, 101, 113
- Four-step process
analysis of variance *F* test, 626
- ANOVA conditions, 635-636
- chi-square test, 567-568
- conditional distributions, 163
- conditions for inference, 608-609
- confidence intervals, 358
- confidence intervals for comparing proportions,
521
- F* test, 28-20-28-23
- inference about prediction, 602-604
- inference for two-way ANOVA, 29-24-29-31
- Kruskal-Wallis test, 26-28
- one-way ANOVA, 29-4-29-6
- parallel regression lines, 28-4-28-5
- rank test ties, 26-14-26-16
- regression inference, 587-589
- regression lines, 125-127
- scatterplot interpretation, 101-104
- scatterplots, 99-100
- signed rank test ties, 26-25
- significance tests for comparing proportions,
522-523, 524-525
- statistical problem organization, 56
- statistical problems, 55
- tests of significance, 379
- General addition rule. *See also* Probability rules
caution, 311
- defined, 312
- examples, 312-313
- Venn diagram, 312
- General multiplication rule. *See also* Probability
rules
defined, 316
- examples, 316-317

- General Social Survey, 213
 Goodness of fit, chi-square test, 572–576
 Graphing calculators
 analysis of variance, 628
 binomial probabilities, 336
 chi-square test, 562
 in graph creation, 53
 least-squares regression, 131
 output, 54
 proportion comparison, 519
 regression inference, 593
 t confidence interval, 447
 t test, 448
 two-sample problems, 475
 Graphs
 bar, 8–10, 26
 boxplots, 46–47
 deviations, 15
 distribution picture, 52
 histogram, 11–19, 26
 interpreting, 17
 overall pattern, 15
 pie chart, 8, 9, 26
 stemplot, 26
 time plot, 23–25
 Groups
 comparison, 29–8
 control, 229
 mean squares for, 641
 rational subgroup, 27–32–27–33, 27–42
 sum of squares for, 26–29
 treatment, 239
 Hennekens, Dr. Charles, 253–254
 Histograms, 11–15. *See also* Quantitative variables
 bar area, 14
 bar graphs versus, 13–14
 binomial distributions, 339
 center, 15, 16
 class choices, 14, 17
 creation example, 12–13
 cross-sectional data, 25
 defined, 11, 26
 drawing, 13
 illustrated, 13, 16, 17, 18, 19
 interpreting, 15–19
 outliers, 15, 16
 overall pattern, 15
 of percents, 18
 right-skewed distribution, 72
 shape, 15, 16
 skewness, 17
 spread, 15, 16
 stemplots versus, 20
 Hypotheses
 alternative, 372, 373, 384
 before data, 373
 example, 373
 Kruskal-Wallis test, 26–29
 null, 372, 384
 population reference, 372
 simultaneous tests of, 29–9
 slope, 613
 stating, 372–374
 testing errors, 406–407
 Wilcoxon test, 26–13–26–14
 Hypothesis testing
 chi-square test, 569, 573
 lack of correlation, 598–600
 multiple comparisons, 556–558
 no linear relationship, 597–598
 null hypothesis, 379, 504
 population correlation between x and y , 613
 for population mean, 378–382
 proportions, 504–505
 P -values in, 374–378
 two proportion comparison, 523
 Type I error, 406–407
 Type II error, 406–407
 Wilcoxon rank sum test, 26–13–26–14
 z statistic, 379, 382, 504
 Income distributions, as right-skewed, 76
 Independent events, 318
 conditional probability and, 318
 defined, 309, 318, 322
 multiplication rule for, 308–311
 Independent observations, 331
 Independent variables. *See also* Explanatory variables
 Indicator variables, 28–6, 28–61
 Individuals
 in data table format, 5
 defined, 3, 25
 in experiments, 226
 how many question, 4
 as rows, 5
 types of, 25
 Inference
 ANOVA, conditions for, 634
 based on Normal distributions, 76
 basic conditions for, 626
 conclusions, 418
 conditions, checking, 359–360
 conditions about a mean, 437–438
 conditions for comparing means, 467
 conditions for practice, 392–395
 from convenience samples, 208
 data analysis before, 408
 data source and, 392
 defined, 1, 208, 351
 exploratory data analysis versus, 197
 methods of, 534
 multiple regression, 28–16–28–26
 one mean, 537
 population mean, 437–463
 population proportion, 493–513, 537–538
 populations, 208–209
 in practice, 391–415
 prediction, 602–604
 probability ideas, 287
 reasoning of, 285, 351, 534
 reasons for studying, 352
 regression, 587–621
 regression parameters, 28–53–28–58
 relationships, 549–550
 standard deviations and, 482
 two-way ANOVA, 29–23–29–32
 use of, 408
 variables, 533–550
 Influential observations, 139–142
 defined, 139, 148
 example, 139–140
 outliers and, 140
 Informed consent. *See also* Data ethics
 defined, 248
 difficulties of, 249
 example, 249–250
 subjects, 248–250
 Institutional review boards, 248
 Interactions
 defined, 28–26, 29–20
 examples, 28–27–28–30
 main effects and, 29–19–29–20
 multiple regression, 28–26–28–32
 two regression lines, 28–27
 Intercept
 defined, 127, 147
 least-squares regression line, 130
 regression inference, 589
 regression line, 128, 147
 Interquartile range (IQR)
 defined, 48
 in rule for outliers, 48–49
 Kruskal-Wallis statistic
 chi-square distribution, 26–35
 defined, 26–30, 26–35
 distribution of, 26–30
 example, 26–30–26–31
 Kruskal-Wallis test. *See also* Nonparametric tests
 conditions for, 26–29
 defined, 26–35
 example, 26–28
 hypotheses for, 26–29
 idea of, 26–29
 null hypothesis, 26–35
 as rank test, 26–29
 sample comparison with, 26–27–26–34
 Labels
 randomized comparative experiments, 230
 SRS, 205, 206, 215
 Lack of realism, 235, 240
 Large-scale confidence intervals. *See also*
 Confidence intervals
 confidence level, 527
 defined, 497, 507

Large-scale confidence intervals (*continued*)

- example, 497–498
- margin of error, 502
- for proportion, 496–499
- for proportion comparison, 517–518, 526

Law of large numbers

- defined, 287, 288, 301
- example, 288–289
- statistical estimation and, 287–290
- uses, 289

Law of Large Numbers applet, 289, 305

Least-squares regression

- defined, 613
- example, 132–133
- explanatory and response variables distinction, 132
- facts about, 132–135
- output, 131
- square of the correlation, 133–134
- as statistical problem, 130

Least-squares regression line, 128–135. *See also*

- Regression lines
- defined, 129, 147
- equation, 130, 589
- illustrated, 129, 136
- intercept, 130
- outliers, 140
- popularity, 129
- slope, 130
- slope and correlation, 133
- standard deviation, 133

Least-squares residuals, 137

Least-squares slope, 598

Leaves, stemplot, 20

Linear association

- measuring, 106–108
- strong, 106

Linear relationships, 106–112

- correlation, 106–112
- direction, 101, 113
- displaying with scatterplots, 99
- form, 101, 113
- measuring, 106–108
- no hypothesis, testing, 597–598
- as regression inference condition, 609
- strength, 101, 109–110, 111, 113

Lurking variables. *See also* Variables

- categorical variables and, 166
- caution, 143
- defined, 143, 148
- example, 143
- in experimental design, 233
- observational studies and, 224
- Simpson's paradox, 167

Main effects

- defined, 29–20, 29–37
- no, 29–18, 29–19, 29–20
- with no interaction, 29–18

Mann-Whitney test, 26–10, 26–35

Margin of error

- assignment to results, 396
- of confidence interval, 364, 395, 408
- confidence level trade-off, 361
- large-scale confidence intervals, 502
- sample proportion, 502
- sample size for, 401–402, 503
- small, getting, 361

Marginal distributions

- bar graph, 161
- calculating, 160–161
- for categorical variables, 161
- defined, 160, 168
- percents display, 160

Matched pairs design

- as block design, 237
- defined, 236, 240, 449
- examples, 236
- principles of, 238–239

Matched pairs t procedures, 451

- defined, 449, 455

example, 449–450

inference procedures, 451

parameters, 449

Matched pairs t test, 26–35

Matched pairs, Wilcoxon signed rank test for,

- 26–19–26–22

Mean comparisons, 623–654

Mean squares

- defined, 640
- for error (MSE), 641, 642
- for groups (MSG), 641

Means

- binomial distributions, 338–340, 344
- chi-square distribution, 571
- comparison, 537
- conditions for inference about, 437–438
- conditions for inference comparing, 467
- defined, 40, 58
- density curves, 74, 89
- estimation in control chart setup, 27–26–27–27
- example, 40
- formula, 40
- least-squares residuals, 137
- median comparison, 42–43
- multiple linear regression model, 28–33
- Normal curves, 75
- Normal distributions, 76
- outliers and, 52
- population, 286, 301, 357–361
- random samples, 298
- regression towards, 127
- sample, 286, 301
- of sample mean, 293
- sampling distribution, 301, 494, 516
- uses, 58

Median

- defined, 41, 58
- density curves, 73, 74, 89

examples, 41–42

- location of, 41
- mean comparison, 42–43
- steps for finding, 41

Midpoint, 15

Minitab

- analysis of variance, 628, 29–10
- binomial probabilities, 337
- cumulative distribution function, 83
- descriptive measures, 53
- inverse cumulative distribution function, 87
- Kruskal-Wallis test, 26–32
- least-squares regression, 131
- multiple regression model, 28–50
- Normal approximation, 26–23
- output, 54
- parallel regression lines, 28–14
- proportion comparison, 519
- regression inference, 594, 605
- regression output, 603
- residuals, 612
- significance test output, 506
- sums of squares, 29–33
- t confidence interval, 447
- t test, 448
- two-sample problems, 475, 476
- two-way ANOVA, 29–25, 29–26, 29–28, 29–29, 29–30
- two-way tables, 164, 554–555, 565, 567

Models

- continuous probability, 271–272, 26–4
- finite probability, 268–271
- infinite probability, 332
- multiple linear regression, 28–33–28–34
- multiple regression, 28–6–28–8
- probability, 264–265, 273–274, 278–279
- regression, 28–55, 28–61

Multiple analyses, 399–400

Multiple comparisons

- contrasts, 29–12–29–16
- defined, 557
- mean, 625
- problem of, 556–558, 29–3
- procedures, 29–11
- statistical methods for, 557
- steps for, 625
- Tukey pairwise, 29–12–29–16, 29–37

Multiple linear regression model

- defined, 28–33
- example, 28–33–28–34
- general, 28–32–28–38

- mean response, 28–33
- standard deviation, 28–33

Multiple regression, 28–3–28–75

- ANOVA F test, 28–54
- automated algorithms, 28–46
- case study, 28–41–28–53
- defined, 28–3, 28–61
- degrees of freedom, 28–8

- estimated model, 28-61
F statistic formula, 28-18
 inference, 28-16–28-26
 inference, checking conditions for, 28-58–28-60
 interaction, 28-26–28-32
 linear model, 28-32–28-38
 model, building, 28-41
 model interpretation, 28-6–28-7
 parallel regression lines, 28-4–28-7, 28-61
 parameter estimation, 28-8–28-13
 regression standard error, 28-8
 residual plots, 28-59
 skills, 28-61
 technology use in, 28-13–28-16
 two regression lines, 28-61
- Multiplication, 316–318
 Multiplication rule for independent events. *See also* Probability rules
 caution, 311
 defined, 309
 examples, 309, 310–311
 extension, 309
- Multistage samples, 209
- Natural tolerances
 defined, 27-33, 27-42
 example, 27-33–27-34
- Negatively associated variables, 102, 108, 113
- Nonparametric tests, 26-3–26-41
 defined, 26-4, 26-34
 Kruskal-Wallis, 26-27–26-34
 rank, 26-4–26-27, 26-34
- Nonresponse, 211–212
- Normal approximation
 accuracy, 340
 to binomial distributions, 340–343
 defined, 340, 344
 example, 340–341
 rank sum statistic, 26-8–26-10
 sampling distribution, 516
 satisfactory dependence, 343
 signed rank statistic, 26-24–26-27
- Normal Approximation to Binomial* applet, 343
- Normal Curve* applet, 83
- Normal curves
 areas under, 81
 characteristics of, 75
 cumulative proportions for, 83
 defined, 75, 89
 illustrated, 75
 locating points on, 87
 standard, 84
 standard deviation, 75
- Normal distributions
 abbreviation, 79
 as ANOVA condition, 634, 645
 central limit theorem and, 296
 chance outcomes, 76
- conditions, checking, 360
 conditions for inference about a mean, 437
 as continuous probability models, 273
 defined, 75, 76, 89
 importance of, 76
 mean, 58, 76
 population, 294
 procedures for, 394
 real data, 78
 regression inference, 589
 sample mean, 301
 68-95-99.7 rule, 77–78, 89
 skills, 178
 smoothing, 83
 standard, 80–81, 89
 standard deviation, 58, 76
 standardized scale, 89
 statistical inference procedures based on, 76
 tests for population mean, 381
- Normal probability model, 273–274
- Normal proportions
 cumulative proportions and, 82
 example, 82
 finding, 81–83
 standard, 83
 using table to find, 84
- Null hypothesis
 chi-square test, 569
 defined, 372, 384
F statistic, 28-19
 Kruskal-Wallis test, 26-35
 significance tests for, 384
 test for, 379
 test statistic for, 504
- Null model, 28-18, 28-61
- Numerical summaries of distributions, 177–178
- Observational studies
 defined, 223, 239
 experiments versus, 223–225
 failure of, 225
 lurking variables and, 224
 survey samples as, 223
- Observations
 independent, 331, 607, 611
 influential, 139–142, 148
 ranking, 26-5
 1.5 × IQR rule, 48–49
- One-sample *t* confidence interval
 defined, 441
 example, 441
- One-sample *t* statistic, 439
 computing, 443
 defined, 455
- One-sample *t* test
 defined, 443
 example, 444–445
P-value, 445
- One-sample *z* test statistic, 379, 384, 455
 One-sided alternative hypothesis, 372, 373, 384
One-Variable Statistical Calculator applet
 histogram function, 14
 splitting stems, 22
- One-way ANOVA, 632, 645, 29-3
 example, 29-4–29-6
F statistic, 29-33
F test, 29-34
 sums of squares, 29-33
 total variation, 29-33–29-34
 two-way ANOVA comparison, 29-35
- One-Way ANOVA* applet, 632, 633
- Outliers
 defined, 15, 26
 example, 16
 finding explanations for, 52
 influential observations and, 140
 least-squares line, 140
 mean and, 52
 not hiding, 40
 1.5 × IQR rule for, 48–49
 in regression, 140
 in scatterplot interpretation, 101
 standard deviation and, 52
 suspected, spotting, 48–49
x direction, 141
 z procedures and, 394
- Out-of-control signals, 27-24–27-25, 27-38, 27-42
- Overall confidence levels, 29-9, 29-37
- Overall patterns, 15
 defined, 26
 density curve as description of, 72
 regression line description, 135
 in scatterplot interpretation, 101
- Overall significance levels, 29-9–29-10
- p* charts
 control limits for, 27-37–27-41
 defined, 27-36–27-37, 27-42
 example, 27-37
 interpretation of, 27-42
 out-of-control signals, 27-38
 with past data, 27-38
- Pairwise differences, 29-9
- Parallel regression lines. *See also* Regression lines
 conditions for inference, 28-17
 defined, 28-61
 example, 28-4–28-5
F test for, 28-18–28-19
 indicator variable, 28-6
 model illustration, 28-7
 scatterplot, 28-21
 technology use in, 28-14
- Parameters
 defined, 285, 301
 estimating, 590–592
 estimating with statistics, 285
 unbiased estimator of, 293

- Pareto charts, 27-7, 27-42
- Percents
- bar graph, 165
 - marginal distributions in, 161
 - two-way tables and, 161
- Permutation tests, 26-3–26-4
- Personal probability, 276–278
- defined, 277
 - example, 277
- Pie charts
- categories, 8
 - for data representation, 9
 - defined, 8, 26
 - illustrated, 8
- Placebo effect, 234
- Placebos
- benefits of, 253
 - controls, 235
 - defined, 234, 240
- Planning studies
- for confidence intervals, 401–402
 - for statistical test, 402–408
- Plus four confidence interval
- for comparing proportions, 527
 - for comparing two proportions, 520
 - defined, 499–500, 507
 - estimate, 500
 - example, 500
 - formula, 500
 - for proportions, 500
- Pooled sample proportions, 523, 527
- Pooled standard deviation, 642
- Population contrast, 29-12–29-16
- Population distributions
- defined, 291, 301
 - Normal, 294
 - shape of, 394
- Population mean, 455
- comparison, 466–469
 - comparison example, 467–469
 - comparison parameters, 467
 - confidence intervals for, 357–361, 364, 455
 - defined, 286, 301
 - inference about, 437–463
 - pairwise differences, 29-9
 - significance tests for, 378–382, 455
- Population regression line
- defined, 613
 - estimating, 590
 - mean response, 28-6
 - residuals, 591
 - slope, 600
 - straight-line relationships, 589
- Populations
- defined, 200
 - hypothesis reference, 372
 - identifying, 200
 - inference about, 208–209
 - proportions, 493–513
- in sample survey, 215
- samples versus, 199–202
- standard deviation, 286
- Positively associated variables, 102, 108, 113
- Power of a Test* applet, 403–404
- Power of statistical test
- with applet, 403–404
 - calculating, 403, 406
 - defined, 403, 409
 - examples, 402–405, 407
 - questions, 402
 - at significance levels, 405
- Prediction, inference about, 602–604
- Prediction intervals
- defined, 604, 613
 - form, 605
 - for individual future responses, 28-61
 - interpretation of, 605
 - meaning of, 604
 - regression models, 28-55
 - for regression response, 605
 - as rough approximations, 608
 - for single observation, 605, 608
- Predictor variables. *See* Explanatory variables
- Probability, 259–283. *See also* Probability rules
- addition rule for disjoint events, 267
 - as area under density curve, 273
 - assigning to sample space, 265
 - binomial, 334–338, 344
 - coin toss examples, 260, 261
 - computing with multiplication rule for independent events, 310–311
 - conditional, 314–315, 322
 - defined, 259, 261, 278, 418
 - disjoint events, 267
 - idea of, 260–262
 - Normal, 274
 - outcomes, 266
 - personal, 276–278
 - proportion of outcomes, 288
 - random numbers, 262–263
 - random variables, 275
 - skills, 420
- Probability* applet, 262
- Probability distribution, 279
- Probability models
- continuous, 271, 278–279
 - defined, 264, 278
 - description basis, 264
 - discrete, 268–271
 - example, 264, 265
 - finite, 268–271, 278
 - Normal, 273–274
 - software random number generator outcome, 272
- Probability rules, 266–268, 307–329
- addition rule for disjoint events, 308, 322
 - conditional probability, 314–315
- defined, 307
- general addition rule, 322
- general multiplication rule, 316–318, 322
- independence, 318
- multiplication rule for independent events, 308–311, 322
- skills, 421
- tree diagrams, 318–322
- Process monitoring
- conditions for, 27-10–27-11
 - in control process, 27-42
 - defined, 27-10–27-23
 - s charts for, 27-16–27-23
 - \bar{x} charts for, 27-10–27-23
- Processes. *See also* Statistical process control
- capability, 27-34–27-36, 27-42
 - cause-and-effect diagrams, 27-4, 27-7
 - common cause variation, 27-9
 - defined, 27-4, 27-42
 - flowcharts, 27-4, 27-6
 - focus on, 27-32
 - Normally distributed, 27-17
 - out of control, 27-42
 - pattern of variation, 27-9
 - skills, 27-43
 - special cause variation, 27-9
 - variation, 27-9
- Proportion comparison, 515–531
- confidence intervals for, 520–522
 - large-sample confidence intervals for, 517–518
- plus four confidence interval for, 520
- sample distribution, 516–517
- significance tests for, 522–526
- standard error, 517, 526
- technology use in, 518–520
- two-sample problems, 515–516
- Proportions
- confidence intervals for, 499–502, 507
 - cumulative, 82, 83, 89
 - finding values and, 86–88
 - large-sample confidence intervals for, 496–499
 - Normal, 81–83
 - plus four confidence interval for, 500
 - pooled sample, 523
 - population, 493–513
 - sample, 494–496
 - sample size selection, 502–504, 507
 - significance tests for, 504–507
 - two-sample problems, 515–516
- P-Value of a Test of Significance* applet, 375
- P-values
- calculating, 384
 - comparison, 377
 - defined, 374, 384, 418
 - examples, 374–375, 376
 - finding, 380

- fixed-standards for, 377
 one-sided *t* test, 445
 one-sided test, 397
 signed rank test, 26-35
 small, 374, 384, 408
 two-sided test, 376, 381
 Wilcoxon rank sum test, 26-35
- Quadratic regression, 28-35–28-37
 Quality
 idea of, 27-3
 satisfactory, 27-34
 Quantitative variables
 correlation and, 111
 defined, 4, 25
 histograms, 11–15
 statistics in summary, 175–176
 stemplots, 20–23
- Quartiles
 calculating, 44
 defined, 44, 58
 examples, 44–45
 first, 44, 58, 88
 third, 44, 58
- R charts, 27-23–27-24
 Random digit dialing. *See also* Simple random sample (SRS)
 computerized, 213
 defined, 207
 example, 206–207
 landline bias, 214
 in national sample surveys, 213–214, 216
 as outdated, 213
- Random digits
 table, 204–205, 239
 use of, 215
- Random numbers
 example, 271–272
 generators, 271–272
 getting, 262–263
 value, 262
- Random samples
 bias elimination, 260
 defined, 203, 215
 designs for, 209
 error, 395
 large versus small, 208
 means, 298
 nonresponse, 211–212, 216
 reasons to use, 208
 simple (SRS), 204, 215
 stratified, 209, 215
 trusting, 208
 undercoverage, 210, 216
 variability, 260
- Random variables
 continuous, 276, 279
 defined, 275, 279
- discrete, 276, 279
 probability distribution, 279
- Randomization
 carrying out, 239
 restricting, 240
 treatment groups, 239
- Randomized block designs, 29-17
- Randomized comparative experiments
 comparative design, 232
 control groups, 229
 defined, 229
 example, 229–230
 labels, 230
 logic of, 232–234
 random assignment, 232
 statistical significance, 233
 tables, 230
- Randomness
 defined, 261, 278
 search for, 262–264
- Rank sum statistic, 26-8–26-10
 Rank tests
 continuous distribution requirement, 26-4
 defined, 26-4, 26-34
 Kruskal-Wallis test, 26-27–26-34
 Mann-Whitney, 26-10
 Normal approximation for W, 26-8–26-10
 skills, 26-36
 technology use in, 26-10–26-12
 ties in, 26-14–26-19
 Wilcoxon rank sum test, 26-4–26-8
 Wilcoxon signed test, 26-19–26-22
- Ranks
 average, 26-14
 defined, 26-5
 skills, 26-35
- Rational subgroups
 defined, 27-32, 27-42
 example, 27-33
- RDD. *See* Random digit dialing
Reasoning of a Statistical Test applet, 369
 Recognition skills, 536–537, 645, 28-61, 29-37
 Regression, 125–157
 cautions, 142–144, 148, 179
 correlation and, 147
 explanatory and response variables distinction, 132
 least-squares, 128–135, 613
 lurking variables and, 143
 multiple, 28-3–28-75
 outliers in, 140
 quadratic, 28-35–28-37
 response, confidence/prediction intervals for, 605
 simple linear, 28-3
 slope, significance test for, 597
 towards the mean, 127
- Regression coefficients, 28-39–28-41
 confidence intervals, 28-55
t tests, 28-55
- Regression inference, 587–621
 about prediction, 602–607
 condition-checking skills, 614
 conditions, checking, 607–612
 conditions for, 589–590, 613
 conditions for inference about a mean, 589–593
 conditions illustration, 590
 constant standard deviation, 611
 example, 587–589
 independent observations, 611
 linear relationship, 609
 Normal residuals, 609
 parameter estimation, 590–592
 residual plots, 607
 slope and intercept parameters, 589
 standard deviation, 589
 straight-line relationship, 589
 technology use in, 593–596
- Regression lines
 defined, 125, 147
 equation from calculators/software programs, 127
 example, 125–127
 illustrated, 126
 intercept, 128, 147
 least-squares, 128–132, 147, 589
 overall pattern description, 135
 parallel, 28-4–28-7
 population, 589, 590, 613, 28-6
 prediction, 147
 skills, 179
 slope, 128, 147
 two, 28-27–28-29, 28-61
 use example, 127
 vertical distances, 128
- Regression slope
 confidence intervals for, 600–602
 significance test for, 597
t test, 613
- Regression standard error. *See also* Standard error
 defined, 591, 28-61
 degrees of freedom, 591, 613
 multiple regression model, 28-8
 as variability measure, 613
- Relationships. *See also* Linear relationships
 alternative hypothesis and, 556
 among variables, 28-43–28-44
 categorical variables, 553
 detecting with chi-square test, 564
 inference about, 549–550
 straight-line, 589
 variable, 99–101
- Residual plots
 defined, 137, 607
 illustrated, 138, 610
 multiple regression, 28-59
 predicted values and explanatory variables, 28-58
 quadratic regression, 28-36

- Residuals, 135–139
 calculation, 137
 data points, 137
 defined, 135, 591, 28–8
 example, 136–137
 least-squares, 137
 Minitab, 612
 negative, 137
 Normal, 608, 609
 standard deviation in, 608
- Resistant measure
 of center, 40
 changes and, 58
- Response bias, 212, 216
- Response variables. *See also* Variables
 in conditional distributions, 165
 correlation and, 108
 defined, 97, 113
 as dependent variables, 98
 examples, 98
 explanatory variable relationship, 125
 identifying, 98
- Robustness
 ANOVA, 634, 645
 defined, 26–3
 t procedures, 452–454
 two-sample t procedures, 477, 483
- Rounding data, 21
- Roundoff errors, 7, 161
- Rows
 individuals as, 5
 totals, 160, 168
 variables, 159, 168
- Runs signal, 27–42
- s charts. *See also* Control charts
 defined, 27–18, 27–42
 example, 27–18–27–20
 with past data, 27–28, 27–29
 in practice, 27–20
 for process monitoring, 27–16–27–23
 special causes, 27–21
- Sample contrast, 29–13–29–14
- Sample mean
 defined, 286, 301
 mean of, 293
 Normal distribution, 301
 sampling distribution of, 293–295
 small differences, 632
 standard deviation of, 293
 standardized, 455
- Sample proportion
 confidence interval, 499
 control charts for, 27–36–27–37
 defined, 494, 507
 example, 495
 formula, 494
 margin of error, 502
 pooled, 523, 527
 sampling distribution of, 494
- Sample size, 293, 398–399
 for confidence intervals, 401–402
 for margin of error, 401–402, 503
 population proportion, 502–504, 507
 population size and, 401
 statistical significance and, 398–399
 in using t procedures, 452
- Sample space
 assigning probability to, 265
 complexity, 264
 defined, 264, 278
 finite, 275
 interval of numbers, 271, 278
- Sample standard deviations, 286
- Sample surveys
 cautions about, 210–213
 defined, 215
 example, 200–201
 measurement item, 200
 national, 213, 216
 planning, 200–201
 population identification, 200
 sampling design, 201
 wording of questions, 212–213, 216
- Samples
 conclusions based on, 200
 convenience, 202
 defined, 200
 large versus small, 293
 multistage, 209
 nonresponse, 211–212, 216
 populations versus, 199–202
 random, 203–208
 random question about, 206
 range, 27–23
 rational subgroups, 27–32, 27–33, 27–42
 response bias, 212
 in sample survey, 215
 undercoverage, 210, 216
 voluntary response, 203, 215
- Sampling, 199–221
 binomial distributions in, 333–334
 cautions, 210–213
 convenience samples, 202
 designs, 209–210
 example, 202
 populations versus samples, 199–202
 random, 215
 simple random sample (SRS), 203–208
 skills, 419–420
 technology impact, 213–215
 voluntary response, 203
 voluntary response samples, 203
- Sampling design, 215
 data analysis from, 210
 defined, 200
- in sample survey, 201
 stratified, 209
- Sampling distributions, 285–305, 421
 center, 292
 central limit theorem and, 295–301
 of counts, 333
 defined, 291, 301
 of difference between proportions, 516–517
 as ideal pattern, 291
 illustrated, 291
 mean, 301, 494, 516
 Normal approximation, 516
 population distributions versus, 291
 of sample mean, 293–295, 421
 of sample proportion, 494
 shape, 292, 293–294
 simulation, 290
 skills, 421
 spread, 292
 standard deviation, 293, 301, 494, 516
- Scatterplots
 adding categorical variables to, 104–106
 defined, 100, 113
 direction, 101, 113
 displaying variable relationships with, 99–101
 example, 99–100
 form, 101, 113
 illustrated, 100, 103, 105, 107
 inference about population correlation, 599
 interpretation examples, 102–104
 interpreting, 101–104
 parallel regression lines, 28–21
 skills, 178–179
 strength, 101, 113
- Shape
 defined, 15, 26
 example, 16
 irregular, 19
 sampling distribution, 292, 293–294
- Signed rank statistic
 defined, 26–20
 Normal approximation for, 26–22
 P -values, 26–35
- Significance
 alternative hypothesis dependence, 397–398
 example, 399
 lack of, 409
 multiple analyses and, 399–400
 sample size and, 398–399
 statistical, 384
 from tables, 382–384
- Significance levels
 defined, 377, 384
 desired power at, 405
 overall, 29–9–29–10
 question, 402
- Significance tests, 369–389
 basic idea, 369
 cautions about, 397–401

- for comparing proportions, 522–526, 527
 defined, 351, 369, 384
 four-step process, 379
 hypotheses, 372–374
 null hypothesis, 384
 for population mean, 378–382, 437, 455
 power of, 402–408
 for proportions, 504–507
 purpose of, 397
 P -value and, 373
 questions, 402
 reasoning of, 370–372, 384
 reasoning of tests of significance, 370–371
 for regression slope, 597
 robust, 452
 skills, 422
 trust of, 392
 Type I error, 406–407
 Type II error, 406–407
 z statistic, 504, 508
- Simple linear regression**, 28–3
- Simple Random Sample** applet, 204, 205–206, 231
- Simple random sample (SRS)**
 choosing, 205
 concept, 204
 conditions, 352
 conditions, checking, 359
 conditions for inference, 393
 conditions for inference about a mean, 437
 defined, 204, 215
 independent, 634, 645
 labels, 205, 206, 215
 random digit dialing, 207
 reasons for use, 204
 table of random digits, 204–205
 tables, 205, 206
 tests for population mean, 381
- Simpson's paradox**
 defined, 167, 169
 example, 166–167
 lurking variables, 167
- Simulations**
 defined, 290
 examples, 290–291
- Simultaneous tests**, 29–9
- Single-peaked distribution**, 17
- 68-95-99.7 rule**
 defined, 77, 89
 example, 77–78
 illustrated, 77
- Skewed distributions**
 defined, 26
 direction, 22
 five-number summary for, 52
 IQR and, 48
 to the left, 16
 mean, 43
 median, 43
 to the right, 16
- sides having different spreads, 51
 in the world, 20
- Skewness**, 17, 452
- Slope**
 confidence intervals for, 613
 defined, 127, 147
 hypothesis, 613
 least-squares, 598
 least-squares regression line, 130
 population regression line, 600
 regression, 597, 600–602, 613
 regression inference, 589
 regression line, 128, 147
- Social science experiments**, 253–254
- Special cause variation**, 27–9, 27–21
- Split stems**, 22
- Spread**, 49–51, 58
 defined, 15, 26
 example, 16
 measures, choosing, 51–53
 measuring, quartiles, 43–45
 measuring, standard deviation, 49–51
 numerical measures of, 52
 overall, 18–19
 quartiles, 43–45
 sampling distribution, 292
 standard deviation, 49–51
 variance, 49
- Spreadsheets**, 5
- Squared correlation**
 defined, 133, 147
 equation, 134
 use example, 134
- Squared deviations**, 51
- Squared multiple correlation coefficient**, 28–15
- SRS**. *See Simple random sample*
- SSG (sum of squares for groups)**, 26–29
- Standard deviations**
 above the mean (UCL), 27–18
 in ANOVA, 634, 645
 avoiding inference about, 482
 below the mean (LCL), 27–18
 binomial distributions, 338–340, 344
 calculation, 49
 calculation example, 50
 defined, 49, 58
 degrees of freedom, 50
 density curves, 74, 89
 estimation in control chart setup, 27–27
 F test for, 482
 least-squares regression line, 133
 multiple linear regression model, 28–33
 Normal curves, 75
 Normal distributions, 76
 outliers and, 52
 pooled, 642
 population, 286
 regression inference, 589, 611
- in residuals, 608
 sample, 286
 sample mean, 293
 sampling distributions, 301, 494, 516
 units of measure, 51
 uses, 58
- Standard error**
 confidence intervals, 496, 517
 defined, 438, 455
 proportion comparison, 517, 526
 regression, 591, 592, 613, 28–8, 28–61
- Standard Normal curve**, 84
- Standard Normal distribution**, 81, 89
- Standard Normal table**, 83–86
- Standardized values**, 107
- Standardized variables**, 89
- Standardizing**
 defined, 80, 89
 example, 80
 symmetric distributions, 81
 z -score and, 80
- Statistical design**
 block, 236–239
 in experiments, 228, 233
 matched pairs, 236–237
 for sample selection, 199
- Statistical estimation**, 352–354
- Statistical inference**. *See Inference*
- Statistical problems**
 conclude, 55, 58
 four-step process, 55
 least-squares regression, 130
 organizing, 55–58
 plan, 55, 58
 process example, 56–57
 real-world setting, 58
 solve, 55, 58
 state, 55, 58
- Statistical process control**, 27–3–27–49
 benefits of, 27–33
 capability versus, 27–34–27–36
 comments on, 27–32–27–34
 common cause variation, 27–9
 control charts, 27–9, 27–10
 defined, 27–9, 27–42
 idea of, 27–9–27–10
 process focus, 27–32
 quality, 27–3
 rational subgroups, 27–32–27–33
 special cause variation, 27–9
- Statistical significance**, 233, 377, 384
- Statistical studies**
 biased, 202
 questions to ask for, 4
- Statistics**
 chi-square, 560–561
 defined, 199, 285
 Kruskal-Wallis, 26–30
 purpose of, 197

Statistics (*continued*)
 signed rank, 26-20, 26-22
 standard error, 438
 in summary, 417, 418, 419, 535
 t , 455
 test, 374
 Wilcoxon rank sum, 26-6
 z , 504, 508
 Stemplots. *See also* Quantitative variables
 back-to-back, 33
 creating, 20
 creation example, 20
 defined, 20, 26
 example, 21-22
 histograms versus, 20
 illustrated, 21, 22
 large data sets and, 20
 leaves, 20
 split stems, 20
 stems, 20
 Straight lines. *See also* Regression lines
 intercept, 127
 relationships, 589
 review, 127
 slope, 127
 Strata
 choosing, 209
 defined, 209, 215
 Stratified random samples
 choosing, 215
 defined, 209
 Strength, 101, 113
 Subjects
 assigning to treatments, 227
 defined, 226, 239
 informed consent, 248-250
 using enough, 233
 Sum of squares for groups (SSG), 26-29
 Sums of squares, 642, 29-33
 Survey samples, as observational study, 223
 Symmetric density curves, 73-74, 89
 Symmetric distributions
 defined, 16, 26
 example, 17
 mean of, 43
 median of, 43
 t approximation
 defined, 480
 degrees of freedom, 481
 details of, 480-481
 example, 480
 t distributions
 defined, 438
 degrees of freedom, 439, 455
 density curves for, 439
 example, 440
 facts about, 439-440
 t procedures
 individual, 28-61

matched pairs, 449-451, 455
 one-sample, 440
 robustness of, 452-454, 455
 use decision, 453
 using, 452
 t statistic
 one-sample, 439, 455
 significance tests based on, 455
 two-sample, 469, 470
 t tests
 for coefficients, 28-19
 individual, 28-19-28-20
 matched pairs, 26-35
 regression coefficients, 28-55
 regression slope, 613
 Table A
 backward use, 87
 examples, 84, 85
 in finding normal proportions, 84-86
 standardize, 84, 85
 state the problem and draw picture, 84, 85, 87
 unstandardize, 87
 using, 84, 85, 87
 Tables
 ANOVA, 642, 645, 28-53
 random digits, 204-205, 239
 randomized comparative experiments, 230
 significance from, 382-384
 SRS, 205, 206
 two-way, 159-173, 553-556, 558-560
 Test statistics
 chi-square, 560-570
 defined, 374, 384
 example, 380
 for null hypothesis, 504
 one-sample z , 379, 384
 Tests of significance. *See* Significance tests
 Third quartile, 44, 58
 Three-sigma control charts. *See also* Control charts
 defined, 27-17, 27-42
 general pattern, 27-17
 Ties, rank test
 dealing with, 26-14-26-19
 example, 26-14-26-16
 Ties, signed rank test
 average of ranks, 26-25
 dealing with, 26-24-26-27
 example, 26-25
 Time plots
 cycles, 24, 26
 defined, 23, 26
 illustrated, 24
 time-series data, 25
 trends, 24, 26
 Time-series data, 25
 Total sum of squares, 29-33, 29-34-29-35
 Treatments
 assigning subjects to, 227, 240
 defined, 226, 239
 effects of, 227
 groups, 239
 Tree diagrams
 defined, 319, 322
 examples, 319-321
 illustrated, 319
 outcome sources, 321
 using, 318-322
 Trends, 24, 26
 Tukey pairwise multiple comparisons
 defined, 29-9, 29-37
 examples, 29-10-29-11
 follow-up analysis, 29-8-29-12
 Tukey simultaneous confidence intervals, 29-9
 Two regression lines model, 28-27-28-29, 28-61
 Two-peaked distribution, 18
 Two-sample problems, 465-491
 data, 482
 defined, 465
 example, 465-466
 inference procedures for, 465
 population mean comparison, 466-469
 proportions, 515-516, 526
 t procedures, 469-474
 technology use in, 474-477
 Two-sample t procedures
 confidence interval, 470
 defined, 470
 examples, 470-471, 472-473
 options, 470
 pooled, avoiding, 481-482
 robustness of, 477, 483
 standard error, 469
 t distribution, 469-470
 Two-sample t statistic
 approximate distribution of, 480
 calculating, 470
 defined, 469, 483
 degrees of freedom, 483
 pooled, 481-482
 Two-sided alternative hypothesis, 372, 375, 384
 Two-variable data, 111
Two-Variable Statistical Calculator applet, 101, 112
 Two-way ANOVA. *See also* Analysis of variance (ANOVA)
 balanced design, 29-17
 conditions for, 29-17, 29-37
 crossed and balanced, 29-37
 defined, 29-6, 29-17, 29-37
 degrees of freedom, 29-35
 details of, 29-32-29-36
 F tests, 29-35-29-36, 29-37
 Factor C, 29-32
 Factor R, 29-32
 independent SRSs, 29-17, 29-37
 inference for, 29-23-29-32

- interactions, 29-19–29-20
- main effects, 29-18, 29-19, 29-21, 29-37
- Normal distribution, 29-17, 29-37
- one-way ANOVA comparison, 29-35
- output, 29-32
- questions, 29-18
- responses for all combinations of values, 29-17
- skills, 29-37–29-38
- standard deviation, 29-17, 29-37
- with strong interaction, 29-28
- sums of squares, 29-33
- total sum of squares, 29-34–29-35
- of variance table, 29-34
- Two-way tables, 159–173
 - arising of, 567
 - column totals, 160, 168
 - column variables, 159, 168
 - conditional distributions, 162–166
 - CrunchIt! output, 164
 - defined, 159, 168, 553
 - example, 159–160, 553–554
 - expected counts in, 558–560, 576
 - marginal distributions, 160–162
 - Minitab output, 164, 565, 567
 - multiple comparisons, 556–558
 - percents calculation, 161
 - row totals, 160, 168
 - row variables, 159, 168
 - Simpson’s paradox, 166–168
 - skills, 577
- Type I error, 406–407
- Type II error, 406–407
- Unbiased estimators, 293, 301
- Undercoverage, 210
- Units of measure, 108
- Values
 - absolute, 26-20
 - categorical variable, 7
- critical, 357, 364, 455
- distribution, 15
- standardized, 107
- Variables
 - association, 144–147
 - categorical, 4, 6–11, 25
 - column, 5, 159, 168
 - confounded, 225, 239
 - in data table format, 5
 - defined, 3, 25
 - displaying relationships with scatterplots, 99–101
 - distribution, 7
 - economic, 43
 - explanatory, 97–99, 113
 - how many question, 4
 - indicator, 28-6, 28-61
 - inference about, 533–550
 - lurking, 143, 148
 - negatively associated, 102, 108, 113
 - positively associated, 102, 108, 113
 - quantitative, 4, 11–15, 20–23, 25, 175–176
 - random, 275–276, 279
 - relationships among, 97, 28-43–28-44
 - response, 97–99, 113
 - row, 159, 168
 - standardized, 89
 - time plots, 23–25
- Variance
 - analysis of, 623–654
 - defined, 49, 58
 - degrees of freedom, 50
- Venn diagrams
 - addition rule, 312
 - defined, 308
 - disjoint events, 308
 - events and probabilities, 313
 - illustrated, 308
- Voluntary response samples, 203, 215
- Web surveys, 214
- Wilcoxon rank sum statistic, 26-6, 26-35
- Wilcoxon rank sum test
 - defined, 26-6, 26-35
 - example, 26-7
 - hypotheses, 26-13–26-14
 - P-values, 26-35
- Wilcoxon signed rank statistic, 26-20, 26-35
- Wilcoxon signed rank test
 - defined, 26-20, 26-35
 - example, 26-21
 - for matched pairs, 26-19–26-22
 - ties, 26-24–26-27
- Wording of questions, 212–213, 216
- \bar{X} charts. *See also* Control charts
 - defined, 27-11, 27-18, 27-42
 - example, 27-18–27-20
 - gradual drift, 27-24
 - illustrated, 27-13
 - interpreting, 27-14–27-15
 - one-point-out, 27-24
 - with past data, 27-28, 27-30
 - in practice, 27-20
 - for process monitoring, 27-10–27-23
 - run, 27-24
 - sample range, 27-23
 - special causes, 27-21
- z procedures
 - defined, 391
 - outliers and, 394
 - use illustration, 418, 419
- z statistic, 523
- z tests, 380, 385
- z -scores
 - defined, 80, 89
 - standard scale, 83

Table entry for C is the critical value t^* required for confidence level C. To approximate one- and two-sided P-values, compare the value of the t statistic with the critical values of t^* that match the P-values given at the bottom of the table.

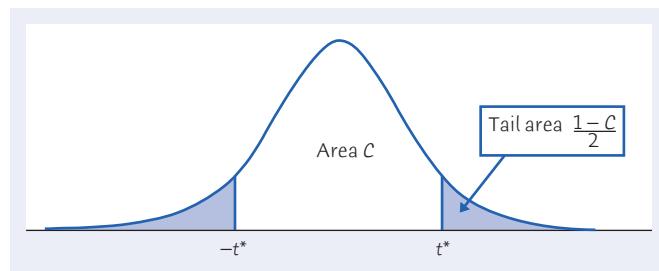


TABLE C *t* distribution critical values

DEGREES OF FREEDOM	CONFIDENCE LEVEL C											
	50%	60%	70%	80%	90%	95%	96%	98%	99%	99.5%	99.8%	99.9%
1	1.000	1.376	1.963	3.078	6.314	12.71	15.89	31.82	63.66	127.3	318.3	636.6
2	0.816	1.061	1.386	1.886	2.920	4.303	4.849	6.965	9.925	14.09	22.33	31.60
3	0.765	0.978	1.250	1.638	2.353	3.182	3.482	4.541	5.841	7.453	10.21	12.92
4	0.741	0.941	1.190	1.533	2.132	2.776	2.999	3.747	4.604	5.598	7.173	8.610
5	0.727	0.920	1.156	1.476	2.015	2.571	2.757	3.365	4.032	4.773	5.893	6.869
6	0.718	0.906	1.134	1.440	1.943	2.447	2.612	3.143	3.707	4.317	5.208	5.959
7	0.711	0.896	1.119	1.415	1.895	2.365	2.517	2.998	3.499	4.029	4.785	5.408
8	0.706	0.889	1.108	1.397	1.860	2.306	2.449	2.896	3.355	3.833	4.501	5.041
9	0.703	0.883	1.100	1.383	1.833	2.262	2.398	2.821	3.250	3.690	4.297	4.781
10	0.700	0.879	1.093	1.372	1.812	2.228	2.359	2.764	3.169	3.581	4.144	4.587
11	0.697	0.876	1.088	1.363	1.796	2.201	2.328	2.718	3.106	3.497	4.025	4.437
12	0.695	0.873	1.083	1.356	1.782	2.179	2.303	2.681	3.055	3.428	3.930	4.318
13	0.694	0.870	1.079	1.350	1.771	2.160	2.282	2.650	3.012	3.372	3.852	4.221
14	0.692	0.868	1.076	1.345	1.761	2.145	2.264	2.624	2.977	3.326	3.787	4.140
15	0.691	0.866	1.074	1.341	1.753	2.131	2.249	2.602	2.947	3.286	3.733	4.073
16	0.690	0.865	1.071	1.337	1.746	2.120	2.235	2.583	2.921	3.252	3.686	4.015
17	0.689	0.863	1.069	1.333	1.740	2.110	2.224	2.567	2.898	3.222	3.646	3.965
18	0.688	0.862	1.067	1.330	1.734	2.101	2.214	2.552	2.878	3.197	3.611	3.922
19	0.688	0.861	1.066	1.328	1.729	2.093	2.205	2.539	2.861	3.174	3.579	3.883
20	0.687	0.860	1.064	1.325	1.725	2.086	2.197	2.528	2.845	3.153	3.552	3.850
21	0.686	0.859	1.063	1.323	1.721	2.080	2.189	2.518	2.831	3.135	3.527	3.819
22	0.686	0.858	1.061	1.321	1.717	2.074	2.183	2.508	2.819	3.119	3.505	3.792
23	0.685	0.858	1.060	1.319	1.714	2.069	2.177	2.500	2.807	3.104	3.485	3.768
24	0.685	0.857	1.059	1.318	1.711	2.064	2.172	2.492	2.797	3.091	3.467	3.745
25	0.684	0.856	1.058	1.316	1.708	2.060	2.167	2.485	2.787	3.078	3.450	3.725
26	0.684	0.856	1.058	1.315	1.706	2.056	2.162	2.479	2.779	3.067	3.435	3.707
27	0.684	0.855	1.057	1.314	1.703	2.052	2.158	2.473	2.771	3.057	3.421	3.690
28	0.683	0.855	1.056	1.313	1.701	2.048	2.154	2.467	2.763	3.047	3.408	3.674
29	0.683	0.854	1.055	1.311	1.699	2.045	2.150	2.462	2.756	3.038	3.396	3.659
30	0.683	0.854	1.055	1.310	1.697	2.042	2.147	2.457	2.750	3.030	3.385	3.646
40	0.681	0.851	1.050	1.303	1.684	2.021	2.123	2.423	2.704	2.971	3.307	3.551
50	0.679	0.849	1.047	1.299	1.676	2.009	2.109	2.403	2.678	2.937	3.261	3.496
60	0.679	0.848	1.045	1.296	1.671	2.000	2.099	2.390	2.660	2.915	3.232	3.460
80	0.678	0.846	1.043	1.292	1.664	1.990	2.088	2.374	2.639	2.887	3.195	3.416
100	0.677	0.845	1.042	1.290	1.660	1.984	2.081	2.364	2.626	2.871	3.174	3.390
1000	0.675	0.842	1.037	1.282	1.646	1.962	2.056	2.330	2.581	2.813	3.098	3.300
z^*	0.674	0.841	1.036	1.282	1.645	1.960	2.054	2.326	2.576	2.807	3.091	3.291
One-sided P	.25	.20	.15	.10	.05	.025	.02	.01	.005	.0025	.001	.0005
Two-sided P	.50	.40	.30	.20	.10	.05	.04	.02	.01	.005	.002	.001

This page intentionally left blank



Nonparametric Tests

The most commonly used methods for inference about the means of quantitative response variables assume that the variables in question have Normal distributions in the population or populations from which we draw our data. In practice, of course, no distribution is exactly Normal. Fortunately, our usual methods for inference about population means (the one-sample and two-sample t procedures and analysis of variance) are quite robust. That is, the results of inference are not very sensitive to moderate lack of Normality, especially when the samples are reasonably large. Practical guidelines for taking advantage of the robustness of these methods appear in Chapters 18, 19, and 25.

What can we do if plots suggest that the data are clearly not Normal, especially when we have only a few observations? This is not a simple question. Here are the basic options:

1. If lack of Normality is due to outliers, it may be legitimate to **remove outliers** if you have reason to think that they do not come from the same population as the other observations. Equipment failure that produced a bad measurement, for example, entitles you to remove the outlier and analyze the remaining data. *But if an outlier appears to be “real data,” you should not arbitrarily remove it.* 
2. In some settings, **other standard distributions** replace the Normal distributions as models for the overall pattern in the population. The lifetimes in service of equipment or the survival times of cancer patients after treatment usually have right-skewed distributions. Statistical studies in these areas use families of right-skewed distributions rather than Normal distributions. There are inference procedures for the parameters of these distributions that replace the t procedures.
3. Modern **bootstrap methods** and **permutation tests** use heavy computing to avoid requiring Normality or any other specific form of sampling

Chapter 26

IN THIS CHAPTER WE COVER...

- Comparing two samples: the Wilcoxon rank sum test
- The Normal approximation for W
- Using technology
- What hypotheses does Wilcoxon test?
- Dealing with ties in rank tests
- Matched pairs: the Wilcoxon signed rank test
- The Normal approximation for W^+
- Dealing with ties in the signed rank test
- Comparing several samples: the Kruskal-Wallis test
- Hypotheses and conditions for the Kruskal-Wallis test
- The Kruskal-Wallis test statistic

distribution. We recommend these methods unless the sample is so small that it may not represent the population well. For an introduction, see Companion Chapter 16 of the somewhat more advanced text *Introduction to the Practice of Statistics*, available online at www.whfreeman.com/ips7e.

4. Finally, there are other **nonparametric methods**, which do not assume any specific form for the distribution of the population. Unlike bootstrap and permutation methods, common nonparametric methods do not make use of the actual values of the observations.

This chapter concerns one type of nonparametric procedure: tests that can replace the *t* tests and one-way analysis of variance when the Normality conditions for those tests are not met. The most useful nonparametric tests are **rank tests** based on the rank (place in order) of each observation in the set of all the data.

Figure 26.1 presents an outline of the standard tests (based on Normal distributions) and the rank tests that compete with them. The rank tests require that the population or populations have *continuous distributions*. That is, each distribution must be described by a *density curve* (Chapter 3, page 71) that allows observations to take any value in some interval of outcomes. The Normal curves are one shape of density curve. Rank tests allow curves of any shape.

The rank tests we will study concern the *center* of a population or populations. When a population has at least roughly a Normal distribution, we describe its center by the mean. The “Normal tests” in Figure 26.1 all test hypotheses about population means. When distributions are strongly skewed, we often prefer the median to the mean as a measure of center. In simplest form, the hypotheses for rank tests just replace mean by median.

FIGURE 26.1

Comparison of tests based on Normal distributions with rank tests for similar settings.

Setting	Normal test	Rank test
One sample	One-sample <i>t</i> test Chapter 18	Wilcoxon signed rank test
Matched pairs	Apply one-sample test to differences within pairs	
Two independent samples	Two-sample <i>t</i> test Chapter 19	Wilcoxon rank sum test
Several independent samples	One-way ANOVA <i>F</i> test Chapter 25	Kruskal-Wallis test

We begin by describing the most common rank test, for comparing two samples. In this setting we also explain ideas common to all rank tests: the big idea of using ranks, the conditions required by rank tests, the nature of the hypotheses tested, and the contrast between exact distributions for use with small samples and Normal approximations for use with larger samples.

COMPARING TWO SAMPLES: THE WILCOXON RANK SUM TEST

Two-sample problems (see Chapter 19) are among the most common in statistics. The most useful nonparametric significance test compares two distributions. Here is an example of this setting.

EXAMPLE 26.1 Weeds among the corn

STATE: Does the presence of small numbers of weeds reduce the yield of corn? Lamb's-quarter is a common weed in corn fields. A researcher planted corn at the same rate in 8 small plots of ground, then weeded the corn rows by hand to allow no weeds in 4 randomly selected plots and exactly 3 lamb's-quarter plants per meter of row in the other 4 plots. Here are the yields of corn (bushels per acre) in each of the plots:¹

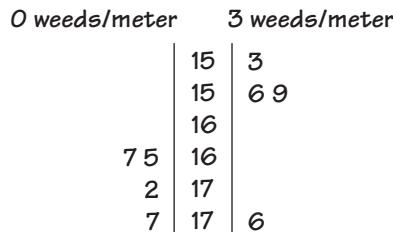


WEEDS3

0 weeds per meter	166.7	172.2	165.0	176.9
3 weeds per meter	158.6	176.4	153.1	156.0

PLAN: Make a graph to compare the two sets of yields. Test the hypothesis that there is no difference against the one-sided alternative that yields are higher when no weeds are present.

SOLVE (first steps): A back-to-back stemplot (Figure 26.2) suggests that yields may be higher when there are no weeds. There is one outlier; because it is correct data, we cannot remove it. The samples are too small to rely on the robustness of the two-sample *t* test. We will now develop a test that does not require Normality. ■

**FIGURE 26.2**

Back-to-back stemplot of corn yields from plots with no weeds and with 3 weeds per meter of row, for Example 26.1. Notice the split stems, with leaves 0 to 4 on the first stem and leaves 5 to 9 on the second stem.

First, arrange all 8 observations from both samples in order from smallest to largest:

153.1 156.0 158.6 **165.0** **166.7** 172.2 176.4 176.9

The boldface entries in the list are the yields with no weeds present. We see that four of the five highest yields come from that group, suggesting that yields are higher with no weeds. The idea of rank tests is to look just at position in this ordered list. To do this, replace each observation by its order, from 1 (smallest) to 8 (largest). These numbers are the *ranks*:

Yield	153.1	156.0	158.6	165.0	166.7	172.2	176.4	176.9
Rank	1	2	3	4	5	6	7	8

RANKS

To rank observations, first arrange them in order from smallest to largest. The **rank** of each observation is its position in this ordered list, starting with rank 1 for the smallest observation.

Moving from the original observations to their ranks retains only the ordering of the observations and makes no other use of their numerical values. Working with ranks allows us to dispense with specific conditions on the shape of the distribution, such as Normality.

If the presence of weeds reduces corn yields, we expect the ranks of the yields from plots without weeds to be larger as a group than the ranks from plots with weeds. Let's compare the sums of the ranks from the two treatments:

Treatment	Sum of ranks
No weeds	23
Weeds	13

These sums measure how much the ranks of the weed-free plots as a group exceed those of the weedy plots. In fact, the sum of the ranks from 1 to 8 is always equal to 36, so it is enough to report the sum for one of the two groups. If the sum of the ranks for the weed-free group is 23, the ranks for the other group must add to 13 because $23 + 13 = 36$. If the weeds have no effect, we would expect the sum of the ranks in either group to be 18 (half of 36). Here are the facts we need in a more general form that takes account of the fact that our two samples need not be the same size.

THE WILCOXON RANK SUM TESTS

Draw an SRS of size n_1 from one population and draw an independent SRS of size n_2 from a second population. There are N observations in all, where $N = n_1 + n_2$. Rank all N observations. The sum W of the ranks for the first sample is the **Wilcoxon rank sum statistic**. If the two populations have the same continuous distribution, then W has mean

$$\mu_W = \frac{n_1(N + 1)}{2}$$

and standard deviation

$$\sigma_W = \sqrt{\frac{n_1 n_2 (N + 1)}{12}}$$

The **Wilcoxon rank sum test** rejects the hypothesis that the two populations have identical distributions when the rank sum W is far from its mean.

In the corn yield study of Example 26.1, we want to test the hypotheses

$$H_0: \text{no difference in distribution of yields}$$

$$H_a: \text{yields are systematically higher in weed-free plots}$$

Our test statistic is the rank sum $W = 23$ for the weed-free plots.

EXAMPLE 26.2 Weeds among the corn: inference

SOLVE: First note that the conditions for the Wilcoxon test are met: the data come from a randomized comparative experiment and the yield of corn in bushels per acre has a continuous distribution.

There are $N = 8$ observations in all, with $n_1 = 4$ and $n_2 = 4$. The sum of ranks for the weed-free plots has mean

$$\begin{aligned}\mu_W &= \frac{n_1(N+1)}{2} \\ &= \frac{(4)(9)}{2} = 18\end{aligned}$$

and standard deviation

$$\begin{aligned}\sigma_W &= \sqrt{\frac{n_1 n_2 (N+1)}{12}} \\ &= \sqrt{\frac{(4)(4)(9)}{12}} = \sqrt{12} = 3.464\end{aligned}$$

Although the observed rank sum $W = 23$ is higher than the mean, it is only about 1.4 standard deviations higher. We now suspect that the data do not give strong evidence that yields are higher in the population of weed-free corn.

The P -value for our one-sided alternative is $P(W \geq 23)$, the probability that W is at least as large as the value for our data when H_0 is true. Software tells us that this probability is $P = 0.1$.

CONCLUDE: The data provide some evidence ($P = 0.1$) that corn yields are lower when weeds are present. There are only 4 observations in each group, so even quite large effects can fail to reach the levels of significance usually considered convincing, such as $P < 0.05$. A larger experiment might clarify the effect of weeds on corn yield. ■

APPLY YOUR KNOWLEDGE

- 26.1 Daily activity and obesity.** Our lead example for the two-sample t procedures in Chapter 19 concerned a study comparing the level of physical activity of lean and mildly obese people who don't exercise. Here are the minutes per day that the subjects spent standing or walking over a 10-day period:  ACTIVITY

Lean subjects	Obese subjects	
511.100	543.388	260.244
607.925	677.188	464.756
319.212	555.656	367.138
584.644	374.831	267.344
578.869	504.700	413.667
		410.631
		347.375
		426.356

The data are a bit irregular but not distinctly non-Normal. Let's use the Wilcoxon test for comparison with the two-sample t test.

- Find the median minutes spent standing or walking for each group. Which group appears more active?
- Arrange all 20 observations in order and find the ranks.
- Take W to be the sum of the ranks for the lean group. What is the value of W ? If the null hypothesis (no difference between the groups) is true, what are the mean and standard deviation of W ?
- Does comparing W with the mean and standard deviation suggest that the lean subjects are more active than the obese subjects?

26.2 How strong are durable press fabrics? Exercise 19.38 (text page 488) describes an experiment comparing the strengths of cotton fabric treated with two “durable press” processes. Here are the breaking strengths in pounds:



Permafresh	29.9	30.7	30.0	29.5	27.6
Hylite	28.8	23.9	27.0	22.1	24.2

There is a mild outlier in the Permafresh group. Perhaps we should use the Wilcoxon test.

- Arrange the breaking strengths in order and find their ranks.
- Find the Wilcoxon statistic W for the Permafresh group, along with its mean and standard deviation under the null hypothesis (no difference between the groups).
- Is W far enough from the mean to suggest that there may be a difference between the groups?

THE NORMAL APPROXIMATION FOR W

To calculate the P -value $P(W \geq 23)$ for Example 26.2, we need to know the sampling distribution of the rank sum W when the null hypothesis is true. This distribution depends on the two sample sizes n_1 and n_2 . Tables are therefore unwieldy. Most statistical software will give you P -values, as well as carry out the ranking and calculate W . However, many software packages give only approximate P -values. You must learn what your software offers.

With or without software, P -values for the Wilcoxon test are often based on the fact that the rank sum statistic W becomes approximately Normal as the two sample sizes increase. We can then form yet another z statistic by standardizing W :

$$\begin{aligned} z &= \frac{W - \mu_W}{\sigma_W} \\ &= \frac{W - n_1(N + 1)/2}{\sqrt{n_1 n_2(N + 1)/12}} \end{aligned}$$

Use standard Normal probability calculations to find P -values for this statistic. Because W takes only whole-number values, an idea called the *continuity correction* improves the accuracy of the approximation.

CONTINUITY CORRECTION

To apply the **continuity correction** in a Normal approximation for a variable that takes only whole-number values, act as if each whole number occupies the entire interval from 0.5 below the number to 0.5 above it.

EXAMPLE 26.3 Weeds among the corn: Normal approximation

The standardized rank sum statistic W in our corn yield example is

$$z = \frac{W - \mu_W}{\sigma_W} = \frac{23 - 18}{3.464} = 1.44$$

We expect W to be larger when the alternative hypothesis is true, so the approximate P -value is (from Table A)

$$P(Z \geq 1.44) = 0.0749$$

We can improve this approximation by using the continuity correction. To do this, act as if the whole number 23 occupies the entire interval from 22.5 to 23.5. Calculate the P -value $P(W \geq 23)$ as $P(W \geq 22.5)$ because the value 23 is included in the range whose probability we want. Here is the calculation:

$$\begin{aligned} P(W \geq 22.5) &= P\left(\frac{W - \mu_W}{\sigma_W} \geq \frac{22.5 - 18}{3.464}\right) \\ &= P(Z \geq 1.30) \\ &= 0.0968 \end{aligned}$$

This is close to the software value, $P = 0.1$. If you do not use the exact distribution of W (from software or tables), you should always use the continuity correction in calculating P -values. ■

APPLY YOUR KNOWLEDGE

26.3 Daily activity and obesity, continued. In Exercise 26.1, you found the Wilcoxon rank sum W and its mean and standard deviation. We want to test the null hypothesis that the two groups don't differ in activity against the alternative hypothesis that the lean subjects spend more time standing and walking.  ACTIVITY

- (a) What is the probability expression for the P -value of W if we use the continuity correction?
- (b) Find the P -value. What do you conclude?

26.4 Strength of durable press fabrics, continued. Use your values of W , μ_W , and σ_W from Exercise 26.2 to see whether fabrics treated with the two processes differ in breaking strength.  FABRICS

- (a) The two-sided P -value is $2P(W \geq ?)$. Using the continuity correction, what number replaces the $?$ in this probability?
 (b) Find the P -value. What do you conclude?



26.5 Tell me a story. A study of early childhood education asked kindergarten students to tell fairy tales that had been read to them earlier in the week. The 10 children in the study included 5 high-progress readers and 5 low-progress readers. Each child told two stories. Story 1 had been read to them; Story 2 had been read and also illustrated with pictures. An expert listened to a recording of the children and assigned a score for certain uses of language. Here are the data.²

Child	Progress	Story 1 score	Story 2 score
1	high	0.55	0.80
2	high	0.57	0.82
3	high	0.72	0.54
4	high	0.70	0.79
5	high	0.84	0.89
6	low	0.40	0.77
7	low	0.72	0.49
8	low	0.00	0.66
9	low	0.36	0.28
10	low	0.55	0.38

Look only at the data for Story 2. Is there good evidence that high-progress readers score higher than low-progress readers? Follow the four-step process as illustrated in Examples 26.1 and 26.2. STORY2

USING TECHNOLOGY

For samples as small as those in the corn yield study of Example 26.1, we prefer software that gives the exact P -value for the Wilcoxon test rather than the Normal approximation. Neither the Excel spreadsheet nor TI graphing calculators have menu entries for rank tests. Minitab offers only the Normal approximation.

EXAMPLE 26.4 Weeds among the corn: software output

Mann-Whitney test

Figure 26.3 displays output from CrunchIt! for the corn yield data. The top panel reports the exact Wilcoxon P -value as $P = 0.1$. The Normal approximation with continuity correction, $P = 0.0968$ in Example 26.3, is quite accurate. There are several differences between the CrunchIt! output and our work in Example 26.3. The most important is that CrunchIt! carries out the **Mann-Whitney test** rather than the Wilcoxon test. The two tests always have the same P -value because the two statistics are related by simple algebra.

The second panel in Figure 26.3 is the two-sample t test from Chapter 19, which does not assume that the two populations have the same standard deviation. It gives $P = 0.0937$, close to the Wilcoxon value. Because the t test is quite robust, it is somewhat unusual for P -values from t and W to differ greatly.

The bottom panel (on page 26-12) shows the result of the “pooled” version of t , now outdated, that assumes equal population standard deviations. You see that its P -value is a bit different from the others. We do not recommend its use in general, despite the reasonable agreement in this example. ■

APPLY YOUR KNOWLEDGE

- 26.6 Strength of durable press fabrics: software.** Use your software to repeat the Wilcoxon test you did in Exercise 26.4. By comparing the results, state how your software finds P -values for W : exact distribution, Normal approximation with continuity correction, or Normal approximation without continuity correction.

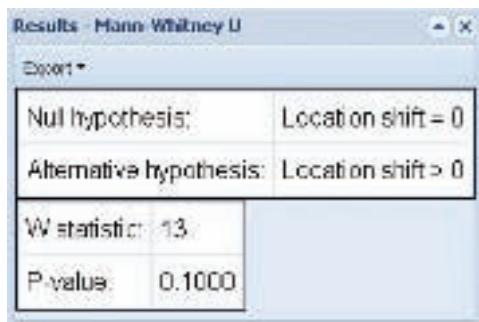
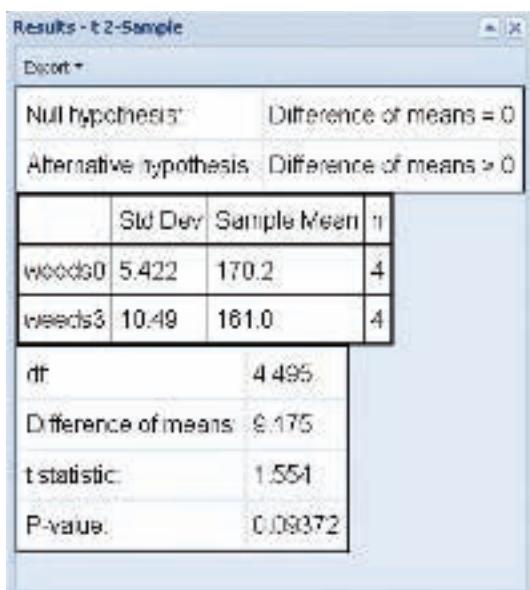
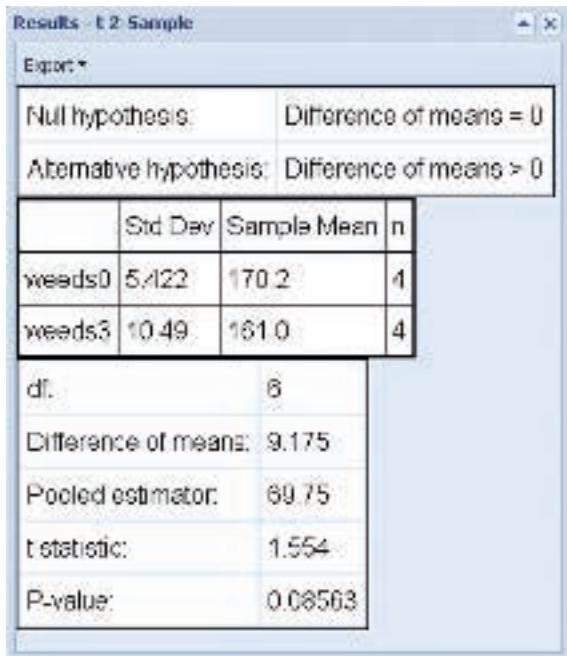


FIGURE 26.3

Output from CrunchIt! for the data of Example 26.1. The output compares the results of three tests that could be used to compare yields for the two groups of corn plots.



**FIGURE 26.3** (Continued)

26.7 Daily activity and obesity: software. Use your software to carry out the one-sided Wilcoxon rank sum test that you did by hand in Exercise 26.3. Use the exact distribution if your software will do it. Compare the software result with your result in Exercise 26.3.

26.8 Weeds among the corn. The corn yield study of Example 26.1 also examined yields in four plots having 9 lamb's-quarter plants per meter of row. The yields (bushels per acre) in these plots were WEEDS9

$$162.8 \quad 142.4 \quad 162.7 \quad 162.4$$

There is a clear outlier, but rechecking the results found that this is the correct yield for this plot. The outlier makes us hesitant to use t procedures because \bar{x} and s are not resistant.

- Is there evidence that 9 weeds per meter reduces corn yields when compared with weed-free corn? Use the Wilcoxon rank sum test with the data above and part of the data from Example 26.1 to answer this question.
- Compare the results from (a) with those from the two-sample t test for these data.
- Now remove the low outlier 142.4 from the data with 9 weeds per meter. Repeat both the Wilcoxon and t analyses. By how much did the outlier reduce the mean yield in its group? By how much did it increase the standard deviation? Did it have a practically important impact on your conclusions?

WHAT HYPOTHESES DOES WILCOXON TEST?

Our null hypothesis is that weeds do not affect yield. The alternative hypothesis is that yields are lower when weeds are present. If we are willing to assume that yields are Normally distributed, or if we have reasonably large samples, we can use the two-sample t test for means. Our hypotheses then have the form

$$H_0: \mu_1 = \mu_2$$

$$H_a: \mu_1 > \mu_2$$

When the distributions may not be Normal, we might restate the hypotheses in terms of population medians rather than means:

$$H_0: \text{median}_1 = \text{median}_2$$

$$H_a: \text{median}_1 > \text{median}_2$$

The Wilcoxon rank sum test provides a test of these hypotheses, but only if an additional condition is met: both populations must have distributions of *the same shape*. That is, the density curve for corn yields with 3 weeds per meter looks exactly like that for no weeds except that it may slide to a different location on the scale of yields. The CrunchIt! output in the top panel of Figure 26.3 states the hypotheses in terms of a location shift, the difference in the medians.

The same-shape condition is too strict to be reasonable in practice. Fortunately, the Wilcoxon test also applies in a more useful setting. It compares any two continuous distributions, whether or not they have the same shape, by testing hypotheses that we can state in words as

$$H_0: \text{the two distributions are the same}$$

$$H_a: \text{one has values that are systematically larger}$$

A more exact statement of the “systematically larger” alternative hypothesis is a bit tricky, so we won’t try to give it here.³ These hypotheses really are “nonparametric” because they do not involve any specific parameter such as the mean or median. If the two distributions do have the same shape, the general hypotheses reduce to comparing medians. Many texts and computer outputs state the hypotheses in terms of medians, sometimes ignoring the same-shape condition. We recommend that you express the hypotheses in words rather than symbols. “Yields are systematically higher in weed-free plots” is easy to understand and is a good statement of the effect that the Wilcoxon test looks for.



APPLY YOUR KNOWLEDGE

- 26.9 Daily activity and obesity: hypotheses.** We could use either two-sample t or the Wilcoxon rank sum to test the null hypothesis that lean and mildly obese people don’t differ in the time they spend standing and walking against the alternative hypothesis that lean people generally spend more time in these activities. Explain carefully what H_0 and H_a are for t and for W .

26.10 Strength of durable press fabrics: hypotheses. We are interested in whether fabrics treated with the Permafresh and Hylite processes have the same breaking strength “on the average.”

- State null and alternative hypotheses in terms of population means. What test would we typically use for these hypotheses? What conditions does this test require?
- State null and alternative hypotheses in terms of population medians. What test would we typically use for these hypotheses? What conditions does this test require?

DEALING WITH TIES IN RANK TESTS

average ranks

We have chosen our examples and exercises to this point rather carefully: they all involve data in which *no two values are the same*. This allowed us to rank all the values. In practice, however, we often find observations tied at the same value. What shall we do? The usual practice is to *assign all tied values the average of the ranks they occupy*. Here is an example with 6 observations:

Observation	153	155	158	158	161	164
Rank	1	2	3.5	3.5	5	6

The tied observations occupy the third and fourth places in the ordered list, so they share rank 3.5.

The exact distribution we have been using for the Wilcoxon rank sum W applies only to data without ties. Moreover, the standard deviation σ_W must be adjusted if ties are present. The Normal approximation can be used after the standard deviation is adjusted. Most statistical software will detect ties and make the necessary adjustment when using the Normal approximation. Although there is



an exact distribution when the data contain tied values, it is more complex and requires specialized software to compute. In practice, be careful using ranks tests for very small sample sizes when ties are present.

Some data have many ties because the scale of measurement has only a few values. Rank tests are often used for such data. Here is an example.



FAIRSAFETY



EXAMPLE 26.5 Food safety at fairs

STATE: Food sold at outdoor fairs and festivals may be less safe than food sold in restaurants because it is prepared in temporary locations and often by volunteer help. What do people who attend fairs think about the safety of the food served? One study asked this question of people at a number of fairs in the Midwest: “How often do you

think people become sick because of food they consume prepared at outdoor fairs and festivals?" The possible responses were

- 1 = very rarely
- 2 = once in a while
- 3 = often
- 4 = more often than not
- 5 = always

In all, 303 people answered the question. Of these, 196 were women and 107 were men.⁴ We suspect that women are more concerned than men about food safety. Is there good evidence for this conclusion?

PLAN: Do data analysis to understand the difference between women and men. Check the conditions required by the Wilcoxon test. If the conditions are met, use the Wilcoxon test for the hypotheses

H_0 : men and women do not differ in their responses

H_a : women give systematically higher responses than men

SOLVE: The responses for the 303 subjects appear in the data file. We can summarize them in a two-way table of counts:

	Response					Total
	1	2	3	4	5	
Female	13	108	50	23	2	196
Male	22	57	22	5	1	107
Total	35	165	72	28	3	303

Comparing row percents shows that the women in the sample do tend to give higher responses (showing more concern):

	Response					Total
	1	2	3	4	5	
Percent of females	6.6	55.1	25.5	11.7	1.0	100
Percent of males	20.6	53.3	20.6	4.7	1.0	100

Are these differences between women and men statistically significant?

The most important condition for inference is that the subjects are a *random sample* of people who attend fairs, at least in the Midwest. The researcher visited 11 different fairs. She stood near the entrance and stopped every 25th adult who passed. Because no personal choice was involved in choosing the subjects, we can reasonably treat the data as coming from a random sample. (As usual, there was some nonresponse, which



Danny Lehman/CORBIS

could create bias.) The Wilcoxon test also requires that responses have *continuous distributions*. We think that the subjects really have a continuous distribution of opinions about how often people become sick from food at fairs. The questionnaire asks them to round off their opinions to the nearest value in the five-point scale. So we are willing to use the Wilcoxon test.

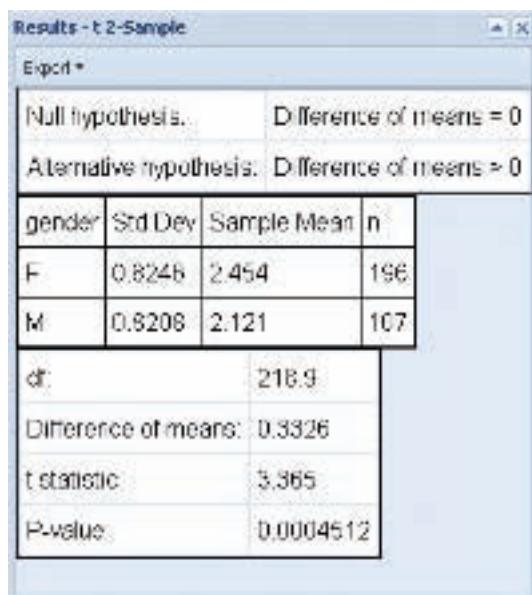
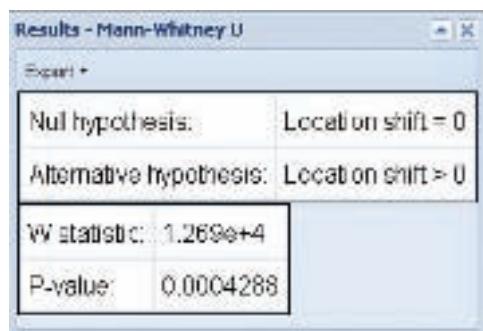
Because the responses can take only five values, there are many ties. All 35 people who chose “very rarely” are tied at 1, and all 165 who chose “once in a while” are tied at 2. Figure 26.4 gives output from CrunchIt! The Wilcoxon (reported as Mann-Whitney) test for the one-sided alternative that women are more concerned about food safety at fairs is highly significant ($P = 0.000429$).

With more than 100 observations in each group and no outliers, we might use the two-sample t test even though responses take only five values. Figure 26.4 shows that $t = 3.365$ with $P = 0.000451$. The one-sided P -value for the two-sample t test is essentially the same as that for the Wilcoxon test.

CONCLUDE: There is very strong evidence ($P = 0.0004$ for the Wilcoxon test) that women are more concerned than men about the safety of food served at fairs. ■

FIGURE 26.4

Output from CrunchIt! for the data of Example 26.5. The Wilcoxon rank sum test and the two-sample t test give similar results.



As is often the case, t and W for the data in Example 26.5 agree closely. There is, however, another reason to prefer the rank test in this example. The t statistic treats the response values 1 through 5 as meaningful numbers. In particular, the possible responses are treated as though they are equally spaced. The difference between “very rarely” and “once in a while” is the same as the difference between “once in a while” and “often.” This may not make sense. The rank test, on the other hand, uses only the order of the responses, not their actual values. The responses are arranged in order from least to most concerned about safety, so the rank test makes sense. Some statisticians avoid using t procedures when there is no fully meaningful scale of measurement.



Because we have a two-way table, we might have applied the chi-square test (Chapter 23), which asks if there is a significant relationship of *any kind* between gender and response. The chi-square test ignores the ordering of the responses and so doesn't tell us whether women are *more* concerned than men about the safety of the food served. This question depends on the ordering of responses from least concerned to most concerned.

APPLY YOUR KNOWLEDGE

Software is required to adequately carry out the Wilcoxon rank sum test in the presence of ties. All the following exercises concern data with ties.

26.11 Perception of life expectancy. Exercise 19.8 (text page 476) compares the perceived life expectancies of men and women. A researcher asked a sample of men and women to indicate their life expectancy. This was compared with values from actuarial tables, and the relative percent difference was computed (perceived life expectancy minus life expectancy from actuarial tables was divided by life expectancy from actuarial tables and converted to a percent). Here are the relative percent differences for all men and women over the age of 70 in the sample:

Men	-28	-23	-20	-19	-14	-13
Women	-20	-19	-15	-12	-10	-8

- (a) What are the null and alternative hypotheses for the Wilcoxon test? For the two-sample t test?
- (b) There are two pairs of tied observations. What ranks do you assign to each observation, using average ranks for ties?
- (c) Apply the Wilcoxon rank sum test to these data. Compare your result with the $P = 0.0528$ obtained from the two-sample t test in Figure 19.5.



26.12 Do birds learn to time their breeding? Exercises 19.43 to 19.45 (text page 489) concern a study of whether supplementing the diet of blue titmice with extra caterpillars will prevent them from adjusting their breeding date the following year to obtain a better food supply. Here are the data (days after the caterpillar peak):



Control	4.6	2.3	7.7	6.0	4.6	-1.2	
Supplemented	15.5	11.3	5.4	16.5	11.3	11.4	7.7

The null hypothesis is no difference in timing; the alternative hypothesis is that the supplemented birds miss the peak by more days because they don't adjust their breeding date.

- (a) There are three sets of ties, at 4.6, 7.7, and 11.3. Arrange the observations in order and assign average ranks to each tied observation.
- (b) Take W to be the rank sum for the supplemented group. What is the value of W ?
- (c) Use software: find the P -value of the Wilcoxon test and state your conclusion.

26.13 Tell me a story, continued. The data in Exercise 26.5 for a story told without pictures (Story 1) have tied observations. Is there good evidence that high-progress readers score higher than low-progress readers when they retell a story they have heard without pictures?

- (a) Make a back-to-back stemplot of the 5 responses in each group. Are any major deviations from Normality apparent?
- (b) Carry out a two-sample t test. State hypotheses and give the two sample means, the t statistic and its P -value, and your conclusion.
- (c) Carry out the Wilcoxon rank sum test. State hypotheses and give the rank sum W for high-progress readers, its P -value, and your conclusion. Do the t and Wilcoxon tests lead you to different conclusions?  STORY1

26.14 Do good smells bring good business? Exercise 19.9 (text page 478) describes an experiment that asked whether background aromas in a restaurant encourage customers to stay longer and spend more. The data on amount spent (in euros) are as follows:

No Odor										
15.9	18.5	15.9	18.5	18.5	21.9	15.9	15.9	15.9	15.9	15.9
15.9	18.5	18.5	18.5	20.5	18.5	18.5	15.9	15.9	15.9	15.9
18.5	18.5	15.9	18.5	15.9	18.5	15.9	25.5	12.9	15.9	15.9
Lavender Odor										
21.9	18.5	22.3	21.9	18.5	24.9	18.5	22.5	21.5	21.9	
21.5	18.5	25.5	18.5	18.5	21.9	18.5	18.5	24.9	21.9	
25.9	21.9	18.5	18.5	22.8	18.5	21.9	20.7	21.9	22.5	

Examine the data and comment on departures from Normality. Is there significant evidence that the lavender odor encourages customers to spend more? Follow the four-step process.  GOODSMELLS

26.15 Cicadas as fertilizer? Exercise 7.46 (text page 193) gives data from an experiment in which some bellflower plants in a forest were “fertilized” with dead cicadas and other plants were not disturbed. The data record the mass of seeds produced by



39 cicada plants and 33 undisturbed (control) plants. Do the data show that dead cicadas increase seed mass? Do data analysis to compare the two groups, explain why you would be reluctant to use the two-sample t test, and apply the Wilcoxon test. Follow the four-step process.



CICADA

26.16 Food safety in restaurants. Example 26.5 describes a study of the attitudes of people attending outdoor fairs about the safety of the food served at such locations. The full data set contains the responses of 303 people to several questions. The variables in this data set are (in order)



FOODSAFETY

subject hfair sfair sfast srest gender

The variable “sfair” contains the responses described in the example concerning safety of food served at outdoor fairs and festivals. The variable “srest” contains responses to the same question asked about food served in restaurants. The variable “gender” contains F if the respondent is a woman, M if he is a man. We saw that women are more concerned than men about the safety of food served at fairs. Is this also true for restaurants? Follow the four-step process in your answer.

26.17 More on food safety. The data file used in Exercise 26.16 contains 303 rows, one for each of the 303 respondents. Each row contains the responses of one person to several questions. We wonder if people are more concerned about safety of food served at fairs than they are about the safety of food served at restaurants. Explain carefully why we *cannot* answer this question by applying the Wilcoxon rank sum test to the variables “sfair” and “srest.”



MATCHED PAIRS: THE WILCOXON SIGNED RANK TEST

We use the one-sample t procedures (Chapter 18) for inference about the mean of one population or for inference about the mean difference in a matched pairs setting. The matched pairs setting is more important because good studies are generally comparative. We will now meet a rank test for this setting.

EXAMPLE 26.6 Tell me a story

STATE: A study of early childhood education asked kindergarten students to tell fairy tales that had been read to them earlier in the week. Each child told two stories. The first had been read to them and the second had been read but also illustrated with pictures. An expert listened to a recording of the children and assigned a score for certain uses of language. Here are the data for five low-progress readers in a pilot study:



STORIES

	Child				
	1	2	3	4	5
Story 2	0.77	0.49	0.66	0.28	0.38
Story 1	0.40	0.72	0.00	0.36	0.55
Difference	0.37	-0.23	0.66	-0.08	-0.17

We wonder if illustrations improve how the children retell a story.

PLAN: We would like to test the hypotheses

H_0 : scores have the same distribution for both stories

H_a : scores are systematically higher for Story 2

SOLVE (first steps): Because this is a matched pairs design, we base our inference on the differences. The matched pairs t test gives $t = 0.635$ with one-sided P -value $P = 0.280$. We cannot assess Normality from so few observations. We would therefore like to use a rank test. ■

absolute value

Positive differences in Example 26.6 indicate that the child performed better telling Story 2. If scores are generally higher with illustrations, the positive differences should be farther from zero in the positive direction than the negative differences are in the negative direction. We therefore compare the **absolute values** of the differences, that is, their magnitudes without a sign. Here they are, with boldface indicating the positive values:

0.37 0.23 **0.66** 0.08 0.17

Arrange these in increasing order and assign ranks, keeping track of which values were originally positive. Tied values receive the average of their ranks. If there are zero differences, discard them before ranking.

Absolute value	0.08	0.17	0.23	0.37	0.66
Rank	1	2	3	4	5

The test statistic is the sum of the ranks of the positive differences. (We could equally well use the sum of the ranks of the negative differences.) This is the **Wilcoxon signed rank statistic**. Its value here is $W^+ = 9$.

THE WILCOXON SIGNED RANK TEST FOR MATCHED PAIRS

Draw an SRS of size n from a population for a matched pairs study and take the differences in responses within pairs. Rank the absolute values of these differences. The sum W^+ of the ranks for the positive differences is the **Wilcoxon signed rank statistic**. If the distribution of the responses is not affected by the different treatments within pairs, then W^+ has mean

$$\mu_{W^+} = \frac{n(n + 1)}{4}$$

and standard deviation

$$\sigma_{W^+} = \sqrt{\frac{n(n + 1)(2n + 1)}{24}}$$

The **Wilcoxon signed rank test** rejects the hypothesis that there are no systematic differences within pairs when the rank sum W^+ is far from its mean.

EXAMPLE 26.7 Tell me a story, continued

SOLVE: In the storytelling study of Example 26.6, $n = 5$. If the null hypothesis (no systematic effect of illustrations) is true, the mean of the signed rank statistic is

$$\mu_{W^+} = \frac{n(n + 1)}{4} = \frac{(5)(6)}{4} = 7.5$$



The standard deviation of W^+ under the null hypothesis is

$$\begin{aligned}\sigma_{W^+} &= \sqrt{\frac{n(n + 1)(2n + 1)}{24}} \\ &= \sqrt{\frac{(5)(6)(11)}{24}} \\ &= \sqrt{13.75} = 3.708\end{aligned}$$

The observed value $W^+ = 9$ is only slightly larger than the mean. We now expect that the data are not statistically significant.

The P -value for our one-sided alternative is $P(W^+ \geq 9)$, calculated using the distribution of W^+ when the null hypothesis is true. Software gives the P -value $P = 0.4063$.

CONCLUDE: The data give no evidence ($P = 0.4$) that scores are higher for Story 2. The data do show an effect, but it fails to be significant because the sample is very small. ■

APPLY YOUR KNOWLEDGE

26.18 Growing trees faster. Exercise 18.39 (text page 459) describes an experiment in which extra carbon dioxide was piped to some plots in a pine forest. Each plot was paired with a nearby control plot left in its natural state. Do trees grow faster with extra carbon dioxide? Here are the average percent increases in base area for trees in the plots:

Pair	Control plot	Treated plot
1	9.752	10.587
2	7.263	9.244
3	5.742	8.675

The investigators used the matched pairs t test. With only 3 pairs, we can't verify Normality. We will try the Wilcoxon signed rank test.

- Find the differences within pairs, arrange them in order, and rank the absolute values. What is the signed rank statistic W^+ ?
- If the null hypothesis (no difference in growth) is true, what are the mean and standard deviation of W^+ ? Does comparing W^+ with this mean lead to a tentative conclusion?

26.19 Fighting cancer. Lymphocytes (white blood cells) play an important role in defending our bodies against tumors and infections. Can lymphocytes be genetically modified to recognize and destroy cancer cells? In one study of this idea, modified cells were infused into 11 patients with metastatic melanoma (serious skin cancer) that had not responded to existing treatments. Here are data for an “ELISA” test for the presence of cells that trigger an immune response, in counts per 100,000 cells before and after infusion.⁵ High counts suggest that infusion had a beneficial effect. 

Patient	1	2	3	4	5	6	7	8	9	10	11
Pre	14	0	1	0	0	0	0	20	1	6	0
Post	41	7	1	215	20	700	13	530	35	92	108

- (a) Examine the differences (post minus pre). Why can't we use the matched pairs t test to see if infusion raised the ELISA counts?
- (b) We will apply the Wilcoxon signed rank test. What are the ranks for the absolute values of the differences in counts? What is the value of W^+ ?
- (c) What would be the mean and standard deviation of W^+ if the null hypothesis (infusion makes no difference) were true? Compare W^+ with this mean (in standard deviation units) to reach a tentative conclusion about significance.

THE NORMAL APPROXIMATION FOR W^+

The distribution of the signed rank statistic when the null hypothesis (no difference) is true becomes approximately Normal as the sample size becomes large. We can then use Normal probability calculations (with the continuity correction) to obtain approximate P -values for W^+ . Let's see how this works in the storytelling example, even though $n = 5$ is certainly not a large sample.

EXAMPLE 26.8 Tell me a story: Normal approximation

For $n = 5$ observations, we saw in Example 26.7 that $\mu_{W^+} = 7.5$ and that $\mu_{W^+} = 3.708$. We observed $W^+ = 9$, so the one-sided P -value is $P(W^+ \geq 9)$. The continuity correction calculates this as $P(W^+ \geq 8.5)$, treating the value $W^+ = 9$ as occupying the interval from 8.5 to 9.5. We find the Normal approximation for the P -value either from software or by standardizing and using the standard Normal table:

$$\begin{aligned} P(W^+ \geq 8.5) &= P\left(\frac{W^+ - 7.5}{3.708} \geq \frac{8.5 - 7.5}{3.708}\right) \\ &= P(Z \geq 0.27) \\ &= 0.394 \blacksquare \end{aligned}$$

Figure 26.5 displays the output of two statistical programs. Minitab uses the Normal approximation and agrees with our calculation $P = 0.394$. We asked CrunchIt! to do two analyses: using the exact distribution of W^+ and using the matched pairs t test. The exact one-sided P -value for the Wilcoxon signed rank test is $P = 0.4063$, as we reported in Example 26.7. The Normal approximation is quite close to this. The t test result is a bit different, $P = 0.28$, but all three tests tell us that this very small sample gives no evidence that seeing illustrations improves the storytelling of low-progress readers.

Minitab

N	for	Wilcoxon	Estimated		
N	Test	Statistic	P	Median	
Diff	5	5	9.0	0.394	0.1000

CrunchIt!

Null hypothesis:	Median = 0
Alternative hypothesis:	Median > 0
W statistic:	9
P-value:	0.4063

Null hypothesis:	Mean difference = 0
Alternative hypothesis:	Mean difference > 0
Mean difference:	0.1100
df:	4
t statistic:	0.6350
P-value:	0.2800

FIGURE 26.5

Output from Minitab and CrunchIt! for the storytelling data of Example 26.6. The CrunchIt! output compares the Wilcoxon signed rank test (with the exact distribution) and the matched pairs t test.



APPLY YOUR KNOWLEDGE

26.20 Growing trees faster: Normal approximation. Continue your work from Exercise 26.18. Use the Normal approximation with continuity correction to find the P -value for the signed rank test against the one-sided alternative that trees grow faster with added carbon dioxide. What do you conclude? TREES

26.21 W^+ versus t . Find the one-sided P -value for the matched pairs t test applied to the tree growth data in Exercise 26.18. The smaller P -value of t relative to W^+ means that t gives stronger evidence of the effect of carbon dioxide on growth. The t test takes advantage of assuming that the data are Normal, a considerable advantage for these very small samples. TREES

26.22 Fighting cancer: Normal approximation. Use the Normal approximation with continuity correction to find the P -value for the test in Exercise 26.19. What do you conclude about the effect of infusing modified cells on the ELISA count? MORECANCER

26.23 Ancient air. Exercise 18.7 (text page 443) reports the following data on the percent of nitrogen in bubbles of ancient air trapped in amber: ANCIENTAIR

63.4 65.0 64.4 63.3 54.8 64.5 60.8 49.1 51.0

We wonder if ancient air differs significantly from the present atmosphere, which is 78.1% nitrogen.

- Graph the data, and comment on skewness and outliers. A rank test is appropriate.
- We would like to test hypotheses about the median percent of nitrogen in ancient air (the population):

$$H_0: \text{median} = 78.1$$

$$H_a: \text{median} \neq 78.1$$

To do this, apply the Wilcoxon signed rank statistic to the differences between the observations and 78.1. (This is the one-sample version of the test.) What do you conclude?



David Sanger Photography/Alamy

DEALING WITH TIES IN THE SIGNED RANK TEST

Ties among the absolute differences are handled by assigning average ranks. A tie *within* a pair creates a difference of zero. Because these are neither positive nor negative, we drop such pairs from our sample. Ties within pairs simply reduce the number of observations, but ties among the absolute differences complicate finding a P -value. Special software is required to use the exact distribution for the signed rank statistic W^+ , and the standard deviation σ_{W^+} must be adjusted for the ties before we can use the Normal approximation. Software will do this. Here is an example.

EXAMPLE 26.9 Golf scores

STATE: Here are the golf scores of 12 members of a college women's golf team in two rounds of tournament play. (A golf score is the number of strokes required to complete the course, so that low scores are better.)

	Player											
	1	2	3	4	5	6	7	8	9	10	11	12
Round 2	94	85	89	89	81	76	107	89	87	91	88	80
Round 1	89	90	87	95	86	81	102	105	83	88	91	79
Difference	5	-5	2	-6	-5	-5	5	-16	4	3	-3	1



Negative differences indicate better (lower) scores on the second round. Based on this sample, can we conclude that this team's golfers perform differently in the two rounds of a tournament?

PLAN: We would like to test the hypotheses that in a tournament play

$$H_0: \text{scores have the same distribution in Rounds 1 and 2}$$

$$H_a: \text{scores are systematically lower or higher in Round 2}$$

SOLVE: A stemplot of the differences (Figure 26.6) shows some irregularity and a low outlier. We will use the Wilcoxon signed rank test.

Figure 26.7 displays CrunchIt! output for the golf score data. The Wilcoxon statistic is $W^+ = 50.5$ with two-sided P -value $P = 0.3843$. The figure also includes the matched pairs t test, for which $P = 0.3716$. The two P -values are once again similar.

CONCLUDE: These data give no evidence for a systematic change in scores between rounds. ■

-1	6
-1	
-0	5 5 5 6
-0	3
0	1 2 3 4
0	5 5

FIGURE 26.6

Stemplot (with split stems) of the differences in scores for two rounds of a golf tournament, for Example 26.9.

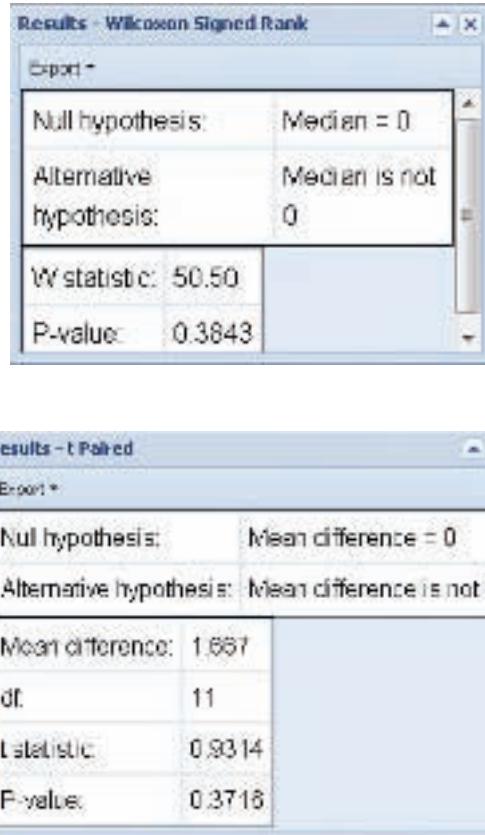
Let's see where the value $W^+ = 50.5$ came from. The absolute values of the differences, with boldface indicating those that were negative, are

5 5 2 6 5 5 5 16 4 3 3 1

Arrange these in increasing order and assign ranks, keeping track of which values were originally negative. Tied values receive the average of their ranks.

Absolute value	1	2	3	3	4	5	5	5	5	6	16
Rank	1	2	3.5	3.5	5	8	8	8	8	11	12

The Wilcoxon signed rank statistic is the sum $W^+ = 50.5$ of the ranks of the negative differences. (We could equally well use the sum of the ranks of the positive differences.)

**FIGURE 26.7**

Output from CrunchIt! for the golf scores data of Example 26.9. Because there are ties, a Normal approximation is used for the Wilcoxon signed rank test.

APPLY YOUR KNOWLEDGE

26.24 Does nature heal best? Exercise 18.33 (text page 458) gives these data on the healing rate (micrometers per hour) for cuts in the hind limbs of 12 newts:

Newt	1	2	3	4	5	6	7	8	9	10	11	12
Control limb	36	41	39	42	44	39	39	56	33	20	49	30
Experimental limb	28	31	27	33	33	38	45	25	28	33	47	23

The electrical field in the experimental limbs was reduced to zero by applying a voltage. The control limbs were not treated, so that they had their natural electrical field. The paired differences include an outlier, so we may choose to use the Wilcoxon signed rank test.

- (a) Find the ranks and give the value of the test statistic W^+ .
 (b) Use software to find the P -value. Give a conclusion. Be sure to include a description of what the data show in addition to the test results. 

26.25 Sweetening colas. Cola makers test new recipes for loss of sweetness during storage. Trained tasters rate the sweetness before and after storage. Here are the sweetness losses (sweetness before storage minus sweetness after storage) found by 10 tasters for one new cola recipe:

2.0 0.4 0.7 2.0 -0.4 2.2 -1.3 1.2 1.1 2.3

Are these data good evidence that the cola lost sweetness?

- (a) These data are the differences from a matched pairs design. State hypotheses in terms of the median difference in the population of all tasters, carry out a test, and give your conclusion.
 (b) The one-sample matched pairs t test had a P -value $P = 0.0123$ for these data. How does this compare with your result from (a)? What are the hypotheses for the t test? What conditions must be met for each of the t and Wilcoxon tests? 

26.26 Fungus in the air. The air in poultry-processing plants often contains fungus spores. Inadequate ventilation can damage the health of the workers. The problem is most serious during the summer. To measure the presence of spores, air samples are pumped to an agar plate, and “colony-forming units (CFUs)” are counted after an incubation period. Here are data from two locations in a plant that processes 37,000 turkeys per day, taken on four days in the summer. The units are CFUs per cubic meter of air.⁶



	Day			
	1	2	3	4
Kill room	3175	2526	1763	1090
Processing	529	141	362	224

Spore counts are clearly much higher in the kill room, but with only 4 pairs of observations, the difference may not be statistically significant. Apply a rank test. 

COMPARING SEVERAL SAMPLES: THE KRUSKAL-WALLIS TEST

We have now considered alternatives to the paired-sample and two-sample t tests for comparing the magnitude of responses to two treatments. To compare mean responses for more than two treatments, we use one-way analysis of variance (ANOVA) if the distributions of the responses to each treatment are at least roughly Normal and have similar spreads. What can we do when these distribution requirements are violated?



EXAMPLE 26.10 Weeds among the corn

STATE: Lamb's-quarter is a common weed that interferes with the growth of corn. A researcher planted corn at the same rate in 16 small plots of ground, then randomly assigned the plots to four groups. He weeded the plots by hand to allow a fixed number of lamb's-quarter plants to grow in each meter of corn row. These numbers were 0, 1, 3, and 9 in the four groups of plots. No other weeds were allowed to grow, and all plots received identical treatment except for the weeds. Here are the yields of corn (bushels per acre) in each of the plots:⁷

Weeds per meter	Corn yield						
0	166.7	1	166.2	3	158.6	9	162.8
0	172.2	1	157.3	3	176.4	9	142.4
0	165.0	1	166.7	3	153.1	9	162.7
0	176.9	1	161.1	3	156.0	9	162.4

Do yields change as the presence of weeds changes?

PLAN: Do data analysis to see how the yields change. Test the null hypothesis “no difference in the distribution of yields” against the alternative that the groups do differ.

SOLVE (first steps): The summary statistics are

Weeds	n	Median	Mean	Std. dev.
0	4	169.45	170.200	5.422
1	4	163.65	162.825	4.469
3	4	157.30	161.025	10.493
9	4	162.55	157.575	10.118

The mean yields do go down as more weeds are added. ANOVA tests whether the differences are statistically significant. Can we safely use ANOVA? Outliers are present in the yields for 3 and 9 weeds per meter. The outliers explain the differences between the means and the medians. They are the correct yields for their plots, so we cannot remove them. Moreover, the sample standard deviations do not quite satisfy our rule of thumb for ANOVA that the largest should not exceed twice the smallest. We may prefer to use a nonparametric test. ■

HYPOTHESES AND CONDITIONS FOR THE KRUSKAL-WALLIS TEST

The ANOVA F test concerns the means of the several populations represented by our samples. For Example 26.10, the ANOVA hypotheses are

$$H_0: \mu_0 = \mu_1 = \mu_3 = \mu_9$$

$$H_a: \text{not all four means are equal}$$

For example, μ_0 is the mean yield in the population of all corn planted under the conditions of the experiment with no weeds present. The data should consist of four independent random samples from the four populations, all Normally distributed with the same standard deviation.

The *Kruskal-Wallis test* is a rank test that can replace the ANOVA F test. The condition about data production (independent random samples from each population) remains important, but we can relax the Normality condition. We assume only that the response has a continuous distribution in each population. The hypotheses tested in our example are

$$H_0: \text{yields have the same distribution in all groups}$$

$$H_a: \text{yields are systematically higher in some groups than in others}$$

If all the population distributions have the same shape (Normal or not), these hypotheses take a simpler form. The null hypothesis is that all four populations have the same *median* yield. The alternative hypothesis is that not all four median yields are equal. The different standard deviations suggest that the four distributions in Example 26.10 do not all have the same shape.

THE KRUSKAL-WALLIS TEST STATISTIC

Recall the analysis of variance idea: we write the total observed variation in the responses as the sum of two parts, one measuring variation among the groups (sum of squares for groups, SSG) and one measuring variation among individual observations within the same group (sum of squares for error, SSE). The ANOVA F test rejects the null hypothesis that the mean responses are equal in all groups if SSG is large relative to SSE.

The idea of the Kruskal-Wallis rank test is to rank all the responses from all groups together and then apply one-way ANOVA to the ranks rather than to the original observations. If there are N observations in all, the ranks are always the whole numbers from 1 to N. The total sum of squares for the ranks is therefore a fixed number no matter what the data are. So we do not need to look at both SSG and SSE. Although it isn't obvious without some unpleasant algebra, the Kruskal-Wallis test statistic is essentially just SSG for the ranks. We give the formula, but you should rely on software to do the arithmetic. When SSG is large, that is evidence that the groups differ.

THE KRUSKAL-WALLIS TEST

Draw independent SRSs of sizes n_1, n_2, \dots, n_I from I populations. There are N observations in all. Rank all N observations and let R_i be the sum of the ranks for the i th sample. The **Kruskal-Wallis statistic** is

$$H = \frac{12}{N(N+1)} \sum \frac{R_i^2}{n_i} - 3(N+1)$$

When the sample sizes n_i are large and all I populations have the same continuous distribution, H has approximately the chi-square distribution with $I - 1$ degrees of freedom.

The **Kruskal-Wallis test** rejects the null hypothesis that all populations have the same distribution when H is large.

We now see that, like the Wilcoxon rank sum statistic, the Kruskal-Wallis statistic is based on the sums of the ranks for the groups we are comparing. The more different these sums are, the stronger is the evidence that responses are systematically larger in some groups than in others.

The exact distribution of the Kruskal-Wallis statistic H under the null hypothesis depends on all the sample sizes n_1 to n_I , so tables are awkward. The calculation of the exact distribution is so time-consuming for all but the smallest problems that even most statistical software uses the chi-square approximation to obtain P -values. As usual, the exact distribution when there are ties among the responses requires special software. We again assign average ranks to tied observations.

EXAMPLE 26.11 Weeds among the corn, continued



SOLVE (inference): In Example 26.10, there are $I = 4$ populations and $N = 16$ observations. The sample sizes are equal, $n_i = 4$. The 16 observations arranged in increasing order, with their ranks, are

Yield	142.4	153.1	156.0	157.3	158.6	161.1	162.4	162.7
Rank	1	2	3	4	5	6	7	8
Yield	162.8	165.0	166.2	166.7	166.7	172.2	176.4	176.9
Rank	9	10	11	12.5	12.5	14	15	16

There is one pair of tied observations. The ranks for each of the four treatments are

Weeds	Ranks					Sum of ranks
0	10	12.5	14	16		52.5
1	4	6	11	12.5		33.5
3	2	3	5	15		25.0
9	1	7	8	9		25.0

The Kruskal-Wallis statistic is therefore

$$\begin{aligned}
 H &= \frac{12}{N(N+1)} \sum \frac{R_i^2}{n_i} - 3(N+1) \\
 &= \frac{12}{(16)(17)} \left(\frac{52.5^2}{4} + \frac{33.5^2}{4} + \frac{25^2}{4} + \frac{25^2}{4} \right) - (3)(17) \\
 &= \frac{12}{272} (1282.125) - 51 \\
 &= 5.56
 \end{aligned}$$

Referring to the table of chi-square critical points (Table D) with $df = 3$, we see that the P -value lies in the interval $0.10 < P < 0.15$.

CONCLUDE: Although this small experiment suggests that more weeds decrease yield, it does not provide convincing evidence that weeds have an effect. ■

Figure 26.8 displays the Minitab output for both ANOVA and the Kruskal-Wallis test, and the CrunchIt! output for the Kruskal-Wallis test. Minitab agrees that $H = 5.56$ and gives $P = 0.135$. Minitab also gives the results of an adjustment that makes the chi-square approximation more accurate when there are ties. CrunchIt! automatically makes a correction for ties as well. For these data, the adjustment has no practical effect. It would be important if there were many ties. A very lengthy computer calculation shows that the exact P -value is $P = 0.1299$. The chi-square approximation is quite accurate.

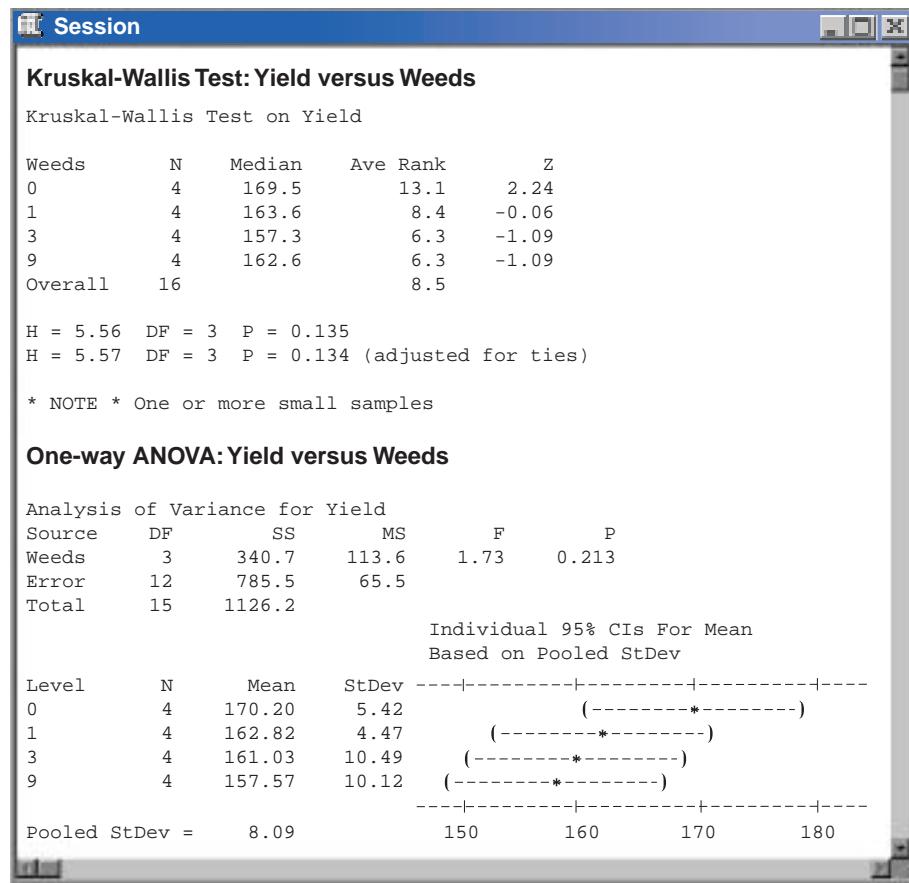
The ANOVA F test gives $F = 1.73$ with $P = 0.213$. Although the practical conclusion is the same, ANOVA and Kruskal-Wallis do not agree closely in this example. The rank test is more reliable for these small samples with outliers.

APPLY YOUR KNOWLEDGE

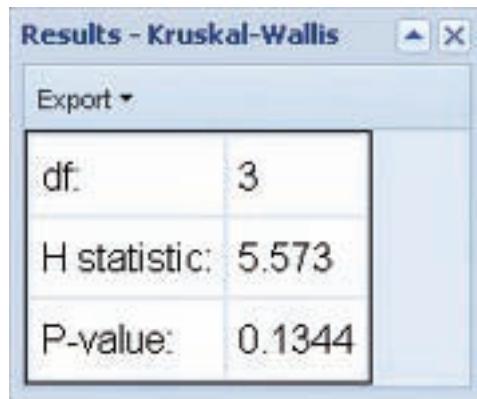
26.27 More rain for California? Exercise 25.31 (text page 649) describes an experiment that examines the effect on plant biomass in plots of California grassland randomly assigned to receive added water in the winter, added water in the spring, or no added water. The experiment continued for several years. Here are data for 2004 (mass in grams per square meter):

Winter	Spring	Control
254.6453	517.6650	178.9988
233.8155	342.2825	205.5165
253.4506	270.5785	242.6795
228.5882	212.5324	231.7639
158.6675	213.9879	134.9847
212.3232	240.1927	212.4862

Minitab



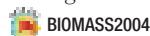
CrunchIt!

**FIGURE 26.8**

Output from Minitab and CrunchIt! for the corn yield data of Example 26.10. The Minitab output compares the Kruskal-Wallis test and one-way ANOVA, while the CrunchIt! output is only for the Kruskal-Wallis test.

The sample sizes are small and the data contain some possible outliers. We will apply a nonparametric test.

- Examine the data. Show that the conditions for ANOVA (text page 634) are not met. What appear to be the effects of extra rain in winter or spring?
- What hypotheses does ANOVA test? What hypotheses does Kruskal-Wallis test?
- What are I , the n_i , and N ? Arrange the counts in order and assign ranks.
- Calculate the Kruskal-Wallis statistic H . How many degrees of freedom should you use for the chi-square approximation to its null distribution? Use the chi-square table to give an approximate P -value. What does the test lead you to conclude?

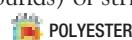


26.28 Logging in the rain forest: species richness. Table 25.2 (text page 530) contains data comparing the number of trees and number of tree species in plots of land in a tropical rain forest that had never been logged with similar plots nearby that had been logged 1 year earlier and 8 years earlier. The third response variable is species richness, the number of tree species divided by the number of trees. There are low outliers in the data, and a histogram of the ANOVA residuals shows outliers as well. Because of lack of Normality and small samples, we may prefer the Kruskal-Wallis test.



- Make a graph to compare the distributions of richness for the three groups of plots. Also give the median richness for the three groups.
- Use the Kruskal-Wallis test to compare the distributions of richness. State hypotheses, the test statistic and its P -value, and your conclusions.

26.29 Does polyester decay? Here are the breaking strengths (in pounds) of strips of polyester fabric buried in the ground for several lengths of time:⁸



2 weeks	118	126	126	120	129
4 weeks	130	120	114	126	128
8 weeks	122	136	128	146	140
16 weeks	124	98	110	140	110

Breaking strength is a good measure of the extent to which the fabric has decayed. Do a complete analysis that compares the four groups. Give the Kruskal-Wallis test along with a statement in words of the null and alternative hypotheses.

26.30 Good weather and tipping. Favorable weather has been shown to be associated with increased tipping. Exercise 25.35 (text page 650) describes a study to investigate whether just the belief that future weather will be favorable can lead to higher tips. The researchers gave 60 index cards to a waitress at an Italian restaurant in New Jersey. Before delivering the bill to each customer, the waitress randomly selected a card and wrote on the bill the same message that was printed on the index card. Twenty of the cards had the message “The weather



is supposed to be really good tomorrow. I hope you enjoy the day!" Another 20 cards contained the message "The weather is supposed to be not so good tomorrow. I hope you enjoy the day anyway!" The remaining 20 cards were blank, indicating that the waitress was not supposed to write any message. Choosing a card at random ensured that there was a random assignment of the diners to the three experimental conditions. Here are the tips as a percent of the total bill for the three messages:⁹

Good weather report	20.8 24.9	18.7 22.3	19.9 27.0	20.6 20.4	22.0 22.2	23.4 24.0	22.8 21.2	24.9 22.1	22.2 22.0	20.3 22.7
Bad weather report	18.0 17.0	19.0 13.6	19.2 17.5	18.8 19.9	18.4 20.2	19.0 18.8	18.5 18.0	16.1 23.2	16.8 18.2	14.0 19.4
No weather report	19.9 18.5	16.0 19.3	15.0 19.3	20.1 19.4	19.3 10.8	19.2 19.1	18.0 19.7	19.2 19.8	21.2 21.3	18.8 20.6

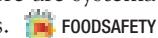
Do a complete analysis that includes a test of significance. Include a statement in words of your null and alternative hypotheses.



26.31 Food safety. Example 26.5 describes a study of the attitudes of people attending outdoor fairs about the safety of the food served at such locations. The data set contains the responses of 303 people to several questions. The variables in this data set are (in order)

subject hfair sfair sfast srest gender

The variable "sfair" contains responses to the safety question described in Example 26.5. The variables "srest" and "sfast" contain responses to the same question asked about food served in restaurants and in fast-food chains. Explain carefully why we *cannot* use the Kruskal-Wallis test to see if there are systematic differences in perceptions of food safety in these three locations.



CHAPTER 26 SUMMARY

CHAPTER SPECIFICS

- **Nonparametric tests** do not require any specific form for the distributions of the populations from which our samples come.
- **Rank tests** are nonparametric tests based on the **ranks** of observations, their positions in the list ordered from smallest (rank 1) to largest. Tied observations receive the average of their ranks. Use rank tests when the data come from random samples or randomized comparative experiments and the populations have continuous distributions.

- The **Wilcoxon rank sum test** compares two distributions to assess whether one has systematically larger values than the other. The Wilcoxon test is based on the **Wilcoxon rank sum statistic W** , which is the sum of the ranks of one of the samples. The Wilcoxon test can replace the **two-sample t test**. Software may perform the **Mann-Whitney test**, another form of the Wilcoxon test.
- **P-values** for the Wilcoxon test are based on the sampling distribution of the rank sum statistic W when the null hypothesis (no difference in distributions) is true. You can find P -values from special tables, software, or a Normal approximation (with continuity correction).
- The **Wilcoxon signed rank test** applies to matched pairs studies. It tests the null hypothesis that there is no systematic difference within pairs against alternatives that assert a systematic difference (either one-sided or two-sided).
- The test is based on the **Wilcoxon signed rank statistic W^+** , which is the sum of the ranks of the positive (or negative) differences when we rank the absolute values of the differences. The **matched pairs t test** is an alternative test in this setting.
- **P-values** for the signed rank test are based on the sampling distribution of W^+ when the null hypothesis is true. You can find P -values from special tables, software, or a Normal approximation (with continuity correction).
- The **Kruskal-Wallis test** compares several populations on the basis of independent random samples from each population. This is the **one-way analysis of variance** setting.
- The null hypothesis for the Kruskal-Wallis test is that the distribution of the response variable is the same in all the populations. The alternative hypothesis is that responses are systematically larger in some populations than in others.
- The **Kruskal-Wallis statistic H** can be viewed in two ways. It is essentially the result of applying one-way ANOVA to the ranks of the observations. It is also a comparison of the sums of the ranks for the several samples.
- When the sample sizes are not too small and the null hypothesis is true, the Kruskal-Wallis test statistic for comparing I populations has approximately the chi-square distribution with $I - 1$ degrees of freedom. We use this approximate distribution to obtain P -values.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Ranks

1. Assign ranks to a moderate number of observations. Use average ranks if there are ties among the observations.
2. From the ranks, calculate the rank sums when the observations come from two or several samples.

B. Rank Test Statistics

1. Determine which of the rank sum tests is appropriate in a specific problem setting.
2. Calculate the Wilcoxon rank sum W from ranks for two samples, the Wilcoxon signed rank sum W^+ for matched pairs, and the Kruskal-Wallis statistic H for two or more samples.
3. State the hypotheses tested by each of these statistics in specific problem settings.
4. Determine when it is appropriate to state the hypotheses for W and H in terms of population medians.

C. Rank Tests

1. Use software to carry out any of the rank tests. Combine the test with data description and give a clear statement of findings in specific problem settings.
2. Use the Normal approximation with continuity correction to find approximate P -values for W and W^+ . Use a table of chi-square critical values to approximate the P -value for H .

LINK IT

The statistical methods in Chapters 18, 19, and 25 were developed for Normal distributions, but are fairly robust against a failure in this assumption. However, when the sample sizes are very small or there are outlying observations, it is important to have alternative techniques. The nonparametric methods described in this chapter provide such alternatives. Nonparametric methods often involve replacing the actual data by their ranks, which makes these methods fairly insensitive to outlying observations and also allows us to use them with minimal assumptions regarding the distributions of the data.

For the paired-data problem, the Wilcoxon signed rank test can be used as an alternative to the paired t test of Chapter 18, while for the two-sample problem, the rank sum test provides an alternative to the two-sample t test described in Chapter 19. The Kruskal-Wallis test can be used in place of the F test for one-way ANOVA. Although we have concentrated on hypothesis testing in this chapter, these nonparametric methods also provide confidence intervals, and in one-way ANOVA there are associated multiple comparisons for determining which of the treatments differ. As with the other supplemental chapters, this chapter is intended to provide only an introduction to a more advanced topic in statistical inference. More advanced courses in statistics will provide additional details for these methods.

CHECK YOUR SKILLS

26.32 A study of the effects of exercise used rats bred to have high or low capacity for exercise. There were 8 high-capacity and 8 low-capacity rats. To compare the blood pressures of the groups, you use the

- (a) Wilcoxon rank sum test.
- (b) Wilcoxon signed rank test.
- (c) Kruskal-Wallis test.

26.33 You interview college students who have done community service and another group of students who have not. To compare the scores of the two groups on a test of attitude toward people of other races, you use the

- (a) Wilcoxon rank sum test.
- (b) Wilcoxon signed rank test.
- (c) Kruskal-Wallis test.

26.34 You interview both the husband and the wife in 64 married couples and give each a test that measures marital satisfaction. To assess whether there is a difference in level of marital satisfaction between husbands and wives, you use the

- (a) Wilcoxon rank sum test.
- (b) Wilcoxon signed rank test.
- (c) Kruskal-Wallis test.

26.35 When some plants are attacked by leaf-eating insects, they release chemical compounds that repel the insects. Here are data on emissions of one compound by plants attacked by leaf bugs and by plants in an undamaged control group:

Control group	14.4	15.2	12.6	11.9	5.1	8.0
Attacked group	10.6	15.3	25.2	19.8	17.1	14.6

The rank sum W for the control group is

- (a) 21.
- (b) 26.
- (c) 52.

26.36 If there is no difference in emissions between the attacked group and the control group, the mean of W in the previous exercise is

- (a) 39.
- (b) 78.
- (c) 6.2.

26.37 Suppose that the 12 observations in Exercise 26.35 were

Control group	14.4	15.2	12.6	11.9	5.1	8.0
Attacked group	14.4	15.3	25.2	19.8	17.1	14.6

The rank sum for the control group is now

- (a) 21.
- (b) 23.
- (c) 25.

26.38 Interview 10 young married couples, wife and husband separately. One question asks how important the attractiveness of their spouse is to them on a scale of 1 to 10. Here are the responses:

	Couple									
	1	2	3	4	5	6	7	8	9	10
Husband	7	7	7	3	9	5	10	6	6	7
Wife	4	2	5	2	2	2	4	7	1	5

The Wilcoxon signed rank statistic W^+ (based on husband's score minus wife's score) is

- (a) 51.
- (b) 53.5.
- (c) 54.

26.39 If husbands and wives don't differ in how important the attractiveness of their spouse is, the mean of W^+ in the previous exercise is

- (a) 27.5.
- (b) 55.
- (c) 105.

26.40 Suppose that the responses in Exercise 26.38 are

	Couple									
	1	2	3	4	5	6	7	8	9	10
Husband	7	7	7	3	9	5	10	6	6	5
Wife	4	2	5	3	2	2	4	7	1	5

The Wilcoxon signed rank statistic W^+ (based on husband's score minus wife's score) is now

- (a) 35.
- (b) 36.
- (c) 52.

26.41 You compare the starting salaries of 7 graduates who majored in accounting, 9 who majored in finance, and 5 who majored in marketing. If the three starting-salary distributions are the same, the Kruskal-Wallis statistic H has approximately a chi-square distribution. The degrees of freedom are

- (a) 1.
- (b) 2.
- (c) 3.

CHAPTER 26 EXERCISES

One of the rank tests discussed in this chapter is appropriate for each of the following exercises. Follow the **Plan**, **Solve**, and **Conclude** parts of the four-step process in your answers.

26.42 Each day I am getting better in math. Table 19.3 (text page 491) gives the before and after scores for two groups of students taking a program to improve their basic mathematics skills. Did the treatment group show significantly greater improvement than the control group?  **SUBLIMINAL**

26.43 Which blue is most blue? The color of a fabric depends on the dye used and also on how the dye is applied. This matters to clothing manufacturers, who want the color of the fabric to be just right. Dye fabric made of ramie with the same “procion blue” die applied in four different ways. Then use a colorimeter to measure the lightness of the color on a scale in which black is 0 and white is 100. Here are the data for 8 pieces of fabric dyed in each way:¹⁰  **BLUEDYE**

Method A	41.72	41.83	42.05	41.44	41.27	42.27	41.12	41.49
Method B	40.98	40.88	41.30	41.28	41.66	41.50	41.39	41.27
Method C	42.30	42.20	42.65	42.43	42.50	42.28	43.13	42.45
Method D	41.68	41.65	42.30	42.04	42.25	41.99	41.72	41.97

Do the methods differ in color lightness?

26.44 The brain responds to sound. Table 18.2 (text page 451) contains data from a study comparing the brain's response to "pure tones" and recognizable sounds. Researchers anesthetized macaque monkeys and fed pure tones and also monkey calls directly to their brains by inserting electrodes. Response to the stimulus was measured by the firing rate (electrical spikes per second) of neurons in various areas of the brain. Researchers suspected that the response to monkey calls would be stronger than the response to a pure tone. Do the data support this idea?  BRAINRESPONSE

26.45 Adolescent obesity. Adolescent obesity is a serious health risk affecting more than 5 million young people in the United States alone. Laparoscopic adjustable gastric banding has the potential to provide a safe and effective treatment. Fifty adolescents between 14 and 18 years old with a body mass index higher than 35 were recruited from the Melbourne, Australia, community for the study.¹¹ Twenty-five were randomly selected to undergo gastric banding, and the remaining twenty-five were assigned to a supervised lifestyle intervention program involving diet, exercise, and behavior modification. All subjects were followed for two years. Here are the weight losses in kilograms for the subjects who completed the study. In the gastric banding group:  GASTRICBANDS

35.6	81.4	57.6	32.8	31.0	37.6	36.5	-5.4
27.9	49.0	64.8	39.0	43.0	33.9	29.7	20.2
15.2	41.7	53.4	13.4	24.8	19.4	32.3	22.0

In the lifestyle intervention group:

6.0	2.0	-3.0	20.6	11.6	15.5	-17.0	1.4	4.0
-4.6	15.8	34.6	6.0	-3.1	-4.3	-16.7	-1.8	-12.8

Does gastric banding result in significantly greater weight loss than a supervised lifestyle intervention program?

26.46 Food safety at fairs and restaurants. Example 26.5 describes a study of the attitudes of people attending outdoor fairs about the safety of the food served at such locations. The full data set contains the responses of 303

people to several questions. The variables in this data set are (in order)

subject hfair sfair sfast srest gender

The variable "sfair" contains responses to the safety question described in Example 26.5.

The variable "srest" contains responses to the same question asked about food served in restaurants. We suspect that restaurant food will appear safer than food served outdoors at a fair. Do the data give good evidence for this suspicion?  FOODSAFETY

26.47 Food safety at fairs and fast-food restaurants.

The food safety survey data described in Example 26.5 also contain the responses of the 303 subjects to the same question asked about food served at fast-food restaurants. These responses are the values of the variable "sfast." Is there a systematic difference between the level of concern about food safety at outdoor fairs and at fast-food restaurants?  FOODSAFETY

26.48 Nematodes and plant growth. A botanist prepares

16 identical planting pots and then introduces different numbers of nematodes (microscopic worms) into the pots. A tomato seedling is transplanted into each pot. Here are data on the increase in height of the seedlings (in centimeters) 16 days after planting:¹²  NEMATODES

Nematodes	Seedling growth			
	0	1,000	5,000	10,000
0	10.8	9.1	13.5	9.2
1,000	11.1	11.1	8.2	11.3
5,000	5.4	4.6	7.4	5.0
10,000	5.8	5.3	3.2	7.5

Do nematodes in soil affect plant growth?

26.49 Mutual fund performance. Mutual funds often

compare their performance with a benchmark provided by an "index" that describes the performance of the class of assets in which the funds invest. For example, the Vanguard International Growth Fund benchmarks its performance against the EAFE (Europe, Australasia, Far East) index. Table 18.4 (text page 461) gives the annual returns (percent) for the fund and the index. Does the fund's performance differ significantly from that of its benchmark?  MUTUALFUND

How does the meeting of large rivers influence the diversity of fish? A study of the Amazon and 13 of its major tributaries concentrated on electric fish, which are common in South America. The researchers trawled in more than 1000 locations in the Amazon above and below each tributary and in the lower part of the tributaries themselves. In all, they found 43 species of electric fish. These distinctive fish can “stand in” for fish in general, which are too numerous to count easily. The researchers concluded that the number of fish species increases when a tributary joins the Amazon, but that the effect is local: there is no steady increase in diversity as we move downstream. Table 26.1 gives the estimated number of electric fish species in the Amazon upstream and downstream from each tributary and in the tributaries themselves just before they flow into the Amazon.¹³ The researchers used nonparametric tests to assess the statistical significance of their results. Exercises 26.50 to 26.52 quote conclusions from the study.



26.50 Downstream versus upstream. “We identified a significant positive effect of tributaries on Amazon mainstem species richness in two respects. First, we found that sample stations downstream of each tributary contained more species than did their respective upstream stations.” Do a test to confirm the statistical significance of this effect and report your conclusion.

26.51 Tributary versus upstream. “Second, we found that species richness within tributaries exceeded that within

TABLE 26.1 Electric fish species in the Amazon

TRIBUTARY	SPECIES COUNTS		
	UPSTREAM	TRIBUTARY	DOWNTSTREAM
Içá	14	23	19
Jutaí	11	15	18
Juruá	8	13	8
Japurá	9	16	11
Coari	5	7	7
Purus	10	23	16
Manacapuru	5	8	6
Negro	23	26	24
Madeira	29	24	30
Trombetas	19	20	16
Tapajós	16	5	20
Xingu	25	24	21
Tocantins	10	12	12

their adjacent upstream mainstem stations.” Again, do a test to confirm significance and report your finding.

26.52 Tributary versus downstream. Species richness “was comparable between tributaries and their adjacent downstream mainstem stations.” Verify this conclusion by comparing tributary and downstream species counts.



EXPLORING THE WEB

26.53 Confidence in the banking system. The General Social Survey (GSS) is a sociological survey used to collect data on demographic characteristics and attitudes of residents of the United States. The survey is conducted by the National Opinion Research Center of the University of Chicago, which interviews face-to-face a randomly selected sample of adults (18 and older). SDA (Survey Documentation and Analysis) is a set of programs that allows you to analyze survey data and includes the GSS as part of its archive. In Exercises 25.43 and 25.44 you used data from the GSS to study the relationship between average respondent age and confidence in the banking system. A one-way ANOVA showed a statistically significant difference between the age of the respondent and his or her confidence in the banking system. Download the data file following the directions in Exercise 25.44, and do the Kruskal-Wallis test to see if the ages are systematically higher for some levels of confidence in the banking system than others. Are your results similar to the one-way ANOVA?

26.54 Who's more liberal? The American National Election Studies (ANES) is the leading academically run national survey of voters in the United States and is conducted before and after every presidential election. SDA (Survey Documentation and Analysis) is a set of programs that allows you to analyze survey data and includes the ANES survey as part of its archive. Go to the Web site sda.berkeley.edu/ and

click on Archive. Go to the 2008 ANES survey. If there is a more recent survey than 2008, you should use it.

- Open the preelection survey data. Under Liberal/Conservative, choose the variable “liberal/conservative self-placement on a 7 point scale.” Use this as your column variable. Under Nonsurvey data, choose “respondent’s gender.” Use this as your row variable. In the details for the table, set Weight to none, and for N of Cases to Display, make sure the unweighted box is checked. For Percentaging, choose row percents. Now click on “run the table.”
- You want to see if either gender rates themselves as more liberal by carrying out the Wilcoxon rank sum test using the same ideas as in Example 26.5. You first need to use the summary table generated in part (a) to re-create the original observations. There are two variables in the data set: gender and liberal/conservative placement. How many observations are there in total? How many observations are extremely liberal males? Each extremely liberal male corresponds to one observation in the data set with the gender variable taking the value male and the liberal/conservative placement variable taking the value 1. Once you have re-created the original observations, use software to perform the Wilcoxon rank sum test. Is your software adjusting for ties? What is your conclusion?
- Should the two-sample *t* test be used to answer this question? Explain.

NOTES AND DATA SOURCES

1. Data provided by Samuel Phillips, Purdue University.
2. Data provided by Susan Stadler, Purdue University.
3. The precise meaning of “yields are systematically larger in plots with no weeds” is that for every fixed value a , the probability that the yield with no weeds is larger than a is at least as great as the same probability for the yield with weeds.
4. Huey Chern Boo, “Consumers’ perceptions and concerns about safety and healthfulness of food served at fairs and festivals,” MS thesis, Purdue University, 1997.
5. Richard A. Morgan et al., “Cancer regression in patients after transfer of genetically engineered lymphocytes,” *Science*, 314 (2006), pp. 126–129. The data appear in the Online Supplementary Material.
6. Michael W. Peugh, “Field investigation of ventilation and air quality in duck and turkey slaughter plants,” MS thesis, Purdue University, 1996.
7. See Note 1.
8. Sapna Aneja, “Biodeterioration of textile fibers in soil,” MS thesis, Purdue University, 1994.
9. Bruce Rind and David Strohmetz, “Effects of beliefs about future weather conditions on restaurant tipping,” *Journal of Applied Social Psychology*, 31 (2001), pp. 2160–2164. We would like to thank the authors for supplying the original data.
10. Yvan R. Germain, “The dyeing of ramie with fiber reactive dyes using the cold pad-batch method,” MS thesis, Purdue University, 1988.

11. Paul E. O'Brien et al., "Laparoscopic adjustable gastric banding in severely obese adolescents," *Journal of the American Medical Association*, 303 (2010), pp. 519–526. We thank the authors for providing the data.
12. Data provided by Matthew Moore.
13. Cristina Cox Fernandes, Jeffrey Podos, and John G. Lundberg, "Amazonian ecology: tributaries enhance the diversity of electric fishes," *Science*, 305 (2004), pp. 1960–1962.



Statistical Process Control

Chapter 27

Organizations are (or ought to be) concerned about the quality of the products and services they offer. A key to maintaining and improving quality is systematic use of *data* in place of intuition or anecdotes. In the words of Stan Sigman, former CEO of Cingular Wireless, "What gets measured gets managed."¹

Because using data is a key to improving quality, statistical methods have much to contribute. Simple tools are often the most effective. A scatterplot and perhaps a regression line can show how the time to answer telephone calls to a corporate call center influences the percent of callers who hang up before their calls are answered. The design of a new product as simple as a multivitamin tablet may involve interviewing samples of consumers to learn what vitamins and minerals they want included and using randomized comparative experiments in designing the manufacturing process. An experiment might discover, for example, what combination of moisture level in the raw vitamin powder and pressure in the tablet-forming press produces the right tablet hardness.

Quality is a vague idea. You may feel that a restaurant serving filet mignon is a higher-quality establishment than a fast-food outlet that serves hamburgers. For statistical purposes we need a narrower concept: *consistently meeting standards appropriate for a specific product or service*. By this definition of quality, the expensive restaurant may serve low-quality filet mignon while the fast-food outlet serves high-quality hamburgers. The hamburgers are freshly grilled, are served at the right temperature, and are the same every time you visit. Statistically minded management can assess quality by sampling hamburgers and measuring the time from order to being served, the temperature of the burgers, and their tenderness.

IN THIS CHAPTER WE COVER...

- Processes
- Describing processes
- The idea of statistical process control
- \bar{x} charts for process monitoring
- s charts for process monitoring
- Using control charts
- Setting up control charts
- Comments on statistical control
- Don't confuse control with capability!
- Control charts for sample proportions
- Control limits for p charts

This chapter focuses on just one aspect of statistics for improving quality: *statistical process control*. The techniques are simple and are based on sampling distributions (Chapter 11), but the underlying ideas are important and a bit subtle.

PROCESSES

In thinking about statistical inference, we distinguish between the *sample* data we have in hand and the wider *population* that the data represent. We hope to use the sample to draw conclusions about the population. In thinking about quality improvement, it is often more natural to speak of *processes* rather than populations. This is because work is organized in processes. Some examples are

- processing an application for admission to a university and deciding whether or not to admit the student;
- reviewing an employee's expense report for a business trip and issuing a reimbursement check;
- hot forging to shape a billet of titanium into a blank that, after machining, will become part of a medical implant for hip, knee, or shoulder replacement.

Each of these processes is made up of several successive operations that eventually produce the output—an admission decision, reimbursement check, or metal component.

PROCESS

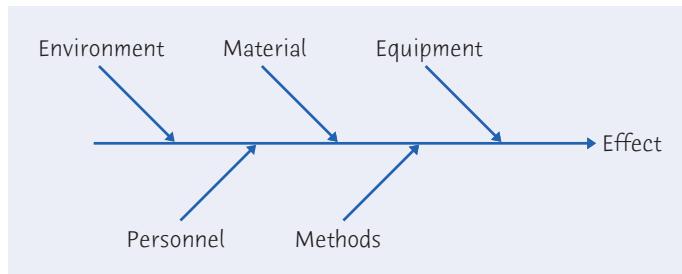
A process is a chain of activities that turns inputs into outputs.

We can accommodate processes in our sample-versus-population framework: think of the population as containing all the outputs that would be produced by the process if it ran forever in its present state. The outputs produced today or this week are a sample from this population. Because the population doesn't actually exist now, it is simpler to speak of a process and of recent output as a sample from the process in its present state.

DESCRIBING PROCESSES

The first step in improving a process is to understand it. Process understanding is often presented graphically using two simple tools: flowcharts and cause-and-effect diagrams. A **flowchart** is a picture of the stages of a process. A **cause-and-effect diagram** organizes the logical relationships between the inputs and stages of a process and an output. Sometimes the output is successful completion of the process task; sometimes it is a quality problem that we hope to solve. A good starting outline for a cause-and-effect diagram appears in Figure 27.1. The main

flowchart
cause-and-effect diagram

**FIGURE 27.1**

An outline for a cause-and-effect diagram. To complete the diagram, group causes under these main headings in the form of branches.

branches organize the causes and serve as a skeleton for detailed entries. You can see why these are sometimes called “fishbone diagrams.” An example will illustrate the use of these graphs.²

EXAMPLE 27.1 Hot forging

Hot forging involves heating metal to a plastic state and then shaping it by applying thousands of pounds of pressure to force the metal into a die (a kind of mold). Figure 27.2 is a flowchart of a typical hot-forging process.³

A process improvement team, after making and discussing this flowchart, came to several conclusions:

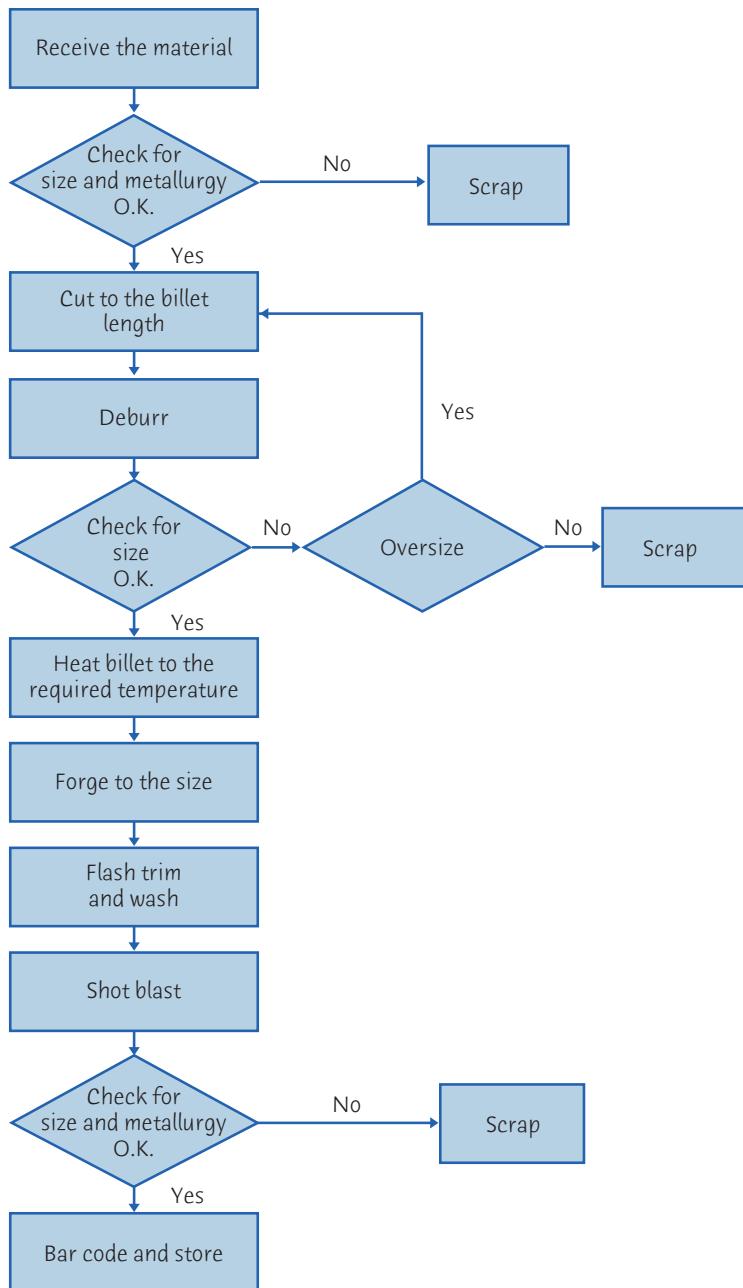
- Inspecting the billets of metal received from the supplier adds no value. We should insist that the supplier be responsible for the quality of the material. The supplier should put in place good statistical process control. We can then eliminate the inspection step.
- Can we buy the metal billets already cut to rough length and ground smooth by the supplier, thus eliminating the cost of preparing the raw material ourselves?
- Heating the metal billet and forging (pressing the hot metal into the die) are the heart of the process. We should concentrate our attention here.

The team then prepared a cause-and-effect diagram (Figure 27.3) for the heating and forging part of the process. The team members shared their specialist knowledge of the causes in their areas, resulting in a more complete picture than any one person could produce. Figure 27.3 is a simplified version of the actual diagram. We have given some added detail for the “Hammer stroke” branch under “Equipment” to illustrate the next level of branches. Even this requires some knowledge of hot forging to understand. Based on detailed discussion of the diagram, the team decided what variables to measure and at what stages of the process to measure them. Producing well-chosen data is the key to improving the process. ■

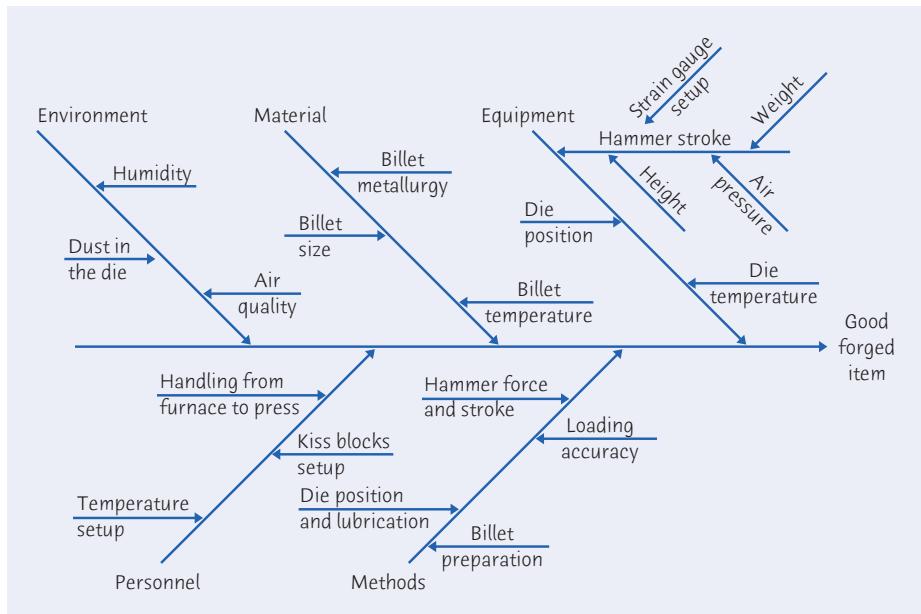
We will apply statistical methods to a series of measurements made on a process. Deciding what specific variables to measure is an important step in quality improvement. Often we use a “performance measure” that describes an output of a process. A company’s financial office might record the percent of errors that outside auditors find in expense account reports or the number of

FIGURE 27.2

Flowchart of the hot-forging process in Example 27.1. Use this as a model for flowcharts: decision points appear as diamonds, and other steps in the process appear as rectangles. Arrows represent flow from step to step.



data entry errors per week. The personnel department may measure the time to process employee insurance claims or the percent of job offers that are accepted. In the case of complex processes, it is wise to measure key steps within the process rather than just final outputs. The process team in Example 27.1 might recommend that the temperature of the die and of the billet be measured just before forging.

**FIGURE 27.3**

Simplified cause-and-effect diagram of the hot-forging process in Example 27.1. Good cause-and-effect diagrams require detailed knowledge of the specific process.

APPLY YOUR KNOWLEDGE

27.1 Describe a process. Choose a process that you know well. If you lack experience with actual business or manufacturing processes, choose a personal process such as ordering something over the Internet, paying a bill online, or recording a TV show on a DVR. Make a flowchart of the process. Make a cause-and-effect diagram that presents the factors that lead to successful completion of the process.

27.2 Describe a process. Each weekday morning, you must get to work or to your first class on time. Make a flowchart of your daily process for doing this, starting when you wake. Be sure to include the time at which you plan to start each step.

27.3 Process measurement. Based on your description of the process in Exercise 27.1, suggest specific variables that you might measure in order to

- assess the overall quality of the process.
- gather information on a key step within the process.

27.4 Pareto charts. Pareto charts are bar graphs with the bars ordered by height. They are often used to isolate the “vital few” categories on which we should focus our attention. Here is an example. A large medical center, financially pressed by restrictions on reimbursement by insurers and the government, looked at losses broken down by diagnosis. Government standards place cases into Diagnostic Related Groups (DRGs). For example, major joint replacements (mostly hip and knee) are DRG 209.⁴ Here is what the hospital found:

Pareto charts

DRG	Percent of losses
104	5.2
107	10.1
109	7.7
116	13.7
148	6.8
209	15.2
403	5.6
430	6.8
462	9.4

What percent of total losses do these 9 DRGs account for? Make a Pareto chart of losses by DRG. Which DRGs should the hospital study first when attempting to reduce its losses?  DRG

27.5 Pareto charts. Continue the study of the process of getting to work or class on time from Exercise 27.2. If you kept good records, you could make a Pareto chart of the reasons (special causes) for late arrivals at work or class. Make a Pareto chart that you think roughly describes your own reasons for lateness. That is, list the reasons from your experience and chart your estimates of the percent of late arrivals each reason explains.

27.6 Pareto charts. A large hospital was concerned about whether it was scheduling its operating rooms efficiently. Operating rooms lying idle may mean loss of potential revenue. Of particular interest was when and for how long the first operation of the day was performed. As a first step in understanding the use of its operating rooms, data were collected on what medical specialties were the first to use one of the rooms for an operation in the morning.⁵ Here is what the hospital found:

Specialty	Percent of all operations
Burns Center	3.7
ENT specialist	7.6
Gynecology	5.9
Ophthalmology	7.2
Orthopedics	12.3
Plastic surgery	21.1
Surgery	30.6
Urology	7.2

What percent of total operations do these 8 specialties account for? Make a Pareto chart of percent of all operations by specialty. Which specialties should the hospital study first when attempting to understand operating-room use?  OPERATIONS

THE IDEA OF STATISTICAL PROCESS CONTROL

The goal of statistical process control is to make a process stable over time and then keep it stable unless planned changes are made. You might want, for example, to keep your weight constant over time. A manufacturer of machine parts wants the critical dimensions to be the same for all parts. “Constant over time” and “the same for all” are not realistic requirements. They ignore the fact that *all processes have variation*. Your weight fluctuates from day to day; the critical dimension of a machined part varies a bit from item to item; the time to process a college admission application is not the same for all applications. Variation occurs in even the most precisely made product due to small changes in the raw material, the adjustment of the machine, the behavior of the operator, and even the temperature in the plant. Because variation is always present, we can’t expect to hold a variable exactly constant over time. The statistical description of stability over time requires that the *pattern of variation* remain stable, not that there be no variation in the variable measured.

STATISTICAL CONTROL

A variable that continues to be described by the same distribution when observed over time is said to be in statistical control, or simply **in control**.

Control charts are statistical tools that monitor a process and alert us when the process has been disturbed so that it is now **out of control**. This is a signal to find and correct the cause of the disturbance.

In the language of statistical quality control, a process that is in control has only **common cause** variation. Common cause variation is the inherent variability of the system, due to many small causes that are always present. When the normal functioning of the process is disturbed by some unpredictable event, **special cause** variation is added to the common cause variation. We hope to be able to discover what lies behind special cause variation and eliminate that cause to restore the stable functioning of the process.

common cause

special cause

EXAMPLE 27.2 Common cause, special cause

Imagine yourself doing the same task repeatedly, say folding an advertising flyer, stuffing it into an envelope, and sealing the envelope. The time to complete the task will vary a bit, and it is hard to point to any one reason for the variation. Your completion time shows only common cause variation.

Now the telephone rings. You answer, and though you continue folding and stuffing while talking, your completion time rises beyond the level expected from common causes alone. Answering the telephone adds special cause variation to the common cause variation that is always present. The process has been disturbed and is no longer in its normal and stable state.

If you are paying temporary employees to fold and stuff advertising flyers, you avoid this special cause by not having telephones present and by asking the employees to turn off their cell phones while they are working. ■

Control charts work by distinguishing the always-present common cause variation in a process from the additional variation that suggests that the process has been disturbed by a special cause. A control chart sounds an alarm when it sees too much variation. The most common application of control charts is to monitor the performance of industrial and business processes. The same methods, however, can be used to check the stability of quantities as varied as the ratings of a television show, the level of ozone in the atmosphere, and the gas mileage of your car. Control charts combine graphical and numerical descriptions of data with use of sampling distributions.



APPLY YOUR KNOWLEDGE

- 27.7 Special causes.** Tayler participates in 10-kilometer races. She regularly runs 15 kilometers over the same course in training. Her time varies a bit from day to day but is generally stable. Give several examples of special causes that might raise Tayler's time on a particular day.
- 27.8 Common causes, special causes.** In Exercise 27.1, you described a process that you know well. What are some sources of common cause variation in this process? What are some special causes that might at times drive the process out of control?
- 27.9 Common causes, special causes.** Each weekday morning, you must get to work or to your first class on time. The time at which you reach work or class varies from day to day, and your planning must allow for this variation. List several common causes of variation in your arrival time. Then list several special causes that might result in unusual variation leading to either early or (more likely) late arrival.

\bar{x} CHARTS FOR PROCESS MONITORING

chart setup
process monitoring

When you first apply control charts to a process, the process may not be in control. Even if it is in control, you don't yet understand its behavior. You will have to collect data from the process, establish control by uncovering and removing special causes, and then set up control charts to maintain control. We call this the **chart setup** stage. Later, when the process has been operating in control for some time, you understand its usual behavior and have a long run of data from the process. You keep control charts to monitor the process because a special cause could erupt at any time. We will call this **process monitoring**.⁶

Although in practice chart setup precedes process monitoring, the big ideas of control charts are more easily understood in the process-monitoring setting. We will start there, then discuss the more complex chart setup setting.

Choose a quantitative variable x that is an important measure of quality. The variable might be the diameter of a part, the number of envelopes stuffed in an hour, or the time to respond to a customer call. Here are the conditions for process monitoring.

PROCESS-MONITORING CONDITIONS

Measure a quantitative variable x that has a **Normal distribution**. The process has been operating in control for a long period, so that we know the **process mean μ** and the **process standard deviation σ** that describe the distribution of x as long as the process remains in control.

In practice, we must of course estimate the process mean and standard deviation from past data on the process. Under the process-monitoring conditions, we have very many observations and the process has remained in control. The law of large numbers tells us that estimates from past data will be very close to the truth about the process. That is, at the process-monitoring stage we can act as if we know the true values of μ and σ . Note carefully that μ and σ describe the center and spread of the variable x only as long as the process remains in control. A special cause may at any time disturb the process and change the mean, the standard deviation, or both.



To make control charts, begin by taking small samples from the process at regular intervals. For example, we might measure 4 or 5 consecutive parts or time the responses to 4 or 5 consecutive customer calls. There is an important idea here: *the observations in a sample are so close together that we can assume that the process is stable during this short period of time*. Variation within the same sample gives us a benchmark for the common cause variation in the process. *The process standard deviation σ refers to the standard deviation within the time period spanned by one sample*. If the process remains in control, the same σ describes the standard deviation of observations across any time period. Control charts help us decide whether this is the case.

We start with the **\bar{x} chart** based on plotting the means of the successive samples. Here is the outline:

\bar{x} chart

1. Take samples of size n from the process at regular intervals. Plot the means \bar{x} of these samples against the order in which the samples were taken.
2. We know that the sampling distribution of \bar{x} under the process-monitoring conditions is Normal with mean μ and standard deviation σ/\sqrt{n} (see text page 294). Draw a solid **center line** on the chart at height μ .
3. The 99.7 part of the 68–95–99.7 rule for Normal distributions (text page 77) says that, as long as the process remains in control, 99.7% of the values of \bar{x} will fall between $\mu - 3\sigma/\sqrt{n}$ and $\mu + 3\sigma/\sqrt{n}$. Draw dashed **control limits** on the chart at these heights. The control limits mark off the range of variation in sample means that we expect to see when the process remains in control.

center line

control limits

If the process remains in control and the process mean and standard deviation do not change, we will rarely observe an \bar{x} outside the control limits. Such an \bar{x} is therefore a signal that the process has been disturbed.



MONITORS

EXAMPLE 27.3 Manufacturing computer monitors

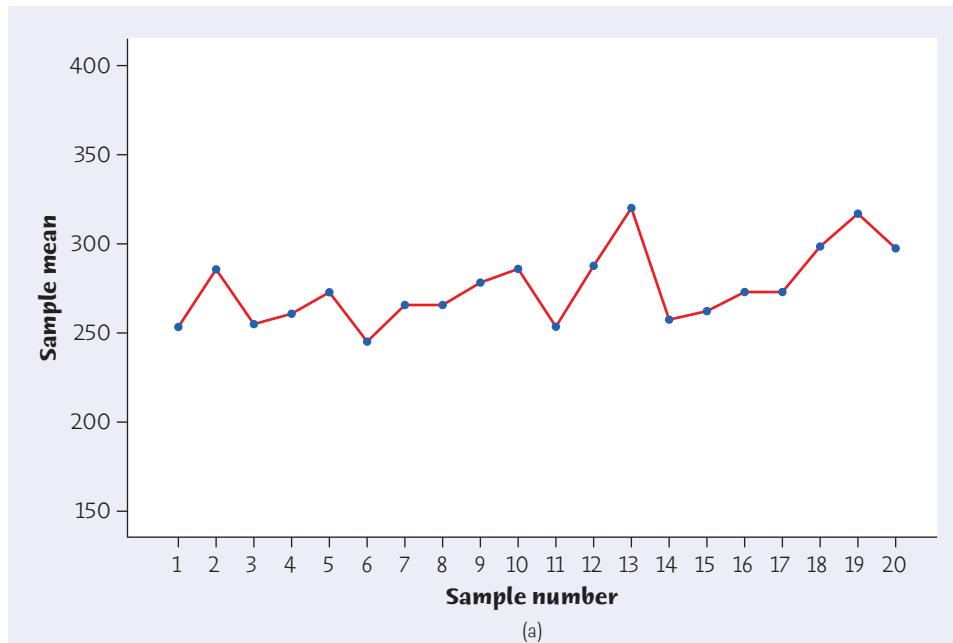
A manufacturer of computer monitors must control the tension on the mesh of fine vertical wires that lies behind the surface of the viewing screen. Too much tension will tear the mesh, and too little will allow wrinkles. Tension is measured by an electrical device with output readings in millivolts (mV). The manufacturing process has been stable with mean tension $\mu = 275$ mV and process standard deviation $\sigma = 43$ mV.

The mean 275 mV and the common cause variation measured by the standard deviation 43 mV describe the stable state of the process. If these values are not satisfactory—for example, if there is too much variation among the monitors—the manufacturer must make some fundamental change in the process. This might involve buying new equipment or changing the alloy used in the wires of the mesh. In fact, the common cause variation in mesh tension does not affect the performance of the monitors. We want to watch the process and maintain its current condition.

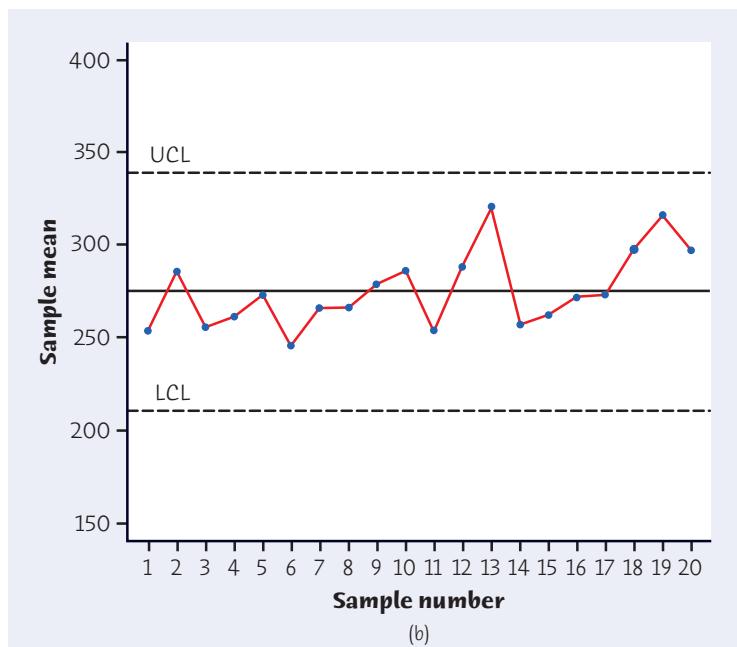
The operator measures the tension on a sample of 4 monitors each hour. Table 27.1 gives the last 20 samples. The table also gives the mean \bar{x} and the standard deviation s for each sample. The operator did not have to calculate these—modern measuring

TABLE 27.1 Twenty control chart samples of mesh tension (in millivolts)

SAMPLE	TENSION MEASUREMENTS					SAMPLE MEAN	STANDARD DEVIATION
1	234.5	272.3	234.5	272.3	253.4	21.8	
2	311.1	305.8	238.5	286.2	285.4	33.0	
3	247.1	205.3	252.6	316.1	255.3	45.7	
4	215.4	296.8	274.2	256.8	260.8	34.4	
5	327.9	247.2	283.3	232.6	272.7	42.5	
6	304.3	236.3	201.8	238.5	245.2	42.8	
7	268.9	276.2	275.6	240.2	265.2	17.0	
8	282.1	247.7	259.8	272.8	265.6	15.0	
9	260.8	259.9	247.9	345.3	278.5	44.9	
10	329.3	231.8	307.2	273.4	285.4	42.5	
11	266.4	249.7	231.5	265.2	253.2	16.3	
12	168.8	330.9	333.6	318.3	287.9	79.7	
13	349.9	334.2	292.3	301.5	319.5	27.1	
14	235.2	283.1	245.9	263.1	256.8	21.0	
15	257.3	218.4	296.2	275.2	261.8	33.0	
16	235.1	252.7	300.6	297.6	271.5	32.7	
17	286.3	293.8	236.2	275.3	272.9	25.6	
18	328.1	272.6	329.7	260.1	297.6	36.5	
19	316.4	287.4	373.0	286.0	315.7	40.7	
20	296.8	350.5	280.6	259.8	296.9	38.8	



(a)



(b)

FIGURE 27.4

(a) Plot of the sample means versus sample number for the mesh tension data of Table 27.1. (b) \bar{x} chart for the mesh tension data of Table 27.1. No points lie outside the control limits.

equipment often comes equipped with software that automatically records \bar{x} and s and even produces control charts. Figure 27.4(a) is a plot of the sample means versus sample number. ■

Figure 27.4(b) is an \bar{x} control chart for the 20 mesh tension samples in Table 27.1. We have plotted each sample mean from the table against its sample number. For example, the mean of the first sample is 253.4 mV, and this is the value

plotted for Sample 1. The center line is at $\mu = 275$ mV. The upper and lower control limits are

$$\mu + 3 \frac{\sigma}{\sqrt{n}} = 275 + 3 \frac{43}{\sqrt{4}} = 275 + 64.5 = 339.5 \text{ mV} \quad (\text{UCL})$$

$$\mu - 3 \frac{\sigma}{\sqrt{n}} = 275 - 3 \frac{43}{\sqrt{4}} = 275 - 64.5 = 210.5 \text{ mV} \quad (\text{LCL})$$

As is common, we have labeled the control limits UCL for upper control limit and LCL for lower control limit.

EXAMPLE 27.4 Interpreting \bar{x} charts

Figure 27.4(b) is a typical \bar{x} chart for a process in control. The means of the 20 samples do vary, but all lie within the range of variation marked out by the control limits. We are seeing the common cause variation of a stable process.

Figures 27.5 and 27.6 illustrate two ways in which the process can go out of control. In Figure 27.5, the process was disturbed by a special cause sometime between Sample 12 and Sample 13. As a result, the mean tension for Sample 13 falls above the upper control limit. It is common practice to mark all out-of-control points with an “x” to call attention to them. A search for the cause begins as soon as we see a point out of control. Investigation finds that the mounting of the tension-measuring device had slipped, resulting in readings that were too high. When the problem was corrected, Samples 14 to 20 were again in control.

Figure 27.6 shows the effect of a steady upward drift in the process center, starting at Sample 11. You see that some time elapses before the \bar{x} for Sample 18 is out of control. Process drift results from gradual changes such as the wearing of a cutting tool or overheating. The one-point-out signal works better for detecting sudden large disturbances than for detecting slow drifts in a process. ■

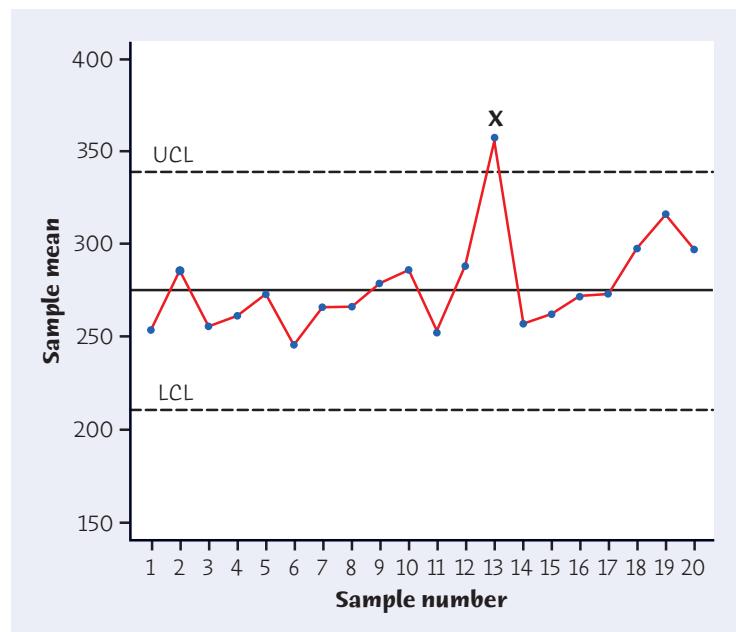
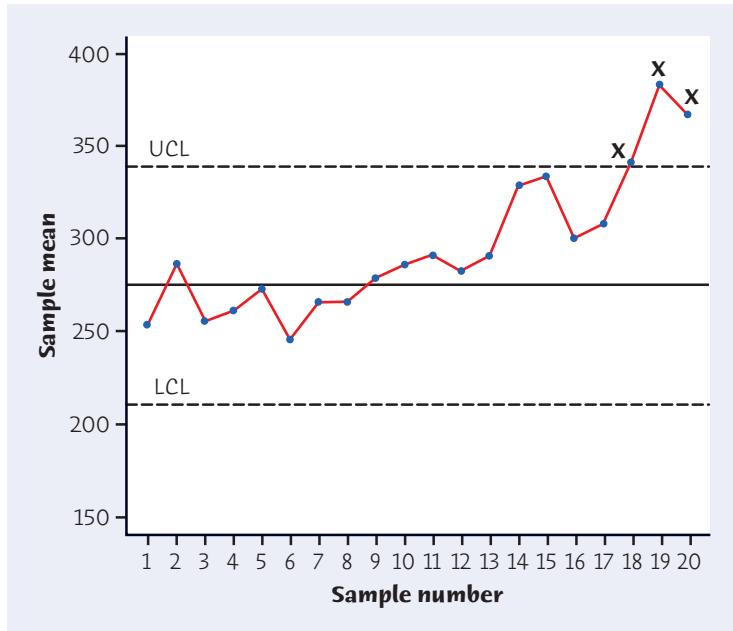


FIGURE 27.5

This \bar{x} chart is identical to that in Figure 27.4(b) except that a special cause has driven \bar{x} for Sample 13 above the upper control limit. The out-of-control point is marked with an x.

**FIGURE 27.6**

The first 10 points on this \bar{x} chart are as in Figure 27.4(b). The process mean drifts upward after Sample 10, and the sample means \bar{x} reflect this drift. The points for Samples 18, 19, and 20 are out of control.



APPLY YOUR KNOWLEDGE

27.10 Dry bleach. The net weight (in ounces) of boxes of dry bleach are monitored by taking samples of 4 boxes from each hour's production. The process mean should be $\mu = 16$ oz. Past experience indicates that the net weight when the process is properly adjusted varies with $\sigma = 0.4$ oz. The mean weight \bar{x} for each hour's sample is plotted on an \bar{x} control chart. Calculate the center line and control limits for this chart.

27.11 Tablet hardness. A pharmaceutical manufacturer forms tablets by compressing a granular material that contains the active ingredient and various fillers. The hardness of a sample from each lot of tablets is measured in order to control the compression process. The process has been operating in control with mean at the target value $\mu = 11.5$ kilograms (kg) and estimated standard deviation $\sigma = 0.2$ kg. Table 27.2 gives three sets of data, each representing \bar{x} for 20 successive samples of $n = 4$ tablets. One set remains in control at the target value. In a second set, the process mean μ shifts suddenly to a new value. In a third, the process mean drifts gradually.

- What are the center line and control limits for an \bar{x} chart for this process?
- Draw a separate \bar{x} chart for each of the three data sets. Mark any points that are beyond the control limits.
- Based on your work in (b) and the appearance of the control charts, which set of data comes from a process that is in control? In which case does the process mean shift suddenly, and at about which sample do you think that the mean changed? Finally, in which case does the mean drift gradually? HARDNESS

TABLE 27.2 Three sets of \bar{x} 's from 20 samples of size 4

SAMPLE	DATA SET A	DATA SET B	DATA SET C
1	11.602	11.627	11.495
2	11.547	11.613	11.475
3	11.312	11.493	11.465
4	11.449	11.602	11.497
5	11.401	11.360	11.573
6	11.608	11.374	11.563
7	11.471	11.592	11.321
8	11.453	11.458	11.533
9	11.446	11.552	11.486
10	11.522	11.463	11.502
11	11.664	11.383	11.534
12	11.823	11.715	11.624
13	11.629	11.485	11.629
14	11.602	11.509	11.575
15	11.756	11.429	11.730
16	11.707	11.477	11.680
17	11.612	11.570	11.729
18	11.628	11.623	11.704
19	11.603	11.472	12.052
20	11.816	11.531	11.905

s CHARTS FOR PROCESS MONITORING

The \bar{x} charts in Figures 27.4(b), 27.5, and 27.6 were easy to interpret because the process standard deviation remained fixed at 43 mV. The effects of moving the process mean away from its in-control value (275 mV) are then clear to see. We know that even the simplest description of a distribution should give both a measure of center and a measure of spread. So it is with control charts. We must monitor both the process center, using an \bar{x} chart, and the process spread, using a control chart for the sample standard deviation s .

The standard deviation s does not have a Normal distribution, even approximately. Under the process-monitoring conditions, the sampling distribution of s is skewed to the right. Nonetheless, control charts for any statistic are based on the “plus or minus three standard deviations” idea motivated by the 68–95–99.7 rule for Normal distributions. Control charts are intended to be practical tools that are easy to use. Standard practice in process control therefore ignores such details as the effect of non-Normal sampling distributions. Here is the general control chart setup for a sample statistic Q (short for “quality characteristic”).

THREE-SIGMA CONTROL CHARTS

To make a **three-sigma** (3σ) control chart for any statistic Q :

1. Take samples from the process at regular intervals and plot the values of the statistic Q against the order in which the samples were taken.
2. Draw a **center line** on the chart at height μ_Q , the mean of the statistic when the process is in control.
3. Draw upper and lower **control limits** on the chart three standard deviations of Q above and below the mean. That is,

$$UCL = \mu_Q + 3\sigma_Q$$

$$LCL = \mu_Q - 3\sigma_Q$$

Here σ_Q is the standard deviation of the sampling distribution of the statistic Q when the process is in control.

4. The chart produces an **out-of-control signal** when a plotted point lies outside the control limits.

We have applied this general idea to \bar{x} charts. If μ and σ are the process mean and standard deviation, the statistic \bar{x} has mean $\mu_{\bar{x}} = \mu$ and standard deviation $\sigma_{\bar{x}} = \sigma/\sqrt{n}$. The center line and control limits for \bar{x} charts follow from these facts.

What are the corresponding facts for the sample standard deviation s ? Study of the sampling distribution of s for samples from a Normally distributed process characteristic gives these facts:

1. The *mean* of s is a constant times the process standard deviation σ , $\mu_s = c_4\sigma$.
2. The *standard deviation* of s is also a constant times the process standard deviation, $\sigma_s = c_s\sigma$.

The constants are called c_4 and c_s for historical reasons. Their values depend on the size of the samples. For large samples, c_4 is close to 1. That is, the sample standard deviation s has little bias as an estimator of the process standard deviation σ . Because statistical process control often uses small samples, we pay attention to the value of c_4 . Following the general pattern for three-sigma control charts:

1. The *center line* of an s chart is at $c_4\sigma$.
2. The *control limits* for an s chart are at

$$UCL = \mu_s + 3\sigma_s = c_4\sigma + 3c_s\sigma = (c_4 + 3c_s)\sigma$$

$$LCL = \mu_s - 3\sigma_s = c_4\sigma - 3c_s\sigma = (c_4 - 3c_s)\sigma$$

That is, the control limits UCL and LCL are also constants times the process standard deviation. These constants are called (again for historical reasons) B_6 and B_5 . We don't need to remember that $B_6 = c_4 + 3c_s$ and $B_5 = c_4 - 3c_s$, because tables give us the numerical values of B_6 and B_5 .

\bar{x} AND s CONTROL CHARTS FOR PROCESS MONITORING⁷

Take regular samples of size n from a process that has been in control with process mean μ and process standard deviation σ . The center line and control limits for an \bar{x} chart are

$$\text{UCL} = \mu + 3 \frac{\sigma}{\sqrt{n}}$$

$$\text{CL} = \mu$$

$$\text{LCL} = \mu - 3 \frac{\sigma}{\sqrt{n}}$$

The center line and control limits for an s chart are

$$\text{UCL} = B_6 \sigma$$

$$\text{CL} = c_4 \sigma$$

$$\text{LCL} = B_5 \sigma$$

The control chart constants c_4 , B_5 , and B_6 depend on the sample size n .

Table 27.3 gives the values of the control chart constants c_4 , c_5 , B_5 , and B_6 for samples of sizes 2 to 10. This table makes it easy to draw s charts. The table has no B_5 entries for samples of size smaller than $n = 6$. The lower control limit for an s chart is zero for samples of sizes 2 to 5. This is a consequence of the fact that s has a right-skewed distribution and takes only values greater than zero. Three standard deviations above the mean (UCL) lies on the long right side of the distribution. Three standard deviations below the mean (LCL) on the short left side is below zero, so we say that $\text{LCL} = 0$.

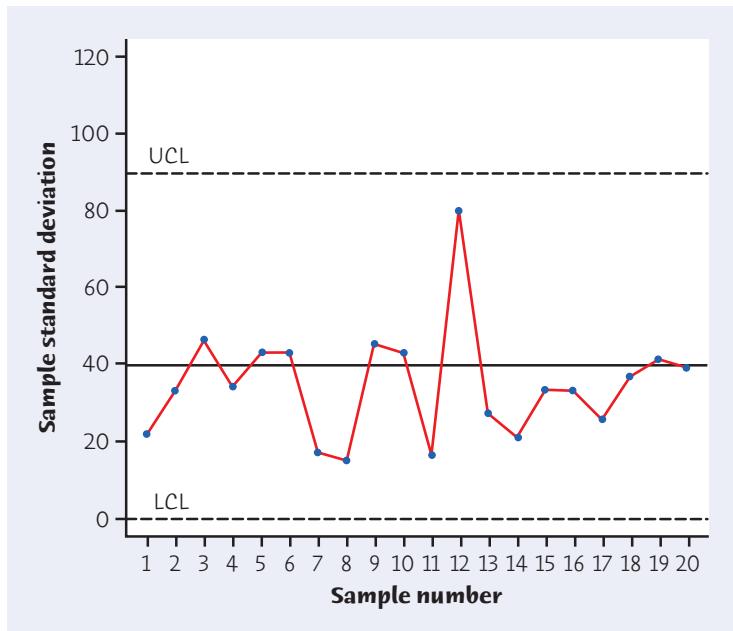
TABLE 27.3 Control chart constants

SAMPLE SIZE n	c_4	c_5	B_5	B_6
2	0.7979	0.6028		2.606
3	0.8862	0.4633		2.276
4	0.9213	0.3889		2.088
5	0.9400	0.3412		1.964
6	0.9515	0.3076	0.029	1.874
7	0.9594	0.2820	0.113	1.806
8	0.9650	0.2622	0.179	1.751
9	0.9693	0.2459	0.232	1.707
10	0.9727	0.2321	0.276	1.669

EXAMPLE 27.5 \bar{x} and s charts for mesh tension

Figure 27.7 is the s chart for the computer monitor mesh tension data in Table 27.1. The samples are of size $n = 4$ and the process standard deviation in control is $\sigma = 43$ mV. The center line is therefore

$$\text{CL} = c_4 \sigma = (0.9213)(43) = 39.6 \text{ mV}$$

**FIGURE 27.7**

s chart for the mesh tension data of Table 27.1. Both the s chart and the \bar{x} chart (Figure 27.4(b)) are in control.

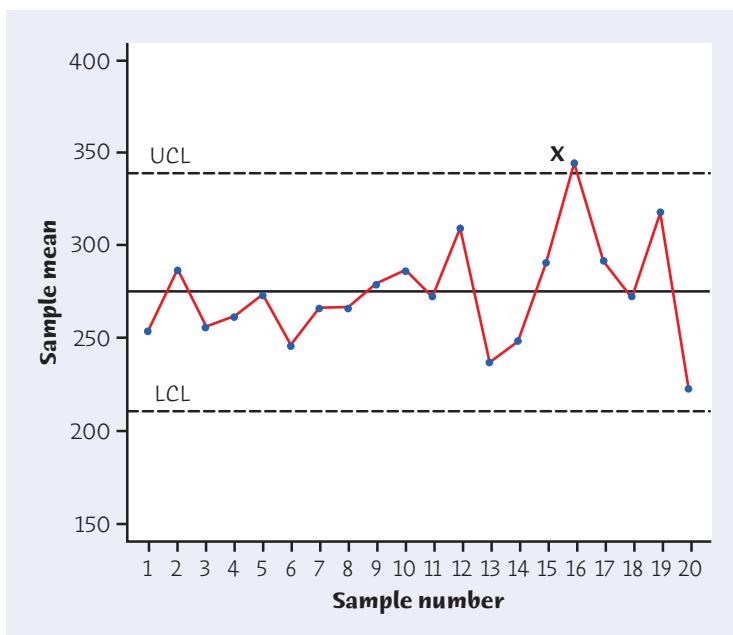
The control limits are

$$UCL = B_6\sigma = (2.088)(43) = 89.8$$

$$LCL = B_5\sigma = (0)(43) = 0$$

Figures 27.4(b) and 27.7 go together: they are \bar{x} and s charts for monitoring the mesh-tensioning process. Both charts are in control, showing only common cause variation within the bounds set by the control limits.

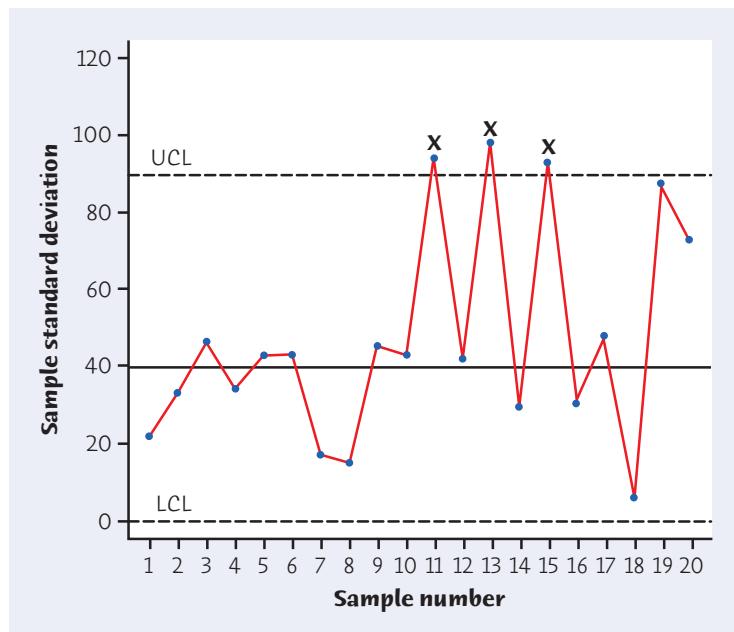
Figures 27.8 and 27.9 are \bar{x} and s charts for the mesh-tensioning process when a new and poorly trained operator takes over between Samples 10 and 11. The new operator

**FIGURE 27.8**

\bar{x} chart for mesh tension when the process variability increases after Sample 10. The \bar{x} chart does show the increased variability, but the s chart is clearer and should be read first.

FIGURE 27.9

s chart for mesh tension when the process variability increases after Sample 10. Increased within-sample variability is clearly visible. Find and remove the s-type special cause before reading the \bar{x} chart.



introduces added variation into the process, increasing the process standard deviation from its in-control value of 43 mV to 60 mV. The \bar{x} chart in Figure 27.8 shows one point out of control. Only on closer inspection do we see that the spread of the \bar{x} 's increases after Sample 10. In fact, the process mean has remained unchanged at 275 mV. The  apparent lack of control in the \bar{x} chart is entirely due to the larger process variation. There is a lesson here: *it is difficult to interpret an \bar{x} chart unless s is in control. When you look at \bar{x} and s charts, always start with the s chart.*

The s chart in Figure 27.9 shows lack of control starting at Sample 11. As usual, we mark the out-of-control points with an "x". The points for Samples 13 and 15 also lie above the UCL, and the overall spread of the sample points is much greater than for the first 10 samples. In practice, the s chart would call for action after Sample 11. We would ignore the \bar{x} chart until the special cause (the new operator) for the lack of control in the s chart has been found and removed by training the operator. ■

Example 27.5 suggests a strategy for using \bar{x} and s charts in practice. First examine the s chart. Lack of control on an s chart is due to special causes that affect the observations *within a sample* differently. New and nonuniform raw material, a new and poorly trained operator, and mixing results from several machines or several operators are typical "s-type" special causes.

Once the s chart is in control, the stable value of the process standard deviation σ means that the variation within samples serves as a benchmark for detecting variation in the level of the process over the longer time periods between samples. The \bar{x} chart, with control limits that depend on σ , does this. The \bar{x} chart, as we saw in Example 27.5, responds to s-type causes as well as to longer-range changes in the process, so it is important to eliminate s-type special causes first. Then the \bar{x} chart will alert us to, for example, a change in process level caused by new raw material that differs from that used in the past or a gradual drift in the process level caused by wear in a cutting tool.

EXAMPLE 27.6 s-type and \bar{x} -type special causes

A large health maintenance organization (HMO) uses control charts to monitor the process of directing patient calls to the proper department or doctor's receptionist. Each day at a random time, 5 consecutive calls are recorded electronically. The first call today is handled quickly by an experienced operator, but the next goes to a newly hired operator who must ask a supervisor for help. The sample has a large s , and lack of control signals the need to train new hires more thoroughly.

The same HMO monitors the time required to receive orders from its main supplier of pharmaceutical products. After a long period in control, the \bar{x} chart shows a systematic shift downward in the mean time because the supplier has changed to a more efficient delivery service. This is a desirable special cause, but it is nonetheless a systematic change in the process. The HMO will have to establish new control limits that describe the new state of the process, with smaller process mean μ . ■

The second setting in Example 27.6 reminds us that a major change in the process returns us to the chart setup stage. In the absence of deliberate changes in the process, process monitoring uses the same values of μ and σ for long periods of time. There is one important exception: careful monitoring and removal of special causes as they occur can permanently reduce the process σ . If the points on the s chart remain near the center line for a long period, it is wise to update the value of σ to the new, smaller value and compute new values of UCL and LCL for both \bar{x} and s charts.

APPLY YOUR KNOWLEDGE

27.12 Making cappuccino. A large chain of coffee shops records a number of performance measures. Among them is the time required to complete an order for a cappuccino, measured from the time the order is placed. Suggest some plausible examples of each of the following.

- Reasons for common cause variation in response time.
- s -type special causes.
- \bar{x} -type special causes.

27.13 Dry bleach. In Exercise 27.10 (page 27-15) you gave the center line and control limits for an \bar{x} chart. What are the center line and control limits for an s chart for this process?

27.14 Tablet hardness. Exercise 27.11 concerns process control data on the hardness of tablets (measured in kilograms) for a pharmaceutical product. Table 27.4 gives data for 20 new samples of size 4, with the \bar{x} and s for each sample. The process has been in control with mean at the target value $\mu = 11.5$ kg and standard deviation $\sigma = 0.2$ kg.  HARDNESS2

- Make both \bar{x} and s charts for these data based on the information given about the process.
- At some point, the within-sample process variation increased from $\sigma = 0.2$ kg to $\sigma = 0.4$ kg. About where in the 20 samples did this happen? What is the effect on the s chart? On the \bar{x} chart?
- At that same point, the process mean changed from $\mu = 11.5$ kg to $\mu = 11.7$ kg. What is the effect of this change on the s chart? On the \bar{x} chart?

TABLE 27.4 Twenty samples of size 4, with \bar{x} and s

SAMPLE	HARDNESS (KILOGRAMS)				\bar{x}	s
1	11.432	11.350	11.582	11.184	11.387	0.1660
2	11.791	11.323	11.734	11.512	11.590	0.2149
3	11.373	11.807	11.651	11.651	11.620	0.1806
4	11.787	11.585	11.386	11.245	11.501	0.2364
5	11.633	11.212	11.568	11.469	11.470	0.1851
6	11.648	11.653	11.618	11.314	11.558	0.1636
7	11.456	11.270	11.817	11.402	11.486	0.2339
8	11.394	11.754	11.867	11.003	11.504	0.3905
9	11.349	11.764	11.402	12.085	11.650	0.3437
10	11.478	11.761	11.907	12.091	11.809	0.2588
11	11.657	12.524	11.468	10.946	11.649	0.6564
12	11.820	11.872	11.829	11.344	11.716	0.2492
13	12.187	11.647	11.751	12.026	11.903	0.2479
14	11.478	11.222	11.609	11.271	11.395	0.1807
15	11.750	11.520	11.389	11.803	11.616	0.1947
16	12.137	12.056	11.255	11.497	11.736	0.4288
17	12.055	11.730	11.856	11.357	11.750	0.2939
18	12.107	11.624	11.727	12.207	11.916	0.2841
19	11.933	10.658	11.708	11.278	11.394	0.5610
20	12.512	12.315	11.671	11.296	11.948	0.5641



Ric Ergenbright/CORBIS

27.15 Dyeing yarn. The unique colors of the cashmere sweaters your firm makes result from heating undyed yarn in a kettle with a dye liquor. The pH (acidity) of the liquor is critical for regulating dye uptake and hence the final color. There are 5 kettles, all of which receive dye liquor from a common source. Twice each day, the pH of the liquor in each kettle is measured, giving samples of size 5. The process has been operating in control with $\mu = 4.22$ and $\sigma = 0.127$.

- (a) Give the center line and control limits for the s chart.
- (b) Give the center line and control limits for the \bar{x} chart.

27.16 Mounting-hole distances. Figure 27.10 reproduces a data sheet from the floor of a factory that makes electrical meters.⁸ The sheet shows measurements on the distance between two mounting holes for 18 samples of size 5. The heading informs us that the measurements are in multiples of 0.0001 inch above 0.6000 inch. That is, the first measurement, 44, stands for 0.6044 inch. All the measurements end in 4. Although we don't know why this is true, it is clear that in effect the measurements were made to the nearest 0.001 inch, not to the nearest 0.0001 inch.

Calculate \bar{x} and s for the first two samples. The data file contains \bar{x} and s for all 18 samples. Based on long experience with this process, you are keeping control charts based on $\mu = 43$ and $\sigma = 12.74$. Make s and \bar{x} charts for the data in Figure 27.10 and describe the state of the process. 

VARIABLES CONTROL CHART (\bar{X} & R)																	
Part name (project) Metal frame				Operation (process) Distance between mounting holes										Part No. 32506	Chart No. 1		
Operator		Machine R-5			Gage			Unit of measure 0.0001"				Zero equals 0.6000"					
Date		3/7			3/8			3/9									
Time		8:30	10:30	11:45	1:30	8:15	10:15	11:45	2:00	3:00	4:00	8:30	10:00	11:45	1:30	2:30	3:30
Sample measurements	1	44	64	34	44	34	34	54	64	24	34	34	54	44	24	54	54
	2	44	44	44	54	14	64	64	34	54	44	44	44	24	24	24	34
	3	44	34	54	54	84	34	34	54	44	44	34	24	34	54	24	74
	4	44	34	44	34	54	44	44	44	34	34	64	54	34	44	44	34
	5	64	54	54	44	44	44	34	44	34	34	34	24	44	44	54	54
Average, \bar{X}																	
Range, R		20	30	20	20	70	30	30	30	30	10	30	30	20	30	40	30

FIGURE 27.10

A process control record sheet kept by operators, for Exercise 27.16. This is typical of records kept by hand when measurements are not automated. We will see in the next section why such records mention \bar{x} and R control charts rather than \bar{x} and s charts.

27.17 Dyeing yarn: special causes. The process described in Exercise 27.15 goes out of control. Investigation finds that a new type of yarn was recently introduced. The pH in the kettles is influenced by both the dye liquor and the yarn. Moreover, on a few occasions a faulty valve on one of the kettles had allowed water to enter that kettle; as a result, the yarn in that kettle had to be discarded. Which of these special causes appears on the s chart and which on the \bar{x} chart? Explain your answer.

USING CONTROL CHARTS

We are now familiar with the ideas that undergird all control charts and also with the details of making \bar{x} and s charts. This section discusses two topics related to using control charts in practice.

\bar{x} and R charts. We have seen that it is essential to monitor both the center and the spread of a process. Control charts were originally intended to be used by factory workers with limited knowledge of statistics in the era before even calculators, let alone software, were common. In that environment, it takes too long to calculate standard deviations. The \bar{x} chart for center was therefore combined with a control chart for spread based on the **sample range** rather than the sample standard deviation. The range R of a sample is just the difference between the largest and smallest observations. It is easy to find R without a calculator. Using R rather than s to measure the spread of samples replaces the s chart with an **R chart**. It also changes the \bar{x} chart because the control limits for \bar{x} use the estimated process spread. So \bar{x} and R charts differ in the details of both charts from \bar{x} and s charts.

Because the range R uses only the largest and smallest observations in a sample, it is less informative than the standard deviation s calculated from all the observations. For this reason, \bar{x} and s charts are now preferred to \bar{x} and R charts. R charts remain common because tradition dies hard and also because it is easier for workers to understand R than s. In this short introduction, we concentrate on the principles of control charts, so we won't give the details of constructing \bar{x} and R charts. These details appear in any text on quality control.⁹ If you meet a set of

sample range

R chart

\bar{x} and R charts, remember that the interpretation of these charts is just like the interpretation of \bar{x} and s charts.

Additional out-of-control signals. So far, we have used only the basic “one point beyond the control limits” criterion to signal that a process may have gone out of control. We would like a quick signal when the process moves out of control, but we also want to avoid “false alarms,” signals that occur just by chance when the process is really in control. The standard 3σ control limits are chosen to prevent too many false alarms, because an out-of-control signal calls for an effort to find and remove a special cause. As a result, \bar{x} charts are often slow to respond to a gradual drift in the process center that continues for some time before finally forcing a reading outside the control limits. We can speed the response of a control chart to lack of control—at the cost of also enduring more false alarms—by adding patterns other than “one-point-out” as signals. The most common step in this direction is to add a *runs signal* to the \bar{x} chart.

OUT-OF-CONTROL SIGNALS

\bar{x} and s or \bar{x} and R control charts produce an out-of-control signal if:

- **One-point-out:** A single point lies outside the 3σ control limits of either chart.
- **Run:** The \bar{x} chart shows 9 consecutive points above the center line or 9 consecutive points below the center line. The signal occurs when we see the 9th point of the run.

EXAMPLE 27.7 Using the runs signal

Figure 27.11 reproduces the \bar{x} chart from Figure 27.6. The process center began a gradual upward drift at Sample 11. The chart shows the effect of the drift—the sample

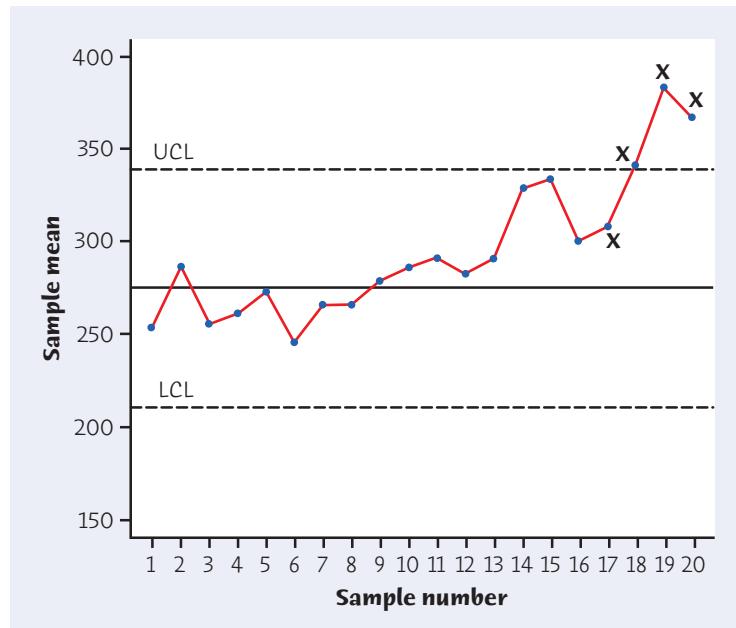


FIGURE 27.11

\bar{x} chart for mesh tension data when the process center drifts upward. The “run of 9” signal gives an out-of-control warning at Sample 17.

means plotted on the chart move gradually upward, with some random variation. The one-point-out signal does not call for action until Sample 18 finally produces an \bar{x} above the UCL. The runs signal reacts more quickly: Sample 17 is the 9th consecutive point above the center line. ■

It is a mathematical fact that the runs signal responds to a gradual drift more quickly (on the average) than the one-point-out signal does. The motivation for a runs signal is that when a process is in control, the probability of a false alarm is about the same for the runs signal as for the one-point-out signal. There are many other signals that can be added to the rules for responding to \bar{x} and s or \bar{x} and R charts. *In our enthusiasm to detect various special kinds of loss of control, it is easy to forget that adding signals always increases the frequency of false alarms.* Frequent false alarms are so annoying that the people responsible for responding soon begin to ignore out-of-control signals. It is better to use only a few signals and to reserve signals other than one-point-out and runs for processes that are known to be prone to specific special causes for which there is a tailor-made signal.¹⁰



APPLY YOUR KNOWLEDGE

27.18 Special causes. Is each of the following examples of a special cause most likely to first result in (i) one-point-out on the s or R chart, (ii) one-point-out on the \bar{x} chart, or (iii) a run on the \bar{x} chart? In each case, briefly explain your reasoning.

- The sharpness of a power saw blade deteriorates as more items are cut.
- Buildup of dirt reduces the precision with which parts are placed for machining.
- A new customer service representative for a Spanish-language help line is not a native speaker and has difficulty understanding customers.
- A nurse responsible for filling out insurance claim forms grows less attentive as her shift continues.

27.19 Mixtures. Here is an artificial situation that illustrates an unusual control chart pattern. Invoices are processed and paid by two clerks, one very experienced and the other newly hired. The experienced clerk processes invoices quickly. The new hire must often refer to a handbook and is much slower. Both are quite consistent, so that their times vary little from invoice to invoice. It happens that each sample of invoices comes from one of the clerks, so that some samples are from one and some from the other clerk. Sketch the \bar{x} chart pattern that will result.

SETTING UP CONTROL CHARTS

When you first approach a process that has not been carefully studied, it is quite likely that the process is not in control. Your first goal is to discover and remove special causes and so bring the process into control. Control charts are an important tool. Control charts for process monitoring follow the process forward in time

to keep it in control. Control charts at the *chart setup* stage, on the other hand, look back in an attempt to discover the present state of the process. An example will illustrate the method.



VISCOSEITY

EXAMPLE 27.8 Viscosity of an elastomer

The viscosity of a material is its resistance to flow when under stress. Viscosity is a critical characteristic of rubber and rubber-like compounds called elastomers, which have many uses in consumer products. Viscosity is measured by placing specimens of the material above and below a slowly rotating roller, squeezing the assembly, and recording the drag on the roller. Measurements are in “Mooney units,” named after the inventor of the instrument.

A specialty chemical company is beginning production of an elastomer that is supposed to have viscosity 45 ± 5 Mooneys. Each lot of the elastomer is produced by “cooking” raw material with catalysts in a reactor vessel. Table 27.5 records \bar{x} and s from samples of size $n = 4$ lots from the first 24 shifts as production begins.¹¹ An s chart therefore monitors variation among lots produced during the same shift. If the s chart is in control, an \bar{x} chart looks for shift-to-shift variation. ■

Estimating μ . We do not know the process mean μ and standard deviation σ . What shall we do? Sometimes we can easily adjust the center of a process by setting some control, such as the depth of a cutting tool in a machining operation or the temperature of a reactor vessel in a pharmaceutical plant. In such cases it is usual to simply take the process mean μ to be the target value, the depth or temperature that the design of the process specifies as correct. The \bar{x} chart then helps us keep the process mean at this target value.

TABLE 27.5 \bar{x} and s for 24 samples of elastomer viscosity (in Mooneys)

SAMPLE	\bar{x}	s	SAMPLE	\bar{x}	s
1	49.750	2.684	13	47.875	1.118
2	49.375	0.895	14	48.250	0.895
3	50.250	0.895	15	47.625	0.671
4	49.875	1.118	16	47.375	0.671
5	47.250	0.671	17	50.250	1.566
6	45.000	2.684	18	47.000	0.895
7	48.375	0.671	19	47.000	0.447
8	48.500	0.447	20	49.625	1.118
9	48.500	0.447	21	49.875	0.447
10	46.250	1.566	22	47.625	1.118
11	49.000	0.895	23	49.750	0.671
12	48.125	0.671	24	48.625	0.895

There is less likely to be a “correct value” for the process mean μ if we are monitoring response times to customer calls or data entry errors. In Example 27.8, we have the target value 45 Mooneys, but there is no simple way to set viscosity at the desired level. In such cases, we want the μ we use in our \bar{x} chart to describe the center of the process as it has actually been operating. To do this, just take the mean of all the individual measurements in the past samples. Because the samples are all the same size, this is just the mean of the sample \bar{x} 's. The overall “mean of the sample means” is therefore usually called $\bar{\bar{x}}$. For the 24 samples in Table 27.5,

$$\begin{aligned}\bar{\bar{x}} &= \frac{1}{24}(49.750 + 49.375 + \cdots + 48.625) \\ &= \frac{1161.125}{24} = 48.380\end{aligned}$$

Estimating σ . It is almost never safe to use a “target value” for the process standard deviation σ because it is almost never possible to directly adjust process variation. We must estimate σ from past data. We want to combine the sample standard deviations s from past samples rather than use the standard deviation of all the individual observations in those samples. That is, in Example 27.8, we want to combine the 24 sample standard deviations in Table 27.5 rather than calculate the standard deviation of the 96 observations in these samples. The reason is that it is the *within-sample* variation that is the benchmark against which we compare the longer-term process variation. Even if the process has been in control, we want only the variation over the short time period of a single sample to influence our value for σ .

There are several ways to estimate σ from the sample standard deviations. In practice, software may use a somewhat sophisticated method and then calculate the control limits for you. We use a simple method that is traditional in quality control because it goes back to the era before software. If we are basing chart setup on k past samples, we have k sample standard deviations s_1, s_2, \dots, s_k . Just average these to get

$$\bar{s} = \frac{1}{k}(s_1 + s_2 + \cdots + s_k)$$

For the viscosity example, we average the s -values for the 24 samples in Table 27.5:

$$\begin{aligned}\bar{s} &= \frac{1}{24}(2.684 + 0.895 + \cdots + 0.895) \\ &= \frac{24.156}{24} = 1.0065\end{aligned}$$

Combining the sample s -values to estimate σ introduces a complication: the samples used in process control are often small (size $n = 4$ in the viscosity example), so s has some bias as an estimator of σ . Recall that $\mu_s = c_4\sigma$. The mean \bar{s} inherits this bias: its mean is also not σ but $c_4\sigma$. The proper estimate of σ corrects this bias. It is

$$\hat{\sigma} = \frac{\bar{s}}{c_4}$$

We get control limits from past data by using the estimates $\bar{\bar{x}}$ and $\hat{\sigma}$ in place of the μ and σ used in charts at the process-monitoring stage. Here are the results.¹²

\bar{x} AND s CONTROL CHARTS USING PAST DATA

Take regular samples of size n from a process. Estimate the process mean μ and the process standard deviation σ from past samples by

$$\hat{\mu} = \bar{\bar{x}} \quad (\text{or use a target value})$$

$$\hat{\sigma} = \frac{\bar{s}}{c_4}$$

The center line and control limits for an \bar{x} chart are

$$\text{UCL} = \hat{\mu} + 3 \frac{\hat{\sigma}}{\sqrt{n}}$$

$$\text{CL} = \hat{\mu}$$

$$\text{LCL} = \hat{\mu} - 3 \frac{\hat{\sigma}}{\sqrt{n}}$$

The center line and control limits for an s chart are

$$\text{UCL} = B_6 \hat{\sigma}$$

$$\text{CL} = c_4 \hat{\sigma} = \bar{s}$$

$$\text{LCL} = B_5 \hat{\sigma}$$

If the process was not in control when the samples were taken, these should be regarded as trial control limits.

We are now ready to outline the chart setup procedure for elastomer viscosity.

Step 1. As usual, we look first at an s chart. For chart setup, control limits are based on the same past data that we will plot on the chart. Calculate from Table 27.5 that

$$\bar{s} = 1.0065$$

$$\hat{\sigma} = \frac{\bar{s}}{c_4} = \frac{1.0065}{0.9213} = 1.0925$$

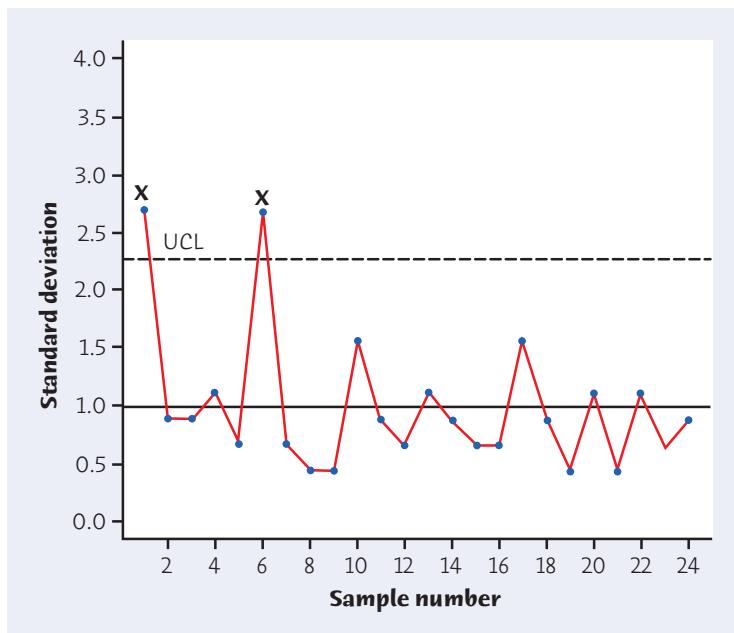
The center line and control limits for an s chart based on past data are

$$\text{UCL} = B_6 \hat{\sigma} = (2.088)(1.0925) = 2.281$$

$$\text{CL} = \bar{s} = 1.0065$$

$$\text{LCL} = B_5 \hat{\sigma} = (0)(1.0925) = 0$$

Figure 27.12 is the s chart. The points for Shifts 1 and 6 lie above the UCL. Both are near the beginning of production. Investigation finds that the reactor operator made an error on one lot in each of these samples. The error changed the viscosity of that lot and increased s for that one sample. The error will not be repeated now that the operators have gained experience. That is, this special cause has already been removed.

**FIGURE 27.12**

s chart based on past data for the viscosity data of Table 27.5. The control limits are based on the same s-values that are plotted on the chart. Points 1 and 6 are out of control.

Step 2. Remove the two values of s that were out of control. This is proper because the special cause responsible for these readings is no longer present. Recalculate from the remaining 22 shifts that $\bar{s} = 0.854$ and $\hat{\sigma} = 0.854/0.9213 = 0.927$. Make a new s chart with

$$UCL = B_6 \hat{\sigma} = (2.088)(0.927) = 1.936$$

$$CL = \bar{s} = 0.854$$

$$LCL = B_3 \hat{\sigma} = (0)(0.927) = 0$$

We don't show the chart, but you can see from Table 27.5 that none of the remaining s-values lies above the new, lower, UCL; the largest remaining s is 1.566. If additional points were now out of control, we would repeat the process of finding and eliminating s-type causes until the s chart for the remaining shifts was in control. In practice, of course, this is often a challenging task.

Step 3. Once s-type causes have been eliminated, make an \bar{x} chart using only the samples that remain after dropping those that had out-of-control s-values. For the 22 remaining samples, we know that $\hat{\sigma} = 0.927$ and we calculate that $\bar{\bar{x}} = 48.4716$. The center line and control limits for the \bar{x} chart are

$$UCL = \bar{\bar{x}} + 3 \frac{\hat{\sigma}}{\sqrt{n}} = 48.4716 + 3 \frac{0.927}{\sqrt{4}} = 49.862$$

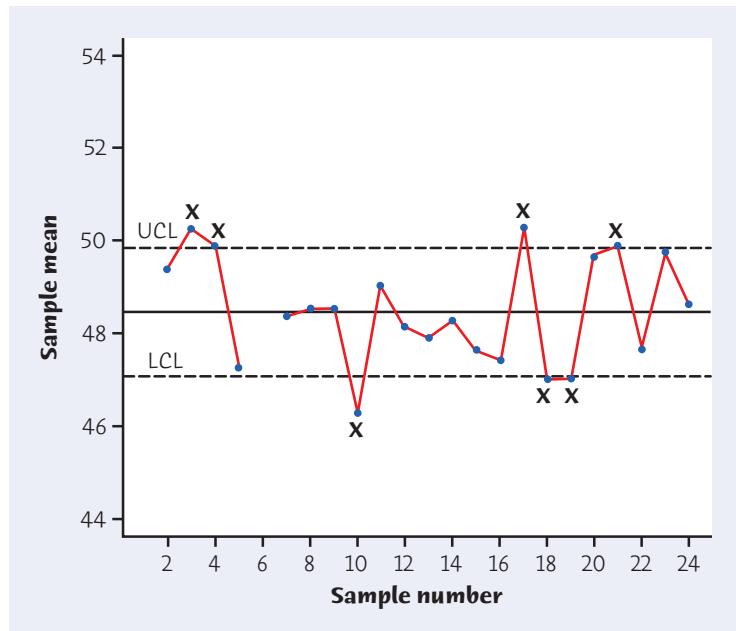
$$CL = \bar{\bar{x}} = 48.4716$$

$$LCL = \bar{\bar{x}} - 3 \frac{\hat{\sigma}}{\sqrt{n}} = 48.4716 - 3 \frac{0.927}{\sqrt{4}} = 47.081$$

Figure 27.13 is the \bar{x} chart. Shifts 1 and 6 have been dropped. Seven of the 22 points are beyond the 3σ limits, four high and three low. Although within-shift

FIGURE 27.13

\bar{x} chart based on past data for the viscosity data of Table 27.5. The samples for Shifts 1 and 6 have been removed because s -type special causes active in those samples are no longer active. The \bar{x} chart shows poor control.



variation is now stable, there is excessive variation from shift to shift. To find the cause, we must understand the details of the process, but knowing that the special cause or causes operate between shifts is a big help. If the reactor is set up anew at the beginning of each shift, that's one place to look more closely.

Step 4. Once the \bar{x} and s charts are both in control (looking backward), use the estimates $\hat{\mu}$ and $\hat{\sigma}$ from the points in control to set tentative control limits to monitor the process going forward. If it remains in control, we can update the charts and move to the process-monitoring stage.

APPLY YOUR KNOWLEDGE

27.20 From setup to monitoring. Suppose that when the chart setup project of Example 27.8 is complete, the points remaining after removing special causes have $\bar{x} = 48.7$ and $\bar{s} = 0.92$. What are the center line and control limits for the \bar{x} and s charts you would use to monitor the process going forward?

27.21 Estimating process parameters. The \bar{x} and s control charts for the mesh-tensioning example (Figures 27.4(b) (page 27-13) and 27.7 (page 27-19)) were based on $\mu = 275$ mV and $\sigma = 43$ mV. Table 27.1 (page 27-12) gives the 20 most recent samples from this process.

- Estimate the process μ and σ based on these 20 samples.
- Your calculations suggest that the process σ may now be less than 43 mV. Explain why the s chart in Figure 27.7 (page 27-19) suggests the same conclusion. (If this pattern continues, we would eventually update the value of σ used for control limits.) 

27.22 Hospital losses. Table 27.6 gives data on the losses (in dollars) incurred by a hospital in treating major joint replacement (DRG 209) patients.¹³ The hospital

TABLE 27.6 Hospital losses (dollars) for 15 samples of DRG 209 patients

SAMPLE	LOSS								SAMPLE MEAN	STANDARD DEVIATION
1	6835	5843	6019	6731	6362	5696	7193	6206	6360.6	521.7
2	6452	6764	7083	7352	5239	6911	7479	5549	6603.6	817.1
3	7205	6374	6198	6170	6482	4763	7125	6241	6319.8	749.1
4	6021	6347	7210	6384	6807	5711	7952	6023	6556.9	736.5
5	7000	6495	6893	6127	7417	7044	6159	6091	6653.2	503.7
6	7783	6224	5051	7288	6584	7521	6146	5129	6465.8	1034.3
7	8794	6279	6877	5807	6076	6392	7429	5220	6609.2	1104.0
8	4727	8117	6586	6225	6150	7386	5674	6740	6450.6	1033.0
9	5408	7452	6686	6428	6425	7380	5789	6264	6479.0	704.7
10	5598	7489	6186	5837	6769	5471	5658	6393	6175.1	690.5
11	6559	5855	4928	5897	7532	5663	4746	7879	6132.4	1128.6
12	6824	7320	5331	6204	6027	5987	6033	6177	6237.9	596.6
13	6503	8213	5417	6360	6711	6907	6625	7888	6828.0	879.8
14	5622	6321	6325	6634	5075	6209	4832	6386	5925.5	667.8
15	6269	6756	7653	6065	5835	7337	6615	8181	6838.9	819.5

has taken from its records a random sample of 8 such patients each month for 15 months.  DRG2

- (a) Make an s control chart using center lines and limits calculated from these past data. There are no points out of control.
- (b) Because the s chart is in control, base the \bar{x} chart on all 15 samples. Make this chart. Is it also in control?

27.23 A cutting operation. A machine tool in your plant is cutting an outside diameter. A sample of 4 pieces is taken near the end of each hour of production. Table 27.7

TABLE 27.7 \bar{x} and s for 21 samples of outside diameter

SAMPLE	\bar{x}	s	SAMPLE	\bar{x}	s
1	-0.14	0.48	12	0.55	0.10
2	0.09	0.26	13	0.50	0.25
3	0.17	0.24	14	0.37	0.45
4	0.08	0.38	15	0.69	0.21
5	-0.17	0.50	16	0.47	0.34
6	0.36	0.26	17	0.56	0.42
7	0.30	0.39	18	0.78	0.08
8	0.19	0.31	19	0.75	0.32
9	0.48	0.13	20	0.49	0.23
10	0.29	0.13	21	0.79	0.12
11	0.48	0.25			

gives \bar{x} and s for the first 21 samples, coded in units of 0.0001 inch from the center of the specifications. The specifications allow a range of ± 0.0002 inch about the center (a range of -2 to $+2$ as coded).

- (a) Make an s chart based on past data and comment on control of short-term process variation.
- (b) Because the data are coded about the center of the specs, we have a given target $\mu = 0$ (as coded) for the process mean. Make an \bar{x} chart and comment on control of long-term process variation. What special \bar{x} -type cause probably explains the lack of control of \bar{x} ?  CUTTING

27.24 The Boston Marathon. The Boston Marathon has been run each year since 1897. Winning times were highly variable in the early years, but control improved as the best runners became more professional. A clear downward trend continued until the 1980s. Rick plans to make a control chart for the winning times from 1950 to the present. The first few times are 153, 148, 152, 139, 141, and 138 minutes. Calculation from the winning times from 1950 to 2011 gives

$$\bar{x} = 134.032 \text{ minutes} \quad \text{and} \quad s = 6.462 \text{ minutes}$$

Rick draws a center line at \bar{x} and control limits at $\bar{x} \pm 3s$ for a plot of individual winning times. Explain carefully why these control limits are too wide to effectively signal unusually fast or slow times.

COMMENTS ON STATISTICAL CONTROL

Having seen how \bar{x} and s (or \bar{x} and R) charts work, we can turn to some important comments and cautions about statistical control in practice.

Focus on the process rather than on the products. This is a fundamental idea in statistical process control. We might attempt to attain high quality by careful inspection of the finished product, measuring every completed forging and reviewing every outgoing invoice and expense account payment. Inspection of finished products can ensure good quality, but it is expensive. Perhaps more important, final inspection comes too late: when something goes wrong early in a process, much bad product may be produced before final inspection discovers the problem. This adds to the expense, because the bad product must then be scrapped or reworked.

The small samples that are the basis of control charts are intended to monitor the process at key points, not to ensure the quality of the particular items in the samples. If the process is kept in control, we know what to expect in the finished product. We want to do it right the first time, not inspect and fix finished product.

Rational subgroups. The interpretation of control charts depends on the distinction between \bar{x} -type special causes and s -type special causes. This distinction in turn depends on how we choose the samples from which we calculate s (or R). We want the variation *within* a sample to reflect only the item-to-item chance variation that (when in control) results from many small common causes. Walter Shewhart, the founder of statistical process control, used the term **rational subgroup** to emphasize that we should think about the process when deciding how to choose samples.

EXAMPLE 27.9 Random sampling versus rational subgroups

A pharmaceutical manufacturer forms tablets by compressing a granular material that contains the active ingredient and various fillers. To monitor the compression process, we will measure the hardness of a sample from each 10 minutes' production of tablets. Should we choose a random sample of tablets from the several thousand produced in a 10-minute period?

A random sample would contain tablets spread across the entire 10 minutes. It fairly represents the 10-minute period, but that isn't what we want for process control. If the setting of the press drifts or a new lot of filler arrives during the 10 minutes, the spread of the sample will be increased. That is, a random sample contains both the short-term variation among tablets produced in quick succession and the longer-term variation among tablets produced minutes apart. We prefer to measure a rational subgroup of 5 consecutive tablets every 10 minutes. We expect the process to be stable during this very short time period, so that variation within the subgroups is a benchmark against which we can see special cause variation. ■

Samples of consecutive items are rational subgroups when we are monitoring the output of a single activity that does the same thing over and over again. Several consecutive items is the most common type of sample for process control. There is no formula for choosing samples that are rational subgroups. You must think about causes of variation in your process and decide which you are willing to think of as common causes that you will not try to eliminate. Rational subgroups are samples chosen to express variation due to these causes and no others. Because the choice requires detailed process knowledge, we will usually accept samples of consecutive items as being rational subgroups.

Why statistical control is desirable. To repeat, if the process is kept in control, we know what to expect in the finished product. The process mean μ and standard deviation σ remain stable over time, so (assuming Normal variation) the 99.7 part of the 68–95–99.7 rule tells us that almost all measurements on individual products will lie in the range $\mu \pm 3\sigma$. These are sometimes called the **natural tolerances** for the product. Be careful to distinguish $\mu \pm 3\sigma$, the range we expect for individual measurements, from the \bar{x} chart control limits $\mu \pm 3\sigma/\sqrt{n}$, which mark off the expected range of sample means.

natural tolerances



EXAMPLE 27.10 Natural tolerances for mesh tension

The process of setting the mesh tension on computer monitors has been operating in control. The \bar{x} and s charts were based on $\mu = 275$ mV and $\sigma = 43$ mV. The s chart in Figure 27.7 and your calculation in Exercise 27.21 suggest that the process σ is now less than 43 mV. We may prefer to calculate the natural tolerances from the recent data on 20 samples (80 monitors) in Table 27.1 (page 27-12). The estimate of the mean is $\bar{\bar{x}} = 275.065$, very close to the target value.

Now a subtle point arises. The estimate $\hat{\sigma} = \bar{s}/c_4$ used for past-data control charts is based entirely on variation *within the samples*. That's what we want for control charts, because within-sample variation is likely to be "pure common cause" variation. Even when the process is in control, there is some additional variation from sample to

sample, just by chance. So the variation in the process output will be greater than the variation within samples. To estimate the natural tolerances, we should estimate σ from all 80 individual monitors rather than by averaging the 20 within-sample standard deviations. The standard deviation for all 80 mesh tensions is

$$s = 38.38$$

(For a sample of size 80, c_4 is very close to 1, so we can ignore it.)

We are therefore confident that almost all individual monitors will have mesh tension

$$\bar{x} \pm 3s = 275.065 \pm (3)(38.38) = 275 \pm 115$$

We expect mesh tension measurements to vary between 160 and 390 mV. You see that the spread of individual measurements is wider than the spread of sample means used for the control limits of the \bar{x} chart. ■

The natural tolerances in Example 27.10 depend on the fact that the mesh tensions of individual monitors follow a Normal distribution. We know that the process was in control when the 80 measurements in Table 27.1 were made, so we can graph them to assess Normality.

APPLY YOUR KNOWLEDGE

27.25 No incoming inspection. The computer makers who buy monitors require that the monitor manufacturer practice statistical process control and submit control charts for verification. This allows the computer makers to eliminate inspection of monitors as they arrive, a considerable cost saving. Explain carefully why incoming inspection can safely be eliminated.

27.26 Natural tolerances. Table 27.6 (page 27-31) gives data on hospital losses for samples of DRG 209 patients. The distribution of losses has been stable over time. What are the natural tolerances within which you expect losses on nearly all such patients to fall? 

27.27 Normality? Do the losses on the 120 individual patients in Table 27.6 appear to come from a single Normal distribution? Make a graph and discuss what it shows. Are the natural tolerances you found in the previous exercise trustworthy? 

DON'T CONFUSE CONTROL WITH CAPABILITY!

A process in control is stable over time. We know how much variation the finished product will show. Control charts are, so to speak, the voice of the process

 telling us what state it is in. *There is no guarantee that a process in control produces products of satisfactory quality.* “Satisfactory quality” is measured by comparing the product to some standard outside the process, set by technical specifications, customer expectations, or the goals of the organization. These external standards are unrelated to the internal state of the process, which is all that statistical control pays attention to.

CAPABILITY

Capability refers to the ability of a process to meet or exceed the requirements placed on it.

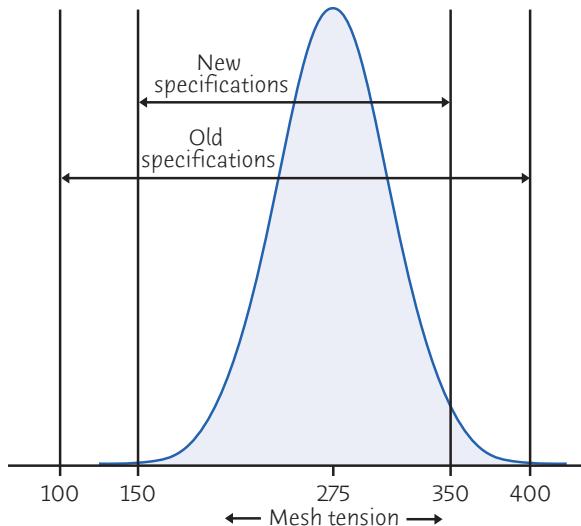
Capability has nothing to do with control—except for the very important point that if a process is not in control, it is hard to tell if it is capable or not.

EXAMPLE 27.11 Capability

The primary customer for our monitors is a large maker of computers. The customer informed us that adequate image quality requires that the mesh tension lie between 100 and 400 mV. Because the mesh-tensioning process is in control, we know (Example 27.10) that almost all monitors will have mesh tension between 160 and 390 mV. The process is capable of meeting the customer's requirement.

Figure 27.14 compares the distribution of mesh tension for individual monitors with the customer's specifications. The distribution of tension is approximately Normal, and we estimate its mean to be very close to 275 mV and the standard deviation to be about 38.4 mV. The distribution is safely within the specifications.

Times change, however. As computer buyers demand better screen quality, the computer maker restudies the effect of mesh tension and decides to require that tension lie between 150 and 350 mV. These new specification limits also appear in Figure 27.14. The process is not capable of meeting the new requirements. The process remains in control. The change in its capability is entirely due to a change in external requirements. ■

**FIGURE 27.14**

Comparison of the distribution of mesh tension (Normal curve) with original and tightened specifications, for Example 27.11. The process in its current state is not capable of meeting the new specifications.

Because the mesh-tensioning process is in control, we know that it is not capable of meeting the new specifications. That's an advantage of control, but the fact remains that control does not guarantee capability. *If a process that is in control does not have adequate capability, fundamental changes in the process are needed.* The process is doing as well as it can and displays only the chance variation that is natural to its present state. Better training for workers, new equipment, or more uniform material may improve capability, depending on the findings of a careful investigation.



APPLY YOUR KNOWLEDGE

- 27.28 Describing capability.** If the mesh tension of individual monitors follows a Normal distribution, we can describe capability by giving the percent of monitors that meet specifications. The old specifications for mesh tension are 100 to 400 mV. The new specifications are 150 to 350 mV. Because the process is in control, we can estimate that tension has mean 275 mV and standard deviation 38.4 mV.
- What percent of monitors meet the old specifications?
 - What percent meet the new specifications?

- 27.29 Improving capability.** The center of the specifications for mesh tension in the previous exercise is 250 mV, but the center of our process is 275 mV. We can improve capability by adjusting the process to have center 250 mV. This is an easy adjustment that does not change the process variation. What percent of monitors now meet the new specifications?

- 27.30 Mounting-hole distances.** Figure 27.10 (page 27-23) displays a record sheet for 18 samples of distances between mounting holes in an electrical meter. The data file adds \bar{x} and s for each sample. In Exercise 27.16, you found that Sample 5 was out of control on the process-monitoring s chart. The special cause responsible was found and removed. Based on the 17 samples that were in control, what are the natural tolerances for the distance between the holes?  MOUNTING HOLES

- 27.31 Mounting-hole distances, continued.** The record sheet in Figure 27.10 gives the specifications as 0.6054 ± 0.0010 inch. That's 54 ± 10 as the data are coded on the record sheet. Assuming that the distance varies Normally from meter to meter, about what percent of meters meet the specifications?

CONTROL CHARTS FOR SAMPLE PROPORTIONS

We have considered control charts for just one kind of data: measurements of a quantitative variable in some meaningful scale of units. We describe the distribution of measurements by its center and spread and use \bar{x} and s or \bar{x} and R charts for process control. There are control charts for other statistics that are appropriate for other kinds of data. The most common of these is the p chart for use when the data are proportions.

***p* CHART**

A p chart is a control chart based on plotting sample proportions \hat{p} from regular samples from a process against the order in which the samples were taken.

EXAMPLE 27.12 p chart settings

Here are two examples of the usefulness of p charts:

- Measure two dimensions of a manufactured part and also grade its surface finish by eye. The part conforms if both dimensions lie within their specifications and the finish is judged acceptable. Otherwise, it is nonconforming. Plot the proportion of nonconforming parts in samples of parts from each shift.
- An urban school system records the percent of its eighth-grade students who are absent three or more days each month. Because students with high absenteeism in eighth grade often fail to complete high school, the school system has launched programs to reduce absenteeism. These programs include calls to parents of absent students, public-service messages to change community expectations, and measures to ensure that the schools are safe and attractive. A p chart will show if the programs are having an effect. ■

The manufacturing example illustrates an advantage of p charts: they can combine several specifications in a single chart. Nonetheless, p charts have been rendered outdated in many manufacturing applications by improvements in typical levels of quality. For example, Delphi, the largest North American auto electronics manufacturer, says that it reduced its proportion of problem parts from 200 per million in 1997 to 20 per million in 2001.¹⁴ At either of these levels, even large samples of parts will rarely contain any bad parts. The sample proportions will almost all be 0, so that plotting them is uninformative. It is better to choose important measured characteristics—voltage at a critical circuit point, for example—and keep \bar{x} and s charts. Even if the voltage is satisfactory, quality can be improved by moving it yet closer to the exact voltage specified in the design of the part.

The school absenteeism example is a management application of p charts. More than 20% of all American eighth-graders miss three or more days of school per month, and this proportion is higher in large cities. A p chart will be useful. Proportions of “things going wrong” are often higher in business processes than in manufacturing, so that p charts are an important tool in business.

CONTROL LIMITS FOR p CHARTS

We studied the sampling distribution of a sample proportion \hat{p} in Chapter 20. The center line and control limits for a 3σ control chart follow directly from the facts stated there, in the box on text page 494. We ought to call such charts “ \hat{p} charts” because they plot sample proportions. Unfortunately, they have always been called p charts in quality control circles. We will keep the traditional name but also keep our usual notation: p is a process proportion and \hat{p} is a sample proportion.

***p* CHART USING PAST DATA**

Take regular samples from a process that has been in control. Estimate the process proportion \bar{p} of “successes” by

$$\bar{p} = \frac{\text{total number of successes in past samples}}{\text{total number of individuals in these samples}}$$

The center line and control limits for a ***p* chart** for future samples of size n are

$$\text{UCL} = \bar{p} + 3\sqrt{\frac{\bar{p}(1 - \bar{p})}{n}}$$

$$\text{CL} = \bar{p}$$

$$\text{LCL} = \bar{p} - 3\sqrt{\frac{\bar{p}(1 - \bar{p})}{n}}$$

Common **out-of-control signals** are one sample proportion \hat{p} outside the control limits or a run of 9 sample proportions on the same side of the center line.

If we have k past samples of the *same* size n , then \bar{p} is just the average of the k sample proportions. In some settings, you may meet samples of unequal size—differing numbers of students enrolled in a month or differing numbers of parts inspected in a shift. The average \bar{p} estimates the process proportion p even when the sample sizes vary. Note that the control limits use the actual size n of a sample.

**ABSENTEEISM****EXAMPLE 27.13 Reducing absenteeism**

Unscheduled absences by clerical and production workers are an important cost in many companies. You have been asked to improve absenteeism in a production facility where 12% of the workers are now absent on a typical day.

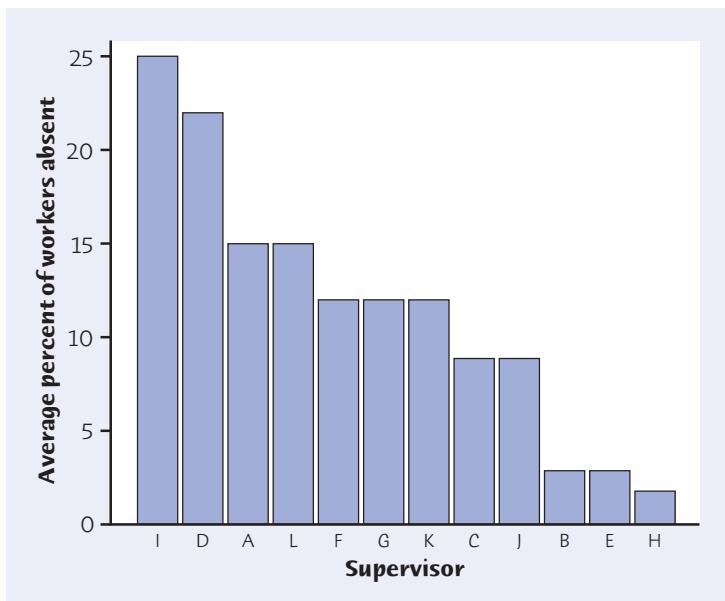
Start with data: the Pareto chart in Figure 27.15 shows that there are major differences among supervisors in the absenteeism rate of their workers. You retrain all the supervisors in human relations skills, using B, E, and H as discussion leaders. In addition, a trainer works individually with supervisors I and D. You also improve lighting and other work conditions.

Are your actions effective? You hope to see a reduction in absenteeism. To view progress (or lack of progress), you will keep a *p* chart of the proportion of absentees. The plant has 987 production workers. For simplicity, you just record the number who are absent from work each day. Only unscheduled absences count, not planned time off such as vacations. Each day you will plot

$$\hat{p} = \frac{\text{number of workers absent}}{987}$$

You first look back at data for the past three months. There were 64 workdays in these months. The total of workdays available for the workers was

$$(64)(987) = 63,168 \text{ person-days}$$

**FIGURE 27.15**

Pareto chart of the average absenteeism rate for workers reporting to each of 12 supervisors, for Example 27.13.

Absences among all workers totaled 7580 person-days. The average daily proportion absent was therefore

$$\bar{p} = \frac{\text{total days absent}}{\text{total days available for work}}$$

$$= \frac{7580}{63,168} = 0.120$$

The daily rate has been in control at this level.

These past data allow you to set up a p chart to monitor future proportions absent:

$$\begin{aligned} \text{UCL} &= \bar{p} + 3\sqrt{\frac{\bar{p}(1 - \bar{p})}{n}} = 0.120 + 3\sqrt{\frac{(0.120)(0.880)}{987}} \\ &= 0.120 + 0.031 = 0.151 \\ \text{CL} &= \bar{p} = 0.120 \\ \text{LCL} &= \bar{p} - 3\sqrt{\frac{\bar{p}(1 - \bar{p})}{n}} = 0.120 - 3\sqrt{\frac{(0.120)(0.880)}{987}} \\ &= 0.120 - 0.031 = 0.089 \end{aligned}$$

Table 27.8 gives the data for the next four weeks. Figure 27.16 is the p chart. ■

Figure 27.16 shows a clear downward trend in the daily proportion of workers who are absent. Days 13 and 19 lie below LCL, and a run of 9 days below the center line is achieved at Day 15 and continues. The points marked “x” are therefore all out of control. It appears that a special cause (the various actions you took) has reduced the absenteeism rate from around 12% to around 10%. The data for the last two weeks suggest that the rate has stabilized at this level. You will update

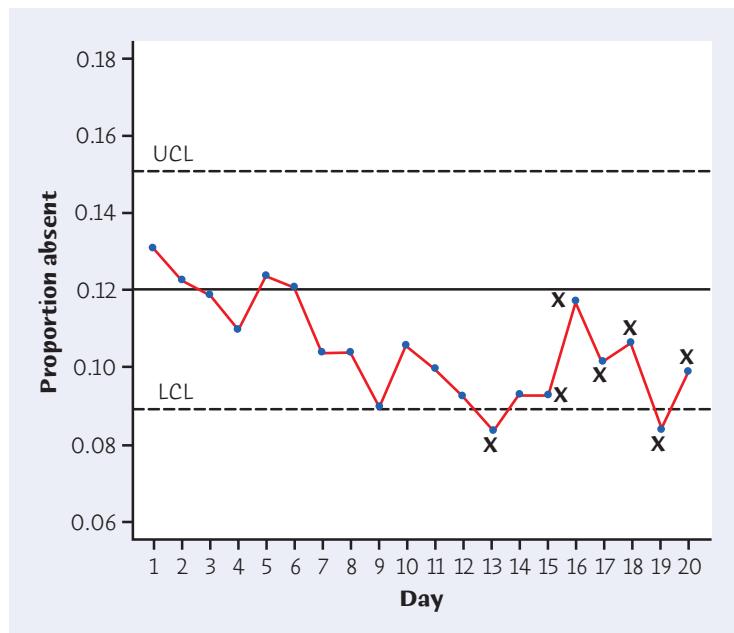
TABLE 27.8 Proportions of workers absent during four weeks

	M	T	W	Th	F	M	T	W	Th	F
Workers absent	129	121	117	109	122	119	103	103	89	105
Proportion \hat{p}	0.131	0.123	0.119	0.110	0.124	0.121	0.104	0.104	0.090	0.106
	M	T	W	Th	F	M	T	W	Th	F
Workers absent	99	92	83	92	92	115	101	106	83	98
Proportion \hat{p}	0.100	0.093	0.084	0.093	0.093	0.117	0.102	0.107	0.084	0.099

the chart based on the new data. If the rate does not decline further (or even rises again as the effect of your actions wears off), you will consider further changes.

Example 27.13 is a bit oversimplified. The number of workers available did not remain fixed at 987 each day. Hirings, resignations, and planned vacations change the number a bit from day to day. The control limits for a day's \hat{p} depend on n , the number of workers that day. If n varies, the control limits will move in and out from day to day. Software will do the extra arithmetic needed for a different n each day, but as long as the count of workers remains close to 987, the greater detail will not change your conclusion.

A single p chart for all workers is not the only, or even the best, choice in this setting. Because of the important role of supervisors in absenteeism, it would be wise to also keep separate p charts for the workers under each supervisor. These charts may show that you must reassess some supervisors.

**FIGURE 27.16**

p chart for daily proportion of workers absent over a four-week period, for Example 27.13. The lack of control shows an improvement (decrease) in absenteeism. Update the chart to continue monitoring the process.


APPLY YOUR KNOWLEDGE

27.32 Setting up a p chart. After inspecting Figure 27.16, you decide to monitor the next four weeks' absenteeism rates using a center line and control limits calculated from the last two weeks of data recorded in Table 27.8. Find \bar{p} for these 10 days and give the new values of CL, LCL, and UCL. (Until you have more data, these are trial control limits. As long as you are taking steps to improve absenteeism, you have not reached the process-monitoring stage.)  **ABSENTEEISM**

27.33 Unpaid invoices. The controller's office of a corporation is concerned that invoices that remain unpaid after 30 days are damaging relations with vendors. To assess the magnitude of the problem, a manager searches payment records for invoices that arrived in the past 10 months. The average number of invoices is 2875 per month, with relatively little month-to-month variation. Of all these invoices, 960 remained unpaid after 30 days.

- What is the total number of invoices studied? What is \bar{p} ?
- Give the center line and control limits for a p chart on which to plot the future monthly proportions of unpaid invoices.

27.34 Lost baggage. The Department of Transportation reports that about 1 of every 208 passengers on domestic flights of the 18 largest U.S. airlines files a report of mishandled baggage. Starting with this information, you plan to sample records for 1000 passengers per day at a large airport to monitor the effects of efforts to reduce mishandled baggage. What are the initial center line and control limits for a chart of the daily proportion of mishandled-baggage reports? (You will find that $LCL < 0$. Because proportions \hat{p} are always 0 or positive, take $LCL = 0$.)

27.35 Aircraft rivets. After completion of an aircraft wing assembly, inspectors count the number of missing or deformed rivets. There are hundreds of rivets in each wing, but the total number varies depending on the aircraft type. Recent data for wings with a total of 37,200 rivets show 220 missing or deformed. The next wing contains 1070 rivets. What are the appropriate center line and control limits for plotting the \hat{p} from this wing on a p chart?

27.36 School absenteeism. Here are data from an urban school district on the number of eighth-grade students with three or more unexcused absences from school during each month of a school year. Because the total number of eighth-graders changes a bit from month to month, these totals are also given for each month.

	Sept.	Oct.	Nov.	Dec.	Jan.	Feb.	Mar.	Apr.	May	June
Students	811	847	839	842	818	820	831	825	802	783
Absent	251	309	324	295	261	282	304	284	263	304

- Find \bar{p} . Because the number of students varies from month to month, also find \bar{n} , the average per month.
- Make a p chart using control limits based on \bar{n} students each month. Comment on control.
- The exact control limits are different each month because the number of students n is different each month. This situation is common in using p charts. What are the exact limits for October and June, the months with the largest and smallest n ? Add these limits to your p chart, using short lines spanning a single month. Do exact limits affect your conclusions?  **ABSENTEEISM2**

CHAPTER 27 SUMMARY

CHAPTER SPECIFICS

- Work is organized in **processes**, chains of activities that lead to some result. Use **flowcharts** and **cause-and-effect diagrams** to describe processes. Other graphs such as **Pareto charts** are often useful.
- All processes have variation. If the pattern of variation is stable over time, the process is **in statistical control**. **Control charts** are statistical plots intended to warn when a process is **out of control**.
- Standard 3σ **control charts** plot the values of some statistic Q for regular samples from the process against the time order of the samples. The **center line** is at the mean of Q . The **control limits** lie three standard deviations of Q above and below the center line. A point outside the control limits is an **out-of-control signal**. For **process monitoring** of a process that has been in control, the mean and standard deviation are based on past data from the process and are updated regularly.
- When we measure some quantitative characteristic of the process, we use **\bar{x} and s charts** for process control. The s chart monitors variation within individual samples. If the s chart is in control, the \bar{x} chart monitors variation from sample to sample. To interpret the charts, always look first at the s chart.
- An **R chart** based on the **range** of observations in a sample is often used in place of an s chart. Interpret \bar{x} and R charts exactly as you would interpret \bar{x} and s charts.
- It is common to use **out-of-control signals** in addition to “one point outside the control limits.” In particular, a **runs signal** for the \bar{x} chart allows the chart to respond more quickly to a gradual drift in the process center.
- **Control charts based on past data** are used at the **chart setup** stage for a process that may not be in control. Start with control limits calculated from the same past data that you are plotting. Beginning with the s chart, narrow the limits as you find special causes, and remove the points influenced by these causes. When the remaining points are in control, use the resulting limits to monitor the process.
- Statistical process control maintains quality more economically than inspecting the final output of a process. Samples that are **rational subgroups** are important to effective control charts. A process in control is stable, so that we can predict its behavior. If individual measurements have a Normal distribution, we can give the **natural tolerances**.
- A process is **capable** if it can meet the requirements placed on it. Control (stability over time) does not in itself improve capability. Remember that control describes the internal state of the process, whereas capability relates the state of the process to external specifications.
- There are control charts for several different types of process measurements. One important type is the **p chart** for sample proportions \hat{p} .
- The interpretation of p charts is very similar to that of \bar{x} charts. The out-of-control signals used are also the same.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Processes

1. Describe the process leading to some desired output using flowcharts and cause-and-effect diagrams.
2. Choose promising targets for process improvement, combining the process description with data collection and tools such as Pareto charts.
3. Demonstrate understanding of statistical control, common causes, and special causes by applying these ideas to specific processes.
4. Choose rational subgroups for control charting based on an understanding of the process.

B. Control Charts

1. Make \bar{x} and s charts using given values of the process μ and σ (usually from large amounts of past data) for monitoring a process that has been in control.
2. Demonstrate understanding of the distinction between short-term (within sample) and longer-term (across samples) variation by identifying possible \bar{x} -type and s -type special causes for a specific process.
3. Interpret \bar{x} and s charts, starting with the s chart. Use both one-point-out and runs signals.
4. Estimate the process μ and σ from recent samples.
5. Set up initial control charts using recent process data, removing special causes, and basing an initial chart on the remaining data.
6. Decide when a p chart is appropriate. Make a p chart based on past data.

C. Process Capability

1. Know the distinction between control and capability and apply this distinction in discussing specific processes.
2. Give the natural tolerances for a process in control, after verifying Normality of individual measurements on the process.

LINK IT

In this chapter we apply and extend many of the ideas discussed previously to understanding processes. Flowcharts, cause-and-effect diagrams, Pareto charts, and control charts are graphical displays for understanding processes, extending the tools discussed in Part I. Choosing rational subgroups reminds one of sample design, discussed in Part II. Evaluating capability of a process to meet or exceed the requirement placed upon it reminds one of Normal probability calculations, discussed in Part II. The use of control limits to determine whether a process is in control reminds one of hypothesis testing for means and proportions, discussed in Parts III and IV. Thus, we see many of the topics discussed previously applied in new ways to help us evaluate processes.

CHECK YOUR SKILLS

27.37 A maker of auto air conditioners checks a sample of 4 thermostatic controls from each hour's production. The thermostats are set at 75°F and then placed in a chamber where the temperature is raised gradually. The temperature at which the thermostat turns on the air conditioner is recorded. The process mean should be $\mu = 75^\circ\text{F}$. Past experience indicates that the response temperature of properly adjusted thermostats varies with $\sigma = 0.5^\circ\text{F}$. The mean response temperature \bar{x} for each hour's sample is plotted on an \bar{x} control chart. The center line for this chart is

- (a) 0.5°F. (b) 75°F. (c) 4.

27.38 In Exercise 27.37, the UCL for the chart is

- (a) 78.0°F. (b) 76.5°F. (c) 75.75°F.

27.39 In Exercise 27.37 suppose the standard deviation s for each hour's sample is plotted on an s control chart. The center line for this chart is

- (a) 0.5°F. (b) 0.46°F. (c) 0.19°F.

27.40 In Exercise 27.37 suppose the standard deviation s for each hour's sample is plotted on an s control chart. The LCL for this chart is

- (a) 0°F. (b) 0.5°F. (c) 1.04°F.

27.41 Samples of size 4 are taken from a manufacturing process every hour. A quality characteristic is measured and \bar{x} and s are measured for each sample. The process is in control and after 25 samples we compute $\bar{\bar{x}} = 48.4$ and $\bar{s} = 5.1$. The LCL for an \bar{x} control chart based on these data is

- (a) 48.4. (b) 33.1. (c) 31.8.

27.42 A process produces rubber fan belts for automobiles. The process is in control and 100 belts are inspected each day for a period of 20 days. The proportion of nonconforming belts found over this 20-day period is $\bar{p} = 0.12$. Based on these data, a p chart for future samples of size 100 would have center line

- (a) 0.12. (b) 12.0. (c) 240.

27.43 Referring to the previous exercise, the UCL for a p chart for future samples of size 100 would be

- (a) 0.22. (b) 0.15. (c) 0.12.

CHAPTER 27 EXERCISES

27.44 Enlighten management. A manager who knows no statistics asks you, "What does it mean to say that a process is in control? Is being in control a guarantee that the quality of the product is good?" Answer these questions in plain language that the manager can understand.

27.45 Special causes. Is each of the following examples of a special cause most likely to first result in (i) a sudden change in level on the s or R chart, (ii) a sudden change in level on the \bar{x} chart, or (iii) a gradual drift up or down on the \bar{x} chart? In each case, briefly explain your reasoning.

(a) An airline pilots' union puts pressure on management during labor negotiations by asking its members to "work to rule" in doing the detailed checks required before a plane can leave the gate.

(b) Measurements of part dimensions that were formerly made by hand are now made by a very accurate laser



David Frazier/Stone/Getty Images

system. (The process producing the parts does not change—measurement methods can also affect control charts.)

(c) Inadequate air conditioning on a hot day allows the temperature to rise during the afternoon in an office that prepares a company's invoices.

27.46 Deming speaks. The quality guru W. Edwards Deming (1900–1993) taught (among much else) that¹⁵

(a) "People work in the system. Management creates the system."

(b) "Putting out fires is not improvement. Finding a point out of control, finding the special cause and removing it, is only putting the process back to where it was in the first place. It is not improvement of the process."

(c) "Eliminate slogans, exhortations and targets for the workforce asking for zero defects and new levels of productivity. Choose one of these sayings. Explain carefully what facts about improving quality the saying attempts to summarize."

27.47 Pareto charts. You manage the customer service operation for a maker of electronic equipment sold to business

customers. Traditionally, the most common complaint is that equipment does not operate properly when installed, but attention to manufacturing and installation quality will reduce these complaints. You hire an outside firm to conduct a sample survey of your customers. Here are the percents of customers with each of several kinds of complaints:



Category	Percent
Accuracy of invoices	27
Clarity of operating manual	6
Complete invoice	25
Complete shipment	16
Correct equipment shipped	15
Ease of obtaining invoice adjustments/credits	34
Equipment operates when installed	5
Meeting promised delivery date	11
Sales rep returns calls	3
Technical competence of sales rep	12

- (a) Why do the percents not add to 100%?
- (b) Make a Pareto chart. What area would you choose as a target for improvement?

27.48 What type of chart? What type of control chart or charts would you use as part of efforts to improve each of the following performance measures in a college admissions office? Explain your choices.

- (a) Time to acknowledge receipt of an application.
- (b) Percent of admission offers accepted.
- (c) Student participation in a healthy meal plan.

27.49 What type of chart? What type of control chart or charts would you use as part of efforts to improve each of the following performance measures in an online business information systems department? Explain your choices.

- (a) Web site availability.
- (b) Time to respond to requests for help.
- (c) Percent of Web site changes not properly documented.

27.50 Purchased material. At the present time, about 4 lots out of every 1000 lots of material arriving at a plant site from outside vendors are rejected because they are incorrect. The plant receives about 250 lots per week. As part of an effort to reduce errors in the system of placing and filling orders, you will monitor the proportion of rejected lots each week. What type of control chart will you use? What are the initial center line and control limits?

27.51 Pareto charts. Painting new auto bodies is a multi-step process. There is an “electrocoat” that resists corrosion, a primer, a color coat, and a gloss coat. A quality study for one paint shop produced this breakdown of the primary problem type for those autos whose paint did not meet the manufacturer’s standards:



Problem	Percent
Electrocoat uneven—redone	4
Poor adherence of color to primer	5
Lack of clarity in color	2
“Orange peel” texture in color	32
“Orange peel” texture in gloss	1
Ripples in color coat	28
Ripples in gloss coat	4
Uneven color thickness	19
Uneven gloss thickness	5
Total	100

Make a Pareto chart. Which stage of the painting process should we look at first?

27.52 Piston rings. The inside diameter of automobile engine piston rings is important to the proper functioning of the engine. The manufacturer checks the control of the piston ring forging process by measuring a sample of 5 consecutive items during each hour’s production. The target diameter for a ring is $\mu = 74.000$ millimeters. The process has been operating in control with center close to the target and $\sigma = 0.015$ millimeters. What center line and control limits should be drawn on the s chart? On the \bar{x} chart?



iStockphoto

27.53 p charts are out of date. A manufacturer of consumer electronic equipment makes full use not only of statistical process control but of automated testing equipment that efficiently tests all completed products. Data from the testing equipment show that finished products have only 3.0 defects per million opportunities.

- (a) What is \bar{p} for the manufacturing process? If the process turns out 4000 pieces per day, how many defects do you expect to see per day? In a typical month of 24 working days, how many defects do you expect to see?
- (b) What are the center line and control limits for a p chart for plotting daily defect proportions?

- (c) Explain why a p chart is of no use at such high levels of quality.

27.54 Manufacturing isn't everything. Because the manufacturing quality in the previous exercise is so high, the process of writing up orders is the major source of quality problems: the defect rate there is 9000 per million opportunities. The manufacturer processes about 600 orders per month.

- (a) What is \bar{p} for the order-writing process? How many defective orders do you expect to see in a month?

- (b) What are the center line and control limits for a p chart for plotting monthly proportions of defective orders? What is the smallest number of bad orders in a month that will result in a point above the upper control limit?

Table 27.9 gives process control samples for a study of response times to customer calls arriving at a corporate call center. A sample of 6 calls is recorded each shift for quality improvement purposes. The time from the first ring until a representative answers the call is recorded. Table 27.9 gives data for 50 shifts, 300 calls total.¹⁶ Exercises 27.55 to 27.57 make use of this setting.

TABLE 27.9 Fifty control chart samples of call center response times (seconds)

SAMPLE	TIME						SAMPLE MEAN	STANDARD DEVIATION
1	59	13	2	24	11	18	21.2	19.93
2	38	12	46	17	77	12	33.7	25.56
3	46	44	4	74	41	22	38.5	23.73
4	25	7	10	46	78	14	30.0	27.46
5	6	9	122	8	16	15	29.3	45.57
6	17	17	9	15	24	70	25.3	22.40
7	9	9	10	32	9	68	22.8	23.93
8	8	10	41	13	17	50	23.2	17.79
9	12	82	97	33	76	56	59.3	32.11
10	42	19	14	21	12	44	25.3	14.08
11	63	5	21	11	47	8	25.8	23.77
12	12	4	111	37	12	24	33.3	39.76
13	43	37	27	65	32	3	34.5	20.32
14	9	26	5	10	30	27	17.8	10.98
15	21	14	19	44	49	10	26.2	16.29
16	24	11	10	22	43	70	30.0	22.93
17	27	10	32	96	11	29	34.2	31.71
18	7	28	22	17	9	24	17.8	8.42
19	15	14	34	5	38	29	22.5	13.03
20	16	65	6	5	58	17	27.8	26.63
21	7	44	14	16	4	46	21.8	18.49
22	32	52	75	11	11	17	33.0	25.88
23	31	8	36	25	14	85	33.2	27.45
24	4	46	23	58	5	54	31.7	24.29
25	28	6	46	4	28	11	20.5	16.34
26	111	6	3	83	27	6	39.3	46.34
27	83	27	2	56	26	21	35.8	28.88
28	276	14	30	8	7	12	57.8	107.20
29	4	29	21	23	4	14	15.8	10.34
30	23	22	19	66	51	60	40.2	21.22
31	14	111	20	7	7	87	41.0	45.82
32	22	11	53	20	14	41	26.8	16.56

TABLE 27.9 (continued)

SAMPLE	TIME						SAMPLE MEAN	STANDARD DEVIATION
33	30	7	10	11	9	9	12.7	8.59
34	101	55	18	20	77	14	47.5	36.16
35	13	11	22	15	2	14	12.8	6.49
36	20	83	25	10	34	23	32.5	25.93
37	21	5	14	22	10	68	23.3	22.82
38	8	70	56	8	26	7	29.2	27.51
39	15	7	9	144	11	109	49.2	60.97
40	20	4	16	20	124	16	33.3	44.80
41	16	47	97	27	61	35	47.2	28.99
42	18	22	244	19	10	6	53.2	93.68
43	43	20	77	22	7	33	33.7	24.49
44	67	20	4	28	5	7	21.8	24.09
45	118	18	1	35	78	35	47.5	43.00
46	71	85	24	333	50	11	95.7	119.53
47	12	11	13	19	16	91	27.0	31.49
48	4	63	14	22	43	25	28.5	21.29
49	18	55	13	11	6	13	19.3	17.90
50	4	3	17	11	6	17	9.7	6.31

27.55 Rational subgroups? The 6 calls each shift are chosen at random from all calls received during the shift. Discuss the reasons behind this choice and those behind a choice to time 6 consecutive calls.  **RESPONSETIMES**

27.56 Chart setup. Table 27.9 also gives \bar{x} and s for each of the 50 samples.  **RESPONSETIMES2**

- (a) Make an s chart and check for points out of control.
- (b) If the s -type cause responsible is found and removed, what would be the new control limits for the s chart? Verify that no points s are now out of control.
- (c) Use the remaining 46 samples to find the center line and control limits for an \bar{x} chart. Comment on the control (or lack of control) of \bar{x} . (Because the distribution of response times is strongly skewed, \bar{s} is large and the control limits for \bar{x} are wide. Control charts based on Normal distributions often work poorly when measurements are strongly skewed.)

27.57 Using process knowledge. Three of the out-of-control values of s in part (a) of the previous exercise are explained by a single outlier, a very long response time to one call in the sample. What are the values of these outliers, and what are the s -values for the 3 samples when the outliers are omitted? (The interpretation of the data is, unfortunately, now clear. Few customers will wait 5 minutes for a call to be answered, as the customer whose call took 333 seconds to answer did. We suspect that other customers hung up before

their calls were answered. If so, response time data for the calls that were answered don't adequately picture the quality of service. We should now look at data on calls lost before being answered to see a fuller picture.)  **RESPONSETIMES2**

27.58 Doctors' prescriptions. A regional chain of retail pharmacies finds that about 2% of prescriptions it receives from doctors are incorrect or illegible. The chain puts in place a secure online system that doctors' offices can use to enter prescriptions directly. It hopes that fewer prescriptions entered online will be incorrect or illegible. A p chart will monitor progress. Use information about past prescriptions to set initial center line and control limits for the proportion of incorrect or illegible prescriptions on a day when the chain fills 80,000 online prescriptions. What are the center line and control limits for a day when only 40,000 online prescriptions are filled?

You have just installed a new system that uses an interferometer to measure the thickness of polystyrene film. To control the thickness, you plan to measure 3 film specimens every 10 minutes and keep \bar{x} and s charts. To establish control, you measure 22 samples of 3 films each at 10-minute intervals. Table 27.10 gives \bar{x} and s for these samples. The units are ten-thousandths of a millimeter. Exercises 27.59 to 27.61 are based on this chart setup setting.

27.59 s chart. Calculate control limits for s , make an s chart, and comment on control of short-term process variation.  **THICKNESS**

TABLE 27.10 \bar{x} and s for 22 samples of film thickness (in ten-thousandths of a millimeter)

SAMPLE	\bar{x}	s	SAMPLE	\bar{x}	s
1	848	20.1	12	823	12.6
2	832	1.1	13	835	4.4
3	826	11.0	14	843	3.6
4	833	7.5	15	841	5.9
5	837	12.5	16	840	3.6
6	834	1.8	17	833	4.9
7	834	1.3	18	840	8.0
8	838	7.4	19	826	6.1
9	835	2.1	20	839	10.2
10	852	18.9	21	836	14.8
11	836	3.8	22	829	6.7

27.60 \bar{x} chart. Interviews with the operators reveal that in Samples 1 and 10 mistakes in operating the interferometer resulted in one high-outlier thickness reading that was clearly incorrect. Recalculate \bar{x} and s after removing Samples 1 and 10. Recalculate UCL for the s chart and add the new UCL to your s chart from the previous exercise. Control for the remaining samples is excellent. Now find the appropriate center line and control limits for an \bar{x} chart, make the \bar{x} chart, and comment on control.  THICKNESS2

27.61 Categorizing the output. Previously, control of the process was based on categorizing the thickness of each film inspected as satisfactory or not. Steady improvement in process quality has occurred, so that just 15 of the last 5000 films inspected were unsatisfactory.

- (a) What type of control chart would be used in this setting, and what would be the control limits for a sample of 100 films?
- (b) The chart in (a) is of little practical value at current quality levels. Explain why.



EXPLORING THE WEB

27.62 Spotting a mass murderer. The Chance Web site discusses an application of statistical process control methods for spotting a mass murderer. Read the article at www.causeweb.org/wiki/chance/index.php/Chance_News_6 and the information found at the links in this article. Write a paragraph summarizing how statistical process control methods might have been used to identify a mass murderer.

27.63 Six sigma. Six Sigma is a methodology used in many companies. Search the Web to learn more about Six Sigma. Write a paragraph explaining what Six Sigma is and how it is related to material discussed in this chapter. Give a list of some companies that use Six Sigma.

NOTES AND DATA SOURCES

1. CNNMoney, "My Golden Rule," at money.cnn.com, November 2005.
2. Texts on quality management give more detail about these and other simple graphical methods for quality problems. The classic reference is Kaoru Ishikawa, *Guide to Quality Control*, Asian Productivity Organization, 1986.

3. The flowchart and a more elaborate version of the cause-and-effect diagram for Example 27.1 were prepared by S. K. Bhat of the General Motors Technical Center as part of a course assignment at Purdue University.
4. For more information and references on DRGs, see the Wikipedia entry “diagnosis-related group.” Search for this term at en.wikipedia.org.
5. Ronald J. M. Does and Thijs M. B. Vermaat, “Reducing start time delays in operating rooms,” *Journal of Quality Technology*, 41 (2009), pp. 95–109.
6. The terms “chart setup” and “process monitoring” are adopted from Andrew C. Palm’s discussion of William H. Woodall, “Controversies and contradictions in statistical process control,” *Journal of Quality Technology*, 32 (2000), pp. 341–350. Palm’s discussion appears in the same issue, pp. 356–360. We have combined Palm’s stages B (“process improvement”) and C (“process monitoring”) when writing for beginners because the distinction between them is one of degree.
7. It is common to call these “standards given” \bar{x} and s charts. We avoid this term because it easily leads to the common and serious error of confusing control limits (based on the process itself) with standards or specifications imposed from outside.
8. Provided by Charles Hicks, Purdue University.
9. See, for example, Chapter 3 of Stephen B. Vardeman and J. Marcus Jobe, *Statistical Quality Assurance Methods for Engineers*, Wiley 1999.
10. The classic discussion of out-of-control signals and the types of special causes that may lie behind special control chart patterns is the *AT&T Statistical Quality Control Handbook*, Western Electric, 1956.
11. The data in Table 27.5 are adapted from data on viscosity of rubber samples appearing in Table P3.3 of Irving W. Burr, *Statistical Quality Control Methods*, Marcel Dekker, 1976.
12. The control limits for the s chart based on past data are commonly given as $B_4\bar{s}$ and $B_3\bar{s}$. That is, $B_4 = B_6/c_4$ and $B_3 = B_5/c_4$. This is convenient for users, but avoiding this notation minimizes the number of control chart constants students must keep straight and emphasizes that process-monitoring and past-data charts are exactly the same except for the source of μ and σ .
13. Simulated data based on information appearing in Arvind Salvekar, “Application of six sigma to DRG 209,” found at the Smarter Solutions Web site, www.smartersolutions.com.
14. Micheline Maynard, “Building success from parts,” *New York Times*, March 17, 2002.
15. The first two Deming quotes are from *Public Sector Quality Report*, December 1993, p. 5. They were found online at [demqtes.txt](http://deming.eng.clemson.edu/pub/den/files/demqtes.txt). The third quote is part of the 10th of Deming’s “14 points of quality management,” from his book *Out of the Crisis*, MIT Press, 1986.
16. The data in Table 27.9 are simulated from a probability model for call pickup times. That pickup times for large financial institutions have median 20 seconds and mean 32 seconds is reported by Jon Anton, “A case study in benchmarking call centers,” Purdue University Center for Customer-Driven Quality, no date.



Multiple Regression*

When a scatterplot shows a linear relationship between a quantitative explanatory variable x and a quantitative response variable y , we fit a regression line to the data to describe the relationship. We can also use the line to predict the value of y for a given value of x . For example, Chapter 5 uses regression lines to describe relationships between

- Fat gain y and nonexercise activity x .
- The brain activity y of women when their partner has a painful experience and their score x on a test measuring empathy.
- The number y of new adults that join a colony of birds and the percent x of adult birds that return from the previous year.

In all these cases, other explanatory variables might improve our understanding of the response y and help us to better predict y .

- Fat gain y depends on nonexercise activity x_1 , time spent daily in exercise activity x_2 , and sex x_3 .
- A woman's brain activity y when her partner has a painful experience may depend on her score x , on a test of empathy and also on her score x_2 on a test of emotional attachment to her partner.
- The number y of new adults in a bird colony depends on the percent x_1 of returning adults and also on the species x_2 of birds we study.

We will now call regression with just one explanatory variable **simple linear regression** to remind us that this is a special case. This chapter introduces the more general case of **multiple regression**, which allows

*The original version of this chapter was written by Professor Bradley Hartlaub of Kenyon College.

Chapter 28

IN THIS CHAPTER WE COVER...

- Parallel regression lines
- Estimating parameters
- Using technology
- Inference for multiple regression
- Interaction
- The general multiple linear regression model
- The woes of regression coefficients
- A case study for multiple regression
- Inference for regression parameters
- Checking the conditions for inference

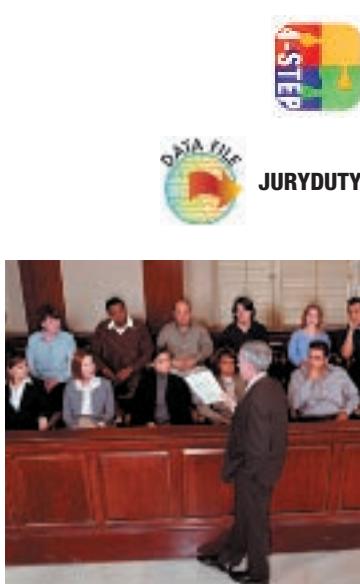
simple linear regression

multiple regression

several explanatory variables to combine in explaining a response variable. The material we discuss will help you understand and interpret the results of a multiple regression analysis. However, there are many issues that we do not discuss that are important to understand if you plan to carry out a multiple regression analysis yourself. Thus, we recommend you take a more advanced course on multiple regression if you plan to use these methods.

PARALLEL REGRESSION LINES

In Chapter 4 we learned how to add a categorical variable to a scatterplot by using different colors or plot symbols to indicate the different values of the categorical variable. Consider a simple case: the categorical variable (call it x_2) takes just two values and the scatterplot seems to show two *parallel* straight-line patterns linking the response y to a quantitative explanatory variable x_1 , one pattern for each value of x_2 . Here is an example.



Michael Kelley/Getty Images

EXAMPLE 28.1 Potential jurors

STATE: Tom Shields, jury commissioner for the Franklin County Municipal Court in Columbus, Ohio, is responsible for making sure that the judges have enough potential jurors to conduct jury trials. Only a small percent of cases go to trial, but potential jurors must be available to serve on short notice. Jury duty for this court is two weeks long, so Tom must bring together a new group of potential jurors twenty-six times a year. Random sampling methods are used to obtain a sample of registered voters in Franklin County every two weeks, and these individuals are sent a summons to appear for jury duty. Not all of the voters who receive a summons actually appear for jury duty. Table 28.1 shows the percent of individuals who reported for jury duty after receiving a summons for two years, 1998 and 2000.¹ The reporting dates vary slightly from year to year, so they are coded in order from 1, the first group to report in January, to 26, the last group to report in December. New efforts were made to increase participation rates in 2000. Is there evidence that these efforts were successful?

PLAN: Make a scatterplot to display the relationship between percent reporting y and reporting date x_1 . Use different colors for the two years. (So year is a categorical variable x_2 that takes two values.) If both years show linear patterns, fit two separate least-squares regression lines to describe them.

SOLVE: Figure 28.1 shows a scatterplot with two separate regression lines, one for 1998 and one for 2000. The slopes of both regression lines are negative, indicating lower participation for those individuals selected later in the year. But the reporting percents in 2000 are higher than the corresponding percents for all but one group in 1998. Software gives these regression lines:

$$\text{For 2000: } \hat{y} = 95.571 - 0.765x_1$$

$$\text{For 1998: } \hat{y} = 76.426 - 0.668x_1$$

The intercepts for the two regression lines are very different, but the slopes are roughly the same. Since the two regression lines are roughly parallel, the difference in the two intercepts gives us an indication of how much better the reporting percents were in 2000

TABLE 28.1 Percents of randomly selected registered voters who appeared for jury duty in Franklin County Municipal Court in 1998 and 2000

REPORTING DATE	1998	2000	REPORTING DATE	1998	2000
1	83.30	92.59	14	65.40	94.40
2	83.60	81.10	15	65.02	88.50
3	70.50	92.50	16	62.30	95.50
4	70.70	97.00	17	62.50	65.90
5	80.50	97.00	18	65.50	87.50
6	81.60	83.30	19	63.50	80.20
7	65.30	94.60	20	75.00	94.70
8	61.30	88.10	21	67.90	76.60
9	62.70	90.90	22	62.00	75.80
10	67.80	87.10	23	71.00	76.50
11	65.00	85.40	24	62.10	80.60
12	64.10	86.60	25	58.50	71.80
13	64.70	88.30	26	50.70	63.70

after taking into account the reporting date. We will soon learn how to formally estimate parameters and make inferences for parallel regression lines. However, our separate regression models clearly indicate an important change in the reporting percents.

CONCLUDE: Our preliminary analysis shows that the reporting percents have improved from 1998 to 2000. We will learn later in this chapter how to formally test if this observed difference is statistically significant. ■

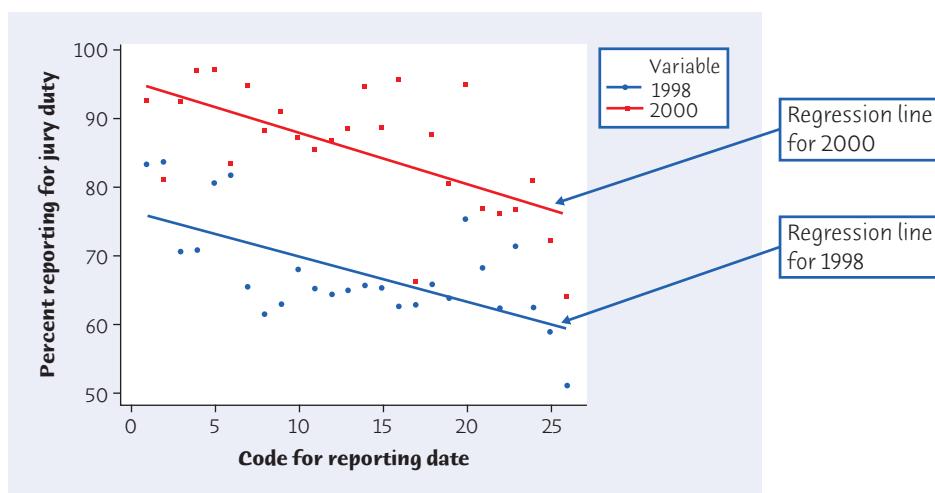


FIGURE 28.1

A scatterplot of percents reporting for jury duty in Franklin County Municipal Court, with two separate regression lines, for Example 28.1.

We now think that the percent of jurors reporting for duty declines at about the same rate in both years, but that the percent increased by a constant amount between 1998 and 2000. We would like to have a single regression model that captures this insight.

To do this, introduce a second explanatory variable x_2 for “year.” We might let x_2 take values 1998 and 2000, but then what would we do if the categorical variable took values “female” and “male”? A better approach is to just use values 0 and 1 to distinguish the two years. Now we have an *indicator variable*

$$x_2 = 0 \text{ for year 1998}$$

$$x_2 = 1 \text{ for year 2000}$$

INDICATOR VARIABLE

An **indicator variable** places individuals into one of two categories, usually coded by the two values 0 and 1.

Indicator variables are commonly used to indicate sex (0 = male, 1 = female), condition of patient (0 = good, 1 = poor), status of order (0 = undelivered, 1 = delivered), and many other characteristics for individuals.

The conditions for inference in simple linear regression (Chapter 24, text page 589) describe the relationship between the explanatory variable x and the mean response μ_y in the population by a *population regression line* $\mu_y = \beta_0 + \beta_1 x$. (The switch in notation from $\mu_y = \alpha + \beta x$ to $\mu_y = \beta_0 + \beta_1 x$ allows an easier extension to other models.) Now we add a second explanatory variable, so that our *regression model* for the population becomes

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

The other conditions for inference are the same as in the simple linear regression setting: for any fixed values of the explanatory variables, y varies about its mean according to a Normal distribution with unknown standard deviation σ that is the same for all values of x_1 and x_2 . We will look in detail at conditions for inference in multiple regression later on.

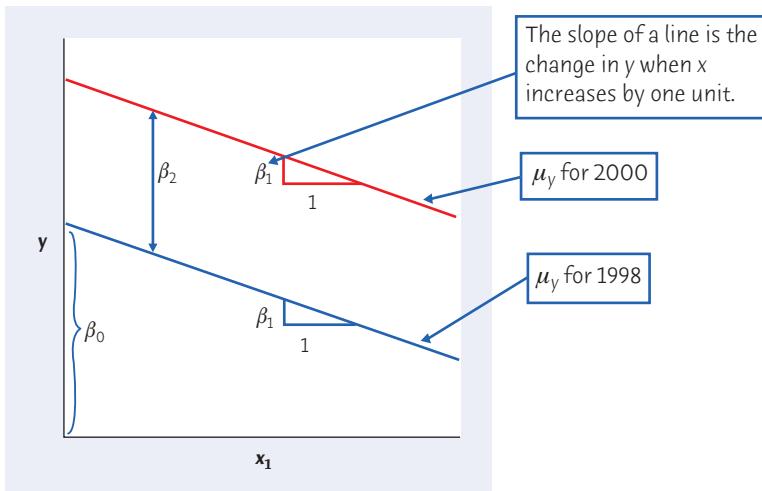
EXAMPLE 28.2 Interpreting a multiple regression model

Multiple regression models are no longer simple straight lines, so we must think a bit harder in order to interpret what they say. Consider our model

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

in which y is the percent of jurors who report, x_1 is the reporting date (1 to 26), and x_2 is an indicator variable for year. For 1998, $x_2 = 0$ and the model becomes

$$\mu_y = \beta_0 + \beta_1 x_1$$

**FIGURE 28.2**

Multiple regression model with two parallel straight lines, for Example 28.2

For 2000, $x_2 = 1$ and the model is

$$\begin{aligned}\mu_y &= \beta_0 + \beta_1 x_1 + \beta_2 \\ &= (\beta_0 + \beta_2) + \beta_1 x_1\end{aligned}$$

Look carefully: the slope that describes how the mean reporting percent changes as the reporting period x_1 runs from 1 to 26 is β_1 in both years. The intercepts differ: β_0 for 1998 and $\beta_0 + \beta_2$ for 2000. So β_2 is of particular interest, because it is the fixed change between 1998 and 2000.

Figure 28.2 is a graph of this model with all three β 's identified. We have succeeded in giving a single model for two parallel straight lines. ■

You will sometimes see indicator variables referred to as *dummy variables*. We have demonstrated how indicator (dummy) variables can be used to represent a categorical variable with two categories. More advanced books on multiple regression discuss how multiple indicator (dummy) variables can be used to represent categorical variables with more than two categories.

APPLY YOUR KNOWLEDGE

- 28.1 Bird colonies.** Suppose (this is too simple to be realistic) that the number y of new birds that join a colony this year has the same straight-line relationship with the percent x_1 of returning birds in colonies of two different bird species. An indicator variable shows which species we observe: $x_2 = 0$ for one and $x_2 = 1$ for the other. Write a population regression model that describes this setting. Explain in words what each β in your model means.

- 28.2 How fast do icicles grow?** We have data on the growth of icicles starting at length 10 centimeters (cm) and at length 20 cm. An icicle grows at the same rate, 0.15 cm per minute, starting from either length. Give a regression model that describes how mean length changes with time x_1 and starting length x_2 . Use numbers, not symbols, for the β 's in your model.

ESTIMATING PARAMETERS

How shall we estimate the β 's in the model $\mu_y = \beta_0 + \beta_1x_1 + \beta_2x_2$? Because we hope to predict y , we want to make the errors in the y direction small. We can't call this the vertical distance from the points to a *line* as we did for a simple linear regression model, because we now have two lines. But we still concentrate on the prediction of y and therefore on the deviations between the observed responses y and the responses predicted by the regression model.

The method of least squares estimates the β 's in the model by choosing the values that minimize the sum of the squared deviations in the y direction,

$$\sum (\text{observed } y - \text{predicted } y)^2 = \sum (y - \hat{y})^2$$

Call the values of the β 's that do this b 's. The least-squares regression model $\hat{y} = b_0 + b_1x_1 + b_2x_2$ estimates the population regression model $\mu_y = \beta_0 + \beta_1x_1 + \beta_2x_2$.

The remaining parameter is the standard deviation σ , which describes the variability of the response y about the mean given by the population regression model. Recall that the **residuals** are the differences between the observed responses y and the responses \hat{y} predicted by the least-squares model. Since the residuals estimate the "left-over variation" about the regression model, the standard deviation s of the residuals is used to estimate σ . The value of s is also referred to as the *regression standard error*.

REGRESSION STANDARD ERROR

The **regression standard error** for the multiple regression model $\hat{y} = b_0 + b_1x_1 + b_2x_2$ is

$$\begin{aligned}s &= \sqrt{\frac{1}{n-3} \sum \text{residual}^2} \\ &= \sqrt{\frac{1}{n-3} \sum (y - \hat{y})^2}\end{aligned}$$

Use s to estimate the standard deviation σ of the responses about the mean given by the population regression model.

degrees of freedom

Notice that instead of dividing by $(n - 2)$, the number of observations less 2, as we did for the simple linear regression model in Chapter 24, we are now dividing by $(n - 3)$, the number of observations less 3. Since we are estimating three β parameters in our population regression model, the degrees of freedom must reflect this change. In general, the **degrees of freedom** for the regression standard error will be the number of data points minus the number of β parameters in the population regression model.

Why do we prefer one regression model with parallel lines to the two separate regressions in Figure 28.1? Simplicity is one reason—why use separate models with four β 's if a single model with three β 's describes the data

well? Looking at the regression standard error provides another reason: the n in the formula for s includes all of the observations in both years. As usual, more observations produce a more precise estimate of σ . (Of course, using one model for both years assumes that σ describes the scatter about the line in both years.)

EXAMPLE 28.3 Potential jurors

Example 28.2 introduced the regression model

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

for predicting reporting percent y from reporting date x_1 and year x_2 . Statistical software gives the least-squares estimate of this model as

$$\hat{y} = 77.1 - 0.717x_1 + 17.8x_2$$

By substituting the two values of the indicator variable into this estimated regression equation, we can obtain a least-squares line for each year. The predicted reporting percents are

$$\hat{y} = 94.9 - 0.717x_1 \text{ for } 2000 (x_2 = 1)$$

and

$$\hat{y} = 77.1 - 0.717x_1 \text{ for } 1998 (x_2 = 0)$$

Comparing these estimated regression equations with the two separate regression lines obtained in Example 28.1, we see that the intercept parameters are very close to one another (95.571 is close to 94.9, and 76.426 is close to 77.1) for both years. The big change, as intended, is that the slope -0.717 is now the same for both lines. In other words, the estimated change in mean reporting percent for a one-unit change in the reporting date is now the same for both models, -0.717 . A closer look reveals that -0.717 is the average of the two slope estimates (-0.765 and -0.668) obtained in Example 28.1.

Finally, the regression standard error $s = 6.709$ indicates the size of the “typical” error. Thus, for a particular reporting date, we would expect approximately 95% of the reporting percents to be within $2 \times 6.709 = 13.418$ of their mean. ■

Table 28.2 provides more data on the reporting percents for randomly selected registered voters who received a summons to appear for jury duty in the Franklin County Municipal Court for 1985 and for 1997 through 2004. Each year 26 different groups of potential jurors are randomly selected to serve two weeks of jury duty. The reporting dates vary slightly from year to year, so they are coded sequentially from 1, the first group to report in January, to 26, the last group to report in December. The jury commissioner and other officials use the data in Table 28.2 to evaluate their efforts to improve turnout from the pool of potential jurors. We will use a variety of models to analyze the reporting percents in exercises and examples throughout this chapter.



Why some men earn more!

Research based on data from the

U.S. Bureau of Labor Statistics and the U.S. Census Bureau suggests that women earn 80 cents for every dollar men earn. While the literature is full of clear and convincing cases of discrimination based on height, weight, race, gender, and religion, new studies suggest that our choices explain a considerable amount of the variation in wages. Earning more often means that you are willing to accept longer commuting times, safety risks, frequent travel, long hours, and other responsibilities that take away from your time at home with family and friends. When choosing between time and money, make sure that you are happy with your choice!

TABLE 28.2 Percents of randomly selected registered voters who appeared for jury duty in Franklin County Municipal Court in 1985 and 1997–2004

REPORTING DATE	1985	1997	1998	1999	2000	2001	2002	2003	2004
1	21.3	38.7	83.30	73.0	92.59	94.0	97.2	89.1	88.50
2	17.3	34.7	83.60	69.1	81.10	87.7	91.4	98.4	88.00
3	21.8	47.3	70.50	67.1	92.50	94.8	90.2	92.0	91.80
4	21.7	43.1	70.70	65.7	97.00	94.2	90.0	83.6	90.84
5	23.5	50.7	80.50	67.6	97.00	71.4	95.2	87.1	81.16
6	15.1	35.1	81.60	65.7	83.30	89.2	92.4	82.8	84.86
7	21.7	33.9	65.30	57.3	94.60	73.1	90.0	82.1	90.91
8	20.0	28.7	61.30	69.6	88.10	68.1	94.0	85.4	81.85
9	21.1	36.6	62.70	64.5	90.90	92.3	95.3	90.5	80.93
10	22.0	29.6	67.80	73.6	87.10	90.3	94.8	98.6	85.70
11	21.7	31.8	65.00	61.6	85.40	76.9	82.4	87.5	78.98
12	20.0	35.2	64.10	75.2	86.60	93.1	90.1	98.5	86.13
13	20.0	23.3	64.70	74.3	88.30	98.5	83.4	89.8	91.50
14	24.4	38.0	65.40	60.0	94.40	92.9	91.4	76.3	85.91
15	14.3	32.8	65.02	59.5	88.50	75.9	84.2	95.8	75.83
16	21.0	40.0	62.30	65.9	95.50	100.0	84.0	87.6	91.14
17	17.9	58.4	62.50	62.5	65.90	88.7	81.9	97.1	80.25
18	26.0	60.1	65.50	65.2	87.50	78.8	78.7	100.0	94.64
19	23.8	52.1	63.50	62.1	80.20	97.0	80.7	86.3	90.84
20	27.6	54.2	75.00	65.8	94.70	95.0	91.0	82.3	86.75
21	29.3	66.6	67.90	69.2	76.60	83.9	98.4	90.5	91.14
22	28.0	88.0	62.00	64.7	75.80	69.6	84.2	80.2	88.27
23	27.0	88.4	71.00	65.7	76.50	70.2	76.4	97.3	90.35
24	21.8	70.3	62.10	58.9	80.60	69.1	70.2	76.5	82.56
25	33.0	71.0	58.50	63.0	71.80	78.2	71.5	91.2	90.66
26	14.2	62.1	50.70	55.5	63.70	n.a.	50.0	n.a.	86.29

Note: n.a. indicates that data are not available.



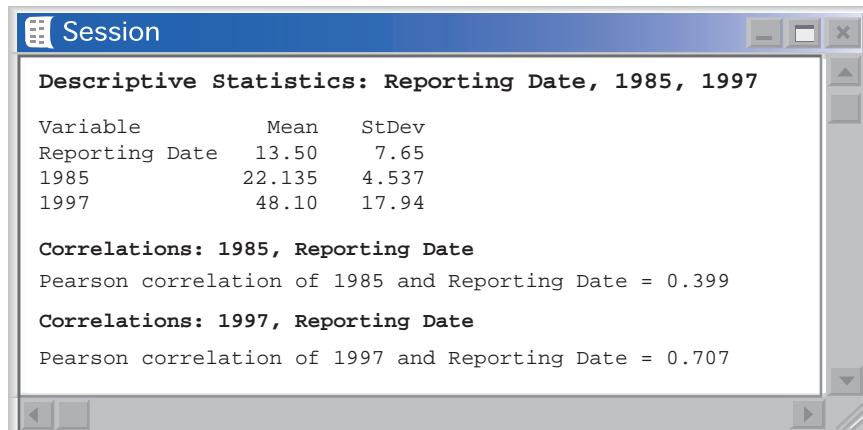
APPLY YOUR KNOWLEDGE

28.3 Potential jurors. On page 28-11 are descriptive statistics and a scatterplot for the reporting percents in 1985 and 1997 from Table 28.2.

- Use the descriptive statistics to compute the least-squares regression line for predicting the reporting percent from the coded reporting date in 1985.
- Use the descriptive statistics to compute the least-squares regression line for predicting the reporting percent from the coded reporting date in 1997.
- Interpret the value of the slope for each of your estimated models.
- Are the two estimated slopes about the same?

- (e) Would you be willing to use the multiple regression model with equal slopes to predict the reporting percents in 1985 and 1997? Explain why or why not. 

Minitab



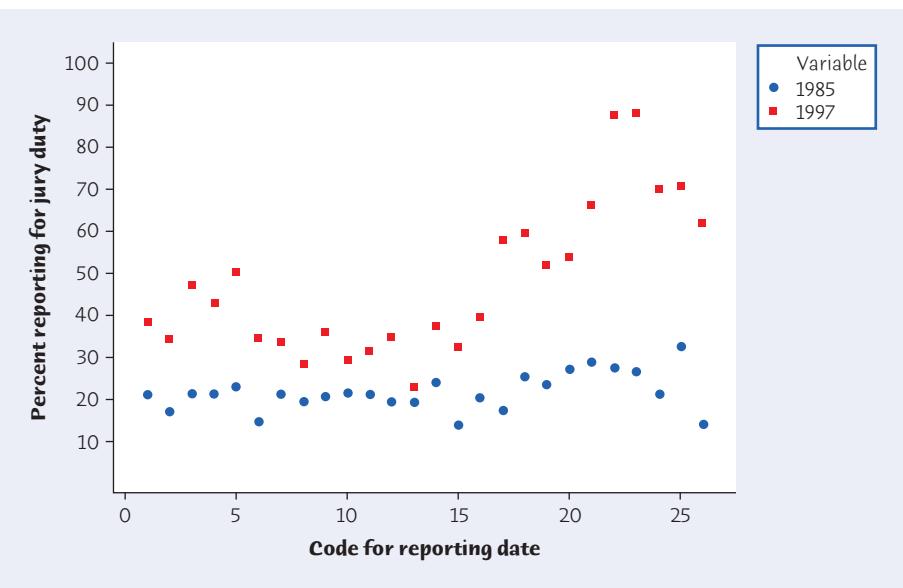
Session

Descriptive Statistics: Reporting Date, 1985, 1997

Variable	Mean	StDev
Reporting Date	13.50	7.65
1985	22.135	4.537
1997	48.10	17.94

Correlations: 1985, Reporting Date
Pearson correlation of 1985 and Reporting Date = 0.399

Correlations: 1997, Reporting Date
Pearson correlation of 1997 and Reporting Date = 0.707



28.4 Potential jurors. In Example 28.3 the indicator variable for year ($x_2 = 0$ for 1998 and $x_2 = 1$ for 2000) was used to combine the two separate regression models from Example 28.1 into one multiple regression model. Suppose that instead of x_2 we use an indicator variable x_3 that reverses the two years, so that $x_3 = 1$ for 1998 and $x_3 = 0$ for 2000. The mean reporting percent is $\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_3$, where x_1 is the code for the reporting date (the value on the x axis in Figure 28.1) and x_3 is an indicator variable to identify the year (different symbols in Figure 28.1). Statistical software now gives the estimated regression model as $\hat{y} = 94.9 - 0.717x_1 - 17.8x_3$.

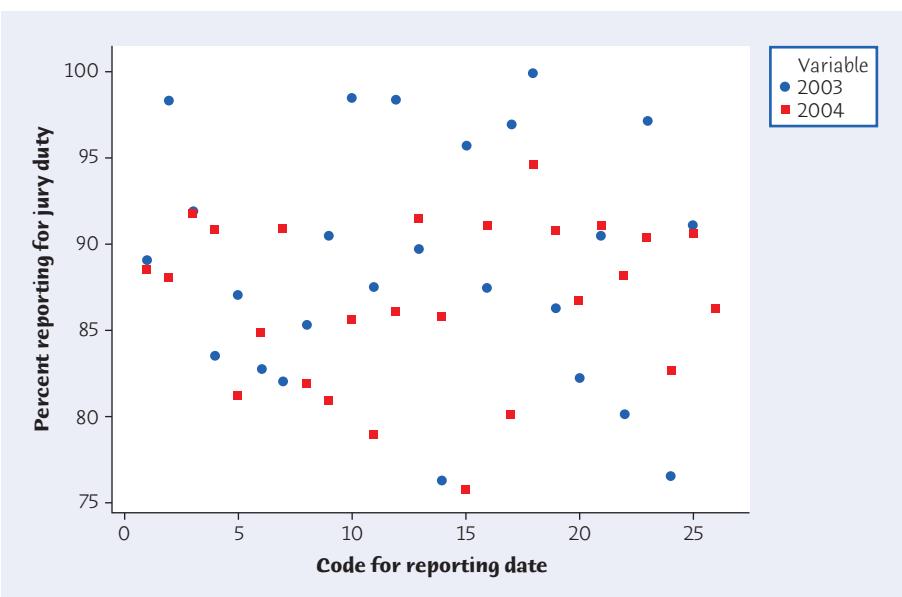
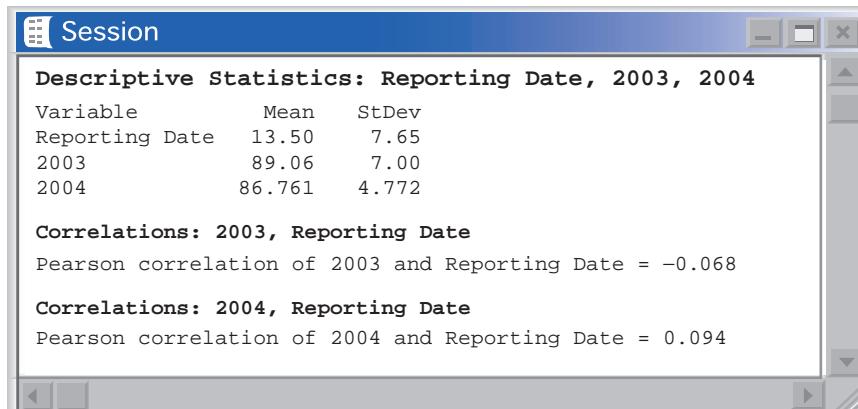
- (a) Substitute the two values of the indicator variable into the estimated regression equation to obtain a least-squares line for each year.

- (b) How do your estimated regression lines in part (a) compare with the estimated regression lines provided for each year in Example 28.3?
- (c) Will the regression standard error change when this new indicator variable is used? Explain.

28.5 Potential jurors. Here are descriptive statistics and a scatterplot for the reporting percents in 2003 and 2004 from Table 28.2.

- (a) Use the descriptive statistics to compute the least-squares regression line for predicting the reporting percent from the coded reporting date in 2003.
- (b) Use the descriptive statistics to compute the least-squares regression line for predicting the reporting percent from the coded reporting date in 2004.
- (c) Interpret the value of the slope for each of your estimated models.
- (d) Are the two estimated slopes about the same?

Minitab



- (e) Would you be willing to use the multiple regression model with equal slopes to predict the reporting percents in 2003 and 2004? Explain why or why not.
- (f) How does the estimated slope in 2003 compare with the estimated slope obtained in Example 28.3 for 1998 and 2000?
- (g) Based on the descriptive statistics and scatterplots provided in Exercise 28.3, Example 28.1, and on page 28-11, do you think that the jury commissioner is happy with the modifications he made to improve the reporting percents?

USING TECHNOLOGY

Table 28.2 provides a compact way to display data in a textbook, but this is not the best way to enter your data into a statistical software package for analysis. The usual format for data files is that each row contains data on one individual and each column contains the values of one variable.

EXAMPLE 28.4 Organizing data

The multiple regression model in Example 28.3 requires three columns. The 52 reporting percents y for 1998 and 2000 appear in a column labeled *Percent*, values of the explanatory variable x_1 make up a column labeled *Group*, and values of the indicator variable x_2 make up a column labeled *Ind2000*. The first five rows of the worksheet are shown below.

Row	Percent	Group	Ind2000
1	83.30	1	0
2	83.60	2	0
3	70.50	3	0
4	70.70	4	0
5	80.50	5	0

To use statistical software, we need only identify the response variable *Percent* and the two explanatory variables *Group* and *Ind2000*. Figure 28.3 shows the regression output from Minitab and CrunchIt! Each package provides parameter estimates, standard errors, *t* statistics, *P*-values, the regression standard error, and R^2 . Minitab also provides an analysis of variance table. We will digest this output one piece at a time: first describe the model, then look at the conditions needed for inference, and finally interpret the results of inference.

EXAMPLE 28.5 Parameter estimates on statistical output

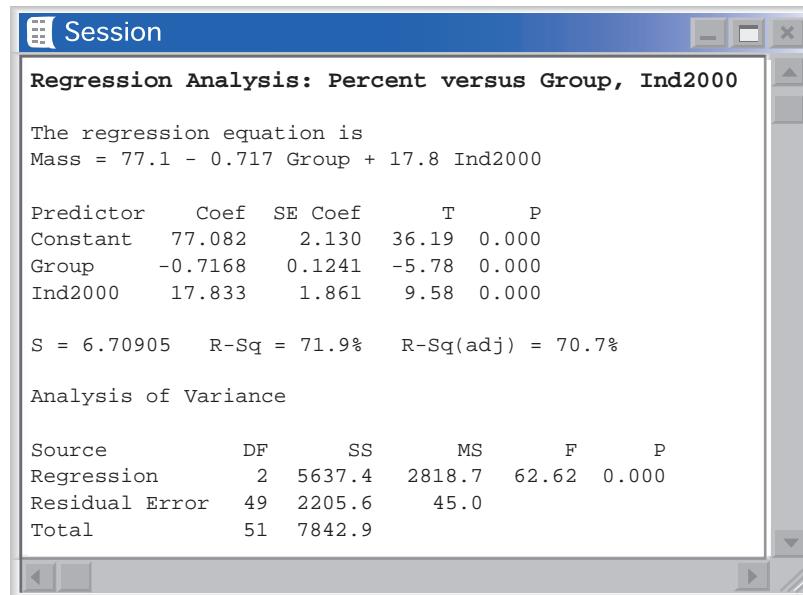
Both outputs give the multiple regression model for predicting the reporting percent (after rounding) as $\hat{y} = 77.082 - 0.717x_1 + 17.833x_2$.

Although the labels differ, the regression standard error is provided by both packages:

CrunchIt!: Root MSE = sigma: 6.709

Minitab: S = 6.70905 ■

Minitab



Minitab Session window output:

```

Session
Regression Analysis: Percent versus Group, Ind2000

The regression equation is
Mass = 77.1 - 0.717 Group + 17.8 Ind2000

Predictor      Coef    SE Coef      T      P
Constant     77.082   2.130   36.19  0.000
Group        -0.7168  0.1241  -5.78  0.000
Ind2000       17.833  1.861   9.58  0.000

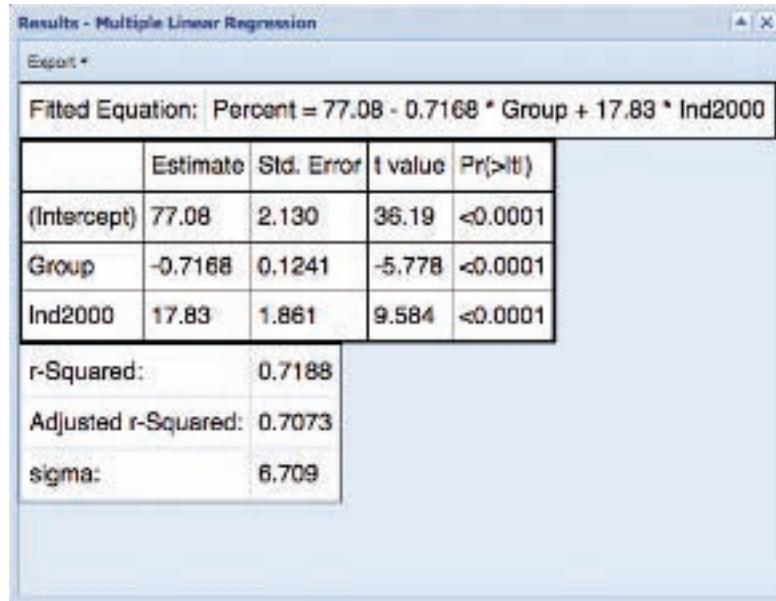
S = 6.70905   R-Sq = 71.9%   R-Sq(adj) = 70.7%

Analysis of Variance

Source          DF      SS      MS      F      P
Regression      2  5637.4  2818.7  62.62  0.000
Residual Error  49  2205.6    45.0
Total           51  7842.9

```

CrunchIt!

**FIGURE 28.3**

Output from Minitab and CrunchIt! for the model with parallel regression lines in Example 28.3.

For simple linear regression models, the square of the correlation coefficient r^2 between y and x measures the proportion of variation in the response variable that is explained by using the explanatory variable. For our multiple regression model with parallel regression lines, we do not have one correlation coefficient. However, by squaring the correlation coefficient between the observed responses y and the predicted responses \hat{y} we obtain the *squared multiple correlation coefficient* R^2 .

The analysis of variance table helps us interpret this new statistic. The sum of squares row in the ANOVA table breaks the total variability in the responses into two pieces. One piece summarizes the variability explained by the model, and the other piece summarizes the “left-over” variability, traditionally called “error.” That is,

$$\text{total sum of squares} = \text{model sum of squares} + \text{error sum of squares}$$

The value of R^2 is the ratio of the model sum of squares to the total sum of squares, so R^2 tells us what proportion of the variation in the response variable y we explained by using the set of explanatory variables in the multiple regression model.

SQUARED MULTIPLE CORRELATION COEFFICIENT

The **squared multiple correlation coefficient** R^2 is the square of the correlation coefficient between the observed responses y and the predicted responses \hat{y} . It is also equal to

$$R^2 = \frac{\text{variability explained by model}}{\text{total variability in } y} = \frac{\text{model sum of squares}}{\text{total sum of squares}}$$

R^2 is almost always given with a regression model to describe the fit of the model to the data.

EXAMPLE 28.6 Using R^2

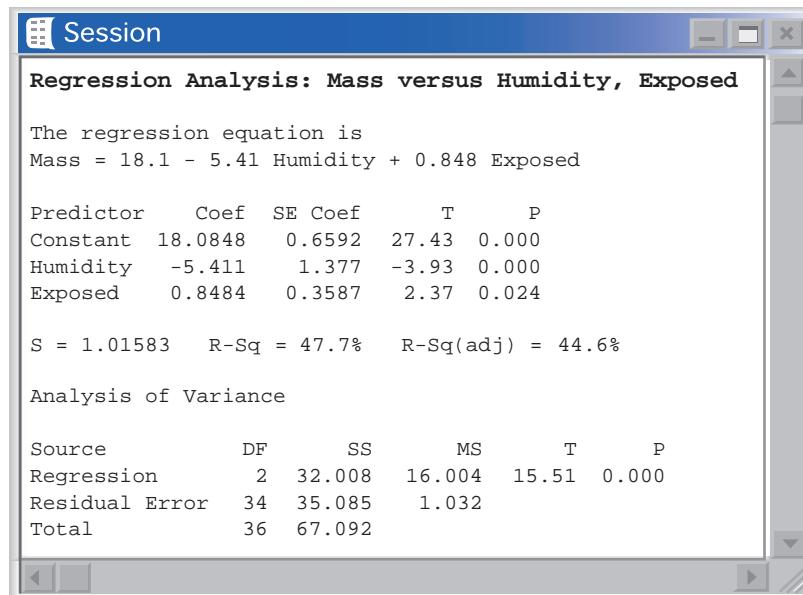
Both outputs in Figure 28.3 give the value $R^2 = 0.7188$ (rounded up to 71.9% in Minitab) for our multiple regression model with parallel lines in Example 28.3. That is, the regression model with explanatory variables *Group* and *Ind2000* explains about 72% of the variation in the response variable *Percent*. ■

APPLY YOUR KNOWLEDGE

28.6 Heights and weights for boys and girls. Suppose you are designing a study to investigate the relationship between height and weight for boys and girls.

- Specify a model with parallel regression lines that could be used to predict height separately for boys and for girls. Be sure to identify all variables and describe all parameters in your model.
- How many columns in a worksheet would be required to fit this model with statistical software? Describe each column.

Minitab



The screenshot shows the Minitab software interface with a title bar "Session". The main window displays a "Regression Analysis: Mass versus Humidity, Exposed". The output includes:

Regression equation: Mass = 18.1 - 5.41 Humidity + 0.848 Exposed

Predictor	Coeff	SE Coef	T	P
Constant	18.0848	0.6592	27.43	0.000
Humidity	-5.411	1.377	-3.93	0.000
Exposed	0.8484	0.3587	2.37	0.024

Statistics: S = 1.01583 R-Sq = 47.7% R-Sq(adj) = 44.6%

Analysis of Variance:

Source	DF	SS	MS	T	P
Regression	2	32.008	16.004	15.51	0.000
Residual Error	34	35.085	1.032		
Total	36	67.092			

28.7 Nestling mass and nest humidity. Researchers investigated the relationship between nestling mass, measured in grams, and nest humidity index, measured as the ratio of total mass of water in the nest divided by nest dry mass, for two different groups of great titmice parents.² One group was exposed to fleas during egg laying and the other was not. Exposed parents were coded as 1, and unexposed parents were coded as 0. Use the output above, obtained by fitting a multiple regression model with parallel lines for the two groups of parents, to answer the following questions. 

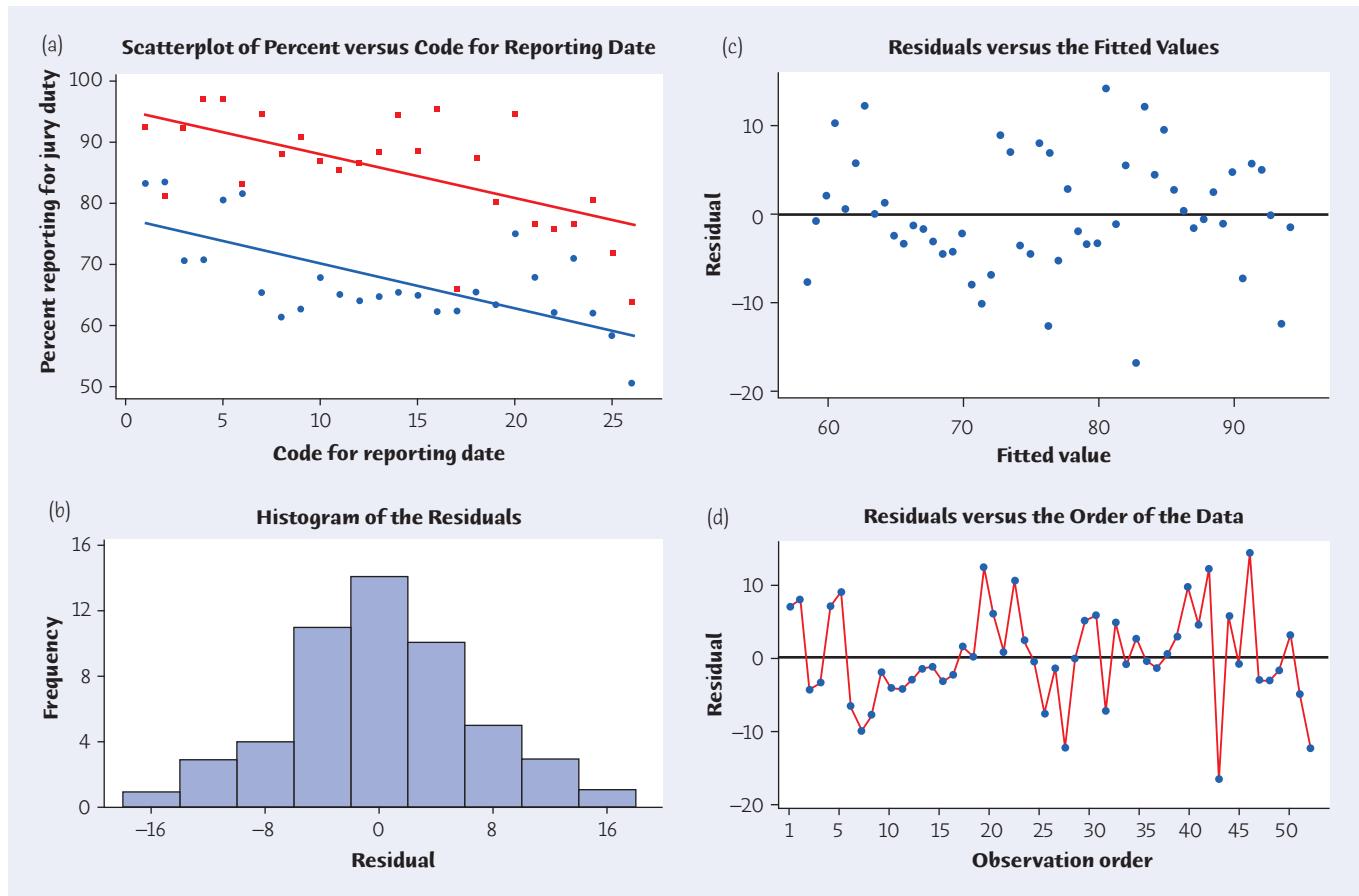
- Identify the regression model for predicting nestling mass from nest humidity index for the two groups of great titmice parents.
- Based on your model, do you think that nestling mass was higher in nests of birds exposed to fleas during egg laying? Explain.
- What is the value of the regression standard error? Interpret this value.
- What is the value of the squared multiple correlation coefficient? Interpret this value.

INFERENCE FOR MULTIPLE REGRESSION

The output in Figure 28.3 (page 28-14) contains a considerable amount of additional information that deals with statistical inference for our multiple regression model with parallel lines. Before taking our first look at inference for multiple regression, we will check the conditions for inference.

EXAMPLE 28.7 Checking the conditions

A scatterplot and residual plots for the multiple regression model with parallel lines in Example 28.3 are shown in Figure 28.4. The conditions for inference are linearity,

**FIGURE 28.4**

Scatterplot, histogram, and residual plots to check the conditions for inference the model with parallel regression lines in Example 28.3.

Normality, constant variance, and independence. We will check these conditions one at a time.

Linear trend: The scatterplot in Figure 28.4(a) shows a parallel linear pattern for the two years, so the model is reasonable.

Normality: The histogram of the residuals in Figure 28.4(b) indicates that the residuals are symmetric about zero and approximately Normal.

Constant variance: The residual plot in Figure 28.4(c) is not a perfectly unstructured horizontal band of points. However, the overall pattern does suggest that the variability in the residuals is roughly constant, with the exception of the residuals for fitted values between 65 and 70. In particular, this residual plot supports the model's condition that a single σ describes the scatter about the 1998 line and the 2000 line.

Independence: Since 26 groups are randomly selected each year, it is reasonable to assume that the reporting percents are independent. The residual plot in Figure 28.4(d) provides a quick check to see if there is a pattern in the residuals based on the order in which they were entered into the worksheet. In some situations, as is the case in this example, the order of entry will correspond to time or some other variable of interest, so plotting the residuals versus order can provide a valuable check of independence.

The plot shows one potentially troubling feature, residuals 7 through 17 are negative, but otherwise there is no systematic pattern. A closer look at the scatterplot in Figure 28.4(a) reveals that the blue points for reporting dates 7 through 17 are all below the line. However, the linear model still provides a reasonable summary of the reporting percents, so we will rely on the fact that multiple regression models are robust to slight departures from the conditions and proceed with inference for this model. ■

To this point we have concentrated on understanding the model, estimating parameters, and verifying the conditions for inference that are part of a regression model. Inference in multiple regression begins with tests that help us decide if a model adequately fits the data and choose between several possible models.

The first inference for a multiple regression model examines the overall model. The ANOVA table summarizes the breakdown of the variability in the response variable. There is one row for each of the three sources of variation: Model, Error, and Total. Each source of variation has a number of degrees of freedom associated with it. These degrees of freedom are listed in a column. Another column provides a sum of squares for the three components. The sums of squares are divided by the degrees of freedom within each row to form a column for the mean sum of squares. Finally, the mean sum of squares for the model is divided by the mean sum of squares for error to form the F statistic for the overall model. This F statistic is used to find out if all of the regression coefficients, except the intercept, are equal to zero.

F STATISTIC FOR REGRESSION MODEL

The analysis of variance F statistic for testing the null hypotheses that all of the regression coefficients (β 's), except β_0 , are equal to zero has the form

$$F = \frac{\text{variation due to model}}{\text{variation due to error}} = \frac{\text{Model mean square}}{\text{Error mean square}}$$

EXAMPLE 28.8 Overall F test for parallel lines

The regression model for the mean reporting percent is $\beta_y = \beta_0 + \beta_1x_1 + \beta_2x_2$, where x_1 is labeled as Group and x_2 is labeled as Ind2000 on the output in Figure 28.3 (page 28-14). The null and alternative hypotheses for the overall F test are

$$H_0: \beta_1 = \beta_2 = 0 \text{ (that is, } \mu_y = \beta_0\text{)}$$

$$H_a: \text{at least one of } \beta_1 \text{ and } \beta_2 \text{ is not } 0$$

null model

The null hypothesis H_0 specifies a model, called the **null model**, where the response variable y is a constant (its mean) plus random variation. In other words, the null model says that x_1 and x_2 together do not help predict y .

The value of the F statistic reported in the ANOVA table in Figure 28.3 is $F = 62.62$. You should check that this value is the mean square for the model divided by the mean square for error. The P -value is obtained from an F distribution with 2 numerator and 49 denominator degrees of freedom. Minitab reports a P -value of 0.000, that

is, zero to 3 decimal places. Since the P -value is less than any reasonable significance level, say $\alpha = 0.01$, we reject the null hypothesis and conclude that at least one of the x 's helps explain the variation in the reporting percent y . ■

Rejecting the null hypothesis with the F statistic tells us that at least one of our β parameters is not equal to zero, but it doesn't tell us which parameters are not equal to zero. We turn to individual tests for each parameter to answer that question.

INDIVIDUAL t TESTS FOR COEFFICIENTS

To test the null hypothesis that one of the β 's in a specific regression model is zero, compute the t statistic

$$t = \frac{\text{parameter estimate}}{\text{standard error of estimate}} = \frac{b}{\text{SE}_b}$$

If the conditions for inference are met, then the t distribution with $(n - 3)$ degrees of freedom can be used to compute confidence intervals and conduct hypothesis tests for β_0 , β_1 and β_2 .

EXAMPLE 28.9 Individual t tests

The output in Figure 28.3 (page 28-14) provides parameter estimates and standard errors for the coefficients β_0 , β_1 and β_2 . The individual t statistic for x_1 (*Group*) tests the hypotheses

$$\begin{aligned} H_0: \beta_1 &= 0 \text{ (that is, } \mu_y = \beta_0 + \beta_2 x_2) \\ H_a: \beta_1 &\neq 0 \end{aligned}$$

We explicitly state the model in the null hypothesis because the bare statement $H_0: \beta_1 = 0$ can be misleading. The hypothesis of interest is that *in this model* the coefficient of x_1 is 0. If the same x_1 is used in a different model with different explanatory variables, the hypothesis $H_0: \beta_1 = 0$ has a different meaning even though we would write it the same way.

Using the CrunchIt! Output we see that the test statistic is (with roundoff)

$$t = \frac{-0.7168}{0.1241} = -5.78$$

The P -value is the area under a t distribution curve with $52 - 3 = 49$ degrees of freedom below -5.78 or above 5.78 . Since this value is very small, CrunchIt! simply reports that the P -value is <0.0001 . Look back at the hypotheses to interpret this result: we have good evidence that reporting date x_1 (*Group*) helps explain the percent reporting y even after we allow year x_2 to explain the reporting percent.

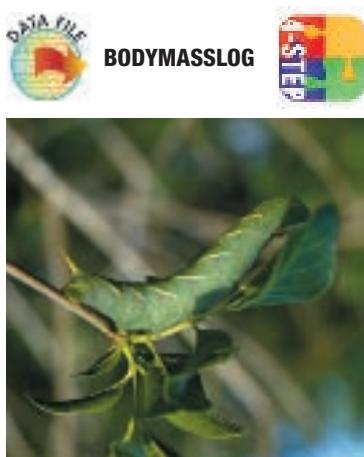
The test statistics for the other two coefficients are

$$t = \frac{77.08}{2.130} = 36.19 \text{ for } \beta_0$$

$$t = \frac{17.83}{1.861} = 9.58 \text{ for } \beta_2$$

The P -values are again obtained using the t distribution with 49 degrees of freedom. Both of the P -values are so small that they are reported by CrunchIt! as being <0.0001 . There is good evidence that the constant term β_0 is not 0, and that year x_2 adds to our ability to explain the reporting percent event after we take reporting date x_1 into account. ■

Example 28.9 is quite straightforward. The overall F test tells us that the two explanatory variables together help explain the response. The individual t tests indicate that each explanatory variable significantly improves the explanation when added to a model that uses only the other explanatory variable. Interpreting the results of individual t tests can get very tricky, so we will return to the more challenging situations later. We end our discussion of the model with parallel regression lines with an example that applies the four-step process.



Grady Harrison/Alamy

EXAMPLE 28.10 Metabolic rate and body mass in caterpillars

STATE: Scientists have long been interested in the question of how body mass (BM) determines physiological characteristics such as metabolic rate (MR). Recent experimental and theoretical research has confirmed the general relationship

$$MR = \alpha(BM)^\beta$$

between basal metabolic rate and body mass.³ However, there is still considerable debate on whether the scaling exponent is $\beta = 2/3$ or $\beta = 3/4$.

A group of researchers investigated the relationship between metabolic rate and body mass for tobacco hornworm caterpillars (*Manduca sexta*). These caterpillars were chosen because they maintain their shape throughout the five stages of larval development and the size of the tracheal system increases at each molt. A subset of the metabolic rates and body masses, after applying the logarithm transformation, is shown in Table 28.3 for caterpillars at the fourth and fifth stages of development.⁴ The complete data set is available in the file BODYMASSLOG.dat. Does the general relationship between metabolic rate and body mass hold for tobacco hornworm caterpillars? Is the relationship the same for the two different stages?

TABLE 28.3 Body masses and metabolic rates, after applying the logarithm transformation, for caterpillars in the fourth and fifth stages of development

LOG OF BODY MASS	LOG OF METABOLIC RATE	STAGE	STAGE INDICATOR
-0.56864	0.90780	4	0
-0.21753	1.24695	4	0
0.05881	1.51624	4	0
0.03342	1.42951	4	0
0.29336	1.56236	5	1
0.65562	1.92571	5	1
0.84757	1.83893	5	1
0.97658	2.03313	5	1

PLAN: To investigate the relationship between MR and BM, transform the data using logarithms so that the linear model

$$\mu_{\log}(\text{MR}) = \log(\alpha) + \beta \log(\text{BM})$$

can be fitted. Since a simple linear regression model can be used to address the first research question, we will leave the details for a review exercise (see Exercise 28.8). To check if the linear relationship is the same for both stages, we will fit a model with parallel regression lines.

SOLVE: Figure 28.5 shows a scatterplot of the transformed metabolic rate, measured in microliters of oxygen per minute ($\mu\text{l}/\text{min}$), against the transformed body mass measured in grams (g). The parallel regression lines on the plot, one for Stage 4 and one for Stage 5, illustrate the predicted model. The overall patterns for each of the two stages appear to be very similar. However, the measurements for Stage 5 (red points on the plot) are shifted up and to the right of those for Stage 4 (blue points on the plot).

The Minitab output (see page 28-22) was obtained by regressing the response variable y (the logarithm of metabolic rate) on two predictor variables, x_1 (the logarithm of body mass) and an indicator variable x_2 , which is 1 for Stage 5 and 0 for Stage 4. Our multiple regression model is $\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$.

The estimated multiple regression model is

$$\hat{y} = 1.24 + 0.698x_1 + 0.147x_2$$

Substituting the values of 0 and 1 for x_2 , we obtain the parallel regression lines

$$\hat{y} = 1.24 + 0.698x_1, \text{ for Stage 4 } (x_2 = 0)$$

$$\hat{y} = 1.387 + 0.698x_1, \text{ for Stage 5 } (x_2 = 1)$$

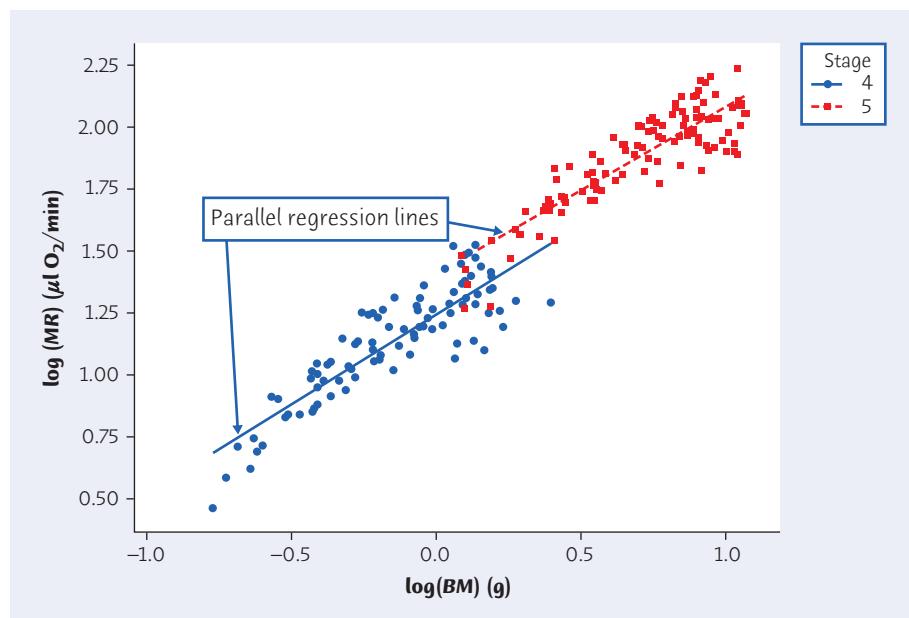


FIGURE 28.5

Scatterplot for the predicted model using parallel regression lines, for Example 28.10.

Minitab

The screenshot shows the Minitab Session window with the following output:

```

Regression Analysis: y versus x1, x2
The regression equation is
y = 1.24 + 0.698 · 1 + 0.147 · 2

Predictor      Coef    SE Coef      T      P
Constant      1.23917  0.01122  110.44  0.000
· 1           0.69828  0.02628   26.57  0.000
· 2           0.14680  0.02658   5.52  0.000

S = 0.100121  R-Sq = 94.5%  R-Sq(adj) = 94.5%

Analysis of Variance

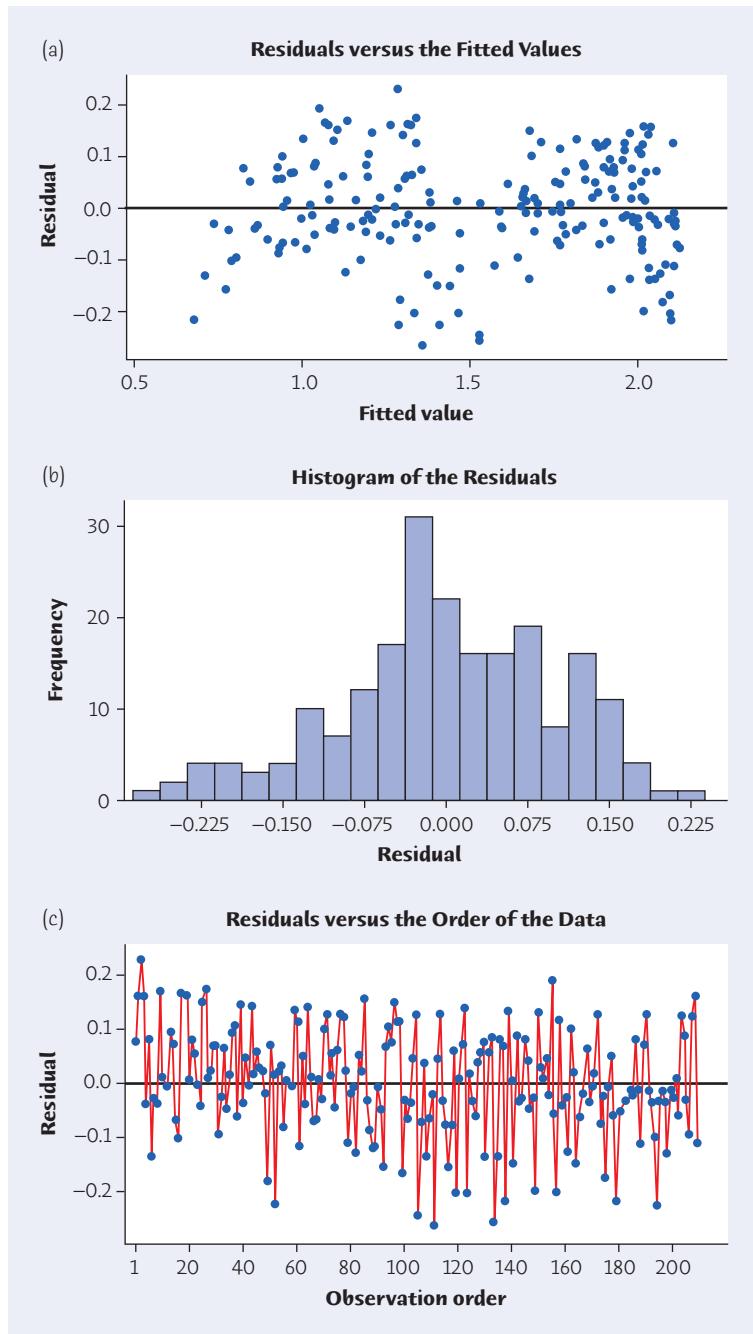
Source        DF      SS      MS      F      P
Regression     2      35.784  17.892  1784.89  0.000
Residual Error 206    2.065   0.010
Total          208    37.849

```

To check the conditions for inference we notice that the scatterplot in Figure 28.5 seems to show a parallel linear pattern, so the model makes sense. The residual plots in Figure 28.6 are used to check the other conditions. The histogram in Figure 28.6(b) indicates that the residuals are approximately symmetric about zero, so the Normality condition is satisfied. The plot of the residuals versus the fitted values in Figure 28.6(a) shows some trends that concerned the researchers. In particular, it appears that a model with some curvature might do a slightly better job because the residuals were always negative for the lowest body mass measurements within each stage. They were also slightly concerned about the constant-variance assumption. The plot of the residuals versus the data order in Figure 28.6(c) shows no systematic change of spread about the model. However, there is a slight curvilinear or “u-shaped” pattern, perhaps suggesting some structure in the process with respect to order.

Since the researchers were interested in comparing their results for caterpillars with the general relationship used by other scientists for a variety of other animals and insects, they decided to proceed with statistical inference for the model parameters. The overall F statistic $F = 1784.89$ and corresponding P -value $P = 0.000$ clearly indicate that at least one of the parameters in the model is not equal to zero. Since the t statistics 110.44 , 26.57 , and 5.52 all have reported P -values of zero, we conclude that all three parameters β_0 , β_1 , and β_2 are significantly different from zero.

CONCLUDE: The researchers were pleased that they were able to explain 94.5% of the variation in the logarithm of the metabolic rates by using a regression model with two parallel lines, one for each stage. The general form of the linear relationship is the same for both stages, with overall slope $b_1 = 0.698$. The major difference in the relationship for the two stages is indicated by an upward shift in the line for the larger caterpillars, which is estimated by $b_2 = 0.147$. ■

**FIGURE 28.6**

Residual plots for the model with parallel regression lines in Example 28.10.

APPLY YOUR KNOWLEDGE

- 28.8 Metabolic rate and body mass for caterpillars.** Does the general relationship between metabolic rate and body mass described in Example 28.10 hold for tobacco hornworm caterpillars? The Minitab output (see page 28-24) was obtained by regressing the response variable $y = \log(\text{MR})$ on $x_1 = \log(\text{BM})$ for the data.

Minitab

The Minitab session window displays the following output:

```

Session
Regression Analysis: y versus x1

The regression equation is
log(MR) = 1.28 + 0.822 log(BM)

 Predictor      Coef    SE Coef        T      P
 Constant     1.28071   0.00890   143.88  0.000
 x1           0.82179   0.01477    55.66  0.000

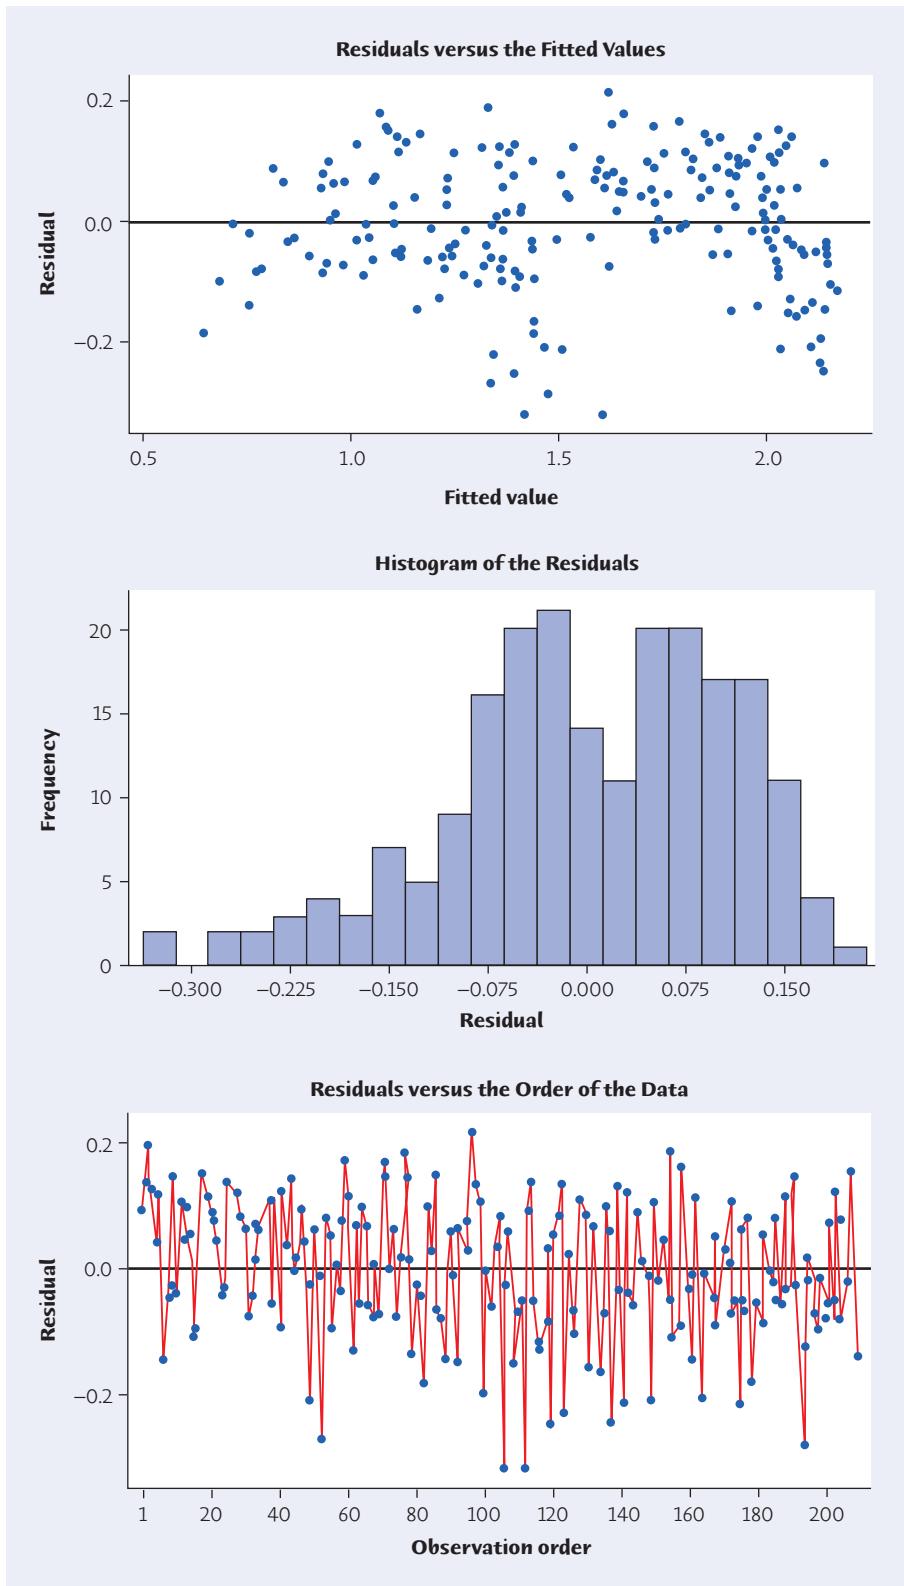
S = 0.107019    R-Sq = 93.7%    R-Sq(adj) = 93.7%

```

- (a) Use the regression equation from the Minitab output to estimate α and β in the general relationship $MR = \alpha(BM)^\beta$, which is the same as $\mu_y = \log(\alpha) + \beta x_1$. The predicted model is $\hat{y} = a + bx_1$, so that a estimates $\log(\alpha)$ and b estimates β in the original model.
- (b) Residual plots for the linear regression model $\mu_y = \alpha + \beta x_1$ are shown on page 28-25. Do you think that the conditions for inference are satisfied?
- (c) Identify the percent of variation in y that is explained by using linear regression with the explanatory variable x_1 .
- (d) Even if you noticed some departures from the conditions for inference, the researchers were interested in making inferences because this model is well known in the field and has been used for a variety of different insects and animals. Find a 95% confidence interval for the slope parameter β .
- (e) Are the values $\beta = 2/3$ and $\beta = 3/4$ contained in your confidence interval?
- (f) Use appropriate values from the Minitab output to test the claim that $\beta = 2/3$.
- (g) Use appropriate values from the Minitab output to test the claim that $\beta = 3/4$.

28.9 Metabolic rate and body mass for caterpillars. Use the output provided in Example 28.10 (page 28-20) to answer the questions below.

- (a) Find a 95% confidence interval for the slope parameter β for caterpillars during Stage 4.
- (b) If you were asked to report a confidence interval for the slope parameter β for caterpillars during Stage 5, would you report the same interval that you calculated in part (a)? Explain why or why not.
- (c) Are the values $\beta = 2/3$ and $\beta = 3/4$ contained in your confidence interval from part (a)?
- (d) How does your confidence interval in part (a) compare with the confidence interval you computed in part (d) of Exercise 28.8?
- (e) Use appropriate values from the output to test the claim that $\beta = 2/3$.
- (f) Use appropriate values from the output to test the claim that $\beta = 3/4$.



28.10 Reporting percents. Use the output in Figure 28.3 (page 28-14) to answer the questions below.

- Is the value of the regression standard error the same on both sets of output? Interpret this values.
- The value of the squared multiple correlation coefficient is reported as 71.9% by Minitab and 0.7188 by CrunchIt! Interpret the value of R^2 for this model.
- Is the value of the slope parameter significantly different from zero?
- Give a 98% confidence interval for the value of the slope parameter.
- Is there a significant difference in the intercepts for the two regression models?

INTERACTION

interaction

Examples with two parallel linear patterns for two values of an indicator variable are rather rare. It's more common to see two linear patterns that are not parallel. To write a regression model for this setting, we need an idea that is new and important: **interaction** between two explanatory variables. Interaction between variables x_1 and x_2 appears as a product term x_1x_2 in the model. The product term means that *the relationship between the mean response and one explanatory variable x_1 changes when we change the value of the other explanatory variable x_2* . Here is an example.



EXAMPLE 28.11 Revisiting state SAT scores

STATE: In Example 4.4 (text page 102) you discovered that states with a higher percent of high school graduates taking the SAT (rather than the ACT) tend to have lower mean SAT scores. You also saw that states fall into two distinct clusters, one for states with 49% or more of high school graduates taking the SAT and the other for states with at most 40% of high school graduates taking the SAT. Is a model with two regression lines helpful in predicting the SAT Math score for the two clusters of states?

PLAN: Fit and evaluate a model with two regression lines for predicting SAT Math score. ■

Let's see how adding an interaction term allows two lines that are not parallel. Consider the model

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

in which y is the SAT Math score, x_1 is the percent of high school students taking the SAT, x_2 is an indicator variable that is 1 if the percent of high school graduates taking the SAT is less than or equal to 40% and 0 otherwise, and $x_1 x_2$ is the interaction term. For states with 49% or more of students taking the SAT, $x_2 = 0$ and the model becomes

$$\mu_y = \beta_0 + \beta_1 x_1$$

For states with at most 40% of the students taking the SAT, $x_2 = 1$ and the model is

$$\begin{aligned}\mu_y &= \beta_0 + \beta_1 x_1 + \beta_2 + \beta_3 x_1 \\ &= (\beta_0 + \beta_2) + (\beta_1 + \beta_3) x_1\end{aligned}$$

A careful look allows us to interpret all four parameters: β_0 and β_1 are the intercept and slope for states with 49% or more of students taking the SAT. The parameters β_2 and β_3 indicate the fixed change in the intercept and slope, respectively, for states with at most 40% of students taking the SAT. Be careful not to interpret β_2 as the intercept and β_3 as the slope for states with a low percent of students taking the SAT. The indicator variable allows us to change the intercept as we did before, and the new interaction term allows us to change the slope.

A MODEL WITH TWO REGRESSION LINES

We have n observations on an explanatory variable x_1 , an indicator variable x_2 coded as 0 for some individuals and as 1 for other individuals, and a response variable y . The mean response μ_y is a linear function of the four parameters β_0 , β_1 , β_2 , and β_3 :

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

EXAMPLE 28.12 Revisiting state SAT scores, continued

SOLVE: Figure 28.7 shows the two regression lines, one for each cluster, for predicting the mean SAT Math score for each state. The fitted model, as shown by the two regression lines in this case, appears to provide a good visual summary for the two clusters.

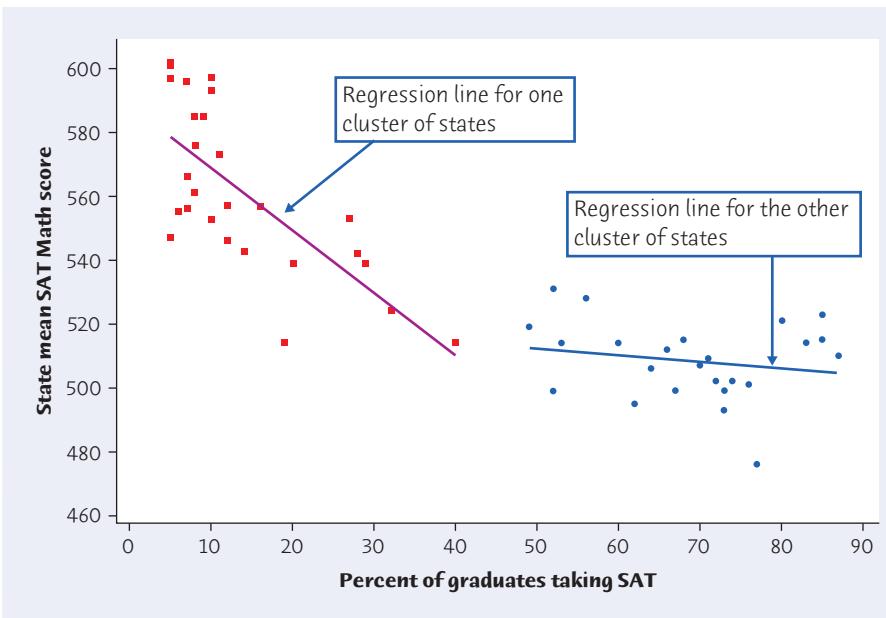


FIGURE 28.7

Model with two regression lines for predicting mean SAT Math score in each state based on the percent of high school graduates who take the SAT, for Example 28.12.

FIGURE 28.8

Output from Minitab for the model with two regression lines in Example 28.12.

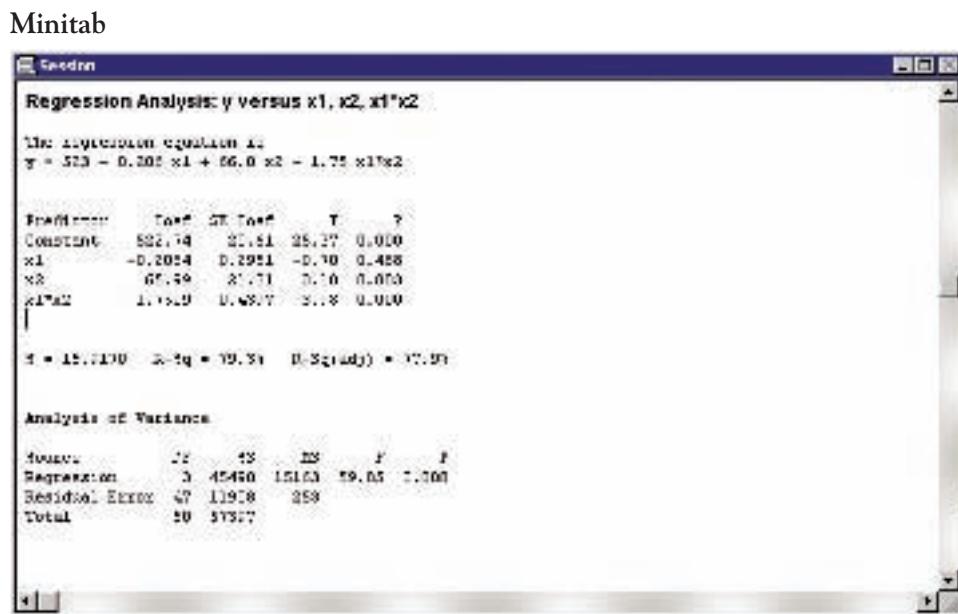


Figure 28.8 provides the regression output from Minitab. By substituting 0 and 1 for the indicator variable x_2 , we can easily obtain the two estimated regression lines. The estimated regression lines are $\hat{y} = 523 - 0.206x_1$ for states with at least 49% of high school graduates taking the SAT and

$$\begin{aligned}\hat{y} &= (523 + 66) - (0.206 + 1.752)x_1 \\ &= 589 - 1.958x_1\end{aligned}$$

for states with at most 40% of high school graduates taking the SAT.

The overall F statistic 59.85 and corresponding P -value in the ANOVA table clearly indicate that at least one of the regression coefficients is significantly different from zero. Thus, at least one of the two explanatory variables or the interaction of both is helpful in predicting the state mean SAT Math scores.

Looking at the individual t tests for the coefficients, we notice that only one coefficient, the coefficient for percent taking the SAT x_1 , is not significantly different from zero. The P -value for x_1 is so large that there is no evidence against $\beta_1 = 0$. The null hypothesis $H_0: \beta_1 = 0$ for this individual t test says that

$$\begin{aligned}\mu_y &= \beta_0 \text{ (a horizontal line) for states with at least 49\%} \\ \mu_y &= (\beta_0 + \beta_2) + \beta_3 x_1 \text{ for states at or below 40\%}\end{aligned}$$

Figure 28.7 shows that the regression line will be close to a horizontal line. The model specified by $H_0: \beta_1 = 0$ is reasonable. There is a clear ACT/SAT state difference but no evidence that percent taking (x_1) affects state SAT Math score once the percent taking is at least 49%. So multiple linear regression has led to an interesting model.

Residual plots (not shown) indicate one very small residual but no major problems with the Normality or constant-variance assumptions.

CONCLUDE: The model with two regression lines, one for each cluster, explains approximately 79.3% of the variation in the mean SAT Math scores. This model provides a better fit than the simple linear regression model that predicts mean SAT Math score from just the percent of high school graduates who take the SAT. ■

Even though we developed models without interaction first, it is best in practice to consider models with interaction terms before going to the more restrictive model with parallel regression lines. If you begin your model fitting with the more restrictive model with parallel regression lines, then you are basically assuming that there is no interaction. We won't discuss model selection formally, but deciding which model to use is an important skill.



EXAMPLE 28.13 Choosing a model

Let's compare three separate models for predicting SAT Math score y using the explanatory variables x_1 , x_2 , and $x_1 x_2$ described in Example 28.11.

Model 1: A simple linear regression model that ignores the two clusters of states

Model 2: The two-line model from Example 28.12

Model 3: A two-line model with 0 slope for states in the right-hand cluster

The predicted response \hat{y} , regression standard error s , and squared multiple correlation coefficient R^2 for the three models are

$$\text{Model 1: } \hat{y} = 575.27 - 0.97x_1 \quad s = 18.03 \quad R^2 = 0.722$$

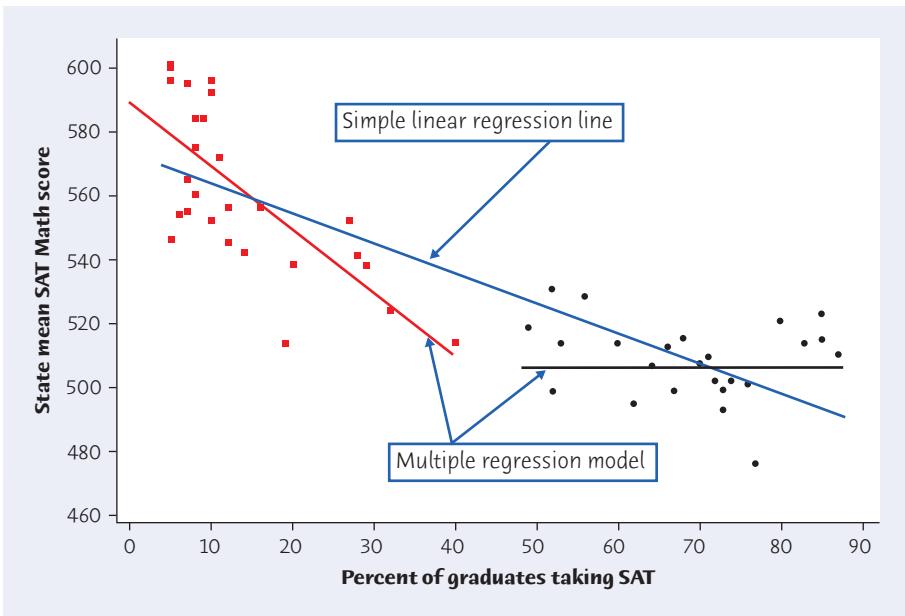
$$\text{Model 2: } \hat{y} = 523 - 0.206x_1 + 66.0x_2 - 1.75x_1x_2 \quad s = 15.92 \quad R^2 = 0.793$$

$$\text{Model 3: } \hat{y} = 509 + 80.2x_2 - 1.96x_1x_2 \quad s = 15.83 \quad R^2 = 0.790$$

We have already seen the fitted lines for Model 2 in Figure 28.7. The fitted lines for Models 1 and 3 appear in Figure 28.9. The blue line shows the simple linear regression model. The red line shows the fitted line for Model 3 in states with at most 40% of their graduates taking the SAT exam. The horizontal black line provides the prediction for Model 3 in all states with at least 49% of their graduates taking the SAT exam. It appears that Model 3 (red and black lines together) does a better job of explaining the variability in the mean SAT Math scores for the two clusters of states. Let's check the statistics.

Comparing Models 1 and 2, we find that Model 2 has the smaller s and the larger R^2 . Thus, the model with two separate regression lines provides a better fit than the simple linear regression model.

Comparing Models 2 and 3, we find that the R^2 -values are essentially the same, but the regression standard error is a bit smaller for Model 3. Therefore, Model 3 (which has one less β) does as good a job as the full two-line model. $P = 0.488$ for the individual t test for x_1 in Model 2 suggested this. ■

**FIGURE 28.9**

Scatterplot for Example 28.13 with two different models for predicting mean SAT Math score in each state based on the percent of high school graduates who take the SAT.

APPLY YOUR KNOWLEDGE

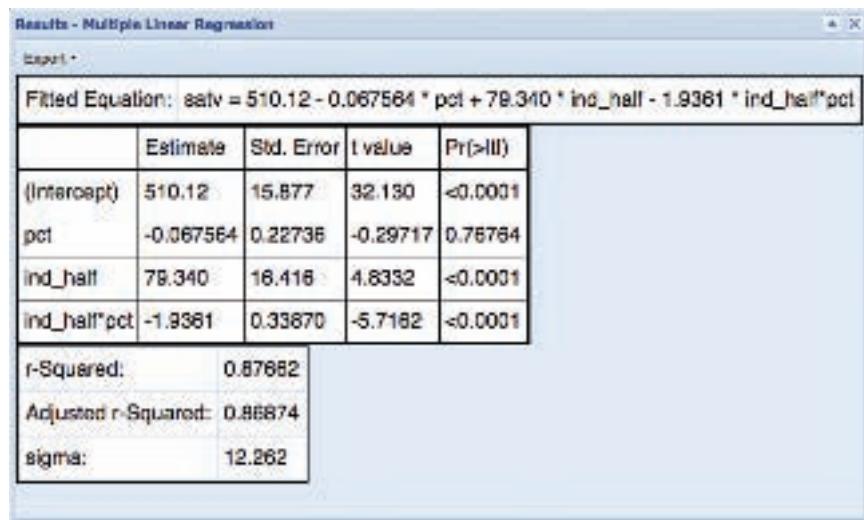
28.11 Bird colonies. Suppose that the number y of new birds that join a colony this year has a straight-line relationship with the percent x_1 of returning birds in colonies of two different bird species. An indicator variable shows which species we observe: $x_2 = 0$ for one and $x_2 = 1$ for the other. Write a population regression model that allows different linear models for the two different bird species. Explain in words what each β in your model means.

28.12 How fast do icicles grow? We have data on the growth of icicles starting at length 10 centimeters (cm) and at length 20 cm. Suppose that icicles which start at 10 cm grow at a rate of 0.15 cm per minute and icicles which start at 20 cm grow at the same rate, 0.15 cm per minute. Give a regression model that describes how mean length changes with time x_1 and starting length x_2 . Use numbers, not symbols, for the β 's in your model.

28.13 Touring battlefields. Suppose that buses complete tours at an average rate of 20 miles per hour and that self-guided cars complete tours at an average rate of 28 miles per hour. Give a regression model that describes how mean time to complete a tour changes with distance x_1 and mode of transportation x_2 . To be realistic, we want the mean time to complete the tour to be zero for both modes of transportation when the distance $x_1 = 0$. Use numbers, not symbols, for the β 's in your model.

28.14 Revisiting state SAT scores. We have examined the relationship between SAT Math scores and the percent of high school graduates who take the SAT.

CrunchIt!



CrunchIt! was used to fit a model with two regression lines, one for each cluster, for predicting SAT Verbal score. Use the CrunchIt! output above to answer the following questions.

- What is the estimated regression line for predicting mean SAT Verbal score for states with more than half of high school graduates taking the SAT?
- What is the estimated regression line for predicting mean SAT Verbal score for states with at most half of high school graduates taking the SAT?
- Interpret the squared multiple correlation.
- A t distribution was used to compute the P-values provided after each t-value in the table. How many degrees of freedom does that t distribution have?
- Identify the value you would use to estimate the standard deviation σ .
- Create a scatterplot containing the estimated regression lines for each cluster.  SATSCORES
- Plot the residuals against the fitted values. Does this plot indicate any serious problems with the conditions for inference?  SATSCORES
- Use a visual display to check the Normality condition for the residuals. Do you think the residuals follow a Normal distribution?  SATSCORES

28.15 World record running times. The table on the next page shows the progress of world record times (in seconds) for the 10,000-meter run for both men and women.  WORLDRECORDS

- Make a scatterplot of world record time against year, using separate symbols for men and women. Describe the pattern for each gender. Then compare the progress of men and women.
- Fit the model with two regression lines, one for women and one for men, and identify the estimated regression lines.

Men				Women	
Record year	Time (seconds)	Record year	Time (seconds)	Record year	Time (seconds)
1912	1880.8	1962	1698.2	1967	2286.4
1921	1840.2	1963	1695.6	1970	2130.5
1924	1835.4	1965	1659.3	1975	2100.4
1924	1823.2	1972	1658.4	1975	2041.4
1924	1806.2	1973	1650.8	1977	1995.1
1937	1805.6	1977	1650.5	1979	1972.5
1938	1802.0	1978	1642.4	1981	1950.8
1939	1792.6	1984	1633.8	1981	1937.2
1944	1775.4	1989	1628.2	1982	1895.2
1949	1768.2	1993	1627.9	1983	1895.0
1949	1767.2	1993	1618.4	1983	1887.6
1949	1761.2	1994	1612.2	1984	1873.8
1950	1742.6	1995	1603.5	1985	1859.4
1953	1741.6	1996	1598.1	1986	1813.7
1954	1734.2	1997	1591.3	1993	1771.8
1956	1722.8	1997	1587.8		
1956	1710.4	1998	1582.7		
1960	1698.8	2005	1577.5		

- (c) Women began running this long distance later than men, so we might expect their improvement to be more rapid. Moreover, it is often said that men have little advantage over women in distance running as opposed to sprints, where muscular strength plays a greater role. Do the data appear to support these claims?

28.16 Heights and weights for boys and girls. Suppose that you are designing a study to investigate the relationship between height and weight for boys and girls. Specify a model with two regression lines that could be used to predict height separately for boys and for girls. Be sure to identify all variables and describe all parameters in your model.

THE GENERAL MULTIPLE LINEAR REGRESSION MODEL

We have seen in a simple but useful case how adding another explanatory variable can fit patterns more complex than the single straight line of simple linear regression. Our examples to this point included two explanatory variables: a quantitative variable x_1 and an indicator variable x_2 . Some of our models added an interaction

term x_1x_2 . Now we want to allow any number of explanatory variables, each of which can be either quantitative or an indicator variable. Here is a statement of the general model that includes the conditions needed for inference.

THE MULTIPLE LINEAR REGRESSION MODEL

We have observations on n individuals. Each observation consists of values of p explanatory variables x_1, x_2, \dots, x_p and a response variable y . Our goal is to study or predict the behavior of y given the values of the explanatory variables.

- For any set of fixed values of the explanatory variables, the response y varies according to a **Normal distribution**. Repeated responses y are **independent** of each other.
- The mean response μ_y has a **linear relationship** given by the **population regression model**

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

The β 's are unknown parameters.

- The **standard deviation** of y (call it σ) is the same for all values of the explanatory variables. The value of σ is unknown.

This model has $p + 2$ parameters that we must estimate from data: the $p + 1$ coefficients $\beta_0, \beta_1, \dots, \beta_p$ and the standard deviation σ .

This is *multiple regression* because there is more than one explanatory variable. Some of the x 's in the model may be interaction terms, products of two explanatory variables. Others may be squares or higher powers of quantitative explanatory variables. So the model can describe quite general relationships.⁵ The main restriction is that the model is *linear regression* because each term is a constant multiple βx . Here are some examples that illustrate the flexibility of multiple regression models.

EXAMPLE 28.14 Two interacting explanatory variables

Suppose we have n observations on two explanatory variables x_1 and x_2 and a response variable y . Our goal is predict the behavior of y for given values of x_1 and x_2 . The mean response is given by

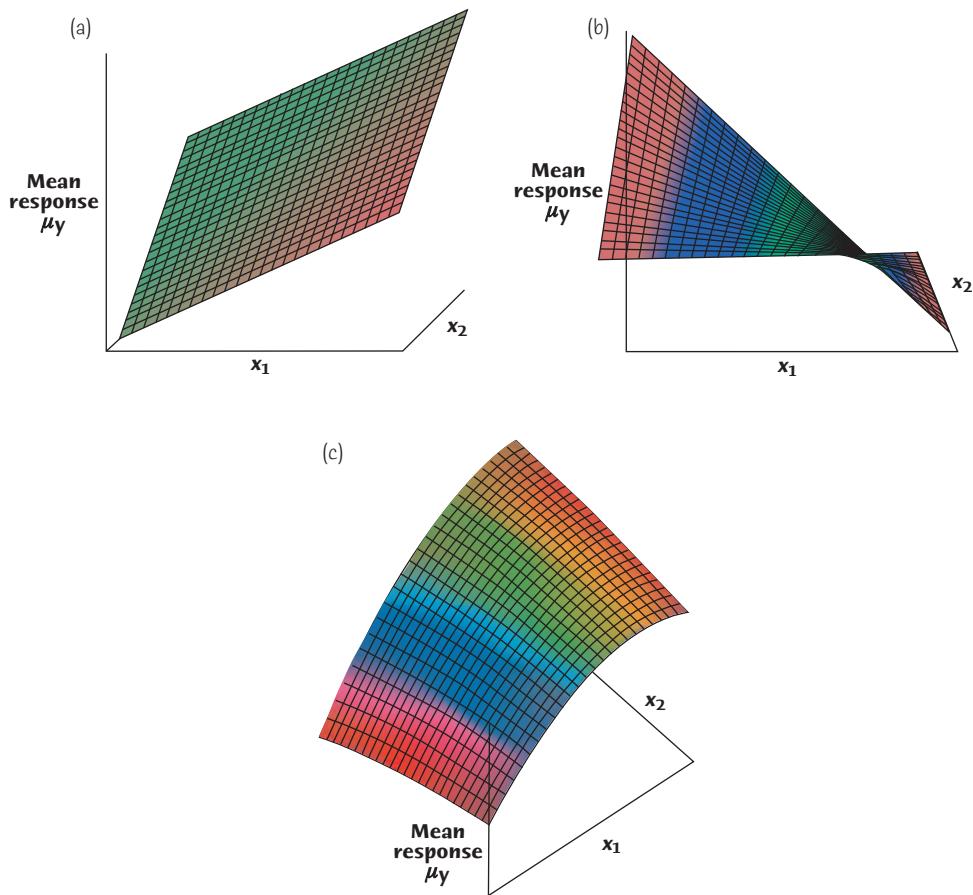
$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

Because there are two explanatory variables x_1 and x_2 , we can graph the relationship of y with x_1 and x_2 in three dimensions. Figure 28.10 shows y vertically above a plane in which x_1 and x_2 take their values. The result is a surface in space. Figure 28.10(a) shows the easiest extension of our simple linear regression model from Chapter 24. Instead of fitting a line to the data, we are now fitting a plane. This figure shows the plane $\mu_y = x_1 + x_2$. The plane is a population model, and when we collect data on our explanatory variables, we will see vertical deviations from the points to the plane. The goal of least-squares regression is to minimize the vertical distances from the points to the plane.

Figure 28.10(b) adds a slight twist. The twist is created by the interaction term in the model. The mean response in Figure 28.10(b) is $\mu_y = 2x_1 + 2x_2 + 10x_1x_2$. The coefficients in front of the explanatory variables indicate part of the effect of a one-unit change on the mean response for each one-unit change in one of the explanatory variables.

FIGURE 28.10

Some possible surfaces for multiple regression models. Figure 28.10(a) shows the plane $\mu_y = x_1 + x_2$. Figure 28.10(b) shows the surface $\mu_y = 2x_1 + 2x_2 + 10x_1x_2$. Figure 28.10(c) shows the surface $\mu_y = 2000 - 20x_1^2 - 2x_1 - 3x_2^2 + 5x_2 + 10x_1x_2$.



But the interpretation of the effect of a one-unit change in the mean response for one variable also depends on the other variable. For example, if $x_2 = 1$, the mean response increases by 12 ($\mu_y = 2 + 12x_1$) for a one-unit increase in x_1 . However, when $x_2 = 2$, the mean response increases by 22 ($\mu_y = 4 + 22x_1$) for a one-unit increase in x_1 . To interpret the parameters in multiple regression models, we think about the impact of one variable on the mean response while all of the other variables are held fixed.

Another way to think about possible multiple regression models for two explanatory variables is take a piece of paper and hold it as shown in Figure 28.10(a). Now begin moving the corners of the paper to get different surfaces. You see that a wide variety of surfaces are possible with only two explanatory variables.

Another possible response surface is shown in Figure 28.10(c). A quick inspection of this figure reveals some curvature in the mean response. To get a curved response surface, add terms for the squares or higher powers of the explanatory variables. The mean response in Figure 28.10(c) is $\mu_y = 2000 - 20x_1^2 - 2x_1 - 3x_2^2 + 5x_2 + 10x_1x_2$. This model has two linear terms, two quadratic terms, and one interaction term. Models of this form are known as second-order polynomial regression models. ■

Software fits the model just as before, estimating the β 's by the least-squares method and estimating σ by the regression standard error based on the residuals. Nothing essential is new, though you will notice different degrees of freedom depending on the number of terms in the model.

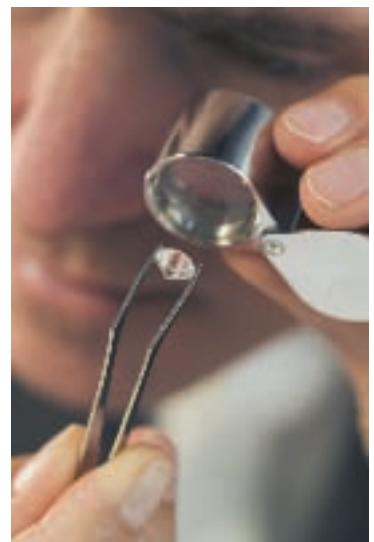
EXAMPLE 28.15 Quadratic regression

If there is a quadratic relationship between a quantitative variable y and another quantitative variable x_1 , the mean response is given by

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_1^2$$

A young couple are shopping for a diamond, so they are interested in learning more about how these gems are priced. They have heard about the 4 C's: carat, color, cut, and clarity. Is there a relationship between these diamond characteristics and the price? Table 28.4 shows records for the first 10 diamonds in a large data base.⁶ The complete data base contains 351 diamonds and is available in the file DIAMONDS.dat. The variables include *Carat*, *Color*, *Clarity*, the *Depth* of the cut, the price per carat *Price/Ct*, and the *Total Price*.

Since the young couple are primarily interested in the price of a diamond, they decide to begin by examining the relationship between *Total Price* and *Carat*. Figure 28.11 shows a scatterplot of *Total Price* versus *Carat*, along with the estimated quadratic



Royalty-Free/CORBIS

TABLE 28.4 Subset of diamond data base

CARAT	COLOR	CLARITY	DEPTH	PRICE/CT	TOTAL PRICE
1.08	E	VS1	68.6	\$6693.3	\$7228.8
0.31	F	VVS1	61.9	3159.0	979.3
0.31	H	VS1	62.1	1755.0	544.1
0.32	F	VVS1	60.8	3159.0	1010.9
0.33	D	IF	60.8	4758.8	1570.4
0.33	G	VVS1	61.5	2895.8	955.6
0.35	F	VS1	62.5	2457.0	860.0
0.35	F	VS1	62.3	2457.0	860.0
0.37	F	VVS1	61.4	3402.0	1258.7
0.38	D	IF	60.0	5062.5	1923.8

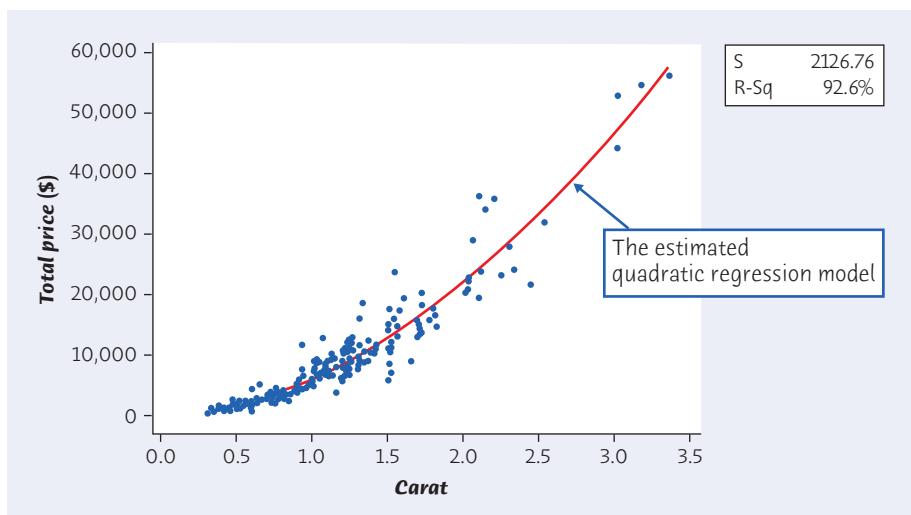
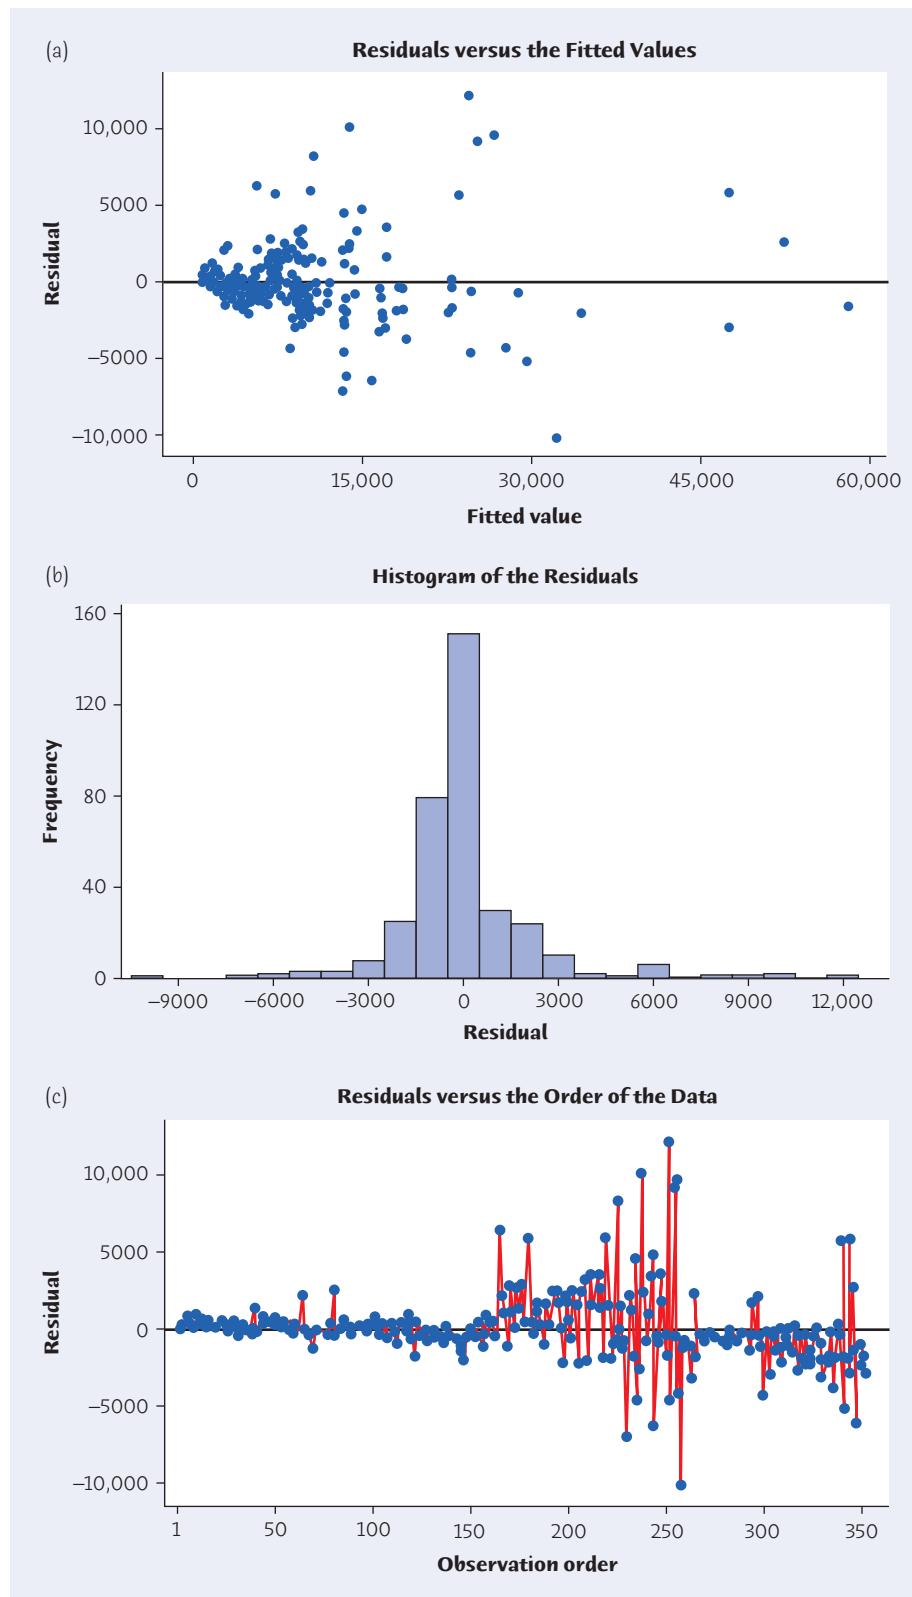


FIGURE 28.11

A scatterplot of *Total Price* versus *Carat* for Example 28.15. The estimated quadratic regression model is also shown.

FIGURE 28.12

Residual plots for the quadratic regression model in Example 28.15.



regression model. Using the quadratic regression model, the couple estimate the mean price of a diamond to be

$$\hat{\mu}_{Price} = -522.7 + 2386Carat + 4498Carat^2$$

The couple are happy because they can explain 92.6% of the variation in the total price of the diamonds in the data base using this quadratic regression model. However, they are concerned because they used explanatory variables that are not independent. An explanatory variable and its square are obviously related to one another. The correlation between $Carat (x_1)$ and $Carat^2 (x_1^2)$ is 0.952.

The residual plots in Figure 28.12 give more reasons for the couple to be concerned. The histogram in Figure 28.12(b) shows that the residuals are roughly symmetric about zero, but the Normal distribution may not be appropriate because of the unusually large and small residuals. The scatterplot of the residuals against the fitted values in Figure 28.12(a) indicates that the variance increases as the fitted value increases up to approximately \$30,000. Finally, the plot of the residuals against order in Figure 28.12(c) does not reveal any troubling pattern, but it does clearly illustrate several unusually large and small residuals.

Having noticed all of the problems with the residual plots, the couple step back and reconsider their objective. They were interested in learning about the relationship between the total price of a diamond and one particular characteristic, carat. The quadratic regression model clearly provides useful information to them even though they will not use this model to make inferences. You will consider additional models to help the couple learn more about diamond pricing in the chapter exercises. ■

APPLY YOUR KNOWLEDGE

28.17 Nest humidity and fleas. In the setting of Exercise 28.7 (page 28-16), researchers showed that the square root of the number of adult fleas y has a quadratic relationship with the nest humidity index x . Specify the population regression model for this situation.

28.18 Diamonds. Specify the population regression model for predicting the total price of a diamond from two interacting variables, $Carat$ and $Depth$ (see Example 28.15 on page 28-35).

28.19 Radioactive decay. An experiment was conducted using a Geiger-Mueller tube in a physics lab. Geiger-Mueller tubes respond to gamma rays and to beta particles (electrons). A pulse that corresponds to each detection of a decay product is produced, and these pulses were counted using a computer-based nuclear counting board. Elapsed time (in seconds) and counts of pulses for a short-lived unstable isotope of silver are shown in Table 28.5 (see page 28-38).⁷  RADIOACTIVITY

- (a) Create a scatterplot of the counts versus time and describe the pattern.
- (b) Since some curvature is apparent in the scatterplot, you might want to consider the quadratic model for predicting counts based on time. Fit the quadratic model and identify the estimated mean response.
- (c) Add the estimated mean response to your scatterplot. Would you recommend the use of the quadratic model for predicting radioactive decay in this situation? Explain.
- (d) Transform the counts using the natural logarithm and create a scatterplot of the transformed variable versus time.

TABLE 28.5 Counts of pulses over time for an unstable isotope of silver

SECONDS	COUNT	SECONDS	COUNT	SECONDS	COUNT	SECONDS	COUNT
20	4611	330	288	640	86	950	13
30	3727	340	331	650	71	960	24
40	3071	350	298	660	77	970	15
50	2587	360	274	670	64	980	13
60	2141	370	289	680	58	990	21
70	1816	380	253	690	48	1000	23
80	1577	390	235	700	58	1010	16
90	1421	400	220	710	57	1020	17
100	1244	410	216	720	55	1030	19
110	1167	420	219	730	50	1040	14
120	992	430	200	740	54	1050	18
130	927	440	170	750	53	1060	10
140	833	450	185	760	38	1070	13
150	811	460	174	770	35	1080	10
160	767	470	163	780	38	1090	11
170	658	480	178	790	28	1100	21
180	656	490	144	800	34	1110	10
190	651	500	147	810	34	1120	10
200	582	510	154	820	32	1130	12
210	530	520	138	830	30	1140	12
220	516	530	140	840	21	1150	11
230	483	540	121	850	33	1160	8
240	500	550	134	860	19	1170	12
250	508	560	105	870	25	1180	13
260	478	570	108	880	30	1190	11
270	425	580	83	890	22	1200	14
280	441	590	104	900	23	1210	11
290	388	600	95	910	28	1220	10
300	382	610	68	920	28	1230	12
310	365	620	85	930	28	1240	8
320	349	630	83	940	19	1250	11

- (e) Fit a simple linear regression model using the natural logarithm of the counts. Provide the estimated regression line, a scatterplot with the estimated regression line, and appropriate residual plots.
- (f) Does the simple linear regression model for the transformed counts fit the data better than the quadratic regression model? Explain.

THE WOES OF REGRESSION COEFFICIENTS

When we start to explore models with several explanatory variables, we quickly meet the big new idea of multiple regression in practice: *the relationship between the response y and any one explanatory variable can change greatly depending on what other explanatory variables are present in the model.* Let's try to understand why this can happen before we illustrate the idea with data.



EXAMPLE 28.16 Coins in your pocket

Let y denote the total amount of change in a person's pocket or purse. Suppose you are interested in modeling this response variable based on two explanatory variables. The first explanatory variable x_1 is the total number of coins in a person's pocket or purse, and the second explanatory variable x_2 is the total number of pennies, nickels, and dimes. Both of these explanatory variables will be positively correlated with the total amount of change in a person's pocket or purse.

Regress y on x_2 alone: we expect the coefficient of x_2 to be positive because the money amount y generally goes up when your pocket has more pennies, nickels, and dimes in it.

Regress y on both x_1 and x_2 : for any fixed x_1 , larger values of x_2 mean fewer quarters in the overall count of coins x_1 , and this means that the money amount y often gets *smaller* as x_2 gets larger. So when we add x_1 to the model, the coefficient of x_2 not only changes but may change sign from positive to negative. ■

The reason for the behavior in Example 28.16 is that the two explanatory variables x_1 and x_2 are related to each other as well as to the response y . When the explanatory variables are correlated, multiple regression models can produce some very odd and counterintuitive results, so we must check carefully for correlation among our potential set of explanatory variables.

For an example with data, let's return to the setting described in Example 28.11 (page 28-26), where we are interested in predicting state average SAT Math scores y based on the percent x_1 of graduates in each state who take the SAT.

EXAMPLE 28.17 Predicting SAT Math scores

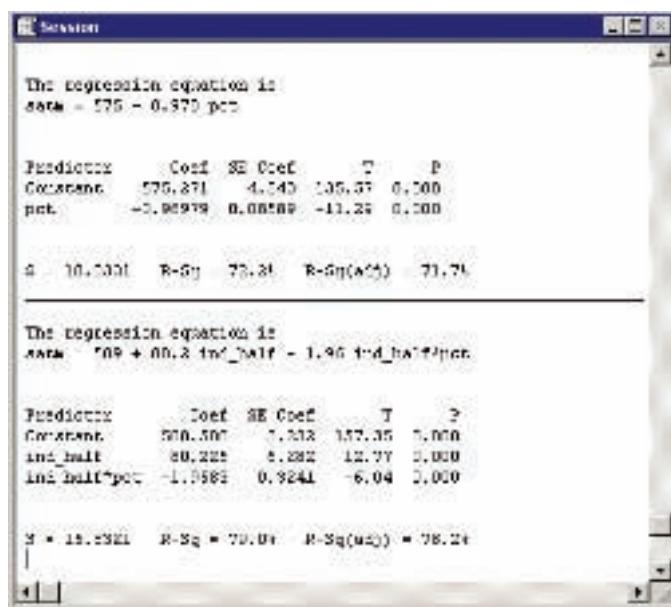
Let's look back at the simple linear regression model (Model 1) and the model with two lines (Model 2) in Example 28.13 (page 28-29). The regression output at the top of Figure 28.13 shows the estimated model $\hat{y} = 575 - 0.97x_1$ and other statistics discussed earlier. Now we focus on the slope, which is close to -1 . The individual t statistic $t = -11.29$ and corresponding P -value (reported as 0.000) clearly indicate that the slope parameter is not equal to 0.

When we add an indicator variable with interaction to fit separate lines for Model 2, the output in Figure 28.8 shows that the individual t statistic for x_1 is $t = -0.70$ and the P -value is 0.488. So we have $b_1 = -0.97$ and highly significant, or $b_1 = -0.2064$ and not at all significant, depending on the other variables present in the model. ■

FIGURE 28.13

Partial regression output for a simple linear regression model and a multiple regression model with indicator variable and interaction term, for Example 28.17.

Minitab



APPLY YOUR KNOWLEDGE

28.20 Predicting SAT Verbal scores. We have been developing models for SAT Math scores for two different clusters of states. Use the SAT data to evaluate similar models for SAT Verbal scores. SATSCORES

- Find the least-squares line for predicting SAT Verbal scores from percent taking the exam.
- Plot SAT Verbal score versus percent taking the exam, and add the least-squares line to your plot.
- Are you happy with the fit of your model? Comment on the value of R^2 and the residual plots.
- Fit a model with two regression lines. Identify the two lines, parameter estimates, t statistics, and corresponding P-values. Does this model improve the fit?
- Specify and fit the model suggested by the inferences for the model in part (d). Identify the two lines, parameter estimates, t statistics, and corresponding P-values. Are you happy with the fit of this model? Explain.

28.21 Body fat for men. You are interested in predicting the amount of body fat on a man y using the explanatory variables waist size x_1 and height x_2 .

- Do you think body fat y and waist size x_1 are positively correlated? Explain.
- For a fixed waist size, height x_2 is negatively correlated with body fat y . Explain why.
- The slope of the simple linear regression line for predicting body fat from height for a sample of men is almost 0, say 0.13. Knowing a man's height does not tell you much about his body fat. Do you think this parameter estimate would become negative if a multiple regression model with height x_2 and waist size x_1 was used to predict body fat? Explain.

28.22 Combining relationships. Suppose that $x_1 = 2x_2 - 4$ so that x_1 and x_2 are positively correlated. Let $y = 3x_2 + 4$ so that y and x_2 are positively correlated.

- Use the relationship between x_1 and x_2 to find the linear relationship between y and x_1 . Are y and x_1 positively correlated?
- Add the equations $x_1 = 2x_2 - 4$ and $y = 3x_2 + 4$ together and solve for y to obtain an equation relating y to both x_1 and x_2 . Are the coefficients of both x 's positive? Combining explanatory variables that are correlated can produce surprising results.

A CASE STUDY FOR MULTIPLE REGRESSION

We will now look at a set of data with several explanatory variables to illustrate the process of arriving at a suitable multiple regression model. In the next section, we will use the model we have chosen for inference, including predicting the response variable.

To build a multiple regression model, first examine the data for outliers and other deviations that might unduly influence your conclusions. Next, use descriptive statistics, especially correlations, to get an idea of which explanatory variables may be most helpful in explaining the response. Fit several models using combinations of these variables, paying attention to the individual t statistics to see if any variables contribute little in any particular model. Always think about the real-world setting of your data and use common sense as part of the process.

EXAMPLE 28.18 Marketing data for a clothing retailer

The data provided in Table 28.6 represent a random sample of 60 customers from a large clothing retailer.⁸ The manager of the store is interested in predicting how much a customer will spend on his or her next purchase.

Our goal is to find a regression model for predicting the amount of a purchase from the available explanatory variables. A short description of each variable is provided below.

Variable	Description
Amount	The net dollar amount spent by customers who made a purchase from this retailer
Recency	The number of months since the last purchase
Freq12	The number of purchases in the last 12 months
Dollar12	The dollar amount of purchases in the last 12 months
Freq24	The number of purchases in the last 24 months
Dollar24	The dollar amount of purchases in the last 24 months
Card	An indicator variable: $Card = 1$ for customers who have a private-label credit card with the retailer, and $Card = 0$ for those who do not



Wide Group/Getty Images

The response variable y is the amount of money spent by a customer. A careful examination of Table 28.6 reveals that the first three values for $Amount$ are zero

TABLE 28.6 Data from clothing retailer

ID	AMOUNT	RECENTY	FREQ12	DOLLAR12	FREQ24	DOLLAR24	CARD
1	0	22	0	0	3	400	0
2	0	30	0	0	0	0	0
3	0	24	0	0	1	250	0
4	30	6	3	140	4	225	0
5	33	12	1	50	1	50	0
6	35	48	0	0	0	0	0
7	35	5	5	450	6	415	0
8	39	2	5	245	12	661	1
9	40	24	0	0	1	225	0
10	45	3	6	403	8	1138	0
11	48	6	3	155	4	262	0
12	50	12	1	42	7	290	0
13	50	5	2	100	8	700	1
14	50	8	3	144	4	202	0
15	50	1	10	562	13	595	1
16	50	2	3	166	4	308	0
17	50	4	4	228	4	228	0
18	50	5	5	322	7	717	1
19	55	13	0	0	6	1050	0
20	55	6	3	244	7	811	0
21	57	20	0	0	2	140	0
22	58	3	4	200	4	818	1
23	60	12	1	70	2	150	0
24	60	3	4	256	7	468	0
25	62	12	1	65	5	255	0
26	64	8	1	70	6	300	0
27	65	2	6	471	8	607	0
28	68	6	2	110	3	150	0
29	70	3	3	222	5	305	0
30	70	6	2	120	4	230	0
31	70	5	3	205	8	455	1
32	72	7	4	445	6	400	0
33	75	6	1	77	2	168	0
34	75	4	2	166	5	404	0
35	75	4	3	210	4	270	0
36	78	8	2	180	7	555	1
37	78	5	3	245	9	602	1
38	79	4	3	225	5	350	0
39	80	3	4	300	6	499	0
40	90	3	5	400	9	723	0
41	95	1	6	650	9	1006	1

TABLE 28.6 (Continued)

ID	AMOUNT	RECENCY	FREQ12	DOLLAR12	FREQ24	DOLLAR24	CARD
42	98	6	2	215	3	333	0
43	100	12	1	100	2	200	0
44	100	2	1	110	4	400	1
45	100	3	3	217	6	605	0
46	100	3	4	330	8	660	1
47	105	2	4	400	7	560	0
48	110	3	4	420	6	570	0
49	125	3	2	270	5	590	1
50	140	6	3	405	6	775	0
51	160	2	2	411	8	706	0
52	180	1	5	744	10	945	1
53	200	1	3	558	4	755	1
54	240	4	4	815	10	1150	1
55	250	3	3	782	10	1500	1
56	300	12	1	250	4	401	0
57	340	1	5	1084	7	1162	1
58	500	4	2	777	3	905	1
59	650	1	4	1493	7	2050	1
60	1,506,000	1	6	5000	11	8000	1

because some customers purchased items and then returned them. We are not interested in modeling returns, so these observations will be removed before proceeding. The last row of Table 28.6 indicates that one customer spent \$1,506,000 in the store. A quick consultation with the manager reveals that this observation is a data entry error, so this customer will also be removed from our analysis. We can now proceed with the cleaned data on 56 customers.

EXAMPLE 28.19 Relationships among the variables

We won't go through all of the expected relationships among the variables, but we would certainly expect the amount of a purchase to be positively associated with the amount of money spent over the last 12 and the last 24 months. Speculating about how the frequency of purchases over the last 12 and 24 months is related to the purchase amount is not as easy. Some customers may buy small amounts of clothing on a regular basis while others may purchase large amounts at less frequent intervals. Yet other people may purchase large amounts on a regular basis.



CLOTHING

Descriptive statistics and a matrix of correlation coefficients for the 6 quantitative variables are shown in Figure 28.14. As expected, *Amount* is strongly correlated with past spending: $r = 0.80368$ with *Dollar12* and $r = 0.67732$ with *Dollar24*. However, the matrix also reveals that these explanatory variables are correlated with one another. Since the variables are dollar amounts in overlapping time periods, there is a strong positive association, $r = 0.82745$, between *Dollar12* and *Dollar24*.

SAS

The CORR Procedure
6 Variables: Amount Recency Freq12 Dollar12 Freq24 Dollar24

Simple Statistics

Variable	N	Mean	Std Dev	Sum	Minimum	Maximum	Label
Amount	56	108.28571	112.18843	6064	30.00000	650.00000	Amount
Recency	56	6.35714	7.29739	356.00000	1.00000	48.00000	Recency
Freq12	56	2.98214	1.86344	167.00000	0	10.00000	Freq12
Dollar12	56	309.26786	283.92915	17319	0	1493	Dollar12
Freq24	56	5.75000	2.74524	322.00000	0	13.00000	Freq24
Dollar24	56	553.55357	379.07941	30999	0	2050	Dollar24

Pearson Correlation Coefficients, N = 56
Prob > |r| under HO: Rho = 0

	Amount	Recency	Freq12	Dollar12	Freq24	Dollar24
Amount	1.00000	-0.22081	0.05160	0.80368	0.10172	0.67732
Amount		0.1020	0.7057	<.0001	0.4557	<.0001
Recency	-0.22081	1.00000	-0.58382	-0.45387	-0.54909	-0.43238
Recency	0.1020		<.0001	0.0004	<.0001	0.0009
Freq12	0.05160	-0.58382	1.00000	0.55586	0.70995	0.42147
Freq12	0.7057	<.0001		<.0001	<.0001	0.0012
Dollar12	0.80368	-0.45387	0.55586	1.00000	0.48495	0.82745
Dollar12	<.0001	0.0004	<.0001		0.0002	<.0001
Freq24	0.10172	-0.54909	0.70995	0.48495	1.00000	0.59622
Freq24	0.4557	<.0001	<.0001	0.0002		<.0001
Dollar24	0.67732	-0.43238	0.42147	0.82745	0.59622	1.00000
Dollar24	<.0001	0.0009	0.0012	<.0001	<.0001	

FIGURE 28.14

Descriptive statistics and correlation coefficients for Example 28.19.

Recency (the number of months since the last purchase) is negatively associated with the purchase amount and with the four explanatory variables that indicate the number of purchases or the amount of those purchases. Perhaps recent customers (low *Recency*) tend to be regular customers and those who have not visited in some time (high *Recency*) include customers who often shop elsewhere. Customers with low *Recency* would then visit more frequently and spend more. ■

One common mistake in modeling is to include too many variables in the multiple regression model, especially variables that are related to one another. A hasty user of statistical software will include all explanatory variables along with some possible interaction terms and quadratic terms. Here's an example to show you what can happen.

EXAMPLE 28.20 Including all explanatory variables

Create the following interaction terms and quadratic terms from the potential explanatory variables:

$$\begin{aligned} \text{Int12} &= \text{Freq12} \times \text{Dollar12} \\ \text{Int24} &= \text{Freq24} \times \text{Dollar24} \\ \text{IntCard12} &= \text{Card} \times \text{Dollar12} \\ \text{Dollar12sq} &= \text{Dollar12} \times \text{Dollar12} \\ \text{Dollar24sq} &= \text{Dollar24} \times \text{Dollar24} \end{aligned}$$



CLOTHING2

Figure 28.15 shows the multiple regression output using all six explanatory variables provided by the manager and the five new variables. Most of the individual *t* statistics

CrunchIt!

Results - Multiple Linear Regression				
Export *				
Fitted Equation:	$\text{Amount} = -0.105244 + 0.913276 * \text{Recency} - 19.8662 * \text{Freq12} + 0.456385 * \text{Dollar12} + 15.0452 * \text{Freq24} + 0.0785828 * \text{Dollar24} - 23.0993 * \text{Card} - 0.0270543 * \text{Int12} - 0.0305059 * \text{Int24} + 0.139565 * \text{IntCard12} - 0.0000486003 * \text{Dollar12sq} + 0.0000700631 * \text{Dollar24sq}$			
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.105244	33.4874	-0.00314281	0.997507
Recency	0.913276	1.09630	0.833056	0.409313
Freq12	-19.8662	10.5040	-1.89131	0.0651779
Dollar12	0.456385	0.105805	4.31345	<0.0001
Freq24	15.0452	7.14569	2.10550	0.0409903
Dollar24	0.0785828	0.0759970	1.03402	0.306775
Card	-23.0993	28.5611	-0.808770	0.422999
Int12	-0.0270543	0.0206868	-1.29529	0.201977
Int24	-0.0305059	0.0106181	-2.87300	0.00623660
IntCard12	0.139565	0.0822976	1.69585	0.0969796
Dollar12sq	-0.0000486003	0.000139699	-0.347893	0.729579
Dollar24sq	0.0000700631	0.0000743761	0.942012	0.351330
r-Squared: 0.916558				
Adjusted r-Squared: 0.895697				
sigma: 36.2324				

FIGURE 28.15

CrunchIt! output for the multiple regression model in Example 28.20.

have P -values greater than 0.2 and only three have P -values less than 0.05. The model is successful at explaining 91.66% of the variation in the purchase amounts, but it is large and unwieldy. Management will have to measure all of these variables to use the model in the future for prediction. This model does set a standard: removing explanatory variables can only reduce R^2 , so no smaller model that uses some of these variables and no new variables can do better than $R^2 = 91.66\%$. But can a simpler model do almost as well? ■

automated algorithms

Some statistical software provides **automated algorithms** to choose regression models. All possible regression algorithms are very useful. On the other hand, algorithms that add or remove variables one at a time often miss good models. We will not illustrate automated algorithms, but will build models by considering and evaluating various possible subsets of models.

EXAMPLE 28.21 Highest correlation

To start, let's look at a simple linear regression model with the single explanatory variable most highly correlated with *Amount*. The correlations in Figure 28.14 show that this explanatory variable is *Dollar12*. The least-squares regression line for predicting the purchase amount y is

$$\hat{y} = 10.0756 + 0.31756Dollar12$$

Figure 28.16 shows the regression output for this simple linear regression model. This simple model has a low R^2 of 64.59%, so we need more explanatory variables. ■

CrunchIt!

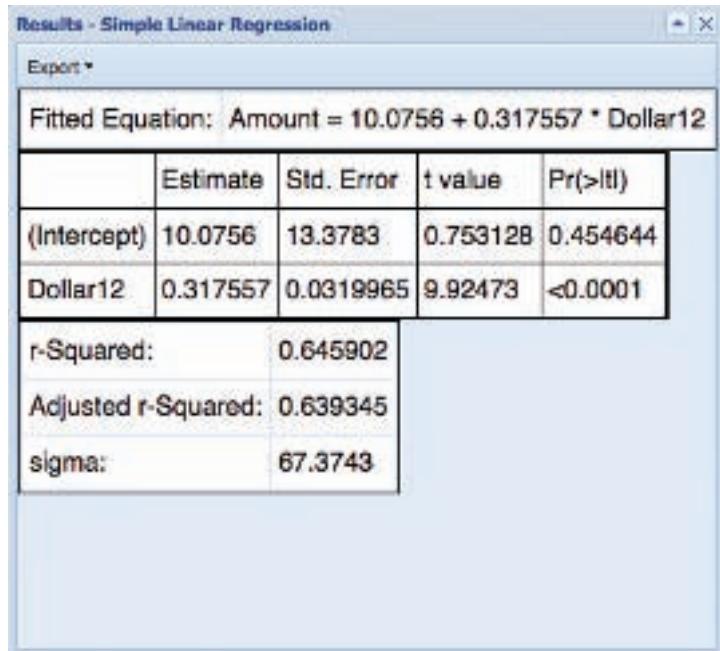


FIGURE 28.16

CrunchIt! output for the simple linear regression model in Example 28.21 using the dollar amount of purchases in the last 12 months (*Dollar12*) as the explanatory variable.

EXAMPLE 28.22 Including other explanatory variables

Of the remaining explanatory variables, *Dollar24* and *Recency* have the strongest associations with the purchase amounts. We will add these variables to try to improve our model. Rather than providing the complete computer output for each model, we will concentrate on the parameter estimates and individual *t* statistics provided in Figure 28.17. The fitted model using both *Dollar12* and *Dollar24* is

$$\hat{y} = 7.63 + 0.30\text{Dollar12} + 0.01\text{Dollar24}$$

The *t* statistic for *Dollar12* has dropped from 9.92 to 5.30, but it is still significant. However, if the amount of the purchases over the last 12 months (*Dollar12*) is already in the model, then adding the amount of purchases over the last 24 months (*Dollar24*) does not improve the model.

CrunchIt!

Results - Multiple Linear Regression				
Export				
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.62619	16.2885	0.468194	0.641566
Dollar12	0.304786	0.0574760	5.30285	<0.0001
Dollar24	0.0115597	0.0430493	0.268523	0.789339
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-17.6985	18.7574	-0.943545	0.349684
Recency	2.78722	1.35728	2.05354	0.0449686
Dollar12	0.350070	0.0348840	10.0353	<0.0001
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	88.7539	16.4472	5.39630	<0.0001
Recency	-1.10472	0.945486	-1.16841	0.247969
Freq12	-36.5015	3.96893	-9.19681	<0.0001
Dollar12	0.437832	0.0237334	18.4479	<0.0001
	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	73.8976	10.4686	7.05898	<0.0001
Freq12	-34.4259	3.56139	-9.66641	<0.0001
Dollar12	0.443146	0.0233735	18.9593	<0.0001

FIGURE 28.17

CrunchIt! parameter estimates and individual *t* statistics for the models in Example 28.22.

Using Recency and *Dollar12* we find the fitted model

$$\hat{y} = -17.7 + 0.35Dollar12 + 2.79Recency$$

Even though the *t* statistics associated with both explanatory variables are significant, the percent of variation in the purchase amounts explained by this model increases only to 67.2%.

The frequency of visits over the last 12 months (*Freq12*) was not strongly associated with the purchase amount, but may be helpful because dollar amount and frequency provide different information. The fitted model using all three explanatory variables is

$$\hat{y} = 88.75 + 0.44Dollar12 - 1.1Recency - 36.5Freq12$$

The *t* statistic for *Dollar12* jumps to 18.45, and the *t* statistic for *Recency* drops to -1.17 , which is not significant. Eliminating *Recency* from the model, we obtain the fitted model

$$\hat{y} = 73.90 + 0.44Dollar12 - 34.43Freq12$$

This model explains 87.51% of the variation in the purchase amounts. That is almost as good as the big clumsy model in Example 28.20, but with only two explanatory variables. We might stop here, but we will take one more approach to the problem. ■

We have used the explanatory variables that were given to us by the manager to fit many different models. However, we have not thought carefully about the data and our objective. Thinking about the setting of the data leads to a new idea.

EXAMPLE 28.23 Creating a new explanatory variable

To predict the purchase amount for a customer, the average purchase over a recent time period might be helpful. We have the total amount and frequency of purchases over 12 months, so we can create a new variable

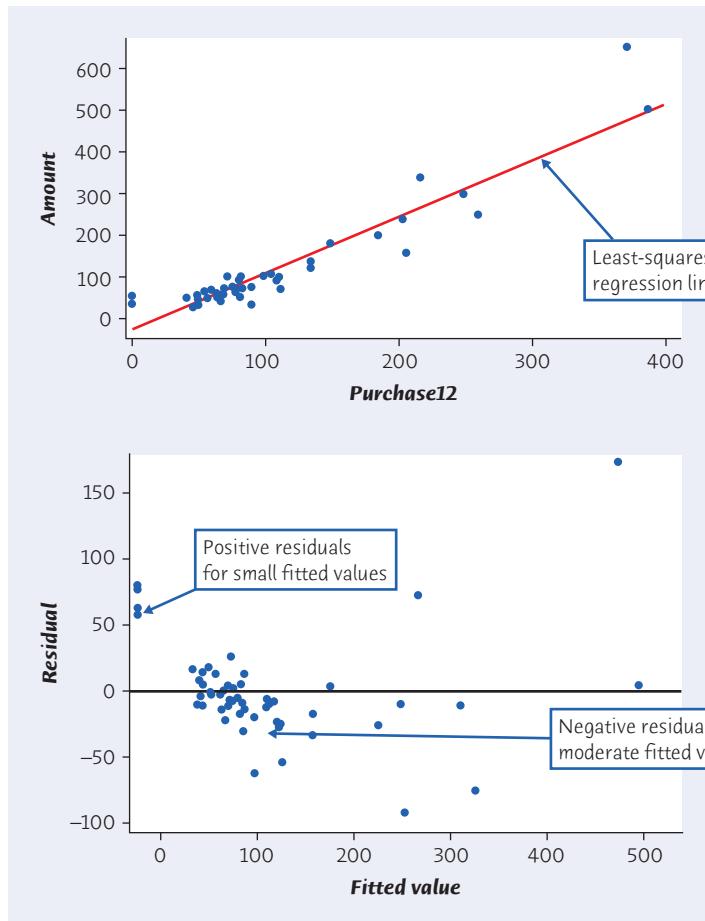
$$Purchase12 = \frac{Dollar12}{Freq12}$$

If no purchases were made in the last 12 months, then *Purchase12* is set to 0. Fitting a simple linear regression model with this new explanatory variable explains 87.64% of the variation in the purchase amounts. This is better than almost all of our previous models. Figure 28.18 shows the fitted model

$$\hat{y} = -22.99 + 1.34Purchase12$$

on a scatterplot of *Amount* versus *Purchase12* and the corresponding residual plot. ■

This new linear model provides a good fit. The residual plot in Figure 28.18 shows that low purchase amounts tend to be above the regression line and moderate purchase amounts tend to be below the line. This suggests that a model with some curvature might improve the fit.

**FIGURE 28.18**

A scatterplot, including the simple linear regression line, and a residual plot for Example 28.23.

EXAMPLE 28.24 A final model

Create the variable *Purchase12sq*, the square of *Purchase12*, to allow some curvature in the model. Previous explorations also revealed that the dollar amount spent depends on how recent the customer visited the store, so an interaction term

$$\text{IntRecency12} = \text{Recency} \times \text{Dollar12}$$

was created to incorporate this relationship into the model. The output for the multiple regression model using the three explanatory variables *Purchase12*, *Purchase12sq*, and *IntRecency12* is shown in Figure 28.19. This model does a great job for the manager by explaining almost 94% of the variation in the purchase amounts. ■

APPLY YOUR KNOWLEDGE

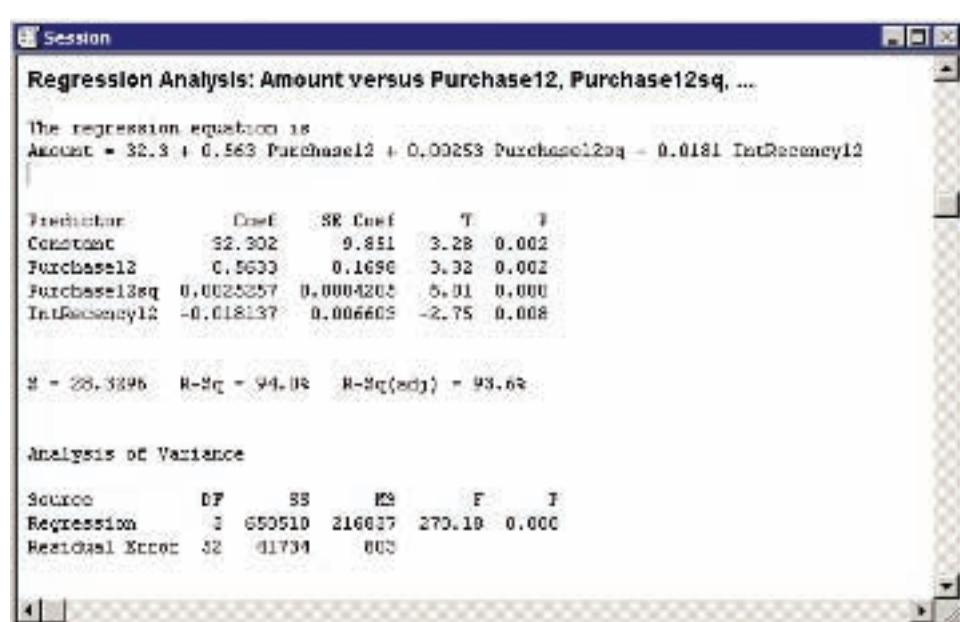
28.23 Diamonds. Suppose that the couple shopping for a diamond in Example 28.15 (page 28-35) had used a quadratic regression model for the other quantitative variable, *Depth*. Use the data in Table 28.4 to answer the following questions. 

- (a) What is the estimated quadratic regression model for mean total price based on the explanatory variable *Depth*?

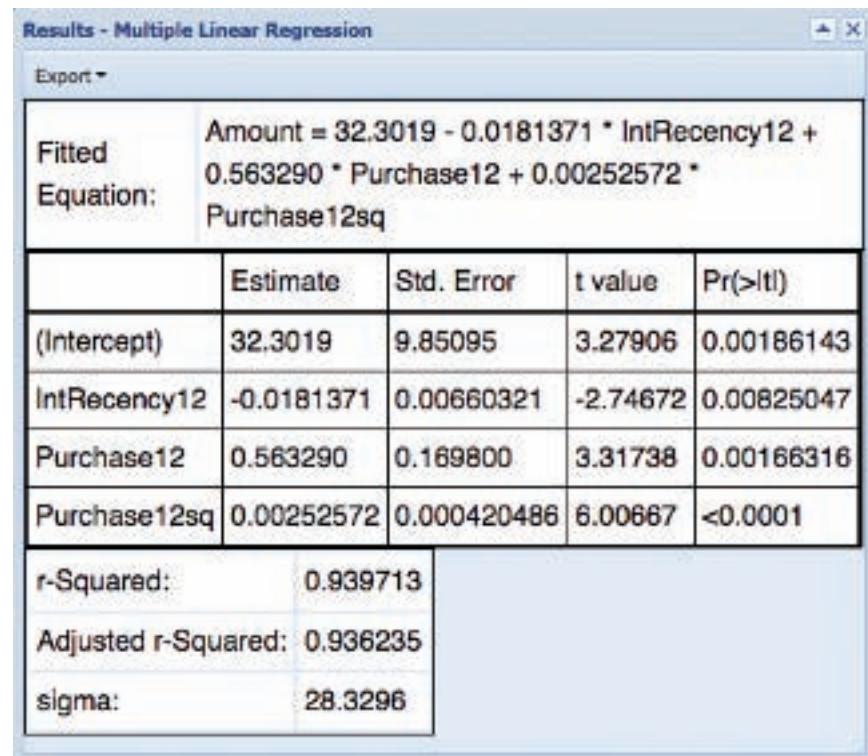
FIGURE 28.19

Minitab output for the multiple regression model in Example 28.24.

Minitab



CrunchIt!



- (b) As you discovered in part (a), it is always possible to fit quadratic models, but we must decide if they are helpful. Is this model as informative to the couple as the model in Example 28.15? What percent of variation in the total price is explained by using the quadratic regression model with *Depth*?

28.24 Tuition and fees at a small liberal arts college. Information regarding tuition and fees at the University of Virginia from 1970 to 2011 is provided in Table 28.7.⁹ Use statistical software to answer the following questions.  TUITION

- Find the simple linear regression equation for predicting tuition and fees from year, and save the residuals and fitted values.
- The value of tuition and fees in 1971 is missing from the data set. Use the least-squares line to estimate this value.
- Does the estimate obtained in part (b) intuitively make sense to you? That is, are you happy with this estimate? Explain.
- Plot the residuals against year. What does the plot tell you about the adequacy of the linear fit?
- Will this linear model overestimate or underestimate the tuition and fees at this college in the 1990s?
- Since the residual plot shows a quadratic trend, it might be helpful to add a quadratic term to this model. Fit the quadratic regression model and provide the estimated model.
- Does the quadratic model provide a better fit than the linear model?
- Would you be willing to make inferences based on the quadratic model? Explain.

TABLE 28.7 Out-of-state tuition and fees (in dollars) at the University of Virginia

YEAR	TUITION AND FEES	YEAR	TUITION AND FEES	YEAR	TUITION AND FEES
1970	1,069	1985	4,886	2000	17,409
1971	n.a.	1986	5,468	2001	18,268
1972	1,372	1987	5,796	2002	19,805
1973	1,447	1988	6,336	2003	21,984
1974	1,569	1989	7,088	2004	22,700
1975	1,619	1990	8,136	2005	24,100
1976	1,819	1991	9,564	2006	25,945
1977	1,939	1992	10,826	2007	27,750
1978	2,024	1993	12,254	2008	29,600
1979	2,159	1994	13,052	2009	31,672
1980	2,402	1995	14,006	2010	33,782
1981	2,646	1996	14,434	2011	36,788
1982	3,276	1997	15,030		
1983	3,766	1998	15,814		
1984	4,336	1999	16,603		

Note: n.a. indicates that data are not available.

28.25 Fish sizes. Table 28.8 contains data on the size of perch caught in a lake in Finland.¹⁰ Use statistical software to help you analyze these data.  PERCH

- Use the multiple regression model with two explanatory variables, length and width, to predict the weight of a perch. Provide the estimated multiple regression equation.
- How much of the variation in the weight of perch is explained by the model in part (a)?
- Does the ANOVA table indicate that at least one of the explanatory variables is helpful in predicting the weight of perch? Explain.
- Do the individual *t* tests indicate that both β_1 and β_2 are significantly different from zero? Explain.

TABLE 28.8 Measurements on 56 perch

OBS. NUMBER	WEIGHT (GRAMS)	LENGTH (CM)	WIDTH (CM)	OBS. NUMBER	WEIGHT (GRAMS)	LENGTH (CM)	WIDTH (CM)
104	5.9	8.8	1.4	132	197.0	27.0	4.2
105	32.0	14.7	2.0	133	218.0	28.0	4.1
106	40.0	16.0	2.4	134	300.0	28.7	5.1
107	51.5	17.2	2.6	135	260.0	28.9	4.3
108	70.0	18.5	2.9	136	265.0	28.9	4.3
109	100.0	19.2	3.3	137	250.0	28.9	4.6
110	78.0	19.4	3.1	138	250.0	29.4	4.2
111	80.0	20.2	3.1	139	300.0	30.1	4.6
112	85.0	20.8	3.0	140	320.0	31.6	4.8
113	85.0	21.0	2.8	141	514.0	34.0	6.0
114	110.0	22.5	3.6	142	556.0	36.5	6.4
115	115.0	22.5	3.3	143	840.0	37.3	7.8
116	125.0	22.5	3.7	144	685.0	39.0	6.9
117	130.0	22.8	3.5	145	700.0	38.3	6.7
118	120.0	23.5	3.4	146	700.0	39.4	6.3
119	120.0	23.5	3.5	147	690.0	39.3	6.4
120	130.0	23.5	3.5	148	900.0	41.4	7.5
121	135.0	23.5	3.5	149	650.0	41.4	6.0
122	110.0	23.5	4.0	150	820.0	41.3	7.4
123	130.0	24.0	3.6	151	850.0	42.3	7.1
124	150.0	24.0	3.6	152	900.0	42.5	7.2
125	145.0	24.2	3.6	153	1015.0	42.4	7.5
126	150.0	24.5	3.6	154	820.0	42.5	6.6
127	170.0	25.0	3.7	155	1100.0	44.6	6.9
128	225.0	25.5	3.7	156	1000.0	45.2	7.3
129	145.0	25.5	3.8	157	1100.0	45.5	7.4
130	188.0	26.2	4.2	158	1000.0	46.0	8.1
131	180.0	26.5	3.7	159	1000.0	46.6	7.6

- (e) Create a new variable, called interaction, that is the product of length and width. Use the multiple regression model with three explanatory variables, length, width, and interaction, to predict the weight of a perch. Provide the estimated multiple regression equation.
 - (f) How much of the variation in the weight of perch is explained by the model in part (e)?
 - (g) Does the ANOVA table indicate that at least one of the explanatory variables is helpful in predicting the weight of perch? Explain.
 - (h) Describe how the individual t statistics changed when the interaction term was added.
-

INFEERENCE FOR REGRESSION PARAMETERS

We discussed the general form of inference procedures for regression parameters earlier in the chapter, using software output. This section provides more details for the analysis of variance (ANOVA) table, the F test, and the individual t statistics for the multiple regression model with p explanatory variables, $\mu_y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \cdots + \beta_px_p$.

Software always provides the ANOVA table. The general form of the ANOVA table is shown below.

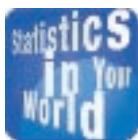
Degrees				
Source	of freedom	Sum of squares	Mean square	F statistic
Model	p	$SSM = \sum (\hat{y} - \bar{y})^2$	$MSM = \frac{SSM}{p}$	$F = \frac{MSM}{MSE}$
Error	$n - p - 1$	$SSE = \sum (y - \hat{y})^2$	$MSE = \frac{SSE}{n - p - 1}$	
Total	$n - 1$	$\sum (y - \bar{y})^2$		

EXAMPLE 28.25 A quick check

The final multiple regression model for the clothing retailer data in Example 28.24 is

$$\mu_y = \beta_0 + \beta_1x_1 + \beta_2x_2 + \beta_3x_3$$

where $x_1 = Purchase12$, $x_2 = Purchase12sq$, and $x_3 = IntRecency$. It is a good idea to check that the degrees of freedom from the ANOVA table on the output match the form above. This verifies that the software is using the number of observations and the number of explanatory variables that you intended. The model degrees of freedom is the number of explanatory variables, 3, and the total degrees of freedom (degrees of freedom for the model plus degrees of freedom for error) is the number of observations minus 1, $56 - 1 = 55$. We usually do not check the other calculations by hand, but knowing that the mean sum of squares is the sum of squares divided by the degrees of freedom and that the F statistic is the ratio of the mean sum of squares for each source helps us understand how the F statistic is formed. ■



Do good looks mean good money?

Experienced researchers who

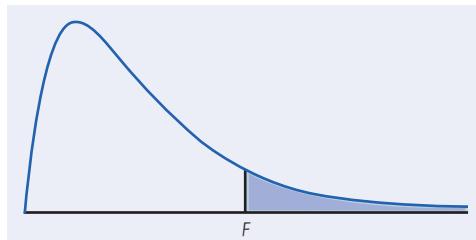
have spent decades studying physical attractiveness suggest that good looks translate into good money. In particular, studies suggest that “plain people earn 5% to 10% less than people of average looks, who in turn earn 3% to 8% less than those deemed good-looking.” Other studies suggest that size is important also, with tall people earning considerably more over their careers than short people. Before you take a look in the mirror, it is important to understand that hiring managers say that the appearance of confidence is more attractive to them than physical beauty.

The first formal test in most multiple regression studies is the ANOVA F test. This test is used to check if the complete set of explanatory variables is helpful in predicting the response variable.

ANALYSIS OF VARIANCE F TEST

The analysis of variance F statistic tests the null hypothesis that all the regression coefficients (β 's) except β_0 are equal to zero. The test statistic is

$$F = \frac{\text{variation due to model}}{\text{variation due to error}}$$



P -values come from the F distribution with p and $n - p - 1$ degrees of freedom.

To give formulas for the numerator and denominator of the F statistic, let \hat{y} stand for predicted values and let \bar{y} be the average of the response observations. The numerator of F is the mean square for the model:

$$\text{variation due to model} = \frac{\Sigma(\hat{y} - \bar{y})^2}{p}$$

The denominator of F is the mean square for error:

$$\text{variation due to error} = \frac{\Sigma(y - \hat{y})^2}{n - p - 1}$$

The P -value for a test of H_0 against the alternative that at least one β parameter is not zero is the area to the right of F under an $F(p, n - p - 1)$ distribution.

EXAMPLE 28.26 Any useful predictors?

The ANOVA table in Figure 28.19 (page 28-50) shows an F statistic of 270.18. The P -value provided on the output is the area to the right of 270.18 under an F distribution with 3 numerator and 52 denominator degrees of freedom. Since this area is so small (<0.001), we reject the hypothesis that the β coefficients associated with the three explanatory variables are all equal to zero. The three explanatory variables together do help predict the response. ■

As we have seen, individual t tests are helpful in identifying the explanatory variables that are useful predictors, but extreme caution is necessary when interpreting the results of these tests. Remember that an individual t assesses the contribution of its variable *in the presence of the other variables in this specific model*. That is, individual t 's depend on the model in use, not just on the direct association between an explanatory variable and the response.

CONFIDENCE INTERVALS AND INDIVIDUAL t TESTS FOR COEFFICIENTS

A level C confidence interval for a regression coefficient β is $b \pm t^*SE_b$.

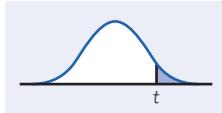
The critical value t^* is obtained from the $t(n - p - 1)$ distribution.

The t statistic for testing the null hypothesis that a regression coefficient β is equal to zero has the form

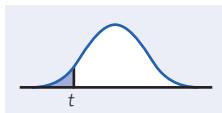
$$t = \frac{\text{parameter estimate}}{\text{standard error of estimate}} = \frac{b}{SE_b}$$

In terms of a random variable T having the $t(n - p - 1)$ distribution, the P -value for a test of H_0 against

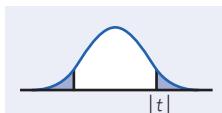
$$H_a: \beta > 0 \text{ is } P(T \geq t)$$



$$H_a: \beta < 0 \text{ is } P(T \leq t)$$



$$H_a: \beta \neq 0 \text{ is } 2P(T \geq |t|)$$



EXAMPLE 28.27 The easiest situation: all predictors are helpful

The individual t statistics and corresponding P -values in Figure 28.19 (page 28-50) indicate that all three of the explanatory variables are useful predictors. All the P -values are below 0.01, which indicates very convincing evidence of statistical significance. The P -values are computed using a t distribution with 52 degrees of freedom. The degrees of freedom for error in the ANOVA table will always tell you which t distribution to use for the individual β coefficients. ■

The main purpose of most regression models is prediction. Construction of *confidence intervals for a mean response* and *prediction intervals for a future observation* with multiple regression models is similar to the methods we used for simple linear regression. The main difference is that we must now specify a list of values for all of the explanatory variables in the model. As we learned in Chapter 24, the additional uncertainty in predicting future observations will result in prediction intervals that are wider than confidence intervals.

CONFIDENCE AND PREDICTION INTERVALS FOR MULTIPLE REGRESSION RESPONSE

A level C confidence interval for the mean response μ_y is $\hat{y} \pm t^*SE_{\hat{\mu}}$.

A level C prediction interval for a single response y is $\hat{y} \pm t^*SE_{\hat{y}}$.

In both intervals, t^* is the critical value for the $t(n - p - 1)$ density curve with area C between $-t^*$ and t^* .

EXAMPLE 28.28 Predicting means and future clothing purchases

Figure 28.20 provides the predicted values, 95% confidence limits for the mean purchase amount, and 95% prediction limits for a future purchase amount for each of the 56 observations in Table 28.6. The values of the explanatory variables don't appear, but they are needed to obtain the predicted values \hat{y} and the endpoints of the intervals. As expected, the prediction intervals for future purchase amounts are always wider than the confidence intervals for the mean purchase amounts. You can also see that predicting future purchase amounts, even with a good model, is not an easy task. Several of the prediction intervals (for Observations 1 to 3, for example) include purchase amounts below zero. The manager will not give customers money for coming to the store, so the lower endpoint of the prediction intervals should be zero for practical purposes. ■

APPLY YOUR KNOWLEDGE

28.26 World record running times. Exercise 28.15 (page 28-31) shows the progress of world record times (in seconds) for the 10,000-meter run for both men and women.  **WORLDRECORDS**

- Provide the ANOVA table for the regression model with two regression lines, one for men and one for women.
- Are all the individual coefficients significantly different from zero? Set up the appropriate hypotheses, identify the test statistics and P -values, and make conclusions in the context of the problem.

28.27 Fish sizes. Use explanatory variables length, width, and interaction from Exercise 28.25 (page 28-52) on the 56 perch to provide 95% confidence intervals for the mean and prediction intervals for future observations. Interpret both intervals for the 10th perch in the data set. What t distribution is used to provide both intervals?  **PERCH**

28.28 Clothing retailer. Since the average purchase amount $Purchase12$ was such a good predictor, the manager would like you to consider another explanatory variable: the average purchase amount from the previous 12 months. Create the new variable  **CLOTHING2**

$$Purchase12b = \frac{Dollar24 - Dollar12}{Freq24 - Freq12}$$

and add it to the final model obtained in Example 28.24 (page 28-49).

- What is R^2 for this model? How does this value compare with R^2 in Example 28.24?
- What is the value of the individual t statistic for this new explanatory variable? How much did the individual t statistics change from their previous values?
- Would you recommend this model over the model in Example 28.24? Explain.

SAS

Output

Observation	Predicted value	Lower confidence limit	Upper confidence limit	Lower prediction limit	Upper prediction limit
1	48.854	38.688	59.020	-8.895	106.603
2	55.898	46.389	65.408	-1.739	113.536
3	32.302	12.535	52.069	-27.884	92.488
4	62.648	44.743	80.552	3.047	122.248
5	57.080	47.212	66.949	-0.618	114.778
6	32.302	12.535	52.069	-27.884	92.488
7	59.603	50.162	69.044	1.977	117.229
8	51.280	41.540	61.019	-6.396	108.956
9	51.274	40.684	61.865	-6.551	109.100
10	57.712	47.957	67.467	0.034	115.391
11	44.265	32.626	55.905	-13.762	102.292
12	61.743	52.592	70.894	4.164	119.323
13	65.182	54.487	75.577	7.392	122.972
14	56.074	47.078	65.071	-1.481	113.630
15	49.852	36.572	63.132	-8.526	108.230
16	32.302	12.535	52.069	-27.884	92.488
17	68.271	57.911	78.631	10.487	126.055
18	32.302	12.535	52.069	-27.884	92.488
19	55.898	46.389	65.408	-1.739	113.536
20	68.873	60.745	77.002	11.447	126.299
21	64.769	56.435	73.102	7.313	122.224
22	65.440	57.153	73.727	7.992	122.888
23	73.951	64.798	83.105	16.372	131.531
24	74.999	66.872	83.126	17.574	132.425
25	58.953	49.973	67.932	1.400	116.505
26	75.737	66.997	84.477	18.221	133.252
27	62.133	53.562	70.704	4.643	119.623
28	63.997	55.576	72.417	6.529	121.464
29	69.730	42.980	96.481	6.903	132.558
30	82.271	71.943	92.599	24.493	140.049
31	84.412	75.037	93.787	26.796	142.027
32	68.873	60.745	77.002	11.447	126.299
33	77.339	67.427	87.251	19.634	135.044
34	72.931	64.150	81.712	15.409	130.453
35	72.432	64.343	80.521	15.012	129.853
36	72.432	64.343	80.521	15.012	129.853
37	71.765	63.073	80.457	14.257	129.273
38	111.178	98.769	123.587	52.992	169.364
39	98.647	88.821	108.472	40.956	156.337
40	92.124	82.895	101.352	34.532	149.715
41	120.835	104.098	137.571	61.575	180.095
42	74.454	65.713	83.195	16.938	131.970
43	78.008	69.779	86.238	20.568	135.448
44	99.378	89.069	109.688	41.604	157.153
45	96.441	86.821	106.061	38.785	154.096
46	139.686	125.150	154.223	81.010	198.363
47	110.304	92.452	128.156	50.719	169.889
48	239.811	218.604	261.019	179.137	300.486
49	158.549	141.885	175.212	99.309	217.788
50	214.333	192.262	236.404	153.351	275.315
51	192.798	168.575	217.022	131.005	254.591
52	308.199	289.938	326.459	248.490	367.907
53	276.571	255.327	297.814	215.883	337.258
54	253.477	233.630	273.324	193.265	313.690
55	575.983	534.751	617.214	505.757	646.208
56	567.343	529.980	604.707	499.316	635.371

FIGURE 28.20

Predicted values, confidence limits for the mean purchase amount, and prediction limits for a future purchase amount, for Example 28.28.

CHECKING THE CONDITIONS FOR INFERENCE

A full picture of the conditions for multiple regression requires much more than a few plots of the residuals. We will present only a few methods here, because regression diagnostics is a subject that could fill an entire book.

Plot the response variable against each of the explanatory variables. These plots help you explore and understand potential relationships. Multiple regression models allow curvature and other interesting features that are not simple to check visually, especially when we get beyond two explanatory variables.

Plot the residuals against the predicted values and against all of the explanatory variables in the model. These plots will allow you to check the condition that the standard deviation of the response about the multiple regression model is the same everywhere. They should show an unstructured horizontal band of points centered at 0. The mean of the residuals is always 0, just as in simple linear regression, so we continue to add a line at 0 to orient ourselves. Funnel or cone shapes indicate that this condition is not met and that the standard deviation of the residuals must be stabilized before making inferences. Other patterns in residual plots can sometimes be fixed by changing the model. For example, if you see a quadratic pattern, then you should consider adding a quadratic term for that explanatory variable.

Look for outliers and influential observations in all residual plots. To check the influence of a particular observation, you can fit your model with and without this observation. If the estimates and statistics do not change much, you can safely proceed. However, if there are substantial changes, you must begin a more careful investigation. Do not simply throw out observations to improve the fit and increase R^2 .

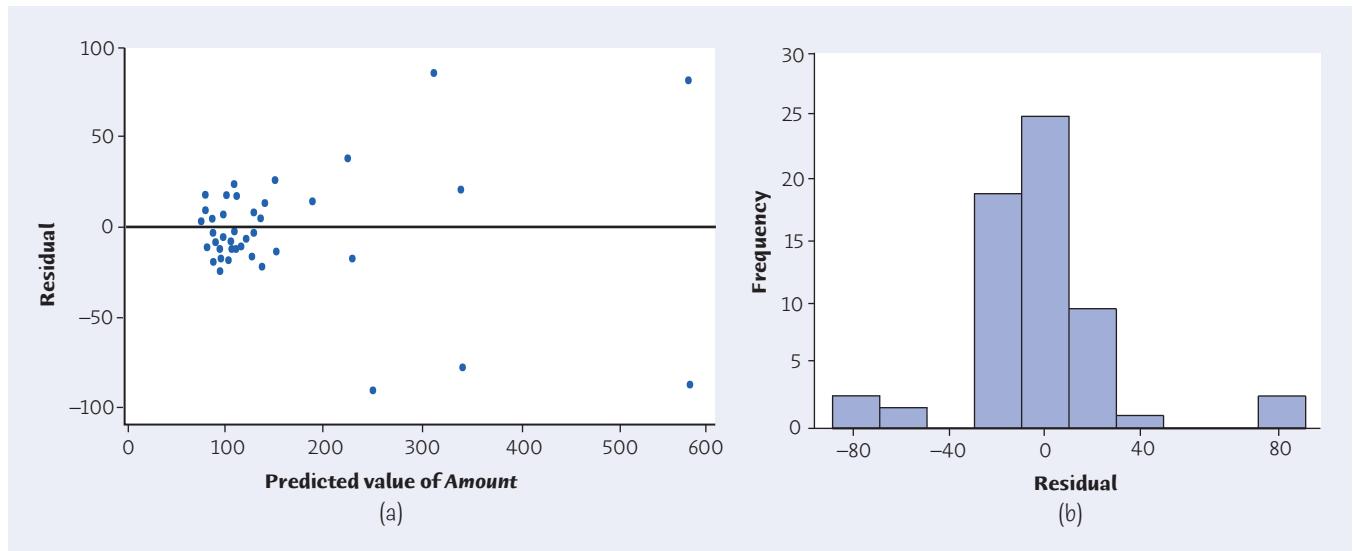
Ideally, we would like all of the explanatory variables to be independent and the observations on the response variable to be independent. As you have seen in this chapter, practical problems include explanatory variables that are not independent. Association between two or more explanatory variables can create serious problems in the model, so use correlations and scatterplots to check relationships.

To check the condition that the response should vary Normally about the multiple regression model, *make a histogram or stemplot of the residuals.* We can rely on the robustness of the regression methods when there is a slight departure from Normality, except for prediction intervals. As in the case of simple linear regression, we view prediction intervals from multiple regression models as rough approximations.

EXAMPLE 28.29 Checking conditions

Figure 28.21 shows residual plots for the final model in Example 28.24 (page 28-49). The scatterplot shows that the variability for the larger predicted values is greater than the variability for the predicted values below 200. The constant-variance condition is not satisfied. Since most of the predicted values are below 200 and the variability is roughly constant in that range, we will not resort to more sophisticated methods to stabilize the variance.

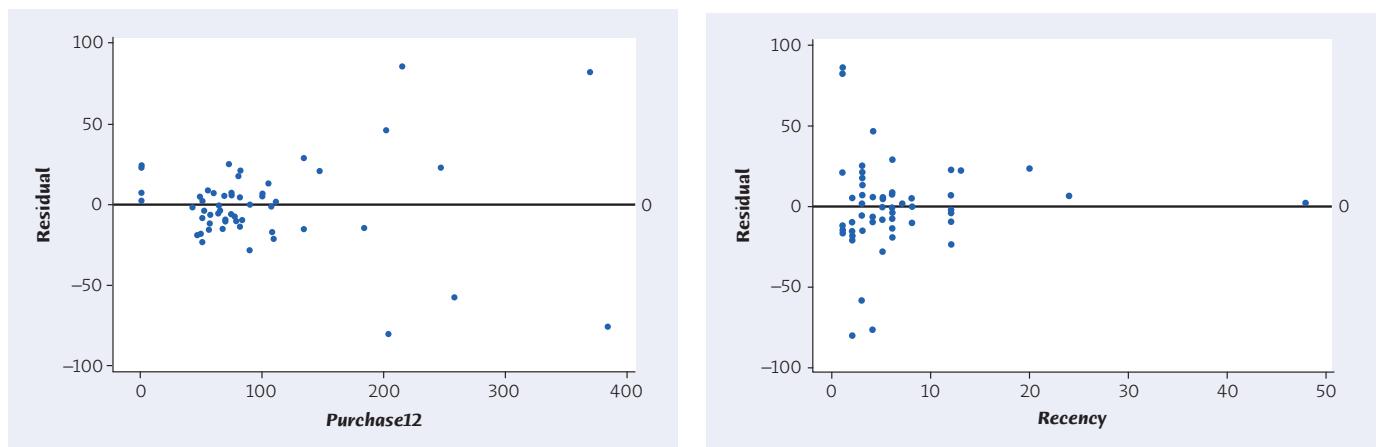
The histogram shows approximate perfect symmetry in the residuals. The residuals above 75 and below -75 are apparent on the scatterplot and the histogram. This is a situation where we need to rely on the robustness of regression inference when there are slight departures from Normality. ■

**FIGURE 28.21**

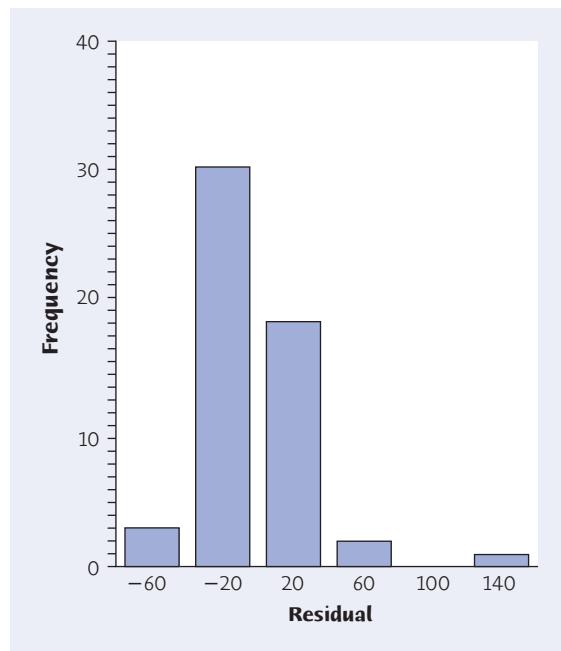
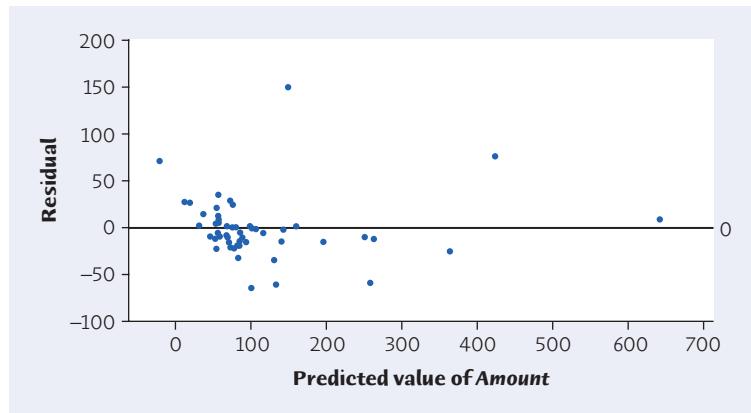
Residual plots for the multiple regression model in Example 28.24. Figure 28.21(a) is a scatterplot of the residuals against the predicted values. Figure 28.21(b) is a histogram of the residuals.

APPLY YOUR KNOWLEDGE

28.29 Final model for the clothing retailer problem. The residual plots below show the residuals for the final model in the clothing retailer problem plotted against *Purchase12* and *Recency*. Do the plots suggest any potential problems with the conditions for inference? Comment.



28.30 The clothing retailer problem. The scatterplot and histogram below show the residuals from the model in Example 28.20 with all explanatory variables, some interaction terms, and quadratic terms. Comment on both plots. Do you see any reason for concern in using this model for inference?



CHAPTER 28 SUMMARY

CHAPTER SPECIFICS

- An **indicator variable** x_2 can be used to fit a regression model with **two parallel lines**. The mean response is

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2$$

where x_1 is an explanatory variable.

- A multiple regression model with **two regression lines** includes an explanatory variable x_1 , an indicator variable x_2 , and an interaction term $x_1 x_2$. The mean response is

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2$$

- The mean response μ_y for a general **multiple regression model** based on p explanatory variables x_1, x_2, \dots, x_p is

$$\mu_y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

- The **estimated regression model** is

$$\hat{y} = b_0 + b_1 x_1 + b_2 x_2 + \cdots + b_p x_p$$

where the b 's are obtained by the method of least squares.

- The **regression standard error** s has $n - p - 1$ degrees of freedom and is used to estimate σ .
- The **sum of squares row** in the **analysis of variance (ANOVA) table** breaks the total variability in the responses into two pieces. One piece summarizes the variability due to the model, and the other piece summarizes the variability due to error:

$$\text{total sum of squares} = \text{model sum of squares} + \text{error sum of squares}$$

- The **squared multiple correlation coefficient R^2** represents the proportion of variability in the response variable y that is explained by the explanatory variables x_1, x_2, \dots, x_p in a multiple regression model.
- To test the hypothesis that all the regression coefficients (β 's), except β_0 , are equal to zero, use the **ANOVA F statistic**. In other words, the **null model** says that the x 's do not help predict y . The alternative is that the explanatory variables as a group are helpful in predicting y .
- **Individual t procedures** in regression inference have $n - p - 1$ degrees of freedom. These individual t procedures depend on the other explanatory variables specified in a multiple regression model. Individual t tests assess the contribution of one explanatory variable in the presence of the other variables in a model. The null hypothesis is written as $H_0: \beta = 0$ but interpreted as "the coefficient of x is 0 in this model."
- **Confidence intervals** for the mean response μ_y have the form $\hat{y} \pm t^* \text{SE}_{\hat{\mu}}$. **Prediction intervals** for individual future responses y have the form $\hat{y} \pm t^* \text{SE}_{\hat{y}}$.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Preliminaries

1. Examine the data for outliers and other deviations that might influence your conclusions.
2. Use descriptive statistics, especially correlations, to get an idea of which explanatory variables may be most helpful in explaining the response.
3. Make scatterplots to examine the relationships between explanatory variables and a response variable.
4. Use software to compute a correlation matrix to explore the relationships between pairs of variables.

B. Recognition

1. Recognize when a multiple regression model with parallel regression lines is appropriate.
2. Recognize when an interaction term needs to be added to fit a multiple regression model with two separate regression lines.
3. Recognize when a multiple regression model with several explanatory variables is appropriate.
4. Recognize the difference between the overall F test and the individual t tests.
5. Recognize that the parameter estimates, t statistics, and P -values for each explanatory variable depend on the specific model.
6. Inspect the data to recognize situations in which inference isn't safe: influential observations, strongly skewed residuals in a small sample, or nonconstant variation of the data points about the regression model.

C. Inference Using Software

1. Use software to find the estimated multiple regression model.
2. Explain the meaning of the regression parameters (β 's) in any specific multiple regression model.
3. Understand the software output for regression. Find the regression standard error, the squared multiple correlation coefficient R^2 , and the overall F test and P -value. Identify the parameter estimates, standard errors, individual t tests, and P -values.
4. Use that information to carry out tests and calculate confidence intervals for the β 's.
5. Use R^2 and residual plots to assess the fit of a model.
6. Choose a model by comparing R^2 -values, regression standard errors, and individual t statistics.
7. Explain the distinction between a confidence interval for the mean response and a prediction interval for an individual response.

LINK IT

Chapters 4, 5, and 24 discuss scatterplots, correlation, and regression. In these chapters we studied how to use a single explanatory variable to predict a response, although in Chapter 4 we saw how to incorporate a categorical variable into a scatterplot. In this chapter, we extend the ideas of Chapters 4, 5, and 24 and learn how to use several

explanatory variables to predict a response. The multiple regression model is similar to the simple linear regression model, but with more explanatory variables. The conditions for inference, the methods for estimating and testing hypotheses about regression coefficients, prediction, and checking the conditions for inference are much like those discussed in Chapter 24. New concepts include the notion of an interaction, deciding which of several candidate regression models is best, and interpreting parameter estimates when several explanatory variables are included. We will encounter some of these new concepts again in Chapter 29.

CHECK YOUR SKILLS

Many exercise bikes, elliptical trainers, and treadmills display basic information like distance, speed, calories burned per hour (or total calories), and duration of the workout. The data in Table 28.9 show the treadmill display's claimed calories per hour by speed for a 175-pound male using a Cybex treadmill at inclines of 0%, 2%, and 4%.

The relationship between speed and calories is different for walking and running, so we need an indicator for slow/fast. The variables created from Table 28.9 are

TABLE 28.9 Cybex treadmill display's claimed calories per hour by speed and incline for a 175-pound man

Mph	INCLINE		
	0%	2%	4%
1.5	174	207	240
2.0	205	249	294
2.5	236	291	347
3.0	267	333	400
3.5	372	436	503
4.0	482	542	607
4.5	592	649	709
5.0	701	756	812
5.5	763	824	885
6.0	825	892	959
6.5	887	960	1032
7.0	949	1027	1105
7.5	1011	1094	1178
8.0	1073	1163	1252
8.5	1135	1230	1325
9.0	1197	1298	1398
9.5	1259	1365	1470
10.0	1321	1433	1544

Calories = calories burned per hour

Mph = speed of the treadmill

Incline = the incline percent (0, 2, or 4)

IndSlow = 1 for Mph ≤ 3 and IndSlow = 0 for Mph > 3.0

Here is part of the Minitab output from fitting a multiple regression model to predict Calories from Mph, IndSlow, and Incline for the Cybex treadmill. Exercises 28.31 to 28.40 are based on this output.

Minitab

```

Session
Predictor      Coef    SE Coef      T      P
Constant     -80.41    18.99   -4.24  0.000
Mph          145.841   2.570   56.74  0.000
IndSlow      -50.01    16.04   -3.12  0.003
Incline       36.264    2.829   12.82  0.000

S = 33.9422    R-Sq = 99.3%    R-Sq(adj) = 99.3%

Analysis of Variance

Source        Df      SS      MS      F      P
Regression    3   8554241  2851414  2475.03  0.000
Residual Error 50   57604   1152
Total         53   8611845

Predicted Values for New Observations
New
Obs      Fit    SE Fit      95% CI           95% PI
1   940.09    5.28  (929.49, 950.69)  (871.09, 1009.08)

Values of Predictors for New Observations
New
Obs      Mph    IndSlow    Incline
1   6.50    0.000000    2.00

```

28.31 The number of parameters in this multiple regression model is

- (a) 4. (b) 5. (c) 6.

28.32 The equation for predicting calories from these explanatory variables is

- (a) $\text{Calories} = -80.41 + 145.84\text{Mph} - 50.01\text{IndSlow} + 36.26\text{Incline.}$
 (b) $\text{Calories} = -4.24 + 56.74\text{Mph} - 3.12\text{IndSlow} + 12.82\text{Incline.}$
 (c) $\text{Calories} = 18.99 + 2.57\text{Mph} + 16.04\text{IndSlow} + 2.83\text{Incline.}$

28.33 The regression standard error for these data is

- (a) 0.993. (b) 33.94. (c) 1152.

28.34 To predict calories when walking ($\text{Mph} \leq 3$) with no incline use the line

- (a) $-80.41 + 145.84\text{Mph.}$
 (b) $(-80.41 - 50.01) + 145.84\text{Mph.}$
 (c) $[-80.41 + (2 \times 36.26)] + 145.84\text{Mph.}$

28.35 To predict calories when running ($\text{Mph} > 3$) with no incline use the line

- (a) $-80.41 + 145.84\text{Mph.}$
 (b) $(-80.41 - 50.01) + 145.84\text{Mph.}$
 (c) $[-80.41 + (2 \times 36.26)] + 145.84\text{Mph.}$

28.36 To predict calories when running on a 2% incline use the line

- (a) $-80.41 + 145.84\text{Mph.}$

- (b) $(-80.41 - 50.01) + 145.84\text{Mph.}$
 (c) $[-80.41 + (2 \times 36.26)] + 145.84\text{Mph.}$

28.37 Is there significant evidence that more calories are burned for higher speeds? To answer this question, test the hypotheses

- (a) $H_0: \beta_0 = 0$ versus $H_a: \beta_0 > 0.$
 (b) $H_0: \beta_1 = 0$ versus $H_a: \beta_1 > 0.$
 (c) $H_0: \beta_1 = 0$ versus $H_a: \beta_1 \neq 0.$

28.38 Confidence intervals and tests for these data use the t distribution with degrees of freedom

- (a) 3. (b) 50. (c) 53.

28.39 Orlando, a 175-pound man, plans to run 6.5 miles per hour for one hour on a 2% incline. He can be 95% confident that he will burn between

- (a) 871 and 1009 calories.
 (b) 929 and 950 calories.
 (c) 906 and 974 calories.

28.40 Suppose that we also have data on a second treadmill, made by LifeFitness. An indicator variable for brand of treadmill, say $Treadmill = 1$ for Cybex and $Treadmill = 0$ for LifeFitness, is created for a new model. If the three explanatory variables above and the new indicator variable $Treadmill$ were used to predict Calories , how many β parameters would need to be estimated in the new multiple regression model?

- (a) 4 (b) 5 (c) 6

CHAPTER 28 EXERCISES

28.41 A computer game. A multimedia statistics learning system includes a test of skill in using the computer's mouse. The software displays a circle at a random location on the computer screen. The subject clicks in the circle with the mouse as quickly as possible. A new circle appears as soon as the subject clicks the old one. Table 5.3 (text page 155) gives data for one subject's trials, 20 with each hand. Distance is the distance from the cursor location to the center of the new circle, in units whose actual size depends on the size of the screen. Time is the time required to click in the new circle, in milliseconds.¹¹  COMPUTERGAME

(a) Specify the population multiple regression model for predicting time from distance separately for each hand. Make sure you include the interaction term that is necessary to allow for the possibility of having different slopes. Explain in words what each β in your model means.

(b) Use statistical software to find the estimated multiple regression equation for predicting time from distance separately for each hand. What percent of variation in the distances is explained by this multiple regression model?

(c) Explain how to use the estimated multiple regression equation in part (b) to obtain the least-squares line for each hand. Draw these lines on a scatterplot of time versus distance.

28.42 Bank wages and length of service. We assume that our wages will increase as we gain experience and become more valuable to our employers. Wages also increase because of inflation. By examining a sample of employees at a given point in time, we can look at part of the picture. How does length of service (LOS) relate to wages? Table 28.10 gives data on the LOS in months and wages for 60 women who work in Indiana banks. Wages are yearly total income divided

TABLE 28.10 Bank wages, length of service, and bank size

WAGES	LOS	SIZE	WAGES	LOS	SIZE	WAGES	LOS	SIZE
48.3355	94	Large	64.1026	24	Large	41.2088	97	Small
49.0279	48	Small	54.9451	222	Small	67.9096	228	Small
40.8817	102	Small	43.8095	58	Large	43.0942	27	Large
36.5854	20	Small	43.3455	41	Small	40.7000	48	Small
46.7596	60	Large	61.9893	153	Large	40.5748	7	Large
59.5238	78	Small	40.0183	16	Small	39.6825	74	Small
39.1304	45	Large	50.7143	43	Small	50.1742	204	Large
39.2465	39	Large	48.8400	96	Large	54.9451	24	Large
40.2037	20	Large	34.3407	98	Large	32.3822	13	Small
38.1563	65	Small	80.5861	150	Large	51.7130	30	Large
50.0905	76	Large	33.7163	124	Small	55.8379	95	Large
46.9043	48	Small	60.3792	60	Large	54.9451	104	Large
43.1894	61	Small	48.8400	7	Large	70.2786	34	Large
60.5637	30	Large	38.5579	22	Small	57.2344	184	Small
97.6801	70	Large	39.2760	57	Large	54.1126	156	Small
48.5795	108	Large	47.6564	78	Large	39.8687	25	Large
67.1551	61	Large	44.6864	36	Large	27.4725	43	Small
38.7847	10	Small	45.7875	83	Small	67.9584	36	Large
51.8926	68	Large	65.6288	66	Large	44.9317	60	Small
51.8326	54	Large	33.5775	47	Small	51.5612	102	Large

by the number of weeks worked. We have multiplied wages by a constant for reasons of confidentiality.¹² 

- (a) Plot wages versus LOS using different symbols for size of the bank. There is one woman with relatively high wages for her length of service. Circle this point and do not use it in the rest of this exercise.
- (b) Would you be willing to use a multiple regression model with parallel slopes to predict wages from LOS for the two different bank sizes? Explain.
- (c) Fit a model that will allow you to test the hypothesis that the slope of the regression line for small banks is equal to the slope of the regression line for large banks. Conduct the test for equal slopes.
- (d) Are the conditions for inference met for your model in part (c)? Construct appropriate residual plots and comment.

28.43 Mean annual temperatures for two California cities. Table 28.11 contains data on the mean annual temperatures (degrees Fahrenheit) for the years 1951 to 2005 at two locations in California: Pasadena and Redding.¹³ 

(a) Plot the temperatures versus year using different symbols for the two cities.

(b) Would you be willing to use a multiple regression model with parallel slopes to predict temperatures from year for the two different cities? Explain.

(c) Fit a model that will allow you to test the hypothesis that the slope of the regression line for Pasadena is equal to the slope of the regression line for Redding. Conduct the test for equal slopes.

(d) Are the conditions for inference met for your model in part (c)? Construct appropriate residual plots and comment.

28.44 Growth of pine trees. The Department of Biology at Kenyon College conducted an experiment to study the growth of pine trees. In April 1990, volunteers planted 1000 white pine (*Pinus strobus*) seedlings at the Brown Family Environmental Center. The seedlings were planted in two grids, distinguished by 10- and 15-foot spacings between the seedlings. Table 28.12 (page 28-67) shows the first 10 rows of a subset of the data collected by students at Kenyon College.¹⁴ 

TABLE 28.11 Mean annual temperatures (°F) in two California cities

MEAN TEMPERATURE			MEAN TEMPERATURE			MEAN TEMPERATURE		
YEAR	PASADENA	REDDING	YEAR	PASADENA	REDDING	YEAR	PASADENA	REDDING
1951	62.27	62.02	1970	64.08	64.30	1989	64.53	61.50
1952	61.59	62.27	1971	63.59	62.23	1990	64.96	62.22
1953	62.64	62.06	1972	64.53	63.06	1991	65.60	62.73
1954	62.88	61.65	1973	63.46	63.75	1992	66.07	63.59
1955	61.75	62.48	1974	63.93	63.80	1993	65.16	61.55
1956	62.93	63.17	1975	62.36	62.66	1994	64.63	61.63
1957	63.72	62.42	1976	64.23	63.51	1995	65.43	62.62
1958	65.02	64.42	1977	64.47	63.89	1996	65.76	62.93
1959	65.69	65.04	1978	64.21	64.05	1997	66.72	62.48
1960	64.48	63.07	1979	63.76	60.38	1998	64.12	60.23
1961	64.12	63.50	1980	65.02	60.04	1999	64.85	61.88
1962	62.82	63.97	1981	65.80	61.95	2000	66.25	61.58
1963	63.71	62.42	1982	63.50	59.14	2001	64.96	63.03
1964	62.76	63.29	1983	64.19	60.66	2002	65.10	63.28
1965	63.03	63.32	1984	66.06	61.72	2003	66.31	63.13
1966	64.25	64.51	1985	64.44	60.50	2004	65.71	63.57
1967	64.36	64.21	1986	65.31	61.76	2005	57.06	62.62
1968	64.15	63.40	1987	64.58	62.94			
1969	63.51	63.77	1988	65.22	63.70			

Variable	Description
Row	Row number in pine plantation
Col	Column number in pine plantation
Hgt90	Tree height at time of planting (cm)
Hgt96	Tree height in September 1996 (cm)
Diam96	Tree trunk diameter in September 1996 (cm)
Grow96	Leader growth during 1996 (cm)
Hgt97	Tree height in September 1997 (cm)
Diam97	Tree trunk diameter in September 1997 (cm)
Spread97	Widest lateral spread in September 1997 (cm)
Needles97	Needle length in September 1997 (mm)
Deer95	Type of deer damage in September 1995: 1 = none, 2 = browsed
Deer97	Type of deer damage in September 1997: 1 = none, 2 = browsed
Cover95	Amount of thorny cover in September 1995: 0 = none, 1 = <1/3, 2 = between 1/3 and 2/3, 3 = >2/3
Fert	Indicator for fertilizer: 0 = no, 1 = yes
Spacing	Distance (in feet) between trees (10 or 15)

- (a) Use tree height at the time of planting ($Hgt90$) and the indicator variable for fertilizer ($Fert$) to fit a multiple regression model for predicting $Hgt97$. Specify the estimated regression model and the regression standard error. Are you happy with the fit of this model? Comment on the value of R^2 and the plot of the residuals against the predicted values.
- (b) Construct a correlation matrix with $Hgt90$, $Hgt96$, $Diam96$, $Grow96$, $Hgt97$, $Diam97$, $Spread97$, and $Needles97$. Which variable is most strongly correlated with the response variable of interest ($Hgt97$)? Does this make sense to you?
- (c) Add tree height in September 1996 ($Hgt96$) to the model in part (a). Does this model do a better job of predicting tree height in 1997? Explain.
- (d) What happened to the individual t statistic for $Hgt90$ when $Hgt96$ was added to the model? Explain why this change occurred.
- (e) Fit a multiple regression model for predicting $Hgt97$ based on the explanatory variables $Diam97$, $Hgt96$, and $Fert$. Summarize the results of the individual t tests. Does this model provide a better fit than the previous models? Explain by comparing the values of R^2 and s for each model.
- (f) Does the parameter estimate for the variable indicating whether a tree was fertilized or not have the sign you expected? Explain. (Experiments can produce surprising results!)

TABLE 28.12 Measurements on pine seedlings at Brown Family Environmental Center

ROW	COL	HGT90	HGT96	DIAM96	GROW96	HGT97	DIAM97	SPREAD97	NEEDLES97	DEER95	COVER95	FERT	SPACING
1	1	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	0	0	15
1	2	14.0	284.0	4.2	96.0	362	6.60	162	66.0	0	1	2	0
1	3	17.0	387.0	7.4	110.0	442	9.30	250	77.0	0	0	1	0
1	4	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	n.a.	0	0	15
1	5	24.0	294.0	3.9	70.0	369	7.00	176	72.0	0	0	2	0
1	6	22.0	310.0	5.6	84.0	365	6.90	215	76.0	0	0	1	0
1	7	18.0	318.0	5.4	96.0	356	7.60	238	74.5	0	0	0	15
1	8	32.0	328.0	5.4	88.0	365	7.70	219	60.5	0	0	1	0
1	9	n.a.	157.0	1.3	64.0	208	2.00	127	56.0	1	1	2	0
1	10	22.0	282.0	4.5	83.0	329	6.10	209	79.5	0	1	2	1

Note: n.a. indicates that data are not available.

(g) Do you think that the model in part (e) should be used for predicting growth in other pine seedlings? Think carefully about the conditions for inference.

28.45 Heating a home. The Sanchez household is about to install solar panels to reduce the cost of heating their house. In order to know how much the solar panels help, they record their consumption of natural gas before the solar panels are installed. Gas consumption is higher in cold weather, so the relationship between outside temperature and gas consumption is important. Here are the data for 16 consecutive months:¹⁵



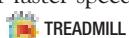
Month	Nov.	Dec.	Jan.	Feb.	Mar.	Apr.	May	June
Degree-days	24	51	43	33	26	13	4	0
Gas used	6.3	10.9	8.9	7.5	5.3	4.0	1.7	1.2
Month	July	Aug.	Sept.	Oct.	Nov.	Dec.	Jan.	Feb.
Degree-days	0	1	6	12	30	32	52	30
Gas used	1.2	1.2	2.1	3.1	6.4	7.2	11.0	6.9

Outside temperature is recorded in degree-days, a common measure of demand for heating. A day's degree-days are the number of degrees its average temperature falls below 65°F. Gas used is recorded in hundreds of cubic feet.

- (a) Create an indicator variable, say *INDwinter*, which is 1 for the months of November, December, January, and February. Make a plot of all the data using a different symbol for winter months.
- (b) Fit the model with two regression lines, one for winter months and one for other months, and identify the estimated regression lines.

(c) Do you think that two regression lines were needed to explain the relationship between gas used and degree-days? Explain.

28.46 Burning calories with exercise. Many exercise bikes, elliptical trainers, and treadmills display basic information like distance, speed, calories burned per hour (or total calories), and duration of the workout. Let's take another look at the data in Table 28.9 (page 28-63) that were used for the Check Your Skills exercises. Scatterplots show different linear relationships for each incline, one for slow speeds and another for faster speeds, so the following indicator variables were created:



$$\text{IndSlow} = 1 \text{ for } \text{Mph} \leq 3 \text{ and } \text{IndSlow} = 0 \text{ for } \text{Mph} > 3.0$$

$$\text{NoIncline} = 1 \text{ for } 0\% \text{ incline and } \text{NoIncline} = 0 \text{ for other inclines}$$

$$\text{2\%Incline} = 1 \text{ for a } 2\% \text{ incline and } \text{2\%Incline} = 0 \text{ for other inclines}$$

Below is part of the Minitab output from fitting a multiple regression model to predict *Calories* from *Mph*, *IndSlow*, *NoIncline*, and *2\%Incline* for the Cybex.

- (a) Use the Minitab output to estimate each parameter in this multiple regression model for predicting calories burned with the Cybex machine. Don't forget to estimate σ .
- (b) How many separate lines are fitted with this model? Do the lines all have the same slope? Identify each fitted line.
- (c) Do you think that this model provides a good fit for these data? Explain.
- (d) Is there significant evidence that more calories are burned for higher speeds? State the hypotheses, identify the test statistic and *P*-value, and provide a conclusion in the context of this question.

Minitab

Session

Regression Analysis: Calories versus Mph, IndSlow, NoIncline, 2%Incline

Predictor	Coef	SE Coef	T	P
Constant	64.75	19.46	3.33	0.002
Mph	145.841	2.596	56.17	0.000
IndSlow	-50.01	16.20	-3.09	0.003
NoIncline	-145.06	11.43	-12.69	0.000
2%Incline	-72.83	11.43	-6.37	0.000

S = 34.2865 R-Sq = 99.3% R-Sq(adj) = 99.3%

Analysis of Variance

Source	Df	SS	MS	F	P
Regression	4	8554242	2138561	1819.18	0.000
Residual Error	49	57603	1176		
Total	53	8611845			

28.47 Burning calories with exercise. Table 28.13 provides data on speed and calories burned per hour for a 175-pound male using two different treadmills (a Cybex and a LifeFitness) at inclines of 0%, 2%, and 4%. 

- (a) Create a scatterplot of calories against miles per hour using six different plotting symbols, one for each combination of incline level and machine.
- (b) Create an indicator variable for brand of treadmill, say *Treadmill* = 1 for Cybex and *Treadmill* = 0 for LifeFitness. Fit a multiple regression model to predict *Calories* from *Mph*, *IndSlow*, *NoIncline*, *2%Incline*, and *Treadmill*.
- (c) Does the model provide a good fit for these data? Explain.
- (d) Is there a significant difference in the relationship between calories and speed for the two different treadmills?

28.48 Metabolic rate and body mass. Metabolic rate, the rate at which the body consumes energy, is important in studies of weight gain, dieting, and exercise. The accompanying table gives data on the lean body mass and resting metabolic rate for 12 women and 7 men who are subjects in a study of dieting. Lean body mass, given in kilograms, is a person's weight leaving out all fat. Metabolic rate is measured in calories burned per 24 hours, the same calories used to describe the energy

content of foods. The researchers believe that lean body mass is an important influence on metabolic rate. 

Subject	Sex	Mass	Rate	Subject	Sex	Mass	Rate
1	M	62.0	1792	11	F	40.3	1189
2	M	62.9	1666	12	F	33.1	913
3	F	36.1	995	13	M	51.9	1460
4	F	54.6	1425	14	F	42.4	1124
5	F	48.5	1396	15	F	34.5	1052
6	F	42.0	1418	16	F	51.1	1347
7	M	47.4	1362	17	F	41.2	1204
8	F	50.6	1502	18	M	51.9	1867
9	F	42.0	1256	19	M	46.9	1439
10	M	48.7	1614				

- (a) Make a scatterplot of the data, using different symbols or colors for men and women. Summarize what you see in the plot.
- (b) Use the model with two regression lines to predict metabolic rate from lean body mass for the different genders. Summarize the results.

TABLE 28.13 Treadmill display's claimed calories per hour by speed for a 175-pound man

<i>Mph</i>	INCLINE			INCLINE		
	CYBEX-0%	CYBEX-2%	CYBEX-4%	LIFE-0%	LIFE-2%	LIFE-4%
1.5	174	207	240	178	212	246
2.0	205	249	294	210	256	301
2.5	236	291	347	243	300	356
3.0	267	333	400	276	343	411
3.5	372	436	503	308	387	466
4.0	482	542	607	341	431	522
4.5	592	649	709	667	718	769
5.0	701	756	812	732	789	845
5.5	763	824	885	797	860	922
6.0	825	892	959	863	930	998
6.5	887	960	1032	928	1015	1075
7.0	949	1027	1105	993	1072	1151
7.5	1011	1094	1178	1058	1143	1228
8.0	1073	1163	1252	1123	1214	1304
8.5	1135	1230	1325	1189	1285	1381
9.0	1197	1298	1398	1254	1356	1457
9.5	1259	1365	1470	1319	1426	1534
10.0	1321	1433	1544	1384	1497	1610

(c) The parameter associated with the interaction term is often used to decide if a model with parallel regression lines can be used. Test the hypothesis that this parameter is equal to zero, and comment on whether or not you would be willing to use the more restrictive model with parallel regression lines for these data.

28.49 Student achievement and self-concept. In order to determine if student achievement is related to self-concept, as measured by the Piers-Harris Children's Self-Concept Scale, data were collected on 78 seventh-grade students from a rural midwestern school. Table 28.14 shows the records for the first 10 students on the following variables:¹⁶



Variable	Description
OBS	Observation number ($n = 78$, some gaps in numbers)
GPA	GPA from school records
IQ	IQ test score from school records
AGE	Age in years, self-reported
GENDER	1 = F, 2 = M, self-reported
RAW	Raw score on Piers-Harris Children's Self-Concept Scale
C1	Cluster 1 within self-concept: behavior
C2	Cluster 2: school status
C3	Cluster 3: physical
C4	Cluster 4: anxiety
C5	Cluster 5: popularity
C6	Cluster 6: happiness

We will investigate the relationship between GPA and only three of the explanatory variables:

- IQ, the student's score on a standard IQ test
- C2, the student's self-assessment of his or her school status
- C5, the student's self-assessment of his or her popularity

Use statistical software to analyze the relationship between students' GPA and their IQ, self-assessed school status (C2), and self-assessed popularity (C5).

(a) One observation is an extreme outlier when all three explanatory variables are used. Which observation is this? Give the observation number and explain how you found it using regression output. Find this observation in the data list. What is unusual about it?

(b) Software packages often identify unusual or influential observations. Have any observations been identified as unusual or influential? If so, identify these points on a scatterplot of GPA versus IQ.

(c) C2 (school status) is the aspect of self-concept most highly correlated to GPA. If we carried out the simple linear regression of GPA on C2, what percent of the variation in students' GPAs would be explained by the straight-line relationship between GPA and C2?

(d) You know that IQ is associated with GPA, and you are not studying that relationship. Because C2 and IQ are positively correlated ($r = 0.547$), a significant relationship between C2 and GPA might occur just because C2 can "stand in" for IQ. Does C2 still contribute significantly to explaining GPA after we have allowed for the relationship between GPA and IQ? (Give a test statistic, its P -value, and your conclusion.)

TABLE 28.14 Student achievement and self-concept scores data for 78 seventh-grade students

OBS	GPA	IQ	AGE	GENDER	RAW	C1	C2	C3	C4	C5	C6
001	7.940	111	13	2	67	15	17	13	13	11	9
002	8.292	107	12	2	43	12	12	7	7	6	6
003	4.643	100	13	2	52	11	10	5	8	9	7
004	7.470	107	12	2	66	14	15	11	11	9	9
005	8.882	114	12	1	58	14	15	10	12	11	6
006	7.585	115	12	2	51	14	11	7	8	6	9
007	7.650	111	13	2	71	15	17	12	14	11	10
008	2.412	97	13	2	51	10	12	5	11	5	6
009	6.000	100	13	1	49	12	9	6	9	6	7
010	8.833	112	13	2	51	15	16	4	9	5	8

- (e) A new student in this class has $IQ = 115$ and $C2 = 14$. What do you predict this student's GPA to be? (Just give a point prediction, not an interval.)

28.50 Children's perception of reading difficulty. Table 28.15 contains measured and self-estimated reading ability data for 60 fifth-grade students randomly sampled from one elementary school.¹⁷ The variables are  **READING**

TABLE 28.15 Measured and self-estimated reading ability data for 60 fifth-grade students randomly sampled from one elementary school

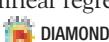
OBS	SEX	LSS	IQ	READ	EST	OBS	SEX	LSS	IQ	READ	EST
1	F	5.00	145	98	4	31	M	7.00	106	55	4
2	F	8.00	139	98	5	32	M	6.00	124	70	4
3	M	6.00	126	90	5	33	M	8.00	115	82	5
4	F	5.33	122	98	5	34	M	8.40	133	94	5
5	F	5.60	125	55	4	35	F	5.00	116	75	4
6	M	9.00	130	95	3	36	F	6.66	102	80	3
7	M	5.00	96	50	4	37	F	5.00	127	85	4
8	M	4.66	110	50	4	38	M	6.50	117	88	5
9	F	4.66	118	75	4	39	F	5.00	109	70	3
10	F	8.20	118	75	5	40	M	5.50	137	80	4
11	M	4.66	101	65	4	41	M	6.66	117	55	4
12	M	7.50	142	68	5	42	M	6.00	90	65	2
13	F	5.00	134	80	4	43	F	4.00	103	30	1
14	M	7.00	124	10	4	44	F	5.50	114	74	5
15	M	6.00	112	67	4	45	M	5.00	139	80	5
16	M	6.00	109	83	3	46	M	6.66	101	70	2
17	F	5.33	134	90	4	47	F	8.33	122	60	4
18	M	6.00	113	90	5	48	F	6.50	105	45	2
19	M	6.00	81	55	3	49	F	4.00	97	45	1
20	F	6.00	113	83	4	50	M	5.50	89	55	4
21	M	6.00	123	65	4	51	M	5.00	102	30	2
22	F	4.66	94	25	3	52	F	4.00	108	10	4
23	M	4.50	100	45	3	53	M	4.66	110	40	1
24	F	6.00	136	97	4	54	M	5.33	128	65	1
25	M	5.33	109	75	4	55	M	5.20	114	15	2
26	F	3.60	131	70	4	56	M	4.00	112	62	2
27	M	4.00	117	23	3	57	F	3.60	114	98	4
28	M	6.40	110	45	3	58	M	6.00	102	52	2
29	F	6.00	127	70	2	59	F	4.60	82	23	1
30	F	6.00	124	85	5	60	M	5.33	101	35	2

Variable	Description
OBS	Observation number for each individual
SEX	Gender of the individual
LSS	Median grade level of student's selection of "best for me to read" (8 repetitions, each with four choices at grades 3, 5, 7, and 9 level)
IQ	IQ score
READ	Score on reading subtest of the Metropolitan Achievement Test
EST	Student's own estimate of his or her reading ability, scale 1 to 5 (1 = low)

- (a) Is the relationship between measured (*READ*) and self-estimated (*EST*) reading ability the same for both boys and girls? Create an indicator variable for gender and fit an appropriate multiple regression model to answer the question.
- (b) Fit a multiple regression model for predicting *IQ* from the explanatory variables *LSS*, *READ*, and *EST*. Are you happy with the fit of this model? Explain.
- (c) Use residual plots to check the appropriate conditions for your model.
- (d) Only two of the three explanatory variables in your model in part (b) have parameters that are significantly different from zero according to the individual *t* tests. Drop the explanatory variable that is not significant, and add the interaction term for the two remaining explanatory variables. Are you surprised by the results from fitting this new model? Explain what happened to the individual *t* tests for the two explanatory variables.

28.51 Florida real estate. The table on text page 614 gives the appraised market values and actual selling prices (in thousands of dollars) of condominium units sold in a beachfront building over a 93-month period. 

- (a) Find the multiple regression model for predicting selling price from appraised market value and month.
- (b) Find and interpret the squared multiple correlation coefficient for your model.
- (c) What is the regression standard error for this model?
- (d) Hamada owns a unit in this building appraised at \$802,600. Use your model to predict the selling price for Hamada's unit.
- (e) Plot the residuals for your model against both explanatory variables and comment on the appearance of these plots.

28.52 Diamonds. Consider the diamond data of which Table 28.4 (page 28-35) is an excerpt. We are interested in predicting the total price of a diamond. Fit a simple linear regression model using *Carat* as the explanatory variable. 

(a) Identify the least-squares line for predicting *Total Price* from *Carat*.

(b) Does the model provide a good fit? Comment on the residual plots. How much variation in price can be explained with this regression line?

(c) Create a new variable *Caratsq* = *Carat* × *Carat*. Fit a quadratic model using *Carat* and *Caratsq* and verify that your estimates for each parameter match those provided in Example 28.15 (page 28-35).

(d) Does the quadratic term *Caratsq* improve the fit of the model? Comment on the residual plots and the value of R^2 .

(e) The individual *t* statistics look at the contribution of each variable when the other variables are in the model. State and test the hypotheses of interest for the quadratic term in your model.

28.53 Diamonds. Use the data in Table 28.4 (page 28-35) to fit the multiple regression model with two explanatory variables, *Carat* and *Depth*, to predict the *Total Price* of diamonds. Don't forget to include the interaction term in your model. 

- (a) Identify the estimated multiple regression equation.
- (b) Conduct the overall *F* test for the model.
- (c) Identify the estimated regression parameters, standard errors, and *t* statistics with *P*-values.
- (d) Prepare residuals plots and comment on whether the conditions for inference are satisfied.
- (e) What percent of variation in *Total Price* is explained by this model?
- (f) Find an estimate for σ and interpret this value.

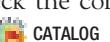
28.54 Catalog spending. This realistic modeling project requires much more time than a typical exercise. Table 28.16 shows catalog-spending data for the first 9 of 200 randomly selected individuals from a very large (over 20,000 households) data base.¹⁸ We are interested in developing a model to predict spending ratio. There are no missing values in the data set, but there are some incorrect entries that must be identified and removed before completing the analysis. Income is coded as an ordinal value, ranging from 1 to 12. Age can be regarded as quantitative, and any value less than 18 is invalid. Length of residence (*LOR*) is a value ranging from zero to someone's age. *LOR* should not be higher than age. All of the catalog variables are represented by indicator variables; either the consumer bought and the variable is coded as 1 or the consumer didn't buy and the variable is coded as 0. The other variables can be viewed as indexes for measuring assets, liquidity, and spending. Find a multiple regression model for predicting the amount of money that consumers will spend on catalog shopping, as measured by spending ratio. Your goal is to identify the best model you can. Remember to check the conditions for inference as you evaluate your models. 

TABLE 28.16 Catalog-spending data for 200 individuals from a very large data base

SPENDING RATIO	AGE	LENGTH OF RESIDENCE		INCOME	TOTAL ASSETS	SECURITY ASSETS	SHORT-TERM LIQUIDITY	LONG-TERM LIQUIDITY	WEALTH INDEX	SPENDING VOLUME	SPENDING VELOCITY
		HOME	ADS								
11.83	0	2	3	122	27	225	422	286	503	285	
16.83	35	3	5	195	36	220	420	430	690	570	
11.38	46	9	5	123	24	200	420	290	600	280	
31.33	41	2	2	117	25	222	419	279	543	308	
1.90	46	7	9	493	105	310	500	520	680	100	
84.13	46	15	5	138	27	340	450	440	440	50	
2.15	46	16	4	162	25	230	430	360	690	180	
38.00	56	31	6	117	27	300	440	400	500	10	
136.28	48	8	5	119	23	250	430	360	610	0	
COLLECTIBLE GIFTS	BRICK/ MORTAR	MARTHA'S HOME	SUNDAY ADS	THEME COLLECTIONS	CUSTOM DECORATING	RETAIL KIDS	TEEN WEAR	CAR LOVERS	COUNTRY COLLECTIONS		
1	0	0	1	0	1	1	1	0	0	1	
0	1	1	0	0	1	1	1	0	0	0	
1	0	0	1	1	1	1	1	0	0	1	
1	0	0	1	1	0	0	0	0	0	1	
0	1	1	0	0	1	0	0	0	0	0	
0	1	1	0	0	1	0	0	0	1	0	
1	0	0	1	0	0	0	0	0	0	1	
1	1	1	0	1	1	1	1	1	1	0	
1	0	1	0	1	1	0	0	0	0	1	



EXPLORING THE WEB

28.55 Are gas prices driving elections? The Chance Web site discusses the use of regression to predict the margin of victory in presidential elections since 1948 from the price of gas (in 2008 dollars). Read the article at www.causeweb.org/wiki/chance/index.php/Chance_News_72. Use the data in the article to do the following.

- Fit a simple linear regression model using gas price to predict margin of victory. Do your results agree with those reported in the article?
- Use the incumbent party as an indicator variable (code Democrats as 1 and Republicans as 0), and add this to your simple linear regression model. What is the value of R^2 ?
- Now add gross domestic product (GDP) to your regression model in (b). What is the value of R^2 ?

28.56 Historical tuition and fees. You can find data on past tuition and fees at several colleges by doing a Google search on “historical tuition and fees.” Select one of the colleges you find and determine whether the data show the same pattern as you observed in Exercise 28.24 (page 28-51). You should try to find data going back at least 20 years (at the time we searched, we were able to find data for Clemson University, Montana State University, Western Washington University, Oregon State University, and University of Tennessee). The data may not be in spreadsheet format, and you may have to enter or cut and paste it into a spreadsheet to carry out your analysis.

NOTES AND DATA SOURCES

- We thank Tom Shields for the data from Franklin County Municipal Court.
- Data were estimated from a scatterplot in Philipp Heeb, Mathias Kolliker, and Heinz Richner, “Bird-ectoparasite interactions, nest humidity, and ectoparasite community structure,” *Ecology*, 81 (2000), pp. 958–968.
- For more details, see H. Hoppeler and E. Weibel, “Scaling functions to body size: theories and facts,” *Journal of Experimental Biology*, 208 (2005), pp. 1573–1574.
- We thank Professor Haruhiko Itagaki and his students Andrew Vreede and Marissa Stearns for providing data on tobacco hornworm caterpillars (*Manduca sexta*).
- For more details, see Michael H. Kutner, Christopher J. Nachtsheim, and John Neter, *Applied Linear Regression Models*, 4th ed., McGraw-Hill, 2004.
- Diamond data base downloaded from AwesomeGems.com on July 28, 2005.
- We thank Terry Klopcek for providing data from a physics lab on radioactive decay.
- We thank David Cameron for providing data from a clothing retailer.
- Found online at www.web.virginia.edu/iaas/data_catalog/institutional/historical/tuition/fees.htm.
- The data in Table 28.8 are part of a larger data set in the *Journal of Statistics Education* archive, accessible via the Internet. The original source is Pekka Brofeldt, “Bidrag till kaennedom on fiskbestonet i vaara sjoeare. Laengelmaevesi,” in T. H. Jaervi, *Finlands fiskeriet*, vol. 4, *Meddelanden utgivna av fiskerifoereringen i Finland*, Helsinki, 1917. The data were put in the archive (with information in English) by Juha Puranen of the University of Helsinki.
- P. Velleman, *ActivStats 2.0*, Addison Wesley Interactive, 1997.
- These data were provided by Professor Shelly MacDermid, Department of Child Development and Family Studies, Purdue University, from a study reported in S. M.

- MacDermid et al., "Is small beautiful? Work-family tension, work conditions, and organizational size," *Family Relations*, 44 (1994), pp. 159–167.
13. Data from the U.S. Historical Climatology Network, archived at www.co2science.org/. (Despite claims made on this site, temperatures at most U.S. locations show a gradual increase over the past century.)
14. I thank Ray and Pat Heithaus for providing data on the pine seedlings at the Brown Family Environmental Center.
15. Data provided by Robert Dale, Purdue University.
16. Darlene Gordon, "The relationships among academic self-concept, academic achievement, and persistence with academic self-attribution, study habits, and perceived school environment," PhD thesis, Purdue University, 1997.
17. James T. Fleming, "The measurement of children's perception of difficulty in reading materials," *Research in the Teaching of English*, 1 (1967), pp. 136–156.
18. I thank David Cameron for providing the random sample of 200 observations from a large catalog-spending data base.



More About Analysis of Variance

Chapter 29

Analysis of variance (ANOVA) is a statistical method for comparing the means of several populations based on independent random samples or on the mean responses to several treatments in a randomized comparative experiment. When we compare just two means, we use the two-sample t procedures described in Chapter 19. ANOVA allows comparison of any number of means. The basic form of ANOVA is one-way ANOVA, which treats the means being compared as mean responses to different values of a single variable. For example, in Chapter 25 we used one-way ANOVA to compare the mean lengths of three varieties of tropical flowers and the mean number of species per plot after each of three logging conditions.

BEYOND ONE-WAY ANOVA

You should recall or review the big ideas of one-way ANOVA from Chapter 25. One-way ANOVA compares the means $\mu_1, \mu_2, \dots, \mu_I$ of I populations based on samples of sizes n_1, n_2, \dots, n_I from these populations.

- The conditions for ANOVA require *independent random samples* from each of the I populations (or a randomized comparative experiment with I treatments); *Normal distributions* for the response variable in each population; and a *common standard deviation σ* in all populations. Fortunately, ANOVA inference is quite robust against moderate violations of the Normality and common standard deviation conditions.
- Using many separate two-sample t procedures to compare many pairs of means is a bad idea because we don't get a P -value or a confidence level for the complete set of comparisons together. This is the problem of **multiple comparisons**.

IN THIS CHAPTER WE COVER...

- Beyond one-way ANOVA
- Follow-up analysis: Tukey pairwise multiple comparisons
- Follow-up analysis: contrasts*
- Two-way ANOVA: conditions, main effects, and interaction
- Inference for two-way ANOVA
- Some details of two-way ANOVA*

conditions for ANOVA

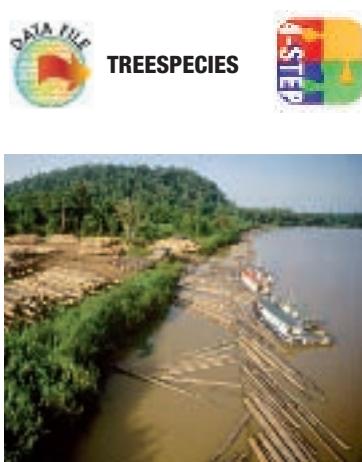
multiple comparisons

- One-way ANOVA gives a single test for the null hypothesis that all the population means are the same against the alternative hypothesis that not all are the same.
- ANOVA works by comparing how far apart the sample means are relative to the variation among individual observations in the same sample. The test statistic is the **ANOVA F statistic**

ANOVA F statistic***F distribution***

- In basic statistical practice, we combine the *F* test with data analysis to check the conditions for ANOVA and to see which means appear to differ and by how much.

Here is an example that illustrates one-way ANOVA.



Frans Lanting/Minden Pictures

EXAMPLE 29.1 Logging in the rain forest

STATE: How does logging in a tropical rain forest affect the forest in later years? Researchers compared forest plots in Borneo that had never been logged (Group A) with similar plots nearby that had been logged 1 year earlier (Group B) and 8 years earlier (Group C). Although the study was not an experiment, the authors explain why we can consider the plots to be randomly selected. Table 29.1 displays data on the number of tree species found in each plot.¹ Is there evidence that the mean numbers of species in the three groups differ?

PLAN: Describe how the three samples appear to differ, check the conditions for ANOVA, and carry out the one-way ANOVA *F* test to compare the three population means.

TABLE 29.1 Counts of tree species in forest plots in Borneo

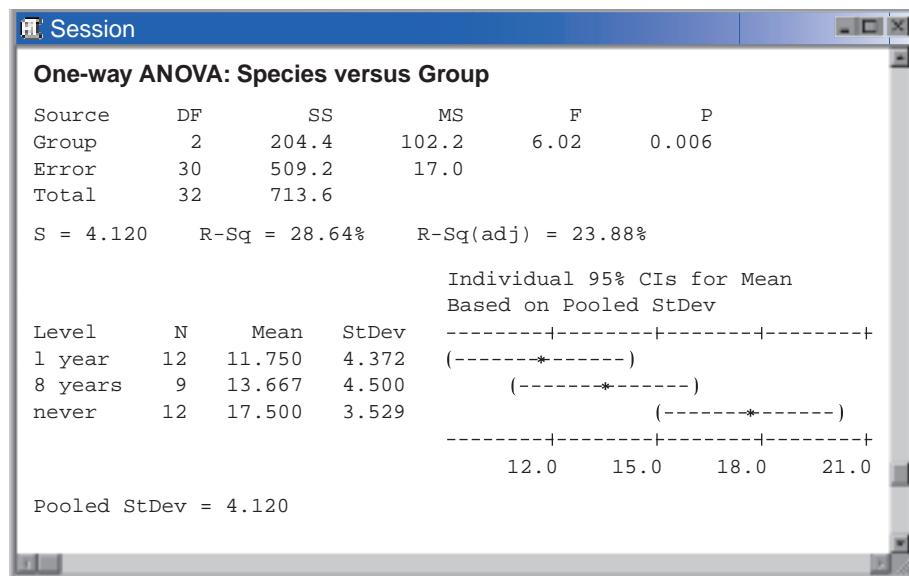
GROUP A NEVER LOGGED	GROUP B LOGGED 1 YEAR AGO	GROUP C LOGGED 8 YEARS AGO
22	11	17
18	11	4
22	14	18
20	7	14
15	18	18
21	15	15
13	15	15
13	12	10
19	13	12
13	2	
19	15	
15	8	

Group A (never)	Group B (1 year)	Group C (8 years)
2	2 0	2
3	3	3
4	4	4 0
5	5	5
6	6	6
7	7 0	7
8	8 0	8
9	9	9
10	10	10 0
11	11 00	11
12	12 0	12 0
13	0 00	13 0
14	14 0	14 0
15	0 0	15 0 0
16	16	16
17	17	17 0
18	0 18 0	18 0 0
19	0 0	19
20	0 20	20
21	0 21	21
22	0 0 22	22

FIGURE 29.1

Side-by-side stemplots comparing the counts of tree species in three groups of forest plots, from Table 29.1.

SOLVE: The stemplots in Figure 29.1 show that the distributions are irregular in shape, as is common for small samples. None of the distributions are strongly skewed, and there are no extreme outliers. The unlogged Group A plots tend to have more species than the two logged groups and also show less variation among plots. The Minitab ANOVA output in Figure 29.2 shows that the standard deviations satisfy our rule of thumb (text page 634) that the largest (4.5) is no more than twice the smallest (3.529). The output

Minitab**FIGURE 29.2**

Minitab ANOVA output for the logging study of Example 29.1.

also shows that the mean number of species in Group A is higher than in Groups B and C. The ANOVA F test is highly significant, with P -value $P = 0.006$.

CONCLUDE: The data provide strong evidence ($F = 6.02$, $P = 0.006$) that the mean number of species per plot is not the same in all three groups. It appears that the mean species count is higher in unlogged plots. ■

This chapter moves beyond basic one-way ANOVA in two directions.

Follow-up analysis. The ANOVA F test in Example 29.1 tells us only that the three population means are not the same. We would like to say which means differ and by how much. For example, do the data allow us to say that the “never logged” population does have a higher mean species count than the “logged 1 year ago” and the “logged 8 years ago” populations of forest plots? This is a *follow-up analysis* to the F test that goes beyond data analysis to confidence intervals and tests of significance for specific comparisons of means.

Two-way ANOVA. Example 29.1, and one-way ANOVA in general, compare mean responses for several values of just one explanatory variable. In Example 29.1, that variable is “how long ago this plot was logged.” Suppose that we have data on *two* explanatory variables: say, how long ago the plot was logged and whether it is located in a river bottom or in the highlands. There are now 6 groups formed by combinations of time since logging and location, as follows:

		Plot Location	
		River bottom	Highlands
Time since logging	Never logged	Group 1	Group 2
	1 year	Group 3	Group 4
	8 years	Group 5	Group 6

One-way ANOVA will still tell us if there is evidence that mean species counts in these 6 types of plot differ. But we want more: Does location matter? Does time since logging matter? And do these two variables *interact*? That is, does the effect of logging change when we move from river bottoms to highlands? Perhaps trees regrow faster in river bottoms, so that logging has less effect there than in the highlands. To answer these questions we must extend ANOVA to take into account the fact that the 6 groups are formed from two explanatory variables. This is *two-way ANOVA*.

We will first discuss follow-up analysis in one-way ANOVA. Fortunately, the distinction between one-way and two-way doesn’t affect the follow-up methods we will present. So once we have mastered these methods in the one-way setting, we can apply them immediately to two-way problems.



APPLY YOUR KNOWLEDGE

These exercises review one-way ANOVA. Follow the **Plan**, **Solve**, and **Conclude** steps of the four-step process, as illustrated in Example 29.1. In the next section, you will do follow-up analysis in all these settings.

29.1 Good weather and tipping. Exercise 25.35 (text page 650) gives the data for a study on the effect of a waitress's weather prediction on the tipping percent. Carry out data analysis and ANOVA to determine whether there are differences among the mean tipping percents for the three experimental conditions.  TIPPING

29.2 Which color attracts beetles best? To detect the presence of harmful insects in farm fields, we can put up boards covered with a sticky material and examine the insects trapped on the boards. Which colors attract insects best? Experimenters placed six boards of each of four colors at random locations in a field of oats and measured the number of cereal leaf beetles trapped. Here are the data:²  BEETLES

Board color	Beetles trapped					
Blue	16	11	20	21	14	7
Green	37	32	20	29	37	32
White	21	12	14	17	13	20
Yellow	45	59	48	46	38	47

Is there evidence that the colors differ in their ability to attract beetles?

29.3 Dogs, friends, and stress. If you are a dog lover, perhaps having your dog along reduces the effect of stress. To examine the effect of pets in stressful situations, researchers recruited 45 women who said they were dog lovers. The EESEE story "Stress among Pets and Friends" describes the results. Fifteen of the subjects were randomly assigned to each of three groups to do a stressful task alone (the control group), with a good friend present, or with their dog present. The subject's mean heart rate during the task is one measure of the effect of stress. Table 29.2 displays

TABLE 29.2 Heart rates (beats per minute) after performing a task under three conditions: P = with pet, F = with friend, C = control group

GROUP	HEART RATE	GROUP	HEART RATE	GROUP	HEART RATE
P	69.169	P	75.985	C	90.015
P	68.862	F	91.354	F	101.062
C	84.738	C	73.277	F	76.908
F	99.692	F	83.400	C	99.046
C	87.231	F	100.877	F	97.046
C	84.877	C	84.523	P	69.538
P	70.169	F	102.154	C	75.477
P	64.169	C	77.800	C	62.646
P	58.692	C	70.877	P	70.077
C	80.369	P	86.446	F	88.015
C	91.754	P	97.538	F	81.600
P	79.662	F	89.815	F	86.985
C	87.446	F	80.277	F	92.492
C	87.785	P	85.000	P	72.262
P	69.231	F	98.200	P	65.446

the data. Are there significant differences among the mean heart rates under the three conditions?  STRESSPETS

FOLLOW-UP ANALYSIS: TUKEY PAIRWISE MULTIPLE COMPARISONS

In Example 29.1 we saw that there is good evidence that the mean number of tree species is not the same for plots that have never been logged, plots logged 1 year ago, and plots logged 8 years ago. The sample means in Figure 29.2 suggest that (as we might expect) the mean species count is highest in never-logged plots and lowest in plots logged just a year ago.

EXAMPLE 29.2 Comparing groups: individual *t* procedures

How much higher is the mean count of tree species in plots that have never been logged than in plots logged a year ago? A 95% confidence interval comparing Groups A and B answers this question. Because the conditions for ANOVA require the population standard deviation to be the same in all three populations of plots, we will use a version of the two-sample *t* confidence interval that also assumes equal standard deviations.

The standard error for the difference of sample means $\bar{x}_A - \bar{x}_B$ estimates the standard deviation of the difference, which is

$$\sqrt{\frac{\sigma^2}{n_A} + \frac{\sigma^2}{n_B}} = \sigma \sqrt{\frac{1}{n_A} + \frac{1}{n_B}}$$

because both populations have the same standard deviation σ . The pooled standard deviation s_p is an estimate of σ based on all three samples (see text page 642). So the standard error of $\bar{x}_A - \bar{x}_B$ is

$$s_p \sqrt{\frac{1}{n_A} + \frac{1}{n_B}}$$

The Minitab output in Figure 29.2 gives $s_p = 4.120$. This estimate has 30 degrees of freedom, the degrees of freedom for “Error” in the ANOVA. A 95% confidence interval for $\mu_A - \mu_B$ uses the critical value of the *t* distribution with 30 degrees of freedom:

$$\begin{aligned} (\bar{x}_A - \bar{x}_B) &\pm t^* s_p \sqrt{\frac{1}{n_A} + \frac{1}{n_B}} \\ &= (17.50 - 11.75) \pm (2.042)(4.120) \sqrt{\frac{1}{12} + \frac{1}{12}} \\ &= 5.75 \pm 3.43 \\ &= 2.32 \text{ to } 9.18 \end{aligned}$$

We are 95% confident that there are on the average between 2.32 and 9.18 more tree species in plots that have never been logged. Because this confidence interval does not contain 0, we can reject the null hypothesis of no difference, $H_0: \mu_A = \mu_B$, in favor of the two-sided alternative at the 5% significance level. ■

Example 29.2 gives a single 95% confidence interval. We would like to estimate all three **pairwise differences** among the population means,

$$\mu_A - \mu_B \quad \mu_A - \mu_C \quad \mu_B - \mu_C$$

Three 95% confidence intervals will not give us 95% confidence that all three simultaneously capture their true parameter values. This is the problem of multiple comparisons, discussed on text page 625.



In general, we want to give confidence intervals for all pairwise differences among the population means $\mu_1, \mu_2, \dots, \mu_I$ of I populations. We want **overall confidence level** (say) 95%. That is, in very many uses of the method, *all* the intervals will simultaneously capture the true differences 95% of the time. To do this, take the number of comparisons into account by replacing the t critical value t^* in Example 29.2 by another critical value based on the distribution of the difference between the largest and smallest of a set of I sample means. We will call this critical value m^* , for multiple comparisons. Values of m^* depend on the number of populations we are comparing and on the total number of observations in the samples, as well as on the confidence level we want. Tables are therefore long and messy, so in practice we rely on software. This method is named after its inventor, John Tukey (1915–2000), who developed the ideas of modern data analysis.

pairwise difference

overall confidence



Alfred Eisenstaedt/Time Life Pictures/Getty Images
John Tukey

TUKEY PAIRWISE MULTIPLE COMPARISONS

In the ANOVA setting, we have independent SRSs of size n_i from each of I populations having Normal distributions with means μ_i and a common standard deviation σ . **Tukey simultaneous confidence intervals** for all pairwise differences $\mu_i - \mu_j$ among the population means have the form

$$(\bar{x}_i - \bar{x}_j) \pm m^* s_p \sqrt{\frac{1}{n_i} + \frac{1}{n_j}}$$

Here \bar{x}_i is the sample mean of the i th sample and s_p is the pooled estimate of σ . The critical value m^* depends on the confidence level C , the number of populations I , and the total number of observations.

To carry out **simultaneous tests** of the hypotheses

$$\begin{aligned} H_0: \mu_i &= \mu_j \\ H_a: \mu_i &\neq \mu_j \end{aligned}$$

for all pairs of population means at fixed significance level $\alpha = 1 - C$, reject H_0 for any pair whose confidence interval does not contain 0.

If all samples are the same size, Tukey simultaneous confidence intervals provide **overall confidence level C** . That is, C is the probability that *all* of the intervals simultaneously capture the true pairwise differences. If the samples differ in size, the true confidence level is at least as large as C , so that conclusions are conservative. Similarly, if all the samples are the same size, the Tukey simultaneous tests have **overall significance level $1 - C$** . That is, $1 - C$ is the probability that when

overall confidence level

overall significance level

all the population means are equal, *any* of the tests incorrectly rejects its null hypothesis. If the samples differ in size, the true significance level is smaller than $1 - C$, so that conclusions are conservative.

EXAMPLE 29.3 Logging in the rain forest: multiple intervals

Figure 29.3 contains more Minitab output for the ANOVA comparing the mean counts of tree species in three groups of forest plots in Borneo. We asked for Tukey multiple comparisons with an overall error rate of 5%. That is, the overall confidence level for the three intervals together is 95%.

The format of the Minitab output takes some study. Be sure you can see that the Tukey confidence intervals are

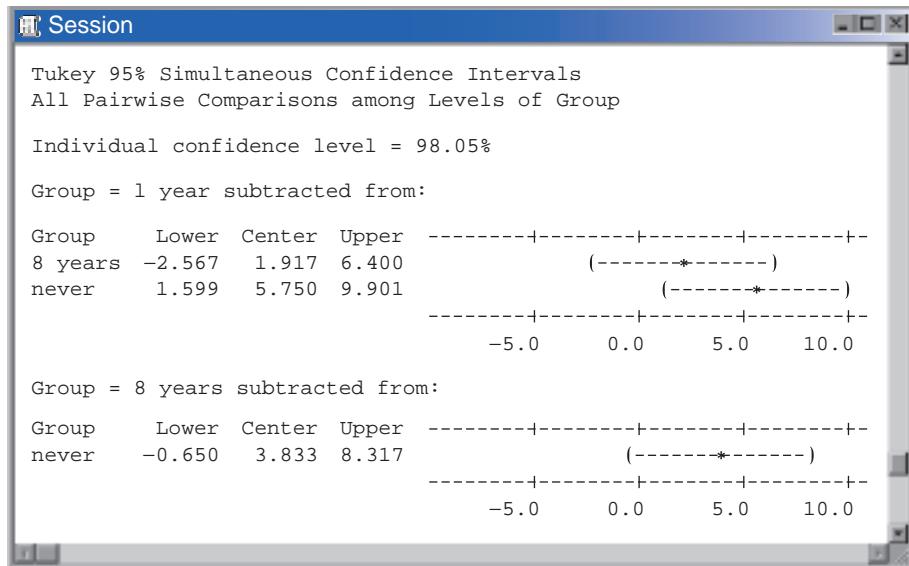
$$\begin{aligned} 1.599 \text{ to } 9.901 & \text{ for } \mu_A - \mu_B \\ -0.650 \text{ to } 8.317 & \text{ for } \mu_A - \mu_C \\ -2.567 \text{ to } 6.400 & \text{ for } \mu_C - \mu_B \end{aligned}$$

The interval for $\mu_A - \mu_B$ is wider than the individual 95% confidence interval in Example 29.2. The wider interval is the price we pay for having 95% confidence not just in one interval but in all three simultaneously. ■

FIGURE 29.3

Additional Minitab ANOVA output for the logging study of Example 29.3, showing Tukey simultaneous confidence intervals.

Minitab



```

Session

Tukey 95% Simultaneous Confidence Intervals
All Pairwise Comparisons among Levels of Group

Individual confidence level = 98.05%

Group = 1 year subtracted from:

Group      Lower   Center   Upper   -----
8 years    -2.567   1.917   6.400   (-----*-----)
never      1.599   5.750   9.901   (-----*-----)
                                         -----
                                         -5.0     0.0     5.0     10.0

Group = 8 years subtracted from:

Group      Lower   Center   Upper   -----
never     -0.650   3.833   8.317   (-----*-----)
                                         -----
                                         -5.0     0.0     5.0     10.0

```

EXAMPLE 29.4 Logging in the rain forest: multiple tests

The ANOVA null hypothesis is that all population means are equal,

$$H_0: \mu_A = \mu_B = \mu_C$$

We know from the output in Figure 29.2 that the ANOVA F test rejects this hypothesis ($F = 6.02$, $P = 0.006$). So we have good evidence that *some* pairs of means are not

the same. Which pairs? Look at the simultaneous 95% confidence intervals in Example 29.3. Which of these intervals do not contain 0? If an interval does not contain 0, we reject the hypothesis that this pair of population means are equal.

The conclusions are

We can reject	$H_0: \mu_A = \mu_B$
We cannot reject	$H_0: \mu_A = \mu_C$
We cannot reject	$H_0: \mu_B = \mu_C$

This Tukey simultaneous test of three null hypotheses has the property that when all three hypotheses are true, there is only 5% probability that *any* of the three tests wrongly rejects its hypothesis. ■

Think for a moment about the conclusion of Example 29.4. At first glance, it appears to say that μ_A equals μ_C and that μ_B equals μ_C , but that μ_A and μ_B are not equal. That sounds like nonsense. Now is the time to recall what a test at a fixed significance level such as 5% tells us: either we do have enough evidence to reject the null hypothesis, or the data do not give enough evidence to allow rejection. There is no contradiction in saying that

- We do have enough evidence to conclude that $\mu_A \neq \mu_B$
- We do not have enough evidence to conclude that $\mu_A \neq \mu_C$
- We do not have enough evidence to conclude that $\mu_B \neq \mu_C$

That is, $\bar{x}_A = 17.500$ and $\bar{x}_B = 11.750$ are far enough apart to conclude that the population means differ, but neither 17.500 nor 11.750 is far enough from $\bar{x}_C = 13.667$. Notice that the Tukey method does not give a *P*-value for the three tests taken together. Rather, we have a set of “reject” or “fail to reject” conclusions with an overall significance level that we fixed in advance, 5% in this example.

There are many other multiple-comparisons procedures that produce various simultaneous confidence intervals with an overall confidence level or simultaneous tests with an overall probability of any false rejection. The Tukey procedures are probably the most useful. If you can interpret results from Tukey, you can understand output from any multiple-comparisons procedure.

APPLY YOUR KNOWLEDGE

29.4 Good weather and tipping. In Exercise 29.1, you carried out basic ANOVA to compare the mean tipping percents for three experimental conditions. 

- Find the Tukey simultaneous 95% confidence intervals for all pairwise differences among the three population means.
- Explain in simple language what “95% confidence” means for these intervals.
- Which pairs of means differ significantly at the overall 5% significance level?

29.5 Which color attracts beetles best? Exercise 29.2 presents data on the numbers of cereal leaf beetles trapped by boards of four different colors. Yellow boards appear most effective. ANOVA gives very strong evidence that the colors differ in their ability to attract beetles. 

- (a) How many pairwise comparisons are there when we compare four colors?
- (b) Which pairs of colors are significantly different when we require significance level 5% for all comparisons as a group? In particular, is yellow significantly better than every other color?

29.6 Dogs, friends, and stress. In Exercise 29.3, the ANOVA F test had a very small P-value, giving good reason to conclude that mean heart rates under stress do differ depending on whether a pet, a friend, or no one is present. Do the means for the two treatments (pet, friend) differ significantly from each other and from the mean for the control group?  STRESSPETS

- (a) What are the three null hypotheses that formulate these questions?
 - (b) We want to be 90% confident that we don't wrongly reject any of the three null hypotheses. Tukey pairwise comparisons can give conclusions that meet this condition. What are the conclusions?
-

FOLLOW-UP ANALYSIS: CONTRASTS*

Multiple-comparisons methods give conclusions about *all* comparisons in some class with a measure of confidence that applies to all the comparisons taken together. For example, Tukey's method gives conclusions about all pairwise differences among a set of population means. These methods are most useful when we did not have any specific comparison in mind before we produced the data.

Multiple comparisons procedures sometimes give tests or confidence intervals for comparisons that don't interest us. And they may leave out comparisons that do interest us. If we have specific questions in mind before we produce data, it is more efficient to plan an analysis that asks these specific questions. Do note that *it is not legitimate to look at the data and then formulate a question suggested by the data*. If data suggest a specific effect, you need new data to give evidence that the effect isn't just chance variation at work.

EXAMPLE 29.5 Which color attracts beetles best?

What color should we use on sticky boards placed in a field of oats to attract cereal leaf beetles? Exercise 29.2 gives data from an experiment in which 24 boards, 6 of each color blue, green, white, and yellow, were placed at random locations in a field. ANOVA (Exercise 29.2) shows that there are significant differences among the mean numbers of beetles trapped by these colors. We might follow ANOVA with Tukey pairwise comparisons (Exercise 29.5).

But in fact we have specific questions in mind: we suspect that warm colors are generally more attractive than cold colors. That is, *before we produce any data*, we suspect that blue and white boards have similar properties, that green and yellow boards are

*This material is optional.

similar, and also that the average beetle count for green and yellow is greater than the average count for blue and white. We therefore want to test three hypotheses:

$$\begin{array}{lll} \text{Hypothesis 1} & \text{Hypothesis 2} & \text{Hypothesis 3} \\ H_0: \mu_B = \mu_W & H_0: \mu_G = \mu_Y & H_0: (\mu_G + \mu_Y)/2 = (\mu_B + \mu_W)/2 \\ H_a: \mu_B \neq \mu_W & H_a: \mu_G \neq \mu_Y & H_a: (\mu_G + \mu_Y)/2 > (\mu_B + \mu_W)/2 \end{array}$$

Two of these hypotheses involve pairwise comparisons. The third does not, and also has a one-sided alternative. ■

We can ask questions about population means by specifying *contrasts* among the means.

CONTRASTS

In the ANOVA setting comparing the means $\mu_1, \mu_2, \dots, \mu_I$ of I populations, a **population contrast** is a combination of the means

$$L = c_1\mu_1 + c_2\mu_2 + \cdots + c_I\mu_I$$

with numerical coefficients that add to 0, $c_1 + c_2 + \cdots + c_I = 0$.

EXAMPLE 29.6 Attracting beetles: contrasts

We can restate the three hypotheses in Example 29.5 in terms of three contrasts:

$$\begin{aligned} L_1 &= \mu_B - \mu_W \\ &= (1)(\mu_B) + (0)(\mu_G) + (-1)(\mu_W) + (0)(\mu_Y) \\ L_2 &= \mu_G - \mu_Y \\ &= (0)(\mu_B) + (1)(\mu_G) + (0)(\mu_W) + (-1)(\mu_Y) \\ L_3 &= (\mu_G + \mu_Y)/2 - (\mu_B + \mu_W)/2 \\ &= (-1/2)(\mu_B) + (1/2)(\mu_G) + (-1/2)(\mu_W) + (1/2)(\mu_Y) \end{aligned}$$



BEETLES

Check that the four coefficients in each line do add to 0. In terms of these contrasts the hypotheses become

$$\begin{array}{lll} \text{Hypothesis 1} & \text{Hypothesis 2} & \text{Hypothesis 3} \\ H_0: L_1 = 0 & H_0: L_2 = 0 & H_0: L_3 = 0 \\ H_a: L_1 \neq 0 & H_a: L_2 \neq 0 & H_a: L_3 > 0 \blacksquare \end{array}$$

Some statistical software will test hypotheses and give confidence intervals for any contrasts you specify. Other software lacks this capability, but it is easy to work with just information from basic ANOVA output. Here's how.

To estimate a population contrast

$$L = c_1\bar{x}_1 + c_2\bar{x}_2 + \cdots + c_I\bar{x}_I$$

use the corresponding **sample contrast**

sample contrast

$$\hat{L} = c_1\bar{x}_1 + c_2\bar{x}_2 + \cdots + c_I\bar{x}_I$$

The sample contrast \hat{L} has standard error (estimated standard deviation)

$$\text{SE}_{\hat{L}} = s_p \sqrt{\frac{c_1^2}{n_1} + \frac{c_2^2}{n_2} + \cdots + \frac{c_I^2}{n_I}}$$

INFERENCE ABOUT A POPULATION CONTRAST

In the ANOVA setting, a level C confidence interval for a population contrast L is

$$\hat{L} \pm t^* \text{SE}_{\hat{L}}$$

where \hat{L} is the corresponding sample contrast and t^* is a critical value from the t distribution with the degrees of freedom for error in the ANOVA.

To test the hypothesis $H_0: L = 0$, use the t statistic

$$t = \frac{\hat{L}}{\text{SE}_{\hat{L}}}$$

with the same degrees of freedom.

For one-way ANOVA, the degrees of freedom for error are $N - I$, where N is the total number of observations and I is the number of populations compared (see text page 639). The box states the result more generally so that it applies to two-way ANOVA as well as one-way. If the contrast is a pairwise difference between means, the contrast confidence interval is exactly the individual confidence interval illustrated in Example 29.2.

EXAMPLE 29.7 Attracting beetles: inference for contrasts

Figure 29.4 displays the Minitab ANOVA output for the study on attracting cereal leaf beetles. The pooled estimate of σ is $s_p = 5.672$, and the number of degrees of freedom for error is 20. Minitab does not offer contrasts, so we must use a calculator.

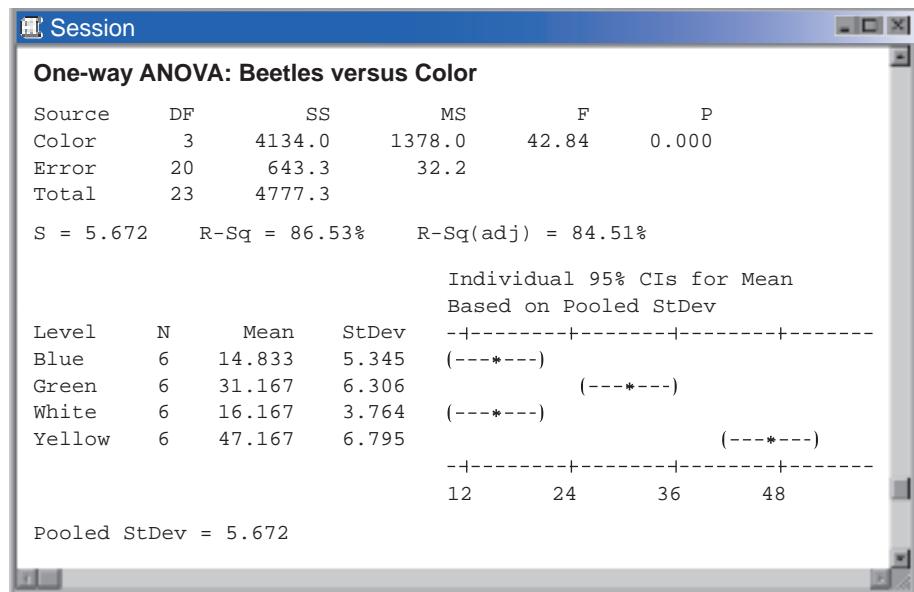
The three sample contrasts and their standard errors are

$$\begin{aligned}\hat{L}_1 &= -1.334 & \text{SE}_1 &= 3.2747 \\ \hat{L}_2 &= -16.000 & \text{SE}_2 &= 3.2747 \\ \hat{L}_3 &= 23.667 & \text{SE}_3 &= 2.3153\end{aligned}$$

Here are the details for the third contrast:

$$\begin{aligned}L_3 &= (-1/2)(\mu_B) + (1/2)(\mu_G) + (-1/2)(\mu_W) + (1/2)(\mu_Y) \\ \hat{L}_3 &= (-1/2)(14.833) + (1/2)(31.167) + (-1/2)(16.167) + (1/2)(47.167) \\ &= 23.667 \\ \text{SE}_3 &= (5.672) \sqrt{\frac{(-1/2)^2}{6} + \frac{(1/2)^2}{6} + \frac{(-1/2)^2}{6} + \frac{(1/2)^2}{6}} \\ &= (5.672)(0.4082) \\ &= 2.3153\end{aligned}$$

Minitab

**FIGURE 29.4**

Minitab ANOVA output for the study on attracting cereal leaf beetles, for Example 29.7.

If you use software, your answers may differ slightly due to roundoff error in the hand calculations. A 95% confidence interval for L_3 uses $t^* = 2.086$ from Table C with $df = 20$,

$$\begin{aligned}\hat{L}_3 \pm t^*SE_3 &= 23.667 \pm (2.086)(2.3153) \\ &= 23.667 \pm 4.830 \\ &= 18.837 \text{ to } 28.497\end{aligned}$$

We are 95% confident that the average number of beetles attracted by green and yellow boards exceeds the average for blue and white boards by between about 18.8 and 28.5 beetles per board.

There is very strong evidence that the population contrast L_3 is greater than 0. The t statistic is

$$t = \frac{\hat{L}_3}{SE_3} = \frac{23.667}{2.3153} = 10.22$$

with 20 degrees of freedom, and $P < 0.0005$ from Table C. The other t tests conclude that μ_B and μ_W do not differ significantly but that there is a significant difference between μ_G and μ_Y . ■

Our confidence intervals and tests for contrasts are individual procedures for each contrast. If we do inference about three contrasts, such as those in Examples 29.6 and 29.7, we face the problem of multiple comparisons again. That is, we do not have an overall confidence level for all three intervals together. There are more advanced multiple-comparisons methods that apply to contrasts just as Tukey's method applies to pairwise differences.


APPLY YOUR KNOWLEDGE

- 29.7 Green versus yellow.** Using the Minitab output in Figure 29.4, verify the values for the sample contrast \hat{L}_2 and its standard error given in Example 29.7. Give a 95% confidence interval for the population contrast L_2 . Carry out a test of the hypothesis $H_0: L_2 = 0$ against the two-sided alternative. Be sure to state your conclusions in the setting of the study.
- 29.8 Logging in the rain forest: contrasts.** Figure 29.2 (page 29-5) gives basic ANOVA output for the study of the effects of logging described in Example 29.1. We might describe the overall effect of logging by comparing the mean species count for unlogged plots (Group A) with the average of the mean counts for the two groups of logged plots (Groups B and C).
- What population contrast L expresses this comparison?
 - Starting from the output in Figure 29.2, give the sample contrast that estimates L and its standard error.
 - Is there good evidence that the mean species count in unlogged plots is higher than the average for the two groups of logged plots? State hypotheses in terms of the population contrast L and carry out a test.
 - How much higher is the mean count in unlogged plots than the average for the two groups of unlogged plots? Give a 95% confidence interval.

TWO-WAY ANOVA: CONDITIONS, MAIN EFFECTS, AND INTERACTION

One-way analysis of variance compares the mean responses from any set of populations or experimental treatments when the responses satisfy the ANOVA conditions. Often, however, a sample or experiment has some design structure that leads to more specific questions than those answered by the one-way ANOVA F test or by Tukey pairwise comparisons. It is common, for example, to compare treatments that are combinations of values of two explanatory variables, two **factors** in the language of experimental design. Here is an example that we already met in Chapter 9.



Andrew Harrer/Bloomberg via Getty Images


EXAMPLE 29.8 Effects of TV advertising

What are the effects of repeated exposure to an advertising message? The answer may depend both on the length of the ad and on how often it is repeated. To investigate this question, assign undergraduate students (the *subjects*) to view a 40-minute television program that includes ads for a digital camera. Some subjects see a 30-second commercial; others, a 90-second version. The same commercial appears either 1, 3, or 5 times during the program.

This experiment has two *factors*: length of the commercial, with 2 values; and repetitions, with 3 values. The 6 combinations of one value of each factor form 6 *treatments*. The subjects are divided at random into 6 groups, one for each treatment:

		Variable C Repetitions		
		1 time	3 times	5 times
Variable R	30 seconds	Group 1	Group 2	Group 3
	90 seconds	Group 4	Group 5	Group 6

This is a *two-way layout*, with values of one factor forming rows and values of the other forming columns. After watching the TV program, all the subjects fill out a questionnaire that produces an “intention to buy” score with values between 0 and 100. This is the *response variable*. ■

To analyze data from such a study, we impose some additional conditions. Here are the **conditions for two-way ANOVA** that will govern our work on this topic:

1. We have *responses for all combinations* of values of the two factors (all 6 cells in Example 29.8). No combinations are missing in our data. In general, call the two explanatory variables R and C (for Row and Column). Variable R has r different values and variable C has c different values. The study compares all rc combinations of these values. Such designs are called **crossed**.

[two-way ANOVA conditions](#)

[crossed design](#)

2. In an observational study, we have *independent SRSs* from each of the rc populations. If the study is an experiment, the available subjects are allocated at random among all rc treatments. That is, we have a *completely randomized design*. We also allow some *randomized block designs* in which we have an SRS from each of r populations and the subjects in each SRS are then separately allocated at random among the same c treatments.³ We met block designs in Chapter 9 (see text page 237).

3. The response variable has a *Normal distribution* in each population. The population mean responses may differ, but all rc populations have a *common standard deviation* σ .

4. We have the *same number of individuals* n in each of the rc treatment groups or samples. Such designs are called **balanced**.

[balanced design](#)

As always, the design of the study is the most important condition for statistical inference. Conditions 1 and 2 describe the designs covered by two-way ANOVA. These designs are very common. Condition 3 lists the usual ANOVA conditions on the distribution of the response variable. ANOVA inference is reasonably robust against violations of these conditions.

When you design a study, you should try to satisfy the fourth condition; that is, you should choose equal numbers of individuals for each treatment. Balanced designs have several advantages in any ANOVA: F tests are most robust against violation of the “common standard deviation” condition when the subject counts are equal or close to equal, and Tukey’s method then gives exact overall confidence or significance levels. In the two-way layout there is an even stronger reason to prefer balanced designs. If the numbers of individuals differ

among treatments (an unbalanced design), several alternative analyses of the data are possible. These analyses answer different sets of questions, and you must decide which questions you want to answer. All the sets of questions and all the analyses collapse to just one in the balanced case. This makes interpreting your data much simpler.

To understand the questions that two-way ANOVA answers, return to the advertising study in Example 29.8. In this section, we will assume that we know the actual population mean responses for each treatment. That is, we deal with an ideal situation in which we don't have to worry about random variation in the mean responses.

EXAMPLE 29.9 Main effects with no interaction

Here again is the two-way layout of the advertising study in Example 29.8. The numbers in the 6 cells are now made-up values of the population mean responses to the 6 treatments:

		Variable C Repetitions		
		1 time	3 times	5 times
Variable R	30 seconds	30	45	50
	90 seconds	40	55	60

The mean “intention to buy” scores increase with more repetitions of the camera ad and also increase with the length of the ad. The means increase *by the same amount* (10 points) when we move from 30 seconds to 90 seconds, *no matter how many times the ad is shown*. Turning to the other variable, the effect of moving from 1 to 3 repetitions is the same (15 points) for both 30-second and 90-second ads, and the effect of moving from 3 to 5 repetitions is also the same (5 points). Because the result of changing the value of one variable is the same for all values of the other variable, we say that there is *no interaction* between the two variables.

Now average the mean responses for 30 seconds and for 90 seconds. The average for 30 seconds is

$$\frac{30 + 45 + 50}{3} = \frac{125}{3} = 41.7$$

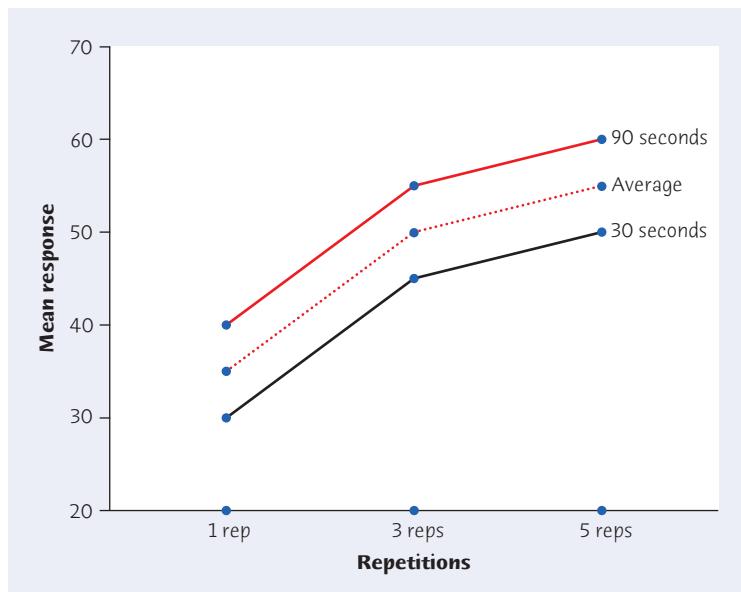
and the average for 90-second ads is

$$\frac{40 + 55 + 60}{3} = \frac{155}{3} = 51.7$$

Because the averages for the two lengths differ, we say that there is a *main effect* for length. Similarly average the mean responses for each number of repetitions. They are 35 for 1 repetition, 50 for 3, and 55 for 5. So changing the number of times the ad is shown has an “on the average” effect on the response. There is a *main effect* for repetitions. ■

Figure 29.5 plots the cell means from Example 29.9. The two solid lines joining 1, 3, and 5 repetitions for 30-second and for 90-second ads are *parallel*. This reflects the fact that the mean response always increases by 10 points when we move from 30-second to 90-second ads, no matter how many times the ad is shown. *Parallel lines in a plot of means show that there is no interaction*. It doesn't matter which variable you choose to place on the horizontal axis.

To see the main effect of repetitions, look at the average response for 30 and 90 seconds at each number of repetitions. This average is the dotted line in the plot. It changes as we move from 1 to 3 to 5 repetitions. A *variable has a main effect when the average response differs for different values of the variable*. "Average" here means averaged over all the values of the other variable. A main effect of repetitions is present in Figure 29.5 because the dotted "average" line is not horizontal. A main effect for length is also present, but it can't be seen directly in the plot.

**FIGURE 29.5**

Plot of the means from Example 29.9. In addition to the means themselves, the plot displays their average for each number of repetitions as a dotted line. The parallel lines show that there is no interaction between the two factors.

EXAMPLE 29.10 Interactions and main effects

Now look at this different set of mean responses for the advertising study:

		Variable C Repetitions		
		1 time	3 times	5 times
Variable R	30 seconds	30	45	50
	90 seconds	40	45	40

FIGURE 29.6

Plot of the means from Example 29.10, along with the average for each number of repetitions. The lines are not parallel, so there is an interaction between length and number of repetitions.

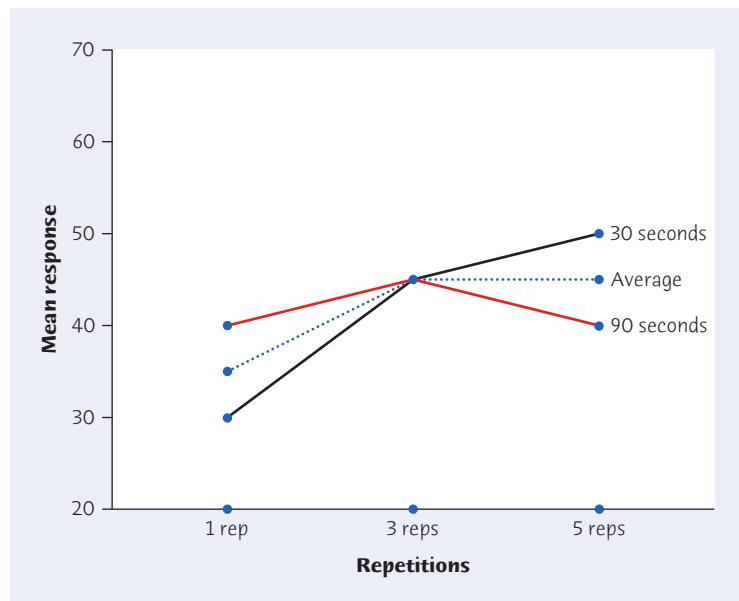


Figure 29.6 plots these means, with the means for 30-second ads connected by solid black lines and those for 90-second ads by solid red lines.

These means reflect the case in which subjects grow annoyed when the ad is both longer and repeated more often. The mean responses for the 30-second ad are the same as in Example 29.9. But showing the 90-second ad 3 times increases the mean response by only 5 points, and showing it 5 times drops the mean back to 40. There is an interaction between length and repetitions: the difference between 30 and 90 seconds changes with the number of repetitions, so that the solid black and red lines in Figure 29.6 are not parallel.

There is still a main effect of repetitions because the average response (dotted line in Figure 29.6) changes as we move from 1 to 3 to 5 repetitions. What about the effect of length? The average over all values of repetitions for 30-second ads is

$$\frac{30 + 45 + 50}{3} = \frac{125}{3} = 41.7$$

For 90-second ads this average is the same,

$$\frac{40 + 45 + 40}{3} = \frac{125}{3} = 41.7$$

On the average over all numbers of repetitions, changing the length of the ad has no effect. There is no *main effect* for length. ■

INTERACTIONS AND MAIN EFFECTS

An **interaction** is present between factors R and C in a two-way layout if the change in mean response when we move between two values of R is different for different values of C. (We can interchange the roles of R and C in this statement.)

A **main effect** for factor R is present if, when we average the responses for a fixed value of R over all values of C, we do not get the same result for all values of R.

Main effects may have little meaning when interaction is present. After all, interaction says that the effect of changing one of the variables is different for different values of the other variable. The main effect, as an “on the average” effect, may not tell us much. In Example 29.10, there is no main effect for the length of the camera ad. But length certainly matters—Figure 29.6 shows that there is little point in paying for 90-second ads if the same ad will be shown several times during the program. That’s the interaction of length with number of repetitions.



EXAMPLE 29.11 Which effect is more important?

There are no simple rules for interpreting results from two-way ANOVA when strong interaction is present. You must look at plots of means and think. Figure 29.7 displays two different mean plots for a study of the effects of classroom conditions on the performance of normal and hyperactive schoolchildren. The two conditions are “quiet” and “noisy,” where the noisy condition is actually the usual environment in elementary school classrooms.

There is an interaction: normal children perform a bit better under noisy conditions, but hyperactive children perform slightly less well under noisy conditions. The interaction is exactly the same size in the two plots of Figure 29.7. To see this, look at the slopes of the “Normal” lines in the two plots: they are the same. The slopes of the two “Hyperactive” lines are also the same. So the size of the gap between normal and hyperactive changes by the same amount when we move from quiet to noisy in both plots even though the gap is much larger in Figure 29.7(b).

In Figure 29.7(a), this interaction is the most important conclusion of the study. Both main effects are small: normal children do a bit better than hyperactive children, for example, but not a great deal better on the average.

In Figure 29.7(b), the main effect of “hyperactive or not” is the big story. Normal children perform much better than hyperactive children in both environments. The interaction is still there, but it is not very important in the face of the large difference in average performance between hyperactive and normal children. ■

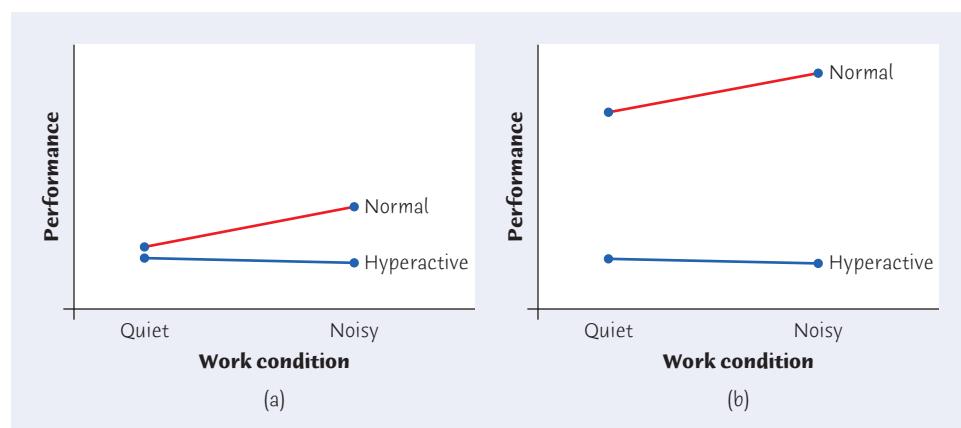


FIGURE 29.7

Plots of two sets of means for a study comparing the performance of normal and hyperactive children under two conditions, for Example 29.11.

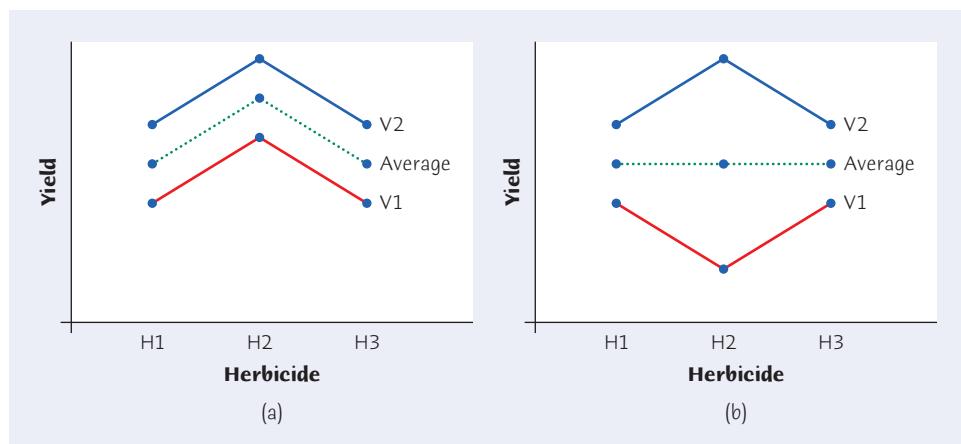
In this section we pretended that we knew the population means so that we could discuss patterns without needing statistical inference. In practice, we don't know the population means. However, plotting the sample means for all groups is an essential part of data analysis for a two-way layout. Look for interaction and main effects just as we did in this section. Of course, in real data you will almost never find exactly parallel lines representing exactly no interaction. Two-way ANOVA inference helps guide you because it assesses whether the interaction in the data is statistically significant. We are now ready to introduce two-way ANOVA inference.

APPLY YOUR KNOWLEDGE

Figure 29.8 shows two different plots of means for a two-way study that compares the yields of two varieties of soybeans (V1 and V2) when three different herbicides (H1, H2, and H3) are applied to the fields. Exercises 29.9 and 29.10 ask you to interpret these plots.

FIGURE 29.8

Plots of two sets of possible means for a study in which the two factors are soybean variety and type of herbicide, for Exercises 29.9 and 29.10. The plots also show the average for each herbicide (dotted line).



29.9 Recognizing effects. Consider the mean responses plotted in Figure 29.8(a).

- Is there an interaction between soybean variety and herbicide type? Why or why not?
- Is there a main effect of herbicide type? Why or why not?
- Is there a main effect of soybean variety? Why or why not?

29.10 Recognizing effects. Consider the mean responses plotted in Figure 29.8(b).

- Is there an interaction between soybean variety and herbicide type? Why or why not?
- Is there a main effect of herbicide type? Why or why not?
- Is there a main effect of soybean variety? Why or why not?

29.11 Angry women, sad men. What are the relationships among the portrayal of anger or sadness, sex, and the degree of status conferred? Sixty-eight subjects were randomly assigned to view a videotaped interview in which either a male or a female professional described feeling either anger or sadness. The people being

interviewed (we'll call them the "targets") wore professional attire and were ostensibly being interviewed for a job. The targets described an incident in which they and a colleague lost an account and, when asked by the interviewer how it made them feel, responded either that the incident made them feel angry or that it made them feel sad. The subjects were divided into four groups; each group evaluated one of the four types of interviews.⁴ After watching the interview, subjects evaluated the target on a composite measure of status conferral that included items assessing how much status, power, and independence the target deserved in his or her future job. The measure of status ranged from 1 = none to 11 = a great deal. Here are the summary statistics:

Treatment	<i>n</i>	\bar{x}	<i>s</i>
Males expressing anger	17	6.47	2.25
Females expressing anger	17	3.75	1.77
Males expressing sadness	17	4.05	1.61
Females expressing sadness	17	5.02	1.80

-
- (a) Display the four treatment means in a two-way layout similar to those given in Exercises 29.9 and 29.10.
 - (b) Plot the means and discuss the interaction and the two main effects.
-

INFERENCE FOR TWO-WAY ANOVA

Inference for two-way ANOVA is in many ways similar to inference for one-way ANOVA. Here is a brief outline:

1. Find and plot the group sample means. Study the plot to understand the interaction and main effects. Do data analysis to check the conditions for ANOVA.
2. Use software for basic ANOVA inference. There are now three F tests with three *P*-values, which answer the questions
 - Is the interaction statistically significant?
 - Is the main effect for variable R statistically significant?
 - Is the main effect for variable C statistically significant?
3. You may wish to carry out a follow-up analysis. For example, Tukey's method makes pairwise comparisons among the means of all *rc* treatment groups.

We will illustrate two-way ANOVA inference with several examples. In the first example, the interaction is both small and insignificant, so that the message is in the main effects.



EXAMPLE 29.12 Computer-assisted instruction

STATE: A study in education concerned computer-aided instruction in the use of “Blissymbols” for communication. Blissymbols are pictographs (think of Egyptian hieroglyphs) sometimes used to help learning-disabled children. Normal-ability schoolchildren were randomly assigned to treatment groups. There are four groups in a two-way layout:

		Learning Style	
		Active	Passive
Placement	Before	Group 1	Group 2
	During	Group 4	Group 3

The two factors are the learning style (active or passive) of the lesson and the placement of the material to be learned (Blissymbols presented before compounds, or Blissymbols and compounds shown during a joint presentation). The response variable is the number of symbols recognized correctly, out of 24, in a test taken after the lesson. Table 29.3 displays the data.⁵

TABLE 29.3 Test scores in an education study

GROUP 1	GROUP 2	GROUP 3	GROUP 4
14	4	11	12
14	4	8	22
16	9	8	9
14	8	7	14
6	15	14	20
10	8	9	15
12	9	7	9
13	13	8	10
12	13	12	11
12	12	10	11
13	12	6	15
9	10	3	6

PLAN: Plot the sample means and discuss interaction and main effects. Check the conditions for ANOVA inference. Use two-way ANOVA F tests to determine the significance of interaction and main effects.

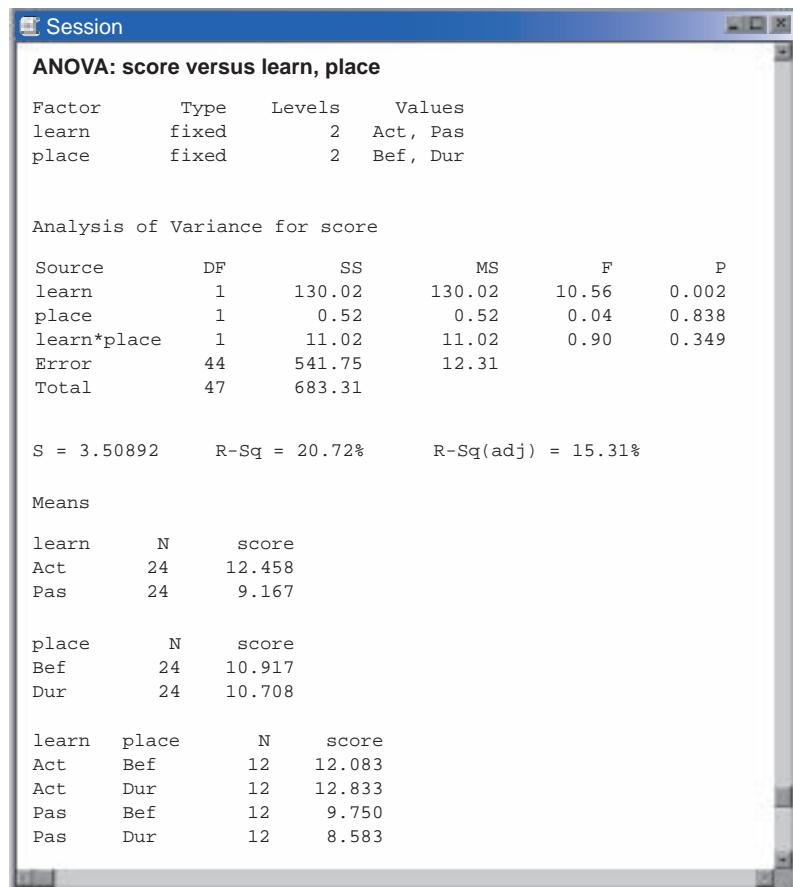
SOLVE: This is a balanced completely randomized experimental design. Figure 29.9 displays side-by-side stemplots of the scores for the 12 subjects in each group. Figure 29.10 shows two-way ANOVA output from Minitab. The stemplots show no departures from Normality strong enough to rule out ANOVA. (One observation, the 6 in Group 1, is slightly outside the range allowed by the $1.5 \times \text{IQR}$ rule. This would not be unusual among 48 Normally distributed observations.) You can check that the group standard deviations satisfy our rule of thumb that the largest (4.648) is no more than twice the smallest (2.678). We can proceed with ANOVA inference.

Group 1	Group 2	Group 3	Group 4
3	3	3 0	3
4	4 00	4	4
5	5	5	5
6 0	6	6 0	6 0
7	7	7 00	7
8	8 00	8 000	8
9 0	9 00	9 0	9 00
10 0	10 0	10 0	10 0
11	11	11 0	11 00
12 000	12 00	12 0	12 0
13 00	13 00	13	13
14 000	14	14 0	14 0
15	15 0	15	15 00
16 0	16	16	16
17	17	17	17
18	18	18	18
19	19	19	19
20	20	20	20 0
21	21	21	21
22	22	22	22 0

FIGURE 29.9

Side-by-side stemplots comparing the counts of correct answers for subjects in the four treatment groups from Example 29.12.

Minitab

**FIGURE 29.10**

Two-way ANOVA output from Minitab for the computer-assisted instruction study, for Example 29.12.

FIGURE 29.11

Minitab plots of the group means from the computer-assisted learning study, for Example 29.12. The two plots use the same four means. They differ only in the choice of which variable to mark on the horizontal axis.

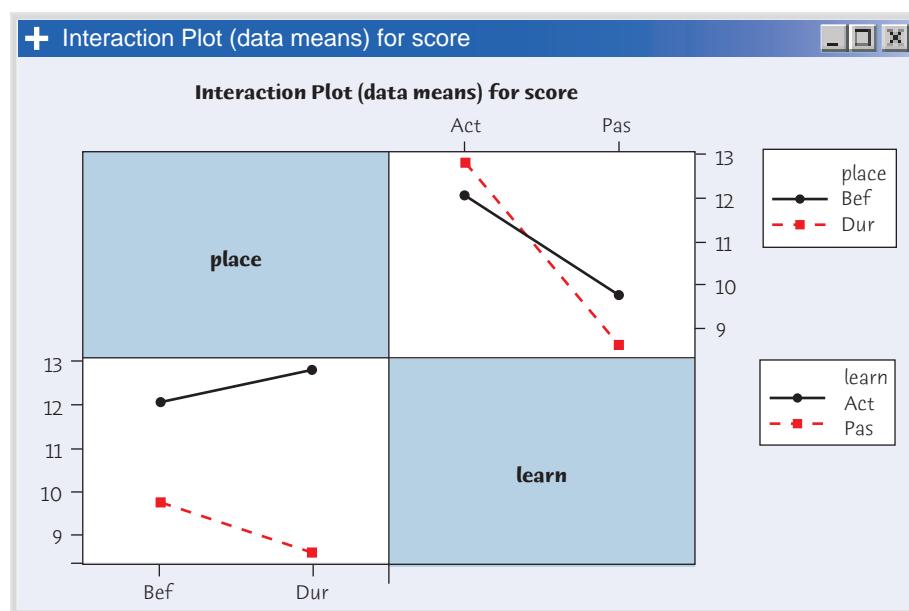
Minitab

Figure 29.11 shows the plots of means produced by Minitab. The two plots display the same 4 sample means; Minitab helpfully gives a plot with each variable marked on the horizontal axis. The plots are easy to interpret: the interaction and the main effect of placement are both small, and the main effect of learning style is quite large. The three F tests in the Minitab output substantiate what the plots of means show: interaction ($P = 0.349$) and placement ($P = 0.838$) are not significant, but learning style ($P = 0.002$) is highly significant.

CONCLUDE: Educators know that active learning almost always beats passive learning. This is the only significant effect that appears in these data. In particular, placement of material has very little effect under either style of learning. ■

The second example illustrates the situation in which there is significant interaction, but main effects are larger and more important. Think of Figure 29.7(b).

**EXAMPLE 29.13 Mycorrhizal colonies and plant nutrition**

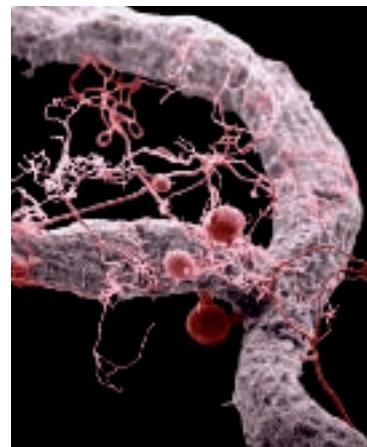
STATE: Mycorrhizal fungi are present in the roots of many plants. This is a symbiotic relationship, in which the plant supplies nutrition to the fungus and the fungus helps the plant absorb nutrients from the soil. An experiment compared the effects of adding nitrogen fertilizer to two genetic types of tomato plants, a normal variety that is susceptible to mycorrhizal colonies and a mutant that is not. Nitrogen was added at rates of 0, 28, or 160 kilograms per hectare (kg/ha). Here is the two-way layout for the 6 treatment combinations:

		Tomato Type	
		Mutant	Normal
Nitrogen	0 kg/ha	Group 1	Group 4
	28 kg/ha	Group 2	Group 5
	160 kg/ha	Group 3	Group 6

Six plants of each type were assigned at random to each amount of fertilizer. The response variables describe the levels of nutrients in a plant after 19 weeks, when the tomatoes are fully ripe. We will look at one response, the percent of phosphorus in the plant. Table 29.4 contains the data.⁶

PLAN: Plot the sample means and discuss interaction and main effects. Check the conditions for ANOVA inference. Use two-way ANOVA F tests to determine the significance of interaction and main effects.

SOLVE: This is a randomized block design. We are willing to assume that each of our two sets of tomato plants is an SRS from its genetic type. We consider the blocks (genetic types) as one of the factors in analyzing the data because we are interested in comparing the two genetic types. In addition to the two-way ANOVA, we might



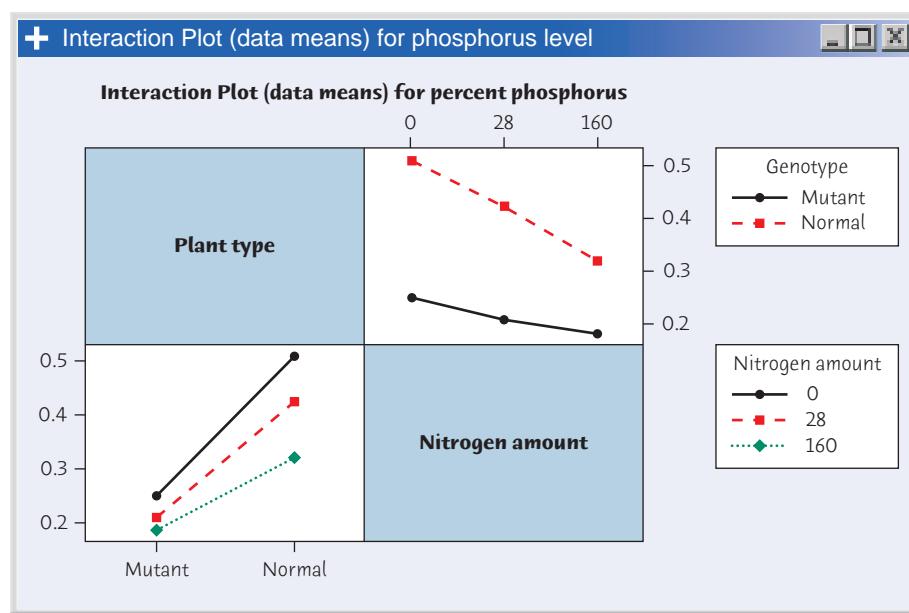
Science VU/M. F. Brown/Visuals Unlimited

TABLE 29.4 Percent of phosphorus in tomato plants

MUTANT			NORMAL		
GROUP	FERT	PCTPHOS	GROUP	FERT	PCTPHOS
1	0	0.29	4	0	0.64
1	0	0.25	4	0	0.54
1	0	0.27	4	0	0.53
1	0	0.24	4	0	0.52
1	0	0.24	4	0	0.41
1	0	0.20	4	0	0.43
2	28	0.21	5	28	0.41
2	28	0.24	5	28	0.37
2	28	0.21	5	28	0.50
2	28	0.22	5	28	0.43
2	28	0.19	5	28	0.39
2	28	0.17	5	28	0.44
3	160	0.18	6	160	0.34
3	160	0.20	6	160	0.31
3	160	0.19	6	160	0.36
3	160	0.19	6	160	0.37
3	160	0.16	6	160	0.26
3	160	0.17	6	160	0.27

FIGURE 29.12

Minitab plots of the group means for the study of phosphorus levels in tomatoes, for Example 29.13.

Minitab

consider separate one-way ANOVAs for mutant and normal types to draw separate conclusions about the effect of fertilizer on each type.

Figure 29.12 displays Minitab's plots of the 6 sample means. The lines are not parallel, so interaction is present. The interaction is rather small compared with the main effects. The main effect of type is expected: normal plants have higher phosphorus levels than the mutants at all levels of fertilization because they benefit from symbiosis with the fungus. The main effect of fertilizer is a bit surprising: phosphorus level goes down as the level of nitrogen fertilizer increases.

Examination of the data (we don't show the details) finds no outliers or strong skewness. But the largest sample standard deviation (0.08329 in Group 4) is much larger than twice the smallest (0.01472 in Group 3). As the number of treatment groups



increases, even samples from populations with exactly the same standard deviation are more likely to produce sample standard deviations that violate our "twice as large" rule of thumb. (Think of comparing the shortest and tallest person among more and more people.) So our rule of thumb is often conservative for two-way ANOVA. Nonetheless, ANOVA inference may not give correct P-values for these data. The P-values for the three two-way ANOVA F tests are $P = 0.008$ for interaction and $P < 0.001$ for both main effects. These agree with the mean plots and are so small that even if not accurate they strongly suggest significance.

CONCLUDE: Normal plants, with their mycorrhizal colonies, have higher phosphorus levels than mutants that lack such colonies. Nitrogen fertilizer actually reduces phosphorus levels in both types of plants. The reduction is stronger for normal plants, but this interaction is not very large in practical terms. ■

Finally, here is an example in which strong interaction makes one of the main effects meaningless. Two-way ANOVA with strong interaction is often difficult to interpret simply, as this example also illustrates.

EXAMPLE 29.14 Better corn for heavier chicks?

STATE: Corn varieties with altered amino acid content can have advantages in feeding animals. Here is an excerpt from a study that compared normal corn (“norm” in the data file) with two altered varieties called opaque-2 (“opaq”) and floury-2 (“flou”).



CATERPILLARS

Nine treatments were arranged in a 3×3 factorial experiment to compare opaque-2, floury-2 and normal corn at dietary protein levels of 20, 16, and 12%. Corn-soybean meal diets containing either opaque-2, floury-2 or normal corn were formulated so that, for a given protein level, an equivalent amount of corn protein was supplied by each corn. Male broiler-type chicks were randomly allotted to treatments at 1 day of age. Feed and water were provided ad libitum. Chicks were weighed at weekly intervals until termination of the experiment at 21 days.⁷

There are 10 chicks in each group. The response variable is the weight in grams after 21 days. Which combinations of corn type and protein content lead to fastest growth?

PLAN: Plot the sample means and discuss interaction and main effects. Check the conditions for ANOVA inference. Use two-way ANOVA F tests to determine the significance of interaction and main effects. If necessary, use Tukey pairwise comparisons to identify significant differences among treatments.

SOLVE: This is a balanced completely randomized experiment with 9 treatments. Figure 29.13 shows Minitab's plots of the sample means. The mean weight of the chicks increases with the percent of protein in the diet, as expected. We are primarily interested in comparing the three types of corn. There are important interaction effects. Normal corn does poorest of the three corn types at 12% protein and best at 20%. Floury is best at both 12% and 16%. Opaque is always inferior to floury and beats normal corn only at 12% protein.

Minitab

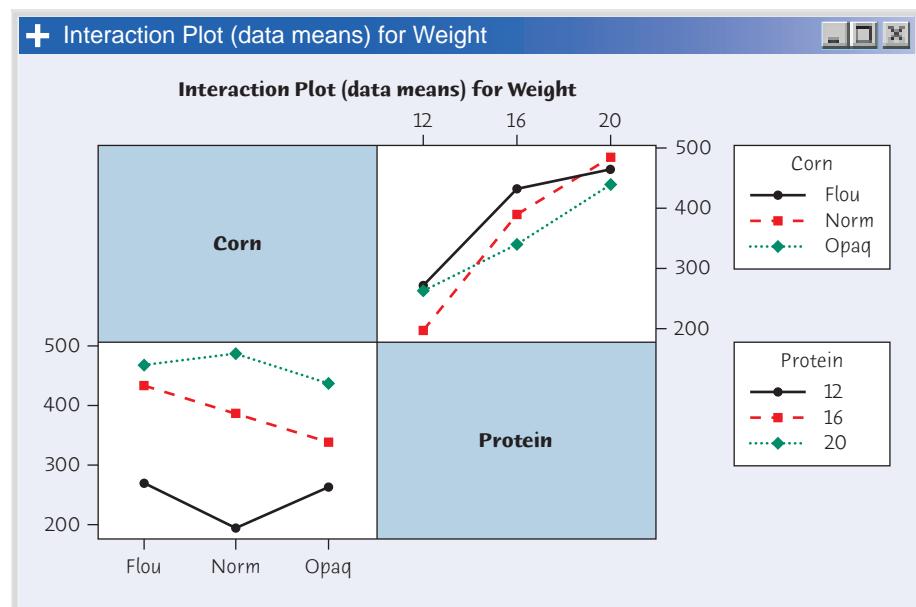


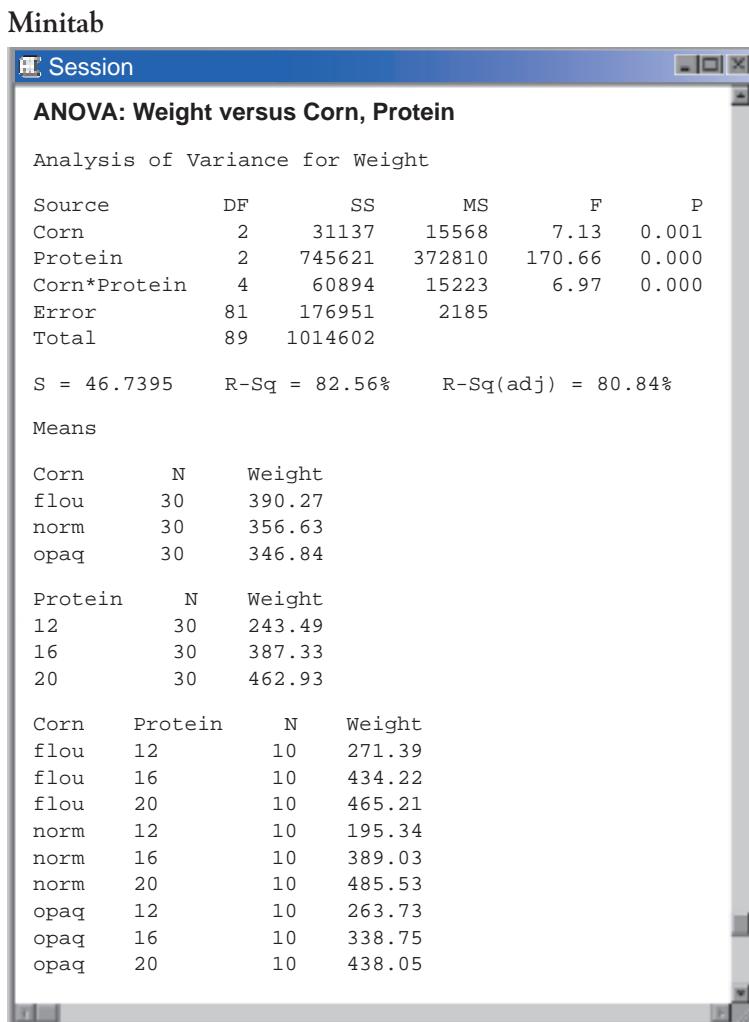
FIGURE 29.13

Minitab plots of the mean weights of 21-day-old chicks fed 9 different diets, for Example 29.14.

Although we don't show the details, ANOVA is justified. There are no outliers or strong skewness. Although the largest sample standard deviation (61.16 for Group 4) is a bit more than twice the smallest (25.99 for Group 6), this is common when we have 9 groups even when all populations really have the same σ .

Figure 29.14 contains Minitab's ANOVA output. We included the means for the 9 groups, for the three types of corn, and for the three levels of protein. The main effects can be seen in the different mean weights for the corn types and for the protein levels. The three F tests are all highly significant. Although there is a substantial main effect for corn (the means range from 346.84 grams for opaque to 390.27 grams for floury), this has little meaning in the light of the interaction that we just described.

When strong interaction makes one or both main effects hard to interpret, it is often useful to find the Tukey pairwise multiple comparisons for all the treatment groups. One way to do this is to do a one-way ANOVA with just "group" as the explanatory variable. There are 36 pairwise comparisons among 9 groups. Minitab's Tukey output is both long and hard to understand in such cases. Here is a condensed version using an idea that some software (though not Minitab) implements. Arrange

**FIGURE 29.14**

Two-way ANOVA output from Minitab for the study of the effect of diet on growth, for Example 29.14.

the 9 groups in the order of their sample means, from smallest to largest. We have identified the groups both by their group number in the data file and by the treatment. Connect with an underline all pairs of treatments that *do not* differ significantly at the overall 5% level:

Treatment:	N12	O12	F12	O16	N16	F16	O20	F20	N20
Group:	4	7	1	8	5	2	9	3	6
	-----	-----	-----	-----	-----	-----	-----	-----	-----

Group 4 has significantly smaller mean weight than any other group. Groups 7 and 1 do not differ significantly but are higher than Group 4 and lower than all other groups. Group 8 is not significantly different from 5 but is higher than 4, 7, and 1 and lower than 2, 9, 3, and 6. And so on. The most interesting finding is that at the high end Groups 2, 9, 3, and 6 do not have significantly different mean weights. Three of these are the three 20% protein groups, but floury corn with 16% protein belongs with these three.

CONCLUDE: More protein clearly helps chicks grow faster. The three types of corn do not differ significantly when the diet has 20% protein. Flurry corn is superior to both opaque and normal corn at middle (16%) and low (12%) protein levels, though not all differences at these levels are statistically significant. ■

APPLY YOUR KNOWLEDGE

29.12 Hooded rats: social play times. How does social isolation during a critical developmental period affect the behavior of hooded rats? Psychology students assigned 24 young female rats at random to either isolated or group housing, then similarly assigned 24 young male rats. This is a randomized block design with the gender of the 48 rats as the blocking variable and housing type as the treatment. Later, the students observed the rats at play in a group setting and recorded data on three types of behavior (object play, locomotor play, and social play).⁸ The data file records the time (in seconds) that each rat devoted to social play during the observation period.  **SOCIALTIMES**

- Make a plot of the 4 group means. Is there a large interaction between gender and housing type? Which main effect appears to be more important?
- Verify that the conditions for ANOVA inference are satisfied.
- Give the complete two-way ANOVA table. What are the F statistics and P -values for interaction and the two main effects? Explain why the test results confirm the tentative conclusions you drew from the plot of means.



Biosphoto/J.M. Labat/P. Rocher/Peter Arnold

29.13 Hooded rats: social play counts. The researchers who conducted the study in the previous exercise also recorded the number of times each of three types of behavior (object play, locomotor play, and social play) occurred. The data file contains the counts of social play episodes by each rat during the observation period. Use two-way ANOVA to analyze the effects of gender and housing.  **SOCIALCOUNTS**



29.14 Metabolic rates in caterpillars. Professor Harihiko Itagaki and his students have been measuring metabolic rates in tobacco hornworm caterpillars (*Manduca sexta*) for years. The researchers do not want the metabolic rates to depend on which analyzer

they use to obtain the measurements. They therefore make 6 repeated measurements on each of 3 caterpillars with each of 3 analyzers.⁹  CATERPILLARS

- Use software to give the two-way ANOVA table. The researchers used caterpillar as a blocking variable because metabolic rates vary from individual to individual. These three caterpillars are not of interest in themselves. They represent the larger population of caterpillars of this species. We should therefore avoid inference about the main effect of caterpillars.
 - Explain why the researchers will be unhappy if there is a significant interaction between analyzer and caterpillar. What does your analysis show about the interaction?
 - Is there a significant effect for analyzer? Do you think the researchers will be happy with this result?
-

SOME DETAILS OF TWO-WAY ANOVA*

All ANOVA F statistics work on the same principle: compare the variation due to the effect being tested with a benchmark level of variation that would be present even if that effect were absent. The three F tests for two-way ANOVA use the same benchmark as the one-way ANOVA F test, namely, the variation among individual responses within the same treatment group.

In two-way ANOVA, we have two factors (explanatory variables) that form treatments in a two-way layout. Factor R has r values and Factor C has c values, so that there are rc treatments. The same number n of subjects are assigned to each treatment. The two-way layout that results is as follows:

		Column Factor C			
		1	2	...	c
Row	1	n subjects	n subjects	...	n subjects
	2	n subjects	n subjects	...	n subjects
	\vdots	\vdots	\vdots		\vdots
Factor R	r	n subjects	n subjects	...	n subjects

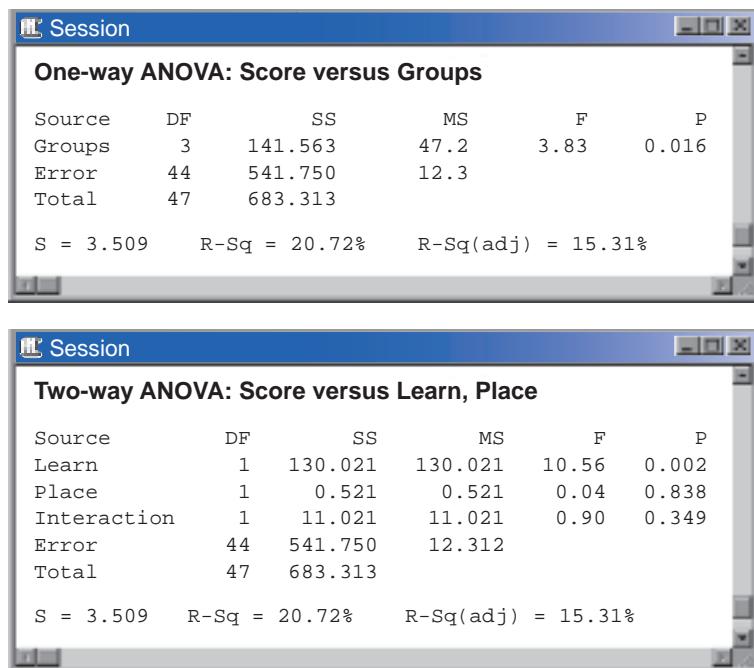
The number of treatments is $I = rc$

The total number of observations is $N = rcn$

Figure 29.15 presents both one-way and two-way ANOVA output for the same data, from the computer-assisted education study in Example 29.12. The one-way analysis just compares the means of rc treatments, ignoring the two-way layout. In discussing one-way ANOVA, we called the number of treatments I . Now $I = rc$. The two-way analysis takes into account that each treatment is formed by combining a value of R with a value of C.

*This optional material requires the optional section on some details of ANOVA on text pages 640–643.

Minitab

**FIGURE 29.15**

Compare the sums of squares in these one-way and two-way ANOVA outputs for the study in Example 29.12.

In both settings, analysis of variance breaks down the overall variation in the observations into several pieces. The overall variation is expressed numerically by the **total sum of squares**

total sum of squares

$$\text{SSTO} = \sum (\text{individual observation} - \text{mean of all observation})^2$$

where the sum runs over all N individual observations. If we divide SSTO by $N - 1$, we get the variance of the observations. So SSTO is closely related to a familiar measure of variability. Because SSTO uses just the N individual observations, it is the same for both one-way and two-way analyses. You can see in Figure 29.15 that SSTO = 683.313 for these data.

One-way ANOVA. We saw in Chapter 25 that the one-way ANOVA F test compares the variation among the I treatment means with the variation among responses to the same treatment. If the means vary more than we would expect based on the variation among subjects who receive the same treatment, that's evidence of a difference among the mean responses in the I populations.

Let's give a bit more detail. One-way ANOVA breaks down the total variation into the sum of two parts:

$$\begin{aligned} \text{total variation among responses} &= \text{variation among treatment means} + \text{variation} \\ &\quad \text{among responses to the same treatment} \end{aligned}$$

$$\text{total sum of squares} = \text{sum of squares for groups}$$

$$+ \text{sum of squares for error}$$

$$\text{SSTO} = \text{SSG} + \text{SSE}$$

Formulas for the sums of squares SSG and SSE appear on text page 641 as the numerators of the mean squares MSG and MSE, but we won't concern ourselves with the algebra. Remember that "error" is the traditional term in ANOVA for variation among observations. It doesn't imply that a mistake has been made. In the one-way output in Figure 29.15, you see that the breakdown for these data is

$$\begin{aligned} \text{SSTO} &= \text{SSG} + \text{SSE} \\ 683.313 &= 141.563 + 541.750 \end{aligned}$$

The one-way ANOVA F test is formed in two stages (text page 641):

1. Divide each sum of squares by its *degrees of freedom* to get the *mean squares* MSG for groups and MSE for error,

$$\text{MSG} = \frac{\text{SSG}}{I - 1} \quad \text{MSE} = \frac{\text{SSE}}{N - I}$$

2. The one-way ANOVA F statistic compares MSG to MSE,

$$F = \frac{\text{MSG}}{\text{MSE}}$$

Find the *P*-value from the *F* distribution with $I - 1$ and $N - I$ degrees of freedom.

ANOVA table

The **analysis of variance table** in the output reports sums of squares, their degrees of freedom, mean squares, and the *F* statistic with its *P*-value.

Two-way ANOVA. Now look at the two-way analysis of variance table in Figure 29.15:

- The total sum of squares and the error sum of squares are the same as in the one-way analysis.
- The sum of squares for groups in one-way is the sum of the three sums of squares for main effects and interaction in two-way.

This is the heart of two-way analysis of variance: break down the variation among the rc groups into three parts: variation due to the main effect of Factor R, variation due to the main effect of Factor C, and variation due to interaction between the two factors. Each type of variation is measured by a sum of squares. The formulas for the two main effects sums of squares are similar to that for the one-way sum of squares for groups, but we will again ignore the algebraic details. The interaction sum of squares is best thought of as what's left over: the variation among treatments that isn't explained by the two main effects. In symbols,

$$\begin{aligned} \text{total sum of squares} &= \text{sum of squares for main effect of Factor R} \\ &\quad + \text{sum of squares for main effect of Factor C} \\ &\quad + \text{sum of squares for interaction between R and C} \\ &\quad + \text{sum of squares for error} \\ \text{SSTO} &= \text{SSR} + \text{SSC} + \text{SSRC} + \text{SSE} \end{aligned}$$

Each of these sums of squares has a degrees of freedom, and these also break down in the same way:

$$\begin{aligned}\text{total df} &= \text{df for main effect of Factor } R \\ &\quad + \text{df for main effect of Factor } C \\ &\quad + \text{df for interaction between } R \text{ and } C \\ &\quad + \text{df for error} \\ rcn - 1 &= (r - 1) + (c - 1) + (r - 1)(c - 1) + rc(n - 1)\end{aligned}$$

You can check that the total degrees of freedom in the line above are $N - 1$ and the degrees of freedom for error are $N - I$, the same as for one-way ANOVA. The one-way degrees of freedom for groups are the sum of the degrees of freedom for the three two-way effects.

EXAMPLE 29.15 Comparing one-way and two-way ANOVA

The data behind Figure 29.15 appear in Table 29.3. The two factors are R = placement of Blissymbol elements and C = learning style. Factor R has $r = 2$ values: Before and During. Factor C has $c = 2$ values: Active and Passive. There are $I = 4$ treatments and $n = 12$ subjects assigned to each treatment, resulting in $N = 48$ observations.

The total degrees of freedom are $N - 1 = 47$. The degrees of freedom for error are $N - I = 48 - 4 = 44$. In the one-way analysis, the degrees of freedom for groups are $I - 1 = 3$. The two-way analysis breaks this into degrees of freedom $r - 1 = 1$ for Factor R , $c - 1 = 1$ for Factor C , and $(r - 1)(c - 1) = 1 \times 1 = 1$ for interaction.

Here are the two breakdowns of the total variation and the degrees of freedom that appear in Figure 29.15:

One-Way		Two-Way			
	Sums of squares	df			
SSG	141.563	3	SSR	130.021	1
SSE	541.750	44	SSC	0.521	1
SSTO	683.313	47	SSRC	11.021	1
			SSE	541.750	44
			SSE	683.313	47

The neat breakdown of SSG into three effects depends on the balance of the two-way layout. It doesn't hold when the counts of observations are not the same for all treatments. That's why two-way ANOVA for unbalanced data is more complicated and harder to interpret than for balanced data.

Two-way ANOVA F tests. Finally, form three F statistics exactly as in the one-way setting.

- Divide each sum of squares by its *degrees of freedom* to get the *mean squares* for the three effects and for error:

$$\text{MSR} = \frac{\text{SSR}}{r-1} \quad \text{MSC} = \frac{\text{SSC}}{c-1} \quad \text{MSRC} = \frac{\text{SSRC}}{(r-1)(c-1)} \quad \text{MSE} = \frac{\text{SSE}}{N-I}$$

- The three *F* statistics compare the mean squares for the three effects with MSE.

TWO-WAY ANOVA F TESTS

The *F* statistics for the three types of treatment effects in two-way ANOVA are

$$\text{For the main effect of Factor } R, \quad F = \frac{\text{MSR}}{\text{MSE}} \quad \text{with dfs } r-1 \text{ and } N-I$$

$$\text{For the main effect of Factor } C, \quad F = \frac{\text{MSC}}{\text{MSE}} \quad \text{with dfs } c-1 \text{ and } N-I$$

$$\text{For the interaction of } R \text{ and } C, \quad F = \frac{\text{MSRC}}{\text{MSE}} \quad \text{with dfs } (r-1)(c-1) \text{ and } N-I$$

In all cases, large values of *F* are evidence against the null hypothesis that the effect is not present in the populations.

APPLY YOUR KNOWLEDGE

29.15 Hooded rats: social play times. In Exercise 29.12 you carried out two-way ANOVA for a study of the effect of social isolation on hooded rats. The response variable is the time (in seconds) that a rat devoted to social play during an observation period. Start your work in this exercise with your two-way ANOVA table from Exercise 29.12.  SOCIALTIMES

- Explain how the sums of squares from the two-way ANOVA table can be combined to obtain the one-way ANOVA sum of squares for the 4 groups (SSG). What is the value of SSG?
- Give the degrees of freedom, mean square (MSG), and *F* statistic for testing for the effect of groups.
- Is there a significant effect of group on the amount of time in play? Give and interpret the *P*-value in the context of this experiment.
- Use software to carry out one-way ANOVA of time on group. Verify that your results in parts (b), (c), and (d) agree with the software output.

29.16 Hooded rats: social play counts. In Exercise 29.13 you carried out two-way ANOVA for a study of the effect of social isolation on hooded rats. The response variable is the count of social play episodes during an observation period. Start your work in this exercise with your two-way ANOVA table from Exercise 29.13.  SOCIALCOUNTS

- How many treatment groups are there in this experiment?
- Starting from the two-way ANOVA table, create a one-way ANOVA table to examine the effect of the 4 groups.
- Is there a significant group effect? Give the appropriate hypotheses, test statistic, *P*-value, and conclusion in the context of this experiment.

CHAPTER 29 SUMMARY

CHAPTER SPECIFICS

- Two-way analysis of variance (ANOVA) compares the means of several populations formed by combinations of two factors R and C in a **two-way layout**.
- The **conditions for ANOVA** state that we have an **independent SRS** from each population (or a completely randomized or randomized block experimental design); that each population has a **Normal distribution**; and that all populations have the **same standard deviation**. In this chapter, we consider only examples that satisfy the additional conditions that the design producing the data is **crossed** (all combinations of the factors are present) and **balanced** (all factor combinations are represented by the same number of individuals).
- A factor has a **main effect** if the mean responses for each value of that factor, averaged over all values of the other factor, are not the same. The two factors **interact** if the effect of moving between two values of one factor is different for different values of the other factor. Plot the **treatment mean responses** to examine main effects and interaction.
- There are three **ANOVA F tests**: for the null hypotheses of no main effect for Factor R, no main effect for Factor C, and no interaction between the two factors.
- **Follow-up analysis** is often helpful in both one-way and two-way ANOVA settings. **Tukey pairwise multiple comparisons** give confidence intervals for all differences among treatment means with an **overall level of confidence**. That is, we can be (say) 95% confident that *all* the intervals simultaneously capture the true population differences between means.

STATISTICS IN SUMMARY

Here are the most important skills you should have acquired from reading this chapter.

A. Recognition

1. Recognize the two-way layout, in which we have a quantitative response to treatments formed by combinations of values of two factors.
2. Recognize when comparing mean responses to the treatments in a two-way layout is helpful in understanding data.
3. Recognize when you can use two-way ANOVA to compare means. Check the data production, the presence of outliers, and the sample standard deviations for the groups you want to compare. Look for data production designs that are crossed and balanced.

B. Interpreting Two-Way ANOVA

1. Plot the sample means for the treatments. Based on your plot, describe the main effects and interaction that appear to be present.
2. Decide which effects are most important in practice. Pay particular attention to whether a large interaction makes one or both main effects less meaningful.
3. Use software to carry out two-way ANOVA inference. From the P-values of the three F tests, learn which effects are statistically significant.

C. Follow-up Analysis

- Decide when it is helpful to know which differences among treatment means are significant.
- Use software to carry out Tukey pairwise multiple comparisons among all the means you want to compare.
- Understand the meaning of the overall confidence level and the overall level of significance provided by Tukey's method for a set of confidence intervals or a set of significance tests.

LINK IT

In a one-way ANOVA, rejection of the null hypothesis is evidence of a difference in the means of the populations, but the result does not tell us *which* differences between the means are statistically significant. Typically, we perform a more detailed *follow-up analysis* to decide which of the means are different and to estimate how large these differences are. This can be done using Tukey's pairwise multiple comparisons. These tell us which populations are significantly different from each other and, through the associated confidence intervals, how large these differences are. Sometimes we have more specific questions about the population means that can be answered using further *contrasts* among the means.

The multiple regression models in Chapter 28 extend the simple linear regression model, allowing us to study the effect of several predictors on the mean of a quantitative response. Similarly, the two-way ANOVA extends the one-way ANOVA described in Chapter 25, allowing us to study the effect of two factors on the mean of a quantitative response. In the two-way ANOVA we have specific questions about the two factors. What are the main effects of each factor, and do the two factors interact? These questions can be answered by examining appropriate plots of the means and the *F* tests in the ANOVA table. In the presence of strong interaction, the main effects may be of little interest. As in the one-way ANOVA, a follow-up analysis may be necessary to obtain more detailed information about the effects of the factors.

CHECK YOUR SKILLS

29.17 In an ANOVA that compares three treatments, how many pairwise comparisons between two of these treatments are there?

- (a) two (b) three (c) four

29.18 After an ANOVA that compares four treatments, you calculate a separate two-sample *t* 95% confidence interval for each pairwise difference formed from the four population means. Are you 95% confident that these intervals capture all the true differences?

- (a) No, because 95% confidence applies to each interval alone, not to all the intervals together.

(b) No, because the ordinary two-sample *t* confidence intervals cannot be used for a difference between two population means when we have data from other populations as well.
 (c) Yes, because under the ANOVA conditions, following the *t* procedure gives correct 95% confidence intervals.

29.19 As part of an ANOVA that compares three treatments, you carry out Tukey pairwise tests at the overall 5% significance level. The Tukey tests find that μ_1 is significantly different from μ_2 but that the other two comparisons are not significant. You can be 95% confident that

- (a) $\mu_1 \neq \mu_2$ and $\mu_1 = \mu_3$ and $\mu_2 = \mu_3$.

(b) just $\mu_1 \neq \mu_2$; there is not enough evidence to draw conclusions about the other pairs of means.

(c) $\mu_1 = \mu_3$ and $\mu_2 = \mu_3$ and this implies that it must also be true that $\mu_1 = \mu_2$.

29.20 The purpose of two-way ANOVA is to learn about

(a) the means of two populations.

(b) the variances of two populations.

(c) the combined effects of two factors on a quantitative response.

29.21 A two-way ANOVA compares two exercise programs (A and B) and two exercise frequencies (twice a week and four times a week). The response variable is weight lost after 8 weeks of exercise. An interaction is present in the data when

(a) the mean weight loss is not the same for Program A and Program B.

(b) the mean weight loss under Program A is different for twice-weekly and four-times-weekly subjects.

(c) the difference between the mean weight losses for Programs A and B is not the same for twice-weekly and four-times-weekly subjects.

29.22 When the F test for interaction in a two-way ANOVA is significant,

(a) the main effects are never meaningful and should be ignored.

(b) the main effects must be interpreted with caution.

(c) interpret each main effect just as you would in a one-way ANOVA.

A student project measured the increase in the heart rates of fellow students when they stepped up and down for three minutes to the

beat of a metronome. The explanatory variables are step height ($Lo = 5.75$ inches, $Hi = 11.5$ inches) and metronome beat (Slow = 14 steps/minute, Med = 21 steps/minute, Fast = 28 steps/minute). The subject's heart rate was measured for 20 seconds before and after stepping. The response variable is the increase in heart rate during exercise.¹⁰ Use the ANOVA table in Figure 29.16 to answer the questions below.

29.23 How many treatment groups were there for this experiment?

(a) 3 (b) 5 (c) 6

29.24 What is the value of the test statistic for interaction between frequency and height?

(a) 218.4 (b) 109.2 (c) 0.97 (d) 0.393

29.25 What is the estimate for the common standard deviation σ ?

(a) 218.4 (b) 109.2 (c) 10.61

29.26 How many students were in the study?

(a) 29 (b) 30

(c) It cannot be determined from the ANOVA table.

29.27 Which statement below provides the best summary of the interaction between frequency and height?

(a) The interaction is significant at the 1% level but not at the 5% level.

(b) The interaction is significant at the 5% level but not at the 1% level.

(c) The interaction is not significant at either the 1% or the 5% level.

Minitab

Session					
Two-way ANOVA: HRinc versus Freq, Height					
Source	DF	SS	MS	F	P
Freq	2	4048.8	2024.4	17.99	0.000
Height	1	1920.0	1920.0	17.07	0.000
Interaction	2	218.4	109.2	0.97	0.393
Error	24	2700.0	112.5		
Total	29	8887.2			
S = 10.61 R-Sq = 69.62% R-Sq(adj) = 63.29%					

FIGURE 29.16

Two-way ANOVA table for Exercises 29.23 to 29.27.


CHAPTER 29 EXERCISES

29.28 Comparing tropical flowers. Example 25.1 (text page 623) describes a study of the lengths of three varieties of the tropical flower *Heliconia*. Table 25.1 gives the data. One-way ANOVA gives strong evidence ($F = 259.12, P < 0.0001$) that the three population mean lengths are not all equal. Software gives these Tukey 95% simultaneous confidence intervals:



- | | |
|------------------|----------------------------------|
| 6.752 to 9.021 | for $\mu_{bihai} - \mu_{red}$ |
| 10.165 to 12.670 | for $\mu_{bihai} - \mu_{yellow}$ |
| 2.375 to 4.688 | for $\mu_{red} - \mu_{yellow}$ |

- (a) How confident are you that all three of these intervals capture the true differences between pairs of population means?
 (b) Write a short summary of the results of the ANOVA, including the multiple comparisons.

29.29 Comparing tropical flowers: a contrast (optional). The data in Table 25.1 contain lengths for the *bihai* variety of *Heliconia* and for two forms (red and yellow) of the *caribaea* variety. We wonder (before looking at the data) whether the population mean for *bihai* differs from the average of the means for the two forms of *caribaea*.

- (a) What contrast expresses this comparison?
 (b) Give the null and alternative hypotheses for this comparison, find the sample contrast, and assess its statistical significance. What do you conclude?

- (c) Give a 90% confidence interval for the population contrast in part (a).

29.30 Does nature heal best? Our bodies have a natural electrical field that helps wounds heal. Might higher or lower levels speed healing? An experiment with newts investigated this question. Newts were randomly assigned to five groups. In four of the groups, an electrode applied to one hind limb (chosen at random) changed the natural field, while the other hind limb was not manipulated. Both limbs in the fifth (control) group remained in their natural state.¹¹

Table 29.5 gives data from this experiment. The “Group” variable shows the field applied as a multiple of the natural field for each newt. For example, “0.5” is half the natural field, “1” is the natural level (the control group), and “1.5” indicates a field 1.5 times natural. “Diff” is the response variable, the difference in the healing rate (in micrometers per hour) of cuts made in the experimental and control limbs of that newt. Negative values mean that the experimental limb healed more slowly. The investigators conjectured that nature heals best, so that changing the field from the natural state (the “1” group) will slow healing.

Carry out a one-way ANOVA to compare the mean healing rates for the five groups. Then perform Tukey multiple comparisons for the 10 pairs of population means. Use the “underline” method illustrated in Example 29.14 to display the complicated results. What do you conclude?

TABLE 29.5 Effect of electrical field on healing rate in newts

GROUP	DIFF								
0	-10	0.5	-1	1	-7	1.25	1	1.5	-13
0	-12	0.5	10	1	15	1.25	8	1.5	-49
0	-9	0.5	3	1	-4	1.25	-15	1.5	-16
0	-11	0.5	-3	1	-16	1.25	14	1.5	-8
0	-1	0.5	-31	1	-2	1.25	-7	1.5	-2
0	6	0.5	4	1	-13	1.25	-1	1.5	-35
0	-31	0.5	-12	1	5	1.25	11	1.5	-11
0	-5	0.5	-3	1	-4	1.25	8	1.5	-46
0	13	0.5	-7	1	-2	1.25	11	1.5	-22
0	-2	0.5	-10	1	-14	1.25	-4	1.5	2
0	-7	0.5	-22	1	5	1.25	7	1.5	10
0	-8	0.5	-4	1	11	1.25	-14	1.5	-4
		0.5	-1	1	10	1.25	0	1.5	-10
		0.5	-3	1	3	1.25	5	1.5	2
				1	6	1.25	-2	1.5	-5
				1	-1				
				1	13				
				1	-8				

29.31 Does nature heal best? A contrast (optional). The researchers who carried out the study of healing in newts described in the previous exercise conjectured that healing is fastest under natural conditions. If this is true, the population mean for the 1 group (natural field level) in Table 29.5 should be larger (or less negative) than the average of the population means for the other four groups.  NEWTSHEAL

- (a) What population contrast expresses this comparison? What null and alternative hypotheses in terms of this contrast should the researchers use to test their conjecture?
- (b) Carry out the test and state your conclusion.
- (c) After you see the five sample means, it is tempting to contrast the average for the 1 and 1.25 groups with the average for the other three groups. Explain carefully why this is not legitimate.

29.32 Exercise and heart rate. A student project measured the increase in the heart rates of fellow students when they stepped up and down for three minutes to the beat of a metronome. The explanatory variables are step height (Lo = 5.75 inches, Hi = 11.5 inches) and metronome beat (Slow = 14 steps/minute, Med = 21 steps/minute, Fast = 28 steps/minute). The subject's heart rate was measured for 20 seconds before and after stepping. The response variable is the increase in heart rate during exercise.¹² The data file has resting and final heart rates as well as the increase in heart rate for this study. If the randomization worked well, there should be no significant differences among the 6 groups in mean resting heart rate (variable HRrest in the data file).  HEARTRATE

- (a) How many pairwise comparisons are there among the means of 6 populations?
- (b) Use Tukey's method to compare these means at the overall 10% significance level.

29.33 Hooded rats: object play times. Exercise 29.12 (page 29-31) describes an experiment to study the effects of social isolation on the behavior of hooded rats. You have analyzed the effects on social play. Now look at another response variable, the time that a rat spends in object play during an observation period. The data file records the time (in seconds) that each rat devoted to object play.  OBJECTTIMES

- (a) Make a plot of the 4 group means. Is there a large interaction between gender and housing type? Which main effect appears to be more important?
- (b) Verify that the conditions for ANOVA inference are satisfied.
- (c) Give the complete two-way ANOVA table. What are the F statistics and P-values for interaction and the two main effects? Explain why the test results confirm the tentative conclusions you drew from the plot of means.

29.34 Hooded rats: object play counts. The researchers who conducted the study in the previous exercise also

 recorded the number of times each of three types of behavior (object play, locomotor play, and social play) occurred. The data file contains the counts of object play episodes for each rat during the observation period. Carry out a complete analysis of the effects of gender and housing type.



29.35 Herbicide and corn hybrids. Genetic engineering has produced new corn hybrids that resist the effects of herbicides. This allows more effective control of weeds, because herbicides don't damage the corn. A study compared the effects of the herbicide glufosinate on a number of corn hybrids. The percents of necrosis (leaf burn) 10 days after application of glufosinate for several application rates (kilograms per hectare) and three corn hybrids, two resistant and one not, are provided in the data file.¹³  CORNHYBRIDS

- (a) Construct a plot of means to examine the effects of application rate and hybrid and their interaction.
- (b) Are the conditions for ANOVA inference satisfied? Explain.

29.36 Girls' cross-country times. Ten runners each completed 10 five-kilometer races for the Thomas Worthington High School girls' cross-country team during the 2004 season. The data file contains each race time in minutes.¹⁴  GIRLSRUN

- (a) Prepare side-by-side boxplots to compare the distributions of race times for the 10 runners. Identify the best runner on the team.
- (b) Use runner as a blocking variable to see if there are significant differences in the overall mean race times for the different meets. Identify the appropriate parameters, state the null and alternative hypotheses, calculate the test statistic and P-value, and state your conclusion in context.
- (c) Should follow-up multiple comparisons be used to identify any significant differences? If so, summarize the appropriate procedure. If not, explain why not.

29.37 Girls' cross-country times: a contrast (optional). Continue your analysis of the data from the previous problem by comparing average race time during the first half of the season with the average for the second half. Identify a population contrast L that makes this comparison, estimate this contrast, and state your conclusion.

29.38 Girls' cross-country times: peak at the end?  Coaches like to see their athletes "peak" at the end of the season. Is there significant evidence that the average finish time for the 10th meet is significantly lower than the average finish time for the first nine meets?

29.39 Comparisons among means for one factor in a two-way analysis. We have illustrated Tukey pairwise comparisons among all treatment means in both one-way and two-way settings. The method can also compare the mean responses to just one of the two factors in a two-way setting:

do a one-way ANOVA on the two-way data listing only one factor as an explanatory variable. This combines data for all levels of the other factor, so it is useful only when interactions are small. Return to the data on phosphorus in tomato plants, Table 29.4. Do a one-way ANOVA that uses all 36 observations with fertilizer type as the only factor. Ask for Tukey pairwise comparisons among the three levels of fertilizer, with overall confidence level 95%.  TOMATOES

- (a) How many observations per group does your analysis use?
- (b) What do you conclude from the F statistic and its P -value?
- (c) Use Tukey: which pairwise differences of means for the 3 fertilizer levels are significant at the overall 5% level?
- (d) Do you think these pairwise comparisons are useful for these data? (*Hint:* What population does each of the three samples represent?)

29.40 Fourth-graders composing music. The Orff xylophone is often used in teaching music to children because it has removable bars that allow the teacher to present different options to the students. An education researcher used the Orff xylophone to examine the effect of tonality (pentatonic or harmonic minor) and number of xylophone bars (5 or 10) on the ability of fourth-graders to compose melodies that they could play repeatedly.¹⁵

- (a) Twelve children were randomly assigned to each combination of tonality and bar count. Give the two-way layout for this experiment.
- (b) Judges listened to tapes of the children's work and assigned scores for several aspects of the melodies. One response variable measured the extent to which children generated new musical ideas in consecutive 5-second intervals of their melodies. Here is the two-way ANOVA table for this variable (the publication does not give the group means):



Argus/Mike Schroeder/Peter Arnold

Source	df	Sum of squares	F	P
Tonality	1	44.08	0.43	0.52
Bar count	1	705.33	6.81	0.01
Interaction	1	50.02	0.48	0.49
Error	44	4556.54		

Comment on the significance of main effects and interaction. Then make a recommendation for teachers who want to use the Orff xylophone to encourage children to generate melodies with new musical ideas.

29.41 Cues for listening comprehension. In speaking, we often signal a shift to the next step in the discussion by saying

 things like "Let's talk about ..." or "Finally, ..." These are "discourse-signaling cues." Do such cues help listeners better comprehend a second language? One study of this question involved 80 Korean learners of English as a foreign language.¹⁶ Half of the 80 learners listened to a lecture in English with discourse-signaling cues, and the other half heard the same lecture without cues. After the lecture, half of each group were asked to summarize information from the lecture and the other half simply to recall information. Here are the group means and two-way ANOVA table for one response variable, a measure of comprehension of high-level information:

Task	Lecture	
	Cues	No cues
Summary	23	17
Recall	19	13

Source	df	Sum of squares	F	P
Lecture	1	649.80	16.582	0.000
Task	1	281.25	7.177	0.009
Interaction	1	0.20	0.005	0.943
Error	76	2978.30		

Use this information to comment on the impact of discourse-signaling cues and the type of task performed. (Assume that the data satisfy the conditions for ANOVA.)

29.42 Fertilization of bromeliads. Does the type of fertilization have a significant effect on leaf development for bromeliads? Researchers compared leaf production and death over a seven-month period for a control group (C), a group fertilized only with nitrogen (N), a group fertilized only with phosphorus (P), and a group that received both nitrogen and phosphorus (NP).¹⁷ In the nitrogen and phosphorus columns in the data file, 0 indicates no fertilization and 1 indicates fertilization. The "new" and "dead" columns give the total number of new or dead leaves produced over the seven months following fertilization. "Change" gives the difference "new" minus "dead."



Douglas Peebles/Stock Connection/IPNstock

- (a) Figure 29.17 displays the one-way ANOVA table and Tukey pairwise comparisons for the number of new leaves. Are there significant differences among the four group means? List the four group means in order from smallest to largest



Minitab

The Minitab session window displays the following output:

One-way ANOVA: New versus treatment

Source	DF	SS	MS	F	P
treatment	3	28.75	9.58	3.77	0.002
Error	28	71.25	2.54		
Total	31	100.00			

S = 1.595 R-Sq = 28.75% R-Sq(adj) = 21.12%

Individual 95% CIs For Mean Based on Poled StDev

Level	N	Mean	StDev	-----+-----+-----+
C	8	13.250	1.909	(-----*-----)
NP	8	15.625	1.685	(-----*-----)
P	8	14.625	1.302	(-----*-----)
	8	13.500	1.414	(-----*-----)
				13.2 14.4 15.6 16.8

Pooled StDev = 1.595

Tukey 95% Simultaneous Confidence Intervals
All Pairwise Comparisons among Levels of treatment
Individual confidence level = 98.92%

treatment = C subtracted from:

treatment	Lower	Center	Upper	-----+-----+-----+
N	0.198	2.375	4.552	(-----*-----)
NP	-0.802	1.375	3.552	(-----*-----)
P	-1.927	0.250	2.427	(-----*-----)
				-2.5 0.0 2.5 5.0

treatment = N subtracted from:

treatment	Lower	Center	Upper	-----+-----+-----+
NP	-3.177	-1.000	1.177	(-----*-----)
P	-4.302	-2.125	0.052	(-----*-----)
				-2.5 0.0 2.5 5.0

treatment = NP subtracted from:

treatment	Lower	Center	Upper	-----+-----+-----+
P	-3.302	-1.125	1.052	(-----*-----)
				-2.5 0.0 2.5 5.0

and give a summary of the Tukey procedure similar to that in Example 29.14.

- (b) Give the two-way ANOVA table for the number of new leaves. How do your conclusions from this analysis differ from your conclusion in part (a)?

29.43 Fertilization of bromeliads, continued. Repeat the analysis of the previous exercise using number of dead leaves as the response variable.

29.44 Fertilization of bromeliads, continued. Repeat the analysis of Exercise 29.42 using change in the number of leaves as the response variable.

FIGURE 29.17

One-way ANOVA table and Tukey pairwise multiple comparisons from Minitab for Exercise 29.42.

29.45 ANOVA with one observation per treatment (optional).

If a two-way layout has just one observation for each treatment, the mean square for error (MSE) is 0 because there is no variation within each treatment group. The usual two-way ANOVA F tests can't be done, because they have MSE in their denominators. But if we are willing to assume that there is no interaction between the two factors, the interaction mean square can be used as the denominator of F statistics for testing the two main effects. Software usually allows you to choose "no interaction." Here is an example.

Does the addition of sodium chloride change brick quality? One measure of quality is "resistance anisotropy," which

combines the results of ultrasound and mechanical tests. Researchers tested two types of bricks, one without sodium chloride (MPW) and one with sodium chloride (MPS), at three firing temperatures (850, 1000, and 1100 degrees Celsius). Here are the data for one brick per treatment:¹⁸



Brick type	Temperature (°C)	Resistance
MPW	850	3.29
MPW	1000	5.78
MPW	1100	4.84
MPS	850	2.57
MPS	1000	3.90
MPS	1100	2.57



EXPLORING THE WEB

29.46 Confidence in the banking system. The General Social Survey (GSS) is a sociological survey used to collect data on demographic characteristics and attitudes of residents of the United States. The survey is conducted by the National Opinion Research Center of the University of Chicago, which interviews face-to-face a randomly selected sample of adults (18 and older). SDA (Survey Documentation and Analysis) is a set of programs that allows you to analyze survey data and includes the GSS as part of its archive. In Exercises 25.43 and 25.44 (text page 654) you used data from the GSS to study the relationship between the average respondent age and confidence in the banking system. A one-way ANOVA showed a statistically significant difference between the age of the respondent and his or her confidence in the banking system. Download the data file following the directions in Exercise 25.43, and do Tukey's multiple comparisons at the 5% overall significance level to compare the mean age for the three levels of confidence in the banking system. Which pairs of means are different?

29.47 Find a two-way ANOVA. Find an example of a two-way ANOVA on the Web. The *Journal of the American Medical Association* (jama.ama-assn.org), *Science Magazine* (www.sciencemag.org), the *Canadian Medical Association Journal* (www.cmaj.ca), and the *Journal of Marketing Research* (www.journals.marketingpower.com/loi/jmkr) are possible sources. To help locate an article, look through the abstracts of articles. If you are having difficulty, consider study 1 in *Psychological Science*, 19 (2008), pp. 268–273. Once you find a suitable article, read the article and then briefly describe the study (including the two factors and the number of levels of each factor) and its conclusions. If the means are given, plot them and discuss the interaction. Whether or not the study is balanced, the test for interaction is interpreted in the same way. If P-values are reported, be sure to discuss them in your summary. Also, be sure to give the reference (either the Web link or the journal, issue, year, title of the paper, authors, and page numbers).

NOTES AND DATA SOURCES

1. We thank Charles Cannon of Duke University for providing the data. The study report is C. H. Cannon, D. R. Peart, and M. Leighton, "Tree species diversity in commercially logged Bornean rainforest," *Science*, 281 (1998), pp. 1366–1367.

(a) Make a plot of the group means (which are the same as the individual observations). Is the interaction between type of brick and firing temperature small? If so, it may be reasonable to assume that *in the population* there is no interaction.

(b) The researchers did fit a two-way ANOVA model without interaction. Use statistical software to obtain this two-way ANOVA table. What do you conclude about the main effects?

2. Modified from M. C. Wilson and R. E. Shade, "Relative attractiveness of various luminescent colors to the cereal leaf beetle and the meadow spittlebug," *Journal of Economic Entomology*, 60 (1967), pp. 578–580.
3. In some cases, we think of the blocks in a block design as a sample from some larger population of potential blocks. For example, a clinical trial may compare several medical therapies (the factor) using patients at several medical centers (the blocks). We are not interested in these particular medical centers; they simply represent all hospitals at which the therapies might be used. Such blocks are called "random." ANOVA with one or more random factors is an advanced topic. In other cases, the specific blocks are of interest in their own right. For example, a clinical trial might compare several therapies (the factor) on both female and male patients (the blocks). We randomly assign therapies among the women and separately among the men, a randomized block design. We are interested in the effect of gender on response to the therapies. Such blocks are called "fixed." In this chapter we allow randomized block designs with fixed blocks, analyzing the data by treating the blocks as the second factor in two-way ANOVA.
4. Victoria L. Brescoll and Eric L. Uhlmann, "Can an angry woman get ahead? Status conferral, gender and expression of emotion in the workplace," *Psychological Science*, 19 (2008), pp. 268–273. The description and data are based on study 1 in this article.
5. Orit E. Hetzroni, "The effects of active versus passive computer-assisted instruction on the acquisition, retention, and generalization of Blissymbols while using elements for teaching compounds," PhD thesis, Purdue University, 1995.
6. Data courtesy of David LeBauer, University of California, Irvine. Provided by Brigitte Baldi.
7. Simulated data, based on data summaries in G. L. Cromwell et al., "A comparison of the nutritive value of opaque-2, floury-2 and normal corn for the chick," *Poultry Science*, 57 (1968), pp. 840–847.
8. We thank Andy Niemiec and Robbie Molden for data from a summer science project at Kenyon College. Provided by Brad Hartlaub.
9. We thank Harihiko Itagaki and Andrew Veerde for data from a summer science project at Kenyon College. Provided by Brad Hartlaub.
10. Data from the EESEE story "Stepping Up Your Heart Rate."
11. Data provided by Drina Iglesia, Purdue University. The data are part of a larger study reported in D. D. S. Iglesia, E. J. Cragoe, Jr., and J. W. Venable, "Electric field strength and epithelialization in the newt (*Notophthalmus viridescens*)," *Journal of Experimental Zoology*, 274 (1996), pp. 56–62.
12. See Note 10.
13. Christopher Eric Mowen, "Use of glufosinate in glufosinate resistant corn hybrids," MS thesis, Purdue University, 1999.
14. We thank Coach Brian Luthy for the data. Provided by Brad Hartlaub.
15. John Kratus, "Effect of available tonality and pitch options on children's compositional processes and products," *Journal of Research in Music Education*, 49 (2001), pp. 294–306.
16. Euen Hyuk Jung, "The role of discourse signaling cues in second language listening comprehension," *Modern Language Journal*, 87 (2003), pp. 562–577.
17. We thank Jacqueline Ngai for providing data from Jacqueline T. Ngai and Diane S. Srivastava, "Predators accelerate nutrient cycling in a bromeliad ecosystem," *Science*, 314 (2006), p. 963. Counts for the last control plant are simulated because a monkey enjoyed snacking on this plant before the researchers were able to obtain all the counts.
18. G. Cultrone et al., "Ultrasound and mechanical tests combined with ANOVA to evaluate brick quality," *Ceramics International*, 27 (2001), pp. 401–406.